Over-parameterization Is Neither Robust nor Fragile in Matrix Sensing: A Geometric Theory of Solution Robustness

Runze Shi, Ziye Ma*

City University of Hong Kong
runzeshi3-c@my.cityu.edu.hk, ziyema@cityu.edu.hk

arXiv:2610.18318v1 [cs.LG] 13 Oct 2026

Abstract

Over-parameterization is ubiquitous in modern machine learning, yet its effect on robustness remains poorly understood. We study this question in matrix sensing, a canonical nonconvex problem that can also be viewed, under an appropriate parameterization, as training a two-layer quadratic neural network. We develop a rank-aware robustness theory for the Burer–Monteiro factorization when the search rank k may exceed the target rank $r_\star$, thereby capturing over-parameterization. Our analysis shows that robustness depends jointly on the degree of over-parameterization and the local geometry of the attained second-order stationary point: increasing k generally worsens sensitivity, but stationary points with the same search rank can exhibit substantially different robustness. Motivated by this observation, we propose a trajectory-aware algorithm that uses spectral balancing and randomized sketching to steer optimization toward less sensitive solutions. In matrix-sensing experiments, the method reduces paired sensitivity by 18–29% for isotropic operators and 87–89% for weak-tail operators, without degrading reconstruction accuracy. These results demonstrate that parameter count alone does not determine robustness and identify solution geometry and the optimization trajectory as important levers for improving it.

1 Introduction

Matrix sensing seeks to recover an unknown low-rank matrix from a collection of linear measurements, typically corrupted by noise. It is a canonical instance of the broader class of low-rank inverse problems, which also includes matrix completion (Candès and Recht 2009), robust principal component analysis and low-rank-plus-sparse decomposition (Candès et al. 2011), lifted formulations of phase retrieval and blind deconvolution (Candès, Strohmer, and Voroninski 2013; Ahmed, Recht, and Romberg 2014), and low-rank quantum state tomography (Gross et al. 2010). Matrix sensing provides a canonical setting for studying nonconvex low-rank estimation: the observation model is simple enough to permit precise analysis, while the rank constraint retains the key non-convex landscape that proves be mathematical challenging in large-scale machine learning (Recht, Fazel, and Parrilo 2010; Bhojanapalli, Neyshabur, and Srebro 2016; Ge, Jin, and Zheng 2017). Specifically, we consider positive-semidefinite matrix sensing,

$$\begin{aligned} f_w(X) &= \tfrac{1}{2}\left\|\mathcal{A}(XX^\top - M_\star) - w\right\|_2^2, \quad X \in \mathbb{R}^{n\times k}, \\ y_w &= \mathcal{A}(M_\star) + w \end{aligned} \tag{1}$$

where $M_\star \succeq 0$ has rank $1 \le r_\star \le k$. For rank-one positive-semidefinite measurements $A_i = a_i a_i^\top$, as in semidefinite liftings of phase retrieval (Candès, Strohmer, and Voroninski 2013), the prediction satisfies $\langle A_i, XX^\top\rangle = \sum_{j=1}^k (a_i^\top x_j)^2$. Thus, this special case is equivalent to training a width-k two-layer quadratic network with fixed output weights (Li, Ma, and Zhang 2018). In the positive-semidefinite setting considered in (1), the Burer–Monteiro factorization writes $M = XX^\top$, replacing the full matrix and eliminating the rank constraint (Burer and Monteiro 2003). The factor width k determines the search rank and may exceed the unknown target rank $r_\star$, indicating a mathematically elegant but representative form of over-parametrization.

Over-parameterization is useful when the target rank is unknown, since choosing $k < r_\star$ cannot recover the true solution, and practitioners normally make k much larger than the conjectured $r_\star$ to achieve a "safe-margin". When the sensing matrices are i.i.d. sampled from Gaussian distribution and enough measurements are made, it is generally known that matrix-sensing problems have no spurious local minima and exhibit directions of escape at saddle points (Bhojanapalli, Neyshabur, and Srebro 2016; Ge, Jin, and Zheng 2017; Zhang, Sojoudi, and Lavaei 2019). However, the requirements on the quantity and geometric uniformity of the sensing data for which such benign landscape exists is quite stringent. Critically, a recent line of work (Zhang 2025a; Ma et al. 2023; Stöger and Soltanolkotabi 2021) has shown that simply increasing search rank k can loosen such requirements. Thus, in terms of optimization, increase in parametrization can serve as useful degrees of freedom rather than merely redundant parameters.

However, the impact of over-parametrization on robustness is a much more complicated one. Existing theory first asks whether all local minima remain accurate after the measurements are corrupted. Under Gaussian noise and a $2r$-RIP constant below $1/10$, Bhojanapalli, Neyshabur, and Srebro (2016) show that every local minimum lies within $20\sigma_w\sqrt{\log(n)/m}$ of the ground truth. For general smooth

*Corresponding author.

objectives in the exactly parameterized regime, Ma and Sojoudi (2023) show that a clean RIP constant $\delta < 1/3$ places every noisy local minimum in a recovery ball whose radius is linear in the noise magnitude; for squared-loss matrix sensing, this radius scales as $O(\|\mathcal{A}\| \|w\|_2 /(1-3\delta))$. Their arbitrary-RIP result instead excludes local minima from a region around the truth. Ma et al. (2023) extend this class-wide analysis to noisy over-parameterized matrix sensing: under suitable RIP conditions, they bound the maximum recovery error of an arbitrary local minimum, while under weaker RIP they establish a close-or-far alternative. *These results therefore certify that robustness can be somehow maintained or expected even in the over-parametrized regime.*

A second line of results asks whether optimization can reach this certified set. The strict-saddle guarantees in Ma and Sojoudi (2023); Ma et al. (2023) show that points outside the recovery region must have either a large gradient or a direction of negative curvature, allowing perturbed gradient descent to find an approximate local minimum in polynomial time. A third approach changes the loss and studies a particular optimization trajectory. For an ℓ_1-loss under Sign-RIP, Ma and Fattahi (2023) show that subgradient descent from a small initialization can recover the truth despite rank overestimation and Gaussian or gross corruptions, even when some global minimizers overfit the noise. This is an algorithm-specific robustness guarantee for a different, robust loss.

A natural next question remains unanswered: is robustness the same for all attained solutions, and if not, how can optimization be steered toward the more robust ones? Existing recovery bounds assign one worst-case radius to the entire set of local minima, while strict-saddle theory guarantees convergence to some member of that set. They can certify that every solution is sufficiently close to the truth, but they cannot compare the sensitivities of two solutions inside the same recovery region. In particular, they do not pair a noisy critical point M_w with its clean counterpart M_0, quantify the solution-specific response $\|M_w - M_0\|_F / \|w\|_2$, or explain why solutions obtained with the same sensing operator, noise level, and search rank can respond differently. This distinction is especially important under over-parameterization because this means that the degree of parametrization is not the only factor influencing robustness. We therefore develop a local, quantitative, and rank-aware theory that identifies the spectral and curvature geometry governing the robustness of an attained stationary point, and then uses those solution-level quantities to steer optimization toward more robust solutions.

This point-wise view also allows for a trajectory-level study of how robustness can be influenced during optimization. Building on small-initialization dynamics and randomized matrix sketching (Stöger and Soltanolkotabi 2021; Tong, Ma, and Chi 2021a; Xu et al. 2023; Halko, Martinsson, and Tropp 2011; Woodruff 2014), we develop a trajectory-aware method that can actively guide trajectories toward more robust solutions. We evaluate the procedure on isotropic and weak-tail sensing operators using both reconstruction accuracy and paired perturbation sensitivity.

Contributions:

- **A geometric theory of robustness.** We develop a point-wise theory that explains how over-parameterization and the local geometry of an attained solution jointly determine its sensitivity to measurement noise. The theory covers both rank-deficient and full-rank solutions and distinguishes solutions that have the same search rank but different robustness.
- **A robust algorithm for matrix sensing and beyond.** We translate the theory into a trajectory-aware method that steers optimization toward less sensitive solutions. Across the sensing operators we test, the method improves paired robustness without sacrificing reconstruction error. In the appendix, we also demonstrate preliminary evidence that such approach can also help improve robustness in LoRA fine-tuning of large langauge models.

2 Related Work

Burer–Monteiro factorization and fixed-rank geometry. The Burer–Monteiro parameterization represents a positive-semidefinite matrix as $M = XX^\top$, enforcing an upper rank bound while reducing the ambient dimension (Burer and Monteiro 2003). Rank-deficient second-order points of smooth factorized semidefinite programs admit global optimality certificates (Boumal, Voroninski, and Bandeira 2016, 2020). Complementarily, fixed-rank manifold optimization provides the tangent-space Hessians and local perturbation tools needed to analyze rank-preserving motion (Absil, Mahony, and Sepulchre 2008; Vandereycken 2013; Boumal 2023). We use these two viewpoints for different rank profiles: rank-deficient points are treated through a matrix-space KKT condition, whereas full-rank points are analyzed through local continuation on a fixed-rank manifold.

Robustness of Matrix Sensing. We first talk about studies of robustness for the classic matrix sensing setting (1) with ℓ_2 loss. Most works in this regime only guarantees global recoverability under noise, but neglects discussion of robustness of non-global solutions. Bhojanapalli, Neyshabur, and Srebro (2016) shows that with Gaussian noise and a restrictive RIP constant threshold of $1/10$, every local minima lies within $20\sigma_w\sqrt{\log(n)/m}$ of the ground truth, which is considered to be the one of the most classic results in this field. Later on, several works have iterated upon this idea, and proved that the RIP constant can be as large as $1/2$ for robust global recovery to still hold (Ma and Sojoudi 2023; Li et al. 2020; Zhu et al. 2021). Specifically, recent works have also extended this analysis to the over-parameterized regime, and has preliminary shown that global solutions can still be robust, but the story is more complicated. Ma et al. (2023) has shown that in the over-parameterized setting global solution can still be recovered to certain accuracy under presence of noise, but this accuracy can be highly fragile and depends on eigenvalues of critical points. Another line of work by Zhang (2024, 2025b) also corroborates this fact and states that over-parameterization is benign under general loss functions. However, the caveat with this line of work is they require the loss function to satisfy RIP property with noise corruptions, which is quite unlikely.

Improving Robustness of Low-rank Optimization A separate line of work considers how to improve robustness of low-rank matrix problems. One of the most commonly considered approach is by adopting a robust loss, like ℓ_1 loss. For a known rank, Li et al. (2020) establishes local exact recovery and convergence guarantees for non-convex robust matrix factorization. With an over-parameterized rank, Ding et al. (2021) show that a subgradient method with diminishing step sizes still achieves exact recovery under sparse Gaussian-measurement corruptions. Ma and Fattahi (2023) further show that small initialization can prevent the redundant factor directions from fitting the noise, yielding global trajectory guarantees under both outlier and Gaussian noise models. Recent work extends this direction to preconditioned asymmetric factorizations (Giampouras, Cai, and Vidal 2025). At the same time, Ma and Fattahi (2025) show that, in over-parameterized robust matrix recovery, the ground-truth factors may be strict saddle points rather than local minima, emphasizing that algorithmic selection cannot be explained by local optimality alone. Despite the success of adopting a robust loss in the matrix space, such practice cannot be applied in larger-scale machine learning problems due to its non-differentiability. Therefore, we still need a more careful, and practically relevant analysis of robustness in differentiable regression losses.

Optimization dynamics of Low-rank Problems. Global landscape guarantees do not determine which region an optimization algorithm selects. However, there are existing works that we can build upon to explain robustness better. For example, in over-parameterized matrix sensing, small initialization induces an implicit spectral bias that first aligns the iterates with the signal subspace and suppresses redundant directions (Stöger and Soltanolkotabi 2021). Gram preconditioning accelerates ill-conditioned low-rank estimation (Tong, Ma, and Chi 2021a), and damped variants extend this mechanism to over-parameterized noisy matrix sensing (Xu et al. 2023). Randomized sketches provide scalable approximations of dominant subspaces and Rayleigh information (Halko, Martinsson, and Tropp 2011; Woodruff 2014). We use these tools for a different purpose: to monitor and modify the local robustness geometry along the optimization trajectory. Because exact RIP certification is generally intractable (Bandeira et al. 2013), our finite-probe RIP quantities are treated as diagnostics rather than formal global certificates.

3 Problem Setup and Notation

Fix positive integers n, m, k with $k \le n$. Let $\mathbb{S}^n$ denote the space of real symmetric matrices equipped with the Frobenius inner product. The sensing operator $\mathcal{A} : \mathbb{S}^n \to \mathbb{R}^m$ is defined by $(\mathcal{A}Z)_i = \langle A_i, Z\rangle$, where $A_i \in \mathbb{S}^n$, with adjoint $\mathcal{A}^* z = \sum_{i=1}^m z_i A_i$. We observe $y_w = \mathcal{A}(M_\star) + w$, where $w \in \mathbb{R}^m$ and $M_\star \succeq 0$ has rank $1 \le r_\star \le k$. The theory treats w deterministically; distributional noise assumptions are used only in the experiments. The case $k > r_\star$ is over-parameterized.

Define the smooth matrix-space loss on $\mathbb{S}^n$ by

$$F_w(M) = \frac{1}{2}\left\|\mathcal{A}(M - M_\star) - w\right\|_2^2. \tag{2}$$

The nonconvex factorized problem studied in this paper is

$$\min_{X\in\mathbb{R}^{n\times k}} f_w(X), \qquad f_w(X) = F_w(XX^\top). \tag{3}$$

We write $H = \mathcal{A}^*\mathcal{A}$, $N_w = \mathcal{A}^* w$, and $G_w(M) = \nabla F_w(M) = H(M - M_\star) - N_w$, where $G_w(M)$ is the ambient Frobenius gradient. Then assume $r_\star \le k$ and $\mathcal{A}$ satisfies $(\delta, 2k)$-RIP for some $\delta \in [0, 1)$, i.e., $\left|\|\mathcal{A}(Z)\|_2^2 - \|Z\|_F^2\right| \le \delta\|Z\|_F^2$ for all $Z \in \mathbb{S}^n$ with $\mathrm{rank}(Z) \le 2k$. We set $s_{\mathrm{RIP}} = 2k$, and, for $M \succeq 0$, write its positive eigenvalues as $\lambda_1(M) \ge \cdots \ge \lambda_{\mathrm{rank}(M)}(M) > 0$.

For $1 \le q \le \mathrm{rank}(M)$, let $P_q(M)$ be a projector onto a top-q eigenspace and set $P_q^\perp = I - P_q$. The tangent space of the rank-q PSD manifold at a matrix with range projector P is

$$\mathcal{T}(P) = \{\eta \in \mathbb{S}^n : P^\perp \eta P^\perp = 0\}. \tag{4}$$

Every $\eta \in \mathcal{T}(P)$ has rank at most $2q$. Let factor width k be the search rank; $q = \mathrm{rank}(M)$ at deficient points, $q = k$ along continuations, and $q < k$ only for diagnostics.

Definition 1 (Tangent gain and normal penalty)**.** For a rank-q projector P, define

$$\alpha_T(P) = \inf_{\eta\in\mathcal{T}(P),\, \|\eta\|_F=1} \|\mathcal{A}(\eta)\|_2^2. \tag{5}$$

For $M \succeq 0$ and $1 \le q \le \mathrm{rank}(M)$, define

$$\rho_q^-(M; w) = \frac{\left\|[P_q^\perp G_w(M) P_q^\perp]_-\right\|_{\mathrm{op}}}{\lambda_q(M)}. \tag{6}$$

Here $[B]_- \succeq 0$ is the negative spectral part of B.Here $[B]_-$ is the negative spectral part of B. The exact Hessian penalty is ρ_q^-.

Under Assumption 1, every rank-q projector with $q \le k$ satisfies

$$1 - \delta \le \alpha_T(P) \le 1 + \delta. \tag{7}$$

4 Stability of Rank-k Critical Points

As per existing over-parametrized matrix sensing theory, any critical point that is rank-deficient is already globally optimal, which is rare to come by (Boumal, Voroninski, and Bandeira 2016). Therefore, the majority of critical points we encounter will be of full rank (rank-k). This section will be dedicated to studying how the robustness of rank-k critical points differ from each other, and in the next section we will extend this discussion to rank-deficient ones. We first show how the normal block of the ambient gradient reduces tangent curvature in the factorized formulation. We then use the resulting curvature margin to study the surrounding space of Riemannian critical points projected on a rank-k manifold.

Tangent curvature in the factorized formulation. Let $M = XX^\top$, where X has full column rank, and suppose $G_w(M)X = 0$. For a factor direction D, let $\eta = XD^\top + DX^\top \in \mathcal{T}(P)$, where P is the range projector of M. The Hessian decomposes as

$$\nabla^2 f_w(X)[D, D] = \|\mathcal{A}(\eta)\|_2^2 + 2\left\langle G_w(M), DD^\top\right\rangle. \tag{8}$$

The first term is the sensing curvature along η. In the second, $G_w(M)X = 0$ removes the range and cross terms, so only $[P^\perp G_w(M)P^\perp]_-$ reduces curvature. Normalized by the active spectral floor $\lambda_k(M)$, this loss is $\rho_k^-(M;w)$ from Definition 1. Hence,

$$\begin{aligned}\nabla^2 f_w(X)[D,D] &\geq \big(\alpha_T(P) - \rho_k^-(M;w)\big)\,\|\eta\|_F^2 \\ &\geq \big(1-\delta-\rho_k^-(M;w)\big)\,\|\eta\|_F^2 .\end{aligned} \tag{9}$$

The first inequality combines the lower tangent gain $\alpha_T(P)$ with the most negative Rayleigh quotient of the normal-normal block, normalized by $\lambda_k(M)$; the second uses the $2k$-RIP lower bound $\alpha_T(P) \geq 1-\delta$. Thus, $\alpha_T(P)-\rho_k^-(M;w)$ lower-bounds $\nabla^2 f_w(X)[D,D]/\,\|\eta\|_F^2$ and is the tangent-curvature margin at M; $1-\delta-\rho_k^-(M;w)$ is its conservative RIP-based lower bound. Because $[P^\perp G_w(M)P^\perp]_-$ can differ across attained points, rank-k critical points with the same search rank and sensing operator can differ in robustness. Details are given in Appendix B.2.

Riemannian critical points on the rank-k manifold. Let

$$\mathcal{M}_k = \{M \succeq 0 : \operatorname{rank}(M) = k\}$$

be the embedded rank-k PSD manifold. At a full-column-rank critical factor, horizontal factor directions lift all tangent directions of $\mathcal{M}_k$. For $\eta \in T_M\mathcal{M}_k$ and a horizontal factor direction D satisfying $\eta = XD^\top + DX^\top$, Luo, Li, and Zhang (2025, Theorem 1) give

$$\langle \eta, \operatorname{Hess}_{\mathcal{M}_k} F_w(M)[\eta]\rangle = \nabla^2 f_w(X)[D,D].$$

Thus, (9) also lower-bounds the Riemannian Hessian on $T_M\mathcal{M}_k$.

For $s \in [0,1]$, write

$$F_s(M) = F_{sw}(M), \qquad G_s(M) = G_{sw}(M).$$

Let $M_1 = M_w \in \mathcal{M}_k$ be a computed rank-k Riemannian critical point of $F_1 = F_w$ satisfying

$$\operatorname{grad}_{\mathcal{M}_k} F_w(M_w) = 0. \tag{10}$$

For the embedded Frobenius metric, $\operatorname{grad}_{\mathcal{M}_k} F_s(M) = 0$ is equivalent to $G_s(M)P = 0$.

Define the RIP-based tangent-curvature margin at M_w by

$$\mu_w = (1-\delta) - \rho_k^-(M_w;w). \tag{11}$$

If $\mu_w > 0$, the Riemannian Hessian at M_w is positive definite on $T_{M_w}\mathcal{M}_k$. The implicit-function theorem gives a C^1 curve for s in a neighborhood of 1, and the small-noise condition in Appendix B.2 continues it over $[0,1]$ while keeping the tangent-Hessian margin and rank-k spectral floor uniformly positive:

$$\begin{aligned} s \mapsto M_s \in \mathcal{M}_k, &\qquad 0 \leq s \leq 1, \\ M_1 = M_w, &\qquad \operatorname{grad}_{\mathcal{M}_k} F_s(M_s) = 0. \end{aligned}$$

Set $M_0 := M_{s=0}$, the clean rank-k Riemannian critical point paired with M_w by this curve. The Hessian equality, complete argument, and uniform bounds are given in Appendix B.2.

Theorem 1 (Pointwise rank-stable perturbation bound)**.** *Suppose Assumption 1 and the local stability conditions at M_w in Assumptions 2–3 hold. Then the noisy rank-k Riemannian critical point M_w and the corresponding clean rank-k Riemannian critical point M_0 satisfy*

$$\|M_w - M_0\|_F \leq \frac{2\sqrt{1+\delta}}{(1-\delta)-\rho_k^-(M_w;w)}\,\|w\|_2 . \tag{12}$$

Theorem 3 makes the comparison among rank-k critical points explicit: a smaller negative normal component gives a larger tangent-curvature margin and a tighter noise-linear perturbation bound, whereas the bound deteriorates as $\rho_k^-(M_w;w)$ approaches $1-\delta$. Thus, the search rank alone does not determine the robustness of an attained point. The bound controls the perturbation between the two rank-k Riemannian critical points M_w and M_0 connected by the continuation curve, rather than the recovery error relative to $M_\star$. In particular, when $r_\star < k$, the rank-k critical point M_0 cannot equal the lower-rank target $M_\star$; direct recovery guarantees are given in Theorem 2(2) and (3). The associated projector-rotation bound is given in Appendix B.2. The result is intrinsic and does not depend on a particular local chart; complementary coordinate expressions are deferred to Appendix B.2.

5 Rank Deficiency and the Three Robustness Cases

We now extend the discussion to both rank-deficient and full-rank SOSPs in this section. First we explain how unused columns provide a matrix-space optimality certificate, and then we will show how the slack spectral tail limits the negative spectral part of the normal block at a full-rank point. These two arguments yield three rank-aware robustness routes centered at the computed noisy point: direct recovery for rank-deficient SOSPs, critical points stability along a full-rank branch, and direct recovery relative to the lower-rank target.

5.1 Rank-Deficient SOSPs

Rank deficiency makes the factorized second-order condition especially useful. If X has an unused column direction, a factor perturbation along that direction has zero first-order matrix variation. The second-order condition then gives a direct condition on the ambient matrix gradient. Since the same argument applies to every matrix direction, it forces

$$G_w(XX^\top) \succeq 0.$$

First-order criticality also gives $G_w(M)X = 0$, where $M = XX^\top$, and therefore $\langle G_w(M), M\rangle = 0$. For every $Z \succeq 0$, these two conditions imply

$$\langle G_w(M), Z - M\rangle = \langle G_w(M), Z\rangle \geq 0.$$

So the rank-deficient SOSP is a global minimizer in matrix space. In the noiseless case, its objective value is zero and the unified $2k$-RIP gives $M = M_\star$. With noise, the same optimality condition and RIP give the noise-linear recovery bound

$$\|M - M_\star\|_F \le \frac{\|w\|_2}{\sqrt{1-\delta}}.$$

Thus, the rank-deficient guarantee is stated directly relative to the target; no clean rank-k critical point is needed. This is the PSD matrix-sensing specialization of the standard rank-deficiency certificate for Burer–Monteiro factorizations (Burer and Monteiro 2003; Boumal, Voroninski, and Bandeira 2016, 2020). The formal statements and complete arguments are deferred to Appendix B.3.

5.2 Full-Rank Points and the Spectral Tail

For a full-rank SOSP, write $M = XX^\top = \sum_{i=1}^k \lambda_i u_i u_i^\top$, with $\lambda_1 \ge \cdots \ge \lambda_k > 0$, and $P_q = \sum_{i=1}^q u_i u_i^\top$, $q \le k$. Second-order stationarity and the unified $2k$-RIP give

$$\left\|[P_q^\perp G_w(M) P_q^\perp]_-\right\|_{\mathrm{op}} \le (1+\delta)\lambda_k(M), \tag{13}$$

$$\rho_q^-(M; w) \le (1+\delta)\frac{\lambda_k(M)}{\lambda_q(M)}. \tag{14}$$

Thus, $\lambda_k(M)/\lambda_q(M)$ bounds the normalized negative spectral part of the normal block, but not the full tail energy $\sum_{i=q+1}^k \lambda_i(M)^2$.

For target recovery, first-order criticality gives $G_w(M_w) = P_w^\perp G_w(M_w) P_w^\perp$, where $P_w = P_k(M_w)$. Define

$$b_w^- := \left\|[P_w^\perp G_w(M_w) P_w^\perp]_-\right\|_{\mathrm{op}} = \lambda_k(M_w)\rho_k^-(M_w; w). \tag{15}$$

Another ingredient controls the portion of the lower-rank target that lies outside $\mathrm{range}(M_w)$. This target-subspace leakage is controlled by the current recovery error. Combining that estimate with b_w^- and the unified $2k$-RIP gives the direct full-rank recovery certificate below. The formal spectral-tail statement, the target-subspace leakage estimate, and the complete normal-leakage argument are deferred to Appendix B.3.

5.3 Main Theoretical Result

Here we present our main theorem, which elaborates three possible robust cases. Given the discussion about rank-deficient solutions above, we only compare it to $M_\star$. For a full-rank critical point under noise, one can either view it as perturbed from a clean and nearby full rank solution, or perturbed from $M_\star$. We believe those three cases offer a rather complete characterization of possible scenarios.

Theorem 2 (Three robustness certificates)**.** *Let X_w be an SOSP of f_w, and write $M_w = X_w X_w^\top$. Under Assumption 1, the robustness of X_w can be characterized as:*

1. ***General Case.*** *Suppose $\mathrm{rank}(X_w) = k$, Assumption 3 additionally holds, and the RIP-based curvature margin is positive:*

$$\mu_w = (1-\delta) - \rho_k^-(M_w; w) > 0. \tag{16}$$

Let $s \mapsto M_s \in \mathcal{M}_k$ be the critical continuation supplied by Lemma 2, with $M_1 = M_w$, and let $M_0 = M_{s=0}$ be its clean rank-k critical point. Then

$$\|M_w - M_0\|_F \le \frac{2\sqrt{1+\delta}}{(1-\delta) - \rho_k^-(M_w; w)} \|w\|_2. \tag{17}$$

For rank $q < k$, Lemma 4 in Appendix B.3 gives deterministic bounds for the best rank-q approximation; direct recovery relative to $M_\star$ is handled separately in part (3).

2. ***Direct Rank-deficient certificate.*** *If $\mathrm{rank}(X_w) < k$, then M_w is a global noisy PSD minimizer and*

$$\|M_w - M_\star\|_F \le \frac{\|w\|_2}{\sqrt{1-\delta}}. \tag{18}$$

3. ***Direct Full-rank Certificate.*** *Suppose $\mathrm{rank}(X_w) = k$ with $P_w = P_k(M_w)$, and let b_w^- be the absolute negative normal leakage defined in* (78)*. Define*

$$\mu_{\mathrm{leak}} = 1 - \delta - \frac{b_w^-}{\lambda_{r_\star}(M_\star)}. \tag{19}$$

If $\mu_{\mathrm{leak}} > 0$, then

$$\|M_w - M_\star\|_F \le \frac{\sqrt{1+\delta}}{\mu_{\mathrm{leak}}} \|w\|_2. \tag{20}$$

This bound is derived using a different technical route than the general case, and can be useful if one is only interested in comparing to $M_\star$.

Remark 1 (Scope of the full-rank continuation certificate). Under standard no-spurious-minimum RIP conditions (Bhojanapalli, Neyshabur, and Srebro 2016; Ge, Jin, and Zheng 2017; Zhang, Sojoudi, and Lavaei 2019), a clean full-rank SOSP distinct from the lower-rank target $M_\star$ may not exist. Since those RIP requirements are usually hard to satisfy, the general case offers a more practical bound under weaker data conditioning and also covers non-global solutions.

6 Anchor-Free Guidance Toward Robust Full-Rank Critical Points

The preceding theory identifies the pointwise margin $(1-\delta) - \rho_k^-$ as a key factor of stability. We now turn this observation into an optimization procedure and ask whether the trajectory can be steered toward a less sensitive critical point. At a high level, our method acts on the numerator and denominator of ρ_k^- during an early guidance window, gradually removes this guidance, and finally polishes the iterate on the original objective f_w. Then we will evaluate the polished critical points under an end-to-end protocol, an initialization-controlled protocol, and comparisons with robust low-rank recovery baselines.

For $M_t = X_t X_t^\top$, let $P_t = P_k(M_t)$ and define

$$\begin{aligned} \rho_{k,t}^- &= \frac{b_t^-}{\lambda_k(M_t)}, \qquad G_t = \mathcal{A}^*\big(\mathcal{A}(M_t) - y_w\big), \\ b_t^- &= \max\big\{0,\, -\lambda_{\min}\big(P_t^\perp G_t P_t^\perp\big)\big\}. \end{aligned} \tag{21}$$

Here b_t^- is the absolute negative normal leakage, while $\lambda_k(M_t)$ is the smallest nonzero eigenvalue $\lambda_k(M_t)$. The method is "anchor-free" because it uses only the current factor, observed residual, applications of $\mathcal{A}$ and $\mathcal{A}^*$, and random probes. It does not use $M_\star$, clean measurements, or a penalty anchoring the trajectory to its initialization.

Algorithm 1: Decayed Anchor-Free Balanced Guidance

Require: Measurements y_w, sensing operator $\mathcal{A}$, search rank k, horizon T, and the optimization and guidance parameters listed in Appendix A.

Spectrally initialize $X_0 \in \mathbb{R}^{n\times k}$.
Set $T_g = \lfloor \theta T \rfloor$.
for $t = 0, \dots, T-1$ **do**
  Compute the residual $e_t = \mathcal{A}(X_t X_t^\top) - y_w$ and $G_t = \mathcal{A}^* e_t$.
  if $t < T_g$ **then**
    Set $c_t = (1 - t/T_g)^3$.
    Compute the scale-normalized inverse-trace direction $g_{\mathrm{inv},t}$ from (42) as defined in Appendix A.
    Draw s normal probes and compute the sampled Rayleigh direction $g_{-,t}$ from (41) using the matrix-free sketch in Appendix A.
  else
    Set $c_t = 0$, $g_{\mathrm{inv},t} = 0$, and $g_{-,t} = 0$.
  end if
  Form

$$g_t = 2G_tX_t + \alpha_0 c_t g_{\mathrm{inv},t} + \beta_0 c_t g_{-,t}.$$

  Set $S_t = X_t^\top X_t$ and $\xi_t = \xi\ \max\{\mathrm{tr}(S_t)/k, \epsilon_0\}$.
  Take the preconditioned step

$$X_{t+1} = X_t - \eta_t g_t (S_t + \xi_t I_k)^{-1}.$$

  if $t < T_g$ and h divides $t+1$ **then**
    Apply the normal-direction correction (43) in Appendix A using a fresh sketch.
  end if
end for
Set all guidance weights to zero and polish on f_w alone.
return The first full-rank polished iterate satisfying (23).

6.1 Certificate-Guided Optimization

Equation (21) exposes two complementary controls. Protecting $\lambda_k(M_t)$ prevents the denominator of $\rho^-_{k,t}$ from collapsing, while reducing b_t^- suppresses negative normal curvature. Since $\rho_k^- = b^-/\lambda_k$, the bounds $\lambda_k(\widetilde{M}) \geq a\lambda_k(M)$ and $b^-(\widetilde{M}) \leq cb^-(M)$ immediately imply $\rho_k^-(\widetilde{M}; w) \leq (c/a)\rho_k^-(M; w)$. Thus both quantities must be controlled. Thus, improving one component alone may not raise the certified margin if the other worsens comparably.

To protect the active floor, we use the $k \times k$ Gram matrix $S_t = X_t^\top X_t$ and the inverse-trace penalty

$$\Omega_{\mathrm{inv}}(X_t) = (\bar\lambda_t + \varepsilon_t)\,\mathrm{tr}\big((S_t + \varepsilon_t I_k)^{-1}\big), \qquad \bar\lambda_t = \frac{\mathrm{tr}(S_t)}{k}.$$

Its scale-normalized gradient $g_{\mathrm{inv},t}$ acts most strongly on weak factor directions and therefore targets the denominator $\lambda_k(M_t)$ without forming M_t.

To control the numerator, we project Gaussian probes onto $\mathrm{range}(X_t)^\perp$ and evaluate their normal Rayleigh quotients against G_t. The direction $g_{-,t}$ penalizes the detected negative quotients and moves them toward zero, providing a matrix-free approximation to reducing b_t^-. Exact gradients, probe construction, and the periodic normal-direction correction are given in Appendix A.

During the guidance window, the two directions are combined with the original gradient:

$$\begin{aligned} g_t &= 2G_tX_t + \alpha_0 c_t g_{\mathrm{inv},t} + \beta_0 c_t g_{-,t}, \\ c_t &= (1 - t/T_g)^3, \qquad T_g = \lfloor \theta T \rfloor. \end{aligned} \tag{22}$$

We take the damped Gram-preconditioned step

$$X_{t+1} = X_t - \eta_t g_t \left(S_t + \xi_t I_k\right)^{-1}, \qquad \xi_t = \xi \max\{\bar\lambda_t, \epsilon_0\}.$$

After T_g, all guidance terms are removed. We then optimize f_w alone until

$$\frac{\|\nabla f_w(X)\|_F}{\max\{1, \|X\|_F\}} \leq \varepsilon_{\mathrm{crit}}. \tag{23}$$

Consequently, guidance affects trajectory and solution selection, but every reported critical point is an approximate critical point of the original objective rather than of a permanently regularized loss.

6.2 Empirical Protocol and Results

We evaluate two questions: whether Algorithm 1 reduces end-to-end critical-point response without degrading reconstruction, and how its robustness–reconstruction trade-off compares with existing low-rank recovery methods under controlled initialization.

Setup. We generate rank-$r_\star$ PSD targets and use five over-parameterized configurations with n ranging from 50 to 150, $r_\star$ from 2 to 5, and k from 5 to 12. Measurements use implicit bilinear sensing with $m = 6nk$. We consider isotropic Gaussian and Rademacher operators, together with a weak-tail operator that attenuates directions orthogonal to the target subspace by a factor 0.35. The latter is a deliberately anisotropic stress test, not a heavy-tailed noise model. We use five seeds and three independent Gaussian measurement perturbations per configuration–seed instance, each with norm $0.08\,\|\mathcal{A}(M_\star)\|_2$. Appendix A gives the exact configurations, seeds, operator construction, and optimization details.

Protocols and metrics. The primary end-to-end protocol initializes the clean and noisy runs from their respective observations, $\mathrm{SpecInit}(\mathcal{A}^* y_0)$ and $\mathrm{SpecInit}(\mathcal{A}^* y_{w_j})$. The conditional shared observed-spectral protocol instead initializes both members of each pair from the same noisy-data initializer $\mathrm{SpecInit}(\mathcal{A}^* y_{w_j})$, thereby isolating trajectory selection from initialization sensitivity. We report results separately.

For protocol $p \in \{\mathrm{E2E}, \mathrm{Shared}\}$, method a, and matched instance s, define

$$\mathsf{R}^a_{p,s} = \underset{j=1,2,3}{\mathrm{median}} \left\| M^a_{\mathrm{crit},p,s}(w_j) - M^a_{\mathrm{crit},p,s}(0) \right\|_F / \|w_j\|_2$$

and $\mathsf{Q}_{p,s} = \mathsf{R}^{\mathrm{A1}}_{p,s}/\mathsf{R}^{\mathrm{Plain}}_{p,s}$. We form $\mathsf{Q}_{p,s}$ for each matched instance and report its median over the 25 instances for each operator. We also report $\mathsf{E}_{\mathrm{rel},s} = \left\|M^{\mathrm{A1}}_{\mathrm{crit},s} - M_\star\right\|_F / \left\|M^{\mathrm{Plain}}_{\mathrm{crit},s} - M_\star\right\|_F$.

We report operator-level medians of $\mathsf{Q}_{p,s}$ over the 25 matched instances. Here, $\mathsf{Q} < 1$ and $\mathsf{E}_{\mathrm{rel}} < 1$ indicate lower paired critical-point sensitivity and reconstruction error than Plain, respectively.

Protocol	Operator	Ratio	Δ_{R}	$\mathsf{E}_{\mathrm{rel}}$	Joint
E2E	Gauss.	0.788	21.2%	0.887	23/25
E2E	Radem.	0.785	21.5%	0.863	25/25
E2E	Weak-tail	0.180	82.0%	0.625	25/25
Shared	Gauss.	0.781	21.9%	0.887	23/25
Shared	Radem.	0.771	22.9%	0.863	25/25
Shared	Weak-tail	0.117	88.3%	0.625	25/25

Table 1: Algorithm 1 relative to Plain. Ratio is the median instancewise response ratio $\mathsf{Q}_{p,s}$, and $\Delta_{\mathsf{R}} = 1 - \mathrm{median}_s\, \mathsf{Q}_{p,s}$. Joint counts instances in which both response and reconstruction improve. Each row contains 25 configuration–seed instances.

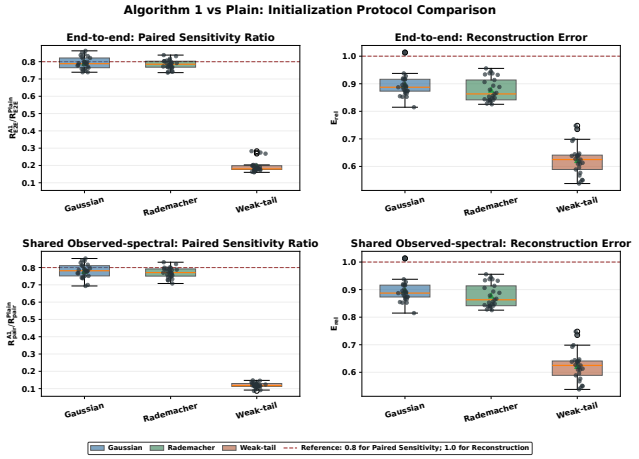

Figure 1: Distributions of the Algorithm 1/Plain Paired Sensitivity Ratio (left) and reconstruction-error ratios (right) under E2E (top) and Shared (bottom) initialization. Smaller values favor Algorithm 1; the dashed lines mark 20% response improvement and reconstruction parity. The right panels coincide because both protocols use the same noisy-data runs.

Main results. Table 1 reports operator-level medians of the paired critical-point sensitivity ratio Q and relative reconstruction error $\mathsf{E}_{\mathrm{rel}}$, while Figure 1 shows their instance-level distributions. Under E2E, the median Q values are 0.788, 0.785, and 0.180, and the corresponding median $\mathsf{E}_{\mathrm{rel}}$ values are 0.887, 0.863, and 0.625 for Gaussian, Rademacher, and weak-tail sensing. Both metrics are below one in $73/75$ matched instances. All reported Algorithm 1 solutions pass the unguided criticality and full-rank checks.

Under Shared initialization, the median Q values are 0.781, 0.771, and 0.117, with the same $\mathsf{E}_{\mathrm{rel}}$ values and joint count. Because Shared conditions out initialization sensitivity, it is treated as a controlled rather than primary result. The weak-tail gains reflect the anisotropic stress test, not an easier sensing problem.

Robust recovery and reconstruction baselines. L1-SubGM (Ma and Fattahi 2023) and ScaledSM (Tong, Ma, and Chi 2021b) are robust matrix-recovery methods using nonsmooth ℓ_1-type losses for corrupted measurements. GNMR (Zilber and Nadler 2022) is included as a strong nonconvex reconstruction baseline, testing whether lower paired sensitivity merely results from improved recovery.

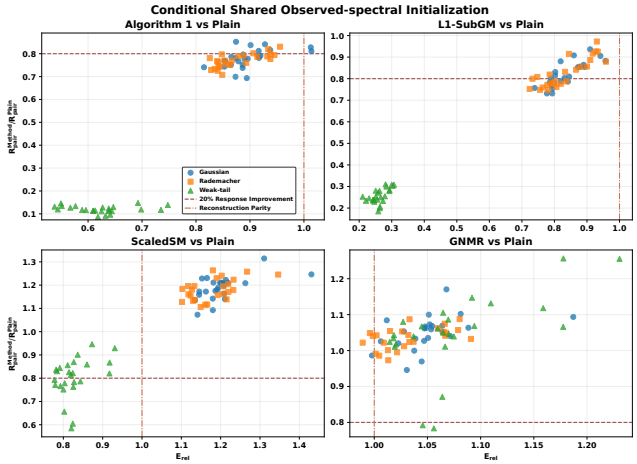

Figure 2: Conditional Shared comparison relative to Plain. Each point is one configuration–seed instance. The horizontal line marks a 20% reduction in paired sensitivity, and the vertical line marks reconstruction parity; the lower-left region improves both metrics.

Figure 2 compares all methods under Shared using the same paired-sensitivity and reconstruction metrics. Algorithm 1 has the smallest pooled median paired-sensitivity ratio (0.751), compared with 0.776 for L1-SubGM, 1.155 for ScaledSM, and 1.049 for GNMR, and improves both metrics in $73/75$ instances. L1-SubGM also gives favorable tradeoffs, especially under weak-tail sensing; ScaledSM improves mainly in that setting, while GNMR gives no consistent sensitivity reduction. All displayed solutions are polished on f_w, and f_w-criticality claims refer to the polished critical points.

Additional experiments. Appendix A reports further E2E, shared-initialization, and large-scale matrix-sensing results with consistent trends. It also presents a preliminary Algorithm 1-inspired LoRA (Hu et al. 2022a) study that fine-tunes RoBERTa-base (Liu et al. 2019) on GLUE/SST-2 (Wang et al. 2018a), improving corrupted-data accuracy by approximately one percentage point.

7 Conclusion

We develop a rank-aware robustness theory for overparameterized PSD matrix sensing. Rank-deficient SOSPs admit direct recovery guarantees, whereas full-rank robustness depends on tangent curvature, the active spectral floor, and negative normal leakage. These quantities motivate a transient, anchor-free method whose polished critical points improve paired robustness across Gaussian, Rademacher, and weak-tail operators without worsening reconstruction. Thus robustness is governed by solution geometry and solution selection, not factor width alone.

References

Absil, P.-A.; Mahony, R.; and Sepulchre, R. 2008. *Optimization Algorithms on Matrix Manifolds*. Princeton, NJ: Princeton University Press. ISBN 9780691132983.

Ahmed, A.; Recht, B.; and Romberg, J. 2014. Blind Deconvolution Using Convex Programming. *IEEE Transactions on Information Theory*, 60(3): 1711–1732.

Bandeira, A. S.; Dobriban, E.; Mixon, D. G.; and Sawin, W. F. 2013. Certifying the Restricted Isometry Property Is Hard. *IEEE Transactions on Information Theory*, 59(6): 3448–3450.

Bhatia, R. 1997. *Matrix Analysis*, volume 169 of *Graduate Texts in Mathematics*. Springer.

Bhojanapalli, S.; Neyshabur, B.; and Srebro, N. 2016. Global Optimality of Local Search for Low Rank Matrix Recovery. In *Advances in Neural Information Processing Systems 29*, 3873–3881.

Boumal, N. 2023. *An Introduction to Optimization on Smooth Manifolds*. Cambridge University Press. ISBN 9781009166171.

Boumal, N.; Voroninski, V.; and Bandeira, A. S. 2016. The Non-Convex Burer–Monteiro Approach Works on Smooth Semidefinite Programs. In *Advances in Neural Information Processing Systems 29*, 2757–2765.

Boumal, N.; Voroninski, V.; and Bandeira, A. S. 2020. Deterministic Guarantees for Burer–Monteiro Factorizations of Smooth Semidefinite Programs. *Communications on Pure and Applied Mathematics*, 73(3): 581–608.

Burer, S.; and Monteiro, R. D. C. 2003. A Nonlinear Programming Algorithm for Solving Semidefinite Programs via Low-Rank Factorization. *Mathematical Programming*, 95(2): 329–357.

Candès, E. J.; Li, X.; Ma, Y.; and Wright, J. 2011. Robust Principal Component Analysis? *Journal of the ACM*, 58(3): 1–37.

Candès, E. J.; and Recht, B. 2009. Exact Matrix Completion via Convex Optimization. *Foundations of Computational Mathematics*, 9(6): 717–772.

Candès, E. J.; Strohmer, T.; and Voroninski, V. 2013. PhaseLift: Exact and Stable Signal Recovery from Magnitude Measurements via Convex Programming. *Communications on Pure and Applied Mathematics*, 66(8): 1241–1274.

Davis, C.; and Kahan, W. M. 1970. The Rotation of Eigenvectors by a Perturbation. III. *SIAM Journal on Numerical Analysis*, 7(1): 1–46.

Ding, L.; Jiang, L.; Chen, Y.; Qu, Q.; and Zhu, Z. 2021. Rank Overspecified Robust Matrix Recovery: Subgradient Method and Exact Recovery. In *Advances in Neural Information Processing Systems*, volume 34, 26767–26778.

Eckart, C.; and Young, G. 1936. The Approximation of One Matrix by Another of Lower Rank. *Psychometrika*, 1(3): 211–218.

Ge, R.; Jin, C.; and Zheng, Y. 2017. No Spurious Local Minima in Nonconvex Low Rank Problems: A Unified Geometric Analysis. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, 1233–1242. PMLR.

Giampouras, P.; Cai, H.; and Vidal, R. 2025. Guarantees of a Preconditioned Subgradient Algorithm for Overparameterized Asymmetric Low-rank Matrix Recovery. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, 19458–19479. PMLR.

Gross, D.; Liu, Y.-K.; Flammia, S. T.; Becker, S.; and Eisert, J. 2010. Quantum State Tomography via Compressed Sensing. *Physical Review Letters*, 105(15): 150401.

Halko, N.; Martinsson, P.-G.; and Tropp, J. A. 2011. Finding Structure with Randomness: Probabilistic Algorithms for Constructing Approximate Matrix Decompositions. *SIAM Review*, 53(2): 217–288.

Hu, E. J.; Shen, Y.; Wallis, P.; Allen-Zhu, Z.; Li, Y.; Wang, S.; Wang, L.; and Chen, W. 2022a. LoRA: Low-Rank Adaptation of Large Language Models. In *International Conference on Learning Representations*.

Hu, E. J.; Shen, Y.; Wallis, P.; Allen-Zhu, Z.; Li, Y.; Wang, S.; Wang, L.; and Chen, W. 2022b. LoRA: Low-Rank Adaptation of Large Language Models. In *International Conference on Learning Representations*.

Li, X.; Zhu, Z.; So, A. M.-C.; and Vidal, R. 2020. Nonconvex Robust Low-Rank Matrix Recovery. *SIAM Journal on Optimization*, 30(1): 660–686.

Li, Y.; Ma, T.; and Zhang, H. 2018. Algorithmic regularization in over-parameterized matrix sensing and neural networks with quadratic activations. In *Conference on learning theory*, 2–47. PMLR.

Liu, Y.; Ott, M.; Goyal, N.; Du, J.; Joshi, M.; Chen, D.; Levy, O.; Lewis, M.; Zettlemoyer, L.; and Stoyanov, V. 2019. RoBERTa: A Robustly Optimized BERT Pretraining Approach. *arXiv preprint arXiv:1907.11692*.

Luo, Y.; Li, X.; and Zhang, A. R. 2025. Nonconvex Factorization and Manifold Formulations Are Almost Equivalent in Low-Rank Matrix Optimization. *INFORMS Journal on Optimization*, 7(1): 38–60.

Ma, J.; and Fattahi, S. 2023. Global Convergence of Subgradient Method for Robust Matrix Recovery: Small Initialization, Noisy Measurements, and Over-parameterization. *Journal of Machine Learning Research*, 24(96): 1–84.

Ma, J.; and Fattahi, S. 2025. Can Learning Be Explained by Local Optimality in Robust Low-Rank Matrix Recovery? *Mathematics of Operations Research*.

Ma, Z.; Bi, Y.; Lavaei, J.; and Sojoudi, S. 2023. Geometric Analysis of Noisy Low-Rank Matrix Recovery in the Exact Parametrized and the Overparametrized Regimes. *INFORMS Journal on Optimization*, 5(4): 356–375.

Ma, Z.; and Sojoudi, S. 2023. Noisy Low-rank Matrix Optimization: Geometry of Local Minima and Convergence Rate. In *International Conference on Artificial Intelligence and Statistics*, 3125–3150. PMLR.

Recht, B.; Fazel, M.; and Parrilo, P. A. 2010. Guaranteed Minimum-Rank Solutions of Linear Matrix Equations via Nuclear Norm Minimization. *SIAM Review*, 52(3): 471–501.

Socher, R.; Perelygin, A.; Wu, J.; Chuang, J.; Manning, C. D.; Ng, A. Y.; and Potts, C. 2013. Recursive Deep Models for Semantic Compositionality Over a Sentiment Treebank. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, 1631–1642. Association for Computational Linguistics.

Stöger, D.; and Soltanolkotabi, M. 2021. Small Random Initialization is Akin to Spectral Learning: Optimization and Generalization Guarantees for Overparameterized Low-Rank Matrix Reconstruction. In *Advances in Neural Information Processing Systems*.

Tong, T.; Ma, C.; and Chi, Y. 2021a. Accelerating Ill-Conditioned Low-Rank Matrix Estimation via Scaled Gradient Descent. *Journal of Machine Learning Research*, 22(150): 1–63.

Tong, T.; Ma, C.; and Chi, Y. 2021b. Low-Rank Matrix Recovery with Scaled Subgradient Methods: Fast and Robust Convergence without the Condition Number. *IEEE Transactions on Signal Processing*, 69: 2396–2409.

Vandereycken, B. 2013. Low-Rank Matrix Completion by Riemannian Optimization. *SIAM Journal on Optimization*, 23(2): 1214–1236.

Wang, A.; Singh, A.; Michael, J.; Hill, F.; Levy, O.; and Bowman, S. 2018a. GLUE: A Multi-Task Benchmark and Analysis Platform for Natural Language Understanding. In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, 353–355. Association for Computational Linguistics.

Wang, A.; Singh, A.; Michael, J.; Hill, F.; Levy, O.; and Bowman, S. R. 2018b. GLUE: A Multi-Task Benchmark and Analysis Platform. In *Proceedings of the 2018 EMNLP Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, 353–355. Association for Computational Linguistics.

Woodruff, D. P. 2014. Sketching as a Tool for Numerical Linear Algebra. *Foundations and Trends in Theoretical Computer Science*, 10(1–2): 1–157.

Xu, X.; Shen, Y.; Chi, Y.; and Ma, C. 2023. The Power of Preconditioning in Overparameterized Low-Rank Matrix Sensing. In Krause, A.; Brunskill, E.; Cho, K.; Engelhardt, B.; Sabato, S.; and Scarlett, J., eds., *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, 38611–38654. PMLR.

Zhang, R. Y. 2024. Improved Global Guarantees for the Nonconvex Burer–Monteiro Factorization via Rank Overparameterization. *Mathematical Programming*, 1-30.

Zhang, R. Y. 2025a. Improved Global Guarantees for the Nonconvex Burer–Monteiro Factorization via Rank Overparameterization. *Mathematical Programming*, 213(1–2): 1009–1038.

Zhang, R. Y. 2025b. Sharp global guarantees for nonconvex low-rank recovery in the noisy overparameterized regime. *SIAM Journal on Optimization*, 35(3): 2128–2154.

Zhang, R. Y.; Sojoudi, S.; and Lavaei, J. 2019. Sharp Restricted Isometry Bounds for the Inexistence of Spurious Local Minima in Nonconvex Matrix Recovery. *Journal of Machine Learning Research*, 20(114): 1–34.

Zhu, Z.; Li, Q.; Tang, G.; and Wakin, M. B. 2021. The global optimization geometry of low-rank matrix optimization. *IEEE Transactions on Information Theory*, 67(2): 1308–1331.

Zilber, P.; and Nadler, B. 2022. GNMR: A Provable One-Line Algorithm for Low Rank Matrix Recovery. *SIAM Journal on Mathematics of Data Science*, 4(2): 909–934.

Appendix

A Algorithmic and Experimental Details

This appendix provides the implementation and diagnostic details of Algorithm 1, followed by the complete experimental protocol and additional matrix-sensing results omitted from Section 6.2.

A.1 Algorithmic and Diagnostic Details

Algorithm 1 is implemented using only factor-space operations, applications of the sensing operator and its adjoint, and inverses of $k \times k$ matrices. It never materializes $M_t = X_t X_t^\top$ or the ambient normal block as a dense $n \times n$ matrix.

Matrix-free normal sketch. Let U_t be an orthonormal basis of $\mathrm{range}(X_t)$. Given $\Xi_t \in \mathbb{R}^{n\times s}$ with i.i.d. Gaussian entries, we form

$$Q_t = \mathrm{orth}\big(\Xi_t - U_t(U_t^\top \Xi_t)\big). \tag{40}$$

Thus, the projection onto $\mathrm{range}(X_t)^\perp$ uses only products with the thin matrix U_t. Writing the columns of Q_t as $q_{t,j}$, define

$$r_{t,j} = \left\langle G_t, q_{t,j} q_{t,j}^\top \right\rangle = \left\langle e_t, \mathcal{A}(q_{t,j} q_{t,j}^\top) \right\rangle,$$

where

$$e_t = \mathcal{A}(X_t X_t^\top) - y_w, \qquad G_t = \mathcal{A}^* e_t.$$

Negative values of $r_{t,j}$ detect negative normal curvature. Let

$$\mathcal{J}_t = \{j : r_{t,j} < 0\}.$$

The sampled normal-curvature penalty is

$$\widehat{\Phi}_{-,t} = \begin{cases} \dfrac{1}{2|\mathcal{J}_t|} \displaystyle\sum_{j\in\mathcal{J}_t} (-r_{t,j})^2, & \mathcal{J}_t \neq \varnothing, \\ 0, & \mathcal{J}_t = \varnothing. \end{cases}$$

The probes are held fixed while differentiating the current update. If

$$H_{t,j} = \mathcal{A}^* \mathcal{A}(q_{t,j} q_{t,j}^\top),$$

the resulting guidance direction is

$$g_{-,t} = -\frac{2}{|\mathcal{J}_t|} \sum_{j\in\mathcal{J}_t} (-r_{t,j}) H_{t,j} X_t \tag{41}$$

when $\mathcal{J}_t \neq \varnothing$, and zero otherwise. The negative sign moves the sampled negative Rayleigh quotients toward zero.

Inverse-trace direction. Let

$$S_t = X_t^\top X_t, \qquad \bar{\lambda}_t = \frac{\mathrm{tr}(S_t)}{k},$$

and set

$$\varepsilon_t = \varepsilon_{\mathrm{frac}} \max\{\bar{\lambda}_t, \epsilon_0\}.$$

With

$$R_t = (S_t + \varepsilon_t I_k)^{-1}, \qquad \tau_t = \bar{\lambda}_t + \varepsilon_t,$$

the scale-normalized inverse-trace penalty is

$$\Omega_{\mathrm{inv}}(X_t) = \tau_t \,\mathrm{tr}(R_t).$$

During the current update, ε_t is treated as fixed. We use the τ_t-normalized gradient direction

$$\begin{aligned} g_{\mathrm{inv},t} &:= \tau_t^{-1} \nabla_X \Omega_{\mathrm{inv}}(X_t) \\ &= 2X_t \left[\frac{\mathrm{tr}(R_t)}{k\tau_t} I_k - R_t^2 \right]. \end{aligned} \tag{42}$$

The R_t^2 term acts most strongly on the weakest factor directions, while the identity term controls the overall factor scale.

Periodic normal correction. Every h iterations during the guidance window, the implementation also applies a low-frequency correction. For a fresh set of normal probes, define

$$a_{t,j} = \frac{(-r_{t,j})_+}{\sum_{\ell=1}^s (-r_{t,\ell})_+}.$$

If all sampled Rayleigh quotients are nonnegative, the correction is skipped. Otherwise, set

$$\widetilde{v}_t = \sum_{j=1}^s a_{t,j} q_{t,j}, \qquad v_t = \frac{\widetilde{v}_t - U_t(U_t^\top \widetilde{v}_t)}{\|\widetilde{v}_t - U_t(U_t^\top \widetilde{v}_t)\|_2}.$$

We approximate the weakest right factor direction $u_t \in \mathbb{R}^k$ using four inverse iterations with a damped version of S_t. The correction is

$$X_{t+1} \leftarrow X_{t+1} + \gamma_t v_t u_t^\top, \tag{43}$$

where

$$\gamma_t = \min\left\{ \gamma_0 c_t \sqrt{\max_j (-r_{t,j})_+},\ \gamma_{\max} \sqrt{\bar{\lambda}_t} \right\}.$$

The cap prevents the correction from dominating the base optimization step.

Decay, preconditioning, and polishing. Let

$$T_g = \lfloor \theta T \rfloor, \qquad c_t = \left(1 - \frac{t}{T_g}\right)^3, \quad 0 \le t < T_g.$$

During this window, the factor-space direction is

$$d_t = 2G_t X_t + \alpha_0 c_t g_{\mathrm{inv},t} + \beta_0 c_t g_{-,t}.$$

We apply the damped Gram-preconditioned update

$$X_{t+1} = X_t - \eta_t d_t \left(S_t + \xi_t I_k\right)^{-1}, \qquad \xi_t = \xi \max\{\bar{\lambda}_t, \epsilon_0\}. \tag{44}$$

The learning rate uses a 25-iteration linear warm-up before reaching its nominal value. For $t \ge T_g$, all guidance terms and periodic corrections are removed, and the method optimizes only

$$f_w(X) = \frac{1}{2} \left\| \mathcal{A}(XX^\top) - y_w \right\|_2^2.$$

After native optimization, the resulting solution is polished on f_w for at most 500 additional steps. A reported critical point must remain full rank and satisfy

$$\frac{\|\nabla f_w(X)\|_F}{\max\{1, \|X\|_F\}} \le \varepsilon_{\mathrm{crit}}. \tag{45}$$

Therefore, guidance affects trajectory and critical-point selection, whereas stationarity is always tested on the original unmodified objective.

Computational cost. The inverse-trace direction requires forming S_t, solving a $k \times k$ linear system, and multiplying an $n \times k$ factor by a $k \times k$ matrix. Its additional cost is

$$O(nk^2 + k^3).$$

The normal guidance uses s rank-one sensing evaluations and adjoint-factor products together with an $n \times s$ QR factorization. Since s and k are small relative to n, the implementation stores only thin factor and sketch matrices and avoids ambient eigendecompositions.

Finite-probe diagnostic. Exact RIP certification is computationally intractable in general. We therefore use a finite-probe diagnostic only as an empirical screen. For each probe, draw an orthonormal matrix $V_j \in \mathbb{R}^{n\times q}$, where $q = \min\{2k, n\}$, and a normalized Gaussian vector $z_j \in \mathbb{R}^q$. Set

$$Z_j = V_j \operatorname{diag}(z_j) V_j^\top, \qquad \|Z_j\|_F = 1.$$

Using $N_{\mathrm{RIP}} = 24$ probes, we compute

$$\widehat{\delta} = \max_{1\le j\le N_{\mathrm{RIP}}} \left|\|\mathcal{A}(Z_j)\|_2^2 - 1\right|, \qquad \widehat{\mu} = 1 - \widehat{\delta} - \rho_k^-.$$

These quantities test only the sampled directions and are not certified RIP constants or theorem-level robustness certificates.

A.2 Matrix-Sensing Experiments

Data and sensing operators. For each instance, we generate

$$M_\star = U_\star \Lambda_\star U_\star^\top, \qquad \operatorname{rank}(M_\star) = r_\star,$$

where $U_\star$ is obtained by orthonormalizing a Gaussian matrix. The nonzero eigenvalues of $M_\star$ are geometrically spaced between 1.5 and 0.55.

We use the implicit bilinear sensing operator

$$\mathcal{A}_i(Z) = \frac{a_i^\top Z b_i}{\sqrt{m}}, \qquad m = 6nk.$$

For Gaussian sensing, a_i and b_i have independent standard normal entries. For Rademacher sensing, their entries are independent random signs. For weak-tail sensing, Gaussian vectors are decomposed into components inside and outside $\operatorname{range}(M_\star)$, and the orthogonal component is scaled by 0.35. This produces an anisotropic stress test with reduced visibility in slack directions. The target subspace is used only to generate the synthetic operator and is not supplied to any optimization method.

The formal grid is

$$(n, r_\star, k) \in \{(50,2,5), (75,3,6), (100,3,8), (125,4,10), (150,5,12)\}.$$

We use seeds $15, 66, 102, 999, 1325$, giving 25 matched configuration–seed instances per operator and 75 in total. Measurement perturbations satisfy

$$y_{w_j} = \mathcal{A}(M_\star) + w_j, \qquad \|w_j\|_2 = 0.08\|\mathcal{A}(M_\star)\|_2,$$

where $j = 1, 2, 3$ and each w_j is an independently sampled Gaussian direction normalized to the prescribed magnitude.

Non-oracle protocols. Both formal protocols construct their initializations only from observed measurements. Under E2E, the clean and perturbed runs use separately constructed observed-data initializations. In particular, Algorithm 1 and Plain use

$$\mathrm{SpecInit}(\mathcal{A}^* y_0) \quad \text{and} \quad \mathrm{SpecInit}(\mathcal{A}^* y_{w_j}),$$

respectively. Hence E2E includes sensitivity introduced by the initialization map.

Under Shared observed-spectral initialization, both members of each clean–perturbed pair use the same initialization constructed from y_{w_j}. This conditions out initialization differences and isolates the subsequent optimization trajectory. E2E is the primary protocol, while Shared is a controlled non-oracle comparison.

Evaluation metrics. Let $M_{p,s}^a(w_j)$ and $M_{p,s}^a(0)$ denote the polished critical points produced by method a, protocol p, and matched instance s. Define

$$d_{p,s,j}^a = \frac{\|M_{p,s}^a(w_j) - M_{p,s}^a(0)\|_F}{\|w_j\|_2}$$

and

$$\mathsf{R}_{p,s}^a = \operatorname{median}_{j=1,2,3} d_{p,s,j}^a.$$

The relative paired-sensitivity ratio is

$$\mathsf{Q}_{p,s}^a = \frac{\mathsf{R}_{p,s}^a}{\mathsf{R}_{p,s}^{\mathrm{Plain}}}.$$

For the primary noisy-data critical point M_s^a, the relative reconstruction error is

$$\mathsf{E}_{\mathrm{rel},s}^a = \frac{\|M_s^a - M_\star\|_F}{\|M_s^{\mathrm{Plain}} - M_\star\|_F}.$$

Values below one favor method a. The primary noisy-data runs are the same under E2E and Shared, so their $\mathsf{E}_{\mathrm{rel}}$ values coincide.

All formal runs are polished on f_w using tolerance $\varepsilon_{\mathrm{crit}} = 0.025$. Every Algorithm 1/Plain matched row passes the paired criticality check for all three perturbations under both protocols. All reported primary solutions also satisfy $\lambda_k(XX^\top) > 10^{-8}$.

Optimization settings. Algorithm 1 and Plain use learning rate 0.035 and at most 220 native iterations. Algorithm 1 uses a 25% guidance window with cubic decay, inverse-trace weight 0.003, Rayleigh-gradient weight 0.02, correction weight 0.08, sketch size 8, correction period 10, and Gram damping fraction 0.05.

L1-SubGM (Ma and Fattahi 2023) uses step size 0.04, decay 0.99, and initialization scale 0.05. ScaledSM (Tong, Ma, and Chi 2021b) uses scale parameter 0.03 and decay 0.99. GNMR (Zilber and Nadler 2022) uses 40 LSQR iterations per Gauss–Newton step. Since these methods have different native objectives, all their reported solutions are subsequently polished on the common squared objective f_w.

Parameter	Value
Measurements	$m = 6nk$
Formal seeds	$15, 66, 102, 999, 1325$
Native iterations	at most 220
Algorithm 1/Plain learning rate	0.035
Relative perturbation norm	0.08
Paired perturbations	3
Guidance fraction	0.25
Guidance decay power	3
Normal sketch size	8
Normal correction period	10
Common polishing steps	at most 500
Criticality tolerance	0.025
Finite-RIP probes	24

Table 2: Common settings for the formal matrix-sensing experiments.

Operator	$(n, r_\star, k)$	$\mathsf{Q}_{\mathrm{E2E}}$	$\mathsf{Q}_{\mathrm{Shared}}$	$\mathsf{E}_{\mathrm{rel}}$
Gauss.	$(50, 2, 5)$	.8323	.8161	.9375
Gauss.	$(75, 3, 6)$	.8269	.8123	.9162
Gauss.	$(100, 3, 8)$	.7883	.7780	.8874
Gauss.	$(125, 4, 10)$	.7651	.7516	.8864
Gauss.	$(150, 5, 12)$	.7483	.7378	.8739
Radem.	$(50, 2, 5)$	.8031	.7991	.9322
Radem.	$(75, 3, 6)$	.8064	.7898	.9338
Radem.	$(100, 3, 8)$	.7884	.7708	.8668
Radem.	$(125, 4, 10)$	.7658	.7504	.8487
Radem.	$(150, 5, 12)$	.7459	.7333	.8347
Weak-tail	$(50, 2, 5)$	.1722	.1171	.5960
Weak-tail	$(75, 3, 6)$	.1768	.1337	.5512
Weak-tail	$(100, 3, 8)$	.1827	.1128	.6250
Weak-tail	$(125, 4, 10)$	.1865	.1119	.6387
Weak-tail	$(150, 5, 12)$	.2683	.1296	.6468

Table 3: Configuration-level medians for Algorithm 1 relative to Plain. Each row aggregates five seeds. Values below one favor Algorithm 1.

Algorithm 1 versus Plain. Table 3 reports configuration-level medians over five seeds. The two Q columns correspond to E2E and Shared observed-spectral initialization. The reconstruction ratio is common to both protocols because their primary noisy-data runs are identical.

All configuration-level medians satisfy $\mathsf{Q} < 1$ and $\mathsf{E}_{\mathrm{rel}} < 1$. At the instance level, both inequalities hold in $23/25$, $25/25$, and $25/25$ cases for Gaussian, Rademacher, and weak-tail sensing, respectively, under both protocols. Shared gives uniformly smaller configuration-level Q values because it conditions out initialization sensitivity. Its weak-tail gains should therefore be read as controlled evidence under the anisotropic stress test, not as evidence that weak-tail sensing is easier.

Figure 3 complements the direct paired measurements with terminal solution geometry. The median ρ^- ratios are 1.476, 1.459, and 2.061, while the median active-floor ratios are 0.755, 0.735, and 0.634 for Gaussian, Rademacher, and weak-tail sensing. Thus, the reduction in Q is not accompanied by uniform improvement in either terminal diagnostic. These quantities describe local geometry at the polished critical point; neither is intended to replace the directly measured paired sensitivity.

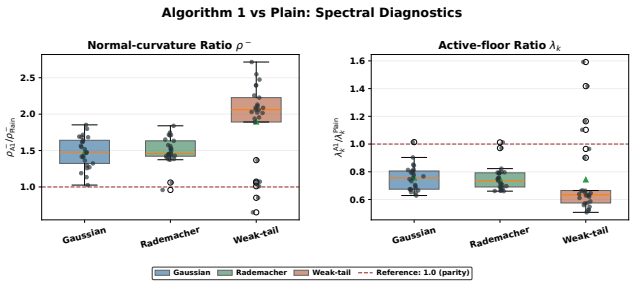

Figure 3: Terminal spectral diagnostics for Algorithm 1 relative to Plain. The panels show $\rho^-_{\mathrm{A1}}/\rho^-_{\mathrm{Plain}}$ and $\lambda_k^{\mathrm{A1}}/\lambda_k^{\mathrm{Plain}}$, respectively. Each operator pools 25 matched instances. Smaller values favor Algorithm 1 for ρ^-, whereas larger values favor it for λ_k. E2E and Shared use the same primary noisy-data runs, so these diagnostics coincide and are shown only once.

External comparators. All external methods are evaluated relative to Plain using the same metrics and common f_w-polishing. E2E recomputes the observed-spectral initialization for each paired perturbation, whereas Shared reuses a common observed-spectral initialization. Consequently, only Q changes between protocols; reconstruction, ρ^-, and λ_k are identical up to numerical precision. Figures 4–6 report the corresponding operator-wise distributions. All displayed paired solutions pass the common f_w-criticality check.

L1-SubGM (Ma and Fattahi 2023) has pooled median Q values 0.779 under E2E and 0.776 under Shared, together with pooled $\mathsf{E}_{\mathrm{rel}} = 0.791$. Both metrics are below one in all 75 instances, with the largest reductions under weak-tail sensing. Its terminal spectral ratios nevertheless move in the opposite directions, especially for weak-tail sensing, illustrating that paired sensitivity is not determined by either diagnostic alone.

ScaledSM (Tong, Ma, and Chi 2021b) has pooled median Q values 1.148 and 1.155 under E2E and Shared, respectively, and pooled $\mathsf{E}_{\mathrm{rel}} = 1.146$. It lowers both paired sensitivity and reconstruction error in all 25 weak-tail instances, but in none of the 50 Gaussian or Rademacher instances. Its benefit is therefore specific to the anisotropic stress test rather than uniform across sensing operators.

GNMR (Zilber and Nadler 2022) has pooled median $\mathsf{Q} \simeq 1.049$ under both protocols and pooled $\mathsf{E}_{\mathrm{rel}} = 1.048$, with only one joint improvement among the 75 instances. Although its median ρ^- ratios are below one and its median active-floor ratios are above one, these favorable terminal directions do not produce a consistent reduction in paired sensitivity.

Taken together, the comparator results show that terminal ρ^- and λ_k ratios do not by themselves rank methods by paired sensitivity. The theoretical certificates use these quan-

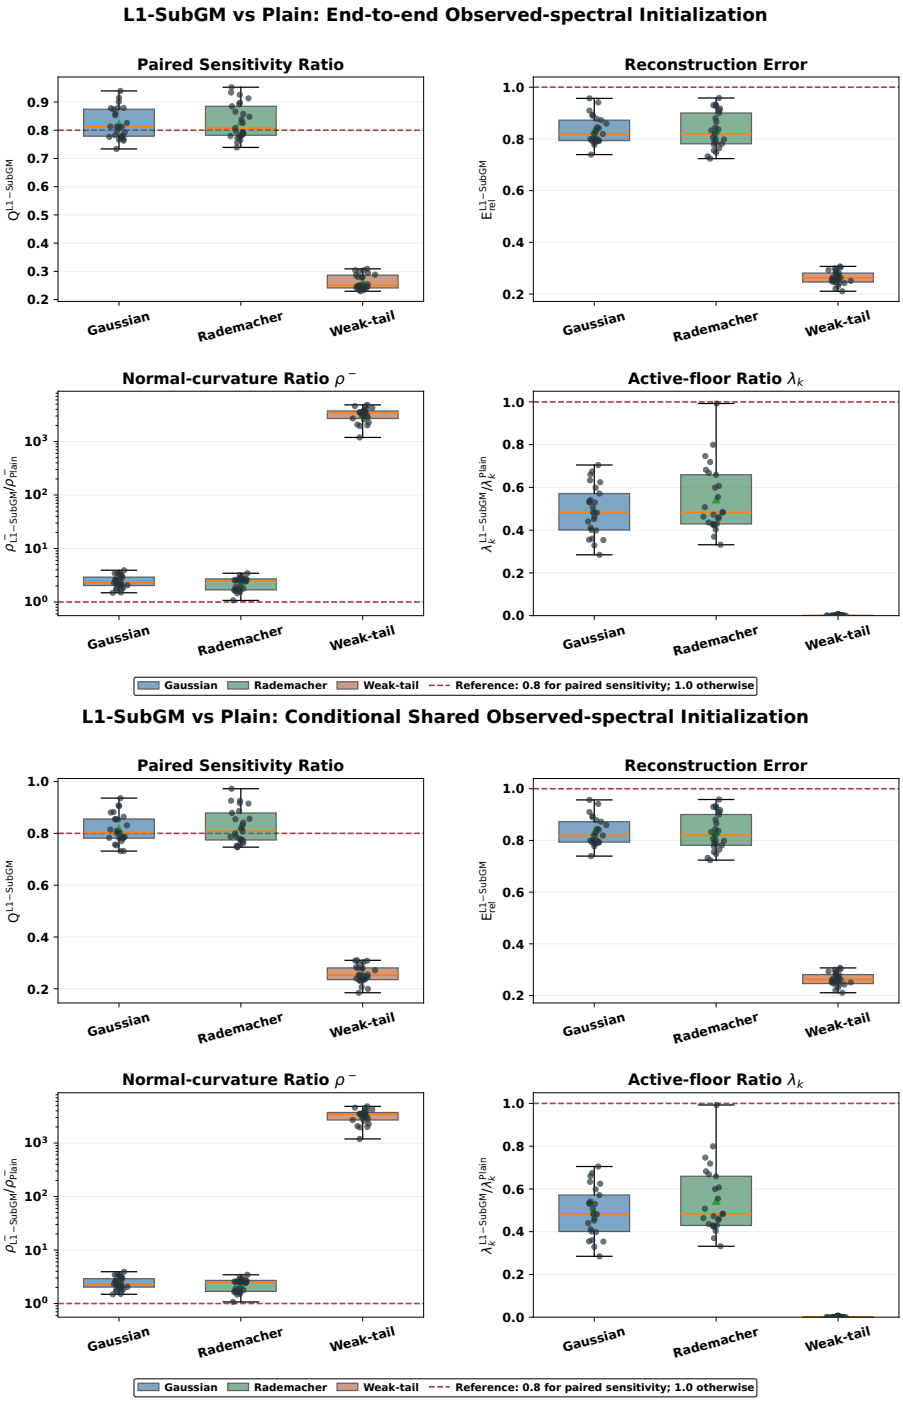

Figure 4: L1-SubGM relative to Plain under E2E (top) and Shared (bottom). From upper left to lower right, the panels show Q, $\mathsf{E}_{\mathrm{rel}}$, the ρ^- ratio, and the λ_k ratio. The dashed references are $\mathsf{Q} = 0.8$ and one for the other metrics. The ρ^- panel uses a logarithmic scale because of the weak-tail range.

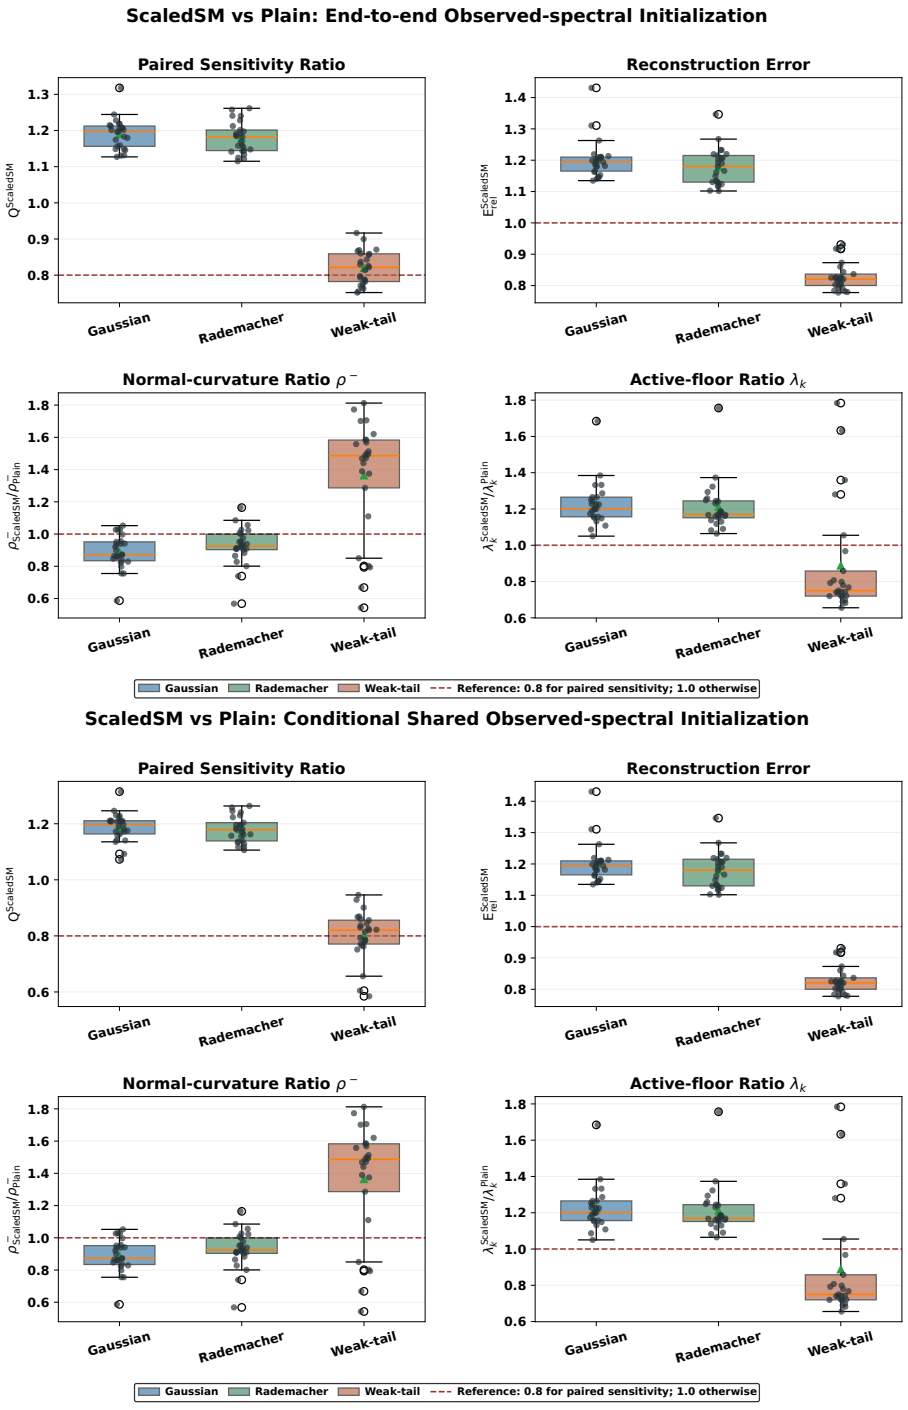

Figure 5: ScaledSM relative to Plain under E2E (top) and Shared (bottom), using the same four diagnostics and reference lines as Figure 4. Each operator contains 25 matched configuration–seed instances.

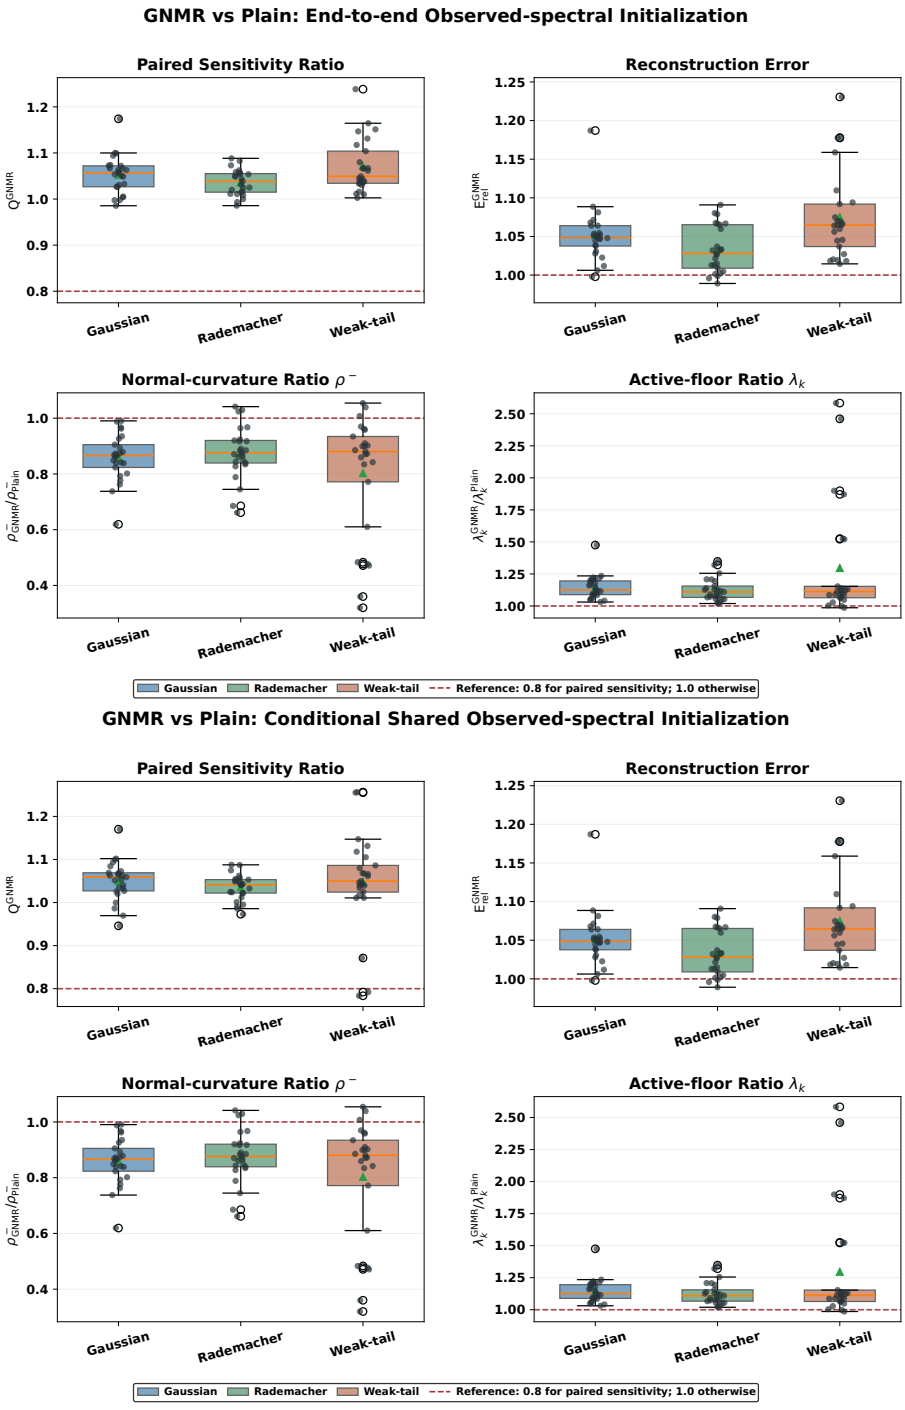

Figure 6: GNMR relative to Plain under E2E (top) and Shared (bottom), using the same four diagnostics and reference lines as Figure 4. Each operator contains 25 matched configuration–seed instances.

tities jointly with the remaining local and operator-dependent conditions; they do not assert cross-method monotonicity in either terminal ratio alone. Accordingly, Q remains the direct experimental measure, while ρ^- and λ_k provide complementary geometric diagnostics.

Large-scale experiments. We evaluate whether the paired-sensitivity improvements persist as the ambient dimension increases. We use

$$(n, r_\star, k) \in \{(250, 4, 10), (500, 5, 12), (1000, 5, 12)\},$$

with Gaussian, Rademacher, and weak-tail sensing, five seeds per configuration, and $m = 6nk$ measurements. The relative measurement-noise level and paired-perturbation norm are both 0.08, and each paired critical-point sensitivity is aggregated over three perturbations. Algorithm 1 and Plain use at most 220 native iterations, followed by the same f_w-polishing procedure for at most 500 iterations. The criticality tolerance is 0.025.

We evaluate both the primary E2E protocol and the controlled Shared observed-spectral protocol. In E2E, the initialization changes with the perturbed observations. Shared instead uses the same observed-spectral initialization for the two objectives in each matched pair, thereby conditioning out initialization sensitivity. The main noisy-data runs are identical under the two protocols; hence their reconstruction and terminal spectral diagnostics coincide, while their relative paired sensitivities Q can differ.

Table 4 shows that the improvements persist through $n = 1000$. Under E2E, the Gaussian and Rademacher median Q values remain between 0.709 and 0.722, giving approximately 28%–29% lower paired critical-point sensitivity than Plain at every dimension. For weak-tail sensing, $\mathsf{Q}_{\mathrm{E2E}}$ increases from 0.307 at $n = 250$ to 0.570 at $n = 1000$, but remains below one and therefore still gives a 43.0% reduction at the largest scale.

Conditioning out initialization sensitivity produces smaller Shared ratios. The Gaussian and Rademacher values decrease to approximately 0.684 at $n = 1000$, while the weak-tail values remain between 0.120 and 0.157. Shared is retained as a controlled diagnostic, rather than replacing the primary E2E result.

Reconstruction also improves uniformly. The median $\mathsf{E}_{\mathrm{rel}}$ values lie between 0.821 and 0.840 for the isotropic operators and between 0.584 and 0.668 for weak-tail sensing. Consequently, both $\mathsf{Q} < 1$ and $\mathsf{E}_{\mathrm{rel}} < 1$ hold in all 45 matched instances under each protocol.

The lower panels of Figures 7 and 8 provide complementary terminal spectral diagnostics. The pooled median normal-curvature ratios are 1.527, 1.454, and 1.489, while the corresponding active-floor ratios are 0.699, 0.704, and 0.619 for Gaussian, Rademacher, and weak-tail sensing. These terminal quantities do not vary monotonically with the directly measured paired sensitivity and are therefore reported as diagnostics rather than substituted for Q.

All main solutions remain full rank, with minimum $\lambda_k = 0.0192$. The largest final relative gradient norm is 0.0249, below the prescribed 0.025 tolerance, and every paired solution used to compute Q passes the same criticality check.

A.3 LoRA Fine-Tuning Experiments

From Matrix-Sensing Guidance to LoRA We evaluate whether the guidance principle developed in Algorithm 1 can improve low-rank adaptation of a pretrained language model. Algorithm 1 is motivated by the matrix-sensing robustness ratio

$$\rho_k^- = \frac{b^-}{\lambda_k(M)},$$

which identifies two complementary mechanisms: protecting the active spectral floor in the denominator and controlling leakage outside the active subspace in the numerator. The LoRA experiment is a practical adaptation of this principle rather than a direct application of the matrix-sensing theorems, since LoRA updates are rectangular matrices and do not possess the same positive-semidefinite manifold geometry.

For layer ℓ, LoRA parameterizes the trainable update as

$$\Delta W_\ell = \frac{\alpha}{r} B_\ell A_\ell,$$

where r is the nominal LoRA rank. Let $s_{\ell,1} \geq \cdots \geq s_{\ell,r}$ denote the singular values of ΔW_ℓ. At the beginning of the guidance window, we freeze a layer-specific active rank q_ℓ as the smallest rank explaining at least 95% of the spectral energy. Layers for which $q_\ell = r$ are excluded from spectral guidance.

For the retained active directions, we use the scale-normalized inverse-trace penalty

$$\Omega_{\ell,\mathrm{floor}} = \left(\frac{1}{q_\ell} \sum_{i=1}^{q_\ell} s_{\ell,i}^2 + \epsilon \right) \left(\frac{1}{q_\ell} \sum_{i=1}^{q_\ell} \frac{1}{s_{\ell,i}^2 + \epsilon} \right) - 1.$$

This term places stronger pressure on weak directions within the frozen active spectrum. To avoid improving the active floor by simultaneously inflating unused directions, we add the normalized tail penalty

$$\Omega_{\ell,\mathrm{tail}} = \frac{\sum_{i=q_\ell+1}^{r} s_{\ell,i}^2}{\mathrm{stopgrad}\left(\sum_{i=1}^{q_\ell} s_{\ell,i}^2 \right) + \epsilon}.$$

The guided objective during the active window is

$$\mathcal{L}_{\mathrm{guided}} = \mathcal{L}_{\mathrm{task}} + c_t \left(10^{-2}\, \overline{\Omega}_{\mathrm{floor}} + 10^{-3}\, \overline{\Omega}_{\mathrm{tail}} \right),$$

where the bars denote averages over guided LoRA layers. Guidance begins after 10% of training, remains active for 20% of the training trajectory, and uses a short fade-in followed by cubic decay. Its contribution is clipped to at most 10% of the detached task loss.

Following the damped Gram preconditioning in Algorithm 1, scaled-active LoRA also preconditions the factor gradients as

$$\begin{aligned} \widetilde{\nabla}_{B_\ell} &= \nabla_{B_\ell} \left(A_\ell A_\ell^\top + \xi I \right)^{-1}, \\ \widetilde{\nabla}_{A_\ell} &= \left(B_\ell^\top B_\ell + \xi I \right)^{-1} \nabla_{A_\ell}. \end{aligned} \tag{46}$$

We use damping $\xi = 0.1$, a damping floor of 10^{-4}, and cap the preconditioning gain at 10. The normal-Rayleigh guidance term from Algorithm 1 is disabled in this experiment,

n	Operator	$\mathsf{Q}_{\mathrm{E2E}}$	Δ_{E2E}	Q_{Sh}	Δ_{Sh}	$\mathsf{E}_{\mathrm{rel}}$	$\rho^{-}_{\mathrm{Plain}}$	ρ^{-}_{A1}	Joint
250	Gauss.	0.715	28.5%	0.710	29.0%	0.821	0.2624	0.4093	5/5
250	Radem.	0.712	28.8%	0.706	29.4%	0.824	0.2543	0.3954	5/5
250	Weak-tail	0.307	69.3%	0.125	87.5%	0.668	0.0077	0.0084	5/5
500	Gauss.	0.713	28.7%	0.703	29.7%	0.823	0.1826	0.2850	5/5
500	Radem.	0.722	27.8%	0.702	29.8%	0.834	0.1951	0.3058	5/5
500	Weak-tail	0.385	61.5%	0.120	88.0%	0.586	0.0029	0.0061	5/5
1000	Gauss.	0.709	29.1%	0.684	31.6%	0.835	0.1287	0.1950	5/5
1000	Radem.	0.720	28.0%	0.684	31.6%	0.840	0.1384	0.1934	5/5
1000	Weak-tail	0.570	43.0%	0.157	84.3%	0.584	0.0031	0.0041	5/5

Table 4: Formal large-scale comparison of Algorithm 1 with Plain over five seeds per $(n, \mathrm{operator})$ cell. The quantities $\mathsf{Q}_{\mathrm{E2E}}$ and Q_{Sh} are the median instancewise relative paired critical-point sensitivities under E2E and Shared observed-spectral initialization. The corresponding sensitivity reductions are $\Delta_{\mathrm{E2E}} = 100(1 - \mathsf{Q}_{\mathrm{E2E}})$ and $\Delta_{\mathrm{Sh}} = 100(1 - \mathsf{Q}_{\mathrm{Sh}})$. The reconstruction ratio is common to both protocols because their main noisy-data runs are identical. The two ρ^{-} columns report the median terminal normal-curvature ratios of Plain and Algorithm 1; these values are likewise common to the two protocols up to numerical precision. Joint counts seeds for which both $\mathsf{Q} < 1$ and $\mathsf{E}_{\mathrm{rel}} < 1$; all cells have $5/5$ under both protocols.

because a rectangular LoRA update has no direct theorem-preserving analogue of the positive-semidefinite normal projector. Consequently, scaled-active LoRA should be interpreted as an active-floor and tail-guided method combined with damped LoRA preconditioning.

The plain baseline uses the same initialization, data order, optimizer, training seed, and objective perturbation, but applies neither spectral guidance nor ScaledLoRA preconditioning. After the guidance window, the spectral penalties are removed and the endpoint is polished on the original training objective. In the reported legacy protocol, the scaled-active branch retains its damped preconditioner during this polish.

Fixed Logit-Objective Perturbation To construct an analogue of matrix-sensing measurement noise, we perturb the training objective rather than the validation inputs. For training example n and perturbation seed s, the perturbed objective is

$$\mathcal{L}_{\varepsilon,s}(\theta) = \frac{1}{N} \sum_{n=1}^{N} \mathrm{CE}\left(f_\theta(x_n) + \varepsilon u_{s,n}, y_n\right),$$

where $u_{s,n}$ is generated deterministically from the pair consisting of the perturbation seed and training-example identifier. It satisfies

$$\mathbf{1}^\top u_{s,n} = 0, \qquad \|u_{s,n}\|_2 = 1,$$

and is fixed across all epochs. Labels are unchanged and no perturbation is applied during validation. Therefore, noisy validation accuracy refers to the clean-validation accuracy of an endpoint trained under the perturbed objective; it is not accuracy under corrupted validation text.

The prespecified experiment uses $\varepsilon = 0.25$ and three fixed perturbation seeds, 10000, 10001, and 10002. We additionally evaluate $\varepsilon \in \{0.05, 0.1, 0.2, 0.3, 0.4, 0.5\}$ without changing any other training setting.

Experimental configuration. We fine-tune `roberta-base` (Liu et al. 2019) on the SST-2 task (Socher et al. 2013), using the version distributed in GLUE (Wang et al. 2018b). Following LoRA (Hu et al. 2022b), adapters are inserted into all attention query and value projections with rank $r = 8$, scaling $\alpha = 16$, and zero dropout. The classification head is trainable in both conditions. We use a maximum sequence length of 128, training batch size 32, evaluation batch size 64, and train for two epochs.

Both methods use SGD with learning rate 10^{-2}, momentum 0.9, zero weight decay, and a learning-rate warm-up ratio of 0.06. The legacy protocol does not use a separate shared warm-up stage. We evaluate training seeds 0, 1, and 2. For each training seed, the plain and scaled-active runs share their initialization and are evaluated using the same three fixed objective-perturbation directions.

For aggregation, noisy validation accuracy is first averaged over the three perturbation directions within each training seed. We then report the mean and sample standard deviation over the three training seeds. The perturbation directions are not treated as independent training repetitions.

Results At the prespecified perturbation magnitude $\varepsilon = 0.25$, plain LoRA achieves a clean validation accuracy of $91.093 \pm 0.331\%$, whereas scaled-active LoRA achieves $92.087 \pm 0.303\%$. This corresponds to a paired clean-accuracy improvement of 0.994 ± 0.066 percentage points.

For endpoints trained under the fixed perturbed objective, plain LoRA achieves $91.106 \pm 0.277\%$ validation accuracy and scaled-active LoRA achieves $92.125 \pm 0.268\%$. The resulting paired noisy-accuracy improvement is 1.019 ± 0.022 percentage points. Scaled-active outperforms plain LoRA for each of the three training seeds.

Figure 9 evaluates sensitivity to the perturbation magnitude. Across the six tested magnitudes, the scaled-active advantage remains between 0.943 and 1.007 percentage points. It is positive for every training-seed and perturbation-magnitude combination. These 18 comparisons reuse the same three training seeds and perturbation directions and are therefore not interpreted as 18 independent replications.

The two accuracy curves are approximately parallel. Thus,

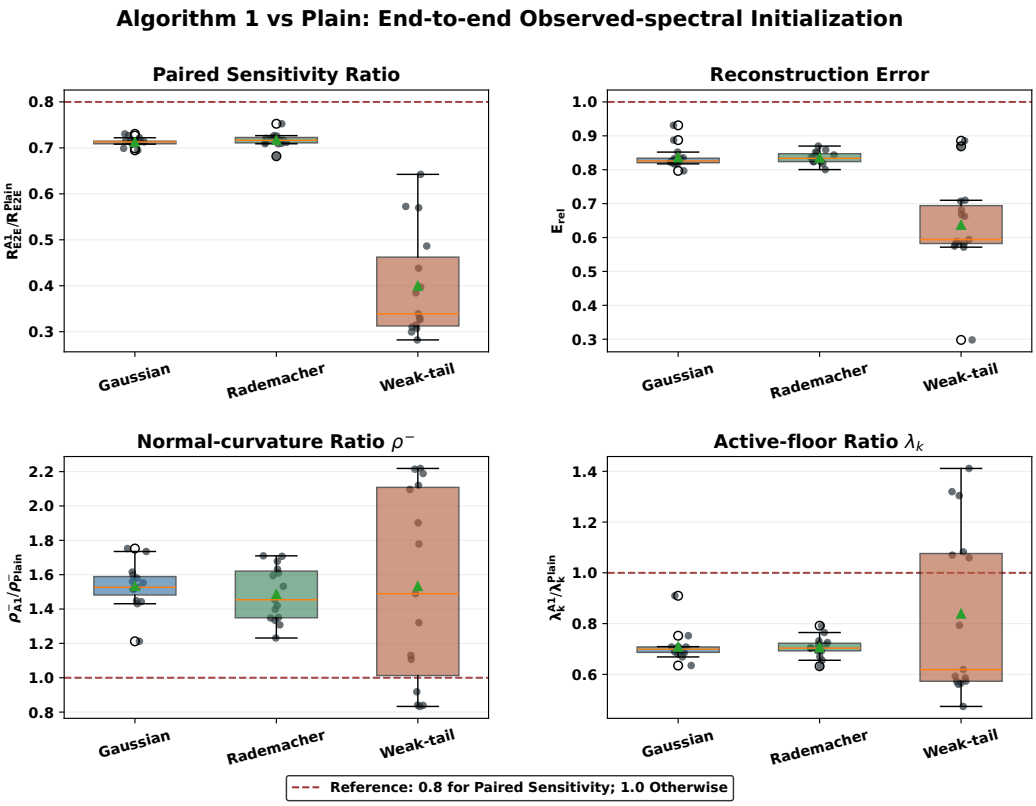

Figure 7: Formal large-scale comparison of Algorithm 1 with Plain. Each row contains the five seeds for one fixed $(n, \text{operator})$ cell. The quantities $\mathsf{Q}_{\mathrm{E2E}}$ and Q_{Sh} are the median instancewise relative paired critical-point sensitivities under E2E and Shared observed-spectral initialization. Their corresponding sensitivity reductions are $\Delta_{\mathrm{E2E}} = 100(1 - \mathsf{Q}_{\mathrm{E2E}})$ and $\Delta_{\mathrm{Sh}} = 100(1 - \mathsf{Q}_{\mathrm{Sh}})$. The reconstruction ratio $\mathsf{E}_{\mathrm{rel}}$ and the terminal ρ^-_{Plain} and ρ^-_{A1} values are common to both protocols because their main noisy-data runs are identical. Joint counts seeds for which both $\mathsf{Q} < 1$ and $\mathsf{E}_{\mathrm{rel}} < 1$; every row has $5/5$ under both protocols. Pooling the three dimensions gives $15/15$ joint improvements for each operator and $45/45$ per protocol.

the experiment supports a consistent improvement in absolute validation accuracy at the perturbed training endpoint, but does not establish a smaller clean-to-noisy accuracy degradation.

B Missing Details and Proofs

B.1 More Definitions

Definition 2 (Second-order stationary point)**.** A point $X_w \in \mathbb{R}^{n\times k}$ is a second-order stationary point of f_w if

$$\nabla f_w(X_w) = 0,$$
$$\nabla^2 f_w(X_w)[D, D] \geq 0 \quad \text{for every } D \in \mathbb{R}^{n\times k}.$$

We write $M_w = X_w X_w^\top$.

Definition 3 (Restricted isometry)**.** The operator $\mathcal{A}$ satisfies (δ, s)-RIP if

$$(1-\delta)\,\|Z\|_F^2 \leq \|\mathcal{A}(Z)\|_2^2 \leq (1+\delta)\,\|Z\|_F^2$$

for every $Z \in \mathbb{S}^n$ with $\operatorname{rank}(Z) \leq s$.

Assumption 1 (Unified $2k$-RIP)**.** Unless stated otherwise, $\mathcal{A}$ satisfies $(\delta, 2k)$-RIP for a fixed $\delta \in [0, 1)$. Since $r_\star \leq k$, this assumption covers all rank-at-most-$2k$ symmetric matrices used below, including differences between rank-at-most-k PSD matrices. We set $s_{\mathrm{RIP}} = 2k$ throughout.

For $M \succeq 0$, let its nonzero eigenvalues be

$$\lambda_1(M) \geq \cdots \geq \lambda_{\operatorname{rank}(M)}(M) > 0.$$

When an effective active rank q is selected, $P_q(M)$ denotes the projector onto the top-q eigenspace and $P_q^\perp = I - P_q$. The tangent space of the rank-q PSD manifold at a rank-q matrix with range projector P is

$$\mathcal{T}(P) = \{\eta \in \mathbb{S}^n : \; P^\perp \eta P^\perp = 0\}. \tag{47}$$

Every $\eta \in \mathcal{T}(P)$ has rank at most $2q$.

B.2 Missing Details of Section 4

Lemma 1 (Factor Hessian identity)**.** *Let $M = XX^\top$. For every $D \in \mathbb{R}^{n\times k}$,*

$$\nabla^2 f_w(X)[D, D] = \left\|\mathcal{A}(XD^\top + DX^\top)\right\|_2^2 + 2\left\langle G_w(M), DD^\top \right\rangle. \tag{48}$$

Moreover, X is first-order critical iff

$$G_w(M)X = 0. \tag{49}$$

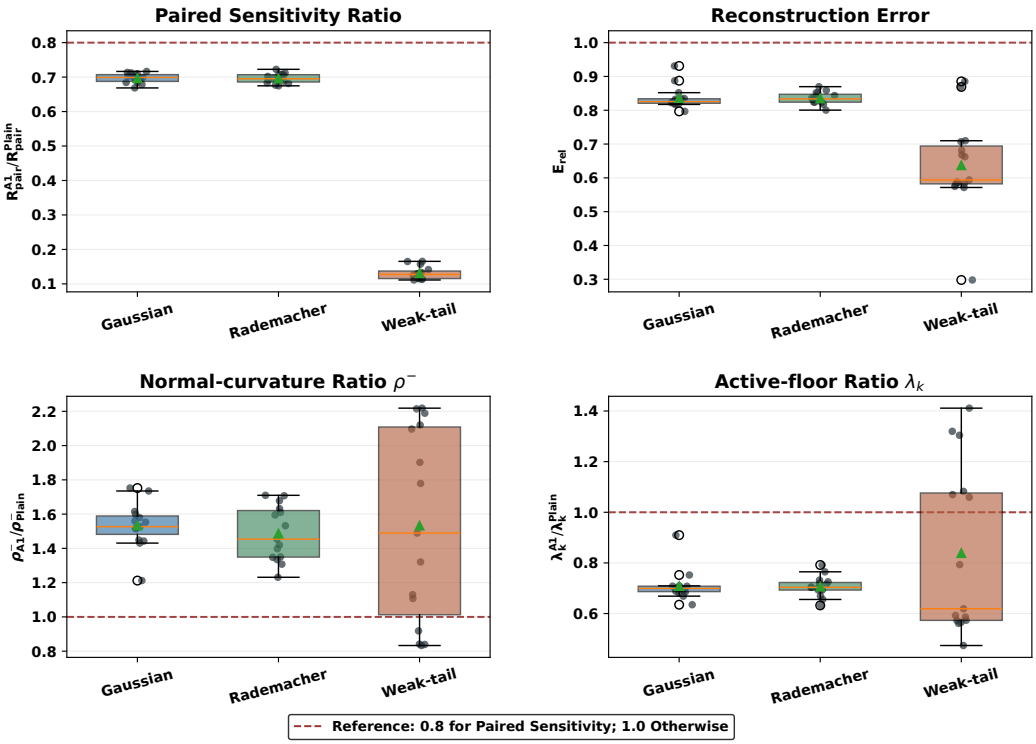

Figure 8: Formal large-scale comparison under Shared observed-spectral initialization. Only the paired-sensitivity panel changes relative to Figure 7; reconstruction and terminal spectral diagnostics coincide because the two protocols use the same main noisy-data solutions.

Proof of Lemma 1. The first and second variations of $XX^\top$ in direction D are

$$\eta = XD^\top + DX^\top, \qquad 2DD^\top.$$

Applying the chain rule to $F_w(XX^\top)$ gives the gradient $2G_w(M)X$ and the quadratic form (48). □

Remark 2 (Order of magnitude of chart constants). The raw chart is well conditioned when the active spectrum is bounded: if $0 < c \le \lambda_k(M_w) \le \lambda_1(M_w) \le C$, then $\kappa_{\mathrm{raw}}(M_w) = O(1)$. If $\lambda_k(M_w) = \varepsilon \ll 1$ while $\lambda_1(M_w) = O(1)$, then $\kappa_{\mathrm{raw}}(M_w) = \Theta(\varepsilon^{-1})$. If $\lambda_1(M_w) = L \gg 1$ and $\lambda_k(M_w) = \Omega(1)$, then $\kappa_{\mathrm{raw}}(M_w) = \Theta(L)$. These scalings are coordinate geometry, not RIP. RIP controls $\mathcal{A}$ on tangent matrices, but it does not control the conditioning of a raw SVD/graph chart near the rank-deficient boundary.

Proposition 1. *Let $X \in \mathbb{R}^{n\times q}$ have rank q, let $M = XX^\top$, let P be the range projector of M, and assume $G_w(M)X = 0$. For any factor direction D producing the tangent matrix*

$$\eta = XD^\top + DX^\top \in \mathcal{T}(P),$$

the factor Hessian obeys

$$\nabla^2 f_w(X)[D,D] \ge \left(\alpha_T(P) - \rho_q^-(M;w)\right)\|\eta\|_F^2. \quad (50)$$

Consequently, under Assumption 1,

$$\nabla^2 f_w(X)[D,D] \ge \left(1 - \delta - \rho_q^-(M;w)\right)\|\eta\|_F^2. \quad (51)$$

Proof. Write $M = U\Lambda U^\top$, where $\Lambda = \mathrm{diag}(\lambda_1,\dots,\lambda_q)$ and $\lambda_q > 0$. Since $XX^\top = M$ and X has full column rank, after a right-orthogonal change of factor coordinates we may take $X = U\Lambda^{1/2}$. Decompose a factor perturbation as

$$D = UA + U_\perp B.$$

Then

$$\eta = U(\Lambda^{1/2}A^\top + A\Lambda^{1/2})U^\top + U\Lambda^{1/2}B^\top U_\perp^\top + U_\perp B\Lambda^{1/2}U^\top.$$

The normal block of $DD^\top$ is

$$P^\perp DD^\top P^\perp = U_\perp BB^\top U_\perp^\top.$$

Since $G_w(M)X = 0$, the range and cross terms of $G_w(M)$ vanish, so

$$\left\langle G_w(M), DD^\top\right\rangle = \left\langle P^\perp G_w(M)P^\perp, U_\perp BB^\top U_\perp^\top\right\rangle.$$

Let $K = B\Lambda^{1/2}$. The off-diagonal block of η is K, hence

$$\|\eta\|_F^2 \ge 2\|K\|_F^2,$$

$$\mathrm{tr}(BB^\top) = \mathrm{tr}(K\Lambda^{-1}K^\top) \le \frac{\|K\|_F^2}{\lambda_q} \le \frac{\|\eta\|_F^2}{2\lambda_q}.$$

Therefore

$$2\left\langle G_w(M), DD^\top\right\rangle \ge -\frac{\left\|[P^\perp G_w(M)P^\perp]_-\right\|_{\mathrm{op}}}{\lambda_q}\|\eta\|_F^2.$$

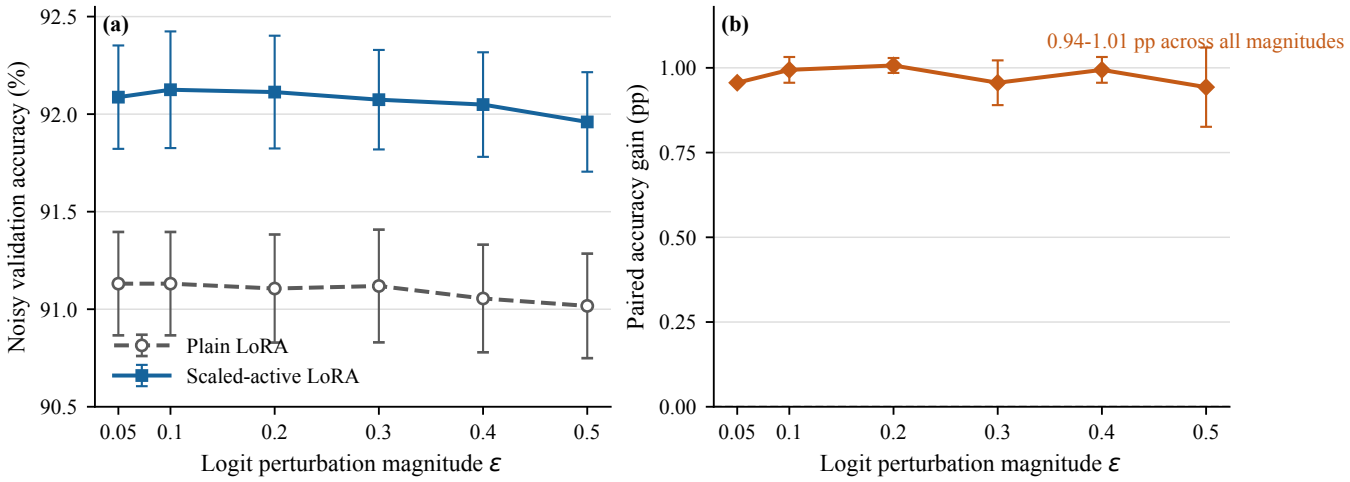

Figure 9: Robust validation performance under fixed unit-ℓ_2 logit-objective perturbations. Panel (a) reports noisy validation accuracy for plain and scaled-active LoRA. Panel (b) reports the paired scaled-active-minus-plain accuracy difference. Each point is the mean over three training seeds, after first averaging three fixed perturbation directions within each seed. Error bars denote sample standard deviations across training seeds.

The first term in (48) is at least $\alpha_T(P)\,\|\eta\|_F^2$, proving (50). The RIP specialization follows from (7). □

Assumption 2 (RIP-certified endpoint Hessian lower margin)**.** Let $M_w \in \mathcal{M}_k$ satisfy (10), and set

$$\mu_w = (1-\delta) - \rho_k^-(M_w; w),$$
$$\lambda_w = \lambda_k(M_w), \qquad P_w = P_k(M_w),$$
$$\gamma_w = \left\| P_w^\perp G_w(M_w) P_w^\perp \right\|_{\mathrm{op}}.$$

Equivalently,

$$\rho_k^-(M_w; w) = \frac{\max\{0, -\lambda_{\min}(P_w^\perp G_w(M_w) P_w^\perp)\}}{\lambda_k(M_w)}. \quad (52)$$

We assume $\mu_w > 0$.

Assumption 3 (Small noise relative to the endpoint scale)**.** With $\mu_w, \lambda_w, \gamma_w$ from Assumption 2, define the computable endpoint variation scale

$$\mathfrak{L}_w = \frac{2(1+\delta+\rho_k^-(M_w; w))}{\lambda_w} + \frac{16\gamma_w}{\lambda_w^2}, \quad (53)$$

with the convention that $\mu_w/(4\mathfrak{L}_w) = +\infty$ if $\mathfrak{L}_w = 0$. The noise obeys

$$\frac{2\sqrt{1+\delta}}{\mu_w}\,\|w\|_2 \le \frac{1}{2} r_w^{\mathrm{stab}}, \qquad r_w^{\mathrm{stab}} = \min\left\{\frac{\lambda_w}{4}, \frac{\mu_w}{4\mathfrak{L}_w}\right\}. \quad (54)$$

Why the margin cannot be dropped. If the endpoint Hessian has a flat eigen-direction, linear robustness can fail even in one dimension. Consider

$$\phi_0(x) = \frac{1}{4}x^4, \qquad \phi_\varepsilon(x) = \frac{1}{4}x^4 - \varepsilon x.$$

The unperturbed critical point is $x_0 = 0$, but $\phi_0''(0) = 0$. The perturbed critical point solves

$$\phi_\varepsilon'(x) = x^3 - \varepsilon = 0, \qquad x_\varepsilon = \varepsilon^{1/3}.$$

Thus $|x_\varepsilon - x_0| = \varepsilon^{1/3}$, which is not bounded by $C\varepsilon$ for any fixed C as $\varepsilon \downarrow 0$. A positive margin is exactly what rules out this flat-bottom behavior and lets the criticality equation be inverted.

Lemma 2 (Small-noise critical-point continuation)**.** *Suppose Assumptions 1–3 hold. Then there exists a C^1 curve*

$$s \mapsto M_s \in \mathcal{M}_k, \qquad 0 \le s \le 1,$$

with $M_1 = M_w$, such that every M_s is a rank-k critical point of F_s:

$$\mathrm{grad}_{\mathcal{M}_k} F_s(M_s) = 0. \quad (55)$$

Moreover, along this branch,

$$\langle \eta, \mathrm{Hess}_{\mathcal{M}_k} F_s(M_s)[\eta] \rangle \ge \frac{\mu_w}{2}\,\|\eta\|_F^2, \qquad \eta \in T_{M_s}\mathcal{M}_k, \quad (56)$$

$$\lambda_k(M_s) \ge \frac{\lambda_w}{2}. \quad (57)$$

We call $M_0 := M_{s=0}$ the paired clean endpoint.

Proof of Lemma 2. Let $\pi(X) = XX^\top$. At every full-column-rank factor X, the differential

$$\mathrm{D}\pi_X[D] = XD^\top + DX^\top$$

is surjective onto $T_M\mathcal{M}_k$, where $M = \pi(X)$. Its kernel is the vertical space generated by the right-orthogonal invariance:

$$\ker(\mathrm{D}\pi_X) = \{X\Omega : \Omega^\top = -\Omega\}.$$

Thus, after removing these representation-induced vertical directions, factor perturbations provide horizontal lifts of all tangent directions on $\mathcal{M}_k$. The factor–manifold Hessian comparison of Luo, Li, and Zhang (2025, Theorem 1 and Corollary 1) then transfers the curvature lower bound in Proposition 1 to the Riemannian Hessian. Assumption 2 therefore implies that $\mathrm{Hess}_{\mathcal{M}_k} F_w(M_w)$ is positive definite

on $T_{M_w}\mathcal{M}_k$. The implicit-function theorem on the embedded manifold now yields a local critical branch in the parameter s (Absil, Mahony, and Sepulchre 2008, Proposition 5.5.6); (Boumal 2023, Proposition 5.45).

It remains to show that the branch cannot leave the endpoint neighborhood before reaching $s = 0$. Let $P = P_k(M)$ and suppose $\|M - M_w\|_F \le r \le \lambda_w/4$. Weyl's inequality gives $\lambda_k(M) \ge 3\lambda_w/4$, and Davis–Kahan gives $\|P - P_w\|_{\mathrm{op}} \le 2r/\lambda_w$ (Davis and Kahan 1970; Bhatia 1997). Since $G_w(M_w)P_w = 0$, the endpoint gradient is purely normal:

$$G_w(M_w) = P_w^{\perp} G_w(M_w) P_w^{\perp}.$$

Using

$$G_s(M) = G_w(M_w) + H(M - M_w) + (1 - s)N_w,$$

RIP and Cauchy–Schwarz give, for every unit vector v,

$$\begin{aligned}\left|v^\top H(M - M_w)v\right| &= \left|\langle \mathcal{A}(vv^\top), \mathcal{A}(M - M_w)\rangle\right| \\ &\le (1 + \delta)\|M - M_w\|_F .\end{aligned}$$

because $vv^\top$ has rank one and $M - M_w$ has rank at most $2k$. Also $\left|v^\top N_w v\right| \le \sqrt{1+\delta}\,\|w\|_2$. Let

$$b_w = \left\|[P_w^{\perp} G_w(M_w) P_w^{\perp}]_-\right\|_{\mathrm{op}} = \lambda_w \rho_k^-(M_w; w).$$

Projector rotation changes the endpoint normal block by at most

$$\begin{aligned}&\left\|P^{\perp} G_w(M_w) P^{\perp} - P_w^{\perp} G_w(M_w) P_w^{\perp}\right\|_{\mathrm{op}} \\ &\le 2\gamma_w \|P - P_w\|_{\mathrm{op}} \le \frac{4\gamma_w}{\lambda_w}\, r.\end{aligned}$$

Moreover $\lambda_k(M) \ge \lambda_w - r$, so

$$\frac{\lambda_w}{\lambda_k(M)} \le 1 + \frac{2r}{\lambda_w}, \qquad \frac{1}{\lambda_k(M)} \le \frac{2}{\lambda_w}$$

for $r \le \lambda_w/4$. Combining the numerator perturbation with these denominator bounds yields the local normal-ratio estimate

$$\rho_k^-(M; sw) \le \rho_k^-(M_w; w) + \mathfrak{L}_w r + \frac{2\sqrt{1+\delta}}{\lambda_w}(1 - s)\|w\|_2 . \tag{58}$$

The constants in $\mathfrak{L}_w$ are intentionally relaxed; the important point is that they are endpoint quantities.

By Assumption 3, the two extra terms in (58) are at most $\mu_w/2$ whenever $r \le r_w^{\mathrm{stab}}$. Hence, throughout this endpoint neighborhood,

$$(1 - \delta) - \rho_k^-(M; sw) \ge \frac{\mu_w}{2},$$

and Proposition 1 again gives (56).

On any interval where (56) holds, differentiating (55) gives

$$\mathrm{Hess}_{\mathcal{M}_k} F_s(M_s)[\dot{M}_s] = \mathrm{Proj}_{T_{M_s}\mathcal{M}_k}(\mathcal{A}^* w),$$

and RIP gives

$$\left\|\mathrm{Proj}_{T_{M_s}\mathcal{M}_k}(\mathcal{A}^* w)\right\|_F \le \sqrt{1+\delta}\,\|w\|_2 .$$

Therefore

$$\left\|\dot{M}_s\right\|_F \le \frac{2\sqrt{1+\delta}}{\mu_w}\,\|w\|_2 . \tag{59}$$

Integrating backward from $s = 1$ gives, on any interval where the branch exists inside the endpoint neighborhood,

$$\|M_s - M_w\|_F \le \frac{2\sqrt{1+\delta}}{\mu_w}(1 - s)\,\|w\|_2 \le \frac{1}{2} r_w^{\mathrm{stab}}.$$

To formalize the continuation, let $I \subset [0, 1]$ be the maximal interval containing 1 on which the implicit-function branch exists and remains in $\{M : \|M - M_w\|_F < r_w^{\mathrm{stab}}\}$. The local implicit-function theorem makes I open relative to $[0, 1]$. The preceding estimate shows that any point of I actually satisfies the stricter bound $\|M_s - M_w\|_F \le r_w^{\mathrm{stab}}/2$; hence limits of branch points remain inside the same neighborhood, the Hessian stays invertible by (56), and the implicit-function theorem extends the branch past the limit. Thus I is also closed, so $I = [0, 1]$. Finally, (57) follows from $r_w^{\mathrm{stab}} \le \lambda_w/4$. □

Theorem 3 (Pointwise rank-stable perturbation bound)**.** *Suppose Assumptions 1–3 hold, and let M_0 be the paired clean endpoint from Lemma 2. Then*

$$\|M_w - M_0\|_F \le \frac{2\sqrt{1+\delta}}{(1 - \delta) - \rho_k^-(M_w; w)}\,\|w\|_2 . \tag{60}$$

The associated rank-k spectral projector $P_s = P_k(M_s)$, with $P_w = P_1$, rotates slowly:

$$\|P_s - P_w\|_{\mathrm{op}} \le \frac{8\sqrt{1+\delta}\,(1 - s)\,\|w\|_2}{\lambda_k(M_w)\big[(1 - \delta) - \rho_k^-(M_w; w)\big]}. \tag{61}$$

Proof of Theorem 3. The rank-k manifold is equipped with the embedded Frobenius metric. By (55), the Riemannian criticality equation is

$$\mathrm{grad}_{\mathcal{M}_k} F_s(M_s) = 0.$$

Differentiate this identity with respect to s. At a critical point, the covariant derivative of the Riemannian gradient in the motion direction $\dot{M}_s$ is the Riemannian Hessian, while the explicit derivative of the matrix gradient $G_s(M) = H(M - M_\star) - sN_w$ with respect to s is $-N_w = -\mathcal{A}^* w$. Hence

$$\begin{aligned}0 &= \frac{D}{ds}\,\mathrm{grad}_{\mathcal{M}_k} F_s(M_s) \\ &= \mathrm{Hess}_{\mathcal{M}_k} F_s(M_s)[\dot{M}_s] - \mathrm{Proj}_{T_{M_s}\mathcal{M}_k}(\mathcal{A}^* w).\end{aligned}$$

and therefore

$$\mathrm{Hess}_{\mathcal{M}_k} F_s(M_s)[\dot{M}_s] = \mathrm{Proj}_{T_{M_s}\mathcal{M}_k}(\mathcal{A}^* w).$$

This is the intrinsic version of differentiating the pullback criticality equation in local coordinates (Boumal 2023, Proposition 5.45); (Absil, Mahony, and Sepulchre 2008, Proposition 5.5.6).

The right-hand side is controlled by RIP. Since every tangent matrix in $T_{M_s}\mathcal{M}_k$ has rank at most $2k$,

$$\left\|\mathrm{Proj}_{T_{M_s}\mathcal{M}_k}(\mathcal{A}^*w)\right\|_F = \sup_{\substack{\eta\in T_{M_s}\mathcal{M}_k \\ \|\eta\|_F=1}} \langle w, \mathcal{A}(\eta)\rangle \le \sqrt{1+\delta}\,\|w\|_2 .$$

By the Hessian stability condition (56), $\mathrm{Hess}_{\mathcal{M}_k} F_s(M_s)$ is self-adjoint and positive definite on $T_{M_s}\mathcal{M}_k$, with inverse norm at most $2/\mu_w$. Therefore

$$\left\|\dot{M}_s\right\|_F \le \frac{2\sqrt{1+\delta}}{\mu_w}\,\|w\|_2 , \qquad \mu_w = (1-\delta)-\rho_k^-(M_w; w).$$

Integrating from $s=0$ to $s=1$ proves (60).

For the projector statement, the spectral gap between the k-th and $(k+1)$-st eigenvalues of M_s is $\lambda_k(M_s)$, because $\mathrm{rank}(M_s) = k$. The differential form of the Davis–Kahan/Kato spectral projector bound gives the relaxed estimate (Davis and Kahan 1970; Bhatia 1997)

$$\left\|\dot{P}_s\right\|_{\mathrm{op}} \le \frac{2\left\|\dot{M}_s\right\|_{\mathrm{op}}}{\lambda_k(M_s)} \le \frac{8\sqrt{1+\delta}}{\lambda_w\mu_w}\,\|w\|_2 ,$$

where the last step uses (57). Integrating from s to 1 gives (61). □

Supporting Chart Diagnostics Theorem 3 is intrinsic: it controls perturbations in the Frobenius geometry of $\mathcal{M}_k$ without choosing local coordinates. We now express the same argument in local coordinates centered at M_w. The lemma below shows that its conditioning is governed by the extreme active eigenvalues of M_w and deteriorates as $\lambda_k(M_w)$ approaches zero.

Lemma 3 (Raw endpoint chart conditioning)**.** *Let*

$$M_w = U_w\Lambda_w U_w^\top, \quad \Lambda_w = \mathrm{diag}(\lambda_1,\dots,\lambda_k), \quad \lambda_1 \ge \cdots \ge \lambda_k > 0.$$

and complete U_w to an orthogonal basis $[U_w\ U_{w,\perp}]$. A raw graph chart centered at M_w is

$$\Phi(B,Z) = Q(Z)BQ(Z)^\top, \quad Q(Z) = (U_w + U_{w,\perp}Z)(I+Z^\top Z)^{-1/2}.$$

Here $B \in \mathbb{S}^k_{++}$, $Z \in \mathbb{R}^{(n-k)\times k}$, $\dot{B} \in \mathbb{S}^k$, and $\dot{Z} \in \mathbb{R}^{(n-k)\times k}$. At the base point $(B,Z) = (\Lambda_w, 0)$, the differential is

$$D\Phi(\Lambda_w,0)[\dot{B},\dot{Z}] = U_w\dot{B}U_w^\top + U_{w,\perp}\dot{Z}\Lambda_w U_w^\top + U_w\Lambda_w\dot{Z}^\top U_{w,\perp}^\top. \tag{62}$$

With the raw coordinate norm

$$\left\|(\dot{B},\dot{Z})\right\|_{\mathrm{raw}}^2 = \left\|\dot{B}\right\|_F^2 + \left\|\dot{Z}\right\|_F^2,$$

the singular values of $D\Phi(\Lambda_w,0)$, viewed as a map into Frobenius tangent matrices, lie between

$$\sigma_- = \min\{1, \sqrt{2}\,\lambda_k(M_w)\}, \quad \sigma_+ = \max\{1, \sqrt{2}\,\lambda_1(M_w)\}. \tag{63}$$

Thus the raw chart condition number is

$$\kappa_{\mathrm{raw}}(M_w) = \frac{\sigma_+}{\sigma_-} = \frac{\max\{1,\sqrt{2}\,\lambda_1(M_w)\}}{\min\{1,\sqrt{2}\,\lambda_k(M_w)\}}. \tag{64}$$

In raw coordinates, the same argument as Theorem 3 gives

$$\left\|\dot{M}_s\right\|_F \le 2\kappa_{\mathrm{raw}}(M_w)^2\frac{\sqrt{1+\delta}}{\mu_w}\,\|w\|_2$$

at the base point. In tangent-scaled coordinates $K = Z\Lambda_w$, this base-point distortion disappears.

Proof of Lemma 3. Let

$$W(Z) = U_w + U_{w,\perp}Z, \qquad R(Z) = I + Z^\top Z,$$

so $Q(Z) = W(Z)R(Z)^{-1/2}$. For a perturbation $\dot{Z}$,

$$DR(Z)[\dot{Z}] = \dot{Z}^\top Z + Z^\top\dot{Z}.$$

At $Z=0$, $DR(0)[\dot{Z}] = 0$, and hence the derivative of the normalization factor $R(Z)^{-1/2}$ vanishes. Therefore

$$DQ(0)[\dot{Z}] = U_{w,\perp}\dot{Z}.$$

Differentiating $\Phi(B,Z) = Q(Z)BQ(Z)^\top$ gives

$$D\Phi(B,Z)[\dot{B},\dot{Z}] = Q\dot{B}Q^\top + \dot{Q}BQ^\top + QB\dot{Q}^\top.$$

Substituting $Q(0) = U_w$, $B = \Lambda_w$, and $\dot{Q} = U_{w,\perp}\dot{Z}$ proves (62). The three blocks in (62) are Frobenius-orthogonal, so

$$\left\|D\Phi(\Lambda_w,0)[\dot{B},\dot{Z}]\right\|_F^2 = \left\|\dot{B}\right\|_F^2 + 2\left\|\dot{Z}\Lambda_w\right\|_F^2.$$

Since

$$\lambda_k(M_w)\left\|\dot{Z}\right\|_F \le \left\|\dot{Z}\Lambda_w\right\|_F \le \lambda_1(M_w)\left\|\dot{Z}\right\|_F,$$

the singular-value bounds (63) follow. The displayed coordinate velocity bound follows from the coordinate version of the proof: the pulled-back noise gradient is bounded by $\sqrt{1+\delta}\sigma_+\,\|w\|_2$, the pulled-back Hessian is lower bounded by $(\mu_w/2)\sigma_-^2$, and matrix velocity conversion costs another σ_+. □

B.3 Missing Details of Section 5

Proposition 2 (Unused columns force PSD matrix gradient)**.** *Let $X \in \mathbb{R}^{n\times k}$ be a second-order point of f_w. If $\mathrm{rank}(X) < k$, then*

$$G_w(XX^\top) \succeq 0. \tag{65}$$

Proof of Proposition 2. Choose a unit $v \in \mathbb{R}^k$ with $Xv = 0$. For any $u \in \mathbb{R}^n$, set $D = uv^\top$. Then $XD^\top + DX^\top = 0$ and $DD^\top = uu^\top$. The second-order condition and Lemma 1 give

$$0 \le \nabla^2 f_w(X)[D,D] = 2u^\top G_w(XX^\top)u.$$

Since u is arbitrary, $G_w(XX^\top) \succeq 0$. □

Theorem 4 (Rank-deficient points are globally optimal in matrix space)**.** *Let $M = XX^\top$ with $\operatorname{rank}(X) < k$, and suppose X is a second-order point of f_w. Then M is a global minimizer of the convex PSD problem*

$$\min_{Z \succeq 0} F_w(Z). \tag{66}$$

Consequently:

1. *If $w = 0$, then $M = M_\star$.*
2. *For general w, if $\operatorname{rank}(M) \le k$, then*

$$\|M - M_\star\|_F \le \frac{\|w\|_2}{\sqrt{1-\delta}}. \tag{67}$$

Proof of Theorem 4. By first-order criticality, $G_w(M)X = 0$, hence $G_w(M)M = 0$ and $\langle G_w(M), M\rangle = 0$. Proposition 2 gives $G_w(M) \succeq 0$. Therefore, for every $Z \succeq 0$,

$$\langle G_w(M), Z - M\rangle = \langle G_w(M), Z\rangle \ge 0.$$

This is the first-order optimality condition for the convex problem (66), so M is a global PSD minimizer.

If $w = 0$, $M_\star$ is feasible and has objective value zero. Hence $\mathcal{A}(M - M_\star) = 0$. Assumption 1 implies $M = M_\star$, since $\operatorname{rank}(M - M_\star) \le k + r_\star \le 2k$.

For general w, use the convex first-order inequality with $Z = M_\star$. Writing $\Delta = M - M_\star$,

$$0 \le \langle G_w(M), M_\star - M\rangle = \langle \mathcal{A}\Delta - w, -\mathcal{A}\Delta\rangle .$$

Thus $\|\mathcal{A}\Delta\|_2^2 \le \langle w, \mathcal{A}\Delta\rangle \le \|w\|_2 \|\mathcal{A}\Delta\|_2$, so $\|\mathcal{A}\Delta\|_2 \le \|w\|_2$. Since $\operatorname{rank}(\Delta) \le k + r_\star \le 2k$, RIP gives $\sqrt{1-\delta}\,\|\Delta\|_F \le \|\mathcal{A}\Delta\|_2$. □

Remark 3 (Relation to Boumal–Voroninski–Bandeira). This result is the PSD matrix-sensing specialization of the standard rank-deficiency certificate for Burer–Monteiro factorizations (Burer and Monteiro 2003; Boumal, Voroninski, and Bandeira 2016, 2020); the RIP structure further yields exact recovery in the noiseless case and a noise-linear recovery bound in the noisy case.

B.4 Full-Rank Points and the Spectral Tail

We now fully discuss the last part of our main theorem. The exact rank-deficiency certificate is no longer available at a full-rank point, but the negative spectral part of the normal block can remain small when the weakest slack direction is small relative to the active spectral floor. Let

$$M = XX^\top = \sum_{i=1}^{k} \lambda_i u_i u_i^\top, \qquad \lambda_1 \ge \cdots \ge \lambda_k > 0,$$

and choose an effective active rank $q \le k$ with $P_q = \sum_{i=1}^{q} u_i u_i^\top$. A negative normal direction can be paired with the weakest active direction. The second-order condition and the unified $2k$-RIP then limit its negative spectral part by

$$\left\|[P_q^\perp G_w(M) P_q^\perp]_-\right\|_{\mathrm{op}} \le (1+\delta)\lambda_k(M). \tag{68}$$

After normalization by the selected active spectral floor, this becomes

$$\rho_q^-(M; w) \le (1+\delta)\frac{\lambda_k(M)}{\lambda_q(M)}. \tag{69}$$

Proposition 3 (Negative normal curvature is limited by the thinnest slack direction)**.** *Let $X \in \mathbb{R}^{n\times k}$ be a full-rank second-order point of f_w, let*

$$M = XX^\top = \sum_{i=1}^{k} \lambda_i u_i u_i^\top, \qquad \lambda_1 \ge \cdots \ge \lambda_k > 0,$$

and choose an effective active rank $q \le k$ with $P_q = \sum_{i=1}^{q} u_i u_i^\top$. Under Assumption 1,

$$\left\|[P_q^\perp G_w(M) P_q^\perp]_-\right\|_{\mathrm{op}} \le (1+\delta)\lambda_k(M). \tag{70}$$

Consequently,

$$\rho_q^-(M; w) \le (1+\delta)\frac{\lambda_k(M)}{\lambda_q(M)}. \tag{71}$$

Proof of Proposition 3. The first-order condition $G_w(M)X = 0$ implies that $G_w(M)$ vanishes on $\operatorname{range}(M)$. Hence a negative eigenvector of $P_q^\perp G_w(M) P_q^\perp$ lies in the orthogonal complement of $\operatorname{range}(M)$. Let u be such a unit vector and let u_k be an eigenvector of M associated with $\lambda_k(M)$. Choose factor coordinates so that $Xe_k = \sqrt{\lambda_k}\,u_k$, and set

$$D = \sqrt{\lambda_k}\,ue_k^\top.$$

Then

$$XD^\top + DX^\top = \lambda_k(u_k u^\top + u u_k^\top),$$
$$\left\|XD^\top + DX^\top\right\|_F^2 = 2\lambda_k^2,$$
$$DD^\top = \lambda_k u u^\top.$$

If

$$u^\top G_w(M) u = -\gamma < 0,$$

the second-order condition and Lemma 1 give

$$0 \le \nabla^2 f_w(X)[D, D] = \left\|\mathcal{A}(XD^\top + DX^\top)\right\|_2^2 - 2\gamma\lambda_k.$$

The matrix $XD^\top + DX^\top$ has rank at most two, so RIP gives

$$\left\|\mathcal{A}(XD^\top + DX^\top)\right\|_2^2 \le (1+\delta)2\lambda_k^2.$$

Thus $\gamma \le (1+\delta)\lambda_k$, proving (70). Dividing by $\lambda_q(M)$ gives (71). □

For $q < k$, the spectral ratio $\lambda_k(M)/\lambda_q(M)$ is therefore an effective-rank diagnostic for normalized negative normal curvature. It controls the weakest slack direction, but not the total discarded tail energy.

Corollary 1 (Effective-rank normal-curvature surrogate)**.** *If M_w is produced by a full-column-rank second-order point of width k, then for any effective active rank $q < k$,*

$$\rho_q^-(M_w; w) \le (1+\delta)\frac{\lambda_k(M_w)}{\lambda_q(M_w)}. \tag{72}$$

This is not the main strict rank-stable theorem, which uses $q = k$ and the directly computed $\rho_k^-(M_w; w)$. It is an effective-rank diagnostic: when the spectral tail ratio $\lambda_k(M_w)/\lambda_q(M_w)$ is small, the normal curvature seen by the active rank-q block is small, but the discarded tail must still be accounted for separately if one truncates to rank q. In particular, λ_k/λ_q is the thinnest-tail ratio; it does not by itself control the whole tail energy.

Lemma 4 (Deterministic spectral-tail bounds)**.** *Let $M \succeq 0$ have rank at most k, eigenvalues $\lambda_1 \geq \cdots \geq \lambda_k \geq 0$, and fix*

$$1 \leq q < k, \qquad \lambda_q(M) > 0.$$

Let

$$M_q = \sum_{i=1}^{q} \lambda_i u_i u_i^\top$$

be its best rank-q approximation. Then

$$\|M - M_q\|_{\mathrm{op}} = \lambda_{q+1}(M), \tag{73}$$

$$\|M - M_q\|_F^2 = \sum_{i>q} \lambda_i(M)^2. \tag{74}$$

If $d = k - q$, then

$$\frac{\|M - M_q\|_F}{\lambda_q(M)} \leq \sqrt{d}\, \frac{\lambda_{q+1}(M)}{\lambda_q(M)}. \tag{75}$$

Also,

$$\frac{\lambda_k(M)}{\lambda_q(M)} \leq \frac{\|M - M_q\|_F}{\sqrt{d}\, \lambda_q(M)}. \tag{76}$$

If, in addition, $\bar{M} \succeq 0$ has rank q and $\|M - \bar{M}\|_{\mathrm{op}} < \lambda_q(\bar{M})$, then Weyl's inequality gives

$$\frac{\lambda_k(M)}{\lambda_q(M)} \leq \frac{\|M - \bar{M}\|_{\mathrm{op}}}{\lambda_q(\bar{M}) - \|M - \bar{M}\|_{\mathrm{op}}}. \tag{77}$$

Proof of Lemma 4. The identities (73) are the PSD form of the Eckart–Young–Mirsky theorem (Eckart and Young 1936). Since the d tail eigenvalues $\lambda_{q+1}, \ldots, \lambda_k$ are all at most λ_{q+1},

$$\|M - M_q\|_F^2 = \sum_{i>q} \lambda_i^2 \leq d\, \lambda_{q+1}^2,$$

which proves (75). Since the same tail eigenvalues are all at least λ_k,

$$d\, \lambda_k^2 \leq \sum_{i>q} \lambda_i^2 = \|M - M_q\|_F^2,$$

which proves (76). If $M = \bar{M} + E$ and $\mathrm{rank}(\bar{M}) = q$, Weyl's inequality yields

$$\lambda_q(M) \geq \lambda_q(\bar{M}) - \|E\|_{\mathrm{op}}, \qquad \lambda_k(M) \leq \|E\|_{\mathrm{op}},$$

because $\lambda_k(\bar{M}) = 0$ for $k > q$. This proves (77). □

For direct recovery relative to the rank-$r_\star$ target $M_\star$, the relevant quantity is the numerator of $\rho_k^-(M_w; w)$, rather than its normalization by the active spectral floor. At a full-rank first-order critical point, $G_w(M_w)X_w = 0$ implies

$$G_w(M_w) = P_w^\perp G_w(M_w) P_w^\perp.$$

We therefore define the absolute negative normal leakage by

$$b_w^- := \left\|[P_w^\perp G_w(M_w) P_w^\perp]_-\right\|_{\mathrm{op}} = \lambda_k(M_w)\rho_k^-(M_w; w). \tag{78}$$

Lemma 5 (Absolute normal leakage)**.** *Let $X_w \in \mathbb{R}^{n\times k}$ be full column rank and first-order critical for f_w, let $M_w = X_w X_w^\top$, and let $P_w = P_k(M_w)$. Then*

$$G_w(M_w) = P_w^\perp G_w(M_w) P_w^\perp.$$

Define the negative normal leakage amplitude

$$b_w^- = \left\|[P_w^\perp G_w(M_w) P_w^\perp]_-\right\|_{\mathrm{op}}. \tag{79}$$

Equivalently,

$$b_w^- = \lambda_k(M_w)\rho_k^-(M_w; w). \tag{80}$$

Proof of Lemma 5. First-order criticality gives $G_w(M_w)X_w = 0$. Since X_w spans $\mathrm{range}(M_w)$, this is equivalent to $G_w(M_w)P_w = 0$. By symmetry also $P_w G_w(M_w) = 0$, so $G_w(M_w)$ is supported on the normal block. Equation (80) is the definition of ρ_k^- in (6) with $q = k$. □

Lemma 5 isolates the adverse part of the normal block at a full-rank critical point: b_w^- measures its absolute spectral amplitude, while ρ_k^- is the corresponding normalization by the weakest active eigenvalue. To convert this local quantity into a recovery estimate, we must also control how much of the target matrix lies outside $\mathrm{range}(M_w)$.The next lemma bounds this target-subspace leakage directly by the reconstruction error.

Lemma 6 (Target subspace leakage)**.** *Let $M_\star \succeq 0$ have rank $1 \leq r_\star \leq k$, let*

$$\lambda_\star = \lambda_{r_\star}(M_\star) > 0,$$

and let $M_w \succeq 0$ have range projector P_w. Then

$$\mathrm{tr}(P_w^\perp M_\star P_w^\perp) \leq \frac{1}{\lambda_\star}\|M_w - M_\star\|_F^2. \tag{81}$$

Proof of Lemma 6. Write

$$M_\star = U_\star \Lambda_\star U_\star^\top$$

and set $E = M_w - M_\star$. Since $P_w^\perp M_w = 0$,

$$\begin{aligned} P_w^\perp U_\star \Lambda_\star^{1/2} &= P_w^\perp M_\star U_\star \Lambda_\star^{-1/2} \\ &= -P_w^\perp E U_\star \Lambda_\star^{-1/2}. \end{aligned}$$

Therefore,

$$\begin{aligned} \mathrm{tr}(P_w^\perp M_\star P_w^\perp) &= \left\|P_w^\perp U_\star \Lambda_\star^{1/2}\right\|_F^2 \\ &= \left\|P_w^\perp E U_\star \Lambda_\star^{-1/2}\right\|_F^2 \\ &\leq \left\|\Lambda_\star^{-1/2}\right\|_{\mathrm{op}}^2 \|E\|_F^2 \\ &= \frac{1}{\lambda_\star}\|M_w - M_\star\|_F^2, \end{aligned}$$

which proves (81). □

Lemmas 5 and 6 now connect the normal geometry to recovery. Writing $E = M_w - M_\star$ and $\lambda_\star = \lambda_{r_\star}(M_\star)$, first-order criticality gives

$$\langle G_w(M_w), E\rangle = -\langle G_w(M_w), M_\star\rangle \le \frac{b_w^-}{\lambda_\star}\|E\|_F^2 .$$

Meanwhile, RIP lower-bounds the data-fit contribution by $(1-\delta)\|E\|_F^2$, up to the measurement-noise term. Hence the normal leakage can be absorbed whenever $(1-\delta)-b_w^-/\lambda_\star > 0$. The following lemma formalizes the resulting recovery bound.

Lemma 7 (Normal leakage recovery bound)**.** *Assume the unified $2k$-RIP in Assumption 1, and let $M_\star \succeq 0$ have rank $1 \le r_\star \le k$. Let X_w be a full-column-rank first-order critical point of f_w, and let $M_w = X_w X_w^\top$. Define b_w^- by* (78)*, set*

$$\mu_{\mathrm{leak}} = 1 - \delta - \frac{b_w^-}{\lambda_{r_\star}(M_\star)}, \tag{82}$$

and assume $\mu_{\mathrm{leak}} > 0$. Then

$$\|M_w - M_\star\|_F \le \frac{\sqrt{1+\delta}}{\mu_{\mathrm{leak}}}\|w\|_2 . \tag{83}$$

Proof of Lemma 7. Let $E = M_w - M_\star$, $G = G_w(M_w)$, and $\lambda_\star = \lambda_{r_\star}(M_\star)$. By first-order criticality, $GX_w = 0$, and hence $GM_w = 0$. Therefore,

$$\langle G, E\rangle = -\langle G, M_\star\rangle .$$

Since G is supported on $P_w^\perp$, Lemmas 5 and 6 give

$$\begin{aligned} -\langle G, M_\star\rangle &\le b_w^- \operatorname{tr}(P_w^\perp M_\star P_w^\perp) \\ &\le \frac{b_w^-}{\lambda_\star}\|E\|_F^2 . \end{aligned}$$

On the other hand,

$$\langle G, E\rangle = \|\mathcal{A}(E)\|_2^2 - \langle w, \mathcal{A}(E)\rangle .$$

Since $\operatorname{rank}(E) \le k + r_\star \le 2k$, RIP and Cauchy–Schwarz imply

$$(1-\delta)\|E\|_F^2 \le \sqrt{1+\delta}\,\|w\|_2\|E\|_F + \frac{b_w^-}{\lambda_\star}\|E\|_F^2 .$$

Rearranging and using (82) gives

$$\mu_{\mathrm{leak}}\|E\|_F^2 \le \sqrt{1+\delta}\,\|w\|_2\|E\|_F ,$$

which proves (83).

□

Another ingredient controls the portion of the lower-rank target that lies outside $\operatorname{range}(M_w)$. This target-subspace leakage is controlled by the current recovery error. Combining that estimate with b_w^- and the unified $2k$-RIP gives the direct full-rank recovery certificate below.

Proof of Theorem 2. Part (1) is Theorem 4. Part (2) is Theorem 3, whose strict full-rank case uses active rank k. Part (3) is Lemma 5 followed by Lemma 7.

□

C Trajectory Case Study: Full-Rank Shadowing from Small Initialization

This section studies robustness along the whole GD path rather than at one critical endpoint. Therefore the clean reference is not a single matrix M_0. It is the time-indexed noiseless trajectory

$$M_t^0 = X_t^0 X_t^{0\top},$$

run from the same initialization as the noisy trajectory

$$M_t^w = X_t^w X_t^{w\top}.$$

The trajectory robustness question is

$$\Delta_t = \left\|M_t^w - M_t^0\right\|_F .$$

Throughout this section we focus on the full-rank-to-full-rank regime with $k \le n$: for the time window under discussion, both X_t^0 and X_t^w remain full column rank, so M_t^0 and M_t^w lie on the rank-k PSD manifold. The role of the normal-curvature ratio is now dynamic. A small $\rho_{k,t}^-$, defined in (85), together with the local Hessian check below, indicates that the rank-k tangent dynamics around M_t^w are well conditioned, so measurement noise cannot separate the noisy and noiseless trajectories too much during the next steps.

C.1 Clean and Noisy GD Paths

Let $N_0 = 0$ and $N_w = \mathcal{A}^* w$. Starting from the same small random initialization $X_0^0 = X_0^w = \alpha G$, define

$$\begin{aligned} X_{t+1}^\zeta &= \left(I + \eta[H(M_\star - M_t^\zeta) + N_\zeta]\right) X_t^\zeta, \\ M_t^\zeta &= X_t^\zeta (X_t^\zeta)^\top, \qquad \zeta \in \{0, w\}. \end{aligned} \tag{84}$$

The clean recursion is the one studied by Stöger and Soltanolkotabi (2021), up to normalization constants. In their notation the target is $M_\star = X_\star X_\star^\top$, the factor condition number is

$$\kappa_X = \frac{\|X_\star\|_{\mathrm{op}}}{\sigma_{\min}(X_\star)} = \sqrt{\frac{\lambda_1(M_\star)}{\lambda_{r_\star}(M_\star)}},$$

the search width is $k > r_\star$, $X_0 = \alpha G$ with entries of G distributed as $N(0, 1/\sqrt{k})$, and the step size satisfies $\eta \lesssim \kappa_X^{-4}\|X_\star\|_{\mathrm{op}}^{-2}$. Equivalently, if $\kappa_M = \lambda_1(M_\star)/\lambda_{r_\star}(M_\star)$, then the step size is of order $\kappa_M^{-2}\lambda_1(M_\star)^{-1}$. Their Theorem 3.3 assumes $k > r_\star$, splits into the regimes $k \ge 2r_\star$ and $r_\star < k < 2r_\star$, and requires rank-$(2r_\star+1)$ RIP with constant of order at most $1/(\kappa_X^4\sqrt{r_\star})$, together with explicit small upper bounds on α. It proves a finite-time early-stopped proximity bound for the clean iterates. In the $k \ge 2r_\star$ case, the normalized error after the prescribed number of iterations is bounded by a polynomial factor times $(\alpha/\|X_\star\|_{\mathrm{op}})^{21/16}$. The exactly parameterized case $k = r_\star$ is treated separately in their Theorem 3.4.

The clean result is used here as a reference path. The new object is the shadowing error Δ_t between (84) with and without w. No critical-point comparison is needed.

C.2 Rank-k Rho and a Shadowing Bound

For a full-rank noisy iterate M_t^w, let $P_t^w = P_k(M_t^w)$. In this section

$$G_w(M) = H(M - M_\star) - N_w$$

is the noisy matrix gradient. Define

$$\rho_{k,t}^- = \frac{\left\| \left[P_t^{w,\perp} G_w(M_t^w) P_t^{w,\perp} \right]_- \right\|_{\mathrm{op}}}{\lambda_k(M_t^w)}, \qquad \mu_t^\rho = (1-\delta) - \rho_{k,t}^-. \tag{85}$$

This is the same full-rank denominator as in the endpoint perturbation theorem, but now it is evaluated along the noisy trajectory. The following deterministic shadowing lemma is the trajectory analogue of the endpoint bound. Because trajectory points are not critical points, $\rho_{k,t}^-$ should be read as a pointwise lower-margin estimate for the local linearized rank-k map, not as an automatic theorem about every GD step.

To make the bridge explicit, let

$$\Psi_t^0(M) = \mathcal{R}_M\big(-\eta \operatorname{grad}_{\mathcal{M}_k} F_0(M)\big)$$

be the noiseless rank-k manifold step in a local chart, with any smooth first-order retraction $\mathcal{R}$. A sufficient check for the contraction used below is for every $\eta_M \in T_{M_t^w}\mathcal{M}_k$,

$$\langle \eta_M, \operatorname{Hess}_{\mathcal{M}_k} F_0(M_t^w)[\eta_M] \rangle \geq \mu_t^\rho \|\eta_M\|_F^2, \qquad \eta L_t \leq 1. \tag{86}$$

where L_t is a local smoothness bound for the charted gradient map. For the local quadratic model, and for the charted map after absorbing higher-order retraction terms into the local check, this gives $\|D\Psi_t^0\| \leq 1 - \eta\mu_t^\rho$. Thus the Hessian Rayleigh quotient is the actual trajectory certificate. The ratio (85) is the rank-k normal-block diagnostic that explains when the check should become favorable; at endpoint-critical or tangent-critical points it reduces to the same normal-curvature margin used in Proposition 1.

Lemma 8 (Full-rank trajectory shadowing controlled by ρ)**.** *Fix a time horizon T. Suppose that the paired clean and noisy trajectories remain in a common full-rank k-dimensional chart for $0 \leq t \leq T$. Assume that the local Hessian/smoothness check* (86) *holds, so the noiseless chart map satisfies*

$$\|D\Psi_t^0\| \leq 1 - \eta\mu_t^\rho.$$

Then the noisy and noiseless matrix iterates obey

$$\Delta_{t+1} \leq (1 - \eta\mu_t^\rho)\Delta_t + \eta\sqrt{1+\delta}\,\|w\|_2, \tag{87}$$

and therefore

$$\Delta_t \leq \sqrt{1+\delta}\,\|w\|_2 \sum_{j=0}^{t-1} \eta \prod_{\ell=j+1}^{t-1} (1 - \eta\mu_\ell^\rho). \tag{88}$$

If $\mu_t^\rho \geq \underline{\mu} > 0$ on the whole window and $\eta\underline{\mu} \leq 1$, then

$$\Delta_t \leq \frac{\sqrt{1+\delta}}{\underline{\mu}} \left(1 - (1-\eta\underline{\mu})^t\right) \|w\|_2 \leq \frac{\sqrt{1+\delta}}{\underline{\mu}} \|w\|_2. \tag{89}$$

Proof. The only first-order difference between the two local rank-k maps at the same matrix is the tangent projection of $N_w = \mathcal{A}^* w$. RIP gives

$$\left\| \operatorname{Proj}_{T_{M_t^w}\mathcal{M}_k} N_w \right\|_F \leq \sqrt{1+\delta}\,\|w\|_2,$$

because tangent matrices have rank at most $2k$. By assumption, the computed rank-k margin $\mu_t^\rho = (1-\delta) - \rho_{k,t}^-$ certifies the one-step tangent contraction. Combining this contraction of the clean map with the additive noise perturbation gives (87). Iterating the scalar recursion gives (88), and the geometric sum gives (89). □

Lemma 8 is the clean trajectory philosophy. The object M_t^0 is the clean iterate at the same time t, not the clean critical endpoint used earlier. The time-varying stability margin is $\mu_t^\rho = (1-\delta) - \rho_{k,t}^-$, where the ratio $\rho_{k,t}^-$ contains the active floor $\lambda_k(M_t^w)$ in its denominator. Early times with a poor rank-k margin contribute large terms to the sum in (88); times with a strong margin damp all previous noise.

C.3 Stöger–Soltanolkotabi Variables and a ρ-Proxy

The remaining question is how the dynamics make the ρ-margin better or worse. The clean analysis of Stöger and Soltanolkotabi (2021) introduces a signal/slack split that is useful for interpreting this. Their "noise" term is an internal over-parameterized slack factor, not observation noise w. Let

$$M_\star = V_\star \Lambda_\star V_\star^\top, \qquad P_\star = V_\star V_\star^\top, \qquad \lambda_\star = \lambda_{r_\star}(M_\star), \qquad \kappa_X = \frac{\|X_\star\|_{\mathrm{op}}}{\sigma_{\min}(X_\star)}.$$

For a clean iterate X_t^0, take the compact SVD

$$V_\star^\top X_t^0 = \widetilde{V}_t \Sigma_t W_t^\top, \qquad W_t \in \mathbb{R}^{k \times r_\star}, \qquad W_t^\top W_t = I, \tag{90}$$

and complete W_t to $[W_t\ W_{t,\perp}]$. Following their Section 5.1, write

$$S_t = X_t^0 W_t, \qquad N_t = X_t^0 W_{t,\perp}, \qquad X_t^0 = S_t W_t^\top + N_t W_{t,\perp}^\top. \tag{91}$$

By Stöger and Soltanolkotabi (2021, Lemma 5.1), the slack factor has column span orthogonal to $V_\star$, and, whenever $V_\star^\top X_t^0$ has full row rank, S_t has rank $r_\star$ while N_t has rank at most $k - r_\star$. Let V_{S_t} denote an orthonormal basis for $\operatorname{col}(S_t)$. Define

$$s_t = \sigma_{\min}(S_t), \quad \nu_t = \|N_t\|_{\mathrm{op}}, \quad \phi_t = \frac{s_t^2}{\lambda_\star}, \qquad \xi_t = \left\|V_\star^\perp V_{S_t}\right\|_{\mathrm{op}}, \quad \chi_t = \frac{\nu_t^2}{s_t^2}, \qquad \varepsilon_t = \frac{\left\|V_\star^\top (M_t^0 - M_\star)\right\|_F}{\lambda_\star}. \tag{92}$$

The angle ξ_t is the clean signal-factor angle; in Stöger and Soltanolkotabi (2021), Lemma 9.1 grows the projected signal singular value, Lemma 9.2 controls the internal slack factor, and Lemma 9.3 gives the direct angle recurrence. The

auxiliary ε_t is used below only to quote their Phase III refinement; it is not a separate term in the four-metric dashboard. With sensing distortion one also uses the top $r_\star$-dimensional spectral space P_H of $H(M_\star)$ and the upper proxy

$$\theta_t = \left\|V_{S_t}^{\perp} P_H\right\|_{\mathrm{op}} + \left\|P_H^{\perp} P_\star\right\|_{\mathrm{op}}, \qquad \xi_t \le \theta_t,$$

where the second term is controlled by Davis–Kahan when $\|H(M_\star) - M_\star\|_{\mathrm{op}} \ll \lambda_\star$.

The strict full-rank robustness ratio uses the smallest rank-k eigenvalue of the noisy iterate. Normalize this floor by the signal scale:

$$\psi_t^w = \frac{\lambda_k(M_t^w)}{\lambda_\star}.$$

The normal block is easiest to read through one subspace metric. For the mathematical bound we use the target alignment

$$m_t^w = \left\|P_t^{w,\perp} P_\star\right\|_{\mathrm{op}}. \tag{93}$$

This measures how much of the true signal space is still visible from the noisy normal space. It is small when the noisy rank-k range contains the target signal space. If $P_\star$ is unavailable and the spectral surrogate P_H is used instead, then

$$m_t^w \le \left\|P_t^{w,\perp} P_H\right\|_{\mathrm{op}} + \left\|P_H^{\perp} P_\star\right\|_{\mathrm{op}},$$

and the second term is controlled by Davis–Kahan under the spectral closeness condition stated above. The normalized external noise level is

$$\eta_w = \frac{\sqrt{1+\delta}\,\|w\|_2}{\lambda_\star}. \tag{94}$$

On a bootstrap window where Lemma 8 keeps $\Delta_t = O(\|w\|_2)$, the clean/noisy path discrepancy is folded into the same relaxed noise scale. The signal-floor/slack part is represented by the measured floor ψ_t^w and the slack tail

$$\tau_t = \sqrt{k - r_\star}\,\chi_t \phi_t = \sqrt{k - r_\star}\,\frac{\nu_t^2}{\lambda_\star}. \tag{95}$$

Thus τ_t is not another signal-growth variable; it is the Frobenius-sized energy of the $k - r_\star$ internal slack directions, normalized by the target spectral floor.

The following lemma is the promised mathematical bound. It has two fixed constants, not two new trajectory metrics. The constant c_{sh} says that the noisy and clean paths are already in the small-noise shadowing window, and c_{sig} is the clean signal-block radius. For example, $c_{\mathrm{sig}} = 5\sqrt{r_\star}\kappa_X^2$ is valid whenever $\|S_t\|_{\mathrm{op}} \le 2\,\|X_\star\|_{\mathrm{op}}$, because then

$$\left\|S_t S_t^\top - M_\star\right\|_F \le 5\sqrt{r_\star}\,\kappa_X^2 \lambda_\star.$$

Lemma 9 (Four-metric upper bound for the trajectory ρ)**.** *Assume $k > r_\star$. Fix a full-rank time t in the trajectory window. Suppose*

$$\Delta_t = \left\|M_t^w - M_t^0\right\|_F \le c_{\mathrm{sh}}\,\|w\|_2,$$
$$\left\|S_t S_t^\top - M_\star\right\|_F \le c_{\mathrm{sig}} \lambda_\star.$$

Then

$$\rho_{k,t}^{-} \le \frac{\kappa_X^2 (m_t^w)^2 + (1+\delta)\tau_t + \delta c_{\mathrm{sig}}}{\psi_t^w} + \frac{\left(1 + c_{\mathrm{sh}}\sqrt{1+\delta}\right)\eta_w}{\psi_t^w}. \tag{96}$$

Equivalently, the margin in Lemma 8 is positive whenever the numerator in (96) *is smaller than $(1-\delta)\psi_t^w$. If the surrogate P_H is used, replace $(m_t^w)^2$ by*

$$\left(\left\|P_t^{w,\perp} P_H\right\|_{\mathrm{op}} + \left\|P_H^{\perp} P_\star\right\|_{\mathrm{op}}\right)^2.$$

Proof. Let $P = P_t^w$ and

$$B = P^{\perp} G_w(M_t^w) P^{\perp}.$$

Since $\|[B]_-\|_{\mathrm{op}} = \sup_{\|v\|=1,\, v\in\mathrm{range}(P^\perp)} [-v^\top B v]_+$, it is enough to bound the negative Rayleigh quotient on a unit normal vector v. Using $M_t^0 = S_t S_t^\top + N_t N_t^\top$,

$$\begin{aligned} -v^\top G_w(M_t^w) v &= -v^\top H(M_t^w - M_t^0) v \\ &\quad - v^\top H(S_t S_t^\top - M_\star) v \\ &\quad - v^\top H(N_t N_t^\top) v \\ &\quad + v^\top N_w v. \end{aligned}$$

The shadowing and observation-noise terms satisfy

$$\begin{aligned} \left|v^\top H(M_t^w - M_t^0) v\right| &\le (1+\delta)\Delta_t \\ &\le c_{\mathrm{sh}}\sqrt{1+\delta}\,\lambda_\star \eta_w, \\ \left|v^\top N_w v\right| &\le \lambda_\star \eta_w. \end{aligned}$$

For the active signal block, restricted polarization gives

$$\begin{aligned} v^\top H(S_t S_t^\top - M_\star) v &\ge v^\top (S_t S_t^\top - M_\star) v \\ &\quad - \delta \left\|S_t S_t^\top - M_\star\right\|_F. \end{aligned}$$

This polarization step uses only the unified $2k$-RIP, because $\mathrm{rank}(S_t S_t^\top - M_\star) \le 2r_\star$ and the trajectory case has $k > r_\star$, so $2r_\star + 1 \le 2k$. Therefore

$$\begin{aligned} \left[-v^\top H(S_t S_t^\top - M_\star) v\right]_+ &\le \left[v^\top (M_\star - S_t S_t^\top) v\right]_+ \\ &\quad + \delta c_{\mathrm{sig}} \lambda_\star \\ &\le v^\top M_\star v + \delta c_{\mathrm{sig}} \lambda_\star \\ &\le \lambda_1(M_\star)\,\|P_\star v\|^2 \\ &\quad + \delta c_{\mathrm{sig}} \lambda_\star \\ &\le \lambda_\star \kappa_X^2 (m_t^w)^2 \\ &\quad + \delta c_{\mathrm{sig}} \lambda_\star. \end{aligned}$$

Finally,

$$\begin{aligned} \left|v^\top H(N_t N_t^\top) v\right| &\le (1+\delta)\left\|N_t N_t^\top\right\|_F \\ &\le (1+\delta)\sqrt{k - r_\star}\,\nu_t^2 \\ &= (1+\delta)\tau_t \lambda_\star. \end{aligned}$$

Combining these four estimates and dividing by $\lambda_k(M_t^w) = \psi_t^w \lambda_\star$ gives (96). □

Combining Lemmas 8 and 9 gives the explicit trajectory perturbation bound. Define the four-metric upper envelope

$$\begin{aligned}\bar{\rho}_{k,t} &= \frac{\kappa_X^2 (m_t^w)^2 + (1+\delta)\tau_t + \delta c_{\mathrm{sig}}}{\psi_t^w} \\ &\quad + \frac{\left(1 + c_{\mathrm{sh}}\sqrt{1+\delta}\right)\eta_w}{\psi_t^w}, \\ \bar{\mu}_t &= (1-\delta) - \bar{\rho}_{k,t}.\end{aligned} \tag{97}$$

On any time window where the hypotheses of Lemma 9 hold at each time and the local Hessian check in (86) holds, if $\bar{\mu}_t > 0$, then $\mu_t^\rho \geq \bar{\mu}_t$ and

$$\Delta_t \leq \sqrt{1+\delta}\,\|w\|_2 \sum_{j=0}^{t-1} \eta \prod_{\ell=j+1}^{t-1} (1 - \eta\bar{\mu}_\ell). \tag{98}$$

In particular, if $\bar{\mu}_t \geq \underline{\mu} > 0$ on the window, then

$$\left\|M_t^w - M_t^0\right\|_F \leq \frac{\sqrt{1+\delta}}{\underline{\mu}}\,\|w\|_2\,. \tag{99}$$

This is the actual robustness statement in the trajectory section: alignment, slack, noise, and the active floor determine $\bar{\rho}_{k,t}$, and the gap $\bar{\mu}_t = (1-\delta) - \bar{\rho}_{k,t}$ determines the perturbation size. The bootstrap form is non-circular. Fix a desired uniform margin $\underline{\mu} > 0$ and set

$$c_{\mathrm{sh}} = \frac{\sqrt{1+\delta}}{\underline{\mu}}.$$

If, with this predetermined value of c_{sh}, the envelope (97) satisfies

$$\bar{\rho}_{k,t} \leq 1 - \delta - \underline{\mu} \qquad \text{for every } 0 \leq t \leq T,$$

then (99) gives

$$\Delta_t \leq c_{\mathrm{sh}}\,\|w\|_2 \qquad \text{for every } 0 \leq t \leq T.$$

Thus the shadowing condition used inside Lemma 9 is proved by induction on the same window, rather than assumed after the fact.

The exact ratio $\rho_{k,t}^-$ remains directly computable from (85). Lemma 9 gives the relaxed four-metric upper bound

$$\widehat{\rho}_{k,t}^{\,\mathrm{traj}} := \frac{\kappa_X^2 (m_t^w)^2 + \tau_t + \eta_w}{\psi_t^w}. \tag{100}$$

This is the readable version of (96) after folding the fixed RIP-bias term δc_{sig} and the fixed shadowing constant c_{sh} into constants. The measured certificate remains $\rho_{k,t}^-$, but now the section has an explicit upper bound:

$$\begin{aligned}\rho_{k,t}^- &\leq \frac{\kappa_X^2 (m_t^w)^2 + (1+\delta)\tau_t + \delta c_{\mathrm{sig}}}{\psi_t^w} \\ &\quad + \frac{\left(1 + c_{\mathrm{sh}}\sqrt{1+\delta}\right)\eta_w}{\psi_t^w} \\ &\approx \widehat{\rho}_{k,t}^{\,\mathrm{traj}}.\end{aligned}$$

The denominator ψ_t^w is the active full-rank floor. The middle term is the slack tail created by the extra $k - r_\star$ columns. Here "active" means active in the rank-k chart, not necessarily signal-aligned with $M_\star$; when $k > r_\star$, the floor may be carried by the thinnest retained slack direction. The remaining noise term η_w absorbs the small perturbation effects we are deliberately not separating in this relaxed dashboard. Even excellent signal variables do not certify strict full-rank robustness if $\lambda_k(M_t^w)$ is still microscopic.

To see (100), set $P = P_t^w$. Then

$$\begin{aligned}P^\perp G_w(M_t^w) P^\perp &= P^\perp H(M_t^w - M_t^0) P^\perp \\ &\quad + P^\perp H(M_t^0 - M_\star) P^\perp \\ &\quad - P^\perp N_w P^\perp.\end{aligned} \tag{101}$$

For a unit vector $v \in \mathrm{range}(P_t^{w,\perp})$, Cauchy–Schwarz and RIP give $\left|v^\top N_w v\right| \lesssim \lambda_\star \eta_w$. The clean residual is controlled by the single misalignment metric m_t^w: larger m_t^w means the noisy normal space still sees signal directions and can therefore contain more negative normal curvature. The slack block has Frobenius norm at most

$$\sqrt{k - r_\star}\,\nu_t^2 = \tau_t \lambda_\star.$$

Dividing the corresponding rigorous estimate by $\lambda_k(M_t^w) = \psi_t^w \lambda_\star$ is precisely Lemma 9. The simplified display (100) only suppresses the fixed constants from (96); the exact robustness certificate is still the measured $\rho_{k,t}^-$ together with the Hessian check (86).

C.4 Clean Phase Dynamics and the Evolution of the ρ-Proxy

The robustness quantity in this section is always the strict full-rank margin $\rho_{k,t}^-$ from (85). The proxy (100) only explains its moving parts. Once the measured ratio gives a positive margin and the local Hessian check passes, Lemma 8 converts that margin into the shadowing bound (88). The phase story below is therefore a story about the dashboard in (100): how the single misalignment metric m_t^w, the slack tail τ_t, the noise level η_w, and the measured rank-k floor ψ_t^w evolve.

Phase I: spectral alignment. The first phase is controlled by the power-method comparison in Stöger and Soltanolkotabi (2021, Lemmas 8.1–8.2). They compare GD to the linearized iterate

$$\widetilde{X}_t = (I + \eta H(M_\star))^t X_0$$

and show that X_t^0 remains close to $\widetilde{X}_t$ while the iterate is small. Their Lemma 8.3 gives SVD perturbation bounds for the approximate power iterate $Z_t X_0 + E_t$; Lemma 8.4 translates top-left-singular-space alignment into properties of $X_t^0 W_t$ and $X_t^0 W_{t,\perp}$; Lemma 8.5 combines these two steps, under spectral closeness $H(M_\star) = M_\star + \widetilde{E}$, small $\left\|\widetilde{E}\right\|_{\mathrm{op}} / \lambda_\star$, and the corresponding small power-method error parameter, into lower bounds on $\sigma_{\min}(X_t^0 W_t)$, upper bounds on $\left\|X_t^0 W_{t,\perp}\right\|$, and angle bounds for $\left\|V_\star^\perp V_{S_t}\right\|_{\mathrm{op}}$. Their Theorem 8.8 supplies the Gaussian smallest-singular-value lower bound for an $r_\star \times k$ Gaussian matrix with entries normalized by $1/\sqrt{k}$:

$$\sigma_{\min}(G) \geq \varepsilon\,\frac{\sqrt{k} - \sqrt{r_\star - 1}}{\sqrt{k}}$$

with probability at least $1-(C\varepsilon)^{k-r_\star+1}-e^{-ck}$. This is the probabilistic source of improved initial overlap under overparameterization.

In the dashboard (100), Phase I primarily reduces the clean component of m_t^w. More precisely, Lemma 8.5 controls the clean subspace angle through ξ_t after the approximate power-method phase, while the small-noise/bootstrap condition keeps the noisy component at the η_w scale. In this relaxed view, Phase I is best viewed as subspace preparation: the trajectory has learned where the signal is, but the full-rank chart may still be fragile.

Phase II: saddle avoidance and signal-floor growth. After alignment, the clean trajectory obeys the signal-floor growth inequality (Stöger and Soltanolkotabi 2021, Lemma 9.1)

$$\sigma_{\min}(V_\star^\top X_{t+1}^0) \geq \left(1+\frac{1}{4}\eta\lambda_\star\right)\sigma_{\min}(V_\star^\top X_t^0) - \eta\,\sigma_{\min}(V_\star^\top X_t^0)^3.$$

under their step-size, angle, norm, and RIP-error conditions. While this projected signal floor is below signal scale, and while $\xi_t = \left\|V_\star^\perp V_{S_t}\right\|_{\mathrm{op}}$ remains small, it is comparable to $\sigma_{\min}(S_t)$ up to the principal-angle factor $\sqrt{1-\xi_t^2}$. Hence a useful reading is

$$s_{t+1} \gtrsim s_t(1+c\eta\lambda_\star), \qquad \phi_{t+1} \gtrsim \phi_t(1+c\eta\lambda_\star)^2. \quad (102)$$

The same analysis controls the slack factor $\nu_t = \|N_t\|_{\mathrm{op}}$ and the signal angle (Stöger and Soltanolkotabi 2021, Lemmas 9.2–9.3). In the proof of Theorem 9.6, starting from a post-spectral time t_0 with initial signal scale γ, Phase II is summarized by inductive bounds of the form

$$\sigma_{\min}(V_\star^\top X_t^0) \gtrsim (1+c\eta\lambda_\star)^{t-t_0}\gamma,$$
$$\nu_t \lesssim (1+c_{\mathrm{slack}}\eta\lambda_\star)^{t-t_0}\gamma,$$
$$\xi_t \lesssim \kappa_X^{-2}.$$

Therefore the clean signal floor grows while the slack ratio $\chi_t = \nu_t^2/s_t^2$ and the angle remain controlled. In the dashboard (100), this is the favorable middle regime: m_t^w is small because the surrogate signal space is aligned, while $\tau_t = \sqrt{k-r_\star}\chi_t\phi_t$ is harmless only if the product $\chi_t\phi_t$ remains small relative to ψ_t^w. The noise level η_w is the fixed floor. The measured denominator ψ_t^w can become meaningful if the smallest rank-k direction is no longer microscopic. The caveat is the slack tail

$$(1+\delta)\tau_t = (1+\delta)\sqrt{k-r_\star}\,\frac{\nu_t^2}{\lambda_\star}.$$

Lemma 9.2 controls slack relative to the signal factor, but the proxy cares about τ_t/ψ_t^w. Thus slack stays harmless only when its normalized energy remains small relative to the measured rank-k floor. For strict full-rank robustness, the decisive test is still whether the computed

$$\rho_{k,t}^- = \frac{\left\|[P_t^{w,\perp}G_w(M_t^w)P_t^{w,\perp}]_-\right\|_{\mathrm{op}}}{\lambda_k(M_t^w)}$$

falls below $1-\delta$. Phase II is the period in which this can happen: the numerator becomes more signal-aligned, while the denominator becomes acceptable if the smallest retained rank-k direction, possibly a slack direction relative to $M_\star$, has grown enough to make the rank-k chart well conditioned.

Phase III: local refinement and slack risk. Once s_t reaches signal scale, the active reconstruction error contracts (Stöger and Soltanolkotabi 2021, Lemma 9.5):

$$\left\|V_\star^\top(M_\star - M_{t+1}^0)\right\|_F \leq (1-c\eta\lambda_\star)\left\|V_\star^\top(M_\star - M_t^0)\right\|_F + C\eta\lambda_\star\left\|N_tN_t^\top\right\|_F.$$

Their Theorem 9.6 then combines Phase II and Phase III. Its post-spectral hypotheses include $\min\{k,n\} > r_\star$, bounded iterate norm, small slack, alignment $\left\|V_\star^\perp V_{S_{t_0}}\right\|_{\mathrm{op}} \lesssim \kappa_X^{-2}$, an initial signal scale

$$\gamma \leq \frac{c\,\sigma_{\min}(X_\star)}{\min\{k,n\}\kappa_X^2},$$

and a rank-$(2r_\star+1)$ RIP constant of order at most $1/(\kappa_X^4\sqrt{r_\star})$. Under these hypotheses, the additional iteration count is

$$O\left(\frac{1}{\eta\lambda_\star}\log\left(\max\left\{1, \frac{\kappa_X r_\star}{\min\{k,n\}-r_\star}\right\}\frac{\|X_\star\|}{\gamma}\right)\right)$$

and the clean error is bounded by a polynomial factor times $(\gamma/\|X_\star\|)^{21/16}$. In the dashboard (100), Phase III no longer introduces a separate active-error metric; the four-metric view keeps η_w as the fixed external noise floor and tracks the remaining risk through τ_t. This is good for clean active reconstruction, but it is not by itself a strict full-rank robustness certificate. The slack tail can still dominate if the extra columns fail to decay relative to ψ_t^w. Then τ_t/ψ_t^w can flatten or increase, and the strict full-rank ρ-margin stops improving.

The trajectory-aware stopping rule is therefore simple: continue while the measured $\rho_{k,t}^-$ decreases or the shadowing bound (88) tightens; stop, damp, or truncate once the full-rank ρ-margin stops improving.