arXiv:2610.21607v1 [cs.LG] 15 Oct 2026

Market Incentives for AI Safety Investment

Nikita Tsoy
INSAIT
Sofia University "St. Kliment Ohridski"
Sofia, Bulgaria
nikita.tsoy@insait.ai

Abstract

Despite their practical significance, modern AI systems pose significant societal risks, including toxicity, malicious use, and misinformation. Their mitigation depends not only on technical feasibility but also on the incentives of firms that develop and deploy these systems. We show how profit maximization shapes safety investment in an AI supply chain with an upstream LLM provider and several competing downstream firms. In our game-theoretic model, we derive the market-driven level of safety investment and compare it with the levels that maximize market welfare and industry profit. We find that downstream competition generally induces downstream firms to invest more in safety than the industry-profit-optimal level. By contrast, the upstream provider may underinvest relative to these benchmarks because it is not fully compensated for the surplus from safety investments. Finally, we analyze how different regulatory interventions affect the participants' incentives and profits, and the resulting levels of model safety.

1 Introduction

Modern large language models (LLMs) pose significant societal risks, including toxicity, malicious use, and misinformation (Mazeika et al., 2024; Xie et al., 2025). As artificial intelligence (AI) systems become increasingly prevalent and economically important, managing these risks has become a central concern in the machine learning (ML) literature. Although technical methods for AI risk management constitute a large and important research direction (Guldimann et al., 2025; Mazeika et al., 2024; Ouyang et al., 2022; Bai et al., 2022), the economic incentives that determine whether such methods are developed and deployed have received relatively less attention.

These incentives, however, shape the behavior of AI ecosystem participants and hence play a central role in AI risk governance. Indeed, understanding economic incentives and the resulting market outcomes is crucial for designing effective AI regulations that balance model safety and usefulness, especially in common supply-chain situations that involve upstream providers and downstream adopter firms (Hopkins et al., 2025). In this context, even well-meaning regulations aimed at improving model safety may backfire, leading to lower levels of safety investment due to divergent incentives among the participants (Laufer et al., 2025).

Contributions In this paper, we theoretically analyze market-driven safety investment in an AI supply chain. We consider a common setting in which one upstream model provider serves an industry with multiple competing downstream firms and demonstrate how profit-maximizing incentives might encourage safety investments by both the provider and the firms. We focus on lower-impact errors that can be described by neural scaling laws, not catastrophic risks. In our model, the firms can access the LLM by paying a fee to the provider. The LLM lowers firms' production costs but also introduces safety-related costs arising from safety accidents. The risk of these accidents can be reduced through safety investments, e.g., by collecting additional data. We use standard economic models of Cournot

40th Conference on Neural Information Processing Systems (NeurIPS 2026).

competition and Nash-in-Nash bargaining to characterize downstream competition and access fees, and model safety-related costs using neural scaling laws established in ML literature.

Our methodological contributions are two-fold. First, we link the benefits of additional LLM alignment to competition outcomes. This link shows how market characteristics shape investment incentives and how safety investments affect market welfare (a notion of welfare centered on economic outcomes). Second, we suggest an economic- and machine-learning-grounded way to determine fees to the provider. Specifically, this description shows how the provider benefits from safety investments even when the provider does not directly suffer reputational harm from the model's safety accidents.

In our game-theoretic model, we show that competition characteristics strongly influence equilibrium levels of safety investment. Specifically, we find that downstream competition generally induces downstream firms to invest more in safety than the industry-profit-optimal level. By contrast, the upstream provider may underinvest because it is not fully compensated for the surplus generated by its safety investments. Additionally, we show that upstream and downstream regulations that enforce a minimum level of safety lead to higher safety investment levels. However, strict regulation reduces the participants' profits, which might slow adoption or cause the model provider to exit the market.

2 Related Work

Safety incentives in AI supply chains Our work analyzes an important problem of safety incentives in AI supply chains. Similarly to us, Laufer et al. (2025) analyze safety investments and regulations in a market with one model provider and one firm. However, they use a stylized quadratic-cost model of safety investments and do not directly model downstream competition and bargaining processes, leading to a backfiring effect of weak safety regulation that we do not observe in our work. Additionally, they do not study how competition affects safety investment and how regulation affects market welfare. Jagadeesan et al. (2025) analyze entry into the LLM market, assuming that an provider can enter the market only if its model is sufficiently safe compared to others. While this work also analyzes safety-related incentives, it focuses on barriers to entry rather than on the market incentives for safety investments for incumbent firms.

Economics of AI regulation Our work contributes to a growing literature on AI regulation. There are several important directions in this literature. Laufer et al. (2024) study the investment incentives in AI supply chains, similarly to our work. However, they do not examine regulations, link AI investments to market characteristics, or incorporate consumer utility in their analysis. Another line of work (Clark and Hadfield, 2025; Falco et al., 2021) analyzes the organization of audits of AI systems. Wu et al. (2025); Qiu et al. (2025); Xu et al. (2025) study the effects of AI models' openness. Hopkins et al. (2025) discuss regulatory challenges in AI supply chains. Taitler and Ben-Porat (2025) study how generative AI affects high-quality data generation. Finally, a growing body of economic work (e.g., Acemoglu and Lensman, 2024; Gans, 2025) examines the optimal rate of AI adoption, given that AI risks are not fully known.

Economics of ML Our work also belongs to a broader literature on the economics of ML, which studies data as a new production asset employed by modern ML systems (Sun et al., 2025). This literature includes data-sharing incentives (Castro Fernandez, 2023), data valuation (Lin et al., 2025), and strategic incentives for data manipulation (Hardt et al., 2016; Karimi et al., 2022). Most closely, our work is related to Tsoy and Konstantinov (2023), who study data-sharing incentives between competing firms. They use the same model of downstream market competition and a similar model of ML impact on production costs, but do not consider investment incentives or AI supply chains.

Economics of information security Our work is also related to the literature on information security (Moore, 2024), as incentives in information safety can be structurally similar to those in information security. In this context, our work is closely related to the work of Gordon and Loeb (2002). Specifically, our parametric form of safety-related costs (Equation (10)) is mathematically identical to one of the parametric forms suggested by Gordon and Loeb (2002) to describe the risk of security breaches, even though the motivation for this parametric form differs materially. This identity makes our Theorem 5.1 similar to some of Gordon and Loeb (2002) 's results. However, the results of Gordon and Loeb (2002) do not establish the relationship between optimal investments and market characteristics. Additionally, we want to note that our approach to safety-related costs

resembles ex post liability in the information security literature (Moore, 2024), while the regulations we study resemble ex ante safety regulation.

AI risk measurement While we use an abstract notion of safety, deploying AI systems requires developing metrics and evaluation protocols to measure AI risk. See Autio et al. (2024); Giudici et al. (2024); Giudici and Kolesnikov (2026) for the overview of risk metrics and risk frameworks.

Industrial organization Our approach to market competition closely follows the standard approaches in industrial organization (Tirole, 1988). Specifically, we use a widely employed model of Cournot competition with quasilinear quadratic utility consumers (Dixit, 1979). Additionally, to model bargaining between the model provider and the firms, we use the notion of Nash-in-Nash bargaining, widely used in applied and theoretical economic work (Collard-Wexler et al., 2019).

3 Framework

In our work, we introduce a game-theoretic framework that describes safety investments in AI supply chains. We focus on a common practical situation in which a provider offers its foundation model as a service to several competing downstream firms.[1] This model streamlines the firms' operations but also introduces safety-related costs. These costs arise from safety accidents (e.g., unqualified advice, toxicity) that lead to reputational losses or regulatory penalties. To reduce these costs and increase profits, both the provider and the firms invest in model safety, e.g., by collecting additional data.

As an example, consider several competing travel agencies that can optimize their operations by introducing an AI chatbot for customer relations. However, the LLM might exhibit unsafe behavior, e.g., giving unqualified legal advice when the user asks about the terms of booking cancellation, leading to reputational losses. Thus, the firms are willing to either train a safer model themselves using their private data or pay the provider for additional safety training on synthetic or public data.

Objectives In our framework, the provider and the firms maximize their profits

$$\Pi_p = \sum_i \text{AccessFee}_i - \text{InferenceCosts} - \text{ProviderSafetyInvestments}. \tag{1}$$

$$\Pi_i = \text{Sales}_i - \text{ProductionCosts}_i - \text{SafetyRelatedCosts}_i - \text{AccessFee}_i - \text{FirmSafetyInvestments}_i. \tag{2}$$

Here, AccessFee_i is the fee for using the model paid by firm i to the provider, InferenceCosts is the costs for model inference incurred by the provider, ProviderSafetyInvestments is the level of provider's safety investments, Sales_i is the sales of firm i, ProductionCosts_i is the production costs of firm i, $\text{SafetyRelatedCosts}_i$ is the safety-related costs incurred by firm i, and $\text{FirmSafetyInvestments}_i$ is the level of firm i's safety investments.

In this framework, we assume only downstream firms incur safety-related costs. This assumption reflects the difficulty of assigning direct responsibility for model behavior to the foundation model provider, given that the provider does not fully control how downstream firms approach additional safety training and model deployment. Hence, first, we assume that the model provider does not incur direct reputational losses. Second, we assume downstream firms settle fines, i.e., the model provider is not liable for safety-related accidents. We discuss an extension of our framework in which the provider becomes liable for safety accidents in Appendix C.1.

Actions Now, we describe the participants' actions. First, both the provider and firms can invest in safety. These investments might take the form of additional data collection, data generation, additional alignment training, etc. We denote by n and k_i the provider's and firm i's investment levels, respectively. Second, each firm and the provider bilaterally bargain over the model access fee, t_i. Third, each firm makes downstream production decisions, q_i.

These actions occur in the following order.

1. The provider enters the market and chooses its safety investment level, n.
2. Firms simultaneously bargain with the provider over access fees, t_i.

[1] We briefly discuss an extension of our framework to the case of several foundation model providers in Appendix B.

3. Each firm chooses its level of additional safety investments, k_i.
4. The firms compete in the market by making their production decisions, q_i.

At each stage, the participants maximize their profits, given the outcomes of all previous stages.

3.1 Game Stages

We now discuss each game stage in more detail. Consistent with the backward induction logic of game analysis, we discuss each stage in time-reversed order.

Competition Classic competition models describe how each firm's sales depend on the production decisions of **all** firms and how each firm's production costs depend on its own production decisions. Additionally, since higher production volume leads to more LLM usage, our framework introduces two new elements that could depend on production decisions: safety-related costs and access fees.

Taken together, the firms optimize the following market profit functions:

$$\max_{q_i} \Pi_i^m \coloneqq \text{Sales}_i(\boldsymbol{q}) - \text{ProductionCosts}_i(q_i) - \text{SafetyRelatedCosts}_i(q_i, k_i, n) - \text{AccessFee}_i(q_i, t_i).$$

Since sales depend on the production decisions of all firms, we assume that the resulting decisions, $q_i^*(\boldsymbol{k}, \boldsymbol{t}, n)$, constitute Nash equilibrium. We denote the resulting market profits by $\Pi_i^{m,*}(\boldsymbol{k}, \boldsymbol{t}, n)$.

Firms' investments To choose safety investment level, the firms optimize their profits:

$$\max_{k_i} \Pi_i^{m,*}(\boldsymbol{k}, \boldsymbol{t}, n) - \text{FirmSafetyInvestments}_i(k_i). \tag{3}$$

We denote the resulting firms' profits by $\Pi_i^*(\boldsymbol{t}, n)$, the resulting safety investment level by $k_i^*(\boldsymbol{t}, n)$, and the resulting provider's profits by

$$\Pi_p^*(\boldsymbol{t}, n) = \sum_i \text{AccessFee}_i(q_i^*, t_i) - \text{InferenceCosts}(\boldsymbol{q}^*) - \text{ProviderSafetyInvestments}(n).$$

Bargaining To determine the access fees, we use the Nash-in-Nash bargaining solution. This concept assumes the provider and each firm negotiate bilaterally over several rounds. At each even round, the provider suggests a fee level. The negotiating firm can accept or refuse this level, thereby limiting the fee size. At each odd round, the firm suggests a fee level, and the provider can either accept or refuse it (see Collard-Wexler et al., 2019, for game-theoretic foundation). This process results in the Nash bargaining solution (Binmore et al., 1986) for each negotiating pair, taking the outcomes of other negotiations as given. The resulting fees should constitute a Nash equilibrium.

Formally, denote $\boldsymbol{t}_{-i} = (t_1, \ldots, t_{i-1}, t_{i+1}, \ldots, t_m)$, a vector of fees without i's entry. According to the Nash-in-Nash concept, each pair maximizes the following objective:

$$\max_{t_i} (\Pi_i^*(\boldsymbol{t}, n) - \Pi_{i,-i}^*(\boldsymbol{t}_{-i}, n))^{1-\phi} (\Pi_p^*(\boldsymbol{t}, n) - \Pi_{p,-i}^*(\boldsymbol{t}_{-i}, n))^{\phi}, \tag{4}$$

where ϕ is the provider's bargaining power and $\Pi_{i,-i}^*$ and $\Pi_{p,-i}^*$ are firm i's and the provider's profits if all firms but i have access to the LLM. We denote the resulting fees by $t_i^*(n)$.

Provider's investments Finally, the provider optimizes profits to choose its safety investments

$$\max_n \Pi_p^*(\boldsymbol{t}^*(n), n).$$

We denote the resulting investment level by n^*.

4 Game-Theoretic Model

To obtain analytical results, we focus on a specific instantiation of the framework. In the following sections, we describe the analytic form of each component in greater detail.

4.1 Production-Related Elements

This section describes production-related elements: competition model, inference costs, and access fees. We focus on the classical Cournot model of competition of Dixit (1979), similarly to Tsoy and Konstantinov (2023). This model, widely used in economic studies (Choné and Linnemer, 2019), describes a market consisting of m firms, where firm i produces its own good i. In our analysis, we assume that $m \geq 2$. Each firm maximizes its profits by tailoring its production volume to consumer demand. To formally describe this model, we need to specify demand (i.e., consumer behavior) and supply (i.e., firm behavior) factors.

Demand There are M representative consumers (i.e., M is the market size), each representing the average behavior of all consumers. Their utility is a quasi-linear quadratic utility function (Choné and Linnemer, 2019), which they maximize given market prices and budget constraint

$$\max_{\boldsymbol{q},g} u(\boldsymbol{q}, g) = \sum_{i=1}^{m} q_i - \Big(\frac{\sum_i q_i^2 + \gamma \sum_{i \neq j} q_i q_j}{2}\Big) + g \quad \text{s.t.} \quad \boldsymbol{q}^\mathsf{T}\boldsymbol{p} + g \leq G, \tag{5}$$

where q_i is the consumed amount of good i, g is the remaining money, $\boldsymbol{p}$ is the vector of prices for each good, G is the budget of the consumer, and $\gamma \in \big(-\frac{1}{m-1}, 1\big)$ is the substitutability between the firms' goods. Higher γ corresponds to more similar goods and, hence, "more intense" competition.

The first-order conditions for this problem yield m demand equations that link prices and quantities.

Supply Firm i produces Mq_i units of good i, resulting in market profits

$$\Pi_i^m = Mq_ip_i - \text{ProductionCosts}_i(q_i) - \text{SafetyRelatedCosts}_i(q_i, k_i, n) - \text{AccessFee}_i(q_i, t_i).$$

We assume that production costs are linear in q_i and in the number of tokens required for production. Additionally, we assume that the number of safety-related accidents grows linearly with the number of tokens (see Infante, 2025, for a similar idea). These assumptions result in the following costs:

$$\text{ProductionCosts}_i(q_i) + \text{SafetyRelatedCosts}_i(q_i, k_i, n) = Mq_ic_i(k_i, n), \tag{6}$$

where c_i includes both production and safety-related costs.

Inference costs and access fees Assuming that inference costs grow linearly in tokens,[2] we get

$$\text{InferenceCosts} = M\Big(\sum_i q_i\Big)w. \tag{7}$$

Under the same assumption, we consider two access fee types: token- and subscription-based,

$$\text{AccessFee}_i = Mq_it_i \text{ (token-based)}, \quad \text{AccessFee}_i = Mq_iw + t_i \text{ (subscription-based)}. \tag{8}$$

Denote $c_i^{\text{eff}} \coloneqq c_i + t_i$ and $c_i^{\text{eff}} \coloneqq c_i + w$ in token- and subscription-based cases, respectively. We get

$$\Pi_i^m = Mq_i(p_i - c_i^{\text{eff}}) \text{ (token-based)}, \ \Pi_i^m = Mq_i(p_i - c_i^{\text{eff}}) - t_i \text{ (subscription-based)}.$$

Token-based fees effectively increase the production costs, reducing production and token use by the firms. By contrast, subscription-based fees do not affect the production decisions. This logic parallels the standard comparison between marginal and lump-sum taxes (Stiglitz and Rosengard, 2015).

Equilibrium profits In Cournot competition, the firms decide how much to produce to maximize their profits. Given these quantities, prices are determined by the demand equations so that markets clear (i.e., quantities produced equal quantities demanded). The following standard lemma (Amir et al., 2017), specialized to our setting, characterizes the resulting equilibrium profits.

Lemma 4.1. *Assume that* $\forall i\ (2-\gamma)(1-c_i^{\text{eff}}) \geq \gamma \sum_{j \neq i}(c_i^{\text{eff}} - c_j^{\text{eff}})$ *and* G *is sufficiently big (e.g.,* $G \geq m$*). Then, firm* i*'s equilibrium profits follow*

$$\Pi_i^{m,*}(\boldsymbol{c}^{\text{eff}}, t_i) = M(q_i^*)^2 \text{ (token-based)}, \quad \Pi_i^{m,*}(\boldsymbol{c}^{\text{eff}}, t_i) = M(q_i^*)^2 - t_i \text{ (subscription-based)},$$

where $q_i^* = \frac{2-\gamma-(2+\gamma(m-2))c_i^{\text{eff}}+\gamma\sum_{j\neq i} c_j^{\text{eff}}}{(2-\gamma)(2+\gamma(m-1))}$.

The first assumption ensures that the consumer demands each good. The second one ensures that the consumer has enough money to spend. As we can see, firm i's profits decrease in its production costs and increase in the production costs of competitors; γ exacerbates this dependency.

[2] See Appendix E.1 for the discussion of nonlinear inference costs.

4.2 Safety-Related Elements

Now, we consider safety-related costs and safety investments. In our work, we focus on safety-related properties that can be improved through additional data collection or generation, along with additional training. These properties include all properties that can be improved through reinforcement learning from human feedback (RLHF) (Ouyang et al., 2022; Bai et al., 2022), adversarial training (Mazeika et al., 2024; Dékány et al., 2025), or supervised fine-tuning (Bianchi et al., 2024), such as avoiding toxicity, refusing to give unqualified advice, and refusing to assist with misconduct.

Safety-related costs We assume that firms have symmetric costs[3] and that, before LLM, they equal

$$c_i = a_0. \tag{9}$$

The LLM access decreases production costs but introduces safety-related costs that grow linearly with the number of safety accidents. We assume that the probability of accidents depends polynomially on the number of data points used for safety training, consistent with the learning-theory logic (Tsybakov, 2004) and neural scaling laws (Hoffmann et al., 2022). So, the final costs follow

$$c_i = a + b/(n_i^{\text{eff}}(k_i, n))^{\beta}, \tag{10}$$

where n_i^{eff} corresponds to the effective number of data points used for safety training of firm i's model, $\beta > 0$ corresponds to learning easiness, a is the sum of production and safety-related costs for a Pareto optimal model, and b corresponds to safety penalty that arises due to generalization error.[4]

In our model, greater safety investments will not decrease the LLM's utility. The incentive for cost minimization already determines the desired balance between safety and utility and accounts for the alignment tax (Ouyang et al., 2022). Additional data collection only reduces the generalization error in finding the cost-minimizing point and, hence, can only be beneficial for the costs (Sukenik and Lampert, 2025).[5] To practically achieve this balance, one could use a reply buffer of utility-enhancing data or techniques to mitigate alignment tax (Niu et al., 2026; Sun et al., 2026).

Investment costs Both the provider and the downstream firms invest in safety. The provider performs upstream training, which benefits all firms. Each firm performs downstream training, which benefits only that firm. We do not assume a specific training pipeline but require that the total pipeline cost, including data collection and training, scales linearly with the amount of training data.

The procedure might be the following. First, the provider aligns the model using generated data. Then, the downstream firms further align the model using their private data, possibly with different training procedures. Our requirement allows each participant to use replay buffers of utility-enhancing data whose sizes scale proportionally with the amount of safety data. The data may be collected through human labeling, RLAIF (Bai et al., 2022), red-teaming (Perez et al., 2022), or other procedures.

We interpret safety investment levels, n and k_i, as the number of points collected for safety training. Since the firms may collect more relevant data for their tasks, we consider their data more effective. Specifically, following the results of Hanneke and Kpotufe (2019, Remark 2), we assume

$$n_i^{\text{eff}} = k_i + (n+1)^{\rho}, \tag{11}$$

where $\rho \in (0, 1)$ corresponds to the transferability of the provider's data for the firms' tasks. We use $n + 1$ instead of n to represent the model's baseline safety capabilities before additional training. To highlight economic meaning, we often refer to n_i^{eff} as cumulative level of safety investments. To simplify the notation, we use the following change of variables in the rest of the paper:

$$n + 1 \mapsto n.$$

Our assumptions about the training pipeline yield the following investment costs:

$$\text{ProviderSafetyInvestments} = \lambda_p(n-1), \quad \text{FirmSafetyInvestments}_i = \lambda k_i. \tag{12}$$

We allow $\lambda_p \neq \lambda$ since the provider's and the firms' training pipelines may differ.

[3] We make costs symmetric to study the model provider's investments analytically. Without this assumption, the provider's investments behave similarly in qualitative terms but become less interpretable; see Appendix D.1.

[4] Our analysis will be qualitatively similar if we use a different form of the cost function; see Appendix D.2.

[5] Using the results of Sukenik and Lampert (2025), we can also extend the framework to analyze several safety properties (see Appendix D.3). This extension will behave similarly to the one presented in the main text.

4.3 Investment Problems and Outcomes

We can now derive the precise safety investment problems the participants face.

Absence of disruption Equation (3), given Theorem 4.1 and Equation (10), results in analytically intractable expressions. To get an analytical answer, we assume

$$(2-\gamma)(1-a-w) \gg (2+\gamma(m-2))(b+\delta),$$

where $\delta \coloneqq a_0 - a - w$. This condition has two implications. First, a small δ implies that access to the LLM does not allow one firm to drive all competitors out of the market, i.e., the LLM is not disruptive. Second, a small b implies that the penalty from the lack of additional safety data is not too severe. Also, note that this condition implies the first condition in Theorem 4.1.[6]

We additionally assume that the introduction of LLM is beneficial for the production process and profitable for all participants. In particular, this assumption implies that there exists a feasible vector of fees $\boldsymbol{t}$ that makes the left- and right-hand sides in Equation (4) non-negative.

Approximation of Equations (3) and (4) The first order approximation of market profits lead to the following approximation of Equation (3) (see Appendix A.2):

$$\max_{k_i} -\frac{Ab}{(k_i+n^\rho)^\beta} - \lambda k_i,\ A \coloneqq \frac{2(2+\gamma(m-2))Mq^*}{(2-\gamma)(2+\gamma(m-1))},\ q^* \coloneqq \frac{1-a-w}{2+\gamma(m-1)}. \tag{13}$$

Here, A is the sensitivity of a firm's profits to changes in its effective cost, and q^* denotes the production level with the Pareto-optimal model. Notice that this equation implies that each firm's private investment does not depend on the private investments of other firms.

Now, we consider Equation (4). We get the following firms' profits:

$$\mathit{\Pi}_i^* = \mathit{\Pi}_i^{m,*}(\boldsymbol{c}^{\text{eff}}, t_i) - \lambda k_i, \quad \mathit{\Pi}_{i,-i}^* = \mathit{\Pi}_i^{m,*}(\boldsymbol{c}_{-i}^{\text{eff}}(a_0), 0),$$

where $\boldsymbol{c}_{-i}^{\text{eff}}(x) \coloneqq (c_1^{\text{eff}}, \dots, c_{i-1}^{\text{eff}}, x, c_{i+1}^{\text{eff}}, \dots, c_m^{\text{eff}})$, a vector of effective costs with ith entry changed. And we get the following provider's profits:

$$\mathit{\Pi}_p^* = \sum_j \text{AccessFee}_j(q_i^*(\boldsymbol{c}^{\text{eff}}), t_j) - Mw\sum_j q_j^*(\boldsymbol{c}^{\text{eff}}) - \lambda_p(n-1),$$

$$\mathit{\Pi}_{p,-i}^* = \sum_{j\neq i} \text{AccessFee}_j(q_j^*(\boldsymbol{c}_{-i}^{\text{eff}}(a_0)), t_j) - Mw\sum_{j\neq i} q_j^*(\boldsymbol{c}_{-i}^{\text{eff}}(a_0)) - \lambda_p(n-1).$$

Thus, Equation (4) becomes

$$\max_{t_i}(\mathit{\Pi}_i^m(\boldsymbol{c}^{\text{eff}}, t_i) - \lambda k_i - \mathit{\Pi}_i^m(\boldsymbol{c}_{-i}^{\text{eff}}(a_0), 0))^{1-\phi}(\mathit{\Pi}_p^* - \mathit{\Pi}_{p,-i}^*)^\phi.$$

And its first-order approximation is (see Appendix A.3)

$$\arg\max_{t_i}\Big(A\delta - \frac{Ab}{(k_i+n^\rho)^\beta} - \lambda k_i - B(t_i-w)\Big)^{1-\phi}(t_i-w)^\phi, \tag{14}$$

where $B = A$ for token-based fees and $B = Mq^*$ for subscription-based fees. For subscription-based fees, t_i is redefined so that $\text{AccessFee}_i = Mq^*t_i$.

Outcomes In our analysis, we track several key metrics: cumulative level of safety investments (n_i^{eff}), industry profit, and market welfare. Industry profit is defined as profits of all participants, $\mathit{\Pi}^{\text{ind}} \coloneqq \mathit{\Pi}_p + \sum_i \mathit{\Pi}_i$. Market welfare is defined as consumer utility plus industry profits, $W^m \coloneqq Mu^* + \mathit{\Pi}^{\text{ind}}$, where $u^* \coloneqq u(\boldsymbol{q}^*, G - \sum_i p_i^* q_i^*)$ is the utility of the consumer after competition. **Note** that market welfare does not directly account for the harms to the consumer from safety accidents.[7]

5 Safety Investments

This section analyzes market-driven investment levels and compares them with industry- and market-welfare-optimal ones.

[6] We briefly discuss how to relax this assumption and model firm exit in Appendix E.2.

[7] We consider a notion of welfare that accounts for harm to consumers in Appendix C.2.

Firms' investments Market-driven firms' investments follow the following proposition.

Proposition 5.1. *The firms' level of investments follows the following formula:*

$$k^* = \max(0, k_m - n^\rho),\ k_m \coloneqq (\beta b A/\lambda)^{1/(1+\beta)}.$$

Note that the firm's safety investment, k^*, is non-increasing in the provider's investment, n, while the cumulative level of investments, $n^{\text{eff},*} = k^* + n^\rho$, is non-decreasing in n. In other words, although firms decrease investments in response to an increase in the provider's investment, this decrease does not reduce cumulative investments.

Provider's investments Equation (14) and Theorem 5.1 give (see Appendix A.5)

$$\Pi_p^* = \frac{m\phi M q^*}{B}\Big(A\delta - \lambda k_m - \frac{Ab}{\max(n^\rho, k_m)^\beta} + \lambda \min(n^\rho, k_m)\Big) - \lambda_p n + \lambda_p. \tag{15}$$

The LLM is profitable for all participants if and only if $\Pi_p^* \geq 0$. Also note that, as predicted, the subscription-based fees yield larger profits since they correspond to the smaller B.

Theorem 5.2. *Denote* $n_s \coloneqq \big(m\frac{\phi M q^*}{B}\frac{\rho\lambda}{\lambda_p}\big)^{1/(1-\rho)}$ *and* $n_b \coloneqq \big(m\frac{\phi M q^*}{B}\frac{\rho\beta b}{\lambda_p}A\big)^{1/(1+\rho\beta)}$. *The provider's and cumulative levels of investments follow the following formulas:*

$$n^* = \max(1, \min(n_b, n_s)),\ (n^*)^\rho + k^* = \max(n_b^\rho, k_m, 1).$$

n_s is increasing in ϕ and λ, is decreasing in λ_p, and does not depend on M and b. n_b is increasing in ϕ, M, and b, is decreasing in λ_p, and does not depend on λ. k_m is increasing in M and b, is decreasing in λ, and does not depend on λ_p and ϕ. k_m is decreasing in m for $\gamma > 0$, while n_s and n_b are increasing in m. n_s does not depend on γ, while the dependence of k_m and n_b on γ is non-monotonic, increasing for small γ and decreasing for big γ.

Thus, we have the following comparative statics Table 1. We see that investments are non-decreasing in market size (M), sensitivity of safety-related costs to investments (b), and the provider's bargaining power (ϕ). Additionally, the investments are generally non-increasing in data collection costs (λ and λ_p) and non-optimizable production costs ($a + w$). At the same time, the dependence on competition-related factors, m and γ, is non-monotonic.

Table 1: Comparative statics

	Firms' level, k^*	Provider's level, n^*	Cumulative level, $k^* + (n^*)^\rho$
Non-decreasing	M and b	M, b, ϕ, λ, and m	M, b, ϕ, and m for big m
Non-increasing	λ, $a + w$, and m	λ_p and $a + w$	λ, λ_p, $a + w$, and m for small m

We also observe several interesting facts. First, the provider can invest to the point where the firms do not invest at all if m is large enough. It happens because the provider's investments benefit all firms (i.e., they exhibit a positive externality). Second, overly fierce downstream competition might lower the resulting cumulative safety, $(n^*)^\rho + k^*$, since k_m and n_b decrease in γ for large γ, and k_m decreases in m. While downstream competition induces competition in safety investments among the firms, as we show below, it also lowers the firms' profits, inhibiting their ability to invest in general. Third, the provider invests more when using subscription-based fees. Thus, subscription fees not only increase the provider's profits but also increase safety investments.

Industry profit and market welfare Now, we derive the optimal investment levels for industry profits and market welfare when participants coordinate their investments and fees.

Proposition 5.3. *The industry, k_i^{ind}, and welfare, k_i^{mw}, optimal firms' investments follow*

$$k_i^x(n) = \max(0, k_m^x - n^\rho),\ k_m^x \coloneqq (\beta b A^x/\lambda)^{1/(1+\beta)},$$

where $x \in \{ind, mw\}$ and

$$A^{ind} \coloneqq \frac{2Mq^*}{2 + \gamma(m-1)},\ A^{mw} \coloneqq \frac{(3 + \gamma(m-1))Mq^*}{2 + \gamma(m-1)}.$$

The optimal firms' investments are generally lower than the equilibrium ones. We can see this by analyzing the behavior of the parameter k_m, which drives firms' investments. If $\gamma > 0$, we get $A^{\text{ind}} < A$ and, hence, $k_m^{\text{ind}} \leq k_m$. If $\gamma^2(m-1)+\gamma > 2$, we get $A^{\text{mw}} < A$ and, hence, $k_m^{\text{mw}} \leq k_m$. This result implies that competition pressure incentivizes the firms to invest more in safety. We also note that $k_m^{\text{ind}} < k_m^{\text{mw}}$ since the welfare also accounts for the consumers' utility.

Proposition 5.4. *If the firms' investments follow $k_i(n) = \max(0, k_0 - n^\rho)$ and $k_0 \geq k_m^x$, the provider's optimal investments are*

$$n^x = \max(1, n_e^x),\ n_e^x \coloneqq \begin{cases} n_s^{opt}, & (n_s^{opt})^\rho < k_0, \\ n_b^x, & (n_b^x)^\rho > k_0, \\ k_0^{1/\rho}, & \textit{otherwise}, \end{cases}$$

where $n_s^{opt} \coloneqq \left(m\frac{\rho\lambda}{\lambda_p}\right)^{1/(1-\rho)}$ and $n_b^x \coloneqq \left(m\frac{\rho\beta b}{\lambda_p}A^x\right)^{1/(1+\rho\beta)}$.

We observe two opposite effects. On the one hand, the absence of bargaining increases the investments (note that the multiplier associated with the presence of bargaining satisfies $\phi M q^*/B < 1$). On the other hand, the absence of competition decreases them ($A^x < A$ for large γ). However, the second mechanism works only if the firms' investments are zero.

The comparison between equilibrium and industry-optimal investment levels reveals a collective-action problem. When the industry-optimal provider's investments are higher than the equilibrium ones, the firms are collectively willing to compensate the provider for the additional investments. However, due to the nature of fee bargaining, this compensation does not naturally occur.

6 Effects of Regulation

In this section, we discuss how regulatory actions can impact our outcome metrics. We assume perfect enforcement of regulations. Specifically, we assume that one can implement a properly designed audit procedure (with sufficiently large fines) that corresponds to fixed administrative costs for the regulator and does not change incentives for safety investments.

Downstream auditing A possible regulatory approach might enforce the safety of the resulting models. Assuming that the enforcement will focus on the probability of safety accidents, it will result in constraint $n_i^{\text{eff}} \geq k_0$.

Proposition 6.1. *Consider the downstream regulation $n_i^{eff} \geq k_0$. If $k_0 \leq k_m$, the investment levels will not change. If $k_0 \geq k_m$, we get*

$$k^* = \max(0, k_0 - n^\rho),\ n^* = \max(1, n_e),\ \textit{where } n_e = \begin{cases} n_s, & n_s^\rho < k_0, \\ n_b, & n_b^\rho > k_0, \\ k_0^{1/\rho}, & \textit{otherwise}. \end{cases}$$

And, we have $n^* \geq n^{*,old}$, where $n^{*,old}$ is the investment level in no regulation regime.

The regulation does not reduce the participants' investments. However, the firms incur additional data-collection costs and may refuse to use the LLM under overly strict regulation (see Equation (15) with $k_m \mapsto k_0$). The effect on industry profit and market welfare is ambiguous. On the one hand, regulation increases firms' investment, which can reduce these metrics when competition is severe, as shown in Theorem 5.3. On the other hand, regulation also increases the provider's investment. These opposing effects imply that regulation increases the metrics when $k_0 \leq \max(k_m^x, n_s^\rho)$, where the metrics either increase by Theorem 5.3 or the firms make zero investments and benefit from higher provider investment. Otherwise, the regulation has a negative effect.

Upstream auditing Another approach is to audit the provider's model, essentially requiring $n \geq n_0$.

Proposition 6.2. *Under the upstream regulation $n \geq n_0$ (where $n_0 \geq 1$), the provider investment level will be $n^* = \max(n_0, \min(n_s, n_b))$ and the firms' response $k^* = \max(0, k_m - n^\rho)$ will not change.*

The regulation does not reduce the cumulative level of investments due to Theorem 5.1. However, again, overly strict regulation might force the provider to exit the market. As for the industry profit and market welfare, in the case of $k_m^x \leq k_m$, the regulation is beneficial if $n_0 \leq n^x$ by Theorem 5.4. Since $n_s < n_s^{\text{opt}}$, the moderate regulation is always beneficial when the firms make non-zero investments. Otherwise, when the firms make zero investments, the regulation is beneficial if the market is not overly competitive or the provider's bargaining power is low.

Indirect regulations The regulator might also increase safety investments through less direct actions. Theorem 5.2 suggests several channels: safety penalty, b, and investment costs, λ and λ_p. By increasing fines for safety accidents, the regulator indirectly increases b (and possibly a, but in this discussion we assume that the increase in a is small). This increase raises cumulative investments (through k_m or n_b). However, as previously, this intervention reduces the participants' profits. By subsidizing safety investments, the regulator decreases λ or λ_p. If the firms make non-zero investments, the subsidy to the firms leads to higher cumulative investments. Otherwise, the subsidy to the provider leads to higher cumulative investments. However, both subsidies require additional funds from the regulator. Also note that the subsidy to the provider might help address the collective action problem identified in Theorem 5.4, allowing the regulator to finance it partially through higher taxes on the participants.

7 Conclusion

We analyzed how profit maximization might incentivize safety investments by the participants of the AI supply chain. In our framework, downstream firms are interested in safer models because they incur safety-related costs. At the same time, the provider is interested in safety because it can charge higher access fees to the firms. We characterized the market-driven level of safety investments and compared it with the industry-profit- and market-welfare-optimal ones. This comparison shows that downstream competition generally leads to higher levels of safety investment.

We outlined several ways the regulator could increase investment in safety. Each of these approaches has upsides and downsides. Auditing and fines do not require additional funds and lead to higher safety investments but reduce profits from the LLM and, hence, might slow AI adoption. Subsidies lead to higher investment levels and profits but require additional funds from the regulator.

We hope that our work contributes to AI governance by highlighting the economic challenges associated with AI safety investment for both policymakers and industry.

Limitations While we demonstrate how profit maximization incentivizes safety investment and obtain analytical expressions for its levels, our analysis has several limitations. First, we assume that the firms are symmetric and, hence, have similar production processes. This assumption might not hold in industries with substantial horizontal differentiation among firms. Second, our analysis does not account for catastrophic risks posed by the models. Instead, we focus on lower-impact errors that can be described by neural scaling laws. Third, our results rely on the absence-of-disruption assumption, which might not hold in more dynamic and innovative industries. Fourth, we assume perfect enforcement in our analysis of regulations and assume that one can design a proper audit procedure by imposing sufficiently large fines. Finally, we assume specific functional forms for the elements of our framework. These forms capture the main conceptual properties of the elements but might differ from real-world relationships.

Future directions We foresee several possible extensions of our model. First, a more general model might analyze asymmetric firms, which is important in horizontally differentiated industries. As a related point, it might be interesting to consider cascading effects in which firms drive competitors out of the market by gaining access to the LLM. Another possible extension is a version of welfare that accounts for the harms to consumers from safety accidents. Finally, it might be interesting to consider a more complex provider market, which might include open-weight models or several competing providers, as we discuss in Appendix B.

Acknowledgments

This research was partially funded from the Ministry of Education and Science of Bulgaria (support for INSAIT, part of the Bulgarian National Roadmap for Research Infrastructure). The authors thank Nikola Konstantinov for helpful comments on the paper's positioning and writing.

References

Acemoglu, D. and Lensman, T. (2024). Regulating Transformative Technologies. *American Economic Review: Insights*, 6(3):359–76.

Amir, R., Erickson, P., and Jin, J. (2017). On the microeconomic foundations of linear demand for differentiated products. *Journal of Economic Theory*, 169:641–665.

Autio, C., Schwartz, R., Dunietz, J., Jain, S., Stanley, M., Tabassi, E., Hall, P., and Roberts, K. (2024). *Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile*. NIST Trustworthy and Responsible AI, National Institute of Standards and Technology, Gaithersburg, MD.

Bai, Y., Kadavath, S., Kundu, S., Askell, A., Kernion, J., Jones, A., Chen, A., Goldie, A., Mirhoseini, A., McKinnon, C., Chen, C., Olsson, C., Olah, C., Hernandez, D., Drain, D., Ganguli, D., Li, D., Tran-Johnson, E., Perez, E., Kerr, J., Mueller, J., Ladish, J., Landau, J., Ndousse, K., Lukosuite, K., Lovitt, L., Sellitto, M., Elhage, N., Schiefer, N., Mercado, N., DasSarma, N., Lasenby, R., Larson, R., Ringer, S., Johnston, S., Kravec, S., Showk, S. E., Fort, S., Lanham, T., Telleen-Lawton, T., Conerly, T., Henighan, T., Hume, T., Bowman, S. R., Hatfield-Dodds, Z., Mann, B., Amodei, D., Joseph, N., McCandlish, S., Brown, T., and Kaplan, J. (2022). Constitutional AI: Harmlessness from AI Feedback. arXiv:2212.08073 [cs.CL].

Bianchi, F., Suzgun, M., Attanasio, G., Rottger, P., Jurafsky, D., Hashimoto, T., and Zou, J. (2024). Safety-Tuned LLaMAs: Lessons From Improving the Safety of Large Language Models that Follow Instructions. In *The Twelfth International Conference on Learning Representations*.

Binmore, K., Rubinstein, A., and Wolinsky, A. (1986). The Nash Bargaining Solution in Economic Modelling. *The RAND Journal of Economics*, 17(2):176–188.

Castro Fernandez, R. (2023). Data-Sharing Markets: Model, Protocol, and Algorithms to Incentivize the Formation of Data-Sharing Consortia. *Proc. ACM Manag. Data*, 1(2).

Choné, P. and Linnemer, L. (2019). The quasilinear quadratic utility model: An overview. CESifo working paper.

Clark, J. and Hadfield, G. K. (2025). Regulatory Markets for AI Safety. arXiv:2001.00078 [cs.CY].

Collard-Wexler, A., Gowrisankaran, G., and Lee, R. S. (2019). "Nash-in-Nash" Bargaining: A Microfoundation for Applied Work. *Journal of Political Economy*, 127(1):163–195.

Dékány, C., Balauca, S., Dimitrov, D. I., Staab, R., and Vechev, M. (2025). MixAT: Combining Continuous and Discrete Adversarial Training for LLMs. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*.

Dixit, A. (1979). A Model of Duopoly Suggesting a Theory of Entry Barriers. *The Bell Journal of Economics*, 10(1):20–32.

Falco, G., Shneiderman, B., Badger, J., Carrier, R., Dahbura, A., Danks, D., Eling, M., Goodloe, A., Gupta, J., Hart, C., et al. (2021). Governing AI safety through independent audits. *Nature Machine Intelligence*, 3(7):566–571.

Gans, J. S. (2025). How learning about harms impacts the optimal rate of artificial intelligence adoption. *Economic Policy*, 40(121):199–219.

Giudici, P., Centurelli, M., and Turchetta, S. (2024). Artificial Intelligence risk measurement. *Expert Systems with Applications*, 235:121220.

Giudici, P. and Kolesnikov, V. (2026). SAFE AI metrics: An integrated approach. *Machine Learning with Applications*, 23:100821.

Gordon, L. A. and Loeb, M. P. (2002). The economics of information security investment. *ACM Trans. Inf. Syst. Secur.*, 5(4):438–457.

Guldimann, P., Spiridonov, A., Staab, R., Jovanović, N., Vero, M., Vechev, V., Gueorguieva, A.-M., Balunović, M., Konstantinov, N., Bielik, P., Tsankov, P., and Vechev, M. (2025). COMPL-AI Framework: A Technical Interpretation and LLM Benchmarking Suite for the EU Artificial Intelligence Act. arXiv:2410.07959 [cs.CL].

Hanneke, S. and Kpotufe, S. (2019). On the Value of Target Data in Transfer Learning. In Wallach, H., Larochelle, H., Beygelzimer, A., d'Alché-Buc, F., Fox, E., and Garnett, R., editors, *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc.

Hardt, M., Megiddo, N., Papadimitriou, C., and Wootters, M. (2016). Strategic Classification. In *Proceedings of the 2016 ACM Conference on Innovations in Theoretical Computer Science*, ITCS '16, pages 111–122, New York, NY, USA. Association for Computing Machinery.

Ho, K. and Lee, R. S. (2019). Equilibrium Provider Networks: Bargaining and Exclusion in Health Care Markets. *American Economic Review*, 109(2):473–522.

Hoffmann, J., Borgeaud, S., Mensch, A., Buchatskaya, E., Cai, T., Rutherford, E., de Las Casas, D., Hendricks, L. A., Welbl, J., Clark, A., Hennigan, T., Noland, E., Millican, K., van den Driessche, G., Damoc, B., Guy, A., Osindero, S., Simonyan, K., Elsen, E., Vinyals, O., Rae, J., and Sifre, L. (2022). An empirical analysis of compute-optimal large language model training. In Koyejo, S., Mohamed, S., Agarwal, A., Belgrave, D., Cho, K., and Oh, A., editors, *Advances in Neural Information Processing Systems*, volume 35, pages 30016–30030. Curran Associates, Inc.

Hopkins, A., Cen, S. H., Struckman, I., Ilyas, A., Videgaray, L., and Mądry, A. (2025). AI Supply Chains: An Emerging Ecosystem of AI Actors, Products, and Services. *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, 8(2):1266–1277.

Infante, A. (2025). When Will AI Transform the Economy? `https://andreinfante.substack.com/p/when-will-ai-transform-the-economy`. Accessed: 23-02-2026.

Jagadeesan, M., Jordan, M. I., and Steinhardt, J. (2025). Safety versus performance: How multi-objective learning reduces barriers to market entry. *Proceedings of the National Academy of Sciences*, 122(42):e2510004122.

Karimi, A.-H., Barthe, G., Schölkopf, B., and Valera, I. (2022). A Survey of Algorithmic Recourse: Contrastive Explanations and Consequential Recommendations. *ACM Comput. Surv.*, 55(5).

Laufer, B., Kleinberg, J., and Heidari, H. (2024). Fine-Tuning Games: Bargaining and Adaptation for General-Purpose Models. In *Proceedings of the ACM Web Conference 2024*, WWW '24, pages 66–76, New York, NY, USA. Association for Computing Machinery.

Laufer, B., Kleinberg, J., and Heidari, H. (2025). The Backfiring Effect of Weak AI Safety Regulation. In *Proceedings of the 5th ACM Conference on Equity and Access in Algorithms, Mechanisms, and Optimization*, EAAMO '25, page 300, New York, NY, USA. Association for Computing Machinery.

Lin, H., Wan, S., Xie, Z., Chen, K., Zhang, M., Shou, L., and Chen, G. (2025). A Comprehensive Study of Shapley Value in Data Analytics. *Proc. VLDB Endow.*, 18(9):3077–3092.

Mazeika, M., Phan, L., Yin, X., Zou, A., Wang, Z., Mu, N., Sakhaee, E., Li, N., Basart, S., Li, B., Forsyth, D., and Hendrycks, D. (2024). HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal. In Salakhutdinov, R., Kolter, Z., Heller, K., Weller, A., Oliver, N., Scarlett, J., and Berkenkamp, F., editors, *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 35181–35224. PMLR.

Moore, T. (2024). Security Economics Knowledge Guide. In Rashid, A., Cherdantseva, Y., Martin, A., and Schneider, S., editors, *CyBOK Knowledge Guides and Topic Guides*. University of Bristol.

Niu, Y., Xiao, H., Liu, D., Chen, N., and Li, J. (2026). Mitigating the Safety Alignment Tax with Null-Space Constrained Policy Optimization. In *The Fourteenth International Conference on Learning Representations*.

Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., Schulman, J., Hilton, J., Kelton, F., Miller, L., Simens, M., Askell, A., Welinder, P., Christiano, P. F., Leike, J., and Lowe, R. (2022). Training language models to follow instructions with human feedback. In Koyejo, S., Mohamed, S., Agarwal, A., Belgrave, D., Cho, K., and Oh, A., editors, *Advances in Neural Information Processing Systems*, volume 35, pages 27730–27744. Curran Associates, Inc.

Perez, E., Huang, S., Song, F., Cai, T., Ring, R., Aslanides, J., Glaese, A., McAleese, N., and Irving, G. (2022). Red Teaming Language Models with Language Models. In Goldberg, Y., Kozareva, Z., and Zhang, Y., editors, *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 3419–3448, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.

Qiu, T., Laufer, B., Kleinberg, J., and Heidari, H. (2025). Modeling the Economic Impacts of AI Openness Regulation. arXiv:2507.14193 [cs.GT].

Stiglitz, J. E. and Rosengard, J. K. (2015). *Economics of the public sector: Fourth international student edition*. WW Norton & Company.

Sukenik, P. and Lampert, C. (2025). Generalization in multi-objective machine learning. *Neural Computing and Applications*, 37(30):24669–24683.

Sun, G., Zhang, S., Wang, L., Zhu, J., Su, H., and Zhong, Y. (2026). Safety Alignment as Continual Learning: Mitigating the Alignment Tax via Orthogonal Gradient Projection. arXiv:2602.07892 [cs.LG].

Sun, H., Wu, Y., Cheng, Y., and Chu, X. (2025). Game Theory Meets Large Language Models: A Systematic Survey. In Kwok, J., editor, *Proceedings of the Thirty-Fourth International Joint Conference on Artificial Intelligence, IJCAI-25*, pages 10669–10677. International Joint Conferences on Artificial Intelligence Organization. Survey Track.

Taitler, B. and Ben-Porat, O. (2025). Selective Response Strategies for GenAI. In Singh, A., Fazel, M., Hsu, D., Lacoste-Julien, S., Berkenkamp, F., Maharaj, T., Wagstaff, K., and Zhu, J., editors, *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 58189–58226. PMLR.

Tirole, J. (1988). *The Theory of Industrial Organization*. MIT Press.

Tsoy, N. and Konstantinov, N. (2023). Strategic Data Sharing between Competitors. In Oh, A., Naumann, T., Globerson, A., Saenko, K., Hardt, M., and Levine, S., editors, *Advances in Neural Information Processing Systems*, volume 36, pages 16483–16514. Curran Associates, Inc.

Tsybakov, A. B. (2004). Optimal aggregation of classifiers in statistical learning. *The Annals of Statistics*, 32(1):135–166.

Wu, Y., Duan, H., Li, X., and Hu, X. (2025). Navigating the Deployment Dilemma and Innovation Paradox: Open-Source versus Closed-source Models. In *Proceedings of the ACM on Web Conference 2025*, WWW '25, pages 1488–1501, New York, NY, USA. Association for Computing Machinery.

Xie, T., Qi, X., Zeng, Y., Huang, Y., Sehwag, U. M., Huang, K., He, L., Wei, B., Li, D., Sheng, Y., Jia, R., Li, B., Li, K., Chen, D., Henderson, P., and Mittal, P. (2025). SORRY-Bench: Systematically Evaluating Large Language Model Safety Refusal. In *The Thirteenth International Conference on Learning Representations*.

Xu, F., Wang, X., Chen, W., and Xie, K. (2025). The Economics of AI Foundation Models: Openness, Competition, and Governance. arXiv:2510.15200 [econ.TH].

Appendices

A Proofs **14**
A.1 Proof of Lemma 4.1 14
A.2 Derivation of Equation (13) 15
A.3 Derivation of Equation (14) 16
A.4 Proof of Proposition 5.1 16
A.5 Derivation of Equation (15) 17
A.6 Proof of Theorem 5.2 17
A.7 Proof of Proposition 5.3 19
A.8 Proof of Proposition 5.4 20
A.9 Proof of Proposition 6.1 20
A.10 Proof of Proposition 6.2 22

B Alternative Provider Structures **22**

C Alternative Liability and Welfare Structures **23**
C.1 Upstream Liability 23
C.2 Additional Welfare Analysis 24

D Safety-Related Costs Extensions **25**
D.1 Asymmetric Costs 25
D.2 Different Form of Costs 26
D.3 Multi-Objective Costs 26

E Extensions **26**
E.1 Nonlinear Inference Costs 27
E.2 Disruption 28

A Proofs

A.1 Proof of Lemma 4.1

Demand Assuming that restrictions $q_i \geq 0$ and $p_i \geq 0$ will hold in the optimum, consumer problem results in the following Lagrange function:

$$L = \sum_{i=1}^{m} q_i - \Big(\frac{\sum_i q_i^2 + \gamma \sum_{i \neq j} q_i q_j}{2} \Big) + g + \nu (G - \sum_{i=1}^{m} p_i q_i - g).$$

Karush-Kuhn-Tucker conditions give

$$\begin{cases} 0 = \frac{\partial L}{\partial q_i} = 1 - q_i - \gamma \sum_{j \neq i} q_j - \nu p_i, \\ 0 = \frac{\partial L}{\partial g} = 1 - \nu, \\ 0 = \nu (G - \sum_{i=1}^{m} p_i q_i - g), \end{cases}$$

where ν is the Lagrange multiplier. This system results in the following equations

$$p_i = 1 - q_i - \gamma \sum_{j \neq i} q_j. \tag{16}$$

Notice that $p_i \le 1$ and $q_i \le 1$, hence, $\sum_i p_i q_i \le m$. Thus, for $G \ge m$, the budget constraint holds. Since function $u(\boldsymbol{q}, g)$ is strongly concave in $\boldsymbol{q}$ and linear in g, this solution is a global maximizer.

Supply Given that price will be determined by Equation (16), the firms get the following production surplus

$$M(p_i q_i - c_i^{\text{eff}} q_i) = M(1 - c_i^{\text{eff}} - q_i - \gamma \sum_{j \ne i} q_j) q_i.$$

Thus, the best response of the firms follow

$$0 = \frac{\partial \Pi_i^m}{\partial q_i} = 1 - c_i^{\text{eff}} - \gamma \sum_{j \ne i} q_j - 2q_i.$$

In Nash equilibrium all actions should follow the best response, resulting in system

$$q_i^* = \frac{1 - c_i^{\text{eff}} - \gamma \sum_{j \ne i} q_j^*}{2}.$$

Therefore,

$$\sum_i q_i^* = \frac{m - \sum_i c_i^{\text{eff}} - \gamma(m-1) \sum_i q_i^*}{2} \implies \sum_i q_i^* = \frac{m - \sum_i c_i^{\text{eff}}}{2 + \gamma(m-1)}.$$

By substituting the sum back,

$$q_i^* = \frac{2 - \gamma - (2 + \gamma(m-2)) c_i^{\text{eff}} + \gamma \sum_{j \ne i} c_j^{\text{eff}}}{(2-\gamma)(2+\gamma(m-1))}.$$

Notice that $\Pi_i^m = M(bq_i - q_i^2)$ for token-based fees, resulting in $\Pi_i^{m,*} = M(q_i^*)^2$. Expression for the one-time fees can be derived similarly.

Non-negativity Finally, we need to verify that, in equilibrium, we get $q_i^* \ge 0$ and $p_i^* \ge 0$. For q_i^*, we get

$$q_i^* \ge 0 \iff 2 - \gamma - (2 + \gamma(m-2)) c_i^{\text{eff}} + \gamma \sum_{j \ne i} c_j^{\text{eff}} \ge 0 \iff (2-\gamma)(1 - c_i^{\text{eff}}) \ge \gamma \sum_{j \ne i} (c_i^{\text{eff}} - c_j^{\text{eff}}).$$

For p_i^*, notice that the best response quantity for the firm is positive only if $p_i \ge 0$. However, we showed that $q_i^* \ge 0$ under the restriction on costs, hence, $p_i^* \ge 0$ under the same restriction.

A.2 Derivation of Equation (13)

The main idea of this approximation is to use a first-order Taylor-expression for all formulas around the point $\boldsymbol{c} = (a + w, a + w, \ldots, a + w)$.

First, notice that the existence of fee vector implies that $\delta \ge 0$, otherwise the introduction of the technology is not profitable for either a firm or the provider.

Now, we consider the case of per-token fees. In this case, any feasible fee vector will satisfy $\forall i \; t_i \in [w, \delta + w]$. Given this restriction, we notice that the effective marginal production costs are close to the point of optimal costs $a + w$. In the case of per-unit fees, we get

$$0 \le c_i^{\text{eff}} - a - w \le \max\Big(\delta, \frac{b}{(n_i^{\text{eff}})^\beta} + t_i - w\Big) \le \max(\delta, b + t_i - w) \le b + \delta.$$

Thus, we can use the first order approximation of profits

$$\Pi_i^{m,*}(\boldsymbol{c}^{\text{eff}}) \approx M(q^*)^2 - \frac{2(2 + \gamma(m-2)) M q^* (c_i^{\text{eff}} - a - w)}{(2-\gamma)(2+\gamma(m-1))} + \frac{2\gamma M q^* \big(\sum_{j \ne i} (c_j^{\text{eff}} - a - w)\big)}{(2-\gamma)(2+\gamma(m-1))}, \tag{17}$$

where $q^* := \frac{1-a-w}{2+\gamma(m-1)}$.

Similarly, in the case of one-time fees, we get

$$0 \leq c_i^{\text{eff}} - a - w \leq \max\Big(\delta, \frac{b}{(n_i^{\text{eff}})^\beta}\Big) \leq \max(\delta, b) \leq b + \delta.$$

And we get the following profits

$$\Pi_i^{m,*}(\boldsymbol{c}^{\text{eff}}) \approx M(q^*)^2 - \frac{2(2+\gamma(m-2))Mq^*(c_i^{\text{eff}} - a - w)}{(2-\gamma)(2+\gamma(m-1))} + \frac{2\gamma Mq^*\big(\sum_{j\neq i}(c_j^{\text{eff}} - a - w)\big)}{(2-\gamma)(2+\gamma(m-1))} - t_i.$$

In both cases, Equation (3) results in the following problem:

$$\max_{k_i} -\frac{Ab}{(k_i + n^\rho)^\beta} - \lambda k_i,$$

where $A \coloneqq \frac{2(2+\gamma(m-2))Mq^*}{(2-\gamma)(2+\gamma(m-1))}$.

A.3 Derivation of Equation (14)

For token-based fees, the market surplus of LLM introduction follows

$$\Pi_i^{m,*}(\boldsymbol{c}^{\text{eff}}) - \Pi_i^{m,*}(\boldsymbol{c}_{-i}^{\text{eff}}(a_0)) \approx A(a_0 - a - w) - A(c_i^{\text{eff}} - a - w) = A\Big(\delta - \frac{b}{(n_i^{\text{eff}})^\beta} + w - t_i\Big).$$

As for the profits of the provider, note

$$\text{AccessFee}_i(q_i, t_i) - Mq_i w = Mq_i(t_i - w) = Mq^*(t_i - w) + M(q_i - q^*)(t_i - w) \approx Mq^*(t_i - w).$$

Thus,

$$\Pi_p^* - \Pi_{p,-i}^* \approx Mq^*(t_i - w).$$

Putting two approximations together, we get the following approximation of Equations (4) and (3):

$$\arg\max_{t_i}\Big(A\delta - \frac{Ab}{(k_i + n^\rho)^\beta} - \lambda k_i - A(t_i - w)\Big)^{1-\phi}(t_i - w)^\phi,\ \arg\max_{k_i} -\frac{Ab}{(k_i + n^\rho)^\beta} - \lambda k_i.$$

For subscription-based fees, we have

$$\text{AccessFee}_i(q_i, t_i) - Mq_i w = t_i.$$

Thus, we get the following approximation of Equations (4) and (3):

$$\arg\max_{t_i}\Big(A\delta - \frac{Ab}{(k_i + n^\rho)^\beta} - \lambda k_i - t_i\Big)^{1-\phi} t_i^\phi,\ \arg\max_{k_i} -\frac{Ab}{(k_i + n^\rho)^\beta} - \lambda k_i.$$

We can get an equivalent problem by replacing $t_i \mapsto Mq^*(t_i - w)$

$$\arg\max_{t_i}\Big(A\delta - \frac{Ab}{(k_i + n^\rho)^\beta} - \lambda k_i - Mq^*(t_i - w)\Big)^{1-\phi}(t_i - w)^\phi,\ \arg\max_{k_i} -\frac{Ab}{(k_i + n^\rho)^\beta} - \lambda k_i.$$

Since $Mq_i w \approx Mq^* w$, this change of variables gives

$$\text{AccessFee}_i \approx Mq^* w + t_i \mapsto Mq^* t_i.$$

A.4 Proof of Proposition 5.1

We get the following Lagrange function

$$L = -\frac{Ab}{(k_i + n^\rho)^\beta} - \lambda k_i + \nu k_i,$$

where ν is the Lagrange multiplier. The first order conditions give

$$\begin{cases} 0 = \frac{\partial L}{\partial k_i} = \frac{\beta Ab}{(k_i + n^\rho)^{\beta+1}} - \lambda + \nu, \\ 0 = \nu k_i, \\ k_i \geq 0,\ \nu \geq 0. \end{cases}$$

In the case $k_i = 0$, we get

$$\frac{\beta A b}{n^{\rho(\beta+1)}} - \lambda = -\nu \leq 0.$$

Otherwise,

$$\frac{\beta A b}{(k_i + n^\rho)^{\beta+1}} - \lambda = 0 \implies \frac{\beta A b}{n^{\rho(\beta+1)}} - \lambda \geq 0.$$

Thus, we get

$$k_i^* = k^* = \max(0, k_m - n^\rho),\ k_m \coloneqq \Big(\frac{\beta A b}{\lambda}\Big)^{1/(1+\beta)}.$$

For the cumulative level, we get

$$k_i^* + n^\rho = \max(k_m, n^\rho).$$

Both k_m and n are non decreasing in n, hence, the cumulative level is non-decreasing in n. For $n^\rho \leq k_m$, the cumulative level equals k_m that does not depend on n.

A.5 Derivation of Equation (15)

Theorem 5.1 gives the following expression for the surplus of the firm:

$$\Pi_i^{m,*} - \lambda k_i^* - \Pi_{i,-i}^{m,*} = A\delta - \frac{Ab}{\max(k_m, n^\rho)^\beta} - \lambda \max(k_m - n^\rho, 0) - B(t_i - w) =: S - B(t_i - w).$$

Then the Nash bargaining concept reduces to

$$\max_{t_i} (S - B(t_i - w))^{1-\phi} (t_i - w)^\phi.$$

The condition for the existence of feasible fees and the LLM's profitability for the firms is

$$S \geq 0 \iff A\delta - \lambda k_m - \frac{Ab}{\max(n^\rho, k_m)^\beta} + \lambda \min(n^\rho, k_m) \geq 0.$$

If feasibility condition holds, the first-order condition is

$$-\frac{(1-\phi)B}{S - B(t_i - w)} + \frac{\phi}{t_i - w} = 0,$$

which gives the following equation for t_i:

$$t_i = \frac{\phi S}{B} + w.$$

This results in the following profits of the provider

$$\begin{aligned}\Pi_p^* &= \sum_{i=1}^m M q^* t_i - m M q^* w - \lambda_p (n-1) \\ &= \frac{m\phi M q^*}{B}\Big(A\delta - \lambda k_m - \frac{Ab}{\max(n^\rho, k_m)^\beta} + \lambda \min(n^\rho, k_m)\Big) - \lambda_p n + \lambda_p.\end{aligned}$$

Note that non-negativity of profits $\Pi_p^* \geq 0$ implies the feasibility condition (since $\lambda_p(n-1) \geq 0$), making the feasibility condition redundant.

A.6 Proof of Theorem 5.2

Maximization of the provider's profits is equivalent to

$$\max_{n \geq 1} f(n) \coloneqq \lambda \min(n^\rho, k_m) - \frac{Ab}{\max(n^\rho, k_m)^\beta} - \frac{\lambda_p B n}{m\phi M q^*}.$$

First, we show that this function is differentiable. Notice that

$$\Big(-\frac{Ab}{n^{\rho\beta}}\Big)' = \frac{\rho\beta A b}{n^{\rho\beta+1}},\ (\lambda n^\rho)' = \rho\lambda n^{\rho-1}.$$

At the point $n^\rho = k_m$, we get

$$\frac{\rho\beta Ab}{n^{\rho\beta+1}} = \frac{\rho\beta Ab}{k_m^{\beta+1}} k_m n^{-1} = \rho\lambda k_m n^{-1} = \rho\lambda n^{\rho-1}.$$

Thus, the left and right derivatives agree at $n^\rho = k_m$ and the function is differentiable at this point. For other points, the function is differentiable as a power function. Moreover,

$$f' = \rho \min\Big(\frac{\beta Ab}{n^{\rho\beta+1}}, \lambda n^{\rho-1}\Big) - \frac{\lambda_p B}{m\phi M q^*},$$

and this derivative is decreasing in n.

Now, the necessary conditions for extremum give

$$\begin{cases} 0 = \rho \min\Big(\frac{\beta Ab}{n^{\rho\beta+1}}, \lambda n^{\rho-1}\Big) - \frac{\lambda_p B}{m\phi M q^*} + \nu, \\ 0 = \nu(n-1), \\ \nu \geq 0, n \geq 1. \end{cases}$$

If $n = 1$, the following should hold:

$$\min(\beta Ab, \lambda) \leq \frac{\lambda_p B}{\rho m \phi M q^*}.$$

Otherwise, $n > 1$ and we have

$$\min\Big(\frac{\beta Ab}{n^{\rho\beta+1}}, \lambda n^{\rho-1}\Big) = \frac{\lambda_p B}{\rho m \phi M q^*}.$$

By solving this equation, we get

$$n^* = \begin{cases} n_s, & n_s^\rho \leq k_m \\ n_b, & n_b^\rho \geq k_m \end{cases} = \min(n_s, n_b).$$

where $n_s \coloneqq \big(\frac{\rho\phi m \lambda M q^*}{\lambda_p B}\big)^{1/(1-\rho)}$, $n_b \coloneqq \big(\frac{\rho\phi m \beta A M q^* b}{\lambda_p B}\big)^{1/(1+\rho\beta)}$.

Monotonicity properties Now, we need to check the properties of the considered values. Most of the properties of n_s are obvious, except for the fact that n_s is increasing in m in the case of per-unit fees. To show this, we note

$$n_s^{1-\rho} \propto \frac{m(2+\gamma(m-1))}{2+\gamma(m-2)} = m + \frac{\gamma}{\frac{2-2\gamma}{m} + \gamma}.$$

Here, clearly both multipliers are increasing in m.

Most of the properties of n_b are also obvious, except for the dependence on m and γ. In the case of one-time fees, we get

$$n_b^{1+\rho\beta} \propto \frac{m(2+\gamma(m-2))}{(2+\gamma(m-1))^2}.$$

The derivative of this expression is

$$\frac{2-2\gamma+2\gamma m}{(2+\gamma(m-1))^2} - \frac{2m(2+\gamma(m-2))\gamma}{(2+\gamma(m-1))^3} = \frac{2(2+\gamma(m-3)+\gamma^2)}{(2+\gamma(m-1))^3} \geq \frac{2(1-\gamma+\gamma m)}{(2+\gamma(m-1))^3} > 0,$$

proving the monotonicity. In the case of per-unit fees, we get

$$n_b^{1+\rho\beta} \propto \frac{m}{(2+\gamma(m-1))} = \frac{1}{\frac{2-\gamma}{m} + \gamma},$$

which is monotonically increasing in m. For the dependence on γ, we note the it is mediated by A and calculate the following derivative:

$$\frac{\partial A}{\partial \gamma} = \frac{4(m-1)(2+\gamma(m-5)-\gamma^2(m-2))Mq^*}{(2-\gamma)^2(2+\gamma(m-1))^2}.$$

The sign of this expression is equal to the sign of $2 + \gamma(m-5) - \gamma^2(m-2)$. The roots of this expression are

$$\gamma = \frac{m - 5 \pm \sqrt{m^2 - 2m + 9}}{2(m-2)}.$$

The smallest root is less than $-1/(m-1)$ since

$$\frac{m - 5 - \sqrt{m^2 - 2m + 9}}{2(m-2)} \leq -\frac{1}{m-1} \iff m - 3 - \frac{2}{m-1} \leq \sqrt{m^2 - 2m + 9},$$

but $m - 3 \leq \sqrt{m^2 - 2m + 9}$. Thus, this expression is positive for small γ and negative for big γ.

Finally, for k_m, most of the properties are again obvious, except the dependence on m and γ. The dependence on γ is the same as for n_b. For m, we have

$$k_m^{1+\beta} \propto \frac{2 + \gamma(m-2)}{(2 + \gamma(m-1))^2}.$$

The derivative of this expression is

$$\frac{\gamma}{(2 + \gamma(m-1))^2} - \frac{2\gamma(2 + \gamma(m-2))}{(2 + \gamma(m-1))^3} = -\frac{\gamma(2 + \gamma m - 3\gamma)}{(2 + \gamma(m-1))^3} < 0,$$

proving the necessary result.

A.7 Proof of Proposition 5.3

The proof of this proposition generally mimics the proof of Theorem 5.1, given that the cost sensitivity of a new objective change.

Industry profits The sum of all profits results in the following expression

$$\begin{aligned}
&\sum_{i=1}^{m} \Pi_i + \Pi_p \\
&= \sum_i \text{Sales}_i - \text{ProductionCosts}_i - \text{ReputationalCosts}_i - \text{FirmSafetyInvestments}_i \\
&\quad - \text{InferenceCosts} - \text{ProviderSafetyInvestments} \\
&\approx \sum_i \Big(M(q^*)^2 - A^{\text{ind}}(c_i - a) - \lambda k_i \Big) - \lambda_p (n-1),
\end{aligned}$$

where we used the fact that InferenceCosts are equivalent to the per-token fee $t_i = w$ and approximation 17.

Market welfare The consumer utility follows

$$u^* = \sum_{i=1}^{m} q_i^* - \frac{\sum_i (q_i^*)^2 + \gamma \sum_{i \neq j} q_i^* q_j^*}{2} - \sum_{i=1}^{m} p_i^* q_i^* + G.$$

So, in the equilibrium, the total market welfare follows

$$W^m = Mu^* + \Pi^{\text{ind}} = M\Big(\sum_{i=1}^{m} (1 - c_i - w) q_i^* - \frac{\sum_i (q_i^*)^2 + \gamma \sum_{i \neq j} q_i^* q_j^*}{2} + G \Big) - \sum_i \lambda k_i - \lambda_p (n-1).$$

So, the first-order approximation gives

$$\begin{aligned}
W^m &\approx const + \sum_i M((a - c_i)q^* + (1 - a - w)(q_i^* - q^*) - (1 + \gamma(m-1))q^*(q_i^* - q^*)) \\
&\quad - \sum_i \lambda k_i - \lambda_p n \\
&= const + \sum_i \Big(-A^{\text{mw}}(c_i - a) - \lambda k_i \Big) - \lambda_p (n-1).
\end{aligned}$$

Private investments Similarly to Theorem 5.1, we get

$$k_i^x = \max(0, k_m^x - n^\rho).$$

A.8 Proof of Proposition 5.4

After substituting the response $k_i(n)$, the optimal investment problem for the provider reduces to

$$\max_n m\Big(\lambda \min(n^\rho, k_0) - \frac{A^x b}{\max(n^\rho, k_0)^\beta}\Big) - \lambda_p n.$$

This equation gives the same problem as Equation (15) for $\phi \mapsto B/Mq^*$ and $A \mapsto A^x$. Then, the proof repeats the proof of Proposition 6.1 (see below). Crucially, the condition $k_0 \geq k_m^x$ ensures that the subgradient of the objective above is non-increasing in n. This property allows us to reduce the number of cases that arises due to the non-differentiability of the objective.

A.9 Proof of Proposition 6.1

The case $k_0 \leq 1$ does not introduce new constraint since $n \geq 1$.

If $k_0 \leq k_m$, the response of downstream firms will not change since their objective is the same. Thus, the objective and response of the provider also will not change.

Now, we consider the case $k_0 > \max(k_m, 1)$.

Downstream firms Similarly to Theorem 5.1, we have $k = \max(0, k_0 - n^\rho)$, if

$$\frac{\beta A b}{\max(k_0^{\beta+1}, n^{\rho(\beta+1)})} - \lambda \leq 0.$$

However, since $k_0 > k_m$, we get the following condition on the derivative of downstream objective

$$\frac{\beta A b}{k_0^{\beta+1}} - \lambda < 0,$$

which implies the condition above and gives $k^* = \max(0, k_0 - n^\rho)$.

Provider's objective The solution for k^* results in the following provider's objective

$$\max_{n \geq 1} f(n) \coloneqq \lambda \min(n^\rho, k_0) - \frac{Ab}{\max(n^\rho, k_0)^\beta} - \frac{\lambda_p B n}{m\phi M q^*}.$$

While this function is not differentiable, its subgradient is decreasing in n in the case $k_0 > k_m$. Thus, the first order conditions would give the same answer as in Theorem 5.2.

First case, $n \leq k_0^{1/\rho}$ We get the following Lagrange function

$$L(n) = \lambda n^\rho - \frac{Ab}{k_0^\beta} - \frac{\lambda_p B n}{m\phi M q^*} + \nu(n-1) + \alpha(k_0^{1/\rho} - n).$$

The necessary conditions for extremum give

$$\begin{cases} 0 = \rho\lambda n^{\rho-1} - \frac{\lambda_p B}{m\phi M q^*} + \nu - \alpha, \\ 0 = \nu(n-1), \\ 0 = \alpha(k_0^{1/\rho} - n), \\ \nu \geq 0, \alpha \geq 0, n \geq 1, n \leq k_0^{1/\rho}. \end{cases}$$

Hence, $n = 1$ if

$$0 \geq \rho\lambda - \frac{\lambda_p B}{m\phi M q^*}.$$

$n = k_0^{1/\rho}$ if

$$0 \leq \rho\lambda k_0^{(\rho-1)/\rho} - \frac{\lambda_p B}{m\phi M q^*}.$$

Otherwise,

$$0 = \rho\lambda n^{\rho-1} - \frac{\lambda_p B}{m\phi M q^*} \implies n = n_s.$$

Second case, $n \geq k_0^{1/\rho}$ Similarly to the first case, $n = \max(1, k_0^{1/\rho})$ if

$$0 \geq \frac{\rho\beta Ab}{\max(k_0^{(\rho\beta+1)/\rho}, 1)} - \frac{\lambda_p B}{m\phi M q^*}.$$

Otherwise,

$$0 = \frac{\rho\beta Ab}{n^{\rho\beta+1}} - \frac{\lambda_p B}{m\phi M q^*} \implies n = n_b.$$

Global solution Now, we need to connect these cases together. For the case $n \leq k_0^{1/\rho}$, we have 3 possible values of n: 1, $k_0^{1/\rho}$, and n_s. For the case $n \geq k_0^{1/\rho}$, we have have 2 possible values: $k_0^{1/\rho}$ and n_s.

If $n = n_b$ in the second case, we should have $n = k_0^{1/\rho}$ in the first case, which implies that the optimal value of n is n_b (since $\forall n \in [1, k_0^{1/\rho}]\ f(n) \leq f(k_0^{1/\rho})$, but $\forall n \geq k_0^{1/\rho}\ f(n) \leq f(n_b)$, which implies $\forall n \geq 1\ f(n) \leq f(n_b)$). To establish this property, notice that $n = n_b$ in the second case implies

$$0 < \frac{\rho\beta Ab}{k_0^{(\rho\beta+1)/\rho}} - \frac{\lambda_p B}{m\phi M q^*} \leq \rho\lambda k_0^{(\rho-1)/\rho} - \frac{\lambda_p B}{m\phi M q^*},$$

giving $n = k_0^{1/\rho}$ in the first case. Note that the last inequality holds since $k_m \leq k_0$.

If $n < k_0^{1/\rho}$ in the first case, we should have $n = k_0^{1/\rho}$ in the second case. To establish this, notice that $n < k_0^{1/\rho}$ in the first case implies

$$0 > \rho\lambda k_0^{(\rho-1)/\rho} - \frac{\lambda_p B}{m\phi M q^*} \geq \frac{\rho\beta Ab}{k_0^{(\rho\beta+1)/\rho}} - \frac{\lambda_p B}{m\phi M q^*},$$

giving $n = k_0^{1/\rho}$ in the second case.

Finally, if $n = k_0^{1/\rho}$ in both cases, we get

$$\rho\lambda k_0^{(\rho-1)/\rho} \geq \frac{\lambda_p B}{m\phi M q^*} \geq \frac{\rho\beta Ab}{k_0^{(\rho\beta+1)/\rho}}.$$

So, we get

$$n^* = \begin{cases} 1, & \frac{\lambda_p B}{m\phi M q^*} \geq \rho\lambda, \\ n_s, & \rho\lambda \geq \frac{\lambda_p B}{m\phi M q^*} \geq \rho\lambda k_0^{(\rho-1)/\rho}, \\ k_0^{1/\rho}, & \rho\lambda k_0^{(\rho-1)/\rho} \geq \frac{\lambda_p B}{m\phi M q^*} \geq \frac{\rho\beta Ab}{k_0^{(\rho\beta+1)/\rho}}, \\ n_b, & \frac{\rho\beta Ab}{k_0^{(\rho\beta+1)/\rho}} \geq \frac{\lambda_p B}{m\phi M q^*}, \end{cases} = \max(1, n_e),$$

where

$$n_e = \begin{cases} n_s, & n_s \leq k_0^{1/\rho}, \\ k_0^{1/\rho}, & n_s \geq k_0^{1/\rho} \geq n_b, \\ n_b, & k_0^{1/\rho} \leq n_b. \end{cases}$$

Comparison with $n^{*,\textbf{old}}$ At the same time, we have

$$n^{*,\text{old}} = \max(1, n_e^{\text{old}}), \text{ where } n_e^{\text{old}} = \begin{cases} n_s, & n_s \leq k_m^{1/\rho}, \\ n_b, & k_m^{1/\rho} \leq n_b. \end{cases}$$

To show that $n^* \geq n^{*,\text{old}}$, we need to show that $n_s \geq n_b$ if $k_m^{1/\rho} \leq n_s \leq k_0^{1/\rho}$ and that $k_0^{1/\rho} \geq n_b$ if $k_m^{1/\rho} \leq n_b \leq k_0^{1/\rho}$. The second inequality is obvious. The first inequality follows from the following:

$$n_b = (n_s^{1-\rho} k_m^{1+\beta})^{1/(1+\rho\beta)} = n_s^{(1-\rho)/(1+\rho\beta)} (k_m^{1/\rho})^{(\rho+\rho\beta)/(1+\rho\beta)}$$

Thus, $n_s \geq n_b \iff n_s \geq k_m^{1/\rho}$, which holds in our case.

A.10 Proof of Proposition 6.2

The action space of downstream firms would not change given the regulation. Thus, their responses would not change.

The final objective of the provider would not change. Hence the derivative of the objective will still decrease in n. Hence, similarly to the proof of Theorem 5.2, we get that the provider will pick $n = n_0$ if

$$\min\left(\frac{\beta Ab}{n_0^{\rho\beta+1}}, \lambda n_0^{\rho-1}\right) \leq \frac{\lambda_p B}{\rho m \phi M q^*}.$$

Otherwise, the provider will pick n_s or n_b similarly to Theorem 5.2.

B Alternative Provider Structures

This section discusses how to extend our analysis to a market with several foundation model providers, a market where another provider could enter, or a market where firms can use open-weight models.

Many providers We suggest the following approach for an extension to a market with several model providers.

First, we add an additional stage that happens before everything else in which model providers match with downstream firms. Conceptually, during this stage, the model providers pitch their model to the downstream firms and make preliminary agreements with them. A simplest extension of this kind is a random assignment (e.g., the sales department of each firm randomly pass or fail during the negotiations with each downstream firm); more sophisticated versions might assume some model of hidden characteristics of firms and model providers.

Second, we replace the Nash-in-Nash solution concept with the Nash-in-Nash with Threat of Replacement solution concept (Ho and Lee, 2019). This new solution concept changes firms' outside options. Previously, the outside option was to stop using AI completely. Now, the firm has two outside options: it can either stop using AI or start using a model of a different provider. The fees to the alternative provider should be set so the alternative provider is no worse off when it starts serving the new firm; i.e., the firm's fee under the alternative deal should equal its "inference costs".

Third, the extension should also add a factor that makes switching between providers difficult. Without such friction, providers will compete on price until each provider receives zero surplus from serving clients and exits the market. In other words, without such friction, the situation reduces to the previous case with one model provider. Such switching frictions can arise from different factors (Tirole, 1988). In the first stage, firms might make model-specific infrastructure investments that make switching costly. Alternatively, each model provider's inference capacity may be bounded and, hence, inference costs for a potential additional client may skyrocket, locking them into their original model provider.

Importantly, under this extension, only each firm's outside option changes. Hence, using the previous notation, only the parameter δ changes. Since δ does not affect the equilibrium and optimal investment levels, *the equilibrium and optimal investment levels will stay the same in the new model*; only the fees to the provider will decrease. This property suggests that the current analysis should not qualitatively depend on the number of upstream model providers.

Threat of entry A similar effect occurs due to the threat of entry. Assume that a potential alternative provider could enter the market. The potential offer from the entrant will again constitute a better outside option for the downstream firms, leading effectively to lower δ.

Specifically, a potential competing provider will enter the market and offer its services if the firm offers a fee that covers its inference and data-collection costs. Formally, we assume that the entrant announces that it will collect n_e points if it enters the market. Then, we get the following surplus for each firm (i.e., the left-hand part of Equation (14)) that corresponds to the new outside option:

$$\min_{k_e} \frac{Ab}{(k_e + n_e^\rho)^\beta} + \lambda k_e + \lambda_p(n_e - 1) - \frac{Ab}{(k_i + n^\rho)^\beta} - \lambda k_i - B(t_i - w),$$

where k_e is the potential number of points a firm collects under the alternative deal.

In this case, only a positive externality from the provider's investments allows for non-zero provider profits. To see this, consider $n_e = n^*$; then, for subscription-based fees, we get the following provider's profits:

$$\Pi_p^* = \frac{m\phi M q^*}{B}\lambda_p(n^* - 1) - \lambda_p(n^* - 1).$$

These profits are positive only if $m\phi M q^*/B > 1$, i.e., when m is big enough.

Open-weight models The presence of open-weight models will also only change firms' outside options and hence will not affect investment levels. Formally, if we assume that the open-weight model behaves as the provider's model, we get the following surplus for each firm:

$$\min_{k_e} \frac{Ab}{(k_e + 1)^\beta} + \lambda k_e - \frac{Ab}{(k_i + n^\rho)^\beta} - \lambda k_i - B(t_i - w),$$

where k_e is the potential number of points collected for the open-weight model. As in Theorem 5.1, the equilibrium investment level in the open-weight case is $k_e^* = \max(k_m - 1, 0)$. This level is generally lower than before because the provider's investments provide no positive externality.

C Alternative Liability and Welfare Structures

This section discusses how to extend our analysis to cases in which the provider is liable for safety-related accidents and how to incorporate consumer harm from such accidents into market welfare.

C.1 Upstream Liability

This section discusses how our analysis changes if we assume the upstream provider is liable for safety-related accidents. For simplicity, we only consider the case of token-based fees.

Structurally, we want to introduce a safety-related cost term related to Equation (1). Using the same linearity assumption as for the term $\text{SafetyRelatedCosts}_i$, we introduce safety-related costs of the provider as

$$\text{SafetyRelatedCosts}_p = \sum_{i=1}^{m} Mc_{p,i}q_i. \tag{18}$$

Here, $Mc_{p,i}q_i$ represents the safety-related costs arising from accidents at firm i; $c_{p,i}$ is the corresponding marginal cost; and Mq_i is the equilibrium quantity produced by firm i. To model marginal costs, we use the same form as in Equation (10), resulting in

$$c_{p,i} = a_p + \frac{b_p}{(k_i + n^\rho)^\beta}, \tag{19}$$

where a_p and b_p are two new constants that describe the scaling of $c_{p,i}$.

Now, Equation (4) in the case of token-based fees require $t_i \in [w + c_{p,i}, a_0 - c_i]$. To study this regime, we reformulate the absence-of-disruption assumption as

$$(2 - \gamma)(1 - a - a_p - w) \gg (2 + \gamma(m - 2))(b + b_p + \delta'), \text{ where } \delta' \coloneqq a_0 - a - a_p - w. \tag{20}$$

Consequently, we approximate the costs with a Taylor expansion around $\boldsymbol{c} = (a + a_p + w, \ldots, a + a_p + w)$ rather than $(a + w, \ldots, a + w)$.

In this setting, Equation (13) will not change significantly, but will use $q'^* \coloneqq \frac{1-a-a_p-w}{2+\gamma(m-1)}$ instead of q^* and $b + b_p$ instead of b. As for Equation (4), we first note that the provider surplus becomes

$$\text{AccessFee}_i - Mq_i w - Mq_i c_{p,i} = Mq_i(t_i - w - c_{p,i}) \approx Mq'^*(t_i - w - c_{p,i}).$$

(Note that $|t_i - w - c_{p,i}| \leq \delta' + b + b_p$, which is small.) Hence,

$$\Pi_p^* - \Pi_{p,-i}^* \approx Mq'^*(t_i - w - c_{p,i}).$$

Then, we will have the following bargaining problem

$$\arg\max_{t_i} \Big(A'\delta' - \frac{A'(b + b_p)}{(k_i + n^\rho)^\beta} - \lambda k_i - A'(t_i - w - c_{p,i}) \Big)^{1-\phi} (t_i - w - c_{p,i})^\phi,$$

where $A' \coloneqq \frac{2(2+\gamma(m-2))Mq'^*}{(2-\gamma)(2+\gamma(m-1))}$. By shifting $t_i \mapsto t_i + c_{p,i}$, we would get the same problem as before, but with new variables $b \mapsto b + b_p, \delta \mapsto \delta', q^* \mapsto q'^*, A \mapsto A'$.

Comparison with the baseline It is interesting to compare the baseline case to the case where one introduces liability to the model provider. In all formulas, introducing upstream liabilities will be reflected through the product Ab. Previously, we had $Ab \propto (1-a-w)b$; now, we have $A'(b+b_p) \propto (1-a-a_p-w)(b+b_p)$. Here we can see two effects. Directly, additional liabilities increase safety investments through $b+b_p$. Indirectly, the additional liabilities reduce safety investments through $(1-a-a_p-w)$. The second effect occurs because additional liabilities reduce the overall profitability of all participants and therefore limit their ability to make safety investments. Which effect dominates depends on the nature of safety-related costs. If costs are well-optimizable (i.e., a_p is small), the first effect dominates; otherwise, the second effect dominates.

At the same time, if we assume that the regulator simply splits the liabilities instead of introducing new ones, i.e., $a = a' + a'_p$ and $b = b' + b'_p$. Then the equilibrium levels of investment and profits for all participants will not change; only the fees to the model provider will increase. This happens because the model provider might easily pass its liability to the downstream firms through higher access fees.

C.2 Additional Welfare Analysis

This section discusses how to incorporate consumer harm from such accidents into market welfare.

Similar to safety-related costs, we assume that the direct harm to consumers is proportional to production scale

$$CH = \sum_{i=1}^{m} Mq_i c_{c,i}, \tag{21}$$

where CH is the direct consumer harm in monetary terms (to avoid double counting, this term only includes the harm not covered by fines paid by firms), Mq_i is the production scale, and $c_{c,i}$ is the marginal increase in harm with the production scale. Use the same logic as for Equation (10), we assume that

$$c_{c,i} = a_c + \frac{b_c}{(k_i + n^\rho)^\beta}. \tag{22}$$

Assuming $b_c/(k_i+n^\rho)^\beta$ is small, the first-order approximation of consumer harm gives

$$\begin{aligned}
CH \approx & mMq^* a_c + \sum_i Mq^*(c_{c,i} - a_c) + M(q_i - q^*)a_c \\
= & mMq^* a_c + \sum_i Mq^*(c_{c,i} - a_c) - \frac{Ma_c(c_i - a)}{2 + \gamma(m-1)} \\
= & mMq^* a_c + \sum_i \frac{Mq^* b_c}{b}(c_i - a) - \frac{Mq^* a_c}{1-a-w}(c_i - a)
\end{aligned}$$

Now, we define the (complete) welfare function as

$$W := W^m - CH \approx const - \sum_i \Big(\Big(A^{\text{mw}} + \frac{Mq^* b_c}{b} - \frac{Mq^* a_c}{1-a-w} \Big)(c_i - a) + \lambda k_i \Big) - \lambda_p(n-1). \tag{23}$$

Define

$$A^{\text{w}} := A^{\text{mw}} + \frac{Mq^* b_c}{b} - \frac{Mq^* a_c}{1-a-w}. \tag{24}$$

Then, we get

$$W = const - \sum_i \Big(A^{\text{w}}(c_i - a) + \lambda k_i \Big) - \lambda_p(n-1).$$

The welfare function has the same form as the market welfare function, but with A^{w} instead of A^{mw}. Thus, the optimal investment levels will also have the same form as before.

Comparison with the baseline Interestingly, accounting for consumer harm does not necessarily lead to higher optimal investments. Specifically, we have two effects. A direct effect, expressed through the first additional term, suggests increasing safety investments to reduce consumer harm. An indirect effect, expressed through the second additional term, suggests reducing safety investments because this action lowers quantities produced and hence reduces the number of safety-related accidents. Which effect dominates depends on the nature of safety-related costs. If costs are well-optimizable (i.e., a_c is small), the first effect dominates; otherwise, the second effect dominates.

D Safety-Related Costs Extensions

This section considers different relaxations of the parametric form of safety-related costs considered in the main text.

D.1 Asymmetric Costs

This section considers relaxation of cost symmetry assumption.

We consider asymmetric cost functions

$$c_i = a_i + \frac{b_i}{(k_i + n^\rho)^\beta} \tag{25}$$

and allow asymmetric initial costs $a_{0,i}$. In this setting, we reformulate the absence-of-disruption assumption as

$$\forall i \, (2-\gamma)(1-a_i-w) + \gamma \sum_j (a_j - a_i) \gg (2+\gamma(m-2))(b_i+\delta_i), \text{ where } \delta_i \coloneqq a_{0,i} - a_i. \tag{26}$$

This assumption will again allow us to use the first-order approximation and get the following asymmetric versions of A and q^*:

$$A_i \coloneqq \frac{2(2+\gamma(m-2))Mq_i^*}{(2-\gamma)(2+\gamma(m-1))}, \; q_i^* \coloneqq \frac{(2-\gamma)(1-w) - (2+\gamma(m-2))a_i + \gamma \sum_{j \neq i} a_j}{(2-\gamma)(2+\gamma(m-1))}. \tag{27}$$

And we get the following equilibrium private investments

$$k_i^* = \max(0, k_{m,i} - n^\rho), \; k_{m,i} \coloneqq (\beta b_i A_i / \lambda)^{1/(1+\beta)}. \tag{28}$$

These results give the following provider's profits:

$$\Pi_p^* = \sum_i \frac{\phi M q_i^*}{B_i} (A_i \delta_i - \lambda k_{m,i} - \frac{A_i b_i}{\max(n^\rho, k_{m,i})^\beta} + \lambda \min(n^\rho, k_{m,i})) - \lambda_p (n-1), \tag{29}$$

where $B_i = Mq_i^*$ for subscription-based fees and $B_i = A_i$ for token-based fees.

The provider's profit function is differentiable, and one can find the provider's investment level by using the first-order conditions. However, this solution is hard to interpret and analyze. To get interpretable expressions, consider two limiting cases: $(n^*)^\rho \leq \min_i k_{m,i}$ and $(n^*)^\rho \geq \max_i k_{m,i}$. These cases correspond to the cases $n^* = \max(1, n_s)$ and $n^* = \max(1, n_b)$ in the main text. In the first case, we get

$$n^* = \max\Big(1, \Big(\frac{\phi M \rho \lambda}{\lambda_p} \sum_i \frac{q_i^*}{B_i}\Big)^{1/(1-\rho)}\Big).$$

This expression resembles n_s. However, instead of symmetric ratios q^*/B, we sum their asymmetric analogs. In the second case, we get

$$n^* = \max\Big(1, \Big(\frac{\phi M \rho \beta}{\lambda_p} \sum_i \frac{b_i A_i q_i^*}{B_i}\Big)^{1/(1+\rho\beta)}\Big).$$

We again see that the result resembles n_b.

D.2 Different Form of Costs

This section relaxes the requirement that the cost follows the parametric form (10). Specifically, we assume that

$$c_i = a + f(k_i + n^\rho), \tag{30}$$

where f is differentiable, monotonically decreasing, and convex. In this setting, the absence-of-disruption assumption becomes

$$(2-\gamma)(1-a-w) \gg (2+\gamma(m-2))(f(1)+\delta). \tag{31}$$

This assumption will again allow us to use the first-order approximation.

We get the following equilibrium private investments:

$$k^* = \min(0, k_m - n^\rho), \text{ where } -f'(k_m) = \frac{\lambda}{A}. \tag{32}$$

Since f is convex, $-f'$ is monotonically decreasing; hence, the behavior of k_m is similar to its behavior in the main text. We get the following provider's profits:

$$\Pi_p^* = \frac{m\phi M q^*}{B}\Big(A\delta - \lambda k_m - Af(\max(n^\rho, k_m)) + \lambda \min(n^\rho, k_m)\Big) - \lambda_p n + \lambda_p.$$

This expression results in the following first-order condition for the equilibrium investments:

$$\min(\lambda\rho n^{\rho-1}, -A\rho n^{\rho-1} f'(n^\rho)) = \frac{\lambda_p B}{m\phi M q^*}. \tag{33}$$

Thus, we get $n^* = \max(1, \min(n_s, n_b))$ as before. Moreover, we get the same expression for n_s, while n_b should satisfy

$$-n_b^{\rho-1} f'(n_b^\rho) = \frac{\lambda_p B}{\rho\phi m A M q^*}. \tag{34}$$

Since $-n^{\rho-1}f'(n^\rho)$ is monotonically decreasing, the behavior of n_b is similar to its behavior in the main text.

D.3 Multi-Objective Costs

This section discusses how to modify the framework to allow optimizing for several properties that follow polynomial scaling laws. Assume we have R properties that influence costs, and that the costs associated with these properties grow linearly with production volume. The results of Sukenik and Lampert (2025) motivate the following structure of marginal cost

$$c_i = a + \sum_{r=1}^{R} \frac{b_r}{(k_{r,i} + n_r^{\rho_r})^{\beta_r}}, \tag{35}$$

where r is the property, b_r, ρ_r, and β_r are scaling constants associated with this property, and $k_{r,i}$ and n_r are firm i's and provider's investments in property r. In this case, we can assume the following version of the absence-of-disruption assumption:

$$(2-\gamma)(1-a-w) \gg (2+\gamma(m-2))(\delta + \sum_r b_r). \tag{36}$$

Since the cost structure is separable across properties and we use a first-order approximation, maximizing profits across investments in all properties reduces to maximizing profits for each property individually. Hence, we get the same results as before for each property individually, and our results will not change qualitatively.

E Extensions

This appendix discusses possible extensions to our framework and theoretical model.

E.1 Nonlinear Inference Costs

This section informally discusses how to incorporate convex nonlinear inference costs into our analysis. In this section, we assume that nonlinearity is small in some sense and clarify what it means below. For simplicity, we focus on token-based fees.

This section again uses the first-order approximation. However, the zero-order approximation of q_i is no longer straightforward because the inference costs are nonlinear. To find this approximation, assume that q^* and t^* are good approximations for q_i and t_i. Then, we should have the following provider's surplus

$$\begin{aligned}
&\Pi_p^* - \Pi_{p,-i}^* \\
&\quad = \sum_j M q_j t_j - \sum_{j\neq i} M q_{j,-i} t_j - C(\sum_j q_j) + C(\sum_{j\neq i} q_j) \\
&\quad \approx m M q^* t^* + \sum_j (M q^*(t_j - t^*) + M(q_j - q^*)t^*) \\
&\qquad - (m-1) M q^* t^* - \sum_{j\neq i} (M q^*(t_j - t^*) + M(q_{j,-i} - q^*)t^*) \\
&\qquad - C(mq^*) - C'(mq^*) \sum_j (q_j - q^*) \\
&\qquad + C((m-1)q^*) + C'((m-1)q^*) \sum_{j\neq i} (q_{j,-i} - q^*)
\end{aligned}$$

Since the first-order expression bounds firms' surplus, the provider's surplus should also be first-order. This logic gives

$$M q^* t^* - C(mq^*) + C((m-1)q^*) = 0 \implies t^* = \frac{C(mq^*) - C((m-1)q^*)}{Mq^*}.$$

Now, the expression for q^* is

$$q^* = \frac{1 - a - t^*}{2 + \gamma(m-1)}.$$

Substituting t^* into this expression gives an equation for q^*.

Now we proceed assuming that q^* and t^* are some fixed expressions. We get the following provider's surplus:

$$\begin{aligned}
\Pi_p^* - \Pi_{p,-i}^* \approx & M q^*(t_i - t^*) + \sum_j M(q_j - q^*)t^* - \sum_{j\neq i} M(q_{j,-i} - q^*)t^* \\
& - C'(mq^*) \sum_j (q_j - q^*) + C'((m-1)q^*) \sum_{j\neq i} (q_{j,-i} - q^*).
\end{aligned}$$

To make this expression tractable, we assume that $Mt^* - C'(mq^*)$ and $C'((m-1)q^*) - Mt^*$ are small. (Even without this assumption, it is possible to proceed with the analysis. However, the resulting expression will differ from the case analyzed in the main text.) These two expressions precisely control the strength of nonlinearity of C on the interval $[(m-1)q^*, mq^*]$. This assumption gives

$$\Pi_p^* - \Pi_{p,-i}^* \approx M q^*(t_i - t^*).$$

Then, we can proceed with our analysis as before.

In summary, we make the following assumptions. First, we assume that C is convex and monotonically increasing. Then, we solve the system

$$\begin{cases} t^* = \frac{C(mq^*) - C((m-1)q^*)}{Mq^*}, \\ q^* = \frac{1-a-t^*}{2+\gamma(m-1)}. \end{cases} \tag{37}$$

Then, we assume that the nonlinearity is small

$$q^* \gg M t^* - C'(mq^*) + C'((m-1)q^*) - Mt^* = C'((m-1)q^*) - C'(mq^*). \tag{38}$$

Finally, we make the following absence-of-disruption assumption:

$$(2-\gamma)(1-a-t^*) \gg (2+\gamma(m-1))(\delta+b), \text{ where } \delta \coloneqq a_0 - a - t^*. \tag{39}$$

Under these assumptions, we get the following approximation for Equation (4):

$$\arg\max_{t_i} \Big(A\delta - \frac{Ab}{(k_i+n^\rho)^\beta} - \lambda k_i - A(t_i - t^*) \Big)^{1-\phi} (t_i - t^*)^\phi, \tag{40}$$

where A is defined as previously, but using new q^*. The analysis of this expression mirrors the analysis presented in the main text.

E.2 Disruption

This section informally discusses possible relaxations of the absence-of-disruption assumption and how the analysis change without this assumption.

Bound on b To relax the bound on b, one can show that investments will always exceed some minimum level k_{min}. Then one needs to upper-bound $b/(k_{\text{min}})^\beta$ instead of b. To establish a minimum level, one can use a "bootstrap" procedure. First, one considers the "worst case of costs", $c_i^{\text{eff}} = a + w$, which incentivizes the safety investments the least, and gets a minimum level of investments in this worst case, $k_{\text{min},1}$. (To show that the considered case is the worst case, one can compute the second derivatives of $\Pi_i^{m,*}$ over $\boldsymbol{k}$.) Using this number, one now bounds the investments from above and gets $k_{\text{max},1}$. The argument is then repeated using the bound implied by $k_{\text{max},1}$ and so forth. After sufficiently many iterations, this yields the desired minimum investment level.

Bound on δ An analytic expression is possible in the opposing limiting case, where δ is large, and the company exits the market if it does not adopt the AI technology. In this case, the outside option in the Nash bargaining problem corresponds to zero profit. In this case, the firms' investment response functions remain the same (since a firm's surplus still depends on k_i in the same way). However, the model provider's investment would change.

Allowing disruption Now, we discuss cases in which firms can drive competitors out of the market. First, if it is easy to make certain firms exit by simply adopting the technology (e.g., costs improve significantly with AI for some firms, but the improvement is small for others, making them uncompetitive), then the analysis mirrors our current analysis but with fewer firms. If, on the other hand, it is possible to make some firms exit the market, but it requires investment to exceed a threshold, then firms' incentives will exhibit a kink (a point where the derivative jumps). Below and above the threshold, the firms' profits remain smooth and follow our analysis. However, crossing the threshold causes some firms to exit, creating a kink in the marginal value of investment (profits will remain continuous). Hence, the investment levels can 1) continue to follow Theorem 5.1, 2) lie at the kink, or 3) fall to zero if the firm exits. The precise structure of these kinks depends on the asymmetry assumptions and is therefore difficult to predict.

NeurIPS Paper Checklist

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: [Yes]

   Justification: The description of the framework is based on Section 3. The description of the results reflect Sections 5 and 6.

   Guidelines:

   - The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
   - The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
   - The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
   - It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: [Yes]

   Justification: We explicitly discuss limitations in Section 7.

   Guidelines:

   - The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
   - The authors are encouraged to create a separate "Limitations" section in their paper.
   - The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
   - The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
   - The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
   - The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
   - If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
   - While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. **Theory assumptions and proofs**

   Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

   Answer: [Yes]

Justification: Appendix contains the proofs of all results.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. **Experimental result reproducibility**

   Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

   Answer: [N/A]

   Justification: The paper is theoretical.

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
   - If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
   - Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general. releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
   - While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
     (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
     (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
     (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
     (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. **Open access to data and code**

   Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [N/A]

Justification: The paper is theoretical.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. **Experimental setting/details**

   Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

   Answer: [N/A]

   Justification: The paper is theoretical.

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
   - The full details can be provided either with the code, in appendix, or as supplemental material.

7. **Experiment statistical significance**

   Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

   Answer: [N/A]

   Justification: The paper is theoretical.

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
   - The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
   - The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
   - The assumptions made should be given (e.g., Normally distributed errors).
   - It should be clear whether the error bar is the standard deviation or the standard error of the mean.

- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. **Experiments compute resources**

   Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

   Answer: [N/A]

   Justification: The paper is theoretical.

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
   - The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
   - The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. **Code of ethics**

   Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics `https://neurips.cc/public/EthicsGuidelines`?

   Answer: [Yes]

   Justification: The authors have read and follow the Code of Ethics.

   Guidelines:

   - The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
   - If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
   - The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. **Broader impacts**

    Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

    Answer: [N/A]

    Justification: We do not foresee any direct negative impact of our work. On the other hand, the positive impact of our work is hard to evaluate.

    Guidelines:

    - The answer [N/A] means that there is no societal impact of the work performed.
    - If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
    - Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. **Safeguards**

    Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

    Answer: [N/A]

    Justification: The paper is theoretical.

    Guidelines:

    - The answer [N/A] means that the paper poses no such risks.
    - Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
    - Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
    - We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. **Licenses for existing assets**

    Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

    Answer: [N/A]

    Justification: The paper is theoretical.

    Guidelines:

    - The answer [N/A] means that the paper does not use existing assets.
    - The authors should cite the original paper that produced the code package or dataset.
    - The authors should state which version of the asset is used and, if possible, include a URL.
    - The name of the license (e.g., CC-BY 4.0) should be included for each asset.
    - For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
    - If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
    - For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.

- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

    Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

    Answer: [N/A]

    Justification: The paper is theoretical.

    Guidelines:

    - The answer [N/A] means that the paper does not release new assets.
    - Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
    - The paper should discuss whether and how consent was obtained from people whose asset is used.
    - At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

    Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

    Answer: [N/A]

    Justification: The paper is theoretical.

    Guidelines:

    - The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
    - Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
    - According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

    Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

    Answer: [N/A]

    Justification: The paper is theoretical.

    Guidelines:

    - The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
    - Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
    - We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
    - For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: The authors used LLMs only for text editing.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.