arXiv:2610.21764v1 [cs.LG] 15 Oct 2026

Q(γ_H): Entropy-Guided Horizon Mixing in Multi-Discount Q-Learning

Arwa Albelaihi[1 [0000−0001−7135−5894]] and Abdulrahman Altahhan[1 [0000−0003−1133−7744]]

School of Computer Science, University of Leeds, Leeds LS2 9JT, UK
scaaalb@leeds.ac.uk, a.altahhan@leeds.ac.uk

Abstract. A fixed discount factor in Q-learning imposes a single planning horizon throughout training, despite the differing demands of early exploration and long-term credit assignment. We propose Q(γ_H), a multi-horizon Q-learning framework that maintains multiple discount-specific heads and adaptively combines them at control time using an entropy-derived confidence signal. The resulting state-dependent mixing weights shift control from shorter- to longer-horizon heads as policy confidence increases, enabling adaptive horizon mixing without modifying individual Bellman updates. Across classical control and deep reinforcement-learning benchmarks, Q(γ_H) shows its strongest improvements on delayed-credit tasks relative to strong fixed-discount baselines. By decoupling horizon adaptation from the underlying learning updates, Q(γ_H) provides a lightweight framework for adaptive planning-horizon control in reinforcement learning, avoiding the need for manual discount-factor selection while remaining compatible with standard value-learning methods.

Keywords: Reinforcement learning · Discount factor · Multi-horizon Q-learning · Entropy · Adaptive planning horizon

1 Introduction

Choosing the discount factor γ in reinforcement learning is challenging. A short horizon can stabilise early learning by limiting bootstrap error propagation, but may prevent solving tasks requiring long-range credit assignment [4,15].

No single fixed γ resolves this tension, and per-task tuning [7] does not address changes in the required planning horizon during training. Neuroscience evidence also suggests that dopaminergic neurons encode reward prediction errors across multiple discount time constants [9].

Recent methods address this through adaptive or multi-horizon discounting [13,1,12,11], but require modified Bellman targets, global rather than state-dependent mixing, reward decomposition, or added deep-RL infrastructure. These approaches do not provide state-dependent control-time mixing across discount-specific Q-heads while preserving their standard Q-learning updates. We propose Q(γ_H), which maintains K heads with discounts $\gamma_k \in \Gamma$ and mixes them using state-dependent confidence. Low confidence favours shorter horizons, while high confidence favours longer ones. For analysis, we define $\bar{\gamma}(s) = \sum_{k=1}^{K} w_k(s)\gamma_k$ as a mixture-weighted discount summary of horizon emphasis, not the discount of an equivalent Bellman operator. Each head retains its fixed-γ_k Q-learning update.

Across tabular, linear, and deep RL settings, Q(γ_H) is strongest on delayed-credit tasks while remaining competitive across the other evaluated environments. Our contributions are: (i) state-dependent horizon mixing that preserves standard Q-learning updates; (ii) a deep RL extension with discount-specific heads; and (iii) controlled and ablation analyses of the entropy-guided mixing mechanism.

2 Related Work

Q(γ_H) builds on multi-horizon value learning, adaptive discounting, and entropy-based uncertainty signals. Unlike prior work, it performs state-dependent mixing across planning horizons only at control time, preserves standard Bellman updates, and extends to deep RL through discount-specific Q-heads.

Parallel value prediction across planning horizons is central to the GVF framework [19] and discount-indexed value learning [4,15]. Related deep-RL methods use multiple discount-indexed Q-heads [1], reward-component-specific discounts [2], multi-discount PPO advantages [12], or time-weighted horizon heads in distributional RL [11]. Q(γ_H) also uses discount-specific heads, but couples them only at action selection: it requires no reward decomposition [2], uses per-state rather than global mixing [12], and avoids distributional RL machinery [11].

Prior work has formalised transition- and state-dependent discounting $\gamma(s, a, s')$ [22,16], used discounts to model termination or continuation values [23,18], and adapted γ online using critic loss or entropy-based uncertainty [13,3]. In contrast, Q(γ_H) does not modify Bellman targets or learn $\gamma(s)$ directly; it performs state-dependent control-time mixing while each head keeps a standard fixed-γ_k update. The resulting $\bar{\gamma}(s) = \sum_k w_k(s)\gamma_k$ is used only as a descriptive summary of the mixing weights.

Entropy is commonly used in RL as an objective [5,24], while ensembles use model diversity for uncertainty estimation and exploration [14,8], and mixture-of-experts methods use input-dependent gating [17]. In contrast, Q(γ_H) uses reference-policy entropy as a lightweight signal for decision ambiguity. Its head diversity is temporal rather than epistemic, and the gate interpolates across planning horizons rather than averaging to reduce variance. Table 1 summarises the closest comparisons.

A complementary motivation comes from neuroscience: dopaminergic neurons have been shown to encode reward prediction errors across multiple discount time constants [9]. This supports the use of parallel fixed-γ estimators, while Q(γ_H) adds an adaptive aggregation mechanism through entropy-guided control-time mixing.

3 Method: Q(γ_H)

3.1 Problem Setup

We consider a finite Markov decision process $\mathcal{M} = (\mathcal{S}, \mathcal{A}, P, r, \gamma)$ with finite action space $\mathcal{A}$ and discount factor $\gamma \in (0, 1)$. Standard Q-learning [21] fixes γ before training, although the appropriate planning horizon may vary across states and learning stages. Q(γ_H) instead maintains K parallel Q-learning heads, each trained with a different fixed discount factor, and mixes their values at control time using an entropy-derived confidence signal.

Table 1. Comparison of Q(γ_H) with closely related methods. State-dep. = state-dependent mixing or adaptation. Bellman changed = whether Bellman targets are modified. Control mix = whether discount-specific heads are mixed at action selection.

Method	Multi-head	State-dep.	Gate signal	Bellman changed	Control mix
Q(γ_H) (ours)	✓	✓	Entropy	No	✓
Ogawa et al. [13]	×	×	Critic loss	Yes	×
Moghimi et al. [11]	✓	×	None	No	×
Nafi et al. [12]	✓	×	None	No	×
Chen et al. [2]	✓	×	None	No	×

3.2 Multi-Horizon Q-Heads

Let $\Gamma = \{\gamma_1 < \cdots < \gamma_K\}$ be a set of discount factors. For each γ_k, Q(γ_H) maintains an independent Q-head $Q^{(k)} : \mathcal{S} \times \mathcal{A} \to \mathbb{R}$:

$$Q^{(k)}(s,a) \approx \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma_k^t r_t \mid s_0 = s, a_0 = a\right]. \tag{1}$$

Each head is parameterised as

$$Q^{(k)}(s,a) = W_k(a,:)^\top \phi(s), \tag{2}$$

with its independent discount factor γ_k and Bellman targets, while sharing the same experience stream and feature representation.

3.3 Reference Q Construction

To compute confidence, we construct a reference Q-vector $q_{\text{ref}} \in \mathbb{R}^{|\mathcal{A}|}$ from the current multi-head estimates. A simple baseline uses one fixed head:

$$q_{\text{ref}} = Q^{(k_{\text{ref}})}(s,\cdot), \tag{3}$$

which we evaluate in the ablation study.

Our default uses all heads through a max–min aggregation:

$$q_{\text{ref}}(a) = \eta \max_k Q^{(k)}(s,a) + (1-\eta) \min_k Q^{(k)}(s,a), \qquad \forall a \in \mathcal{A}, \tag{4}$$

where $\eta \in [0,1]$ balances optimistic and conservative estimates. We use $\eta = 0.5$ by default and assess alternatives in the ablation and sensitivity studies.

3.4 Entropy-Based Confidence

Given q_{ref}, we compute a Boltzmann policy with temperature $\tau > 0$:

$$\pi_{\text{ref}}(a \mid s) = \frac{\exp(q_{\text{ref}}(a)/\tau)}{\sum_{a'} \exp(q_{\text{ref}}(a')/\tau)}. \tag{5}$$

The normalised Shannon entropy of this policy is:

$$H_n(s) = -\frac{1}{\log|\mathcal{A}|}\sum_{a\in\mathcal{A}}\pi_{\text{ref}}(a \mid s)\log\pi_{\text{ref}}(a \mid s) \tag{6}$$

Normalisation by $\log|\mathcal{A}|$ ensures $H_n(s) = 1$ when all actions appear equally valuable and $H_n(s) \to 0$ when one action is clearly preferred. The agent's *confidence* is then:

$$c(s) = 1 - H_n(s). \tag{7}$$

Thus, high entropy implies low confidence, while low entropy implies high confidence.

3.5 Adaptive Horizon Mixing

We map confidence $c(s) \in [0, 1]$ to a distribution over the K heads using a soft proximity weighting. Each head k is assigned a fixed position on the unit interval:

$$u_k = \frac{k-1}{K-1}, \qquad k = 1, \ldots, K, \tag{8}$$

so that $u_1 = 0$ corresponds to the shortest-horizon head and $u_K = 1$ to the longest-horizon head. Mixing weights are computed as a softmax over negative distances from the confidence value:

$$w_k(s) \;=\; \frac{\exp\big(-\beta\,|u_k - c(s)|\big)}{\sum_{j=1}^{K}\exp\big(-\beta\,|u_j - c(s)|\big)}, \tag{9}$$

where β controls selection sharpness. Low confidence gives more weight to short-horizon heads, while high confidence favours longer horizons.

The head values are aggregated for action selection using the confidence-modulated weights:

$$Q_{\text{agg}}(s, a) = \sum_{k=1}^{K} w_k(s)Q^{(k)}(s, a). \tag{10}$$

For analysis, we additionally define the mixture-weighted discount summary

$$\bar{\gamma}(s) = \sum_{k=1}^{K} w_k(s)\gamma_k. \tag{11}$$

This quantity summarises whether the mixing weights emphasise shorter- or longer-horizon heads; it is not the discount factor of an equivalent Bellman operator.

3.6 Learning Rule

Actions are selected ε-greedily with respect to Q_{agg}. After observing (s, a, r, s'), each head is updated independently:

$$\delta_k = r + \mathbf{1}_{\{s'\notin\mathcal{T}\}}\gamma_k \max_{a'} Q^{(k)}(s', a') - Q^{(k)}(s, a). \tag{12}$$

$$W_k(a, :) \leftarrow W_k(a, :) + \alpha\delta_k\phi(s). \tag{13}$$

The mixing weights affect only action selection; each head keeps its own standard fixed-γ_k Q-learning update. Algorithm 1 summarises the full procedure.

Algorithm 1 Q(γ_H): Entropy-Guided Horizon Mixing

Require: $\Gamma = \{\gamma_1, \ldots, \gamma_K\}$, $\alpha, \tau, \beta, \eta, \varepsilon$

Initialise $W_k \leftarrow \mathbf{0}$ for all k; set $u_k \leftarrow (k-1)/(K-1)$
for each episode **do**
  Observe initial state s
  while s not terminal **do**
    Compute $Q^{(k)}(s,\cdot) \leftarrow W_k \phi(s)$ for all k
    Build reference: $q_{\text{ref}}(a) \leftarrow \eta \max_k Q^{(k)}(s,a) + (1-\eta) \min_k Q^{(k)}(s,a)$
    Derive $\pi_{\text{ref}}(\cdot \mid s)$ from q_{ref} via softmax with τ; compute entropy $H_n(s)$
    Set $c(s) \leftarrow 1 - H_n(s)$; compute weights $w_k(s)$ via softmax over $|u_k - c(s)|$
    $Q_{\text{agg}}(s,\cdot) \leftarrow \sum_k w_k(s)\, Q^{(k)}(s,\cdot)$
    Select a ε-greedily w.r.t. $Q_{\text{agg}}(s,\cdot)$; execute; observe r, s'
    $\delta_k = r + \mathbf{1}_{\{s' \notin \mathcal{T}\}}\, \gamma_k \max_{a'} Q^{(k)}(s',a') - Q^{(k)}(s,a)$ for all k
    $W_k(a,:) \leftarrow W_k(a,:) + \alpha\, \delta_k\, \phi(s)$ for all k
    $s \leftarrow s'$
  end while
  Decay ε
end for

3.7 Complexity and Practical Considerations

Each head in Q(γ_H) uses a standard off-policy Q-learning update under its own Bellman operator, where sufficient exploration is the key practical requirement [21,20]. The mixing weights are bounded, non-negative, and sum to one, so Q_{agg} is a well-defined convex combination of the heads:

$$Q_{\text{agg}}(s,a) = \sum_k w_k(s) Q^{(k)}(s,a).$$

Thus, state-dependent mixing affects action selection but not the individual Bellman targets.

The overhead of Q(γ_H) is linear in the number of heads K. Maintaining K heads requires $O(K|\mathcal{A}|d)$ memory, and dense linear evaluation and updates cost $O(K|\mathcal{A}|d)$ per step. The reference, entropy, and weighting computations add only $O(K|\mathcal{A}|)$ operations. In tabular and linear settings, no replay buffer, target network, or neural component is required. In the deep-RL extension, the same confidence-gated mixing is implemented on top of standard DQN and DDQN machinery using discount-specific output heads.

4 Experiments

We evaluate Q(γ_H) in tabular/linear and deep-RL settings. The tabular/linear suite contains eight benchmarks spanning delayed-credit, navigation, stochastic, and risk-sensitive tasks. The deep learning setting evaluates CartPole-v1, Acrobot-v1, and MountainCar-v0 using a shared-trunk multi-head DDQN architecture with one head per $\gamma \in \Gamma$. Table 2 summarises all environments.

In the tabular and linear settings, we compare Q(γ_H) with Q-learning, SARSA, Expected SARSA, and Double Q-learning. Each baseline uses an oracle discount γ^*

Table 2. Experimental environments used in evaluations. TC denotes tile coding, IHT denotes index hash table, and FC denotes a fully connected neural network.

Environment	Type	$\|\mathcal{A}\|$	Repr.	Key challenge
Tabular / linear				
MountainCar	Continuous	3	TC (IHT 2048)	Delayed credit
Acrobot	Continuous	3	TC (IHT 4096)	Delayed credit
LunarLander	Continuous	4	TC (IHT 8192)	Dense reward
Gridworld	Navigation	4	Tabular	Deterministic
WindyGridworld	Navigation	8	Tabular	Stochastic wind
CliffWalking	Navigation	4	Tabular	Cliff penalty
Stoch. CliffWalking	Navigation	4	Tabular	Slip p=0.10
FrozenLake-8×8	Stoch. nav.	4	One-hot	Slippery, sparse
Deep learning (normalised state)				
CartPole-v1	Control	2	FC [128, 128, 64]	Easy baseline
Acrobot-v1	Control	3	FC [128, 128, 64]	Delayed credit
MountainCar-v0	Control	3	FC [128, 128, 64]	Hard exploration

selected post hoc from Γ using best performance, while Q(γ_H) uses all discounts in Γ and adapts its horizon mixture online.

In the deep-RL setting, we compare Deep Q(γH) against five baselines: DDQN [6] and DQN [10] (single-head, fixed γ); Hyperbolic DQN [4] (K heads with hyperbolic spacing γ_k=$1/(1+kh)$, uniform mixing, no confidence gate); Ogawa-Adapt [13] (single γ adjusted by a loss-based PI controller); and Scheduled-γ (γ grown linearly from 0.80 to 0.99, no learned signal).

Table 3 lists all hyperparameters. Shared parameters are fixed across environments and both settings. Unless otherwise stated, experiments use N=10 seeds. For MountainCar, the default Q(γ_H) configuration uses the same 10 canonical runs wherever the identical setting appears; alternative configurations use separate 10-seed runs, with 20-seed results reported where stated. We report EarlyMean (first 50 episodes), TailMean (last 50 episodes), and AUC (full run). Episode length is the performance metric (lower is better) except for LunarLander (return, higher is better).

4.1 Main Results

Table 4 shows that Q(γ_H) is strongest on delayed-credit tasks. It achieves the best mean AUC on MountainCar and Acrobot, with reductions of 6.0% and 2.1%, respectively, relative to the next-best fixed-γ baselines. Across the remaining environments, Q(γ_H) remains competitive and often improves early learning, although fixed-γ or on-policy baselines can perform better when one horizon is sufficient or risk-sensitive behaviour is required.

Table 5 extends the comparison to deep RL. Deep Q(γH)-DDQN achieves the best AUC on all three environments. On Acrobot-v1 and MountainCar-v0, it outperforms DDQN by 8.5% and 7.0%, respectively. The largest margins are against adaptive baselines: Scheduled-γ is outperformed by 48.0% on Acrobot (p<0.0001) and 7.5% on

Table 3. Default hyperparameters used in the experiments.

Parameter	Symbol	Value	Notes
Shared (tabular/linear & deep)			
Initial exploration	ε_0	0.10	ε-greedy
Exploration decay	δ_ε	0.985	Per-episode
Min. exploration	$\varepsilon_{\min}$	0.05	
Tabular/linear only			
Discount set	Γ	$\{0.70, 0.80, 0.90, 0.97, 0.99\}$	$K = 5$ heads
Learning rate	α	0.025	TD step size
Episodes	T	200	
Max steps	`max_t`	5,000	
Temperature	τ	0.25	
Sharpness	β	3.0	
Mixing	η	0.5	
Deep learning only			
Discount set	Γ	$\{0.90, 0.95, 0.99\}$	$K = 3$ heads
Learning rate	α	10^{-4}	Adam optimiser
Episodes	T	300	
Replay buffer	—	50,000	Uniform sampling
Batch size	—	64	
Target sync	—	Every 100 steps	Hard update
Entropy temp.	τ	0.1	Sharper than tabular
Entropy warmup	—	200 steps	
Network	—	FC [128, 128, 64]	Shared trunk

MountainCar (p=0.0007), while Ogawa-Adapt is outperformed by 21.4% (p=0.0029) and 7.6% (p=0.0008), respectively. Within these classic-control tasks, these results support the benefit of confidence-guided horizon mixing.

4.2 Mechanism Analysis

Figure 1 examines the internal mixing signals on MountainCar. As training progresses, normalised entropy decreases from 0.960 to 0.800, while $\bar{\gamma}$ increases from 0.782 to 0.816, indicating a shift in mixing weight toward longer-horizon heads. The entropy–episode length relationship shows a positive association ($r = 0.306$), with higher entropy associated with longer episodes and poorer performance, consistent with its role as a lightweight confidence signal for horizon mixing.

To examine state dependence directly, we evaluated the learned gate over the MountainCar position–velocity state space using the same 10 canonical training seeds. In the 10-seed mean state-space evaluation, H_n ranged from 0.042 to 1.000 and $\bar{\gamma}$ from 0.773 to 0.951. The weights of the $\gamma = 0.70$ and $\gamma = 0.99$ heads ranged from 0.031–0.540 and 0.027–0.479, respectively, showing that the learned horizon mixture varies substantially across states.

Table 4. Q(γ_H) vs. oracle-tuned fixed-γ baselines. Values are means over 10 seeds. Steps ↓ except LunarLander return ↑.

Environment	Method	γ^*	Early	Tail	AUC
Delayed-credit environments					
MountainCar	Q(γ_H)	—	**366.36**	136.45	**203.58**
	Q-Learning	0.97	402.42	**136.16**	217.12
	SARSA	0.97	388.81	143.95	216.51
	Exp. SARSA	0.97	401.04	139.79	218.13
	Double Q	0.97	580.43	146.54	280.95
Acrobot	Q(γ_H)	—	602.97	**127.88**	**192.78**
	Q-Learning	0.97	535.94	144.53	200.92
	SARSA	0.97	543.22	144.99	200.32
	Exp. SARSA	0.99	**521.68**	143.68	196.96
	Double Q	0.97	636.53	163.37	228.16
General environments					
Gridworld	Q(γ_H)	—	**65.88**	19.60	37.60
	Q-Learning	0.99	66.14	19.74	**36.96**
	SARSA	0.97	68.68	**19.32**	37.16
	Exp. SARSA	0.97	66.16	19.96	36.98
	Double Q	0.99	77.70	31.74	48.90
WindyGridworld	Q(γ_H)	—	**319.23**	93.78	162.93
	Q-Learning	0.70	369.04	**68.37**	**160.50**
	SARSA	0.80	347.49	73.59	162.65
	Exp. SARSA	0.80	354.09	72.73	163.10
	Double Q	0.70	594.94	108.61	268.34
FrozenLake	Q(γ_H)	—	**197.40**	22.34	76.70
	Q-Learning	0.97	197.72	22.16	76.60
	SARSA	0.99	198.22	**21.02**	77.13
	Exp. SARSA	0.80	197.97	21.36	**72.80**
	Double Q	0.80	198.43	25.90	82.24
LunarLander ↑	Q(γ_H)	—	**−162.43**	−83.53	−122.56
	Q-Learning	0.99	−164.36	**−74.74**	**−116.09**
	SARSA	0.99	−168.46	−102.92	−137.53
	Exp. SARSA	0.99	−166.03	−105.40	−138.27
	Double Q	0.99	−171.01	−93.53	−130.96
Risk-sensitive environments					
CliffWalking	Q(γ_H)	—	128.06	**24.29**	58.29
	Q-Learning	0.99	**118.57**	24.31	**55.17**
	SARSA	0.97	120.60	25.05	57.03
	Exp. SARSA	0.70	131.43	24.30	59.42
	Double Q	0.99	159.47	48.80	88.31
Stoch. Cliff	Q(γ_H)	—	130.81	29.04	61.85
	Q-Learning	0.70	148.85	27.82	68.94
	SARSA	0.99	**122.11**	**27.64**	**58.74**
	Exp. SARSA	0.80	130.94	27.34	61.85
	Double Q	0.97	145.96	39.64	79.14

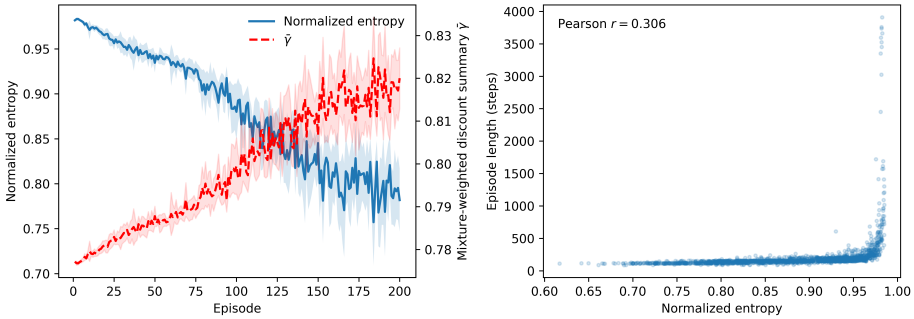

Fig. 1. Mechanism analysis on MountainCar using the same 10 canonical training seeds. Left: normalised entropy decreases while $\bar{\gamma}$ increases; shaded bands denote 95% confidence intervals across seeds. Right: normalised entropy versus episode length across all seed–episode observations ($N = 2{,}000$, Pearson $r = 0.306$).

Table 5. Deep Q(γH)-DDQN vs. all baselines (10 seeds, 300 episodes). AUC ↑ higher is better. Gain is computed as $100 \times (\text{AUC}_{Q(\gamma H)} - \text{AUC}_{\text{base}})/|\text{AUC}_{\text{base}}|$; positive values indicate that Q(γH)-DDQN performs better. $^{*}p<0.05$, $^{**}p<0.01$, $^{***}p<0.001$ (paired t-test, 10 seeds). **Bold** = best per column.

Method	**CartPole-v1 ↑**		**Acrobot-v1 ↑**		**MountainCar-v0 ↑**	
	AUC	Gain	AUC	Gain	AUC	Gain
Q(γH)-DDQN	**40,476**	—	**−34,660**	—	**−55,185**	—
Q(γH)-DQN	38,804	+4.3%	−40,242	+13.9%	−58,312	+5.4%*
DDQN	38,084	+6.3%	−37,900	+8.5%	−59,314	+7.0%**
DQN	26,851	+50.7%	−35,634	+2.7%	−59,112	+6.6%**
Hyperbolic [4]	34,616	+16.9%	−38,169	+9.2%**	−58,025	+4.9%*
Ogawa-Adapt [13]	36,999	+9.4%	−44,102	+21.4%**	−59,702	+7.6%**
Scheduled-γ	40,226	+0.6%	−66,684	+48.0%***	−59,638	+7.5%**

4.3 Head Disagreement as Uncertainty Proxy

We also measure head disagreement as an independent value-side diagnostic:

$$d(s,a) = \max_k Q^{(k)}(s,a) - \min_k Q^{(k)}(s,a). \tag{14}$$

Table 6 shows that disagreement increases during learning while entropy decreases and $\bar{\gamma}$ increases. The negative disagreement–entropy correlations ($r = -0.883$ on MountainCar and $r = -0.969$ on Acrobot) are consistent with increasingly decisive reference policies as horizon-specific value estimates separate. Positive disagreement–$\bar{\gamma}$ correlations indicate a concurrent shift in mixing weight toward longer-horizon heads.

4.4 Ablation Study

We analyse the contribution of the mixing gate, reference construction, gating signal, and number of heads.

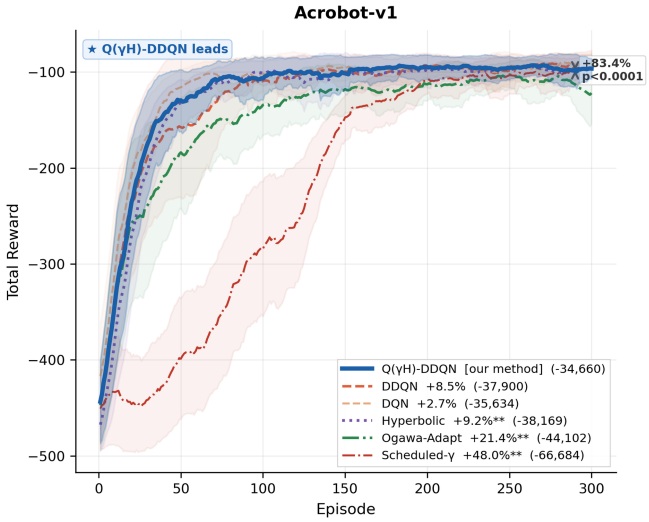

Fig. 2. Learning curves on Acrobot-v1 (10 seeds, 300 episodes, shaded = 95% CI, reward ↑). Q(γH)-DDQN (solid blue) converges faster and achieves the best cumulative reward (−34,660). It outperforms Scheduled-γ by 48.0% (p<0.001), Ogawa-Adapt by 21.4% (p<0.01), and DDQN by 8.5%. The annotation marks the gain over Scheduled-γ relative to the episode-wise peak gap. Since Acrobot assigns −1 per timestep, values closer to zero indicate faster task completion.

To isolate the effect of entropy-guided mixing from the general benefit of maintaining multiple heads, we compare Q(γ_H) with a parameter-matched uniform multi-head baseline using the same five discounts, representation, training procedure, and seeds, but fixed weights $w_k = 1/K$. Q(γ_H) improves EarlyMean from 397.63 to 366.36, TailMean from 141.83 to 136.45, and AUC from 216.50 to 203.58. Paired t-tests show significant differences for EarlyMean ($p < 0.001$), TailMean ($p = 0.007$), and AUC ($p < 0.001$). These results show that entropy-guided mixing improves both learning efficiency and final performance beyond the general benefit of maintaining multiple discount-specific heads.

We compare fixed-head and aggregated reference constructions on MountainCar. As shown in Table 7, MaxMin linear achieves the best EarlyMean, TailMean, and AUC. Mean Q and Max Q also outperform the best fixed-head variant in AUC, whereas the pessimistic Min Q construction does not. MaxMin linear improves over the best fixed-head variant by 2.3% AUC.

We compare three confidence signals on MountainCar using the MaxMin linear reference: entropy $c(s) = 1 - H_n(s)$, variance $c(s) = 1/(1 + v(s))$, and disagreement $c(s) = 1/(1 + d(s))$, where $v(s)$ and $d(s)$ are the mean Q-value variance and range across heads, respectively.

Table 8 shows similar performance across the three signals, with AUC spanning 203.58–206.41, suggesting that the mechanism is not tied to one gating signal. Entropy

Table 6. Signal dynamics and correlations on delayed-credit tasks (10 seeds). Pearson r is computed episode-wise over 200 episodes.

	MountainCar	**Acrobot**
Signal dynamics (first 50 → last 50 episodes)		
Mean head disagreement	15.53 → 27.77	1.03 → 8.15
Normalised entropy H_n	0.960 → 0.800	0.969 → 0.878
Discount summary $\bar{\gamma}$	0.782 → 0.816	0.780 → 0.800
Pearson r (episode-wise, 200 episodes)		
Disagreement → H_n	−0.883	−0.969
Disagreement → steps	−0.608	−0.318
Disagreement → $\bar{\gamma}$	+0.888	+0.970

Table 7. Reference construction comparison on MountainCar (10 seeds, steps ↓). Final $\bar{\gamma}$ and H_n are means over the last 50 episodes.

Method	**Type**	**Early ↓**	**Tail ↓**	**AUC ↓**	$\bar{\gamma}$	H_n
MaxMin linear	Aggregated	**366.36**	**136.45**	**203.58**	0.816	0.800
Mean Q	Aggregated	373.1	137.7	206.4	0.802	0.863
Max Q optimistic	Aggregated	370.7	138.9	206.8	0.789	0.926
Fixed middle	Fixed head	379.1	139.9	208.3	0.798	0.882
Fixed short	Fixed head	375.4	139.2	208.5	0.789	0.924
Min Q pessimistic	Aggregated	379.1	140.7	209.0	0.842	0.670
Fixed long	Fixed head	383.7	140.5	211.6	0.844	0.658

is retained as the default as a bounded measure of action ambiguity; its scale sensitivity is examined in the sensitivity analysis.

Table 9 compares the default $K = 5$ setting with a smaller $K = 3$ variant on MountainCar. In the 10-seed comparison, $K = 5$ achieves slightly lower EarlyMean, TailMean, and AUC than $K = 3$, although the differences in EarlyMean and AUC are small. In the separate 20-seed analysis, $K = 3$ performs marginally better. Overall, the small differences between the two configurations suggest limited sensitivity to the number of heads in this setting. Both configurations increase $\bar{\gamma}$ over training, indicating a shift in mixing weight toward longer-horizon heads. We retain $K = 5$ as the default to provide finer horizon granularity.

4.5 Sensitivity Analysis

Table 10 summarises sensitivity on MountainCar. Q(γ_H) is most sensitive to the sharpness parameter β: too small a value spreads weight too diffusely and slows early learning. It is less sensitive to τ and η, with robust performance for $\tau \in [0.25, 0.50]$ and $\eta \in [0.5, 0.9]$. We therefore use $\beta = 3.0$, $\tau = 0.25$, and $\eta = 0.5$ as defaults.

Table 8. Gating signal ablation on MountainCar (10 seeds, MaxMin linear reference, steps ↓). Best value per column in **bold**.

Signal	Early ↓	Tail ↓	AUC ↓	$\bar{\gamma}$ first 50	$\bar{\gamma}$ last 50
Entropy	**366.36**	**136.45**	**203.58**	0.782	0.816
Variance	372.79	138.70	205.25	0.795	0.784
Disagreement	375.35	138.47	206.41	0.796	0.787

Table 9. Head count ablation on MountainCar (steps ↓), 10 and 20 seeds. Best value per column in **bold**.

Seeds	Config	Early ↓	Tail ↓	AUC ↓	$\bar{\gamma}$ first/last 50
10	K=3	367.20	139.39	204.16	0.755 / 0.790
	K=5	**366.36**	**136.45**	**203.58**	0.782 / 0.816
20	K=3	**367.36**	**139.94**	**204.62**	0.755 / 0.790
	K=5	370.36	140.60	205.60	0.780 / 0.811

5 Discussion and Limitations

The results show a task-dependent pattern across tabular, linear, and deep settings: Q(γ_H) is most useful when the planning horizon is difficult to fix in advance.

The strongest gains occur in MountainCar and Acrobot, where short horizons support stable early learning but longer horizons are needed for delayed credit assignment. Q(γ_H) addresses this trade-off by shifting control from shorter- to longer-horizon heads as confidence increases, producing a gradual shift in horizon emphasis within a single run.

When a single fixed discount is already effective, as in CartPole and LunarLander, horizon mixing provides limited benefit. In the cliff environments, SARSA remains competitive because its on-policy updates encourage safer behaviour near hazardous states; horizon mixing does not remove this on-policy advantage.

The deep Q(γH) results demonstrate that the same control-time mixing mechanism can be applied with neural function approximation, with the largest gains on MountainCar-v0 and Acrobot-v1 and little improvement on CartPole-v1. However, the current deep-RL evaluation is limited to three classic discrete-action control tasks and a shared-trunk multi-head architecture, so these results establish applicability rather than broad deep-RL generality. Further work should evaluate larger benchmark suites, separate-head architectures, continuous-action actor–critic variants, and the theoretical properties of the coupled multi-head system.

The tabular and linear experiments isolate the horizon-mixing mechanism without replay buffers, target networks, or neural optimisation effects. The comparison is conservative because fixed-γ baselines use oracle-selected γ^*, whereas Q(γ_H) adapts its horizon mixture online. In the deep setting, DDQN uses $\gamma = 0.99$, the strongest fixed discount considered.

The mechanism analyses make horizon mixing observable through the joint evolution of confidence, $\bar{\gamma}$, and head disagreement, while the state-wise analysis confirms that the

Table 10. Sensitivity analysis on MountainCar (10 seeds, steps ↓). Robust ranges are determined from the corresponding AUC results. Best value per block in **bold**.

Param	Value	Early ↓	Tail ↓	Default	Robust range	Sensitivity
	3.0	**366.36**	**136.45**			
β	1.5	388.69	140.0	3.0	[1.5, 3.0]	Moderate
	0.5	403.03	141.9			
	0.25	**366.36**	**136.45**			
τ	0.50	368.12	137.8	0.25	[0.25, 0.50]	Low
	0.10	369.92	140.1			
	0.5	**366.36**	**136.45**			
	0.7	368.84	138.33			
η	0.9	373.02	137.21	0.5	[0.5, 0.9]	Low–moderate
	0.3	374.86	140.68			
	0.1	377.00	139.70			

learned mixture also varies across the MountainCar state space. Here, $\bar{\gamma}$ is used only as a descriptive summary of the relative weight assigned to shorter- and longer-horizon heads.

This pattern is also consistent with neuroscience evidence that reward prediction errors are represented across multiple discount timescales [9], although we make no claim of mechanistic correspondence.

A limitation of the current entropy gate is its sensitivity to reward and Q-value scale under a fixed Softmax temperature. The reward-scaling experiment shows that this sensitivity can materially affect performance at large scale changes. Scale-adaptive temperatures or normalised confidence signals are therefore important directions for future work.

6 Conclusion

We introduced Q(γ_H), a multi-horizon Q-learning framework that avoids choosing a single fixed discount factor by maintaining discount-specific Q-heads and mixing them at control time using a state-dependent confidence signal. The resulting mixing weights adapt horizon emphasis during training while each head retains a standard fixed-γ_k Q-learning update. For analysis, $\bar{\gamma}(s) = \sum_k w_k(s)\gamma_k$ provides a descriptive summary of this horizon mixture.

Across eight tabular and linear benchmarks, Q(γ_H) is most effective on delayed-credit tasks, achieving the best mean AUC on MountainCar and Acrobot against oracle-tuned fixed-discount baselines. Mechanism and ablation analyses show state-dependent horizon mixing and improved learning efficiency from the entropy-guided gate, while sensitivity analysis identifies reward/Q-value scale as a limitation of the current fixed-temperature confidence signal.

The deep-RL extension, Deep Q(γH), shows the same qualitative pattern across the three evaluated classic-control tasks, with the largest gains on MountainCar-v0 and

Acrobot-v1 and little improvement on CartPole-v1. Overall, the results suggest that adaptive horizon mixing is a lightweight and interpretable alternative to manual selection of a single discount factor, especially when delayed credit assignment makes horizon choice difficult. Future work will examine larger benchmarks, continuous-action settings, alternative multi-head architectures, and scale-robust confidence mechanisms.

Disclosure of Interests. The authors have no competing interests to declare.

References

1. Ali, R.F., Duong, K., Nafi, N.M., Hsu, W.: Multi-horizon learning in procedurally-generated environments for off-policy reinforcement learning (student abstract). In: Proceedings of the AAAI Conference on Artificial Intelligence. vol. 37, pp. 16150–16151 (2023)
2. Chen, P., Liu, Q., Li, Y., Yan, K., Ma, S.: Multi-discounting reinforcement learning based on reward decomposition. IEEE Transactions on Emerging Topics in Computing **14**(1), 1–14 (2026)
3. Cho, H., Kim, H.: Entropy-guided distributional reinforcement learning with controlling uncertainty in robotic tasks. Applied Sciences **15**(5), 2773 (2025)
4. Fedus, W., Gelada, C., Bengio, Y., Bellemare, M.G., Larochelle, H.: Hyperbolic discounting and learning over multiple horizons. arXiv preprint arXiv:1902.06865 (2019)
5. Haarnoja, T., Zhou, A., Hartikainen, K., Tucker, G., Ha, S., Tan, J., Kumar, V., Zhu, H., Gupta, A., Abbeel, P., et al.: Soft actor-critic algorithms and applications. arXiv preprint arXiv:1812.05905 (2018)
6. Hasselt, H.v., Guez, A., Silver, D.: Deep reinforcement learning with double q-learning. Proceedings of the AAAI Conference on Artificial Intelligence **30**(1), 2094–2100 (2016)
7. Henderson, P., Islam, R., Bachman, P., Pineau, J., Precup, D., Meger, D.: Deep reinforcement learning that matters. In: Proceedings of the AAAI conference on artificial intelligence. vol. 32 (2018)
8. Lakshminarayanan, B., Pritzel, A., Blundell, C.: Simple and scalable predictive uncertainty estimation using deep ensembles. In: Advances in Neural Information Processing Systems. vol. 30 (2017)
9. Masset, P., Tano, P., Kim, H.R., Malik, A.N., Pouget, A., Uchida, N.: Multi-timescale reinforcement learning in the brain. Nature (2025). https://doi.org/10.1038/s41586-025-08929-9
10. Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A.A., Veness, J., Bellemare, M.G., Graves, A., Riedmiller, M., Fidjeland, A.K., Ostrovski, G., et al.: Human-level control through deep reinforcement learning. nature **518**(7540), 529–533 (2015)
11. Moghimi, M., Coache, A., Ku, H.: Decoupling time and risk: Risk-sensitive reinforcement learning with general discounting. arXiv preprint arXiv:2602.04131 (2026)
12. Nafi, N.M., Ali, R.F., Hsu, W., Duong, K., Vick, M.: Policy optimization using horizon regularized advantage to improve generalization in reinforcement learning. In: Proceedings of the 23rd International Conference on Autonomous Agents and Multi-Agent Systems (AAMAS) (2024)
13. Ogawa, K., et al.: Adaptive discount factor for accelerating policy learning considering long-term returns in reinforcement learning with non-stationary environments. Journal of Computational Science (2026)
14. Osband, I., Blundell, C., Pritzel, A., Van Roy, B.: Deep exploration via bootstrapped dqn. Advances in neural information processing systems **29** (2016)

15. Romoff, J., Henderson, P., Touati, A., Brunskill, E., Pineau, J., Ollivier, Y.: Separating value functions across time-scales. In: International Conference on Machine Learning. pp. 5468–5477. PMLR (2019)
16. Sharma, A., Gupta, R., Lakshmanan, K., Gupta, A.: Transition based discount factor for model free algorithms in reinforcement learning. Symmetry **13**(7) (2021). https://doi.org/10.3390/sym13071197, https://www.mdpi.com/2073-8994/13/7/1197
17. Shazeer, N., Mirhoseini, A., Maziarz, K., Davis, A., Le, Q., Hinton, G., Dean, J.: Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. arXiv preprint arXiv:1701.06538 (2017)
18. Stachurski, J., Zhang, J.: Dynamic programming with state-dependent discounting. Journal of Economic Theory **192**, 105190 (2021)
19. Sutton, R.S., Modayil, J., Delp, M., Degris, T., Pilarski, P.M., White, A., Precup, D.: Horde: A scalable real-time architecture for learning knowledge from unsupervised sensorimotor interaction. In: The 10th International Conference on Autonomous Agents and Multiagent Systems-Volume 2. pp. 761–768 (2011)
20. Tsitsiklis, J.N.: Asynchronous stochastic approximation and q-learning. Machine Learning **16**(3), 185–202 (1994)
21. Watkins, C.J.C.H., Dayan, P.: Q-learning. Machine Learning **8**(3–4), 279–292 (1992)
22. White, M.: Unifying task specification in reinforcement learning. In: International Conference on Machine Learning. pp. 3742–3750. PMLR (2017)
23. Yoshida, N., Uchibe, E., Doya, K.: Reinforcement learning with state-dependent discount factor. In: 2013 IEEE third joint international conference on development and learning and epigenetic robotics (ICDL). pp. 1–6. IEEE (2013)
24. Ziebart, B.D., Maas, A., Bagnell, J.A., Dey, A.K.: Maximum entropy inverse reinforcement learning. In: Proceedings of the 23rd AAAI Conference on Artificial Intelligence. pp. 1433–1438 (2008)