arXiv:2610.21640v1 [cs.LG] 15 Oct 2026

ReFlow: Replay-Driven Decoupling for Efficient RL Post-Training

Yukun Tian[1,3], Junwei Su[2], Haofeng Xu[1,3], Lansong Diao[3], Zhengping Qian[3], Weibo Zhao[3], Wenchao Zhou[3], Ruiting Zhou[4], Chuan Wu[1]*

[1]*The University of Hong Kong* [2]*University of Science and Technology of China*
[3]*Alibaba Group* [4]*Southeast University*

Abstract

Reinforcement learning (RL) has been a critical paradigm for improving large language model (LLM) reasoning and agentic capabilities, but its efficiency remains a major bottleneck. Existing RL systems commonly rely on asynchronous execution to overlap rollout generation with policy optimization for enhanced efficiency. However, this approach faces inherent limitations: mild policy staleness offers only limited training efficiency gains, whereas aggressive asynchrony can degrade model quality and yields diminishing efficiency gains. The dynamic nature of LLM RL and multi-turn agentic workloads further challenges RL system design, as resource demands shift over time, making static resource allocation and generation strategy inefficient.

This paper presents ReFlow, a system-algorithm co-designed RL framework for efficient LLM post-training. Its key insight is that, unlike existing RL frameworks which consume each trajectory only once, experience replay of recent trajectories enables effective sample reuse, thereby amortizing the high cost of rollout generation without incurring excessive policy staleness. To fully exploit this opportunity, ReFlow adopts a decoupled and adaptive architecture with three key designs: a dual replay data cycle supports efficient experience reuse while balancing training throughput and optimization stability; an elastic resource pool enables dynamic resource placement to accommodate evolving RL workloads; an adaptive scheduler coordinates generation strategies and resource allocation to improve resource utilization and boost generation throughput. Experiments across representative single-turn reasoning and multi-turn agentic RL tasks show that ReFlow improves training throughput by up to 2.06× over state-of-the-art synchronous and asynchronous RL systems, while maintaining comparable or better model performance.

*Corresponding author.

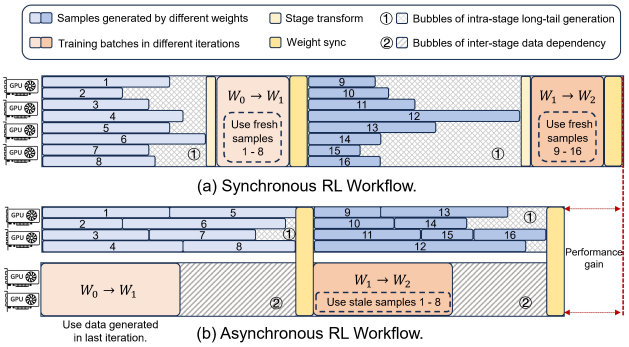

Figure 1: Typical LLM RL workflow: (a) synchronous RL; (b) asynchronous RL. W denotes different versions of weights.

1 Introduction

Reinforcement learning (RL) has been widely adopted in post-training large language models (LLMs), enabling model alignment with human preferences [31, 39], enhanced reasoning [11, 30] and interactive agentic capabilities [4, 10]. In RL post-training, the pretrained LLM iteratively generates responses to prompts, receives reward signals from human or automated feedback, and updates its policy to maximize expected return. In practice, this workflow is typically decomposed into two primary stages—rollout generation and policy training—with data dependency (Fig. 1 (a)). During rollout generation, rollout workers use the behavior policy, i.e., the LLM, to produce responses and collect batches of trajectories; subsequently, actor workers perform policy optimization on the accumulated fresh on-policy samples and update the behavior policy with new parameters.

Despite its effectiveness, this RL workflow introduces significant system efficiency challenges in large-scale RL training. Rollout generation is particularly expensive for LLMs, as responses are generated autoregressively and can account for around 80% of total training time [37], with multi-turn agentic RL further amplifying this cost. Moreover, response lengths and interaction turns vary substantially across prompts, re-

sulting in severe GPU underutilization in synchronous RL training (① in Fig. 1 (a), where trainer needs to collect a batch of on-policy data to start policy optimization). To mitigate these inefficiencies, recent systems increasingly adopt asynchronous RL (Fig. 1 (b)), where rollout generation and policy training proceed concurrently on disaggregated resources, and the trainer updates the policy with stale samples generated by rollout workers using earlier policy versions. By enabling trainers and rollout workers to operate in parallel, asynchronous RL execution can significantly reduce long-tail generation bubbles and improve hardware utilization. As LLM RL workloads continue to scale in terms of model size, context length, and interaction complexity, such asynchronous designs are rapidly evolving from a performance optimization into a practical requirement [10, 25].

Introducing asynchrony alone does not guarantee efficient or stable RL training, as policy staleness can degrade model performance. Therefore, effective asynchronous RL systems must jointly balance execution efficiency, algorithmic stability, and resource utilization, raising key challenges at the intersection of distributed systems and RL algorithms.

Challenge #1: Residual data dependencies still limit system efficiency. Asynchronous RL allows the trainer to use trajectories generated from slightly stale policies, reducing resource under-utilization from waiting for long-tail generation. However, existing asynchronous RL systems still incur undesirable GPU idling bubbles due to residual coupling of training and generation. As shown in Fig. 1 (b), the trainer may still block until enough rollout generation data is collected to form a batch. This phenomenon is particularly pronounced in LLM RL, where rollout time dominates overall execution time and varies substantially due to diverse response lengths and potentially complex environmental interactions.

Challenge #2: Limited system–algorithm co-optimization fails to meet the joint requirements of efficiency and stability. Effective asynchronous RL requires jointly optimizing system efficiency and algorithmic stability. Stale-policy rollouts can destabilize learning if not properly controlled, leading to a *staleness dilemma*: small staleness cannot hide long-tail generation latency [13, 23, 43, 51], while excessive staleness causes severe policy mismatch and training instability [6, 29]. Although recent algorithmic correction methods can improve RL robustness, they often overlook system-level efficiency and yield only limited performance gain [29, 45].

Challenge #3: Non-adaptive system designs fail to accommodate dynamic and heterogeneous RL workloads. Compared with conventional distributed training, LLM RL workloads are more dynamic and heterogeneous. Rollout generation is dominated by autoregressive decoding and environmental interaction, whose costs vary significantly across samples due to diverse response lengths and unpredictable tool-execution latency. In contrast, cost of policy training is relatively regular and batch-oriented. Moreover, these costs evolve during training and vary across models, tasks, and hardware platforms, causing shifting resource demands of different RL stages over time. Non-adaptive resource allocation designs in existing asynchronous RL systems [6, 23, 24] are therefore inherently inefficient: fixed configurations like resource allocation and generation concurrency can cause severe GPU load imbalance and resource under-utilization. On the other hand, naive runtime resource reconfiguration is also expensive, as it requires cold-starting generation and training engines on new devices and rebuilding communication groups, often leaving all GPUs idle for several minutes.

Our Approach. We propose ReFlow, a fully decoupled and adaptive asynchronous RL framework for efficient LLM post-training. ReFlow aims to accelerate training by reducing the demand for expensive rollouts through system–algorithm co-design. Our key insight is that, beyond widely-adopted policy asynchrony, *experience replay* [21] offers an additional, underexplored opportunity to improve LLM RL efficiency. **By carefully reusing recently generated samples, we enable experience replay in RL post-training to significantly reduce the demand for costly new rollouts without introducing excessive policy mismatch for model optimization.** This is particularly valuable in multi-turn agentic RL, where trajectories involve long-term interactions. We also provide theoretical guarantee that experience replay can reduce gradient variance and stabilize training. Together, these properties enable more flexible and efficient asynchronous execution.

Based on this insight, ReFlow is designed to fully exploit experience replay to achieve efficient LLM RL post-training:

First, ReFlow introduces a carefully designed dataflow named *dual replay data cycle* that turns experience replay into improvements in both system throughput and training convergence. ReFlow redesigns the dataflow and component interactions of existing asynchronous RL systems, decoupling rollout generation from policy training. In its replay-driven cyclic dataflow, training batches are constructed from a mixture of fresh, stale, and selected replayed samples, reducing the need for costly new rollouts while keeping policy mismatch under control. It also adjusts the rollout budget and training batch composition accordingly at each iteration to balance efficiency and stability. To better manage the dataflow, ReFlow inserts a replay buffer as an intermediate data layer. The buffer enables global staleness control over fresh and replayed data, and exposes extensible sampling and filtering interfaces for customized data selection, forming another replay data cycle inside experience buffer.

Second, to address dynamic resource demands of LLM RL and new workload patterns introduced by experience replay, ReFlow employs an elastic resource pool for efficient runtime resource scaling. By shifting the overhead of resource reallocation from system reinitialization to lightweight worker state transitions and reassignment, ReFlow allows compute resources to be reallocated between rollout generation and policy training within seconds, ensuring efficient fresh-sample

generation to reduce policy mismatch and preserve training stability. ReFlow further incorporates an *adaptive scheduler* that continuously optimizes generation strategies and resource allocation based on replay requirements and runtime system metrics. By adapting to diverse and evolving system demands in replay-driven RL, the scheduler sustains both high resource utilization and stable fresh-sample supply.

We deploy ReFlow on clusters with up to 128 GPUs and evaluate it on representative single-turn reasoning and multi-turn agentic RL tasks against state-of-the-art RL systems. Results show that ReFlow improves training throughput by up to 2.06× while maintaining comparable or even better model performance, demonstrating its scalability and efficiency.

2 Background and Motivation

2.1 RL for LLM Post-Training

Reinforcement learning (RL) for LLM post-training adopts RL algorithms such as PPO [35] and its variants [19, 36, 49] to optimize the policy LLM using feedback from reward models [2, 31], rule-based evaluators [11, 36], or environment interactions [4, 28]. A typical LLM RL workflow consists of two stages: *generation* and *training*. During generation, the policy model generates responses or trajectories for a batch of prompts; in agentic settings, this stage may further involve interactions with tools, simulators, or external environments. During training, the generated trajectories are evaluated to compute learning signals, such as rewards and advantages, which are then used to update the policy model through distributed optimization [38].

Despite algorithmic differences, modern LLM RL systems share a common execution pattern: large-scale rollout generation using many GPU workers followed by gradient-based policy updates. This creates a heterogeneous workload: rollout generation requires high-throughput autoregressive inference over large prompt batches, while training relies on distributed backpropagation with intensive GPU computation and communication; reward evaluation and trajectory processing add further compute and data-movement overhead.

Recent LLM RL systems increasingly adopt asynchronous execution that allows trainers to consume trajectories produced by rollout workers using earlier policy versions. They run rollout generation and policy training on disaggregated resources in parallel, reducing long-tail generation bubbles [6, 13, 16, 29, 37, 51].

2.2 Limitations of Existing Asynchronous RL Systems

Strict data dependencies between training and generation. In RL frameworks, rollout generation and training usually have strict data dependencies.

(1) *Intra-stage data dependencies* within rollout generation stage, in that the same prompt batch is processed collectively and generated trajectories can only be consumed after the entire input batch has finished generation. Response lengths often vary substantially across prompts: generation of some trajectories terminates quickly, while others require much longer decoding. Consequently, rollout workers (GPUs) that finish early must wait for long-tail generation, resulting in significant GPU usage bubbles (① in Figure 1).

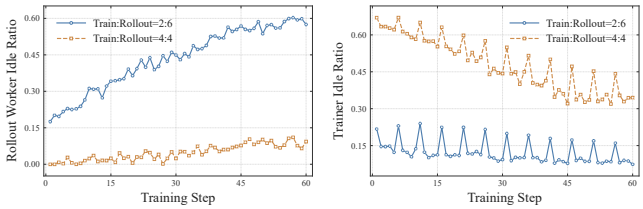

(a) Rollout idle ratio. (b) Trainer idle ratio.

Figure 2: GPU idle ratio (idle time / step time) under different resource allocations for Qwen3-4B training on 8 GPUs.

(2) *Inter-stage data dependencies*, that the training stage typically requires a complete batch of generated trajectories (composed from one or multiple rollout workers) for policy optimization. When rollout generation lags behind or produces insufficient samples, training GPUs are forced to stall, waiting for additional data (② in Figure 1). Conversely, when the system bottleneck lies on the trainer side, rollout workers must wait for the trainers before advancing to the next generation batch. Figure 2 shows that the GPU idle time due to data dependencies can occupy up to 60% of the training iteration time, severely limiting overall system utilization.

Existing asynchronous RL systems address this inefficiency mainly by optimizing rollout generation to mitigate the long-tail generation effect, through load balancing across rollout workers [16, 33, 37], generation-order scheduling [9, 52], and dynamic batching [37, 51]. These techniques often involve additional profiling or coordination overhead which may offset their performance gains, or rely on assumptions that limit their generality. StreamRL [51] separates long and short samples generation into different resource groups, requiring response-length profiling and prediction. Rollpacker and APRIL [9, 52] defer long samples to form dedicated batches, which may alter the original data distribution and break iteration consistency for training [37]. Seer [33] estimates request lengths from other requests in the same group, suitable mostly for group-based RL algorithms with highly similar responses and incompatible with more general RL methods such as PPO. More broadly, systems that use response length as a proxy for rollout workload [33, 51] are less effective in agentic RL, where tool execution time can dominate request latency.

Overall, the existing asynchronous RL systems optimize the rollout stage but retain the strict data dependencies between rollout workers and trainers. This fundamentally limits their end-to-end training efficiency.

Limited system–algorithm co-optimization. In LLM post-

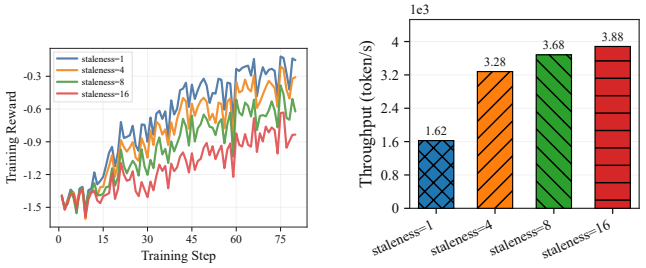

(a) Convergence. (b) Throughput.

Figure 3: Training model performance under different staleness settings when training Qwen3-1.7B on 8 GPUs, where staleness denotes the number of training iterations between two consecutive parameter synchronizations.

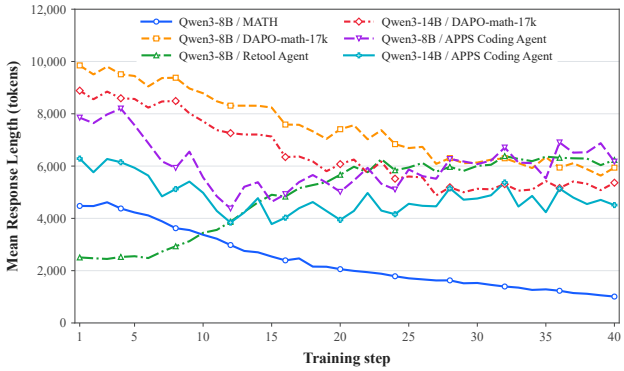

Figure 4: Diverse and evolving generation characteristics under different training configurations, demonstrating the highly dynamic workload pattern of LLM RL.

training, a fundamental goal is to achieve the best model performance under limited system time/resource budget. This requires jointly optimizing two complementary dimensions: *system efficiency*, in terms of how many training samples can be generated and consumed per unit time, and *algorithmic convergence*, in that how effectively each sample improves the policy. We formalize the overall RL training performance of an execution design $\mathcal{S}$ as

$$\text{Perf}(\mathcal{S}) = \underbrace{\text{Throughput}(\mathcal{S})}_{\text{system efficiency}} \cdot \underbrace{\text{DataEfficiency}(\mathcal{S})}_{\text{algorithmic convergence}} \quad (1)$$

$\text{Throughput}(\mathcal{S})$ measures the number of valid samples generated and consumed per unit system time, while $\text{DataEfficiency}(\mathcal{S})$ measures the training improvement achieved per unit data.

Asynchronous RL systems are facing the *staleness dilemma* (Figure 3): increasing staleness can better decouple training and generation, thereby reducing rollout bubbles and improving their execution overlap; on the other hand, larger staleness amplifies policy mismatch, which can reduce data efficiency and destabilize training, offsetting the system-level benefits [6, 37].

To balance this trade-off, many asynchronous systems adopt a conservative *one-step-off* strategy [13, 23, 43, 51], that the generation policy lags behind the training policy by only one version. Some further reduce policy staleness by mixing one-step-off stale samples with on-policy samples for training [24]. While this maintains relatively fresh training data, the limited staleness window provides little room to hide long-tail rollout latency, resulting in only modest improvements in system efficiency [37]. Allowing larger staleness can further decouple the training pipeline and improve resource utilization, but the system benefit of increasing staleness quickly diminishes: beyond a certain point, the additional overlap between generation and training becomes marginal, limiting further throughput gains.

There is another line of work focusing on improving the algorithmic robustness of asynchronous RL under high asynchrony [18, 22, 44, 45, 50]. These approaches may stabilize training, but often overlook system efficiency implications. For instance, GAC [45] requires additional gradient storage to handle stale updates, which significantly increases GPU memory consumption. Other works do not consider system-level optimization opportunities [29, 44, 50]. Consequently, it remains unclear whether such techniques can translate algorithmic improvements into meaningful end-to-end training performance gains in LLM RL.

Overall, neither system- nor algorithm-oriented approaches are sufficient for building an effective LLM RL framework. The key challenge is to develop a co-design that jointly accounts for system throughput and learning dynamics.

Non-adaptive system designs not accommodating dynamic RL workloads well. Existing RL frameworks often adopt static execution strategies, in which resource allocation and key system configurations are fixed during RL training, not adaptive to the highly dynamic workloads.

First, LLM RL post-training exhibits substantial *runtime dynamics*. As the policy evolves, the distribution of rollout trajectories changes continuously, often resulting in shifting response length patterns (Figure 4). Early and late training steps can exhibit markedly different rollout behaviors, including variations in response length, reasoning structure, and computational intensity. Consequently, the workload composition across decoding, memory usage, and compute demand evolves throughout training, making initially selected static system configurations increasingly suboptimal over time, as shown in Figure 2.

Second, RL workloads are highly sensitive to diverse system and algorithmic settings. The LLM post-training ecosystem is rapidly evolving across model architectures, datasets, optimization algorithms, and hardware platforms. For example, critic-free methods such as GRPO and GSPO [36, 49] eliminate the need for a separate critic model, but require multiple responses per prompt, further shifting the bottleneck to rollout generation. Meanwhile, datasets vary from short-answer benchmarks such as GSM8K [3] to long-chain

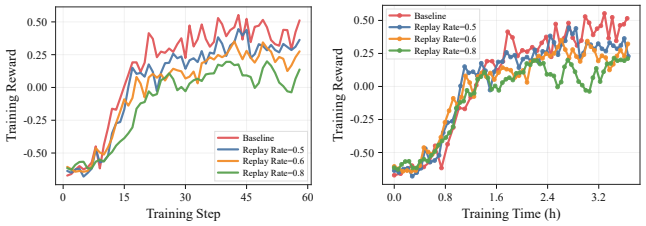

(a) Performance vs. step. (b) Performance vs. time.

Figure 5: Qwen3-4B on DAPO-MATH-17K. Naive experience replay slows convergence and degrades performance.

reasoning datasets such as DAPO-Math-17K [47], leading to substantially different rollout patterns. Emerging *agentic RL* further amplifies this variability, as rollouts may involve tool calls, multi-step reasoning, and interactions with external environments, resulting in dynamic execution paths and unpredictable latency [8, 37]. Consequently, optimized designs tuned to a specific stage or workload can quickly become ineffective as workload patterns shift. This makes it difficult to determine suitable configurations before training, such as GPU allocation between rollout and training or the number of concurrent generation requests.

Taken together, these factors make it difficult for a single static execution strategy to remain efficient throughout RL training. Adaptive system designs are therefore needed to dynamically adjust execution strategies as workload evolve.

2.3 Experience Replay in Asynchronous RL

Experience replay is widely used in conventional deep RL (e.g., DQN [26]) to improve data efficiency and stabilize training [7, 12, 20, 21, 34]. By storing past trajectories and reusing them for multiple policy updates, it partially decouples data collection from policy optimization and amortizes the cost of expensive environmental interactions.

Opportunity. In LLM RL post-training, where rollout generation is often dominated by long autoregressive decoding, experience replay provides a natural way to improve system efficiency. When newly generated trajectories are insufficient to form a batch, the trainer can continue optimization with replayed trajectories, thereby relaxing the strict dependency between rollout generation and training and mitigating long-tail rollout delays. Unlike simply increasing policy staleness in rollout generation, whose benefits quickly diminish [6], increasing the replay ratio directly reduces the demand for new rollouts and lowers generation workload. Moreover, we further show that introducing replay samples can help reduce gradient variance and stabilize training, which is particularly important for LLM RL. This makes experience replay particularly attractive for asynchronous LLM RL, with several challenges introduced.

Challenge #1: Experience replay alone does not guarantee better end-to-end training performance. Similar to stale

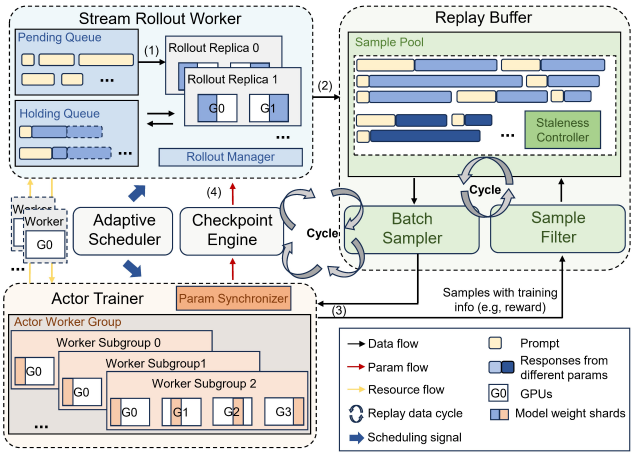

Figure 6: System overview of ReFlow.

samples, replayed trajectories introduce off-policy data, potentially causing distributional bias and hinder convergence. Thus, whether experience replay can improve LLM RL efficiency without compromising model convergence remains unclear. Recent studies provide only partial evidence. RePO [17] can outperform on-policy GRPO in some cases, but its additional updates on replayed data increase training time by around 15%, undermining its efficiency gains. Another concurrent work [1] analyzes how experience replay can reduce computation, but lacks system-level optimization and is evaluated at limited scale and scope. Our experiments show that naively reusing a fixed fraction of samples from the previous step can slow convergence and degrade final model performance (Fig. 5). These highlight that replay cannot be treated as a simple source of training data. Realizing its efficiency while preserving training stability remains an open challenge.

Challenge #2: Experience replay reshapes workload patterns and introduces additional dynamics into LLM RL post-training. Increasing the proportion of replayed trajectories in each training batch reduces the rollout workload and shifts the system bottleneck toward the trainer. Combined with the evolving and highly dynamic workload characteristics described in Sec. 2.2, it is even more difficult to decide optimized training configurations and maintain high efficiency throughout RL training. The optimized settings can vary significantly over time, depending on both the current system conditions and the data replay strategy.

3 ReFlow System Overview

We introduce ReFlow, an efficient LLM RL post-training framework that leverages experience replay to substantially reduce the need for costly new rollouts without incurring excessive policy mismatch. This introduces a new optimization dimension beyond conventional model asynchrony and expands the system design space. Building on this insight, ReFlow adopts a fully decoupled and adaptive architecture

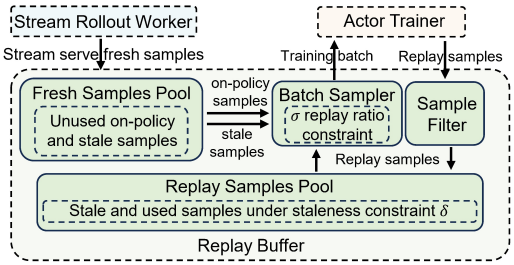

Figure 7: Dataflow of ReFlow.

centered on a dual replay data cycle, with system–algorithm co-design to jointly improve system efficiency and training stability. Figure 6 presents the overall architecture of ReFlow, which consists of four main components.

Stream Rollout. The rollout worker continuously generates trajectories using the behavior policy model and interacts with the environment in agentic tasks. It consists of a rollout manager for request scheduling and rollout coordination, and a set of rollout workers that run serving engines such as vLLM [15] for inference. Unlike conventional batched rollout, ReFlow adopts sample-level streaming generation, which continuously produces trajectories and pushes them to the replay buffer. We also enable support for partial rollouts [41]. This design reduces GPU idle time caused by long-tail response generation and improves rollout throughput.

Replay Buffer. The replay buffer is the central data layer of ReFlow. It stores generated trajectories and supports replay-driven cyclic dataflow by reusing them across training iterations. It consists of a staleness-aware sample pool, a batch sampler, and a sample filter that determines reuse eligibility. ReFlow exposes modular interfaces for customizing the sampler and replay filter, enabling flexible replay policies.

Trainer. The trainer consumes trajectories from the replay buffer and updates the policy model using RL optimization algorithms. It also handles other computations associated with experience consumption, such as reference model evaluation and critic model updates. The trainer can run on multiple GPUs in a distributed training setting.

Adaptive Scheduler. The adaptive scheduler orchestrates the system by dynamically controlling rollout generation and resource allocation. Based on runtime metrics and replay requirements, it periodically adjusts generation scheduling and resource placement between rollout workers and trainers. This enables ReFlow to maintain high GPU utilization while controlling policy staleness across diverse and evolving LLM RL workloads, in the presence of experience replay.

4 Dual Replay Data Cycle

4.1 Replay Data Cycle in Decoupled Architecture

ReFlow leverages experience replay to further decouple training from generation and mitigate the rollout bottleneck. To exploit replay while addressing the stability challenges discussed in Section 2.3, we design a replay-driven dataflow built on a fully decoupled architecture (Fig. 7). Table 1 summarizes the key hyperparameters of our replay data cycle.

Table 1: Hyperparameters in dual replay data cycle.

Symbol	Description
σ	Maximum fraction of replay data in a training batch.
τ	Loss weight controlling the relative contribution of replayed trajectories.
δ	Maximum data policy lag allowed for training.

(1) Request Dispatch. Input prompts are first packed and dispatched to the rollout workers as sample-level requests and then inserted into a pending queue to be served in a streaming manner. We set an upper bound $B_{\max}$ to cap the number of samples dispatched to rollout workers at each step. A larger $B_{\max}$ allows the inference engine to process more generation requests in parallel and better fill long-tail bubbles, but may also introduce more stale samples into the next policy-optimization step.

(2) Rollout Generation. Rollout workers continuously generate trajectories in a sample-level streaming fashion, agnostic to training steps, batch size and model versions. Completed trajectories are asynchronously pushed to the replay buffer.

(3) Replay-Driven Training. The trainer continuously samples trajectories from the replay buffer to perform policy updates. Since the buffer contains both newly generated samples from different model versions (unused fresh on-policy and stale samples) and replay samples (trajectories used in former iteration), the trainer can construct batches from various sources, reducing its dependence on new rollouts.

As shown in Figure 7, rather than fixing the data replay ratio, we argue that the replay ratio should adapt to the relative rates of rollout generation and training consumption to accommodate fluctuations in sample availability. However, unrestricted replay can exacerbate policy mismatch and hinder convergence. We therefore introduce a *data refeed boundary* σ, which caps the replay fraction while allowing it to vary with runtime conditions. This bounded flexibility balances system efficiency and training stability, with each batch of size b consuming between $b(1-\sigma)$ and b unused samples.

Moreover, to mitigate potential policy shifts caused by replayed samples, we reweight its policy loss:

$$w_i = \begin{cases} \omega_i, & i \in \mathcal{D}_{\text{unused}}, \\ \tau\omega_i, & i \in \mathcal{D}_{\text{replay}}, \end{cases} \qquad \tilde{w}_i = \frac{w_i}{Z}, \quad Z = \frac{\sum_j w_j n_j}{\sum_j n_j}.$$

Here, $\mathcal{D}_{\text{unused}}$ and $\mathcal{D}_{\text{replay}}$ denote sets of unused and replay data, $\tau \geq 0$ is the replay loss weight, ω_i is the importance weight, and n_i is the number of valid tokens. The normalized weight $\tilde{w}_i$ preserves the policy-loss scale while keeping replayed samples relatively down-weighted.

(4) Parameter Synchronization. To control policy staleness, the trainer triggers parameter synchronization with rollout workers through a checkpoint engine after policy updates

every few training steps. Due to the stream dispatch design in Step (1), rollout workers may still hold several partially generated samples when receiving model synchronization signals. These requests are moved to a holding queue and continued after parameter synchronization, forming partial rollout [41].

Overall, our replay-driven workflow combines bounded sample reuse, normalized loss reweighting, and explicit staleness control. It reduces reliance on new rollouts without simply enlarging the staleness window, enabling system–algorithm co-optimization.

Theoretical Justification. ReFlow combines two complementary mechanisms: replay reduces rollout demand, while normalized reweighting controls replay's contribution to policy updates. The following local analysis provides a theoretical basis for this design. Consider a fixed replay fraction $\lambda_t \in (0, \sigma]$, with $\sigma < 1$, and equal importance-weighted token mass per sample. Let $\hat{g}_t^F$ and $\hat{g}_t^R$ denote gradient estimates from unused and replayed data. For $0 \le \tau \le 1$, normalized reweighting gives

$$\hat{g}_t^{\mathrm{rw}} = (1-\alpha_t)\hat{g}_t^F + \alpha_t \hat{g}_t^R,$$
$$\alpha_t = \frac{\tau\lambda_t}{1-\lambda_t+\tau\lambda_t} \le \frac{\tau\sigma}{1-\sigma+\tau\sigma}. \tag{2}$$

Thus, ReFlow separately controls *how much data is reused* through λ_t and *how strongly replay influences the update* through α_t, allowing efficient reuse with moderated gradient influence.

Controlling gradient error. Let $g_t^\star = \nabla J_t(\theta_t)$ be the local on-policy gradient. In this simplified model, suppose $\hat{g}_t^F$ is unbiased for $g_t^\star$, while $\hat{g}_t^R$ has bias norm at most Δ_t. Assume their centered errors are uncorrelated, with variances v_t^F and v_t^R. All moments condition on the learner state, batch composition, and replay weight. Then

$$\mathbb{E}_t \|\hat{g}_t^{\mathrm{rw}} - g_t^\star\|^2 \le \mathcal{E}_t(\alpha_t), \tag{3}$$
$$\mathcal{E}_t(a) := \underbrace{(1-a)^2 v_t^F + a^2 v_t^R}_{\text{variance}} + \underbrace{a^2\Delta_t^2}_{\text{squared bias bound}}.$$

For an L-smooth surrogate, the update $\theta_{t+1} = \theta_t + \eta_t \hat{g}_t^{\mathrm{rw}}$ with $0 < \eta_t \le 1/L$ satisfies

$$\mathbb{E}_t[J_t(\theta_{t+1}) - J_t(\theta_t)] \ge \frac{\eta_t}{2}\left(\|g_t^\star\|^2 - \mathcal{E}_t(\alpha_t)\right). \tag{4}$$

This directly connects replay reweighting to optimization quality: a smaller gradient-error bound yields a stronger certified one-step progress bound.

Benefit of weighted replay. For fixed source estimates and $v_t^F > 0$, the error bound is minimized over $a \in [0, \lambda_t]$ by

$$\alpha_t^{\mathrm{opt}} = \min\left\{\lambda_t, \frac{v_t^F}{v_t^F + v_t^R + \Delta_t^2}\right\}. \tag{5}$$

Importantly, $\alpha_t^{\mathrm{opt}} > 0$ and $\mathcal{E}_t(\alpha_t^{\mathrm{opt}}) < \mathcal{E}_t(0) = v_t^F$. Under the stated assumptions, appropriately weighted replay therefore achieves a strictly tighter gradient-error bound than discarding replay, without requiring additional unused samples. Intuitively, a small replay contribution reduces the unused-data variance term to first order, while replay variance and squared bias enter at second order. When $\alpha_t^{\mathrm{opt}} < \lambda_t$, down-weighting also strictly improves the bound over uniform weighting. These results support retaining replay while adjusting its influence to balance variance reduction and policy mismatch.

Support for our co-design. ReFlow's controls address complementary aspects of this balance. The replay cap σ guarantees at least $b(1-\sigma)$ unused samples per batch, while τ moderates replay influence without changing batch composition. For example, $\sigma = \tau = 0.5$ permits half the samples to be replayed while limiting α_t to $1/3$ in the equal-mass model. Synchronization and the policy-lag limit restrict trajectory age, helping control Δ_t when per-update policy drift is bounded.

4.2 Replay Data Cycle in Experience Buffer

Experience replay introduces a substantially more complex dataflow and data-scheduling space than existing RL systems. ReFlow's experience buffer represents a central data layer enabling explicit lifecycle management for training data under experience replay. Especially, the experience buffer is designed with two key considerations: (1) it should allow flexible data reuse, controllable staleness, and adaptive sample selection, to fully exploit replay opportunities; (2) it should be easily customizable, exposing sufficient interfaces to support algorithm-specific sampling and filtering strategies. To this end, the buffer implements a replay-driven cyclic dataflow with structured data management.

Sample Pool with Staleness Control. The sample pool maintains fresh samples from on-policy and stale model versions, and replayed trajectories available for training. To prevent excessive policy mismatch caused by stale data, ReFlow enforces a global staleness δ on stored trajectories, i.e., only samples generated in the recent δ steps are available. We enable this by associating each sample with additional metadata including its policy version for generation and user-customized training information (e.g, advantages).

Batch Sampling. The sampler selects trajectories from the sample pool to construct training batches, with replay proportion in each batch determined by runtime conditions and sampling strategies and is capped by σ (Sec. 4.1). Users can implement customized sampling strategies to further improve training stability. For example, a medium-reward sampler selects only trajectories with intermediate rewards, i.e., trajectories that are neither too easy nor too hard [47].

Replay Filtering and Reinsertion. After being consumed by training, trajectories are passed through a replay filter that determines whether they should be reinserted into the sample pool for future reuse. The filtering policy can incorporate quality signals such as rewards, advantages, or other algorithm-

Settings	Cost
8B / 32 GPUs	228.20s
14B / 64 GPUs	286.31s
32B / 128 GPUs	368.64s

(a) Total overhead.

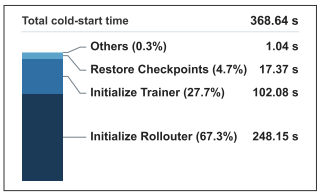

(b) Overhead breakdown.

Figure 8: Cold-start overhead. (a) Total overhead under Qwen3 models. (b) Breakdown for Qwen3-32B on 128 GPUs.

dependent metrics. Similar to sampling, replay filtering is orthogonal to the system design and can be customized using different algorithms and replay strategies. To ensure a fair comparison with other RL systems, we use random sampling and accept all replay samples without filtering rather than carefully designed strategies in our experiments.

Together, the sample pool, sampler, and replay filter form a closed-loop replay data cycle. Newly generated trajectories enter the sample pool, are sampled for training, and may be reinserted for later reuse.

5 Dynamic Resource Placement

5.1 Design Consideration

Dynamic RL workloads and experience replay jointly reshape the resource demands of rollout generation and policy training. To accommodate these evolving bottlenecks, ReFlow requires an elastic resource pool that efficiently reallocates GPUs between the two stages. Runtime scaling is expensive in asynchronous RL systems due to substantial system overhead. Rollout workers host inference engines with large KV caches, whereas training workers maintain optimizer states, gradients, and training model shards. Training and rollout often rely on different parallelism strategies, operator implementations, and fixed communication groups. Resource scaling under disaggregated placement is fundamentally more complex than the fixed training-generation context switch in synchronous RL systems such as VeRL [38]. Rather than switching context under a stable topology, the distributed execution structure itself needs to be changed, requiring communication-group reconfiguration, resource remapping, and materialization of missing states such as up-to-date optimizer states. Reallocating GPUs across stages requires expensive reinitialization and incurs significant cold-start overhead, including rebuilding inference engines, reconstructing training meshes, and relaunching distributed jobs—operations that can consume minutes of GPU time, as shown in Figure 8.

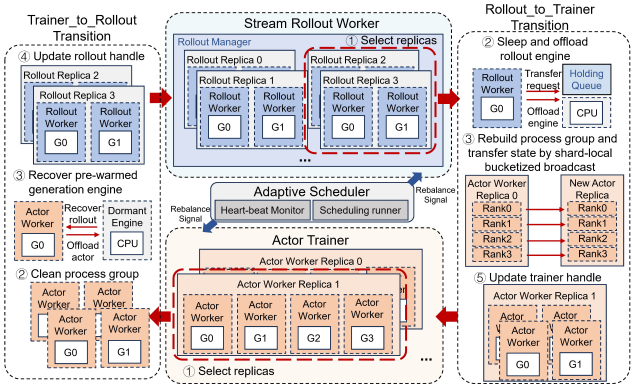

Figure 9: Workflow of ReFlow's resource rebalance. We use TP=2 for the rollout worker and FSDP group size=4 for the trainer as an example.

5.2 Elastic Resource Pool Using Unified AsyncHybrid Worker

To eliminate heavy cold-start overhead in runtime resource placement, ReFlow introduces the *Unified AsyncHybrid Worker*, a unified execution abstraction that decouples worker lifetime from its assigned role. Instead of maintaining separate rollout and training workers like existing systems [6, 23, 24, 37], our Unified AsyncHybrid Worker encapsulates necessary runtime modules for both inference and training, and can dynamically switch roles on demand, enabling flexible resource reallocation. To mitigate the additional memory overhead for maintaining dual capabilities, ReFlow employs a selective offloading strategy, by which inactive modules (e.g., inference model when training) are temporarily offloaded to CPU memory, thereby preserving valuable GPU memory for the active workload.

ReFlow treats workers (each holding a GPU) as the basic unit of resource management. Resource pools of high-level abstractions, such as the Stream Rollout Component and the Trainer, are defined as collections of Unified AsyncHybrid Workers, and runtime scaling is achieved by adjusting GPU assignments to run the workers and switching execution states. By replacing heavyweight system reinitialization with lightweight worker reassignment and state transitions, ReFlow enables efficient and flexible elastic scaling under evolving workload demands.

5.3 Resource Rebalancing

ReFlow builds an elastic resource pool on top of the Unified AsyncHybrid Worker abstraction to support efficient resource rebalancing between rollout generation and policy training. The rebalancing workflow is given in Figure 9, whose design follows two principles. *First*, resources are reassigned in complete replica groups to preserve valid parallelism configurations. Let s_T and s_R denote the GPUs per trainer and rollout

replica. For a pool of W GPUs,

$$A = n_T s_T, \qquad R = n_R s_R, \qquad A + R = W, \tag{6}$$

where n_T and n_R are the corresponding replica counts. *Second*, shard partitions are preserved when possible, allowing retained trainers to keep local model and optimizer shards and avoiding costly full-state gathering or repartitioning.

Pre-Warm Initialization. To avoid cold-start reinitialization overhead, we perform a pre-warm stage before RL training for expensive runtime contexts such as inference engines. At startup of training, workers are initialized as hybrid workers capable of serving both trainer and rollout roles. Each worker may hold both the training and inference engines, while inactive modules are offloaded or put into a dormant state to reduce GPU memory usage. This amortizes expensive setup costs during resource rebalance such as process initialization, CUDA graph capture, and inference engine construction.

Role Transition. Lightweight bidirectional role transitions between trainer and rollout pools are enabled, by preserving shard layouts and reusing pre-warmed components. When some resources are shifted from trainer workers to rollout workers, shifting workers activate their dormant rollout modules, while the remaining trainer workers keep their local parameter and optimizer shards. Since only the replica dimension changes, no full model or optimizer state needs to be gathered or redistributed. Newly added rollout workers then receive up-to-date policy model weights. When resources are shifted from rollout workers to trainer workers, ReFlow similarly avoids full-state materialization by initializing new trainer replicas through shard-local broadcast: existing trainer replicas serve as sources, and each worker transfers only its local parameter and optimizer shard to the corresponding worker in the new replica. We implement bucketized CUDA broadcasts optimization, which is inspired by Pytorch DDP [32] to reduce collective launch overhead. These shard-local transfers proceed in parallel and avoid full gather, CPU staging, or global redistribution. Together, pre-warmed rollout reuse and shard-local trainer initialization enable efficient resource rebalancing as local state changes rather than costly subsystem restart.

Moreover, we typically trigger resource rebalancing upon periodical parameter synchronization (Sec. 4.1), avoiding conflicts between inference and policy update and extra parameter updates for new rollout workers. After the transition completes, ReFlow updates the trainer and rollout worker-group handles and resumes normal parameter synchronization.

6 Adaptive Scheduling for Dynamic RL Workloads

We further design an adaptive scheduler for short-term generation scheduling among rollout workers and long-term resource rebalancing across trainer and rollout workers, adapting to both generation dynamics and resource-demand shifts.

6.1 Adaptive Generation Scheduling

Our sample-level streaming rollout largely removes the long-tail generation bottleneck, shifting the main challenge to *generation concurrency control* (which decides how many requests a rollout worker receives at a time) and *load balancing* among rollout workers. Low concurrency under-utilizes rollout GPUs, while aggressive concurrency for long-context workloads can create KV-cache pressure and excessive preemption [33]. Our adaptive generation scheduling keeps rollout workers highly utilized while avoiding KV-cache congestion, exploiting: a slow feedback controller that adjusts the real-time effective concurrency cap of each rollout worker every few seconds, and a fast global request dispatcher that admits and routes incoming requests to rollout workers subject to these caps.

Adaptive Cap Control. The controller periodically collects KV-cache utilization, waiting queue length, generation throughput, number of preemption requests and active generation tasks of each inference server. For agentic workloads, only active LLM generations are counted toward concurrency. Each worker i maintains EWMA-smoothed KV utilization ($\hat{u}_{i,t}$) and preemption signals ($\widehat{\Delta P}_{i,t}$):

$$\hat{u}_{i,t} = \alpha u_{i,t} + (1-\alpha)\hat{u}_{i,t-1}, \quad \widehat{\Delta P}_{i,t} = \alpha \Delta P_{i,t} + (1-\alpha)\widehat{\Delta P}_{i,t-1}.$$

The effective sequence cap $C_{i,t}$ is updated using AIMD: it increases additively when KV utilization is below the target kv cache usage band and demand exists, and decreases multiplicatively when congestion appears, such as sustained KV preemption growth or KV utilization reaching a hard limit.

Token-aware Request Admission and Routing. Within the adaptive caps, the global dispatcher admits requests based on projected KV-token pressure rather than request count. A request r can be assigned to worker i only if

$$\frac{T_{i,t}^{\text{pred}} + \widehat{T}(r)}{Budget_i} \le \rho,$$

where $T_{i,t}^{\text{pred}}$ is the predicted active token load on worker i, using the average response length of recent requests to estimate workload, $\widehat{T}(r)$ is the estimated prompt-plus-response footprint and $Budget_i$ is the worker token budget. Requests are not routed to a rollout worker if observed KV utilization exceeds the admission limit of the worker, even when available request slots remain. Among feasible rollout workers, the dispatcher chooses the one with the minimum cost:

$$\text{cost}_i(r) = \max\left(\frac{T_{i,t}^{\text{pred}} + \widehat{T}(r)}{Budget_i}, \hat{u}_{i,t}\right) + \lambda_p \widehat{\Delta P}_{i,t}.$$

Here λ_p is a hyper-parameter. This favors workers with lower KV load, shorter queues, and lower preemption pressure. We

Algorithm 1: Rollout/Train Resource Scheduler

Input: Current resource placement (A, R) of Actor Trainer A and Stream Rollout Worker R, feasible placement set $\mathcal{P}$, total workers W, Runtime statistics: number of produced and consumed samples $(\Delta N_{\text{prod}}, \Delta N_{\text{cons}})$, trainer and rollout worker idle ratio $i_{\text{rollout}}, i_{\text{train}}$, time window Δt

Input: Target refeed ratio σ, small constant ε

Output: Updated rollout and trainer worker placement

Estimate active-time rates using idle ratio correction:

$$\lambda_{\text{r}} \leftarrow \frac{\Delta N_{\text{prod}}/\Delta t}{1 - i_{\text{rollout}} + \varepsilon}, \qquad \lambda_{\text{d}} \leftarrow \frac{\Delta N_{\text{cons}}/\Delta t}{1 - i_{\text{train}} + \varepsilon}.$$

Set $q^\star \leftarrow 1 - \sigma$;

Select the best feasible placement:

$$(A^\star, R^\star) \leftarrow \arg \min_{\substack{(A',R') \in \mathcal{P} \\ A'+R'=W}} \left| \log \frac{\frac{R'}{R}\lambda_{\text{r}}}{\frac{A'}{A}\lambda_{\text{d}}(q^\star + \varepsilon) + \varepsilon} \right|.$$

if *the cooldown gate is active* **then**
  return (A, R);

return $(A^\star, R^\star)$;

further explore prompt affinity, by placing sibling requests of the same prompt on the same worker when possible to exploit prefix cache, while allowing spillover under overload.

Overall, this scheduling raises concurrency when utilization is low, prevents KV overcommitment through token-aware admission, and balances load across rollout workers under diverse, dynamic workloads.

6.2 Rollout/Trainer Resource Allocation

Dynamic resource allocation between trainer and rollout workers is necessary to tackle time-varying workloads in asynchronous RL training. Experience replay further introduces a tradeoff between generating fresh samples and reusing buffered data. We design a lightweight resource allocation scheme adapting rollout/train resources according to the observed producer–consumer dynamics.

Producer-consumer Modeling. As detailed in Algorithm 1, we consider rollout workers as sample producers and trainers as sample consumers. The scheduler measures production, consumption, and GPU idle ratios over a recent time window (i.e., a training step) and estimates the effective rollout capacity and trainer demand. This correction prevents the scheduler from misinterpreting blocked execution caused by sample starvation or staleness constraints as intrinsic throughput changes.

Scheduling Strategy. As given in Algorithm 1, the scheduler searches over a discrete feasible placement set that respects global worker limits and system constraints such as tensor parallelism and trainer shard grouping. Each candidate placement is evaluated by extrapolating production and consumption rates from the current allocation. The goal is to match the predicted fresh-data coverage of the new placement plan to the target coverage induced by the desired refeed ratio σ. We use a logarithmic mismatch metric to make the decision invariant to multiplicative scale and symmetrically penalizes under- and over-supply of fresh samples.

To avoid frequent resource reallocation, a new placement is applied only when permitted by a cooldown gate, which enforces a minimum number of training steps between two consecutive rebalancing operations.

7 Evaluation

7.1 Implementation

We implement ReFlow on top of VeRL [38] with a total codebase of 10.6K lines of Python.

Fully decoupled RL architecture. Following VeRL, we use RAY [27] as distributed framework. The core implementation (6.6K LoC) realizes a fully decoupled RL architecture with a dual replay data cycle, separating rollout generation and policy training into independent pipelines. The stream rollout worker is built on VLLM [15] and the trainer is implemented with PYTORCH using FSDP [48], while the overall design remains compatible with alternative inference and training engines. We use NCCL with a customized *CheckpointEngine* [40] for efficient parameter synchronization.

Adaptive scheduler is implemented with 1.3K LoC and constructed when initializing trainer and rollout workers.

7.2 Experiment Setup

Testbed. We deploy ReFlow on an H800 cluster with 16 nodes (128 GPUs in total). Each node is equipped with eight H800 GPUs connected by NVLink, and nodes are interconnected with 8×400 Gbps bandwidth. The key software versions include CUDA 12.6, PyTorch 2.7.1, and NCCL 2.26.2.

Models and Datasets. We run RL training on Qwen3-4B, 8B, 14B and 32B [46], which are widely used in both research and industry. For single-turn reasoning RL, we follow [37,47] to use DAPO-Math-17K [47], a representative benchmark for mathematical reasoning. For multi-turn agentic workloads, we follow DiDPO [42] and train coding agents on APPS [14], where the agent iteratively develops a solution program for challenging competitive-programming problems. We further follow ReTool [5] to evaluate tool-interactive reasoning on DAPO-Math-17K, where the model invokes a code sandbox during problem solving.

Baselines. We compare ReFlow with several state-of-the-art open-source LLM RL systems. For the synchronous RL baseline, we use VeRL [38], a widely adopted LLM post-training framework, with its official implementation. For asyn-

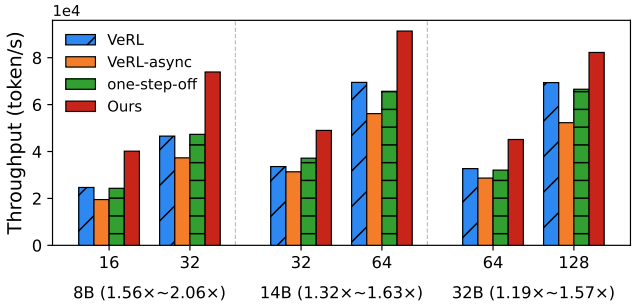

(a) Single-turn reasoning RL on different number of GPUs.

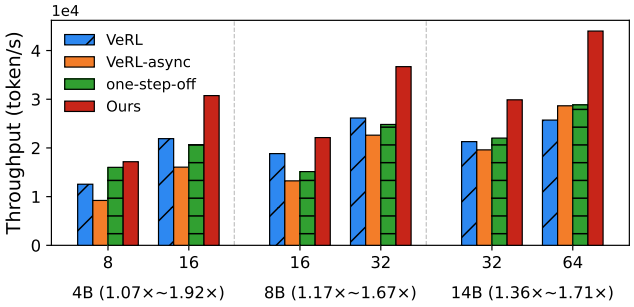

(b) Multi-turn coding agentic RL on different number of GPUs.

Figure 10: System throughput under different RL workloads.

chronous RL baselines, many concurrent systems do not release code [16, 37], provide limited hardware support [13], or lack mature system optimizations [29]. Therefore, following [16, 37], we use two state-of-the-art open-source asynchronous RL systems, one-step-off [23] and VeRL-async [24], as baselines. We follow [37] and official running scripts of baselines to set the hyper-parameters (i.e., parallelism). To ensure fair comparison, all systems use the same software environment, with vLLM as the generation backend and FSDP2 as the training backend.

Settings. For design parameters in ReFlow, we set the maximal generated samples B_{max} equal to the training batch size b. For training refeed boundary σ and replay loss weight τ, we set both to 0.5 to reduce the policy mismatch introduced by experience replay. The global data staleness δ is set to 1, allowing only samples generated using the latest and one-version-early models to be kept in the replay buffer.

Metrics. We measure system efficiency using throughput (tokens per second), computed as the number of tokens processed by the trainer in one step divided by the step time. We report the average throughput over five measured steps after five warm-up steps. For coding agent experiments, where complex interactions lead to higher performance variance, we average throughput over ten measured steps after ten warm-up steps. We also evaluate average training reward during the training steps, which serves as an indicator of training stability and algorithmic effectiveness of the RL framework.

7.3 Training Throughput

We evaluate the training throughput of ReFlow against various baselines across single-turn reasoning and multi-turn agentic workloads, varying both model size and cluster scale.

Single-Turn Reasoning. As shown in Figure 10 (a), ReFlow consistently achieves the highest throughput for Qwen3-8B, 14B, and 32B on 16–128 GPUs, delivering 1.19×–2.06× speedups over the evaluated baselines. The speedups are 1.56×–2.06×, 1.32×–1.63×, and 1.19×–1.57× for 8B, 14B, and 32B, respectively.

Multi-Turn Agentic Tasks. These gains extend to workloads involving repeated generation and environment interaction.

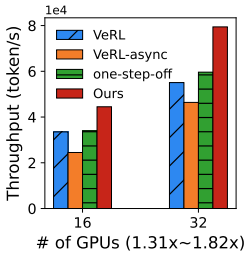

(a) Training throughput

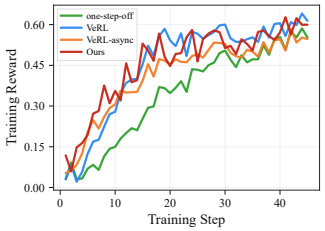

(b) Training reward

Figure 11: The training performance on Retool agentic tasks for Qwen3-8B.

For APPS coding tasks, Figure 10 (b) shows that ReFlow outperforms all baselines for Qwen3-4B, 8B, and 14B on 8–64 GPUs, achieving speedups of 1.07×–1.92×, 1.17×–1.67×, and 1.36×–1.71×, respectively. For ReTool tool-interaction tasks, ReFlow further achieves 1.31×–1.82× higher throughput with Qwen3-8B on 16 and 32 GPUs (Figure 11 (a)). Together, these results demonstrate the effectiveness of ReFlow across both single-turn and multi-turn RL workloads.

These improvements come from two complementary designs. The dual replay data cycle reduces fresh-rollout demand through trajectory reuse and mitigates trainer stalls caused by variable rollout completion times, while dynamic resource placement and adaptive scheduling reduce pipeline idle time under changing workloads. We further analyze their contributions in Sec. 7.5.

7.4 Training Convergence

We evaluate reward progression over both training steps and wall-clock time to assess whether ReFlow accelerates training while preserving convergence.

Single-Turn Reasoning. Figure 12 shows that ReFlow achieves final training rewards comparable to synchronous and asynchronous baselines for both Qwen3-8B and Qwen3-14B, while maintaining competitive early-stage progress. These results demonstrate that trajectory reuse does not compromise convergence under our replay-driven dataflow. More-

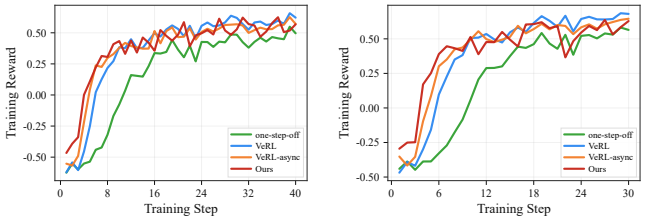

(a) Qwen3-8B on 32 GPUs (b) Qwen3-14B on 64 GPUs

Figure 12: Training reward with respect to step.

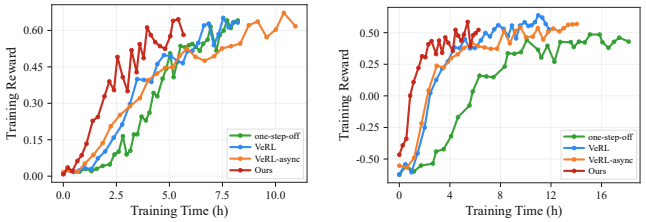

(a) Qwen3-4B on 8 GPUs (b) Qwen3-8B on 32 GPUs

Figure 13: Training reward with respect to wall clock time.

over, the Qwen3-8B results in Figure 13 (b) show that ReFlow reaches comparable reward levels substantially earlier, confirming that its system efficiency gains translate into shorter training time.

Multi-Turn Agentic Tasks. The benefits also extend to multi-turn workloads with environment interactions. For the Qwen3-4B coding agent, Figure 13 (a) shows that ReFlow improves reward substantially faster in wall-clock time, reaching reward levels comparable to the baselines within a shorter training period. For ReTool, Figure 11 (b) shows that ReFlow maintains reward progression comparable to VeRL and VeRL-async over training steps and consistently outperforms one-step-off throughout most of training. Despite fluctuations during optimization, reward continues to improve without sustained degradation, supporting the stability of our design in multi-turn agentic RL.

Together, these results show that ReFlow preserves stable learning across single-turn and multi-turn tasks while translating system throughput gains into faster time to target reward.

7.5 Performance Analysis

Resource efficiency of ReFlow. Figure 14 compares the average idle ratios of trainers and rollout workers across two model scales. We define the idle ratio as the fraction of each step time during which a worker remains idle, and report the average across all steps in the convergence experiments. Compared with VeRL-async, ReFlow reduces trainer idle ratios from 55.8% to 15.5% for Qwen3-8B on 32 GPUs and from 46.5% to 15.8% for Qwen3-14B on 64 GPUs. Rollout worker idle ratios are further reduced from 3.5% to 0.2% and from 8.4% to 0.4%, respectively.

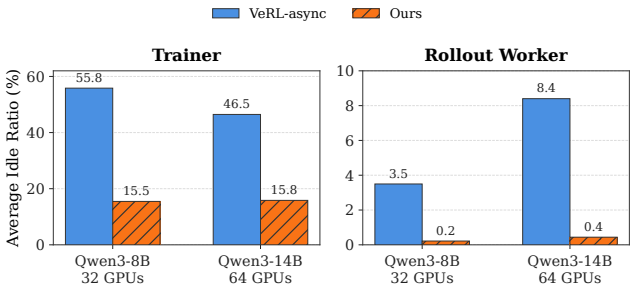

Figure 14: Trainer and rollout worker idle ratio comparison.

Table 2: Dynamic resource placement overhead in ReFlow.

Rollout : Train Setting	Direction	Cost (s)
32 GPUs (16:16 → 24:8)	train to rollout	5.72
32 GPUs (24:8 → 16:16)	rollout to train	9.20

Table 3: Effect of adaptive generation scheduling in ReFlow.

Setting	Average Generation Throughput (token/s)
w/o adaptive scheduling	1585.1 (+0.0%)
w/ adaptive scheduling	1749.8 (+10.4%)

These results show that ReFlow effectively reduces cross-stage waiting: replay-driven decoupling mitigates trainer stalls from data dependency, while keeping rollout workers nearly continuously utilized. This improved resource utilization directly contributes to the throughput gains in Figure 10.

Overhead of dynamic resource placement. Table 2 reports the overhead of reallocating GPUs between rollout and training on a 32-GPU setup. Reconfiguring from 16:16 to 24:8 and back to 16:16 takes only 5.72s and 9.20s, respectively. These results show that ReFlow can adapt its rollout–training resource split at runtime with only lightweight overhead. This is enabled by the Unified AsyncHybrid Worker design, which allows workers to switch roles without costly system-level reinitialization, full worker teardown, or long execution stalls. As a result, ReFlow can support adaptive resource placement during long-running RL workloads while keeping rebalancing overhead small relative to the overall training process.

Efficiency of adaptive generation scheduling. Table 3 evaluates the impact of adaptive generation scheduling when training Qwen3-4B on 4 GPUs. Compared with the case without generation scheduling, enabling the adaptive scheduling improves generation throughput from 1585.1 to 1749.8 tokens per second, achieving a 10.4% gain. This reveals effectively improved runtime generation efficiency by better coordinating generation workloads, without introducing additional resource requirements.

8 Conclusion

We present ReFlow, a fully decoupled and adaptive framework for efficient LLM RL post-training. Introducing a dual replay-driven cyclic dataflow, ReFlow decouples rollout generation and training while enabling trajectory reuse, realizing co-optimization of system efficiency and training stability. Combined with efficient dynamic resource placement and adaptive resource scheduling, ReFlow effectively improves system efficiency under dynamic and evolving RL workloads. Experimental results on up to 128 GPUs show that ReFlow achieves up to 2.06× throughput gains while maintaining comparable or even better model performance. By substantially reducing the cost of policy-improvements, ReFlow opens up opportunities for more compute-intensive and iterative post-training paradigms, such as recursive self-improvement.

Availability

We are currently organizing the code and awaiting the necessary approvals. The code will be released once the approval process is complete.

References

[1] Charles Arnal, Vivien Cabannes, Taco Cohen, Julia Kempe, and Remi Munos. Efficient rl training for llms with experience replay, 2026.

[2] Paul Christiano, Jan Leike, Tom B. Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep reinforcement learning from human preferences, 2023.

[3] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.

[4] DeepSeek-AI, Aixin Liu, Aoxue Mei, Bangcai Lin, Bing Xue, Bingxuan Wang, Bingzheng Xu, Bochao Wu, Bowei Zhang, Chaofan Lin, et al. Deepseek-v3.2: Pushing the frontier of open large language models, 2025.

[5] Jiazhan Feng, Shijue Huang, Xingwei Qu, Ge Zhang, Yujia Qin, Baoquan Zhong, Chengquan Jiang, Jinxin Chi, and Wanjun Zhong. Retool: Reinforcement learning for strategic tool use in llms, 2025.

[6] Wei Fu, Jiaxuan Gao, Xujie Shen, Chen Zhu, Zhiyu Mei, Chuyi He, Shusheng Xu, Guo Wei, Jun Mei, Jiashu Wang, Tongkai Yang, Binhang Yuan, and Yi Wu. Areal: A large-scale asynchronous reinforcement learning system for language reasoning, 2025.

[7] Scott Fujimoto, Herke van Hoof, and David Meger. Addressing function approximation error in actor-critic methods. In *International Conference on Machine Learning*, pages 1587–1596, 2018.

[8] Jiaxuan Gao, Wei Fu, Minyang Xie, Shusheng Xu, Chuyi He, Zhiyu Mei, Banghua Zhu, and Yi Wu. Beyond ten turns: Unlocking long-horizon agentic search with large-scale asynchronous rl, 2025.

[9] Wei Gao, Yuheng Zhao, Dakai An, Tianyuan Wu, Lunxi Cao, Shaopan Xiong, Ju Huang, Weixun Wang, Siran Yang, Wenbo Su, Jiamang Wang, Lin Qu, Bo Zheng, and Wei Wang. Rollpacker: Mitigating long-tail rollouts for fast, synchronous rl post-training, 2025.

[10] GLM-5-Team, Aohan Zeng, Xin Lv, Zhenyu Hou, Zhengxiao Du, Qinkai Zheng, Bin Chen, Da Yin, Chendi Ge, Chenghua Huang, et al. Glm-5: from vibe coding to agentic engineering, 2026.

[11] Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, Runxin Xu, Ruoyu Zhang, Shirong Ma, Xiao Bi, et al. Deepseek-r1 incentivizes reasoning in llms through reinforcement learning. *Nature*, 645(8081):633–638, September 2025.

[12] Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, and Sergey Levine. Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor. In *International Conference on Machine Learning*, pages 1861–1870, 2018.

[13] Zhenyu Han, Ansheng You, Haibo Wang, Kui Luo, Guang Yang, Wenqi Shi, Menglong Chen, Sicheng Zhang, Zeshun Lan, Chunshi Deng, Huazhong Ji, Wenjie Liu, Yu Huang, Yixiang Zhang, Chenyi Pan, Jing Wang, Xin Huang, Chunsheng Li, and Jianping Wu. Asyncflow: An asynchronous streaming rl framework for efficient llm post-training, 2025.

[14] Dan Hendrycks, Steven Basart, Saurav Kadavath, Mantas Mazeika, Akul Arora, Ethan Guo, Collin Burns, Samir Puranik, Horace He, Dawn Song, and Jacob Steinhardt. Measuring coding challenge competence with apps. *NeurIPS*, 2021.

[15] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the ACM SIGOPS 29th Symposium on Operating Systems Principles*, 2023.

[16] Haoyang Li, Sheng Lin, Fangcheng Fu, Yuming Zhou, Xiaodong Ji, Yanfeng Zhao, Lefeng Wang, Jie Jiang, and Bin Cui. Unleashing efficient asynchronous rl post-training via staleness-constrained rollout coordination, 2026.

[17] Siheng Li, Zhanhui Zhou, Wai Lam, Chao Yang, and Chaochao Lu. Repo: Replay-enhanced policy optimization, 2025.

[18] Yingru Li, Jiacai Liu, Jiawei Xu, Yuxuan Tong, Ziniu Li, Qian Liu, and Baoxiang Wang. Trust region masking for long-horizon llm reinforcement learning. *arXiv preprint arXiv:2512.23075*, 2025.

[19] Ziniu Li, Tian Xu, Yushun Zhang, Zhihang Lin, Yang Yu, Ruoyu Sun, and Zhi-Quan Luo. Remax: a simple, effective, and efficient reinforcement learning method for aligning large language models. In *Proceedings of the 41st International Conference on Machine Learning*, ICML'24. JMLR.org, 2024.

[20] Timothy P. Lillicrap, Jonathan J. Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver, and Daan Wierstra. Continuous control with deep reinforcement learning. *arXiv preprint arXiv:1509.02971*, 2015.

[21] Long-Ji Lin. Self-improving reactive agents based on reinforcement learning, planning and teaching. *Machine Learning*, 8:293–321, 1992.

[22] Jiacai Liu, Yingru Li, Yuqian Fu, Jiawei Wang, Qian Liu, and Yu Shen. When speed kills stability: Demystifying RL collapse from the training-inference mismatch, September 2025.

[23] meituan-search. Recipe: One Step Off Policy Async Trainer. https://verl.readthedocs.io/en/latest/advance/one_step_off.html, July 2025. Last updated: July 17, 2025.

[24] meituan-search. Recipe: Fully Async Policy Trainer. https://verl.readthedocs.io/en/latest/advance/fully_async.html, February 2026. Last updated: February 5, 2026.

[25] MiniMax. Forge: Scalable agent rl framework and algorithm. https://www.minimax.io/news/forge-scalable-agent-rl-framework-and-algorithm, February 2026. Technical blog.

[26] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, Stig Petersen, Charles Beattie, Amir Sadik, Ioannis Antonoglou, Helen King, Dharshan Kumaran, Daan Wierstra, Shane Legg, and Demis Hassabis. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533, 2015.

[27] Philipp Moritz, Robert Nishihara, Stephanie Wang, Alexey Tumanov, Richard Liaw, Eric Liang, Melih Elibol, Zongheng Yang, William Paul, Michael I. Jordan, and Ion Stoica. Ray: A distributed framework for emerging AI applications. In *13th USENIX Symposium on Operating Systems Design and Implementation (OSDI 18)*, pages 561–577, Carlsbad, CA, October 2018. USENIX Association.

[28] Reiichiro Nakano, Jacob Hilton, Suchir Balaji, Jeff Wu, Long Ouyang, Christina Kim, Christopher Hesse, Shantanu Jain, Vineet Kosaraju, William Saunders, Xu Jiang, Karl Cobbe, Tyna Eloundou, Gretchen Krueger, Kevin Button, Matthew Knight, Benjamin Chess, and John Schulman. Webgpt: Browser-assisted question-answering with human feedback, 2022.

[29] Michael Noukhovitch, Shengyi Huang, Sophie Xhonneux, Arian Hosseini, Rishabh Agarwal, and Aaron Courville. Asynchronous rlhf: Faster and more efficient off-policy rl for language models, 2025.

[30] OpenAI, Aaron Jaech, Adam Kalai, Adam Lerer, Adam Richardson, Ahmed El-Kishky, Aiden Low, Alec Helyar, Aleksander Madry, Alex Beutel, et al. Openai o1 system card, 2026.

[31] Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback. In *Proceedings of the 36th International Conference on Neural Information Processing Systems*, NIPS '22, Red Hook, NY, USA, 2022. Curran Associates Inc.

[32] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Köpf, Edward Yang, Zach DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. *PyTorch: an imperative style, high-performance deep learning library*. Curran Associates Inc., Red Hook, NY, USA, 2019.

[33] Ruoyu Qin, Weiran He, Weixiao Huang, Yangkun Zhang, Yikai Zhao, Bo Pang, Xinran Xu, Yingdi Shan, Yongwei Wu, and Mingxing Zhang. Seer: Online context learning for fast synchronous llm reinforcement learning, 2025.

[34] Tom Schaul, John Quan, Ioannis Antonoglou, and David Silver. Prioritized experience replay. In *International Conference on Learning Representations*, 2016.

[35] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms, 2017.

[36] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. Deepseekmath: Pushing the limits of mathematical reasoning in open language models, 2024.

[37] Guangming Sheng, Yuxuan Tong, Borui Wan, Wang Zhang, Chaobo Jia, Xibin Wu, Yuqi Wu, Xiang Li, Chi Zhang, Yanghua Peng, Haibin Lin, Xin Liu, and Chuan Wu. Laminar: A scalable asynchronous rl post-training framework, 2025.

[38] Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng, Haibin Lin, and Chuan Wu. Hybridflow: A flexible and efficient rlhf framework. In *Proceedings of the Twentieth European Conference on Computer Systems*, EuroSys '25, page 1279–1297. ACM, March 2025.

[39] Nisan Stiennon, Long Ouyang, Jeff Wu, Daniel M. Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford, Dario Amodei, and Paul Christiano. Learning to summarize from human feedback. In *Proceedings of the 34th International Conference on Neural Information Processing Systems*, NIPS '20, Red Hook, NY, USA, 2020. Curran Associates Inc.

[40] Kimi Team, Yifan Bai, Yiping Bao, Guanduo Chen, Jiahao Chen, Ningxin Chen, Ruijue Chen, Yanru Chen, Yuankun Chen, Yutian Chen, et al. Kimi k2: Open agentic intelligence, 2025.

[41] Kimi Team, Angang Du, Bofei Gao, Bowei Xing, Changjiu Jiang, Cheng Chen, Cheng Li, Chenjun Xiao, Chenzhuang Du, Chonghua Liao, et al. Kimi k1.5: Scaling reinforcement learning with llms. *arXiv preprint arXiv:2501.12599*, 2025.

[42] Xucong Wang, Zhe Zhao, Liheng Yu, Di Wu, Xiaofeng Cao, and Pengkun Wang. Didpo: Diff-in-diff policy optimization for coding agent training, 2026.

[43] Bo Wu, Sid Wang, Yunhao Tang, Jia Ding, Eryk Helenowski, Liang Tan, Tengyu Xu, Tushar Gowda, Zhengxing Chen, Chen Zhu, Xiaocheng Tang, Yundi Qian, Beibei Zhu, and Rui Hou. Llamarl: A distributed asynchronous reinforcement learning framework for efficient large-scale llm training, 2025.

[44] Zhiheng Xi, Xin Guo, Yang Nan, Enyu Zhou, Junrui Shen, Wenxiang Chen, Jiaqi Liu, Jixuan Huang, Zhihao Zhang, Honglin Guo, Xun Deng, Zhikai Lei, Miao Zheng, Guoteng Wang, Shuo Zhang, Peng Sun, Rui Zheng, Hang Yan, Tao Gui, Qi Zhang, and Xuanjing Huang. Bapo: Stabilizing off-policy reinforcement learning for llms via balanced policy optimization with adaptive clipping, 2025.

[45] Haofeng Xu, Junwei Su, Yukun Tian, Lansong Diao, Zhengping Qian, and Chuan Wu. Gac: Stabilizing asynchronous rl training for llms via gradient alignment control, 2026.

[46] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report, 2025.

[47] Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaohong Liu, Lingjun Liu, Xin Liu, Haibin Lin, Zhiqi Lin, Bole Ma, Guangming Sheng, Yuxuan Tong, Chi Zhang, Mofan Zhang, Wang Zhang, Hang Zhu, Jinhua Zhu, Jiaze Chen, Jiangjie Chen, Chengyi Wang, Hongli Yu, Yuxuan Song, Xiangpeng Wei, Hao Zhou, Jingjing Liu, Wei-Ying Ma, Ya-Qin Zhang, Lin Yan, Mu Qiao, Yonghui Wu, and Mingxuan Wang. Dapo: An open-source llm reinforcement learning system at scale, 2025.

[48] Yanli Zhao, Andrew Gu, Rohan Varma, Liang Luo, Chien-Chin Huang, Min Xu, Less Wright, Hamid Shojanazeri, Myle Ott, Sam Shleifer, Alban Desmaison, Can Balioglu, Pritam Damania, Bernard Nguyen, Geeta Chauhan, Yuchen Hao, Ajit Mathews, and Shen Li. Pytorch fsdp: Experiences on scaling fully sharded data parallel. *Proc. VLDB Endow.*, 16(12):3848–3860, August 2023.

[49] Chujie Zheng, Shixuan Liu, Mingze Li, Xiong-Hui Chen, Bowen Yu, Chang Gao, Kai Dang, Yuqiong Liu, Rui Men, An Yang, Jingren Zhou, and Junyang Lin. Group sequence policy optimization, 2025.

[50] Haizhong Zheng, Jiawei Zhao, and Beidi Chen. Prosperity before collapse: How far can off-policy rl reach with stale data on llms?, 2025.

[51] Yinmin Zhong, Zili Zhang, Xiaoniu Song, Hanpeng Hu, Chao Jin, Bingyang Wu, Nuo Chen, Yukun Chen, Yu Zhou, Changyi Wan, Hongyu Zhou, Yimin Jiang, Yibo Zhu, and Daxin Jiang. Streamrl: Scalable, heterogeneous, and elastic rl for llms with disaggregated stream generation, 2025.

[52] Yuzhen Zhou, Jiajun Li, Yusheng Su, Gowtham Ramesh, Zilin Zhu, Xiang Long, Chenyang Zhao, Jin Pan, Xiaodong Yu, Ze Wang, Kangrui Du, Jialian Wu, Ximeng Sun, Jiang Liu, Qiaolin Yu, Hao Chen, Zicheng Liu, and Emad Barsoum. April: Active partial rollouts in reinforcement learning to tame long-tail generation, 2025.