arXiv:2610.26437v1 [stat.ML] 19 Oct 2026

R2-GP: Variance-Calibrated Priors for Additive Gaussian Processes

Javier Enrique Aguilar[2], Soham Mukherjee[1], and David Kohns[1,*]

[1]Aalto University, Finland
[2]TU Dortmund, RC TRUST (UA Ruhr), Germany
[*]Corresponding author: david.kohns@aalto.fi

Abstract

We develop R2-GP, an R2D2-type prior for Gaussian process (GP) generalised additive models (GAMs) that calibrates explained variance (R^2) for Matérn covariance kernels. The function variance for standard Matérn kernels depends jointly on the effect magnitude and length-scale. As such, smooth and rough components with identical amplitudes contribute very different amounts of variation. With independent priors on component amplitudes, this variation results in variance overestimation, especially in higher dimensions. We overcome this by taking direct control over each component's expected contribution to prior explained variance. We do this by rescaling each component based on its length scale and observed inputs, and modelling each rescaled variance as a share of the function variance induced by a global R^2 prior. This allows R2-GP to scale well to many components. To also scale to large sample sizes, we extend R2-GP to Hilbert space approximation methods, a framework for kernels exploiting their spectral decompositions for fast and efficient inference. Our simulations and real-world case studies demonstrate superior generalisation and parameter recovery over independent variance priors.

1 Introduction

Gaussian processes (GPs) are a class of probabilistic methods for modelling non-linear functions. They do so by assuming a multivariate Gaussian prior distribution over a finite set of these functions and are specified through a mean function and a covariance kernel (Williams and Rasmussen, 1995). Common choices of covariance kernel involve the Matérn class of functions (Rasmussen and Williams, 2006), defined using a length scale parameter and a magnitude parameter (also called GP marginal variance). The former controls the correlation decay rate between two arbitrary inputs while the latter controls prior variance of the functions at a single input.

GP priors have been specifically used in generalised additive models (GAMs; Hastie and Tibshirani, 1986, 1990) as an approach to model smooth functions (Wood, 2017). In GAMs, a response $y \in \mathbb{R}^N$ is modelled using multiple covariates $X \in \mathbb{R}^{N \times D}$ such that

$$y_i = \mu_0 + \sum_{d=1}^{D} f_d(x_{id}) + \varepsilon_i, \qquad \varepsilon_i \sim N(0, \sigma^2), \tag{1}$$

for $i = 1, \dots, N$, where μ_0 is the intercept, ε_i is the residual error and f_d is a GP specific to the d-th covariate. Additive GPs arise as a special case of the general class of models presented in Duvenaud et al. (2011).

In GAMs, modelling a response with many covariates introduces a variance explosion problem familiar from high-dimensional linear regression. If each of the D components receives an independent variance prior with positive mass on non-zero values, the total signal variance accumulates as D grows. Consequently, the implied prior on the proportion of explained variance R^2 pathologically concentrates near one. This may lead to overfitting and contradict the modeller's domain knowledge about R^2, rendering the model predictively inconsistent (Riha et al., 2026).

Preventing this variance accumulation requires explicitly controlling the model's total variance. In regression contexts, a natural approach is to place a prior on R^2 and partition it across the components, analogous to the R2D2 prior for linear models (Zhang et al., 2022). The mapping between a component's magnitude parameter and the observed function variance one wishes to control is however blurred by the simultaneous impact of the length scale and the input design. Because of this entanglement, smooth and rough components with identical magnitudes contribute vastly different amounts to observed function variation (Figure 1).

We overcome this structural issue with our R2-GP prior. By rescaling each component with a normaliser which is derived from its length scale and the input design, R2-GP forces the component's magnitude to exactly match its expected variation across the data. This calibration allows a global R^2 prior which is partitioned with a simplex prior that forces components to compete for a share of the variance budget.

A drawback of this normalisation is the computational burden it adds to posterior inference. Since the normaliser depends on the length scale, it must be recomputed at every gradient evaluation when using MCMC methods like Hamiltonian Monte Carlo (Neal, 2011; Hoffman and Gelman, 2014). For an exact GP, this requires $O(N^2)$, added to the standard $O(N^3)$ cost of factorising the exact covariance matrix.

We solve this scalability issue by extending the R2-GP prior to Hilbert space approximation methods (HSGPs) (Solin and Särkkä, 2020; Riutort-Mayol et al., 2023; Mukherjee et al., 2026). HSGPs exploit the spectral decomposition of the kernel using basis functions M for large sample scalability and have been extended to GAMs efficiently (Timonen and Lähdesmäki, 2025). With the extension to HSGPs based on $M << N$ basis functions, computing the normaliser simplifies to a single pre-computation of $O(NM)$ and $O(M)$.

Code to replicate the experiments is available at https://github.com/davkoh/r2-gp.

Contributions.

1. We develop R2-GP by extending the R2D2 family of priors to GPs for GAMs. We derive the necessary normalising approach of the kernel such that a prior over the GP's R^2 maps analytically to its expected sample variance. A Dirichlet decomposition of R^2 into contributions of its GAM components allows its application to high dimensions.

2. We derive R2-GP prior for the Hilbert space approximation and show that the normaliser admits an exact finite-sample expression requiring $O(NM)$ preparation and $O(M)$ and scales to high dimensions.

3. Through simulations and real-world data, we show that the R2-GP prior outperforms independent priors for additive HSGPs in prediction and parameter recovery and sampling efficiency, and is competitive with alternate shrinkage priors.

Related Work. Priors placed on R^2, since the original Zhang et al. (2022), have been extended to multilevel models (Aguilar and Bürkner, 2023), generalised linear mixed models (Yanchenko et al., 2025), autoregressions (Kohns et al., 2025), spatial models (Yanchenko et al., 2024) and grouped

predictors (Aguilar et al., 2025). The present paper takes the variance allocation mechanism from that literature and addresses the question it leaves on how to extend this to GPs.

Closest to our normalisation approach is the standardisation procedure of Ferrari and Ventrucci (2025), who scale variance partitioning priors in latent Gaussian models by the variance of a zero-mean-constrained effect over an assumed covariate distribution, evaluated at fixed values of unknown kernel parameters such as the range of a Matérn process. Our normaliser is instead a function of the length scale and is recomputed at every evaluation inside the sampler, which is what makes the $O(M)$ identity of Lemma A.1 particularly efficient. Penalised complexity priors for Gaussian random fields (Simpson et al., 2017; Fuglstad et al., 2019) shrink towards a simpler base model, whereas the normalisation here fixes the units in which the scale parameter is expressed, and the two could in principle be combined.

The remainder of the paper is organised as follows. Section 2 derives the normalised additive prior and its properties, among them the implied prior on explained variance. Section 3 presents our simulation studies and Section 4 demonstrates a real-world application. Proofs and additional results are collected in Appendix.

2 Methods

In our framework, the d-th column of the design matrix X holds the observed values $x_{1d}, \ldots, x_{Nd}$ of covariate d. We use GAM GP prior with independent processes per covariate,

$$f_d = \alpha_d \, g_d, \qquad g_d \sim \mathcal{GP}\big(0,\, k(\cdot,\cdot\,;\rho_d)\big), \tag{2}$$

where k is a stationary kernel of the Matérn family, which includes the squared exponential (SE) kernel as a limiting case, with $k(x, x; \rho_d) = 1$, so that g_d has unit marginal variance. Parameter α_d^2 is the *magnitude* of the component, which equals its pointwise *marginal variance*, $\mathrm{Var}(f_d(x)) = \alpha_d^2$ at every input.[1] In contrast, the *signal variance* is the expected variation that the function exhibits across the inputs, $\mathbb{E}_{f_d}[\mathrm{Var}_x(f_d(x) \mid f_d)]$. We denote by $\mathbb{E}_x$ and Var_x the expectation and variance operators with respect to x. In the following, we first derive the mathematical disconnect between these two variances, then provide the necessary concepts of the Hilbert space (HS) approximation to finally define an efficient finite-sample normaliser, which we ultimately use to construct the R2-GP prior.

2.1 Calibrating GPs for Explained Variance

In linear regression, placing a prior directly on the proportion of explained variance $R^2 = \mathrm{Var}(f(x))/\mathrm{Var}(y)$ provides an interpretable way to specify regularisation on variance parameters of regression coefficients. For GAM GPs the R^2 identity is,

$$R^2_{\text{marg}} = \frac{\sum_{d=1}^{D} \alpha_d^2}{\sum_{d=1}^{D} \alpha_d^2 + \sigma^2}, \tag{3}$$

where a prior π_{R^2} maps to prior π_{α_d}, conditional on σ^2. We define Equation 3 as the marginal R^2, since it's marginal of the effect of the design and any other kernel parameters. The magnitude α_d^2, however is not equal to the signal variance. To see this, consider the total variance of component

[1] The GP literature often calls α_d^2 itself the signal variance (Rasmussen and Williams, 2006). Section 2.1 shows that the two differ.

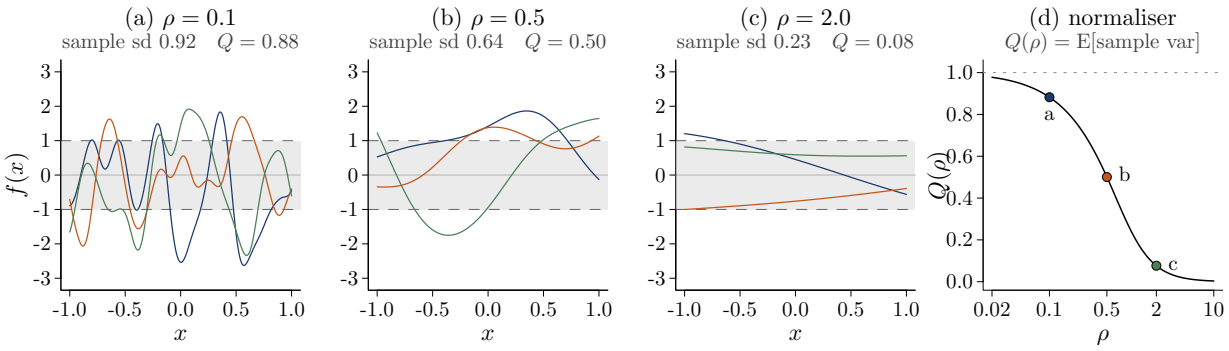

Figure 1: GP Realisations and Normaliser $Q(\rho)$. Realisations of a zero-mean GP with an SE kernel and unit marginal variance on $x \in [-1, 1]$. The shaded band is ± 1 marginal standard deviation. Panels (a) to (c) hold the marginal variance fixed and vary the length scale, and report the sample standard deviation averaged over 4000 realisations. As ρ grows, the signal variance falls even though the marginal variance is unchanged. Panel (d) shows the resulting normaliser $Q(\rho)$ of (5).

$f_d(x)$,

$$\mathrm{Var}(f_d(x)) = \alpha_d^2 = \underbrace{\mathbb{E}_{f_d}\big[\mathrm{Var}_x(f_d(x) \mid f_d)\big]}_{\text{signal variance}} + \underbrace{\mathrm{Var}_{f_d}\big(\mathbb{E}_x[f_d(x) \mid f_d]\big)}_{\text{variance of the level}}. \tag{4}$$

The first term is the expected variance of the function across the observed inputs, while the second term captures the variance of the average height of function realisations (See Appendix A.1).

The second term is precisely the point of departure between a correlated GP and a linear regression. If $f(x)$ were a linear term with mean-centred inputs, the level, $\mathbb{E}_x[f_d(x) \mid f_d]\big)$, is deterministically zero, such that marginal variance equals signal variance. In contrast, the average level of a realised GP acts as a random constant shift across all input observations. Because this shift is constant with respect to x, it does not contribute to the function's variation over the inputs (the signal variance), and the level itself is completely absorbed by the intercept μ_0 in (1).

The size of the level variance depends on the length scale ρ_d and on the input design. Denoting by $\bar{k}_d = \mathbb{E}_{x,x'}[k(x, x'; \rho_d)]$ the average kernel correlation between two independent inputs x and x', the level variance reduces to $\alpha_d^2 \bar{k}_d$, and the signal variance is therefore

$$\begin{aligned} \mathbb{E}_{f_d}\big[\mathrm{Var}_x(f_d(x) \mid f_d)\big] &= \alpha_d^2 Q_d(\rho_d), \\ Q_d(\rho_d) &= 1 - \bar{k}_d. \end{aligned} \tag{5}$$

For the Matérn family, whose correlations are positive and increase with the length scale, $Q_d(\rho_d) \in (0, 1)$ (unless all inputs coincide). As Figure 1 illustrates, smoother realisations (larger ρ_d) carry a larger part of their variance in the level, so that their signal variance falls well below the magnitude α_d^2.

Hence, we define the *realised* R^2 as the modelling target:

$$R^2_{\text{real}} = \frac{\sum_{d=1}^{D} \alpha_d^2 Q_d(\rho_d)}{\sum_{d=1}^{D} \alpha_d^2 Q_d(\rho_d) + \sigma^2}, \tag{6}$$

which, whenever $Q_d(\rho_d) < 1$, implies $R^2_{\text{real}} < R^2_{\text{marg}}$. Furthermore, because Q_d depends on ρ_d, the implied prior on explained variance depends on the length scales, and so does the share of each component in the total signal variance.

This motivates replacing (2) with a normalised prior,

$$f_d = \frac{\alpha_d}{\sqrt{Q_d(\rho_d)}}\, g_d. \tag{7}$$

Given ρ_d, this is a deterministic rescaling that leaves f_d Gaussian and its correlation structure unchanged. By (5), it makes the signal variance of f_d equal to α_d^2 for every length scale, while the marginal variance becomes $\alpha_d^2/Q_d(\rho_d)$. Consequently, R^2_{marg} and R^2_{real} agree, and, as in the linear case, we can place a prior on a global R^2 and partition it across components with a Dirichlet distribution,

$$\begin{aligned} R^2 \sim \text{Beta}\big(\mu_R\phi_R,\ (1-\mu_R)\phi_R\big), \qquad \alpha^2 = \sigma^2\frac{R^2}{1-R^2}, \\ \alpha_d^2 = \psi_d\,\alpha^2, \qquad \psi \sim \text{Dirichlet}(a\mathbf{1}_D), \end{aligned} \tag{8}$$

where $\mu_R \in (0,1)$ and $\phi_R > 0$ parameterise the Beta prior through its mean and precision (Aguilar and Bürkner, 2023). We refer to (7) together with (8) as the R2-GP prior. Its total signal variance is α^2, and the Dirichlet weight ψ_d is the share of component d, for any length scale value. Since the proportions in a Dirichlet distribution have strictly negative pairwise correlation, the GAM components intrinsically compete for their share of the total explained variance.

2.2 The Hilbert Space Representation

The Hilbert space (HS) approximation methods is based on the eigenvalue decomposition of the Gram matrix for a chosen kernel using its spectral density along with eigenvalues and eigenfunctions based on the GP inputs. Following Solin and Särkkä (2020), we use a compact cover $[-L_d, L_d]$ containing the observed values of covariate d. Based on the Dirichlet boundary conditions, we get closed form solutions for eigenvalues ω and eigenfunctions ϕ of the kernel such that

$$\begin{aligned} \phi_{dm}(x) &= \frac{1}{\sqrt{L_d}}\sin\big(\omega_{dm}(x+L_d)\big), \\ \omega_{dm} &= \frac{m\pi}{2L_d}, \qquad m = 1,\dots,M, \end{aligned} \tag{9}$$

where M is the number of basis functions. The matrix $\Phi_d \in \mathbb{R}^{N\times M}$ contains the eigenfunction values at the observed inputs, such that $[\Phi_d]_{im} = \phi_{dm}(x_{id})$, with $\phi_{dm} \in \mathbb{R}^N$ for its m-th column. We obtain the spectral density $S_{dm} = S(\omega_{dm};\rho_d) > 0$ of our chosen kernel evaluated at the m-th frequency, collected in $S_d = (S_{d1},\dots,S_{dM})^\mathsf{T}$. For example, the spectral density of the SE kernel $k(r;\rho) = \exp\{-r^2/(2\rho^2)\}$ is of the form $S(\omega;\rho) = \sqrt{2\pi}\,\rho\exp(-\rho^2\omega^2/2)$ obtained by solving the Fourier integral with respect to the inputs. The spectral forms of other Matérn kernels are provided in Riutort-Mayol et al. (2023).

Lemma 2.1 (Finite-design covariance)**.** *Let $z_d \sim N(0, I_M)$ and $g_d = \Phi_d\,\text{diag}(\sqrt{S_d})z_d \in \mathbb{R}^N$. Then, conditionally on X and ρ_d,*

$$g_d \sim N\big(0, \widetilde{K}_d\big), \qquad \widetilde{K}_d = \Phi_d\,\text{diag}(S_d)\Phi_d^\mathsf{T}, \tag{10}$$

and $\widetilde{K}_d$ is symmetric positive semi-definite with $\text{rank}\,\widetilde{K}_d \le \min(N, M)$.

The matrix $\widetilde{K}_d$ is the finite-basis approximation to the exact kernel matrix K_d with $[K_d]_{ij} = k(x_{id} - x_{jd};\rho_d)$. Under a sufficiently large boundary and increasing M it converges to K_d We use this convergence as established and refer to Solin and Särkkä (2020) for the underlying expansion. For practical guidance on choosing L_d and M, we follow Riutort-Mayol et al. (2023). Since $k(0) = 1$, the diagonal of $\widetilde{K}_d$ is approximately one, which is the marginal normalisation that Section 2.1 identified as insufficient for modelling the signal variance of the function.

2.3 The Finite-sample Normaliser

In HS form, we define the normaliser of component d as the expected sample variance of the unit-magnitude component g_d of Lemma 2.1, evaluated at the observed inputs,

$$Q_d(\rho_d) \;:=\; \mathbb{E}_z\big[\widehat{\mathrm{Var}}_N(g_d) \mid X, \rho_d\big], \tag{11}$$

where $\widehat{\mathrm{Var}}_N(v) = (N-1)^{-1}\sum_{i=1}^N (v_i - \bar{v})^2$ is the sample variance of a vector $v \in \mathbb{R}^N$ and $\mathbb{E}_z$ is the expectation over the coefficients z_d. This is the finite-sample counterpart of (5), with the sample variance over the observed inputs in place of the variance operator Var_x and the HSGP covariance $\widetilde{K}_d$ in place of the exact kernel. Evaluating (11) through the $N \times N$ matrix $\widetilde{K}_d$ would cost $O(N^2M)$ operations each time ρ_d changes. However, because the coefficients in z_d are independent with unit variance, the cross terms between different basis functions have expectation zero, and the normaliser simplifies to the sum of the sample variances of the basis functions, weighted by the spectral density (Lemma A.1 in Appendix A.2),

$$Q_d(\rho_d) = \sum_{m=1}^{M} S_{dm}\, \widehat{\mathrm{Var}}_N(\phi_{dm}). \tag{12}$$

This identity allows the R2-GP prior to scale well with N. The sample variances $\widehat{\mathrm{Var}}_N(\phi_{dm})$ of the basis columns depend on the design alone and not on any parameter, so they are precomputed once at cost $O(NM)$. Inside the sampler, the normaliser then reduces to an inner product of length M between the spectral density S_d and these precomputed variances, which requires only $O(M)$ work each time the length scale ρ_d changes.[2]

Expanding the sample variance in (11) gives the HSGP counterpart of the decomposition (4),

$$Q_d(\rho_d) = \frac{1}{N-1}\sum_{i=1}^{N}[\widetilde{K}_d]_{ii} \;-\; \frac{N}{N-1}\,\mathrm{Var}_z(\bar{g}_d), \tag{13}$$

where $\bar{g}_d$ is the sample mean of the realisation g_d. The first term aggregates the marginal variances, each approximately one, and the second is the prior variance of the realisation's level, which the intercept absorbs. Together they equal the average marginal variance minus the average covariance between two distinct observed inputs, the sample analogue of $1 - \bar{k}_d$ in (5).

2.4 The Normalised Additive Prior in HS Form

We construct the normalised HSGPs, using α_d^2 from (8) as follows

$$\begin{aligned} f_d &= \frac{\alpha_d}{\sqrt{Q_d(\rho_d)}}\, g_d = \Phi_d\big(c_d \odot z_d\big), \\ c_d &= \alpha_d\sqrt{\frac{S_d}{Q_d(\rho_d)}}, \qquad f = \mu_0 + \sum_{d=1}^{D} f_d, \end{aligned} \tag{14}$$

where $\odot$ is the elementwise product. The square root and division are elementwise. Together with (8), (14) is the HSGP form of the R2-GP prior.

[2] The normaliser is strictly positive whenever $M \geq 2$ and the observed inputs are not all equal (Lemma A.1), so that the division by $\sqrt{Q_d(\rho_d)}$ in Section 2.4 is well defined.

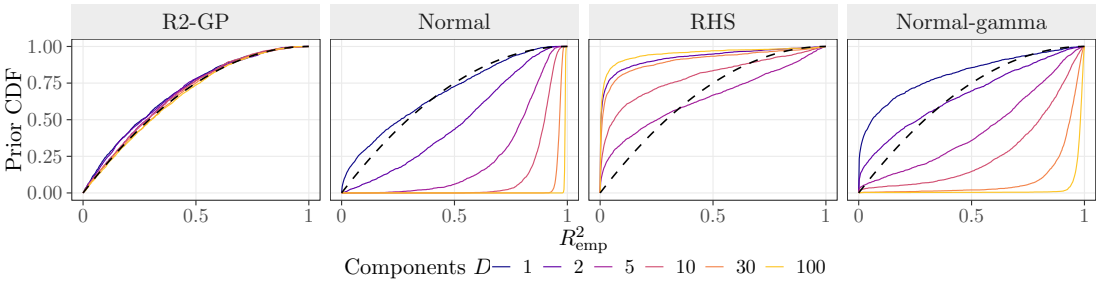

Figure 2: Implied Prior on R^2_{emp}. Empirical CDF of $R^2_{\text{emp}} = \widehat{\text{Var}}_N(f)/\big(\widehat{\text{Var}}_N(f) + \sigma^2\big)$ for additive HSGP GAMs under four priors on α_d, as the number of components D grows. CDFs are generated from 2000 prior draws: $N = 200$, $M = 100$, $\sigma = 1$ and $\rho_d \sim \text{InvGamma}(5, 4)$. These figures are insensitive to the prior ρ_d. RHS and normal-gamma are specified as in Section 3, with $p_0 = 5$. The dashed curve is the target $\text{Beta}(1, 2)$. Only R2-GP is stable in D.

As for the exact GP in Section 2.1, the signal variance of each component is then its allocated variance,

$$\mathbb{E}_z\big[\widehat{\text{Var}}_N(f_d) \mid X, \rho_d, \alpha_d\big] = \alpha_d^2, \qquad d = 1, \ldots, D, \tag{15}$$

and, for independent $z_1, \ldots, z_D$, the signal variances add up to the total budget,

$$\mathbb{E}_z\Big[\widehat{\text{Var}}_N\Big(\sum_{d=1}^{D} f_d\Big)\Big| X, \rho, \alpha\Big] = \sum_{d=1}^{D} \alpha_d^2 = \alpha^2. \tag{16}$$

Three comments are in order that highlight the usefulness of the R2 prior for GPs. Firstly, by writing $f_d = \Phi_d \beta_d^{\text{eff}}$ with $\beta_d^{\text{eff}} = c_d \odot z_d$, the proposed prior is equivalent to independent centred priors on the basis weights

$$\beta_{dm}^{\text{eff}} \sim N\Big(0,\ \alpha_d^2 S_{dm}/Q_d(\rho_d)\Big), \qquad m = 1, \ldots, M. \tag{17}$$

Computationally, the z_d remains standard normal and the scale is applied through a deterministic transformation, and thus, equivalent to the non-centred parameterisation (Papaspiliopoulos et al., 2007) which is well known to keep the posterior geometry well conditioned when a signal variance is small.

Second, because the deterministic rescaling in (7) applies a single global scalar $1/\sqrt{Q_d(\rho_d)}$ to all basis functions of component d equally, it preserves the relative weighting of the spectral densities S_{dm} that govern the smoothness of the function. Furthermore, because this scalar multiplier is independent of the basis index m, it does not interact with the truncation of the series. Consequently, the normalised HSGP inherits the same convergence guarantees as the standard HSGP: as the number of basis functions M grows, the approximate prior converges to the exact normalised GP prior.

Third, the prior on the total signal variance does not change as components are added. By (16) and (8), it equals $\sigma^2 R^2/(1 - R^2)$ for every D, so that its prior is the one placed on R^2. If instead each α_d has a prior second moment s^2 that is free of D, as under independent priors, it has prior mean Ds^2. Figure 2 shows the consequence for prior draws for the empirical proportion of explained variance, $R^2_{\text{emp}} = \widehat{\text{Var}}_N(f)/\big(\widehat{\text{Var}}_N(f) + \sigma^2\big)$ under different priors. Under the R2-GP prior with $R^2 \sim \text{Beta}(1, 2)$ and $a = 1$ (panel R2-GP), its implied prior stays close to the target for every D between 1 and 100. Under independent half-normal priors with unit scale, the median rises from

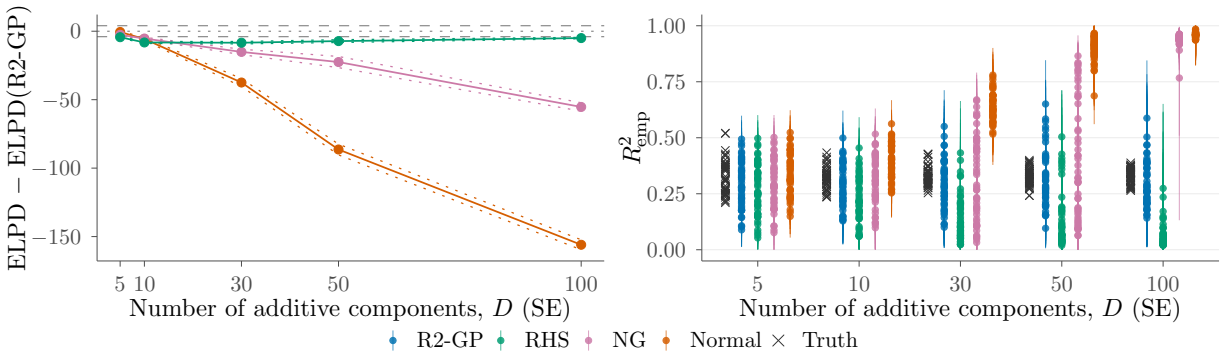

Figure 3: Dense DGP, SE Kernel, $\rho_0 = 0.3$. Left: mean Δelpd over 50 simulations (solid) with ± 1 Monte Carlo standard error (dotted). Dashed lines mark ± 4 and negative values favour R2-GP. Right: posterior means (points) and central 90% intervals (bars) of R^2_{emp} for each simulation. Crosses mark the realised R^2_{true}.

0.26 at $D = 1$ to 0.99 at $D = 100$. As competing shrinkage priors, which are further introduced in Section 3, the regularised horseshoe (Piironen and Vehtari, 2017) and normal-gamma (Griffin and Brown, 2010) may shrink each α_d but leave the prior on their total to shrink or grow with D.

Notice that calibration holds for $\widehat{\text{Var}}_N(f)$ but not for the empirical proportion of explained variance, which is a nonlinear function of it. The map $t \mapsto t/(t+\sigma^2)$ is strictly concave on $t > 0$, so Jensen's inequality gives

$$\mathbb{E}\left[R^2_{\text{emp}} \mid X, \sigma, R^2\right] \;\leq\; \frac{\alpha^2}{\alpha^2 + \sigma^2} \;=\; R^2, \tag{18}$$

with strict inequality whenever $\widehat{\text{Var}}_N(f)$ is non-degenerate. This accounts for the slight shift to the left for small D in Figure 2 for the R2-GP prior. This is not a deficiency of the proposed prior, since no choice of Q_d calibrates both $\mathbb{E}[\widehat{\text{Var}}_N(f)]$ and $\mathbb{E}[h(\widehat{\text{Var}}_N(f))]$ for a nonlinear h.

Finally, to ensure the component's prior expected marginal variance remains finite, the prior on the length scale ρ_d must be sufficiently light-tailed. As $\rho_d \to \infty$, the normaliser for the exact kernel decays such that $1/Q_d(\rho_d) \asymp \rho_d^\gamma$, where $\gamma = 1$ for the exponential kernel and $\gamma = 2$ for smoother Matérn variants (Rasmussen and Williams, 2006) . Therefore, the exact marginal variance $\alpha_d^2/Q_d(\rho_d)$ has a finite prior expectation if and only if $\mathbb{E}[\rho_d^\gamma] < \infty$. While the HSGP approximation inherently bounds this variance for finite M, preserving the properties of the exact-kernel limit rules out heavy-tailed priors, such as an Inverse-Gamma$(0.1, 0.1)$ distribution.

3 Simulation study

We assess finite-sample predictive and parameter recovery performance of the R2-GP prior, that is, (8) together with (14), as the number of GAM components D increases. The competitors are priors on the component scales that carry popular shrinkage approaches over to (1). These have not yet been considered for HSGP GAMs to our knowledge, but pose straightforward extensions based off our methodology.

3.1 Setup

Data Generation. We generate data from (1) with $\mu_0 = 0$ and $\sigma = 1$. Each f_d is a draw from the normalised prior (14) with a common length scale $\rho_d = \rho_0$ and signal variance

$$\alpha_d^2 = \frac{\sigma^2}{D} \frac{R_0^2}{1 - R_0^2}, \qquad R_0^2 = \frac{1}{3}, \tag{19}$$

for $d = 1, \ldots, D$. By (16), the total signal variance is $\sigma^2 R_0^2/(1 - R_0^2) = 1/2$ for every D, while the signal variance of a single component decays at rate D^{-1}. We call (19) the dense data-generating process (DGP). We also consider a sparse DGP in Appendix B.2, in which the first five components share the same total signal variance and $\alpha_d = 0$ for $d > 5$. We vary $D \in \{5, 10, 30, 50, 100\}$, $\rho_0 \in \{0.3, 0.5, 0.7\}$ and the kernel choice that includes SE, Matérn 3/2 and Matérn 5/2. For each condition, we simulate 50 replications. A data set has $N = 120$ training and 500 test observations. Inputs are iid uniform on $[-1, 1]$. Following Riutort-Mayol et al. (2023), we set $L_d = \max(1.5,\ 4.5\rho_0)$ and $M = \lceil 1.2\,\kappa L_d/\rho_0 \rceil$ with $\kappa = 1.75$, 3.42 and 2.65 for the three kernels, so that $10 \leq M \leq 21$. The data are generated with $4M$ basis functions, so that the f is generated close as a close representation to the desired GP. Since (16) holds in expectation over z, the realised value R^2_{true}, that is, R^2_{emp} at the generating f and σ, varies across data sets around the nominal R_0^2.

Priors. All models use the normalised components (14) and differ only in the prior on $\alpha_1, \ldots, \alpha_D$. R2-GP is (8) with $\mu_R = 1/3$, $\phi_R = 3$ and $a = 1$, so that $R^2 \sim \text{Beta}(1, 2)$, as in Section 2.4. We compare it with three alternatives, where a superscript $+$ denotes truncation to $(0, \infty)$. (i) The *normal* prior, $\alpha_d \sim N^+(0, 1)$ independently, is the analogue of ridge regression and represents independent priors on the component scales. (ii) The regularised horseshoe (RHS; Piironen and Vehtari, 2017), $\alpha_d \mid \lambda_d, \tau, \zeta \sim N^+(0, \tau^2\tilde{\lambda}_d^2)$ with $\tilde{\lambda}_d^2 = \zeta^2\lambda_d^2/(\zeta^2 + \tau^2\lambda_d^2)$, $\lambda_d \sim t_3^+(0, 1)$, $\tau \sim C^+(0, \tau_0)$ and $\zeta^2 \sim \text{InvGamma}(2, 8)$, represents sparsity-favouring priors. We set $\tau_0 = p_0(D - p_0)^{-1}\sigma N^{-1/2}$ as in Piironen and Vehtari (2017), with a prior guess of $p_0 = \min(5, D - 1)$ active components. While $p_0 = 5$ is miscalibrated for the dense DGP, τ is nevertheless learned from the data, and p_0 is well calibrated for the sparse DGP. (iii) The normal-gamma prior (NG; Griffin and Brown, 2010), $\alpha_d \mid \xi_d \sim N^+(0, \sigma^2\xi_d)$ with $\xi_d \sim \text{Gamma}(\vartheta, \vartheta/\eta)$ in shape-rate form, $\vartheta \sim \text{Exp}(1)$ and $\eta \sim \text{InvGamma}(2, 1)$, represents global-local shrinkage priors that are popular in other domains. The remaining priors are common to all models, namely $\rho_d \sim N^+(\rho_0, 0.1^2)$, $\sigma \sim N^+(0, 1)$ and $\mu_0 \sim N(0, 2^2)$. The comparisons presented here isolate the effect of the prior on α_d. We show in Appendix B.3 that the qualitative behaviour of the competing priors is the same also without the normalisation.

Expected Behaviour. The total signal variance is independent of D under R2-GP, whereas its mean is D under normal, and $D\sigma^2$ under NG, respectively. Since the generating value is 1/2 throughout, expect normal and NG to overestimate the signal variance and to underestimate σ by a margin that grows with D, in line with Figure 2. RHS is a sparsity favouring prior which for $p_0 = 5$ has $\tau_0 = O(D^{-1})$, hence, decays faster than the generating $\alpha_d = O(D^{-1/2})$ which is expected to lead to excessive shrinkage for large D.

Evaluation. Predictive performance is measured by the log predictive density summed over the test observations, which estimates the expected log predictive density (elpd; Vehtari et al., 2017), relative to R2-GP,

$$\Delta\text{elpd}_m^{(s)} = \text{elpd}_m^{(s)} - \text{elpd}_{\text{R2-GP}}^{(s)}, \tag{20}$$

for prior m and replication s, so that negative values favour R2-GP. We report its mean over replications and, following Sivula et al. (2025) and McLatchie et al. (2025), regard $|\Delta\text{elpd}| > 4$ as a meaningful difference. For $\theta \in \{\alpha, \rho, \sigma\}$ we report the coverage of the central 90% posterior intervals (Tables 2, 5 and 8) and the relative root-mean-squared error, that is relative to the RMSE of the R2-GP (RMSE; Tables 3, 6 and 9)

Values above one indicate worse recovery than under R2-GP. Lastly, we compare the posterior of R^2_{emp} with R^2_{true}. All models are implemented in Stan (Carpenter et al., 2017) and fitted with four chains of 2000 iterations, with the half as warm-up at a target acceptance rate of 0.95. Figure 3 shows the predictive performance and the recovery of R^2_{true} for the SE kernel with $\rho_0 = 0.3$. We show figures for all other conditions in Appendix B. See Tables 4, 7 and 10 for sampling diagnostics.

3.2 Results

Predictive Performance. R2-GP attains the highest mean elpd in each of the 45 conditions, against every competitor. At $D = 5$ the priors are close, however normal and NG deteriorate as D grows. This is as a consequence of the lack of joint shrinkage of the GAM components, which leads to overestimation of R^2_{true}. RHS stays closest to R2-GP at higher D, with mean $\Delta\text{elpd} \in [-13.5, -4.0]$ for all $D \geq 10$. As expected, with $p_0 = 5$ RHS overshrinks as D increases, which can be seen as underestimation of R^2_{true} in Figure 3. Due to the small true variance shares as D increases, this results in less of a penalty in predictive performance than the overestimation of R^2_{true} seen for normal and NG.

Recovery. Recovery mirrors the predictive results for the dense DGP. R2-GP has the lowest or close to the lowest RMSE for (α, σ) in all conditions. Interval coverage is close to or above the nominal level, with its advantage widening with D. (Tables 2 and 3 for the SE kernel, Appendix B for the Matérn kernels).

Sampling Efficiency. R2-GP also improves the sampling efficiency markedly. On average over replications, its largest $\widehat{R}$ is at most 1.01 and its smallest bulk effective sample size is at least 448 in every condition. It has the fewest divergent transitions in all conditions and the shortest sampling time in all but one. At $D = 100$, by contrast, the normal prior and NG fail to converge. The results of the normal and NG priors at large D therefore reflect the prior and the sampler jointly. The RHS prior converges, with $\widehat{R} \leq 1.03$, but needs up to 1.6 times the sampling time.

Sparse Allocation. Under the sparse DGP (Appendix B.2), RHS with the calibration $p_0 = 5$ attains a higher mean elpd than R2-GP for $D \geq 30$, however by small margins. The normal prior underperforms as in the dense DGP, and NG does so at $D = 100$. Relative to RHS, R2-GP's RMSE is slightly lower for the non-zero α_d in all conditions, however higher for the true zero α_d. This is as expected, since for $a = 1$ the allocation ψ is uniform on the simplex, whereas concentration near zero requires a small a (Zhang et al., 2022; Aguilar and Bürkner, 2023; Kohns et al., 2025).

Summary. For dense signals R2-GP performs at least as well as every competing prior in terms of prediction and recovers component scales most accurately. In particular, our method recovers the component scales most accurately in all conditions, with a margin over the normal and NG priors that grows with D. Sampling wise, our prior is the most reliable prior to sample. The patterns are almost identical across the three kernels and the three length scales.

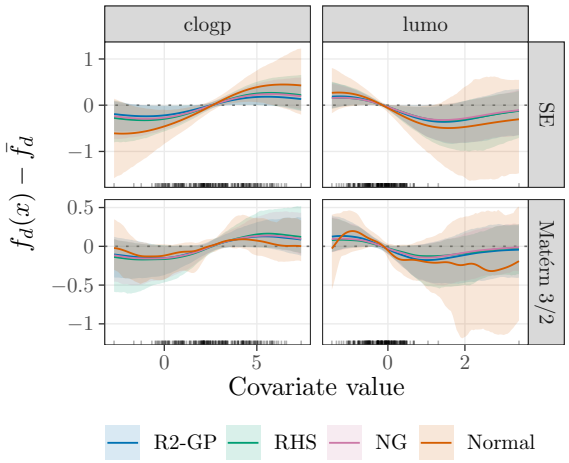

Figure 4: Covariate Effects, BloodBrain Data. Top 2 covariates are shown based on their contribution to explained variance.

4 Real-world Case Study

Data. In addition to our simulation scenarios, we demonstrate our proposed methods in analysing the BloodBrain data (Mente and Lombardo, 2005) from the R package *caret* (Kuhn, 2008). The dataset describes 208 chemical compounds using 134 molecular descriptors based on the compounds structure. The response y is the logarithm of the ratio of the compound's concentration in the brain to the blood, a standard procedure to measure how readily a chemical compound crosses the blood-brain barrier (with *logBBB* as response). For simplicity, we consider only continuous covariates. The data has $N = 208$ samples and $D = 101$ covariates.

Model Fit. We estimate the same set of models as in Section 3, except for the RHS prior for which $p_0 = 10$. We consider the SE and Matérn 3/2 kernels. We set boundary conditions with $c = 2$ and compute the number of basis functions accordingly. To compare predictive performance, the data is split into 80/20 train/test ratio with 10 random replications. Each fit uses 4 chains with 1000 warm-up and 1000 sample draws.

Results. The covariate *clogp* describing lipophilicity is a biologically relevant covariate for modelling *logBBB* (Mente and Lombardo, 2005) in the BloodBrain data. Additionally, we also investigate here *lumo* describing the energy component. The effects of these two covariates on response *logBBB* for all fitted models are shown in Figure 4. Across both kernels and all models, *logBBB* increases with an increase in *clogp*, indicating lipophilic compounds cross the blood-brain-barrier readily. For *lumo*, we see a decreasing effect on *logBBB*. The shrinkage priors agree while maintaining tight posterior credible intervals while the Normal model remains unstable. See Appendix C for the effect of more covariates. Our proposed R2-GP prior provides superior predictive performance as well as sampling speed and diagnostics (Appendix Table 20) compared to the Normal model. RHS and NG provide comparable predictive performance R2-GPs, but worse MCMC convergence, the latter also reflected in the simulation study. Overall, using R2-GP is computationally and inferentially more robust for this data while providing the benefit of a prior based off of first principles.

5 Discussion

We introduce the R2-GP prior for GAM GPs under which the magnitude of a component equals its signal variance, the expected variation across the inputs. We model magnitudes based off of an interpretable prior on the model's proportion of explained variance and a Dirichlet partition of that variance across components. Thus, components compete for explained variance, preventing variance explosion with many components. In HS form, the normaliser costs $O(M)$ per gradient evaluation, so the approach also scales well for large sample sizes. Our simulations and real-data case study demonstrate superior predictive performance of R2-GP over independent normal priors. Although competitive in predictive performance with other shrinkage priors, R2-GP outperforms in terms of sampling efficiency, especially when many components are active. Future works should include richer kernel structures, such as interactions between covariates and as well as applications in other GP extensions.

AI Use Statement

In this work, we used generative AI tools for code prototyping of parts of the code (Claude Opus 5 High from Anthropic). Further, we used the ChatGPT 5.6 model from OpenAI for minor language edits, aiming to enhance readability. LLM-generated code and text was verified, modified and tested for correctness by the authors. We take responsibility for the final content of this work, including text, claims, or artifacts produced with the aid of generative AI.

References

Javier Enrique Aguilar and Paul-Christian Bürkner. Intuitive joint priors for Bayesian linear multilevel models: The R2D2M2 prior. *Electronic Journal of Statistics*, 17(1):1711–1767, 2023. doi: 10.1214/23-EJS2136.

Javier Enrique Aguilar, David Kohns, Aki Vehtari, and Paul-Christian Bürkner. R2 priors for grouped variance decomposition in high-dimensional regression. *arXiv preprint arXiv:2507.11833*, 2025.

Bob Carpenter, Andrew Gelman, Matthew D. Hoffman, Daniel Lee, Ben Goodrich, Michael Betancourt, Marcus Brubaker, Jiqiang Guo, Peter Li, and Allen Riddell. Stan: A probabilistic programming language. *Journal of Statistical Software*, 76(1):1–32, 2017. doi: 10.18637/jss.v076.i01.

David Duvenaud, Hannes Nickisch, and Carl Edward Rasmussen. Additive Gaussian processes. In *Advances in Neural Information Processing Systems*, volume 24, 2011. Page range to be confirmed.

Luisa Ferrari and Massimo Ventrucci. A standardization procedure to incorporate variance partitioning based priors in latent Gaussian models. *arXiv preprint arXiv:2501.16057*, 2025.

Geir-Arne Fuglstad, Daniel Simpson, Finn Lindgren, and Håvard Rue. Constructing priors that penalize the complexity of Gaussian random fields. *Journal of the American Statistical Association*, 114(525):445–452, 2019. doi: 10.1080/01621459.2017.1415907.

Jim E. Griffin and Philip J. Brown. Inference with normal-gamma prior distributions in regression problems. *Bayesian Analysis*, 5(1):171–188, 2010. doi: 10.1214/10-BA507.

Trevor Hastie and Robert Tibshirani. Generalized additive models. *Statistical Science*, 1(3):297–310, 1986.

Trevor J. Hastie and Robert J. Tibshirani. *Generalized Additive Models*. Chapman and Hall, London, 1990.

Matthew D. Hoffman and Andrew Gelman. The No-U-Turn sampler: Adaptively setting path lengths in Hamiltonian Monte Carlo. *Journal of Machine Learning Research*, 15:1593–1623, 2014.

David Kohns, Noa Kallioinen, Yann McLatchie, and Aki Vehtari. The ARR2 prior: Flexible predictive prior definition for Bayesian auto-regressions. *Bayesian Analysis*, 2025. doi: 10.1214/25-BA1512. Advance publication.

Max Kuhn. Building predictive models in R using the caret package. *Journal of Statistical Software*, 28(5):1–26, 2008.

Yann McLatchie, Sölvi Rögnvaldsson, Frank Weber, and Aki Vehtari. Advances in projection predictive inference. *Statistical Science*, 40(1):128–147, 2025. doi: 10.1214/24-STS949.

Scot R. Mente and Franco Lombardo. A recursive-partitioning model for blood–brain barrier permeation. *Journal of Computer-Aided Molecular Design*, 19(7):465–481, 2005.

Soham Mukherjee, Manfred Claassen, and Paul-Christian Bürkner. Hilbert space methods for approximating multi-output latent variable gaussian processes. *Statistics and Computing, 36:110*, 2026. doi: https://doi.org/10.1007/s11222-026-10869-x.

Radford M. Neal. MCMC using Hamiltonian dynamics. In Steve Brooks, Andrew Gelman, Galin L. Jones, and Xiao-Li Meng, editors, *Handbook of Markov Chain Monte Carlo*, pages 113–162. Chapman and Hall/CRC, 2011.

Omiros Papaspiliopoulos, Gareth O. Roberts, and Martin Sköld. A general framework for the parametrization of hierarchical models. *Statistical Science*, 22(1):59–73, 2007. doi: 10.1214/088342307000000014.

Juho Piironen and Aki Vehtari. Sparsity information and regularization in the horseshoe and other shrinkage priors. *Electronic Journal of Statistics*, 11(2):5018–5051, 2017. doi: 10.1214/17-EJS1337SI.

Carl Edward Rasmussen and Christopher K. I. Williams. *Gaussian Processes for Machine Learning*. MIT Press, Cambridge, MA, 2006.

Anna Elisabeth Riha, Leevi Lindgren, David Kohns, Paul-Christian Bürkner, and Aki Vehtari. To select or not to select: predictively consistent priors instead of model selection. *arXiv preprint arXiv:2606.22850*, 2026.

Gabriel Riutort-Mayol, Paul-Christian Bürkner, Michael R. Andersen, Arno Solin, and Aki Vehtari. Practical Hilbert space approximate Bayesian Gaussian processes for probabilistic programming. *Statistics and Computing*, 33(1):17, 2023. doi: 10.1007/s11222-022-10167-2.

Daniel Simpson, Håvard Rue, Andrea Riebler, Thiago G. Martins, and Sigrunn H. Sørbye. Penalising model component complexity: A principled, practical approach to constructing priors. *Statistical Science*, 32(1):1–28, 2017. doi: 10.1214/16-STS576.

Tuomas Sivula, Måns Magnusson, Asael Alonzo Matamoros, and Aki Vehtari. Uncertainty in Bayesian leave-one-out cross-validation based model comparison. *Bayesian Analysis*, 2025. doi: 10.1214/25-BA1569. Advance publication.

Arno Solin and Simo Särkkä. Hilbert space methods for reduced-rank Gaussian process regression. *Statistics and Computing*, 30(2):419–446, 2020. doi: 10.1007/s11222-019-09886-w.

Juho Timonen and Harri Lähdesmäki. Scalable Mixed-Domain Gaussian Process Modeling and Model Reduction for Longitudinal Data. *Bayesian Analysis*, pages 1 – 28, 2025. doi: 10.1214/25-BA1539. URL https://doi.org/10.1214/25-BA1539.

Aki Vehtari, Andrew Gelman, and Jonah Gabry. Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. *Statistics and Computing*, 27(5):1413–1432, 2017. doi: 10.1007/s11222-016-9696-4.

Christopher Williams and Carl Rasmussen. Gaussian processes for regression. In D. Touretzky, M. C. Mozer, and M. Hasselmo, editors, *Advances in Neural Information Processing Systems*, volume 8. MIT Press, 1995.

Simon N. Wood. *Generalized Additive Models, An Introduction with R, Second Edition*. Chapman and Hall/CRC, 2017.

Eric Yanchenko, Howard D Bondell, and Brian J Reich. Spatial regression modeling via the r2d2 framework. *Environmetrics*, 35(2):e2829, 2024.

Eric Yanchenko, Howard D Bondell, and Brian J Reich. The r2d2 prior for generalized linear mixed models. *The American Statistician*, 79(1):40–49, 2025.

Yan Dora Zhang, Brian P. Naughton, Howard D. Bondell, and Brian J. Reich. Bayesian regression using a prior on the model fit: The R2-D2 shrinkage prior. *Journal of the American Statistical Association*, 117(538):862–874, 2022. doi: 10.1080/01621459.2020.1825449.

Supplementary Material

Overview. This supplement has four parts. Appendix A gives the derivations and proofs for Section 2. Appendix B gives the full results of the simulation study of Section 3. Appendix C gives further results for the case study of Section 4. Appendix D describes the computing infrastructure. Table 1 lists the content of each part and the section of the main text that it supports.

Table 1: Overview of the Supplementary Material

Appendix	Content	Figures	Tables	Section
A.1	Variance decomposition (4) and normaliser (5) for the exact GP			2.1
A.2	Proofs of Lemmas 2.1 and A.1, of the decomposition (13) and of the calibration results (15) and (16)			2.2–2.4
B.1	Dense DGP: Δelpd and posterior of R^2_{emp} (figures), interval coverage, relative RMSE and sampling diagnostics (tables), for all kernels and length scales	5–10	2–10	3.2
B.2	Sparse DGP: the same figures and tables	11–16	11–19	3.2
B.3	Competing priors without the normalisation: Δelpd for both DGPs	17–22		3.1
C	BloodBrain data: effects of the 12 covariates with the largest share of explained variance, predictive performance and sampling diagnostics	23–24	20	4
D	Hardware, software versions and run times			3, 4

A Proofs

A.1 Derivations for Section 2.1

Throughout, $\mathbb{E}_x$ and Var_x are the expectation and variance operators with respect to random variable x, and x and x' are independent inputs.

Decomposition (4). Apply the law of total variance to $f_d(x)$ under the joint distribution of x and f_d. Conditioning on x gives

$$\mathrm{Var}\big(f_d(x)\big) = \mathbb{E}_x\big[\mathrm{Var}_{f_d}(f_d(x))\big] + \mathrm{Var}_x\big(\mathbb{E}_{f_d}[f_d(x)]\big) = \alpha_d^2,$$

since under (2) the prior mean is zero and $\mathrm{Var}_{f_d}(f_d(x)) = \alpha_d^2\, k(x, x; \rho_d) = \alpha_d^2$ at every input. Conditioning on f_d instead gives $\mathrm{Var}(f_d(x)) = \mathbb{E}_{f_d}[\mathrm{Var}_x(f_d(x) \mid f_d)] + \mathrm{Var}_{f_d}(\mathbb{E}_x[f_d(x) \mid f_d])$, and the two expressions together are (4).

Linear Term. The same two steps hold for any zero-mean prior on a function f of x, with $\mathbb{E}_x[\mathrm{Var}_f(f(x))]$ on the left-hand side. For $f(x) = x\beta$ with centred inputs, $\mathbb{E}_x[x] = 0$, the level is $\mathbb{E}_x[x]\,\beta = 0$ for every β. Its variance is therefore zero, and the marginal variance, averaged over the inputs, equals the signal variance, $\mathrm{Var}(\beta)\,\mathbb{E}_x[x^2]$.

Variance of the Level and Q_d. Under (2), the level has mean zero, and x, x' and f_d are independent, so that

$$\mathrm{Var}_{f_d}\big(\mathbb{E}_x[f_d(x) \mid f_d]\big) = \mathbb{E}_{f_d}\big[\mathbb{E}_x[f_d(x) \mid f_d]\,\mathbb{E}_{x'}[f_d(x') \mid f_d]\big] = \mathbb{E}_{x,x'}\big[\mathbb{E}_{f_d}[f_d(x)f_d(x')]\big] = \alpha_d^2\,\mathbb{E}_{x,x'}\big[k(x,x';\rho_d)\big],$$

which is $\alpha_d^2\,\bar{k}_d$ and, with (4), gives (5). A Matérn kernel has the form $k(x,x';\rho_d) = \kappa(|x-x'|/\rho_d)$, where κ is positive and strictly decreasing with $\kappa(0)=1$. Hence $0 < k(x,x';\rho_d) < 1$ for $x \neq x'$, and $k(x,x';\rho_d)$ increases with ρ_d. Unless all inputs coincide, it follows that $\bar{k}_d \in (0,1)$ increases with ρ_d, so that $Q_d(\rho_d) \in (0,1)$ decreases with ρ_d. For a general stationary kernel, $\alpha_d^2\bar{k}_d$ is still a variance that is at most α_d^2, so that $Q_d \in [0,1]$ without any assumption on the sign of the correlations.

Finite-sample Counterpart. Section 2.3 replaces Var_x by the sample variance $\widehat{\mathrm{Var}}_N$ over the observed inputs $x_{1d},\dots,x_{Nd}$. The proof of (24) holds for any covariance matrix. For the exact kernel matrix K_d in place of $\widetilde{K}_d$, with $[K_d]_{ii}=1$, it gives

$$\mathbb{E}_{f_d}\big[\widehat{\mathrm{Var}}_N(f_d)\big] = \alpha_d^2\Big\{1 - \frac{1}{N(N-1)}\sum_{i\neq j} k(x_{id},x_{jd};\rho_d)\Big\}$$

under (2), so that the expectation over two independent inputs in $\bar{k}_d$ becomes the average over pairs of distinct observed inputs.

A.2 Proofs for Sections 2.2 to 2.4

Throughout, $\mathbb{E}_z$ denotes expectation over the coefficient vectors $z_1,\dots,z_D$ only, holding X, ρ and α fixed. Two elementary identities are used repeatedly and are stated here once. Let $P_N = I_N - N^{-1}\mathbf{1}_N\mathbf{1}_N^\mathsf{T}$ be the centring matrix, which is symmetric and idempotent with $P_N\mathbf{1}_N = 0$ and $\mathrm{tr}\,P_N = N-1$. Subtracting the sample mean is an orthogonal projection, so for any $v \in \mathbb{R}^N$

$$\widehat{\mathrm{Var}}_N(v) = \frac{1}{N-1}\|P_N v\|^2 = \frac{1}{N-1}\,v^\mathsf{T}P_N v, \tag{21}$$

and for a random vector u with mean zero and covariance Σ and a fixed matrix A, writing the scalar $u^\mathsf{T}Au$ as $\mathrm{tr}(Auu^\mathsf{T})$ and using linearity of the trace gives

$$\mathbb{E}\big[u^\mathsf{T}Au\big] = \mathrm{tr}(A\Sigma). \tag{22}$$

Proof of Lemma 2.1. Write $g_d = Az_d$ with the fixed matrix $A = \Phi_d\,\mathrm{diag}(\sqrt{S_d})$. A linear image of a Gaussian vector is Gaussian, with $\mathbb{E}_z[g_d] = A\mathbb{E}_z[z_d] = 0$ and

$$\mathbb{E}_z[g_d g_d^\mathsf{T}] = A\,\mathbb{E}_z[z_d z_d^\mathsf{T}]\,A^\mathsf{T} = AA^\mathsf{T} = \Phi_d\,\mathrm{diag}(\sqrt{S_d})\,\mathrm{diag}(\sqrt{S_d})\Phi_d^\mathsf{T} = \Phi_d\,\mathrm{diag}(S_d)\Phi_d^\mathsf{T},$$

using $\mathrm{diag}(\sqrt{S_d})^2 = \mathrm{diag}(S_d)$, valid since the matrices are diagonal and $S_{dm} > 0$. Symmetry is immediate, and $v^\mathsf{T}\widetilde{K}_d v = \|\mathrm{diag}(\sqrt{S_d})\Phi_d^\mathsf{T}v\|^2 \geq 0$ gives positive semi-definiteness. The rank bound follows from $\mathrm{rank}(AA^\mathsf{T}) = \mathrm{rank}(A) \leq \min(N,M)$. □

Lemma A.1 (Equivalent expressions for the normaliser)**.** *Let g_d and $\widetilde{K}_d$ be as in Lemma 2.1. Then the normaliser (11) satisfies*

$$\begin{aligned} Q_d(\rho_d) &= \frac{1}{N-1}\,\mathrm{tr}\big(P_N\widetilde{K}_d\big) && (23)\\ &= \sum_{m=1}^{M} S_{dm}\,\widehat{\mathrm{Var}}_N(\phi_{dm}) && (12)\\ &= \frac{1}{N-1}\Big\{\mathrm{tr}(\widetilde{K}_d) - \tfrac{1}{N}\mathbf{1}_N^\mathsf{T}\widetilde{K}_d\mathbf{1}_N\Big\}. && (24) \end{aligned}$$

Moreover, $Q_d(\rho_d) > 0$ whenever $M \geq 2$ and the observed inputs of covariate d lie in $(-L_d, L_d)$ and are not all equal.

Proof. By (21), $\widehat{\mathrm{Var}}_N(g_d) = (N-1)^{-1} g_d^{\mathsf{T}} P_N g_d$. Conditionally on X and ρ_d the matrix P_N is fixed and g_d has mean zero with covariance $\widetilde{K}_d$ by Lemma 2.1, so (22) gives (23).

For (12), substitute $\widetilde{K}_d = \Phi_d \operatorname{diag}(S_d) \Phi_d^{\mathsf{T}}$ and use cyclicity together with the fact that $\operatorname{tr}(\operatorname{diag}(S)B) = \sum_m S_m B_{mm}$ for any square B:

$$\begin{aligned}\operatorname{tr}\big(P_N \Phi_d \operatorname{diag}(S_d) \Phi_d^{\mathsf{T}}\big) &= \operatorname{tr}\big(\operatorname{diag}(S_d)\, \Phi_d^{\mathsf{T}} P_N \Phi_d\big) \\ &= \sum_{m=1}^{M} S_{dm} \big[\Phi_d^{\mathsf{T}} P_N \Phi_d\big]_{mm} = \sum_{m=1}^{M} S_{dm}\, \phi_{dm}^{\mathsf{T}} P_N \phi_{dm}.\end{aligned}$$

Dividing by $N-1$ and applying (21) to each column turns $\phi_{dm}^{\mathsf{T}} P_N \phi_{dm}/(N-1)$ into $\widehat{\mathrm{Var}}_N(\phi_{dm})$. The off-diagonal entries of $\Phi_d^{\mathsf{T}} P_N \Phi_d$, which are the sample covariances between different basis columns, do not enter, because the coefficients in z_d are uncorrelated.

For (24), expand $P_N = I_N - N^{-1}\mathbf{1}_N\mathbf{1}_N^{\mathsf{T}}$ inside the trace and use $\operatorname{tr}(\mathbf{1}_N\mathbf{1}_N^{\mathsf{T}}\widetilde{K}_d) = \mathbf{1}_N^{\mathsf{T}}\widetilde{K}_d\mathbf{1}_N$.

For strict positivity, $S_{dm} > 0$ and $\widehat{\mathrm{Var}}_N(\phi_{dm}) \geq 0$, so that (12) is zero only if every basis column is constant across the observed inputs. Let $u_i = (x_{id} + L_d)/(2L_d)$, which lies in $(0,1)$ by assumption, so that $\phi_{dm}(x_{id}) = L_d^{-1/2}\sin(m\pi u_i)$ by (9). The first column is constant only if every u_i equals u^* or $1-u^*$ for some $u^* \in (0, 1/2]$. If both values occur and $u^* < 1/2$, the second column takes the two distinct values $\pm L_d^{-1/2}\sin(2\pi u^*)$ and is not constant. Otherwise all inputs are equal. The condition $M \geq 2$ cannot be dropped: for $M = 1$, a design with two distinct input values that are symmetric about the centre of the interval gives $Q_d = 0$. □

Proof of the decomposition (13). Write $\widehat{\mathrm{Var}}_N(g_d) = (N-1)^{-1}\big\{\sum_{i=1}^N [g_d]_i^2 - N\bar{g}_d^2\big\}$ with $\bar{g}_d = N^{-1}\mathbf{1}_N^{\mathsf{T}} g_d$. By Lemma 2.1, $[g_d]_i$ and $\bar{g}_d$ have mean zero, with $\mathbb{E}_z\big[[g_d]_i^2\big] = [\widetilde{K}_d]_{ii}$ and $\mathbb{E}_z[\bar{g}_d^2] = \mathrm{Var}_z(\bar{g}_d) = N^{-2}\mathbf{1}_N^{\mathsf{T}}\widetilde{K}_d\mathbf{1}_N$. Taking the expectation term by term gives (13), in agreement with (24). Splitting $\mathbf{1}_N^{\mathsf{T}}\widetilde{K}_d\mathbf{1}_N$ into its diagonal and off-diagonal sums gives the equivalent form

$$Q_d(\rho_d) = \frac{1}{N}\sum_{i=1}^{N}[\widetilde{K}_d]_{ii} - \frac{1}{N(N-1)}\sum_{i\neq j}[\widetilde{K}_d]_{ij},$$

the average marginal variance minus the average covariance between two distinct observed inputs. When the diagonal of $\widetilde{K}_d$ is one, $Q_d < 1$ if and only if this average covariance is positive. Without a condition on the sign of the covariances, positive semi-definiteness of $\widetilde{K}_d$ gives only $Q_d \leq \operatorname{tr}(\widetilde{K}_d)/(N-1)$, which is $N/(N-1)$ for a unit diagonal. □

Proof of the component calibration (15). Conditionally on X, ρ_d and α_d, the factor $\alpha_d/\sqrt{Q_d}$ is a deterministic scalar, and $\widehat{\mathrm{Var}}_N$ is homogeneous of degree two, since $\widehat{\mathrm{Var}}_N(cv) = (N-1)^{-1}(cv)^{\mathsf{T}}P_N(cv) = c^2\widehat{\mathrm{Var}}_N(v)$. Hence

$$\mathbb{E}_z\big[\widehat{\mathrm{Var}}_N(f_d)\big] = \frac{\alpha_d^2}{Q_d}\,\mathbb{E}_z\big[\widehat{\mathrm{Var}}_N(g_d)\big] = \frac{\alpha_d^2}{Q_d}\cdot Q_d = \alpha_d^2$$

by the definition (11). Strict positivity of Q_d holds under the conditions of Lemma A.1. □

Proof of the additive calibration (16). Let $\widehat{\mathrm{Cov}}_N(u,v) = (N-1)^{-1}u^{\mathsf{T}}P_N v$, which is bilinear and symmetric with $\widehat{\mathrm{Var}}_N(u) = \widehat{\mathrm{Cov}}_N(u,u)$. Bilinearity gives the exact decomposition

$$\widehat{\mathrm{Var}}_N\Big(\sum_d f_d\Big) = \sum_{d=1}^{D}\widehat{\mathrm{Var}}_N(f_d) + 2\sum_{1\leq d<e\leq D}\widehat{\mathrm{Cov}}_N(f_d, f_e).$$

For $d \neq e$, cyclicity of the trace as in (22) gives

$$\mathbb{E}_z\big[\widehat{\mathrm{Cov}}_N(f_d, f_e)\big] = \frac{1}{N-1}\operatorname{tr}\big(P_N\,\mathbb{E}_z[f_e f_d^{\mathsf{T}}]\big) = 0,$$

since given X, ρ and α the vectors f_d and f_e are functions of the independent z_d and z_e alone and each has mean zero, so $\mathbb{E}_z[f_e f_d^{\mathsf{T}}] = \mathbb{E}_z[f_e]\mathbb{E}_z[f_d]^{\mathsf{T}} = 0$. Applying (15) to each diagonal term gives $\sum_d \alpha_d^2$, which equals $\alpha^2\sum_d \psi_d = \alpha^2$ because ψ lies in the simplex. □

B Additional Simulation Results

This appendix gives the full results of the simulation study of Section 3. The settings are those of Section 3.1. In the figures and tables, ρ in a panel title or in the first column of a table is the generating length scale ρ_0, and RHS without a qualifier and RHS (p_0=5) are the calibration with $p_0 = \min(5, D-1)$.

Tables. For each DGP and kernel, three tables give the coverage of the central 90% posterior intervals, the RMSE of the posterior mean relative to R2-GP, and the sampling diagnostics. All entries are means over 50 replications, with one column for each D. For α and ρ, coverage and RMSE are taken across the components. A relative RMSE above one indicates worse recovery than under R2-GP. For each fit, $\widehat{R}$ is the largest value and ESS is the smallest bulk effective sample size over all α_d, ρ_d, σ and the log density, and divergences are counted among the 4000 post-warm-up draws. Bold marks the best value in each column for each ρ_0.

B.1 Dense DGP

Figures 5 to 7 show the mean Δelpd and Figures 8 to 10 the posterior of R^2_{emp}, for the three kernels and the three length scales. Figure 3 is the condition with the SE kernel and $\rho_0 = 0.3$. Tables 2 to 10 give the parameter recovery and the sampling diagnostics. Averaged over replications, the posterior mean of R^2_{emp} minus R^2_{true} lies in $[-0.05, 0.00]$ under R2-GP in all conditions. At $D = 100$, it lies in $[0.62, 0.64]$ under the normal prior, in $[0.23, 0.62]$ under NG and in $[-0.27, -0.21]$ under RHS (p_0=5). The recovery of ρ is governed by its informative prior: its coverage is 100% after rounding and its RMSE is below 0.03 under all priors.

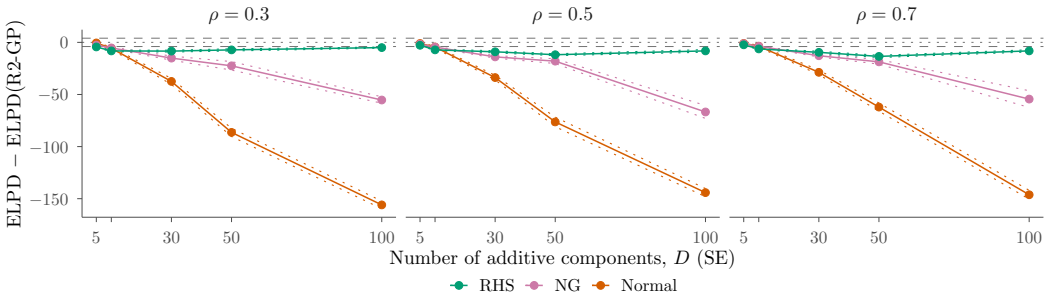

Figure 5: Predictive Performance, Dense DGP, SE Kernel. Mean Δelpd over 50 replications (solid) with ± 1 Monte Carlo standard error (dotted), for each generating length scale ρ_0 (panels). Dashed lines mark ± 4 and negative values favour R2-GP.

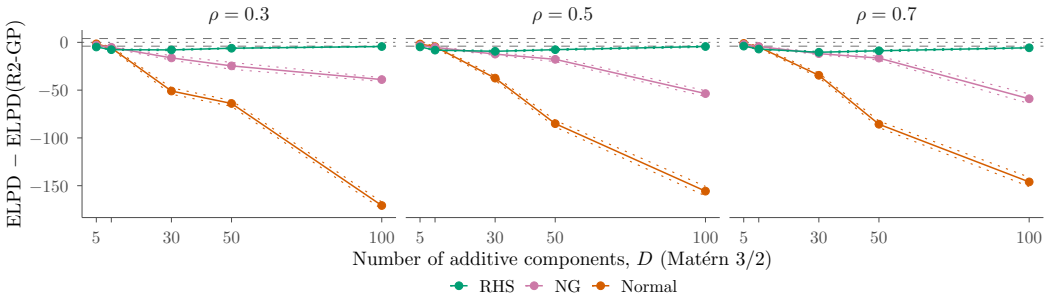

Figure 6: Predictive Performance, Dense DGP, Matérn 3/2 Kernel. As in Figure 5.

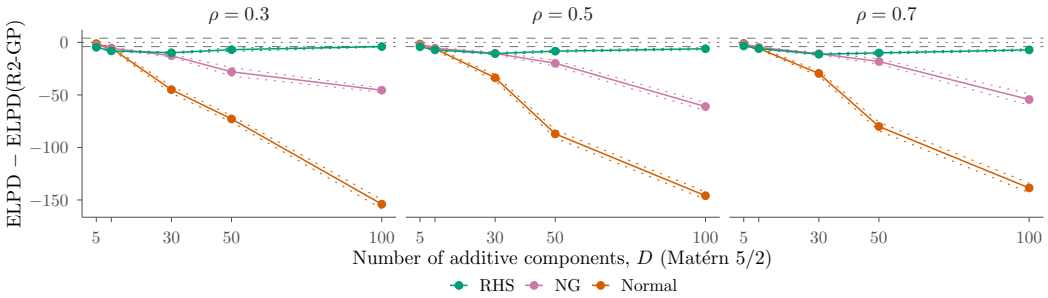

Figure 7: Predictive Performance, Dense DGP, Matérn 5/2 Kernel. As in Figure 5.

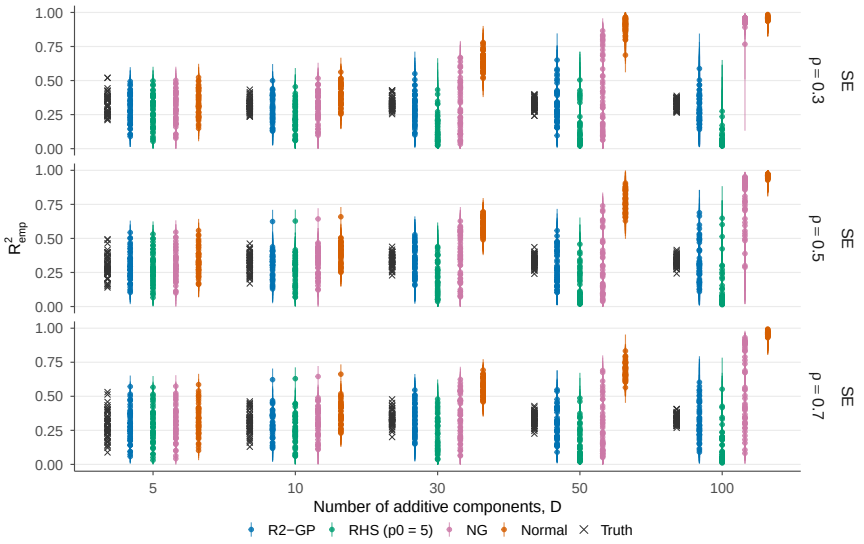

Figure 8: Explained Variance, Dense DGP, SE Kernel. Posterior means (points) and central 90% intervals (bars) of R^2_{emp} for each replication, for each generating length scale ρ_0 (rows). Crosses mark R^2_{true}.

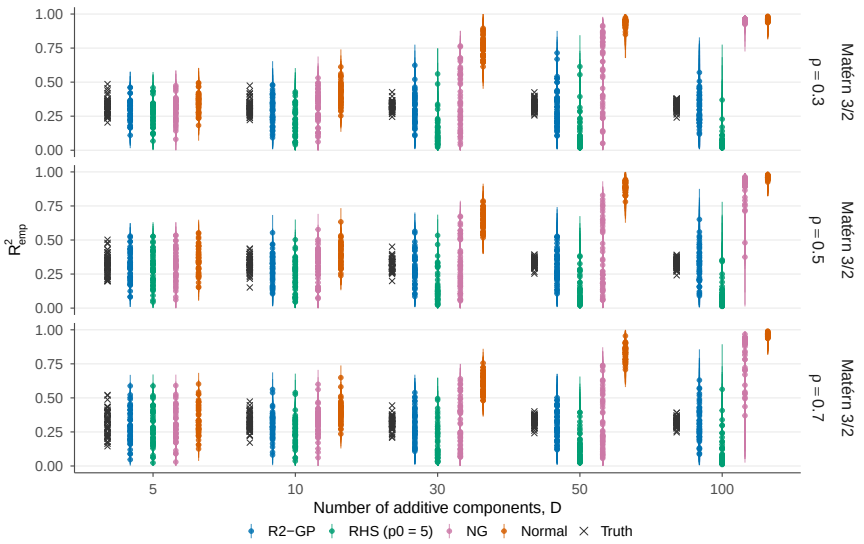

Figure 9: Explained Variance, Dense DGP, Matérn 3/2 Kernel. As in Figure 8.

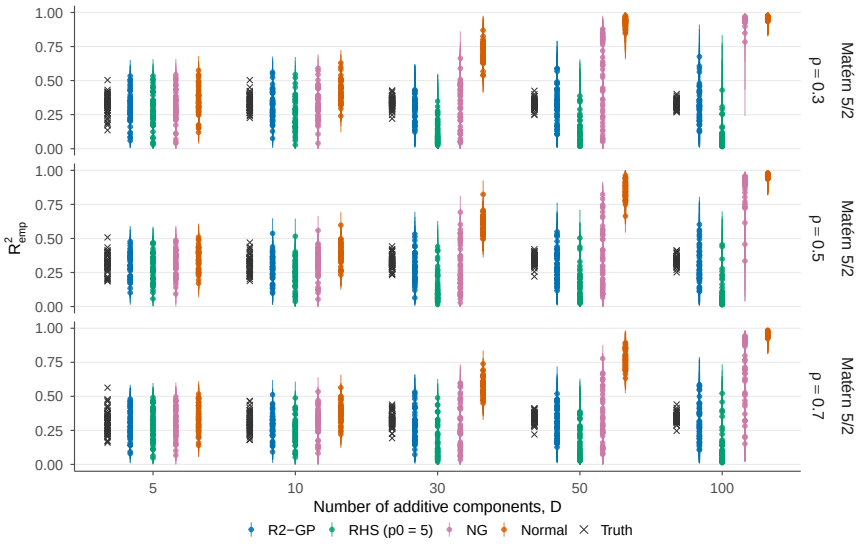

Figure 10: Explained Variance, Dense DGP, Matérn 5/2 Kernel. As in Figure 8.

Table 2: Interval Coverage (%), Dense DGP, SE Kernel. Bold marks the coverage closest to 90%.

		α					ρ					σ				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	97	100	100	99	100	**100**	**100**	**100**	**100**	**100**	92	**82**	**92**	**94**	**98**
	Normal	95	96	**95**	**96**	100	**100**	**100**	**100**	**100**	**100**	**90**	80	38	2	0
	NG	**91**	97	80	77	**99**	**100**	**100**	**100**	**100**	**100**	92	80	72	68	2
	RHS (p_0=5)	84	**85**	63	47	18	**100**	**100**	**100**	**100**	**100**	92	74	64	54	42
0.5	R2-GP	99	99	100	100	100	**100**	**100**	**100**	**100**	**100**	94	**86**	**88**	**90**	**92**
	Normal	95	95	94	**94**	**99**	**100**	**100**	**100**	**100**	**100**	**92**	84	60	8	0
	NG	94	97	**90**	64	**99**	**100**	**100**	**100**	**100**	**100**	94	84	72	66	34
	RHS (p_0=5)	**88**	**88**	82	52	33	**100**	**100**	**100**	**100**	**100**	94	84	70	58	48
0.7	R2-GP	98	99	100	100	99	**100**	**100**	**100**	**100**	**100**	92	**90**	**88**	**88**	**96**
	Normal	95	95	94	**95**	96	**100**	**100**	**100**	**100**	**100**	**90**	84	72	24	0
	NG	98	98	**93**	68	**88**	**100**	**100**	**100**	**100**	**100**	92	86	74	72	58
	RHS (p_0=5)	**87**	**89**	85	57	40	**100**	**100**	**100**	**100**	**100**	92	86	74	68	54

Table 3: Relative RMSE, Dense DGP, SE Kernel. The R2-GP rows give its RMSE in italics.

		α					ρ					σ				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	***0.097***	***0.062***	***0.033***	***0.025***	***0.017***	*0.022*	*0.016*	*0.008*	*0.006*	*0.003*	*0.054*	*0.084*	***0.079***	***0.082***	***0.078***
	Normal	1.57	2.30	4.94	7.59	11.61	1.05	1.22	1.84	2.50	3.13	**0.98**	**0.95**	2.65	7.73	9.97
	NG	1.52	1.88	2.74	2.92	3.13	0.94	0.98	1.12	1.30	3.13	1.01	1.02	1.66	2.30	9.26
	RHS (p_0=5)	1.64	1.97	2.64	2.93	3.46	**0.87**	**0.76**	**0.57**	**0.54**	**0.58**	1.11	1.30	1.72	1.90	2.22
0.5	R2-GP	***0.092***	***0.061***	***0.030***	***0.029***	***0.018***	*0.016*	*0.012*	*0.006*	*0.004*	*0.003*	*0.053*	***0.065***	***0.077***	***0.078***	***0.097***
	Normal	1.96	2.85	6.65	8.05	13.26	1.01	1.07	1.60	2.50	3.65	**0.98**	1.02	1.83	4.79	7.92
	NG	1.66	2.02	3.05	2.89	3.35	0.93	0.95	1.06	1.13	2.70	1.00	1.03	1.49	1.65	5.26
	RHS (p_0=5)	1.75	2.02	2.61	2.51	3.08	**0.88**	**0.81**	**0.68**	**0.66**	**0.81**	1.07	1.21	1.48	1.71	1.88
0.7	R2-GP	***0.095***	***0.061***	***0.031***	***0.027***	***0.018***	*0.011*	*0.008*	*0.004*	*0.003*	*0.002*	*0.053*	***0.065***	***0.074***	***0.075***	***0.082***
	Normal	2.07	3.11	7.12	8.98	15.52	0.94	1.08	1.45	1.98	6.32	**0.98**	**1.00**	1.41	3.23	9.53
	NG	1.71	2.05	3.02	3.04	3.50	0.92	0.96	1.01	1.27	1.82	1.00	1.03	1.41	1.52	3.34
	RHS (p_0=5)	1.79	2.06	2.48	2.55	3.04	**0.87**	**0.84**	**0.70**	**0.71**	**0.78**	1.05	1.17	1.42	1.72	1.97

Table 4: Sampling Diagnostics, Dense DGP, SE Kernel. Time is the sampling time in seconds.

		$\widehat{R}$					ESS					Divergences					Time (s)				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1078**	**998**	**814**	**661**	498	**0.0**	**0.0**	**0.2**	**0.6**	**1.9**	**5.7**	**9.1**	**26**	**51**	**129**
	Normal	**1.00**	1.01	**1.01**	1.26	1.37	951	868	480	43	18	1.2	1.4	5.8	283	235	7.7	16	76	458	1634
	NG	1.02	1.04	1.17	1.40	1.12	569	222	39	25	62	14	10	116	116	78	12	32	204	350	903
	RHS (p_0=5)	**1.00**	1.01	**1.01**	**1.01**	**1.01**	896	770	551	578	**708**	2.3	3.6	7.2	15	12	8.4	14	36	63	134
0.5	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1176**	**1148**	**1019**	**946**	**607**	**0.6**	**0.0**	**0.0**	**0.3**	**1.4**	**7.2**	**12**	**32**	**56**	191
	Normal	**1.00**	**1.00**	**1.01**	1.05	1.37	1043	973	756	217	19	2.0	1.6	1.1	80	233	9.3	20	97	349	1292
	NG	1.01	1.03	1.15	1.25	1.43	605	328	51	24	32	8.9	18	69	199	49	13	31	175	395	961
	RHS (p_0=5)	**1.00**	1.01	**1.01**	**1.01**	**1.01**	1035	911	578	624	568	4.1	4.6	17	16	32	11	16	45	69	**171**
0.7	R2-GP	**1.00**	**1.00**	**1.00**	**1.01**	**1.01**	**1289**	**1266**	**1099**	**999**	**684**	**0.7**	**0.4**	**0.0**	**0.0**	**1.7**	**8.1**	**14**	**38**	**65**	**196**
	Normal	**1.00**	**1.00**	1.01	**1.01**	1.48	1111	1078	919	531	14	2.9	1.9	0.8	8.2	247	13	26	128	324	1457
	NG	1.01	1.03	1.15	1.23	1.61	609	334	54	22	20	28	34	89	193	58	17	41	205	439	1125
	RHS (p_0=5)	**1.00**	1.01	1.01	**1.01**	**1.01**	1093	930	632	636	587	6.6	10	12	18	21	13	21	57	82	199

Table 5: Interval Coverage (%), Dense DGP, Matérn 3/2 Kernel. As in Table 2.

		α					ρ					σ				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	98	99	100	100	**100**	**100**	**100**	**100**	**100**	**100**	**92**	**88**	**96**	**90**	**100**
	Normal	94	96	95	98	**100**	**100**	**100**	**100**	**100**	**100**	98	84	2	0	0
	NG	**92**	**95**	**87**	**91**	**100**	**100**	**100**	**100**	**100**	**100**	96	82	**84**	52	0
	RHS (p_0=5)	87	78	60	42	16	**100**	**100**	**100**	**100**	**100**	82	82	70	50	46
0.5	R2-GP	98	99	98	100	**98**	**100**	**100**	**100**	**100**	**100**	**88**	94	**94**	**88**	**92**
	Normal	94	**94**	**95**	**96**	100	**100**	**100**	**100**	**100**	**100**	86	**88**	36	0	0
	NG	93	95	82	77	99	**100**	**100**	**100**	**100**	**100**	86	86	80	66	14
	RHS (p_0=5)	**89**	81	63	52	29	**100**	**100**	**100**	**100**	**100**	86	80	68	44	54
0.7	R2-GP	97	97	98	100	**98**	**100**	**100**	**100**	**100**	**100**	88	**92**	**98**	**88**	**94**
	Normal	**93**	94	**95**	**94**	99	**100**	**100**	**100**	**100**	**100**	**90**	84	50	0	0
	NG	95	96	84	75	99	**100**	**100**	**100**	**100**	**100**	88	86	**82**	70	30
	RHS (p_0=5)	**87**	**87**	67	61	35	**100**	**100**	**100**	**100**	**100**	86	86	74	60	58

Table 6: Relative RMSE, Dense DGP, Matérn 3/2 Kernel. As in Table 3.

		α					ρ					σ				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	***0.087***	***0.065***	***0.034***	***0.025***	***0.016***	*0.019*	*0.012*	*0.006*	*0.005*	*0.003*	*0.057*	***0.068***	***0.067***	***0.094***	***0.087***
	Normal	1.60	1.90	4.92	6.26	11.18	1.07	1.17	2.09	2.18	3.42	**0.93**	**1.00**	5.67	7.54	8.93
	NG	1.58	1.75	2.73	2.82	2.99	0.96	0.93	1.22	1.70	2.44	0.99	1.05	1.93	3.13	8.62
	RHS (p_0=5)	1.70	1.95	2.72	3.01	3.65	**0.87**	**0.71**	**0.55**	**0.57**	**0.74**	1.16	1.39	2.03	1.78	2.13
0.5	R2-GP	***0.096***	***0.066***	***0.035***	***0.024***	***0.016***	*0.012*	*0.008*	*0.004*	*0.003*	*0.002*	***0.065***	*0.068*	***0.064***	***0.089***	***0.098***
	Normal	1.77	2.19	4.85	8.62	12.48	0.97	1.03	1.50	2.60	4.45	1.02	**0.94**	3.40	6.66	7.85
	NG	1.60	1.84	2.61	3.35	3.38	0.93	0.91	1.05	1.39	2.46	1.02	1.01	1.76	2.09	6.61
	RHS (p_0=5)	1.67	1.93	2.52	3.07	3.47	**0.88**	**0.77**	**0.63**	**0.67**	**0.96**	1.09	1.31	1.91	1.98	1.74
0.7	R2-GP	***0.097***	***0.062***	***0.036***	***0.024***	***0.017***	*0.007*	*0.005*	*0.003*	*0.002*	***0.002***	*0.064*	*0.066*	***0.059***	***0.079***	***0.093***
	Normal	1.81	2.62	5.09	9.39	13.07	0.92	1.00	1.32	2.08	6.92	**0.96**	**0.97**	2.75	5.92	8.48
	NG	1.58	1.92	2.57	3.31	3.31	0.89	0.91	1.15	1.27	2.11	0.99	0.99	1.60	1.89	5.85
	RHS (p_0=5)	1.66	1.99	2.40	2.96	3.26	**0.84**	**0.77**	**0.75**	**0.91**	1.11	1.08	1.24	1.86	2.06	1.75

Table 7: Sampling Diagnostics, Dense DGP, Matérn 3/2 Kernel. As in Table 4.

		$\widehat{R}$					ESS					Divergences					Time (s)				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1053**	**990**	**716**	**581**	475	**0.0**	**0.0**	**0.1**	**2.4**	**6.1**	**7.7**	**13**	**38**	**72**	**215**
	Normal	**1.00**	1.01	1.07	1.25	1.37	925	807	171	23	16	0.6	0.4	114	243	247	9.1	20	177	738	2401
	NG	1.02	1.05	1.19	1.39	1.08	512	197	37	28	58	2.3	12	69	66	53	14	49	272	450	1678
	RHS (p_0=5)	1.01	1.01	**1.01**	**1.01**	1.02	861	748	541	579	**678**	1.2	2.3	6.9	18	28	11	19	54	99	273
0.5	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1179**	**1120**	**966**	**775**	540	**0.0**	**0.1**	**0.1**	**0.4**	**7.2**	**8.7**	**15**	**42**	**73**	**265**
	Normal	**1.00**	1.01	**1.01**	1.20	1.39	1016	988	545	42	17	1.7	0.8	6.5	239	273	12	25	138	690	2543
	NG	1.03	1.04	1.18	1.33	1.21	650	276	38	23	50	6.1	28	110	118	73	16	47	294	483	1760
	RHS (p_0=5)	**1.00**	1.01	**1.01**	**1.01**	1.02	940	833	614	559	**634**	2.2	8.4	8.1	17	40	14	22	59	112	317
0.7	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1228**	**1179**	**1099**	**856**	**610**	**1.0**	**0.0**	**0.1**	**0.4**	**2.0**	**11**	**17**	**49**	**90**	**282**
	Normal	**1.00**	**1.00**	**1.01**	1.09	1.45	1057	1026	754	112	16	2.1	1.0	2.3	130	238	14	30	148	574	2364
	NG	1.01	1.03	1.15	1.28	1.37	679	317	41	21	31	18	24	122	118	59	20	49	294	568	1924
	RHS (p_0=5)	**1.00**	1.01	**1.01**	1.02	1.02	1018	921	643	530	561	4.2	5.5	14	33	37	16	27	72	127	349

Table 8: Interval Coverage (%), Dense DGP, Matérn 5/2 Kernel. As in Table 2.

		α					ρ					σ				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	92	97	100	100	**100**	**100**	**100**	**100**	**100**	**100**	82	**88**	96	**92**	**98**
	Normal	93	97	**95**	97	**100**	**100**	**100**	**100**	**100**	**100**	**88**	86	20	0	0
	NG	**89**	**95**	**85**	**84**	**100**	**100**	**100**	**100**	**100**	**100**	84	86	**86**	48	2
	RHS (p_0=5)	81	80	59	49	22	**100**	**100**	**100**	**100**	**100**	82	86	68	64	58
0.5	R2-GP	98	98	98	99	100	**100**	**100**	**100**	**100**	**100**	**90**	92	**96**	**92**	**94**
	Normal	94	**95**	**95**	**94**	100	**100**	**100**	**100**	**100**	**100**	**90**	88	48	2	0
	NG	94	97	82	76	**99**	**100**	**100**	**100**	**100**	**100**	88	**90**	80	70	14
	RHS (p_0=5)	**88**	80	63	56	27	**100**	**100**	**100**	**100**	**100**	86	88	62	60	50
0.7	R2-GP	98	98	98	100	100	**100**	**100**	**100**	**100**	**100**	**88**	94	**90**	**94**	**94**
	Normal	93	**95**	**95**	**94**	99	**100**	**100**	**100**	**100**	**100**	**88**	**90**	60	8	0
	NG	95	97	84	77	**95**	**100**	**100**	**100**	**100**	**100**	**88**	92	78	76	44
	RHS (p_0=5)	**90**	**85**	66	61	38	**100**	**100**	**100**	**100**	**100**	86	88	66	66	54

Table 9: Relative RMSE, Dense DGP, Matérn 5/2 Kernel. As in Table 3.

		α					ρ					σ				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	***0.100***	***0.066***	***0.031***	***0.025***	***0.017***	*0.020*	*0.016*	*0.007*	*0.005*	*0.003*	***0.066***	***0.071***	***0.056***	***0.082***	***0.078***
	Normal	1.58	2.03	5.16	6.79	10.64	1.03	1.14	1.94	2.37	2.89	**1.00**	1.11	5.11	8.49	9.87
	NG	1.53	1.76	2.81	3.02	2.78	0.93	0.99	1.04	1.78	2.63	1.02	1.06	1.66	3.29	9.56
	RHS (p_0=5)	1.62	1.89	2.91	2.96	3.36	**0.87**	**0.81**	**0.54**	**0.57**	**0.80**	1.09	1.15	2.55	1.88	1.86
0.5	R2-GP	***0.099***	***0.067***	***0.034***	***0.024***	***0.018***	*0.013*	*0.009*	*0.004*	*0.003*	*0.002*	***0.054***	***0.063***	***0.070***	***0.086***	***0.087***
	Normal	1.81	2.20	5.18	8.96	11.94	0.95	1.09	1.58	2.44	4.71	1.02	1.03	2.31	6.07	8.97
	NG	1.59	1.80	2.67	3.37	3.25	0.92	0.93	0.99	1.33	3.00	1.02	1.02	1.58	1.90	7.14
	RHS (p_0=5)	1.67	1.89	2.55	3.01	3.19	**0.85**	**0.78**	**0.64**	**0.67**	**0.85**	1.04	1.13	1.82	1.73	1.75
0.7	R2-GP	***0.101***	***0.064***	***0.036***	***0.024***	***0.017***	*0.009*	*0.006*	*0.003*	*0.002*	*0.002*	***0.053***	***0.063***	***0.072***	***0.079***	***0.087***
	Normal	1.90	2.52	5.49	9.80	13.47	0.91	0.99	1.45	1.87	5.34	1.02	**1.00**	1.67	4.58	8.83
	NG	1.62	1.89	2.62	3.42	3.38	0.89	0.90	1.07	1.36	2.01	1.01	1.03	1.40	1.75	4.77
	RHS (p_0=5)	1.71	1.92	2.42	2.91	3.15	**0.86**	**0.79**	**0.71**	**0.73**	**0.95**	1.02	1.10	1.63	1.79	1.81

Table 10: Sampling Diagnostics, Dense DGP, Matérn 5/2 Kernel. As in Table 4.

		$\widehat{R}$					ESS					Divergences					Time (s)				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1053**	**984**	**765**	**635**	448	**0.0**	**0.0**	**0.2**	**0.4**	**4.6**	**7.2**	**12**	**33**	**59**	**178**
	Normal	**1.00**	1.01	1.02	1.28	1.34	932	839	320	30	17	0.8	0.6	38	233	263	9.2	18	135	662	2012
	NG	1.01	1.03	1.21	1.32	1.10	547	265	27	29	58	15	15	104	102	65	16	40	275	397	1400
	RHS (p_0=5)	1.01	1.01	**1.01**	**1.01**	1.03	860	736	584	558	**575**	2.1	2.7	7.3	12	47	11	18	49	84	242
0.5	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1169**	**1158**	**996**	**809**	618	**0.0**	**0.0**	**0.1**	**0.7**	**1.6**	**9.0**	**14**	**39**	**69**	**207**
	Normal	**1.00**	1.01	**1.01**	1.14	1.38	990	975	677	84	17	1.8	2.0	1.9	136	212	12	23	128	529	2034
	NG	1.01	1.03	1.17	1.28	1.22	644	266	38	23	48	7.2	28	106	104	84	16	44	283	525	1315
	RHS (p_0=5)	**1.00**	1.01	**1.01**	**1.01**	**1.01**	971	896	654	583	**657**	2.3	4.0	15	22	29	13	21	57	99	238
0.7	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1241**	**1231**	**1114**	**902**	**659**	**0.2**	**0.1**	**0.1**	**0.6**	**2.5**	**9.9**	**17**	**48**	**85**	**239**
	Normal	**1.00**	**1.00**	**1.01**	1.03	1.36	1102	1069	862	279	15	2.2	1.4	2.1	53	228	15	31	150	500	2106
	NG	1.01	1.03	1.15	1.25	1.50	666	330	37	19	29	12	35	119	122	69	18	51	281	542	1488
	RHS (p_0=5)	1.01	1.01	**1.01**	**1.01**	**1.01**	1010	964	717	565	593	10	3.9	14	16	28	15	24	68	128	310

B.2 Sparse DGP

In the sparse DGP only the first five components are active, with $\alpha_d^2 = \sigma^2 R_0^2/\{5(1-R_0^2)\}$ for $d \leq 5$ and $\alpha_d = 0$ otherwise, so that the total signal variance equals that of the dense DGP (19). All other settings are those of Section 3.1. The priors are R2-GP, normal, NG and RHS with $p_0 = \min(5, D-1)$, which equals the number of active components for $D \geq 10$. Figures 11 to 16 and Tables 11 to 19 give the results. The tables report the RMSE of α separately for the active and the inactive components. For $D \geq 30$, the intervals of R2-GP contain the generating value of an active α_d in at most 27% of cases, and those of RHS in at most 49%. This agrees with the uniform allocation for $a = 1$ that Section 3.2 discusses. The recovery of σ and of R^2_{true} under R2-GP stays accurate.

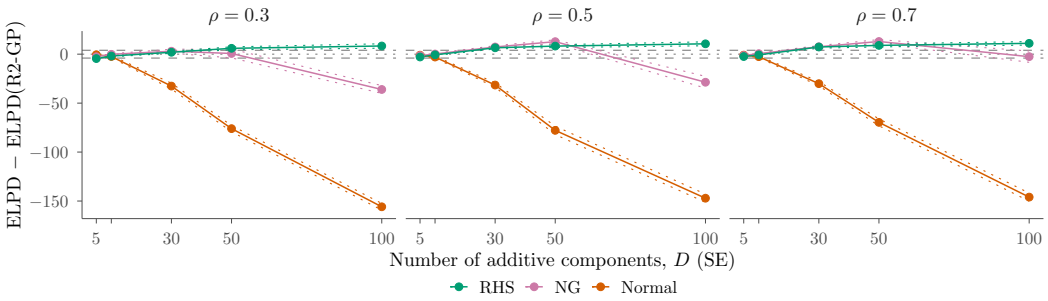

Figure 11: Predictive Performance, Sparse DGP, SE Kernel. As in Figure 5.

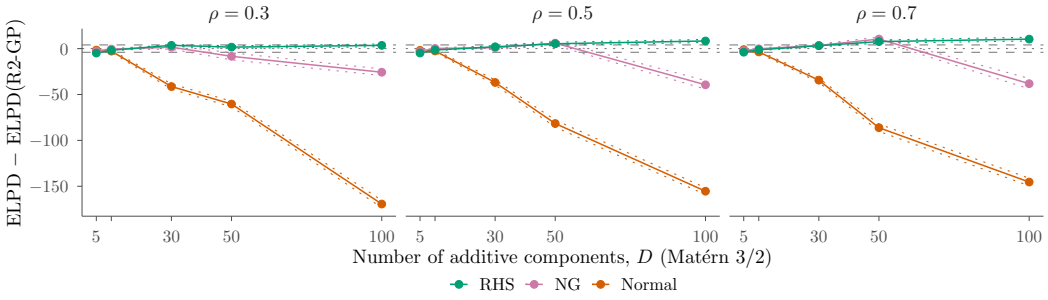

Figure 12: Predictive Performance, Sparse DGP, Matérn 3/2 Kernel. As in Figure 5.

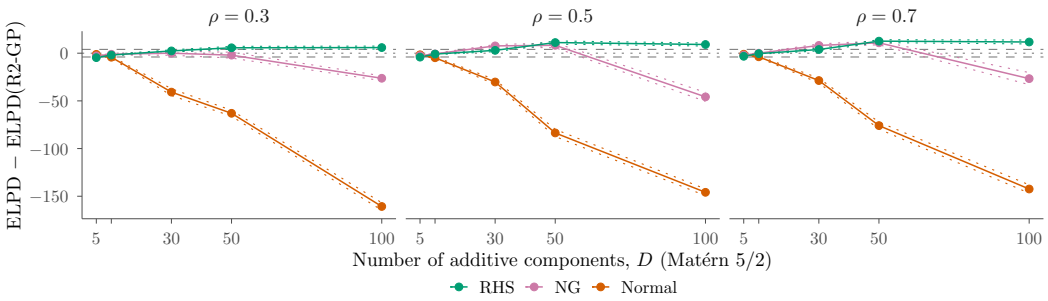

Figure 13: Predictive Performance, Sparse DGP, Matérn 5/2 Kernel. As in Figure 5.

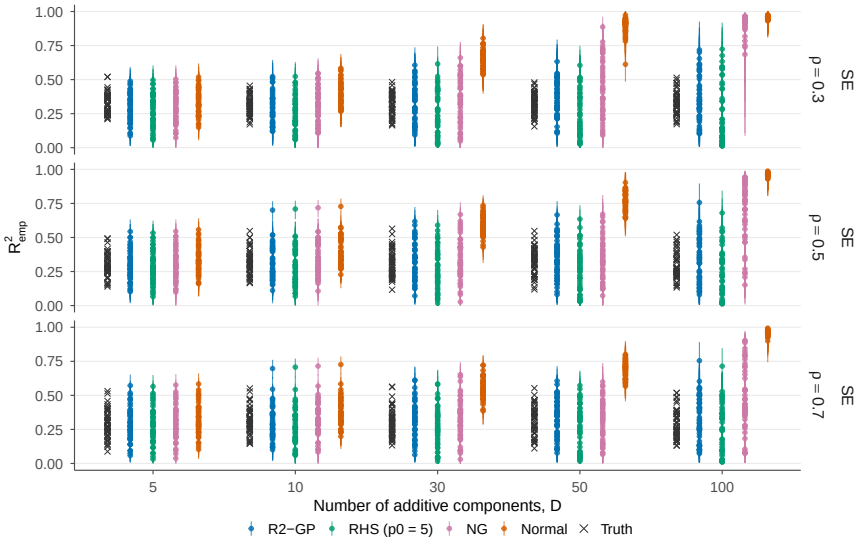

Figure 14: Explained Variance, Sparse DGP, SE Kernel. As in Figure 8.

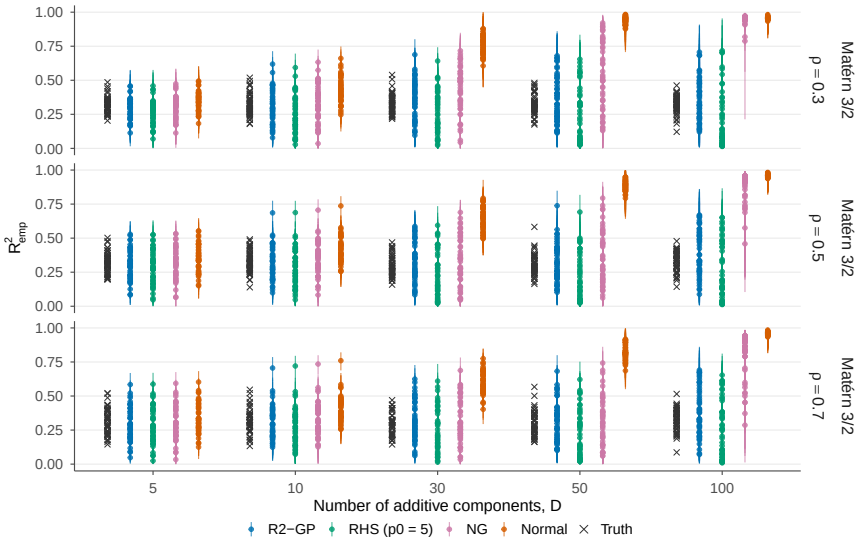

Figure 15: Explained Variance, Sparse DGP, Matérn 3/2 Kernel. As in Figure 8.

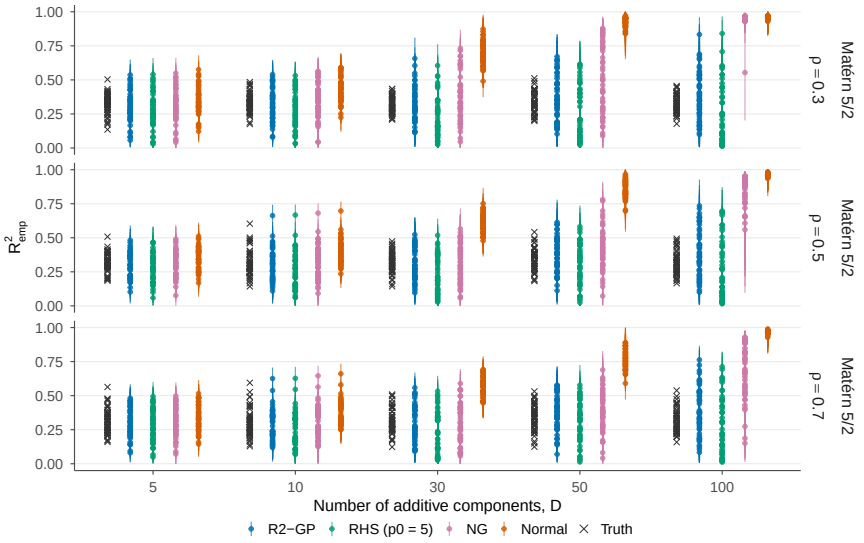

Figure 16: Explained Variance, Sparse DGP, Matérn 5/2 Kernel. As in Figure 8.

Table 11: Interval Coverage (%), Sparse DGP, SE Kernel. As in Table 2, for the active components. The interval of a positive parameter cannot contain the zero of an inactive component.

		α					ρ					σ				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	97	77	23	11	1	**100**	**100**	**100**	**100**	**100**	92	**82**	**84**	94	**94**
	Normal	95	96	**95**	**95**	**100**	**100**	**100**	**100**	**100**	**100**	**90**	78	30	2	0
	NG	**92**	**85**	64	67	60	**100**	**100**	**100**	**100**	**100**	92	80	80	80	8
	RHS (p_0=5)	85	66	44	41	23	**100**	**100**	**100**	**100**	**100**	92	78	80	**90**	84
0.5	R2-GP	99	84	21	9	0	**100**	**100**	**100**	**100**	**100**	94	**84**	**92**	**90**	**90**
	Normal	95	95	**97**	**96**	**98**	**100**	**100**	**100**	**100**	**100**	**92**	78	52	8	0
	NG	94	**90**	66	63	64	**100**	**100**	**100**	**100**	**100**	94	80	**88**	80	46
	RHS (p_0=5)	**88**	75	47	40	25	**100**	**100**	**100**	**100**	**100**	94	**84**	94	82	94
0.7	R2-GP	98	82	18	6	1	**100**	**100**	**100**	**100**	**100**	**92**	**88**	82	**90**	**90**
	Normal	94	95	**97**	**94**	**99**	**100**	**100**	**100**	**100**	**100**	**92**	82	68	18	0
	NG	98	**91**	66	56	55	**100**	**100**	**100**	**100**	**100**	**92**	82	86	**90**	76
	RHS (p_0=5)	**88**	76	45	38	26	**100**	**100**	**100**	**100**	**100**	**92**	84	**88**	86	**90**

Table 12: Relative RMSE, Sparse DGP, SE Kernel. As in Table 3, with the RMSE of α taken separately across the active and the inactive components, and the RMSE of ρ across the active components. At $D = 5$ all components are active.

		α, active					α, inactive					ρ					σ				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	***0.097***	***0.118***	*0.179*	*0.203*	*0.242*	—	*0.158*	*0.103*	*0.086*	*0.060*	*0.022*	*0.022*	*0.017*	*0.016*	*0.013*	*0.053*	*0.082*	***0.084***	***0.077***	***0.102***
	Normal	1.57	1.30	**0.95**	**0.80**	**0.35**	—	1.40	2.34	2.98	4.35	1.06	1.17	1.39	1.42	0.95	**0.98**	**0.93**	2.66	7.94	7.56
	NG	1.52	1.39	1.21	1.01	0.72	—	0.74	0.56	0.70	1.65	0.93	1.00	1.07	1.18	1.67	1.00	1.05	1.24	1.97	6.88
	RHS (p_0=5)	1.64	1.49	1.23	1.14	1.08	—	**0.54**	**0.47**	**0.44**	**0.39**	**0.88**	**0.87**	**0.82**	**0.78**	**0.83**	1.11	1.15	1.18	1.23	1.15
0.5	R2-GP	***0.092***	***0.112***	***0.180***	*0.205*	*0.243*	—	*0.163*	*0.103*	*0.083*	*0.058*	*0.016*	*0.014*	*0.012*	*0.011*	*0.008*	*0.053*	***0.067***	***0.073***	***0.075***	***0.081***
	Normal	1.96	1.64	**1.00**	**0.90**	**0.52**	—	1.49	2.60	3.55	5.09	0.99	1.02	1.11	1.16	2.08	**0.97**	1.03	2.03	4.92	9.85
	NG	1.65	1.41	1.17	1.13	0.77	—	0.82	0.62	0.48	1.44	0.92	0.91	0.92	0.95	1.70	1.00	1.04	**1.00**	1.04	5.30
	RHS (p_0=5)	1.76	1.50	1.17	1.11	1.06	—	**0.58**	**0.52**	**0.46**	**0.42**	**0.87**	**0.81**	**0.83**	**0.81**	**0.92**	1.07	1.11	1.01	1.02	1.04
0.7	R2-GP	***0.095***	***0.118***	***0.183***	*0.208*	*0.243*	—	*0.160*	*0.100*	*0.080*	*0.057*	*0.011*	*0.010*	*0.009*	*0.008*	*0.006*	*0.053*	***0.066***	*0.069*	*0.069*	*0.074*
	Normal	2.07	1.64	1.07	**0.97**	**0.69**	—	1.67	2.89	3.90	5.74	0.95	0.98	1.05	1.04	2.52	**0.98**	**1.00**	1.64	3.56	10.69
	NG	1.70	1.34	1.18	1.18	0.94	—	0.85	0.60	**0.46**	0.95	0.91	0.91	0.89	0.93	1.36	1.00	1.01	1.02	**0.91**	2.52
	RHS (p_0=5)	1.79	1.47	1.14	1.09	1.05	—	**0.59**	**0.53**	0.47	**0.44**	**0.86**	**0.80**	**0.88**	**0.81**	**0.89**	1.05	1.11	**0.97**	0.98	**0.95**

Table 13: Sampling Diagnostics, Sparse DGP, SE Kernel. As in Table 4.

		$\widehat{R}$					ESS					Divergences					Time (s)				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1040**	**1003**	**843**	**645**	463	**0.0**	**0.0**	**0.0**	**0.5**	**11**	**6.2**	**9.1**	**26**	**49**	**124**
	Normal	**1.00**	1.01	**1.01**	1.18	1.29	946	901	503	51	20	1.4	1.7	3.9	251	253	8.0	15	75	437	1377
	NG	1.01	1.03	1.15	1.32	1.24	524	243	34	17	43	5.7	28	95	110	81	12	38	212	356	904
	RHS (p_0=5)	**1.00**	1.01	**1.01**	1.02	1.02	904	832	627	417	**487**	3.0	3.4	5.0	12	21	8.6	13	35	65	156
0.5	R2-GP	**1.00**	**1.00**	**1.00**	**1.01**	**1.01**	**1191**	**1186**	**1087**	**960**	**648**	**0.1**	**0.0**	**0.0**	**0.1**	**1.9**	**7.3**	**11**	**32**	**53**	**137**
	Normal	**1.00**	**1.00**	1.01	1.02	1.51	1054	1047	836	250	12	2.4	1.5	2.3	70	222	10	19	87	303	1331
	NG	1.03	1.02	1.11	1.20	1.59	594	293	42	23	19	11	20	121	146	92	13	33	224	435	867
	RHS (p_0=5)	**1.00**	1.01	1.01	**1.01**	1.03	1010	925	689	625	547	6.5	4.2	11	16	40	10	16	40	72	169
0.7	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1290**	**1266**	**1204**	**1115**	**811**	**0.7**	**0.1**	**0.1**	**0.1**	**1.4**	**8.3**	**13**	**37**	**64**	**167**
	Normal	**1.00**	**1.00**	**1.01**	**1.01**	1.50	1094	1108	949	562	11	4.0	1.7	2.7	9.2	181	12	26	123	308	1303
	NG	1.01	1.03	1.11	1.22	1.59	674	292	46	22	9	25	53	161	175	89	16	44	260	467	1047
	RHS (p_0=5)	**1.00**	1.01	**1.01**	**1.01**	**1.01**	1067	1027	828	763	612	9.1	5.8	16	11	21	12	18	50	85	199

Table 14: Interval Coverage (%), Sparse DGP, Matérn 3/2 Kernel. As in Table 11.

		α					ρ					σ				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	98	80	27	11	0	**100**	**100**	**100**	**100**	**100**	**94**	**88**	98	**88**	**94**
	Normal	95	**95**	**94**	**97**	**100**	**100**	**100**	**100**	**100**	**100**	98	86	2	0	0
	NG	95	84	70	69	64	**100**	**100**	**100**	**100**	**100**	98	86	86	64	0
	RHS (p_0=5)	**88**	68	49	33	24	**100**	**100**	**100**	**100**	**100**	82	86	**90**	66	80
0.5	R2-GP	97	80	18	6	0	**100**	**100**	**100**	**100**	**100**	**88**	**90**	**88**	**92**	**90**
	Normal	93	**93**	**96**	**98**	**100**	**100**	**100**	**100**	**100**	**100**	84	86	26	0	0
	NG	94	85	60	59	62	**100**	**100**	**100**	**100**	**100**	84	86	**92**	78	18
	RHS (p_0=5)	**88**	70	38	33	30	**100**	**100**	**100**	**100**	**100**	86	88	94	82	82
0.7	R2-GP	98	80	19	7	0	**100**	**100**	**100**	**100**	**100**	**88**	**86**	**96**	96	**86**
	Normal	94	93	**96**	**96**	**100**	**100**	**100**	**100**	**100**	**100**	**88**	82	40	0	0
	NG	94	**88**	59	58	64	**100**	**100**	**100**	**100**	**100**	**88**	84	**96**	**90**	22
	RHS (p_0=5)	**87**	71	40	34	32	**100**	**100**	**100**	**100**	**100**	84	84	**96**	80	82

Table 15: Relative RMSE, Sparse DGP, Matérn 3/2 Kernel. As in Table 12.

		α, active					α, inactive					ρ					σ				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	***0.087***	***0.121***	*0.176*	*0.207*	*0.244*	—	*0.163*	*0.109*	*0.085*	*0.060*	*0.019*	*0.016*	*0.015*	*0.012*	*0.010*	*0.056*	***0.066***	***0.079***	***0.106***	***0.103***
	Normal	1.61	1.14	**0.93**	**0.66**	**0.27**	—	1.35	2.21	2.72	4.15	1.07	1.15	1.36	1.38	1.15	**0.94**	1.03	4.79	6.91	7.46
	NG	1.56	1.34	1.11	0.94	0.70	—	0.77	0.72	0.93	1.62	0.97	0.98	1.09	1.19	1.68	0.97	1.05	1.49	2.45	7.20
	RHS (p_0=5)	1.70	1.51	1.24	1.17	1.10	—	**0.55**	**0.48**	**0.43**	**0.36**	**0.87**	**0.83**	**0.82**	**0.72**	**0.83**	1.17	1.20	1.08	1.21	1.36
0.5	R2-GP	***0.096***	***0.116***	*0.187*	*0.213*	*0.235*	—	*0.163*	*0.100*	*0.081*	*0.064*	*0.012*	*0.011*	*0.009*	*0.007*	*0.008*	***0.065***	*0.067*	***0.066***	***0.079***	***0.123***
	Normal	1.78	1.43	**0.79**	**0.70**	**0.38**	—	1.36	2.45	3.35	4.13	0.98	0.94	1.16	1.28	1.25	1.02	1.00	3.54	7.64	6.29
	NG	1.60	1.44	1.14	1.06	0.73	—	0.75	0.63	0.63	1.54	0.94	0.90	0.93	1.05	1.40	1.02	**0.99**	1.30	1.48	5.17
	RHS (p_0=5)	1.68	1.55	1.20	1.13	1.07	—	**0.55**	**0.49**	**0.43**	**0.42**	**0.89**	**0.80**	**0.78**	**0.79**	**0.89**	1.08	1.13	1.10	1.42	1.02
0.7	R2-GP	***0.097***	***0.117***	*0.187*	*0.213*	*0.235*	—	*0.162*	*0.099*	*0.079*	*0.063*	*0.007*	*0.007*	*0.006*	*0.005*	*0.005*	*0.063*	*0.066*	***0.054***	***0.068***	*0.114*
	Normal	1.82	1.53	**0.86**	**0.81**	**0.48**	—	1.45	2.57	3.67	4.49	0.90	0.95	1.12	1.28	1.79	**0.97**	0.99	3.30	6.83	6.83
	NG	1.59	1.45	1.15	1.11	0.74	—	0.78	0.60	0.55	1.49	0.87	0.91	1.00	0.98	1.47	1.00	**0.97**	1.13	1.24	4.51
	RHS (p_0=5)	1.66	1.53	1.18	1.11	1.06	—	**0.55**	**0.49**	**0.44**	**0.43**	**0.84**	**0.82**	**0.83**	**0.83**	**0.88**	1.09	1.11	**1.00**	1.37	**0.96**

Table 16: Sampling Diagnostics, Sparse DGP, Matérn 3/2 Kernel. As in Table 4.

		$\widehat{R}$					ESS					Divergences					Time (s)				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1049**	**1029**	**759**	**555**	**426**	**0.0**	**0.0**	**0.1**	**4.1**	**6.2**	**7.9**	**14**	**45**	**86**	**282**
	Normal	**1.00**	1.01	1.06	1.28	1.37	925	840	209	25	14	0.5	0.6	59	248	233	9.8	21	188	847	3109
	NG	1.01	1.03	1.15	1.36	1.15	518	240	40	23	41	2.8	29	61	95	77	16	53	312	559	1934
	RHS (p_0=5)	1.01	1.01	**1.01**	**1.01**	1.03	849	750	458	475	416	2.0	0.9	16	11	41	11	18	58	107	345
0.5	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1181**	**1156**	**1020**	**750**	**483**	**0.1**	**0.0**	**0.1**	**0.1**	**6.0**	**9.1**	**15**	**48**	**89**	324
	Normal	**1.00**	**1.00**	**1.01**	1.19	1.33	1041	1000	576	40	18	0.8	1.4	9.2	201	223	12	25	146	738	2664
	NG	1.01	1.04	1.14	1.24	1.32	599	271	47	22	35	5.6	35	114	151	88	18	53	333	654	1844
	RHS (p_0=5)	**1.00**	1.01	**1.01**	1.02	1.02	984	888	704	542	419	2.3	3.5	9.1	28	39	13	23	60	107	**323**
0.7	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1239**	**1239**	**1121**	**904**	**581**	**0.5**	**0.0**	**0.1**	**0.4**	**3.4**	**11**	**18**	**53**	**101**	**372**
	Normal	**1.00**	1.01	**1.01**	1.07	1.37	1069	1057	782	112	17	2.1	1.0	2.4	170	204	15	32	162	666	2537
	NG	1.01	1.02	1.12	1.20	1.41	693	310	43	25	21	17	40	106	165	108	22	65	382	720	1967
	RHS (p_0=5)	**1.00**	1.01	**1.01**	**1.01**	1.03	1014	964	748	633	463	5.1	5.7	8.2	15	43	15	25	68	130	466

Table 17: Interval Coverage (%), Sparse DGP, Matérn 5/2 Kernel. As in Table 11.

		α					ρ					σ				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	92	82	23	14	1	**100**	**100**	**100**	**100**	**100**	82	**92**	94	98	**92**
	Normal	92	96	**96**	**96**	**100**	**100**	**100**	**100**	**100**	**100**	**88**	86	12	0	0
	NG	**90**	**86**	66	68	65	**100**	**100**	**100**	**100**	**100**	84	84	84	66	2
	RHS (p_0=5)	81	67	42	40	26	**100**	**100**	**100**	**100**	**100**	82	86	**90**	**92**	74
0.5	R2-GP	98	74	18	13	0	**100**	**100**	**100**	**100**	**100**	**90**	**92**	94	**90**	**86**
	Normal	94	**93**	**97**	**92**	**100**	**100**	**100**	**100**	**100**	**100**	88	**88**	46	4	0
	NG	94	84	63	64	64	**100**	**100**	**100**	**100**	**100**	86	**92**	**90**	86	10
	RHS (p_0=5)	**88**	66	43	44	28	**100**	**100**	**100**	**100**	**100**	88	**88**	80	94	78
0.7	R2-GP	98	76	19	10	0	**100**	**100**	**100**	**100**	**100**	88	94	94	**90**	86
	Normal	93	94	**95**	**93**	**99**	**100**	**100**	**100**	**100**	**100**	**90**	**90**	52	10	0
	NG	94	**87**	62	61	70	**100**	**100**	**100**	**100**	**100**	88	**90**	**88**	86	34
	RHS (p_0=5)	**89**	70	42	45	32	**100**	**100**	**100**	**100**	**100**	86	92	86	**90**	**88**

Table 18: Relative RMSE, Sparse DGP, Matérn 5/2 Kernel. As in Table 12.

		α, active					α, inactive					ρ					σ				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	***0.100***	***0.114***	*0.183*	*0.201*	*0.240*	—	*0.162*	*0.103*	*0.087*	*0.062*	*0.020*	*0.018*	*0.015*	*0.014*	*0.011*	***0.067***	***0.070***	***0.069***	***0.085***	***0.124***
	Normal	1.57	1.33	**0.84**	**0.79**	**0.30**	—	1.32	2.31	2.76	4.04	1.01	1.17	1.26	1.41	1.05	**1.00**	1.12	4.31	8.10	6.19
	NG	1.54	1.39	1.12	0.99	0.69	—	0.77	0.67	0.83	1.57	0.92	1.01	0.98	1.33	1.77	1.02	1.05	1.43	2.47	5.98
	RHS (p_0=5)	1.62	1.51	1.22	1.16	1.09	—	**0.57**	**0.46**	**0.44**	**0.40**	**0.88**	**0.91**	**0.78**	**0.85**	**0.80**	1.09	1.08	1.25	1.21	1.03
0.5	R2-GP	***0.099***	***0.125***	*0.186*	*0.195*	*0.237*	—	*0.158*	*0.098*	*0.089*	*0.061*	*0.013*	*0.012*	*0.010*	*0.010*	*0.008*	***0.055***	***0.060***	***0.064***	*0.077*	*0.114*
	Normal	1.83	1.28	**0.84**	**0.99**	**0.39**	—	1.46	2.55	3.13	4.42	0.94	1.02	1.17	1.32	1.43	1.02	1.06	2.68	6.89	6.87
	NG	1.60	1.34	1.15	1.16	0.72	—	0.77	0.52	0.62	1.59	0.90	0.90	0.93	1.08	1.48	1.01	1.03	1.19	1.40	5.28
	RHS (p_0=5)	1.68	1.47	1.19	1.11	1.07	—	**0.54**	**0.46**	**0.51**	**0.42**	**0.85**	**0.81**	**0.74**	**0.87**	**0.85**	1.04	1.09	1.23	**0.97**	**0.99**
0.7	R2-GP	***0.100***	***0.126***	*0.188*	***0.197***	*0.234*	—	*0.156*	*0.096*	*0.087*	*0.062*	*0.009*	*0.008*	*0.007*	*0.007*	*0.006*	***0.053***	***0.061***	***0.063***	*0.074*	*0.102*
	Normal	1.92	1.34	**0.92**	1.05	**0.55**	—	1.56	2.76	3.36	4.74	0.90	0.94	1.06	0.94	2.26	1.01	**1.00**	2.07	4.93	7.88
	NG	1.62	1.32	1.16	1.20	0.80	—	0.80	0.52	0.56	1.27	0.89	0.89	0.92	0.87	1.16	**1.00**	1.01	1.17	1.15	3.79
	RHS (p_0=5)	1.72	1.45	1.17	1.10	1.04	—	**0.55**	**0.47**	**0.51**	**0.44**	**0.86**	**0.83**	**0.76**	**0.81**	**0.86**	1.01	1.05	1.20	**0.92**	**0.98**

Table 19: Sampling Diagnostics, Sparse DGP, Matérn 5/2 Kernel. As in Table 4.

		$\widehat{R}$					ESS					Divergences					Time (s)				
ρ	Prior	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100	5	10	30	50	100
0.3	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.02**	**1033**	**1032**	**745**	**582**	**386**	**0.0**	**0.0**	**0.1**	**2.1**	**20**	**7.6**	**12**	**37**	**75**	**234**
	Normal	**1.00**	1.01	1.02	1.22	1.34	930	867	368	30	18	1.0	0.5	32	228	231	9.7	20	130	768	2202
	NG	1.01	1.04	1.14	1.31	1.13	536	239	42	23	46	10	13	65	81	43	20	47	304	484	1628
	RHS (p_0=5)	1.01	1.01	**1.01**	1.02	**1.02**	879	768	482	432	376	2.2	1.8	9.7	14	26	11	18	48	95	295
0.5	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1163**	**1171**	**1036**	**812**	**570**	**0.2**	**0.0**	**0.1**	**0.3**	**12**	**8.9**	**14**	**41**	**81**	**236**
	Normal	**1.00**	1.01	**1.01**	1.16	1.38	1004	1026	706	77	17	1.8	1.3	2.7	129	189	13	25	131	606	2350
	NG	1.01	1.04	1.14	1.22	1.39	623	233	41	20	25	6.7	43	137	134	74	17	53	364	590	1465
	RHS (p_0=5)	**1.00**	1.01	**1.01**	**1.01**	1.02	977	912	695	548	425	2.6	5.5	7.9	15	31	13	19	56	110	323
0.7	R2-GP	**1.00**	**1.00**	**1.01**	**1.01**	**1.01**	**1234**	**1262**	**1168**	**1006**	**701**	**0.2**	**0.0**	**0.1**	**0.1**	**1.1**	**11**	**18**	**51**	**95**	**311**
	Normal	**1.00**	1.01	**1.01**	1.05	1.56	1069	1083	898	298	12	5.0	2.1	1.4	48	185	15	30	161	573	2116
	NG	1.01	1.05	1.11	1.19	1.65	653	270	44	25	14	13	51	164	146	56	19	67	380	716	1724
	RHS (p_0=5)	**1.00**	1.01	**1.01**	**1.01**	1.02	1010	1018	827	666	496	7.0	7.9	9.2	14	30	15	25	64	131	374

B.3 Non-normalised Competing Priors

All comparisons above apply the competing priors to the normalised components (14). Here we place the normal, NG and RHS priors on α_d in the usual HSGP parameterisation, $f_d = \alpha_d g_d$, in which α_d^2 is the marginal variance. R2-GP remains the reference, and all other settings are unchanged. Since α_d is then on a different scale from the generating value, we compare predictive performance. Figures 17 to 22 show that the qualitative behaviour of the competing priors is the same as with the normalisation.

In the dense DGP, R2-GP attains the highest mean elpd in each condition against every non-normalised prior. At $D = 100$ the mean Δelpd lies between -149 and -105 for the normal prior, between -47 and -33 for NG, and between -9.1 and -4.5 for RHS. Since $Q_d < 1$ for the kernels considered, a given prior on α_d implies less signal variance without the normalisation. This lessens the overestimation under the normal and NG priors, whose mean elpd improves by up to 37 and 24, but does not remove it. For RHS the mean elpd is lower by at most 1.4 in both DGPs. In the sparse DGP, RHS again attains a higher mean elpd than R2-GP for $D \geq 30$, by at most 12.1, and the normal prior again underperforms, with a mean Δelpd between -150 and -106 at $D = 100$. At $D = 100$ the bias in R^2_{emp} has the same sign and a similar range for each prior as with the normalisation, and the normal and NG priors again fail to converge (not shown).

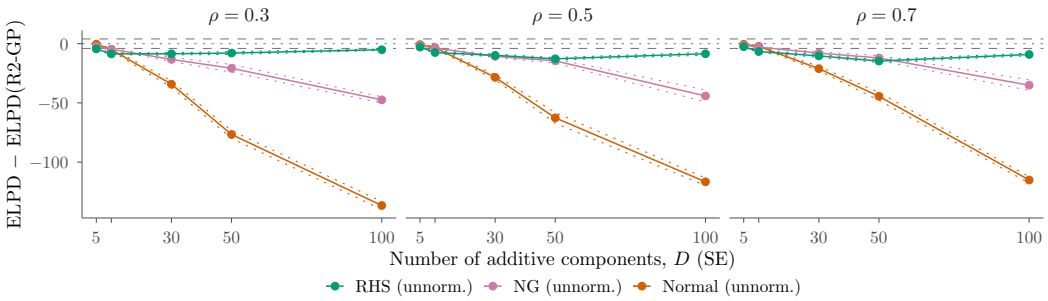

Figure 17: Non-normalised Competing Priors, Dense DGP, SE Kernel. As in Figure 5, with the normal, NG and RHS priors on α_d in $f_d = \alpha_d g_d$.

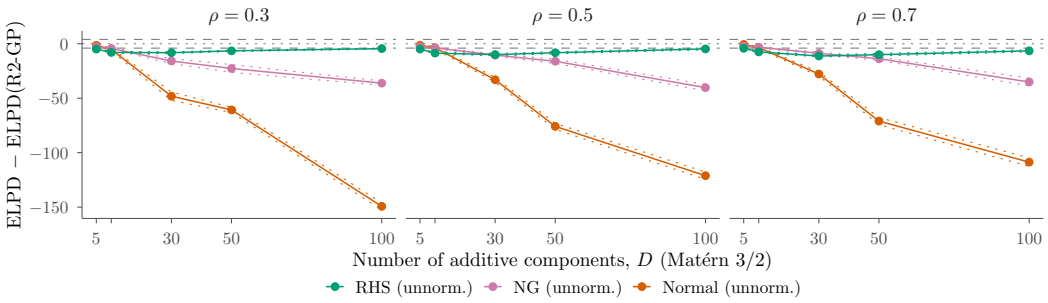

Figure 18: Non-normalised Competing Priors, Dense DGP, Matérn 3/2 Kernel. As in Figure 17.

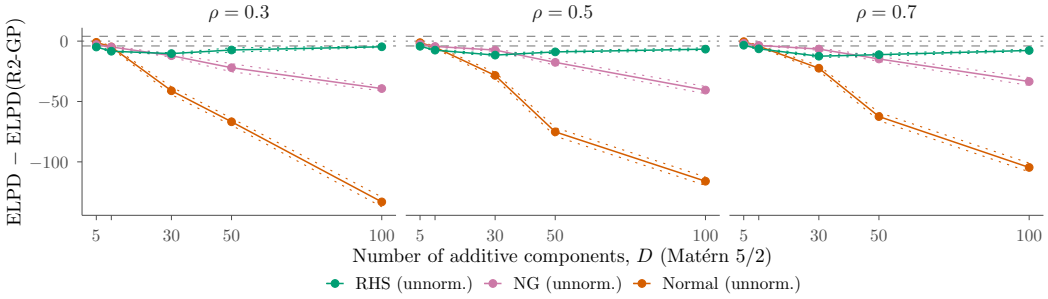

Figure 19: Non-normalised Competing Priors, Dense DGP, Matérn 5/2 Kernel. As in Figure 17.

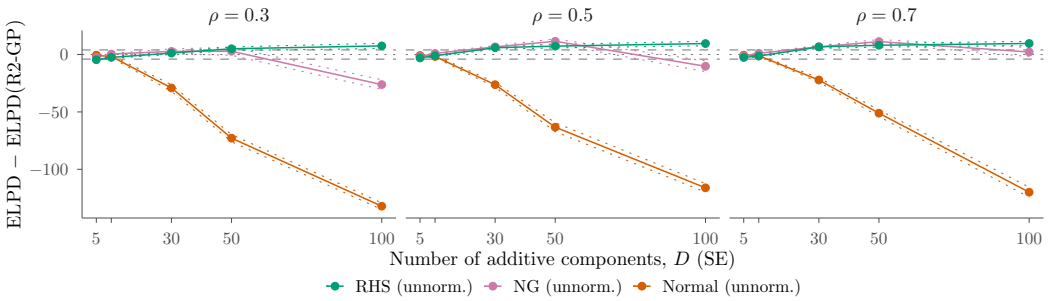

Figure 20: Non-normalised Competing Priors, Sparse DGP, SE Kernel. As in Figure 17.

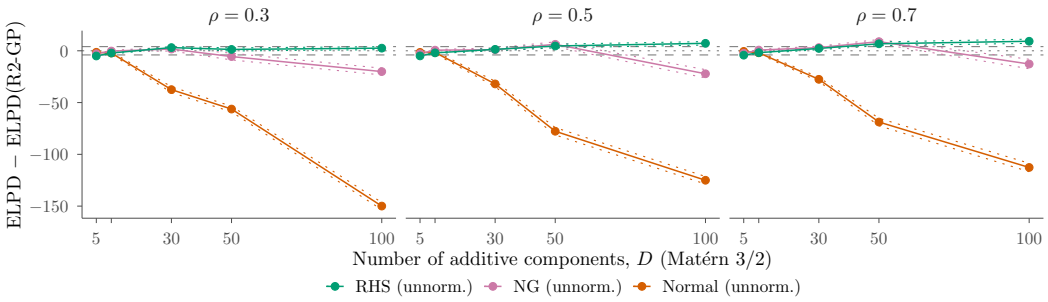

Figure 21: Non-normalised Competing Priors, Sparse DGP, Matérn 3/2 Kernel. As in Figure 17.

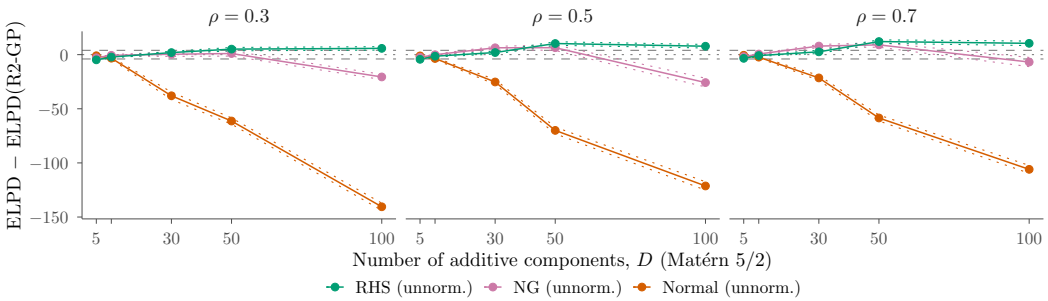

Figure 22: Non-normalised Competing Priors, Sparse DGP, Matérn 5/2 Kernel. As in Figure 17.

C Additional Case Study Results

Figures 23 and 24 extend Figure 4 to 12 covariates, for the SE and the Matérn 3/2 kernel. The covariates are ranked by their posterior mean share of explained variance under R2-GP with the SE kernel, and both figures show the same covariates. For these figures, the models of Section 4 are fitted to all 208 compounds. Table 20 gives the predictive performance and the sampling diagnostics over the 10 train/test splits.

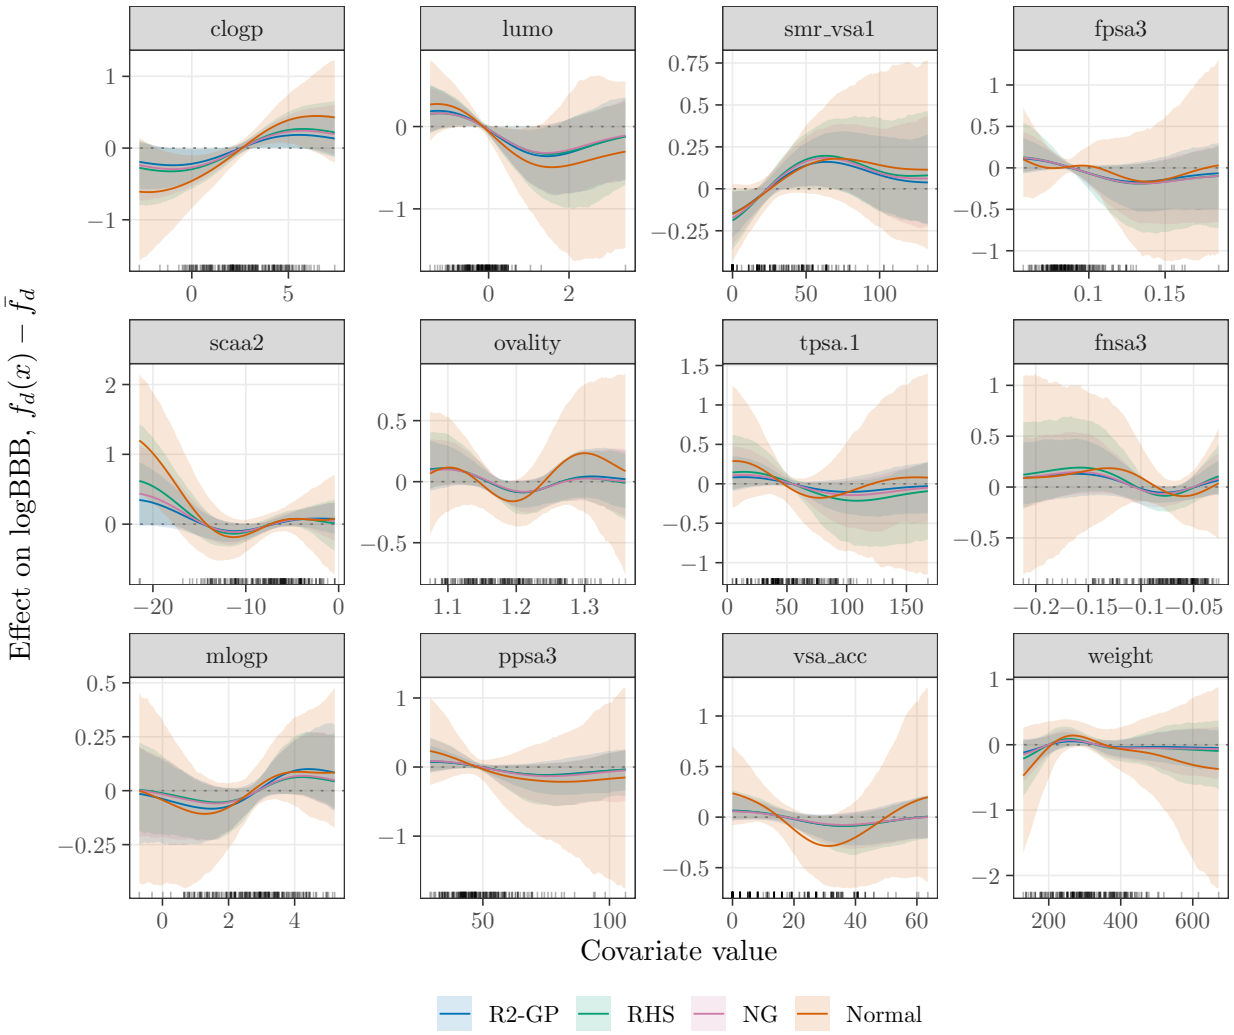

Figure 23: Covariate Effects, BloodBrain Data, SE Kernel. Posterior means (lines) and pointwise 90% intervals (bands) of the centred effects $f_d(x) - \bar{f}_d$ for the 12 covariates with the largest share of explained variance. Ticks mark the observed covariate values.

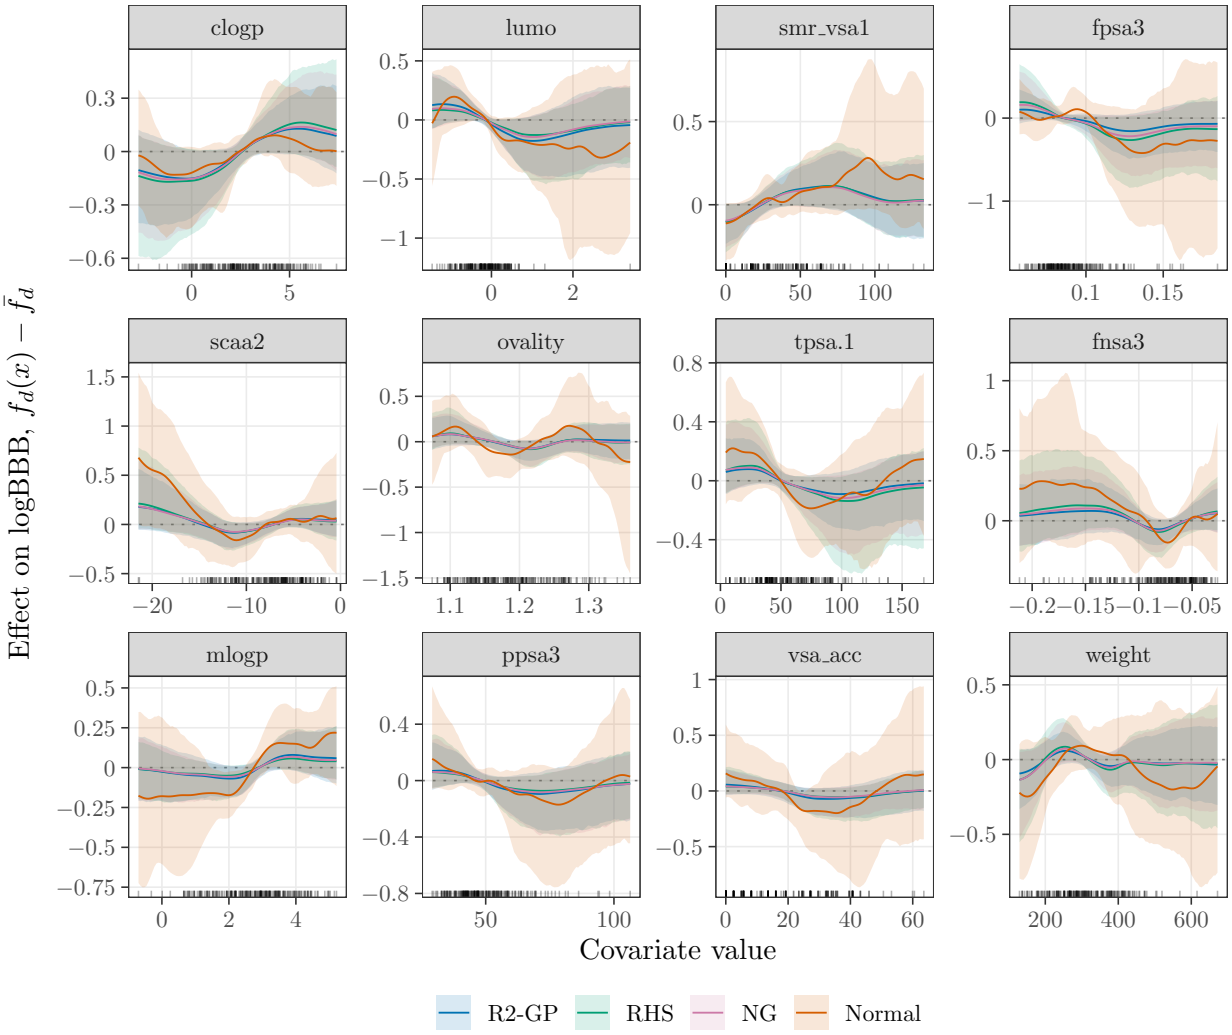

Figure 24: Covariate Effects, BloodBrain Data, Matérn 3/2 Kernel. As in Figure 23, for the same covariates.

Table 20: Case Study Predictive and Sampling Diagnostics

	SE							Matérn 3/2						
Prior	ELPD	R^2	R^2_{oos}	$\widehat{R}$	ESS	Div.	Time	ELPD	R^2	R^2_{oos}	$\widehat{R}$	ESS	Div.	Time
R2-GP	−40.9	0.73	0.56	1.008	529	0.0	51	−40.2	0.77	0.56	1.016	479	0.0	126
RHS	−41.4	0.72	0.55	1.023	178	0.2	72	−41.0	0.76	0.55	1.023	150	0.0	176
NG	−41.3	0.70	0.55	1.232	14	0.0	61	−40.7	0.74	0.55	1.037	71	0.0	126
Normal	−46.0	0.95	0.56	2.117	5	2.5	100	−48.9	1.00	0.55	3.462	4	23.8	257

Note: Predictive performance and sampling behaviour of fitted models on the BloodBrain data over 10 random train/test split replications. ELPD and posterior R^2 (in-sample and out-of-sample) are reported as means over the replications. $\widehat{R}$ and ESS are reported as the worst case over all model parameters across replications; Div. is the mean number of divergent transitions per fit and time the median sampling time per fit in minutes.

D Computing Infrastructure

All computations ran on CPUs; no GPUs were used. The simulation study was run on an institutional high-performance computing cluster managed by Slurm, on nodes with Intel Xeon (Haswell to Cascade Lake) and AMD EPYC (Milan) processors. Each fit was allocated four cores, one per chain, and 3 GB of memory, and the workflow was orchestrated with Snakemake 8.18. Models were fitted with CmdStan 2.35.0 (Carpenter et al., 2017) through the R package *cmdstanr* 0.9.0 in R 4.5.3, compiled with GCC 12.3.

The simulation comprises 57 750 model fits, all of which completed. Sampling took 1.3 minutes per fit at the median and 5.9 minutes on average; 99% of the fits finished within about an hour and the slowest took 4.7 hours. In total this is about 5 700 hours of four-core time, or 23 000 core-hours. The case study adds 80 fits on the data splits, which took between 27 minutes and 4.5 hours each (median 1.7 hours), about 650 core-hours in total, and eight fits to the full data.