arXiv:2610.35280v1 [cs.AI] 26 Oct 2026

Self-Supervised Pretraining for PHM: What Transfers, What Does Not, and What Federating It Costs

José Sanz-Durán David Solís-Martín
Department of Computer Science and Artificial Intelligence
Universidad de Sevilla, Seville, Spain
jsandur@outlook.com dsolis@us.es

Abstract

Self-supervised pretraining has produced capable time-series foundation models, but its value for Prognostics and Health Management (PHM) has not been measured under leakage-controlled conditions. We measure it on a corpus of 47 public PHM datasets, harmonized to a single tensor contract and partitioned so that 11 datasets are withheld from pretraining in full. Transfer turns out to be task-dependent: a compact channel-independent Transformer encoder (801,808 parameters) pretrained with masked patch prediction raises fault-classification macro-F1 from 0.350 to 0.812 on CWRU and from 0.569 to 0.951 on HSG18, yet on CMAPSS remaining-useful-life estimation no pretrained checkpoint reliably improves upon a from-scratch baseline. We pretrain under two regimes—centralized, and simulated cross-silo federated with ten clients grouped by physical domain—holding architecture, objective, source data and optimization budget fixed. Two further findings are negative. At matched budget the centralized encoder leads by 15.5–20.2 macro-F1 points in four of four comparisons, although the comparison is confounded by a learning-rate schedule that favours the centralized arm and rests on a single pretraining run per regime; on RUL the two regimes are indistinguishable. And the intrinsic SSL loss is not a valid selection criterion: the centralized encoder more than halves the federated reconstruction loss without a corresponding gap on RUL, and a server-momentum sweep moves downstream validation macro-F1 by up to 0.29 while the SSL loss varies by 0.021. Transfer also hinges on the backbone learning rate: the same checkpoint yields 0.169 or 0.829 macro-F1 on CWRU depending on a single order of magnitude. We release the corpus pipeline, the protocol and the code.

Keywords: prognostics and health management; self-supervised learning; time-series foundation models; federated learning; transfer learning; negative results.

1 Introduction

Industrial maintenance has shifted from corrective and calendar-based strategies towards condition-based and predictive policies. Prognostics and Health Management (PHM) provides the analytical frame for that shift: it turns sensor signals into diagnosis of the current health state and prognosis of the future one, typically instantiated as fault classification, remaining useful life (RUL) estimation and anomaly detection [1, 2]. The practical bottleneck is well known. Raw condition-monitoring signals are abundant; labelled failures are scarce, expensive and often ambiguous across organizations. That asymmetry—abundant unlabelled signals, scarce labels, high heterogeneity—is precisely the setting in which self-supervised pretraining has succeeded in language and vision, and more recently in general time series [5, 6, 7].

Two obstacles separate that promise from industrial practice. The first is heterogeneity. A PHM corpus mixes vibration sampled at tens of kilohertz with per-cycle counters, one channel

with three hundred, and run-to-failure trajectories with steady-state logs; target semantics, in addition, are not comparable across benchmarks. A model trained dataset by dataset does not scale to that diversity; the question is whether a *single* temporal representation can be learned across it and reused. The second obstacle is governance. Condition data belong to different operators and frequently cannot be centralized for reasons of confidentiality, ownership or regulation. Federated learning addresses this in principle [11], but its combination with multi-dataset self-supervised pretraining for PHM has, to our knowledge, not been evaluated empirically under a controlled protocol.

There is also a measurement problem specific to this literature. Recent work has documented that time-series foundation model evaluations are vulnerable to information leakage, both through reuse of datasets across pretraining and testing and through overlapping windows split at random [18]. When windows are generated with stride smaller than their length—the norm in PHM—a random split places nearly identical samples on both sides of the train/test boundary and inflates results. Any claim about transfer therefore depends as much on the evaluation protocol as on the model.

This paper is an empirical study, not an architecture proposal. We deliberately fix a small, standard backbone and vary only what we want to measure. Our contributions are:

1. **A leakage-controlled multi-domain PHM corpus and protocol.** We download 53 public datasets, audit 51, harmonize 47 to a common channel-independent tensor contract, and partition them so that 11 datasets are withheld *entirely* from pretraining. Splits are made by physical unit before windowing, normalization is per instance, and automated anti-leakage and round-trip checks pass on 47 of 47 datasets (Section 3).
2. **A budget-matched comparison of centralized and federated self-supervised pretraining.** Both regimes share architecture, objective, source data and optimization budget; only the optimization regime differs. This isolates the cost of not centralizing (Sections 4 and 6).
3. **Three empirical findings, reported with their limits.** Transfer is task-dependent and, on RUL, absent; federated pretraining retains part of the representation but does not match the centralized encoder in classification; and the intrinsic pretraining loss does not predict downstream utility, which invalidates it as a model-selection signal. We also quantify how strongly the adaptation learning rate governs whether transfer survives at all.

We consider the negative results to be the more useful half of this paper. Published evidence on when multi-dataset self-supervised pretraining *fails* to transfer in industrial settings is thin, and the failure modes we observe—task dependence, metric decoupling, and destruction of the representation by an overly aggressive fine-tuning rate—are specific, reproducible and actionable.

2 Related work

Deep learning for PHM. Surveys agree on both the capability and the fragmentation of the field [3, 2, 1, 4]. Architectures, preprocessing and splits are typically designed per task and evaluated on a single benchmark with intra-dataset protocols. This produces strong specialists and weak evidence about reusable representations: it is difficult to tell whether a reported gain reflects a transferable inductive bias or specialization to one controlled environment. The problem of building systems that transfer across machines, domains and sensor configurations remains open.

Time-series foundation models. PatchTST [5] introduced two design choices that matter here: *patching*, which reduces the effective sequence length and provides a more semantic token than an individual sample, and *channel independence*, which applies a shared encoder to each

channel separately so that no layer depends on the number of sensors. MOMENT [6] showed that masked patch prediction over a large multi-dataset corpus yields a general-purpose encoder evaluated under limited supervision. TimesFM [7] and Time-LLM [8] broaden the design space towards forecasting-oriented decoders and reprogrammed language models. Surveys organize the field along modality, backbone, pretraining objective and adaptation mechanism [9, 10]. None of these corpora or benchmarks is built for PHM, and none addresses federated pretraining; we borrow their mechanisms rather than proposing new ones.

Federated learning under heterogeneity. FedAvg [11] averages locally trained weights and degrades under non-IID client distributions through client drift. FedProx [12] adds a proximal term that penalizes divergence from the global model; SCAFFOLD [13] corrects drift with control variates; server-side momentum (FedAvgM) [14] smooths the aggregation trajectory. Applications to PHM remain narrow in scope, often partitioning a single dataset artificially, which understates the heterogeneity of a realistic cross-silo setting where each silo is a distinct physical domain [15, 16].

Evaluation and leakage. TSFM-Bench [17] argues for standardized splits and adaptation protocols. Meyer et al. [18] document leakage in TSFM evaluation, both from dataset reuse across phases and from temporally overlapping splits. Our protocol targets both failure modes, and Section 3 reports the outcome of the automated checks.

Positioning. The gap we address is the intersection: a temporal self-supervised encoder trained over a genuinely heterogeneous PHM corpus, compared across centralized and federated regimes, and evaluated on datasets withheld in full under a protocol that controls leakage and calibrates adaptation. We do not claim a new architecture, nor state-of-the-art results per benchmark; we claim a controlled measurement of when this recipe works and when it does not.

3 Corpus and harmonization

3.1 Corpus construction and partition

We start from PHMD [19], a library that unifies access to public PHM datasets. Of 55 candidate time-series datasets, 53 were downloaded (59.77 GiB compressed); 51 could be audited (two failed to load because of external library defects) and 47 satisfied the structural contract. Four datasets were dropped by explicit, pre-registered criteria—median trajectory length below the window size, fewer than 100 windows, or tabular behaviour with extreme padding—and none was dropped because of downstream performance.

The 47 accepted datasets are partitioned into **36 source** datasets, used *only* for self-supervised pretraining, and **11 target** datasets, withheld in full and used *only* for downstream evaluation (Figure 1). This partition underpins every claim in the paper (Figure 2): a target dataset never participates in pretraining under either regime, so a downstream gain cannot be attributed to having already seen the domain.

The audit inspects the downloaded data rather than catalogue metadata, and records per dataset: shape and unit structure, missing values and constant channels, target candidates and coverage, temporal granularity, and post-windowing statistics (windows, valid patches, padding ratio, estimated size). Aggregate figures for the harmonized corpus are given in Table 1.

3.2 Tensor contract

Every dataset is mapped to a single interface. A series of a physical unit, (T, C) with C channels, is segmented into windows of $W = 512$ samples with stride 256; each window is divided into $N = 32$

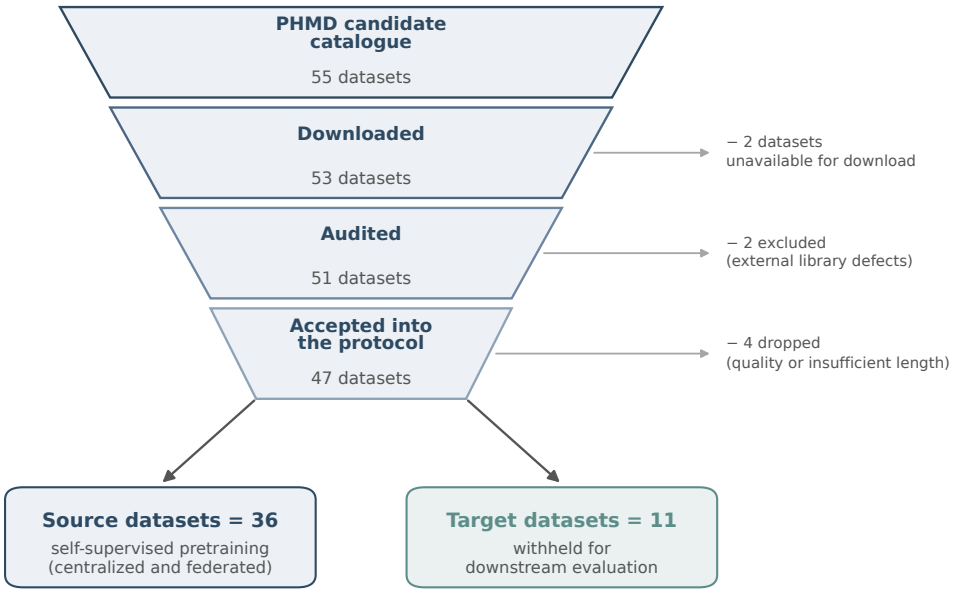

Figure 1: From the PHMD catalogue to experimental roles. Of 55 candidate datasets, 47 satisfy the structural contract and are partitioned into 36 source datasets (pretraining only) and 11 target datasets (evaluation only). No dataset was dropped on the basis of downstream performance.

patches of $P = 16$ samples. Samples are stored as (C, N, P) and batched as (B, C, N, P). Before the encoder the batch is reshaped to $(B{\cdot}C, N, P)$, so the patch embedding $\mathrm{Linear}(P \to d_{\mathrm{model}})$ and every subsequent layer are shared across channels and independent of C. This is what allows one encoder to process a single-channel disk-drive log and a 317-channel industrial record without changing a parameter.

Two details carry methodological weight. First, **validity masks**. A window may be partially filled when a trajectory ends; the tail is padded and two masks are stored, `valid_time_mask` at sample resolution and `valid_patch_mask` at patch resolution. The pretraining loss is evaluated only where a patch is both selected for masking and composed of real samples, so padding never contributes a gradient. Because the last patch of a padded window mixes real and filler samples, patch-level masking alone is insufficient; the loss combines both masks at sample resolution, a behaviour covered by dedicated unit tests. Second, **tail policy**. We retain the final partial window with padding rather than discarding it. In degradation data the end of a trajectory is where failure information lives; discarding tails removed 23.2 % of the raw rows of one target dataset (PHM18) in a preliminary comparison.

Normalization is per window and per channel, computed ignoring padding. This is compatible with the federated setting—it requires no statistics shared across clients—and avoids a leakage channel, at the cost of discarding absolute amplitude, which can itself be diagnostic in PHM. We record the per-window mean and standard deviation on disk so that this choice can be revisited without reprocessing.

3.3 Splits and leakage control

Within each dataset, train, validation and test partitions are formed *by physical unit or trajectory, before windowing*. With $W = 512$ and stride 256 consecutive windows overlap by half, so splitting after windowing would place near-duplicate samples on both sides of the boundary. Where a benchmark provides an official split we respect it, and we construct a global unit identifier from dataset, original split and unit so that identifier collisions across splits cannot occur. Automated

Table 1: Harmonized corpus. Source totals are the quantities that enter pretraining; the target subset is reserved for downstream evaluation only.

Quantity	Value
Candidate datasets in catalogue	55
Downloaded	53
Audited	51
Accepted into the protocol	47
Source (pretraining only)	36
Target (evaluation only)	11
Dropped / excluded	4 / 2
Total windows	5,686,667
Total channel patches	938,903,623
Processed size	158.97 GB
Source valid temporal patches	175,544,903
Source valid channel patches	922,665,279
Anti-leakage checks passed	47 / 47
Round-trip read/write checks passed	47 / 47

checks verify that no unit appears in more than one split, that no trajectory is divided, and that normalization statistics never use test data. These checks pass for all 47 datasets.

4 Model and pretraining

4.1 Backbone

The encoder, which we refer to as PatchTSTPhm, follows PatchTST [5] in its channel-independent patch formulation and MOMENT [6] in its pretraining objective. Each patch is projected by a shared Linear($16 \to 128$), a learnable positional embedding is added, and the sequence passes through a four-layer Transformer encoder with four attention heads, $d_{\text{model}} = 128$ and pre-normalization. A linear head reconstructs patches during pretraining and is replaced by a task head downstream. The model has 801,808 trainable parameters.

The model is deliberately small, for three reasons. It is trainable within the compute budget available to this study; it is cheap to transmit once per federated round, which keeps the federated arm feasible; and, being fixed and compact, it isolates the variable under study—the origin of the weights—from capacity effects. We note explicitly that a 0.8M-parameter encoder is two to three orders of magnitude smaller than published time-series foundation models, and we therefore make no claim about scale.

4.2 Self-supervised objective

Pretraining uses masked patch prediction. For each window, 30 % of the *valid* patches are replaced by a learnable mask token and the model reconstructs their content from the visible context; padding is never masked and never scored. The objective requires no labels, which is what makes the 36 source datasets usable despite their incompatible target semantics.

4.3 Centralized regime

Sampling is weighted by channel patches, capped so that no dataset exceeds 0.10 and no client group exceeds 0.25 of the mixture, with a floor of 0.005 to guarantee that the smallest group participates. The full run executed 99,961 optimizer steps out of 100,000 requested (39 steps skipped by the AMP gradient scaler on non-finite gradients, a rate of 0.04 %) in 3.58 h on a single

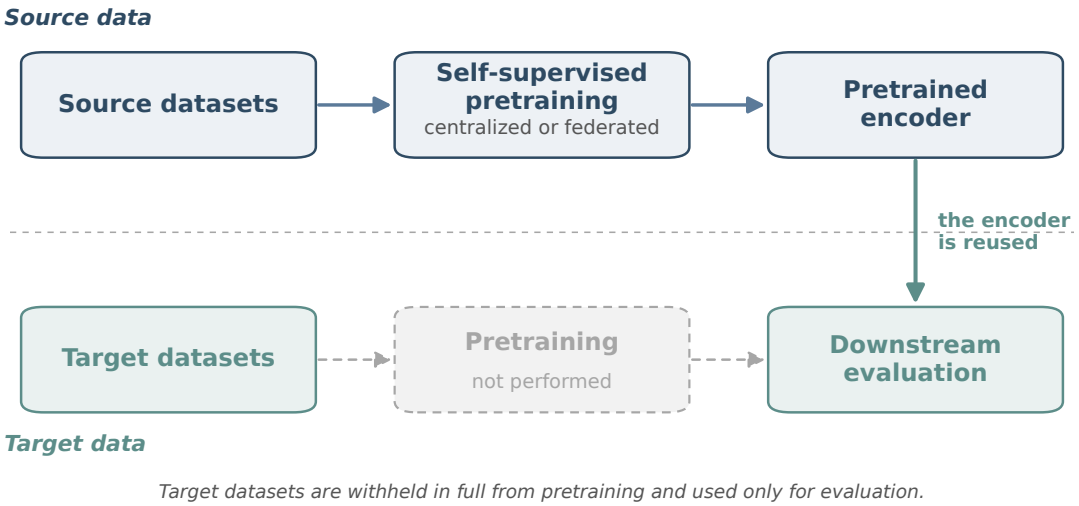

Figure 2: Separation between source and target data. Pretraining, centralized or federated, sees only source datasets; the resulting encoder is then adapted and evaluated on target datasets that never entered pretraining. The greyed-out path is the one the protocol forbids.

A100. Training loss fell from 0.7974 to 0.4167, a reduction of 47.7 %, and the empirical dataset mixture matched the planned caps to within 0.0018. All 36 source datasets were seen.

Central-25k, the budget-matched comparator for the federated arm, is the step-25,000 snapshot of this same run rather than an independent training run. This has a consequence we return to in Section 8: its cosine schedule was dimensioned for 100,000 steps, so at step 25,000 the checkpoint sits in the high-learning-rate phase of a schedule the federated arm does not have. The federated arm uses a constant learning rate of $3{\cdot}10^{-4}$ without a scheduler, a choice made to isolate the effect of server momentum. The two arms therefore match in architecture, objective, source data and number of local optimizer steps, but *not* in optimizer schedule, and the mismatch favours the centralized arm.

4.4 Federated regime

The federated arm simulates a cross-silo setting with ten clients formed by physical domain (bearings, batteries, aero engines, hard disk drives, wind, gearboxes, CNC milling, and two deliberately heterogeneous groups). This is a simulation: there are no real organizations, no network latency, no client dropout and no formal privacy mechanism. What it does reproduce is the property under study, namely genuine non-IID heterogeneity—clients differ in sensors, scales, lengths and failure semantics, not in a random partition of one dataset.

Each round distributes the global model, each client performs 50 local steps, and the server aggregates. With 50 rounds and ten clients this amounts to 25,000 local optimizer steps, matching Central-25k. Aggregation weights inherit the same caps as the centralized sampler, so the two regimes see statistically equivalent data mixtures. We evaluate FedAvg and FedAvgM with server momentum β=0.7, both retained for downstream evaluation, and FedProx with μ=0.01, which was screened and excluded for reasons reported in Section 6.3.

Table 2 reports the intrinsic outcome. At matched budget the centralized encoder reaches less than half the federated validation reconstruction loss. This is expected—the server averages once per round, each client optimizes a partial distribution, and the caps redistribute weight away from dominant domains—but, as we show next, it is a poor predictor of what happens

Table 2: Intrinsic pretraining outcome. SSL validation loss is a diagnostic of the reconstruction objective, not evidence of downstream utility. Central-100k was not evaluated under this protocol, and we do not estimate it. All values are computed on ten batches drawn from five of the 36 source datasets (three of ten clients) and are comparable only among themselves.

Checkpoint	Regime	Budget (steps)	SSL val. loss
Central-100k	centralized	100,000	—
Central-25k	centralized	25,000	0.3467
FedAvg	federated	25,000	0.7844
FedAvgM β=0.7	federated	25,000	0.7636

downstream.

5 Downstream evaluation protocol

We evaluate four pretrained checkpoints—the *checkpoint bank*: Central-100k, Central-25k, FedAvg and FedAvgM β=0.7—on three held-out tasks: fault classification on CWRU (rolling-element bearings, 2 channels, 140,134 windows, 153 units) and HSG18 (hard disk drives, 1 channel, 38,896 windows, 17 units), measured by macro-F1; and RUL regression on CMAPSS (24 channels, comprising 21 sensors and 3 operating-regime variables, 11,795/3,131/6,610 train/validation/test windows, target capped at 125 cycles).

Composition of the evaluation splits. Because splits are formed by physical unit, their class composition is a consequence of the unit structure rather than a design choice, and it is uneven. We state it explicitly, since it bounds the interpretation of every classification number in this paper. The CWRU test split contains 3,331 windows (2.4 % of the dataset) drawn from three of the four fault classes: the fourth class has 6,629 windows in training and none in test. Reported CWRU macro-F1 is therefore an average over the three classes present in test, and the class distribution of that split is markedly skewed (477/2378/476). The HSG18 test split contains 4,576 windows and is exactly balanced (2288/2288). We report macro-F1 throughout for comparability across tasks; on HSG18 it coincides with accuracy and balanced accuracy to five decimal places, whereas on CWRU the three quantities differ and the per-class breakdown, given in Appendix C, is more informative than any single scalar.

We consider four adaptation modes. *From scratch* trains a randomly initialized encoder on the target data alone. *Linear probing* freezes the encoder and trains a single layer on top, which measures representation quality directly. *Full fine-tuning* retrains the whole model with a small backbone step size. For RUL we add a *two-layer MLP head* over a frozen encoder, since a single linear layer proved insufficient for continuous regression.

Three rules govern every number we report. **Selection on validation:** all configuration choices, including the backbone learning rate and the epoch, are made on the validation split; the test split is read once, for traceability only, and never used to choose. **Calibrated learning rate:** the backbone step size is selected on validation *per dataset and per checkpoint* from $\{1{\cdot}10^{-6}, 3{\cdot}10^{-6}, 1{\cdot}10^{-5}\}$, because, as Section 6.4 shows, a single global value silently destroys some checkpoints. **Three seeds:** principal comparisons are repeated with seeds 42, 43 and 44 and reported as mean ± sample standard deviation.

One limitation of the protocol conditions the reading of every table that follows: the *from-scratch* baselines were run with a single seed, whereas the pretrained checkpoints were run with three. Comparisons between them are therefore indicative of effect size but not of statistical reliability.

Table 3: Fault classification, test macro-F1 (higher is better). Pretrained checkpoints: mean ± sample standard deviation over three seeds, full fine-tuning with the backbone learning rate calibrated on validation. *From scratch* and *linear probing* rows are single-seed and are reported as reference points, not as matched comparators.

Configuration	Backbone LR	CWRU	HSG18
From scratch (1 seed)	—	0.350	0.569
Linear probing, Central-100k (1 seed)	frozen	0.705	0.906
Central-100k	$3{\cdot}10^{-6}$ / $1{\cdot}10^{-5}$	0.812 ± 0.033	0.951 ± 0.002
Central-25k	$3{\cdot}10^{-6}$ / $1{\cdot}10^{-5}$	0.734 ± 0.022	0.911 ± 0.008
FedAvg	$3{\cdot}10^{-6}$ / $1{\cdot}10^{-5}$	0.572 ± 0.119	0.742 ± 0.048
FedAvgM β=0.7	$1{\cdot}10^{-6}$ / $1{\cdot}10^{-5}$	0.532 ± 0.079	0.756 ± 0.068

Table 4: Budget-matched comparisons (Central-25k versus federated), test macro-F1, Welch's t-test. The direction favours the centralized encoder in all four cases. Two are nominally significant, but none survives a Bonferroni correction for four comparisons ($\alpha = 0.0125$). The dispersion being tested comes from downstream adaptation seeds over a single pretraining run per arm (Section 8), so these p-values quantify adaptation noise, not the regime.

Dataset	Comparison	Δ	t	df	p
CWRU	Central-25k − FedAvg	+0.162	2.31	2.14	0.139
CWRU	Central-25k − FedAvgM	+0.202	4.25	2.30	0.040
HSG18	Central-25k − FedAvg	+0.169	6.01	2.11	0.023
HSG18	Central-25k − FedAvgM	+0.155	3.92	2.06	0.057

6 Results

6.1 Fault classification

Table 3 reports test macro-F1 under full fine-tuning with calibrated learning rates.

Pretraining transfers to fault classification. On CWRU the centralized encoder more than doubles the from-scratch reference (0.350 to 0.812); on HSG18 it moves from 0.569 to 0.951. Linear probing alone—an encoder frozen entirely, with a single trainable layer—already reaches 0.705 and 0.906. That linear probing captures most of the gain indicates that the improvement comes from the pretrained representation rather than from the additional optimization afforded by fine-tuning. The effect reproduces across two physically unrelated domains, rolling-element bearings and disk-drive telemetry, which suggests that the encoder captures structure common to industrial signals rather than a domain-specific shortcut.

Federated pretraining does not close the gap. The budget-matched comparison is Central-25k against the better federated checkpoint. The centralized encoder leads by 16.2 points on CWRU and 15.5 on HSG18, and it leads in all four pairwise comparisons (two datasets × two federated variants). The direction is therefore consistent across all four comparisons. Statistical reliability is another matter, and we report it plainly in Table 4. Two of the four comparisons are nominally significant and none survives correction for four tests. More importantly, the quantity being tested is not the one the claim needs: our three seeds vary the *downstream adaptation*, while each regime is represented by a single pretraining run, so the test measures adaptation noise rather than variability of the regime itself (Section 8). We therefore state the classification result as a large and consistent effect, not as a statistically established one.

Table 5: CMAPSS RUL, RMSE in cycles (lower is better), mean ± sample standard deviation over three seeds. The from-scratch reference is single-seed and uses a linear head, so it is *not* a matched comparator; it is shown because it bounds the practical value of pretraining on this task.

Checkpoint	Adaptation	Validation	Test
From scratch (1 seed)	linear head, full training	39.08	**43.67**
Central-100k (1 seed)	linear probing	43.21	53.86
Central-100k (1 seed)	full fine-tuning	39.54	44.33
Central-100k	MLP head, frozen	38.84 ± 0.07	44.16 ± 0.21
Central-25k	MLP head, frozen	39.07 ± 0.06	44.44 ± 0.58
FedAvg	MLP head, frozen	40.39 ± 0.06	42.72 ± 0.15
FedAvgM β=0.7	MLP head, frozen	38.50 ± 0.06	44.04 ± 1.02
Central-100k	full fine-tuning	39.53 ± 0.28	44.51 ± 1.13
Central-25k	full fine-tuning	39.85 ± 0.27	43.83 ± 0.71
FedAvg	full fine-tuning	39.39 ± 0.04	44.79 ± 0.20
FedAvgM β=0.7	full fine-tuning	39.63 ± 0.06	45.54 ± 1.30

The CWRU ordering is not robust to the learning-rate rule. Our protocol selects the backbone learning rate on validation, and on CWRU that rule decides between two values whose test consequences differ enormously. Central-25k attains validation 0.7794 at $3{\cdot}10^{-6}$ and 0.7716 at $1{\cdot}10^{-5}$—a difference of 0.008, well inside seed noise—while the corresponding test scores are 0.7535 and 0.5551. FedAvg, at those same two rates, tests at 0.6749 and 0.7152. Under the validation rule we obtain the gap reported above; under a single global $1{\cdot}10^{-5}$—the value validation selects for every checkpoint on HSG18—the CWRU ordering would reverse, with FedAvg ahead by roughly 16 points. On HSG18 no such fragility appears: the centralized encoder leads at all three grid values. The classification claim therefore rests on HSG18, where it is stable across the grid, and must be read as protocol-dependent on CWRU.

That instability is itself informative. A standard deviation of 0.119 on CWRU test macro-F1, against 0.022 for the centralized checkpoint at the same budget, indicates that averaging models trained on strongly non-IID clients produces representations whose downstream behaviour depends heavily on the adaptation seed.

Budget matters, but does not explain the gap. Central-100k outperforms Central-25k on both datasets (0.812 versus 0.734; 0.951 versus 0.911), so additional pretraining continues to help. Crucially, the federated deficit persists against the budget-matched Central-25k, so it cannot be attributed to compute alone.

6.2 Remaining useful life

Results on CMAPSS are markedly different. Table 5 reports them in full.

No reliable transfer. Under the protocol's own rule—select on validation—the best checkpoint is FedAvgM with 38.50, and its test RMSE is 44.04, statistically indistinguishable from Central-100k's 44.16 (t=0.20, df≈2.2). Meanwhile the from-scratch reference attains 43.67 on test, better than three of the four pretrained checkpoints under the MLP head. Validation and test therefore order the models differently, and the ordering that the protocol commits to does not produce a model that beats training from scratch.

Why the comparison is even weaker than it looks. The from-scratch run uses a linear output head and trains the full network for a single seed; the multi-seed rows use a two-layer MLP head over a frozen encoder. Head capacity is therefore confounded with the pretraining

Table 6: CMAPSS, single-seed basic modes: the two metrics disagree. The asymmetric score is the benchmark's official criterion; lower is better for both columns.

Mode	Test RMSE	Asymmetric score	R^2
From scratch	**43.67**	754,754	−0.40
Linear probing	53.86	1,002,485	−1.13
Full fine-tuning	44.33	**645,389**	−0.44

variable, and no matched from-scratch baseline with the MLP head exists. Table 5 therefore does not show that federated pretraining wins on RUL; it shows that this experiment establishes no benefit of pretraining on CMAPSS.

Under the benchmark's own metric the sign reverses. RMSE weighs early and late predictions equally. CMAPSS ships an asymmetric score that penalizes late predictions—those that overestimate remaining life, and are therefore the operationally dangerous ones—far more heavily than early ones. Under that score the single-seed comparison inverts: full fine-tuning of the centralized encoder scores 645,389 against the from-scratch baseline's 754,754, an improvement of 14.5 %, while its RMSE is 0.65 cycles worse (Table 6). We had computed this metric and initially omitted it; reporting only RMSE would have concealed the fact that the conclusion on this task depends on which error model one adopts. The honest summary is that pretraining does not improve symmetric error on CMAPSS and does appear to reduce the cost of late predictions, on one seed.

The test split is structurally different. The coefficient of determination on test is negative for every mode ($R^2 = -0.40$ from scratch, -0.44 full fine-tuning, -1.13 linear probing), meaning every model performs worse than predicting the test mean. The cause is a distribution shift intrinsic to the benchmark: CMAPSS test trajectories are truncated before failure, so the test target distribution differs from the training one. Absolute prognostic quality is consequently not interpretable here, and only relative comparisons between models evaluated under the same protocol carry information. Relative comparisons under a fixed protocol remain informative, so the numbers are reported with that restriction in mind.

Adaptation mode matters more than provenance. Linear probing is by far the worst mode on RUL (53.86 test RMSE against 43.67 from scratch). A frozen embedding projected by a single linear layer is not sufficient for continuous degradation regression, whereas the same frozen embedding with a two-layer head reaches the 42.7–44.4 band. The practical lesson is that conclusions about "whether a foundation model transfers" are contingent on the adaptation head, and that reporting a single mode invites the wrong conclusion.

6.3 The intrinsic loss does not predict downstream utility

The evidence here is weaker than the two preceding subsections and we present it as such.

A large intrinsic gap that does not map onto downstream utility. At matched budget the centralized encoder more than halves the federated reconstruction loss (0.3467 versus 0.7844, Table 2). If the intrinsic loss were a good proxy for representation quality, a factor of 2.3 should be visible downstream. It is not: on CMAPSS the two regimes are indistinguishable (Table 5), and on CWRU the ordering depends on the adaptation learning rate (Section 6.4). A monotone relation between the two metrics is therefore not supported by our data.

Table 7: Server-momentum sweep for FedAvgM (single seed, selected on validation). The intrinsic reconstruction loss is nearly constant while downstream utility varies by an order of magnitude more.

	β=0.0	β=0.3	β=0.5	β=0.7	β=0.9
SSL validation loss (↓)	0.784	0.778	0.765	**0.764**	0.768
CWRU, linear probing (↑)	0.163	0.160	0.217	**0.371**	0.273
HSG18, full fine-tuning (↑)	0.713	0.464	**0.750**	0.724	0.462

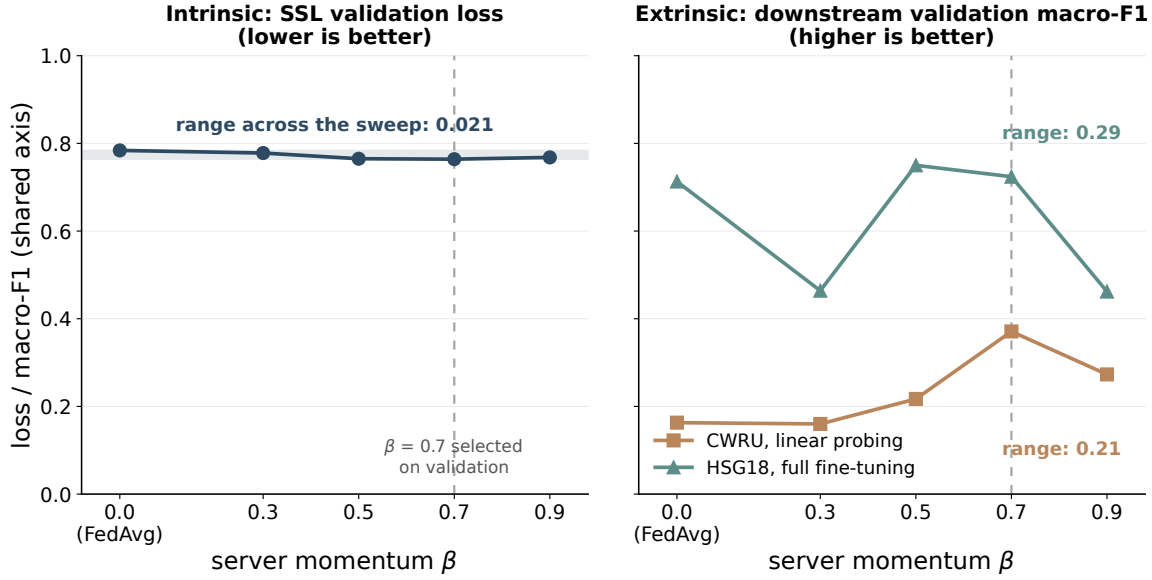

Figure 3: The intrinsic metric is flat where the extrinsic one is not. Server-momentum sweep for FedAvgM (data of Table 7; single seed, screening). Both panels share the vertical axis. The SSL validation loss moves within 0.021 across the sweep, while downstream validation macro-F1 moves by 0.21 (CWRU, linear probing) and 0.29 (HSG18, full fine-tuning), including the collapse at β=0.9.

A counter-example to our own claim: FedProx. It would be convenient to add FedProx here, and the thesis version of this work did. In a matched 5,000-step screening run FedProx reduced the aggregated reconstruction loss by 11.57 % against FedAvg's 7.55 % (0.7423 versus 0.7670). Its downstream behaviour in that same screening, however, was *also* mostly better than FedAvg's—it won on HSG18 linear probing by a wide margin and lost clearly only on CWRU fine-tuning. FedProx was excluded from the final bank because its advantage did not persist when the budget was scaled to 25,000 steps, where FedAvgM dominated it, not because its pretraining loss misled us. At the screening budget, FedProx is evidence *for* a correspondence between intrinsic and extrinsic quality, not against it. We report this because it is the clearest case in our study where a convenient narrative did not survive checking the underlying runs.

Server-momentum sweep. Table 7 and Figure 3 sweep the FedAvgM server momentum. The SSL validation loss varies within a band of 0.021 across the whole sweep—essentially flat—while downstream validation macro-F1 varies by 0.21 on CWRU and by 0.29 on HSG18, including a collapse at β=0.9. The two quantities are not merely imperfectly correlated; over this sweep the intrinsic metric carries almost no information about the extrinsic one.

The consequence for practice is concrete: a time-series foundation model cannot be validated on its pretraining curves. Checkpoint selection requires evaluation on target tasks, with selection performed on a validation split.

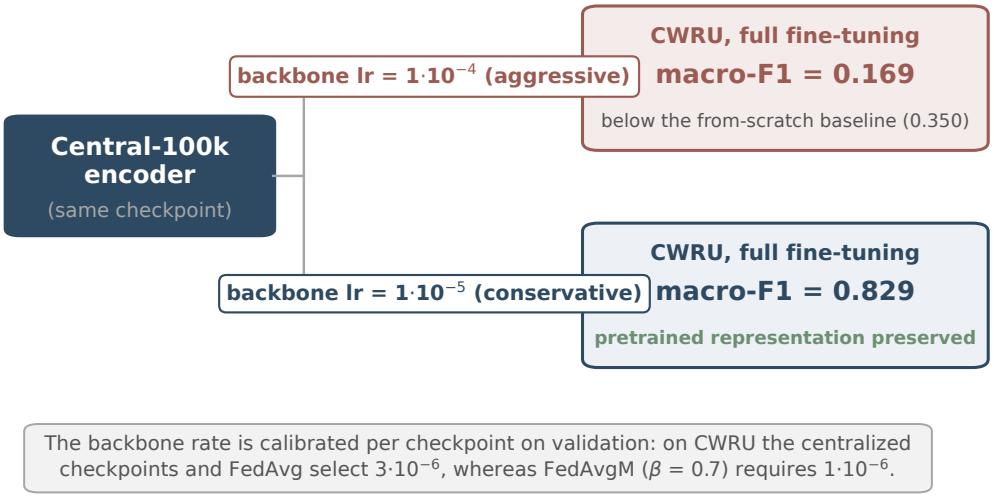

Figure 4: The same checkpoint, two outcomes. Adapting Central-100k to CWRU with a backbone learning rate of $1{\cdot}10^{-4}$ overwrites the pretrained representation and lands below the from-scratch baseline; $1{\cdot}10^{-5}$ preserves it. Both runs use a single seed; the rates selected on validation for each checkpoint are given in the text.

6.4 The backbone learning rate governs whether transfer survives

Adapting the *same* centralized checkpoint to CWRU with a backbone step size of $1{\cdot}10^{-4}$ yields 0.169 test macro-F1—worse than training from scratch—whereas $1{\cdot}10^{-5}$ yields 0.829 (Figure 4). The aggressive rate overwrites the pretrained representation, a catastrophic-forgetting effect in which fine-tuning is actively harmful.

The optimum is not a global constant. On CWRU, Central-100k, Central-25k and FedAvg select $3{\cdot}10^{-6}$ on validation, whereas FedAvgM requires $1{\cdot}10^{-6}$; the validation difference between FedAvgM's optimum and the canonical $1{\cdot}10^{-5}$ is nearly 0.10 macro-F1. Had we fixed a single rate across the bank, FedAvgM would have appeared substantially worse than it is, and we would have reported a conclusion about federated representations that would in fact have been a conclusion about an uncalibrated hyperparameter. On HSG18 and on CMAPSS full fine-tuning, by contrast, $1{\cdot}10^{-5}$ is optimal for all four checkpoints.

This is an evaluation-design result rather than a modelling one: comparisons between pretrained checkpoints are only meaningful when the adaptation rate is calibrated per checkpoint on validation.

7 Discussion

Why classification and RUL diverge. Fault classification on these benchmarks is largely separable by local morphology—impulses, frequency signatures, transient shapes—which is precisely the structure masked patch prediction must model in order to reconstruct hidden patches. RUL is different: it demands a monotone mapping from window-level evidence to a global degradation coordinate, information that is distributed across a trajectory rather than contained in a window. A window-level encoder pretrained on reconstruction has little reason to encode it. We offer this as an interpretation consistent with the evidence, not as a demonstrated mechanism; testing it would require architectures with explicit trajectory-level aggregation.

What the federated deficit does and does not show. At matched budget the federated encoder retains a substantial part of the representation—clearly above from scratch in classification—but does not reach the centralized encoder, and it is markedly less stable across seeds. Both effects are consistent with client drift under strong non-IID conditions: each client optimizes a partial and physically distinct distribution, the server averages once per round, and the caps deliberately move weight away from the dominant domains. Our experiment does not separate these causes; it measures their joint cost, which is the quantity a practitioner deciding whether to centralize actually faces.

Scale. Our encoder has 0.8M parameters. It is possible that the qualitative conclusions change with two or three additional orders of magnitude—in particular, larger encoders may transfer to RUL where ours does not. What our study establishes is that, at this scale and under this protocol, the pattern is task-dependent transfer, an unclosed federated gap and metric decoupling. We consider the decoupling finding the most likely to persist across scales, since it concerns the relationship between two metrics rather than the capacity of the model.

8 Threats to validity and limitations

The first four items below are not peripheral caveats: each of them limits a headline claim, and two of them we consider serious enough that the corresponding result should be treated as preliminary.

- **The regimes differ in optimizer schedule, not only in regime.** Central-25k is a mid-cosine snapshot of the 100,000-step run, whereas the federated arm uses a constant learning rate without a scheduler. The mismatch favours the centralized arm, and we cannot separate schedule from regime. The classification comparison should be read as an upper bound on the cost of federating, and a matched Central-25k trained with the federated arm's constant rate is the single experiment most needed to settle it. Its configuration is committed in the repository (`training/configs/ssl_central_25k_constlr_seed*.yaml`) and validated by dry run; it was not executed for this version.
- **One pretraining run per regime.** The three seeds vary downstream adaptation, not pretraining. The unit of randomization implied by a claim about regimes is the pretraining run, and there $n{=}1$ for each arm. Consequently the p-values in Table 4 do not license inference about the regimes; multi-seed federated pretraining checkpoints exist in our artefacts and repeating the comparison over them is straightforward future work; the corresponding evaluation configuration (`downstream_paper_v2_regime.yaml`) is committed alongside the pretraining configuration above.
- **Uncalibrated baselines.** Every from-scratch baseline was run with a single backbone learning rate of $1{\cdot}10^{-4}$ and no grid, while each pretrained checkpoint received a validation-selected rate from a three-point grid. The baselines are therefore lower bounds rather than matched comparators, and the reported classification improvements are upper bounds on the true effect. A learning-rate grid for the baselines is likewise committed (`downstream_paper_v2_fromscratch.yaml`) and not executed.
- **CMAPSS departs from the corpus contract.** The RUL dataset is rebuilt by a dedicated pipeline with causal windows and stride 5, not the stride-256 contract used elsewhere, and its median trajectory (184 cycles) is shorter than the 512-sample window, so every window is padded. The RUL null may reflect an encoder operating far outside its pretraining regime rather than an absence of transferable structure; a $W \in \{64, 128\}$ ablation is required before attributing it to pretraining.
- **Statistical power.** Three seeds, with federated standard deviations up to 0.119 (Table 4); the from-scratch baselines are single-seed (Section 5).

- **Class coverage in CWRU.** One of the four fault classes is absent from the test split, which contains 2.4 % of the dataset; reported macro-F1 averages three classes. Our anti-leakage checks verify unit disjointness but not class coverage, an omission we recommend others avoid.
- **Physical window length.** At CWRU's sampling rates, a 512-sample window spans between roughly 1.3 and 0.3 revolutions of the shaft. Bearing fault signatures that rely on amplitude modulation at shaft frequency are therefore only marginally observable, which may bound achievable performance independently of pretraining.
- **Evaluation coverage.** Eleven target datasets were reserved; the multi-seed evaluation covers three. The remaining eight are documented with their status in Appendix A, including one negative case (PBCP16) and one dataset excluded from the body because it was run with an uncapped batch configuration (PHM18).
- **No external baselines.** We compare regimes within a fixed architecture, which isolates the variable of interest but does not situate the results against specialized PHM methods or against existing time-series foundation models. A linear-probing comparison against a published TSFM on the same held-out targets is the single most valuable missing experiment.
- **Federated simulation.** Ten clients on one machine, with no latency, no partial participation, no secure aggregation and no differential privacy. The heterogeneity is real; the systems dimension is not modelled.
- **Fixed window in samples.** W=512 is a computational contract, not a physical duration: it spans milliseconds in a high-rate vibration channel and hours in a per-cycle counter. Most PHMD datasets do not expose a recoverable sampling rate, which makes a physically matched window infeasible without additional metadata.
- **Instance normalization.** Per-window normalization discards absolute amplitude, which can be diagnostic. The statistics are retained on disk so this can be revisited.
- **RUL confound.** As discussed, the from-scratch RUL baseline uses a different head from the multi-seed rows, so head capacity is confounded with the pretraining variable on that task.

9 Conclusion

We measured, under a leakage-controlled protocol on 47 harmonized PHM datasets, what self-supervised multi-dataset pretraining buys and what federating it costs. Pretraining transfers strongly to fault classification across two unrelated physical domains, with linear probing alone capturing most of the gain. It does not transfer reliably to remaining-useful-life estimation on CMAPSS, where a from-scratch baseline remains competitive and validation and test disagree. Federated pretraining at matched budget retains part of the representation but does not close the classification gap, and produces markedly less stable downstream behaviour. Finally, the self-supervised loss is not a valid criterion for selecting among pretrained checkpoints: a regime that more than halves the reconstruction loss is indistinguishable from the other on RUL, and a hyperparameter sweep can leave the intrinsic metric flat while downstream utility varies by an order of magnitude more.

The practical recommendations that follow are narrow but firm. Withhold evaluation datasets in full and split by physical unit before windowing. Calibrate the adaptation learning rate per checkpoint on validation, or risk attributing to the representation what belongs to the optimizer. Report linear probing alongside fine-tuning, since their divergence is informative. And never select a pretrained encoder on its pretraining curve.

Reproducibility

Code, configurations, corpus construction and audit artefacts are available at `https://github.com/sanzzjose/fm-predictive-maintenance`. Every reported number is traceable to a versioned run record containing the configuration hash, the seed, the commit and the resulting

metrics. Raw datasets are redistributed by PHMD [19] and are not mirrored in the repository.

This article extends the first author's master's thesis at Universidad de Sevilla (2026). Relative to the thesis, the present version adds the significance analysis of Table 4, the re-examination of FedProx in Section 6.3, and the first three items of Section 8. The three experiments that those items identify as most needed are specified as committed, dry-run-validated configurations, so that they can be executed without further design decisions.

A Coverage of the held-out target datasets

Eleven datasets were withheld from pretraining. Table 8 documents the status of each, so that the three used in the body can be read against the full reserve rather than as a selection.

Table 8: Status of the eleven held-out target datasets.

Dataset	**Status**	**Reason**
CWRU	final, n=3	classification, body of the paper
HSG18	final, n=3	classification, body of the paper
CMAPSS_RUL	final, n=3	RUL, body of the paper
PBCP16	diagnostic, n=1	negative case, Appendix B
PHM18	excluded from body	run with uncapped batch; pending re-execution
PHMAP23	documented negative	extreme class imbalance, noise floor
CALCE_CS2	blocked	target semantics (RUL/SOH) unresolved
IEEE14	not evaluated	RUL target; no dedicated pipeline (only CMAPSS)
CBM14	not evaluated	secondary tier
CNCMILL18	not evaluated	secondary tier
PHME20	not evaluated	secondary tier

B A negative case: PBCP16

PBCP16 is small (1,975 windows, 5 balanced classes, one channel). A from-scratch model reaches 0.907 macro-F1; linear probing reaches 0.829 and full fine-tuning 0.889. Pretraining does not help because there is no headroom: the task is already solved from scratch. We report this because it bounds the claim. Self-supervised pretraining is valuable where the target task is hard relative to its available supervision, not universally, and a study that reported only CWRU and HSG18 would overstate its scope.

C Per-class breakdown on CWRU

Table 9 gives the per-class decomposition of the best CWRU configuration (centralized checkpoint, full fine-tuning, $1{\cdot}10^{-5}$, single seed, test macro-F1 0.8292). Two facts are invisible in the scalar. First, class 3 has no test support, so the macro average is over three classes. Second, class 0 is recalled almost perfectly but with precision 0.480: 477 windows of class 1 are predicted as class 0. Whether that matters depends on which class encodes the healthy state—an operational question that a macro-F1 of 0.83 does not answer.

D Hyperparameters

References

[1] Y. Huang et al. Prognostics and health management for predictive maintenance: a review. *Journal of Manufacturing Systems*, 2024.

Table 9: CWRU per-class performance, best configuration. Confusion matrix rows are true classes, columns predicted.

Class	Support	Precision	Recall	F1
0	477	0.480	0.998	0.648
1	2378	0.997	0.799	0.887
2	476	0.998	0.910	0.952
3	0	—	—	undefined
Macro-F1 over the three supported classes				0.829

Confusion matrix

	pred. 0	pred. 1	pred. 2	pred. 3
true 0	476	0	1	0
true 1	477	1901	0	0
true 2	38	5	433	0
true 3	0	0	0	0

Table 10: Principal configuration. Identical across regimes unless stated.

Parameter	Value
Window / stride / patch	512 / 256 / 16
Patches per window (N)	32
Mask ratio	0.30 of valid patches
d_{model} / layers / heads	128 / 4 / 4
Normalization	pre-norm; per-window, per-channel instance norm
Trainable parameters	801,808
Sampling caps (dataset / client / floor)	0.10 / 0.25 / 0.005
Central budget	100,000 steps (full), 25,000 (matched)
Federated budget	50 rounds $\times$ 50 local steps $\times$ 10 clients
Federated algorithms	FedAvg, FedAvgM (β=0.7), FedProx (μ=0.01)
Downstream head LR	$1{\cdot}10^{-3}$
Backbone LR grid	$\{1{\cdot}10^{-6}, 3{\cdot}10^{-6}, 1{\cdot}10^{-5}\}$
Seeds	42, 43, 44
Hardware	single NVIDIA A100

[2] O. Fink, Q. Wang, M. Svensen, P. Dersin, W.-J. Lee, and M. Ducoffe. Potential, challenges and future directions for deep learning in prognostics and health management applications. *Engineering Applications of Artificial Intelligence*, 92:103678, 2020.

[3] L. Zhang, J. Lin, B. Liu, Z. Zhang, X. Yan, and M. Wei. A review on deep learning applications in prognostics and health management. *IEEE Access*, 7:162415–162438, 2019.

[4] C. Ferreira and G. Gonçalves. Remaining useful life prediction and challenges: a literature review on the use of machine learning methods. *Journal of Manufacturing Systems*, 63:550–562, 2022.

[5] Y. Nie, N. H. Nguyen, P. Sinthong, and J. Kalagnanam. A time series is worth 64 words: long-term forecasting with transformers. In *International Conference on Learning Representations*, 2023.

[6] M. Goswami, K. Szafer, A. Choudhry, Y. Cai, S. Li, and A. Dubrawski. MOMENT: a family of open time-series foundation models. In *International Conference on Machine Learning*, 2024.

[7] A. Das, W. Kong, R. Sen, and Y. Zhou. A decoder-only foundation model for time-series forecasting. In *International Conference on Machine Learning*, 2024.

[8] M. Jin et al. Time-LLM: time series forecasting by reprogramming large language models. In *International Conference on Learning Representations*, 2024.

[9] Y. Liang et al. Foundation models for time series analysis: a tutorial and survey. In *ACM SIGKDD*, 2024.

[10] S. Kottapalli et al. A survey of time series foundation models. *arXiv preprint*, 2025.

[11] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas. Communication-efficient learning of deep networks from decentralized data. In *AISTATS*, 2017.

[12] T. Li, A. K. Sahu, M. Zaheer, M. Sanjabi, A. Talwalkar, and V. Smith. Federated optimization in heterogeneous networks. In *MLSys*, 2020.

[13] S. P. Karimireddy, S. Kale, M. Mohri, S. Reddi, S. Stich, and A. T. Suresh. SCAFFOLD: stochastic controlled averaging for federated learning. In *International Conference on Machine Learning*, 2020.

[14] T.-M. H. Hsu, H. Qi, and M. Brown. Measuring the effects of non-identical data distribution for federated visual classification. *arXiv preprint arXiv:1909.06335*, 2019.

[15] W. Zhuang, C. Chen, and L. Lyu. When foundation model meets federated learning. *arXiv preprint*, 2023.

[16] C. Ren et al. Advances and open challenges in federated foundation models. *arXiv preprint*, 2024.

[17] Z. Li et al. TSFM-Bench: a comprehensive and unified benchmark of foundation models for time series forecasting. *arXiv preprint*, 2024.

[18] J. Meyer et al. On information leakage in the evaluation of time series foundation models. *arXiv preprint*, 2025.

[19] D. Solís-Martín, J. Galán-Páez, and J. Borrego-Díaz. PHMD: a unified library for prognostics and health management datasets. *SoftwareX*, 2025.