TimeWorld: Towards Action-Conditioned Multimodal World Models for Time Series Forecasting

Hang Ni
The Hong Kong University of Science and Technology (Guangzhou)
Guangzhou, China
hni017@connect.hkust-gz.edu.cn

Weijia Zhang✉
The Hong Kong University of Science and Technology (Guangzhou)
Guangzhou, China
wzhang411@connect.hkust-gz.edu.cn

Hao Liu✉
The Hong Kong University of Science and Technology (Guangzhou)
Guangzhou, China
liuh@ust.hk

arXiv:2610.34797v1 [cs.LG] 26 Oct 2026

Abstract

Time series forecasting supports decision-making in web applications such as online retail and network management, where planned actions can reshape future dynamics. World models provide a natural paradigm for predicting such action-dependent changes beyond extrapolating historical patterns. However, existing time series world models (TSWMs) typically rely on fixed, domain-specific covariate schemas that limit flexible intervention specification, while multimodal forecasters treat text primarily as passive predictive context. To investigate these gaps, we first introduce **TSWM-Bench**, a cross-domain benchmark for *multimodal TSWMs* spanning seven real-world datasets, with textual interventions grounded in recorded action–outcome data. Our evaluation reveals that conventional multimodal forecasters underperform history-only time series foundation models (TSFMs) despite dataset-specific training, whereas simple text-prefix adaptations improve forecasting across all evaluated TSFM backbones. These findings suggest that pretrained TSFMs encode implicit world knowledge of temporal dynamics, offering a strong foundation for action-conditioned modeling beyond conventional multimodal fusion. Building on this insight, we further propose **TimeWorld**, a plug-in framework that adapts pretrained TSFMs through explicit action-response modeling. Specifically, the Prototype-Guided Latent Modulation (PLM) module grounds textual actions in temporal dynamics by composing representative response prototypes to modulate future representations based on the action and historical patterns. The Action Contrastive Learning (ACL) component strengthens action-specific response modeling by contrasting these updates under alternative interventions sharing the same history, distinguishing action types while preserving relative intensity relationships. Experiments across diverse domains and TSFM backbones demonstrate consistent improvements over strong baselines.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
Conference acronym 'XX, Woodstock, NY
© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-1-4503-XXXX-X/2018/06
https://doi.org/XXXXXXX.XXXXXXX

CCS Concepts

• **Mathematics of computing** → **Time series analysis**; • **Computing methodologies** → **Machine learning**; • **Applied computing** → *Forecasting*.

Keywords

Time Series World Models, Textual Interventions, Time Series Foundation Models, Multimodal Learning

ACM Reference Format:
Hang Ni, Weijia Zhang, and Hao Liu. 2018. TimeWorld: Towards Action-Conditioned Multimodal World Models for Time Series Forecasting. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 13 pages. https://doi.org/XXXXXXX.XXXXXXX

1 Introduction

Time series forecasting (TSF) underpins web-related applications such as retail planning [33], network management [18], and financial analytics [49]. Recent time series foundation models (TSFMs) have further improved forecasting by learning transferable priors over temporal dynamics through large-scale pretraining [8]. However, forecasting in these applications often requires anticipating how planned actions alter future trajectories, beyond extrapolating historical patterns [27]. For example, an online retailer often needs to estimate demand under alternative discount campaigns before allocating its promotional budget. Motivated by the world-model paradigm of predicting system evolution under different actions [5, 12], we study *action-conditioned time series world models* (TSWMs) that forecast future trajectories from historical observations and a specified action [46].

Existing TSWMs commonly encode interventions as numerical or categorical covariates [27, 32, 46]. Such predefined, domain-specific schemas struggle to capture complex intervention semantics (*e.g.*, buy-one-get-one offers), limiting cross-domain generalization [35]. Moreover, their reliance on stepwise future controls makes it difficult to express high-level temporal intervention patterns (*e.g.*, phasing out discounts) without prescribing exact control sequences. Instead, natural language provides a flexible interface to specify action semantics, intensity, and temporal scope without fixed schemas or stepwise controls [15]. However, recent studies of multimodal time series forecasting typically use text to describe temporal patterns [9, 20], provide domain metadata [16, 22], or supply contextual signals such as news and events [21, 38–40]. In these methods, text provides predictive context rather than specifying

interventions whose effects must be modeled. We therefore investigate *multimodal TSWMs* that condition future temporal dynamics on textual actions.

Early efforts in multimodal TSWMs rely on synthetic data for counterfactual ground truth [11] or limit evaluation to a single domain [44]. To address these gaps, we first introduce **TSWM-Bench**, a benchmark for multimodal TSWMs spanning seven real-world datasets across economic, physical, and healthcare domains, with interventions ranging from price adjustments to medical treatments. Each instance pairs a univariate history and its observed future with textual action and context descriptions derived from source records. Table 1 compares its coverage with existing benchmarks. We evaluate history-only TSFMs, representative multimodal forecasters, and covariate-aware TSWMs to examine their capabilities for action-conditioned prediction. Our results reveal that existing multimodal forecasting methods consistently underperform history-only TSFMs despite dataset-specific training. By contrast, simple text-prefix adaptations of TSFMs improve forecasting accuracy across all evaluated backbones and even outperform covariate-aware TSWMs. These findings motivate a new route to world modeling: *reusing pretrained world-dynamics priors while learning how textual actions reshape future states.*

However, effectively adapting pretrained TSFMs to multimodal TSWMs presents two critical challenges: *(1) Action Grounding.* Incorporating textual actions requires explicitly translating intervention semantics into state-dependent changes to the predictive dynamics encoded by pretrained TSFMs. Existing multimodal forecasters typically fuse textual and temporal representations directly, complicating cross-modal interpretation [30] while leaving fine-grained action effects on local temporal patterns implicit. Consequently, capturing meaningful intervention responses and their state-dependent evolution remains challenging. *(2) Action Discrimination.* Even with action grounding, forecasts may remain dominated by pretrained dynamics, obscuring action-specific responses. Forecasting loss alone provides limited supervision for action–response alignment, as it evaluates factual outcomes without comparing responses to alternative actions under the same history. This requires additional supervision for action discrimination to ensure that predictions faithfully reflect the specified intervention.

To address these challenges, we further propose **TimeWorld**, a plug-in framework for building action-conditioned time series world models from pretrained TSFMs. For action grounding, *Prototype-Guided Latent Modulation (PLM)* constructs a bank of representative temporal response prototypes and adaptively composes them according to the specified action, forecast interval, and temporal state. The resulting responses modulate future hidden states within a compact latent subspace, preserving the pretrained temporal components. For action discrimination, *Action Contrastive Learning (ACL)* uses large language models (LLMs) to construct alternative interventions under the same history and applies contrastive supervision to their latent updates. This strengthens action–response alignment while preserving relative intensity relationships. *Context-enhanced Temporal Modeling (CTM)* further incorporates textual system context through lightweight prefix conditioning. Extensive experiments demonstrate the effectiveness and generality of TimeWorld across diverse domains and TSFM backbones.

Table 1: Comparison of benchmarks.

Study	Action-Conditioned	Action Format	Real-world Action	Multi-domain
Multimodal TSF				
Time-MMD [21]	✗	–	–	✔
CiK [39]	✗	–	–	✔
Structured TSWM				
MobiWM [32]	✔	Covariates	✗	✗
CEDAR [27]	✔	Covariates	✔	✗
Zhang et al. [46]	✔	Covariates	✔	✔
Multimodal TSWM				
TADiff [11]	✔	Text	✗	✔
ReasonCast [44]	✔	Text	✔	✗
TSWM-Bench	✔	Text	✔	✔

Our main contributions are threefold: (1) We introduce TSWM-Bench, the first cross-domain benchmark for multimodal TSWMs with textual interventions grounded in real-world action–outcome records. (2) We propose TimeWorld, a general framework that augments pretrained TSFM dynamics through PLM-based action grounding and ACL-based discrimination. (3) Extensive experiments demonstrate consistent improvements over strong baselines across diverse domains and TSFM backbones.

2 Related Work

In this section, we review time series foundation models, multimodal forecasting, and time series world models.

Time Series Foundation Models. Large-scale pretraining enables TSFMs to acquire transferable world knowledge of temporal evolution across datasets [8, 29]. Existing approaches primarily follow two forecasting paradigms. *Autoregressive models* generate future patches sequentially, allowing flexible prediction horizons but potentially accumulating errors over long rollouts [19, 23, 34]. *Masked-future models* instead predict future patches jointly from historical observations and masked placeholders, enabling parallel inference without recursive error propagation [1, 14, 17]. Despite this implicit knowledge of system dynamics, these models lack mechanisms to ground textual actions in state-dependent changes to future trajectories.

Multimodal Time Series Forecasting. Multimodal forecasting enriches numerical observations with textual context, exploiting complementary semantic information to improve predictive accuracy. Text serves varied roles: early efforts use *task descriptions* to formulate forecasting as language generation [42], summarize *temporal patterns* to expose numerical structure [9, 20], or provide *domain metadata* to contextualize heterogeneous series [16, 22, 50]. Beyond describing the task or series, *predictive context* supplies supplementary information correlated with future observations but unavailable in numerical histories, including retrospective reports and news [21, 38], event announcements [48], and prospective conditions or constraints [11, 39]. Multimodal foundation models further learn transferable numerical–textual representations through shared pretraining, extending context-aware forecasting across datasets [6, 7, 36, 40]. However, predominantly using text

as passive predictive context leaves the causal effects of actively planned interventions insufficiently modeled, motivating the shift toward action-conditioned simulation.

Time Series World Models. World models for control [12] and interactive video generation [5] motivate modeling temporal evolution for prediction and planning. Inspired by video joint-embedding prediction [2], *latent predictive modeling* forecasts future representations before decoding observations, but does not condition state transitions on actions [24, 43]. Although multimodal extensions incorporate contextual information [7, 52], they still do not explicitly model how controllable actions alter temporal state transitions. In contrast, early *action-conditioned* approaches use structured treatments for counterfactual prediction or clinical policy learning [26, 37, 41]. Recent TSWMs encode actions as numerical or categorical covariates: base-station power and antenna settings [32], advertising budgets [27], or engineering controls and treatment doses [46]. These formulations restrict actions to domain-specific schemas and may omit intervention details. Language offers a flexible interface without prescribing a unique control sequence, motivating *text-conditioned multimodal TSWMs*. For example, TADiff [11] forecasts under alternative textual conditions through text-attributive diffusion, but its real-data evaluation concerns descriptive conditions and semantic consistency rather than recorded action effects, with ground-truth counterfactual trajectories confined to synthetic data. ReasonCast [44] translates textual actions into semantic interventions through agentic reasoning, but focuses on e-commerce demand. A comprehensive multimodal TSWM benchmark grounded in real action–outcome records across domains remains missing. We fill this gap with TSWM-Bench and introduce TimeWorld, which extends pretrained predictive dynamics to textual interventions through prototype-guided latent updates and contrastive discrimination.

3 TSWM-Bench

This section introduces TSWM-Bench, a benchmark for evaluating action-conditioned multimodal time series forecasting.

3.1 Task Formulation

Given L historical observations of a univariate time series $\mathbf{x} = (x_{t-L+1}, \ldots, x_t) \in \mathbb{R}^L$, textual context c describing the system, and a textual action a specifying an intervention and its temporal scope, the task is to learn a world model $\mathcal{F}$ that predicts the next H values $\mathbf{y} = (x_{t+1}, \ldots, x_{t+H}) \in \mathbb{R}^H$ under the specified action, formulated as $\widehat{\mathbf{y}} = \mathcal{F}(\mathbf{x}, c, a)$.

3.2 Datasets

TSWM-Bench comprises seven real-world datasets spanning three significant domains: economics, physical systems, and healthcare. In *economics*, M5 [25] pairs daily item-level retail sales with recorded changes in weekly average selling prices. MOF [28] links Japan's foreign-exchange interventions, specifying traded currencies, transaction direction, and amount, to subsequent yen-per-dollar exchange rates. FOMC [4] pairs U.S. monetary policy statements with subsequent two-year Treasury yields. In *physical systems*, Causal Chambers [10] records downstream air pressure in a wind tunnel under changes in inlet fan load. DaISy [3] pairs changes in liquid flow through a heat exchanger with its outlet temperature measurements. In *healthcare*, Shanghai Diabetes [51] links continuous glucose measurements to clinical records of meals, non-insulin medication, and insulin administration. For HUPA-UCM [13], we use raw FreeStyle glucose records at approximately 15-minute intervals, paired with recorded carbohydrate intake and insulin doses in people with type 1 diabetes.

We construct each instance by aligning recorded interventions with preceding time series histories and subsequent observations, consolidating simultaneous operations into a single action where applicable. We then generate textual action descriptions from source records: FOMC retains its original policy statements, while the other datasets use standardized templates that preserve intervention details and timing. System-level context is constructed separately using only information available at the forecast origin. These components are organized into a unified sample format across datasets, providing a consistent interface for textual action conditioning. Finally, dataset-specific preprocessing and validation ensure temporal consistency, filter incomplete action records and invalid observations, and produce well-aligned action–outcome samples.

Table 2: Averaged forecasting performance across seven datasets. Rank is the mean rank over the three metrics. The best and second-best results in each column are shown in red and blue, respectively.

Model	nMAE ↓	nRMSE ↓	MASE ↓	Rank ↓
TimesFM 3.0	1.0540	1.2087	3.5023	8.33
Chronos-2	1.0237	1.1695	3.4948	5.00
Toto 2.0	1.0257	1.1730	3.4923	5.33
Moirai 2.0	1.0329	1.1759	3.4996	7.00
Time-LLM	1.4213	1.5774	4.6371	12.67
MM-TSFlib	1.1301	1.2530	3.8602	11.00
TimeCMA	1.4410	1.6115	5.3578	14.33
TADiff	1.4515	1.6465	5.0358	14.67
ChatTime	1.3114	1.4571	4.6900	12.33
Aurora	1.1108	1.2301	3.7131	10.00
CovTSWM	1.0583	1.1930	3.6517	8.67
TSFM++ / TimesFM 3.0	**0.9983**	**1.1116**	**3.3634**	**1.00**
TSFM++ / Chronos-2	1.0140	**1.1135**	**3.4107**	**2.33**
TSFM++ / Toto 2.0	1.0310	1.1603	3.4827	4.67
TSFM++ / Moirai 2.0	**1.0121**	1.1298	3.4428	2.67

3.3 Evaluation Protocol

Baselines. We evaluate four classes of baselines, covering zero-shot forecasting and dataset-specific training: (1) *Time Series Foundation Models (TSFMs)*: TimesFM 3.0 [14], Chronos-2 [1], Toto 2.0 [17], and Moirai 2.0 [19] provide zero-shot forecasts from numerical histories, establishing the performance of pretrained temporal models without action information. (2) *Multimodal Time Series Models*: Time-LLM [16], MM-TSFlib from Time-MMD [21], TimeCMA [20], and TADiff [11] are trained on each dataset to evaluate task-specific numerical–textual integration for action-conditioned forecasting. (3) *Multimodal TSFMs*: ChatTime [36] and Aurora [40] are evaluated zero-shot to assess whether pretrained multimodal representations

transfer to textual action conditioning. (4) *Covariate-aware TSWM*: CovTSWM (our name for the method in [46]) is trained on each dataset with an adapted covariate schema that encodes recorded actions as structured controls. We use the best-performing module combinations for the modular frameworks: MM-TSFlib and CovTSWM, with implementation details in Appendix A. All text-capable baselines receive equivalent context and action information through their respective input interfaces.

Metrics. We report normalized mean absolute error (nMAE) and normalized root mean squared error (nRMSE), scaled by the standard deviation of each input history, and mean absolute scaled error (MASE), scaled by its mean absolute lag-one difference, with lower values indicating better forecasts. All scales use observed historical values only, and zero-scale instances are excluded rather than assigned zero error (Appendix B). Additional implementation details appear in Section 5.1.

3.4 Benchmark Findings

Table 2 shows that both dataset-trained multimodal models and zero-shot multimodal foundation models have higher aggregate errors than history-only TSFMs, suggesting that access to textual actions alone does not ensure effective forecasting. CovTSWM achieves lower errors than these multimodal baselines across all three metrics, supporting the effectiveness of explicit action modeling. Yet it remains less accurate than the strongest TSFMs, underscoring the predictive value of pretrained world-dynamics knowledge. These findings raise a central question for world modeling: *can pretrained TSFMs build on their implicit world dynamics to explicitly model how textual actions shape future trajectories*?

We thereby introduce TSFM++, a dataset-specific adaptation applicable to each evaluated TSFM that prepends projected context and action embeddings to its time series tokens (illustrated in Section 4.1). We denote backbone-specific variants by TSFM++ $_{/\,\text{backbone}}$. As shown in Table 2, action-prefix adaptation generally reduces aggregate errors, although the gains vary across backbones and Toto 2.0 does not improve in nMAE. These results suggest that textual actions can refine pretrained predictive dynamics, motivating TimeWorld's explicit action-dependent latent updates, as detailed in Section 4. Additional results in Appendix F highlight the advantage of textual action conditioning, with TSFM++ outperforming TimesFM 3.0 using native structured covariates under both zero-shot and fine-tuning settings.

4 TimeWorld

Figure 1 illustrates the overall architecture of TimeWorld. First, *Context-enhanced Temporal Modeling (CTM)* incorporates textual context into a pretrained TSFM to obtain context-aware temporal representations. Next, *Prototype-Guided Latent Modulation (PLM)* learns a bank of temporal response prototypes and adaptively uses them to update the TSFM hidden states according to the specified action. The original TSFM prediction head then decodes these action-adjusted states into future values. To complement the forecasting loss, *Action Contrastive Learning (ACL)* further guides the latent updates to distinguish action types and preserve the ordering of action intensities.

4.1 Context-enhanced Temporal Modeling

To leverage implicit world knowledge of temporal dynamics, we build TimeWorld upon a pretrained TSFM consisting of three main components: a patch encoder Enc, a temporal backbone Backbone, and a prediction head Dec.

Given a historical time series $\mathbf{x}$ of length L, the patch encoder divides it into patches of length L_{in} and maps them into embeddings $\mathbf{E} = \text{Enc}(\mathbf{x}) \in \mathbb{R}^{N\times D}$, where $N = \lceil L/L_{\text{in}} \rceil$ is the number of input patches and D is the hidden dimension. To capture domain-specific dynamics that may not be identifiable from numerical observations alone, we additionally introduce textual context c. Specifically, a pretrained text embedding model TextEnc^c encodes c into a vector, which is projected into the TSFM hidden space by a lightweight adapter Adapt^c comprising two linear layers and a GELU activation: $\mathbf{c} = \text{Adapt}^c(\text{TextEnc}^c(c)) \in \mathbb{R}^D$. We prepend this context embedding as an additional token to the numerical patch sequence, obtaining $\mathbf{E}^c = [\mathbf{c}; \mathbf{E}] \in \mathbb{R}^{(N+1)\times D}$, which is processed by the temporal backbone using its native forecasting mechanism.

For a forecasting horizon of H time steps, the TSFM predicts $N_{\text{f}} = \lceil H/L_{\text{out}} \rceil$ future patches of length L_{out}, each associated with a hidden state $\mathbf{g}_k \in \mathbb{R}^D$, where $k = 1, \ldots, N_{\text{f}}$ indexes future patches. Different TSFM architectures obtain these states through distinct forecasting paradigms. Autoregressive models generate future patches sequentially, using previously predicted patches to compute the next hidden state: $\mathbf{g}_k = \text{Backbone}([\mathbf{E}^c; \text{Enc}(\widehat{\mathbf{y}}_{<k})])_{\text{last}}$, where $\widehat{\mathbf{y}}_{<k}$ denotes the preceding predicted patches and the subscript last selects the final hidden state. In contrast, masked-future models obtain all future hidden states jointly by appending N_{f} masked tokens $\mathbf{E}^{\text{mask}} \in \mathbb{R}^{N_{\text{f}}\times D}$, which is formulated as $[\mathbf{g}_1; \ldots; \mathbf{g}_{N_{\text{f}}}] = \text{Backbone}([\mathbf{E}^c; \mathbf{E}^{\text{mask}}])_{\text{future}}$, where future selects the hidden states at the masked future positions. In either paradigm, the TSFM yields context-aware but action-unaware future states $\mathbf{g}_k$, which our subsequent *PLM* module updates into action-conditioned states $\widetilde{\mathbf{g}}_k$. They are finally decoded by the prediction head as $\widehat{\mathbf{y}}_k = \text{Dec}(\widetilde{\mathbf{g}}_k)$.

For comparison, TSFM++, the action-prefix baseline in Section 3.4, directly prepends adapted embeddings of both context and action to the numerical patch sequence, without *PLM* or *ACL* modules.

4.2 Prototype-Guided Latent Modulation

To ground textual interventions in pretrained world dynamics, *PLM* composes temporal response prototypes as building blocks to modulate the TSFM hidden states. It consists of two components: *(1) Temporal Prototype Construction*, which builds a bank of representative temporal patterns, and *(2) Patch-wise Modulation*, which adaptively combines these prototypes to produce action-dependent latent updates for future forecasts.

4.2.1 Temporal Prototype Construction. To capture diverse intervention effects, we construct N_{p} reference curves representing typical temporal responses (*e.g.*, steps, ramps, pulses) with variations in onset, duration, and rate. Each curve is normalized and resampled to the input patch length L_{in}, yielding $\mathbf{x}_m^{\text{ref}} \in \mathbb{R}^{L_{\text{in}}}$. We encode each curve as a single patch using TSFM's frozen patch encoder and temporal backbone, followed by a trainable two-layer projection:

$$\mathbf{p}_m = \text{Proj}\Big(\text{Backbone}(\text{Enc}(\mathbf{x}_m^{\text{ref}}))\Big) \in \mathbb{R}^R, \qquad m = 1, \ldots, N_{\text{p}}, \quad (1)$$

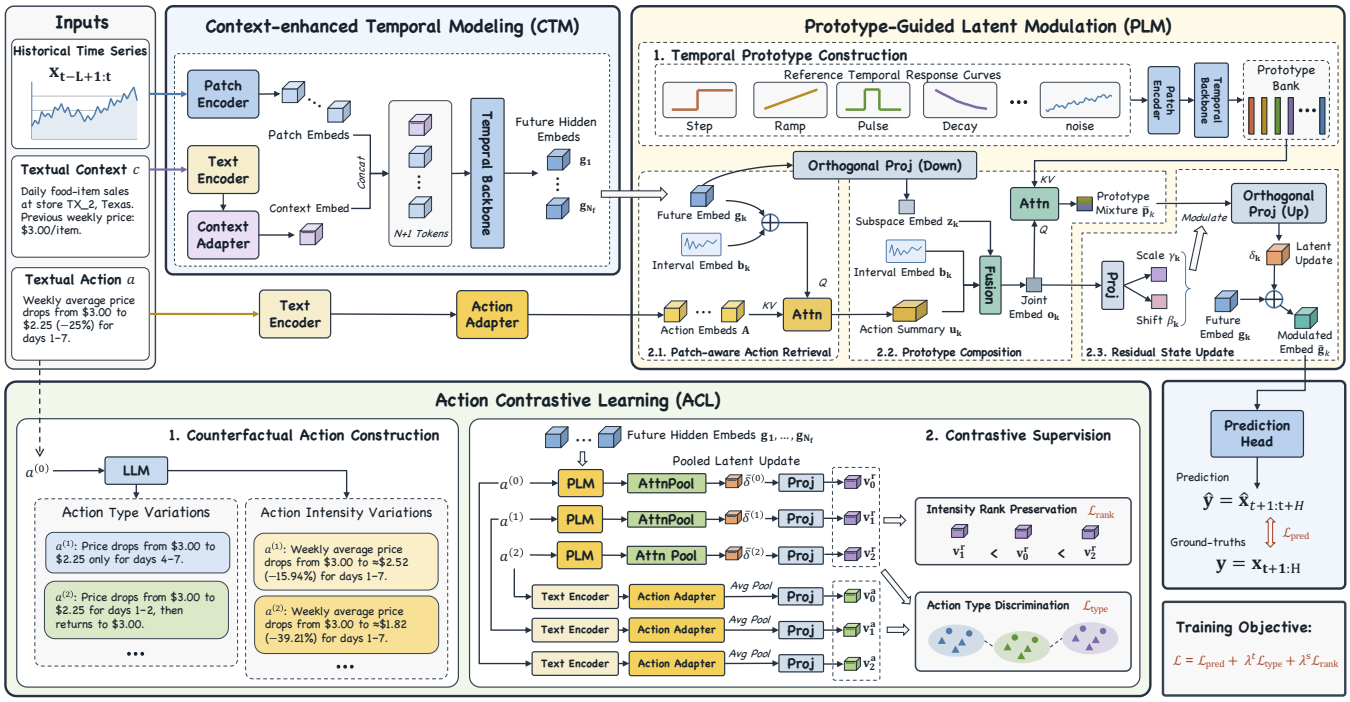

Figure 1: Overview of TimeWorld.

where Backbone returns the single-patch hidden state, and $R < D$ is the dimension of the latent update subspace introduced below. The resulting prototype bank $\mathbf{P} = [\mathbf{p}_1; \ldots; \mathbf{p}_{N_\mathrm{p}}] \in \mathbb{R}^{N_\mathrm{p} \times R}$ provides reusable temporal prototype representations for action-conditioned future modulation.

4.2.2 Patch-wise Modulation. To preserve pretrained knowledge of temporal evolution while accommodating action-induced deviations, we constrain latent updates to an R-dimensional learnable subspace. This subspace is spanned by an orthonormal basis matrix $\mathbf{Q} \in \mathbb{R}^{D \times R}$ satisfying $\mathbf{Q}^\top \mathbf{Q} = \mathrm{Id}_R$, where Id_R is the R-dimensional identity matrix. The basis is obtained through differentiable reduced QR decomposition of a trainable matrix $\mathbf{\Theta} \in \mathbb{R}^{D \times R}$. For each future patch, modulation involves three steps: retrieving relevant action information, composing prototypes based on the action and temporal state, and applying a residual update.

Patch-aware Action Retrieval. An action description may specify operations, magnitudes, and timing that affect different forecast intervals. To selectively retrieve such information, we retain the token-level representations $\mathrm{TextEnc}^a(a)$ from the pretrained text encoder, rather than the pooled context representation used in *CTM*. A two-layer adapter Adapt^a produces TSFM-aligned action embeddings $\mathbf{A} = \mathrm{Adapt}^a(\mathrm{TextEnc}^a(a)) \in \mathbb{R}^{N_\mathrm{a} \times D}$, where N_a is the number of action tokens.

For future patch k, we encode its start and end steps $(\mathrm{start}_k, \mathrm{end}_k)$ relative to the forecast origin using sinusoidal features and a linear projection, obtaining an interval embedding $\mathbf{b}_k \in \mathbb{R}^D$. Combining this embedding with the normalized temporal state $\tilde{\mathbf{g}}_k = \mathrm{LayerNorm}(\mathbf{g}_k)$, we retrieve relevant action information through cross-attention:

$$\mathbf{u}_k = \mathrm{Attn}(\mathrm{Linear}(\tilde{\mathbf{g}}_k) + \mathbf{b}_k, \mathbf{A}) \in \mathbb{R}^D, \tag{2}$$

where Linear is a trainable D-dimensional projection, and Attn denotes cross-attention with trainable query, key, and value projections. The resulting $\mathbf{u}_k$ summarizes action information relevant to the prediction interval of patch k.

Prototype Composition. Since the response to an action also depends on the current temporal dynamics, we combine the retrieved action information with the hidden state to select relevant prototypes. We first project $\mathbf{g}_k$ onto the learned subspace, obtaining its coordinates $\mathbf{z}_k = \mathbf{Q}^\top \mathbf{g}_k \in \mathbb{R}^R$ in the orthonormal basis $\mathbf{Q}$. A two-layer MLP then integrates the subspace coordinates, action information, and interval embedding: $\mathbf{o}_k = \mathrm{MLP}([\mathbf{z}_k; \mathbf{u}_k; \mathbf{b}_k]) \in \mathbb{R}^D$. Using this joint representation as the query, we retrieve a weighted mixture of temporal prototypes:

$$\bar{\mathbf{p}}_k = \mathrm{Attn}(\mathbf{o}_k, \mathbf{P}) \in \mathbb{R}^R, \tag{3}$$

where a separate cross-attention layer adaptively combines the prototypes according to the action and temporal state.

Residual State Update. To flexibly adjust the selected response patterns, two linear heads map $\mathbf{o}_k$ to a signed scale $\boldsymbol{\gamma}_k \in \mathbb{R}^R$ and an additive shift $\boldsymbol{\beta}_k \in \mathbb{R}^R$. The resulting update is mapped back to the TSFM hidden space and added to the original state:

$$\boldsymbol{\delta}_k = \mathbf{Q}(\boldsymbol{\gamma}_k \odot \bar{\mathbf{p}}_k + \boldsymbol{\beta}_k), \qquad \widetilde{\mathbf{g}}_k = \mathbf{g}_k + \boldsymbol{\delta}_k, \tag{4}$$

where $\odot$ denotes element-wise multiplication. The signed scale controls the strength and direction of prototype contributions, while the shift provides additional state-dependent adjustments. Since $\boldsymbol{\delta}_k$ lies in the subspace spanned by $\mathbf{Q}$, the update preserves the

orthogonal component of the original hidden state. The adjusted states $\{\widetilde{\mathbf{g}}_k\}_{k=1}^{N_\mathrm{f}}$ are decoded by the original prediction head Dec, giving final predictions $\{\widehat{\mathbf{y}}_k\}_{k=1}^{N_\mathrm{f}}$.

4.3 Action Contrastive Learning

While *PLM* achieves fine-grained action grounding, learning from factual actions alone may still be dominated by history patterns and overlook the specific characteristics of the given action. To strengthen action sensitivity, *ACL* explicitly contrasts the latent updates induced by alternative actions under the same history and context. It consists of two components: *(1) Counterfactual Action Construction*, which generates variations in action type and intensity, and *(2) Contrastive Supervision*, which distinguishes action-specific responses and preserves their relative intensity relationships.

4.3.1 Counterfactual Action Construction. We construct two groups of alternative actions that vary in type or intensity, both including the factual action. To distinguish these candidates, we index the actions and write the original factual action a as $a^{(0)}$. For type variations, an LLM first summarizes the action categories in each dataset's training split according to their operations or temporal policies. Given a factual action, it then selects N_t distinct categories other than the factual category and jointly generates one alternative per category, forming $\{a^{(i)}\}_{i=0}^{N_\mathrm{t}}$. For intensity variations, we randomly assign the factual action a rank $\rho^{(0)} \in \{0, \ldots, N_\mathrm{s}\}$ among $N_\mathrm{s}+1$ levels ordered from weakest to strongest, and sample a percentage change between adjacent levels. Given this rank and percentage, the LLM jointly generates N_s alternatives by progressively decreasing intervention magnitudes below the factual rank and increasing them above it, while retaining the action type. Together with the unchanged factual action, these alternatives form $\{a^{(i)}\}_{i=0}^{N_\mathrm{s}}$. In both cases, generation preserves the factual description's format and all content unrelated to the intended variation. Further construction and validation details are provided in Appendix C.

4.3.2 Contrastive Supervision. Since alternative actions share the same historical observations and context, their hidden states may be dominated by common temporal information rather than action-specific effects. We therefore apply contrastive supervision directly to the latent updates from *PLM*, defined as $\boldsymbol{\delta}_k^{(i)} = \widetilde{\mathbf{g}}_k^{(i)} - \mathbf{g}_k^{(i)}$ for the i-th action and k-th future patch.

Latent Response Aggregation. To capture each action's overall effect across the forecast horizon, we aggregate its patch-level latent updates through attention-weighted pooling:

$$\bar{\boldsymbol{\delta}}^{(i)} = \sum_{k=1}^{N_\mathrm{f}} \omega_k^{(i)} \boldsymbol{\delta}_k^{(i)}, \quad \omega_k^{(i)} = \mathrm{softmax}_k(\mathbf{w}^\top \boldsymbol{\delta}_k^{(i)} + \log \kappa_k), \tag{5}$$

where $\mathbf{w} \in \mathbb{R}^D$ is trainable, κ_k denotes the valid fraction of patch k excluding padding, and $\mathrm{softmax}_k$ normalizes weights across future patches. We also obtain an action representation $\bar{\mathbf{a}}^{(i)}$ by averaging the valid token embeddings in $\mathbf{A}^{(i)}$. The pooled action and latent response representations are projected into a shared embedding space and ℓ_2-normalized:

$$\mathbf{v}_i^\mathrm{a} = \mathrm{Norm}(\mathrm{Proj}_\mathrm{a}(\bar{\mathbf{a}}^{(i)})), \qquad \mathbf{v}_i^\mathrm{r} = \mathrm{Norm}(\mathrm{Proj}_\mathrm{r}(\bar{\boldsymbol{\delta}}^{(i)})), \tag{6}$$

where Proj_a and Proj_r are trainable linear projections to the same dimension, and Norm denotes ℓ_2 normalization.

Action Type Discrimination. To distinguish the effects of different action types, we align each action embedding with its corresponding latent response while separating it from responses to alternative actions. For the $N_\mathrm{t}+1$ actions contrasted by type, we define the action-response similarity as $\xi_{i,i'}^\mathrm{t} = (\mathbf{v}_i^\mathrm{a})^\top \mathbf{v}_{i'}^\mathrm{r}/\tau^\mathrm{t}$, where i indexes the action, i' indexes the response, and $\tau^\mathrm{t} > 0$ is the temperature. We adopt a bidirectional contrastive objective:

$$\mathcal{L}_\mathrm{type} = -\frac{1}{2(N_\mathrm{t}+1)} \sum_{i=0}^{N_\mathrm{t}} \log \frac{\exp(2\xi_{i,i}^\mathrm{t})}{(\sum_{i'=0}^{N_\mathrm{t}} \exp(\xi_{i,i'}^\mathrm{t}))(\sum_{i'=0}^{N_\mathrm{t}} \exp(\xi_{i',i}^\mathrm{t}))}. \tag{7}$$

The two terms enforce action-to-response and response-to-action matching, respectively, using other action types under the same observations as negatives.

Intensity Rank Preservation. Unlike distinct action types, intensity variations have an inherent ordering that should be reflected in their latent responses. Inspired by Rank-N-Contrast [45], we encourage responses with closer intensity ranks to exhibit greater similarity. For the $N_\mathrm{s}+1$ actions ordered by intensity, let $\rho^{(i)}$ denote the intensity rank of $a^{(i)}$, and define the rank distance $\mathrm{dist}_{i,i'} = |\rho^{(i)} - \rho^{(i')}|$ and response similarity $\xi_{i,i'}^\mathrm{s} = -\|\mathbf{v}_i^\mathrm{r} - \mathbf{v}_{i'}^\mathrm{r}\|_2/\tau^\mathrm{s}$, where $\tau^\mathrm{s} > 0$ is the temperature. For each anchor i and candidate i', the comparison set $\mathcal{S}_{i,i'} = \{i'' \in \{0, \ldots, N_\mathrm{s}\} \setminus \{i\} : \mathrm{dist}_{i,i''} \geq \mathrm{dist}_{i,i'}\}$ contains candidates with equal or greater rank distances. The ranking objective is

$$\mathcal{L}_\mathrm{rank} = -\frac{1}{(N_\mathrm{s}+1)N_\mathrm{s}} \sum_{i=0}^{N_\mathrm{s}} \sum_{\substack{i'=0 \\ i' \neq i}}^{N_\mathrm{s}} \log \frac{\exp(\xi_{i,i'}^\mathrm{s})}{\sum_{i'' \in \mathcal{S}_{i,i'}} \exp(\xi_{i,i''}^\mathrm{s})}. \tag{8}$$

This preserves relative intensity relationships in the latent updates without assuming a fixed direction of state change.

4.4 Model Optimization

For a single instance, we jointly optimize factual forecasting accuracy and the two counterfactual contrastive objectives:

$$\mathcal{L} = \mathcal{L}_\mathrm{pred} + \lambda^\mathrm{t}\mathcal{L}_\mathrm{type} + \lambda^\mathrm{s}\mathcal{L}_\mathrm{rank}, \tag{9}$$

where $\lambda^\mathrm{t}, \lambda^\mathrm{s} \geq 0$ balance the contrastive objectives; the full training objective averages $\mathcal{L}$ over all training instances. The forecasting loss is $\mathcal{L}_\mathrm{pred} = \frac{1}{H\sigma_\mathbf{x}^2}\sum_{j=1}^{H}(\widehat{y}_j - y_j)^2$, where j indexes forecast steps and $\sigma_\mathbf{x}$ is the historical standard deviation. The pretrained TSFM and text encoders $\mathrm{TextEnc}^c$, $\mathrm{TextEnc}^a$ remain frozen, while the adapters, *PLM*, and contrastive projections are jointly optimized. Appendix G further examines whether fine-tuning the TSFM yields additional gains. Counterfactual branches and contrastive projections are used only during training, adding no inference overhead. The computational complexity is detailed in Appendix D.

5 Experiments

5.1 Experimental Setup

Datasets. We use fixed chronological train/validation/test splits targeting a 60%/ 20%/20% ratio. History lengths L and forecast horizons H reflect each dataset's sampling frequency and intervention timescale. Except for Causal Chambers, each dataset has two (L, H) settings, whose performance is averaged. Detailed descriptions of datasets are provided in Appendix E.

Table 3: Comparison results of four TSFM backbones. Vanilla performs zero-shot forecasting by TSFMs, and Overall averages the 13 settings across seven datasets. The best and second-best results of each metric are shown in red and blue, respectively. Underlining marks the best result for each TSFM backbone.

Dataset	Metric	TimesFM 3.0			Chronos-2			Toto 2.0			Moirai 2.0		
		Vanilla	TSFM++	TimeWorld	Vanilla	TSFM++	TimeWorld	Vanilla	TSFM++	TimeWorld	Vanilla	TSFM++	TimeWorld
M5	nMAE	1.1612	0.9658	0.9803	0.8784	1.0204	0.9829	**0.8705**	0.9255	0.9917	**0.8602**	0.9788	1.0219
	nRMSE	1.7597	1.3084	1.2997	1.4231	**1.2914**	**1.2991**	1.4190	1.3742	1.3102	1.4071	1.3476	1.3002
	MASE	2.4538	2.0766	2.1151	1.9042	2.2255	2.1294	**1.8886**	1.9649	2.0875	**1.8606**	2.1473	2.1837
MOF	nMAE	0.7793	0.7787	0.8153	0.8019	0.7569	**0.7347**	0.7673	0.7908	0.7638	**0.7235**	0.7414	0.7527
	nRMSE	0.8066	0.8012	0.8402	0.8286	0.7845	**0.7604**	0.7958	0.8178	0.7923	**0.7488**	0.7677	0.7772
	MASE	3.2255	3.2568	3.4272	3.3110	3.0944	**3.0290**	3.1208	3.1972	3.1621	**3.0477**	3.1069	3.1953
FOMC	nMAE	0.9576	0.9530	0.9458	0.9239	0.9177	0.9143	**0.8857**	0.9008	**0.8986**	0.9293	0.9223	0.9059
	nRMSE	0.9810	0.9750	0.9718	0.9480	0.9400	0.9382	**0.9132**	0.9266	**0.9247**	0.9533	0.9455	0.9316
	MASE	2.5115	2.5069	2.4799	2.3817	2.3699	2.3408	**2.2863**	2.3335	2.3314	2.3935	2.3728	**2.3112**
Causal Chambers	nMAE	4.4275	4.4563	**3.2116**	4.5003	4.6659	4.4761	4.7793	4.7962	**4.0175**	4.8220	4.5824	4.1415
	nRMSE	4.7647	4.8067	**3.6509**	4.8322	4.9809	4.7993	5.1062	5.1089	**4.3950**	5.1251	4.9205	4.4478
	MASE	13.2580	13.3632	**9.3159**	13.8178	14.2030	13.6584	14.6407	14.6188	**12.1117**	14.6040	14.1977	12.4113
DaISy	nMAE	0.4025	0.2971	**0.2289**	0.4191	0.2843	0.2437	0.3957	0.3536	**0.2430**	0.4212	0.3348	0.2506
	nRMSE	0.4025	0.2971	**0.2289**	0.4191	0.2843	0.2437	0.3957	0.3536	**0.2430**	0.4212	0.3348	0.2506
	MASE	0.9123	0.6732	**0.5183**	0.9495	0.6439	0.5522	0.8966	0.8018	**0.5507**	0.9544	0.7595	0.5661
Shanghai Diabetes	nMAE	0.7221	0.6725	0.6662	0.7368	0.6703	**0.6418**	0.7383	0.7054	0.6866	0.7212	0.6907	**0.6563**
	nRMSE	0.7978	0.7435	0.7360	0.8129	0.7405	**0.7091**	0.8166	0.7774	0.7565	0.7971	0.7624	**0.7240**
	MASE	3.9093	3.6565	3.6637	3.9866	3.6276	**3.4871**	3.9903	3.8338	3.7984	3.8950	3.7245	**3.5913**
HUPA-UCM	nMAE	0.6148	**0.5939**	0.6057	0.6441	0.6083	0.6023	0.6197	0.6276	**0.5923**	0.6474	0.6193	0.6048
	nRMSE	0.7264	**0.6969**	0.7181	0.7540	0.7069	0.7022	0.7308	0.7380	**0.6972**	0.7536	0.7256	0.7063
	MASE	3.1235	**3.0106**	3.0610	3.2746	3.1067	3.0698	3.1974	3.1969	**3.0060**	3.2944	3.1684	3.1117
Overall	nMAE	1.0540	0.9983	**0.8997**	1.0237	1.0140	0.9781	1.0257	1.0310	**0.9515**	1.0329	1.0121	0.9635
	nRMSE	1.2087	1.1116	**1.0185**	1.1695	1.1135	1.0850	1.1730	1.1603	1.0648	1.1759	1.1298	**1.0636**
	MASE	3.5023	3.3634	**3.0651**	3.4948	3.4107	3.2981	3.4923	3.4827	**3.2295**	3.4996	3.4428	3.2562

Hyperparameters. We retain each TSFM's native input and output patch lengths L_{in}, L_{out} and hidden dimension D. We set the modulation subspace dimension to $R = 32$ and use $N_{\text{p}} = 24$ temporal prototypes, with $N_{\text{t}} = N_{\text{s}} = 5$ alternatives for action type and intensity, respectively. The type and intensity contrastive objectives use temperatures $\tau^{\text{t}} = 0.1$ and $\tau^{\text{s}} = 0.2$ and loss weights $\lambda^{\text{t}} = \lambda^{\text{s}} = 0.01$. The text adapters have a hidden width of 256, while the contrastive projections map to a shared 128-dimensional space.

Training Configurations. We encode text with Qwen3-Embedding-4B [47] and generate counterfactual actions with GPT-6 Luna [31]. TSFM++ and TimeWorld are trained per dataset setting with frozen TSFMs and text encoders, using AdamW with a default learning rate of 10^{-4}, weight decay of 0.01, and an effective batch size of 16. We train for up to ten epochs, with early stopping after three epochs without validation improvement, and retain the checkpoint with the lowest validation forecasting loss. All runs use NVIDIA H200 GPUs and a fixed seed of 42 to control training randomness and facilitate reproducibility.

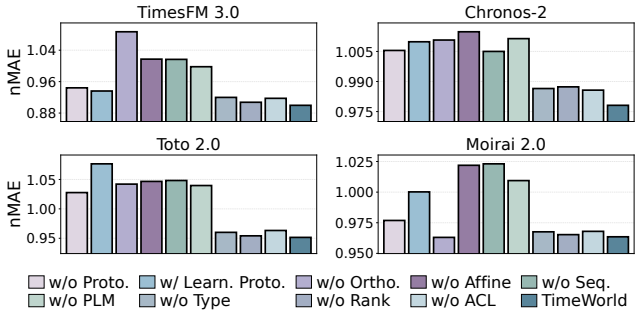

Figure 2: Ablation results of TimeWorld.

5.2 Comparison Results

Table 3 compares TimeWorld with vanilla TSFMs and TSFM++, revealing three main observations: (1) *TimeWorld consistently improves aggregate forecasting performance across backbones.* It reduces all three overall errors relative to both alternatives, with nMAE improvements of 3.5%–9.9% over TSFM++. These gains support structured action-dependent updates to pretrained world dynamics beyond direct text-prefix conditioning. (2) *Stronger history-only forecasting does not imply greater benefits from action conditioning.* TimesFM 3.0 has the highest overall errors among vanilla backbones but achieves the lowest errors with both TSFM++ and TimeWorld. Its larger gains, compared with the more modest improvements of Chronos-2, indicate that the ability to exploit textual actions varies across pretrained backbones. (3) *The benefits of action modeling vary across domains.* TimeWorld improves all three metrics over both alternatives on Causal Chambers and DaISy for every backbone. Healthcare forecasting also generally improves, though TSFM++ remains stronger with TimesFM on HUPA-UCM. Economic results are less consistent: improvements on FOMC are small, MOF results depend on the backbone, and M5 often favors vanilla forecasting

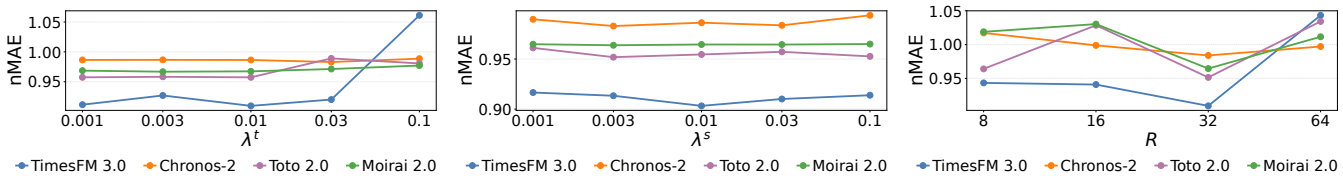

(a) Type weight λ^t. **(b) Intensity weight λ^s.** **(c) Subspace dimension R.**

Figure 3: Parameter sensitivity of TimeWorld across four TSFM backbones, measured by nMAE.

in nMAE and MASE despite reductions in nRMSE. Confidence intervals for the gains over vanilla TSFM forecasting are reported in Appendix H.

5.3 Ablation Study

In this section, we evaluate nine variants to isolate the contributions of *PLM* and *ACL*. Within PLM, "w/o Proto." directly modulates the subspace coordinates without prototypes, "w/ Learn. Proto." replaces encoded prototypes with randomly initialized learnable vectors, "w/o Ortho." replaces the orthogonal subspace projections with MLP mappings, "w/o Affine" removes the final scale and shift, and "w/o Seq." replaces token-level action attention with a single action embedding. "w/o PLM" uses action-prefix conditioning instead of latent modulation, while "w/o Type", "w/o Rank", and "w/o ACL" remove type discrimination, intensity ranking, and both contrastive objectives, respectively. Figure 2 yields three observations: (1) *Structured latent modulation improves action grounding.* Removing prototypes, affine adjustment, or token-level retrieval increases nMAE for every backbone, as does replacing PLM with action-prefix conditioning. Learnable prototypes also worsen all four backbones, supporting the value of temporal response priors encoded by pretrained TSFMs. These results support composing temporal response patterns and adapting them to the action and current state. (2) *The subspace parameterization matters, but its effect depends on the backbone.* Switching it degrades TimesFM 3.0, Chronos-2, and Toto 2.0, while Moirai 2.0 changes only marginally in the opposite direction. (3) *Both contrastive objectives contribute to forecasting accuracy.* Removing either objective or ACL as a whole worsens all four backbones, supporting complementary supervision of action type and intensity.

5.4 Parameter Analysis

In this section, we examine the contrastive weights λ^t and λ^s, which balance action discrimination against forecasting accuracy, and the subspace dimension R, which controls the capacity of latent modulation. Each contrastive weight is varied independently over $\{0.001, 0.003, 0.01, 0.03, 0.1\}$, and R over $\{8, 16, 32, 64\}$, keeping the remaining settings fixed. Figure 3 shows that: (1) *Moderate type weights provide stable performance, but excessive weighting can be harmful.* TimesFM is relatively stable up to $\lambda^t = 0.03$, whereas increasing it to 0.1 raises nMAE by 16.7% relative to 0.01. This may reflect overly strong pressure to separate action types at the expense of factual prediction. (2) *Performance is particularly robust to the intensity weight.* Across the full range of λ^s, each backbone's nMAE remains within 1.5% of its best value, indicating limited sensitivity to the strength of rank supervision. (3) *An intermediate subspace dimension balances capacity and constraint.* All four backbones achieve their lowest nMAE at $R = 32$, supporting our chosen dimension. Smaller subspaces may restrict useful action-dependent updates, while larger ones permit unnecessary changes to pretrained representations.

5.5 Action Consistency Analysis

Forecast errors measure agreement with observed outcomes but do not reveal whether a world model distinguishes alternative actions. We therefore examine representation similarity as a proxy for counterfactual action sensitivity, comparing TimeWorld with w/o ACL under the same history and context. In Table 4, *Same* and *Different* measure cosine similarity between action embeddings and responses to the corresponding or alternative action types. *Adjacent* and *Distant* compare response embeddings at neighboring and maximally separated intensity ranks. Both models use a shared readout, and scores are averaged across backbones and datasets. We observe that ACL increases the separation between matching and mismatched action types while making responses to nearby intensity ranks more similar than those to distant ranks. This supports more distinguishable counterfactual responses beyond improvements in factual forecasting.

Table 4: Action-conditioned representation similarity.

Model	Action Type		Action Intensity	
	Same	Different	Adjacent	Distant
TimeWorld	0.1077	0.0533	0.9815	0.9380
w/o ACL	0.0293	0.0271	0.9311	0.9284

We further examine the effectiveness of prototype composition in Appendix I and provide a qualitative case study in Appendix J.

6 Conclusion

In this work, we study time series world models that specify interventions through textual actions beyond fixed covariate schemas, and introduce TSWM-Bench for evaluation across seven real-world datasets. Our benchmark findings suggest that pretrained TSFMs encode implicit world knowledge of temporal evolution, providing a foundation for action-conditioned world modeling. Building on this insight, we propose TimeWorld, a plug-in framework that adapts pretrained TSFMs through explicit action grounding and discrimination. Prototype-guided latent modulation translates textual

actions into state- and interval-dependent updates, while action contrastive learning distinguishes responses to different action types and preserves relative intensity relationships under shared histories. Experiments across four TSFM backbones demonstrate consistent improvements in aggregate forecasting accuracy over history-only and action-prefix baselines. Ablation and parameter analyses substantiate the effectiveness of latent modulation and contrastive supervision, while action consistency analysis further reveals stronger action-type discrimination and more coherent intensity-dependent responses.

References

[1] Abdul Fatir Ansari, Oleksandr Shchur, Jaris Küken, Andreas Auer, Boran Han, Pedro Mercado, Syama Sundar Rangapuram, Huibin Shen, Lorenzo Stella, Xiyuan Zhang, Mononito Goswami, Shubham Kapoor, Danielle C. Maddix, Pablo Guerron, Tony Hu, Junming Yin, Nick Erickson, Prateek Mutalik Desai, Hao Wang, Huzefa Rangwala, George Karypis, Yuyang Wang, and Michael Bohlke-Schneider. 2025. Chronos-2: From Univariate to Universal Forecasting. arXiv:2510.15821 [cs.LG] https://arxiv.org/abs/2510.15821

[2] Adrien Bardes, Quentin Garrido, Jean Ponce, Xinlei Chen, Michael Rabbat, Yann LeCun, Mahmoud Assran, and Nicolas Ballas. 2024. Revisiting Feature Prediction for Learning Visual Representations from Video. arXiv:2404.08471 [cs.CV] https://arxiv.org/abs/2404.08471

[3] S. Bittanti and L. Piroddi. 1997. Nonlinear identification and control of a heat exchanger: A neural network approach. *Journal of the Franklin Institute* 334, 1 (1997), 135–153. doi:10.1016/S0016-0032(96)00059-2

[4] Board of Governors of the Federal Reserve System. [n. d.]. Meeting Calendars, Statements, and Minutes. Federal Open Market Committee. https://www.federalreserve.gov/monetarypolicy/fomccalendars.htm

[5] Jake Bruce, Michael D Dennis, Ashley Edwards, Jack Parker-Holder, Yuge Shi, Edward Hughes, Matthew Lai, Aditi Mavalankar, Richie Steigerwald, Chris Apps, Yusuf Aytar, Sarah Maria Elisabeth Bechtle, Feryal Behbahani, Stephanie C.Y. Chan, Nicolas Heess, Lucy Gonzalez, Simon Osindero, Sherjil Ozair, Scott Reed, Jingwei Zhang, Konrad Zolna, Jeff Clune, Nando De Freitas, Satinder Singh, and Tim Rocktäschel. 2024. Genie: Generative Interactive Environments. In *Proceedings of the 41st International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 235)*. PMLR, 4603–4623. https://proceedings.mlr.press/v235/bruce24a.html

[6] Hanyin Cheng, Linfeng Wang, Zhengbo Qu, Yang Shu, Zhongwen Rao, Meng Wang, Yijie Li, Xin Jiang, Bin Yang, and Chenjuan Guo. 2026. QiYao-M: Multimodal Time Series Foundation Model with Role-Aware Modeling of Endogenous and Exogenous Modalities. arXiv:2609.34842 [cs.LG] https://arxiv.org/abs/2609.34842

[7] Xilin Dai, Hongzhou Chen, Yifan Hu, Yiding Liu, Zewei Dong, and Jiang-Ming Yang. 2026. Pythia: Toward Foundation World Models for Multimodal Time Series. arXiv:2610.05240 [cs.LG] https://arxiv.org/abs/2610.05240

[8] Abhimanyu Das, Weihao Kong, Rajat Sen, and Yichen Zhou. 2024. A decoder-only foundation model for time-series forecasting. In *Proceedings of the 41st International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 235)*. PMLR, 10148–10167. https://proceedings.mlr.press/v235/das24c.html

[9] Kuiye Ding, Fanda Fan, Yao Wang, Ruijie Jian, Xiaorui Wang, Luqi Gong, Yishan Jiang, Chunjie Luo, and Jianfeng Zhan. 2025. DualSG: A Dual-Stream Explicit Semantic-Guided Multivariate Time Series Forecasting Framework. In *Proceedings of the 33rd ACM International Conference on Multimedia (MM '25)*. ACM, 508–517. doi:10.1145/3746027.3755458

[10] Juan L. Gamella, Jonas Peters, and Peter Bühlmann. 2025. Causal chambers as a real-world physical testbed for AI methodology. *Nature Machine Intelligence* 7, 1 (2025), 107–118. doi:10.1038/s42256-024-00964-x

[11] Shuqi Gu, Yongxiang Zhao, Baoyu Jing, and Kan Ren. 2026. What if Tomorrow is the World Cup Final? Counterfactual Time Series Forecasting with Textual Conditions. In *Proceedings of the 43rd International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 306)*, Tong Zhang, Miroslav Dudik, Martin Jaggi, Alekh Agarwal, Sharon Li, Dale Schuurmans, Jerry Zhu, Felix Berkenkamp, Hanze Dong, and Alberto Bietti (Eds.). PMLR, 37376–37394. https://proceedings.mlr.press/v306/gu26r.html

[12] Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, and Timothy Lillicrap. 2025. Mastering diverse control tasks through world models. *Nature* 640 (2025), 647–653. doi:10.1038/s41586-025-08744-2

[13] J. Ignacio Hidalgo, Jorge Alvarado, Marta Botella, Aranzazu Aramendi, J. Manuel Velasco, and Oscar Garnica. 2024. HUPA-UCM Diabetes Dataset. doi:10.17632/3hbcscwz44.1

[14] Ayush Jain and Rajat Sen. 2026. TimesFM-3: A zero-shot foundation model for multivariate forecasting. Google Research Blog. Published August 31, 2026; accessed October 9, 2026. https://research.google/blog/timesfm-3-a-zero-shot-foundation-model-for-multivariate-forecasting/

[15] Joel Jang, Seonghyeon Ye, Zongyu Lin, Jiannan Xiang, Johan Bjorck, Yu Fang, Fengyuan Hu, Spencer Huang, Kaushil Kundalia, Yen-Chen Lin, Loïc Magne, Ajay Mandlekar, Avnish Narayan, You Liang Tan, Guanzhi Wang, Jing Wang, Qi Wang, Yinzhen Xu, Xiaohui Zeng, Kaiyuan Zheng, Ruijie Zheng, Ming-Yu Liu, Luke Zettlemoyer, Dieter Fox, Jan Kautz, Scott Reed, Yuke Zhu, and Linxi Fan. 2025. DreamGen: Unlocking Generalization in Robot Learning through Video World Models. In *Proceedings of The 9th Conference on Robot Learning (Proceedings of Machine Learning Research, Vol. 305)*, Joseph Lim, Shuran Song, and Hae-Won Park (Eds.). PMLR, 5170–5194. https://proceedings.mlr.press/v305/jang25a.html

[16] Ming Jin, Shiyu Wang, Lintao Ma, Zhixuan Chu, James Zhang, Xiaoming Shi, Pin-Yu Chen, Yuxuan Liang, Yuan-Fang Li, Shirui Pan, and Qingsong Wen. 2024. Time-LLM: Time Series Forecasting by Reprogramming Large Language Models. In *International Conference on Learning Representations*, B. Kim, Y. Yue, S. Chaudhuri, K. Fragkiadaki, M. Khan, and Y. Sun (Eds.), Vol. 2024. 23857–23880. https://proceedings.iclr.cc/paper_files/paper/2024/file/680b2a8135b9c71278a09cafb605869e-Paper-Conference.pdf

[17] Emaad Khwaja, Chris Lettieri, Gerald Woo, Eden Belouadah, Marc Cenac, Guillaume Jarry, Enguerrand Paquin, Xunyi Zhao, Viktoriya Zhukov, Othmane Abou-Amal, Chenghao Liu, Ameet Talwalkar, and David Asker. 2026. Toto 2.0: Time Series Forecasting Enters the Scaling Era. arXiv:2605.20119 [cs.LG] https://arxiv.org/abs/2605.20119

[18] Aggelos Lazaris and Viktor K. Prasanna. 2019. An LSTM Framework For Modeling Network Traffic. In *2019 IFIP/IEEE Symposium on Integrated Network and Service Management (IM)*. IEEE. https://ieeexplore.ieee.org/document/8717920/

[19] Chenghao Liu, Taha Aksu, Juncheng Liu, Xu Liu, Hanshu Yan, Quang Pham, Silvio Savarese, Doyen Sahoo, Caiming Xiong, and Junnan Li. 2025. Moirai 2.0: When Less Is More for Time Series Forecasting. arXiv:2511.11698 [cs.LG] https://arxiv.org/abs/2511.11698

[20] Chenxi Liu, Qianxiong Xu, Hao Miao, Sun Yang, Lingzheng Zhang, Cheng Long, Ziyue Li, and Rui Zhao. 2025. TimeCMA: Towards LLM-Empowered Multivariate Time Series Forecasting via Cross-Modality Alignment. *Proceedings of the AAAI Conference on Artificial Intelligence* 39, 18 (April 2025), 18780–18788. doi:10.1609/aaai.v39i18.34067

[21] Haoxin Liu, Shangqing Xu, Zhiyuan Zhao, Lingkai Kong, Harshavardhan Kamarthi, Aditya B. Sasanur, Megha Sharma, Jiaming Cui, Qingsong Wen, Chao Zhang, and B. Aditya Prakash. 2024. Time-MMD: Multi-Domain Multimodal Dataset for Time Series Analysis. In *Advances in Neural Information Processing Systems*, A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang (Eds.), Vol. 37. Curran Associates, Inc., 77888–77933. doi:10.52202/079017-2476

[22] Xu Liu, Junfeng Hu, Yuan Li, Shizhe Diao, Yuxuan Liang, Bryan Hooi, and Roger Zimmermann. 2024. UniTime: A Language-Empowered Unified Model for Cross-Domain Time Series Forecasting. In *Proceedings of the ACM Web Conference 2024 (WWW '24)*. ACM, 4095–4106. doi:10.1145/3589334.3645434

[23] Yong Liu, Haoran Zhang, Chenyu Li, Xiangdong Huang, Jianmin Wang, and Mingsheng Long. 2024. Timer: Generative Pre-trained Transformers Are Large Time Series Models. In *Proceedings of the 41st International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 235)*, Ruslan Salakhutdinov, Zico Kolter, Katherine Heller, Adrian Weller, Nuria Oliver, Jonathan Scarlett, and Felix Berkenkamp (Eds.). PMLR, 32369–32399. https://proceedings.mlr.press/v235/liu24cb.html

[24] Yihang Lu, Yangyang Xu, Qitao Qin, and Xianwei Meng. 2025. TimeCapsule: Solving the Jigsaw Puzzle of Long-Term Time Series Forecasting with Compressed Predictive Representations. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2 (KDD '25)*. ACM, 1987–1998. doi:10.1145/3711896.3737157

[25] Spyros Makridakis, Evangelos Spiliotis, and Vassilios Assimakopoulos. 2022. The M5 competition: Background, organization, and implementation. *International Journal of Forecasting* 38, 4 (2022), 1325–1336. doi:10.1016/j.ijforecast.2021.07.007

[26] Valentyn Melnychuk, Dennis Frauen, and Stefan Feuerriegel. 2022. Causal Transformer for Estimating Counterfactual Outcomes. In *Proceedings of the 39th International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 162)*, Kamalika Chaudhuri, Stefanie Jegelka, Le Song, Csaba Szepesvari, Gang Niu, and Sivan Sabato (Eds.). PMLR, 15293–15329. https://proceedings.mlr.press/v162/melnychuk22a.html

[27] Junjie Meng, Ranxu Zhang, Zi-an Zhang, Shujun Liu, Xiaoning Qi, Xiaozhou Xu, Yanyong Zhang, Hui Xiong, and Chao Wang. 2026. CEDAR: Controlled and Event-Driven Demand Forecasting via Residual Decomposition. In *Proceedings of the 32nd ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2*. ACM, 7797–7808. doi:10.1145/3770855.3818338

[28] Ministry of Finance, Japan. [n. d.]. Foreign Exchange Intervention Operations. Historical intervention data. https://www.mof.go.jp/english/policy/international_policy/reference/feio/foreign_exchange_intervention_operations.csv

[29] Hang Ni, Weijia Zhang, and Hao Liu. 2026. UniExtreme: A Universal Foundation Model for Extreme Weather Forecasting. In *Proceedings of the 32nd ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.1*. ACM, 1101–1112. doi:10.1145/3770854.3780172

[30] Hang Ni, Weijia Zhang, Fei Wang, Zezhi Shao, and Hao Liu. 2026. From Consistency to Complementarity: Aligned and Disentangled Multi-modal Learning for Time Series Understanding and Reasoning. arXiv:2601.21436 [cs.LG] https://arxiv.org/abs/2601.21436

[31] OpenAI. 2026. GPT-6 Luna. OpenAI API model documentation. Accessed October 14, 2026. https://developers.openai.com/api/docs/models/gpt-6-luna

[32] Xiaoqian Qi, Haoye Chai, Yue Wang, and Yong Li. 2026. Beyond Static Forecasting: Unleashing the Power of World Models for Mobile Traffic Extrapolation. arXiv:2604.08199 [cs.NI] https://arxiv.org/abs/2604.08199

[33] David Salinas, Valentin Flunkert, Jan Gasthaus, and Tim Januschowski. 2020. DeepAR: Probabilistic forecasting with autoregressive recurrent networks. *International Journal of Forecasting* 36, 3 (2020), 1181–1191. doi:10.1016/j.ijforecast.2019.07.001

[34] Xiaoming Shi, Shiyu Wang, Yuqi Nie, Dianqi Li, Zhou Ye, Qingsong Wen, and Ming Jin. 2025. Time-MoE: Billion-Scale Time Series Foundation Models with Mixture of Experts. In *International Conference on Learning Representations*, Y. Yue, A. Garg, N. Peng, F. Sha, and R. Yu (Eds.), Vol. 2025. 34635–34667. https://proceedings.iclr.cc/paper_files/paper/2025/file/558d48c1f08675daa636e09bfe94a89e-Paper-Conference.pdf

[35] Qwen Team. 2026. Qwen-RobotWorld Technical Report: Unifying Embodied World Modeling through Language-Conditioned Video Generation. (2026).

[36] Chengsen Wang, Qi Qi, Jingyu Wang, Haifeng Sun, Zirui Zhuang, Jinming Wu, Lei Zhang, and Jianxin Liao. 2025. ChatTime: A Unified Multimodal Time Series Foundation Model Bridging Numerical and Textual Data. *Proceedings of the AAAI Conference on Artificial Intelligence* 39, 12 (April 2025), 12694–12702. doi:10.1609/aaai.v39i12.33384

[37] Haotian Wang, Haoxuan Li, Hao Zou, Haoang Chi, Long Lan, Wanrong Huang, and Wenjing Yang. 2025. Effective and Efficient Time-Varying Counterfactual Prediction with State-Space Models. In *International Conference on Learning Representations*, Y. Yue, A. Garg, N. Peng, F. Sha, and R. Yu (Eds.), Vol. 2025. 15000–15025. https://proceedings.iclr.cc/paper_files/paper/2025/file/26f16a11f4dbba3eef876b571d9f3200-Paper-Conference.pdf

[38] Xinlei Wang, Maike Feng, Jing Qiu, Jinjin Gu, and Junhua Zhao. 2024. From News to Forecast: Integrating Event Analysis in LLM-Based Time Series Forecasting with Reflection. In *Advances in Neural Information Processing Systems*, A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang (Eds.), Vol. 37. Curran Associates, Inc., 58118–58153. doi:10.52202/079017-1853

[39] Andrew Robert Williams, Arjun Ashok, Étienne Marcotte, Valentina Zantedeschi, Jithendaraa Subramanian, Roland Riachi, James Requeima, Alexandre Lacoste, Irina Rish, Nicolas Chapados, and Alexandre Drouin. 2025. Context is Key: A Benchmark for Forecasting with Essential Textual Information. In *Proceedings of the 42nd International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 267)*, Aarti Singh, Maryam Fazel, Daniel Hsu, Simon Lacoste-Julien, Felix Berkenkamp, Tegan Maharaj, Kiri Wagstaff, and Jerry Zhu (Eds.). PMLR, 66887–66944. https://proceedings.mlr.press/v267/williams25a.html

[40] Xingjian Wu, Jianxin Jin, Wanghui Qiu, Peng Chen, Yang Shu, Bin Yang, and Chenjuan Guo. 2026. Aurora: Towards Universal Generative Multimodal Time Series Forecasting. In *International Conference on Learning Representations*, C. Vondrick, B. Hariharan, C. Raffel, L. Pinto, D. Yang, and A. Faust (Eds.), Vol. 2026. 103239–103277. https://proceedings.iclr.cc/paper_files/paper/2026/file/a8e18adf6489fd3a27417d128660eaa9-Paper-Conference.pdf

[41] Qianyi Xu, Gousia Habib, Feng Wu, Dilruk Perera, and Mengling Feng. 2026. medDreamer: Model-Based Reinforcement Learning with Latent Imagination on Complex EHRs for Clinical Decision Support. In *Proceedings of the 32nd ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.1*. ACM, 1693–1704. doi:10.1145/3770854.3780292

[42] Hao Xue and Flora D. Salim. 2024. PromptCast: A New Prompt-Based Learning Paradigm for Time Series Forecasting. *IEEE Transactions on Knowledge and Data Engineering* 36, 11 (Nov. 2024), 6851–6864. doi:10.1109/tkde.2023.3342137

[43] Jie Yang, Yifan Hu, Yuante Li, Kexin Zhang, Kaize Ding, and Philip S. Yu. 2026. From Observations to States: Latent Time Series Forecasting. In *Proceedings of the 43rd International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 306)*, Tong Zhang, Miroslav Dudik, Martin Jaggi, Alekh Agarwal, Sharon Li, Dale Schuurmans, Jerry Zhu, Felix Berkenkamp, Hanze Dong, and Alberto Bietti (Eds.). PMLR, 144796–144819. https://proceedings.mlr.press/v306/yang26y.html

[44] Ziyue Yang, Chaolin Xu, Yijing Wang, Tiankai Gu, Hui Yang, Yanhong Lin, Kaiyuan Liu, and Fei Xiao. 2026. ReasonCast: Agentic Demand Forecasting with Selective Semantic Reasoning. arXiv:2608.15291 [cs.AI] https://arxiv.org/abs/2608.15291

[45] Kaiwen Zha, Peng Cao, Jeany Son, Yuzhe Yang, and Dina Katabi. 2023. Rank-N-Contrast: Learning Continuous Representations for Regression. In *Advances in Neural Information Processing Systems*, Vol. 36. https://papers.nips.cc/paper_files/paper/2023/hash/39e9c5913c970e3e49c2df629daff636-Abstract-Conference.html

[46] Haochen Zhang, Jiaheng Guo, Zhen Xu, Zachary Plotkin, Nicholas Konz, Zhen Tan, and Tianlong Chen. 2026. On the Divergence of Accuracy and Mechanism Consistency in Time Series World Models. arXiv:2610.01842 [cs.AI] https://arxiv.org/abs/2610.01842

[47] Yanzhao Zhang, Mingxin Li, Dingkun Long, Xin Zhang, Huan Lin, Baosong Yang, Pengjun Xie, An Yang, Dayiheng Liu, Junyang Lin, Fei Huang, and Jingren Zhou. 2025. Qwen3 Embedding: Advancing Text Embedding and Reranking Through Foundation Models. *arXiv preprint arXiv:2506.05176* (2025). https://arxiv.org/abs/2506.05176

[48] Yang Zhang, Wenbo Yang, Jun Wang, Qiang Ma, and Jie Xiong. 2025. CAMEF: Causal-Augmented Multi-Modality Event-Driven Financial Forecasting by Integrating Time Series Patterns and Salient Macroeconomic Announcements. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2 (KDD '25)*. ACM, 3867–3878. doi:10.1145/3711896.3736872

[49] Zihao Zhang, Stefan Zohren, and Stephen Roberts. 2019. DeepLOB: Deep Convolutional Neural Networks for Limit Order Books. *IEEE Transactions on Signal Processing* 67, 11 (2019), 3001–3012. doi:10.1109/TSP.2019.2907260

[50] Siru Zhong, Weilin Ruan, Ming Jin, Huan Li, Qingsong Wen, and Yuxuan Liang. 2025. Time-VLM: Exploring Multimodal Vision-Language Models for Augmented Time Series Forecasting. In *Proceedings of the 42nd International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 267)*, Aarti Singh, Maryam Fazel, Daniel Hsu, Simon Lacoste-Julien, Felix Berkenkamp, Tegan Maharaj, Kiri Wagstaff, and Jerry Zhu (Eds.). PMLR, 78478–78497. https://proceedings.mlr.press/v267/zhong25a.html

[51] Jinhao Zhu. 2022. Diabetes Datasets, ShanghaiT1DM and ShanghaiT2DM. doi:10.6084/m9.figshare.20425518.v5

[52] Yuhan Zhu, Xiangfei Qiu, Hanyin Cheng, Wangmeng Shen, Chenjuan Guo, Bin Yang, Jilin Hu, and Christian S. Jensen. 2026. WorldTS: World Modeling for Multimodal Covariate-aware Time Series Forecasting. arXiv:2609.31162 [cs.LG] https://arxiv.org/abs/2609.31162

A Baseline Implementation Details

MM-TSFlib [21] uses iTransformer for temporal modeling and GPT-2 for text encoding. CovTSWM [46] combines TiDE with a frozen autoencoder and output gating to condition forecasts on structured actions. We adapt its covariate schema to each dataset's recorded controls and align their values, magnitudes, and event indicators with the forecast timeline. For FOMC, these covariates encode the explicitly announced target interest rate when available.

B Evaluation Metric Details

For evaluation across test instances, we add an instance index v to the notation in Section 3.1, yielding history $\mathbf{x}_v$, target $\mathbf{y}_v$, and prediction $\widehat{\mathbf{y}}_v$. For each instance, L and H denote its history length and forecast horizon. Here, t locally indexes chronological history positions in $x_{v,t}$ $(t = 1, \ldots, L)$ and forecast steps in $y_{v,t}$ and $\widehat{y}_{v,t}$ $(t = 1, \ldots, H)$. The comma separates instance and temporal indices. The prediction errors are

$$\mathrm{MAE}_v = \frac{1}{H}\sum_{t=1}^{H} |\widehat{y}_{v,t} - y_{v,t}|, \ \mathrm{RMSE}_v = \sqrt{\frac{1}{H}\sum_{t=1}^{H} (\widehat{y}_{v,t} - y_{v,t})^2}.$$

The history mean $\mu_{\mathbf{x}_v}$, population standard deviation $\sigma_{\mathbf{x}_v}$, and mean absolute lag-one difference $\mathrm{MAE}_v^{\mathrm{naive}}$ are computed separately from each instance's history:

$$\mu_{\mathbf{x}_v} = \frac{1}{L}\sum_{t=1}^{L} x_{v,t}, \ \ \sigma_{\mathbf{x}_v} = \sqrt{\frac{1}{L}\sum_{t=1}^{L} (x_{v,t} - \mu_{\mathbf{x}_v})^2},$$

$$\mathrm{MAE}_v^{\mathrm{naive}} = \frac{1}{L-1}\sum_{t=2}^{L} |x_{v,t} - x_{v,t-1}|.$$

These expressions assume fully observed windows. With missing values, each sum includes only observed targets, observed

history values, or adjacent history pairs with both values originally observed, as appropriate, and its denominator is replaced by the corresponding valid count. The normalized and scaled errors are $\mathrm{nMAE}_\nu = \mathrm{MAE}_\nu/\sigma_{\mathbf{x}_\nu}$, $\mathrm{nRMSE}_\nu = \mathrm{RMSE}_\nu/\sigma_{\mathbf{x}_\nu}$, $\mathrm{MASE}_\nu = \mathrm{MAE}_\nu/\mathrm{MAE}_\nu^{\mathrm{naive}}$, where MASE uses the nonseasonal naive scale computed from each instance's history. Instances with no valid observations or a zero normalization scale are excluded from the corresponding metric average. We average each metric over valid test instances for each (L, H) configuration, taking the RMSE square root before averaging. Dataset-level results average configurations within each dataset. Overall results weight all configurations across datasets equally.

C Counterfactual Action Construction Details

Each generation request to the LLM includes the factual action and specifies which attribute may change, so type alternatives vary the operation while intensity alternatives retain the operation and follow the assigned magnitude ordering. All candidates preserve unrelated context, units, and temporal information to keep the contrast focused on the intended action variation. The consistent and regular action descriptions in TSWM-Bench facilitate checks of format, required fields, and the requested variations. Programmatic validation of all generated actions, supplemented by manual inspection of sampled examples, revealed no formatting errors. However, handling irregular, unconstrained action text remains a nontrivial challenge beyond the scope of this paper, motivating future benchmarks and solutions for such inputs.

D Computational Complexity

TimeWorld augments a frozen TSFM with one context token and lightweight latent modulation. Beyond text encoding and native TSFM computation, adapters and PLM cost $O\big(N_\mathrm{a}D^2 + N_\mathrm{f}(D^2 + N_\mathrm{a}D + DR + N_\mathrm{p}R)\big)$ per forecast, assuming projection widths of $O(D)$. Reduced QR uses PyTorch's built-in differentiable `torch.linalg.qr` and costs $O(DR^2)$, while prototype projection costs $O(N_\mathrm{p}DR)$ given reusable frozen reference encodings. Both can be cached at inference. With $R = 32$ and $N_\mathrm{p} = 24$, PLM remains small and linear in N_f for fixed action length. Training additionally evaluates $1 + N_\mathrm{t} + N_\mathrm{s}$ action-conditioned branches for ACL. After pooling and projection, type matching costs $O((N_\mathrm{t}+1)^2D)$, while intensity ranking costs $O((N_\mathrm{s}+1)^2D + (N_\mathrm{s}+1)^3)$ under direct evaluation of the comparison-set denominators. The cubic term is limited to six intensity levels because $N_\mathrm{t} = N_\mathrm{s} = 5$. These extra branches and contrastive losses are training-only and incur no inference cost.

E Dataset Details

The history lengths and forecast horizons reflect each domain's sampling frequency and action timescale, as summarized in Table 5.

M5. M5 [25] contains daily item–store sales and weekly average prices. Actions describe recorded price changes, including the previous and updated prices and their effective interval. We use $L = 112$ days and $H \in \{7, 28\}$ days to capture weekly and four-week demand responses, with a fixed 1% sample shared by all methods.

MOF. MOF [28] pairs Japan's foreign-exchange interventions with yen-per-dollar observations from FRED. Actions specify the traded currencies, direction, and amount, combining operations on the same date. We use $L = 60$ market observations and $H \in \{1, 5\}$ to examine immediate and short-term responses, aligning observations to separate pre-intervention history from subsequent outcomes.

FOMC. FOMC [4] pairs monetary policy statements with U.S. two-year Treasury yields. The full statement serves as the textual action, retaining policy information beyond the announced rate. We use $L = 60$ and $H \in \{1, 5\}$ market observations, with targets beginning after the announcement date.

Causal Chambers. Causal Chambers [10] provides controlled wind-tunnel experiments that vary inlet fan load while holding other actuators fixed. We forecast downstream absolute pressure with $(L, H) = (100, 100)$. Lengths count native observations because sampling intervals are irregular.

DaISy. The DaISy heat-exchanger data [3] record liquid flow and outlet temperature every second. We use $L \in \{120, 300\}$ and $H = 1$, reflecting two or five minutes of history and the per-step changes in flow. The target is the temperature at the current input step, with history ending at the preceding step.

Shanghai Diabetes. Shanghai Diabetes [51] combines 15-minute glucose measurements with records of meals, medication, and insulin administration. We consolidate simultaneous actions while preserving food quantities and medication details. We use $L = 96$ and $H \in \{4, 8\}$, corresponding to a 24-hour history and forecast horizons of one or two hours. History excludes the glucose measurement contemporaneous with the action.

HUPA-UCM. HUPA-UCM [13] records glucose, carbohydrate intake, and insulin doses for people with type 1 diabetes. We use the raw FreeStyle historical glucose records at approximately 15-minute intervals, rather than the released 5-minute interpolated series. We set $L = 96$ (approximately 24 hours) and $H \in \{4, 8\}$. We exclude instances with long sampling gaps and participants whose pump records cannot be parsed consistently.

Table 5: History–horizon configurations and split sizes.

Dataset	(L, H)	Train	Val.	Test
M5	(112, 7)	438	147	147
M5	(112, 28)	437	142	147
MOF	(60, 1)	57	15	21
MOF	(60, 5)	57	15	21
FOMC	(60, 1)	117	39	40
FOMC	(60, 5)	116	39	40
Causal Chambers	(100, 100)	54	18	18
DaISy	(120, 1)	2208	656	776
DaISy	(300, 1)	1920	440	740
Shanghai Diabetes	(96, 4)	1687	487	678
Shanghai Diabetes	(96, 8)	1687	487	678
HUPA-UCM	(96, 4)	274	92	90
HUPA-UCM	(96, 8)	273	92	85

In terms of data ethics, all datasets were obtained from publicly available sources through their official distribution channels. Our study uses existing records and involves no new data collection from human participants.

F Covariate Conditioning Study

We compare textual action modeling with TimesFM 3.0's native support for structured covariates. Using the CovTSWM action encoding

Table 6: Comparison with covariate action conditioning.

Method	nMAE ↓	nRMSE ↓	MASE ↓
Vanilla	1.0540	1.2087	3.5023
Covariates (zero-shot)	1.1097	1.2519	3.9601
Covariates (full tuning)	1.0207	1.1340	3.4507
Covariates (head + LoRA)	1.1035	1.2438	3.9267
TSFM++	**0.9983**	**1.1116**	**3.3634**

described in Appendix A, we evaluate native covariate conditioning zero-shot and with either full-parameter or parameter-efficient fine-tuning on the same data splits. Both fine-tuning variants start from pretrained TimesFM 3.0 weights: the former updates all parameters at a learning rate of 10^{-5}, while the latter tunes the output head and rank-8 attention-layer LoRA at 10^{-6}, keeping the remaining weights frozen. Both use AdamW for up to ten epochs and early stopping after three epochs without validation improvement. Table 6 shows that zero-shot covariate inputs increase aggregate errors relative to vanilla forecasting, while full fine-tuning improves all three metrics over vanilla. Head and LoRA tuning yields only modest gains over zero-shot covariate conditioning and remains less accurate than vanilla. Neither fine-tuning strategy closes the gap to TSFM++'s lightweight textual action adaptation.

Table 7: Performance against variants that fine-tune TSFMs.

Method	nMAE ↓	nRMSE ↓	MASE ↓
TimeWorld (frozen TSFM)	**0.8997**	**1.0185**	**3.0651**
+ Full tuning	0.9082	1.0224	3.0916
+ Head and LoRA	0.9003	1.0188	3.0682

G TSFM Fine-Tuning Study

We investigate whether fine-tuning the TSFM improves forecasting after action conditioning has been learned. The first stage trains TimeWorld with a frozen TimesFM 3.0 backbone. Starting from the selected checkpoint, the second stage freezes the learned action-conditioning modules and either fine-tunes all TSFM parameters or tunes only the output head and rank-8 attention-layer LoRA. Both variants minimize factual forecasting loss at a learning rate of 10^{-6} for up to ten epochs, stopping after three epochs without validation improvement. As shown in Table 7, both strategies slightly increase all three aggregate errors, with a larger degradation under full fine-tuning. These results suggest that adapting pretrained TSFM parameters to action-dependent dynamics remains a non-trivial challenge, motivating fine-tuning strategies that better align temporal knowledge with action effects.

H Forecasting Uncertainty Analysis

To assess whether forecasting gains persist under dataset resampling, we compute 95% confidence intervals (CIs) for ΔnRMSE, defined as each method's overall nRMSE minus that of its corresponding vanilla backbone. We perform 10,000 paired bootstrap resamples of the seven datasets, jointly sampling each dataset's history–horizon configurations and preserving their weights in the overall average. The 2.5th and 97.5th percentiles define the interval bounds. In Table 8, all four intervals for TimeWorld fall below zero, supporting lower nRMSE than vanilla forecasting, whereas the TSFM++ interval for Toto includes zero. These results support the consistency of TimeWorld improvements over vanilla forecasting under dataset resampling across all four TSFMs.

Table 8: 95% confidence intervals (CIs) for ΔnRMSE (method minus vanilla). Negative values indicate improvement.

Method	TimesFM	Chronos	Toto	Moirai
TSFM++	[−0.227, −0.010]	[−0.099, −0.001]	[−0.034, 0.008]	[−0.089, −0.014]
TimeWorld	[−0.427, −0.028]	[−0.124, −0.044]	[−0.246, −0.027]	[−0.243, −0.031]

Table 9: Performance against prototype composition variant.

Method	nMAE ↓	nRMSE ↓	MASE ↓
TimeWorld	**0.8997**	**1.0185**	**3.0651**
Random reassignment	0.9215	1.0438	3.1586

I Prototype Composition Analysis

To assess the importance of input-dependent prototype composition, we compare TimeWorld with a random-reassignment variant of TimesFM 3.0 across datasets. At each forward pass, this variant replaces every future patch's mixture weights with a vector sampled from a fixed bank of weights extracted by the original model on training instances. We retrain its learnable modules from scratch using the original training protocol, with random reassignment applied throughout training, validation, and testing. Table 9 shows that input-dependent composition achieves lower errors on all three metrics, supporting the value of matching prototype mixtures to the current input.

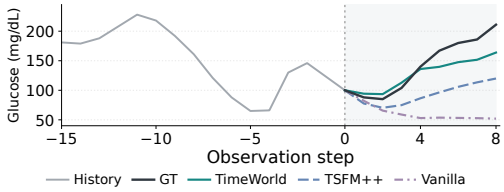

Figure 4: Forecasting example on HUPA-UCM with Toto.

J Case Study

To examine how action information helps predict a departure from recent observations, we study a Toto-based forecast on HUPA-UCM in Figure 4. The input comprises 96 glucose observations and a recorded event combining carbohydrate intake of 2.5 portions with 2 IU of rapid-acting insulin. The next eight observations briefly decline before rising from 85 to 211 mg/dL, reversing the recent

downward trajectory. This reversal makes continuation of the historical trend inadequate and provides a concrete test of forecasting under the recorded action context. Vanilla continues downward, while TSFM++ anticipates a recovery but substantially underestimates its magnitude. TimeWorld follows the initial rise more closely and reduces MAE by 59.41% relative to TSFM++, although it still underestimates the later increase. This case shows that TimeWorld better captures the observed trend reversal under the recorded action conditions.