arXiv:2611.03552v1 [math.ST] 2 Nov 2026

NONPARAMETRIC INTENSITY ESTIMATION FOR COVARIATE-DRIVEN POISSON PROCESSES USING DEEP LEARNING

Atsutomo Yara
Graduate School of Engineering Science,
the University of Osaka
a-yara@sigmath.es.osaka-u.ac.jp

Yoshikazu Terada
Graduate School of Engineering Science,
the University of Osaka
RIKEN Center for Advanced Intelligence Project (AIP)
terada.yoshikazu.es@osaka-u.ac.jp

ABSTRACT

We consider nonparametric maximum likelihood estimation of the intensity function of a Poisson process, allowing the intensity to depend on observed covariates. Such models describe how event rates vary with external conditions, including conditions under which the intensity becomes arbitrarily small or vanishes. However, risk analysis based on the Kullback–Leibler (KL) divergence typically requires the true intensity to be bounded away from zero, a restrictive assumption in applications. We develop a unified approach that avoids this assumption and derive an oracle inequality for the covariate-averaged Hellinger distance between the true intensity and the nonparametric maximum likelihood estimator (NPMLE). As an application of the oracle inequality, we establish convergence rates for NPMLEs based on deep neural networks (DNNs). Our results show that, as in nonparametric regression, the ability of DNNs to capture underlying structure leads to improved convergence rates. Specifically, DNN-based NPMLEs mitigate the curse of dimensionality when the true intensity has a compositional structure and automatically adapt to a low-dimensional Riemannian manifold structure of the event space. For intensity functions with a compositional structure, DNN-based NPMLEs achieve nearly minimax-optimal convergence rates without explicitly incorporating this structure into the network architecture. Simulations demonstrate that DNN-based estimators outperform kernel estimators in most settings, while manifold kernels remain competitive in some designs. Finally, we analyze stimulus-evoked neural spiking using data from the Allen Brain Observatory Visual Coding Neuropixels project. DNN-based NPMLEs provide better out-of-sample performance than a conditional kernel estimator, as evaluated by held-out log-likelihood.

***K*eywords** point process · intensity estimation · deep neural networks · convergence rate

1 Introduction

Poisson (point) processes are fundamental probabilistic models for random events occurring in time and space, and they have been widely applied in fields such as epidemiology, ecology, geology, seismology, and urban informatics. Poisson processes (and their generalizations) have been used, for example, to model the occurrence of earthquakes Ogata (1998), the spatial distribution of disease cases as a function of environmental exposure Diggle et al. (1997), and the frequency of taxi demand in urban areas Moreira-Matias et al. (2013). A Poisson process is characterized by its intensity function, which represents the mean occurrence rate of events at any point in time or space. Estimating this intensity function from event data is an important problem, as it provides insights into the underlying event-generation mechanism and enables the prediction of future events.

Methods for estimating the intensity function can be broadly categorized according to the type of data available. In some applications, only a single realization of a Poisson process is observed; that is, we can observe only one sequence of event times generated by the underlying process. Examples include the recorded occurrence times of earthquakes in a region over a century or data on the failure time of a system observed only once. In such cases, because only a single

sample path of the underlying process is available, estimation of the intensity function is challenging without some structural assumptions. In this situation, the intensity function is often modeled parametrically, and its parameters are estimated by maximum likelihood estimation. Since such parametric approaches are interpretable, we can gain insights into the underlying structure. Moreover, it is well known that the estimators have desirable theoretical properties, such as consistency and asymptotic normality, which can be established under suitable regularity conditions (Kutoyants (2012)). However, when modeling complex phenomena in real-world applications, the true intensity function is often too complicated to be captured by simple parametric models. In such situations, nonparametric estimation of the intensity function becomes important. Classic examples of nonparametric estimators include kernel smoothing methods Diggle (1985). Recently, more sophisticated estimators have been proposed, such as those based on reproducing kernel Hilbert spaces Flaxman et al. (2017) and estimators for Poisson processes supported on Riemannian manifolds Ward et al. (2023).

In other settings, one may observe multiple independent realizations of a Poisson process obtained under comparable experimental or environmental conditions. Typical examples include spike sequences collected from repeated experimental trials or temporal event data from multiple patients. Because these settings provide substantially more observations than the single-realization case, nonparametric intensity estimation becomes considerably more powerful. A wide range of estimators has been proposed for this data type, including those based on thresholding rules Reynaud-Bouret and Rivoirard (2010), kernel smoothing techniques Chiang et al. (2005), and hierarchical Bayesian formulations Kirichenko and van Zanten (2015), accompanied by analyses of their convergence rates as the number of observed processes $n \to \infty$.

Covariate-driven point process models provide a flexible framework for describing heterogeneity across time, space, or experimental conditions. In such models, the intensity function is allowed to depend on observed covariates, and this dependence is often interpreted through the framework of Cox processes, or doubly stochastic Poisson processes. A number of estimation methods have been developed for covariate-dependent intensity functions, including parametric estimators (Lawless, 1987; O'Sullivan, 1993), kernel-smoothing-based methods Klutchnikoff and Massiot (2016), and more recent deep-learning-based approaches (Okawa et al., 2019; Wang et al., 2020). From a theoretical perspective, Giordano et al. (2026) recently studied nonparametric Bayesian estimation for covariate-driven inhomogeneous Poisson point processes under growing-domain asymptotics. They considered the setting in which a single spatial point pattern is observed together with a spatial covariate field Z, and the intensity is modeled as $\lambda(x) = \rho(Z(x))$. Under suitable stationarity and ergodicity conditions on the covariate field, they established posterior contraction rates for several classes of priors, including Gaussian process priors, Gaussian mixture priors, and Pólya-tree-type priors. In another related direction, Gyotoku et al. (2025) investigated deep learning for point processes whose intensities depend on mixing covariate processes. Their framework covers Cox-type models, marked point processes, and multivariate point processes, and they derived an oracle inequality and convergence rate for the risk function. In this work, we study the problem of estimating the true conditional intensity function $S_0(\cdot \mid \star)$ from i.i.d. observations of the form $(Z_1, \{x_{1,1}, \ldots, x_{1,N_1}\}), \ldots, (Z_n, \{x_{n,1}, \ldots, x_{n,N_n}\})$, where each Poisson process is driven by the conditional intensity $S_0(\cdot \mid Z_i)$ given covariates Z_i.

Most existing theoretical studies of nonparametric intensity estimation focus on estimators with explicit forms, such as kernel, spline, or Bayesian estimators. However, theoretical analysis remains limited for cases where the estimator lacks an explicit representation, such as deep-learning-based intensity estimation, despite its strong predictive performance in practice. In particular, in nonparametric maximum likelihood estimation, standard generalization error analysis typically investigates the convergence of the Kullback-Leibler (KL) divergence between the true model and the estimated model. However, such an approach could fail in our setting because the KL divergence diverges when the support of the estimated intensity function does not cover that of the true intensity function. As a result, standard KL-based analysis becomes difficult unless additional assumptions are imposed on the true intensity function. For example, it is often assumed that the true intensity function is bounded away from zero (e.g., Kirichenko and van Zanten (2015)). Under this assumption, the KL divergence reduces to the squared error, yielding convergence rates of the same form as in regression problems. On the other hand, from the viewpoint of real-world applications, convergence in the KL divergence is not necessarily required. Instead, it is often sufficient to establish consistency with respect to an appropriate metric. Motivated by this, we apply the techniques of van de Geer (2000) and derive an oracle inequality to evaluate the Hellinger distance rather than the KL divergence. Using this oracle inequality, we can establish the consistency of the nonparametric maximum likelihood estimator (NPMLE) for the intensity function under realistic assumptions. Moreover, as an analog of the α-SVB condition for conditional probability estimation, which is introduced in Bos and Schmidt-Hieber (2022), we show that when the distribution of the covariates and the class of the true intensity functions satisfy suitable conditions, the approximation error can be improved even if the true intensity function is not bounded away from zero. As an important application of the oracle inequality, we derive convergence rates for the maximum likelihood estimator of the intensity function based on deep learning in the following two cases: (i) when the true intensity function has a compositional structure, and (ii) when the true intensity function is supported on a

Riemannian manifold. These results reveal that deep learning can avoid the curse of dimensionality in case (i), and can automatically adapt to the Riemannian manifold structure in case (ii), even without being given explicit information such as the Riemannian metric. Furthermore, in case (i), we show that deep learning achieves a convergence rate that is nearly minimax optimal.

The main contributions of this study can be summarized as follows:

- Theoretical analysis of the NPMLE is difficult within the standard framework of generalization error analysis, due to the unboundedness of the KL divergence, unless one imposes assumptions on the true intensity function. By adapting the density estimation argument of van de Geer (2000), we derive an oracle inequality that evaluates the performance of the NPMLE for the intensity function of a Poisson process under mild assumptions.
- As an analog of the α-SVB condition in conditional probability estimation, we introduce a corresponding condition for conditional intensity estimation and show that, under this condition, the approximation error can be improved.
- When the true intensity function has a compositional structure, we derive convergence rates for the NPMLE based on deep learning. This shows that, in this setting, deep learning can avoid the curse of dimensionality. We further show that the resulting convergence rates are nearly minimax optimal.
- When the underlying Poisson process is supported on a low-dimensional Riemannian manifold, we show that the deep-learning-based NPMLE can automatically adapt to the low-dimensional structure.

2 Preliminaries

2.1 Notation

In this section, we introduce the notation used throughout the paper. First, let $\mathbb{N}$ denote the set of natural numbers, and let $\mathbb{R}$ denote the set of real numbers. For a metric space $(\mathcal{X}, d)$, we write $\mathcal{B}(\mathcal{X})$ for the Borel σ-algebra on $\mathcal{X}$. For two measures λ and μ defined on the same measurable space, the notation $\lambda \ll \mu$ means that λ is absolutely continuous with respect to μ, that is, for every measurable set B, $\mu(B) = 0$ implies $\lambda(B) = 0$. In this case, we write $d\lambda/d\mu$ for the density of λ with respect to μ, namely, the Radon–Nikodym derivative. Following Leskelä (2024), for two σ-finite measures λ and μ on the same measurable space, we define the Hellinger distance between λ and μ by

$$H(\lambda, \mu) \coloneqq \left\{ \frac{1}{2} \int (\sqrt{f} - \sqrt{g})^2 \, d\nu \right\}^{1/2},$$

where $f \coloneqq d\lambda/d\nu$ and $g \coloneqq d\mu/d\nu$ are the densities of λ and μ, respectively, with respect to a σ-finite measure ν. For two probability measures P and Q on the same probability space, we define the Kullback–Leibler (KL) divergence by

$$\mathrm{KL}(P \parallel Q) \coloneqq \begin{cases} \displaystyle\int \log\left(\frac{dP}{dQ}\right) dP & \text{if } P \ll Q, \\ \infty & \text{otherwise.} \end{cases}$$

We also adopt the conventions $\log(0) \coloneqq -\infty$, $0 \log(0) \coloneqq 0$, $1/0 \coloneqq \infty$, and $0/0 \coloneqq 0$. For real numbers $a, b > 0$, we write $a \vee b$ for the maximum of a and b. For a vector $v = (v_1, \ldots, v_d)^\top \in \mathbb{R}^d$ and a univariate function $f : \mathbb{R} \to \mathbb{R}$, we write

$$f(v) \coloneqq (f(v_1), \ldots, f(v_d))^\top.$$

Throughout the paper, operations and inequalities involving vectors and vector-valued functions are interpreted componentwise. For a vector $v \in \mathbb{R}^d$, we define its ℓ_p norm by

$$\|v\|_p \coloneqq \left(\sum_{i=1}^{d} |v_i|^p \right)^{1/p}.$$

Moreover, for a measure μ on $\mathbb{R}^d$ and a real-valued function $f : D \subset \mathbb{R}^d \to \mathbb{R}$, we define the L^p norm by

$$\|f\|_{L^p(D,\mu)} \coloneqq \left(\int_D |f|^p \, d\mu \right)^{1/p}.$$

When the domain D or the measure μ is clear from the context, we abbreviate this notation as $\|f\|_{L^p(D)}$, $\|f\|_{L^p(\mu)}$, or simply $\|f\|_{L^p}$. We also define $\|f\|_\infty \coloneqq \sup_{x \in D} |f(x)|$. Finally, for a real number $x \in \mathbb{R}$, we write $\lceil x \rceil$ for the smallest integer greater than or equal to x, and $\lfloor x \rfloor$ for the largest integer less than or equal to x.

2.2 Setting

In this section, following Last and Penrose (2017), we introduce the general definitions of Poisson processes and Cox processes. Let $(\Omega, \mathcal{F}, \mathbb{P})$ be a probability space, and let $(\mathcal{X}, \mathcal{B})$ be a measurable space. We denote by $\mathcal{M}_{<\infty} = \mathcal{M}_{<\infty}(\mathcal{X})$ the space of all nonnegative integer-valued measures on $\mathcal{X}$ such that $\mu(B) < \infty$ for all $B \in \mathcal{B}$. We further denote by $\mathcal{M} = \mathcal{M}(\mathcal{X})$ the space of all measures that can be written as countable sums of measures in $\mathcal{M}_{<\infty}$, and by $\mathcal{N} = \mathcal{N}(\mathcal{X})$ the smallest σ-algebra on $\mathcal{M}$ such that the mapping $\mu \mapsto \mu(B)$ is measurable for every $B \in \mathcal{B}$. We then define a point process on $\mathcal{X}$ as follows.

Definition 2.1 (Point process). A point process on $\mathcal{X}$ is a random element of $(\mathcal{M}, \mathcal{N})$. That is, it is a measurable mapping $X : \Omega \to \mathcal{M}$.

Intuitively, a point process describes a random collection of points on the space $\mathcal{X}$. Next, we define a Poisson process.

Definition 2.2 (Poisson (point) process). Let λ be a σ-finite measure on $\mathcal{X}$. A Poisson (point) process on $\mathcal{X}$ with intensity measure λ is a point process X on $\mathcal{X}$ satisfying the following two properties:

(i) For every $B \in \mathcal{B}$, the random variable $X(B)$ follows the Poisson distribution with intensity parameter $\lambda(B)$.

(ii) For every $N \in \mathbb{N}$ and every collection of pairwise disjoint sets $B_1, \ldots, B_N \in \mathcal{B}$, the random variables $X(B_1), \ldots, X(B_N)$ are independent.

Hereafter, we write P_λ for the distribution of a Poisson process X with intensity measure λ. That is, for any $N \in \mathcal{N}$, we define $P_\lambda(N) \coloneqq \mathbb{P}(X^{-1}(N))$. For a Poisson process X on $\mathcal{X}$ with intensity measure λ, and for any σ-finite measure ν satisfying $\lambda \ll \nu$, the Radon–Nikodym derivative

$$f(x) \coloneqq \frac{d\lambda}{d\nu}(x)$$

is called the intensity function of X with respect to ν. Intuitively, the intensity function describes how likely points are to occur at each location in $\mathcal{X}$. Next, we define a Cox process.

Definition 2.3 (Cox process). Let ξ be a random measure on $\mathcal{X}$. A point process η is called a Cox process directed by ξ if

$$\mathbb{P}(\eta \in N \mid \xi) = P_\xi(N), \quad \mathbb{P}\text{-a.s.}, \ N \in \mathcal{N}.$$

In this case, ξ is called the directing random measure of η.

In this study, given a covariate Z taking values in $[0,1]^{d_z}$, we assume that $X(\cdot \mid Z)$ follows a Poisson process on a compact set $A \in \mathcal{B}(\mathbb{R}^{d_x})$ in the d_x-dimensional Euclidean space $\mathbb{R}^{d_x}$, and we consider the problem of estimating the intensity function of $X(\cdot \mid Z)$. That is, $X(\cdot \mid \cdot)$ is a Cox process whose directing random measure can be written as $\xi(\omega) = \xi(Z(\omega))$ for $\omega \in \Omega$, so that the randomness affects the process only through the covariate.

Hereafter, we set $d \coloneqq d_z + d_x$ and write $\lambda_0(\cdot \mid Z = z) \coloneqq \xi(z)$ for the directing random measure given the covariate value $Z = z$. Using i.i.d. observations $(Z_1, X_1), \ldots, (Z_n, X_n)$ from the Poisson process $X = X(\cdot \mid Z)$ with covariate Z and intensity measure $\lambda_0(\cdot \mid Z)$, we study the problem of estimating the conditional intensity function of X given Z, defined by

$$S_0(x \mid z) \coloneqq \frac{d\lambda_0(\cdot \mid z)}{d\nu}(x).$$

Here, ν is a measure on A, and we assume that $\lambda_0 \ll \nu$. Throughout this paper, ν is basically taken to be the d_x-dimensional Lebesgue measure, while if A is a low-dimensional manifold embedded in $\mathbb{R}^{d_x}$, then ν is taken to be the volume element on that manifold. We assume that $0 \le \lambda_0(A \mid z) < \infty$. If $\lambda_0(A \mid z) = 0$, the corresponding Poisson process is the empty point pattern almost surely. Throughout, dx denotes the d_x-dimensional Lebesgue measure. Moreover, the observation corresponding to each X_i given z_i, $i = 1, \ldots, n$ is a random set of points $\{x_{i1}, \ldots, x_{im_i}\} \subset A$, where the number of points $m_i = X_i(A \mid z_i)$ is a realization of a Poisson random variable with intensity parameter $\lambda_0(A \mid z_i)$. We consider, as the true intensity measure $\lambda_0(\cdot \mid z)$ given the covariate z, a finite measure, and we assume that the true intensity function is bounded. Furthermore, we assume that the covariate Z has a positive density on $[0,1]^{d_z}$.

Remark 2.1. The term "conditional intensity function" usually refers to the intensity function given the information available up to the previous time in the context of self-exciting stochastic processes. In this paper, however, we use the term "conditional intensity function" to mean the intensity function given the covariate.

We briefly explain the significance of modeling a point process given covariates. For example, suppose that we model the number of taxi pickups in a certain region by a Poisson process. Then it is natural to think that the corresponding intensity function depends on covariates such as temperature and precipitation. In this case, one could consider modeling

the covariates such as temperature and precipitation themselves as part of a Poisson process, that is, modeling X as a Poisson process on the Cartesian product of the two-dimensional spatial domain of taxi pickup locations and the covariate space. However, this would imply that the data-generating mechanism produces not only taxi pickup locations but also temperature and precipitation as count data, which is clearly unnatural. Therefore, in such situations, it is natural and practically useful to model the count data as a Poisson process given the covariates.

2.3 Nonparametric Maximum Likelihood Estimation

2.3.1 Formulation of Nonparametric Maximum Likelihood Estimation

In this section, we formulate the nonparametric maximum likelihood estimation of the intensity function. We begin by reviewing the expression for the likelihood ratio between two Poisson processes.

Lemma 2.1 (Theorem 1.3 of Kutoyants, 2012)**.** *Let λ and ν be finite intensity measures satisfying $\lambda \ll \nu$. Then the corresponding Poisson process distributions satisfy $P_\lambda \ll P_\nu$. Moreover, the likelihood ratio is given by*

$$\frac{dP_\lambda}{dP_\nu}(\eta) = \exp\left\{\int \log(\phi)\, d\eta - \int (\phi - 1)\, d\nu\right\},$$

where $\phi = d\lambda/d\nu$.

Now consider a Poisson process on a compact set $A \subset \mathbb{R}^{d_x}$, and apply Lemma 2.1. Then, writing the conditional intensity function given the covariate z as

$$S(x \mid z) = \frac{d\lambda(\cdot \mid z)}{d\nu}(x),$$

we obtain

$$\frac{dP_\lambda}{dP_\nu}(\eta) = \exp\left\{\int_A \log(S(x \mid z))\, d\eta - \int_A (S(x \mid z) - 1)\, d\nu(x)\right\}. \tag{1}$$

Based on the likelihood ratio (1), for a given conditional intensity function S, we define the negative log-likelihood function for an observation (X_i, z_i) by

$$\ell(X_i, z_i \mid S) \coloneqq -\int_A \log(S(x \mid z_i))\, dX_i(x \mid z_i) + \int_A S(x \mid z_i)\, d\nu(x) = -\sum_{j=1}^{N_i} \log(S(x_{ij} \mid z_i)) + \int_A S(x \mid z_i)\, d\nu(x).$$

Here, $\{x_{i1}, \ldots, x_{iN_i}\} \subset A$ denotes the observed points corresponding to X_i given the covariate z_i, and $N_i = X_i(A \mid z_i)$. Therefore, for i.i.d. observations $(X_1, z_1), \ldots, (X_n, z_n)$ from the covariate Z and the Poisson process $X(\cdot \mid Z)$ given Z, the negative log-likelihood function normalized by $1/n$ is given by

$$\begin{aligned}
\ell_n(X_1, \ldots, X_n, z_1, \ldots, z_n \mid S) &\coloneqq \frac{1}{n}\sum_{i=1}^{n} \ell(X_i, z_i \mid S) \\
&= \frac{1}{n}\sum_{i=1}^{n} \left\{-\sum_{j=1}^{N_i} \log(S(x_{ij} \mid z_i)) + \int_A S(x \mid z_i)\, d\nu(x)\right\}. \qquad (2) \\
\ell_n(S) &\coloneqq \ell_n(X_1, \ldots, X_n, z_1, \ldots, z_n \mid S).
\end{aligned}$$

In nonparametric maximum likelihood estimation, the estimator is defined as a function that minimizes the negative log-likelihood (2) over a function space (hypothesis space) $\mathcal{F}_n$, which is not necessarily parameterized by a finite-dimensional parameter. The resulting estimator is called the nonparametric maximum likelihood estimator (NPMLE). That is, the NPMLE $\hat{S}_n(x \mid z)$ of the conditional intensity function $S_0(x \mid z)$ of the Poisson process $X(\cdot \mid z)$ given the covariate z is defined by

$$\hat{S}_n \in \operatorname*{arg\,min}_{S \in \mathcal{F}_n} \ell_n(S). \tag{3}$$

2.3.2 Difficulty of Evaluating the Excess Risk

In the context of standard generalization error analysis, the performance of the NPMLE is typically evaluated by the excess risk

$$R(\hat{S}_n) \coloneqq \mathbb{E}[\ell(\hat{S}_n) \mid (Z_1, X_1), \ldots, (Z_n, X_n)] - \inf_{S \in \mathcal{F}_0} \mathbb{E}[\ell(S)],$$

where $\mathcal{F}_0$ denotes the set of all nonnegative bounded measurable functions $f : A \times [0,1]^{d_z} \to [0,\infty)$. The excess risk depends on the estimator $\hat{S}_n$. Thus, it is a random variable. When no fonfusion is likely to arise from the context, we abbreviate $\mathbb{E}[\ell(\hat{S}_n) \mid (Z_1, X_1), \ldots, (Z_n, X_n)]$ as $\mathbb{E}[\ell(\hat{S}_n)]$. The following lemma gives an explicit expression for the excess risk. Here, we write $\hat{\lambda}_n(\cdot \mid z)$ for the intensity measure corresponding to the estimator $\hat{S}_n(\cdot \mid z)$.

Lemma 2.2. *For the excess risk $R(\hat{S}_n)$, we have*

$$\begin{aligned} R(\hat{S}_n) &= \mathbb{E}_Z\left[\mathrm{KL}\Big(P_{\lambda_0(\cdot|Z)} \parallel P_{\hat{\lambda}_n(\cdot|Z)}\Big)\right] \\ &= \mathbb{E}_Z\left[\int_A S_0(x \mid Z) \log\left(\frac{S_0(x \mid Z)}{\hat{S}_n(x \mid Z)}\right) d\nu(x) + \int_A \big(\hat{S}_n(x \mid Z) - S_0(x \mid Z)\big)\, d\nu(x)\right], \end{aligned} \tag{4}$$

where $\mathbb{E}_Z$ denotes expectation with respect to Z.

Equation (4) diverges unless $\lambda_0(\cdot \mid z) \ll \hat{\lambda}_n(\cdot \mid z)$ on A for every z. Therefore, in order to evaluate the excess risk directly, one would need to assume, for example, that there exists a constant $c > 0$ such that $\hat{S}_n > c$ for all estimators $\hat{S}_n$. However, this is a rather strong and unrealistic assumption when discussing the convergence of the estimator to the true intensity function. Moreover, from the viewpoint of real-world applications, convergence of the excess risk is not necessarily required; it is often sufficient to establish consistency with respect to an appropriate metric. Therefore, by adapting the techniques of van de Geer (2000), we evaluate a lower bound of the excess risk $R(\hat{S}_n)$ instead of the excess risk itself, and establish consistency of that lower bound under realistic assumptions. The following lemma shows that the Hellinger distance between intensity measures appears as a lower bound of the excess risk. For notational simplicity, we also write the Hellinger distance between intensity measures in terms of the corresponding intensity functions, for example,

$$H(S_1, S_2) \coloneqq H(\lambda_1, \lambda_2).$$

We also define the expectation, with respect to the covariate, of the squared Hellinger distance between conditional intensity functions by

$$\widetilde{H}^2(S_1, S_2) \coloneqq \mathbb{E}_Z\big[H^2(S_1(\cdot \mid Z), S_2(\cdot \mid Z))\big].$$

Lemma 2.3. *For the excess risk $R(\hat{S}_n)$, we have*

$$R(\hat{S}_n) \geq 2\widetilde{H}^2(\hat{S}_n, S_0).$$

Remark 2.2. Since $x + 1 \leq e^x$ for any $x \in \mathbb{R}$, we have

$$\widetilde{H}^2(\hat{S}_n, S_0) \geq \mathbb{E}_Z\big[1 - e^{-H^2(\hat{S}_n, S_0)}\big].$$

The quantity inside the expectation on the right-hand side is the Hellinger distance between the Poisson processes $P_{\hat{\lambda}_n}$ and P_{λ_0} (see Leskelä, 2024, Corollary 5.3). Thus, convergence in the Hellinger distance between intensity measures is stronger than convergence in the Hellinger distance between the corresponding Poisson processes.

3 Oracle Inequality

In this section, by adapting the argument of van de Geer (2000), we establish an oracle inequality for evaluating the Hellinger distance between intensity measures. We impose the following assumption on the true intensity function and the estimation model $\mathcal{F}_n$.

Assumption 3.1. There exist a sequence of elements $S_n^* \in \mathcal{F}_n$ and a constant $c_0 \geq 0$ such that

$$\forall n \in \mathbb{N},\ \forall x \in A,\ \forall z \in [0,1]^{d_z};\ \frac{S_0(x \mid z)}{S_n^*(x \mid z)} \leq c_0^2.$$

Remark 3.1. Assumption 3.1 means that, if we suitably choose the hypothesis space $\mathcal{F}_n$, we do not need to impose additional conditions on the true conditional intensity function S_0. We later show that the estimation model $\mathcal{F}_n$ is taken to be a suitable class of deep neural networks; this assumption is automatically satisfied. See the proof of Theorem 4.2 for details.

For a measure Q and a measurable function f, we use the notation $Qf \coloneqq \int f\, dQ$. We also write the population distribution as $P \coloneqq P_{\lambda_0|Z} \otimes P_Z$, and denote by P_n the empirical distribution corresponding to the observations $(X_1, Z_1), \ldots, (X_n, Z_n)$. We first show, using the property of the maximum likelihood estimator, that the Hellinger distance between intensity measures can be bounded by a certain empirical process. The following lemma is a key to deriving the oracle inequality.

Lemma 3.1 (Basic Inequality)**.** *Suppose that Assumption 3.1 holds. Then, for any $n \in \mathbb{N}$ and any NPMLE $\hat{S}_n$, we have*

$$\begin{aligned}\widetilde{H}^2\left(\frac{\hat{S}_n + S_n^*}{2}, S_n^*\right) \leq (P_n - P)\left\{\int_{A_n^*(Z)} \frac{1}{2}\log\left(\frac{\hat{S}_n(x \mid Z) + S_n^*(x \mid Z)}{2S_n^*(x \mid Z)}\right) dX(x \mid Z)\right. \\ + \frac{1}{2}\int_A \left(S_n^*(x \mid Z) - \frac{\hat{S}_n(x \mid Z) + S_n^*(x \mid Z)}{2}\right) d\nu(x) \\ \left. + 2(c_0+1)\widetilde{H}\left(\frac{\hat{S}_n + S_n^*}{2}, S_n^*\right)\widetilde{H}\left(S_n^*, S_0\right),\right.\end{aligned}$$

where $A_n^*(z) \coloneqq A \cap \{x \in A : S_n^*(x \mid z) > 0\}$ for $z \in [0,1]^{d_z}$.

We denote the index sets of the empirical processes by

$$\begin{aligned}\mathcal{G}_n &\coloneqq \left\{g_S = \int_{A_n^*(Z)} \frac{1}{2}\log\left(\frac{S(x \mid Z) + S_n^*(x \mid Z)}{2S_n^*(x \mid Z)}\right) dX(x \mid Z) : S \in \mathcal{F}_n\right\}, \\ \mathcal{G}_n' &\coloneqq \left\{g_S' = \int_A \frac{1}{2}\left(S_n^*(x \mid Z) - \frac{S(x \mid Z) + S_n^*(x \mid Z)}{2}\right) d\nu(x) : S \in \mathcal{F}_n\right\}.\end{aligned}$$

The oracle inequality is derived by bounding the empirical process on the right-hand side of Lemma 3.1 using a Bernstein-type concentration inequality (Theorem B.2). When applying Theorem B.2, a key difficulty, compared with density estimation in van de Geer (2000), is that the function class $\mathcal{G}_n$ indexing the empirical process is not uniformly bounded from below. In density estimation, the corresponding term is like

$$\frac{1}{2}\log\left(\frac{S(x \mid z) + S_n^*(x \mid z)}{2S_n^*(x \mid z)}\right),$$

which is uniformly bounded from below. In that setting, this fact, together with Lemma B.3, is used to show that the Bernstein condition (14) holds. In contrast, in the estimation of the intensity function of a Poisson process, the lower bound of $\mathcal{G}_n$ involves a random variable following a Poisson distribution. Since such a quantity is generally not bounded from below, Lemma B.3 cannot be applied directly. To overcome this difficulty, we exploit the fact that, given the number of points of the Poisson process X in A_n^*, namely $X(A_n^*) = N$, the individual points are independently distributed according to the distribution with density given by the normalized intensity function; see Proposition 3.8 of Last and Penrose (2017). Using this property, we show that the Bernstein condition (14) is satisfied.

Before presenting the oracle inequality, we introduce several definitions and pieces of notation. For a given hypothesis space $\mathcal{F}_n$ and $\delta > 0$, define

$$\bar{\mathcal{F}}_n^{1/2}(\delta) \coloneqq \left\{\sqrt{\frac{S + S_n^*}{2}} : \widetilde{H}^2\left(\frac{S + S_n^*}{2}, S_n^*\right) \leq \delta^2,\ S \in \mathcal{F}_n\right\}.$$

Next, for $p \geq 1$ and a measure Q, we define the bracketing number of a function class $\mathcal{F}$ with respect to the $L^p(Q)$ metric as follows.

Definition 3.1 (Bracketing Number for the $L^p(Q)$ Metric). Let $N_{p,B}(\delta, \mathcal{F}, Q)$ denote the smallest integer N such that there exist pairs of functions $\{(f_j^L, f_j^U)\}_{j=1}^N$ satisfying $\|f_j^U - f_j^L\|_{L^p(Q)} \leq \delta$ for all $j \in \{1, \ldots, N\}$, and such that for every $f \in \mathcal{F}$, there exists an index $j = j(f) \in \{1, \ldots, N\}$ for which

$$f_j^L \leq f \leq f_j^U.$$

The quantity $N_{p,B}(\delta, \mathcal{F}, Q)$ is called the δ-bracketing number of $\mathcal{F}$.

Using the two definitions above, we define the bracketing entropy integral associated with the hypothesis space $\mathcal{F}_n$ and S_n^* in Assumption 3.1 by

$$J_B(\delta, \bar{\mathcal{F}}_n^{1/2}(\delta)) = \int_{\delta^2/(2^{16}C_\rho)}^{\delta} \sqrt{\log N_{2,B}(u, \bar{\mathcal{F}}_n^{1/2}(\delta), \nu \otimes P_Z)}\, du \vee \delta, \tag{5}$$

where an integral with lower limit greater than its upper limit is set to zero, and the definition of C_ρ can be found in the statement of Lemma B.5. The quantity $J_B(\delta, \bar{\mathcal{F}}_n^{1/2}(\delta))$ measures the local complexity of the hypothesis space in a neighborhood of S_n^*.

We are now ready to state the oracle inequality.

Theorem 3.2 (Oracle Inequality)**.** *Consider the problem of estimating the conditional intensity function $S_0(\cdot \mid Z)$ of a Poisson process $X(\cdot \mid Z)$. Let $\hat{S}_n$ be the NPMLE defined in* (3)*, and suppose that Assumption 3.1 holds. Assume further that $\sup_{f\in\mathcal{F}_n}\|f\|_\infty \le F$ for some $F > 0$. For $\delta > 0$, let $\Psi(\delta)$ satisfy $\Psi(\delta) \ge J_B(\delta, \bar{\mathcal{F}}_n^{1/2}(\delta))$ and assume that $\Psi(\delta)/\delta^2$ is a nonincreasing function of δ. Then there exists a constant $c \ge 1$ depending only on $c_0, F, \nu(A)$ such that, for any δ_n satisfying*

$$\sqrt{n}\,\delta_n^2 \ge c\Psi(\delta_n), \quad n \in \mathbb{N}, \tag{6}$$

we have for $\delta \ge \delta_n$,

$$\mathbb{P}\left(\widetilde{H}^2(\hat{S}_n, S_0) > 1154(c_0+1)^2\big(\delta^2 + \widetilde{H}^2(S_n^*, S_0)\big)\right) \le c\exp\left(-\frac{n\delta^2}{c^2}\right). \tag{7}$$

Furthermore, by integrating both sides of (7), we obtain the following corollary concerning the expectation with respect to the observed data.

Corollary 3.3. *Under the same conditions as in Theorem 3.2, for every $\delta \ge \delta_n$, the following inequality holds:*

$$\mathbb{E}\left[\widetilde{H}^2(\hat{S}_n, S_0)\right] \le 1154(c_0+1)^2\left(\delta^2 + \widetilde{H}^2(S_n^*, S_0) + \frac{c^3}{n}\right).$$

Here, the expectation is taken with respect to the observed data $(X_1, z_1), \ldots, (X_n, z_n)$.

Remark 3.2. The oracle inequalities in Theorem 3.2 and Corollary 3.3 hold for any reference intensity $S_n^* \in \mathcal{F}_n$ satisfying Assumption 3.1. This condition should be viewed as a uniform lower-domination requirement: S_n^* must be positive wherever the true conditional intensity is positive and must dominate S_0 up to a constant factor. For instance, if the model class contains a conditional intensity function that is uniformly bounded away from zero, then the assumption is satisfied because S_0 is bounded. Nevertheless, the choice of S_n^* is not merely technical. The resulting oracle bound also contains the approximation error $\widetilde{H}^2(S_n^*, S_0)$, and a poorly chosen reference intensity may therefore lead to a loose bound. To derive sharp convergence rates, it is necessary to choose S_n^* so that the domination condition and the approximation requirement are both satisfied. In the DNN analysis below, we construct such a sequence explicitly by combining a neural network approximation of S_0 with an appropriate lower clipping. This construction ensures Assumption 3.1 and yields a sufficiently small approximation error, which can then be substituted into Corollary 3.3.

Theorem 3.2 and Corollary 3.3 imply that, once we can control the bracketing number representing the complexity of the hypothesis space $\mathcal{F}_n$ (or, under suitable conditions, a covering number that provides an upper bound on the bracketing number) together with the approximation error $\widetilde{H}^2(S_n^*, S_0)$, we can derive convergence rates for the NPMLE with respect to $\widetilde{H}^2$.

4 Maximum Likelihood Estimation of Intensity Functions Using Deep Learning

In this section, we consider maximum likelihood estimation for conditional intensity functions using deep learning. That is, we take the hypothesis space $\mathcal{F}_n$ to be a class of deep neural networks (DNNs). By applying Corollary 3.3, we derive convergence rates for the DNN-based NPMLE. We first give a mathematical definition of the deep neural networks in Section 4.1, and then present the main results on convergence rates in Section 4.2.

4.1 Deep Neural Networks

In this section, we give a mathematical definition of DNNs. We use the ReLU (rectified linear unit) activation function $\sigma : \mathbb{R} \to \mathbb{R}, \quad x \mapsto x \vee 0$. For a bias vector $v = (v_1, \ldots, v_r)^\top \in \mathbb{R}^r$, we define the activation function with bias by

$$\sigma_v\begin{pmatrix} y_1 \\ \vdots \\ y_r \end{pmatrix} = \begin{pmatrix} \sigma(y_1 - v_1) \\ \vdots \\ \sigma(y_r - v_r) \end{pmatrix}.$$

Given a final-layer activation function ψ, we define a neural network for modeling an intensity function with depth $L \in \mathbb{N}$ and width $r \in \mathbb{N}$ at each hidden layer as follows:

$$f : \mathbb{R}^d \to \mathbb{R}, \quad x \mapsto f(x) = \psi W_L \sigma_{v_L} W_{L-1}\sigma_{v_{L-1}} \cdots W_1 \sigma_{v_1} W_0 x. \tag{8}$$

Here, $W_0 \in \mathbb{R}^{r\times d}$, $W_L \in \mathbb{R}^{1\times r}$, and $W_i \in \mathbb{R}^{r\times r}$, $i = 1, \ldots, L-1$, are weight matrices, and $v_i \in \mathbb{R}^r$, $i = 1, \ldots, L$, are bias vectors. Although the widths may vary across layers in general, we use a common width r for all hidden

layers for simplicity. We assume that all network parameters are bounded in absolute value by $B > 0$. Specifically, we consider the following class of DNNs:

$$\mathcal{F}_\psi(L, r, B, F) \coloneqq \left\{ f \text{ of the form (8)} : \max_{j=0,\dots,L} (\|W_j\|_\infty \vee \|v_j\|_\infty) \le B, \|f\|_\infty \le F \right\},$$

where we set $v_0 = 0$. Since intensity functions are nonnegative, we use the ReLU activation function in the final layer, that is, $\psi = \sigma$. We also define the class of DNNs with a restricted number of nonzero parameters by

$$\mathcal{F}_\psi(L, r, B, F, s) \coloneqq \left\{ f \in \mathcal{F}_\psi(L, r, B, F) : \sum_{j=0}^{L} (\|W_j\|_0 + \|v_j\|_0) \le s \right\}.$$

Here, $\|W_j\|_0$ and $\|v_j\|_0$ denote the number of nonzero components of W_j and v_j, respectively.

4.2 Convergence Rates

In this section, we take the hypothesis space $\mathcal{F}_n$ to be the class of DNNs defined in the previous subsection, and apply Corollary 3.3 to derive convergence rates for the NPMLE of the conditional intensity function based on DNNs. To discuss convergence rates, we need to impose certain restrictions on the function class to which the true conditional intensity function S_0 belongs. We assume several structural conditions on the true conditional intensity function.

We first introduce an assumption common to all settings. An interesting phenomenon occurs when the true intensity function takes values close to zero. To quantify this behavior, we introduce below a quantity that measures the size of the region where the intensity function is small. This condition is a modification, for conditional intensity functions, of the α-SVB condition introduced for conditional probabilities in Definition 3.1 of Bos and Schmidt-Hieber (2022).

Definition 4.1 (Small value bound for conditional intensity functions). Let $\alpha \ge 0$, and let $\mathcal{H}$ be a class of conditional intensity functions. We say that $\mathcal{H}$ is α-small value bounded, or α-SVB, with respect to the distribution of the covariate Z, if there exists a constant $C > 0$ such that, for every $S \in \mathcal{H}$,

$$\int_{[0,1]^{d_z}} \int_A \mathbb{1}\left(S(x \mid z) \le t\right) \, d\nu(x) dP_Z(z) \le Ct^\alpha, \quad \text{for all } t \in (0, \infty).$$

By Tonelli's theorem, the above condition can be rewritten as

$$\int_A \mathbb{P}_Z\left(S(x \mid Z) \le t\right) \, d\nu(x) \le Ct^\alpha, \quad \text{for all } t \in (0, \infty).$$

Therefore, this condition always holds with $\alpha = 0$ and $C = \nu(A)$. Thus, the condition is automatically satisfied with $\alpha = 0$. For $\alpha > 0$, however, it imposes a restriction on the size of regions where the intensity is close to zero. On the other hand, if the set of points $x \in A$ satisfying $\mathbb{P}_Z(S(x \mid Z) = 0) > 0$ has positive ν-measure, then this condition cannot hold for any $\alpha > 0$. If every function in $\mathcal{H}$ is bounded below by a constant $B_0 > 0$ ($\mathbb{P}_Z - a.s.$), then the condition holds for any $\alpha \ge 0$ by taking $C = \nu(A) B_0^{-\alpha}$. Thus, the α-SVB condition can be interpreted as quantifying, through α, how sharply the true conditional intensity function decays near zero after averaging over the covariates. We illustrate this interpretation with a simple example.

Let $d_z = d_x = 1$, and let ν be the Lebesgue measure on $A = [0, 1]$. Let Z be any random variable supported on $[0, 1]$. For $\alpha > 0$, define

$$S(x \mid z) = z + x^{1/\alpha}.$$

Then, for any $t \in (0, 2)$, $\{x \in A : S(x \mid z) \le t\} = \{x \in [0, 1] : z + x^{1/\alpha} \le t\}$. If $z > t$, this set is empty, while if $0 \le z \le t$, it is equal to $[0, (t - z)^\alpha]$. Hence,

$$\nu\{x \in A : S(x \mid z) \le t\} = \{(t - z) \vee 0\}^\alpha,$$

Therefore,

$$\int_{[0,1]^{d_z}} \int_A \mathbb{1}\{S(x \mid z) \le t\} \, d\nu(x) dP_Z(z) = \mathbb{E}_Z[((t - Z) \vee 0)^\alpha] \le t^\alpha.$$

Thus, the α-SVB condition holds with $C = 1$.

The following Theorem 4.1 shows that a larger value of α improves the approximation error. The proof is almost identical to that of Theorem 3.2 in Bos and Schmidt-Hieber (2022), but we provide the details in Appendix C for completeness.

Theorem 4.1. *Suppose that the function class to which the true conditional intensity function S_0 belongs is α-SVB. Then, for any approximating function S satisfying*

$$\|S - S_0\|_\infty \le \frac{C_1}{M}, \qquad \inf_{z\in[0,1]^{d_z}} \inf_{x\in A} S(x \mid z) \ge \frac{1}{M}$$

for some constant $C_1 \ge 0$ and some $M \ge C_1 + 1$, we have

$$\widetilde{H}^2(S, S_0) \le (C+1)\frac{(C_1+1)^{2+(\alpha\wedge 1)}}{M^{1+(\alpha\wedge 1)}}\left(\nu(A) + \frac{\mathbb{1}(\alpha<1)}{1-\alpha} + \log(M)\right).$$

Here, C is the constant in the SVB condition.

4.2.1 Convergence Rates for Composition Structured Functions

In this section, we discuss convergence rates when the true conditional intensity function S_0 belongs to a class of functions called composition-structured functions. When the true function has such a compositional structure, various studies show that deep learning can avoid the curse of dimensionality in regression and conditional probability estimation (e.g., Schmidt-Hieber, 2020; Bos and Schmidt-Hieber, 2022; Yara and Terada, 2026).

We now define composition-structured functions. Suppose that the true conditional intensity function S_0 can be represented as the composition of $q+1$ functions:

$$S_0 = g_q \circ g_{q-1} \circ \cdots \circ g_1 \circ g_0.$$

Here, for $i = 0, \dots, q$, let

$$g_i : [a_i, b_i]^{d^{(i)}} \to [a_{i+1}, b_{i+1}]^{d^{(i+1)}}, \qquad a_i, b_i \in \mathbb{R}, \quad d^{(i)} \in \mathbb{N},$$

where $d^{(0)} = d$ and $d^{(q+1)} = 1$. We write each component of g_i as

$$g_i = (g_{i1}, \dots, g_{i,d^{(i+1)}})^\top,$$

and denote by t_i the maximum number of variables on which each component function g_{ij}, $j = 1, \dots, d^{(i+1)}$, depends. For example, consider

$$f(x_1, x_2, x_3) = g_{11}\big(g_{01}(x_1, x_2), g_{02}(x_2, x_3)\big) = \exp\{-(x_1^2 + x_2^2)\} + \exp\{-(x_2^2 + x_3^2)\}.$$

This function has a compositional representation with $q = 1$, $g_{01}(x_1, x_2) = \exp\{-(x_1^2 + x_2^2)\}$, $g_{02}(x_2, x_3) = \exp\{-(x_2^2 + x_3^2)\}$, and $g_{11}(u_1, u_2) = u_1 + u_2$. In this case, we have $d^{(0)} = 3$, $t_0 = 2$, $d^{(1)} = t_1 = 2$, $d^{(2)} = 1$.

We further assume that each component function g_{ij} belongs to a Hölder class. For a domain D, a radius $K > 0$, and a smoothness parameter $\beta > 0$, the Hölder class is defined by

$$\mathcal{C}^\beta(D, K) = \left\{ f : D \to \mathbb{R} : \sum_{\gamma:\|\gamma\|_1<\beta} \|\partial^\gamma f\|_\infty + \sum_{\gamma:\|\gamma\|_1=\lfloor\beta\rfloor} \sup_{x,y\in D,\, x\ne y} \frac{|\partial^\gamma f(x) - \partial^\gamma f(y)|}{\|x-y\|_\infty^{\beta-\lfloor\beta\rfloor}} \le K \right\}.$$

Here, we use multi-index notation. That is, for $\gamma = (\gamma_1, \dots, \gamma_r) \in \mathbb{N}^r$, we write $\partial^\alpha = \partial^{\alpha_1} \cdots \partial^{\alpha_r}$.

Thus, the function class to which the true intensity function belongs is defined by

$$\mathcal{G}(q, t, \beta, K) \coloneqq \big\{ S = g_q \circ \cdots \circ g_0 : S \ge 0,\, g_i = (g_{ij})_j : [a_i, b_i]^{d^{(i)}} \to [a_{i+1}, b_{i+1}]^{d^{(i+1)}},$$
$$g_{ij} \in \mathcal{C}^{\beta_i}([a_i, b_i]^{t_i}, K),\ \text{for some } |a_i|, |b_i| \le K \big\},$$

where $t = (t_0, \cdots, t_q)$, and $\beta = (\beta_0, \dots, \beta_q)$.

As an illustrative example, consider the conditional intensity function of a covariate-driven Poisson process describing taxi pickup requests in an urban area. The covariates may include daily weather conditions such as precipitation and temperature. It is natural to expect the intensity to increase under weather conditions that make walking less comfortable, such as heavy rain, extreme heat, or extreme cold. A typical spatial pattern of taxi demand may consist of several peaks centered around major facilities, such as department stores, railway stations, or taxi stands. These peaks can be modeled using Gaussian-density components whose locations and shapes reflect the city's spatial structure. The effect of covariates can then be incorporated by allowing the height, scale, or anisotropy of each Gaussian-type component to vary with them. For instance, under rainy conditions, the peaks around major facilities may become higher and more

concentrated, whereas under mild weather conditions, they may become lower and more diffuse. A conditional intensity function of this type can be represented schematically as

$$S_0(x \mid z) = \sum_{l=1}^{L} a_l(z) \exp\left(-\|B_l(z)(x-\mu_l)\|_2^2\right),$$

where x denotes the spatial location, z denotes the covariates, μ_l represents the location of the l-th major facility, and $a_l(z)$ and $B_l(z)$ control the covariate-dependent height and shape of the corresponding peak. This is a Gaussian-mixture-type conditional intensity function whose mixture components are modulated by covariates. Such a model naturally has a compositional structure. Indeed, the inner functions compute low-dimensional quantities such as covariate-dependent weights, transformed distances from major facilities, and Gaussian-type local peaks, while the outer function aggregates these components by summation. Since the summation map is infinitely differentiable, the effective smoothness of this outer component is high. Consequently, this representation can lead to faster convergence rates than those suggested by the ambient dimension alone, reflecting the advantage of deep learning for composition-structured functions.

For fixed P_Z, $\alpha \geq 0$, and $C > 0$, let $\mathcal{G}_\alpha(q, t, \beta, K, C)$ denote the subclass of $\mathcal{G}(q, t, \beta, K)$ consisting of all functions S satisfying

$$\int_{[0,1]^{d_z}} \int_A \mathbb{1}\{S(x \mid z) \leq u\}\, d\nu(x)\, dP_Z(z) \leq Cu^\alpha \qquad \text{for every } u > 0.$$

We assume that S_0 belongs to this subclass and derive convergence rates with respect to $\widetilde{H}^2$ for the DNN-based NPMLE.

Theorem 4.2 (Convergence Rates for Composition Structured Functions Using Deep Learning)**.** *Consider the problem of estimating the conditional intensity function S_0 of the Poisson process $X \mid Z$, where the covariate satisfies $Z \sim P_Z$. Suppose that $S_0 \in \mathcal{G}_\alpha(q, t, \beta, K, C)$. Let $\hat{S}_n$ be the NPMLE defined in* (3)*, and let the hypothesis space $\mathcal{F}_n$ be the class of DNNs $\mathcal{F}_\sigma(L, r, B, F)$ with $F \geq K$ satisfying*

$$\text{(i) } L \asymp \log(n), \qquad \text{(ii) } r \asymp \sqrt{n\phi_n}, \qquad \text{(iii) } B \asymp \max_{i=0,\ldots,q} n^{\frac{2(\beta_i+1)}{(1+(\alpha\wedge 1))\beta_i^* + t_i}}.$$

Then there exists a constant $C' > 0$, depending only on $A, \nu, q, t, \beta, K, F, \alpha, C, c_0$, and $d^{(0)}, \ldots, d^{(q)}$, such that

$$\mathbb{E}\left[\widetilde{H}^2(\hat{S}_n, S_0)\right] \leq C'\phi_n \log(n)^3.$$

Here,

$$\phi_n \coloneqq \max_{i=0,\ldots,q} n^{-\frac{(1+(\alpha\wedge 1))\beta_i^*}{(1+(\alpha\wedge 1))\beta_i^* + t_i}}, \qquad \beta_i^* \coloneqq \beta_i \prod_{l=i+1}^{q} (\beta_l \wedge 1).$$

Remark 4.1. In this theorem (also in Theorem 4.4), we assume that the NPMLE defined in (3) exists and is computable. More precisely, this estimator can be interpreted as a regularized estimator because the magnitudes of the parameters are bounded by B. Since we can easily extend this theorem to the case where we cannot compute the exact NPMLE (see Schmidt-Hieber (2020), for example), we state the ideal case in this theorem for simplicity.

Remark 4.2. The network budgets in Theorems 4.2 and 4.4 are calibrated using unknown smoothness and structural parameters of the true intensity, as well as the SVB exponent. The resulting bounds characterize the performance of these calibrated network classes; they do not, by themselves, establish adaptation to the unknown parameters. In the manifold setting, the dependence on intrinsic dimension enters the choice of network budgets, while the networks receive ambient coordinates and do not require intrinsic coordinates or chart maps as inputs. A similar issue is discussed in Kohler and Langer (2021), which suggests selecting network depth and width via sample splitting. Such a selection procedure is a natural extension of our approach. Establishing that a data-driven choice retains the stated Hellinger rates, however, requires an additional model-selection argument and is not asserted here.

The convergence rate ϕ_n in Theorem 4.2 does not depend on the input dimension d directly. Instead, it depends on the intrinsic dimensions t_i of the component functions g_i, $i = 0, \ldots, q$, appearing in the compositional representation. Therefore, if

$$i^* \in \operatorname*{arg\,max}_{i=0,\ldots,q} \left\{-\frac{\beta_i^*}{\beta_i^* + t_i}\right\}$$

satisfies $t_{i^*} \ll d$, then deep learning can mitigate the curse of dimensionality.

4.2.2 Minimax Optimality for Composition Structured Functions

In this section, we show that the convergence rate for deep learning obtained in the previous section is minimax optimal up to logarithmic factors when $\alpha = 0$. The following theorem gives this result.

Theorem 4.3 (Minimax Optimality)**.** *Let $A = [0,1]^{d_x}$ and let ν be d_x-dimensional Lebesgue measure. Consider the unrestricted small-value class ($\alpha = 0$) of conditional intensity functions. Suppose that the covariate Z has a density with respect to the Lebesgue measure, and that this density is bounded above and below by positive constants. Then, for any nonnegative integer q, any smoothness vector β, any dimensions $d^{(0)}, \ldots, d^{(q)}$ and t satisfying*

$$t_i \le \min\{d^{(0)}, \ldots, d^{(i-1)}\}, \qquad i = 1, \ldots, q,$$

and any sufficiently large $K > 0$, there exists a positive constant c such that, for all sufficient large n,

$$\inf_{\hat{S}_n} \sup_{S_0 \in \mathcal{G}(q,t,\beta,K)} \mathbb{E}\left[\widetilde{H}^2(\hat{S}_n, S_0)\right] \ge c \max_{i=0,\ldots,q} n^{-\frac{\beta_i^*}{\beta_i^* + t_i}}.$$

Here, the infimum is taken over all estimators, and the expectation is taken with respect to the distribution of the data determined by S_0 and P_Z.

The lower bound differs from the classical $n^{-2\beta/(2\beta+d)}$ rate for densities or intensities bounded away from zero. This is because the present class allows intensities to approach or equal zero. This leads to the slower rate.

4.2.3 Convergence Rates for Poisson Processes on Riemannian Manifolds

In this section, we consider the case where the true Poisson process is supported on a Riemannian manifold. That is, the region in which events occur is a low-dimensional Riemannian manifold embedded in the ambient space where the events are observed. We show that, in this setting, the convergence rate of the NPMLE of the intensity function based on deep learning does not depend on the dimension d_x of the original event-observation space, but on the intrinsic dimension $d_{\mathcal{M}}$ of the low-dimensional Riemannian manifold embedded in $\mathbb{R}^{d_x}$.

As a motivating example, point processes supported on Riemannian manifolds arise in cellular biology and microbiology, where super-resolution microscopy records the spatial arrangement of proteins and other molecules on cellular membranes. Since such membranes are curved surfaces embedded in an ambient Euclidean space, the support of the point process is more naturally modeled as a low-dimensional Riemannian manifold than as the full ambient Euclidean space (Ward et al., 2023). Related theoretical work on point processes in manifolds also emphasizes that, when data are embedded in a high-dimensional Euclidean space but lie on a lower-dimensional manifold, intrinsic dimension and manifold learning become central objects of analysis (Penrose and Yukich, 2013). Spherical point processes provide another concrete example, where the support is the sphere, a canonical Riemannian manifold, and the goal is to estimate the intensity function of a non-homogeneous Poisson process on that manifold (Ng and Zammit-Mangion, 2022).

Before stating the convergence rates, we introduce several definitions. Let $\mathcal{M} \subset \mathbb{R}^{d_x}$ be a fixed compact C^∞ embedded manifold without boundary, of dimension $d_{\mathcal{M}}$, and let ν be its Riemannian volume measure, with $0 < \nu(\mathcal{M}) < \infty$. The likelihood uses this known reference measure.

We fix a finite atlas $\{(V_i, P_i)\}_{i=1}^J$ of $\mathcal{M}$ in which each P_i is the restriction of an affine ambient map and $P_i(V_i) = \Omega_i$ is an open ball in $\mathbb{R}^{d_{\mathcal{M}}}$. Such charts are obtained by restricting projections onto tangent spaces. Fix also a smooth partition of unity $\{\rho_i\}_{i=1}^J$ with $\mathrm{supp}(\rho_i) \subset V_i$. All geometric choices remain fixed as the intensity and the sample size vary.
Definition 4.2 (Hölder class for a fixed atlas). For $i = 1, \ldots, J$, write

$$D_i = \Omega_i \times [0,1]^{d_z}, \qquad S_i(u,z) = S(P_i^{-1}(u), z).$$

For a radius $K > 0$, output bound $F > 0$, and a smoothness parameter $\beta > 0$, the Hölder class on $\mathcal{M} \times [0,1]^{d_z}$ is defined by

$$\mathcal{H}^\beta_{\mathcal{M}}(K,F) \coloneqq \left\{S : \mathcal{M} \times [0,1]^{d_z} \to \mathbb{R} : 0 \le S \le F,\ S_i \in C^\beta(D_i, K)\right\}.$$

For fixed P_Z, $\alpha \ge 0$, and $C > 0$, let $\mathcal{H}^\beta_{\mathcal{M},\alpha}(K,F,C)$ denote the subclass of $\mathcal{H}^\beta_{\mathcal{M}}(K,F)$ consisting of all functions S satisfying

$$\int_{[0,1]^{d_z}} \int_{\mathcal{M}} \mathbb{1}\{S(x \mid z) \le u\}\, d\nu(x)\, dP_Z(z) \le C u^\alpha \qquad \text{for every } u > 0.$$

The next theorem derives the convergence rates for manifold-supported intensities.

Theorem 4.4 (Convergence Rates for Manifold-Supported Intensities Using Deep Learning)**.** *Consider the problem of estimating the conditional intensity function S_0 of the Poisson process $X \mid Z$, where $Z \sim P_Z$. Fix the manifold, atlas, partition, and reference measure specified above. Suppose that $\beta > 0$, $\alpha \geq 0$, $K, F, C > 0$, and $S_0 \in \mathcal{H}^{\beta}_{\mathcal{M},\alpha}(K, F, C)$. Let $\hat{S}_n$ be the NPMLE defined in* (3)*, and let the hypothesis space $\mathcal{F}_n$ be the class of DNNs $\mathcal{F}_\sigma(L, r, B, F)$ satisfying*

$$\text{(i) } L \asymp \log(n), \qquad \text{(ii) } r \asymp \sqrt{n\phi'_n}, \qquad \text{(iii) } B \asymp n^{\frac{2(\beta+1)}{(1+(\alpha\wedge 1))\beta + d_{\mathcal{M}} + d_z}}.$$

Then there exists a constant $C' > 0$, depending only on the fixed geometry, $\beta, \alpha, K, F, C, P_Z$, and the chosen architecture constants, such that, for all sufficiently large n,

$$\mathbb{E}[\widetilde{H}^2(\hat{S}_n, S_0)] \leq C'\phi'_n \log(n)^3.$$

Here,

$$\phi'_n := n^{-\frac{(1+(\alpha\wedge 1))\beta}{(1+(\alpha\wedge 1))\beta + d_{\mathcal{M}} + d_z}}.$$

The exponent depends on the intrinsic input dimension $d_{\mathcal{M}} + d_z$ rather than $d_x + d_z$. The networks nevertheless take ambient coordinates as inputs; the charts are used to prove approximation within that network class. The likelihood uses the exact volume measure.

5 Simulation Studies

We conducted simulation studies to examine how estimation accuracy depends on sample size, covariate dimension, and the structure of the intensity. Three designs were considered: a compositional intensity, an intensity with a deep valley whose location varies with the covariates, and intensities supported on a circle or sphere. We compared DNN-based estimators with kernel estimators, including a pooled estimator to assess the effect of ignoring covariate dependence.

5.1 Simulation design

We generated independent pairs $Z_i \sim \mathrm{Unif}([0,1]^{d_z})$ and $X_i \mid Z_i \sim \mathrm{PP}\{S_0(\cdot \mid Z_i)\, d\nu\}$, $i = 1, \ldots, n$, with $d_z \in \{0, 1, 5\}$ and $n \in \{100, 316, 1000, 3162, 10000\}$. Each configuration was repeated 20 times. Here n is the number of point patterns. We scaled the intensities to give approximately 30 events per pattern on average. The Euclidean supports were $[0, 1]$ and $[0, 1]^2$; on the unit circle and sphere, ν was arc length and surface area, respectively.

The compositional design combines low-dimensional functions of event coordinates and covariates. The second design has a covariate-dependent valley with a positive floor of 10^{-4} before rescaling, allowing us to examine estimation in regions of low intensity. In both designs, at most the first two covariates enter the intensity, so the five-covariate cases also assess sensitivity to irrelevant coordinates. The manifold designs introduce periodic variation with angle and covariates. The spherical intensity includes a longitude-dependent term that is discontinuous at the poles; this design extends the numerical comparison beyond the smooth class in Theorem 4.4. Section G.1 gives the generating functions and simulation procedure.

Both DNN-based estimators have ReLU hidden activations. Their output functions are Softplus and ReLU, each with 10^{-8} added to ensure positivity for an initial stage of the optimization; we refer to them as DNN–Softplus and DNN–ReLU. We trained the networks by minimizing the negative log-likelihood using AdamW and stochastic quadrature. A random 80% of the patterns was used for parameter updates, and the remaining 20% for early stopping. Each update used a batch of 64 patterns and 128 integration points per pattern. Training was limited to 200 epochs. Network depth and width increased with n according to the rule in Section G.2.

For Euclidean supports, we used a Gaussian conditional kernel estimator and a pooled kernel estimator. The pooled estimator ignores covariates; when $d_z = 0$, the two estimators coincide. Bandwidths were selected by five-fold cross-validation, using at most 512 patterns and 128 events per pattern for this calculation, and the selected estimator was then fitted to all n patterns. For manifold supports, the kernel estimator used the known geodesic distance and volume measure. The main DNN comparison used a reconstructed integration mesh, obtained from independent samples on the support together with event locations. Thus, the mesh construction used support information in addition to the observed point patterns. We also fitted DNNs using the known manifold measure; Section G.4 compares these versions.

Accuracy was measured by the covariate-averaged squared Hellinger loss

$$\widetilde{H}^2(\widehat{S}, S_0) = \frac{1}{2}\mathbb{E}_Z\left[\int \left(\sqrt{\widehat{S}(x \mid Z)} - \sqrt{S_0(x \mid Z)}\right)^2 d\nu(x)\right].$$

Table 1: Mean squared Hellinger loss at $n = 10000$, with Monte Carlo standard errors in parentheses, based on 20 repetitions. Here d is the event-support dimension. "Kernel" denotes the conditional kernel for Euclidean supports and the geodesic kernel for manifold supports; the manifold DNNs use reconstructed integration meshes. A dash indicates that the pooled comparator was not used.

Design/support	d	d_z	Softplus	ReLU	Kernel	Pooled
Composition	1	1	0.00926 (0.00073)	0.0104 (0.00096)	0.0329 (0.0048)	0.474 (0.007)
Composition	1	5	0.0194 (0.0011)	0.0203 (0.0011)	0.144 (0.01)	0.544 (0.011)
Composition	2	1	0.0205 (0.00071)	0.0213 (0.00099)	0.0441 (0.0055)	0.398 (0.0052)
Composition	2	5	0.0346 (0.0013)	0.0329 (0.0012)	0.182 (0.0062)	0.466 (0.0067)
Near zero	1	1	0.0092 (0.0012)	0.0229 (0.0023)	0.0442 (0.00036)	2.61 (0.044)
Near zero	1	5	0.0179 (0.001)	0.0359 (0.0032)	0.597 (0.0051)	2.67 (0.031)
Near zero	2	1	0.00788 (0.00056)	0.0108 (0.00085)	0.0455 (0.0035)	0.516 (0.012)
Near zero	2	5	0.0296 (0.0011)	0.0408 (0.0022)	0.408 (0.0082)	0.951 (0.014)
Circle	1	1	0.0057 (0.0005)	0.00643 (0.00051)	0.0052 (0.00028)	–
Circle	1	5	0.0105 (0.00066)	0.0108 (0.00057)	0.0244 (0.00093)	–
Sphere	2	1	0.0102 (0.00058)	0.00997 (0.00057)	0.0205 (0.00084)	–
Sphere	2	5	0.0184 (0.00077)	0.0182 (0.00075)	0.0392 (0.0014)	–

We evaluated this loss using numerical integration and covariate draws independent of the fitting sample. Within each repetition, all methods used the same simulated data and evaluation points. Table 1 summarizes the results at $n = 10000$. Monte Carlo standard errors (MCSEs) are the sample standard deviations divided by $\sqrt{20}$. Learning curves and results for all sample sizes are reported below and in Section G.

5.2 Euclidean supports

Figure 1 shows the results for the compositional intensity. The conditional kernel performed well at small sample sizes. With one event coordinate, one covariate, and $n = 100$, its mean loss was 0.120, compared with 0.728 for DNN–Softplus. At $n = 1000$, the corresponding losses were 0.0554 and 0.0503. With two event coordinates and one covariate, the kernel still had a smaller mean loss at $n = 1000$: 0.0777, compared with 0.0869 for DNN–Softplus. The unconditional results showed a similar small-sample advantage for kernel smoothing: for example, with one event coordinate and $n = 100$, the pooled kernel and DNN–Softplus had mean losses of 0.0682 and 0.214, respectively.

As n increased, the DNN estimators became more competitive. At $n = 10000$, a DNN attained the smallest mean loss in each of the four covariate-dependent compositional configurations. The difference was largest with five covariates: the smaller of the two DNN mean losses was 0.0194 in one event dimension and 0.0329 in two event dimensions, compared with 0.144 and 0.182 for the conditional kernel. These comparisons indicate that neural estimation can be advantageous when several covariates are present, even though only a subset affects the intensity. The relative performance of the two output functions varied across configurations. The pooled kernel retained appreciable error at large n, reflecting the covariate variation that it cannot represent.

For the intensity with a moving valley, DNN–Softplus had a smaller mean loss than DNN–ReLU in 25 of the 30 configurations (Figure 3). It also attained the smallest mean loss in all four covariate-dependent configurations at $n = 10000$. With two event coordinates and five covariates, its mean loss was 0.0296 (MCSE 0.0011), compared with 0.0408 (0.0022) for DNN–ReLU and 0.408 (0.0082) for the conditional kernel. Thus, Softplus performed particularly well in this design with regions of very low intensity. Because the positive floor was held fixed, these results concern finite-sample estimation near zero rather than comparisons among the small-value exponents appearing in the theory.

5.3 Manifold supports

The geodesic kernel remained competitive on manifold supports (Figure 4). On the circle with one covariate and $n = 10000$, its mean loss was 0.00520, slightly smaller than 0.00570 for the mesh-based DNN–Softplus. With five covariates, the ordering reversed: the losses were 0.0244 and 0.0105, respectively. On the sphere with five covariates, the kernel had a smaller mean loss at $n = 1000$ (0.0480 versus 0.0646 for DNN–Softplus), whereas at $n = 10000$ the DNN losses were approximately 0.018, compared with 0.0392 for the kernel. As in the Euclidean experiments, the DNN advantage was most apparent at the largest sample size with several covariates.

Using the known manifold measure did not uniformly improve the DNN estimates. For example, with five covariates and $n = 10000$, the mesh-based and known-measure DNN–Softplus losses were 0.0105 and 0.0102 on the circle, and 0.0184 and 0.0185 on the sphere. These versions also use different dimensions in the network-width rule, so their

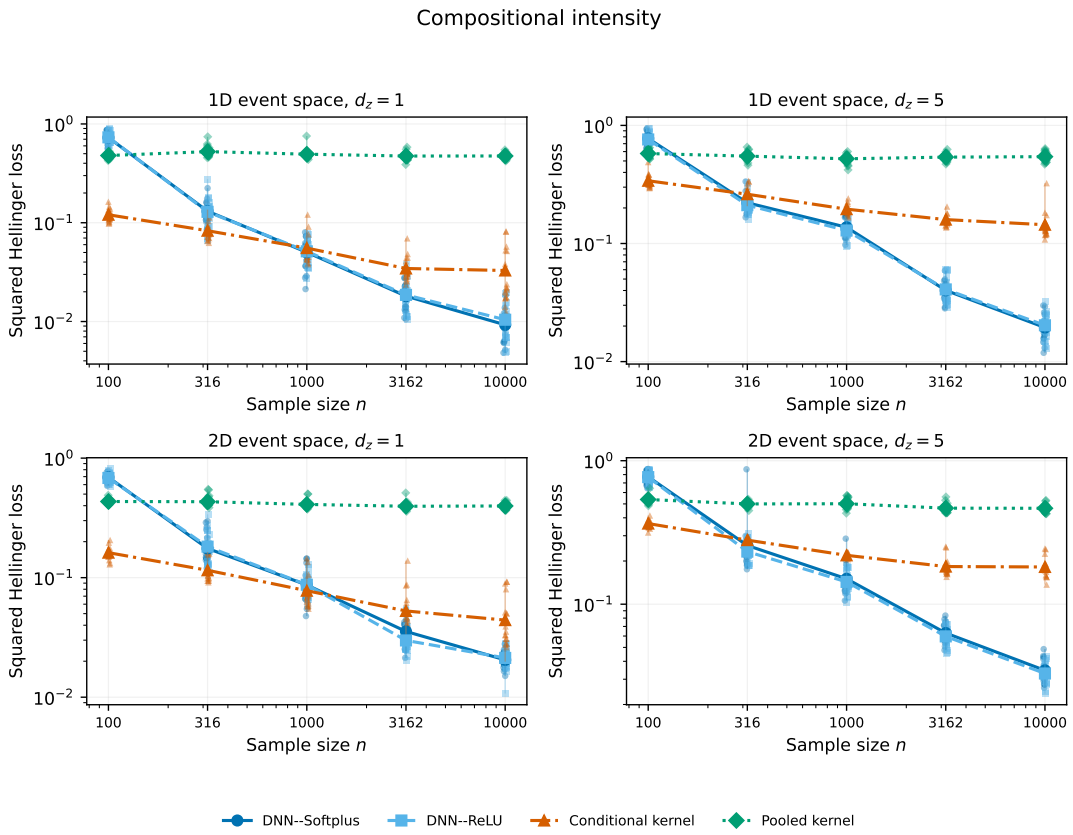

Figure 1: Squared Hellinger loss for the compositional intensity. Rows correspond to event dimension and columns to covariate dimension. Points represent individual repetitions, connected markers are means, and vertical segments span the minimum and maximum over 20 repetitions.

comparison reflects both architecture and integration measure. The full results in Section G document the variation across sample sizes and include the no-covariate cases.

6 Stimulus-Dependent Neural Spike Intensity

Visual stimuli can affect both the number and timing of neural spikes. We analyzed repeated Gabor presentations from the Allen Brain Observatory Visual Coding Neuropixels project (Siegle et al., 2021) to estimate how the temporal firing-rate profile varies with stimulus position and orientation. The main question was whether a neural intensity estimator could improve prediction beyond conditional kernel smoothing. A pooled temporal estimator provided a second comparison to quantify the benefit of accounting for the stimulus.

6.1 Data and analysis design

We used session 715093703, selected as the smallest session identifier among sessions with an available recording file and at least 20 eligible primary visual cortex (VISp) units. After excluding 620 presentations that overlapped invalid recording intervals and one with a different duration, 3,024 of the 3,645 Gabor presentations remained. Each presentation lasted $T = 0.25$ seconds. There were 243 stimulus conditions, formed by 81 screen positions and three axial orientations.

The covariate vector contained horizontal and vertical stimulus position, scaled using development-data extrema, and $\sin(2\theta)$ and $\cos(2\theta)$ for orientation θ. The doubled-angle encoding accounts for the equivalence of orientations separated by π. We split presentations within stimulus condition, assigning 2,417 to development and 607 to testing. Both subsets contained all 243 conditions. This design evaluates prediction for new presentations of conditions represented in the development sample.

Of the 62 VISp units satisfying the recording-quality criteria, we selected the 12 with the highest development-sample spike counts and estimated a separate intensity for each unit. Quality thresholds and preprocessing are detailed in Appendix H. Test results for one of these units, 950931315, had been examined in a preliminary analysis. We therefore report results both for the 12-unit panel and for the 11 units excluding this previously examined unit. All models reported here were refitted using the development-sample preprocessing described below.

6.2 Estimation and predictive assessment

We compared DNN–Softplus, DNN–ReLU, a conditional Gaussian kernel estimator, and a pooled temporal kernel estimator. Both networks had three hidden layers of width 64 with ReLU activations and a positive output offset of 10^{-8}. We selected the number of training epochs using 1,935 development presentations for fitting and 482 for validation, then refitted each network from initialization on all 2,417 development presentations for the selected number of epochs. Kernel bandwidths were chosen by five-fold presentation-level Poisson cross-validation on the development sample. The conditional kernel used the same stimulus covariates as the networks, whereas the pooled kernel used only spike times. Additional fitting details appear in Appendix H.

For unit v, let t_{vij} be spike j in presentation i, with m_{vi} spikes and covariate vector z_i. We fitted intensities on normalized time $u = t/T$ and converted them to spikes per second through $\widehat{S}_{v,t}(t \mid z) = \widehat{S}_{v,u}(t/T \mid z)/T$. Predictive performance was assessed by the Poisson negative log-likelihood (NLL) on the test presentations,

$$L_{vi}(\widehat{S}) = \int_0^T \widehat{S}_{v,t}(t \mid z_i)\, dt - \sum_{j=1}^{m_{vi}} \log \widehat{S}_{v,t}(t_{vij} \mid z_i), \tag{9}$$

with parameter-independent terms omitted. Integrals were evaluated on a 1,024-point trapezoidal grid. Scores are in nats per presentation, with time measured in seconds; paired score differences are invariant to the choice of time unit.

For a comparison of methods a and b, we averaged $L_{vi}(a) - L_{vi}(b)$ over test presentations within each unit, then averaged equally over units. A negative difference favors method a. We also report the equal-unit mean information gain in bits per spike, obtained by dividing each unit's total log-likelihood improvement by its test spike count and $\log 2$. To quantify uncertainty across stimulus conditions, we drew 2,000 paired bootstrap samples of the 243 conditions, retaining all presentations within a sampled condition and using the same draws for all methods and units. The resulting percentile intervals are pointwise 95% intervals conditional on the fitted models and selected units; they do not describe variation across neuronal populations or recording sessions.

6.3 Results

Accounting for the stimulus substantially improved prediction. Relative to the pooled temporal kernel, mean NLL decreased by 0.38926 nats per presentation for the conditional kernel, 0.41827 for DNN–Softplus, and 0.41006 for DNN–ReLU. Both neural estimators improved on the pooled estimator in all 12 units. The fitted profiles in Figure 8 illustrate this stimulus dependence: for unit 950931315, the covariate-dependent estimators distinguish the transient response at the preferred condition from the weaker responses at the other two conditions. The pooled estimator gives the same temporal profile for every stimulus.

The DNN estimators provided a smaller additional improvement over the conditional kernel (Table 2). For the 12-unit panel, the mean NLL differences were -0.02901 for DNN–Softplus and -0.02080 for DNN–ReLU, with 95% intervals $[-0.04042, -0.01729]$ and $[-0.03092, -0.01026]$, respectively. The corresponding information gains were 0.01073 and 0.00600 bits per spike. Excluding the previously examined unit gave similar results: the mean differences were -0.02610 and -0.01954, with both intervals remaining below zero.

Performance varied across units (Figure 2). DNN–Softplus had a lower mean NLL than the conditional kernel for 10 of the 12 units, and DNN–ReLU for nine. Both had their largest gain in unit 950932563, whereas the conditional kernel performed better in unit 950931458. Many unit-level intervals included zero. Although DNN–Softplus had the lower panel-average NLL, its difference from DNN–ReLU was -0.00821, with interval $[-0.01728, 0.00092]$. Thus, the NLL comparison did not clearly distinguish the two output functions.

Changing the integration grid had a negligible effect on the reported score differences. Supplementary analyses examine network initialization, width, running speed, prediction at held-out stimulus conditions, and temporal variation. The main comparison indicates that most of the gain over a pooled temporal estimate comes from modeling stimulus dependence, with a further modest improvement from the DNNs. This finding concerns repeated presentations in one session and a selected set of highly active units. The Poisson likelihood serves as a criterion for first-order intensity prediction; the time-rescaling diagnostics in Appendix H show remaining departures from a full Poisson description of the spike trains.

Table 2: Predictive performance relative to the conditional kernel. NLL differences and information gains are averaged equally over units; brackets give pointwise 95% paired condition-bootstrap intervals. "Improved" is the number of units with a negative NLL difference. The 11-unit analysis excludes unit 950931315, examined in the pilot.

Units	DNN	NLL difference (nats/presentation)	Improved	Information gain (bits/spike)
12	Softplus	-0.02901	10/12	0.01073
		$[-0.04042, -0.01729]$		$[0.00606, 0.01591]$
12	ReLU	-0.02080	9/12	0.00600
		$[-0.03092, -0.01026]$		$[0.00175, 0.01001]$
11	Softplus	-0.02610	9/11	0.00944
		$[-0.03791, -0.01378]$		$[0.00473, 0.01461]$
11	ReLU	-0.01954	8/11	0.00526
		$[-0.02991, -0.00933]$		$[0.00101, 0.00930]$

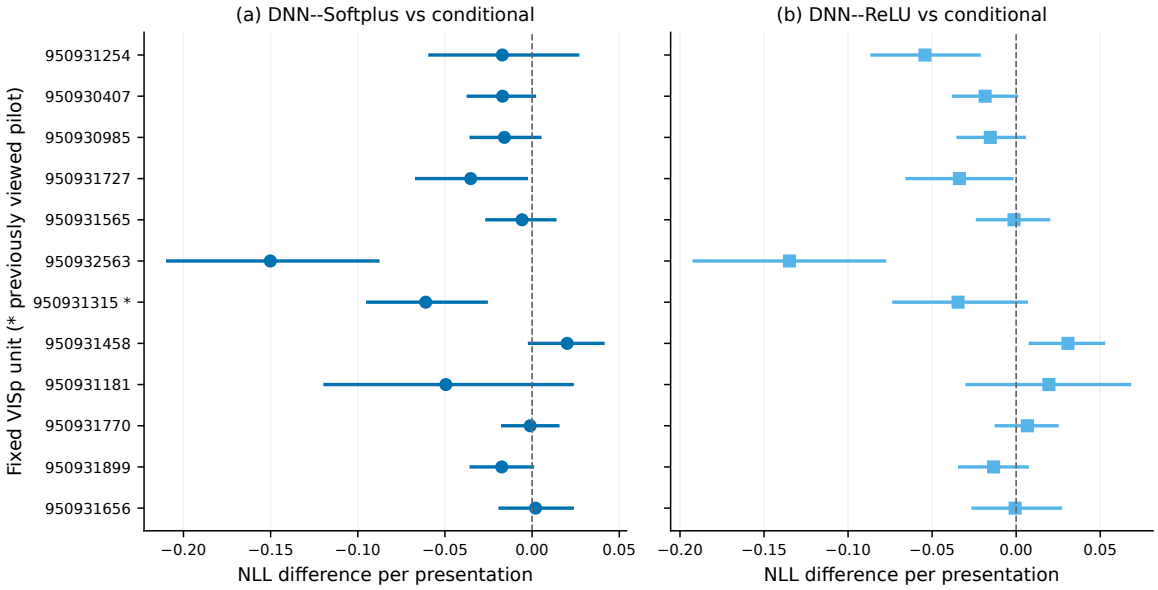

Figure 2: Unit-specific mean test NLL differences from the conditional kernel, in nats per presentation. Negative values favor the DNN. Horizontal segments give pointwise 95% paired condition-bootstrap intervals. The asterisk denotes unit 950931315, examined in the pilot.

7 Conclusion and Discussion

Likelihood-based intensity estimation can be analyzed without requiring the entire candidate class to be uniformly bounded away from zero. The oracle inequality makes the relevant trade-off explicit: an approximating intensity must both dominate the truth and approximate it under covariate-averaged Hellinger loss, while the local entropy controls estimation error. The small-value condition describes how the behavior of the true intensity near zero affects this trade-off. For compositional functions, the resulting rate reflects the dimensions of the component maps. The minimax comparison applies to the unrestricted small-value class on the specified Euclidean domain; it does not establish optimality for every positive small-value exponent.

The simulations show that the relative performance of neural and kernel estimation depends on sample size and covariate dimension. Kernel estimators often perform better in small samples, whereas the DNNs are more accurate in the largest covariate-dependent Euclidean settings and in several higher-dimensional manifold settings. Softplus performs particularly well for the intensity with a moving low-intensity region, but neither output function is uniformly preferable. These findings complement the approximation theory by showing the finite-sample trade-offs under the fitting procedures used here.

In the neural-recording application, modeling stimulus dependence accounts for most of the improvement over a pooled temporal estimate. The DNNs provide an additional, more modest improvement over conditional kernel smoothing,

with variation across units and no clear ordering of the two output functions by mean NLL. The analysis concerns repeated presentations within one session and a selected set of highly active units. Validation across sessions and animals would be needed to assess broader predictive performance.

Several extensions would strengthen the connection between theory and practice. First, the current theory concerns an exact constrained ReLU likelihood estimator with exact integration. A computational guarantee would need to account jointly for optimization error, quadrature error, and the output transformation used by the fitted networks. Second, learning an unknown support manifold requires control of the estimated integration measure in addition to function approximation on that support. The manifold theorem establishes approximation and risk bounds on a fixed smooth support for every $\beta > 0$, with the chartwise Hölder convention of Definition 4.2; it does not solve support estimation or adaptation to unknown intrinsic dimension. Third, repeated-session validation and chronological holdout would address forms of generalization not measured by the current neural-recording split. Modeling spike-history effects would also permit a direct assessment of departures from the conditional Poisson working model.

A Proof of lemmas in Section 2.3.2

A.1 Proof of Lemma 2.2

Proof. We first compute $\inf_{S\in\mathcal{F}_0}\mathbb{E}[\ell(S)]$. Suppose that, $\log(S(\cdot \mid z)) \in L^1(\lambda_0(\cdot \mid z)),\ P_Z\text{-}a.s.$ Then, by Campbell's formula (Last and Penrose, 2017, Proposition 2.7),

$$\begin{aligned}
\mathbb{E}[\ell(S)] &= \mathbb{E}\left[-\int_A \log(S(x \mid Z))\, dX(x \mid Z) + \int_A S(x \mid Z)\, d\nu(x)\right] \\
&= \mathbb{E}_Z\left[-\mathbb{E}\left[\int_A \log(S(x \mid Z))\, dX(x \mid Z)\middle| Z\right] + \int_A S(x \mid Z)\, d\nu(x)\right] \\
&= \mathbb{E}_Z\left[-\int_A \log(S(x \mid Z)) S_0(x \mid Z)\, d\nu(x) + \int_A S(x \mid Z)\, d\nu(x)\right] \\
&= \mathbb{E}_Z\left[\int_A \{S(x \mid Z) - S_0(x \mid Z)\log(S(x \mid Z))\}\ d\nu(x)\right].
\end{aligned}$$

Since, for $a \geq 0$ and $y \geq 0$, the function $y - a\log y$ is minimized at $y = a$, it follows that $\inf_{S\in\mathcal{F}_0}\mathbb{E}[\ell(S)] = \mathbb{E}[\ell(S_0)]$.

Indeed, for sufficiently small $\epsilon > 0$,

$$\begin{aligned}
&\int_A |\log(S_0(x \mid z))|\, d\lambda_0(x \mid z) = \int_A |\log(S_0(x \mid z))| S_0(x \mid z)\, d\nu(x) \\
&= \int_A |\log(S_0(x \mid z))| S_0(x \mid z) \mathbb{1}(S_0(x \mid z) \geq \epsilon)\, d\nu(x) + \int_A |\log(S_0(x \mid z))| S_0(x \mid z) \mathbb{1}(S_0(x \mid z) < \epsilon)\, d\nu(x) \\
&\leq \log\left(\frac{1}{\epsilon}\right)\int_A S_0(x \mid z)\mathbb{1}(S_0(x \mid z) \geq \epsilon)\, d\nu(x) + \int_A \log\left(\frac{1}{S_0(x \mid z)}\right) S_0(x \mid z)\mathbb{1}(S_0(x \mid z) < \epsilon)\, d\nu(x) \\
&\leq \log\left(\frac{1}{\epsilon}\right)\int_A S_0(x \mid z)\mathbb{1}(S_0(x \mid z) \geq \epsilon)\, d\nu(x) + \int_A (1 - S_0(x \mid z))\mathbb{1}(S_0(x \mid z) < \epsilon)\, d\nu(x).
\end{aligned}$$

Since S_0 is bounded and A is compact, we obtain $\log(S_0(\cdot \mid z)) \in L^1(\lambda_0(\cdot \mid z)),\ P_Z$-a.s. Here, in the first inequality, we used the boundedness of S_0 and the fact that ϵ is sufficiently small. In the last inequality, we used $\log x \leq x - 1$ for $x > 0$. On the other hand, if $\log(S(\cdot \mid z)) \notin L^1(\lambda_0(\cdot \mid z))$ for some $\mathcal{Z} \subset [0,1]^{d_z}$ with $P_Z(\mathcal{Z}) > 0$, then $\mathbb{E}[\ell(S)] = \infty$. To see this, first suppose that $S(\cdot \mid z)$ is a simple function, that is,

$$S(x \mid z) = S^{(m)}(x) \coloneqq \sum_{j=1}^m S_j \mathbb{1}(x \in A_j),$$

where $S_j \geq 0$ are constants, $m < \infty$, and $A_j \in \mathcal{B}(A)$ are pairwise disjoint. Then

$$-\mathbb{E}\left[\int_A \log(S(x \mid z)) dX(x \mid z)\right] = -\sum_{j=1}^m \log(S_j)\lambda_0(A_j \mid z) = -\int_A \log(S(x \mid z))\, d\lambda_0(x \mid z).$$

Next, consider a general nonnegative, bounded, measurable function. By measurability, there exists a sequence of nonnegative simple functions $(S^{(m)})$ such that, for every $x \in A$,

$$S^{(m)}(x \mid z) \to S(x \mid z) \ \text{ as } m \to \infty.$$

Hence, by the monotonicity and continuity of the logarithm, the boundedness of S, and the monotone convergence theorem, for any $\omega \in \Omega$,

$$\begin{aligned}
-\int_A \log(S^{(m)}(x \mid z))dX(x \mid z) &= \int_A \{-\log(S^{(m)}(x \mid z)) \vee 0\}dX(x \mid z) - \int_A \{\log(S^{(m)}(x \mid z)) \vee 0\}dX(x \mid z) \\
&\nearrow \int_A \{-\log(S(x \mid z)) \vee 0\}dX(x \mid z) - \int_A \{\log(S(x \mid z)) \vee 0\}dX(x \mid z) \\
&= -\int_A \log(S(x \mid z))dX(x \mid z) \ \text{ as } \ m \to \infty.
\end{aligned}$$

Therefore, again by the monotonicity and continuity of the logarithm, the boundedness of S, and the monotone convergence theorem,

$$\begin{aligned}
-\mathbb{E}\left[\int_A \log(S(x \mid z))dX(x \mid z)\right] &= \mathbb{E}\left[\lim_{m\to\infty}\left\{-\int_A \log(S^{(m)}(x \mid z))dX(x \mid z)\right\}\right] \\
&= \lim_{m\to\infty} \mathbb{E}\left[-\int_A \log(S^{(m)}(x \mid z))dX(x \mid z)\right] \\
&= -\lim_{m\to\infty}\int_A \log(S^{(m)}(x \mid z))\, d\lambda_0(x \mid z) \\
&= -\int_A \log(S(x \mid z))\, d\lambda_0(x \mid z).
\end{aligned}$$

Since $\log(S(\cdot \mid z)) \notin L^1(\lambda_0(\cdot \mid z))$ and S is bounded, we must have

$$-\int_A \log(S(x \mid z))\, d\lambda_0(x \mid z) = \infty.$$

Hence, together with the assumption that the covariate Z has a positive density on its domain, this implies that if $\log(S(\cdot \mid z)) \notin L^1(\lambda_0(\cdot \mid z))$, then $\mathbb{E}[\ell(S)] = \infty$, and therefore

$$\inf_{S\in\mathcal{F}_0} \mathbb{E}[\ell(S)] = \mathbb{E}[\ell(S_0)].$$

Next, we compute $R(\hat{S}_n)$. If $\lambda_0(\cdot \mid z) \ll \hat{\lambda}_n(\cdot \mid z)$, $P_Z - a.s.$, then by the formula of the KL divergence (Leskelä, 2024, Corollary 5.2) and Campbell's formula, we have

$$\begin{aligned}
R(\hat{S}_n) &= \mathbb{E}[\ell(\hat{S}_n)] - \mathbb{E}[\ell(S_0)] \\
&= \mathbb{E}\left[\int_A \log\left(\frac{S_0(x \mid Z)}{\hat{S}_n(x \mid Z)}\right) dX(x \mid Z) + \int_A \left(\hat{S}_n(x \mid Z) - S_0(x \mid Z)\right)\, d\nu(x)\right] \\
&= \mathbb{E}_Z\left[\mathbb{E}\left[\int_A \log\left(\frac{S_0(x \mid Z)}{\hat{S}_n(x \mid Z)}\right) dX(x \mid Z)\middle| Z\right] + \int_A \left(\hat{S}_n(x \mid Z) - S_0(x \mid Z)\right)\, d\nu(x)\right] \\
&= \mathbb{E}_Z\Big[\mathrm{KL}\Big(P_{\lambda_0(\cdot|Z)} \parallel P_{\hat{\lambda}_n(\cdot|Z)}\Big)\Big].
\end{aligned}$$

On the other hand, if there exists some $\mathcal{Z} \subset [0,1]^{d_z}$ with $P_Z(\mathcal{Z}) > 0$ such that $\lambda_0(\cdot \mid z) \ll \hat{\lambda}_n(\cdot \mid z)$ is not satisfied for all $z \in \mathcal{Z}$, then there exists $B_z \in \mathcal{B}(A)$ such that $\lambda_0(B_z \mid z) > 0$ and $\hat{\lambda}_n(B_z \mid z) = 0$. Hence, $\hat{S}_n(x \mid z) = 0$ almost everywhere on B_z, and there exists a subset $B_0 \subset B_z$ of positive measure such that $S_0(x \mid z) > 0$ for every $x \in B_0$.

Using this fact and Proposition 3.8 of Last and Penrose (2017), we obtain

$$\begin{aligned}
&\mathbb{E}\left[\int_A \log\left(\frac{S_0(x \mid z)}{\hat{S}_n(x \mid z)}\right) dX(x \mid z)\right] \\
&= \mathbb{E}_N\left[\mathbb{E}_{x_1:x_N}\left[\sum_{j=1}^N \log\left(\frac{S_0(x_j \mid z)}{\hat{S}_n(x_j \mid z)}\right)\middle| N\right]\right] \\
&= \mathbb{E}_N\left[\int_A \cdots \int_A \sum_{j=1}^N \log\left(\frac{S_0(x_j \mid z)}{\hat{S}_n(x_j \mid z)}\right) \prod_{j=1}^N \frac{S_0(x_j \mid z)}{\lambda_0(A \mid z)}\, d\nu(x_1)\cdots d\nu(x_N)\right] \\
&\geq \mathbb{E}_N\left[\int_{B_0} \cdots \int_{B_0} \sum_{j=1}^N \log\left(\frac{S_0(x_j \mid z)}{\hat{S}_n(x_j \mid z)}\right) \prod_{j=1}^N \frac{S_0(x_j \mid z)}{\lambda_0(A \mid z)}\, d\nu(x_1)\cdots d\nu(x_N)\right] = \infty.
\end{aligned}$$

Here, N is a Poisson random variable with intensity parameter $\lambda_0(A \mid z)$. Therefore, together with the assumption that the covariate Z has a positive density on its domain, the claim follows. □

A.2 Proof of Lemma 2.3

Proof. By Lemma 2.2 and the inequality $\log x \leq x - 1$ for any $x > 0$,

$$\begin{aligned}
R(\hat{S}_n) &= \mathbb{E}_Z\left[\int_A S_0(x \mid Z) \log\left(\frac{S_0(x \mid Z)}{\hat{S}_n(x \mid Z)}\right) d\nu(x) + \int_A \left(\hat{S}_n(x \mid Z) - S_0(x \mid Z)\right) d\nu(x)\right] \\
&= \mathbb{E}_Z\left[-2\int_A S_0(x \mid Z) \log\left(\sqrt{\frac{\hat{S}_n(x \mid Z)}{S_0(x \mid Z)}}\right) d\nu(x) + \int_A \left(\hat{S}_n(x \mid Z) - S_0(x \mid Z)\right) d\nu(x)\right] \\
&\geq \mathbb{E}_Z\left[-2\int_A S_0(x \mid Z) \left(\sqrt{\frac{\hat{S}_n(x \mid Z)}{S_0(x \mid Z)}} - 1\right) d\nu(x) + \int_A \left(\hat{S}_n(x \mid Z) - S_0(x \mid Z)\right) d\nu(x)\right] \\
&= \mathbb{E}_Z\left[\int_A \left(S_0(x \mid Z) - 2\sqrt{S_0(x \mid Z)\hat{S}_n(x \mid Z)} + \hat{S}_n(x \mid Z)\right) d\nu(x)\right] = 2\widetilde{H}^2(\hat{S}_n, S_0).
\end{aligned}$$

□

B Proof of Theorem 3.2 and Corollary 3.3

In this section, we provide the proof of Theorem 3.2, Corollary 3.3, and technical lemmas to prove the theorem.

B.1 Proof of Lemma 3.1

Proof. Let

$$\bar{S}_n \coloneqq \frac{\hat{S}_n + S_n^*}{2}.$$

First, by the concavity of the logarithm, for any $n \in \mathbb{N}$, $z \in [0,1]^{d_z}$, and $x \in A_n^*(z)$,

$$\frac{1}{2}\log\left(\frac{\bar{S}_n(x \mid z)}{S_n^*(x \mid z)}\right) = \frac{1}{2}\log\left(\frac{\hat{S}_n(x \mid z) + S_n^*(x \mid z)}{2S_n^*(x \mid z)}\right) \geq \frac{1}{4}\log\left(\frac{\hat{S}_n(x \mid z)}{S_n^*(x \mid z)}\right). \tag{10}$$

Since $\hat{S}_n$ is the maximum likelihood estimator and $S_n^* \in \mathcal{F}_n$, we have

$$0 \leq P_n\left\{\int_A \log\left(\frac{\hat{S}_n(x \mid Z)}{S_n^*(x \mid Z)}\right) dX(x \mid Z) - \int_A \left(\hat{S}_n(x \mid Z) - S_n^*(x \mid Z)\right) d\nu(x)\right\}. \tag{11}$$

By Assumption 3.1,

$$S_n^*(x \mid z) = 0 \quad \Longrightarrow \quad S_0(x \mid z) = 0 \quad \text{for } \nu\text{-a.e. } x.$$

Hence,

$$\mathbb{E}\left[X_i\left(A \setminus A_n^*(z_i) \mid z_i\right) \mid Z_i = z_i\right] = \int_{A \setminus A_n^*(z_i)} S_0(x \mid z_i)\, d\nu(x) = 0.$$

Therefore, since $X_i(A \setminus A_n^*(z_i) \mid z_i)$ is nonnegative random variable,

$$X_i\left(A \setminus A_n^*(z_i) \mid z_i\right) = 0 \quad \text{a.s.}$$

Consequently, the stochastic integral in (11) may be restricted to $A_n^*(Z)$ almost surely. Moreover, on $A \setminus A_n^*(Z)$, we have $S_n^*(x \mid Z) = 0$. Thus,

$$\begin{aligned} 0 \le P_n \Bigg\{ & \int_{A_n^*(Z)} \log\left(\frac{\hat{S}_n(x \mid Z)}{S_n^*(x \mid Z)}\right) dX(x \mid Z) \\ & + \int_{A_n^*(Z)} \left(S_n^*(x \mid Z) - \hat{S}_n(x \mid Z)\right) d\nu(x) - \int_{A \setminus A_n^*(Z)} \hat{S}_n(x \mid Z)\, d\nu(x) \Bigg\}. \end{aligned} \tag{12}$$

Multiplying both sides by $1/4$ and using (10), we obtain

$$\begin{aligned} 0 \le P_n \Bigg\{ & \int_{A_n^*(Z)} \frac{1}{2} \log\left(\frac{\bar{S}_n(x \mid Z)}{S_n^*(x \mid Z)}\right) dX(x \mid Z) \\ & + \frac{1}{2} \int_{A_n^*(Z)} \left(S_n^*(x \mid Z) - \bar{S}_n(x \mid Z)\right) d\nu(x) \; - \frac{1}{2} \int_{A \setminus A_n^*(Z)} \bar{S}_n(x \mid Z)\, d\nu(x) \Bigg\}. \end{aligned} \tag{13}$$

Here we used

$$\frac{1}{4}\left(S_n^* - \hat{S}_n\right) = \frac{1}{2}\left(S_n^* - \bar{S}_n\right) \quad \text{on } A_n^*(Z),$$

and, since $S_n^* = 0$ on $A \setminus A_n^*(Z)$,

$$\frac{1}{4}\hat{S}_n = \frac{1}{2}\bar{S}_n \quad \text{on } A \setminus A_n^*(Z).$$

Define

$$\begin{aligned} G_n(Z, X) \coloneqq & \int_{A_n^*(Z)} \frac{1}{2} \log\left(\frac{\bar{S}_n(x \mid Z)}{S_n^*(x \mid Z)}\right) dX(x \mid Z) \\ & + \frac{1}{2} \int_{A_n^*(Z)} \left(S_n^*(x \mid Z) - \bar{S}_n(x \mid Z)\right) d\nu(x) - \frac{1}{2} \int_{A \setminus A_n^*(Z)} \bar{S}_n(x \mid Z)\, d\nu(x). \end{aligned}$$

Then (13) gives

$$0 \le P_n G_n = (P_n - P)G_n + PG_n.$$

By Campbell's formula,

$$\begin{aligned} PG_n = \mathbb{E}_Z \Bigg\{ & \int_{A_n^*(Z)} \frac{1}{2} \log\left(\frac{\bar{S}_n(x \mid Z)}{S_n^*(x \mid Z)}\right) S_0(x \mid Z)\, d\nu(x) \\ & + \frac{1}{2} \int_{A_n^*(Z)} \left(S_n^*(x \mid Z) - \bar{S}_n(x \mid Z)\right) d\nu(x) \; - \frac{1}{2} \int_{A \setminus A_n^*(Z)} \bar{S}_n(x \mid Z)\, d\nu(x) \Bigg\}. \end{aligned}$$

The application of Campbell's formula is justified as follows. Since $\log(\bar{S}_n / S_n^*) \ge -\log 2$ on $A_n^*(z)$, the inequality $\log x \le x - 1$ for $x > 0$, together with Assumption 3.1, implies that, for any $z \in [0,1]^{d_z}$,

$$\begin{aligned} & \int_{A_n^*(z)} \left|\log\left(\frac{\bar{S}_n(x \mid z)}{S_n^*(x \mid z)}\right)\right| S_0(x \mid z)\, d\nu(x) \\ & = \int_{A_n^*(z)} \left|\log\left(\frac{\bar{S}_n(x \mid z)}{S_n^*(x \mid z)}\right)\right| \frac{S_0(x \mid z)}{S_n^*(x \mid z)} S_n^*(x \mid z)\, d\nu(x) \\ & \le c_0^2 \int_{A_n^*(z)} \left[\log(2) S_n^*(x \mid z) + \left|\bar{S}_n(x \mid z) - S_n^*(x \mid z)\right|\right] d\nu(x) < \infty. \end{aligned}$$

Next, for any $z \in [0,1]^{d_z}$, using $\log x \le x - 1$ with $x = \sqrt{\bar{S}_n / S_n^*}$ gives

$$\int_{A_n^*(z)} \frac{1}{2} \log\left(\frac{\bar{S}_n(x \mid z)}{S_n^*(x \mid z)}\right) S_0(x \mid z)\, d\nu(x) \le -\int_{A_n^*(z)} \left(1 - \sqrt{\frac{\bar{S}_n(x \mid z)}{S_n^*(x \mid z)}}\right) S_0(x \mid z)\, d\nu(x).$$

We decompose S_0 as $S_0 = S_n^* + (S_0 - S_n^*)$ on $A_n^*(z)$. Then

$$\begin{aligned}
&-\int_{A_n^*(z)} \left(1 - \sqrt{\frac{\bar{S}_n(x \mid z)}{S_n^*(x \mid z)}}\right) S_0(x \mid z)\, d\nu(x) \\
&= -\int_{A_n^*(z)} \left(S_n^*(x \mid z) - \sqrt{S_n^*(x \mid z)}\sqrt{\bar{S}_n(x \mid z)}\right) d\nu(x) \\
&\quad -\int_{A_n^*(z)} \left(\sqrt{S_n^*(x \mid z)} - \sqrt{\bar{S}_n(x \mid z)}\right) \left(\sqrt{\frac{S_0(x \mid z)}{S_n^*(x \mid z)}} + 1\right) \left(\sqrt{S_0(x \mid z)} - \sqrt{S_n^*(x \mid z)}\right) d\nu(x).
\end{aligned}$$

By Assumption 3.1,

$$\sqrt{\frac{S_0(x \mid z)}{S_n^*(x \mid z)}} \le c_0 \quad \text{on } A_n^*(z).$$

Hence, by the Cauchy–Schwarz inequality,

$$\begin{aligned}
&-\int_{A_n^*(z)} \left(\sqrt{S_n^*(x \mid z)} - \sqrt{\bar{S}_n(x \mid z)}\right) \left(\sqrt{\frac{S_0(x \mid z)}{S_n^*(x \mid z)}} + 1\right) \left(\sqrt{S_0(x \mid z)} - \sqrt{S_n^*(x \mid z)}\right) d\nu(x) \\
&\le 2(c_0+1) H\left(\bar{S}_n(\cdot \mid z), S_n^*(\cdot \mid z)\right) H\left(S_n^*(\cdot \mid z), S_0(\cdot \mid z)\right).
\end{aligned}$$

Therefore,

$$\begin{aligned}
&\int_{A_n^*(z)} \frac{1}{2} \log\left(\frac{\bar{S}_n(x \mid z)}{S_n^*(x \mid z)}\right) S_0(x \mid z)\, d\nu(x) \\
&\le -\int_{A_n^*(z)} \left(S_n^*(x \mid z) - \sqrt{S_n^*(x \mid z)}\sqrt{\bar{S}_n(x \mid z)}\right) d\nu(x) \\
&\quad + 2(c_0+1) H\left(\bar{S}_n(\cdot \mid z), S_n^*(\cdot \mid z)\right) H\left(S_n^*(\cdot \mid z), S_0(\cdot \mid z)\right).
\end{aligned}$$

Combining this bound with the second integral term in the definition of G_n, we obtain

$$\begin{aligned}
&-\int_{A_n^*(z)} \left(S_n^*(x \mid z) - \sqrt{S_n^*(x \mid z)}\sqrt{\bar{S}_n(x \mid z)}\right) d\nu(x) + \frac{1}{2}\int_{A_n^*(z)} \left(S_n^*(x \mid z) - \bar{S}_n(x \mid z)\right) d\nu(x) \\
&= -\frac{1}{2}\int_{A_n^*(z)} \left(\sqrt{\bar{S}_n(x \mid z)} - \sqrt{S_n^*(x \mid z)}\right)^2 d\nu(x).
\end{aligned}$$

Moreover, since $S_n^*(x \mid z) = 0$ on $A \setminus A_n^*(z)$,

$$-\frac{1}{2}\int_{A \setminus A_n^*(z)} \bar{S}_n(x \mid z)\, d\nu(x) = -\frac{1}{2}\int_{A \setminus A_n^*(z)} \left(\sqrt{\bar{S}_n(x \mid z)} - \sqrt{S_n^*(x \mid z)}\right)^2 d\nu(x).$$

Hence,

$$\begin{aligned}
PG_n &\le -\mathbb{E}_Z\left[H^2\left(\bar{S}_n(\cdot \mid Z), S_n^*(\cdot \mid Z)\right)\right] + 2(c_0+1)\mathbb{E}_Z\left[H\left(\bar{S}_n(\cdot \mid Z), S_n^*(\cdot \mid Z)\right) H\left(S_n^*(\cdot \mid Z), S_0(\cdot \mid Z)\right)\right] \\
&= -\widetilde{H}^2\left(\bar{S}_n, S_n^*\right) + 2(c_0+1)\mathbb{E}_Z\left[H\left(\bar{S}_n(\cdot \mid Z), S_n^*(\cdot \mid Z)\right) H\left(S_n^*(\cdot \mid Z), S_0(\cdot \mid Z)\right)\right].
\end{aligned}$$

By Hölder's inequality,

$$\mathbb{E}_Z\left[H\left(\bar{S}_n(\cdot \mid Z), S_n^*(\cdot \mid Z)\right) H\left(S_n^*(\cdot \mid Z), S_0(\cdot \mid Z)\right)\right] \le \widetilde{H}\left(\bar{S}_n, S_n^*\right) \widetilde{H}\left(S_n^*, S_0\right).$$

Since $S_n^*(x \mid Z) = 0$ on $A \setminus A_n^*(Z)$, the integral terms in G_n can be combined as

$$\frac{1}{2}\int_{A_n^*(Z)} \left(S_n^*(x \mid Z) - \bar{S}_n(x \mid Z)\right) d\nu(x) - \frac{1}{2}\int_{A \setminus A_n^*(Z)} \bar{S}_n(x \mid Z) d\nu(x) = \frac{1}{2}\int_A \left(S_n^*(x \mid Z) - \bar{S}_n(x \mid Z)\right) d\nu(x).$$

This completes the proof. □

B.2 Technical lemmas

In this section, we provide several definitions, technical lemmas, and their proof.

Lemma B.1. *For any conditional intensity functions S_1 and S_2, we have*

$$\widetilde{H}^2(S_1, S_2) \leq 16\widetilde{H}^2\left(\frac{S_1+S_2}{2}, S_2\right).$$

Proof. By the subadditivity of the square-root function, for any $a, b \geq 0$, we have

$$\sqrt{\frac{a+b}{2}} + \sqrt{b} \leq 2\left(\sqrt{a}+\sqrt{b}\right).$$

Therefore, by simple calculation, we have

$$|\sqrt{a}-\sqrt{b}| \leq \frac{2|(a+b)/2-b|}{\sqrt{a}+\sqrt{b}} = 2\cdot\frac{\sqrt{(a+b)/2}+\sqrt{b}}{\sqrt{a}+\sqrt{b}}\left|\sqrt{\frac{a+b}{2}}-\sqrt{b}\right| \leq 4\left|\sqrt{\frac{a+b}{2}}-\sqrt{b}\right|.$$

Substituting $a = S_1$, $b = S_2$, squaring both sides, integrating over A, and taking expectation with respect to Z gives the claim. □

Definition B.1 (Bernstein size). Let P be a probability measure on a measurable space $(S, \mathcal{S})$, and let $K > 0$ be a constant. For a measurable function f, define

$$\rho_K^2(f) = 2K^2\mathbb{E}\left[e^{|f(X)|/K} - 1 - \frac{|f(X)|}{K}\right],$$

where X is a random variable with distribution P. The positive square root of $\rho_K^2(f)$ is called the Bernstein size of f.

Using this Bernstein size, we define the generalized bracketing number as follows.

Definition B.2 (Generalized bracketing number). Let $\mathcal{F}$ be a class of functions, let $\delta > 0$, and let P be a probability measure. Let $\mathcal{N}_{B,K}(\delta, \mathcal{F}, P)$ be the smallest integer N for which there exists a collection of pairs of functions $\{(f_j^L, f_j^U)\}_{j=1}^N$ satisfying the following conditions:

- For every $j = 1, \ldots, N$, we have $\rho_K(f_j^U - f_j^L) \leq \delta$.
- For every $f \in \mathcal{F}$, there exists an index $j = j(f) \in \{1, \ldots, N\}$ such that $f_j^L \leq f \leq f_j^U$.

If no such N exists, we set $\mathcal{N}_{B,K}(\delta, \mathcal{F}, P) = \infty$. We call $\mathcal{N}_{B,K}(\delta, \mathcal{F}, P)$ the δ-generalized bracketing number of $\mathcal{F}$.

The following Bernstein-type uniform inequality will be used to control the empirical process.

Theorem B.2 (Theorem 5.11 of van de Geer, 2000; Theorem 3.5.21 of Giné and Nickl, 2021)**.** *Let $\mathcal{G}$ be a class of functions. Suppose that, for some constants K and R, the class $\mathcal{G}$ satisfies*

$$\sup_{g\in\mathcal{G}} \rho_K(g) \leq R. \tag{14}$$

If

$$a \leq C_1\sqrt{n}R^2/K, \tag{15}$$

$$a \leq 8\sqrt{n}R, \tag{16}$$

$$a \geq C_0\left(\int_{a/(2^6\sqrt{n})}^{R} \sqrt{\log \mathcal{N}_{B,K}(u, \mathcal{G}, P)}\, du \vee R\right), \tag{17}$$

$$C_0^2 \geq C^2(C_1+1), \tag{18}$$

then

$$\mathbb{P}\left(\sqrt{n}\sup_{g\in\mathcal{G}}|(P_n - P)g| \geq a\right) \leq C\exp\left(-\frac{a^2}{C^2(C_1+1)R^2}\right).$$

Lemma B.3 (Lemma 7.1 of van de Geer, 2000)**.** *For any $x \geq -\Gamma$ with $\Gamma > 0$, we have*

$$2(e^{|x|} - 1 - |x|) \leq c_\Gamma^2(e^x - 1)^2, \quad c_\Gamma = \frac{2(e^\Gamma - 1 - \Gamma)}{(e^{-\Gamma}-1)^2}.$$

Lemma B.4. *For any $n \in \mathbb{N}$ and any bounded real numbers $x_1, \ldots, x_n, y_1, \ldots, y_n$ satisfying $|x_i| \le M$ and $|y_i| \le M$ for some $M \ge 0$, we have*

$$\left(\prod_{i=1}^n x_i - \prod_{i=1}^n y_i\right)^2 \le M^{2n-2} n \sum_{i=1}^n (x_i - y_i)^2.$$

Proof. We first prove, by induction, that

$$\left|\prod_{i=1}^n x_i - \prod_{i=1}^n y_i\right| \le M^{n-1} \sum_{i=1}^n |x_i - y_i|.$$

The claim is obvious when $n = 1$. Suppose that $n \ge 2$ and that the claim holds for $n-1$. Then, by the boundedness of x_i and y_i, the triangle inequality, and the induction hypothesis,

$$\begin{aligned}
\left|\prod_{i=1}^n x_i - \prod_{i=1}^n y_i\right| &= \left|x_n \prod_{i=1}^{n-1} x_i - y_n \prod_{i=1}^{n-1} x_i + y_n \prod_{i=1}^{n-1} x_i - y_n \prod_{i=1}^{n-1} y_i\right| \\
&\le |x_n - y_n| \prod_{i=1}^{n-1} |x_i| + |y_n| \left|\prod_{i=1}^{n-1} x_i - \prod_{i=1}^{n-1} y_i\right| \le M^{n-1}|x_n - y_n| + M^{n-1} \sum_{i=1}^{n-1} |x_i - y_i| \\
&\le M^{n-1} \sum_{i=1}^n |x_i - y_i|.
\end{aligned}$$

This proves the desired bound. Therefore, using the inequality

$$\left(\sum_{j=1}^n a_j\right)^2 \le n \sum_{j=1}^n a_j^2$$

for arbitrary nonnegative real numbers $a_j \ge 0$, $j = 1, \ldots, n$, we obtain

$$\left(\prod_{i=1}^n x_i - \prod_{i=1}^n y_i\right)^2 \le M^{2n-2} n \sum_{i=1}^n (x_i - y_i)^2.$$

□

Lemma B.5. *Suppose that Assumption 3.1 holds and $\sup_{S \in \mathcal{F}_n} \|S\|_\infty \le F < \infty$. Then, for any $g_S \in \mathcal{G}_n$, we have*

$$\rho_1^2(g_S) \le 2^6 C_\rho^2 \widetilde{H}^2 \left(\frac{S + S_n^*}{2}, S_n^*\right).$$

Here, the constant C_ρ is given by

$$C_\rho = \exp\left(2c_0^2(\nu(A) \vee 1)(F \vee 1)\right).$$

Proof. Fix z and put $A^* = A_n^*(z)$, $\lambda = \lambda_0(A^* \mid z)$, $V = \nu(A^*)$, and $H_z^2 = H^2(\bar{S}(\cdot \mid z), S_n^*(\cdot \mid z))$. If $\lambda = 0$, then $X(A^*) = 0$ almost surely and $g_S = 0$. Suppose henceforth that $\lambda > 0$. Conditional on $N = X(A^*) \ge 1$, the locations have density S_0/λ and $g_S \ge -N \log 2/2$. With $\Gamma = N \log 2/2$, the preceding two lemmas and $S_0 \le c_0^2 S_n^*$ therefore give

$$\begin{aligned}
&\mathbb{E}[2(e^{|g_S|} - 1 - |g_S|) \mid Z = z, N] \\
&\le \frac{c_\Gamma^2 c_0^{2N}}{\lambda^N} \int_{(A^*)^N} \left(\prod_{j=1}^N \sqrt{\bar{S}(x_j \mid z)} - \prod_{j=1}^N \sqrt{S_n^*(x_j \mid z)}\right)^2 d\nu^{\otimes N} \\
&\le \frac{2c_\Gamma^2 c_0^{2N} N^2 (FV)^{N-1}}{\lambda^N} H_z^2.
\end{aligned}$$

For $b = e^\Gamma > 1$, $\log b \ge 1 - b^{-1}$ implies

$$c_\Gamma = \frac{2(b - 1 - \log b)}{(1 - b^{-1})^2} \le 2b, \qquad c_\Gamma^2 \le 4\, 2^N.$$

Put $M = (\nu(A) \vee 1)(F \vee 1) \geq 1$ and $x = 2c_0^2 M$. The contribution of $N = 0$ is zero. Summing the preceding bound over $N \sim \text{Poisson}(\lambda)$ yields

$$\begin{aligned}\mathbb{E}[2(e^{|g_S|} - 1 - |g_S|) \mid Z = z] &\leq \frac{8}{M} e^{-\lambda} H_z^2 \sum_{N \geq 1} \frac{N^2 x^N}{N!} \\ &= \frac{8}{M}(x^2 + x) e^{x-\lambda} H_z^2 \leq 8 e^{2x} H_z^2 \leq 64 C_\rho^2 H_z^2.\end{aligned}$$

Here $(x^2 + x) \leq e^x$ for $x \geq 0$; for example this follows from $e^x \geq 1 + x + x^2/2 + x^3/6$ and $1 - x^2/2 + x^3/6 \geq 1/3$. Also $e^{2x} = C_\rho^2$. The last coefficient is independent of z. Taking expectation over Z now proves the result without separating dependent factors. □

Lemma B.6. *Suppose Assumption 3.1 holds and* $\sup_{S \in \mathcal{F}_n} \|S\|_\infty = F < \infty$. *Let* S_L *and* S_U *satisfy* $0 \leq S_L(x \mid z) \leq S_U(x \mid z) \leq F$. *pointwise. Define*

$$\begin{aligned}g_L &= \int_{A_n^*(Z)} \frac{1}{2} \log\left(\frac{S_L(x \mid Z) + S_n^*(x \mid Z)}{2 S_n^*(x \mid Z)}\right) dX(x \mid Z), \\ g_U &= \int_{A_n^*(Z)} \frac{1}{2} \log\left(\frac{S_U(x \mid Z) + S_n^*(x \mid Z)}{2 S_n^*(x \mid Z)}\right) dX(x \mid Z).\end{aligned}$$

Then

$$\rho_1^2(g_U - g_L) \leq 2 C_\rho^2 \widetilde{H}^2\left(\frac{S_U + S_n^*}{2}, \frac{S_L + S_n^*}{2}\right).$$

Here, C_ρ *is the same constant as in Lemma B.5.*

Proof. Put $\bar{S}_j = (S_j + S_n^*)/2$ for $j \in \{L, U\}$ and $h = g_U - g_L \geq 0$. Fix z and write

$$A^* = A_n^*(z), \qquad \lambda = \lambda_0(A^* \mid z), \qquad V = \nu(A^*), \qquad H_z^2 = H^2\big(\bar{S}_U(\cdot \mid z), \bar{S}_L(\cdot \mid z)\big).$$

If $\lambda = 0$, then $h = 0$ almost surely conditional on $Z = z$. Suppose that $\lambda > 0$ and condition on $N = X(A^*) \geq 1$. Since $2(e^u - 1 - u) \leq (e^u - 1)^2$ for $u \geq 0$ and $S_0/\bar{S}_L \leq 2c_0^2$ on A^*, the conditional distribution of the points and Lemma B.4 give

$$\begin{aligned}&\mathbb{E}\big[2(e^h - 1 - h) \mid Z = z, N\big] \\ &\leq \frac{(2c_0^2)^N}{\lambda^N} \int_{(A^*)^N} \left(\prod_{j=1}^N \sqrt{\bar{S}_U(x_j \mid z)} - \prod_{j=1}^N \sqrt{\bar{S}_L(x_j \mid z)}\right)^2 d\nu^{\otimes N} \\ &\leq \frac{2(2c_0^2)^N N^2 (FV)^{N-1}}{\lambda^N} H_z^2.\end{aligned}$$

Let $M = (\nu(A) \vee 1)(F \vee 1)$ and $x = 2c_0^2 M$. The contribution of $N = 0$ is zero. Since

$$\sum_{N \geq 1} \frac{N^2 x^N}{N!} = (x^2 + x) e^x \quad \text{and} \quad x^2 + x \leq e^x \quad (x \geq 0),$$

summing over $N \sim \text{Poisson}(\lambda)$ yields

$$\begin{aligned}\mathbb{E}\big[2(e^h - 1 - h) \mid Z = z\big] &\leq \frac{2}{M} e^{-\lambda} H_z^2 \sum_{N \geq 1} \frac{N^2 x^N}{N!} \\ &= \frac{2}{M}(x^2 + x) e^{x-\lambda} H_z^2 \\ &\leq 2 e^{2x} H_z^2 = 2 C_\rho^2 H_z^2.\end{aligned}$$

The coefficient is independent of z. Taking expectation over Z proves the result. □

Lemma B.7 (Lemma 3.5.18 of Giné and Nickl, 2021)**.** *The Bernstein size* $\rho_K(\cdot)$ *satisfies the following properties:*

- $\rho_K(f)$ *is nonincreasing in* K.

- *If* $\|f\|_\infty \le K$ *and* $\mathbb{E}[f(X)^2] \le R^2$*, then* $\rho_{2K}^2(f) \le 2R^2$.

Lemma B.8. *Suppose that* $\sup_{S\in\mathcal{F}_n}\|S\|_\infty = F < \infty$*. For any* $g'_S \in \mathcal{G}'_n$*, we have*

$$\rho^2_{F\nu(A)}(g'_S) \le 4F\nu(A)\widetilde{H}^2\left(S_n^*, \frac{S+S_n^*}{2}\right).$$

Proof. Since $0 \le S_n^* \le F$ and $0 \le (S+S_n^*)/2 \le F$, we have

$$|g'_S| \le \frac{1}{2}\int_A \left|S_n^*(x \mid Z) - \frac{S(x \mid Z)+S_n^*(x \mid Z)}{2}\right| d\nu(x) \le \frac{F\nu(A)}{2}.$$

Moreover, by Jensen's inequality,

$$\mathbb{E}[(g'_S)^2] \le \frac{\nu(A)}{4}\mathbb{E}_Z\left[\int_A \left(S_n^*(x \mid Z) - \frac{S(x \mid Z)+S_n^*(x \mid Z)}{2}\right)^2 d\nu(x)\right].$$

Since $0 \le S_n^*, (S(x \mid Z)+S_n^*(x \mid Z))/2 \le F$, we have

$$\begin{aligned}\left(S_n^* - \frac{S(x \mid Z)+S_n^*(x \mid Z)}{2}\right)^2 &= \left(\sqrt{S_n^*} - \sqrt{\frac{S(x \mid Z)+S_n^*(x \mid Z)}{2}}\right)^2\left(\sqrt{S_n^*} + \sqrt{\frac{S(x \mid Z)+S_n^*(x \mid Z)}{2}}\right)^2 \\ &\le 4F\left(\sqrt{S_n^*} - \sqrt{\frac{S(x \mid Z)+S_n^*(x \mid Z)}{2}}\right)^2.\end{aligned}$$

Hence,

$$\mathbb{E}[(g'_S)^2] \le 2F\nu(A)\widetilde{H}^2(S_n^*, \bar{S}).$$

By Lemma B.7, with $K = F\nu(A)/2$, we obtain

$$\rho^2_{F\nu(A)}(g'_S) \le 4F\nu(A)\widetilde{H}^2(S_n^*, \bar{S}).$$

□

Lemma B.9. *Suppose that* $\sup_{S\in\mathcal{F}_n}\|S\|_\infty \le F < \infty$*. For any* $0 \le S_L \le S_U \le F$*, define*

$$g'_L \coloneqq \frac{1}{2}\int_A \left(S_n^*(x \mid Z) - \frac{S_U(x \mid Z)+S_n^*(x \mid Z)}{2}\right) d\nu(x),$$

and

$$g'_U \coloneqq \frac{1}{2}\int_A \left(S_n^*(x \mid Z) - \frac{S_L(x \mid Z)+S_n^*(x \mid Z)}{2}\right) d\nu(x).$$

Then

$$\rho^2_{F\nu(A)}(g'_U - g'_L) \le 4F\nu(A)\widetilde{H}^2\left(\frac{S_U+S_n^*}{2}, \frac{S_L+S_n^*}{2}\right).$$

Proof. We have

$$g'_U - g'_L = \frac{1}{4}\int_A (S_U(x \mid Z) - S_L(x \mid Z))\, d\nu(x).$$

Hence,

$$0 \le g'_U - g'_L \le \frac{F\nu(A)}{4}.$$

Moreover, by Jensen's inequality and the inequality

$$\left(\frac{S_U+S_n^*}{2} - \frac{S_L+S_n^*}{2}\right)^2 \le 4F\left(\sqrt{\frac{S_U+S_n^*}{2}} - \sqrt{\frac{S_L+S_n^*}{2}}\right)^2,$$

we obtain

$$\mathbb{E}[(g'_U - g'_L)^2] \le 2F\nu(A)\widetilde{H}^2\left(\frac{S_U+S_n^*}{2}, \frac{S_L+S_n^*}{2}\right).$$

Therefore, Lemma B.7 gives

$$\rho^2_{F\nu(A)}(g'_U - g'_L) \le 4F\nu(A)\widetilde{H}^2\left(\frac{S_U+S_n^*}{2}, \frac{S_L+S_n^*}{2}\right).$$

□

B.3 Main proof

Proof. (Proof of Theorem 3.2) Put

$$\bar{S}_n \coloneqq \frac{\hat{S}_n + S_n^*}{2}.$$

For $S \in \mathcal{F}_n$, define

$$\bar{S} \coloneqq \frac{S + S_n^*}{2},$$

and let

$$g_S \coloneqq \int_{A_n^*(Z)} \frac{1}{2} \log\left(\frac{\bar{S}(x \mid Z)}{S_n^*(x \mid Z)}\right) dX(x \mid Z),$$

and

$$g'_S \coloneqq \frac{1}{2} \int_A \left(S_n^*(x \mid Z) - \bar{S}(x \mid Z)\right) d\nu(x).$$

By Lemma 3.1, we have

$$\widetilde{H}^2(\bar{S}_n, S_n^*) \le (P_n - P)g_{\hat{S}_n} + (P_n - P)g'_{\hat{S}_n} + 2(c_0 + 1)\widetilde{H}(\bar{S}_n, S_n^*)\widetilde{H}(S_n^*, S_0).$$

Denote $r_n \coloneqq \widetilde{H}(\bar{S}_n, S_n^*)$, $h_n \coloneqq \widetilde{H}(S_n^*, S_0)$.

First consider the case in which

$$\max\{(P_n - P)g_{\hat{S}_n}, (P_n - P)g'_{\hat{S}_n}, 2(c_0 + 1)r_n h_n\} = 2(c_0 + 1)r_n h_n.$$

Then, we obtain

$$r_n^2 \le 6(c_0 + 1)r_n h_n.$$

If $r_n > 0$, this implies

$$r_n^2 \le 36(c_0 + 1)^2 h_n^2.$$

If $r_n = 0$, the desired bound is immediate. Therefore, by the triangle inequality and Lemma B.1,

$$\begin{aligned}
\widetilde{H}^2(\hat{S}_n, S_0) &\le 2\left\{\widetilde{H}^2(\hat{S}_n, S_n^*) + \widetilde{H}^2(S_n^*, S_0)\right\} \\
&\le 2\left\{16\widetilde{H}^2(\bar{S}_n, S_n^*) + \widetilde{H}^2(S_n^*, S_0)\right\} \\
&\le 1154(c_0 + 1)^2 \widetilde{H}^2(S_n^*, S_0).
\end{aligned}$$

Hence the theorem follows in this case.

It remains to consider the cases in which either I or II is the largest term. If

$$\max\{(P_n - P)g_{\hat{S}_n}, (P_n - P)g'_{\hat{S}_n}, 2(c_0 + 1)r_n h_n\} = (P_n - P)g_{\hat{S}_n},$$

then

$$r_n^2 \le 3(P_n - P)g_{\hat{S}_n}.$$

If

$$\max\{(P_n - P)g_{\hat{S}_n}, (P_n - P)g'_{\hat{S}_n}, 2(c_0 + 1)r_n h_n\} = (P_n - P)g'_{\hat{S}_n},$$

then

$$r_n^2 \le 3(P_n - P)g'_{\hat{S}_n}.$$

Let

$$E_\delta = \left\{\widetilde{H}^2(\hat{S}_n, S_0) > 1154(c_0 + 1)^2(\delta^2 + h_n^2)\right\}, \qquad D(S) = \widetilde{H}^2\left(\frac{S + S_n^*}{2}, S_n^*\right).$$

By the triangle inequality and Lemma B.1,

$$\widetilde{H}^2(\hat{S}_n, S_0) \le 32r_n^2 + 2h_n^2.$$

Consequently, on E_δ we have $r_n^2 > \delta^2/16$. Moreover, the preceding case analysis shows that the approximation term cannot be the largest term on E_δ. Hence, on this event, at least one of

$$(P_n - P)g_{\hat{S}_n} \ge r_n^2/3, \qquad (P_n - P)g'_{\hat{S}_n} \ge r_n^2/3$$

holds.

Since $D(S) \le F\nu(A)$ for every $S \in \mathcal{F}_n$, the event E_δ is empty if $\delta^2 \ge 16F\nu(A)$. Otherwise, set

$$T = \min\left\{t \in \{0,1,2,\dots\} : 2^{2t-2}\delta^2 \ge F\nu(A)\right\}.$$

Partitioning the range $D(S) > \delta^2/16$ into the intervals $(2^{2t-4}\delta^2, 2^{2t-2}\delta^2]$, $t = 0,\dots,T$, and using $2^{2t-4}/3 \ge 2^{2t-6}$, we obtain

$$\mathbb{P}(E_\delta) \le \sum_{t=0}^{T} \mathbb{P}\left(\sup_{\substack{S\in\mathcal{F}_n:\\ D(S)\le 2^{2t-2}\delta^2}} \sqrt{n}(P_n - P)g_S \ge 2^{2t-6}\sqrt{n}\,\delta^2 \right) + \sum_{t=0}^{T} \mathbb{P}\left(\sup_{\substack{S\in\mathcal{F}_n:\\ D(S)\le 2^{2t-2}\delta^2}} \sqrt{n}(P_n - P)g'_S \ge 2^{2t-6}\sqrt{n}\,\delta^2 \right). \tag{19}$$

We first control the terms involving g_S. For each $t = 0,\dots,T$, set

$$a_t \coloneqq 2^{2t-6}\sqrt{n}\,\delta^2.$$

This quantity will be used as the parameter a in Theorem B.2. We verify the conditions of the theorem with

$$a = a_t, \qquad K = 1, \qquad R = 2^{t+2}C_\rho\delta.$$

Let

$$\mathcal{G}_{n,t} \coloneqq \left\{ g_S : S \in \mathcal{F}_n,\ \widetilde{H}^2\left(\frac{S+S_n^*}{2}, S_n^*\right) \le 2^{2t-2}\delta^2 \right\}.$$

We first check the Bernstein-size condition. By Lemma B.5, for every $g_S \in \mathcal{G}_{n,t}$,

$$\rho_1^2(g_S) \le 2^6 C_\rho^2 \widetilde{H}^2\left(\frac{S+S_n^*}{2}, S_n^*\right) \le 2^{2t+4}C_\rho^2\delta^2.$$

Hence, we have

$$\sup_{g\in\mathcal{G}_{n,t}} \rho_1(g) \le 2^{t+2}C_\rho\delta = R,$$

which verifies condition (14).

Next we check condition (15). Since

$$\frac{\sqrt{n}R^2}{K} = \sqrt{n}\left(2^{t+2}C_\rho\delta\right)^2 = 2^{2t+4}C_\rho^2\sqrt{n}\,\delta^2,$$

we have

$$a_t = 2^{2t-6}\sqrt{n}\,\delta^2 \le C_1 2^{2t+4}C_\rho^2\sqrt{n}\,\delta^2 = C_1\frac{\sqrt{n}R^2}{K},$$

provided $C_1 C_\rho^2 \ge 2^{-10}$. Since $C_\rho^2 \ge 1$, this is satisfied, for instance, by taking $C_1 = 15$. Thus condition (15) holds.

We next verify condition (16). We need $a_t \le 8\sqrt{n}R$. Substituting the definitions of a_t and R, this is equivalent to

$$2^{t-11}\delta \le C_\rho.$$

It is enough to show that

$$2^{2t-4}\delta^2 \le 2^{18}C_\rho^2.$$

Since $t \le T$, by the definition of T, we have

$$2^{2t-4}\delta^2 \le 2^{2(T-1)-2}\delta^2 \le F\nu(A).$$

Moreover, by enlarging c_0 if necessary, we may assume $C_\rho^2 \ge F\nu(A) \vee 1$. Therefore condition (16) holds after changing only universal constants.

It remains to verify the entropy condition (17). By Lemma B.6, if

$$\left\| \sqrt{\frac{S_U + S_n^*}{2}} - \sqrt{\frac{S_L + S_n^*}{2}} \right\|_{L^2(\nu\otimes P_Z)} \le u,$$

then

$$\rho_1(g_U - g_L) \le C_\rho u.$$

Consequently, for every $u > 0$,

$$\mathcal{N}_{B,1}\left(C_\rho u, \mathcal{G}_{n,t}, P\right) \le N_{2,B}\left(u, \bar{\mathcal{F}}_n^{1/2}(2^{t-1}\delta), \nu \otimes P_Z\right).$$

Equivalently,

$$\log \mathcal{N}_{B,1}\left(v, \mathcal{G}_{n,t}, P\right) \le \log N_{2,B}\left(\frac{v}{C_\rho}, \bar{\mathcal{F}}_n^{1/2}(2^{t-1}\delta), \nu \otimes P_Z\right).$$

We now bound the entropy integral in condition (17). Since $a_t/(2^6\sqrt{n}) = 2^{2t-12}\delta^2$ and $R = 2^{t+2}C_\rho\delta$, we have

$$\begin{aligned}
&\int_{a_t/(2^6\sqrt{n})}^{R} \sqrt{\log \mathcal{N}_{B,1}\left(u, \mathcal{G}_{n,t}, P\right)}\, du \\
&\le C_\rho \int_{2^{2t-12}\delta^2/C_\rho}^{2^{t+2}\delta} \sqrt{\log N_{2,B}\left(v, \bar{\mathcal{F}}_n^{1/2}(2^{t-1}\delta), \nu \otimes P_Z\right)}\, dv,
\end{aligned}$$

where we used the change of variables $v = u/C_\rho$. By the definition of J_B and the assumption $\Psi(s) \ge J_B(s, \bar{\mathcal{F}}_n^{1/2}(s))$, we have

$$\begin{aligned}
&C_\rho \int_{2^{2t-12}\delta^2/C_\rho}^{2^{t+2}\delta} \sqrt{\log N_{2,B}\left(v, \bar{\mathcal{F}}_n^{1/2}(2^{t-1}\delta), \nu \otimes P_Z\right)}\, dv, \\
&\le C_\rho \int_{2^{2t-12}\delta^2/C_\rho}^{2^{t+2}\delta} \sqrt{\log N_{2,B}\left(v, \bar{\mathcal{F}}_n^{1/2}(2^{t+2}\delta), \nu \otimes P_Z\right)}\, dv, \le C_\rho \Psi(2^{t+2}\delta).
\end{aligned}$$

Since $\Psi(s)/s^2$ is nonincreasing and $\delta \ge \delta_n$, the critical inequality $\sqrt{n}\,\delta_n^2 \ge c\Psi(\delta_n)$ implies $\sqrt{n}\,\delta^2 \ge c\Psi(\delta)$. Moreover, again using the monotonicity of $\Psi(s)/s^2$, we have $\Psi(2^{t+2}\delta) \le 2^{2t+4}\Psi(\delta)$. Hence,

$$C_\rho \Psi(2^{t+2}\delta) \le \frac{C_\rho}{c} 2^{2t+4}\sqrt{n}\,\delta^2.$$

Then, we obtain

$$\int_{a_t/(2^6\sqrt{n})}^{R} \sqrt{\log \mathcal{N}_{B,1}\left(u, \mathcal{G}_{n,t}, P\right)}\, du \le \frac{2^{10}C_\rho}{c} a_t.$$

Furthermore, by the same critical inequality and the fact that $\Psi(\delta) \ge J_B(\delta, \bar{\mathcal{F}}_n^{1/2}(\delta)) \ge \delta$, enlarging c if necessary gives

$$R = 2^{t+2}C_\rho\delta \le \frac{2^{10}C_\rho}{c} a_t.$$

Therefore, with $C_0 = c/(2^{10}C_\rho)$,

$$a_t \ge C_0\left(\int_{a_t/(2^6\sqrt{n})}^{R} \sqrt{\log \mathcal{N}_{B,1}\left(u, \mathcal{G}_{n,t}, P\right)}\, du \vee R\right),$$

which verifies condition (17).

Finally, condition (18) is satisfied by choosing $C = C_0/4$ sufficiently large so that

$$C_0^2 \ge C^2(C_1 + 1).$$

Thus all the conditions of Theorem B.2 are verified for the class $\mathcal{G}_{n,t}$ with

$$a = a_t, \qquad K = 1, \qquad R = 2^{t+2}C_\rho\delta.$$

Therefore,

$$\mathbb{P}\left(\sqrt{n}\sup_{g \in \mathcal{G}_{n,t}} |(P_n - P)g| \ge a_t\right) \le C\exp\left(-\frac{a_t^2}{C^2(C_1+1)R^2}\right).$$

Since

$$\frac{a_t^2}{R^2} = \frac{2^{4t-12}n\delta^4}{2^{2t+4}C_\rho^2\delta^2} = \frac{2^{2t-16}n\delta^2}{C_\rho^2},$$

we obtain

$$\mathbb{P}\left(\sup_{g\in\mathcal{G}_{n,t}}(P_n-P)g\geq 2^{2t-6}\delta^2\right)\leq C\exp\left(-\frac{2^{2t}n\delta^2}{2^{20}C^2C_\rho^2}\right),$$

where the universal constants have been enlarged if necessary. Summing over t gives

$$\sum_{t=0}^{T}\mathbb{P}\left(\sup_{g\in\mathcal{G}_{n,t}}(P_n-P)g\geq 2^{2t-6}\delta^2\right)\leq\sum_{t=0}^{\infty}\mathbb{P}\left(\sup_{g\in\mathcal{G}_{n,t}}(P_n-P)g\geq 2^{2t-6}\delta^2\right)\leq c\exp\left(-\frac{n\delta^2}{c^2}\right).$$

We next control the terms involving g'_S. For each $t=0,\dots,T$, we use the same value

$$a_t=2^{2t-6}\sqrt{n}\,\delta^2$$

as the parameter a in Theorem B.2. Define

$$\mathcal{G}'_{n,t}\coloneqq\left\{g'_S: S\in\mathcal{F}_n,\ \widetilde{H}^2\left(\frac{S+S_n^*}{2},S_n^*\right)\leq 2^{2t-2}\delta^2\right\}.$$

We verify the conditions of Theorem B.2 with

$$a=a_t,\qquad K=F\nu(A),\qquad R=2^t\sqrt{F\nu(A)}\,\delta.$$

First, by Lemma B.8, for every $g'_S\in\mathcal{G}'_{n,t}$,

$$\rho^2_{F\nu(A)}(g'_S)\leq 4F\nu(A)\widetilde{H}^2\left(\frac{S+S_n^*}{2},S_n^*\right)\leq F\nu(A)2^{2t}\delta^2=R^2$$

This verifies condition (14).

Next, we verify condition (15). Since

$$\frac{\sqrt{n}R^2}{K}=\frac{\sqrt{n}\left(2^t\sqrt{F\nu(A)}\,\delta\right)^2}{F\nu(A)}=2^{2t}\sqrt{n}\,\delta^2,$$

we have

$$a_t=2^{2t-6}\sqrt{n}\,\delta^2\leq C_1 2^{2t}\sqrt{n}\,\delta^2=C_1\frac{\sqrt{n}R^2}{K}$$

as long as $C_1\geq 2^{-6}$. Thus condition (15) holds, for example with $C_1=15$.

We now check condition (16). We need

$$2^{t-9}\delta\leq\sqrt{F\nu(A)}.$$

It is enough to show

$$2^{2t-18}\delta^2\leq F\nu(A).$$

By the definition of T, for $t\leq T$, the quantity $2^{2t-4}\delta^2$ is bounded by $F\nu(A)$. Therefore condition (16) follows.

We next verify the entropy condition (17). By Lemma B.9, if

$$\left\|\sqrt{\frac{S_U+S_n^*}{2}}-\sqrt{\frac{S_L+S_n^*}{2}}\right\|_{L^2(\nu\otimes P_Z)}\leq u,$$

then

$$\rho_{F\nu(A)}(g'_U-g'_L)\leq\sqrt{2F\nu(A)}\,u.$$

Hence, for every $u>0$,

$$\mathcal{N}_{B,F\nu(A)}\left(\sqrt{2F\nu(A)}\,u,\mathcal{G}'_{n,t},P\right)\leq N_{2,B}\left(u,\bar{\mathcal{F}}_n^{1/2}(2^{t-1}\delta),\nu\otimes P_Z\right).$$

Equivalently,

$$\log\mathcal{N}_{B,F\nu(A)}\left(v,\mathcal{G}'_{n,t},P\right)\leq\log N_{2,B}\left(\frac{v}{\sqrt{2F\nu(A)}},\bar{\mathcal{F}}_n^{1/2}(2^{t-1}\delta),\nu\otimes P_Z\right).$$

Since $a_t/(2^6\sqrt{n}) = 2^{2t-12}\delta^2$ and $R = 2^t\sqrt{F\nu(A)}\,\delta$, we obtain, by the change of variables $v = u/\sqrt{2F\nu(A)}$,

$$\begin{aligned}
&\int_{a_t/(2^6\sqrt{n})}^{R} \sqrt{\log \mathcal{N}_{B,F\nu(A)}\left(u, \mathcal{G}'_{n,t}, P\right)}\, du \\
&\leq \sqrt{2F\nu(A)} \int_{2^{2t-12}\delta^2/\sqrt{2F\nu(A)}}^{2^{t-1/2}\delta} \sqrt{\log N_{2,B}\left(v, \bar{\mathcal{F}}_n^{1/2}(2^{t-1}\delta), \nu \otimes P_Z\right)}\, dv \\
&\leq C_\rho \int_{2^{2t-12}\delta^2/C_\rho}^{2^{t-1/2}\delta} \sqrt{\log N_{2,B}\left(v, \bar{\mathcal{F}}_n^{1/2}(2^{t-1}\delta), \nu \otimes P_Z\right)}\, dv.
\end{aligned}$$

Here, we used that we may take $C_\rho \geq \sqrt{2F\nu(A)}$ by setting c_0 sufficiently large if necessary. Remember the definition of J_B and the condition $\Psi(s) \geq J_B(s, \bar{\mathcal{F}}_n(s))$, we have

$$\begin{aligned}
&C_\rho \int_{2^{2t-12}\delta^2/C_\rho}^{2^{t-1/2}\delta} \sqrt{\log N_{2,B}\left(v, \bar{\mathcal{F}}_n^{1/2}(2^{t-1}\delta), \nu \otimes P_Z\right)}\, dv, \\
&\leq C_\rho \int_{2^{2t-12}\delta^2/C_\rho}^{2^{t+2}\delta} \sqrt{\log N_{2,B}\left(v, \bar{\mathcal{F}}_n^{1/2}(2^{t+2}\delta), \nu \otimes P_Z\right)}\, dv, \leq C_\rho \Psi(2^{t+2}\delta).
\end{aligned}$$

We note that, by setting c_0 sufficiently large if necessary, $2^{t+2}C_\rho \geq 2^t\sqrt{F\nu(A)}$. Therefore, by repeating almost the same argument as controlling the terms involving g_S, we have

$$\sum_{t=0}^{T} \mathbb{P}\left(\sup_{g\in\mathcal{G}_{n,t}} (P_n - P)g \geq 2^{2t-6}\delta^2\right) \leq \sum_{t=0}^{\infty} \mathbb{P}\left(\sup_{g\in\mathcal{G}_{n,t}} (P_n - P)g \geq 2^{2t-6}\delta^2\right) \leq c\exp\left(-\frac{n\delta^2}{c^2}\right).$$

Combining the two bounds in (19), we obtain

$$\mathbb{P}\left(\widetilde{H}^2(\hat{S}_n, S_0) > 1154(c_0+1)^2\left(\delta^2 + \widetilde{H}^2(S_n^*, S_0)\right)\right) \leq c\exp\left(-\frac{n\delta^2}{c^2}\right).$$

This completes the proof. □

Proof. (Proof of Corollary 3.3) For any nonnegative random variable Y, we have

$$\mathbb{E}[Y] = \int_0^\infty \mathbb{P}(Y \geq t)\, dt.$$

Using this identity together with Theorem 3.2, we obtain

$$\begin{aligned}
\mathbb{E}\left[\widetilde{H}^2(\hat{S}_n, S_0)\right] &= \int_0^\infty \mathbb{P}\left(\widetilde{H}^2(\hat{S}_n, S_0) \geq t\right) dt \\
&= \int_0^{1154(c_0+1)^2(\delta^2+\widetilde{H}^2(S_n^*,S_0))} \mathbb{P}\left(\widetilde{H}^2(\hat{S}_n, S_0) \geq t\right) dt \\
&\quad + \int_{1154(c_0+1)^2(\delta^2+\widetilde{H}^2(S_n^*,S_0))}^\infty \mathbb{P}\left(\widetilde{H}^2(\hat{S}_n, S_0) \geq t\right) dt \\
&\leq 1154(c_0+1)^2\left(\delta^2 + \widetilde{H}^2(S_n^*, S_0)\right) \\
&\quad + 1154(c_0+1)^2 \int_{\delta^2}^\infty \mathbb{P}\left(\widetilde{H}^2(\hat{S}_n, S_0) \geq 1154(c_0+1)^2\left(t + \widetilde{H}^2(S_n^*, S_0)\right)\right) dt \\
&\leq 1154(c_0+1)^2\left(\delta^2 + \widetilde{H}^2(S_n^*, S_0) + c\int_0^\infty \exp\left(-\frac{nt}{c^2}\right) dt\right) \\
&= 1154(c_0+1)^2\left(\delta^2 + \widetilde{H}^2(S_n^*, S_0) + \frac{c^3}{n}\right).
\end{aligned}$$

□

C Proof of Theorem 4.1

Lemma C.1 gives an upper bound on the integral of the reciprocal of a conditional intensity function, which appears in the proof of Theorem 4.1.

Lemma C.1. *Suppose that there exist $\alpha \geq 0$ and a constant $C > 0$ such that, for a conditional intensity function S,*

$$\int_{[0,1]^{d_z}} \int_A \mathbb{1}\left(S(x \mid z) \leq t\right) \, d\nu(x) dP_Z(z) \leq Ct^\alpha, \quad \textit{for all } t \in (0, \infty).$$

Let $H \in [0,1]$. Then

$$\int_{[0,1]^{d_z}} \int_A \frac{1}{S(x \mid z)} \mathbb{1}\left(S(x \mid z) \geq H\right) \, d\nu(x) dP_Z(z) \leq \begin{cases} (C \vee \nu(A)) \dfrac{H^{\alpha-1}}{1-\alpha}, & \textit{if } \alpha \in [0,1), \\ (C \vee 1)\left(\nu(A) - \log(H)\right), & \textit{if } \alpha \geq 1. \end{cases}$$

Proof. For any nonnegative function h and random variable Y, we have

$$\mathbb{E}[h(Y)] = \int_0^\infty \mathbb{P}_Y(h(Y) \geq u) \, du.$$

Combining this identity with Tonelli's theorem gives

$$\begin{aligned} &\int_{[0,1]^{d_z}} \int_A \frac{1}{S(x \mid z)} \mathbb{1}\left(S(x \mid z) \geq H\right) \, d\nu(x) dP_Z(z) = \int_A \int_0^\infty \mathbb{P}_Z\left(\frac{1}{S(x \mid Z)} \mathbb{1}\left(S(x \mid Z) \geq H\right) \geq u\right) du \, d\nu(x) \\ &\leq \int_A \int_0^\infty \mathbb{P}_Z\left(H \leq S(x \mid Z) \leq \frac{1}{u}\right) du \, d\nu(x) = \int_A \int_0^{1/H} \mathbb{P}_Z\left(H \leq S(x \mid Z) \leq \frac{1}{u}\right) du \, d\nu(x) \\ &\leq \int_0^{1/H} \int_A \mathbb{P}_Z\left(S(x \mid Z) \leq \frac{1}{u}\right) d\nu(x) \, du. \end{aligned}$$

If $\alpha = 0$, then $\mathbb{P}_Z(S(x \mid Z) \leq t) \leq 1$ for any $t \geq 0$, and hence

$$\int_0^{1/H} \int_A \mathbb{P}_Z\left(S(x \mid Z) \leq \frac{1}{u}\right) d\nu(x) \, du \leq \frac{\nu(A)}{H}.$$

If $0 < \alpha < 1$, then the assumption gives

$$\int_0^{1/H} \int_A \mathbb{P}_Z\left(S(x \mid Z) \leq \frac{1}{u}\right) d\nu(x) \, du \leq C \int_0^{1/H} u^{-\alpha} \, du = \frac{CH^{\alpha-1}}{1-\alpha}.$$

If $\alpha \geq 1$ and $C \vee 1 \leq H^{-1}$, then using

$$\int_A \mathbb{P}_Z(S(x \mid Z) \leq t) \, d\nu(x) \leq Ct$$

for any $t \in [0,1]$, we obtain

$$\begin{aligned} &\int_0^{1/H} \int_A \mathbb{P}_Z\left(S(x \mid Z) \leq \frac{1}{u}\right) d\nu(x) \, du \\ &\leq \int_0^{C \vee 1} \int_A \mathbb{P}_Z\left(S(x \mid Z) \leq \frac{1}{u}\right) d\nu(x) \, du + \int_{C \vee 1}^{1/H} \int_A \mathbb{P}_Z\left(S(x \mid Z) \leq \frac{1}{u}\right) d\nu(x) \, du \\ &\leq \int_0^{C \vee 1} \nu(A) \, du + C \int_{C \vee 1}^{1/H} \frac{1}{u} \, du \\ &= (C \vee 1)\nu(A) + C\left(-\log(H) - \log(C \vee 1)\right). \end{aligned}$$

On the other hand, if $\alpha \geq 1$ and $C \vee 1 \geq H^{-1}$, then

$$\int_0^{1/H} \int_A \mathbb{P}_Z\left(S(x \mid Z) \leq \frac{1}{u}\right) d\nu(x) \, du \leq \int_0^{C \vee 1} \nu(A) \, du = (C \vee 1)\nu(A).$$

Combining these bounds and simplifying the resulting expression yields the claim. □

Now, we prove Theorem 4.1.

Proof. (Proof of Theorem 4.1) From the condition $\|S - S_0\|_\infty \le C_1/M$, we have, for any $x \in A$ and $z \in [0,1]^{d_z}$,

$$S(x \mid z) \ge S_0(x \mid z) - \frac{C_1}{M}.$$

Combining this with

$$\inf_{z\in[0,1]^{d_z}} \inf_{x\in A} S(x \mid z) \ge \frac{1}{M},$$

we obtain, for any $x \in A$ and $z \in [0,1]^{d_z}$,

$$S(x \mid z) \ge \left(S_0(x \mid z) - \frac{C_1}{M}\right) \vee \frac{1}{M} \ge \frac{S_0(x \mid z)}{C_1+1} \vee \frac{1}{M}.$$

In the last inequality, we used the fact that if $S_0(x \mid z) \ge (C_1+1)/M$, then

$$S_0(x \mid z) - \frac{C_1}{M} \ge S_0(x \mid z)\left(1 - \frac{C_1}{C_1+1}\right) = \frac{S_0(x \mid z)}{C_1+1}.$$

Using this bound, we obtain

$$\begin{aligned}\left(\sqrt{S(x \mid z)} - \sqrt{S_0(x \mid z)}\right)^2 &\le \frac{\left(S(x \mid z) - S_0(x \mid z)\right)^2}{\left(\sqrt{S(x \mid z)} + \sqrt{S_0(x \mid z)}\right)^2} \\ &\le \frac{C_1^2}{M}\mathbb{1}\left(S_0(x \mid z) \le \frac{C_1+1}{M}\right) \\ &\quad + \frac{C_1^2}{M^2}\cdot\frac{C_1+1}{(\sqrt{C_1+1}+1)^2}\cdot\frac{1}{S_0(x \mid z)}\mathbb{1}\left(S_0(x \mid z) \ge \frac{C_1+1}{M}\right).\end{aligned}$$

Integrating both sides with respect to ν and taking expectation with respect to Z gives

$$\begin{aligned}\widetilde{H}^2(S, S_0) \le{}& \frac{C_1^2}{M}\int_{[0,1]^{d_z}}\int_A \mathbb{1}\left(S_0(x \mid z) \le \frac{C_1+1}{M}\right) d\nu(x) dP_Z(z) \\ &+ \frac{C_1^2}{M^2}\int_{[0,1]^{d_z}}\int_A \frac{\mathbb{1}\left(S_0(x \mid z) \ge (C_1+1)/M\right)}{S_0(x \mid z)} d\nu(x) dP_Z(z).\end{aligned}$$

For the first term, the assumption gives

$$\frac{C_1^2}{M}\int_{[0,1]^{d_z}}\int_A \mathbb{1}\left(S_0(x \mid z) \le \frac{C_1+1}{M}\right) d\nu(x) dP_Z(z) \le \frac{C_1^2 C}{M}\left(\frac{C_1+1}{M}\right)^{\alpha\wedge 1}.$$

For the second term, applying Lemma C.1 with $H = (C_1+1)/M$, we obtain, for $\alpha < 1$,

$$\frac{C_1^2}{M^2}\int_{[0,1]^{d_z}}\int_A \frac{\mathbb{1}\left(S_0(x \mid z) \ge (C_1+1)/M\right)}{S_0(x \mid z)} d\nu(x) dP_Z(z) \le \frac{(C_1+1)^2}{M^{1+\alpha}}\cdot\frac{C \vee \nu(A)}{1-\alpha},$$

while for $\alpha \ge 1$,

$$\frac{C_1^2}{M^2}\int_{[0,1]^{d_z}}\int_A \frac{\mathbb{1}\left(S_0(x \mid z) \ge (C_1+1)/M\right)}{S_0(x \mid z)} d\nu(x) dP_Z(z) \le \frac{C_1^2}{M^2}(C \vee 1)\left(\nu(A) + \log(M)\right).$$

Combining these bounds and simplifying the expression yields the conclusion. □

D Proof of Theorem 4.2

In this section, we provide the proof of Theorem 4.2. We first state several definitions and lemmas used to evaluate the approximation error $\widetilde{H}^2(S_n^*, S_0)$ and the bracketing entropy integral in (5), which appear in Theorem 3.2 and Corollary 3.3.

Theorem D.1 (Theorem 2 of Kohler and Langer, 2021). *Let $d \in \mathbb{N}$, and let $f \in \mathcal{C}^\beta([0,1]^d, K)$ for some $\beta, K \in \mathbb{R}$. Let $M \in \mathbb{N}$ be sufficiently large, and suppose that*

$$M \geq 2 \quad \textit{and} \quad M^{2\beta} \geq c_1(1 \vee K)^{4(\lfloor\beta\rfloor+1)}$$

hold for some sufficiently large constant $c_1 \geq 1$.

Let $L, r \in \mathbb{N}$ and $B \in \mathbb{R}$ satisfy

$$\text{(i)} \quad L \geq 5 + \left\lceil \log_4(M^{2\beta}) \right\rceil \left(\lceil \log_2(\lfloor\beta\rfloor \vee d) \rceil + 1 \right),$$

$$\text{(ii)} \quad r \geq 64 \cdot 2^d \binom{d+\lfloor\beta\rfloor}{d} d^2(\lfloor\beta\rfloor+1) M^d,$$

$$\text{(iii)} \quad B \geq M^{4\beta+4} \vee 2Ke^d \vee 4.$$

Then there exists a neural network $\hat{f} \in \mathcal{F}_{\mathrm{id}}(L, r, B)$ such that

$$\|f - \hat{f}\|_\infty \leq c_2(1 \vee K)^{4(\lfloor\beta\rfloor+1)} M^{-2\beta},$$

Here, id *denotes the identity map, and c_2 is a constant depending only on β, K, and d.*

Remark D.1. In the original statement of Theorem 2 of Kohler and Langer (2021), an upper bound on the network parameters is not explicitly stated. In this study, however, such a bound is needed to evaluate the covering number of the DNN class. Therefore, following the proof of Theorem 2 of Kohler and Langer (2021), we calculate and include the parameter bound B. For simplicity, we also take the original parameter a to be $a = 1$.

Lemma D.2 (Lemma 3 of Schmidt-Hieber, 2020). *Let $f = g_q \circ \cdots \circ g_0$, and suppose that, for each $g_i = (g_{ij})_{j=1,\ldots,d^{(i+1)}}$, we have $g_{ij} \in \mathcal{C}^{\beta_i}([a_i, b_i]^{t_i}, K_i)$. Assume further that $K_i \geq 1$. For $i = 1, \ldots, q-1$, define*

$$h_0(x) := \frac{g_0(x)}{2K_0} + \frac{1}{2}, \qquad h_i(x) := \frac{g_i(2K_{i-1}x - K_{i-1})}{2K_i} + \frac{1}{2}, \qquad h_q(x) := g_q(2K_{q-1}x - K_{q-1}).$$

Then

$$f = g_q \circ \cdots \circ g_0 = h_q \circ \cdots \circ h_0.$$

Moreover, h_{0j} takes values in $[0,1]$, and

$$h_{0j} \in \mathcal{C}^{\beta_0}([0,1]^{t_0}, 1), \qquad h_{ij} \in \mathcal{C}^{\beta_i}([0,1]^{t_i}, (2K_{i-1})^{\beta_i}),$$

for $i = 1, \ldots, q-1$, while

$$h_{q1} \in \mathcal{C}^{\beta_q}([0,1]^{t_q}, K_q(2K_{q-1})^{\beta_q}).$$

Furthermore, for any $\tilde{h}_i = (\tilde{h}_{ij})_{j=1,\ldots,d^{(i+1)}}$ such that $\tilde{h}_{ij} : [0,1]^{t_i} \to [0,1]$ for $i = 0, \ldots, q-1$, we have

$$\left\| h_q \circ \cdots \circ h_0 - \tilde{h}_q \circ \cdots \circ \tilde{h}_0 \right\|_\infty \leq K_q \prod_{l=0}^{q-1} (2K_l)^{\beta_{l+1}} \sum_{i=0}^{q} \left\| h_i - \tilde{h}_i \right\|_\infty^{\prod_{k=i+1}^{q} (\beta_k \wedge 1)}.$$

Definition D.1 (Covering number with respect to the uniform norm). Let $\mathcal{F}$ be a class of functions, and let $\delta > 0$. We define $N(\delta, \mathcal{F}, \|\cdot\|_\infty)$ as the smallest integer N for which there exists a collection of functions $\{f_j\}_{j=1}^N$ such that

$$\sup_{f \in \mathcal{F}} \min_{j=1,\ldots,N} \|f - f_j\|_\infty \leq \delta.$$

If no such N exists, we set $N(\delta, \mathcal{F}, \|\cdot\|_\infty) = \infty$. We call $N(\delta, \mathcal{F}, \|\cdot\|_\infty)$ the δ-covering number of $\mathcal{F}$ with respect to the uniform norm.

Lemma D.3 (Lemma A.7 of Yara and Terada, 2026). *Let $\mathcal{F}$ be a class of functions on $\mathcal{X}$, and let μ be a finite measure on $\mathcal{X}$. Then, for any $\delta > 0$,*

$$N_{p,B}(\delta, \mathcal{F}, \mu) \leq N\left(\frac{\delta}{2\mu(\mathcal{X})^{1/p}}, \mathcal{F}, \|\cdot\|_\infty \right).$$

Lemma D.4 (Lemma 3 of Suzuki, 2019). *For any $\delta > 0$,*

$$\log N\left(\delta, \mathcal{F}_{\mathrm{id}}(L-1, r, B, \infty, s), \|\cdot\|_\infty\right) \leq 2sL \log((B \vee 1)(r+1)) + s \log(\delta^{-1} L).$$

In particular, for fully connected (non-sparse) neural networks,

$$\log N\left(\delta, \mathcal{F}_{\mathrm{id}}(L-1, r, B, \infty), \|\cdot\|_\infty\right) \leq 2r^2L^2 \log((B \vee 1)(r+1)) + r^2 L \log(\delta^{-1} L),$$

where id *denotes the identity map.*

Remark D.2. Note that the statement of Lemma D.4 differs slightly from that in Suzuki (2019), because the convention for counting the number of layers in our definition of DNNs differs by one from that used in Suzuki (2019).

We are now ready to derive the convergence rate of the DNN-based maximum likelihood estimator when the true intensity function belongs to $\mathcal{G}_\alpha(q, \mathsf{t}, \beta, K, C)$.

Proof. (Proof of Theorem4.2) For each $i = 0, \dots, q$, let

$$M_i = \left\lceil n^{\frac{1}{2\left((1+(\alpha\wedge 1))\beta_i^* + t_i\right)}} \right\rceil, \qquad \xi_n = \max_{i=0,\dots,q} n^{-\frac{\beta_i^*}{(1+(\alpha\wedge 1))\beta_i^* + t_i}} \qquad \phi_n = \max_{i=0,\dots,q} n^{-\frac{(1+(\alpha\wedge 1))\beta_i^*}{(1+(\alpha\wedge 1))\beta_i^* + t_i}}.$$

We write the hypothesis space as $\mathcal{F}_n \coloneqq \mathcal{F}_\sigma(L, r, B, F)$.

We first evaluate the approximation error $\widetilde{H}^2(S_n^*, S_0)$. Since $S_0 \in \mathcal{G}(q, \mathsf{d}, \mathsf{t}, \beta, K)$, we can write

$$S_0 = g_q \circ \cdots \circ g_0, \qquad g_{ij} \in \mathcal{C}^{\beta_i}([a_i, b_i]^{t_i}, K_i),$$

for $i = 0, \dots, q$ and $j = 1, \dots, d^{(i+1)}$. Without loss of generality, we may assume that $K_i \geq 1$. By the first part of Lemma D.2, we can rewrite

$$S_0 = h_q \circ \cdots \circ h_0,$$

where

$$h_{0j} \in \mathcal{C}^{\beta_0}([0,1]^{t_0}, 1), \qquad h_{ij} \in \mathcal{C}^{\beta_i}([0,1]^{t_i}, (2K_{i-1})^{\beta_i}), \quad i = 1, \dots, q-1,$$

and

$$h_{q1} \in \mathcal{C}^{\beta_q}([0,1]^{t_q}, K_q(2K_{q-1})^{\beta_q}).$$

Hence, by Theorem D.1, for sufficiently large n, there exist DNNs $\tilde{h}_{ij} \in \mathcal{F}_{\mathrm{id}}(L_i, r_i, B_i, \infty)$ such that

$$\begin{aligned}
&\|\tilde{h}_{0j} - h_{0j}\|_\infty \leq c_2 M_0^{-2\beta_0}, \\
&\|\tilde{h}_{ij} - h_{ij}\|_\infty \leq c_2 (2K_{i-1})^{4\beta_i(\lfloor\beta_i\rfloor+1)} M_i^{-2\beta_i}, \quad i = 1, \dots, q-1, \\
&\|\tilde{h}_{q1} - h_{q1}\|_\infty \leq c_2 \left(K_q(2K_{q-1})^{\beta_q}\right)^{4(\lfloor\beta_q\rfloor+1)} M_q^{-2\beta_q},
\end{aligned} \tag{20}$$

where id denotes the identity map. Here,

$$\begin{aligned}
L_i &= 5 + \left\lceil \log_4(M_i^{2\beta_i}) \right\rceil \left(\lceil \log_2(\lfloor\beta_i\rfloor \vee t_i)\rceil + 1\right), \\
r_i &= 64 \cdot 2^{t_i} \binom{t_i + \lfloor\beta_i\rfloor}{t_i} t_i^2 (\lfloor\beta_i\rfloor + 1) M_i^{t_i}, \\
B_i &= M_i^{4\beta_i+4} \vee 2Ke^{t_i}.
\end{aligned}$$

To apply the second part of Lemma D.2, we restrict the range of $\tilde{h}_{ij}$, $i = 0, \dots, q-1$, to $[0,1]$. Define

$$h_{ij}^* = \sigma(\tilde{h}_{ij}) - \sigma(\tilde{h}_{ij} - 1).$$

Then

$$h_{ij}^* = (\tilde{h}_{ij} \vee 0) \wedge 1, \qquad \sigma(h_{ij}^*) = h_{ij}^*.$$

This clipping operation is represented by one additional hidden layer with two active neurons, output weights $(1, -1)$, and biases $(0, 1)$. Since $r_i \geq 2$ and $B_i \geq 1$, zero padding gives

$$h_{ij}^* \in \mathcal{F}_{\mathrm{id}}(L_i + 1, r_i, B_i, \infty).$$

Because h_{ij} takes values in $[0,1]$,

$$\|\sigma(h_{ij}^*) - h_{ij}\|_\infty \leq \|\tilde{h}_{ij} - h_{ij}\|_\infty. \tag{21}$$

Let $h_i^* = (h_{ij}^*)_{j=1,\dots,d^{(i+1)}}$. By computing the DNNs h_{ij}^* in parallel, we obtain $h_i^* \in \mathcal{F}_{\mathrm{id}}(L_i + 1, d^{(i+1)} r_i, B_i, \infty)$. Define

$$\widetilde{S}_n = \tilde{h}_{q1} \circ \sigma(h_{q-1}^*) \circ \cdots \circ \sigma(h_0^*).$$

Then, $\widetilde{S}_n \in \mathcal{F}_{\mathrm{id}}(\widetilde{L}, \tilde{r}, \widetilde{B}, \infty)$, where

$$\widetilde{L} \coloneqq 2q + 1 + \sum_{i=0}^{q} L_i, \qquad \tilde{r} \coloneqq \left(\max_{i=0,\dots,q-1} d^{(i+1)} r_i\right) \vee r_q, \qquad \widetilde{B} \coloneqq \max_{i=0,\dots,q} B_i.$$

By Lemma D.2, (20), and (21), we obtain

$$\begin{aligned}\|\widetilde{S}_n - S_0\|_\infty &\le K_q \prod_{l=0}^{q-1}(2K_l)^{\beta_{l+1}} \left(\sum_{i=0}^{q-1}\|\sigma(h_i^*) - h_i\|_\infty^{\prod_{k=i+1}^{q}(\beta_k\wedge 1)} + \|\tilde{h}_{q1} - h_{q1}\|_\infty\right) \\ &\le K_q \prod_{l=0}^{q-1}(2K_l)^{\beta_{l+1}} \left(\sum_{i=0}^{q-1}\|\tilde{h}_i - h_i\|_\infty^{\prod_{k=i+1}^{q}(\beta_k\wedge 1)} + \|\tilde{h}_{q1} - h_{q1}\|_\infty\right) \\ &\le C' \max_{i=0,\ldots,q} M_i^{-2\beta_i \prod_{k=i+1}^{q}(\beta_k\wedge 1)} \le C'\xi_n. \end{aligned} \tag{22}$$

Here, C' is a constant depending only on q, t, β, K, and $d^{(0)}, \ldots, d^{(q)}$.

For all sufficiently large n, we have $0 < \xi_n \le F$ and $\widetilde{B} \ge F \vee 1$. Define

$$S_n^* = \xi_n + \sigma(\widetilde{S}_n - \xi_n) - \sigma(\widetilde{S}_n - F) = (\widetilde{S}_n \vee \xi_n) \wedge F.$$

The additive constant is represented by a hidden neuron with zero incoming weights and bias -1, followed by output weight ξ_n. Together with the two clipping neurons, this adds one hidden layer and uses three active neurons. Since $\tilde{r} \ge 3$, zero padding preserves the prescribed width. The resulting output lies in $[\xi_n, F]$, so the prescribed final ReLU leaves it unchanged. Consequently,

$$S_n^* \in \mathcal{F}_\sigma(\widetilde{L} + 1, \tilde{r}, \widetilde{B}, F).$$

From the definition of L_i, there exists a constant E independent of M_i's such that

$$2q + 2 + \sum_{i=0}^{q} L_i = E + \sum_{i=0}^{q} \left(\lceil \log_2(t_i \vee \beta_i)\rceil + 1\right)\lceil 2\beta_i \log_4(M_i)\rceil.$$

Since $\lceil x\rceil < x + 1$, for sufficiently large n, we have

$$\widetilde{L} + 1 \le \sum_{i=0}^{q} 2\beta_i \left(\log_2(\beta_i \vee t_i) + 2\right)\log_4(M_i).$$

Therefore, by the choices of L, r, and B in the statement of the theorem, we have $S_n^* \in \mathcal{F}_n = \mathcal{F}_\sigma(L, r, B, F)$.

We now evaluate the approximation error. If $S_n^*(x \mid z) = \widetilde{S}_n(x \mid z) \wedge F$, then (22) and $F \ge K$ gives

$$|S_n^*(x \mid z) - S_0(x \mid z)| \le C'\xi_n.$$

If $S_n^*(x \mid z) = \xi_n$, then $\widetilde{S}_n(x \mid z) \le \xi_n$. Hence, by (22) and the nonnegativity of functions in $\mathcal{G}_\alpha(q, t, \beta, K, C)$, we obtain

$$0 \le S_0(x \mid z) \le (C' + 1)\xi_n.$$

Thus, in this case as well,

$$|S_n^*(x \mid z) - S_0(x \mid z)| \le (C' + 2)\xi_n.$$

Consequently,

$$\|S_n^* - S_0\|_\infty \le (C' + 2)\xi_n.$$

Applying Theorem 4.1 with $C_1 = C' + 2$, we obtain

$$\widetilde{H}^2(S_n^*, S_0) \lesssim \phi_n \log(n).$$

More explicitly, the hidden constant depends only on $q, t, \beta, K, \alpha, C, F$, and $d^{(0)}, \ldots, d^{(q)}$.

Moreover, for any $x \in A$ and $z \in [0,1]^{d_z}$,

$$\frac{S_0(x \mid z)}{S_n^*(x \mid z)} \le \frac{(C' + 2)\xi_n}{S_n^*(x \mid z)} + 1 \le C' + 3.$$

Hence, Assumption 3.1 holds with

$$c_0 \ge \sqrt{C' + 3}.$$

We next evaluate the bracketing entropy integral (5) in order to apply Corollary 3.3. For the above $S_n^* \in \mathcal{F}_n$, define

$$\bar{\mathcal{F}}_n^{1/2} := \left\{\sqrt{\frac{S + S_n^*}{2}} : S \in \mathcal{F}_n\right\}.$$

Since $\bar{\mathcal{F}}_n^{1/2}(\delta) \subset \bar{\mathcal{F}}_n^{1/2}$, for any $u > 0$,

$$N_{2,B}\left(u, \bar{\mathcal{F}}_n^{1/2}(\delta), \nu \otimes P_Z\right) \leq N_{2,B}\left(u, \bar{\mathcal{F}}_n^{1/2}, \nu \otimes P_Z\right).$$

Let $\{(f_j^L, f_j^U)\}_{j=1}^N$ be a minimal $2\sqrt{2\xi_n}\,u$-bracketing set for $\mathcal{F}_n$ with respect to the $L^2(\nu \otimes P_Z)$ norm. Since $S_n^* \geq \xi_n$, we have

$$\mathbb{E}_Z\left[\int_A \left(\sqrt{\frac{f_j^U(x \mid Z) + S_n^*(x \mid Z)}{2}} - \sqrt{\frac{f_j^L(x \mid Z) + S_n^*(x \mid Z)}{2}}\right)^2 d\nu(x)\right]$$
$$\leq \mathbb{E}_Z\left[\int_A \frac{\left(f_j^U(x \mid Z) - f_j^L(x \mid Z)\right)^2}{8S_n^*(x \mid Z)} d\nu(x)\right] \leq u^2.$$

Therefore, using Lemma D.3 together with the fact that ReLU is 1-Lipschitz, we obtain

$$\begin{aligned} N_{2,B}\left(u, \bar{\mathcal{F}}_n^{1/2}(\delta), \nu \otimes P_Z\right) &\leq N_{2,B}\left(2\sqrt{2\xi_n}\,u, \mathcal{F}_n, \nu \otimes P_Z\right) \\ &\leq N\left(\frac{\sqrt{2\xi_n}}{\sqrt{\nu(A)}}u, \mathcal{F}_n, \|\cdot\|_\infty\right) \\ &\leq N\left(\frac{\sqrt{2\xi_n}}{\sqrt{\nu(A)}}u, \mathcal{F}_{\mathrm{id}}(L, r, B, \infty), \|\cdot\|_\infty\right). \end{aligned}$$

By Lemma D.4,

$$\log N_{2,B}\left(u, \bar{\mathcal{F}}_n^{1/2}(\delta), \nu \otimes P_Z\right) \leq 2s(L+1)\log\left(\frac{(L+1)(r+1)(B \vee 1)\sqrt{\nu(A)}}{\sqrt{2\xi_n}}u^{-1}\right),$$

where s denotes the total number of parameters of the DNN.

Set

$$D \coloneqq \frac{(L+1)(r+1)(B \vee 1)\sqrt{\nu(A)}}{\sqrt{2\xi_n}}.$$

Then

$$\int_{\delta^2/(2^{16}C_\rho)}^{\delta} \sqrt{\log N_{2,B}\left(u, \bar{\mathcal{F}}_n^{1/2}(\delta), \nu \otimes P_Z\right)}\, du \leq \sqrt{2s(L+1)}\int_0^\delta \sqrt{\log\left(\frac{D}{u}\right)}\, du.$$

By the change of variables $x = \sqrt{\log(D/u)}$, we obtain

$$\int_0^\delta \sqrt{\log\left(\frac{D}{u}\right)}\, du = 2D\int_{\sqrt{\log(D/\delta)}}^\infty x^2 e^{-x^2}\, dx = \delta\sqrt{\log(D/\delta)} + D\int_{\sqrt{\log(D/\delta)}}^\infty e^{-x^2}\, dx.$$

For the second term, let Y be a centered normal random variable with variance $1/2$. Then its density is $\pi^{-1/2}e^{-x^2}$, and therefore,

$$\int_{\sqrt{\log(D/\delta)}}^\infty e^{-x^2}\, dx = \frac{\sqrt{\pi}}{2}\mathbb{P}\left(|Y| \geq \sqrt{\log(D/\delta)}\right).$$

Using the standard Gaussian tail bound,

$$\mathbb{P}\left(|Y| \geq a\right) \leq 2e^{-a^2}, \qquad a > 0,$$

with $a = \sqrt{\log(D/\delta)}$, we have

$$D\int_{\sqrt{\log(D/\delta)}}^\infty e^{-x^2}\, dx \leq D\frac{\sqrt{\pi}}{2} \cdot 2\exp\left\{-\log(D/\delta)\right\} = \sqrt{\pi}\,\delta.$$

Combining the above estimates yields

$$\int_0^\delta \sqrt{\log\left(\frac{D}{u}\right)}\, du \le \delta\left(\sqrt{\log(D/\delta)} + \sqrt{\pi}\right).$$

Therefore,

$$J_B(\delta, \bar{\mathcal{F}}_n^{1/2}(\delta)) \le \delta\sqrt{2s(L+1)}\left(\sqrt{\log(D/\delta)} + \sqrt{\pi}\right).$$

In Corollary 3.3, take

$$\Psi(\delta) = \delta\sqrt{2s(L+1)}\left(\sqrt{\log(D/\delta)} + \sqrt{\pi}\right).$$

To satisfy the critical inequality (6), set

$$\delta_n = \frac{\sqrt{2s(L+1)}\left(\sqrt{\log(\sqrt{n}D)} + \sqrt{\pi}\right)}{\sqrt{n}}.$$

Then Corollary 3.3 yields

$$\mathbb{E}\left[\widetilde{H}^2(\hat{S}_n, S_0)\right] \le 1154(1+c_0)^2\left(\frac{2s(L+1)\left(\sqrt{\log(\sqrt{n}D)} + \sqrt{\pi}\right)^2}{n} + \widetilde{H}^2(S_n^*, S_0) + \frac{1}{n}\right).$$

Combining this bound with

$$\widetilde{H}^2(S_n^*, S_0) \lesssim \phi_n \log(n)$$

and using $s \lesssim r^2 L$, $L \asymp \log(n)$, and $r \asymp \sqrt{n\phi_n}$, we obtain that there exists a constant $C' > 0$, depending only on $q, t, \beta, K, \alpha, C, c_0$, and $d^{(0)}, \ldots, d^{(q)}$, such that

$$\mathbb{E}\left[\widetilde{H}^2(\hat{S}_n, S_0)\right] \le C'\phi_n \log(n)^3.$$

□

E Proof of Theorem 4.3

Proof. Let v and V denote lower and upper bounds, respectively, of the density function p_Z of Z with respect to the Lebesgue measure. Moreover, we denote $\epsilon_n := \max_{i=0,\ldots,q} n^{-\beta_i^*/(\beta_i^*+t_i)}$ By Theorem 2.7 of Tsybakov (2008), it suffices to construct functions

$$S_0, \ldots, S_M \in \mathcal{G}(q, t, \beta, K)$$

for some integer $M > 1$ and some $\kappa > 0$ such that

$$\text{(i)} \quad \widetilde{H}^2(S_j, S_k) \ge \kappa\epsilon_n \quad \text{for all } 0 \le j < k \le M,$$

and

$$\text{(ii)} \quad \sum_{j=1}^M \mathrm{KL}(P_j \parallel P_0) \le \frac{M\log M}{9}.$$

Indeed, if such functions exist, then there exists a positive constant $c = c(\kappa, \gamma)$ such that

$$\inf_{\hat{S}_n} \sup_{S_0 \in \mathcal{G}(q,t,\beta,K)} \mathbb{E}\left[\widetilde{H}^2(\hat{S}_n, S_0)\right] \ge c\epsilon_n.$$

Here, the infimum is taken over all estimators, the expectation is taken with respect to the observed data, and P_j denotes the distribution of the data when the conditional intensity function is S_j.

We now construct functions $S_0, \ldots, S_M \in \mathcal{G}(q, t, \beta, K)$ satisfying (i) and (ii). Let

$$i^* \in \operatorname*{arg\,min}_{i=0,\ldots,q} \frac{\beta_i^*}{\beta_i^* + t_i}.$$

For notational convenience, write $\beta^* := \beta_{i^*}$, $\beta^{**} := \beta_{i^*}^*$, $t^* := t_{i^*}$. Choose a nonzero nonnegative function $\xi \in C_c^\infty((0,1))$. Multiplying it by a sufficiently small positive constant, we may assume that $\xi \in \mathcal{C}^{\beta^*}(\mathbb{R}, 1)$. Then

$\xi \in L^1(\mathbb{R})$, and $0 < \|\xi^B\|_1 < \infty$ for the positive exponent B defined below. Set $m_n \coloneqq \lfloor \rho n^{1/(\beta^{**}+t^*)} \rfloor$, $h_n \coloneqq 1/m_n$, where the constant $\rho \geq 1$ is chosen so that

$$nh_n^{\beta^{**}+t^*} \leq \frac{\log 2}{72V\|\xi^B\|_1^{t^*}}.$$

Here,

$$B \coloneqq \prod_{l=i^*+1}^{q} (\beta_l \wedge 1).$$

Let $\mathcal{U}_n \coloneqq \{(u_1, \ldots, u_{t^*}) : u_i \in \{0, h_n, 2h_n, \ldots, (m_n - 1)h_n\}\}$. For each $u = (u_1, \ldots, u_{t^*}) \in \mathcal{U}_n$, define

$$\psi_u(x_1, \ldots, x_{t^*}) \coloneqq h_n^{\beta^*} \prod_{j=1}^{t^*} \xi\left(\frac{x_j - u_j}{h_n}\right).$$

Since $\xi \in \mathcal{C}^{\beta^*}(\mathbb{R}, 1)$, for any multi-index γ satisfying $\|\gamma\|_1 < \beta^*$,

$$\|\partial^\gamma \psi_u\|_\infty = h_n^{\beta^*} \prod_{j=1}^{t^*} \frac{\|\xi^{(\gamma_j)}\|_\infty}{h_n^{\gamma_j}} \leq 1.$$

Moreover, for any multi-index γ with $\|\gamma\|_1 = \lfloor \beta^* \rfloor$, the triangle inequality and the fact that $\xi \in \mathcal{C}^{\beta^*}(\mathbb{R}, 1)$ imply

$$h_n^{\beta^* - \lfloor \beta^* \rfloor} \frac{\left|\prod_{j=1}^{t^*} \xi^{(\gamma_j)}\left(\frac{x_j - u_j}{h_n}\right) - \prod_{j=1}^{t^*} \xi^{(\gamma_j)}\left(\frac{y_j - u_j}{h_n}\right)\right|}{\max_{i=1,\ldots,t^*} |x_i - y_i|^{\beta^* - \lfloor \beta^* \rfloor}} \leq t^*.$$

Hence,

$$\psi_u \in \mathcal{C}^{\beta^*}([0,1]^{t^*}, C_{\beta^*, t^*}).$$

Here $C_{\beta^*, t^*} > 0$ is a fixed constant, independent of n, u, and w. For $w = (w_u)_{u \in \mathcal{U}_n} \in \{0,1\}^{|\mathcal{U}_n|}$, define

$$\varphi_w \coloneqq \sum_{u \in \mathcal{U}_n} w_u \psi_u.$$

By construction, for distinct $u, u' \in \mathcal{U}_n$, the functions ψ_u and $\psi_{u'}$ have disjoint supports. Therefore,

$$\varphi_w \in \mathcal{C}^{\beta^*}([0,1]^{t^*}, 2C_{\beta^*, t^*}).$$

For $i < i^*$, define

$$g_i(x) \coloneqq (x_1, \ldots, x_{d_i \wedge d_{i+1}}, 0, \ldots, 0)^\top.$$

For $i = i^*$, define

$$g_{i^*,w}(x) \coloneqq (\varphi_w(x_1, \ldots, x_{t^*}), 0, \ldots, 0)^\top.$$

For $i > i^*$, define

$$g_i(x) \coloneqq (x_1^{\beta_i \wedge 1}, 0, \ldots, 0)^\top.$$

Since $t_i \leq \min\{d_0, \ldots, d_{i-1}\}$, and since the functions ψ_u have disjoint supports, noting that $\beta^{**} = \beta^* B$, we define

$$\begin{aligned} f_w(x) &\coloneqq g_q \circ \cdots \circ g_{i^*+1} \circ g_{i^*,w} \circ g_{i^*-1} \circ \cdots \circ g_0(x) \\ &= \varphi_w^B(x_1, \ldots, x_{t^*}) \\ &= \sum_{u \in \mathcal{U}_n} w_u \psi_u^B(x_1, \ldots, x_{t^*}). \end{aligned}$$

Taking K sufficiently large ensures that $f_w \in \mathcal{G}(q, t, \beta, K)$. For any $u \in \mathcal{U}_n$, we have

$$\|\psi_u^B\|_1 = h_n^{\beta^* B} \prod_{j=1}^{t^*} \int \xi^B\left(\frac{x_j - u_j}{h_n}\right) dx_j = h_n^{\beta^{**}+t^*} \prod_{j=1}^{t^*} \int \xi^B(x_j)\, dx_j = h_n^{\beta^{**}+t^*} \|\xi^B\|_1^{t^*}.$$

Let

$$\mathrm{Ham}(w, w') \coloneqq \sum_{u \in \mathcal{U}_n} \mathbb{1}(w_u \neq w'_u)$$

denote the Hamming distance. Since the supports of ψ_u are pairwise disjoint,

$$\begin{aligned}
\widetilde{H}^2(f_w, f_{w'}) &= \mathbb{E}_Z\left[\frac{1}{2}\int_A \left(\sqrt{f_w(x \mid Z)} - \sqrt{f_{w'}(x \mid Z)}\right)^2 dx\right] \\
&= \mathbb{E}_Z\left[\frac{1}{2}\int_A \left(\sqrt{\sum_{u\in\mathcal{U}_n} w_u \psi_u^B(x \mid Z)} - \sqrt{\sum_{u\in\mathcal{U}_n} w'_u \psi_u^B(x \mid Z)}\right)^2 dx\right] \\
&= \mathbb{E}_Z\left[\frac{1}{2}\int_A \sum_{u\in\mathcal{U}_n} \left(\sqrt{w_u \psi_u^B(x \mid Z)} - \sqrt{w'_u \psi_u^B(x \mid Z)}\right)^2 dx\right] \\
&= \mathbb{E}_Z\left[\frac{1}{2}\int_A \sum_{u\in\mathcal{U}_n} \mathbb{1}(w_u \neq w'_u)\psi_u^B(x \mid Z)dx\right] \\
&\geq \frac{v}{2}\mathrm{Ham}(w, w')h_n^{\beta^{**}+t^*}\|\xi^B\|_1^{t^*}.
\end{aligned}$$

Here, among the variables $x_1, \ldots, x_{t^*}$, we have relabeled those corresponding to the support of the Poisson process as x and those corresponding to the covariates as Z. By Lemma 2.9 of Tsybakov (2008) and $m_n^{t^*} = |\mathcal{U}_n|$, there exists a subset $\mathcal{W} \subset \{0,1\}^{m_n^{t^*}}$ containing the all-one vector such that $|\mathcal{W}| \geq 2^{m_n^{t^*}/8}$, and, for any distinct $w, w' \in \mathcal{W}$,

$$\mathrm{Ham}(w, w') \geq \frac{m_n^{t^*}}{8}.$$

Therefore, taking

$$\kappa = \frac{v^{1/2}\|\xi^B\|_1^{t^*/2}}{4\rho^{\beta^{**}/2}},$$

we obtain, for any distinct $w, w' \in \mathcal{W}$,

$$\widetilde{H}^2(f_w, f_{w'}) \geq \frac{v}{16}h_n^{\beta^{**}}\|\xi^B\|_1^{t^*} \geq \kappa^2\epsilon_n.$$

We next bound the KL divergence for the data distributions induced by the intensity functions constructed above. Let P_w denote the population distribution when the intensity function is f_w, and let f_1, P_1 denote the conditional intensity function and population distribution corresponding to the all-one vector $(1, \ldots, 1)^\top$, respectively. Using the disjointness of the supports of ψ_u, the KL-divergence formula for Poisson processes (Leskelä, 2024, Corollary 5.2), and the definitions of h_n and ρ, we obtain

$$\begin{aligned}
\mathrm{KL}\left(P_w^{1:n} \parallel P_1^{1:n}\right) &= n\mathbb{E}_Z\left[\int\left\{f_w(x \mid Z)\log\left(\frac{f_w(x \mid Z)}{f_1(x \mid Z)}\right) + f_1(x \mid Z) - f_w(x \mid Z)\right\}dx\right] \\
&\leq n\mathbb{E}_Z\left[\int\frac{(f_w(x \mid Z) - f_1(x \mid Z))^2}{f_1(x \mid Z)}dx\right] \\
&\leq nVm_n^{t^*}\|\psi_u^B\|_1 \\
&\leq nVm_n^{t^*}h_n^{\beta^{**}+t^*}\|\xi^B\|_1^{t^*} \\
&\leq \frac{\log|\mathcal{W}|}{9}.
\end{aligned}$$

Here, for a probability distribution P, the notation $P^{1:n}$ denotes the joint distribution of n i.i.d. observations from P. This proves the claim. □

F Proof of Theorem 4.4

We first supply uniform local extensions, then construct the reference network and apply the oracle inequality. Constants in this appendix do not depend on the member S of the fixed Hölder class or on the approximation index. Throughout, $q = \lceil\beta\rceil - 1$ and $\theta = \beta - q \in (0, 1]$, including the Lipschitz endpoint.

F.1 Charts, covariate boundaries, and local extensions

The required geometric choices exist for every fixed compact smooth embedded manifold without boundary. At a point of $\mathcal{M}$, projection onto its tangent space has invertible derivative on the tangent space. The inverse function theorem gives a chart; shrinking its image gives an open ball. Compactness gives a finite cover by smaller chart domains whose closures lie inside such charts. For completeness, the partition can be chosen with smooth ambient extensions. Choose nonnegative ambient bumps b_i supported in neighborhoods whose intersections with $\mathcal{M}$ lie in V_i, with $\sum_i b_i > 0$ on $\mathcal{M}$. On a neighborhood O of $\mathcal{M}$ where this sum is positive, take an ambient smooth cutoff ζ supported compactly in O, equal to one near $\mathcal{M}$. The functions

$$R_i = \zeta \frac{b_i}{\sum_j b_j} \quad \text{on } O, \qquad R_i = 0 \quad \text{outside } O$$

are smooth, take values in $[0,1]$, and restrict to the required ρ_i on $\mathcal{M}$. Refining the finite cover if necessary ensures $\operatorname{supp}\rho_i \Subset V_i$. In the product domain we use $\rho_i(x,z) = \rho_i(x)$.

Lemma F.1 (Uniform local Euclidean extensions)**.** *Under Definition 4.2, there exist fixed affine bijections $A_i : [0,1]^k \to Q_i$, where Q_i are rectangular boxes, and functions $f_{i,S} \in C^{q,\theta}([0,1]^k)$ such that*

$$f_{i,S}\big(A_i^{-1}(P_i x, z)\big) = S(x,z) \quad \textit{if } x \in \operatorname{supp}\rho_i,\ z \in [0,1]^{d_z}, \qquad \max_i \|f_{i,S}\|_{C^{q,\theta}([0,1]^k)} \le C_{\mathrm{ext}} K.$$

The constant is uniform over $S \in \mathcal{H}^{\beta}_{\mathcal{M}}(K)$.

Proof. Write $u_i(v,z) = S(P_i^{-1}v, z)$ and $E_i = P_i(\operatorname{supp}\rho_i) \subset \Omega_i$. The covariate faces prevent $E_i \times [0,1]^{d_z}$ from being an interior subset of an open product chart, so we first extend in those coordinates. Let $J_q = q+1$ and choose the unique coefficients $a_1, \dots, a_{J_q}$ satisfying the Vandermonde system

$$\sum_{j=1}^{J_q} a_j (-j)^r = 1, \qquad r = 0, \dots, q. \tag{23}$$

Set $\delta = 1/(2J_q)$. In one covariate coordinate, with all remaining coordinates denoted by w, use

$$(Tu)(w,t) = \begin{cases} \sum_{j=1}^{J_q} a_j u(w, -jt), & -\delta < t < 0, \\ u(w,t), & 0 \le t \le 1, \\ \sum_{j=1}^{J_q} a_j u(w, 1 - j(t-1)), & 1 < t < 1+\delta. \end{cases}$$

Every argument on the right belongs to $[0,1]$. Equation (23) matches derivatives of order r in t at both endpoints. It also matches every mixed derivative of total order at most q, because differentiating the other coordinates leaves the same identities. On each of the three pieces, the chain rule bounds all derivative suprema and highest-order Hölder seminorms by a constant depending only on q, θ times the original norm. For points in different pieces, split the line segment at the intervening faces. The derivatives agree at those faces, and $\sum_{\ell=1}^{N} h_\ell^\theta \le N^{1-\theta}(\sum_{\ell=1}^{N} h_\ell)^\theta$ bounds the resulting seminorm for every $0 < \theta \le 1$. At $\theta = 1$ the right-hand side is simply the sum of the segment lengths, so the same estimate is Lipschitz and has no singular endpoint constant. The domain is convex, so the segment stays in the domain. Applying this construction successively to the d_z coordinates gives $\bar{u}_i$ on $W_i = \Omega_i \times (-\delta, 1+\delta)^{d_z}$, with

$$\bar{u}_i = u_i \text{ on } \Omega_i \times [0,1]^{d_z}, \qquad \|\bar{u}_i\|_{C^{q,\theta}(W_i)} \le C_{\beta,d_z} K.$$

The same segment argument across at most $2d_z$ faces controls corners. For $\beta = 1$, $q = 0$ and $a_1 = 1$: this is ordinary reflection of a Lipschitz function, and only continuity is matched at each face. When $d_z = 0$, this step is the identity.

Choose fixed smooth cutoffs $\chi_i \in C_c^\infty(\Omega_i)$ and $\eta \in C_c^\infty((-\delta, 1+\delta)^{d_z})$, equal to one near E_i and $[0,1]^{d_z}$, respectively; omit η when $d_z = 0$. Define $F_{i,S} = \chi_i \eta \bar{u}_i$ on W_i and zero outside. Leibniz's rule and $[fg]_\theta \le \|f\|_\infty [g]_\theta + \|g\|_\infty [f]_\theta$ give the $C^{q,\theta}$ bound on W_i. Lower-order derivatives are Lipschitz on this bounded convex domain, with bounds supplied by the next derivatives, and hence have bounded θ-seminorms. The product vanishes in a fixed neighborhood of ∂W_i. Its zero extension is consequently $C^{q,\theta}$: for a nonzero derivative value inside and a point outside W_i, the distance is at least the fixed distance of the cutoff support from the boundary. This bounds the cross-boundary seminorm also at $\theta = 1$; no derivative of order $q+1$ of the target is required. Therefore

$$\|F_{i,S}\|_{C^{q,\theta}(\mathbb{R}^k)} \le C_i K, \qquad F_{i,S}(P_i x, z) = S(x,z) \quad (x \in \operatorname{supp}\rho_i).$$

Take a fixed box Q_i containing $E_i \times [0,1]^{d_z}$ in its interior and a fixed affine parametrization A_i of that box. Then $f_{i,S} = F_{i,S} \circ A_i$ has the asserted properties by the chain rule. Taking the maximum of the finitely many constants proves uniformity. The possibly signed extension is an approximation target; clipping below imposes the intensity bounds. No general extension theorem is used. □

F.2 A reference network with intrinsic width

We use Theorem D.1 in its asymptotic form: for fixed dimension d, smoothness $s > 0$, and Hölder radius, an error $O(m^{-2s})$ is attainable with depth $O(\log m)$, width $O(m^d)$, and parameters $O(m^{4s+4})$. The constants are uniform on that Hölder ball. In fixed dimension, the maximum-based norm and Euclidean distance in Definition 4.2 control the sum-based norm and maximum distance used in that theorem up to a fixed constant, which is absorbed into the Hölder radius. At integer smoothness this is exactly the $C^{q,1}$ class used in the imported approximation result; no smoothing or loss of smoothness is needed.

We record the elementary network operations used below. For fixed $U > 0$,

$$\mathrm{clip}_{[0,U]}(v) = \sigma(v) - \sigma(v - U), \qquad \mathrm{clip}_{[\varepsilon,U]}(v) = \varepsilon + \sigma(v - \varepsilon) - \sigma(v - U).$$

An additive constant is represented by a hidden neuron with zero incoming weights and a negative bias. Thus the absence of an output bias in (4.1) causes no restriction. To compose subnetworks without multiplying two dense weight matrices, propagate a scalar linear output v as the two nonnegative channels $\sigma(v), \sigma(-v)$ and recover v in the next affine map. Parallel branches are padded using identity maps on nonnegative hidden activations. Each scalar connection adds a fixed number of layers and channels. Padding equalizes branch depths at their maximum, and the finite number of successive stages keeps the total depth of order $\log m$ below. Widths and parameter bounds change by fixed factors only.

Multiplication on $[0, F] \times [0, 1]$ can be implemented with error ε, constant width, depth $O(\log(1/\varepsilon))$, and parameters bounded by a constant depending only on F. Here is an explicit construction, following the squaring argument of Yarotsky (2017). Put

$$T(t) = 2\sigma(t) - 4\sigma(t - 1/2) + 2\sigma(t - 1), \qquad Q_R(t) = t - \sum_{\ell=1}^{R} 4^{-\ell} T^{\circ\ell}(t).$$

On $[0, 1]$, Q_R is the linear interpolant of t^2 on the dyadic mesh of spacing 2^{-R}, as follows by induction on its midpoint values. Hence $\|Q_R - t^2\|_\infty \le 4^{-R}$. A network propagating $t, T^{\circ\ell}(t)$, and the accumulated partial sum has constant width, $O(R)$ layers, and bounded parameters. For $a = h/F$ and $b = v$, define

$$\mathrm{Mult}_{F,\varepsilon}(h, v) = \frac{F}{2}\{4Q_R((a + b)/2) - Q_R(a) - Q_R(b)\}.$$

The error from hv is at most $3F4^{-R}$, so $R \ge \max\{1, \lceil \log_4(3F/\varepsilon)\rceil\}$ suffices. Signed channels as above implement its possibly signed output.

Lemma F.2 (Uniform manifold approximation)**.** *For all sufficiently large integers m and every $S \in \mathcal{H}^\beta_{\mathcal{M}}(K)$ with $0 \le S \le F$, there is a ReLU network S^*_m with*

$$m^{-2\beta} \le S^*_m \le F \text{ on } \mathbb{R}^D, \qquad \|S^*_m - S\|_{\infty,K_0} \le C_A m^{-2\beta},$$

and depth at most $C_L \log m$, width at most $C_r m^k$, and parameter bound at most $C_B m^{4\beta+4}$. All constants are uniform over S.

Proof. Set $\varepsilon = m^{-2\beta}$. Lemma F.1 and the Euclidean approximation theorem give networks g_i approximating $f_{i,S}$ to error $C\varepsilon$, with depth $O(\log m)$, width $O(m^k)$, and parameter bound $O(m^{4\beta+4})$. The map $y = (x, z) \mapsto A_i^{-1}(P_i x, z)$ is affine on the ambient space. Its input dimension is fixed, so absorbing it into the first affine layer changes the parameter bound by a fixed factor. Set

$$h_i(y) = \mathrm{clip}_{[0,F]}\{g_i(A_i^{-1}(P_i x, z))\}.$$

Then $0 \le h_i \le F$ everywhere and $|h_i(y) - S(y)| \le C\varepsilon$ whenever $x \in \mathrm{supp}\, \rho_i$. No approximation outside that support is needed.

Fix a noninteger $\gamma > \beta D/k$ and a bounded ambient box containing K_0. Approximate the fixed smooth ambient extensions $R_i(x)$ on this box in dimension D, after a fixed affine rescaling, using $m_\rho = \lceil m^{\beta/\gamma}\rceil$. Clip to $[0, 1]$ to obtain $\widehat{\rho}_i$ with error $C\varepsilon$ on K_0 and budgets

$$L_\rho \lesssim \log m, \qquad r_\rho \lesssim m^{\beta D/\gamma} \le m^k, \qquad B_\rho \lesssim m^{4\beta+4\beta/\gamma} \le m^{4\beta+4}.$$

The last inequality uses $\beta/\gamma < k/D \le 1$. Now form

$$\widetilde{S}_m(y) = \sum_{i=1}^{J} \mathrm{Mult}_{F,\varepsilon}(h_i(y), \widehat{\rho}_i(y)).$$

Since the partition is applied exactly once, for $y \in K_0$,

$$|\widetilde{S}_m(y) - S(y)| \le J\varepsilon + \sum_i |h_i(y)|\,|\widehat{\rho}_i(y) - \rho_i(y)| + \sum_i \rho_i(y)|h_i(y) - S(y)| \le C\varepsilon.$$

The middle sum is controlled by global clipping. In the last sum, a nonzero partition weight guarantees that the local approximation holds. Finite parallelization, the signed-channel connections, multiplication, and summation have the claimed budgets. Finally set $S_m^* = \mathrm{clip}_{[\varepsilon,F]}(\widetilde{S}_m)$, for m sufficiently large that $\varepsilon \le F$. For $S(y) \in [0,F]$, clipping increases the approximation error by at most ε. The final output is nonnegative, so applying the prescribed output ReLU leaves it unchanged. This proves all assertions, including the global output bound and uniformity. □

F.3 Main proof

Proof of Theorem 4.4. Set $\xi_n = n^{-\beta/(a\beta+k)}$ and take $m_n = \lceil n^{1/[2(a\beta+k)]} \rceil$ in Lemma F.2. For the final clipping, use ξ_n in place of $m_n^{-2\beta}$. Since $m_n^{-2\beta} \le \xi_n \le 2^{2\beta} m_n^{-2\beta}$, the resulting network, denoted by S_n^*, satisfies

$$S_n^* \ge \xi_n, \qquad \|S_n^* - S_0\|_{\infty,K_0} \le C_A \xi_n.$$

Its depth, width, and parameter bound are dominated by L, r, B when the implicit lower constants in (i)–(iii) are chosen sufficiently large. Zero-neuron padding and identity layers on nonnegative activations embed it into the exact architecture without changing its output or parameter orders. Thus $S_n^* \in \mathcal{F}_\sigma(L, r, B, F)$ uniformly over the asserted class.

Let $\mu = \nu \otimes P_Z$ and use the notation a, k, ϕ_n' from Theorem 4.4. Since $S_0/S_n^* \le 1 + |S_0 - S_n^*|/S_n^* \le 1 + C_A$, Assumption 3.1 holds with a constant c_0 independent of n. Apply Theorem 4.1 with $M = \xi_n^{-1}$ and $C_1 = \max\{1, C_A\}$. Its lower-threshold condition holds for all sufficiently large n, and it gives

$$\widetilde{H}^2(S_n^*, S_0) \lesssim \xi_n^{1+(\alpha\wedge 1)}(1 + \log(1/\xi_n)) \lesssim \phi_n' \log n.$$

For completeness, consider the entropy term. Every network in $\mathcal{F}_n$ takes values in $[0,F]$. Bracket endpoints may therefore be clipped to this interval without increasing their width. Since $S_n^* \ge \xi_n$, for $u \ge v \ge 0$,

$$\left| \sqrt{\frac{u + S_n^*}{2}} - \sqrt{\frac{v + S_n^*}{2}} \right| \le \frac{|u - v|}{2\sqrt{2\xi_n}}.$$

All covering norms below are taken on a fixed bounded ambient box containing K_0. An affine normalization to $[0,1]^D$ changes the input-layer parameter bound by a fixed factor, so Lemma D.4 applies with fixed-domain constants. Combining its bound with covering-to-bracketing conversion and the preceding inequality yields, for $0 < t \le 1$,

$$\log N_{2,B}\left(t, \bar{\mathcal{F}}_n^{1/2}(\delta), \mu\right) \le C s_n (L+1) \log(D_n/t),$$

where the total number of network parameters is $s_n = (L-1)r^2 + (D+L+1)r$, and one may choose

$$D_n = e \vee \left\{ C(L+1)(r+1)(B \vee 1)\sqrt{\nu(\mathcal{M})/\xi_n} \right\}.$$

Here and below constants may depend on the fixed input dimension and the fixed domain. The architecture assumptions imply

$$s_n \lesssim r^2 L \lesssim n\phi_n' \log n, \qquad \log D_n \lesssim \log n.$$

As in the entropy calculation for the composition theorem, an admissible majorant is

$$\Psi_n(\delta) = C\delta\sqrt{s_n(L+1)} \left\{ \sqrt{\log\left(\frac{D_n}{\delta \wedge 1}\right)} + 1 \right\}, \qquad \delta > 0.$$

For radii greater than one, extend the localized class by the same defining formula and use monotonicity of bracketing numbers in the bracket width. Integrating the bound with t replaced by $t \wedge 1$ from zero to δ verifies this majorant on the whole positive axis. Its ratio to δ^2 is nonincreasing. For a sufficiently large fixed C_0, the radius $\delta_n = C_0\{s_n(L+1)(\log n)/n\}^{1/2}$ satisfies the critical inequality for all sufficiently large n. In particular, $\delta_n^2 \lesssim \phi_n' \log^3 n$ and $\delta_n \to 0$. The constants and the threshold for n are uniform in S_0. Corollary 3.3 therefore yields

$$\sup_{S_0 \in \mathcal{H}^\beta_{\mathcal{M},\alpha}(K,F,C)} \mathbb{E}[\widetilde{H}^2(\hat{S}_n, S_0)] \lesssim \delta_n^2 + \phi_n' \log n + n^{-1} \lesssim \phi_n' \log^3 n.$$

□

G Supplementary Simulation Studies

G.1 Generating intensities

We used Lebesgue measure on the Euclidean supports, arc length on the unit circle, and surface area on the unit sphere. In the formulas below, c is a scaling constant chosen separately for each intensity, support, and covariate dimension to give $\mathbb{E}_Z \int S_0(x \mid Z)\, d\nu(x) \approx 30$. We calculated c by numerical integration with resolution 1024 and 128 uniform covariate draws, using seed 314159. When $d_z = 0$, the covariate average reduces to a single integral. Covariate coordinates absent from a design are set to zero.

For the compositional design, we set

$$\begin{aligned} S_0(x \mid z) &= c\left[0.2 + \{0.6 + 0.4\sin(2\pi h_1(x,z))\}^2 + 0.5\exp\{-10(h_2(x,z) - 0.5)^2\}\right], \\ h_1(x,z) &= x_1 + 0.3z_1 + 0.2z_2, \qquad h_2(x,z) = x_2 + 0.5\sin(2\pi z_1). \end{aligned} \tag{24}$$

On $[0,1]$, we replace x_2 by $x_1/2$. Although the input dimension varies, the intensity is formed by composing functions of a small number of arguments. Only z_1 and z_2 enter the formula, leaving three irrelevant covariates when $d_z = 5$.

The second design is

$$S_0(x \mid z) = c\left[10^{-4} + \{1 - \exp(-18D^2(x,z))\}^2\right], \tag{25}$$

where

$$D^2(x,z) = \begin{cases} (x_1 - 0.25 - 0.5z_1)^2, & d_x = 1, \\ (x_1 - 0.25 - 0.5z_1)^2 + \{x_2 - 0.5 - 0.25\sin(2\pi z_2)\}^2, & d_x = 2. \end{cases}$$

The intensity is small near $D(x,z) = 0$, and the center of this region moves with z. The additive constant gives a strictly positive minimum of $10^{-4}c$ in every experiment.

On the unit circle, parameterized by angle θ, we use

$$S_0(\theta \mid z) = c\left[\{1 + 0.5\sin(3\theta + z_1) + 0.3\cos(2\theta - z_2)\}_+ + 10^{-4}\right]. \tag{26}$$

On the unit sphere, with colatitude θ and longitude ϕ, we use

$$S_0(\theta, \phi \mid z) = c\left[\{1 + 0.5\sin(2\theta)\cos(3\phi + z_1) + 0.3\sin(\phi + z_2)\}_+ + 10^{-4}\right]. \tag{27}$$

The expressions inside the positive-part operators are bounded below by 0.2, so the operators do not alter these intensities. The term $\sin(\phi + z_2)$ in (27) has no common limit at either pole. Any choice of values at the poles defines the same intensity measure, but the resulting function is not globally continuous. The sphere example therefore concerns estimation of a bounded measurable intensity outside the smooth manifold class used in the rate theorem.

We simulated point patterns by thinning uniform candidate locations. For each covariate vector, we evaluated the intensity on a grid with resolution 1024 and set an envelope B to 1.25 times the grid maximum, with a small additive safeguard. We drew a Poisson number of candidates with mean $B\nu(A)$ and retained a candidate at x with probability $\min\{S_0(x \mid z)/B, 1\}$. This numerical envelope introduces a possible simulation approximation if the intensity between grid points exceeds B.

G.2 DNN and kernel estimators

For a candidate intensity S, the neural fitting criterion is

$$\ell_n(S) = \frac{1}{n}\sum_{i=1}^{n}\left\{\int S(x \mid Z_i)\, d\nu(x) - \sum_{x \in X_i} \log S(x \mid Z_i)\right\}.$$

The two output functions are $u \mapsto \log(1 + e^u) + 10^{-8}$ and $u \mapsto u_+ + 10^{-8}$. We initialized hidden layers by Kaiming initialization and output weights near zero, with the output bias chosen to give approximately 30 expected events on the support. AdamW was used with learning rate 10^{-3}, weight decay 10^{-4}, and gradient-norm clipping at 5. Each update used 64 patterns and 128 newly sampled integration points per pattern. The epoch limits were 150 for Euclidean supports and 300 for manifold supports. We stopped after 15 epochs without improvement in validation loss and restored the best validation fit. The random 80%/20% split was made at the pattern level; the networks were not subsequently refitted on the validation patterns. Network training and prediction used single precision. We did not impose hard parameter bounds or an upper bound on the fitted intensity.

All hidden layers have common width W_n, with depth and width given by

$$L_n = \left\lceil 3\frac{\log n}{\log 1000} \right\rceil, \qquad W_n = 8\left\lceil \frac{64(n/1000)^{(1-\gamma_{\mathrm{arch}})/2}}{8} \right\rceil. \tag{28}$$

For the compositional and moving-valley designs, respectively, we set

$$\gamma_{\mathrm{arch}} = \frac{2}{2 + \min(d_x + d_z, 4)}, \qquad \gamma_{\mathrm{arch}} = \frac{4}{4 + d_x + d_z}.$$

For mesh-based integration on a manifold, we used $\gamma_{\mathrm{arch}} = 2/(2 + d_{\mathrm{amb}} + d_z)$, where $d_{\mathrm{amb}} = 2$ for the circle and 3 for the sphere. For integration with the known manifold measure, we replaced d_{amb} by $d_{\mathcal{M}} = 1$ or 2. In both versions, the networks take ambient coordinates as inputs. These architecture choices were specified using the generating design rather than selected from the data.

To construct an integration mesh, we combined 4,096 independent uniform support samples with a reproducible subset of at most 4,096 event locations. On the circle, convex-hull edges define a polygon; on the sphere, triangular convex-hull faces define a polyhedral surface. Integration weights are geometric edge lengths or planar face areas, rather than event frequencies. The known-measure version instead uses uniform samples on the true support and its volume. The comparison between these versions consequently changes both the integration measure and the network architecture.

For Euclidean supports, the conditional kernel estimator is

$$\widehat{S}(x \mid z) = \frac{\sum_{i=1}^{n} K_{h_z}(z - Z_i) \sum_{u \in X_i} K^A_{h_x}(x, u)}{\max\{\sum_{i=1}^{n} K_{h_z}(z - Z_i), 10^{-6}\}},$$

where K_{h_z} is a Gaussian covariate kernel and $K^A_{h_x}$ is a Gaussian event kernel normalized for the support boundary. This estimator follows the conditional-kernel approach to Poisson intensity estimation (Klutchnikoff and Massiot, 2016). The pooled estimator uses equal weights for all patterns. We truncated Gaussian contributions beyond four bandwidths. Five-fold Poisson cross-validation used pilot bandwidths proportional to $N_{\mathrm{events,train}}^{-1/(4+d_x)}$ for event coordinates and $n_{\mathrm{train}}^{-1/(4+d_z)}$ for covariates, followed by a five-value coarse search and a three-value refinement. To limit the cost of bandwidth selection, this calculation used at most 512 patterns and 128 events per pattern. The final kernel fit used all n patterns.

On manifold supports, we used a geodesic kernel estimator motivated by Ward et al. (2023), with bandwidth selected by five-fold Poisson cross-validation over a geodesic bandwidth grid. The estimator uses the known geometry and does not include local or global shape corrections.

G.3 Loss evaluation and Monte Carlo summaries

For each repetition, we generated one dataset and one set of evaluation points shared by all methods, with separate random seeds for fitting. The loss in (5.1) was averaged over 64 independent uniform covariate vectors when $d_z > 0$. For Euclidean supports, we used 512 integration points generated by scrambled Sobol quadrature when available. On the circle, we used 64 arc-length-weighted points. On the sphere, eight colatitude subdivisions and 16 longitude subdivisions formed triangles; their projected centroids were the evaluation locations and their planar areas the weights. Neural predictions were calculated in single precision and converted to double precision for loss evaluation; kernel predictions and loss accumulation used double precision.

The supplementary tables give the mean, median, standard deviation, minimum, and maximum of the 20 repetitions for each method and configuration. MCSEs are $\mathrm{SD}/\sqrt{20}$ and describe variation across repetitions. They do not include systematic quadrature error; the simulation losses were evaluated at a single integration resolution.

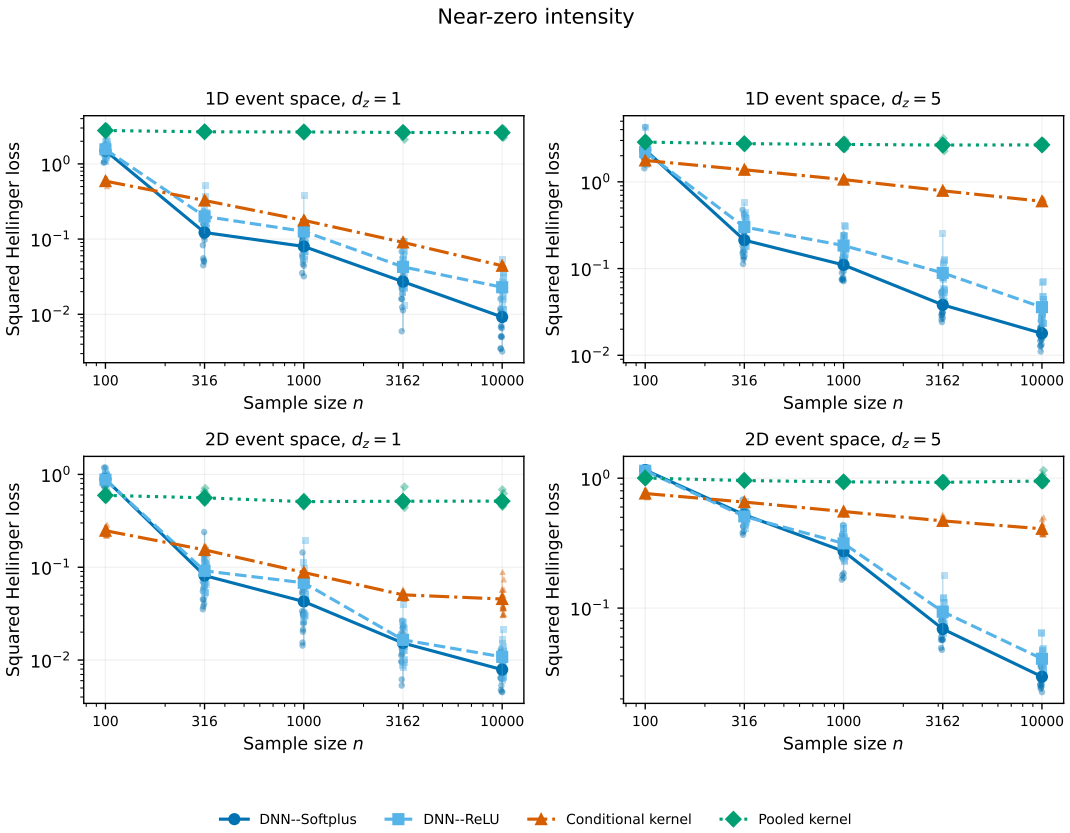

Figure 3: Squared Hellinger loss for the intensity with a moving valley. The intensity floor is 10^{-4} before rescaling. Points, connected markers, and vertical segments show individual repetitions, means, and ranges, as in Figure 1.

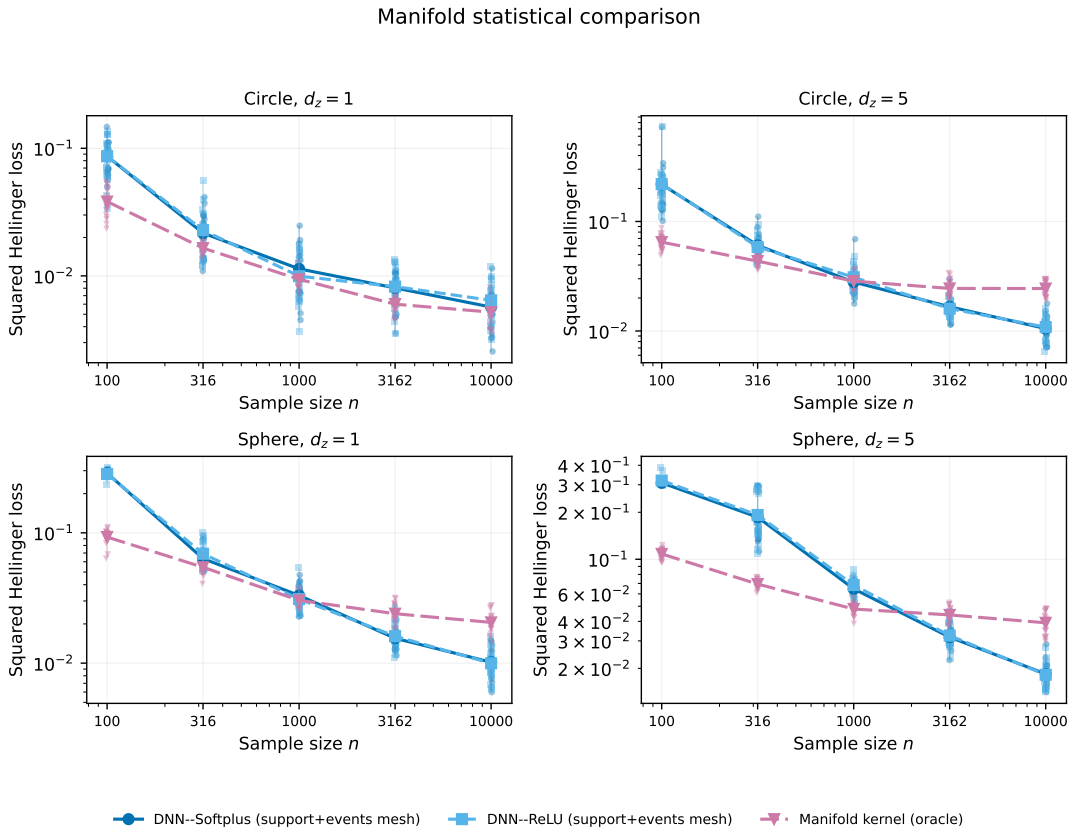

Figure 4: Squared Hellinger loss on the circle and sphere. The DNNs use ambient coordinates and reconstructed integration meshes; the kernel uses known geodesic distances and geometry. Graphical summaries follow Figure 1.

G.4 No-covariate cases and manifold integration

Figure 5 shows the results without covariates. These cases complement the main comparisons by separating estimation of an ordinary intensity from estimation of its covariate dependence. For Euclidean supports, the pooled and conditional kernel estimators coincide, and we fitted only the pooled version.

Figure 6 compares the mesh-based and known-measure neural estimators. For DNN–Softplus at $n = 10000$, the respective mean losses on the circle were 0.00570 and 0.00615 with one covariate, and 0.0105 and 0.0102 with five covariates. On the sphere, the corresponding pairs were 0.0102 and 0.00967, and 0.0184 and 0.0185. Neither version uniformly dominated the other. As described above, the architecture rule also differs between the two versions, so these differences cannot be attributed to the integration measure alone.

G.5 Computation times

Figure 7 and Table 3 summarize computation times. At $n = 10000$, median total times ranged from 26.4 to 33.8 seconds across the six DNN variants, compared with 198 seconds for the manifold kernel, 207 seconds for the pooled kernel, and 459 seconds for the conditional kernel. Total time includes data generation, fitting, evaluation, and file output. The configurations included and computing environments can differ across methods, so these times are descriptive rather than a benchmark under common hardware and workloads. Neural fitting used CUDA when available, with CPU fallback otherwise.

Table 3: Median computation times in seconds at $n = 10000$, pooled over the configurations available for each method. M and O denote reconstructed-mesh and known-measure integration. Total time includes data generation, fitting, evaluation, and file output; component medians need not sum to the median total.

Method	Runs	Total	Fit	Simulation	Evaluation
DNN–Softplus	240	31.3	26.4	2.03	0.714
DNN–ReLU	240	30.5	26.3	2	0.0455
Conditional kernel	160	459	198	0.751	252
Pooled kernel	240	207	68.9	0.762	135
Manifold kernel O	120	198	173	0.885	21.9
DNN–Softplus M	120	33	28.9	2.18	0.0408
DNN–ReLU M	120	33.8	29.3	2.19	0.0397
DNN–Softplus O	120	27.7	23.4	2.39	0.0409
DNN–ReLU O	120	26.4	22.1	2.34	0.0402

H Supplementary Neural Spike-Train Analysis

H.1 Data selection and preprocessing

We required an available NWB recording file and at least 20 VISp units with `isi_violations` < 0.5, `amplitude_cutoff` < 0.1, and `presence_ratio` > 0.9, together with finite spike times. Session 715093703 was the smallest eligible identifier, and 62 units met these criteria. Here a unit is a spike-sorted recording unit, which need not correspond to a perfectly isolated neuron (Allen Institute for Brain Science, n.d.). We removed presentations with missing stimulus covariates, overlap with an invalid recording interval, or a duration different from the modal duration rounded to one millisecond. The exclusions comprised 620 interval overlaps and one duration mismatch. Figure 9 summarizes presentation timing, orientations, and exclusions.

We assigned presentations to development and test samples using within-condition deterministic hashes with seed 20260902. The test sample contained the rounded 20% of presentations in each condition, subject to retaining observations in both samples when possible. The 12-unit selection used development spike counts, with unit identifier breaking ties. Unit 950931315, whose development count was nearest the panel median, was used to illustrate fitted intensities and for the single-unit supplementary analyses. Its earlier pilot test results had already been viewed, as disclosed in the main text. That pilot used normalization based on a reduced sample; all models reported here were refitted with normalization based on the full development sample before their test scores were calculated.

Horizontal and vertical positions were mapped to $[0, 1]$ using the development-sample extrema of -40 and 40 degrees. We retained the original $[-1, 1]$ scale of the doubled-angle sine and cosine coordinates. Both development and test samples contained all 243 combinations of position and orientation.

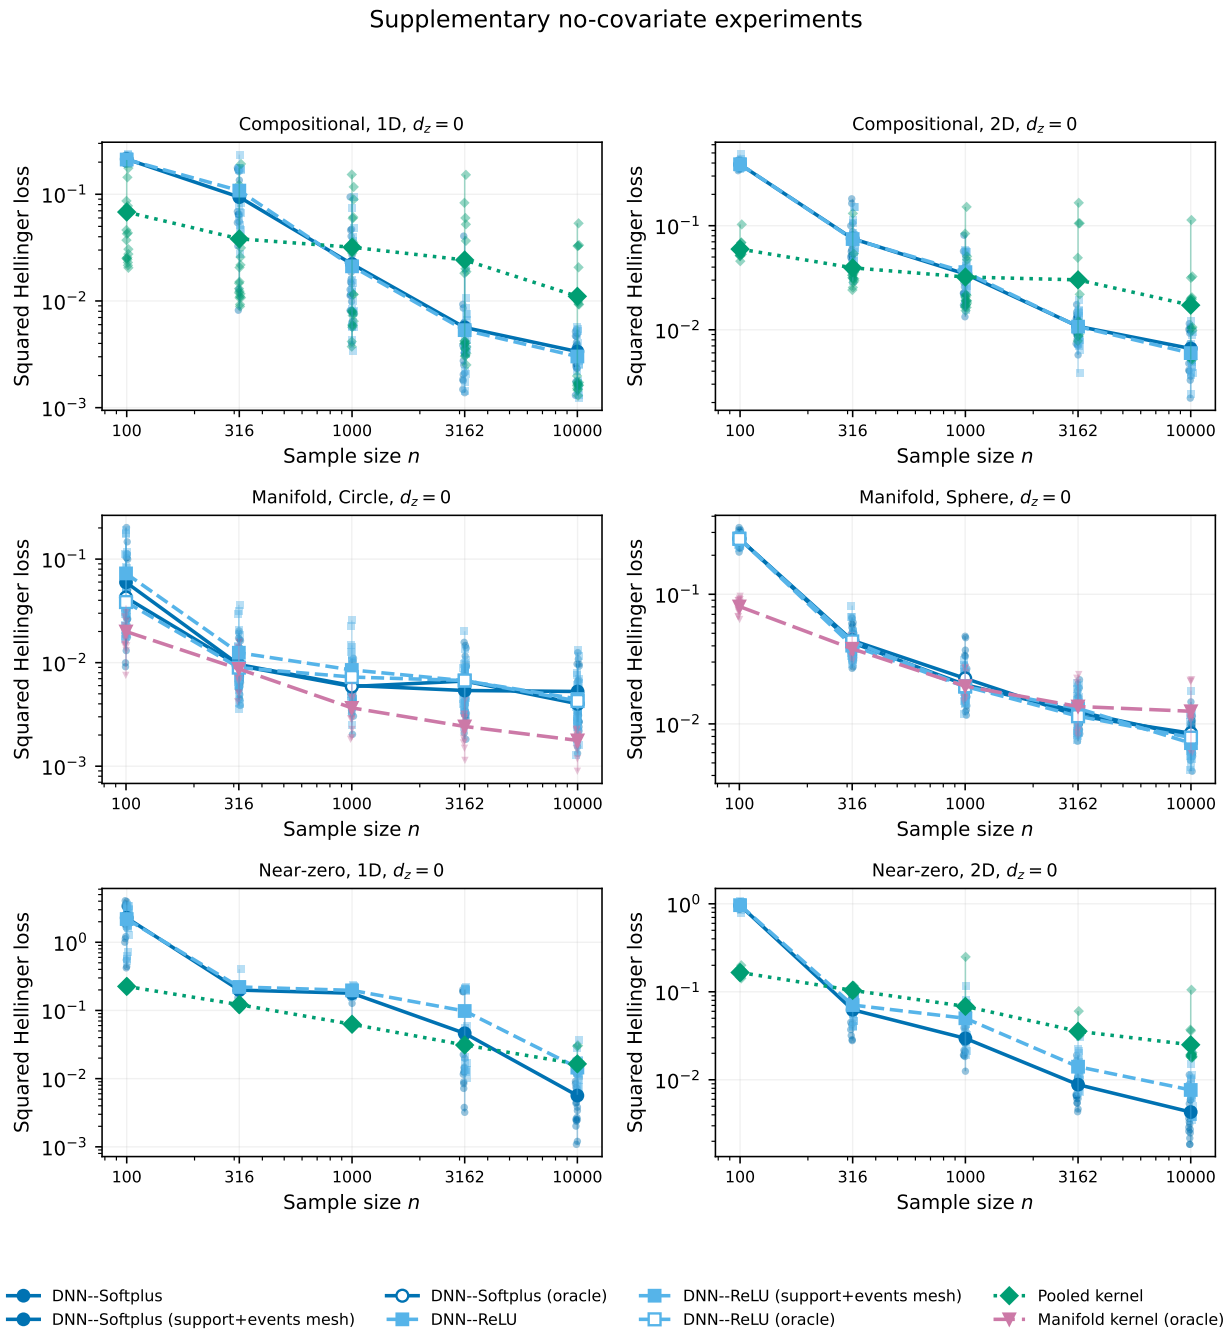

Figure 5: Squared Hellinger loss without covariates. The Euclidean kernel estimates the unconditional intensity. Points, connected markers, and vertical segments show individual repetitions, means, and ranges.

H.2 Fitting procedures

The DNNs used AdamW with learning rate 10^{-3}, weight decay 10^{-4}, batches of 64 presentations, and gradient-norm clipping at 5. Training integrals were approximated by stochastic quadrature with 256 points. We allowed at most 200 epochs and stopped after 20 epochs without improvement in validation NLL. After selecting the epoch count using the 1,935/482 fit/validation split, we refitted from initialization on all 2,417 development presentations for that number of epochs. Optimization and quadrature seeds were 20260903 and 20260904, respectively. As in the simulations, no hard parameter bound or fixed upper intensity bound was imposed.

The kernel estimators used Gaussian smoothing with boundary renormalization on normalized time. Temporal bandwidths were selected from $\{0.02, 0.035, 0.05, 0.075, 0.10, 0.15, 0.225\}$; for the conditional estimator, covariate bandwidths were selected from $\{0.10, 0.15, 0.225, 0.35, 0.50, 0.75, 1.00\}$. Five-fold presentation-level Poisson cross-validation evaluated the full candidate grid using all development presentations and spikes. The denominator in the conditional estimator was bounded below by 10^{-6}.

H.3 Score aggregation and uncertainty

Table 4 gives all pairwise comparisons, and Table 5 reports unit-specific scores and count prediction errors. In addition to the equal-unit, or macro, information gain used in the main text, we calculated a micro gain by pooling log-likelihood improvements and spike counts across units before converting to bits per spike. The 12-unit test sample contained 38,823 spikes. Relative to the conditional kernel, the micro gains were 0.00785 bits per spike for DNN–Softplus and 0.00563 for DNN–ReLU. Differences between macro and micro gains reflect the unequal numbers of spikes across units.

For each of the 2,000 bootstrap draws, we sampled 243 stimulus-condition labels with replacement, using seed 20260905. A presentation received the multiplicity of its condition. We recomputed the presentation-weighted mean score difference within each unit and then the equal-unit mean. Using common multiplicities across methods and units preserves their pairing. The percentile intervals hold the selected units and fitted models fixed and describe uncertainty under condition resampling; they do not account for refitting, between-session variation, or serial dependence across conditions. Intervals are pointwise rather than simultaneous over comparisons.

Table 4: Pairwise predictive comparisons for the 12-unit and 11-unit analyses. Negative mean NLL differences and positive information gains favor the first method. Micro information gain pools score differences and spike counts across units. Brackets give pointwise 95% paired condition-bootstrap intervals.

Units	First vs second method	Macro ΔNLL 95% interval	Micro bits/spike 95% interval
12	DNN–Softplus vs Conditional kernel	-0.029006 [-0.040423, -0.017292]	0.0078515 [0.0047153, 0.010883]
11	DNN–Softplus vs Conditional kernel	-0.026096 [-0.037913, -0.013776]	0.0068535 [0.0036527, 0.0099081]
12	DNN–ReLU vs Conditional kernel	-0.020796 [-0.030918, -0.010262]	0.0056291 [0.0028146, 0.0082415]
11	DNN–ReLU vs Conditional kernel	-0.019545 [-0.029905, -0.0093311]	0.005133 [0.0024795, 0.0077339]
12	DNN–Softplus vs DNN–ReLU	-0.0082103 [-0.017285, 0.00092019]	0.0022224 [-0.00025334, 0.0047148]
11	DNN–Softplus vs DNN–ReLU	-0.0065513 [-0.016178, 0.0035349]	0.0017206 [-0.00092349, 0.0043033]
12	DNN–Softplus vs Pooled kernel	-0.41827 [-0.52087, -0.32659]	0.11322 [0.090735, 0.13701]
11	DNN–Softplus vs Pooled kernel	-0.38835 [-0.48738, -0.29823]	0.10199 [0.080291, 0.12445]
12	DNN–ReLU vs Pooled kernel	-0.41006 [-0.51219, -0.31736]	0.11099 [0.088519, 0.13499]
11	DNN–ReLU vs Pooled kernel	-0.3818 [-0.48114, -0.29125]	0.10027 [0.078504, 0.12314]
12	Conditional kernel vs Pooled kernel	-0.38926 [-0.4869, -0.29954]	0.10536 [0.083373, 0.12835]
11	Conditional kernel vs Pooled kernel	-0.36226 [-0.45632, -0.27519]	0.095138 [0.073976, 0.11694]

Table 5: Unit-specific prediction scores for the four estimators. Mean NLL is in nats per test presentation; count deviance is summed over the 607 test presentations. Count MAE and RMSE are in spikes per presentation.

Unit	Method	Mean NLL	Count deviance	Count MAE	Count RMSE
950930407	DNN–ReLU	-32.091	733.63	2.8705	3.5959
950930407	DNN–Softplus	-32.089	734.56	2.8668	3.596
950930407	Conditional kernel	-32.072	742.38	2.8862	3.6215
950930407	Pooled kernel	-32.009	791.48	2.9592	3.752
950930985	DNN–ReLU	-27.301	931.75	3.0347	3.7733
950930985	DNN–Softplus	-27.302	929.87	3.0328	3.769
950930985	Conditional kernel	-27.286	929.48	3.0395	3.7682
950930985	Pooled kernel	-27.247	967.47	3.1041	3.8542

Unit	Method	Mean NLL	Count deviance	Count MAE	Count RMSE
950931181	DNN–ReLU	-5.4646	2292.4	2.4919	3.5381
950931181	DNN–Softplus	-5.5336	2236.2	2.246	3.4521
950931181	Conditional kernel	-5.4841	2307.2	2.4948	3.5727
950931181	Pooled kernel	-4.5327	3291.6	3.3046	4.4735
950931254	DNN–ReLU	-36.991	1055.6	3.6299	4.503
950931254	DNN–Softplus	-36.954	1077.8	3.614	4.5745
950931254	Conditional kernel	-36.937	1063.6	3.6721	4.5482
950931254	Pooled kernel	-36.594	1243.1	3.9453	4.9961
950931315	DNN–ReLU	-6.707	1272.1	1.9907	2.6299
950931315	DNN–Softplus	-6.7334	1254.6	2.0037	2.6105
950931315	Conditional kernel	-6.6724	1303.2	2.0731	2.6873
950931315	Pooled kernel	-5.9861	1952.2	2.6699	3.4474
950931458	DNN–ReLU	-5.6187	571.92	1.3685	1.763
950931458	DNN–Softplus	-5.6294	562.17	1.3511	1.7467
950931458	Conditional kernel	-5.6495	559.75	1.3422	1.757
950931458	Pooled kernel	-5.441	734.37	1.5685	2.0679
950931565	DNN–ReLU	-6.6409	800.71	1.7104	2.1462
950931565	DNN–Softplus	-6.6454	794.45	1.7068	2.1375
950931565	Conditional kernel	-6.6396	791.37	1.7167	2.138
950931565	Pooled kernel	-6.5978	838.68	1.7835	2.2117
950931656	DNN–ReLU	-3.1643	485.06	1.0063	1.2609
950931656	DNN–Softplus	-3.1617	495.25	1.0088	1.2732
950931656	Conditional kernel	-3.1637	493.63	1.0084	1.2855
950931656	Pooled kernel	-2.9916	633.3	1.1747	1.5097
950931727	DNN–ReLU	-9.8035	471.29	1.4317	1.9249
950931727	DNN–Softplus	-9.805	475.97	1.4351	1.9465
950931727	Conditional kernel	-9.7698	477.8	1.4584	1.9598
950931727	Pooled kernel	-9.3176	883.66	2.0054	2.7506
950931770	DNN–ReLU	-4.5335	552.61	1.2958	1.6014
950931770	DNN–Softplus	-4.5413	532.31	1.231	1.5694
950931770	Conditional kernel	-4.5402	540.2	1.2628	1.5851
950931770	Pooled kernel	-4.4079	651.68	1.3496	1.7743
950931899	DNN–ReLU	-3.8556	556.82	1.1241	1.4201
950931899	DNN–Softplus	-3.8596	549.86	1.1253	1.4085
950931899	Conditional kernel	-3.8422	553.41	1.1316	1.4259
950931899	Pooled kernel	-3.6331	776.51	1.416	1.7515
950932563	DNN–ReLU	-7.5589	1803.2	2.3126	2.9565
950932563	DNN–Softplus	-7.5743	1792.8	2.3058	2.9641
950932563	Conditional kernel	-7.4241	1953.7	2.5046	3.2094
950932563	Pooled kernel	-6.0513	3390.6	3.5601	4.7588

H.4 Numerical and modeling sensitivity

We first checked the numerical integration used to score all 48 unit–method fits. Relative to the 1,024-point grid, grids of 512 and 2,048 points changed the mean NLL by at most 1.393×10^{-4} and 3.481×10^{-5}, respectively (Figure 11). These changes were small compared with the panel-average DNN–kernel differences.

The remaining analyses in this subsection used unit 950931315, the previously examined pilot unit, and were exploratory. With optimization seeds 20260903, 20260913, and 20260923, DNN–Softplus gave mean test NLLs of -6.73342, -6.68916, and -6.71965; DNN–ReLU gave -6.70696, -6.71506, and -6.69711. The ordering of the output functions therefore depended on initialization. For Softplus networks with widths 32, 64, and 128, the scores were -6.70779, -6.73342, and -6.71669. The main results use width 64 and seed 20260903 throughout, without selection among these supplementary fits.

We also included mean signed running speed during the presentation as an additional covariate. The resulting NLLs were -6.72147 for DNN–Softplus, -6.71394 for DNN–ReLU, and -6.51896 for the conditional kernel. Relative to the stimulus-only fits, the DNN changes were small and differed in direction, whereas the conditional-kernel score increased. The pooled estimator was unchanged because it does not use covariates. Running speed was available for all retained presentations; pupil area was unavailable and was omitted.

To examine prediction at new stimulus conditions, we partitioned the 243 position–orientation combinations into five folds and withheld all presentations of each fold in turn. Normalization and epoch selection were repeated using the other four folds. Both DNNs had lower NLL than the conditional kernel in all five held-out folds (Figure 13). This exploratory analysis repartitions all presentations, including the primary test sample; it is not an independent replication. Table 6 summarizes these comparisons.

Table 6: Sensitivity analyses for unit 950931315. Mean NLL is in nats per test presentation, and count deviance is summed over the test presentations. The seed and width analyses include the corresponding primary fit for reference. Condition-fold scores use different held-out stimuli.

Analysis	Method	Presentations	Mean NLL	Count deviance
Seed 20260903	DNN–Softplus	607	-6.7334	1254.6
Seed 20260903	DNN–ReLU	607	-6.707	1272.1
Seed 20260913	DNN–Softplus	607	-6.6892	1318.5
Seed 20260913	DNN–ReLU	607	-6.7151	1269.9
Seed 20260923	DNN–Softplus	607	-6.7197	1262.4
Seed 20260923	DNN–ReLU	607	-6.6971	1275.9
Width 32	DNN–Softplus	607	-6.7078	1286.7
Width 64	DNN–Softplus	607	-6.7334	1254.6
Width 128	DNN–Softplus	607	-6.7167	1271.7
Running speed	DNN–Softplus	607	-6.7215	1254.1
Running speed	DNN–ReLU	607	-6.7139	1253.6
Running speed	Pooled kernel	607	-5.9861	1952.2
Running speed	Conditional kernel	607	-6.519	1412.7
Condition fold 1	DNN–Softplus	620	-6.8944	1507.3
Condition fold 1	DNN–ReLU	620	-6.8842	1509
Condition fold 1	Pooled kernel	620	-6.3286	1956.2
Condition fold 1	Conditional kernel	620	-6.8342	1517.2
Condition fold 2	DNN–Softplus	610	-5.4715	1477.9
Condition fold 2	DNN–ReLU	610	-5.4734	1474.6
Condition fold 2	Pooled kernel	610	-5.1138	1801.1
Condition fold 2	Conditional kernel	610	-5.4171	1497.4
Condition fold 3	DNN–Softplus	595	-6.6679	1382.4
Condition fold 3	DNN–ReLU	595	-6.6394	1408.3
Condition fold 3	Pooled kernel	595	-5.9596	2016.4
Condition fold 3	Conditional kernel	595	-6.6126	1443.8
Condition fold 4	DNN–Softplus	597	-8.3756	1327.9
Condition fold 4	DNN–ReLU	597	-8.39	1309.1
Condition fold 4	Pooled kernel	597	-7.4918	2140.2
Condition fold 4	Conditional kernel	597	-8.2677	1413.7
Condition fold 5	DNN–Softplus	602	-6.3423	1403.3
Condition fold 5	DNN–ReLU	602	-6.3328	1407.7
Condition fold 5	Pooled kernel	602	-5.644	1983.6
Condition fold 5	Conditional kernel	602	-6.2842	1455.4

H.5 Fitted response profiles and model diagnostics

Figure 8 displays fitted temporal intensities for unit 950931315 at three stimulus conditions. We selected these conditions using development counts to represent preferred, intermediate, and low-response stimuli. The conditional estimates vary in both response magnitude and temporal shape, whereas the pooled temporal kernel returns one curve for all conditions.

For count calibration, we pooled test unit–presentation pairs and formed deciles of predicted count separately for each method. The left panel of Figure 12 compares average observed and predicted counts in these groups. This display assesses aggregate count prediction rather than calibration separately for every unit and stimulus condition.

We assessed within-presentation timing for unit 950931315 by integrating the fitted intensity on a 2,048-point grid. Transformed intervals ending at a spike were treated as completed observations, and the remaining interval to the end of each presentation was treated as right-censored; for a zero-spike presentation, the whole window was censored. The right panel of Figure 12 compares Kaplan–Meier survival curves of these integrated-hazard increments with the unit-exponential survival function. The visible deviations indicate features of the spike trains not fully captured by the conditional Poisson model, despite improvements in predictive NLL.

Finally, we examined temporal variation using the original fitted models. Figure 13 reports scores by chronological quartile of the test presentations. Differences in absolute NLL across quartiles can arise from spike counts and stimulus composition as well as predictive fit. We also calculated unit-level paired intervals using noncircular moving blocks of 20 or 50 ordered test presentations, with 2,000 resamples and concatenated blocks truncated to the original sample size. Because these blocks follow the ordered test subset rather than contiguous presentations of the full recording, this analysis is a sensitivity check for dependence within the test sample, not an assessment of prediction beyond the training period.

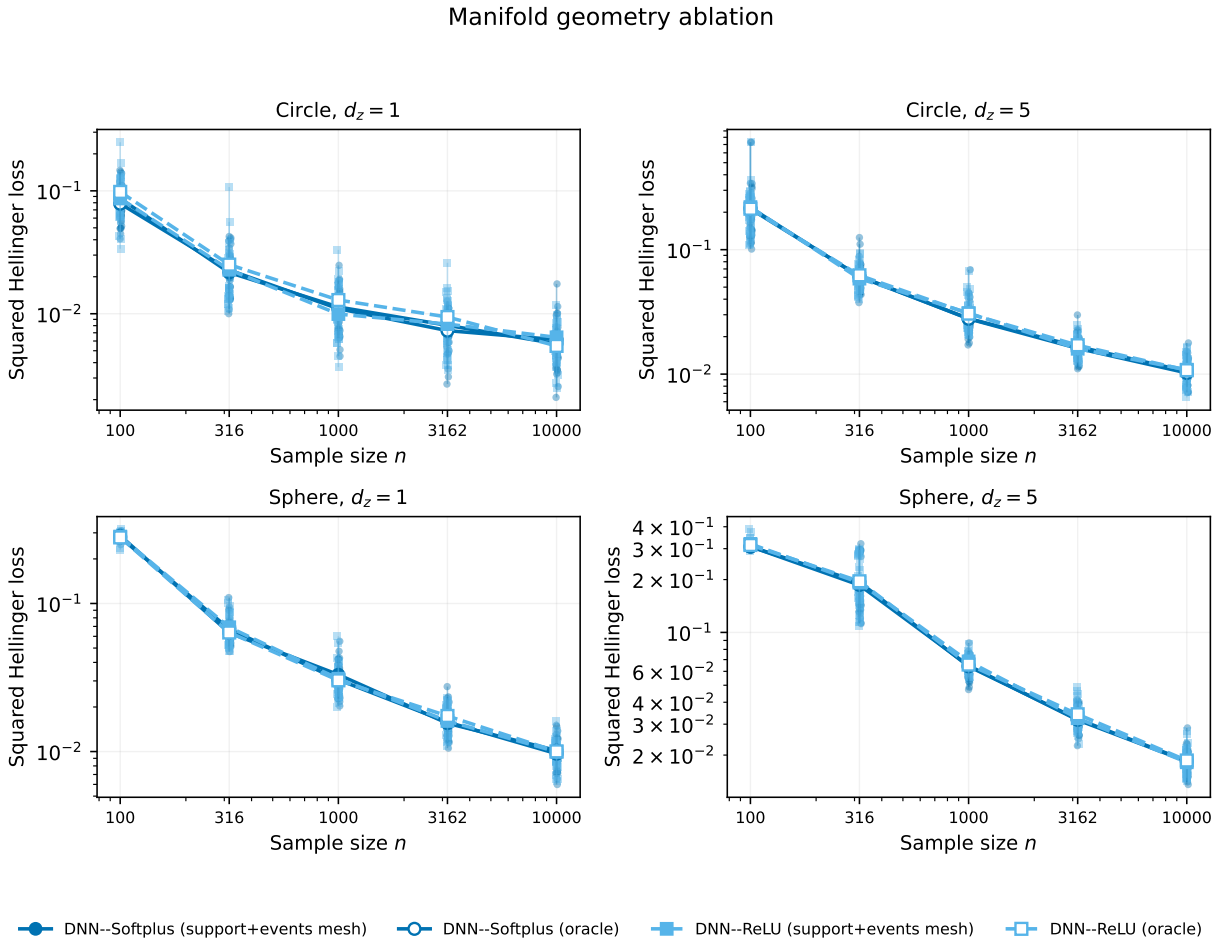

Figure 6: Neural estimators with reconstructed-mesh (M) and known-measure (O) integration. Both use ambient inputs, with different dimensions in the architecture rule (28).

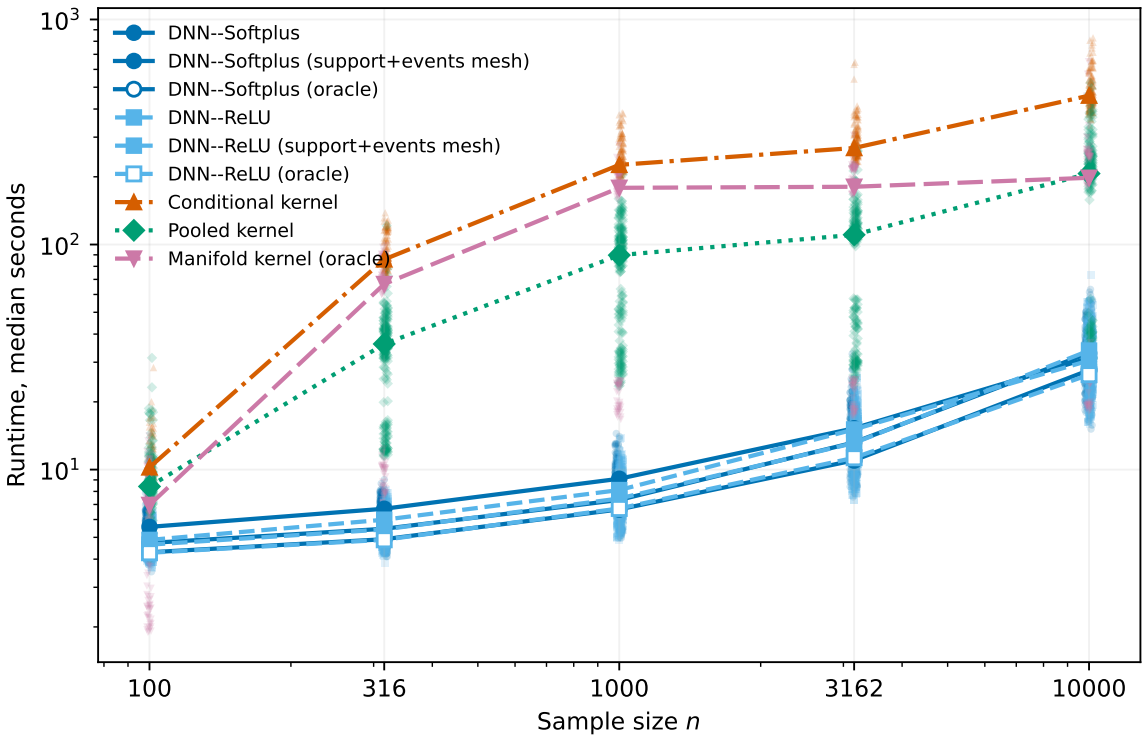

Figure 7: Median total computation time by sample size, pooled over the configurations and repetitions available for each method. Totals include data generation, fitting, evaluation, and output.

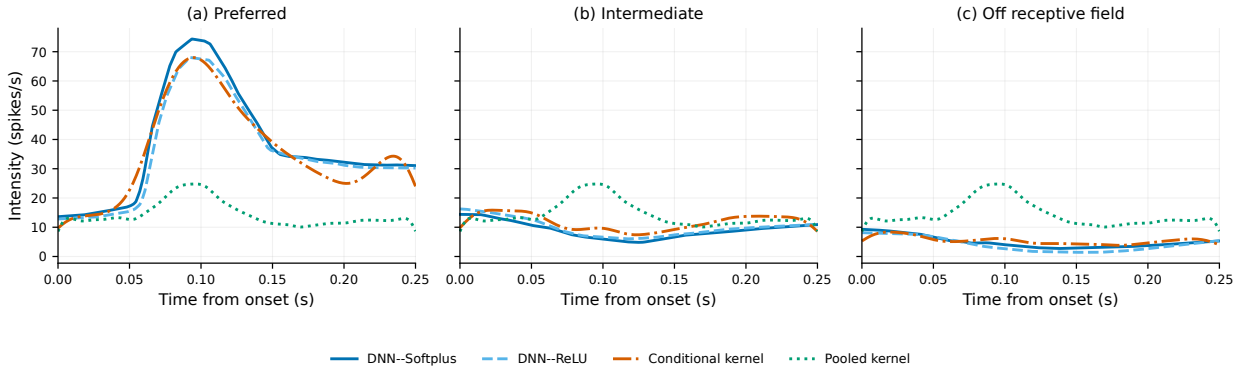

Figure 8: Estimated temporal firing rates for unit 950931315 at preferred, intermediate, and low-response stimulus conditions selected using development counts. Time is measured from stimulus onset and intensity is in spikes per second. Curves are fitted intensities.

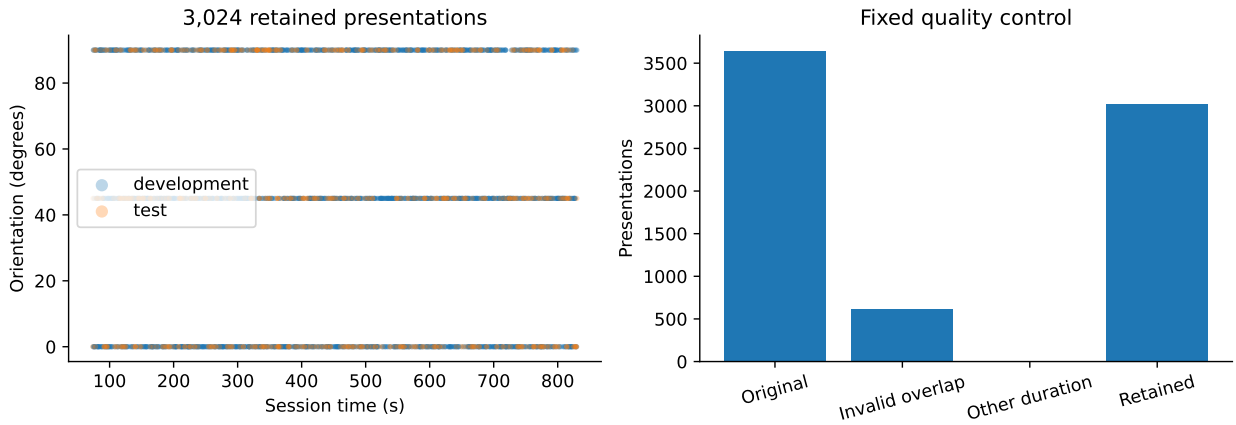

Figure 9: Gabor presentation timing and orientation in the development and test samples (left), and presentation counts before and after exclusions (right). The split was stratified by position–orientation combination.

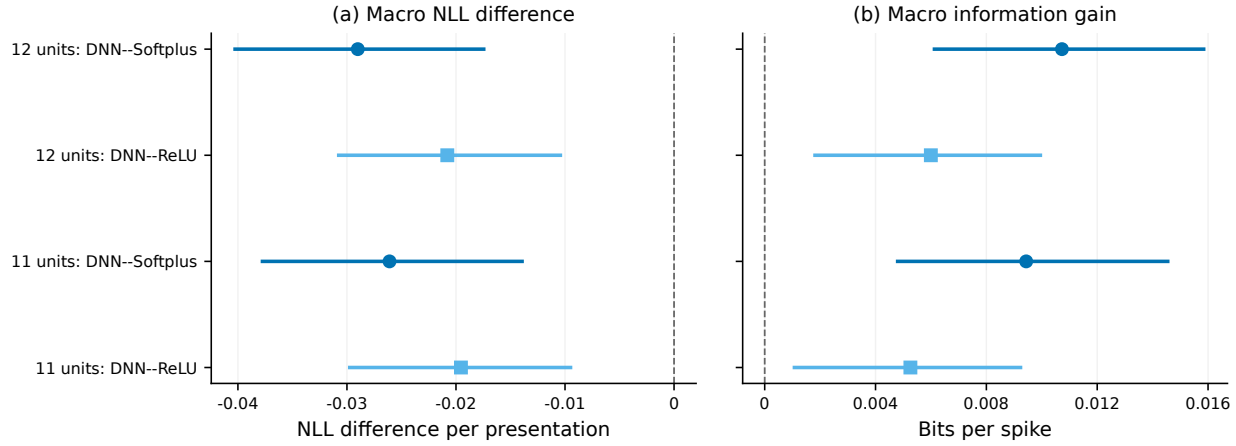

Figure 10: Panel-average DNN performance relative to the conditional kernel for all 12 selected units and the 11 units excluding the pilot. Negative NLL differences and positive information gains favor the DNN. Segments are pointwise 95% paired condition-bootstrap intervals.

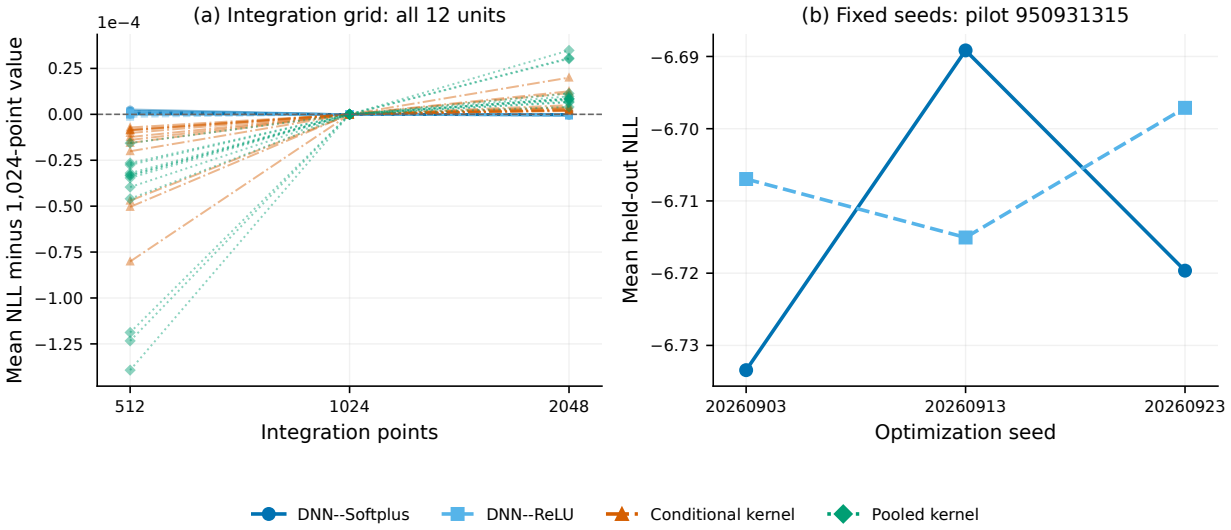

Figure 11: Sensitivity to integration resolution and network initialization. Left: changes in mean NLL from the 1,024-point score for each of the 48 unit–method fits. Right: mean test NLL at three optimization seeds for unit 950931315.

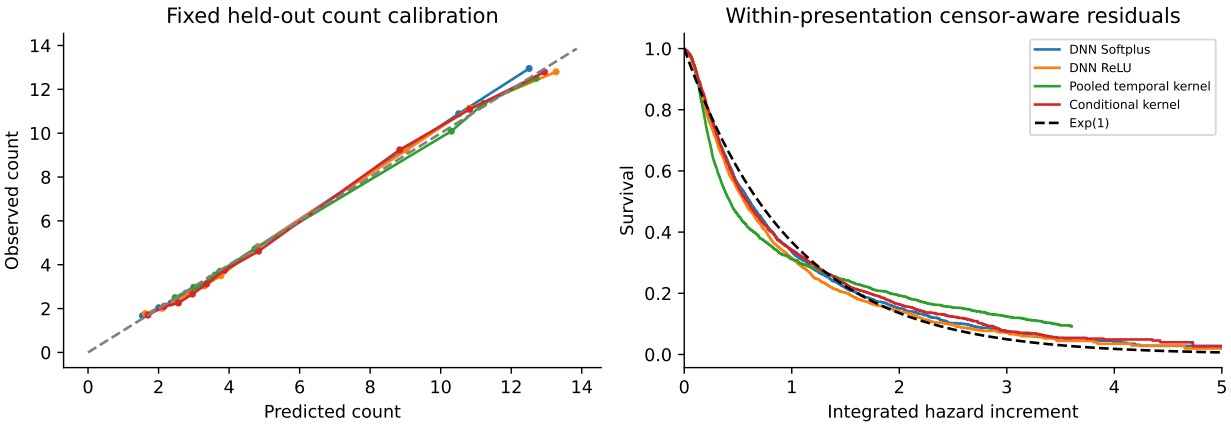

Figure 12: Count and timing diagnostics. Left: observed and predicted counts in method-specific deciles of predicted count, pooled over test unit–presentation pairs. Right: Kaplan–Meier survival curves of integrated-hazard increments for unit 950931315, with right censoring at presentation end, and the unit-exponential reference.

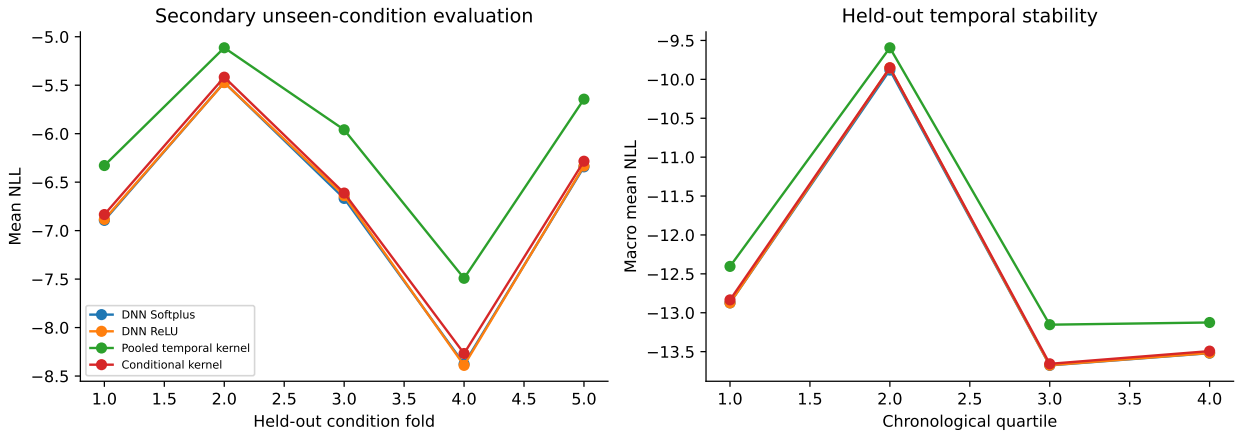

Figure 13: Prediction across stimulus conditions and recording time. Left: mean NLL for unit 950931315 in five condition-held-out folds. Right: equal-unit mean NLL by chronological quartile of the primary test presentations, using the original fits.

References

Allen Institute for Brain Science. Visual coding—neuropixels: AllenSDK documentation. `https://allensdk.readthedocs.io/en/v2.16.2/visual_coding_neuropixels.html`, n.d. Version 2.16.2; accessed September 20, 2026.

Thijs Bos and Johannes Schmidt-Hieber. Convergence rates of deep ReLU networks for multiclass classification. *Electron. J. Stat.*, 16:2724–2773, 2022. doi: 10.1214/22-EJS2011.

Chin-tsang Chiang, Mei-cheng Wang, and Chiung-yu Huang. Kernel estimation of rate function for recurrent event data. *Scandinavian journal of statistics*, 32(1):77–91, 2005.

Peter Diggle. A kernel method for smoothing point process data. *Journal of the Royal Statistical Society: Series C (Applied Statistics)*, 34(2):138–147, 1985.

Peter Diggle, Sara Morris, Paul Elliott, and Gavin Shaddick. Regression modelling of disease risk in relation to point sources. *Journal of the Royal Statistical Society Series A: Statistics in Society*, 160(3):491–505, 1997.

Seth Flaxman, Yee Whye Teh, and Dino Sejdinovic. Poisson intensity estimation with reproducing kernels. In *Artificial Intelligence and Statistics*, pages 270–279. PMLR, 2017.

Evarist Giné and Richard Nickl. *Mathematical foundations of infinite-dimensional statistical models*. Cambridge university press, 2021.

Matteo Giordano, Alisa Kirichenko, and Judith Rousseau. Nonparametric bayesian intensity estimation for covariate-driven inhomogeneous point processes. *Bernoulli*, 32(2):1020–1044, 2026.

Yoshihiro Gyotoku, Ioane Muni Toke, and Nakahiro Yoshida. Deep learning of point processes for modeling high-frequency data. *arXiv preprint arXiv:2504.15944*, 2025.

Alisa Kirichenko and Harry van Zanten. Optimality of poisson processes intensity learning with gaussian processes. *J. Mach. Learn. Res.*, 16(1):2909–2919, 2015.

Nicolas Klutchnikoff and Gaspar Massiot. Kernel estimation of the intensity of cox processes. *arXiv preprint arXiv:1605.06703*, 2016.

Michael Kohler and Sophie Langer. On the rate of convergence of fully connected deep neural network regression estimates. *Ann. Statist.*, 49(4):2231–2249, 2021. doi: 10.1214/20-AOS2034.

Yu A Kutoyants. *Statistical inference for spatial Poisson processes*. Number 134. Springer Science & Business Media, 2012.

Günter Last and Mathew Penrose. *Lectures on the Poisson process*, volume 7. Cambridge University Press, 2017.

Jerald Franklin Lawless. Regression methods for poisson process data. *Journal of the American Statistical Association*, 82(399):808–815, 1987.

Lasse Leskelä. Information divergences and likelihood ratios of poisson processes and point patterns. *arXiv preprint arXiv:2404.00294*, 2024.

Luis Moreira-Matias, Joao Gama, Michel Ferreira, Joao Mendes-Moreira, and Luis Damas. Predicting taxi–passenger demand using streaming data. *IEEE Transactions on Intelligent Transportation Systems*, 14(3):1393–1402, 2013.

Tin Lok James Ng and Andrew Zammit-Mangion. Spherical poisson point process intensity function modeling and estimation with measure transport. *Spatial Statistics*, 50:100629, 2022.

Yosihiko Ogata. Space-time point-process models for earthquake occurrences. *Annals of the Institute of Statistical Mathematics*, 50:379–402, 1998.

Maya Okawa, Tomoharu Iwata, Takeshi Kurashima, Yusuke Tanaka, Hiroyuki Toda, and Naonori Ueda. Deep mixture point processes: Spatio-temporal event prediction with rich contextual information. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, pages 373–383, 2019.

Finbarr O'Sullivan. Nonparametric estimation in the cox model. *The Annals of Statistics*, pages 124–145, 1993.

Mathew D Penrose and Joseph E Yukich. Limit theory for point processes in manifolds. *The Annals of Applied Probability*, pages 2161–2211, 2013.

Patricia Reynaud-Bouret and Vincent Rivoirard. Near optimal thresholding estimation of a poisson intensity on the real line. *Electron. J. Stat.*, 2010.

Johannes Schmidt-Hieber. Nonparametric regression using deep neural networks with ReLU activation function. *Annals of Statistics*, 48:1875 – 1897, 2020.

Joshua H. Siegle, Xiaoxuan Jia, Séverine Durand, Sam Gale, Corbett Bennett, Nile Graddis, et al. Survey of spiking in the mouse visual system reveals functional hierarchy. *Nature*, 592:86–92, 2021. doi: 10.1038/s41586-020-03171-x. URL https://www.nature.com/articles/s41586-020-03171-x.

Taiji Suzuki. Adaptivity of deep ReLU network for learning in besov and mixed smooth besov spaces: optimal rate and curse of dimensionality. In *International Conference on Learning Representations*, 2019.

Alexandre B. Tsybakov. *Introduction to Nonparametric Estimation*. Springer Publishing Company, Incorporated, 2008.

Sara van de Geer. *Empirical Processes in M-estimation*, volume 6. Cambridge university press, 2000.

Ruixin Wang, Prateek Jaiswal, and Harsha Honnappa. Estimating stochastic poisson intensities using deep latent models. In *2020 Winter Simulation Conference (WSC)*, pages 596–607. IEEE, 2020.

S Ward, HS Battey, and EAK Cohen. Nonparametric estimation of the intensity function of a spatial point process on a riemannian manifold. *Biometrika*, 110(4):1009–1021, 2023.

Atsutomo Yara and Yoshikazu Terada. Nonparametric logistic regression with deep learning. *Bernoulli*, 32(2):952–977, 2026.

Dmitry Yarotsky. Error bounds for approximations with deep relu networks. *Neural networks*, 94:103–114, 2017.