arXiv:2611.06228v1 [cs.LG] 4 Nov 2026

Separation Is a Riesz Energy: How Many Environments Does ICA Need?

Yuda Bi [*1] **Yubo Guo** [*2] **Jiaxiao Cai** [3] **Xiaonai Li** [4] **Zhuolin Li** [5] **Vince D. Calhoun** [1]

Abstract

Can independent components be recovered when all of them are Gaussian? Yes, if their variances change across environments; but how many environments, and how many samples as the number of components d grows? We answer both for the joint diagonalization of K covariances. Each component has a *signature*, its variances across the environments. Separating component j costs the Riesz 2-energy E_j of this signature, a sum, over the other components, of inverse pairwise Fisher informations. For Gaussian sources no estimator in $\mathrm{O}(d)$ beats $cE_j/(n+E_j)$ at any n and d, and above an explicit sample size a spectral estimator attains E_j/n up to a constant, for every mixing and without the signatures (given a lower bound on the pairwise information). Since signatures live in $K-1$ dimensions, the rate obeys a dimension law: d^2/n, $d\log d/n$ and d/n for $K=2$, 3 and $K\geq 4$ well-spread environments. What counts is the dimension the signatures span, not their number, and $\log d$ well-designed environments make $d\log d$ samples suffice. Projecting the covariances onto one linear combination pays d^2/n at first order instead. In simulations up to $d=1000$ the predictions hold and joint diagonalization with $K\geq 4$ attains the rate once n/d is in the hundreds; on task fMRI the energy says which components are hard to separate.

1. Introduction

Independent component analysis (ICA; Hyvärinen et al., 2001) recovers independent sources from a linear mixture $x=As$, with $A\in\mathbb{R}^{d\times d}$ invertible. Non-Gaussian sources are identified from i.i.d. samples by higher-order statistics (Comon, 1994). Gaussian sources are not: i.i.d. samples determine A only up to a rotation. *Diversity* rescues them (Tong et al., 1991; Adalı et al., 2014). Suppose the source variances change across K environments, be they segments of a nonstationary recording (Matsuoka et al., 1995; Pham & Cardoso, 2001), time lags (Belouchrani et al., 1997) or datasets (Anderson et al., 2014). The covariances $\Sigma_k=A\Lambda_kA^\top$ then share the mixing and differ in the variances, and A is identifiable from second-order statistics alone unless two sources share a variance profile. The same idea underlies nonlinear ICA with auxiliary variables and causal representation learning, where environments or interventions must be varied enough ("sufficient variability") (Hyvärinen & Morioka, 2016; Hyvärinen et al., 2019; Khemakhem et al., 2020; Squires et al., 2023; Jin & Syrgkanis, 2024). Such conditions are qualitative: they say when identification is possible, not how many samples it takes or which environments to collect (finite-sample results: Appendix I).

For the non-Gaussian route, Auddy & Yuan (2025) showed that the minimax rate in dimension d is d/n for the squared angle, and that no low-degree polynomial estimator succeeds with $n\lesssim d^2$. For the diversity route, classical theory gives asymptotic variances at fixed dimension (Pham & Cardoso, 2001; Tichavský & Yeredor, 2009; Yeredor, 2010; Miettinen et al., 2016; Afsari, 2008), and Chang et al. (2018) analyse one reduction in growing dimension. How the rate depends on d, K and the sources' spread across environments was open. The question is practical. Group ICA pipelines in neuroimaging estimate $d=100$–300 components (Calhoun et al., 2001; Calhoun & Adalı, 2012; Du et al., 2020; Iraji et al., 2023), and in time series many lags (SOBI) are preferred to one (AMUSE) (Belouchrani et al., 1997; Miettinen et al., 2016); lags and independent vector analysis (IVA; Kim et al., 2006) are outside our theorems.

This paper. Second-order statistics see source j only through its *signature* $\mathbf{p}_j=(\mathbf{p}_j(1),\dots,\mathbf{p}_j(K))$, the vector of its variances across environments. We show that the diversity route is governed by the geometry of these d points.

[*]Equal contribution [1]Tri-institutional Center for Translational Research in Neuroimaging and Data Science (TReNDS), Georgia State University, Georgia Institute of Technology, Emory University, Atlanta, GA, USA [2]Huazhong University of Science and Technology, Wuhan, China [3]University of Pennsylvania, Philadelphia, PA, USA [4]Lancaster University, Lancaster, UK [5]Georgia Institute of Technology, Atlanta, GA, USA. Correspondence to: Yuda Bi <ybi3@gsu.edu>.

Preprint. November 5, 2026.

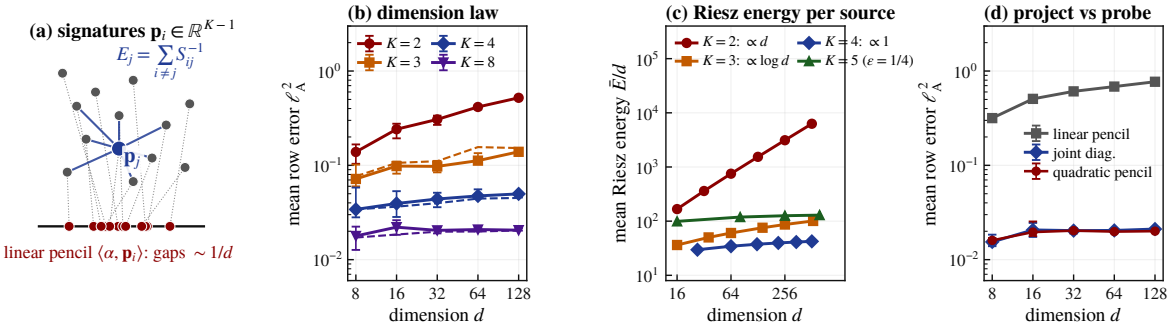

Figure 1. **Separation is a Riesz energy. (a)** Each source is a point $\mathbf{p}_i$ in the $(K-1)$-dimensional space of signatures; separating source j costs the inverse-square energy of its interactions (blue), which a quadratic pencil centred at $\mathbf{p}_j$ measures, while a linear pencil $\sum_k \alpha_k \Sigma_k$ projects the signatures onto a line (red), where gaps are $\sim 1/d$. **(b)** *The dimension law, simulated:* mean row error of joint diagonalization for random signatures at $n = 400d$; dashed: $\bar{E}^{\mathrm{w}}/n$ (Proposition 2). **(c)** $\bar{E}/d$ for lattice signatures grows like d ($K = 2$) and $\log d$ ($K = 3$) and saturates ($K \geq 4$). **(d)** *Projecting versus probing:* on identical data ($K = 8$, $n = 400d$) the eigenvectors of one random linear pencil degrade with d; joint diagonalization and the quadratic pencil of Theorem 2 do not. Details: Appendix J.

Let

$$S_{ij} = \frac{1}{K} \sum_{k=1}^{K} \frac{\big(\mathbf{p}_i(k) - \mathbf{p}_j(k)\big)^2}{\mathbf{p}_i(k)\, \mathbf{p}_j(k)}, \qquad E_j = \sum_{i \neq j} S_{ij}^{-1}. \tag{1}$$

Here S_{ij} is the Fisher information per sample, averaged over the environments, about the rotation that mixes sources i and j. E_j is the Riesz 2-energy of the point $\mathbf{p}_j$ in the configuration of signatures. A toy case shows why the dimension of the configuration matters. With two environments the signatures are d numbers in an interval. Some neighbours are then $1/d$ apart, and E_j is of order d^2. With four, they are d points in a 3-dimensional body. Neighbours are then $d^{-1/3}$ apart and the many far sources dominate the inverse-square sum, so E_j is of order d.

Contributions.

1. **Separation costs a Riesz energy (Section 3).** For Gaussian sources and every n and d, no estimator in $\mathrm{O}(d)$ recovers the j-th mixing direction with worst-case squared-angle risk below $c\, E_j/(n + E_j)$ (Theorem 1). Conversely, a spectral estimator attains E_j/n up to a constant for every mixing, uniformly in d, once $n \gtrsim K(d + \log n)/S_{\min}$, where $S_{\min}$ is the smallest divergence (Theorem 2). It needs the signatures, and it may estimate the whitening from the data. The estimator is the bottom eigenvector of a sample version of $\sum_k w_k(\Sigma_k - \mathbf{p}_j(k)I)^2$, a matrix that vanishes only along source j. A grid of such matrices also finds the signatures, given a lower bound s_0 on $S_{\min}$ (Theorem 3). In this regime the energy is the minimax rate (Corollary 1).

2. **Dimension, not number (Section 4).** Counting how many signatures lie at each distance from $\mathbf{p}_j$ gives $E_j \asymp d^2$, $d \log d$ and d for $K = 2$, 3 and $K \geq 4$ well-spread environments (Theorem 4). Combined with Theorem 1, the lower bounds on E_j make $n \gtrsim d^2$, $d \log d$ and d samples necessary, for every configuration (Corollary 2). With $K \asymp \log d$ environments in which every pair of sources differs by a fixed factor in a constant fraction of them, $n \gtrsim d \log d$ samples are sufficient for the estimator of Theorem 2. Eight environments whose signatures lie on a line behave like two (Remark 1, Figure 2(a)).

3. **Projecting versus probing (Section 5).** The eigenvectors of one combination $\sum_k \alpha_k \hat{\Sigma}_k$ (AMUSE, generalised eigenproblems, random pencils) have a first-order error whose mean over sources is $\gtrsim d^2/n$ for every K (Theorem 5): the combination projects the signatures onto a line. Probing one signature at a time avoids this. The first-order error of joint diagonalization is again of Riesz type (Proposition 2).

In simulations up to $d = 1000$ the predicted errors match the observed ones in the first-order regime. Joint diagonalization attains the Riesz rate with n/d nearly constant (Section 6, Conjecture 1). On task fMRI the energy says which components are hard to separate, but not how large their error is.

Relation to prior work. Inverse-square error formulas are classical at fixed d: for eigenvectors (Anderson, 1963) and, pair by pair, for joint diagonalization (Cardoso, 1994; Miettinen et al., 2016). The structure is visible in Kuleshov et al. (2015, Lemma 1), and Afsari (2007; 2008) showed that nearly collinear signatures make joint diagonalization ill-conditioned. What was missing is a quantitative role for the *dimension* of the configuration, and a minimax theory that holds as d grows. Matrices $\sum_k (X_k - \lambda_k I)^2$ are a known numerical device for approximate joint eigenvectors (Loring & Vides, 2022; Cerjan et al., 2023); we supply their

statistics. Table 1 lists every result with its conditions, a reader's guide follows the notation table of the appendix, and Appendix I discusses further related work.

Why this model. We work in the linear Gaussian model with K environments, the common-principal-components model of Flury (1984). First, it is the least informative case. Its lower bounds transfer to every class of sources containing the Gaussian ones, and within a fixed non-Gaussian family the Fisher information of a rotation is at least as large (Appendix A). Second, it is the case in which upper and lower bounds match (Corollary 1), as in sparse-PCA minimax theory (Vu & Lei, 2013; Cai et al., 2013), whose lower bounds are also proved in the Gaussian model. Non-Gaussian sources and real data: Section 6, Appendix J.

2. Setting and signature geometry

Model. There are K environments. In environment k we observe $n_k = n/K$ independent copies of $x^{(k)} = A\, s^{(k)}$, with $A \in \mathbb{R}^{d\times d}$ invertible and independent, centred coordinates $s_j^{(k)}$ of variance $\mathbf{p}_j(k)$. We normalise $\frac{1}{K}\sum_k \mathbf{p}_j(k) = 1$ and assume $\mathbf{p}_j(k) \in [1/B, B]$, so the pooled covariance is $AA^\top$. After whitening, $A = Q \in \mathrm{O}(d)$ and $\Sigma_k = Q\Lambda_k Q^\top$ with $\Lambda_k = \mathrm{diag}(\mathbf{p}_1(k), \dots, \mathbf{p}_d(k))$; $\hat\Sigma_k$ is the sample covariance of environment k.

Scope. Lower bounds: Gaussian sources, and estimators in $\mathrm{O}(d)$ that may use the signatures. Non-asymptotic upper bounds (Theorems 2 and 3): Gaussian sources and equal sizes $n_k = n/K$; each theorem has a part for known and a part for estimated whitening. First-order results: independent sources with finite fourth moments. Unequal sizes $n_k = \pi_k n$ replace $\frac{1}{K}\sum_k$ by $\sum_k \pi_k$ in the lower bounds, the first-order results and the dimension law. An environment with very few samples (two in Appendix A's example) can, however, break the estimators of the upper bounds.

Loss. Estimators take values in the orthogonal group $\mathrm{O}(d)$, as every orthogonal joint diagonalizer does, and may be randomised. For an estimate $\hat Q \in \mathrm{O}(d)$ with columns $\hat q_a$ and the truth Q with columns q_j, let

$$\rho_j(\hat Q) = 1 - \max_a \langle \hat q_a, q_j\rangle^2, \qquad \ell_{\mathrm{A}}^2 = \frac{1}{d}\sum_{j=1}^{d}\rho_j, \tag{2}$$

and $\ell_{\mathrm{M}}^2 = \max_j \rho_j$. So ρ_j is the squared sine of the angle between q_j and the closest estimated direction, in whitened coordinates, and ℓ_{A}^2 is permutation invariant but decomposes by source. When the signatures are known, an estimator may also be *labelled*: one unit vector $\hat q_j$ per source, with $\rho_j = 1 - \langle \hat q_j, q_j\rangle^2$. Values of ℓ_{A}^2 above $1/2$ are only lower bounds for a permutation-matched loss, since two sources may then share their closest estimate (other losses: Appendix A).

Notation: $\mathrm{sym}(X) = \frac{1}{2}(X + X^\top)$; $a \lesssim b$ means $a \le Cb$ with C depending only on the stated constants, and $a \asymp b$ means both $a \lesssim b$ and $b \lesssim a$; $\Rightarrow$ is convergence in law; a polar factor is a nearest orthogonal matrix.

Signatures. The signature $\mathbf{p}_j \in \mathbb{R}^K$ lies on the hyperplane $\sum_k \mathbf{p}_j(k) = K$, a $(K-1)$-dimensional space. Each summand of (1) equals $r + 1/r - 2$ with $r = \mathbf{p}_i(k)/\mathbf{p}_j(k)$, the symmetrised Itakura–Saito divergence of the two variances, so under bounded ratios

$$\frac{\|\mathbf{p}_i - \mathbf{p}_j\|^2}{KB^2} \le S_{ij} \le \frac{B^2\|\mathbf{p}_i - \mathbf{p}_j\|^2}{K}\,: \tag{3}$$

S_{ij} is comparable to a squared Euclidean distance, and E_j to a Euclidean Riesz 2-energy. Each term $r+1/r-2$ is at most $B^2 + B^{-2} - 2$, so $E_j \ge (d-1)/(B^2 + B^{-2} - 2)$ for every configuration and every K. Moreover, $S_{ij} = 0$ iff the two signatures coincide, which is exactly the non-identifiable case of Pham & Cardoso (2001). We write $\bar E = \frac{1}{d}\sum_j E_j$, $E_{\max} = \max_j E_j$, $S_{\mathrm{nn}}(j) = \min_{i\ne j} S_{ij}$ for a fixed source j, and $S_{\min} = \min_j S_{\mathrm{nn}}(j)$. For Gaussian sources this is exact: the Kullback–Leibler divergence between the laws under Q and QR, with R a rotation by θ in the plane (i,j), is $\frac{n}{2}\sin^2\theta\, S_{ij}$, and the information is diagonal across planes and orthogonal to the variances (Lemma 1, Appendix A). Figure 1(a) shows the geometry, and a shell count shows why its dimension matters. The number of signatures at distance t from $\mathbf{p}_j$ grows like t^{K-2}, each contributes t^{-2}, and $\int t^{K-4}\,dt$ converges at zero iff $K \ge 4$.

3. The Riesz rate

3.1. A lower bound for every n and d

Theorem 1 (Riesz rate). *For Gaussian sources and $K \ge 2$ there is $c > 0$, depending only on B, such that for every n and d, every source j and every estimator $\hat Q \in \mathrm{O}(d)$, which may use the signatures,*

$$\sup_{Q\in\mathrm{O}(d)} \mathbb{E}\,\rho_j(\hat Q) \;\ge\; c\,\frac{E_j}{n + E_j}. \tag{4}$$

The same bound holds for ℓ_{A}^2 and for ℓ_{M}^2, with $E_j/(n+E_j)$ replaced by its average and by its maximum over j. If $S_{ij} = 0$ for some $i \ne j$, then $E_j = \infty$ and the bound reads c. (Full statement: Theorem B.1, Appendix B.)

The risk is additive over pairs: E_j/n sums the variances of the rotation angles in the planes (j,i). Since $E_j/(n+E_j) \ge \frac{1}{2}\min(1, E_j/n)$, source j cannot be recovered from fewer than order E_j samples. The constant is qualitative ($c = 3/(4096B^4)$ in the proof). At fixed d, with $S_{ij} > 0$ for all pairs, the rate E_j/n is attained with constant one: for the Gaussian maximum-likelihood joint diagonalizer, $n\rho_j \Rightarrow \sum_{i\ne j} S_{ij}^{-1} Z_i^2$ with Z_i i.i.d. standard

Table 1. What is known about ℓ_{A}^2 and under which conditions (T: Theorem, C: Corollary, P: Proposition). Orders in d for well-spread signatures (row 2: every configuration); constants depend on K, B and the spreading constants. The upper bounds (rows 3–4) are for equal environment sizes. Row 2: no estimator in $\mathrm{O}(d)$ has risk below $c\min(1, f_K(d)/n)$, $f_K(d)$ the listed order; †any mixing for labelled estimators and under the matched loss (Corollary 1). Rows 2–4 are non-asymptotic, 5–6 first order at fixed d, 7 simulation ($\bar{E}$: Section 2; $\bar{E}^{\mathrm{w}}$: Section 5).

	$K=2$	$K=3$	$K\geq 4$	$K\asymp\log d$, sep.	sources	whitening	signatures
energy E_j (T4; Sec. 4)	d^2	$d\log d$	d	d	(geometry of the signatures only)		
necessary: $n\gtrsim$ (T1, C2)	d^2	$d\log d$	d	d	Gaussian	known†	may be used
sufficient: $\asymp \bar{E}/n$ for $n\gtrsim$ (T2, C1)	d^3	d^2	$d^{\frac{K+1}{K-1}}$	$d\log d$	Gaussian	known or est.	known
without signatures (T3)	d^3	d^2	$d^{\frac{K+1}{K-1}}$	$d\log d$	Gaussian	known or est.	$s_0\leq S_{\min}$
one linear pencil (T5)			$\gtrsim d^2/n$ for every K		4th mom.	known	not used
joint diagonalization (P2)			$\asymp \bar{E}/n$		4th mom.	known or est.	not used
Jacobi (simulation, $n\geq 200d$)			$\leq 1.5\,\bar{E}^{\mathrm{w}}/n$, $d\leq 1000$		Gaussian	estimated	not used

Algorithm 1 Quadratic pencil at a probe $g\in\mathbb{R}^K$ (Theorem 2(b); $g=\mathbf{p}_j$ for known signatures)

split each environment's sample into three equal parts 0, 1, 2; $\hat{C}\leftarrow$ pooled covariance of part 0
$\hat{\Sigma}_k^{(i)}\leftarrow$ covariance of part $i\in\{1,2\}$ in environment k, whitened by $\hat{C}^{-1/2}$
$\hat{H}\leftarrow\sum_k g(k)^{-2}\,\mathrm{sym}\big[(\hat{\Sigma}_k^{(1)}-g(k)I)(\hat{\Sigma}_k^{(2)}-g(k)I)\big]$
$\hat{q}\leftarrow$ unit eigenvector of $\hat{H}$ for its smallest eigenvalue (Lanczos: $O(d^2)$ per iteration)
return the unmixing direction $\hat{b}=\hat{C}^{-1/2}\hat{q}$ (with known whitening: no part 0, return $\hat{q}$)

normal (Theorem B.1(b); Flury, 1986; Pham & Cardoso, 2001). This fixes the constant but not the d-dependence. The proof (Appendix B) applies Assouad's lemma to a hypercube of rotations in all planes (j,i), ordered by decreasing S_{ji}. Neighbouring hypotheses then differ by a rotation that only moves sources at least as close to $\mathbf{p}_j$ as the flipped one, and no truncation is needed. For one covariance matrix, lower bounds of this form are due to Wahl (2022; 2023). Here the information is averaged over environments, and the loss matches sources to estimates (Appendix I). The average bound follows by Haar symmetrisation, and all three bounds hold for labelled estimators.

3.2. Attainment by a quadratic pencil

To attain the bound uniformly in d we probe one signature at a time. For source j, let $H_j=\sum_k w_k(\Sigma_k-\mathbf{p}_j(k)I)^2$. Then $H_j=Q\,\mathrm{diag}(\lambda_a)Q^\top$ with $\lambda_a=\sum_k w_k(\mathbf{p}_a(k)-\mathbf{p}_j(k))^2$, zero at $a=j$. With $w_k=\mathbf{p}_j(k)^{-2}$, $B^{-2}KS_{aj}\leq\lambda_a\leq B^2KS_{aj}$: the gap at q_j is the nearest-neighbour divergence, not a spacing on a line. Algorithm 1 computes a sample version from two parts of the data and estimates the whitening on a third.

Theorem 2 (Attainment, uniformly in d; simplified)**.** *Let the sources be Gaussian and the signatures known. Let $\hat{q}_j$ be a unit eigenvector for the smallest eigenvalue of*

$$\hat{H}_j=\sum_k w_k\,\mathrm{sym}\big[(\hat{\Sigma}_k^{(1)}-\mathbf{p}_j(k)I)(\hat{\Sigma}_k^{(2)}-\mathbf{p}_j(k)I)\big],$$

with $w_k=\mathbf{p}_j(k)^{-2}$ and $\hat{\Sigma}_k^{(1)},\hat{\Sigma}_k^{(2)}$ the covariances of two disjoint parts of each environment's sample (halves in (a), thirds in (b)). There are $c,C>0$, depending only on B, such that the following hold if $\gamma_j=K(d+6\log n)/(nS_{\mathrm{nn}}(j))\leq c$.

(a) (Known whitening.) For $A=Q\in\mathrm{O}(d)$, $\mathbb{E}\big[1-\langle\hat{q}_j,q_j\rangle^2\big]\leq C\,E_j/n$.

(b) (Estimated whitening.) Let A be invertible. Split the sample into three equal parts, form the pencil from parts 1 and 2 whitened by the pooled covariance $\hat{C}$ of part 0 (Algorithm 1), and let $\hat{b}_j=\hat{C}^{-1/2}\hat{q}_j$. Then $\mathbb{E}\rho_j\leq C\,E_j/n$, where ρ_j is the share of the pooled variance of the estimated source $\hat{b}_j^\top x$ due to the other sources.

(Full statement, with the first-order coefficients and the cost of estimating the whitening: Theorem C.1, Appendix C.)

The proof rests on three facts. First, the two halves remove the bias of the squared noise. Second, the noise is small *relative* to H_j (cf. Jirak & Wahl, 2018): a direction at divergence λ_a from source j receives noise of order $\sqrt{\lambda_a\nu}$ with $\nu\lesssim K^2(d+\log n)/n$. The noise therefore only has to be small against the gap to the nearest signature, not against the spread of the λ_a. Third, the first-order error in direction a is the (a,j) noise entry divided by λ_a, with variance of order $K/(n\lambda_a)$. The squared errors thus sum to E_j/n, not to $E_j^{(4)}=\sum_{i\neq j}S_{ij}^{-2}$, and the second-order terms have moments that are again energies. Any probe within $1/(8B^4)$ of the nearest-neighbour distance from $\mathbf{p}_j$, independent of the halves, gives the same rate (Lemma 3); Theorem 3 exploits this. In (b), whitening acts on the pencil by a congruence, so the risk does not depend on the mixing. The constants are qualitative: the first-order coefficient E_j^{qp} of Theorem C.1 satisfies $E_j\leq E_j^{\mathrm{qp}}\leq\kappa_B E_j$ with $\kappa_B=\frac{1}{4}(B^2+B^{-2})^2$,

and for the random signatures of Table 2 ($B = 2$) it averages 1.08–1.24 times E_j over sources, against $\kappa_B = 4.5$.

Theorem 3 (Unknown signatures; simplified)**.** *Let the sources be Gaussian, and let only B and a lower bound $s_0 \le S_{\min}$ be known. Algorithm 2 (Appendix D) scores every point g of a grid in signature space by $\lambda_{\min}\big(\sum_k (\hat{\Sigma}_k - g(k)I)^2\big)$ on one part of the sample, and keeps the points below a threshold, which, under the condition below, form one cluster per source. It then runs Algorithm 1 on the other parts, at one probe per cluster. There are $c, C > 0$, depending only on B, such that the following hold for every configuration with $S_{\min} \ge s_0$ if $K(d + 6\log n)/(ns_0) \le c$.*

(a) (Known whitening.) For $A = Q \in \mathrm{O}(d)$, the polar factor $\hat{Q}$ of the d vectors satisfies $\sup_Q \mathbb{E}\ell_{\mathrm{A}}^2(\hat{Q}) \le C\,\bar{E}/n$.

(b) (Estimated whitening.) With the whitening estimated on a further part, $\mathbb{E}\ell_{\mathrm{A}}^2 \le C\,\bar{E}/n$ for every invertible A, where each source is matched to its closest estimate: $\ell_{\mathrm{A}}^2 = \frac{1}{d}\sum_j \min_a(1 - z_{a,j}^2/\|z_a\|^2)$ with $z_a = A^\top \hat{b}_a$.

The grid has at most $(C/\sqrt{s_0})^K$ points. (Full statement: Theorem D.1.)

The grid has $O(d^{K/(K-1)})$ points for well-spread signatures and fixed K, and $d^{O(1)}$ points for separated signatures with $K \asymp \log d$. Algorithm 2 is a construction for polynomial-time attainability, with a qualitative constant, not a practical algorithm. In practice, a coarse grid with local refinement works, with s_0 taken from a Jacobi pilot and reduced by a safety factor (Appendix J).

Corollary 1 (Minimax rate)**.** *Let the sources be Gaussian and $\gamma = K(d + 6\log n)/(nS_{\min}) \le c$. Then, for every configuration, the minimax risk is $\asymp \bar{E}/n$ with constants depending only on B, in three settings. (i) Labelled estimators and any invertible mixing, with ℓ_{A}^2 the average of the ρ_j of Theorem 2(b). (ii) Known whitening, with labelled estimators or estimators in $\mathrm{O}(d)$. (iii) Any invertible mixing and estimators that return d unmixing directions, with the matched loss of Theorem 3(b). The upper bounds are attained by Theorem 2(b) and by the polar factor of the vectors of Theorem 2(a). Without the signatures, the estimator of Theorem 3 with $s_0 = S_{\min}$ attains the bounds of (ii) in $\mathrm{O}(d)$ and of (iii) with the matched loss.*

The lower bound is the average bound of Theorem 1, since $E_{\max} \le d/S_{\min} \le n$ under $\gamma \le c$. The polar factor loses at most a factor 8 (Appendix C). In (iii) the estimates need not be orthogonal; in whitened coordinates their polar factor loses at most a factor 16 in the matched loss (Lemma 5, Appendix D).

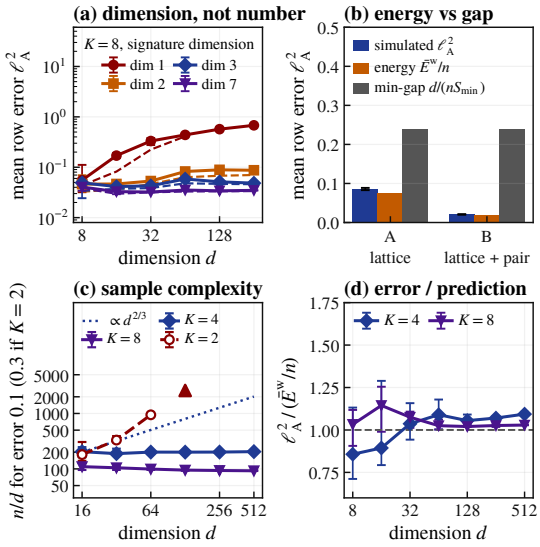

Figure 2. **The mechanism. (a)** *Dimension, not number:* $K = 8$, signatures well-spread (Halton points, fixed scale) in an affine subspace of dimension 1, 2, 3 or 7; $n = 1600d$; dashed: $\bar{E}^{\mathrm{w}}/n$. **(b)** *Energy beats the smallest gap:* two configurations with the same $S_{\min}$ ($K = 4$, $d = 27$, $n = 400d$), a lattice (A) and a coarser lattice plus one close pair (B): simulated ℓ_{A}^2 with standard errors, $\bar{E}^{\mathrm{w}}/n$ and $(d/n)/S_{\min}$. **(c)** *Sample complexity:* n/d at which the median ℓ_{A}^2 reaches 0.1 ($K = 4, 8$) or 0.3 ($K = 2$; triangle: not reached by $n = 1600d$), from Figure 8, with 90% bootstrap intervals; dotted: the growth $d^{2/(K-1)}$ of the proved threshold, $K = 4$. **(d)** *First order explains the curves:* error divided by $\bar{E}^{\mathrm{w}}/n$ at $n = 400d$, per draw (draws: Appendix J).

4. The dimension law

Riesz energies $\sum_{i \ne j} \|\mathbf{p}_i - \mathbf{p}_j\|^{-\beta}$ of N points on a D-dimensional set have three regimes (Hardin & Saff, 2004; 2005; Borodachov et al., 2019). The total energy grows like N^2 for $\beta < D$, like $N^2 \log N$ for $\beta = D$, and like $N^{1+\beta/D}$ for $\beta > D$, where nearest neighbours dominate (the shell count of Section 2). Here $\beta = 2$, $D = K - 1$ and $N = d$.

Definition 1 (Well-spread signatures)**.** Let $r_d = d^{-1/(K-1)}$ and let Ω be a fixed convex body in $\{\sum_k \mathbf{p}(k) = K\} \cap [1/B, B]^K$ containing the signatures. The signatures are *well-spread* with constants (c_1, C_1) if they are $c_1 r_d$-separated and every point of Ω is within $C_1 r_d$ of one of them.

The covering condition rules out voids, which would make the per-source lower bound fail for $K \le 3$; lattices in a box satisfy it.

Theorem 4 (Dimension law)**.** *Fix K. For well-spread signatures, uniformly in j and with constants depending only*

on (K, B, c_1, C_1, Ω),

$$E_j \asymp \begin{cases} d^2, & K = 2, \\ d \log d, & K = 3, \\ d, & K \geq 4. \end{cases} \tag{5}$$

The upper bounds need only the separation $c_1 r_d$*, and the lower bounds for* $\bar{E}$ *hold for* every *configuration, by minimal-energy asymptotics (Hardin & Saff, 2005).*

The constants depend on K (Figure 1(c) shows the three regimes on lattices): a lattice needs m^{K-1} points, so Definition 1 needs $2^{K-1} \lesssim d$. For $K \geq 4$ the lower bound is the floor $E_j \geq (d-1)/(B^2 + B^{-2} - 2)$, that is, parameter counting; the geometry enters the upper bound, and the lower bounds only for $K = 2, 3$.

Remark 1 (Dimension, not number)**.** If the signatures lie in an affine subspace of dimension D', Theorem 4 holds with D' in place of $K - 1$ (Definition 1 taken in that subspace, $r_d = d^{-1/D'}$). Corollary 2 then holds with $D' + 1$ for K. Appendix E reads an effective dimension at each scale.

Combined with Theorem 1, the lower bounds give necessary sample sizes for every configuration.

Corollary 2 (Necessary sample size)**.** *Let* $d \geq 2$ *and* $f_K(d) = d^2$*,* $d \log d$*,* d *for* $K = 2$*,* 3*,* ≥ 4*. For Gaussian sources there is* $c > 0$ *depending only on* (K, B) *such that, for every configuration of signatures and every estimator* $\hat{Q} \in \mathrm{O}(d)$*,*

$$\sup_{Q} \mathbb{E}\, \ell_{\mathrm{A}}^2(\hat{Q}) \;\geq\; c \min\big(1,\; f_K(d)/n\big).$$

No estimator recovers an average source consistently unless $n/f_K(d) \to \infty$*, however the environments are chosen (with variance ratios in* $[1/B, B]$*).*

The proof (Appendix E) shows that in every configuration at least half of the sources have $E_j \gtrsim f_K(d)$. For well-spread signatures $\bar{E} \asymp f_K(d)$, so the bound has the order of the maximum-likelihood limit $\bar{E}/n$. For i.i.d. random signatures with a bounded density, $\bar{E}$ has the same orders for $K \geq 3$ with high probability (d^3 for $K = 2$). The worst source, however, needs five environments to be linear: $E_{\max} \asymp d^{4/3}$ for $K = 4$ (Proposition 1, Appendix E).

Sufficient sample sizes. Corollary 1 gives the other side. Well-spread signatures have $S_{\min} \asymp d^{-2/(K-1)}$, so the minimax risk is $\asymp \bar{E}/n$, of order d/n for $K \geq 4$, once $n \gtrsim d^{(K+1)/(K-1)}$. Separated signatures, with $S_{\min} \geq s_0$ for a constant s_0 (unlike Definition 1's $c_1 r_d$-separation, whose scale shrinks with d), need $K \gtrsim \log d$ by packing. And $K \asymp \log d$ suffices. With balanced binary codes as signatures, every pair of sources differs by a fixed factor in a constant fraction of the environments ($K \geq 32 \log d$, Appendix C). Then $E_j \asymp d$, and the minimax risk is $\asymp d/n$ once $n \gtrsim K(d + \log n)/s_0$. Thus $\log d$ environments make the sample size linear in d up to a logarithm: $n \gtrsim d$ is necessary and $n \gtrsim d \log d$ sufficient (Corollary 2, Theorems 2–3).

Necessary versus sufficient. For fixed K, the proved sufficient size $d^{(K+1)/(K-1)}$ (Table 1) exceeds the necessary $f_K(d)$ by a power of d. The sufficient side is the condition $\gamma_j \leq c$ of Theorem 2. For well-spread signatures and $K \geq 4$ it says that E_j/n is below a fraction of $S_{\mathrm{nn}}(j)/K$ (the factor K partly reflects the two halves' noise product, Appendix H). By (4) no estimator is that accurate with fewer samples: a risk below $\eta S_{\mathrm{nn}}(j)$ at source j needs $n \geq cE_j/(2\eta S_{\mathrm{nn}}(j))$ (Proposition 3, Appendix H). For $K \geq 4$ and $\eta \asymp 1/K$ this is Corollary 1's threshold, which suggests why Theorem 2's operator-norm argument stops there.

Open problem. For well-spread signatures and fixed K, is the Riesz rate attained between the necessary sample size $f_K(d)$ and the proved sufficient size $d^{(K+1)/(K-1)}$? Below that threshold the quadratic pencil drifts (Appendix H), while joint diagonalization shows no such threshold in simulations (Section 6). For $K \geq 4$ we conjecture (for $K = 2, 3$ we make no conjecture):

Conjecture 1 (Linear regime)**.** *For well-spread signatures and* $K \geq 4$*, the output* $\hat{U}$ *of Jacobi sweeps from the identity attains the Riesz rate,* $\mathbb{E}\, \ell_{\mathrm{A}}^2(\hat{U}) \leq C\, \bar{E}/n$*, as soon as* $n \geq C'\, d\, \mathrm{polylog}(d)$*.*

5. Projecting versus probing

A common way to use several covariances is to take the eigenvectors of one matrix $\hat{M}_\alpha = \sum_k \alpha_k \hat{\Sigma}_k$, for a fixed or random $\alpha \in \mathbb{R}^K$. Two-environment generalised eigenproblems, AMUSE, and the random-pencil initialisations of He & Kressner (2024a;b) are of this form.

Theorem 5 (A linear pencil pays d^2 at first order)**.** *Assume known whitening. Fix* $\alpha \in \mathbb{R}^K$*, or draw it independently of the data, with the eigenvalues* $\lambda_j = \langle \alpha, \mathbf{p}_j \rangle$ *of* $M_\alpha = Q\, \mathrm{diag}(\lambda) Q^\top$ *pairwise distinct. Let* $\hat{q}_j$ *be the eigenvector of* $\hat{M}_\alpha$ *for its* j*-th eigenvalue in the order of the* λ_j*, and let* $\rho_j(\hat{q}_j) = 1 - \langle \hat{q}_j, q_j \rangle^2$*. As* $n \to \infty$ *at fixed* d*,*

$$n\, \rho_j(\hat{q}_j) \Rightarrow \sum_{i \neq j} \frac{v_{ij}(\alpha)}{(\lambda_i - \lambda_j)^2}\, Z_i^2,$$

with $v_{ij}(\alpha) = K \sum_k \alpha_k^2\, \mathbf{p}_i(k) \mathbf{p}_j(k)$ *and* Z_i *i.i.d. standard normal. For every configuration and every* K*, the means of these limit laws average over* j *to at least* $(d-1)^3/(dB^4)$*, of order* d^2 *uniformly in* K *and* α*. (The sharper bound (22), in terms of the range of the* λ_j*, is in Appendix F.)*

The proof (Appendix F) combines the eigenvector expansion of Anderson (1963) with a power-mean inequality on the spacings of d points in an interval. A linear pencil collapses the configuration onto the line spanned by α, where the Riesz 2-energy is always hypersingular ($\beta = 2 > D = 1$). The barrier is the projection, not the use of a pencil. For $K = 2$ nothing is lost; for well-spread signatures the projection loses a factor of order at least $d/\log d$ ($K = 3$) and d ($K \geq 4$). The theorem concerns the limit law of $n\rho_j$ at fixed d, not $\mathbb{E}\rho_j$, and one combination that is fixed or drawn independently of the data; adaptively chosen or iterated combinations are not covered. The quadratic pencil of Theorem 2 probes instead. Its eigenvalues are the divergences of all signatures from $\mathbf{p}_j$, in all $K-1$ directions, and $E_j/B^2 \leq K \operatorname{tr} H_j^+ \leq B^2 E_j$ for its pseudo-inverse H_j^+.

Joint diagonalization. Orthogonal joint diagonalization (Bunse-Gerstner et al., 1993; Cardoso & Souloumiac, 1996) minimises $\hat{f}(U) = \sum_{k=1}^{K} \| \operatorname{off}(U^\top \hat{\Sigma}_k U)\|_F^2$ over $U \in \mathrm{O}(d)$, $\operatorname{off}(X)$ the off-diagonal part of X. The weighted criterion of Tichavský & Yeredor (2009) and the Gaussian likelihood of Pham & Cardoso (2001) replace the Frobenius norm by Fisher weights. For independent sources with finite fourth moments and $S_{ij} > 0$, the global minimiser $\hat{U}$ satisfies $n\rho_j(\hat{U}) \Rightarrow \sum_{i \neq j} v_{ij} Z_i^2$ at fixed d. The coefficients are explicit and, for known and for estimated whitening alike, satisfy $S_{ij}^{-1} \leq v_{ij} \leq C_B S_{ij}^{-1}$ with C_B depending only on B. With known whitening, Fisher weights give $v_{ij} = S_{ij}^{-1}$ (Proposition 2, Appendix G). Both facts are classical at fixed d (Cardoso, 1994; Yeredor, 2000; Pham & Cardoso, 2001; Tichavský & Yeredor, 2009). What is new is that the coefficients are of Riesz type, so the dimension law applies to them. We write $\bar{E}^{\mathrm{w}}/n$ for the resulting first-order prediction with estimated whitening, the pipeline of Section 6.

6. Experiments

Setup. All experiments use the model of Section 2 with a Haar-distributed A, n/K samples per environment and covariances drawn from their Wishart law. Details, draw counts, code and raw results: Appendix J and the supplement. The default pipeline whitens with the pooled sample covariance $\hat{C}$ and runs Jacobi sweeps (Cardoso & Souloumiac, 1996) from the identity (Algorithm 3; at most 100 sweeps). The estimators of Theorems 2 and 3 use the true whitening unless stated. The loss is ℓ_{A}^2 of the unit-normalised rows of $\hat{U}^\top \hat{C}^{-1/2}$ against the columns of A, in observation coordinates. Predictions use v^{w} of (24) (Appendix G). The range $\bar{E}^{\mathrm{w}}/n \lesssim 0.1$, where Jacobi follows this prediction, is the first-order regime. Random signatures are i.i.d. uniform on $\{\mathbf{p} \in [\frac{1}{2}, \frac{3}{2}]^K : \sum_k \mathbf{p}(k) = K\}$ ($B = 2$).

The error follows the dimension of the signatures, not the number of environments. At $n = 400d$, from $d = 8$ to 128, the error grows from 0.14 to 0.52 for $K = 2$ and from 0.072 to 0.14 for $K = 3$, but stays between 0.018 and 0.022 for $K = 8$ (Figure 1(b)). The rise for $K = 4$ is first order: 0.86–1.09 times $\bar{E}^{\mathrm{w}}/n$ up to $d = 512$ (Figure 2(d); 1.02–1.14 for $K = 8$). With $K = 8$, $n = 1600d$ and signatures in a subspace of dimension 1, 2, 3 or 7, the error follows the dimension. From $d = 8$ to 256 it grows from 0.055 to 0.68 in dimension 1, and stays near 0.049 and 0.035 in dimensions 3 and 7 (Figure 2(a)).

Energy, not the smallest gap, predicts the error. Two configurations with equal $S_{\min}$ and energies 3.8 times apart have errors 4.1 times apart, within 8% and 1% of $\bar{E}^{\mathrm{w}}/n$ (Figure 2(b)).

For $K \geq 4$, joint diagonalization stays on its prediction down to $n \approx 200d$, and the n/d it needs barely grows with d. The sample size at which ℓ_{A}^2 reaches 0.1 is $92d$–$110d$ for $K = 8$ and $188d$–$205d$ for $K = 4$, from $d = 16$ to 512 (Figure 2(c)). The proved threshold $d^{(K+1)/(K-1)}$ would instead multiply n/d by 2.7 and 10.1. On fixed $K = 4$ lattices, with only the mixing and the data redrawn, the mean error reaches 0.1 at $n \approx 360d$–$490d$ from $d = 27$ to 1000 (Appendix J, Figure 7); their energy per source is about three times that of the random signatures of Figure 2(c). There n/d grows by 1.4, against 1.5 for $\bar{E}^{\mathrm{w}}/d$ and 11.1 for the proved threshold. Since 0.1 lies in the first-order regime, this is the evidence for Conjecture 1. For $K = 2$, the n needed for an error of 0.3 grows with d (Figure 2(c)). Below $100d$ ($\bar{E}^{\mathrm{w}}/n \gtrsim 0.15$) longer runs and restarts do not change the error. Pham's likelihood algorithm is 28%–34% more accurate at $K = 8$, $n = 50d$, yet above $\bar{E}^{\mathrm{w}}/n$ (Figure 4(d), Appendix J).

The estimator of Theorem 2 attains its rate, also with estimated whitening; a practical variant of the grid of Theorem 3 finds the signatures. With $K = 16$ code signatures, the estimator of Theorem 2 gives the same error at the same n/d for $d = 64$ to 512. It approaches its first-order prediction (1.06–1.07 times it at $n = 800d$; Figure 4(b)). With the whitening estimated on a third of the sample, the error (the loss ρ_j of Theorem 2(b)) is the same for every mixing. It is within 2% of its first-order prediction per source at $\gamma \leq 0.03$ (Appendix J, Table 6). Without the signatures, a coarse grid with three refinement steps finds all d probes in every run with $n \geq 6400d$ ($K = 4$ lattices, $d \leq 216$). The estimator is then within 4% of the same estimator at the true signatures (Appendix J, Table 7). Below its proved regime, the estimator of Theorem 2 stays within a small factor of $\bar{E}/n$ (Appendix J, Table 5).

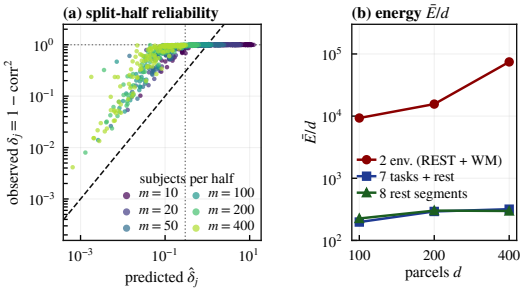

Figure 3. **Real data (HCP task fMRI, $d = 200$). (a)** Split-half discrepancy $\delta_j = 1 - \mathrm{corr}^2$ of each component between two disjoint groups of m subjects, against its first-order prediction (colour: subjects per half; dashed: $y = x$; dotted: $\hat{\delta} = 0.3$ and $\delta = 1$). **(b)** $\bar{E}/d$ against d for two environments (rest and working memory; noise-limited), the seven tasks and rest, and eight rest segments.

One linear pencil degrades with d; every joint diagonalizer sits on its first-order prediction. On identical data ($K = 8$, $n = 400d$; Figures 1(d) and 4(a), Table 2), Pham's algorithm (Pham, 2001) is within a factor 1.00–1.15 of its Cramér–Rao bound, while one linear pencil degrades from 0.47 to 0.84. As an initialisation, however, the pencil costs nothing. Three Jacobi sweeps from its eigenvectors come within 13% (in median, $d \leq 256$) of the converged run, which is the same minimiser as from the identity (Appendix J). Appendix J adds $B = 10$, unequal sizes, non-Gaussian sources, departures from the model and, outside the theorems, the lag route: on MA(8) sources AMUSE degrades with d, SOBI with ≥ 3 lags does not.

Real data: the energy says which fMRI components are hard to separate, not how large their error is. We separate the cortical fMRI of 813 Human Connectome Project (HCP) subjects (Van Essen et al., 2013; Barch et al., 2013) into $d = 200$ temporally independent components (Schaefer parcels, Schaefer et al., 2018; cf. Smith et al., 2012). The seven tasks and rest are the $K = 8$ environments, each with 352 frames per subject (Appendix J). Reliability is the split-half discrepancy δ_j of a component between two disjoint groups of m subjects (0: identical, 1: unrelated). The first-order theory predicts it from the estimated signatures, with an autocorrelation correction ($\hat{\delta}_j$, Appendix J). The signatures vary little across tasks ($\bar{E}/d = 292$), so the prediction is in its first-order range, below 0.3, for only 26% of the components with 50 subjects per half. Where it is, it ranks the observed discrepancies (Spearman 0.92), and the energy ranks them better than the nearest-neighbour gap alone (0.89 for E_j, 0.72 for $1/S_{\mathrm{nn}}(j)$). It does not predict their size: with 400 subjects per half the mean discrepancy is 0.723 against a predicted 0.118. Fitted over m, the discrepancies have a floor (median 0.59) that carries most of the ranking, and two joint diagonalizers disagree on near-degenerate pairs (Appendix J). A common mixing of independent sources thus describes these components only roughly. With rest and working memory as the only environments the energy is about 53 times larger ($\bar{E}/d = 15608$, noise-limited). Eight rest segments, whose signatures vary mainly along one direction (participation ratio 1.54 against 4.69), carry about as much energy as the seven tasks and rest. The energy depends on the nearest-neighbour spread, not on the global one (Appendix E; Figure 3).

7. Discussion

Diversity versus non-Gaussianity. In the non-Gaussian route the information cost is of order d, but low-degree estimators pay d^2 (Auddy & Yuan, 2025). In the diversity route it is the Riesz energy: d^2 for two environments and d for four well-spread ones. It is attained in polynomial time for fixed K once $n \gtrsim d^{(K+1)/(K-1)}$, given the signatures or a separation bound (below that: the open problem of Section 4). Two environments suffice for identifiability (Pham & Cardoso, 2001); for efficiency the signatures' dimension counts. Our codewords are analogous to the $K \asymp \log d$ unknown interventions that identify the model in linear causal representation learning (Lee et al., 2026), under different assumptions and with a sample size attached.

Practical takeaways. (1) Read the signatures off a Jacobi pilot and compute $\bar{E}^{\mathrm{w}}/n$; in our simulations joint diagonalization follows it below about 0.1, so it is a sample-size diagnostic. (2) Choose environments or lags that spread the signatures in new directions; eight environments on a line are worth two. (3) $K \asymp \log d$ conditions in which every pair of sources differs by a fixed factor in a constant fraction of them make $n \gtrsim d \log d$ samples enough for the pencil of Theorem 2. (4) Do not take the eigenvectors of one linear combination of the covariances as the final estimate in high dimension; as an initialisation they cost nothing in our simulations. (5) Autocorrelation rescales the error, a shared factor or subject-specific mixings add a floor, and the ranking of the sources survives (Appendix J).

Limitations. The scope is deliberate (Section 1). Theorems 2–3 assume Gaussian sources and equal environment sizes (Section 2). Theorem 3 needs the oracle s_0 (not adaptively), and its grid is polynomial in d only when $K \log(C/s_0) = O(\log d)$. For fixed K, the gap between the necessary and the proved sufficient sample size is open (Section 4, Conjecture 1). Practical algorithms are covered at first order only (Proposition 2). Independence excludes lags and spatially correlated parcels. The fMRI analysis assumes a mixing shared across subjects, corrects autocorrelation only through Bartlett factors, may place relatives on both sides of a split, and tests the energy as a diagnostic, not the theorems.

Impact Statement

We foresee no negative societal consequences specific to this methodological work on source separation; the fMRI data are the open-access HCP data, used under their data use terms, and no attempt is made to identify subjects.

Acknowledgements

Data were provided by the Human Connectome Project, WU-Minn Consortium (Principal Investigators: David Van Essen and Kamil Ugurbil; 1U54MH091657) funded by the 16 NIH Institutes and Centers that support the NIH Blueprint for Neuroscience Research, and by the McDonnell Center for Systems Neuroscience at Washington University.

References

Acartürk, E., Varıcı, B., Shanmugam, K., and Tajer, A. Sample complexity of interventional causal representation learning. In *Advances in Neural Information Processing Systems*, 2024.

Adalı, T., Anderson, M., and Fu, G.-S. Diversity in independent component and vector analyses: Identifiability, algorithms, and applications in medical imaging. *IEEE Signal Processing Magazine*, 31(3):18–33, 2014.

Afsari, B. What can make joint diagonalization difficult? In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2007.

Afsari, B. Sensitivity analysis for the problem of matrix joint diagonalization. *SIAM Journal on Matrix Analysis and Applications*, 30(3):1148–1171, 2008.

Amari, S.-i. and Cardoso, J.-F. Blind source separation—semiparametric statistical approach. *IEEE Transactions on Signal Processing*, 45(11):2692–2700, 1997.

Anderson, M., Fu, G.-S., Phlypo, R., and Adalı, T. Independent vector analysis: Identification conditions and performance bounds. *IEEE Transactions on Signal Processing*, 62(17):4399–4410, 2014.

Anderson, T. W. Asymptotic theory for principal component analysis. *The Annals of Mathematical Statistics*, 34(1): 122–148, 1963.

Auddy, A. and Yuan, M. Large-dimensional independent component analysis: Statistical optimality and computational tractability. *The Annals of Statistics*, 53(2):477–505, 2025.

Barch, D. M., Burgess, G. C., Harms, M. P., et al. Function in the human connectome: Task-fMRI and individual differences in behavior. *NeuroImage*, 80:169–189, 2013.

Bartlett, M. S. Some aspects of the time-correlation problem in regard to tests of significance. *Journal of the Royal Statistical Society*, 98(3):536–543, 1935.

Belouchrani, A., Abed-Meraim, K., Cardoso, J.-F., and Moulines, E. A blind source separation technique using second-order statistics. *IEEE Transactions on Signal Processing*, 45(2):434–444, 1997.

Bhaskara, A., Charikar, M., Moitra, A., and Vijayaraghavan, A. Smoothed analysis of tensor decompositions. In *Proceedings of the 46th Annual ACM Symposium on Theory of Computing (STOC)*, 2014.

Borodachov, S. V., Hardin, D. P., and Saff, E. B. *Discrete energy on rectifiable sets*. Springer, 2019.

Buchholz, S., Rajendran, G., Rosenfeld, E., Aragam, B., Schölkopf, B., and Ravikumar, P. Learning linear causal representations from interventions under general nonlinear mixing. In *Advances in Neural Information Processing Systems*, volume 36, 2023.

Bunse-Gerstner, A., Byers, R., and Mehrmann, V. Numerical methods for simultaneous diagonalization. *SIAM Journal on Matrix Analysis and Applications*, 14(4):927–949, 1993.

Cai, T. T., Ma, Z., and Wu, Y. Sparse PCA: Optimal rates and adaptive estimation. *The Annals of Statistics*, 41(6): 3074–3110, 2013.

Calhoun, V. D. and Adalı, T. Multisubject independent component analysis of fMRI: A decade of intrinsic networks, default mode, and neurodiagnostic discovery. *IEEE Reviews in Biomedical Engineering*, 5:60–73, 2012.

Calhoun, V. D., Adalı, T., Pearlson, G. D., and Pekar, J. J. A method for making group inferences from functional MRI data using independent component analysis. *Human Brain Mapping*, 14(3):140–151, 2001.

Cardoso, J.-F. Perturbation of joint diagonalizers. Technical Report 94D027, Télécom Paris, 1994.

Cardoso, J.-F. and Souloumiac, A. Jacobi angles for simultaneous diagonalization. *SIAM Journal on Matrix Analysis and Applications*, 17(1):161–164, 1996.

Cerjan, A., Loring, T. A., and Vides, F. Quadratic pseudospectrum for identifying localized states. *Journal of Mathematical Physics*, 64(2):023501, 2023.

Chang, J., Guo, B., and Yao, Q. Principal component analysis for second-order stationary vector time series. *The Annals of Statistics*, 46(5):2094–2124, 2018.

Chen, A. and Bickel, P. J. Efficient independent component analysis. *The Annals of Statistics*, 34(6):2825–2855, 2006.

Comon, P. Independent component analysis, a new concept? *Signal Processing*, 36(3):287–314, 1994.

Damian, A., Pillaud-Vivien, L., Lee, J. D., and Bruna, J. Computational-statistical gaps in Gaussian single-index models. In *Conference on Learning Theory*, 2024.

Davis, C. and Kahan, W. M. The rotation of eigenvectors by a perturbation. III. *SIAM Journal on Numerical Analysis*, 7(1):1–46, 1970.

Diakonikolas, I., Kane, D. M., and Stewart, A. Statistical query lower bounds for robust estimation of high-dimensional Gaussians and Gaussian mixtures. In *IEEE Symposium on Foundations of Computer Science*, 2017.

Diakonikolas, I., Hopkins, S. B., Pensia, A., and Tiegel, S. Sum-of-squares lower bounds for non-Gaussian component analysis. *arXiv preprint arXiv:2410.21426*, 2024.

Du, Y., Fu, Z., Sui, J., Gao, S., Xing, Y., Lin, D., Salman, M., Abrol, A., Rahaman, M. A., Chen, J., Hong, L. E., Kochunov, P., Osuch, E. A., and Calhoun, V. D. NeuroMark: An automated and adaptive ICA based pipeline to identify reproducible fMRI markers of brain disorders. *NeuroImage: Clinical*, 28:102375, 2020.

Flury, B. N. Common principal components in k groups. *Journal of the American Statistical Association*, 79(388): 892–898, 1984.

Flury, B. N. Asymptotic theory for common principal component analysis. *The Annals of Statistics*, 14(2):497–507, 1986.

Genc, E. I., Demir, S., and Dogan, Z. Learnability and competition in high-dimensional multi-component ICA. *arXiv preprint arXiv:2605.08552*, 2026.

Glasser, M. F., Sotiropoulos, S. N., Wilson, J. A., Coalson, T. S., Fischl, B., Andersson, J. L., Xu, J., Jbabdi, S., Webster, M., Polimeni, J. R., Van Essen, D. C., and Jenkinson, M. The minimal preprocessing pipelines for the Human Connectome Project. *NeuroImage*, 80:105–124, 2013.

Goyal, N., Vempala, S., and Xiao, Y. Fourier PCA and robust tensor decomposition. In *Proceedings of the 46th Annual ACM Symposium on Theory of Computing (STOC)*, 2014.

Hälvä, H. and Hyvärinen, A. Hidden Markov nonlinear ICA: Unsupervised learning from nonstationary time series. In *Proceedings of the 36th Conference on Uncertainty in Artificial Intelligence*, volume 124 of *PMLR*, pp. 939–948, 2020.

Hardin, D. P. and Saff, E. B. Discretizing manifolds via minimum energy points. *Notices of the AMS*, 51(10): 1186–1194, 2004.

Hardin, D. P. and Saff, E. B. Minimal Riesz energy point configurations for rectifiable d-dimensional manifolds. *Advances in Mathematics*, 193(1):174–204, 2005.

He, H. and Kressner, D. Randomized joint diagonalization of symmetric matrices. *SIAM Journal on Matrix Analysis and Applications*, 45(1):661–684, 2024a.

He, H. and Kressner, D. A randomized algorithm for simultaneously diagonalizing symmetric matrices by congruence. *arXiv preprint arXiv:2402.16557*, 2024b.

Hoeffding, W. Probability inequalities for sums of bounded random variables. *Journal of the American Statistical Association*, 58(301):13–30, 1963.

Hyvärinen, A. and Morioka, H. Unsupervised feature extraction by time-contrastive learning and nonlinear ICA. In *Advances in Neural Information Processing Systems*, 2016.

Hyvärinen, A., Karhunen, J., and Oja, E. *Independent Component Analysis*. Wiley, New York, 2001.

Hyvärinen, A., Sasaki, H., and Turner, R. Nonlinear ICA using auxiliary variables and generalized contrastive learning. In *Artificial Intelligence and Statistics*, 2019.

Hyvärinen, A., Khemakhem, I., and Morioka, H. Nonlinear independent component analysis for principled disentanglement in unsupervised deep learning. *Patterns*, 4(10), 2023.

Ipsen, I. C. F. Relative perturbation results for matrix eigenvalues and singular values. *Acta Numerica*, 7:151–201, 1998.

Iraji, A., Fu, Z., Faghiri, A., Duda, M., Chen, J., Rachakonda, S., DeRamus, T., Kochunov, P., et al. Identifying canonical and replicable multi-scale intrinsic connectivity networks in 100k+ resting-state fMRI datasets. *Human Brain Mapping*, 44(17):5729–5748, 2023.

Jiang, Y. Finite-sample analysis of nonlinear independent component analysis: Sample complexity and identifiability bounds. *arXiv preprint arXiv:2604.08850*, 2026.

Jin, J. and Syrgkanis, V. Learning causal representations from general environments: Identifiability and intrinsic ambiguity. In *Advances in Neural Information Processing Systems*, volume 37, 2024.

Jirak, M. and Wahl, M. Relative perturbation bounds with applications to empirical covariance operators. *arXiv preprint arXiv:1802.02869*, 2018.

Khemakhem, I., Kingma, D., Monti, R., and Hyvärinen, A. Variational autoencoders and nonlinear ICA: A unifying framework. In *Artificial Intelligence and Statistics*, 2020.

Kim, T., Eltoft, T., and Lee, T.-W. Independent vector analysis: An extension of ICA to multivariate components. In *Independent Component Analysis and Blind Signal Separation (ICA 2006)*, volume 3889 of *Lecture Notes in Computer Science*, pp. 165–172. Springer, 2006.

Kuijlaars, A. B. J. and Saff, E. B. Asymptotics for minimal discrete energy on the sphere. *Transactions of the American Mathematical Society*, 350(2):523–538, 1998.

Kuleshov, V., Chaganty, A., and Liang, P. Tensor factorization via matrix factorization. In *Artificial Intelligence and Statistics*, pp. 507–516, 2015.

Lee, I., Jin, T., and Aragam, B. Beyond identifiability: Learning causal representations with few environments and finite samples. *arXiv preprint arXiv:2603.25796*, 2026.

Loring, T. A. and Vides, F. Computing truncated joint approximate eigenbases for model order reduction. *arXiv preprint arXiv:2201.05928*, 2022.

Matsuoka, K., Ohya, M., and Kawamoto, M. A neural net for blind separation of nonstationary signals. *Neural Networks*, 8(3):411–419, 1995.

Miettinen, J., Nordhausen, K., Oja, H., and Taskinen, S. Statistical properties of a blind source separation estimator for stationary time series. *Statistics & Probability Letters*, 82(11):1865–1873, 2012.

Miettinen, J., Illner, K., Nordhausen, K., Oja, H., Taskinen, S., and Theis, F. J. Separation of uncorrelated stationary time series using autocovariance matrices. *Journal of Time Series Analysis*, 37(3):337–354, 2016.

Nordhausen, K. and Ruiz-Gazen, A. On the usage of joint diagonalization in multivariate statistics. *Journal of Multivariate Analysis*, 188:104844, 2022.

Pham, D.-T. Joint approximate diagonalization of positive definite Hermitian matrices. *SIAM Journal on Matrix Analysis and Applications*, 22(4):1136–1152, 2001.

Pham, D.-T. and Cardoso, J.-F. Blind separation of instantaneous mixtures of nonstationary sources. *IEEE Transactions on Signal Processing*, 49(9):1837–1848, 2001.

Rathod, A., Belsare, S., Nautiyal, S. K., Laad, D., and Kumaraguru, P. Sample complexity of causal identification with temporal heterogeneity. *arXiv preprint arXiv:2602.06899*, 2026.

Ricci, F., Bardone, L., and Goldt, S. Feature learning from non-Gaussian inputs: the case of independent component analysis in high dimensions. *arXiv preprint arXiv:2503.23896*, 2025.

Schaefer, A., Kong, R., Gordon, E. M., Laumann, T. O., Zuo, X.-N., Holmes, A. J., Eickhoff, S. B., and Yeo, B. T. T. Local-global parcellation of the human cerebral cortex from intrinsic functional connectivity MRI. *Cerebral Cortex*, 28(9):3095–3114, 2018.

Smith, S. M., Fox, P. T., Miller, K. L., Glahn, D. C., Fox, P. M., Mackay, C. E., Filippini, N., Watkins, K. E., Toro, R., Laird, A. R., and Beckmann, C. F. Correspondence of the brain's functional architecture during activation and rest. *Proceedings of the National Academy of Sciences*, 106(31):13040–13045, 2009.

Smith, S. M., Miller, K. L., Moeller, S., Xu, J., Auerbach, E. J., Woolrich, M. W., Beckmann, C. F., Jenkinson, M., Andersson, J., Glasser, M. F., Van Essen, D. C., Feinberg, D. A., Yacoub, E. S., and Ugurbil, K. Temporally-independent functional modes of spontaneous brain activity. *Proceedings of the National Academy of Sciences*, 109(8):3131–3136, 2012.

Squires, C., Seigal, A., Bhate, S. S., and Uhler, C. Linear causal disentanglement via interventions. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *PMLR*, pp. 32540–32560, 2023.

Tichavský, P. and Yeredor, A. Fast approximate joint diagonalization incorporating weight matrices. *IEEE Transactions on Signal Processing*, 57(3):878–891, 2009.

Tong, L., Liu, R.-W., Soon, V. C., and Huang, Y.-F. Indeterminacy and identifiability of blind identification. *IEEE Transactions on Circuits and Systems*, 38(5):499–509, 1991.

Tripuraneni, N., Jin, C., and Jordan, M. I. Provable meta-learning of linear representations. In *International Conference on Machine Learning*, 2021.

Troedsson, E., Falkowski, D., Lidgren, C.-F., Wendt, H., and Carlsson, M. On joint eigen-decomposition of matrices. *arXiv preprint arXiv:2409.10292*, 2024.

Tsybakov, A. B. *Introduction to nonparametric estimation*. Springer Series in Statistics. Springer, 2009.

Usevich, K., Li, J., and Comon, P. Approximate matrix and tensor diagonalization by unitary transformations: convergence of Jacobi-type algorithms. *SIAM Journal on Optimization*, 30(4):2998–3028, 2020.

Van Essen, D. C., Smith, S. M., Barch, D. M., Behrens, T. E. J., Yacoub, E., and Ugurbil, K. The WU-Minn

Human Connectome Project: An overview. *NeuroImage*, 80:62–79, 2013.

Vershynin, R. Introduction to the non-asymptotic analysis of random matrices. In Eldar, Y. C. and Kutyniok, G. (eds.), *Compressed Sensing: Theory and Applications*, pp. 210–268. Cambridge University Press, 2012.

von Kügelgen, J., Besserve, M., Wendong, L., Gresele, L., Kekić, A., Bareinboim, E., Blei, D. M., and Schölkopf, B. Nonparametric identifiability of causal representations from unknown interventions. In *Advances in Neural Information Processing Systems*, volume 36, 2023.

Vu, V. Q. and Lei, J. Minimax sparse principal subspace estimation in high dimensions. *The Annals of Statistics*, 41(6):2905–2947, 2013.

Wahl, M. Lower bounds for invariant statistical models with applications to principal component analysis. *Annales de l'Institut Henri Poincaré, Probabilités et Statistiques*, 58 (3):1565–1589, 2022.

Wahl, M. Van Trees inequality, group equivariance, and estimation of principal subspaces. In *Foundations of Modern Statistics*, Springer Proceedings in Mathematics & Statistics, pp. 301–321. Springer, 2023.

Yeredor, A. Blind separation of Gaussian sources via second-order statistics with asymptotically optimal weighting. *IEEE Signal Processing Letters*, 7(7):197–200, 2000.

Yeredor, A. Non-orthogonal joint diagonalization in the least-squares sense with application in blind source separation. *IEEE Transactions on Signal Processing*, 50(7): 1545–1553, 2002.

Yeredor, A. Blind separation of Gaussian sources with general covariance structures: Bounds and optimal estimation. *IEEE Transactions on Signal Processing*, 58(10): 5057–5068, 2010.

Yu, Y., Wang, T., and Samworth, R. J. A useful variant of the Davis–Kahan theorem for statisticians. *Biometrika*, 102(2):315–323, 2015.

Ziehe, A., Laskov, P., Nolte, G., and Müller, K.-R. A fast algorithm for joint diagonalization with non-orthogonal transformations and its application to blind source separation. *Journal of Machine Learning Research*, 5:777–800, 2004.

Notation

$\mathbf{p}_j \in \mathbb{R}^K$, $\mathbf{p}_j(k) \in [1/B, B]$	signature of source j (variances across environments), $\frac{1}{K}\sum_k \mathbf{p}_j(k) = 1$
S_{ij} (1); $S_{\min}$, $S_{\mathrm{nn}}(j)$	Fisher information of the (i,j) rotation; its minimum over pairs, over i for fixed j
$E_j = \sum_{i\neq j} S_{ij}^{-1}$; $\bar{E}$, $E_{\max}$	Riesz 2-energy of $\mathbf{p}_j$; its average and maximum over sources
$\delta_a = \mathbf{p}_a - \mathbf{p}_j$; $r_j = \min_{i\neq j}\|\delta_i\|$	signature differences from source j; distance to the nearest signature
δ_j, $\hat{\delta}_j$ (Section 6, Appendix J only)	split-half discrepancy of fMRI component j and its prediction
H_j, λ_a, λ_*	quadratic pencil at $\mathbf{p}_j$, its eigenvalues, its gap (Section 3.2, Appendix C)
v_{ij}^{qp}; E_j^{qp}, $\bar{E}^{\mathrm{qp}}$; $\hat{E}_j^{\mathrm{qp}}$	first-order coefficients of the estimator of Theorem 2; with a perturbed probe (Lemma 3)
v_{ij}^{wh}; E_j^{wh}	first-order coefficient of the estimated whitening in Theorem C.1(b)
γ_j, γ	$K(d+6\log n)/(nS_{\mathrm{nn}}(j))$ and $K(d+6\log n)/(nS_{\min})$
s_0	known lower bound on $S_{\min}$ (Theorem 3)
v_{ij}^{F}, v_{ij}^{w}; $\bar{E}^{\mathrm{F}}$, $\bar{E}^{\mathrm{w}}$	first-order coefficients of joint diagonalization, known/estimated whitening (23)–(24)
v_{ij}^{PC}; $\bar{E}^{\mathrm{PC}}$	Cramér–Rao coefficient of the non-orthogonal problem (Appendix G)
ρ_j, ℓ_{A}^2, ℓ_{M}^2	squared sine of the angle for source j; average and maximal loss
$f_K(d)$	d^2, $d\log d$, d for $K = 2, 3, \geq 4$

Reader's guide

Each result of the main text with its full statement, its conditions and its proof. "Gaussian": Gaussian sources; "4th mom.": independent sources with finite fourth moments; "equal" and "any": the environment sizes n_k ("any" with the weighted divergence of Appendix A); $\gamma_j = K(d+6\log n)/(nS_{\mathrm{nn}}(j))$, $\gamma = \max_j \gamma_j$ and $\gamma_0 = K(d+6\log n)/(ns_0)$.

Result (main text)	Full statement; proof	Sources; whitening; sizes	Signatures	Regime	Not covered
Theorem 1, lower bound	Theorem B.1(a); App. B	Gaussian; known; any	may be used; estimators in $\mathrm{O}(d)$ or labelled; d directions under the matched loss, average bound only (Lemma 5)	every n, d	the per-source and maximal losses ρ_j, ℓ_{M} for non-orthogonal estimators
Theorem 1, attainment at fixed d	Theorem B.1(b); App. B	Gaussian; known; any	$S_{ij} > 0$	$n \to \infty$, fixed d	the d-dependence
Theorem 2(a), pencil, known whitening	Theorem C.1(a); App. C	Gaussian; known; equal	known	$\gamma_j \leq c$	unequal sizes (App. A)
Theorem 2(b), estimated whitening	Theorem C.1(b); App. C, Lemma 2	Gaussian; estimated on a third; equal	known	$\gamma_j \leq c$	unequal sizes
Lemmas 3, 4, perturbed probes	App. C, D	Gaussian; known / estimated; equal	a probe within η/B^4 of the nearest-neighbour distance ($\eta \leq 1/8$), independent of the halves	$\gamma_j \leq c$	
Theorem 3(a), grid, known whitening	Theorem D.1(a); App. D	Gaussian; known; equal	only B and $s_0 \leq S_{\min}$	$\gamma_0 \leq c$	adaptivity in s_0; the grid is exponential in K
Theorem 3(b), estimated whitening	Theorem D.1(b); App. D, Lemma 4	Gaussian; estimated on a fourth; equal	only B and $s_0 \leq S_{\min}$	$\gamma_0 \leq c$	
Corollary 1, minimax rate	App. C; (iii): App. D, Lemma 5	Gaussian; known or estimated; equal	known; $s_0 = S_{\min}$ for Theorem 3	$\gamma \leq c$	
Theorem 4, dimension law	App. E	geometry only	well-spread (Definition 1); lower bounds for $\bar{E}$: every configuration	fixed K	constants depend on K
Corollary 2, necessary size	App. E	Gaussian; known; any	every configuration	every n, d	
Proposition 1, random signatures	App. E	geometry only	i.i.d., bounded density	w.h.p.	
Theorem 5, linear pencil	App. F	4th mom.; known; any	not used; α independent of the data	$n \to \infty$, fixed d	adaptive or iterated α; the pencil as an initialisation (App. J)
Proposition 2, joint diagonalization	App. G	4th mom.; known or estimated; any	not used; global minimiser	$n \to \infty$, fixed d	Jacobi's output (Conjecture 1)
Proposition 3, barrier	App. H	Gaussian; known; any	–	every n, d	

Result (main text)	Full statement; proof	Sources; sizes	whitening;	Signatures	Regime	Not covered
Conjecture 1, linear regime	Section 6, App. J	simulation; equal	estimated;	well-spread, $K \geq 4$ (not used by Jacobi)	$n \geq 200d$, a proof $d \leq 1000$	

A. Fisher information of a rotation (proof of Lemma 1), unequal sizes and non-Gaussian sources

Lemma 1 (Fisher information of a rotation)**.** *For Gaussian sources and $R \in \mathrm{O}(d)$, with $Q' = QR$ and P_Q the law of the full sample,*

$$\mathrm{KL}\left(P_Q \,\|\, P_{Q'}\right) = \frac{n}{2K} \sum_{k=1}^{K} \sum_{a \neq b} R_{ab}^2 \Big(\frac{\mathbf{p}_a(k)}{\mathbf{p}_b(k)} - 1 \Big). \tag{6}$$

If R is a rotation by θ in the coordinate plane (i, j), the divergence equals $\frac{n}{2} \sin^2\theta \, S_{ij}$ exactly. Hence S_{ij} is the Fisher information per sample, averaged over environments, for the rotation angle in plane (i, j). The information is diagonal across planes and orthogonal to the variance parameters.

For centred Gaussians, $\mathrm{KL}(N(0,\Sigma)\|N(0,\Sigma')) = \frac{1}{2}[\mathrm{tr}(\Sigma'^{-1}\Sigma) - d - \log\det(\Sigma'^{-1}\Sigma)]$. With $\Sigma_k = Q\Lambda_k Q^\top$ and $\Sigma'_k = QR\Lambda_k R^\top Q^\top$ the determinants agree and $\mathrm{tr}(\Sigma_k'^{-1}\Sigma_k) = \mathrm{tr}(R\Lambda_k^{-1}R^\top\Lambda_k) = \sum_{a,b} R_{ab}^2 \, \mathbf{p}_a(k)/\mathbf{p}_b(k)$. Since $\sum_b R_{ab}^2 = 1$, subtracting $d = \sum_{a,b} R_{ab}^2$ gives $\sum_{a\neq b} R_{ab}^2(\mathbf{p}_a(k)/\mathbf{p}_b(k) - 1)$, and summing the $n_k = n/K$ samples of each environment gives (6). With $n_k = \pi_k n$ the same computation gives (6) with $\frac{1}{K}\sum_k$ replaced by $\sum_k \pi_k$, hence $S_{ij} = \sum_k \pi_k(\mathbf{p}_i(k)/\mathbf{p}_j(k) + \mathbf{p}_j(k)/\mathbf{p}_i(k) - 2)$. For a planar rotation in (i, j), $R_{ij}^2 = R_{ji}^2 = \sin^2\theta$ and all other off-diagonal entries vanish, so the sum equals $\sin^2\theta\,(\mathbf{p}_i(k)/\mathbf{p}_j(k) + \mathbf{p}_j(k)/\mathbf{p}_i(k) - 2)$, whose average over k is S_{ij}; the Fisher information is the coefficient of $\theta^2/2$ in KL per sample, i.e. S_{ij}. For the orthogonality statements, the score for the rotation angle in plane (i, j) at $\theta = 0$ is $\sum_t s_{i,t}s_{j,t}(\mathbf{p}_j(k_t)^{-1} - \mathbf{p}_i(k_t)^{-1})$, the score for the variance $\mathbf{p}_l(k)$ is a function of $s_{l,t}^2$ alone, and the cross-expectations $\mathbb{E}[s_i s_j s_{i'} s_{j'}]$, $\mathbb{E}[s_i s_j (s_l^2 - \mathbf{p}_l)]$ vanish by independence unless $\{i, j\} = \{i', j'\}$. □

Unequal sizes. With $n_k = \pi_k n$, the proofs of Theorems 1 and 4, Corollary 2 and Propositions 3 and 1 use only Lemma 1 and (3) with $\|\cdot\|^2$ replaced by $\sum_k \pi_k(\cdot)_k^2$ and K by 1, so they hold with the weighted divergence. The comparison of the weighted norm with $\|\cdot\|^2/K$ makes the constants of the upper bounds in Theorem 4 and Proposition 1 depend also on $\min_k K\pi_k$ (an environment with $\pi_k \to 0$ collapses one dimension); their lower bounds and Corollary 2 hold uniformly in the sizes, since they need only $S_{ij} \leq B^2 \sum_k \pi_k(\mathbf{p}_i(k) - \mathbf{p}_j(k))^2 \leq B^2\|\mathbf{p}_i - \mathbf{p}_j\|^2$. The first-order results hold as well: Theorem B.1(b) through the weighted Fisher information, and Theorem 5 and Proposition 2, whose proofs use only $\mathrm{Var}(N_{k,ab}) = \mathbf{p}_a(k)\mathbf{p}_b(k)/n_k$, for the π-weighted pencils and criteria; Appendix G gives the coefficients used in Figure 6(c). The final bound $(d-1)^3/(dB^4)$ of Theorem 5 holds for every π, since then $v_{ij}(\alpha) = \sum_k \alpha_k^2 \mathbf{p}_i(k)\mathbf{p}_j(k)/\pi_k \geq B^{-2}\|\alpha\|_1^2$ and $L_\alpha \leq B\|\alpha\|_1$.

The non-asymptotic upper bounds (Theorems 2 and 3, Lemma 3, Corollary 1) are proved for equal sizes only, and the π-weighted estimators fail for arbitrary sizes. Take $K_0 \asymp \log d$ environments with m samples each and separated code signatures (Section 4), and add one environment k_0 with $\mathbf{p}_j(k_0) = 1$ for every j and two samples, so that $n = K_0 m + 2$ and $\pi_{k_0} = 2/n$. In the split pencil of Theorem 2 with weights $\pi_k w_k$, whose halves of k_0 have one sample each, this environment contributes $\frac{2}{n}\left[\frac{1}{2}x^\top y\,(xy^\top + yx^\top) - xx^\top - yy^\top + I\right]$ with x, y independent $N(0, I_d)$, whose smallest eigenvalue is about $-|x^\top y|\,\|x\|\|y\|/n \asymp -d^{3/2}/n$ with probability bounded below, while the population pencil has its gap of order one and receives nothing from k_0. With $n = CK_0 d$ the condition $\gamma_j \leq c$ of Theorem 2, with the weighted divergence, holds for large C, but this eigenvalue has size of order $\sqrt{d}/(C\log d) \to \infty$; so, with probability bounded below, the bottom eigenvector lies near $\mathrm{span}\{x, y\}$, which is isotropic, so that $\langle \hat{q}_j, q_j\rangle^2 = O_P(1/d)$ and $\rho_j \to 1$, whereas $E_j/n \asymp 1/\log d$. With six samples in k_0 the same happens to the estimator of Theorem 3. The proof of Theorem 2 needs the concentration bound of Step 2 in every environment, which fails here; how far a lower bound on $\min_k n_k$ suffices we leave open. Whether the minimax rate of Corollary 1 survives arbitrary sizes is also open; in this example it does, since discarding k_0 leaves an equal-size problem whose divergences differ from the weighted ones by the factor $n/(n-2)$. In practice, environments with $n_k \lesssim d + \log n$ can be dropped or merged before the pencil is formed. We expect Theorem 2 to hold whenever $\min_k K\pi_k \geq c > 0$, with constants depending on c, and for independent sub-Gaussian sources: Step 2 of Appendix C has a sub-Gaussian version (Vershynin, 2012, Thm. 5.39) and Step 5 and Lemma 2 use only independence and fourth moments. We have not checked the constants and claim neither.

Non-Gaussian sources. This paragraph concerns the parametric model in which the source densities are known; for the semiparametric problem with unknown densities, in the i.i.d. model, see Amari & Cardoso (1997); Chen & Bickel (2006). Let source j in environment k have density $\mathbf{p}_j(k)^{-1/2} g_j(s/\sqrt{\mathbf{p}_j(k)})$, where g_j has mean 0 and variance 1 and is absolutely continuous with score $\psi_j = -(\log g_j)'$, finite location information $\kappa_j = \mathbb{E}\psi_j(u)^2$ and finite scale information $\mathbb{E}[u^2\psi_j(u)^2]$, and $s\, g_j(s) \to 0$ as $|s| \to \infty$. Integration by parts gives $\mathbb{E}[\psi_j(u)u] = 1$, so $\kappa_j \geq 1$ by Cauchy–Schwarz, with equality iff g_j is Gaussian. The score for the rotation angle in plane (i, j) at $\theta = 0$ is $\psi_i(s_i)s_j - \psi_j(s_j)s_i$ (scores at the actual scales); since $\mathbb{E}\psi_i(s_i)^2 = \kappa_i/\mathbf{p}_i(k)$, its variance in environment k is $\kappa_i \mathbf{p}_j(k)/\mathbf{p}_i(k) + \kappa_j \mathbf{p}_i(k)/\mathbf{p}_j(k) - 2$, and it is uncorrelated with the scores of the variances and of the other planes, as in Lemma 1. The per-sample information is therefore

$$S_{ij}^{\kappa} = \frac{1}{K} \sum_k \Big[\kappa_i \frac{\mathbf{p}_j(k)}{\mathbf{p}_i(k)} + \kappa_j \frac{\mathbf{p}_i(k)}{\mathbf{p}_j(k)} \Big] - 2 \;\geq\; S_{ij},$$

strictly unless both sources are Gaussian, and $S_{ij}^{\kappa} = \kappa S_{ij} + 2(\kappa - 1)$ when $\kappa_i = \kappa_j = \kappa$: within this parametric model, Gaussian sources are the least informative. The term $2(\kappa - 1)$ caps the contribution of each pair to the energy $E_j^{\kappa} = \sum_{i \neq j} (S_{ij}^{\kappa})^{-1}$ at $1/(2(\kappa - 1))$, so $E_j^{\kappa} \leq (d - 1)/(2(\kappa - 1))$ whatever the geometry, while for weakly non-Gaussian sources ($\kappa - 1 \ll S_{\mathrm{nn}(j)}$) the geometry of Theorem 4 takes over. This is a statement about Fisher information; minimax lower bounds for non-Gaussian classes, and the low-degree barrier of the non-Gaussian route (Auddy & Yuan, 2025), are separate questions.

Optimality among second-order estimators. Theorem B.1(b) is the efficiency of maximum likelihood for Gaussian sources. For other source laws the Riesz rate remains the first-order limit of every estimator that uses the data through the sample covariances in a regular way. Let the sources be independent with finite fourth moments and let $\hat{Q} = T(\hat{\Sigma}_1, \ldots, \hat{\Sigma}_K)$, where T is differentiable at the population covariances and Fisher consistent, $T(\Sigma_1(\theta'), \ldots, \Sigma_K(\theta')) = Q(\theta')$ for every parameter $\theta' = (Q', \{\Lambda_k'\})$ in a neighbourhood; every joint diagonalizer of Proposition 2 is of this form. In the basis of Q the entries $N_{k,ab} = (\hat{\Sigma}_k - \Sigma_k)_{ab}$, $a \leq b$, are uncorrelated across pairs and environments, with variance $\mathbf{p}_a(k)\mathbf{p}_b(k)/n_k$ for $a < b$, since $\mathbb{E}[s_a s_b s_{a'} s_{b'}] = 0$ unless $\{a, b\} = \{a', b'\}$ and $\mathbb{E}[s_a s_b (s_l^2 - \mathbf{p}_l)] = 0$ for $a \neq b$; only the variances of the diagonal entries involve the fourth moments. A rotation by Ξ moves the mean of $N_{k,ab}$, $a < b$, by $\Xi_{ab}(\mathbf{p}_a(k) - \mathbf{p}_b(k)) + o(\|\Xi\|)$ and the means of the diagonal entries only at second order, while a change of $\mathbf{p}_l(k)$ moves only the mean of the diagonal entry (k, ll). By the delta method and the chain rule applied to the identity $T(\Sigma(\theta')) = Q(\theta')$, the rotation coordinates of $\hat{Q}$ have asymptotic covariance $LVL^{\top}$ for a matrix L with $L\dot{\Sigma} = [\, I \;\; 0\,]$, where V is the asymptotic covariance of the sample covariances and $\dot{\Sigma}$ the derivative of their means in θ. By the two facts just stated, $\dot{\Sigma}$ is block diagonal: the rotation coordinates enter only the off-diagonal means at first order, the variance coordinates only the diagonal ones. So the restriction L_{off} of L to the off-diagonal entries satisfies $L_{\mathrm{off}}\dot{\Sigma}_{\mathrm{off}} = I$, and since diagonal and off-diagonal entries are uncorrelated, $LVL^{\top} \succeq L_{\mathrm{off}} V_{\mathrm{off}} L_{\mathrm{off}}^{\top} \succeq (\dot{\Sigma}_{\mathrm{off}}^{\top} V_{\mathrm{off}}^{-1} \dot{\Sigma}_{\mathrm{off}})^{-1}$ by the Gauss–Markov inequality, where V_{off} is diagonal with the entries $\mathbf{p}_a(k)\mathbf{p}_b(k)/n_k$ above and $\dot{\Sigma}_{\mathrm{off}}^{\top} V_{\mathrm{off}}^{-1} \dot{\Sigma}_{\mathrm{off}}$ is diagonal with entries $\sum_k n_k(\mathbf{p}_a(k) - \mathbf{p}_b(k))^2/(\mathbf{p}_a(k)\mathbf{p}_b(k))$, which is nS_{ab} for $n_k = n/K$. Hence every rotation angle has asymptotic variance at least $1/(nS_{ab})$, whatever the fourth moments (only the off-diagonal block of V enters, and it is nonsingular), and the limit of $n\rho_j$ has mean at least E_j, which Proposition 2(ii) attains. Regularity cannot be dropped. For symmetric two-valued sources, $s_a = \pm\sqrt{\mathbf{p}_a(k)}$, the diagonal of $Q^{\top}\hat{\Sigma}_k Q$ equals Λ_k exactly, V is singular on the diagonal entries, and a non-regular estimator can exploit them: for $d = 2$ and known signatures (where the bound above holds for regular T as well), the common solutions θ' of $\mathrm{diag}(R_{\theta'}^{\top}\hat{\Sigma}_k R_{\theta'}) = \Lambda_k$, $k = 1, 2$, with $R_{\theta'}$ the rotation by θ', are the true angle and, only on an event of probability $O(n^{-1/2})$ (the off-diagonal sample moments $\nu_k = n_k^{-1}\sum_t s_{1t}s_{2t}$ are lattice variables), one further angle at distance $O(|\nu_k|)$ from it; an estimator that returns a common solution therefore has $\mathbb{E}[n\rho_j] \to 0$, below the regular bound $1/S_{12}$. Estimators that use higher-order statistics are not covered; for them the information S_{ij}^{κ} above applies.

The loss. ℓ_{A}^2 and ℓ_{M}^2 coincide with the squared average and maximal losses of Auddy & Yuan (2025) whenever the matching is a permutation, and are never larger, so the lower bounds of Theorem 1 transfer to their losses. Like the Amari index, ℓ_{A}^2 is permutation invariant, but it decomposes by source.

B. Proof of Theorem 1

Theorem B.1 (Theorem 1, full statement)**.** *For Gaussian sources and $K \geq 2$:*

(a) (Lower bound.) There is a constant $c > 0$ depending only on B such that, for every n and d, every source j and every

estimator $\hat{Q} \in \mathrm{O}(d)$ (which may use the signatures),

$$\sup_{Q\in\mathrm{O}(d)} \mathbb{E}\,\rho_j(\hat{Q}) \;\geq\; c\,\frac{E_j}{n+E_j} \quad \textit{(this is (4))},$$

$$\sup_{Q\in\mathrm{O}(d)} \mathbb{E}\,\ell_{\mathrm{A}}^2(\hat{Q}) \;\geq\; \frac{c}{d}\sum_{j=1}^{d}\frac{E_j}{n+E_j}, \tag{7}$$

$$\sup_{Q\in\mathrm{O}(d)} \mathbb{E}\,\ell_{\mathrm{M}}^2(\hat{Q}) \;\geq\; c\,\max_j \frac{E_j}{n+E_j}. \tag{8}$$

If $S_{ij} = 0$ for some $i \neq j$, then $E_j = \infty$ and $E_j/(n+E_j)$ is read as 1.

(b) (Attainment at fixed d.) Let $S_{ij} > 0$ for all $i \neq j$. For fixed d and $n \to \infty$, the Gaussian maximum-likelihood joint diagonalizer $\hat{Q}_{\mathrm{ML}}$ satisfies $n\,\rho_j(\hat{Q}_{\mathrm{ML}}) \Rightarrow \sum_{i\neq j} S_{ij}^{-1} Z_i^2$ with Z_i i.i.d. standard normal, a limit law with mean E_j.

Part (b) is the classical efficiency of maximum likelihood (Flury, 1986; Pham & Cardoso, 2001), proved at the end of this appendix with the information of Lemma 1. All three bounds of (a) hold for labelled estimators (Appendix C, proof of Corollary 1).

Fix j and a reference $Q_0 \in \mathrm{O}(d)$. Write $C_{\mathrm{tri}} = 2B^4$. By (3), $S_{ab} \leq \frac{B^2}{K}\|\mathbf{p}_a - \mathbf{p}_b\|^2 \leq \frac{2B^2}{K}\big(\|\mathbf{p}_a - \mathbf{p}_j\|^2 + \|\mathbf{p}_j - \mathbf{p}_b\|^2\big) \leq C_{\mathrm{tri}}(S_{ja} + S_{jb})$ for all a, b, with $S_{jj} = 0$.

The hypercube. Order the other sources as $i_1, \ldots, i_m$, $m = d-1$, so that $t_l = S_{ji_l}$ is non-increasing: $t_1 \geq t_2 \geq \cdots \geq t_m > 0$. Let $E = E_j = \sum_l t_l^{-1}$, $\eta = 1/(64B^4)$ and

$$\theta_l^2 = \frac{\eta}{(n+E)\,t_l}, \qquad \Theta = \sum_l \theta_l^2 = \frac{\eta E}{n+E} \leq \eta \leq \frac{1}{64}.$$

For $\varepsilon \in \{\pm 1\}^m$ put

$$Q_\varepsilon = Q_0\, R_1(\varepsilon_1\theta_1) R_2(\varepsilon_2\theta_2) \cdots R_m(\varepsilon_m\theta_m),$$

where $R_l(\theta)$ is the rotation by θ in the coordinate plane (j, i_l), with $R_l(\theta)e_j = \cos\theta\, e_j + \sin\theta\, e_{i_l}$. Hypercubes of planar rotations are standard in minimax bounds for principal subspaces (Vu & Lei, 2013; Cai et al., 2013); the ordering by S_{ji}, which keeps the divergence between neighbours bounded without truncation, is the new ingredient.

Geometry of column j. Since R_l fixes every $e_{i_{l'}}$ with $l' \neq l$, applying the rotations from right to left gives

$$q_j(\varepsilon) = Q_0\Big[\Big(\prod_l \cos\theta_l\Big)e_j + \sum_l \varepsilon_l \sin\theta_l \Big(\prod_{l'>l}\cos\theta_{l'}\Big) e_{i_l}\Big].$$

The coordinate of $q_j(\varepsilon)$ along $Q_0 e_{i_l}$ is $\varepsilon_l \tau_l$ with $\tau_l = \sin\theta_l \prod_{l'>l}\cos\theta_{l'} \geq \theta_l/2$, because $\Theta \leq 1/64$ gives $\sin\theta_l \geq 0.99\,\theta_l$ and $\prod_{l'}\cos\theta_{l'} \geq 1 - \Theta/2$. Also $\sin^2\angle(q_j(\varepsilon), Q_0 e_j) = 1 - \prod_l \cos^2\theta_l \leq \Theta$.

Reduction to a Hamming loss. The reduction must use a function of the data alone. Let $\hat{q}$ be the column of $\hat{Q} \in \mathrm{O}(d)$ with the largest $|\langle \hat{q}_a, Q_0 e_j\rangle|$, signed so that $\langle \hat{q}, Q_0 e_j\rangle \geq 0$, and put $\hat{\varepsilon}_l = \operatorname{sign}\langle \hat{q}, Q_0 e_{i_l}\rangle$ (with $\operatorname{sign} 0 = 1$). We claim that for every ε and every $\hat{Q} \in \mathrm{O}(d)$,

$$\sum_l \tau_l^2\, \mathbf{1}[\hat{\varepsilon}_l \neq \varepsilon_l] \leq 2\rho_j(\hat{Q}). \tag{9}$$

If $\rho_j(\hat{Q}) \geq 1/8$ the left side is at most $\sum_l \tau_l^2 \leq \Theta \leq 1/64 \leq 2\rho_j$. If $\rho_j(\hat{Q}) < 1/8$, let q^* be the column closest to $q_j(\varepsilon)$, at angle $\alpha < \arcsin(1/\sqrt{8}) < 0.362$. Since $\angle(q_j(\varepsilon), Q_0 e_j) \leq \arcsin\sqrt{\Theta} < 0.13$, q^* is within angle $0.5 < \pi/4$ of $Q_0 e_j$, while every other column, being orthogonal to q^*, is at angle more than $\pi/4$ from $Q_0 e_j$. Hence $\hat{q} = \pm q^*$, and with the sign convention $\langle \hat{q}, q_j(\varepsilon)\rangle = \cos\alpha > 0$. If $\hat{\varepsilon}_l \neq \varepsilon_l$ then $|\langle \hat{q} - q_j(\varepsilon), Q_0 e_{i_l}\rangle| \geq \tau_l$, so

$$\sum_l \tau_l^2\, \mathbf{1}[\hat{\varepsilon}_l \neq \varepsilon_l] \leq \|\hat{q} - q_j(\varepsilon)\|^2 = 2(1-\cos\alpha) \leq 2\sin^2\alpha = 2\rho_j(\hat{Q}).$$

Neighbouring hypotheses. Let $\varepsilon, \varepsilon'$ differ only in coordinate l. Then

$$R := Q_\varepsilon^\top Q_{\varepsilon'} = T^\top R_l(2\varepsilon_l'\theta_l)\, T, \qquad T = R_{l+1}(\varepsilon_{l+1}\theta_{l+1}) \cdots R_m(\varepsilon_m\theta_m),$$

an exact rotation by $\delta = 2\varepsilon_l'\theta_l$ in the plane spanned by the orthonormal vectors $u = T^\top e_j$ and $w = e_{i_l}$. The vector u is supported on $J_l = \{j, i_{l+1}, \dots, i_m\}$, and $u_{i_l} = 0$. Write $R = I + (\cos\delta - 1)(uu^\top + ww^\top) + \sin\delta\,(wu^\top - uw^\top)$. Two facts follow.

(1) *R only moves sources at least as close to $\mathbf{p}_j$ as i_l.* $R_{ab} \neq 0$ for $a \neq b$ only if $a, b \in J_l \cup \{i_l\}$. By the ordering, every $a \in J_l \cup \{i_l\}$ other than j has $S_{ja} \leq t_l$, so by the quasi-triangle inequality $S_{ab} \leq 2C_{\mathrm{tri}} t_l$ for all a, b in this set.

(2) *$R_{ab}^2 = R_{ba}^2$ for $a \neq b$.* The antisymmetric part $\sin\delta\,(w_a u_b - u_a w_b)$ of R_{ab} vanishes unless a or b equals i_l; in that case the symmetric part $(\cos\delta - 1)(u_a u_b + w_a w_b)$ vanishes, because $u_{i_l} = 0$ and w is zero off i_l.

By (2), the two terms (a, b) and (b, a) of (6) combine into $R_{ab}^2\big(\mathbf{p}_a(k)/\mathbf{p}_b(k) + \mathbf{p}_b(k)/\mathbf{p}_a(k) - 2\big)$, whose average over k is $R_{ab}^2 S_{ab}$. Hence, by Lemma 1, (1), and $\sum_{a\neq b} R_{ab}^2 \leq \|R - I\|_F^2 = 4(1 - \cos\delta) = 8\sin^2\theta_l$,

$$\mathrm{KL}\left(P_{Q_\varepsilon} \,\|\, P_{Q_{\varepsilon'}}\right) = \frac{n}{2}\sum_{a<b} R_{ab}^2 S_{ab} \leq \frac{n}{2} \cdot 2C_{\mathrm{tri}} t_l \cdot 4\sin^2\theta_l \leq 4C_{\mathrm{tri}}\, n\, t_l\, \theta_l^2 = 8B^4\eta\,\frac{n}{n+E} \leq \frac{1}{8}.$$

By Pinsker's inequality the total variation between neighbouring hypotheses is at most $1/4$.

Assouad's lemma. Average (9) over ε uniform on the cube. For each l, pairing ε with its neighbour in coordinate l gives $\mathbb{P}_\varepsilon(\hat{\varepsilon}_l \neq \varepsilon_l) + \mathbb{P}_{\varepsilon'}(\hat{\varepsilon}_l \neq \varepsilon_l') \geq 1 - \mathrm{TV} \geq 3/4$. Hence (Tsybakov, 2009, weighted form of Thm. 2.12)

$$\sup_Q \mathbb{E}\rho_j(\hat{Q}) \;\geq\; \frac{1}{2}\sum_l \tau_l^2 \cdot \frac{3}{8} \;\geq\; \frac{3}{64}\sum_l \theta_l^2 \;=\; \frac{3\eta}{64} \cdot \frac{E_j}{n + E_j},$$

which is (4) with $c = 3\eta/64 = 3/(4096B^4)$. The argument bounds a Bayes risk under the uniform prior on the cube; since the Bayes risk of a randomised estimator is an average of Bayes risks of deterministic ones, (4) holds for randomised estimators as well. If some $S_{ji} = 0$, then $\mathbf{p}_i = \mathbf{p}_j$ and rotations in the plane (j, i) do not change the law of the data. The truths Q_0 and $Q_0 R(\pi/4)$ have j-th columns at angle $\pi/4$, and every $\hat{Q} \in \mathrm{O}(d)$ has $\rho_j^{(1)} + \rho_j^{(2)} \geq 2\sin^2(\pi/8)$: whether one column serves both truths or two orthogonal columns do, the two angles sum to at least $\pi/4$. If both are at most $\pi/4$, convexity of $\sin^2$ on $[0, \pi/4]$ gives the claim; otherwise one term alone exceeds $\frac{1}{2} > 2\sin^2(\pi/8)$. Hence $\sup_Q \mathbb{E}\rho_j \geq \sin^2(\pi/8) > 1/8$, which is (4) with $E_j = \infty$.

Average and maximal losses (7)–(8). The bound for ℓ_{M}^2 follows from $\ell_{\mathrm{M}}^2 \geq \rho_j$ for every j. For ℓ_{A}^2 we symmetrise. Fix an estimator $\hat{Q}$, write X for the full sample (a $d \times n$ matrix whose columns are the observations of all environments), and let H be a Haar-distributed element of $\mathrm{O}(d)$, independent of X. Define the randomised estimator

$$\tilde{Q}(X) = H^\top \hat{Q}(HX).$$

Under the truth Q, with known signatures, HX has the law of the sample under the truth HQ. The loss is invariant under a common rotation: $\langle H^\top \hat{q}_a(HX), q_j\rangle = \langle \hat{q}_a(HX), Hq_j\rangle$ and Hq_j is the j-th column of HQ, so $\rho_j(\tilde{Q}(X); Q) = \rho_j(\hat{Q}(HX); HQ)$. Taking expectations, with $r_j(G) = \mathbb{E}_{Y\sim P_G}\,\rho_j(\hat{Q}(Y); G)$ the risk function of $\hat{Q}$,

$$\mathbb{E}_Q\,\rho_j(\tilde{Q}) = \mathbb{E}_H\, r_j(HQ) = \int_{\mathrm{O}(d)} r_j\, d\mu =: \bar{r}_j,$$

because HQ is Haar-distributed for every fixed Q. The right side does not depend on Q. Applying (4) to $\tilde{Q}$ gives $\bar{r}_j \geq cE_j/(n + E_j)$ for every j. Finally, for the original estimator,

$$\sup_Q \mathbb{E}_Q\,\ell_{\mathrm{A}}^2(\hat{Q}) \;\geq\; \int_{\mathrm{O}(d)} \mathbb{E}_G\,\ell_{\mathrm{A}}^2(\hat{Q})\, d\mu(G) \;=\; \frac{1}{d}\sum_j \bar{r}_j \;\geq\; \frac{c}{d}\sum_j \frac{E_j}{n + E_j}. \qquad \square$$

Attainment (b). The block-Gaussian model is a regular parametric family in $(Q, \{\Lambda_k\})$; by Lemma 1 the Fisher information for the rotation angle in plane (j, i) is nS_{ij}, different planes are orthogonal at first order, and the variances are orthogonal nuisance parameters. With $S_{ij} > 0$ for all pairs the model is identifiable up to signs and permutations, and the maximum-likelihood estimator is asymptotically normal with inverse-information covariance (Flury, 1986; Pham & Cardoso, 2001; Yeredor, 2010). In the local coordinates $\hat{Q}_{\mathrm{ML}} = Q\exp(\Xi)$ the angles $\sqrt{n}\,\Xi_{ji}$ converge jointly to independent $N(0, S_{ij}^{-1})$, and $\rho_j = \sum_{i\neq j}\Xi_{ji}^2 + O(\|\Xi\|^3)$ gives the limit law, whose mean is E_j. $\square$

C. Proof of Theorem 2, Lemma 3 and Corollary 1

Theorem C.1 (Theorem 2, full statement)**.** *Let the sources be Gaussian and the signatures known, and put $w_k = \mathbf{p}_j(k)^{-2}$ and $\gamma_j = K(d + 6\log n)/(nS_{\mathrm{nn}}(j))$. There are $c, C > 0$ depending only on B such that, if $\gamma_j \le c$, the following hold.*

(a) (Known whitening.) Let $A = Q \in \mathrm{O}(d)$ and $2K \mid n$, split each environment's sample into halves with covariances $\hat\Sigma_k^{(1)}, \hat\Sigma_k^{(2)}$, and let $\hat q_j$ be a unit eigenvector for the smallest eigenvalue of

$$\hat H_j = \sum_k w_k \,\mathrm{sym}\big[(\hat\Sigma_k^{(1)} - \mathbf{p}_j(k)I)(\hat\Sigma_k^{(2)} - \mathbf{p}_j(k)I)\big].$$

Then

$$\mathbb{E}\big[1 - \langle \hat q_j, q_j\rangle^2\big] \;\le\; \frac{E_j^{\mathrm{qp}}}{n}\big(1 + C\sqrt{\gamma_j}\big)^2 + \frac{1}{n^2}, \tag{10}$$

where the first-order coefficient E_j^{qp} (defined below) satisfies $E_j \le E_j^{\mathrm{qp}} \le \kappa_B E_j$, $\kappa_B = \frac14(B^2 + B^{-2})^2$.

(b) (Estimated whitening.) Let A be any invertible matrix and $3K \mid n$. Split each environment's sample into three equal parts, whiten parts 1 and 2 by $\hat C^{-1/2}$, $\hat C$ the pooled sample covariance of part 0, compute $\hat q_j$ as in (a) from them, and let $\hat b_j = \hat C^{-1/2}\hat q_j$. The estimated source is $\hat b_j^\top x = z^\top s$ with $z = A^\top \hat b_j$, and $\rho_j = 1 - z_j^2/\|z\|^2$ satisfies, for every A,

$$\mathbb{E}\rho_j \;\le\; \frac{3}{2n}\big(E_j^{\mathrm{qp}} + 2E_j^{\mathrm{wh}}\big)\big(1 + C\sqrt{\gamma_j}\big)^2 + \frac{1}{n^2}, \tag{11}$$

where the coefficient E_j^{wh} of the estimated whitening (defined in the proof of part (b)) satisfies $E_j^{\mathrm{wh}} \le B^4 E_j$.

The simplified statement of Section 3.2 follows: under $\gamma_j \le c$, $(1+C\sqrt{\gamma_j})^2 \le (1+C\sqrt{c})^2$, $E_j^{\mathrm{qp}} + 2E_j^{\mathrm{wh}} \le (\kappa_B + 2B^4)E_j$, and $n^{-2} \le (B^2 + B^{-2} - 2)E_j/n$ because $E_j \ge (d-1)/(B^2 + B^{-2} - 2)$ and $d \ge 2$. In (b), $\hat b_j^\top x = z^\top s$ has pooled variance $\|z\|^2$ (the sources have pooled variance one), of which $\|z\|^2 - z_j^2$ comes from the sources other than j, so ρ_j is the share of the pooled variance of the estimated source that comes from the other sources.

Throughout, j is fixed, C denotes a constant depending only on B that may change from line to line, and we work in the basis of Q (the estimator is equivariant: $\hat H_j$ computed from Qx is $Q\hat H_j Q^\top$). The sample covariances are uncentred, $\hat\Sigma_k^{(i)} = \frac{1}{m}\sum xx^\top$ over the $m = n/(2K)$ samples of half i (we assume $2K \mid n$), $N_k^{(i)} = \hat\Sigma_k^{(i)} - \Lambda_k$, $\bar N_k = \frac12(N_k^{(1)} + N_k^{(2)})$ (the noise of the full sample of environment k), $\tilde\Lambda_k = \Lambda_k - \mathbf{p}_j(k)I$ and $\delta_a = \mathbf{p}_a - \mathbf{p}_j$. Expanding the product,

$$\hat H_j = H + P_1 + P_2, \qquad H = \sum_k w_k\tilde\Lambda_k^2 = \mathrm{diag}(\lambda_a), \quad P_1 = \sum_k w_k(\tilde\Lambda_k\bar N_k + \bar N_k\tilde\Lambda_k), \quad P_2 = \sum_k w_k\,\mathrm{sym}(N_k^{(1)}N_k^{(2)}),$$

with $\lambda_a = \sum_k w_k\delta_a(k)^2$ and $\lambda_j = 0$. With $r_k = \mathbf{p}_a(k)/\mathbf{p}_j(k) \in [B^{-2}, B^2]$ we have $w_k\delta_a(k)^2 = (r_k - 1)^2$ and $KS_{aj} = \sum_k (r_k - 1)^2/r_k$, so $B^{-2}KS_{aj} \le \lambda_a \le B^2 KS_{aj}$ and $\lambda_* := \min_{a\ne j}\lambda_a \ge KS_{\mathrm{nn}}(j)/B^2$. A prime denotes restriction to $e_j^\perp$: $H' = \mathrm{diag}(\lambda_a)_{a\ne j}$, $P' = P_1' + P_2'$, and $h = (\hat H_j e_j)_{-j} = h^{(1)} + h^{(2)}$ with $h_a^{(1)} = \sum_k w_k\delta_a(k)\bar N_{k,aj}$ and $h^{(2)} = (P_2 e_j)_{-j}$. The first-order error vector is $g = H'^{-1}h$, and $\mathbb{E}\|H'^{-1}h^{(1)}\|^2 = \sum_{a\ne j}\sum_k w_k^2\delta_a(k)^2\mathbf{p}_a(k)\mathbf{p}_j(k)/(n_k\lambda_a^2) = E_j^{\mathrm{qp}}/n$ with $E_j^{\mathrm{qp}} = \sum_{i\ne j} v_{ij}^{\mathrm{qp}}$ and, writing $\delta = \mathbf{p}_i - \mathbf{p}_j$, $v_{ij}^{\mathrm{qp}} = K\sum_k w_k^2\delta_k^2\mathbf{p}_i(k)\mathbf{p}_j(k)/(\sum_k w_k\delta_k^2)^2$. The bounds $S_{ij}^{-1} \le v_{ij}^{\mathrm{qp}} \le \kappa_B S_{ij}^{-1}$ follow from Cauchy–Schwarz, $(\sum_k w_k\delta_k^2)^2 \le \sum_k w_k^2\delta_k^2\mathbf{p}_i\mathbf{p}_j \cdot \sum_k \delta_k^2/(\mathbf{p}_i\mathbf{p}_j)$, and from Kantorovich's inequality applied to $v_{ij}^{\mathrm{qp}}S_{ij} = \sum(r-1)^2 r \cdot \sum(r-1)^2 r^{-1}/(\sum(r-1)^2)^2$ with $r_k \in [B^{-2}, B^2]$, which gives $v_{ij}^{\mathrm{qp}}S_{ij} \le \frac14(B^2 + B^{-2})^2 =: \kappa_B$.

Step 1: relative boundedness. Let $\nu = \sum_k w_k\|\bar N_k\|^2$ and $\pi_2 = \|P_2\|$ (operator norms). For $x \perp e_j$,

$$|x^\top P_1 x| = 2\Big|\sum_k w_k(\tilde\Lambda_k x)^\top(\bar N_k x)\Big| \le 2\Big(\sum_k w_k\|\tilde\Lambda_k x\|^2\Big)^{1/2}\Big(\sum_k w_k\|\bar N_k x\|^2\Big)^{1/2} \le 2\sqrt{x^\top Hx}\,\sqrt{\nu}\,\|x\|,$$

because $\sum_k w_k\|\tilde\Lambda_k x\|^2 = \sum_a \lambda_a x_a^2$. With $x = H'^{-1/2}y$ this gives $\|E_1\| \le 2\sqrt{\nu/\lambda_*}$ for $E_1 = H'^{-1/2}P_1'H'^{-1/2}$: the noise is small relative to the nearest-neighbour gap λ_*, not relative to the spread of the λ_a. Similarly, $H'^{-1/2}h^{(1)} = \sum_k M_k y_k$ with $M_k = \sqrt{w_k}H'^{-1/2}\tilde\Lambda_k'$ and $y_k = \sqrt{w_k}(\bar N_k e_j)_{-j}$; since $\sum_k M_k M_k^\top = I$, $\|H'^{-1/2}h^{(1)}\|^2 \le \sum_k w_k\|\bar N_k e_j\|^2 \le \nu$. Finally $\|H'^{-1/2}h^{(2)}\| \le \lambda_*^{-1/2}\pi_2$.

Step 2: the good event. For a $d \times m$ standard Gaussian matrix Z and $s > 0$, all singular values of $Z/\sqrt{m}$ lie in $1 \pm (\sqrt{d}+s)/\sqrt{m}$ with probability at least $1 - 2e^{-s^2/2}$ (Vershynin, 2012, Cor. 5.35), so $\|ZZ^\top/m - I\| \le 3(\sqrt{d}+s)/\sqrt{m}$ when the latter is at most one. Since $N = \Lambda_k^{1/2}(ZZ^\top/m - I)\Lambda_k^{1/2}$, with $s = \sqrt{6\log n}$ the event $\mathcal{G}$ that $\|N_k^{(i)}\| \le B\zeta$ and $\|\bar{N}_k\| \le B\zeta$ for all k and i, where $\zeta^2 = 36K(d + 6\log n)/n$, has $\mathbb{P}(\mathcal{G}^c) \le 4Kn^{-3} \le n^{-2}$ (the bound for $\bar{N}_k$ follows from those for the halves). The condition $(\sqrt{d}+s)/\sqrt{m} \le 1$ holds because $9(\sqrt{d}+s)^2/m \le \zeta^2 = 36\gamma_j S_{\mathrm{nn}}(j)$ and $S_{\mathrm{nn}}(j) \le B^2 + B^{-2} - 2$, once $c \le 1/(36(B^2 + B^{-2} - 2))$; then also $n \ge 36K \cdot 6\log n \ge 4K$. On $\mathcal{G}$, since $w_k \le B^2$, $\nu \le KB^4\zeta^2$ and $\pi_2 \le KB^4\zeta^2$, hence $\nu/\lambda_* \le 36B^6\gamma_j$ and $\pi_2/\lambda_* \le 36B^6\gamma_j$.

Step 3: Schur complement. Let $\mu_0 = (\hat{H}_j)_{jj} = (P_2)_{jj}$, so $|\mu_0| \le \pi_2$, and $A_0 = (\hat{H}_j)_{-j,-j} = H' + P'$. By Step 1, $A_0 - \mu_0 \succeq H'^{1/2}(I + E_1)H'^{1/2} - 2\pi_2 I \succeq (1 - e_0)H'$ with $e_0 = 2\sqrt{\nu/\lambda_*} + 2\pi_2/\lambda_*$, and $e_0 \le 1/4$ on $\mathcal{G}$ once $\gamma_j \le c$. The smallest eigenvalue μ of $\hat{H}_j$ satisfies $\mu \le e_j^\top \hat{H}_j e_j = \mu_0$, so $A_0 - \mu \succ 0$. Any eigenvector for μ then has a nonzero j-th coordinate (otherwise its other coordinates v' would satisfy $(A_0 - \mu)v' = 0$), so the eigenvector is unique and equals $(1, w)/\sqrt{1 + \|w\|^2}$ with $w = -(A_0 - \mu)^{-1}h$. Hence

$$\rho_j := 1 - \langle \hat{q}_j, q_j \rangle^2 = \frac{\|w\|^2}{1 + \|w\|^2} \le \|w\|^2.$$

From $\mu = \mu_0 - h^\top (A_0 - \mu)^{-1} h$ and $(A_0 - \mu)^{-1} \preceq (A_0 - \mu_0)^{-1} \preceq (1 - e_0)^{-1} H'^{-1}$ we get $|\mu| \le \pi_2 + 2\|H'^{-1/2}h\|^2 \le \pi_2 + 4\nu + 4\pi_2^2/\lambda_*$, so $|\mu|/\lambda_* \le C\gamma_j$ on $\mathcal{G}$.

Step 4: expansion. Write $A_0 - \mu = H'^{1/2}(I + E)H'^{1/2}$ with $E = E_1 + H'^{-1/2}(P_2' - \mu)H'^{-1/2}$; on $\mathcal{G}$, $\|E\| \le e := 2\sqrt{\nu/\lambda_*} + (\pi_2 + |\mu|)/\lambda_* \le C\sqrt{\gamma_j} \le \frac{1}{2}$. The identity $(I + E)^{-1} = I - E + E(I + E)^{-1}E$ gives

$$w = -g + H'^{-1}(P' - \mu)H'^{-1}h - H'^{-1/2}E(I + E)^{-1}EH'^{-1/2}h.$$

Let $T_2 = H'^{-1}P_1'H'^{-1}h^{(1)}$ and $T_3 = H'^{-1/2}P_1'H'^{-1}h^{(1)}$, the two second-order terms that are polynomials of the noise alone. Splitting $h = h^{(1)} + h^{(2)}$ and $P' - \mu = P_1' + (P_2' - \mu)$ and using Step 1 for the remaining pieces, on $\mathcal{G}$,

$$\|w\| \le \big(1 + C\gamma_j\big)\|g\| + \|T_2\| + C\sqrt{\gamma_j/\lambda_*}\,\big(\|T_3\| + \|H'^{-1/2}h^{(2)}\|\big).$$

For instance $\|H'^{-1}P_1'H'^{-1}h^{(2)}\| = \|H'^{-1/2}E_1H'^{-1/2}h^{(2)}\| \le \lambda_*^{-1/2}\|E_1\|\,\|H'^{-1/2}h^{(2)}\|$, $\|H'^{-1}(P_2' - \mu)H'^{-1}h\| \le \lambda_*^{-1}(\pi_2 + |\mu|)\|g\|$, and the last term of w has norm at most $2e\lambda_*^{-1/2}\|H'^{-1/2}(P' - \mu)H'^{-1}h\|$, which is bounded in the same way.

Step 5: moments. Off-diagonal entries of the noise are uncorrelated across pairs and environments, with $\mathrm{Var}(\bar{N}_{k,ab}) = \mathbf{p}_a(k)\mathbf{p}_b(k)/n_k$ and $\mathrm{Var}(N_{k,ab}^{(i)}) = \mathbf{p}_a(k)\mathbf{p}_b(k)/m$ for a half, $m = n/(2K)$. For one environment with m' Gaussian samples and $a, b, b' \ne j$ with $a \notin \{b, b'\}$, expanding the four sample averages over time indices,

$$\mathbb{E}\big[N_{ab}N_{bj}N_{ab'}N_{b'j}\big] = \mathbf{1}[b = b']\Big(1 - \frac{1}{m'}\Big)\,\mathrm{Var}(N_{ab})\,\mathrm{Var}(N_{bj}) + \frac{\chi_{bb'}}{m'^3}, \qquad |\chi_{bb'}| \le 3B^4,$$

(the only pairing of time indices that survives is $\{ab, ab'\}\{bj, b'j\}$, and coincident indices contribute $m'^{-3}\mathbb{E}[s_a^2 s_b^2 s_{b'}^2 s_j^2]$, i.e. $\chi_{bb'} = \mathbf{p}_a\mathbf{p}_b\mathbf{p}_{b'}\mathbf{p}_j$ for $b \ne b'$ and $3\mathbf{p}_a\mathbf{p}_b^2\mathbf{p}_j$ for $b = b'$); entries of different environments are independent, and $\mathbb{E}[N_{ab}N_{bj}] = 0$. We use $w_k \le B^2$, $\|\delta_a\|^2 \le B^2\lambda_a$, $\lambda_a \ge KS_{aj}/B^2$, $\|\delta_a\|^2 \le KB^2 S_{aj}$ and $v_{aj}^{\mathrm{qp}} \le \kappa_B S_{aj}^{-1}$, which give, with $E_j^{(4)} = \sum_{i \ne j} S_{ij}^{-2}$,

$$\sum_{a \ne j} \frac{\|\delta_a\|^2}{\lambda_a^2} \le \frac{B^4 E_j}{K}, \qquad \sum_{a \ne j} \frac{1}{\lambda_a^2} \le \frac{B^4 E_j^{(4)}}{K^2}, \qquad \sum_{a \ne j} \frac{1}{\lambda_a} \le \frac{B^2 E_j}{K}, \qquad \sum_{b \ne j} \|\delta_b\|^2 v_{bj}^{\mathrm{qp}} \le \kappa_B B^2 K(d-1). \tag{12}$$

The term T_2. $(T_2)_a = \lambda_a^{-1}\sum_{b \ne j}(P_1')_{ab}\lambda_b^{-1}h_b^{(1)}$, where $(P_1')_{ab} = \sum_k w_k(\delta_a + \delta_b)(k)\bar{N}_{k,ab}$ for $b \ne a$ and $h_b^{(1)} = \sum_k w_k\delta_b(k)\bar{N}_{k,bj}$. In $\mathbb{E}[(P_1')_{ab}(P_1')_{ab'}h_b^{(1)}h_{b'}^{(1)}]$ each factor is a sum over environments; by independence across environments and $\mathbb{E}[N_{ab}N_{bj}] = 0$, the pairings of two different environments vanish unless they are $\{ab, ab'\}\{bj, b'j\}$, which

forces $b = b'$, and one environment gives the display above. Hence, for the terms with $b \neq a$ and up to the χ terms,

$$\mathbb{E}\|T_2\|^2 \le \sum_{a\neq j} \frac{1}{\lambda_a^2} \sum_{b\neq j} \frac{\mathrm{Var}((P_1')_{ab})\,\mathbb{E}(h_b^{(1)})^2}{\lambda_b^2} \le \frac{2B^6K}{n^2}\Big[\sum_a \frac{\|\delta_a\|^2}{\lambda_a^2}\sum_b v_{bj}^{\mathrm{qp}} + \sum_a \frac{1}{\lambda_a^2}\sum_b \|\delta_b\|^2 v_{bj}^{\mathrm{qp}}\Big] \le \frac{C}{n^2}\big(E_j^2 + dE_j^{(4)}\big),$$

by $\mathrm{Var}((P_1')_{ab}) \le (2B^6K/n)(\|\delta_a\|^2 + \|\delta_b\|^2)$, $\mathbb{E}(h_b^{(1)})^2 = \lambda_b^2 v_{bj}^{\mathrm{qp}}/n$ and (12): the factor K of the noise variance cancels against the factors $1/K$ and $1/K^2$ of the sums. *The term* $T_3 = H'^{-1/2}P_1'H'^{-1}h^{(1)}$. The same pairing gives

$$\mathbb{E}\|T_3\|^2 \le \frac{2B^6K}{n^2}\Big[\sum_a \frac{\|\delta_a\|^2}{\lambda_a}\sum_b v_{bj}^{\mathrm{qp}} + \sum_a \frac{1}{\lambda_a}\sum_b \|\delta_b\|^2 v_{bj}^{\mathrm{qp}}\Big] \le \frac{CK}{n^2}\,dE_j.$$

The diagonal entries $(P_1')_{aa} = 2\sum_k w_k\delta_a(k)\bar{N}_{k,aa}$ ($b = a$) have $\mathbb{E}[N_{aa}^2 N_{aj}^2] = (1 + 4/m')\,\mathrm{Var}(N_{aa})\,\mathrm{Var}(N_{aj})$ and $\mathbb{E}[N_{aa}N_{aj}N_{ab'}N_{b'j}] = 2\mathbf{p}_a^2\mathbf{p}_{b'}\mathbf{p}_j/m'^3$, and since $\mathrm{Var}((P_1')_{aa}) \le 8B^4K\lambda_a/n$ they contribute at most $CE_j^{(4)}/n^2$ to $\mathbb{E}\|T_2\|^2$ and CKE_j/n^2 to $\mathbb{E}\|T_3\|^2$. The χ terms add

$$\frac{CK}{n^3}\big(d^2E_j^{(4)} + dE_j^2\big) \text{ to } \mathbb{E}\|T_2\|^2, \qquad \frac{CK^2}{n^3}\,d^2E_j \text{ to } \mathbb{E}\|T_3\|^2,$$

through $|w_k(\delta_a + \delta_b)(k)| \le B^2(\|\delta_a\| + \|\delta_b\|)$, $\sum_k |w_k\delta_b(k)\,w_k\delta_{b'}(k)| \le B^4\|\delta_b\|\|\delta_{b'}\|$ (so the sum over environments of the factors $1/m'^3 = (K/n)^3$ adds no further K), $\lambda_b \ge \|\delta_b\|^2/B^2$, $\sum_b \|\delta_b\|^{-2} \le B^2E_j/K$ and $\sum_b(\|\delta_a\| + \|\delta_b\|)/\|\delta_b\| \le d + \|\delta_a\|\sqrt{d\sum_b\|\delta_b\|^{-2}}$, with $\|\delta_a\|^2$ absorbed into $\lambda_a^{-2} \le B^4\|\delta_a\|^{-4}$ (respectively λ_a^{-1}); note also $dE_j^2 \le d^2E_j^{(4)}$. *The term* $h^{(2)}$. $h_a^{(2)} = \sum_k w_k \sum_b \frac{1}{2}\big(N_{k,ab}^{(1)}N_{k,bj}^{(2)} + N_{k,ab}^{(2)}N_{k,bj}^{(1)}\big)$ with independent halves; different k, and different b, are uncorrelated, and $\mathbb{E}[N_{ab}N_{bj}] = 0$, so

$$\mathbb{E}(h_a^{(2)})^2 = \sum_k w_k^2\,\frac{1}{2}\sum_b \mathrm{Var}(N_{k,ab}^{(1)})\,\mathrm{Var}(N_{k,bj}^{(2)}),$$

which lies between cK^3d/n^2 and CK^3d/n^2; hence $\mathbb{E}\|H'^{-1}h^{(2)}\|^2 \le CKdE_j^{(4)}/n^2$ and $\mathbb{E}\|H'^{-1/2}h^{(2)}\|^2 \le CK^2dE_j/n^2$ by (12). Since $E_j \le d/S_{\mathrm{nn}}(j)$ and $E_j^{(4)} \le E_j/S_{\mathrm{nn}}(j)$, every one of these bounds, after the factors of Step 4, is at most $C\gamma_j\,E_j/n$: for example $(E_j^2 + dE_j^{(4)})/n^2 \le 2(E_j/n)\,d/(nS_{\mathrm{nn}}(j))$, $(\gamma_j/\lambda_*)\,\mathbb{E}\|T_3\|^2 \le C\gamma_j^2E_j/n$ and $KdE_j^{(4)}/n^2 \le \gamma_jE_j/n$.

Step 6: conclusion. On $\mathcal{G}^c$, $\rho_j \le 1$. By Minkowski's inequality on $\mathcal{G}$ and $E_j \le E_j^{\mathrm{qp}}$,

$$\mathbb{E}\rho_j \le \Big(\big(1 + C\gamma_j\big)\big(\mathbb{E}\|g\|^2\big)^{1/2} + C\sqrt{\gamma_j}\big(E_j/n\big)^{1/2}\Big)^2 + \mathbb{P}(\mathcal{G}^c) \le \frac{E_j^{\mathrm{qp}}}{n}\big(1 + C\sqrt{\gamma_j}\big)^2 + \frac{1}{n^2},$$

using $\mathbb{E}\|g\|^2 \le \mathbb{E}\|H'^{-1}h^{(1)}\|^2 + \mathbb{E}\|H'^{-1}h^{(2)}\|^2 \le (E_j^{\mathrm{qp}}/n)(1 + C\gamma_j)$ (the cross term vanishes because $h^{(2)}$ has mean zero given either half while $h^{(1)}$ is linear in the halves). □

Computation. $\hat{H}_j$ combines the $2K + 1$ matrices $\mathrm{sym}(\hat{\Sigma}_k^{(1)}\hat{\Sigma}_k^{(2)})$, $\hat{\Sigma}_k^{(1)} + \hat{\Sigma}_k^{(2)}$ and I, formed once in $O(Kd^3)$, so full eigendecompositions cost $O(d^4)$ for all sources. On $\mathcal{G}$ the gap of $\hat{H}_j$ above its bottom eigenvalue is at least $\frac{3}{4}\lambda_*$ (Step 3) and its largest eigenvalue is $O(K)$ with a constant depending on B, so a Lanczos solver for the bottom eigenvector, at $O(d^2)$ per iteration, needs about $S_{\mathrm{nn}}(j)^{-1/2}\log(d/\varepsilon)$ iterations from a random start, up to factors of B: $O(Kd^3 + d^3S_{\min}^{-1/2}\log(d/\varepsilon))$ for all sources.

Without splitting. The split pencil averaged over all balanced splits of each environment is the U-statistic $\sum_k w_k\frac{n_k}{n_k-1}\big[(\hat{\Sigma}_k - \mathbf{p}_j(k)I)^2 - n_k^{-2}\sum_t(x_tx_t^\top - \mathbf{p}_j(k)I)^2\big]$, the sum over the n_k observations x_t of environment k, with $\hat{\Sigma}_k$ their covariance: the full-sample pencil with the bias of the squared noise removed by a correction computable from the data. We analyse one split because its halves are independent; the U-statistic would need a decoupling argument.

Estimated whitening: proof of part (b). *Reduction.* Let $N_k^{(i)}$ ($i = 0, 1, 2$) be the noise of the sample covariance of part i in environment k, in source coordinates, so that part i has covariances $A(\Lambda_k + N_k^{(i)})A^\top$, and let $\bar{E} = \frac{1}{K}\sum_k N_k^{(0)}$. Since the parts have equal sizes and $\frac{1}{K}\sum_k \Lambda_k = I$, $\hat{C} = A(I + \bar{E})A^\top$. Put $M = \hat{C}^{-1/2}A$. Then $M(I + \bar{E})M^\top = I$, so $M^\top M = (I + \bar{E})^{-1}$, and since $\mathbf{p}_j(k)I = \mathbf{p}_j(k)M(I + \bar{E})M^\top$,

$$\hat{C}^{-1/2}A(\Lambda_k + N_k^{(i)})A^\top\hat{C}^{-1/2} - \mathbf{p}_j(k)I = MY_k^{(i)}M^\top, \qquad Y_k^{(i)} = \tilde{\Lambda}_k + N_k^{(i)} - \mathbf{p}_j(k)\bar{E}.$$

The pencil of (a) on the whitened parts is therefore $MXM^\top$ with $X = \sum_k w_k \operatorname{sym}[Y_k^{(1)}(I + \bar{E})^{-1}Y_k^{(2)}]$, and y is a bottom eigenvector of $MXM^\top$ iff $z = M^\top y$ solves $Xz = \mu(I + \bar{E})z$, with μ the smallest generalised eigenvalue. As $\hat{b}_j^\top x = y^\top Ms$, this is the z of the theorem, and $\rho_j = 1 - z_j^2/\|z\|^2$ is a function of X and $\bar{E}$, that is, of the noise in source coordinates: its law is the same for every invertible A. (For $\bar{E} = 0$ and $A = Q$ this is part (a).) We work in the source basis, where $H = \operatorname{diag}(\lambda_a)$ as before, with parts 1 and 2 as the halves.

Expansion and first order. With $(I + \bar{E})^{-1} = I - \bar{E} + \bar{E}(I + \bar{E})^{-1}\bar{E}$,

$$X = H + P_1 + P_2 + W_1 + W_2, \qquad W_1 = -\sum_k w_k\big[\mathbf{p}_j(k)(\bar{E}\tilde{\Lambda}_k + \tilde{\Lambda}_k\bar{E}) + \tilde{\Lambda}_k\bar{E}\tilde{\Lambda}_k\big],$$

where H, P_1, P_2 are those of (a) and W_2 consists of $\sum_k w_k\mathbf{p}_j(k)^2\bar{E}^2 = K\bar{E}^2$, $-\sum_k w_k\mathbf{p}_j(k)\operatorname{sym}(\bar{E}N_k^{(2)} + N_k^{(1)}\bar{E})$, $-\sum_k w_k \operatorname{sym}[(N_k^{(1)} - \mathbf{p}_j(k)\bar{E})\bar{E}\tilde{\Lambda}_k + \tilde{\Lambda}_k\bar{E}(N_k^{(2)} - \mathbf{p}_j(k)\bar{E}) + (N_k^{(1)} - \mathbf{p}_j(k)\bar{E})\bar{E}(N_k^{(2)} - \mathbf{p}_j(k)\bar{E})]$ and $\sum_k w_k \operatorname{sym}[Y_k^{(1)}\bar{E}(I + \bar{E})^{-1}\bar{E}Y_k^{(2)}]$, each of second order in the noise. Since $\tilde{\Lambda}_k e_j = 0$, $(W_1)_{jj} = 0$ and $(W_1e_j)_a = -\beta_a\bar{E}_{aj}$ with $\beta_a = \sum_k w_k\mathbf{p}_j(k)\delta_a(k) = \sum_k \delta_a(k)/\mathbf{p}_j(k)$. The entries $\bar{E}_{aj}$, $a \neq j$, are uncorrelated, independent of parts 1 and 2, with variance $\frac{1}{Kn_0}\sum_k \mathbf{p}_a(k)\mathbf{p}_j(k)$, $n_0 = n/3$; so the first-order part $H'^{-1}(h^{(1)} + h^{\mathrm{w}})$ of the error vector, $h^{\mathrm{w}} = (W_1e_j)_{-j}$, has second moment $E_j^{\mathrm{qp}}/n_{12} + E_j^{\mathrm{wh}}/n_0$ with $n_{12} = 2n/3$, $E_j^{\mathrm{wh}} = \sum_{a\neq j} v_{aj}^{\mathrm{wh}}$ and $v_{aj}^{\mathrm{wh}} = \beta_a^2\,\frac{1}{K}\sum_k \mathbf{p}_a(k)\mathbf{p}_j(k)/\lambda_a^2$, which is the leading term of (11). By Cauchy–Schwarz, $\beta_a^2 \le \sum_k w_k\delta_a(k)^2 \cdot \sum_k w_k\mathbf{p}_j(k)^2 = K\lambda_a$, so $v_{aj}^{\mathrm{wh}} \le B^2K/\lambda_a \le B^4S_{aj}^{-1}$; $\beta_a = 0$ for every a when $\mathbf{p}_j$ is constant, since $\sum_k \delta_a(k) = 0$.

Steps 1–4. Let $\mathcal{G}$ also require $\|N_k^{(0)}\| \le B\zeta$ for every k, with ζ defined from parts of $n/3$ observations (this changes only constants); then $\|\bar{E}\| \le B\zeta \le \frac{1}{4}$. For every x, since $\sum_k w_k\mathbf{p}_j(k)^2 = K$ and $\tilde{\Lambda}_k e_j = 0$,

$$|x^\top W_1 x| \le 2\sqrt{K}\,\|\bar{E}\|\,\sqrt{x^\top Hx}\,\|x\| + \|\bar{E}\|\,x^\top Hx,$$

so W_1 is relatively bounded like P_1, with ν replaced by $\nu + K\|\bar{E}\|^2$ and an additional $\|\bar{E}\|$ in the bound on $\|E_1\|$; also $\|H'^{-1/2}h^{\mathrm{w}}\|^2 \le K\|\bar{E}\|^2$. Each term of W_2 is bounded in the same way, $|x^\top(\cdot)x| \le C\zeta^2\big(x^\top Hx + \sqrt{K}\sqrt{x^\top Hx}\,\|x\|\big) + CK\zeta^2\|x\|^2$; since $\|H\| \le CK$, also $\|W_2\| \le CK\zeta^2$, which plays the role of π_2 in Steps 3 and 4. Writing $z = (1, w)$, with 1 in the j-th coordinate, the generalised eigenvalue equation gives

$$w = -\big(A_0 - \mu(I + \bar{E})'\big)^{-1}(h - \mu\bar{e}), \qquad \mu(1 + \bar{E}_{jj}) = X_{jj} - (h - \mu\bar{e})^\top\big(A_0 - \mu(I + \bar{E})'\big)^{-1}(h - \mu\bar{e}),$$

with $A_0 = X_{-j,-j}$, $h = (Xe_j)_{-j}$ and $\bar{e} = (\bar{E}e_j)_{-j}$, and $\rho_j \le \|w\|^2$ as in Step 3. On $\mathcal{G}$, $X \succeq -CK\zeta^2 I$ and $X_{jj} = (P_2 + W_2)_{jj} \le CK\zeta^2$, so $|\mu| \le CK\zeta^2 \le C\gamma_j\lambda_*$. By the relative bounds, $A_0 - \mu(I + \bar{E})' \succeq (1 - e)H'$ with $e \le C\sqrt{\gamma_j} \le \frac{1}{2}$, so as in Step 3 the bottom generalised eigenvector is unique and has $z_j \neq 0$; Steps 3 and 4 go through with $A_0 - \mu$ replaced by $A_0 - \mu(I + \bar{E})'$ and h by $h - \mu\bar{e}$: the first change adds at most $|\mu|(1 + \|\bar{E}\|)/\lambda_* \le C\gamma_j$ to the bounds on $\|E\|$, the second adds $\|H'^{-1}\mu\bar{e}\| \le C\gamma_j\|\bar{e}\|$ to $\|w\|$, with $\mathbb{E}\|\bar{e}\|^2 \le 2B^2(d - 1)/n_0 \le CE_j/n$.

Step 5: the remainder. Write $Q_1 = P_1' + W_1'$ for the first-order part of $X' - H'$, $h^{\mathrm{w}2} = (W_2e_j)_{-j}$, and $g = g_1 + g_2$ with $g_1 = H'^{-1}(h^{(1)} + h^{\mathrm{w}})$ and $g_2 = H'^{-1}(h^{(2)} + h^{\mathrm{w}2})$. Step 4 bounds $\|w\|$ by $(1 + C\gamma_j)\|g\|$, the analogues $H'^{-1}Q_1H'^{-1}(h^{(1)} + h^{\mathrm{w}})$ and $H'^{-1/2}Q_1H'^{-1}(h^{(1)} + h^{\mathrm{w}})$ of T_2 and T_3, and $C\sqrt{\gamma_j/\lambda_*}\,\|H'^{-1/2}(h^{(2)} + h^{\mathrm{w}2})\|$. The following lemma, the only place where part (b) needs new moment bounds, controls every term that involves $\bar{E}$; the terms without $\bar{E}$ are those of part (a).

Lemma 2 (Remainder of part (b))**.** *In the setting of Theorem 2(b) there is C depending only on B such that, uniformly in d and K: $(W_1)_{ab} = -c_{ab}\bar{E}_{ab}$ with $c_{ab}^2 \le CK(\|\delta_a\|^2 + \|\delta_b\|^2)$ and $\mathbb{E}(h_b^{\mathrm{w}})^2 \le CK\lambda_b/n$, so that the four vectors $H'^{-1}Q_1H'^{-1}h$ and $H'^{-1/2}Q_1H'^{-1}h$ with $h \in \{h^{(1)}, h^{\mathrm{w}}\}$ obey the bounds of T_2 and T_3 in Step 5 of part (a); and*

$$\mathbb{E}\big[\mathbf{1}_{\mathcal{G}}\|H'^{-1}h^{\mathrm{w}2}\|^2\big] \le C\gamma_j\,\frac{E_j}{n}, \qquad \mathbb{E}\big[\mathbf{1}_{\mathcal{G}}\|H'^{-1/2}h^{\mathrm{w}2}\|^2\big] \le C\gamma_j\lambda_*\,\frac{E_j}{n}. \tag{13}$$

Proof. Each part has $m_0 = n/(3K)$ observations per environment, so $\mathrm{Var}(N^{(i)}_{k,ab}) \le 2B^2/m_0$ and $\mathrm{Var}(\bar E_{ab}) \le 2B^2/n_0$, $n_0 = n/3$, for all entries and $i \in \{0,1,2\}$; moreover $\mathbb{E}\|\bar E\|_F^2 \le 2B^2d^2/n_0$, $\mathbb{E}\|N^{(i)}_k e_j\|^2 \le CKd/n$, $\mathbb{E}\|\bar E e_j\|^2 \le Cd/n$, and on $\mathcal{G}$, $\|N^{(i)}_k\| \le B\zeta$, $\|\bar E\| \le B\zeta \le \frac14$ and $\zeta^2 \le C\gamma_j S_{\mathrm{nn}}(j)$. We use three devices. (D1) *Independence across environments:* for vectors c_k that are functions of $\bar E$ and of the other half, $\mathbb{E}\big[(\sum_k w_k (N_k c_k)_a)^2 \,\big|\, c\big] = \sum_k w_k^2 \sum_b c_{k,b}^2 \,\mathrm{Var}(N_{k,ab}) \le \frac{CK}{n}\sum_k \|c_k\|^2$ for every a, since $\mathbb{E}[N_{k,ab}N_{k,ab'}] = 0$ for $b \ne b'$ and different environments are independent. (D2) *Weights:* if $\mathbb{E}v_a^2 \le \sigma^2$ for every $a \ne j$, then by (12) $\mathbb{E}\|H'^{-1}v\|^2 \le \sigma^2 B^4 E_j^{(4)}/K^2$ and $\mathbb{E}\|H'^{-1/2}v\|^2 \le \sigma^2 B^2 E_j/K$; if instead $\mathbb{E}v_a^2 \le \sigma^2\lambda_a$, then $\mathbb{E}\|H'^{-1}v\|^2 \le \sigma^2 B^2 E_j/K$ and $\mathbb{E}\|H'^{-1/2}v\|^2 \le \sigma^2 d$. (D3) *The regime:* $Kd/n \le \gamma_j S_{\mathrm{nn}}(j)$, $E_j^{(4)} \le E_j/S_{\mathrm{nn}}(j)$, $d \le CE_j$, $S_{\mathrm{nn}}(j) \le C$ and $\lambda_* \ge KS_{\mathrm{nn}}(j)/B^2$; hence $\sigma^2 = CK^2 d\zeta^{2r}/n^2$ with $r \in \{0,1,2\}$ in the first case of (D2), and $\sigma^2 = CKd\zeta^{2r}/n^2$ in the second, give exactly (13). For instance, in the first case $\mathbb{E}\|H'^{-1}v\|^2 \le CdE_j^{(4)}/n^2 \le C(d/(nS_{\mathrm{nn}}(j)))E_j/n \le C\gamma_j E_j/n$ and $\mathbb{E}\|H'^{-1/2}v\|^2 \le CKdE_j/n^2 \le C\gamma_j\lambda_* E_j/n$ because $\gamma_j\lambda_* \ge K^2 d/(nB^2)$.

W_1 and h^{w}. From $W_1 = -\sum_k w_k[\mathbf{p}_j(k)(\bar E\tilde\Lambda_k + \tilde\Lambda_k\bar E) + \tilde\Lambda_k\bar E\tilde\Lambda_k]$, $c_{ab} = \sum_k w_k[\mathbf{p}_j(k)(\delta_a+\delta_b)(k) + \delta_a(k)\delta_b(k)]$, and Cauchy–Schwarz gives $c_{ab}^2 \le CK(\|\delta_a\|^2 + \|\delta_b\|^2)$; $h^{\mathrm{w}}_b = -\beta_b\bar E_{bj}$ with $\beta_b^2 \le K\lambda_b$. So $\mathrm{Var}((W_1)_{ab})$ and $\mathbb{E}(h^{\mathrm{w}}_b)^2$ are bounded like $\mathrm{Var}((P_1')_{ab})$ and $\mathbb{E}(h_b^{(1)})^2$ with n_0 in place of n, and the entries of $\bar E$, a sum of independent environment noises, have the pairing structure of Step 5 of part (a) (the only surviving pairing of two entries of $\bar E$ against two entries of $\bar E$, $N^{(1)}$ or $N^{(2)}$ is again $\{ab, ab'\}\{bj, b'j\}$); the computation of $\mathbb{E}\|T_2\|^2$ and $\mathbb{E}\|T_3\|^2$ therefore applies verbatim to the four mixed vectors.

W_2. Expanding $X = \sum_k w_k \,\mathrm{sym}[Y_k^{(1)}(I-\bar E+R)Y_k^{(2)}]$ with $R = \bar E(I+\bar E)^{-1}\bar E$, $\|R\| \le 2B^2\zeta^2$ and $\|R\|_F \le 2B\zeta\|\bar E\|_F$ on $\mathcal{G}$, and using $\tilde\Lambda_k e_j = 0$ and $Y_k^{(i)}e_j = b_k^{(i)} := N_k^{(i)}e_j - \mathbf{p}_j(k)\bar E e_j$, the vector W_2e_j is the sum of five classes of terms: (i) $K\bar E^2 e_j$; (ii) $-\frac12\sum_k w_k\mathbf{p}_j(k)\big[N_k^{(1)}\bar E e_j + \bar E N_k^{(1)}e_j + \bar E N_k^{(2)}e_j + N_k^{(2)}\bar E e_j\big]$; (iii) $-\frac12\sum_k w_k\tilde\Lambda_k\bar E\big(b_k^{(1)} + b_k^{(2)}\big)$ and (iv) $-\frac12\sum_k w_k\big[(N_k^{(1)} - \mathbf{p}_j(k)\bar E)\bar E b_k^{(2)} + (N_k^{(2)} - \mathbf{p}_j(k)\bar E)\bar E b_k^{(1)}\big]$, which together are $-\sum_k w_k\,\mathrm{sym}[Y_k^{(1)}\bar E Y_k^{(2)}]e_j$; and (v) $\frac12\sum_k w_k\big[Y_k^{(1)}Rb_k^{(2)} + Y_k^{(2)}Rb_k^{(1)}\big]$, whose $\tilde\Lambda_k$-terms are (iii) with $\bar E$ replaced by R and whose other terms are (iv) with the inner $\bar E$ replaced by R. We bound each class, always restricted to the coordinates $a \ne j$.

(i) $(K\bar E^2)_{aj} = K\sum_b \bar E_{ab}\bar E_{bj}$, and the pairing identity of Step 5 gives $\mathbb{E}(K\bar E^2)_{aj}^2 \le K^2\big[\sum_b \mathrm{Var}(\bar E_{ab})\,\mathrm{Var}(\bar E_{bj}) + Cd^2/n_0^3\big] \le CK^2d/n^2$ for every a (the two terms with $b \in \{a,j\}$ involve diagonal entries, whose moments are bounded as at the end of Step 5 of part (a)): (D2) with $r=0$.

(ii) By (D1), conditionally on $\bar E$, $\mathbb{E}(\sum_k w_k\mathbf{p}_j(k)(N_k^{(i)}\bar E e_j)_a)^2 \le \frac{CK}{n}\cdot KB^2\|\bar E e_j\|^2$ and $\mathbb{E}(\sum_k w_k\mathbf{p}_j(k)(\bar E N_k^{(i)}e_j)_a)^2 = \sum_k w_k^2\mathbf{p}_j(k)^2\sum_b \bar E_{ab}^2\,\mathrm{Var}(N^{(i)}_{k,bj}) \le \frac{CK^2}{n}\|\bar E_{a\cdot}\|^2$; taking expectations, $\mathbb{E}v_a^2 \le CK^2d/n^2$ for every a: (D2) with $r=0$.

(iii) The a-th coordinate is $-\frac12\sum_k w_k\delta_a(k)u_{k,a}$ with $u_{k,a} = (\bar E N_k^{(2)})_{aj} + (N_k^{(1)}\bar E)_{ja} - 2\mathbf{p}_j(k)(\bar E^2)_{aj}$. For the first two terms (D1), applied with the factor $\delta_a(k)$ inside the sum over k, gives $\mathbb{E}(\sum_k w_k\delta_a(k)(\bar E N_k^{(2)})_{aj})^2 = \sum_k w_k^2\delta_a(k)^2\sum_b \bar E_{ab}^2\,\mathrm{Var}(N^{(2)}_{k,bj}) \le B^2\lambda_a\cdot\frac{CK}{n}\|\bar E_{a\cdot}\|^2$, and the same for $(N_k^{(1)}\bar E)_{ja}$; the third term is $\beta_a(\bar E^2)_{aj}$ with $\beta_a^2 \le K\lambda_a$ and $\mathbb{E}(\bar E^2)_{aj}^2 \le Cd/n_0^2$. Hence $\mathbb{E}v_a^2 \le CKd\lambda_a/n^2$: the second case of (D2) with $r=0$, which gives $\mathbb{E}\|H'^{-1}v\|^2 \le CdE_j/n^2 \le C\gamma_jE_j/n$ and $\mathbb{E}\|H'^{-1/2}v\|^2 \le CKd^2/n^2 \le C\gamma_j\lambda_*E_j/n$.

(iv) Split $(N_k^{(1)} - \mathbf{p}_j(k)\bar E)\bar E b_k^{(2)} = N_k^{(1)}(\bar E b_k^{(2)}) - \mathbf{p}_j(k)\bar E^2 b_k^{(2)}$. For the first part (D1), conditionally on $(\bar E, N^{(2)})$, gives $\mathbb{E}v_a^2 \le \frac{CK}{n}\sum_k\|\bar E b_k^{(2)}\|^2 \le \frac{CK}{n}B^2\zeta^2\sum_k\|b_k^{(2)}\|^2$ on $\mathcal{G}$, and $\mathbb{E}\sum_k\|b_k^{(2)}\|^2 \le CK^2d/n$; so $\mathbb{E}v_a^2 \le CK^3d\zeta^2/n^2$, and (D2) with $\sigma^2 = CK^3d\zeta^2/n^2$ gives $\mathbb{E}\|H'^{-1}v\|^2 \le CKd\zeta^2E_j^{(4)}/n^2 \le C\gamma_j(Kd/n)E_j/n \le C\gamma_jE_j/n$ and $\mathbb{E}\|H'^{-1/2}v\|^2 \le CK^2d\zeta^2E_j/n^2 \le C\gamma_j\lambda_*E_j/n$, by (D3). For the second part write $\bar E^2\sum_k w_k\mathbf{p}_j(k)b_k^{(2)} = \bar E^2 z_0 - K\bar E^3 e_j$ with $z_0 = \sum_k w_k\mathbf{p}_j(k)N_k^{(2)}e_j$, whose coordinates are centred and uncorrelated given $\bar E$ with $\mathbb{E}z_{0,b}^2 \le CK^2/n$; then $\mathbb{E}\|\bar E^2 z_0\|^2 \le \frac{CK^2}{n}\mathbb{E}\|\bar E^2\|_F^2 \le \frac{CK^2}{n}B^2\zeta^2\frac{Cd^2}{n}$, and $\|H'^{-1}\bar E^2z_0\| \le \lambda_*^{-1}\|\bar E^2z_0\|$ gives $\mathbb{E}\|H'^{-1}\bar E^2z_0\|^2 \le Cd^2\zeta^2/(n^2S_{\mathrm{nn}}(j)^2) \le C\gamma_j^2E_j/n$ and $\mathbb{E}\|H'^{-1/2}\bar E^2z_0\|^2 \le CKd^2\zeta^2/(n^2S_{\mathrm{nn}}(j)) \le C\gamma_j\lambda_*E_j/n$, again by (D3). The remaining term is a function of $\bar E$ alone: $\mathbb{E}[\mathbf{1}_{\mathcal{G}}K^2\|\bar E^3e_j\|^2] \le K^2B^4\zeta^4\,\mathbb{E}\|\bar E e_j\|^2 \le CK^2\zeta^4d/n$, which after the factor λ_*^{-2} is at most $C\zeta^4d/(nS_{\mathrm{nn}}(j)^2) \le C\gamma_j^2d/n \le C\gamma_j^2E_j/n$, and after λ_*^{-1} at most $CK\zeta^4d/(nS_{\mathrm{nn}}(j)) \le C\gamma_j\lambda_*E_j/n$. The terms with 1 and 2 exchanged are identical.

(v) The $\tilde\Lambda_k$-terms are (iii) with $\bar E$ replaced by R. For the first two parts of (iii) the per-coordinate bound becomes $\mathbb{E}[\mathbf{1}_{\mathcal{G}}v_a^2] \le$

$B^2\lambda_a\frac{CK}{n}\|R_{a\cdot}\|^2 \le CK\zeta^4\lambda_a/n$, and the second case of (D2) with $\sigma^2 = CK\zeta^4/n$ gives $\mathbb{E}\|H'^{-1}v\|^2 \le C\zeta^4 E_j/n \le C\gamma_j^2 E_j/n$ and $\mathbb{E}\|H'^{-1/2}v\|^2 \le CKd\zeta^4/n \le C\gamma_j\lambda_* E_j/n$. The third part, $\beta_a(R\bar E)_{aj}$, is summed over a instead: $\sum_a \beta_a^2(R\bar E)_{aj}^2/\lambda_a^2 \le (K/\lambda_*)\|R\bar E e_j\|^2$ and $\sum_a \beta_a^2(R\bar E)_{aj}^2/\lambda_a \le K\|R\bar E e_j\|^2$, with $\mathbb{E}[\mathbf{1}_{\mathcal G}\|R\bar E e_j\|^2] \le 4B^4\zeta^4\,Cd/n$, which gives $C\gamma_j^2 S_{\rm nn}(j)d/n \le C\gamma_j^2 E_j/n$ and $CK\zeta^4 d/n \le C\gamma_j\lambda_* E_j/n$. The other terms are (iv) with the inner $\bar E$ replaced by R, bounded in the same way with $\|R\| \le 2B^2\zeta^2$ in place of $\|\bar E\| \le B\zeta$, $\|\bar E R\|_F \le \|\bar E\|\,\|R\|_F$ in place of $\|\bar E^2\|_F$ and $K\bar E R\bar E e_j$ in place of $K\bar E^3 e_j$: the resulting bounds are those of (iv) multiplied by $\zeta^2 \le C$. Every class therefore satisfies (13), and summing the five classes proves the lemma. □

Step 6. By Lemma 2 and Step 5 of part (a), $\mathbb{E}[\mathbf{1}_{\mathcal G}\|g_2\|^2] \le C\gamma_j E_j/n$, the second moment of g_1 is exactly $\frac{3}{2n}(E_j^{\rm qp} + 2E_j^{\rm wh})$, and the T_2-, T_3-type terms of Step 4 are at most $C\gamma_j E_j/n$ in mean square after the factors of Step 4. Unlike in (a), some cross terms between g_1 and g_2 need not vanish (for instance $h^{\rm w}$ against $K(\bar E^2)_{\cdot j}$), so we do not expand $\mathbb{E}\|g\|^2$: Minkowski's inequality, as in Step 6 with $E_j \le E_j^{\rm qp}$, bounds $(\mathbb{E}\rho_j\mathbf{1}_{\mathcal G})^{1/2}$ by $(1+C\gamma_j)(\mathbb{E}\|g_1\|^2)^{1/2} + C(\gamma_j E_j/n)^{1/2}$, which gives (11). □

Perturbed probes. The estimator uses $\mathbf{p}_j$ only through the shifts $\mathbf{p}_j(k)I$, which commute with every Λ_k, and through the weights. The next lemma replaces both by a probe $\hat{\mathbf{p}}$.

Lemma 3 (Perturbed probes)**.** *In the setting of Theorem 2(a), let $\hat{\mathbf{p}} \in [1/B, B]^K$ be independent of the two halves, and let $\hat q_j$ be a unit bottom eigenvector of $\hat H_j$ with $\mathbf{p}_j$ replaced by $\hat{\mathbf{p}}$ in the shifts and in the weights, $\hat w_k = \hat{\mathbf{p}}(k)^{-2}$. Write $\hat\lambda_a = \sum_k \hat w_k(\mathbf{p}_a(k) - \hat{\mathbf{p}}(k))^2$ and $\hat\delta_a = \delta_a - 2(\hat{\mathbf{p}} - \mathbf{p}_j)$, and suppose that*

$$\|\hat{\mathbf{p}} - \mathbf{p}_j\|^2 \;\le\; \eta^2 K\, S_{\rm nn}(j)/B^6. \tag{14}$$

There are $c, C > 0$ depending only on B such that, if $\eta \le 1/8$ and $\gamma_j \le c$,

$$\mathbb{E}\big[1 - \langle \hat q_j, q_j\rangle^2 \,\big|\, \hat{\mathbf{p}}\big] \le \frac{\hat E_j^{\rm qp}}{n}\big(1 + C\sqrt{\gamma_j}\big)^2 + \frac{1}{n^2}, \qquad \hat E_j^{\rm qp} = \sum_{a\neq j}\frac{K\sum_k \hat w_k^2\hat\delta_a(k)^2\mathbf{p}_a(k)\mathbf{p}_j(k)}{(\hat\lambda_a - \hat\lambda_j)^2},$$

where $E_j \le \hat E_j^{\rm qp} \le B^8\big(\frac{1+2\eta}{1-2\eta}\big)^2 E_j \le 3B^8 E_j$, and $\hat E_j^{\rm qp} \le (1+C\eta)E_j^{\rm qp}$ if moreover $\max_k|\hat{\mathbf{p}}(k)/\mathbf{p}_j(k) - 1| \le \eta$.

Proof. Condition on $\hat{\mathbf{p}}$. Put $\varpi_k = \sqrt{\hat w_k}\,(\hat{\mathbf{p}}(k) - \mathbf{p}_j(k))$ and, for $a \neq j$, $u_a(k) = \sqrt{\hat w_k}\,\delta_a(k)$. Since $\hat{\mathbf{p}} \in [1/B, B]^K$, $\hat w_k \in [B^{-2}, B^2]$, so (3) gives $\|u_a\|^2 \ge B^{-2}\|\delta_a\|^2 \ge KS_{aj}/B^4$, while $\|\varpi\|^2 \le B^2\|\hat{\mathbf{p}} - \mathbf{p}_j\|^2 \le \eta^2 KS_{\rm nn}(j)/B^4 \le \eta^2\|u_a\|^2$. Also $\|\hat{\mathbf{p}} - \mathbf{p}_j\| \le \eta\|\delta_a\|/B^2$ for every $a \neq j$.

(i) The population pencil. $\sum_k \hat w_k(\Lambda_k - \hat{\mathbf{p}}(k)I)^2 = \mathrm{diag}(\hat\lambda_a)$ with $\hat\lambda_a = \|u_a - \varpi\|^2$ for $a \neq j$ and $\hat\lambda_j = \|\varpi\|^2$. Subtracting $\hat\lambda_j I$, which does not change eigenvectors, the estimator is the bottom eigenvector of $\tilde H + \hat P_1 + \hat P_2$ with $\tilde H = \mathrm{diag}(\tilde\lambda_a)$, $\tilde\lambda_a = \hat\lambda_a - \hat\lambda_j = \|u_a\|^2 - 2\langle u_a, \varpi\rangle \in [(1-2\eta)\|u_a\|^2, (1+2\eta)\|u_a\|^2]$, $\tilde\lambda_j = 0$, $\hat P_1 = \sum_k \hat w_k(\hat\Lambda_k\bar N_k + \bar N_k\hat\Lambda_k)$ with $\hat\Lambda_k = \Lambda_k - \hat{\mathbf{p}}(k)I$, and $\hat P_2 = \sum_k \hat w_k\,\mathrm{sym}(N_k^{(1)}N_k^{(2)})$. Hence $\tilde\lambda_* := \min_{a\neq j}\tilde\lambda_a \ge (1-2\eta)KS_{\rm nn}(j)/B^4$ and $\hat\lambda_j \le \eta^2\tilde\lambda_*/(1-2\eta)$.

(ii) Steps 1–4. Let $\hat\nu = \sum_k \hat w_k\|\bar N_k\|^2$ and $\hat\pi_2 = \|\hat P_2\|$; on $\mathcal G$, $\hat\nu \le KB^4\zeta^2$ and $\hat\pi_2 \le KB^4\zeta^2$ as before, so $\hat\nu/\tilde\lambda_*$ and $\hat\pi_2/\tilde\lambda_*$ are at most $C\gamma_j$. For $x \perp e_j$, the computation of Step 1 gives $|x^\top\hat P_1 x| \le 2(\sum_k \hat w_k\|\hat\Lambda_k x\|^2)^{1/2}\sqrt{\hat\nu}\,\|x\|$ with $\sum_k \hat w_k\|\hat\Lambda_k x\|^2 = x^\top\tilde H x + \hat\lambda_j\|x\|^2 \le (1+2\eta^2)\,x^\top\tilde H x$, so $\|\tilde H'^{-1/2}\hat P_1'\tilde H'^{-1/2}\| \le 3\sqrt{\hat\nu/\tilde\lambda_*}$. The entries of the first-order vector are $h_a^{(1)} = \sum_k \hat w_k\hat\delta_a(k)\bar N_{k,aj}$, because $(\mathbf{p}_a - \hat{\mathbf{p}}) + (\mathbf{p}_j - \hat{\mathbf{p}}) = \hat\delta_a$, and $\tilde H'^{-1/2}h^{(1)} = \sum_k \hat M_k y_k$ with $\hat M_k = \sqrt{\hat w_k}\,\tilde H'^{-1/2}\,\mathrm{diag}(\hat\delta_a(k))_{a\neq j}$; since $\sqrt{\hat w}\,\hat\delta_a = u_a - 2\varpi$, $\sum_k \hat M_k\hat M_k^\top = \mathrm{diag}(\|u_a - 2\varpi\|^2/\tilde\lambda_a) \preceq (1+2\eta)^2(1-2\eta)^{-1}I$, and $\|\tilde H'^{-1/2}h^{(1)}\|^2 \le (1+C\eta)\hat\nu$. The diagonal entry is now $\mu_0 = (\hat P_1)_{jj} + (\hat P_2)_{jj}$ with $(\hat P_1)_{jj} = 2\sum_k \hat w_k(\mathbf{p}_j(k) - \hat{\mathbf{p}}(k))\bar N_{k,jj}$, so $|\mu_0| \le \hat\pi_2 + 2\sqrt{\hat\lambda_j\hat\nu} \le \hat\pi_2 + 3\eta\sqrt{\hat\nu/\tilde\lambda_*}\,\tilde\lambda_*$. With λ_*, ν, π_2 replaced by $\tilde\lambda_*$, $(1+C\eta)\hat\nu$ and $\hat\pi_2 + 2\sqrt{\hat\lambda_j\hat\nu}$, Steps 3 and 4 go through with two changes, both because μ_0 now contains the linear term $(\hat P_1)_{jj}$: Step 3 gives $|\mu|/\tilde\lambda_* \le C(\gamma_j + \eta\sqrt{\gamma_j})$ instead of $C\gamma_j$, and the factor $1 + C\gamma_j$ in front of $\|g\|$ in Step 4 becomes $1 + C\sqrt{\gamma_j}$. Step 6 needs only the latter, and on $\mathcal G$ every ratio that Steps 3 and 4 bound is still at most $C\sqrt{\gamma_j}$.

(iii) Steps 5–6. The moment computations of Step 5 have the same form, with δ_a replaced by $\hat\delta_a$ in $h^{(1)}$ and by $\mathbf{p}_a - \hat{\mathbf{p}}$ in the entries of $\hat P_1$, w by $\hat w$, and λ_a by $\tilde\lambda_a$. They used only $w_k \le B^2$, $\lambda_a \ge B^{-2}\|\delta_a\|^2$, $KS_{aj}/B^2 \le \|\delta_a\|^2 \le KB^2 S_{aj}$,

$|w_k(\delta_a+\delta_b)(k)| \le B^2(\|\delta_a\|+\|\delta_b\|)$, and $v_{bj}^{\mathrm{qp}} \le \kappa_B S_{bj}^{-1}$ through $\mathbb{E}(h_b^{(1)})^2 = \lambda_b^2 v_{bj}^{\mathrm{qp}}/n$; the last is replaced by $\hat v_{bj} \le 3B^8 S_{bj}^{-1}$, the bound on each term of $\hat E_j^{\mathrm{qp}}$ proved below. Since $\|\hat{\mathbf{p}}-\mathbf{p}_j\| \le \eta\|\delta_a\|/B^2$, the vectors $\hat\delta_a$ and $\mathbf{p}_a-\hat{\mathbf{p}}$ have norms within factors $1\pm 2\eta$ of $\|\delta_a\|$, and $\tilde\lambda_a \ge (1-2\eta)B^{-2}\|\delta_a\|^2$, so every bound of Step 5 holds with a constant depending only on B. Step 6 gives the claim, with $\mathbb{E}\|\tilde H'^{-1}h^{(1)}\|^2 = \hat E_j^{\mathrm{qp}}/n$.

The coefficient. Write $\hat v_{aj}$ for the a-th term of $\hat E_j^{\mathrm{qp}}$. Since $\tilde\lambda_a = \sum_k \hat w_k \delta_a(k)\hat\delta_a(k)$, Cauchy–Schwarz gives $\tilde\lambda_a^2 \le \sum_k \hat w_k^2 \hat\delta_a(k)^2 \mathbf{p}_a(k)\mathbf{p}_j(k) \cdot \sum_k \delta_a(k)^2/(\mathbf{p}_a(k)\mathbf{p}_j(k))$, that is $\hat v_{aj} \ge S_{aj}^{-1}$, so $\hat E_j^{\mathrm{qp}} \ge E_j$ exactly as in Theorem 2; this is what Step 6 uses. Since $\hat w_k^2 \mathbf{p}_a(k)\mathbf{p}_j(k) \le B^4 \hat w_k$ and $\sqrt{\hat w}\,\hat\delta_a = u_a - 2\varpi$, the numerator of $\hat v_{aj}$ is at most $KB^4\|u_a - 2\varpi\|^2 \le KB^4(1+2\eta)^2\|u_a\|^2$, and the denominator $\tilde\lambda_a^2$ is at least $(1-2\eta)^2\|u_a\|^4$; with $\|u_a\|^2 \ge KS_{aj}/B^4$ this gives $\hat v_{aj} \le B^8\big(\frac{1+2\eta}{1-2\eta}\big)^2 S_{aj}^{-1}$, which is at most $3B^8 S_{aj}^{-1}$ for $\eta \le 1/8$. If also $|\hat{\mathbf{p}}(k)/\mathbf{p}_j(k) - 1| \le \eta$ for every k, then $\hat w_k/w_k \in [(1+\eta)^{-2}, (1-\eta)^{-2}]$, and since $\|\hat\delta_a - \delta_a\| \le 2\eta\|\delta_a\|/B^2$, each term of $\hat E_j^{\mathrm{qp}}$ is within a factor $1+C\eta$ of the corresponding term of E_j^{qp}. □

Signatures known to relative accuracy $\eta\sqrt{S_{\mathrm{nn}}(j)}$ and clipped to $[1/B, B]^K$ (clipping does not increase their distance to $\mathbf{p}_j$, which lies in the box) satisfy (14) with η replaced by $B^4\eta$. In general (14) holds for a probe within η/B^4 of the distance from $\mathbf{p}_j$ to its nearest neighbour, since $KS_{\mathrm{nn}}(j) \ge \min_{i\ne j}\|\mathbf{p}_i - \mathbf{p}_j\|^2/B^2$ by (3); Theorem 3 provides such probes. The lemma is conservative: with $K = 16$ codes ($B = 2$, $d \in \{128, 256\}$, $n \in \{100d, 400d\}$), signatures read off a Jacobi pilot computed on an independent batch change the error by at most 1.4%, although their relative error reaches $0.57\sqrt{S_{\mathrm{nn}}(j)}$, far beyond (14).

Corollary 1. Theorem B.1(a) holds for labelled estimators with the same proof: the reduction (9) only uses the vector assigned to q_j (if its angle to $q_j(\varepsilon)$ is below 0.362, the sign convention makes $\langle \hat q, q_j(\varepsilon)\rangle = \cos\alpha$), and the Haar symmetrisation preserves labels. Under $\gamma \le c$ every source has $\gamma_j \le c$ and $E_{\max} \le d/S_{\min} \le n$, so the lower bound (7) is at least $\frac{c}{2}\bar E/n$, and averaging (10) over j gives the upper bound with $C' = 2\kappa_B$, after decreasing c if necessary so that $(1+C\sqrt{c})^2 \le \frac{3}{2}$ (the term n^{-2} is absorbed because $\gamma \le c$ forces $n\bar E \ge d$). For (i) the lower bound holds a fortiori, since the orthogonal mixings $A = Q$ are invertible and for them $\rho_j = 1 - \langle \hat b_j, q_j\rangle^2/\|\hat b_j\|^2$ is the labelled loss; averaging (11) over j, with $E_j^{\mathrm{qp}} \le \kappa_B E_j$ and $E_j^{\mathrm{wh}} \le B^4 E_j$, gives the upper bound, uniformly in A. *Orthogonal estimators.* Let U be the orthogonal polar factor of $\hat V = [\hat q_1, \dots, \hat q_d]$, its nearest orthogonal matrix in Frobenius norm. For the sign matrix D with $\langle D_{jj}\hat q_j, q_j\rangle \ge 0$, UD is the polar factor of $\hat V D$, so $\|UD - Q\|_F \le \|UD - \hat V D\|_F + \|\hat V D - Q\|_F \le 2\|\hat V D - Q\|_F$. Since $1 - \langle u, q\rangle^2 \le \|u - q\|^2$ and $\|D_{jj}\hat q_j - q_j\|^2 = 2(1 - |\langle \hat q_j, q_j\rangle|) \le 2\rho_j$, $\ell_{\mathrm{A}}^2(U) \le \frac{1}{d}\|UD - Q\|_F^2 \le 8\,\ell_{\mathrm{A}}^2(\hat q)$ deterministically: the same rate holds over $\mathrm{O}(d)$, with $8C'$. Theorem 3 gives the same upper bound over $\mathrm{O}(d)$ (not over labelled estimators) for an estimator that uses only B and a lower bound s_0 on $S_{\min}$; with $s_0 = S_{\min}$ its condition is $\gamma \le c$ for a smaller c. For well-spread signatures $S_{\min} \asymp d^{-2/(K-1)}$. For separated signatures $E_j \le (d-1)/s_0$, while every configuration has $E_j \ge (d-1)/(B^2 + B^{-2} - 2)$, because each summand of (1) is at most $B^2 + B^{-2} - 2$; hence $\bar E \asymp d$ with constants depending only on (B, s_0) and not on K. Setting (iii), the matched loss of Theorem 3(b) for estimators that return d directions, is proved in Appendix D after Lemma 5.

Separated configurations exist with $K \asymp \log d$. Let K be even and let $y_1, \dots, y_d \in \{\pm 1\}^K$ have $K/2$ entries equal to $+1$ each and pairwise Hamming distances at least $K/4$; for independent uniform balanced words the Hamming distance is twice a hypergeometric variable with mean $K/4$, so by Hoeffding's inequality for sampling without replacement (Hoeffding, 1963) a pair violates the condition with probability at most $e^{-K/16}$, and d such words exist once $K \ge 32\log d$. Put $\mathbf{p}_j = \mathbf{1} + \epsilon y_j$ with $\epsilon \in (0,1)$. Then $\sum_k \mathbf{p}_j(k) = K$, $\mathbf{p}_j(k) \in [1-\epsilon, 1+\epsilon]$, and each coordinate where y_i and y_j differ contributes $(2\epsilon)^2/(1-\epsilon^2)$ to KS_{ij}, so $S_{ij} = \frac{4\epsilon^2}{1-\epsilon^2} \cdot \frac{\#\{k : y_i(k) \ne y_j(k)\}}{K} \ge \frac{\epsilon^2}{1-\epsilon^2}$. Conversely, d signatures with $S_{\min} \ge s_0$ are, by (3), $\sqrt{Ks_0}/B$-separated in a $(K-1)$-dimensional body of diameter at most $B\sqrt{K}$, so a volume count gives $d \le (1 + 2B^2/\sqrt{s_0})^{K-1}$, i.e. $K \gtrsim \log d$. □

D. Unknown signatures: proof of Theorem 3

Theorem D.1 (Theorem 3, full statement)**.** *Let the sources be Gaussian and let only B and a lower bound $s_0 \le S_{\min}$ be known. There are $c, C > 0$ depending only on B such that, if $K(d + 6\log n)/(ns_0) \le c$, the following hold for every configuration with $S_{\min} \ge s_0$.*

(a) (Known whitening.) Let $A = Q \in \mathrm{O}(d)$ and $3K \mid n$. Algorithm 2 selects probes on a third of the sample by thresholding $\lambda_{\min}\big(\sum_k (\hat\Sigma_k - g(k)I)^2\big)$ over a grid of points g, runs the pencil of Theorem 2(a) on the other two thirds at each kept probe, and returns the polar factor $\hat Q$ of the d vectors. Then $\sup_Q \mathbb{E}\ell_{\mathrm{A}}^2(\hat Q) \le C\,\bar E/n$; the grid has at most $(C/\sqrt{s_0})^K$ points.

(b) (Estimated whitening.) Let A be any invertible matrix and $4K \mid n$. Split the sample into four equal parts. The second, whitened by its own pooled covariance, selects the probes as in (a), with the mesh and threshold given below; at each kept probe the pencil of Theorem 2(b) is formed from the last two, whitened by the pooled covariance $\hat C$ of the first, giving $\hat b = \hat C^{-1/2}\hat q$. Then $\mathbb{E}\ell_{\mathrm{A}}^2 \le C\,\bar E/n$ for every A, where $\ell_{\mathrm{A}}^2 = \frac{1}{d}\sum_j \min_a \big(1 - z_{a,j}^2/\|z_a\|^2\big)$ with $z_a = A^\top \hat b_a$ is the loss ρ_j of Theorem 2(b) with each source matched to its closest estimate.

The proof takes $c = \min\{B^{-8}/27648,\, 2c_L/3\}$ in part (a), with c_L the constant of Lemma 3, and a smaller qualitative constant in part (b).

Throughout, $\Omega_B = \{\mathbf{p} \in \mathbb{R}^K : \sum_k \mathbf{p}(k) = K\} \cap [1/B, B]^K$, which contains every signature, c_L and C_L are the constants c, C of Lemma 3 (which holds for every $\eta \le 1/8$), and Π is the Euclidean projection onto the convex set Ω_B. We assume $3K \mid n$.

Algorithm 2 Grid probes (the estimator of Theorem 3(a))

input: B and $s_0 \le S_{\min}$; the sample, split into parts 0, 1, 2 of $n/3$ observations, $n/(3K)$ per environment, with covariances $\hat\Sigma_k^{(0)}, \hat\Sigma_k^{(1)}, \hat\Sigma_k^{(2)}$
$\eta \leftarrow 1/8$, $t \leftarrow \eta\sqrt{K s_0}/B^3$, $h \leftarrow t/4$, $\bar\varepsilon \leftarrow 54K^2B^2(d + 6\log n)/n$, $\tau \leftarrow 2h^2 + 2\bar\varepsilon$, $\underline{r} \leftarrow \sqrt{K s_0}/B$
$\mathcal{N} \leftarrow \{\Pi(x) : x \in \{1/B + 2hm/\sqrt{K} : m = 0, 1, \dots\}^K \cap [1/B, B + 2h/\sqrt{K}]^K\}$
accept $g \in \mathcal{N}$ if $\lambda_{\min}\big(\hat H^{(0)}(g)\big) \le \tau$, where $\hat H^{(0)}(g) = \sum_k (\hat\Sigma_k^{(0)} - g(k)I)^2$
cluster the accepted points by single linkage at distance $\underline{r}/2$; if there are not d clusters, return I_d
from each cluster keep one point $\hat{\mathbf{p}}$ (say, with the smallest $\lambda_{\min}$)
for each kept $\hat{\mathbf{p}}$: $\hat q \leftarrow$ unit bottom eigenvector of $\sum_k \hat w_k \,\mathrm{sym}\big[(\hat\Sigma_k^{(1)} - \hat{\mathbf{p}}(k)I)(\hat\Sigma_k^{(2)} - \hat{\mathbf{p}}(k)I)\big]$, $\hat w_k = \hat{\mathbf{p}}(k)^{-2}$
return a polar factor $\hat Q$ of the $d \times d$ matrix of these vectors

The grid. Every $\mathbf{p} \in \Omega_B$ has a point x of the cubic grid with $\|x - \mathbf{p}\| \le \frac{2h}{\sqrt{K}} \cdot \frac{\sqrt{K}}{2} = h$, and $\|\Pi(x) - \mathbf{p}\| = \|\Pi(x) - \Pi(\mathbf{p})\| \le h$ because Π is non-expansive. So $\mathcal{N} \subset \Omega_B$ is an h-net of Ω_B. Its size is at most $\big((B - 1/B)\sqrt{K}/(2h) + 2\big)^K \le \big(16B^4/\sqrt{s_0} + 2\big)^K$. Computing $\sum_k (\hat\Sigma_k^{(0)})^2$ once, each $\hat H^{(0)}(g)$ costs $O(Kd^2)$ and its smallest eigenvalue $O(d^3)$.

Two facts for all probes at once. Work in the basis of Q and let $N_k = \hat\Sigma_k^{(0)} - \Lambda_k$ and $\varepsilon = \sum_k \|N_k\|^2$. For every $g \in \mathbb{R}^K$ and every source j,

$$\lambda_{\min}\big(\hat H^{(0)}(g)\big) \le e_j^\top \hat H^{(0)}(g) e_j = \sum_k \big\|(\mathbf{p}_j(k) - g(k))e_j + N_k e_j\big\|^2 \le 2\|\mathbf{p}_j - g\|^2 + 2\varepsilon, \tag{15}$$

$$\lambda_{\min}\big(\hat H^{(0)}(g)\big) \ge \tfrac{1}{2}\min_a \|\mathbf{p}_a - g\|^2 - \varepsilon. \tag{16}$$

For (16), write $\hat H^{(0)}(g) = \sum_k (D_k + N_k)^2$ with $D_k = \Lambda_k - g(k)I$. For symmetric X, Y, $(X+Y)^2 \succeq \frac{1}{2}X^2 - Y^2$, because $(X+Y)^2 - \frac{1}{2}X^2 + Y^2 = (X/\sqrt{2} + \sqrt{2}Y)^2 \succeq 0$. Hence $\hat H^{(0)}(g) \succeq \frac{1}{2}\sum_k D_k^2 - \sum_k N_k^2 \succeq \frac{1}{2}\,\mathrm{diag}(\|\mathbf{p}_a - g\|^2) - \varepsilon I$.

The noise event. Let $\mathcal{G}_0$ be the event that $\|N_k\| \le B\omega$ for all k, with $\omega = 3(\sqrt{d} + s)/\sqrt{n/(3K)}$ and $s = \sqrt{6\log n}$. As in Step 2 of Appendix C, $\mathbb{P}(\mathcal{G}_0^c) \le 2Ke^{-s^2/2} \le n^{-2}$, provided $(\sqrt{d} + s)^2 \le n/(3K)$, which the assumption below implies. On $\mathcal{G}_0$, $\varepsilon \le KB^2\omega^2 \le 54K^2B^2(d + 6\log n)/n = \bar\varepsilon$.

Selection. Assume

$$\frac{K(d + 6\log n)}{n s_0} \;\le\; \frac{1}{27648\, B^8} = \frac{\eta^2}{432 B^8}, \qquad \text{equivalently} \qquad \bar\varepsilon \le t^2/8. \tag{17}$$

On $\mathcal{G}_0$: every signature $\mathbf{p}_j$ has a point $g \in \mathcal{N}$ with $\|g - \mathbf{p}_j\| \le h$, and by (15) it is accepted. By (16), every accepted g is within $\sqrt{2(\tau+\varepsilon)} \le \sqrt{4h^2 + 6\bar{\varepsilon}} \le \sqrt{t^2/4 + 3t^2/4} = t$ of some signature. Since $S_{ij} \le B^2\|\mathbf{p}_i - \mathbf{p}_j\|^2/K$, signatures are at least $\underline{r}$ apart, and $t = \eta\underline{r}/B^2 \le \underline{r}/8$. Accepted points near the same signature are therefore within $2t \le \underline{r}/4$ of each other, and points near different signatures are more than $\underline{r} - 2t \ge \frac{3}{4}\underline{r}$ apart. Single linkage at $\underline{r}/2$ returns exactly one cluster per source, and every kept $\hat{\mathbf{p}}$ lies in Ω_B within t of the signature of its cluster.

Estimation. Let j be the source of a cluster and $\hat{\mathbf{p}}$ its kept point. Since $S_{\mathrm{nn}}(j) \ge S_{\min} \ge s_0$, $\|\hat{\mathbf{p}} - \mathbf{p}_j\|^2 \le t^2 = \eta^2 K s_0/B^6 \le \eta^2 K S_{\mathrm{nn}}(j)/B^6$, which is (14). Parts 1 and 2 are independent of part 0, and their $2n/3$ observations give $\gamma_j' = K(d + 6\log(2n/3))/((2n/3)S_{\mathrm{nn}}(j)) \le \frac{3}{2}K(d + 6\log n)/(ns_0) \le c_L$ under the condition of the theorem with $c \le \frac{2}{3}c_L$. Lemma 3, applied conditionally on part 0 with parts 1 and 2 as the halves, gives on $\mathcal{G}_0$

$$\mathbb{E}\big[1 - \langle \hat{q}, q_j\rangle^2 \,\big|\, \text{part } 0\big] \le \frac{3\hat{E}_j^{\mathrm{qp}}}{2n}\big(1 + C_L\sqrt{c_L}\big)^2 + \frac{9}{4n^2} \le \frac{CE_j}{n},$$

using $\hat{E}_j^{\mathrm{qp}} \le 3B^8 E_j$ and $E_j \ge (d-1)/(B^2 + B^{-2} - 2)$.

Conclusion. On $\mathcal{G}_0$ the clusters label the vectors by sources, and the polar argument of Appendix C (after permuting the columns, a polar factor of the permuted matrix is the permuted polar factor) gives $\ell_{\mathrm{A}}^2(\hat{Q}) \le 8 \cdot \frac{1}{d}\sum_j (1 - \langle \hat{q}_{(j)}, q_j\rangle^2)$, where $\hat{q}_{(j)}$ is the vector of the cluster of $\mathbf{p}_j$. Off $\mathcal{G}_0$ the loss is at most one. Hence $\mathbb{E}\ell_{\mathrm{A}}^2(\hat{Q}) \le 8C\bar{E}/n + n^{-2} \le C'\bar{E}/n$. The law of N_k, the event $\mathcal{G}_0$ and Lemma 3 do not depend on Q, so the bound holds for every Q, which gives the supremum. Part (a) holds with $c = \min\{1/(27648B^8),\, 2c_L/3\}$. □

Estimated whitening: proof of part (b). The four parts $W, 0, 1, 2$ have $n/4$ observations each and $n/(4K)$ per environment; ζ (the radius called ω in part (a)), $\bar{\varepsilon}$ and the event $\mathcal{G}_0$ are defined as above for parts of this size, $\mathcal{G}_0$ now the event that $\|N_k^{(0)}\| \le B\zeta$ for all k, so that the pooled noise $\bar{E}_0 = \frac{1}{K}\sum_k N_k^{(0)}$ of part 0 satisfies $\|\bar{E}_0\| \le B\zeta \le \frac{1}{4}$ on $\mathcal{G}_0$ (the noise of parts W, 1, 2 is handled inside the lemma below); constants C_0, C depend only on B.

Selection with self-whitening. Part 0 is whitened by its own pooled covariance $\hat{C}_0 = A(I + \bar{E}_0)A^\top$. With $M_0 = \hat{C}_0^{-1/2}A$, $M_0^\top M_0 = (I + \bar{E}_0)^{-1}$, and $g(k)I = g(k)M_0(I + \bar{E}_0)M_0^\top$,

$$\hat{H}^{(0)}(g) = \sum_k \big(\hat{C}_0^{-1/2}\hat{\Sigma}_k^{(0)}\hat{C}_0^{-1/2} - g(k)I\big)^2 = M_0 Z(g) M_0^\top, \qquad Z(g) = \sum_k Y_k(I + \bar{E}_0)^{-1}Y_k,$$

with $Y_k = \Lambda_k + N_k^{(0)} - g(k)(I + \bar{E}_0)$. Since $\lambda_{\min}(M_0 Z M_0^\top) = \min_z z^\top Z z/z^\top(I + \bar{E}_0)z$ and $Z \succeq 0$, it lies between $\lambda_{\min}(Z)/(1 + \|\bar{E}_0\|)$ and $\lambda_{\min}(Z)/(1 - \|\bar{E}_0\|)$. With $\varepsilon' = \sum_k \|N_k^{(0)}\|^2 + KB^2\|\bar{E}_0\|^2$, the two facts (15)–(16) become $\lambda_{\min}(Z) \le e_j^\top Z e_j \le (1 - \|\bar{E}_0\|)^{-1}\sum_k \|Y_k e_j\|^2 \le (3\|\mathbf{p}_j - g\|^2 + 3\varepsilon')/(1 - \|\bar{E}_0\|)$, from $Y_k e_j = (\mathbf{p}_j(k) - g(k))e_j + N_k^{(0)}e_j - g(k)\bar{E}_0 e_j$, and $\lambda_{\min}(Z) \ge (\frac{1}{2}\min_a \|\mathbf{p}_a - g\|^2 - 2\varepsilon')/(1 + \|\bar{E}_0\|)$, from $(X+Y)^2 \succeq \frac{1}{2}X^2 - Y^2$ with $X = \Lambda_k - g(k)I$ and $Y = N_k^{(0)} - g(k)\bar{E}_0$. On $\mathcal{G}_0$ therefore $\lambda_{\min}(\hat{H}^{(0)}(g)) \le 6\|\mathbf{p}_j - g\|^2 + 6\varepsilon'$ for every j and $\lambda_{\min}(\hat{H}^{(0)}(g)) \ge \frac{8}{25}\min_a \|\mathbf{p}_a - g\|^2 - \frac{32}{25}\varepsilon'$, with $\varepsilon' \le \bar{\varepsilon}' := C_0 K^2 B^2 (d + 6\log n)/n$. Algorithm 2 is run with mesh $h = t/8$ and threshold $\tau = 6h^2 + 6\bar{\varepsilon}'$: every signature has an accepted grid point within h, and every accepted point is within $\big(\frac{75}{4}h^2 + \frac{91}{4}\bar{\varepsilon}'\big)^{1/2} \le t$ of a signature once $\bar{\varepsilon}' \le t^2/64$, that is, once $K(d + 6\log n)/(ns_0) \le \eta^2/(64C_0B^8)$. The clustering step is unchanged, so on $\mathcal{G}_0$ the kept points $\hat{\mathbf{p}}$ are within t of their signatures, one per source, and they are functions of part 0 alone. The grid has at most $(32B^4/\sqrt{s_0} + 2)^K$ points.

Estimation. Parts 1 and 2 are whitened by $\hat{C}^{-1/2}$, $\hat{C}$ the pooled covariance of part W, and the pencil of Theorem 2(a) is formed at the probe $\hat{\mathbf{p}}$. The following lemma combines Lemma 3 with the proof of Theorem 2(b).

Lemma 4 (Perturbed probes with estimated whitening)**.** *In the setting of Theorem 2(b), with parts $W, 1, 2$ of $n/4$ observations each, let $\hat{\mathbf{p}} \in [1/B, B]^K$ be independent of parts W, 1, 2 and satisfy (14) with $\eta \le 1/8$. Let $\hat{q}$ be a unit bottom eigenvector of the pencil of Theorem 2(a) at the probe $\hat{\mathbf{p}}$ (shifts $\hat{\mathbf{p}}(k)I$, weights $\hat{w}_k = \hat{\mathbf{p}}(k)^{-2}$) computed from the whitened parts 1 and 2, and let $z = A^\top\hat{C}^{-1/2}\hat{q}$. There are $c, C > 0$ depending only on B such that, if $\gamma_j \le c$, then $\mathbb{E}\big[1 - z_j^2/\|z\|^2 \,\big|\, \hat{\mathbf{p}}\big] \le CE_j/n$ for every invertible A.*

Proof. Condition on $\hat{\mathbf{p}}$ and write $\epsilon_j = \mathbf{p}_j - \hat{\mathbf{p}}$, so that $\hat{\Lambda}_k := \Lambda_k - \hat{\mathbf{p}}(k)I$ has $\hat{\Lambda}_k e_j = \epsilon_j(k)e_j$ and, by the proof of Lemma 3, $|\epsilon_j(k)| \le \eta\|\delta_a\|/B^2$ for every $a \ne j$. The reduction of Theorem 2(b) gives, with $M = \hat{C}^{-1/2}A$, $\bar{E} = \bar{E}_W$

and $M^\top M = (I+\bar E)^{-1}$, that the pencil is $M\hat X M^\top$ with $\hat X = \sum_k \hat w_k \,\mathrm{sym}[\hat Y_k^{(1)}(I+\bar E)^{-1}\hat Y_k^{(2)}]$, $\hat Y_k^{(i)} = \hat\Lambda_k + N_k^{(i)} - \hat{\mathbf p}(k)\bar E$, that z solves $\hat X z = \mu(I+\bar E)z$ for the smallest generalised eigenvalue μ, and that the law of $1 - z_j^2/\|z\|^2$ does not depend on A. Subtracting $\hat\lambda_j(I+\bar E)$, $\hat\lambda_j = \sum_k \hat w_k \epsilon_j(k)^2$, does not change the generalised eigenvectors, and $\tilde X = \hat X - \hat\lambda_j(I+\bar E) = \tilde H + \hat P_1 + \hat P_2 + \hat W_1 + \hat W_2$ with $\tilde H = \mathrm{diag}(\tilde\lambda_a)$, $\tilde\lambda_j = 0$, as in Lemma 3, $\hat P_1, \hat P_2$ as there, and $\hat W_1, \hat W_2$ the terms of W_1, W_2 with $(\tilde\Lambda_k, \mathbf p_j(k), w_k)$ replaced by $(\hat\Lambda_k, \hat{\mathbf p}(k), \hat w_k)$, together with $-\hat\lambda_j\bar E$. Three changes against the two proofs being combined. (1) The first-order whitening term is $(\hat W_1 e_j)_a = -\hat\beta_a \bar E_{aj}$ with $\hat\beta_a = \sum_k \hat w_k\big(\mathbf p_a(k)\mathbf p_j(k) - \hat{\mathbf p}(k)^2\big) + \hat\lambda_j$, and $\hat\beta_a^2 \le CK\tilde\lambda_a$ by Cauchy–Schwarz, since $\mathbf p_a\mathbf p_j - \hat{\mathbf p}^2 = \mathbf p_j(\mathbf p_a - \hat{\mathbf p}) + \hat{\mathbf p}(\mathbf p_j - \hat{\mathbf p})$, $\|\sqrt{\hat w}(\mathbf p_a - \hat{\mathbf p})\|^2 \le (1+\eta)^2\tilde\lambda_a/(1-2\eta)$ and $\|\sqrt{\hat w}\,\epsilon_j\|^2 \le \eta^2\tilde\lambda_a/(1-2\eta)$; hence $\mathbb E\|\tilde H'^{-1}\hat W_1' e_j\|^2 \le CE_j/n$. (2) The relative bounds of Step 1 hold for $\hat W_1$ with $\sum_k \hat w_k\|\hat\Lambda_k x\|^2 \le (1+2\eta^2)x^\top\tilde H x$ for $x \perp e_j$, as in Lemma 3(ii), and $\hat\lambda_j\|\bar E\|\,\|x\|^2 \le \eta^2\|\bar E\|\,x^\top\tilde H x/(1-2\eta)$; $\|\hat W_2\| \le CK\zeta^2$ as before. The bound on μ needs the Schur identity. The diagonal entry $(\tilde X)_{jj} = 2\sum_k \hat w_k\epsilon_j(k)\bar N_{k,jj} - 2\bar E_{jj}\big(\sum_k \hat w_k\hat{\mathbf p}(k)\epsilon_j(k) + \hat\lambda_j\big) + (\hat P_2 + \hat W_2)_{jj}$ now contains first-order terms; by Cauchy–Schwarz and $\sum_k \hat w_k\epsilon_j(k)^2 = \hat\lambda_j \le \eta^2\tilde\lambda_*/(1-2\eta)$, on $\mathcal G$ it satisfies $|(\tilde X)_{jj}| \le C(\eta\sqrt{\gamma_j} + \gamma_j)\tilde\lambda_*$. The relative bounds of $\hat P_1$ and of $\hat W_1$ without its term $-\hat\lambda_j\bar E$ hold against $\mathrm{diag}(\hat\lambda_a)$ for every x, so, as in part (b), $\tilde X \succeq -CK\zeta^2 I$ and $\tilde X \succeq -\big(CK\zeta^2 + \hat\lambda_j(1+\|\bar E\|)\big)I$; with $\mu \le (\tilde X)_{jj}/(1+\bar E_{jj})$ this gives only the a priori bound $|\mu| \le C(\gamma_j+\eta^2)\tilde\lambda_* \le C\tilde\lambda_*$. The identity $\mu(1+\bar E_{jj}) = (\tilde X)_{jj} - (h-\mu\bar e)^\top(A_0 - \mu(I+\bar E)')^{-1}(h-\mu\bar e)$ of part (b), with $A_0 = \tilde X_{-j,-j}$ and $h = (\tilde X e_j)_{-j}$, sharpens it. Since $-\mu(I+\bar E)' \succeq -2|(\tilde X)_{jj}|\,I$, the relative bounds give $A_0 - \mu(I+\bar E)' \succeq (1-e_0)\tilde H'$ with $e_0 \le C\sqrt{\gamma_j} \le \frac14$; moreover $\|\tilde H'^{-1/2}h\|^2 \le CK\zeta^2$, from $\|\tilde H'^{-1/2}h^{(1)}\|^2 \le (1+C\eta)\hat\nu$, $\|\tilde H'^{-1/2}(\hat W_1 e_j)_{-j}\|^2 \le CK\|\bar E\|^2$ and $\|\hat P_2 + \hat W_2\| \le CK\zeta^2$, and $\|\tilde H'^{-1/2}\bar e\|^2 \le B^2\zeta^2/\tilde\lambda_*$. Hence $|\mu|(1-\|\bar E\|) \le |(\tilde X)_{jj}| + CK\zeta^2 + C\zeta^2\mu^2/\tilde\lambda_*$, and by the a priori bound the last term is at most $C\zeta^2|\mu| \le |\mu|/4$ for $\gamma_j \le c$. So $|\mu|/\tilde\lambda_* \le C(\gamma_j + \eta\sqrt{\gamma_j})$ in place of $C\gamma_j$, as in Lemma 3. (3) Lemma 2 applies to the second-order terms with $(\hat\delta_a, \hat\Lambda_k, \hat w_k, \tilde\lambda_a)$ in place of $(\delta_a, \tilde\Lambda_k, w_k, \lambda_a)$, since its proof uses only $w_k \le B^2$, the pairing structure of the noise, $\lambda_a \asymp \|\delta_a\|^2$ up to factors of B, which holds for $\tilde\lambda_a$ and $\|\hat\delta_a\|$ within factors $1 \pm 2\eta$, and the sums (12); the terms created by $\hat\Lambda_k e_j = \epsilon_j(k)e_j$, which vanished for $\tilde\Lambda_k$, obey the same bounds: $\sum_k \hat w_k\epsilon_j(k)N_k^{(i)}\bar E e_j$ and $\big(\sum_k \hat w_k\epsilon_j(k)\hat{\mathbf p}(k)\big)\bar E^2 e_j$ are the classes (ii) and (i) with $\epsilon_j(k)$ in place of $\mathbf p_j(k)$ and $\sum_k \hat w_k\epsilon_j(k)\hat{\mathbf p}(k)$ in place of K, and $\sum_k \hat w_k\epsilon_j(k)^2 = \hat\lambda_j \le \eta^2\tilde\lambda_a/(1-2\eta)$ for every $a \ne j$ keeps their bounds; the same terms with R in place of $\bar E$ obey smaller bounds, since $\|Re_j\| \le 2\|\bar E\|\,\|\bar E e_j\|$ on $\mathcal G$; and the term $\sum_k \hat w_k\epsilon_j(k)\hat\Lambda_k Re_j$, which has one factor $\bar E$ fewer than class (v), has a-th coordinate $R_{aj}\sum_k \hat w_k\epsilon_j(k)(\mathbf p_a(k) - \hat{\mathbf p}(k))$ with $|\sum_k \hat w_k\epsilon_j(k)(\mathbf p_a(k) - \hat{\mathbf p}(k))| \le C\eta\tilde\lambda_a$ by Cauchy–Schwarz, so that $\|\tilde H'^{-1}v\|^2 \le C\eta^2\|Re_j\|^2$ and $\|\tilde H'^{-1/2}v\|^2 \le C\eta^2 K\|Re_j\|^2$, and $\mathbb E[\mathbf 1_{\mathcal G}\|Re_j\|^2] \le C\zeta^2 d/n$ gives the two bounds of (13). Steps 3–6 of Theorem 2(b) then give $\mathbb E[1 - z_j^2/\|z\|^2 \mid \hat{\mathbf p}] \le (2\hat E_j^{\mathrm{qp}} + CE_j)(1 + C\sqrt{\gamma_j})^2/n + n^{-2} \le CE_j/n$ (parts 1 and 2 hold $n/2$ observations and part W holds $n/4$), with $\hat E_j^{\mathrm{qp}} \le 3B^8E_j$ from Lemma 3. □

Conclusion of (b). The kept probes are functions of part 0 and are, on $\mathcal G_0$, within t of their signatures, which is (14) with $\eta = 1/8$ as in part (a); parts W, 1, 2 are independent of part 0, and their sizes give $\gamma_j' \le 2K(d + 6\log n)/(ns_0) \le c$ under the condition of the theorem. By Lemma 4, conditionally on part 0 and on $\mathcal G_0$, the vector of the cluster of source j has $\mathbb E[1 - z_j^2/\|z\|^2] \le CE_j/n$, and the matched loss of the theorem is at most this labelled loss. Off $\mathcal G_0$ the loss is at most one. Hence $\mathbb E\ell_{\mathrm A}^2 \le C\bar E/n + n^{-2} \le C'\bar E/n$ for every invertible A, since the laws of the noises do not depend on A or Q. □

The matched loss is minimax too: Corollary 1(iii). Theorem 3(b) returns d unmixing directions that need not be orthogonal in the true whitened coordinates, and its loss matches every source to its closest estimate, so two sources may be matched to the same estimate. The lower bound of Theorem B.1(a) is for estimators in $\mathrm O(d)$. The next lemma transfers it.

Lemma 5 (Matched loss and the polar factor)**.** *Let $Q \in \mathrm O(d)$ have columns q_j and let $V = [v_1, \dots, v_d]$ have unit columns, not necessarily orthogonal. Put $\varrho_j = 1 - \max_a\langle q_j, v_a\rangle^2$, and let $U \in \mathrm O(d)$ be an orthogonal matrix nearest to V in Frobenius norm (for nonsingular V, its polar factor $V(V^\top V)^{-1/2}$). Then*

$$\sum_{j=1}^d \rho_j(U) \le 16\sum_{j=1}^d \varrho_j, \qquad \rho_j(U) = 1 - \max_a\langle q_j, u_a\rangle^2.$$

Proof. Let $J = \{j : \varrho_j < \frac12\}$ and, for $j \in J$, let $a(j)$ be a column with $\langle q_j, v_{a(j)}\rangle^2 = 1 - \varrho_j$. The map a is injective on J: if $a(i) = a(j) = a$ with $i \ne j$, then $\langle q_i, v_a\rangle^2 + \langle q_j, v_a\rangle^2 \le \|v_a\|^2 = 1$, because $q_i \perp q_j$, so $r_i + \varrho_j \ge 1$, which contradicts $r_i, \varrho_j < \frac12$. Extend a to a bijection π of $\{1, \dots, d\}$ by assigning the sources outside J to the unused columns in any order.

Let $s_j \in \{\pm 1\}$ be the sign of $\langle q_j, v_{\pi(j)}\rangle$ ($s_j = 1$ if it vanishes), and let $P = \sum_j s_j e_{\pi(j)} e_j^\top$, a signed permutation matrix, so that $(VP)e_j = s_j v_{\pi(j)}$ and $(UP)e_j = s_j u_{\pi(j)}$. For $j \in J$, with $x = \langle q_j, v_{\pi(j)}\rangle$,

$$\|s_j v_{\pi(j)} - q_j\|^2 = 2 - 2|x| \le 2(1 - x^2) = 2\varrho_j,$$

since $1 - |x| \le 1 - x^2$ for $|x| \le 1$; for $j \notin J$, $\|s_j v_{\pi(j)} - q_j\|^2 \le 2 \le 4\varrho_j$. Hence $\|V - QP^\top\|_F^2 = \|VP - Q\|_F^2 \le 4\sum_j \varrho_j$. The matrix $QP^\top$ is orthogonal, so the choice of U gives $\|U - V\|_F \le \|QP^\top - V\|_F$, and by the triangle inequality $\|U - QP^\top\|_F \le 2\|V - QP^\top\|_F$. Finally, for $x = \langle q_j, s_j u_{\pi(j)}\rangle \in [-1, 1]$, $\rho_j(U) \le 1 - x^2 \le 2 - 2x = \|s_j u_{\pi(j)} - q_j\|^2$. Summing over j,

$$\sum_j \rho_j(U) \le \|UP - Q\|_F^2 = \|U - QP^\top\|_F^2 \le 4\,\|V - QP^\top\|_F^2 \le 16 \sum_j \varrho_j. \qquad \square$$

Proof of Corollary 1(iii). Let an estimator return nonzero directions $\hat{b}_1, \dots, \hat{b}_d$ (a zero vector may be replaced by any unit vector, which does not increase the matched loss); it may use the signatures and be randomised. Restrict the supremum over invertible mixings to $A = Q \in \mathrm{O}(d)$. Then $z_a = Q^\top \hat{b}_a$, $z_{a,j}^2/\|z_a\|^2 = \langle q_j, v_a\rangle^2$ with $v_a = \hat{b}_a/\|\hat{b}_a\|$, and the matched loss is $\ell_{\mathrm{A}}^2 = \frac{1}{d}\sum_j \varrho_j$ in the notation of Lemma 5. Let U be a nearest orthogonal matrix to $V = [v_1, \dots, v_d]$, chosen as a measurable function of V (a measurable choice exists, as $\mathrm{O}(d)$ is compact and the distance is continuous). Then U is an estimator in $\mathrm{O}(d)$ that may use the signatures, and Theorem B.1(a) applies to it. By (7) and the lemma,

$$\sup_{Q \in \mathrm{O}(d)} \mathbb{E}\,\ell_{\mathrm{A}}^2(\hat{b}) \;\ge\; \frac{1}{16} \sup_{Q \in \mathrm{O}(d)} \mathbb{E}\,\ell_{\mathrm{A}}^2(U) \;\ge\; \frac{c}{16d} \sum_{j=1}^{d} \frac{E_j}{n + E_j} \;\ge\; \frac{c}{32} \cdot \frac{\bar{E}}{n},$$

the last step because $E_{\max} \le d/S_{\min} \le n$ under $\gamma \le c$. The supremum over all invertible A is at least the supremum over $\mathrm{O}(d)$. The upper bound is Theorem D.1(b) with $s_0 = S_{\min}$, whose condition is $\gamma \le c$ for a smaller c; with the signatures, Theorem C.1(b) gives it as well, since the matched loss is at most the labelled loss $\frac{1}{d}\sum_j \rho_j$. For the same reason the lower bound of (iii) implies the lower bound of (i) again, with the constant $c/32$ instead of $c/2$. $\square$

Cost and practice. For well-spread signatures and fixed K, with $s_0 \asymp S_{\min} \asymp d^{-2/(K-1)}$, $|\mathcal{N}| = O(d^{K/(K-1)})$ and the total cost is $O(d^{3+K/(K-1)})$. A grid of the hyperplane itself, with mesh $2h/\sqrt{K-1}$ in an orthonormal basis of $\{\sum_k x(k) = 0\}$ around $\mathbf{1}$ and projected onto Ω_B, is also an h-net of Ω_B, with at most $(32B^4\sqrt{K}/\sqrt{s_0} + 2)^{K-1} = O(d)$ points for fixed K, and the proof is unchanged. For separated signatures with $K \le a \log d$, $|\mathcal{N}| \le (16B^4/\sqrt{s_0} + 2)^{a \log d} = d^{\,a \log(16B^4/\sqrt{s_0}+2)}$, a polynomial whose degree grows with a, B and $1/s_0$; for the codes of Section 4 ($a \ge 32$) it cannot be computed in practice, and the grid is a device for the proof rather than a practical algorithm. The proof uses two facts only: the acceptance rule finds a probe near every signature, and every accepted probe is near some signature. Any procedure with these two properties can replace the grid. Appendix J reports a coarse grid followed by alternating refinement, and Appendix C (after Lemma 3) a Jacobi pilot.

E. Proof of Theorem 4, Corollary 2 and Proposition 1

By (3) it suffices to consider the Euclidean energy $\tilde{E}_j = K\sum_{i \ne j} \|\mathbf{p}_i - \mathbf{p}_j\|^{-2}$ in the $(K-1)$-dimensional normalised signature space, with $r_{\min} = c_1 d^{-1/(K-1)}$. *Upper bound.* The balls of radius $r_{\min}/2$ around the signatures are disjoint, so the number of signatures at distance in $[m r_{\min}, (m+1) r_{\min})$ from $\mathbf{p}_j$ is at most $C m^{K-2}$ for $m \ge 1$ (the volume of the shell divided by that of a ball of radius $r_{\min}/2$), and

$$\tilde{E}_j \le \frac{CK}{r_{\min}^2} \sum_{m=1}^{M} m^{K-4}, \qquad M \le \frac{\mathrm{diam}}{r_{\min}} \asymp d^{1/(K-1)}. \tag{18}$$

For $K = 2$ the series converges, $\sum m^{-2} = \pi^2/6$, and $\tilde{E}_j \le C'/r_{\min}^2 \asymp d^2$; for $K = 3$ it is $\sum m^{-1} \le \log M + 1 \asymp \log d$, giving $d \log d$; for $K \ge 4$ it is $\asymp M^{K-3}$, giving $M^{K-3}/r_{\min}^2 \asymp d^{(K-3)/(K-1)} d^{2/(K-1)} = d$. *Lower bound.* Since Ω is a convex body containing $\mathbf{p}_j$, the shell $\{x \in \Omega :\ m r_d \le \|x - \mathbf{p}_j\| < (m+1) r_d\}$ has volume at least $c\, m^{K-2} r_d^{K-1}$ for $1 \le m \le cM$, with c depending on Ω (a cone with apex $\mathbf{p}_j$ of fixed aperture lies in Ω up to distance $\asymp \mathrm{inradius}(\Omega)$). Every point of this shell is within $C_1 r_d$ of a signature, so the signatures at distance in $[(m - C_1) r_d, (m + 1 + C_1) r_d)$ from $\mathbf{p}_j$

cover the shell with balls of radius $C_1 r_d$ and number at least $c' m^{K-2}$. Summing over m with the shells grouped so that the counts do not overlap, the same sums bound $\tilde{E}_j$ from below, uniformly in j. For arbitrary configurations, all of which lie in the compact convex set $\Omega_B = \{\sum_k \mathbf{p}(k) = K\} \cap [1/B, B]^K$ of positive $(K-1)$-dimensional volume, the lower bound $\bar{E} \gtrsim d$ follows from $S_{ij} \le B^2 \mathrm{diam}^2$ termwise, and the lower bounds for $K \in \{2, 3\}$ are the minimal-energy asymptotics on Ω_B, $N^{1+\beta/D}$ for $\beta > D$ and $N^2 \log N$ for $\beta = D$ (Kuijlaars & Saff, 1998; Hardin & Saff, 2005; Borodachov et al., 2019), applied to $\beta = 2$, $D = K - 1$, $N = d$ and divided by N.

Proof of Corollary 2. Put $f(N) = N^2$, $N \log(N+1)$ and N for $K = 2, 3$ and $K \ge 4$; since $d \log d \le d \log(d+1)$, it suffices to prove the bound with this f. Every configuration lies in Ω_B, and the minimal-energy bounds just used give $c_* > 0$, depending only on (K, B), such that any $N \ge 2$ points of Ω_B satisfy $\sum_{a \ne b} \|\mathbf{p}_a - \mathbf{p}_b\|^{-2} \ge c_* N f(N)$. For $K \ge 4$ this holds because all $N(N-1)$ distances are at most $\mathrm{diam}\, \Omega_B$. For $K = 2$ the $N - k$ distances between points k apart in the sorted order sum to at most $k \,\mathrm{diam}\, \Omega_B$, so Jensen's inequality for $k = 1$ already gives order N^3. For $K = 3$ it is the case $\beta = D$ of the asymptotics, together with positivity of the minimal energy for each fixed N. The largest energy within a subset is at least the average, so by (3) any $N \ge 2$ of the signatures contain one, j, with $E_j \ge (K/B^2) c_* f(N) =: c_3 f(N)$. Let $m = \max(2, \lceil d/2 \rceil) \le d$ and $\mathcal{A} = \{j : E_j < c_3 f(m)\}$. If $|\mathcal{A}| \ge m$, then any m members of $\mathcal{A}$ contain one with $E_j \ge c_3 f(m)$, which is impossible. Hence at least $d - m + 1 \ge d/2$ sources have $E_j \ge c_3 f(m) \ge c_3 f(d)/4$. Since $x \mapsto x/(n+x)$ is increasing, (7) gives $\sup_Q \mathbb{E}\ell_{\mathrm{A}}^2 \ge \frac{c}{2}\, x/(n+x)$ with $x = c_3 f(d)/4$, and $x/(n+x) \ge \frac{1}{2} \min(1, c_3/4) \min(1, f(d)/n)$. □

Proposition 1 (Random signatures)**.** *Let the normalised signatures be i.i.d. from a density bounded above and below on a convex body in the hyperplane. For every $\epsilon > 0$ there are constants such that, with probability at least $1 - \epsilon$,*

$$\bar{E} \asymp d^3,\ d \log d,\ d \quad (K = 2,\ 3,\ \ge 4),$$
$$E_{\max} \asymp d^4,\ d^2,\ d^{4/3},\ d \quad (K = 2,\ 3,\ 4,\ \ge 5).$$

Random signatures: proof of Proposition 1. Write $D = K - 1$. By (3) it suffices to prove the statements for the Euclidean energies, now without the fixed factor K: $\tilde{E}_j = \sum_{i \ne j} \|\mathbf{p}_i - \mathbf{p}_j\|^{-2}$, $\tilde{\bar{E}} = \frac{1}{d} \sum_j \tilde{E}_j$ and $\tilde{E}_{\max}$. Let the signatures be i.i.d. with density g on a convex body Ω of the hyperplane, $0 < g_{\min} \le g \le g_{\max}$. For $x \in \Omega$ and $t \le \mathrm{diam}\, \Omega$, convexity gives $c\, t^D \le \mathrm{vol}(B(x,t) \cap \Omega) \le \omega_D t^D$, so $\pi(x,t) = \mathbb{P}(\|\mathbf{p} - x\| \le t) \in [c_\pi t^D, C_\pi t^D]$ with $c_\pi = c g_{\min}$, $C_\pi = \omega_D g_{\max}$. Constants C, C_ϵ below depend on $(D, g_{\min}, g_{\max}, \Omega)$ and on ϵ, and change from line to line. Let $N_j(t)$ be the number of signatures $i \ne j$ within distance t of $\mathbf{p}_j$, and $N_{\mathrm{P}}(t)$ the number of unordered pairs within distance t. The layer-cake formula gives

$$\tilde{E}_j = \int_0^\infty 2t^{-3} N_j(t)\, dt, \qquad \sum_{i<j} \|\mathbf{p}_i - \mathbf{p}_j\|^{-2} = \int_0^\infty 2t^{-3} N_{\mathrm{P}}(t)\, dt. \tag{19}$$

Step 1: bulk of every centre. Conditionally on $\mathbf{p}_j$, $N_j(t) \sim \mathrm{Bin}(d-1, \pi(\mathbf{p}_j, t))$. Let $t_1 = (A \log d/d)^{1/D}$ with A large. For $t \ge t_1$, $\mathbb{E}[N_j(t) \mid \mathbf{p}_j] \ge c_\pi d t^D/2 \ge (c_\pi A/2) \log d$, and the Chernoff bound gives $\mathbb{P}(N_j(t) \ge 2\mathbb{E}[N_j(t) \mid \mathbf{p}_j]) \le d^{-c_\pi A/6}$. On the dyadic grid $t_m = 2^m t_1$, $m \le \log_2 d$, a union bound over j and m shows that with probability at least $1 - d^{-2}$ (for A large), for all j and all $t \in [t_1, \mathrm{diam}\, \Omega]$, $N_j(t) \le N_j(t_{m+1}) \le 2^{D+1} C_\pi d t^D$. On this event the part of (19) above t_1 is at most $2^{D+2} C_\pi d \int_{t_1}^{\mathrm{diam}} t^{D-3} dt$, which is at most Cd for $D \ge 3$, $Cd \log d$ for $D = 2$ and $Cd/t_1 \le Cd^2/\log d$ for $D = 1$, uniformly in j; and every centre has at most $M_1 = 2^{D+1} C_\pi A \log d$ neighbours below t_1.

Step 2: near neighbours of every centre. Let $r_{(k)}(j)$ be the distance from $\mathbf{p}_j$ to its k-th nearest neighbour. Since $N_j(t) \ge k$ implies $N_j(t)(N_j(t)-1)\cdots(N_j(t)-k+1) \ge k!$, and the falling factorial moments of a binomial are $(d-1)^{(k)} \pi^k \le (C_\pi d t^D)^k$,

$$\mathbb{P}\big(\exists j :\ r_{(k)}(j) \le t\big) \le d\, (C_\pi d t^D)^k / k!.$$

Choose t_k by $d(C_\pi d t_k^D)^k/k! = \epsilon 2^{-k-2}$. Then $\sum_k \mathbb{P}(\exists j : r_{(k)}(j) \le t_k) \le \epsilon/4$, and $(k!)^{1/k} \ge k/e$ gives $t_k^{-2} \le (2eC_\pi d/k)^{2/D} (4d/\epsilon)^{2/(kD)}$. On the complementary event, for every j the part of (19) below t_1 is at most

$$\sum_{k \le M_1} t_k^{-2} \le C_\epsilon \Big(d^{4/D} + d^{3/D} + d^{8/(3D)} \sum_{3 \le k \le M_1} k^{-2/D} \Big) \le C_\epsilon \big(d^{4/D} + d^{8/(3D)} \log d \big) \le C_\epsilon d^{4/D}.$$

With Step 1, with probability at least $1 - \epsilon/2$,

$$\tilde{E}_{\max} \le C_\epsilon \big(d^{4/D} + d\, [1,\ \log d,\ d/\log d]_{D \ge 3,\ D=2,\ D=1} \big) \le C_\epsilon\, (d^4,\ d^2,\ d^{4/3},\ d,\ d) \quad (D = 1, 2, 3, 4, \ge 5). \tag{20}$$

Step 3: the average. By (19), $\tilde{\tilde{E}} = \frac{2}{d}\int_0^\infty 2t^{-3}N_{\mathrm{P}}(t)\,dt$. Here $\mathbb{E}N_{\mathrm{P}}(t) = \binom{d}{2}\mathbb{E}\pi(\mathbf{p},t) \in [c_\pi d^2 t^D/4,\ C_\pi d^2 t^D/2]$ for $t \le \operatorname{diam}\Omega$, and, since indicators of disjoint pairs are independent and at most d^3 ordered pairs of pairs share a point, $\operatorname{Var} N_{\mathrm{P}}(t) \le \mathbb{E}N_{\mathrm{P}}(t) + d^3 C_\pi^2 t^{2D} \le \mathbb{E}N_{\mathrm{P}}(t)\,(1 + Cdt^D)$. Let $t_2 = (M/(c_\pi d^2))^{1/D}$ with $M = 64/\epsilon$. On the dyadic grid $t_m = 2^m t_2$, Chebyshev's inequality gives $\mathbb{P}(N_{\mathrm{P}}(t_m) \ge 2\mathbb{E}N_{\mathrm{P}}(t_m)) \le (1 + Cdt_m^D)/\mathbb{E}N_{\mathrm{P}}(t_m) \le 4\cdot 2^{-mD}/M + C/d$; summed over the grid this is at most $8/M + C\log d/d \le \epsilon/4$ for $d \ge d_0(\epsilon)$. Hence, with probability at least $1-\epsilon/4$, $N_{\mathrm{P}}(t) \le 2^{D+1}C_\pi d^2 t^D$ for all $t \ge t_2$, and the part of the pair energy above t_2 is at most $Cd^2\int_{t_2}^{\operatorname{diam}} t^{D-3}dt \le Cd^2\,[1,\ \log d,\ 1/t_2]$, i.e. at most $C_\epsilon(d^2,\ d^2\log d,\ d^4)$ for $D \ge 3$, $D = 2$, $D = 1$. Below t_2 there are at most $N_{\mathrm{P}}(t_2) \le 2\mathbb{E}N_{\mathrm{P}}(t_2) \le M_2 := (C_\pi/c_\pi)M$ pairs. For the k-th smallest pair distance $r_{[k]}$ and fixed $k \le M_2$, $\mathbb{P}(r_{[k]} \le t) \le \mathbb{E}[N_{\mathrm{P}}(t)^{(k)}]/k!$, and the falling factorial moment is a sum over ordered k-tuples of distinct pairs. A k-tuple spans a graph with v vertices and c components, has probability at most $(C_\pi t^D)^{v-c}$ (place the vertices along a spanning forest), and there are at most $d^v(2k)^{2k}$ tuples with given (v,c). The disjoint tuples ($v = 2k$, $c = k$) contribute at most $(C_\pi d^2 t^D)^k$; every other tuple has $v - c - k \le 0$ and $2c - v \le -1$, so at a scale t with $C_\pi d^2 t^D = \kappa$ its contribution is at most $(2k)^{2k}\max(1,\kappa)^k d^{-1}$. Choosing t_k by $2(C_\pi d^2 t_k^D)^k/k! = \epsilon 2^{-k-2}$ (so that $\kappa_k = C_\pi d^2 t_k^D$ depends only on k and ϵ), $\sum_{k\le M_2}\mathbb{P}(r_{[k]} \le t_k) \le \epsilon/4$ for $d \ge d_0(\epsilon)$, and on the complementary event the pair energy below t_2 is at most $\sum_{k\le M_2} t_k^{-2} \le C_\epsilon d^{4/D}$. Dividing by $d/2$, with probability at least $1-\epsilon$,

$$\tilde{\tilde{E}} \le C_\epsilon\big(d^{4/D-1} + [d,\ d\log d,\ d^3]_{D\ge 3,\ D=2,\ D=1}\big) \le C_\epsilon\,(d^3,\ d\log d,\ d) \quad (D = 1,\ 2,\ \ge 3). \tag{21}$$

Lower bounds. The minimal-energy asymptotics of Theorem 4 hold for every configuration, so $\tilde{\tilde{E}} \ge c\,(d^2,\ d\log d,\ d)$ for $D = 1, 2, \ge 3$ surely, and $\tilde{E}_{\max} \ge \tilde{\tilde{E}}$. For the closest pair, with $t = (M'/(c_\pi d^2))^{1/D}$ and $M' = 64/\epsilon$, Chebyshev gives $\mathbb{P}(N_{\mathrm{P}}(t) = 0) \le \operatorname{Var} N_{\mathrm{P}}(t)/(\mathbb{E}N_{\mathrm{P}}(t))^2 \le 4(1 + CM'/d)/M' \le \epsilon/4$ for $d \ge d_0(\epsilon)$. Hence, with probability at least $1-\epsilon/4$, some pair is within distance t, its two members have energy at least $t^{-2} \ge c_\epsilon d^{4/D}$, and $\tilde{\tilde{E}} \ge \frac{2}{d}t^{-2} \ge c_\epsilon d^{4/D-1}$. For $D = 1$ this gives d^3 and d^4; for $D = 2, 3$ it gives $\tilde{E}_{\max} \ge c_\epsilon d^2,\ c_\epsilon d^{4/3}$; for $D \ge 4$, $\tilde{E}_{\max} \ge \tilde{\tilde{E}} \ge cd$. Together with (20)–(21) this proves the proposition: the exceptional events have total probability at most $\epsilon + d^{-2}$, and replacing ϵ by $\epsilon/2$ gives the statement. □

Effective dimension at a scale. The identity $\sum_{i\ne j}\|\mathbf{p}_i - \mathbf{p}_j\|^{-2} = \int_0^\infty 2t^{-3}N_j(t)\,dt$ of (19) holds for every configuration, with $N_j(t)$ the number of other signatures within distance t of $\mathbf{p}_j$. The energy therefore depends on the configuration only through the growth of N_j between the nearest-neighbour distance and the diameter, and an effective dimension can be read at each scale as the local growth exponent of N_j. For instance, signatures within distance τ of a D'-dimensional affine subspace and well spread there at scales above τ have $N_j(t) \asymp d\,(t/\operatorname{diam})^{D'}$ for $\max(\tau, r_{\mathrm{nn}}) \lesssim t \le \operatorname{diam}$, r_{nn} the nearest-neighbour distance; these scales contribute to the energy as in dimension D', and Remark 1 applies approximately when τ is below the nearest-neighbour distance.

F. Proof of Theorem 5

Theorem 5 states that the means of the limit laws average over j to at least $(d-1)^3/(dB^4)$. We prove the sharper bound

$$\frac{1}{d}\sum_j\sum_{i\ne j}\frac{v_{ij}(\alpha)}{(\lambda_i - \lambda_j)^2} \ \ge\ \frac{K\|\alpha\|^2}{B^2}\cdot\frac{(d-1)^3}{d\,L_\alpha^2}, \tag{22}$$

where $L_\alpha \le \|\alpha\|\operatorname{diam}\{\mathbf{p}_j\}$ is the length of the range of the λ_j; since $\operatorname{diam}^2\{\mathbf{p}_j\} \le KB^2$, the right-hand side is at least $(d-1)^3/(dB^4)$.

Write $\hat{M}_\alpha = M_\alpha + N$ with $N = \sum_k \alpha_k(\hat{\Sigma}_k - \Sigma_k)$. In the basis of Q, $N' = Q^\top N Q$ has entries that are uncorrelated across pairs, with $\operatorname{Var}(N'_{ij}) = \sum_k \alpha_k^2 \operatorname{Var}((\hat{\Sigma}_k)_{ij}) = \sum_k \alpha_k^2 \mathbf{p}_i(k)\mathbf{p}_j(k)/n_k = v_{ij}(\alpha)/n$. With distinct eigenvalues, first-order eigenvector perturbation gives $\hat{q}_j = q_j + \sum_{i\ne j}\frac{N'_{ij}}{\lambda_j - \lambda_i}q_i + O(\|N\|^2)$, and $\sqrt{n}\,N'_{ij}$ are asymptotically independent $N(0, v_{ij}(\alpha))$ by the central limit theorem, hence the limit law of $n\rho_j$ (Anderson, 1963). For (22), $v_{ij}(\alpha) \ge K\|\alpha\|^2/B^2$. Let $g_1, \dots, g_{d-1}$ be the consecutive spacings of the sorted λ's, $\sum_m g_m \le L_\alpha$. The nearest-neighbour distance of a point is the smaller of its two adjacent spacings, so $\sum_j\sum_{i\ne j}(\lambda_i - \lambda_j)^{-2} \ge \sum_j \max(g_{\mathrm{left}}(j)^{-2}, g_{\mathrm{right}}(j)^{-2}) \ge \frac{1}{2}\sum_j\big(g_{\mathrm{left}}(j)^{-2} + g_{\mathrm{right}}(j)^{-2}\big) = \sum_m g_m^{-2}$, and by the power-mean inequality $\frac{1}{d-1}\sum_m g_m^{-2} \ge \big(\frac{1}{d-1}\sum_m g_m\big)^{-2} \ge (d-1)^2/L_\alpha^2$. Finally $\operatorname{diam}^2\{\mathbf{p}_j\} \le \sum_k (B - 1/B)^2 \le KB^2$. With estimated whitening the noise N is replaced

by $\sum_k \alpha_k N_k^{\mathrm{w}}$ with $N_{k,ij}^{\mathrm{w}} = N_{k,ij} - \frac{1}{2}(\mathbf{p}_i(k) + \mathbf{p}_j(k))\bar{N}_{ij}$, whose variance is $K\sum_k(\alpha_k - \bar{\alpha}_{ij})^2\mathbf{p}_i(k)\mathbf{p}_j(k)/n$ with $\bar{\alpha}_{ij} = \frac{1}{2K}\sum_k \alpha_k(\mathbf{p}_i(k)+\mathbf{p}_j(k))$; since $\sum_k(\alpha_k - \bar{\alpha}_{ij})^2\mathbf{p}_i\mathbf{p}_j \geq \|\alpha - \bar{\alpha}_{ij}\mathbf{1}\|^2/B^2 \geq \mathrm{dist}(\alpha, \mathbb{R}\mathbf{1})^2/B^2$, the same argument gives (22) with $\|\alpha\|$ replaced by $\mathrm{dist}(\alpha, \mathbb{R}\mathbf{1})$. □

G. Joint diagonalization: proof of Proposition 2

Proposition 2 (Joint diagonalization). *Let $\hat{U}$ be a global minimiser of $\hat{f}$ over $\mathrm{O}(d)$, let the sources be independent with finite fourth moments, and let $S_{ij} > 0$ for all $i \neq j$.*

(i) (First-order variance, distribution-free.) As $n \to \infty$ at fixed d, $n\,\rho_j(\hat{U}) \Rightarrow \sum_{i\neq j} v_{ij} Z_i^2$, Z_i i.i.d. standard normal, with explicit coefficients v_{ij}, one formula for known whitening and one for the unmixing rows with estimated whitening ((23) and (24) below), both satisfying $S_{ij}^{-1} \leq v_{ij} \leq C_B S_{ij}^{-1}$ with $C_B = 2B^4(1+B^6)$.

(ii) (Efficiency of weighting.) With known whitening, weights $1/(\hat{\mathbf{p}}_a(k)\hat{\mathbf{p}}_b(k))$ on the (a,b) entries, with consistent $\hat{\mathbf{p}}$, give $v_{ij} = S_{ij}^{-1}$, attaining Theorem B.1(b).

For a pair (i,j) write $\delta = \mathbf{p}_i - \mathbf{p}_j \in \mathbb{R}^K$ and $\bar{\delta}_j = \frac{1}{K}\sum_k \mathbf{p}_j(k)\delta_k$. The coefficients of part (i) are

$$v_{ij}^{\mathrm{F}} = \frac{K\sum_k \delta_k^2\,\mathbf{p}_i(k)\mathbf{p}_j(k)}{\|\delta\|^4}, \tag{23}$$

$$v_{ij}^{\mathrm{w}} = \frac{K\sum_k (\delta_k - \bar{\delta}_j)^2\,\mathbf{p}_i(k)\mathbf{p}_j(k)}{\|\delta\|^4}, \tag{24}$$

the second for the loss of the unmixing row of source j against the columns of A; $\bar{\delta}_j$ uses the source whose row is estimated, so $v_{ij}^{\mathrm{w}} \neq v_{ji}^{\mathrm{w}}$ in general. We write $E_j^{\mathrm{F}} = \sum_{i\neq j} v_{ij}^{\mathrm{F}}$, $E_j^{\mathrm{w}} = \sum_{i\neq j} v_{ij}^{\mathrm{w}}$, and $\bar{E}^{\mathrm{F}}$, $\bar{E}^{\mathrm{w}}$ for their averages over j. With environment proportions π_k and the π-weighted criterion and pooled covariance, the same derivation gives the variance of $\sum_k \beta_k N_{k,ab}$ with $\beta_k = \omega_k\delta_k/\Omega_{ab} - \pi_k c_{ab}$, $\Omega_{ab} = \sum_k \omega_k\delta_k^2$, $c_{ab} = \frac{1}{2}\big(\sum_k \omega_k\delta_k(\mathbf{p}_a(k) + \mathbf{p}_b(k))/\Omega_{ab} + 1\big)$, for weights $\omega_k = \pi_k$ (Frobenius) or $\omega_k = \pi_k/(\mathbf{p}_a(k)\mathbf{p}_b(k))$ (Fisher); for the non-orthogonal problem the efficient coefficient is the Cramér–Rao bound $v_{ij}^{\mathrm{PC}} = \omega_{ji}/(\omega_{ij}\omega_{ji} - 1)$ with $\omega_{ij} = \sum_k \pi_k\mathbf{p}_i(k)/\mathbf{p}_j(k)$, which satisfies $v_{ij}^{\mathrm{PC}} \geq S_{ij}^{-1}$ because $(\omega_{ji} - 1)^2 \geq 0$.

(i) First-order variance, known whitening. At fixed d the global minimiser is consistent: the population criterion has the true Q, up to signs and permutations, as its unique minimiser when all $S_{ij} > 0$, and $\hat{f} \to f$ uniformly on $\mathrm{O}(d)$. So we may write $V = Q^\top\hat{U} = \exp(\Xi)$ with Ξ antisymmetric and small. In the basis of Q, $\hat{\Sigma}_k$ becomes $\Lambda_k + N_k$ with $N_k = Q^\top(\hat{\Sigma}_k - \Sigma_k)Q$, and for $a \neq b$,

$$(V^\top(\Lambda_k + N_k)V)_{ab} = N_{k,ab} + \Xi_{ab}\big(\mathbf{p}_a(k) - \mathbf{p}_b(k)\big) + O(\|\Xi\|^2 + \|\Xi\|\|N_k\|).$$

The criterion decouples at first order into independent pair problems $\min_{\Xi_{ab}} \sum_k (N_{k,ab} + \Xi_{ab}\delta_k)^2$ with $\delta = \mathbf{p}_a - \mathbf{p}_b$, solved by $\Xi_{ab} = -\sum_k \delta_k N_{k,ab}/\|\delta\|^2$, i.e. $\Xi_{ba} = \sum_k \delta_k N_{k,ab}/\|\delta\|^2$. For independent centred sources, $\mathbb{E}[s_a s_b s_{a'} s_{b'}]$ vanishes unless $\{a,b\} = \{a',b'\}$ and $\mathbb{E}[s_a^2 s_b^2] = \mathbf{p}_a\mathbf{p}_b$, so the $N_{k,ab}$ are uncorrelated across pairs and environments with $\mathrm{Var}(N_{k,ab}) = \mathbf{p}_a(k)\mathbf{p}_b(k)/n_k$, whatever the law of the sources. The central limit theorem for the $N_{k,ab}$ and $\rho_j = \sum_{i\neq j}\Xi_{ji}^2 + O(\|\Xi\|^3)$ give the limit law, whose mean is (23). The lower bound $v_{ij}^{\mathrm{F}} \geq S_{ij}^{-1}$ is Cauchy–Schwarz, $\|\delta\|^4 \leq \sum_k \delta_k^2\mathbf{p}_i\mathbf{p}_j \cdot \sum_k \delta_k^2/(\mathbf{p}_i\mathbf{p}_j)$, with equality iff $\mathbf{p}_i(k)\mathbf{p}_j(k)$ is constant over the k with $\mathbf{p}_i(k) \neq \mathbf{p}_j(k)$; the upper bound $v_{ij}^{\mathrm{F}} \leq B^4 S_{ij}^{-1}$ uses $\mathbf{p}_i\mathbf{p}_j \in [B^{-2}, B^2]$.

(i) Estimated whitening. Let $\bar{N} = \frac{1}{K}\sum_k N_k$ be the noise of the pooled covariance in the basis of Q; the pooled population covariance is the identity by the normalisation. The estimated whitening matrix is $(I + \bar{N})^{-1/2} = I - \frac{1}{2}\bar{N} + O(\|\bar{N}\|^2)$, and the whitened environment covariances are

$$(I - \tfrac{1}{2}\bar{N})(\Lambda_k + N_k)(I - \tfrac{1}{2}\bar{N}) = \Lambda_k + N_k - \tfrac{1}{2}(\bar{N}\Lambda_k + \Lambda_k\bar{N}) + O(\|N\|^2),$$

so the pair problem is solved by $\Xi_{ba} = \sum_k \delta_k N_{k,ab}^{\mathrm{w}}/\|\delta\|^2$ with $N_{k,ab}^{\mathrm{w}} = N_{k,ab} - \frac{1}{2}(\mathbf{p}_a(k) + \mathbf{p}_b(k))\bar{N}_{ab}$, as in the complete formula of Miettinen et al. (2016, Thm. 2). The a-th unmixing row is the a-th column of $(I - \frac{1}{2}\bar{N})V$, whose coefficient along

the b-th column of A is $\Xi_{ba} - \frac{1}{2}\bar{N}_{ab} + O(\|N\|^2)$. Since $\sum_k \delta_k \frac{1}{2}(\mathbf{p}_a(k) + \mathbf{p}_b(k)) = \frac{1}{2}(\|\mathbf{p}_a\|^2 - \|\mathbf{p}_b\|^2)$ and $\bar{N} = \frac{1}{K}\sum_k N_k$, a short computation gives

$$\Xi_{ba} - \tfrac{1}{2}\bar{N}_{ab} = \frac{1}{\|\delta\|^2}\sum_k \Big(\delta_k - \frac{1}{K}\sum_{k'} \mathbf{p}_a(k')\delta_{k'}\Big) N_{k,ab},$$

i.e. the weights δ_k are replaced by $\delta_k - \bar{\delta}_a$ with $\bar{\delta}_a$ the $\mathbf{p}_a$-weighted mean of δ. Its variance is (24) (with $j = a$, $i = b$). The loss of the whitened-space rotation V against Q, and of the estimated mixing columns $(I + \frac{1}{2}\bar{N})V$ against A, have the same form with $\bar{\delta}_a$ replaced by the $\frac{1}{2}(\mathbf{p}_a + \mathbf{p}_b)$- and $\mathbf{p}_b$-weighted means. The two-sided comparison: $v^{\mathrm{w}}_{ij} \geq S^{-1}_{ij}$ by Cauchy–Schwarz, $\|\delta\|^4 = \big(\sum_k (\delta_k - \bar{\delta}_j)\delta_k\big)^2 \leq \sum_k (\delta_k - \bar{\delta}_j)^2 \mathbf{p}_i\mathbf{p}_j \cdot \sum_k \delta_k^2/(\mathbf{p}_i\mathbf{p}_j)$, where $\sum_k \delta_k = 0$ by the normalisation; and $\sum_k (\delta_k - \bar{\delta}_j)^2 \mathbf{p}_i\mathbf{p}_j \leq 2\sum_k \delta_k^2 \mathbf{p}_i\mathbf{p}_j + 2\bar{\delta}_j^2 \sum_k \mathbf{p}_i\mathbf{p}_j \leq (2 + 2B^6)\sum_k \delta_k^2 \mathbf{p}_i\mathbf{p}_j$, using $\bar{\delta}_j^2 \leq B^2\|\delta\|^2/K$, $\sum_k \mathbf{p}_i\mathbf{p}_j \leq KB^2$ and $\sum_k \delta_k^2 \mathbf{p}_i\mathbf{p}_j \geq \|\delta\|^2/B^2$; hence $v^{\mathrm{w}}_{ij} \leq (2 + 2B^6) v^{\mathrm{F}}_{ij} \leq C_B S^{-1}_{ij}$ with $C_B = 2B^4(1 + B^6)$. The constants can differ a lot: for $K = 2$, $\mathbf{p}_i = (1.90, 0.10)$ and $\mathbf{p}_j = (1.89, 0.11)$ the coefficients are $v^{\mathrm{F}}_{ij} = 18010$, $v^{\mathrm{w}}_{ij} = 414$ for the row of source j and $v^{\mathrm{w}}_{ji} = 378$ for that of source i, all above the information bound $S^{-1}_{ij} = 219$; simulations with $n = 2 \cdot 10^5$ give $n\,\mathbb{E}\rho = 421 \pm 11$ and 385 ± 10. For the random signatures of Section 6 the two formulas differ by a few percent.

(ii) With weights $1/(\mathbf{p}_a(k)\mathbf{p}_b(k))$ the pair problem becomes $\min \sum_k (N_{k,ab} + \Xi_{ab}\delta_k)^2/(\mathbf{p}_a\mathbf{p}_b)$, whose solution has variance $K/(n\sum_k \delta_k^2/(\mathbf{p}_a\mathbf{p}_b)) = 1/(nS_{ab})$; consistently estimated weights change this by $o(1/n)$.

H. Below the threshold

Proposition 3 (Accuracy at the nearest-neighbour scale)**.** *Let c be the constant of Theorem 1. If an estimator satisfies $\sup_Q \mathbb{E}\,\rho_j(\hat{Q}) \leq \eta\, S_{\mathrm{nn}}(j)$ with $\eta S_{\mathrm{nn}}(j) \leq c/2$, then $n \geq c\,E_j/(2\eta S_{\mathrm{nn}}(j))$. If $\sup_Q \mathbb{E}\,\ell^2_{\mathrm{M}}(\hat{Q}) \leq \eta\, S_{\min} \leq c/2$, then $n \geq c\,E_{\max}/(2\eta S_{\min})$.*

Indeed, (4) gives $cE_j/(n + E_j) \leq \eta S_{\mathrm{nn}}(j)$, so $n \geq E_j(c/(\eta S_{\mathrm{nn}}(j)) - 1) \geq cE_j/(2\eta S_{\mathrm{nn}}(j))$, and (8) gives the second claim. □

The threshold of Theorem 2. Steps 1–4 of Appendix C need $\nu/\lambda_* \leq c$, with $\nu \asymp K^2(d + \log n)/n$ and $\lambda_* \asymp K S_{\mathrm{nn}}(j)$. When $E_j \asymp d$, as for well-spread signatures with $K \geq 4$ and for separated ones, this says that the first-order error E_j/n is below a fraction of $S_{\mathrm{nn}}(j)/K$, and by Proposition 3 no estimator is that accurate unless $n \gtrsim KE_j/S_{\mathrm{nn}}(j)$. This suggests, but does not prove, that arguments comparing the noise with the nearest-neighbour gap in operator norm cannot go further. Below the threshold the worst-case error of a source is $\gg S_{\mathrm{nn}}(j)/K$ (Proposition 3) but spread over $d - 1$ directions, each pair contributing about $1/(nS_{ij})$; we expect that a proof there must average these contributions like independent noise, as a local law for the joint-diagonalization resolvent would.

The factor K in γ_j. Part of it comes from the operator-norm bounds of Step 1, which use Cauchy–Schwarz over environments and could be sharpened, but not all of it. By Step 5 the error vector $g = H'^{-1}h$ of the estimator contains $H'^{-1}h^{(2)}$, the product of the two halves' noise, uncorrelated with $H'^{-1}h^{(1)}$ and with second moment of exact order $KdE^{(4)}_j/n^2$ (up to factors of B): each environment's halves have $n/(2K)$ observations, d terms add up in each entry, and so do the K environments. For separated signatures $E^{(4)}_j \asymp d$ and $E_j \asymp d$, so this term is $(Kd/n)\,E_j/n$, and $\mathbb{E}\|g\|^2$ stays within a constant of E_j/n only once $n \gtrsim Kd$; for $K \asymp \log d$ this is the $d \log d$ of Section 4. A simulation agrees (balanced codes, $d = 32$, known whitening, 150 draws): for $K = 8$, 16, 32, $n\,\mathbb{E}\ell^2_{\mathrm{A}}/\bar{E}^{\mathrm{qp}}$ of the labelled vectors is $1.16, 1.32, 1.67$ at $n = 200d$ and $1.03, 1.07, 1.14$ at $n = 800d$, against $1 + n\,\mathbb{E}\|H'^{-1}h^{(2)}\|^2/\bar{E}^{\mathrm{qp}}$ (averaged over sources) $= 1.17, 1.31, 1.59$ and $1.04, 1.08, 1.15$, and it is about the same at equal Kd/n (1.16 for $K = 8$ at $n = 200d$, 1.14 for $K = 32$ at $n = 800d$). We therefore expect this estimator to need the factor K; reaching $n \asymp d$ with $\log d$ separated environments would need another estimator, for instance one that does not multiply the per-environment covariances.

A second-order drift of the quadratic pencil. Under $\gamma_j \leq c$ no second-order term of the proof of Theorem 2 matters; below the threshold one of them has a deterministic part. In the basis of Q at the exact probe $\mathbf{p}_j$, let i be a near source and F a set of far sources, $\lambda_b \gg \lambda_i$ for $b \in F$. Eliminating F by a Schur complement, the expected diagonal entries of the effective matrix on $\{i, j\}$ are, to second order in the noise, $\lambda_a - \sum_{b\in F} \mathbb{E}(P_1)^2_{ab}/\lambda_b$, since $\mathbb{E}P_1 = \mathbb{E}P_2 = 0$. With

$\mathbb{E}(P_1)^2_{ab} = \frac{K}{n}\sum_k w_k^2(\delta_a(k)+\delta_b(k))^2\mathbf{p}_a(k)\mathbf{p}_b(k)$ and $\delta_j = 0$, expanding to first order in δ_i gives the effective gap

$$\lambda_i - \frac{K}{n}\langle g_j, \delta_i\rangle, \qquad g_j(k) = \sum_{b\in F} \frac{w_k^2\,\mathbf{p}_b(k)\big(\mathbf{p}_b(k)^2 - \mathbf{p}_j(k)^2\big)}{\lambda_b}. \tag{25}$$

The far sources repel the two levels by different amounts. Moving the probe to $\mathbf{p}_j + \eta_j$ changes the population gap by $-2\sum_k w_k\delta_i(k)\eta_j(k)$, so $\eta_j(k) = -\frac{K}{2n}g_j(k)/w_k$ undoes the shift at first order. Since $|\mathbf{p}_b^2 - \mathbf{p}_j^2| \le 2B|\delta_b|$ and $\lambda_b \ge B^{-2}\|\delta_b\|^2$, $|g_j(k)| \le 2B^8\sum_b\|\delta_b\|^{-1}$, of order d for well-spread signatures with $K \ge 3$, so $\|\eta_j\| = O(K^{3/2}d/n)$. As $|\langle g_j, \delta_i\rangle| \le \sqrt{K}\max_k|g_j(k)|\,\|\delta_i\|$, the shift is $O(K^{3/2}d\,r_j/n)$ against a gap $\lambda_i \asymp r_j^2$, so the gap can close unless $n \gtrsim K^{3/2}d/r_j$, which is $d^{K/(K-1)}$ for well-spread signatures and fixed K. Under $\gamma_j \le c$ the relative shift is $O(r_j)$. The expansion (25) is a second-order computation, not a theorem.

Two simulations confirm it: $K = 5$, $d = 300$, one near source at $\mathbf{p}_j \pm 0.05\,z$ for a unit vector z in the hyperplane, and $d-2$ sources around $\mathbf{p}_j + Rz$; the effective gap is measured on the Schur complement of the far block. With $\mathbf{p}_j = \mathbf{1}$ ($w_k = 1$), $R = 0.5$, $n = 1000d$ and 400 draws, the population gap is 0.0025 for both signs, the predicted gaps are 0.0014 and 0.0036, the measured ones 0.0013 and 0.0036 (standard errors at most 0.0001), with $\mathbb{E}\rho_j = 0.31$ and 0.10; at the corrected probe (weights unchanged) they are 0.0024 and 0.0026, with $\mathbb{E}\rho_j = 0.18$ and 0.17. A tilted $\mathbf{p}_j \propto (1.3, 0.8, 1.1, 0.9, 0.9)$ ($R = 0.4$, 300 draws, population gap 0.00279) gives the same picture: at $n = 1000d$ the measured gaps are 0.00127 and 0.00411 against predicted 0.00147 and 0.00411, and the corrected ones 0.00241 and 0.00279 (2.4 standard errors apart), and at $n = 4000d$ the corrected ones are 0.00276 and 0.00278 (standard errors at most 0.00003): the correction removes the asymmetry at first order.

Towards the linear regime. A proof of Conjecture 1, or of its analogue for the quadratic pencil at the corrected probe, would need, we expect, a multiscale replacement of Step 1 that compares the noise with the gap shell by shell, $\{i : \|\mathbf{p}_i - \mathbf{p}_j\| \asymp 2^{-l}\}$; a shell at distance r holds about $d\,r^{K-1}$ sources and contributes about $d\,r^{K-3}$ to the energy. In simulations below the proved threshold, Jacobi sweeps follow the first-order prediction wherever $\bar{E}^{\mathrm{w}}/n \lesssim 0.1$ (Section 6, Figure 7), and below $100d$ they are stable under initialisation (Appendix J, runs below $n = 100d$).

I. Further related work

The Gaussian model with a common Q and diagonal Λ_k is the common-principal-components model of Flury (1984), whose maximum-likelihood asymptotics (Flury, 1986) give the information S_{ij}. Fixed-dimension theory of second-order BSS goes back to Cardoso (1994); Belouchrani et al. (1997); Pham & Cardoso (2001), with optimal weightings and Cramér–Rao bounds (Yeredor, 2000; Tichavský & Yeredor, 2009; Yeredor, 2010), sensitivity analyses (Afsari, 2007; 2008) and complete variance formulas for SOBI and AMUSE (Miettinen et al., 2012; 2016); see Nordhausen & Ruiz-Gazen (2022). Pham's algorithm (Pham, 2001) computes the Gaussian likelihood estimator and AC-DC (Yeredor, 2002) and FFDIAG (Ziehe et al., 2004) are non-orthogonal least-squares methods. Kuleshov et al. (2015) bound joint diagonalization by Euclidean signature gaps and remove the smallest gap by random projections, without counting the dimension of the configuration or proving lower bounds; He & Kressner (2024a;b) reduce it to a random pencil plus refinement; their guarantees bound the joint-diagonalization residual, and when signatures are close a small residual does not imply a small direction error, which is what Theorem 5 quantifies at first order. Random combinations of slices are also the first step of simultaneous diagonalization for tensors; smoothed and Fourier analyses lower-bound the eigenvalue gaps of such a combination in probability (Bhaskara et al., 2014; Goyal et al., 2014), while (22) bounds the first-order risk of its eigenvectors for every combination. In growing dimension Chang et al. (2018) analyse a one-matrix reduction; we found no sample-size results for full-rank common principal components. Step 1 of Appendix C bounds the noise relative to H_j, so that it must beat the nearest-neighbour gap rather than the spread of the spectrum, in the spirit of relative perturbation bounds for eigenvalues and eigenvectors (Ipsen, 1998) and for empirical covariance operators (Jirak & Wahl, 2018); Davis–Kahan bounds (Davis & Kahan, 1970; Yu et al., 2015) compare the noise in operator norm with the eigengap. For a single covariance matrix, Wahl (2022; 2023) proves non-asymptotic lower bounds for principal subspaces of the same form as Theorem 1, a sum over pairs of inverse Fisher informations $\lambda_i\lambda_j/(\lambda_i - \lambda_j)^2$ each truncated at a constant, from a van Trees inequality for orthogonally equivariant experiments and Haar priors; the general inequality of the second paper covers the Haar symmetrisation of Appendix B. Three things differ here. The K covariances share their eigenvectors and the information is averaged over environments, so the quantity is the signature divergence S_{ij} rather than an eigengap of one matrix; the sources have no natural ordering, so the loss matches each source to the closest estimated direction, and a bound for one specified

direction does not transfer to it (Appendix B uses the orthogonality of the other columns to control the matching); and ordering the hypercube by S_{ji} removes the truncation, so that (4) holds for every n. The matrix $\sum_k (X_k - \lambda_k I)^2$ of Theorem 2 is the quadratic composite operator of numerical linear algebra and mathematical physics: Cerjan et al. (2023) use its smallest singular value, as a function of the probe λ, to locate approximate joint eigenvectors of non-commuting observables (the quadratic pseudospectrum), and Loring & Vides (2022) take its bottom eigenvectors before a Jacobi joint diagonalization. Both are deterministic; Theorems 2 and 3 add the statistics, a sample-split form that removes the bias of the squared noise, and the grid of probes. Landscape and convergence are studied by Usevich et al. (2020); Troedsson et al. (2024). For the non-Gaussian route see Auddy & Yuan (2025); Ricci et al. (2025); Genc et al. (2026) and hardness results (Diakonikolas et al., 2017; 2024; Damian et al., 2024). Identifiability from several environments is qualitative in nonlinear ICA (Hyvärinen & Morioka, 2016; Hyvärinen et al., 2019; Khemakhem et al., 2020; Hyvärinen et al., 2023), also when the segmentation into environments is itself learned (Hälvä & Hyvärinen, 2020), and in causal representation learning, where environments are interventions on a latent causal model, linearly mixed (Squires et al., 2023; Jin & Syrgkanis, 2024) or not (Buchholz et al., 2023; von Kügelgen et al., 2023); our results give sample sizes for independent latents and a common linear mixing. Jiang (2026) bounds nonlinear ICA by an environment-level diversity, the smallest KL divergence between two environments, which does not resolve which pairs of sources are hard. Acartürk et al. (2024) give finite-sample upper bounds for interventional causal representation learning, without lower bounds, and Lee et al. (2026) show that $K \asymp \log d$ unknown multi-node interventions forming a separating system identify a linear causal model; our codewords play an analogous role, under different assumptions, for the sample size of second-order separation. Rathod et al. (2026) observe eigenvalue crowding in a pencil-based causal discovery method, which Theorem 5 may explain. In multi-task linear representation learning the benefit of many tasks is governed by a task-diversity condition number (Tripuraneni et al., 2021); here the analogous quantity is the dimension of the signature configuration, not the number of environments.

J. Experimental details

The supplementary material contains the code, the raw results and three tables: the runs below $n = 100d$ (means per setting; Figure 4(d) plots the individual runs), the mean risk on fixed lattices (Figure 7(a), and $n = 50d$) and the values of Figure 6(c). Its README maps every figure, table and number of this appendix to the script that produced it.

Algorithms. Algorithm 3 is implemented in C and agrees with a NumPy implementation to machine precision. Each sweep visits the pairs in lexicographic order with the Cardoso–Souloumiac angle, and a run stops when the largest $|\sin\theta|$ of a sweep is below 10^{-10}, or after 100 sweeps; in the sample-complexity runs with $n \geq 100d$, 662 of 684 converged, with a median of 13 sweeps. *Pham's algorithm* (Pham, 2001) minimises $\sum_k \pi_k[\log\det \operatorname{diag}(T\hat\Sigma_k T^\top) - \log\det(T\hat\Sigma_k T^\top)]$ over unmixing matrices T, the Gaussian likelihood of the non-orthogonal model, by pairwise updates (the update of the reference implementation in `pyriemann`, implemented in C; raw covariances, at most 300 sweeps, rows of T normalised). *FFDIAG* (Ziehe et al., 2004) updates $V \leftarrow (I + W)V$ with W solving the 2×2 systems of the Frobenius criterion on the whitened covariances, step capped at $\|W\|_F \leq 0.9$. The *Fisher step* is the weighted least-squares step defined under "Attainment" below, applied to the covariances with estimated whitening. Appendix G gives the predictions for unequal sizes and for Fisher weights with estimated whitening, and the Cramér–Rao coefficient v^{PC} of the non-orthogonal problem.

Algorithm 3 Orthogonal joint diagonalization (Jacobi sweeps)

whiten by the pooled covariance; form $\hat\Sigma_1, \dots, \hat\Sigma_K$; $V \leftarrow I_d$
repeat
  for each pair $a < b$ **do**
    $g_k \leftarrow \big((\hat\Sigma_k)_{aa} - (\hat\Sigma_k)_{bb},\ 2(\hat\Sigma_k)_{ab}\big)$ for $k \leq K$
    $G \leftarrow \sum_k g_k g_k^\top$
    $\theta \leftarrow \frac{1}{4}\operatorname{atan2}\big(2G_{12},\ G_{11} - G_{22}\big)$ ($\theta = 0$ iff $G_{12} = 0 \leq G_{11} - G_{22}$)
    rotate rows/columns (a, b) of every $\hat\Sigma_k$ and columns (a, b) of V by θ
  end for
until all angles are below tolerance

Simulation and draws. For Gaussian sources the environment covariances are sufficient statistics: $n_k\hat\Sigma_k = A\Lambda_k^{1/2} G G^\top \Lambda_k^{1/2} A^\top$ with G a $d \times n_k$ standard Gaussian matrix, and $GG^\top$ is drawn through the Bartlett decomposition when $n_k \geq d$ (exact in law, $O(d^3)$) and formed from G otherwise. The halves of Theorem 2 are two independent such

Table 2. Baselines on identical data, $n = 400d$ (means over draws, $\times 10^3$); the estimator of Theorem 2 uses the known signatures and whitening, and the column gives the loss of its labelled vectors (Table 5 gives their polar factor).

K	d	Jacobi	FFDIAG	+step	Pham	Thm 2	pencil	$\bar{E}^{\mathrm{w}}/n$	$\bar{E}^{\mathrm{PC}}/n$	$\bar{E}/n$
8	16	21.0	21.0	20.1	20.9	34.6	471	19.8	18.2	17.6
8	32	18.4	18.5	16.7	16.6	32.6	545	18.1	16.7	16.0
8	64	20.3	20.3	18.4	18.3	32.1	666	19.8	18.2	17.5
8	128	20.1	20.2	18.5	18.4	32.0	779	19.5	17.9	17.3
8	256	20.6	20.6	18.9	18.8	33.2	842	20.4	18.8	18.1
4	16	37.4	37.5	36.4	36.7	72.5	366	34.1	31.9	31.0
4	32	61.4	61.5	56.3	52.6	115.7	483	42.1	39.2	38.0
4	64	46.1	46.2	41.8	40.6	128.8	579	40.3	37.0	35.9
4	128	51.3	51.5	47.3	43.6	170.9	677	43.9	40.4	39.2
4	256	49.9	50.0	45.8	44.6	214.7	764	46.0	42.1	40.9

draws with $n_k/2$ samples; the other methods use their average, the full-sample covariance. Each draw uses new signatures (when random), a Haar-distributed A and new data. Curves show medians and quartiles over draws; Figures 5(c,d) and 7 show means over draws of A and the data for fixed signature configurations. Values of ℓ_{A}^2 above $1/2$ (for example in Figures 2(a), 4(d), 7 and 8) are lower bounds for a permutation-matched loss (Section 2).

Signatures. *Random*: i.i.d. uniform on $\{\mathbf{p} \in [\frac{1}{2}, \frac{3}{2}]^K : \sum_k \mathbf{p}(k) = K\}$, drawn by sampling the first $K-1$ coordinates uniformly and accepting if the last one lies in $[\frac{1}{2}, \frac{3}{2}]$ (a linear map, so the accepted law is uniform; $B = 2$). *Lattice*: $\mathbf{p}(k) = 1 + \epsilon u_k$ ($k < K$), u on a regular grid of m^{K-1} points in $(-\frac{1}{2}, \frac{1}{2})^{K-1}$, and $\mathbf{p}(K) = 1 - \epsilon \sum_{k<K} u_k$. With $\epsilon = \frac{1}{2}$, used for $K \le 4$, every variance lies in $(\frac{1}{4}, \frac{7}{4})$, so $B < 4$ and Definition 1 holds with fixed constants; for $K = 5$, where $\mathbf{p}(K)$ would approach 0, Figure 1(c) uses $\epsilon = \frac{1}{4}$ (variances in $(\frac{1}{2}, \frac{3}{2})$), and the $K = 5$ lattices of Figure 6(a), with $\epsilon = \frac{1}{2}$, have $B = 2$ and 3 at $d = 16$ and 81. *Halton*: $\mathbf{p}_j = \mathbf{1} + sVz_j$ with z_j the first d points of the Halton sequence in $[-\frac{1}{2}, \frac{1}{2}]^{D'}$, V a random orthonormal $K \times D'$ basis orthogonal to $\mathbf{1}$, and a fixed scale s that keeps every coordinate in $[\frac{1}{2}, \frac{3}{2}]$; the separation $\min_{i \neq j} \|z_i - z_j\|$ and the covering radius in the cube, multiplied by $d^{1/D'}$, stay in $[0.22, 0.81]$ and $[0.99, 1.91]$ for $d = 8$ to 256, so the configurations are well spread in the range simulated. *Codes*: $\mathbf{p}_j = \mathbf{1} + \frac{1}{2}y_j$ with balanced $y_j \in \{\pm 1\}^{16}$ at pairwise Hamming distance at least 4, built greedily from random balanced words ($S_{\min} \ge 0.33$). *Equal gap, different energy* (Figure 2(b)): configuration A is the $K = 4$ lattice with $d = 27$; configuration B takes 26 points of a coarser $K = 4$ lattice ($\epsilon = 1$, without the point whose last variance is 0; variances in $[\frac{1}{3}, 2]$, $B = 3$) and adds one point at Fisher divergence $S_{\min}(A)$ from one of them.

Figure details. Figure 1(b): medians and quartiles over sixteen draws (eight at $d = 128$); the well-spread case, lattices, is Figure 6(a). Figure 1(c): see "Signatures". Figure 1(d): known signatures, whitening by the pooled covariance of the same data, polar factor; six draws per d. Figure 2(a): medians over eight draws, four for $d \ge 128$; (b): twelve draws, with standard errors; (c): see Figure 8; (d): medians and quartiles over draws.

Lags and runs below $n = 100d$. The lag-signature of an MA(8) source $s_t = \sum_{i=0}^{8} b_i \epsilon_{t-i}$ ($b \sim N(0, I_9)$, normalised to unit variance) is $(\sum_i b_i b_{i+\tau})_{\tau \le L}$; AMUSE takes the eigenvectors of the whitened, symmetrised lag-1 covariance, and SOBI applies Algorithm 3 to the whitened, symmetrised covariances at lags $1, \dots, L$. Lag-signatures have no sum constraint, so L lags give a configuration of dimension up to L, and the first-order variance of SOBI (Miettinen et al., 2016, Thm. 2) has their Euclidean gap in its denominator: the dimension count suggests $n \gtrsim d^2$ for AMUSE and $n \gtrsim d$ for SOBI with $L \ge 3$ lags whose lag-signatures are well spread, a heuristic outside our theorems, since lags are not independent environments. At $n = 200d$, AMUSE degrades from 0.23 to 0.50 as d goes from 16 to 128, SOBI with two lags from 0.047 to 0.074, and SOBI with $L \ge 3$ lags stays between 0.007 and 0.020 (Figure 4(c)). Below $n = 100d$ (Figure 4(d); 54 runs, $n/d \in \{12, 25, 50\}$, 3 draws per setting), where all 54 runs stop at the 100-sweep cap, we keep the data and the whitening, continue Jacobi to 1000 sweeps, restart it from four Haar-random rotations (at most 300 sweeps each) and run Pham's algorithm. Continuing changes the error by at most 3% (median 0%), the restarts end within 1% of each other (median; at most 14%), their criterion values are within 0.19% of the continued run, and 37 of the 54 continued runs converge: Jacobi is stable under initialisation, which does not show that the landscape is benign. Pham's algorithm is 28%–34% more accurate at $K = 8$, $n = 50d$, 5%–13% in the other settings with $n \ge 25d$, and about as accurate at $n = 12d$, where $\bar{E}^{\mathrm{w}}/n$ is about 0.93 (the first-order prediction is not meaningful above about 0.3).

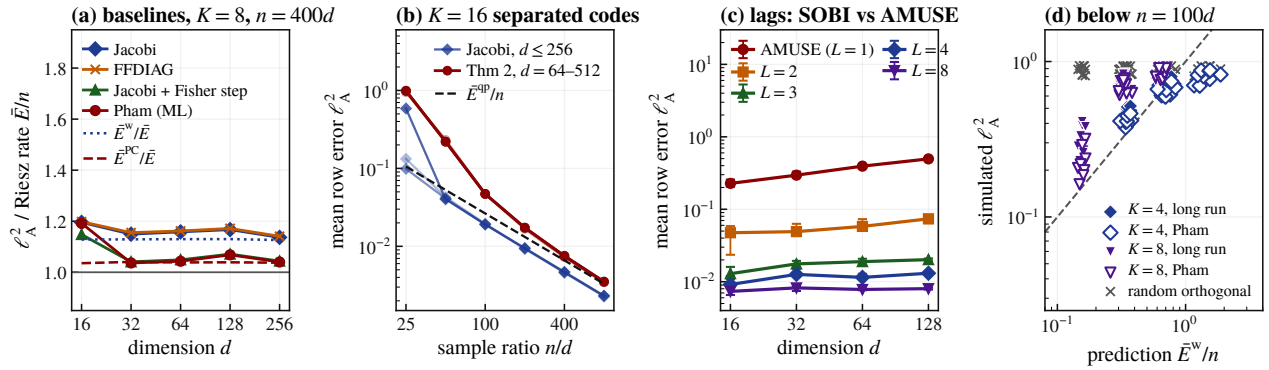

Figure 4. **Algorithms, attainment and the lag route. (a)** $K = 8$, $n = 400d$, identical data for all methods (medians over three to six draws), error divided by the Riesz rate $\bar{E}/n$ (solid line at 1); dashed: $\bar{E}^{\mathrm{PC}}/\bar{E}$, the Cramér–Rao coefficient with estimated whitening; dotted: $\bar{E}^{\mathrm{w}}/\bar{E}$. **(b)** Separated signatures ($K = 16$ balanced codes, $S_{\min} \ge 0.33$): the estimator of Theorem 2 with known whitening (red) and Jacobi with estimated whitening (blue, $d \le 256$) against n/d for $d = 64$ to 512 (light to dark); dashed: $\bar{E}^{\mathrm{qp}}/n$. **(c)** SOBI with L lags on MA(8) sources, $n = 200d$ (medians and quartiles); AMUSE is $L = 1$. **(d)** Runs below $n = 100d$ ($K \in \{4, 8\}$, $d \in \{32, 64, 128\}$): Jacobi continued to 1000 sweeps (filled) and Pham's algorithm (open) against $\bar{E}^{\mathrm{w}}/n$; crosses: a random orthogonal matrix.

The linear pencil as an initialisation. Theorem 5 concerns the eigenvectors of one combination as an estimate. As a starting point for Algorithm 3 they cost nothing ($K = 8$ random signatures, $n = 400d$, estimated whitening, $d \in \{16, 32, 64, 128, 256\}$, six draws each). The eigenvectors of one random combination $\sum_k \alpha_k \hat{\Sigma}_k$, $\alpha \sim N(0, I_K)$, have median error 0.43 at $d = 16$ and 0.84 at $d = 256$. One Jacobi sweep from them brings it to 0.066 and 0.39, two sweeps to 0.024 and 0.078, and three sweeps to within 13% of the converged run from the identity (0.020–0.024) at every d. Run to convergence, the two starting points give the same minimiser: the errors differ by less than 10^{-10} in relative terms and the squared sines between matched columns are below 10^{-13} (all runs converged, within 17 sweeps). Starting from the eigenvectors V_0 is the same as running the sweeps from the identity on the rotated covariances $V_0^{\top} \hat{\Sigma}_k V_0$.

Robustness. *Strong contrast and unequal sizes* (Figure 6(c), means over six draws): for $B = 10$ the random signatures are uniform on the slice of $[0.1, 1.9]^K$ (ratios bounded by 10 up to the renormalisation for unequal sizes); unequal sizes have π_k geometric from 1 to 10, normalised; Jacobi uses the π-weighted criterion and pooled covariance and is compared with the π-weighted $\bar{E}^{\mathrm{w}}/n$, and Pham's algorithm uses the weights π_k. Over the $B = 10$ runs and the unequal-size runs, the largest deviations of Jacobi from $\bar{E}^{\mathrm{w}}/n$ (28% at $B = 10$, 41% at $B = 2$) are both for $K = 4$ with unequal sizes at $n = 400d$, where the smallest environment has only $22d$ samples; these two mean errors have relative standard errors 14% and 18%. The equal-size $B = 2$ point far below the diagonal ($K = 4$, $d = 16$) has $\bar{E}^{\mathrm{w}}/n \approx 0.24$, outside the first-order regime. From $n = 200d$ on, Laplace and uniform sources give the error of Gaussian ones (Figure 5(c)). *Departures from the model* (Table 3; $K = 8$ random signatures, $d = 64$, $n/d \in \{200, 800\}$, eight draws of signatures, mixing and data per cell): with AR(1) sources of coefficient ϕ in every environment, the variance of an off-diagonal sample covariance, and with it the prediction, is multiplied by $\tau = (1 + \phi^2)/(1 - \phi^2)$ (Bartlett's formula; Bartlett, 1935); at $\phi = 0.5$ the error is 1.02–1.17 times $\tau \bar{E}^{\mathrm{w}}/n$, at $\phi = 0.8$ ($\tau = 4.6$) 1.11 times at $n = 800d$ and 2.65 times at $n = 200d$, where the effective sample size is $43d$ and zero of the eight runs converge. Mixings $A(I + \epsilon G_k/\sqrt{d})$ that differ across environments (G_k Gaussian) give at most 1.22 times the prediction for $\epsilon \le 0.03$ and leave a floor of about $22 \cdot 10^{-3} \approx 2.2\,\epsilon^2$ at $\epsilon = 0.1$ (fitted from the two sample sizes as $a/n + f$); a shared factor with correlation r between all sources leaves a floor of $4 \cdot 10^{-3}$ at $r = 0.05$ and $6 \cdot 10^{-3}$ at $r = 0.1$. At $B \in \{2, 4, 10\}$ the ratio is 0.99–1.09. In every cell the predicted per-source errors E_j^{w}/n rank the observed ones (Spearman correlation 0.54–0.78, mean over draws): the diagnostic still says which source is hard when its constant is off.

How many environments. One simulated recording of $n = 800d$ samples, $d = 32$, has sources whose log-variances are random combinations of three low-frequency cosines (eight draws of trajectories, mixing and data); cutting it into $K \in \{2, 4, \ldots, 256\}$ equal segments gives K environments, with the average variance in a segment as the signature (Table 4). Two segments put the signatures on a line and fail ($\ell_{\mathrm{A}}^2 = 0.37$); from $K = 4$ the error follows $\bar{E}^{\mathrm{w}}(K)/n$ (ratio 1.07 at $K = 4$, within 1.04 for $8 \le K \le 64$), which stops decreasing once the segments resolve the dimension of the trajectories (participation ratio of the centred signatures about 4.8), and the error leaves the prediction only for short segments: 1.17

Table 3. The first-order prediction under departures from the model ($K = 8$, $d = 64$, eight draws per cell, means): mean ℓ_{A}^2 ($\times 10^3$), its ratio to the prediction (multiplied by τ for AR(1) sources), and the Spearman correlation between the predicted and the observed per-source errors.

	$n = 200d$			$n = 800d$		
departure	error	error/pred.	Spearman	error	error/pred.	Spearman
AR(1), ϕ =0	45.0	1.12	0.73	10.4	1.04	0.69
AR(1), ϕ =0.5	76.9	1.17	0.65	16.5	1.02	0.70
AR(1), ϕ =0.8	479.1	2.65	0.54	51.0	1.11	0.66
mixing, ϵ =0	43.3	1.09	0.71	10.5	1.02	0.73
mixing, ϵ =0.01	46.6	1.17	0.70	10.4	1.05	0.67
mixing, ϵ =0.03	45.5	1.16	0.61	12.1	1.22	0.68
mixing, ϵ =0.1	74.4	1.80	0.72	35.4	3.53	0.67
shared factor, r =0	43.6	1.09	0.69	9.9	1.01	0.67
shared factor, r =0.05	46.1	1.19	0.69	14.4	1.44	0.68
shared factor, r =0.1	48.9	1.22	0.72	16.6	1.72	0.74
contrast, B =2	43.2	1.09	0.65	10.5	1.02	0.70
contrast, B =4	18.0	1.06	0.72	4.4	1.01	0.72
contrast, B =10	12.3	1.04	0.74	2.9	0.99	0.78

Table 4. Segments of one nonstationary recording as environments ($d = 32$, $n = 800d$, eight draws): mean ℓ_{A}^2 ($\times 10^3$) with its standard error, the median first-order prediction $\bar{E}^{\mathrm{w}}(K)/n$ ($\times 10^3$; off scale for $K = 2$), the median per-draw ratio, the participation ratio of the centred signature matrix, and the runs that met the tolerance.

K	$n/(Kd)$	mean ℓ_{A}^2	s.e.	$\bar{E}^{\mathrm{w}}(K)/n$	median ratio	participation ratio	converged
2	400	371.1	22.2	–	0.01	1.0	8/8
4	200	24.6	3.0	22.0	1.07	2.7	8/8
8	100	6.7	0.4	6.5	0.99	4.8	8/8
16	50	5.8	0.3	5.7	1.00	5.0	8/8
32	25	5.7	0.3	5.5	1.01	5.1	8/8
64	12.5	5.9	0.3	5.5	1.04	5.1	8/8
128	6.25	6.6	0.5	5.5	1.17	5.1	8/8
256	3.125	9.7	2.2	5.5	1.34	5.1	8/8

times at $K = 128$ ($n/K \approx 6.2d$) and 1.34 times at $K = 256$ ($n/K \approx 3.1d$).

Mean risk on fixed configurations. The theorems bound expected risks, while most curves show medians, so we fix one configuration per d, the $K = 4$ lattice, and redraw only the mixing and the data (Figure 7(a); at least 6 draws per point; all values in the supplementary table). At $n = 400d$ the mean error is 1.08–1.13 times $\bar{E}^{\mathrm{w}}/n$ from $d = 27$ to 1000, with standard errors at most 3% of the mean and 90% quantiles within 16% of it, and at $n = 800d$ ($d = 64$ to 1000) it is 1.03–1.04 times. At $n = 200d$, where $\bar{E}^{\mathrm{w}}/n$ is 0.15–0.23 and most runs stop at 100 sweeps, the ratio falls from 1.47 to 1.22 as d grows, and the sweep cap does not explain the excess: continued to 1000 sweeps, the 90 capped runs at $n = 200d$ (all for $d \le 216$, the first four at $d = 343$) change their loss by at most 5%, the mean of the continued runs at each d by at most 0.1%, and 89 of them then meet the tolerance. At $n = 100d$ ($d = 64$ to 1000), where $\bar{E}^{\mathrm{w}}/n$ is 0.36–0.46, outside the first-order regime, the ratio is 1.67–1.78, and 1.77–1.78 for $d \ge 512$; at $n = 50d$ ($d = 64$ to 343) the error is 0.78–0.91, and there its ratio to the prediction, which is not meaningful above about 0.3, says nothing; at $n \le 100d$ zero of 146 runs meet the tolerance within 100 sweeps. The crossings of Figure 7(d) interpolate $\log$ of the mean linearly in $\log n$ between neighbouring sample sizes: from $d = 27$ to 1000 the n/d at which the mean error reaches 0.1 grows by 1.4 (Section 6) and the one for 0.2 by 1.2, against 1.5 for $\bar{E}^{\mathrm{w}}/d$ and 11.1 for the proved threshold $d^{5/3}$. A run at $d = 1000$ takes 6–21 minutes on one core. One fixed draw of $K = 4$ random signatures with $B = 2$ and $K = 16$ balanced codes with $S_{\min} \ge 0.33$ behave alike for $d = 64$ to 512 (Figure 7(b,c)): at $n = 400d$ their mean errors are 1.06–1.12 and 1.01 times $\bar{E}^{\mathrm{w}}/n$, and the n/d for a target error of 0.1 (0.01 for the codes) grows by 1.0 (1.0), against 1.1 (1.0) for $\bar{E}^{\mathrm{w}}/d$.

Attainment. With known whitening, the Fisher step after Algorithm 3 takes $M_k = \hat{V}^\top \hat{\Sigma}_k \hat{V}$, $\hat{\mathbf{p}}_a(k) = (M_k)_{aa}$ and $\hat{\delta}_k = \hat{\mathbf{p}}_a(k) - \hat{\mathbf{p}}_b(k)$, estimates the residual angle of each pair by weighted least squares, $\hat{\Xi}_{ab} = \sum_k w_k \hat{\delta}_k (M_k)_{ab} / \sum_k w_k \hat{\delta}_k^2$ with $w_k = 1/(\hat{\mathbf{p}}_a(k)\hat{\mathbf{p}}_b(k))$, and returns $\hat{V}\exp(-\hat{\Xi})$; by Proposition 2(ii) its first-order variance per pair is S_{ab}^{-1} (Figure 5(d)).

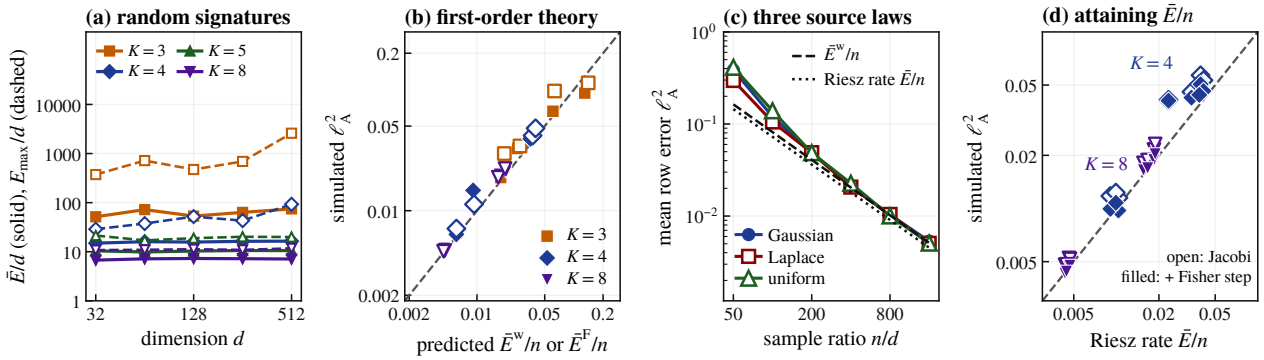

Figure 5. **Verification. (a)** Energy per source for random signatures (medians over eleven draws): the average $\bar{E}/d$ (solid) saturates from $K = 4$ on, the worst source $E_{\max}/d$ (dashed) only from $K = 5$ on (Proposition 1); $K = 2$ is off scale. **(b)** Simulated mean row error against the first-order prediction for $K \in \{3, 4, 8\}$, $d \in \{16, 32\}$ and $n \in \{400d, 1600d\}$, one signature draw and six mixings each: filled markers use estimated whitening and $\bar{E}^{\mathrm{w}}/n$, open markers known whitening and $\bar{E}^{\mathrm{F}}/n$; dashed: identity. Points off the line are draws with a pair outside the first-order regime. **(c)** Gaussian, Laplace and uniform sources ($d = 32$, $K = 8$, one signature draw, eight mixings each) give the same error from $n = 200d$ on, on top of $\bar{E}^{\mathrm{w}}/n$ (dashed; ratios 1.03, 1.01 and 1.11 at $n = 400d$) and above the Riesz rate $\bar{E}/n$ (dotted). **(d)** Attaining the Riesz rate (known whitening, $K \in \{4, 8\}$, $d \in \{16, \dots, 128\}$, $n \in \{400d, 1600d\}$): Algorithm 3 (open) sits at $\bar{E}^{\mathrm{F}}/n$, above the Riesz rate; one Fisher-weighted step (filled) lands on it.

Table 5. The estimator of Theorem 2 on random signatures, $n = 400d$, three draws per row (means, $\times 10^3$; γ: median): labelled vectors, their polar factor, and one Fisher step from it, with the first-order coefficient of the labelled vectors, the Riesz rate, and the ratio of the polar factor's loss to the Riesz rate.

K	d	γ	labelled	polar	+step	$\bar{E}^{\mathrm{qp}}/n$	$\bar{E}/n$	polar$/(\bar{E}/n)$
8	64	1.4	35.7	20.7	18.7	22.2	18.3	1.13
8	128	2.7	32.7	19.5	18.1	21.5	17.5	1.11
8	256	3.0	33.1	19.7	18.4	21.7	17.8	1.11
4	64	18.1	143.5	71.8	46.2	43.0	40.0	1.79
4	128	27.0	173.3	78.0	41.9	42.3	39.3	1.98
4	256	55.1	186.3	85.0	48.9	42.9	39.9	2.13

The estimator of Theorem 2 below its proved regime (Table 5; known signatures and whitening, random signatures, $n = 400d$, $\gamma = Kd/(nS_{\min})$ from 1.3 to 4 for $K = 8$ and up to 100 for $K = 4$ over the draws; Table 5 reports medians and means): over the draws, for $K = 8$ its labelled vectors are 1.82–2.13 times $\bar{E}/n$, their polar factor 1.09–1.18 times (it averages the two estimates of each rotation angle, from the pencils at $\mathbf{p}_i$ and $\mathbf{p}_j$), and one Fisher step from it 1.00–1.06 times; for $K = 4$, where random signatures have close pairs, the three are 3.13–4.84, 1.40–2.19 and 0.99–1.42 times $\bar{E}/n$.

Estimated whitening (Theorem 2(b)). Table 6 runs the estimator of Theorem 2 on $K = 4$ lattices with known signatures and the sample split into thirds: on parts 1 and 2 with the true whitening, and on the same parts whitened by the pooled sample covariance of part 0, scored by ρ_j of Theorem 2(b); the comparison is with known whitening on the same two thirds. Mixing the same sources by $I + G/\sqrt{d}$ (G Gaussian, $d = 27$, condition number 41) instead of $A = I$ changes no ρ_j by more than 10^{-13}, as the reduction in the proof says. On these lattices $\bar{E}^{\mathrm{wh}} = 0.011\,\bar{E}$, so estimating the whitening costs 2.1% at first order; the measured excess is 4%–6% at $n = 25600d$ ($\gamma \le 0.08$) and 10%–15% at $n = 6400d$ ($\gamma = 0.22$–0.30), and for $d = 27$ it falls to 2.9% and 2.3% at $n = 102400d$ and $409600d$: a second-order term, which the factor $(1 + C\sqrt{\gamma_j})^2$ of (11) absorbs. To resolve the coefficient $2E_j^{\mathrm{wh}}$ we take $K = 4$, $d = 8$, $n = 2 \cdot 10^6 d$ and 30,000 draws on a configuration chosen for a large E^{wh} (the largest $\bar{E}^{\mathrm{wh}}/\bar{E}$ among 4000 draws on $\{\mathbf{p} \in [\frac{1}{4}, \frac{7}{4}]^K : \sum_k \mathbf{p}(k) = K\}$ with $S_{\min} \ge 0.05$; E_j^{wh}/E_j up to 0.15, $\gamma = 3 \cdot 10^{-4}$): $n\mathbb{E}\rho_j/[\frac{3}{2}(E_j^{\mathrm{qp}} + 2E_j^{\mathrm{wh}})]$ is 0.99–1.00 for every source (standard errors at most 0.005), whereas E_j^{wh} in place of $2E_j^{\mathrm{wh}}$ would miss by up to 11%; on a random configuration as in Section 6 ($\bar{E}^{\mathrm{wh}} = 0.007\,\bar{E}$, $\gamma = 0.03$) the ratio is 1.00–1.02.

Unknown signatures. The grid of Algorithm 2 is a device for the proof; Table 7 tests a practical variant on $K = 4$ lattices with $d = m^3$, known whitening and the sample split into thirds. On part 0, $\lambda_{\min}\big(\sum_k (\hat{\Sigma}_k^{(0)} - g(k)I)^2\big)$ is computed on a

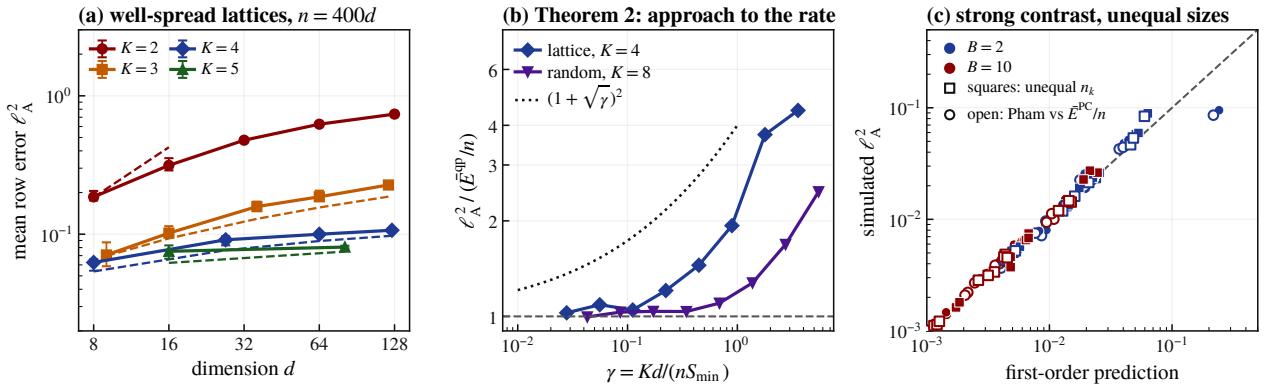

Figure 6. **(a)** Lattice signatures (Definition 1), $n = 400d$, eight draws: the dimension law with the assumptions of Theorems 1 and 4. **(b)** The estimator of Theorem 2 (known whitening) approaches its prediction $\bar{E}^{qp}/n$ as $\gamma = Kd/(nS_{min})$ decreases ($d = 64$, four draws per point; diamonds: $K = 4$ lattice, with ratio 1.03 to $\bar{E}^{qp}/n$ at $\gamma = 0.028$; triangles: $K = 8$ random signatures); dotted: the shape $(1 + \sqrt{\gamma})^2$ of (10). **(c)** Strong contrast ($B = 10$, red) and unequal sizes (squares): Jacobi against $\bar{E}^{w}/n$ (filled) and Pham against $\bar{E}^{PC}/n$ (open).

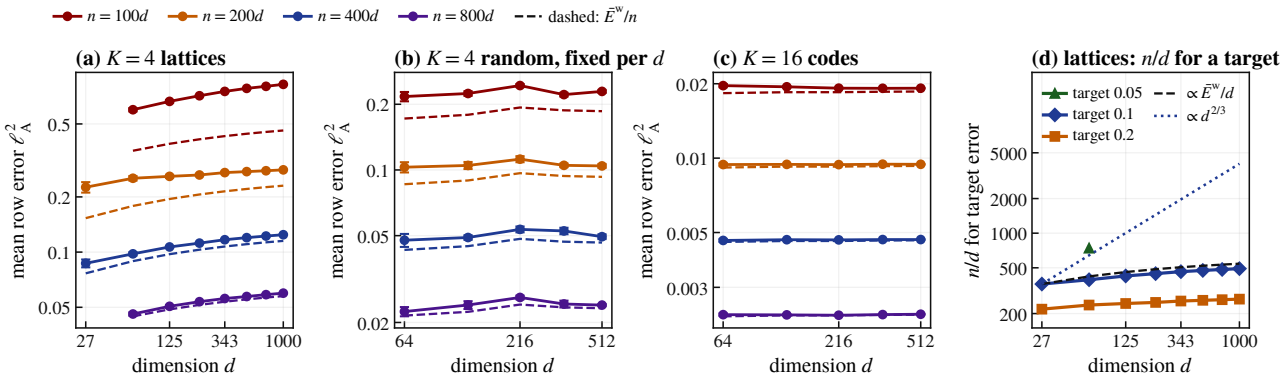

Figure 7. **Mean risk on fixed configurations, larger d.** Mean ℓ_A^2 with 95% intervals against d at $n/d \in \{100, 200, 400, 800\}$; dashed: $\bar{E}^{w}/n$. **(a)** $K = 4$ lattices, $d = 27$ to 1000. **(b)** One fixed draw of $K = 4$ random signatures per d. **(c)** $K = 16$ balanced codes. **(d)** n/d at which the mean error of the lattices reaches 0.1 and 0.2 (and 0.05, reached within $n \le 800d$ only at $d = 64$); dashed: the growth of $\bar{E}^{w}/d$; dotted: the growth $d^{2/3}$ of the proved threshold.

cubic grid of mesh $r_0/4$ in the hyperplane within $[1/2.6, 2.6]^K$ ($[1/2.7, 2.7]^K$ at $d = 216$, whose lattice has $B = 8/3$), r_0 the smallest distance between signatures; the grid points below all their face neighbours are refined by three alternating steps ($u \leftarrow$ bottom eigenvector, $g(k) \leftarrow u^\top \hat{\Sigma}_k^{(0)} u$), refined probes closer than $r_0/2$ are merged, and the d with the smallest λ_{min} are kept (a global threshold, as in the proof, needs larger n). The estimator of Theorem 2 then runs on parts 1 and 2 at each probe, and its polar factor is compared with the same estimator at the true signatures. From $n = 6400d$ on, all d probes are found in every run, the largest probe error is 0.20 of the nearest-neighbour distance, and the two estimators differ by at most 4%; the selection takes 4 s at $d = 27$ and about 9 minutes at $d = 216$ (210,464 grid points) on one core. At $n = 1600d$ the selection misses up to 10 of the d probes, and the loss over all runs, failures included (the estimator at every returned probe, orthonormalised, each source scored against its closest returned vector), is 1.8–2.6 times the oracle. These probes are far outside the accuracy that Lemma 3 assumes ($1/(8B^4)$ of the nearest-neighbour distance), so the lemma is conservative. *Misspecified* s_0. In the variant the mesh and the merge radius stand for s_0, which in Algorithm 2 also sets the acceptance threshold (the variant has none) and the sample size needed: replacing r_0 by cr_0 corresponds to $s_0 \approx c^2 S_{min}$ up to factors of B. On the same data ($d \in \{27, 64\}$, $n = 6400d$), $c = 1/2$ multiplies the number of grid points by 7.9 and the selection time by about 8 and returns the same probes as $c = 1$ in all six runs, while $c = 2$ merges neighbouring sources and no run succeeds: too small an s_0 costs time, too large an s_0 fails. A Jacobi pilot on an additional independent batch of $n/3$ samples (rotated by a Haar matrix, so that the sweeps do not start at the truth) estimates r_0 by the smallest distance $\hat{r}_0$ between its

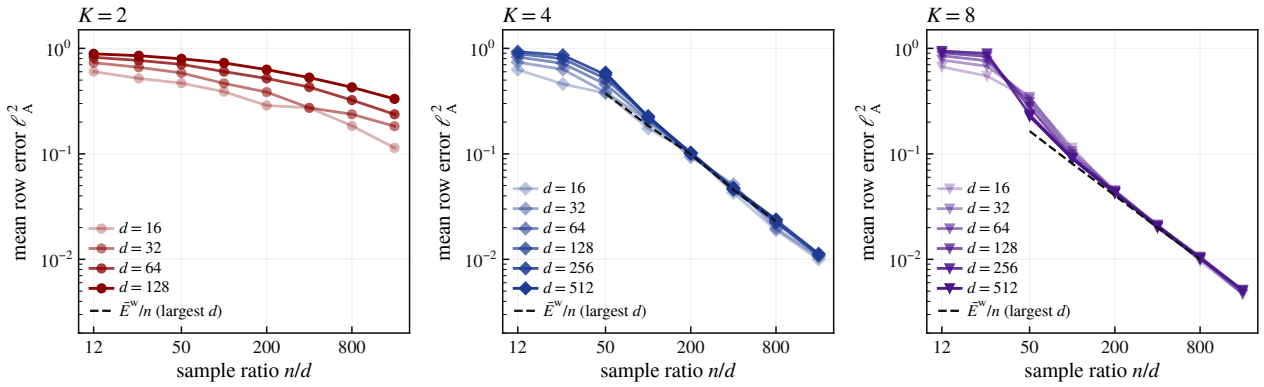

Figure 8. **Collapse in n/d.** Median error against n/d for random signatures, $K \in \{2, 4, 8\}$ and d from 16 (light) to 512 (dark; $K = 2$ up to 128), ten draws per point ($d \le 128$), six ($d = 256$) or three ($d = 512$); dashed: $\bar{E}^w/n$ for the largest d. For $K = 4$ and $K = 8$ the curves coincide; for $K = 2$ they move up with d.

Table 6. Theorem 2(b) in simulation: $K = 4$ lattices, known signatures, $n/3$ observations per part. The estimator of Theorem 2 on parts 1 and 2 with the true whitening (known; loss $1 - \langle \hat{q}_j, q_j \rangle^2$) and with the whitening estimated on part 0 (estimated; loss ρ_j of Theorem 2(b)): means of ℓ_A^2 over the stated number of draws ($\times 10^3$), their first-order predictions with $n_B = 2n/3$, and the ratio of the two means; $\gamma = \max_j \gamma_j$.

d	n/d	draws	γ	known	$\bar{E}^{qp}/n_B$	estimated	$(\bar{E}^{qp} + 2\bar{E}^{wh})/n_B$	estimated/known
27	6400	40	0.22	7.99	7.28	8.81	7.43	1.103
27	25600	40	0.060	1.87	1.82	1.95	1.86	1.042
27	102400	40	0.016	0.464	0.455	0.478	0.465	1.029
27	409600	40	0.0043	0.114	0.114	0.117	0.116	1.023
64	6400	16	0.25	9.41	8.39	10.8	8.57	1.146
64	25600	16	0.065	2.15	2.10	2.25	2.14	1.050
125	6400	8	0.30	10.6	9.10	12.3	9.30	1.154
125	25600	8	0.077	2.38	2.28	2.53	2.32	1.063

signatures: in six runs ($d \in \{27, 64\}$), $\hat{r}_0/r_0$ lies between 0.80 and 0.89 and $\hat{S}_{\min}/S_{\min}$ between 0.67 and 0.90 (biased low, the safe side, because the lattice has many pairs at distance r_0; with a unique closest pair the bias can be upward, which the safety factor must absorb). With $\hat{r}_0/2$ in place of r_0, a safety factor of four on s_0, the selection finds all probes in all six runs, and the estimator is within 3% of the oracle.

Real data. *Data.* HCP S1200 (Van Essen et al., 2013): the 813 subjects with all fourteen task runs (seven tasks, two runs each; Barch et al., 2013) and the first resting-state run, 799 of them with four resting-state runs. Grayordinate time series (MSMAll, from the minimal preprocessing pipelines; Glasser et al., 2013) were averaged within Schaefer parcels (Schaefer et al., 2018) at $d \in \{100, 200, 400\}$, converted to percent signal change, cleaned per run by regressing out a trend, twelve motion parameters and a discrete-cosine high-pass basis (0.008 Hz), and demeaned. *Environments*: the seven tasks and rest, 176 frames per task run, i.e. 352 per subject and environment, and a contiguous rest segment of the same length (286176 frames per environment in the group fit). *Group fit*: pooled whitening and Algorithm 3 (at most 1000 sweeps) on the covariances summed over subjects; the signatures are $\hat{\mathbf{p}}_j(k) = \hat{b}_j^\top \hat{\Sigma}_k \hat{b}_j / \hat{b}_j^\top \hat{C} \hat{b}_j$, normalised to mean one. Only the eight-task fit at $d = 100$ reached the tolerance; at $d = 200$ the pairs are nearly degenerate ($S_{\min} = 0.0000570$) and Pham's algorithm (Pham, 2001) differs from Jacobi by a matching loss of up to 0.97 for such pairs (the energies below are those of the Jacobi solution). *Autocorrelation*: $\tau_{ab} = \sum_t \rho_a(t)\rho_b(t)$ (Bartlett, 1935) from the estimated sources of 50 subjects (Bartlett window of $L = 30$ frames; median $\tau_{aa} = 1.19$; an AR(2) fit agrees to within 2.3% in median). *Split halves*: 50 random disjoint pairs of halves with $m \in \{10, \dots, 400\}$ subjects each (no family information was available; at $m = 400$ the two halves use 2×400 of the 813 subjects, so the repeats are far from independent), each fitted with at most 300 sweeps (none reached the tolerance), labelled against the group solution and matched by the Hungarian algorithm; $\delta_j = 1 - \mathrm{corr}^2$ between the two estimates of source j under the pooled covariance, against $\hat{\delta}_j = 2 \sum_{a \neq j} v_{aj}^w \tau_{aj} / n_{\mathrm{half}}$. The first-order region is $\hat{\delta}_j < 0.3$: $\hat{\delta}_j$ adds the errors of the two halves, so this is a per-half first-order error below 0.15, about the

Table 7. A practical variant of Algorithm 2 (not a test of Theorem 3): $K = 4$ lattices, three draws per row (two for $d = 216$). All probes: runs in which the selection returns one probe per source; probe error: largest distance of a probe from its signature over the signature's nearest-neighbour distance (median over those runs); grid and all runs: ℓ^2_{A} with the selected probes over those runs and over every run, failures included; oracle: true signatures, on the same data; losses and predictions $\times 10^3$, with $n_B = 2n/3$ the size of parts 1 and 2; selection: time on one core.

d	n/d	all probes	probe error	grid	all runs	oracle	$\bar{E}^{\mathrm{qp}}/n_B$	$\bar{E}/n_B$	selection (s)
27	1600	1/3	0.40	36.13	70.37	30.51	29.12	27.82	4
27	6400	3/3	0.18	6.86	6.86	6.84	7.28	6.95	3
27	25600	3/3	0.10	1.81	1.81	1.81	1.82	1.74	4
64	1600	0/3	–	–	62.01	34.79	33.58	32.29	23
64	6400	3/3	0.18	8.37	8.37	8.29	8.39	8.07	22
64	25600	3/3	0.08	2.06	2.06	2.06	2.10	2.02	22
125	1600	0/3	–	–	99.46	38.13	36.42	35.15	100
125	6400	3/3	0.17	9.21	9.21	9.17	9.10	8.79	100
125	25600	3/3	0.07	2.17	2.17	2.16	2.28	2.20	103
216	6400	2/2	0.18	9.38	9.38	9.33	9.60	9.28	533

regime of Section 6; with $\hat{\hat{\delta}}_j < 0.1$ the fractions below become 12% ($m = 50$) and 46% ($m = 400$). *Results*: the fraction of components in the first-order region grows from 8% ($m = 10$) to 64% ($m = 200$) and 199 of 200 ($m = 400$); within it the Spearman correlation between $\hat{\delta}_j$ and δ_j is 0.92 at $m = 50$ and 0.88 at $m = 400$ (pooled over m: 0.84, slope 0.731 on log scales), and $\delta_j/\hat{\delta}_j$ has median 6.1. The energy ranks better than the nearest-neighbour gap alone: at $m = 50$ the Spearman correlation with δ_j is 0.89 for E_j and 0.72 for $1/S_{\mathrm{nn}}(j)$ within the first-order region (52 components), and 0.84 against 0.78 over all 200 components, although the two quantities rank the components alike (Spearman 0.94); the nearest neighbour contributes a median 5% of E_j (2% to 9% between the deciles). Fitting $\delta_j \approx a_j/n_{\mathrm{half}} + f_j$ over m gives, for the 128 components with at least two sample sizes in the region, a floor f_j with median 0.59 and coefficients a_j that, with two to six points each, do not rank with the prediction (Spearman 0.073), so the ranking of δ_j within each m is carried mostly by the floor. Most of these discrepancies are close to their random level (mean 0.723 at $m = 400$ against a predicted 0.118), so the fitted floor largely measures saturation. The perturbed mixings of Table 3 bias every estimate in the same way, so their floor does not show in a split-half comparison. Of the 70 subsets of four environments, the five with the smallest energy ($\bar{E}/d \approx 440$) give $\delta = 0.884$ at $m = 50$ against 0.909 for the five with the largest ($\bar{E}/d \approx 2769$), and eight rest segments (participation ratio 1.54) give 0.894 against 0.860 for the seven tasks and rest, whose signatures differ little from one another (cf. the correspondence of task and rest architectures, Smith et al., 2009) (same m and number of repeats): at this m all of them are close to saturation. In Figure 3(b), with two environments the nearest-neighbour distances of the estimated signatures are below their estimation noise, so those energies, and their growth with d, are rough. Figure 9 shows the components with the smallest and the largest predicted error and the hardest pair.

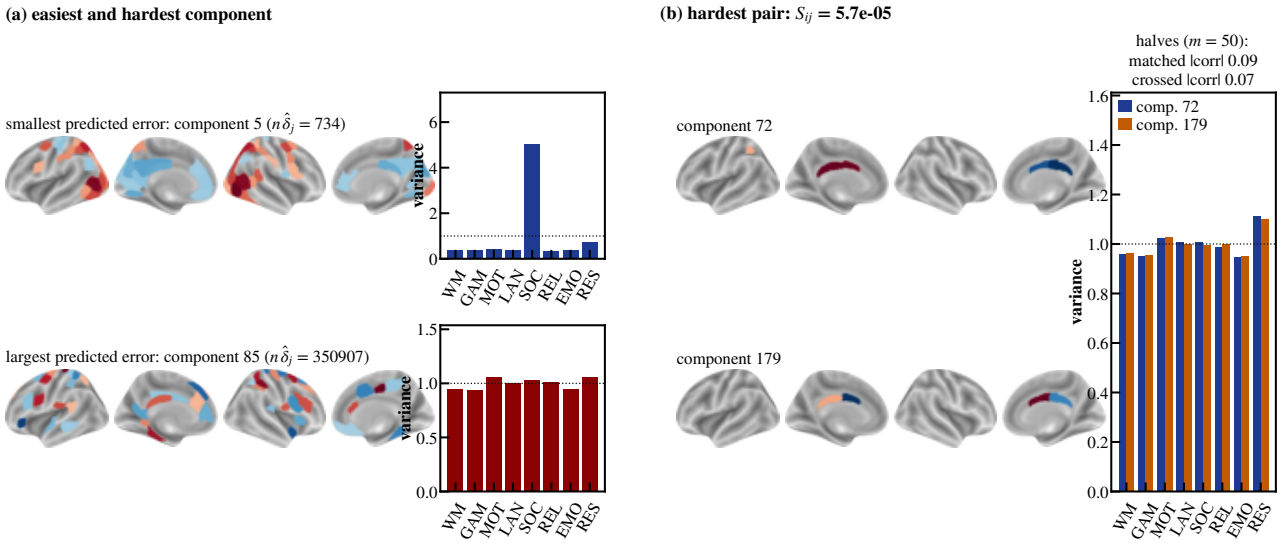

Figure 9. **Real data: components ($d = 200$). (a)** The components with the smallest and the largest predicted error ($n\hat{\delta}_j = 734$ and 350907): parcel loadings on the inflated cortical surface and task-variance signatures. **(b)** The hardest pair (components 72 and 179, $S_{ij} = 0.0000570$) has nearly identical signatures; between halves of $m = 50$ subjects the matched $|\mathrm{corr}|$ of their estimates is 0.091 and the crossed one 0.071.