arXiv:2611.06828v1 [cs.LG] 4 Nov 2026

Systolic Linear Attention: Co-Designing Tile-Grouped Gated Delta Attention for TPUs

Gagik Amirkhanyan
Google
agagik@google.com

Nuojin Cheng
Google
chengnuojin@google.com

Abstract

Gated-delta linear attention layers such as Gated DeltaNet (GDN), Kimi Delta Attention (KDA), and GDN-2 are tuned for GPUs, where small tensor-core tiles and fast vector registers make per-channel decay rescaling and sequential in-block solves cheap. On Google TPUs, which concentrate compute in wide 256×256 systolic matrix units backed by a slower vector unit, porting these layers directly leaves half of each matrix tile empty and bottlenecks on vector operations. We present **Systolic Linear Attention (SLA)**, which reshapes gated delta attention around the matrix unit by widening keys to 256 with grouped value heads, grouping decay gates into four 64-channel tiles, and computing the chunk solve entirely with block matrix products in a single fused kernel. On TPU v7x, a full `SLA-R4` layer runs $1.6\times$–$2.5\times$ faster than the open-source Tokamax KDA kernel ($1.4\times$–$2.4\times$ with exact chunk inversion), driven primarily by head sharing and kernel fusion. Across a 2,048-chip TPU v7x pre-training study of sparse Mixture-of-Experts (MoE) models from 6.5B to 63.1B parameters, SLA trains 11%–24% faster than KDA at similar perplexity, and reaches slightly lower perplexity than KDA in smaller matched-seed runs.

1 Introduction

Linear attention removes the quadratic compute and key-value (KV) cache of softmax attentions [15, 63] by storing the context in a fixed-size state matrix $S_t \in \mathbb{R}^{d_k \times d_v}$ (for key width d_k and value width d_v), updated at each token t with its key $k_t \in \mathbb{R}^{d_k}$ and value $v_t \in \mathbb{R}^{d_v}$ as [33]:

$$S_t = S_{t-1} + k_t v_t^\top.$$

Because this sum grows without bound, the state saturates on long sequences. Gated models such as Mamba-2 [14] and GLA [66] add a data-dependent forget (decay) gate $F_t \in (0,1]^{d_k \times d_k}$,

$$S_t = F_t S_{t-1} + k_t v_t^\top,$$

but decay fades all stored keys together and hurts associative recall [5].

To overwrite individual keys selectively, DeltaNet [50, 67] introduced the delta rule with unit-norm key k_t ($\|k_t\|_2 = 1$), scalar write gate β_t, and $d_k \times d_k$ identity I_{d_k}:

$$S_t = (I_{d_k} - \beta_t k_t k_t^\top) S_{t-1} + \beta_t k_t v_t^\top.$$

Viewing such updates as online gradient descent, prior work also added heavy-ball momentum so that a single noisy update cannot overwrite established memories [9, 44]. Combining decay with the delta rule produced **Gated DeltaNet (GDN)** [68], adopted in Qwen3-Next [48]. **Kimi Delta Attention (KDA)** [71] then expanded GDN's scalar forget gate into a per-channel diagonal gate $F_t = \mathrm{diag}(f_t)$

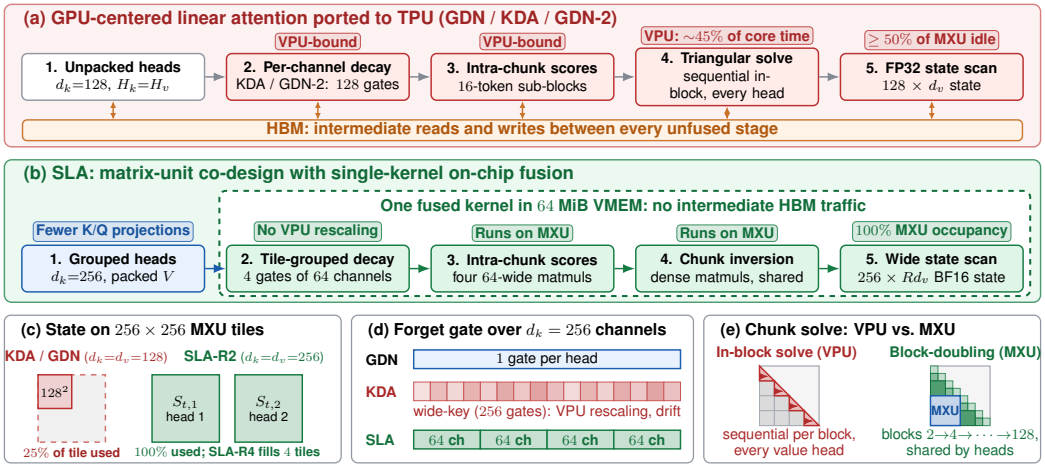

Figure 1: **Overview of SLA. (a)** Direct ports of GPU-tuned gated-delta layers (GDN [68], KDA [71], GDN-2 [25]) run as an unfused pipeline on TPUs: intermediates round-trip through HBM, per-channel decay rescaling in KDA and GDN-2 and in-block triangular substitution run on the VPU, and 128-wide states leave at least half of each 256×256 MXU tile padded. **(b)** SLA fuses stages 2–5 into one VMEM kernel. **(c)** A head with key and value widths $d_k = d_v = 128$ fills a quarter of an MXU tile, while SLA's wide packed state fills whole tiles. **(d)** Decay gates are grouped into four 64-channel tiles, so intra-chunk scores are four 64-wide MXU matmuls with no VPU rescaling. **(e)** The chunk solve uses only block matrix products on the MXU.

with channel-wise decay vector $f_t \in (0,1]^{d_k}$, and **GDN-2** [25] decoupled per-channel erase and write gates. Section 5 surveys further variants.

However, these architectures and their chunkwise WY training algorithms [10, 66, 68] were developed and tuned on NVIDIA GPUs (`fla` [65]). They map less well onto 2D systolic-array accelerators such as Google TPUs [18, 31, 32], for three reasons:

Small tiles and unpacked heads. GPU Tensor Cores operate on small fragments, so $d_k = d_v = 128$ with one key head per value head already reaches high occupancy. The TPU v6e/v7x Matrix Multiply Unit (MXU) is a 256×256 systolic array: a 128-wide state pads half of every tile, and projections, convolutions, and chunk solves are repeated for every head.

Per-channel decay gates. In KDA and GDN-2, each key channel decays at its own rate. Kernels must rescale keys elementwise in 16-token sub-blocks to avoid float32 overflow. GPUs absorb this vector work cheaply; on TPUs, where matrix throughput exceeds vector throughput by roughly $128\times$, it stalls the kernel.

Sequential triangular solves. The chunkwise WY algorithm [10, 68] partitions training sequences into chunks of length C and solves a $C \times C$ unit lower-triangular system per head and chunk. Existing kernels [19, 65] substitute sequentially within diagonal blocks; on TPUs this runs on the Vector Processing Unit (VPU) and takes about half of the core kernel time in our measurements.

Kernel tuning alone does not close this gap: fusing a GPU-shaped KDA layer into on-chip memory speeds it up by only $1.03\times$ on TPU v7x (Table 5), because the time goes to VPU solves and padded MXU tiles rather than High-Bandwidth Memory (HBM) traffic. We therefore change the layer itself. **Systolic Linear Attention (SLA)** applies a block-diagonal, tile-grouped forget gate $F_t \in \mathbb{R}^{d_k \times d_k}$ to S_{t-1} and then performs the delta update with key $k_t \in \mathbb{R}^{d_k}$, packed value $v_t \in \mathbb{R}^{Rd_v}$ (stacking $R \in \{2, 4\}$ value heads per key head), and a write gate $\beta_t \in (0, 2)$ shared across the packed value heads:

$$S_t = F_t S_{t-1} + \beta_t k_t \big(v_t - (F_t S_{t-1})^\top k_t\big)^\top. \tag{1}$$

SLA also smooths chunk updates with a chunk-boundary momentum term with coefficient $\mu = 0.25$ that stays in on-chip SRAM (VMEM). SLA combines three changes (Section 2):

Wide keys shared across value heads. SLA widens keys to $d_k = 256$ and shares each key head and gate β_t across $R \in \{2, 4\}$ value heads (`SLA-R2`, `SLA-R4`), following grouped value attention (GVA) [47]. The packed state fills whole MXU tiles, and key-query projections, convolutions, and chunk solves are computed $R\times$ less often.

Tile-grouped forget gate. SLA splits the 256 key channels into 4 tiles of 64 channels, each with one scalar decay rate. Intra-chunk scores then become four 64-wide MXU matmuls scaled by decay masks, with no VPU rescaling. This is a middle point between GDN's single rate per head and KDA's rate per channel.

Chunk inversion on the MXU. Because β_t is shared, SLA needs one $C \times C$ solve per key head. It computes the inverse only with block matrix products on the MXU: an exact recursive block-doubling inverse by default, or an optional 1-step Neumann approximation for faster solves, fused with the rest of the layer in VMEM.

We evaluate SLA on TPU v7x at the kernel, layer, serving, and context-parallel level, and across a 2,048-chip TPU v7x pre-training study of sparse Mixture-of-Experts (MoE) models [16, 52] from 6.5B to 63.1B parameters on Dolma [54, 60] (Section 3). We find that most of SLA's speedup comes from head sharing and fusion, which the tile-grouped gate and MXU chunk inversion make possible. On matched seeds, SLA improves on standard KDA and stays within about 0.1 perplexity of a wide-key per-channel KDA control while avoiding per-channel long-context drift; the optional 1-step Neumann inverse trades about 0.2 perplexity for speed. SLA is therefore best understood as a hardware-efficient operating point rather than a more expressive update rule; it is complementary to modeling advances such as GDN-2, which reaches lower perplexity per token (Section 3.5).

In the following sections, Section 2 reviews gated delta attention, its TPU bottlenecks, and SLA's design; Section 3 evaluates SLA from kernels and serving to MoE pre-training; Section 4 details optional low-precision key whitening; and Section 5 covers related work.

2 Hardware Bottlenecks and the SLA Design

To motivate SLA, we first review gated delta recurrences and their chunkwise WY training algorithm (Section 2.1), identify three hardware bottlenecks that GPU-centric designs hit on TPUs (Section 2.2), and then present SLA's mathematical formulation and chunk-boundary momentum (Section 2.3), its three algorithmic pillars and mathematical derivations (Section 2.4), and its single-kernel VMEM backward pass (Section 2.5), followed by empirical evaluation in Section 3.

2.1 Preliminaries: Gated Delta Recurrences and Chunkwise WY Parallelism

Given an input sequence $X = (x_1, \ldots, x_T) \in \mathbb{R}^{T\times D}$ of length T and model width D, a gated-delta linear attention layer with H_v value heads and H_k query-key heads maps each token $x_t \in \mathbb{R}^D$ ($t \in \{1, \ldots, T\}$) to L_2-normalized queries and keys $q_t, k_t \in \mathbb{R}^{d_k}$ ($\|q_t\|_2 = \|k_t\|_2 = 1$), values $v_t \in \mathbb{R}^{d_v}$, and head output $o_t \in \mathbb{R}^{d_v}$ (with $H_k = H_v$ in unpacked baselines). Rather than caching t past tokens, each head maintains a fixed-size associative state matrix $S_t \in \mathbb{R}^{d_k\times d_v}$ updated by applying a diagonal forget gate $F_t \in (0, 1]^{d_k\times d_k}$ followed by a rank-one delta update [9, 50, 59, 67] with scalar step size β_t:

Prior Work: Gated Delta Recurrence in GDN and KDA

$$
\begin{aligned}
S_t &= F_t S_{t-1} + \beta_t k_t \big(v_t - (F_t S_{t-1})^\top k_t\big)^\top \\
&= \big(I_{d_k} - \beta_t k_t k_t^\top\big) F_t S_{t-1} + \beta_t k_t v_t^\top, \\
o_t &= S_t^\top q_t.
\end{aligned} \tag{2}
$$

Here $(F_t S_{t-1})^\top k_t$ is the value currently retrieved for key k_t, and β_t sets how far it is moved toward v_t. Existing linear attentions differ only in (F_t, β_t) in Eq. (2), where we characterize F_t by its number of **decay groups** G (distinct decay rates across the d_k key channels of a head). Unweighted

Linear Attention [33] sets $F_t = I_{d_k}$ without erasure ($S_t = S_{t-1} + k_t v_t^\top$), while DeltaNet [50, 67] and test-time learning [59, 64] retain $F_t = I_{d_k}$ with the rank-one delta update. Recent gated-delta architectures instantiate (F_t, β_t) using the logistic sigmoid $\sigma(\cdot)$ on scalar write-gate logit $\tilde{\beta}_t \in \mathbb{R}$ and $\mathrm{softplus}(u) = \ln(1 + e^u)$ on forget-gate logits $\tilde{f}_t \in \mathbb{R}^G$:

Gated DeltaNet (GDN). **GDN** [68] uses one scalar decay $\gamma_t \in (0, 1]$ per head ($G = 1$) together with a unit-interval scalar step size:

$$F_t = \gamma_t I_{d_k}, \qquad \beta_t = \sigma(\tilde{\beta}_t) \in (0, 1).$$

Kimi Delta Attention (KDA). **KDA** [71] assigns an independent decay rate to every key channel ($G = d_k$, $f_t = \exp(-\mathrm{softplus}(\tilde{f}_t))$):

$$F_t = \mathrm{diag}(f_t), \qquad f_t \in (0, 1]^{d_k}.$$

While standard GDN and KDA use $\beta_t = \sigma(\tilde{\beta}_t) \in (0, 1)$, negative-eigenvalue extensions [20] widen the step-size range to

$$\beta_t = 2\sigma(\tilde{\beta}_t) \in (0, 2),$$

so that $I_{d_k} - \beta_t k_t k_t^\top$ has eigenvalue $1 - \beta_t \in (-1, 1)$ along k_t (adopted by SLA in Section 2.3).

GDN-2. **GDN-2** [25] retains per-channel decay ($G = d_k$) and further replaces the scalar β_t with separate channel-wise erase and write gates.

Chunkwise WY parallelization. To saturate matrix units during training, chunkwise WY parallelization [10, 67, 68] partitions the sequence into T/C contiguous chunks of length C (Algorithm 1, written unified across sharing ratio $R = H_v/H_k$ and gate scale $s_\beta \in \{1, 2\}$, where unpacked GDN/KDA set $R = 1, s_\beta = 1$). Op 1 projects X to pre-convolution activations $\tilde{Q}, \tilde{K}, \tilde{V}$, gate logits $(\tilde{\beta}_t, \tilde{f}_t)_{t=1}^T$, and output-gate logits $(z_t)_{t=1}^T$ ($z_t \in \mathbb{R}^{Rd_v}$ per key head), followed by causal short convolutions, L_2 normalization into Q, K, and gate activations (Op 2). Within chunk $c \in \{1, \ldots, T/C\}$ starting from boundary state $S_{c-1} \in \mathbb{R}^{d_k \times Rd_v}$ ($S_0 = 0$), let local token indices satisfy $1 \le j \le t \le C$, let $Q_c, K_c \in \mathbb{R}^{C \times d_k}$ and $V_c, O_c \in \mathbb{R}^{C \times Rd_v}$ stack the chunk's queries, keys, values, and readouts as rows, let $\beta_c = (\beta_1, \ldots, \beta_C)^\top \in \mathbb{R}^C$ stack its write gates, and let I_C denote the $C \times C$ identity. Define

$$P_{t \leftarrow j} = \prod_{s=j+1}^{t} F_s \quad (P_{t \leftarrow t} = I_{d_k}), \qquad F_c = P_{C \leftarrow 0} = \prod_{s=1}^{C} F_s$$

as the cumulative diagonal decay from local token j to t and the total chunk decay, and let $\bar{K}_c, \bar{Q}_c, \hat{K}_c \in \mathbb{R}^{C \times d_k}$ denote the decayed key and query matrices whose t-th rows are $(P_{t \leftarrow 0} k_t)^\top$, $(P_{t \leftarrow 0} q_t)^\top$, and $(P_{C \leftarrow t} k_t)^\top$. Within a chunk, each token's write depends on what earlier tokens in the same chunk have already written, so the C writes cannot be computed independently. Collect them as rows of $E_c \in \mathbb{R}^{C \times Rd_v}$. Each row is what the token would write on its own, $\mathrm{diag}(\beta_c)(V_c - \bar{K}_c S_{c-1})$, minus the earlier tokens' writes, weighted by how strongly their decayed keys overlap with its own key. Using indicator $\mathbf{1}[\cdot] \in \{0, 1\}$, these overlaps define the strictly lower-triangular Gram matrix $M_c \in \mathbb{R}^{C \times C}$ (Op 3a, alongside lower-triangular intra-chunk query-key score matrix $L_c \in \mathbb{R}^{C \times C}$, Op 3b), so the chunk writes satisfy (derived in Section 2.4.2):

$$E_c = \mathrm{diag}(\beta_c)\,(I_C + M_c)^{-1}(V_c - \bar{K}_c S_{c-1}).$$

Inverting $A_c = (I_C + M_c)^{-1}$ (Op 4) and forming the WY factors U_c, W_c (Op 5) runs in parallel across chunks, leaving a sequential scan over the T/C chunk boundaries for S_c and optional velocity buffer $Z_c \in \mathbb{R}^{d_k \times Rd_v}$ ($Z_0 = 0$, momentum $\mu \ge 0$; Op 6), parallel chunk readouts O_c (rows $o_t^\top$; Op 7a), and gated root-mean-square normalization ($\mathrm{RMSNorm}$, with elementwise product $\odot$) followed by the output projection to $Y \in \mathbb{R}^{T \times D}$ (Op 7b–c).

2.2 Three Hardware Bottlenecks on TPUs

TPUs pair a wide systolic MXU with a much slower VPU. On TPU v7x the MXU tile is 256×256 and the MXU:VPU throughput ratio is about $128\times$, compared with about $15\times$ on H100 (Table 1). This skew creates three bottlenecks for Algorithm 1:

Algorithm 1 Chunkwise WY Parallel Forward Pass

Require: Sequence $X \in \mathbb{R}^{T\times D}$, chunk size C, head ratio R, gate scale s_β, momentum μ ($(R, s_\beta, \mu) = (1, 1, 0)$ for GDN/KDA; $(R, 2, 0.25)$ with $R \in \{2, 4\}$ for SLA)
Ensure: Output $Y \in \mathbb{R}^{T\times D}$, final state $S_{T/C} \in \mathbb{R}^{d_k \times Rd_v}$

$\tilde{Q}, \tilde{K}, \tilde{V}, (\tilde{\beta}_t, \tilde{f}_t, z_t)_{t=1}^{T} \leftarrow \mathrm{LinearProj}(X)$ ▷ Op 1 (MXU)
$Q, K \leftarrow \mathrm{L2Norm}(\mathrm{Conv1D}(\tilde{Q}, \tilde{K}))$ ▷ Op 2a (VPU)
$V \leftarrow \mathrm{SiLU}(\mathrm{Conv1D}(\tilde{V}))$
$\beta_t \leftarrow s_\beta \sigma(\tilde{\beta}_t)$; $F_t \leftarrow \mathrm{diag}(e^{-\mathrm{softplus}(\tilde{f}_t)})$ ▷ Op 2b (VPU)
Partition Q, K, V, β, F into T/C chunks; $S_0 \leftarrow 0$, $Z_0 \leftarrow 0 \in \mathbb{R}^{d_k \times Rd_v}$
for all $c \in \{1, \ldots, T/C\}$ **in parallel do**
  $P_{t\leftarrow j} \leftarrow \prod_{s=j+1}^{t} F_s \;\; \forall 0 \le j \le t \le C \quad (P_{t\leftarrow t} = I_{d_k})$
  $(\bar{K}_c, \bar{Q}_c, \hat{K}_c)_t \leftarrow P_{t\leftarrow 0} k_t, P_{t\leftarrow 0} q_t, P_{C\leftarrow t} k_t$ ▷ each row t
  $(M_c)_{t,j} \leftarrow \mathbf{1}[t > j]\, \beta_j k_t^\top P_{t\leftarrow j} k_j$ ▷ Op 3a: Gram
  $(L_c)_{t,j} \leftarrow \mathbf{1}[t \ge j]\, q_t^\top P_{t\leftarrow j} k_j$ ▷ Op 3b: Score
  $A_c \leftarrow (I_C + M_c)^{-1} \in \mathbb{R}^{C\times C}$ ▷ Op 4: Solve
  $U_c \leftarrow \mathrm{diag}(\beta_c) A_c V_c \in \mathbb{R}^{C\times Rd_v}$ ▷ Op 5a (MXU)
  $W_c \leftarrow \mathrm{diag}(\beta_c) A_c \bar{K}_c \in \mathbb{R}^{C\times d_k}$ ▷ Op 5b (MXU)
end for
for $c = 1$ **to** T/C **do**
  $E_c \leftarrow U_c - W_c S_{c-1}$ ▷ Op 6a (MXU)
  $Z_c \leftarrow \mu Z_{c-1} + \hat{K}_c^\top E_c$ ▷ Op 6b (MXU)
  $S_c \leftarrow F_c S_{c-1} + Z_c$
end for
for all $c \in \{1, \ldots, T/C\}$ **in parallel do**
  $O_c \leftarrow \bar{Q}_c S_{c-1} + L_c E_c$ ▷ Op 7a (MXU)
end for
$Y \leftarrow \mathrm{OutProj}\big((\sigma(z_t) \odot \mathrm{RMSNorm}(o_t))_{t=1}^{T}\big)$ ▷ Op 7b–c
return $Y, S_{T/C}$

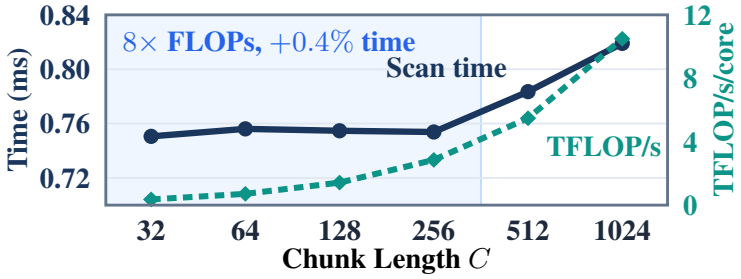

Figure 2: **Sequential state-scan cost on TPU v7x.** A 64-step scan of $(C, 256) \times (256, 256)$ matmuls takes nearly constant time for $C \in \{32, \ldots, 256\}$, so a fixed per-step overhead, not matmul FLOPs, dominates the scan.

Bottleneck 1: Narrow states and a fixed per-step scan cost. Op 6 carries the state across T/C chunk boundaries sequentially. On TPU v7x, each scan step costs a nearly fixed $\approx 12\,\mu$s regardless of chunk size: $8\times$ more FLOPs per step adds under 1% time (Figure 2). Scan cost is thus set by the number of steps and heads, not by matmul size. A 256-wide tile also processes a $d_k = 128$ state at the same cost as a $d_k = 256$ state, so 128-wide states leave half of each tile idle while paying the full per-step cost for every head.

Bottleneck 2: Per-channel decay. KDA [71] and GDN-2 [25] give every key channel its own forget gate ($F_t = \mathrm{diag}(f_t)$). Let $\Lambda_{t,i} = \sum_{s=1}^{t} \log f_{s,i}$ be the cumulative decay of channel i within a chunk. Each Gram entry then weights channels differently:

$$(M_c)_{t,j} = \mathbf{1}[t > j]\, \beta_j \sum_i e^{\Lambda_{t,i} - \Lambda_{j,i}} k_{t,i} k_{j,i}.$$

Table 1: **Documented hardware specifications per chip** for NVIDIA H100 SXM and four Google TPU generations [6, 18, 32]. MXU:VPU is the ratio of peak BF16 matrix throughput to FP32 vector throughput; Ridge is the BF16 arithmetic-intensity ridge point (FLOPs/byte). *8-bit peak is FP8 on H100 and TPU v7x and INT8 on TPU v4–v6e.

Accelerator	Cores	Array Tile	BF16 TF/s	8-bit TF/s*	VPU TF/s	MXU:VPU	HBM (GB)	HBM GB/s	Ridge
H100 SXM	132 SM	16×16	989	1,979	67.0	14.8×	80	3,350	295
TPU v4	2	128×128	275	275	3.2	85.9×	32	1,200	229
TPU v5p	2	128×128	459	918	5.4	85.0×	95	2,765	166
TPU v6e	1	256×256	918	1,836	7.2	127.5×	32	1,600	574
TPU v7x	2	256×256	2,307	4,614	18.0	128.2×	192	7,373	313

To use the MXU, kernels fold $e^{\Lambda_{t,i}}$ and $e^{-\Lambda_{j,i}}$ into the keys, but $e^{-\Lambda_{j,i}}$ overflows FP32 for long chunks. KDA kernels therefore split each chunk into 16-token sub-blocks and re-center the keys on the VPU before every sub-block matmul, which stalls the kernel. In our KDA runs, about 1% of channels also learn near-unit forget gates, which hurts length extrapolation unless the gates are clamped (Section 3.5). GDN's scalar gate ($G = 1$) avoids both issues, at the cost of a single memory horizon per head.

Bottleneck 3: Sequential triangular solves. Existing kernels [19, 65] compute $A_c = (I_C + M_c)^{-1}$ (Op 4) by sequential substitution within small diagonal blocks, which runs on the VPU. Repeated for every chunk and each of $H_k = H_v = 8$ heads, together with the Gram matrices (Ops 3–5) it accounts for about 45% of the core intra-chunk and scan time in our measurements (Table 5), keeping the best chunk size near 64–128.

2.3 SLA: Addressing the Three Bottlenecks

Guided by the three systolic bottlenecks above, we now formulate SLA, which reshapes the state geometry, decay structure, and chunk inversion around the 256×256 systolic tile:

Definition 2.1 (SLA Recurrence)**.** Consider a linear attention layer with H_v value heads of width d_v and $H_k = H_v/R$ query-key heads of width $d_k = 256$ indexed by $h_k \in \{1, \dots, H_k\}$, where each query-key head is shared across $R \in \{2, 4\}$ value heads. At token t, each query-key head computes:

- **Shared query, key, and write gate:** L_2-normalized $q_t, k_t \in \mathbb{R}^{256}$ with $\|q_t\|_2 = \|k_t\|_2 = 1$, and a scalar step size $\beta_t = 2\sigma(\tilde{\beta}_t) \in (0, 2)$ ($s_\beta = 2$) that doubles the $(0, 1)$ range of standard GDN/KDA;
- **Packed values and output gates:** R value vectors $v_{t,1}, \dots, v_{t,R} \in \mathbb{R}^{d_v}$ stacked as $v_t = [v_{t,1}; \dots; v_{t,R}] \in \mathbb{R}^{Rd_v}$, and output-gate vector $z_t = [z_{t,1}; \dots; z_{t,R}] \in \mathbb{R}^{Rd_v}$ with per-value-head slices $z_{t,m} \in \mathbb{R}^{d_v}$ ($m \in \{1, \dots, R\}$);
- **Four-tile forget gate:** a block-diagonal gate $F_t \in (0, 1]^{256 \times 256}$ that splits the 256 key channels into $G = 4$ contiguous tiles of $b = 64$ channels, with forget-gate logit $\tilde{f}_t = (\tilde{f}_{t,1}, \dots, \tilde{f}_{t,4}) \in \mathbb{R}^4$, all-ones vector $\mathbf{1}_{64} \in \mathbb{R}^{64}$, and one scalar decay $\gamma_{t,g} = \exp(-\operatorname{softplus}(\tilde{f}_{t,g})) \in (0, 1]$ on each tile $g \in \{1, 2, 3, 4\}$:

$$F_t = \operatorname{diag}(\gamma_{t,1}\mathbf{1}_{64},\ \gamma_{t,2}\mathbf{1}_{64},\ \gamma_{t,3}\mathbf{1}_{64},\ \gamma_{t,4}\mathbf{1}_{64}). \tag{3}$$

(i) Intra-chunk token recurrence. Maintaining a packed state $S_t = [S_{t,1}, \dots, S_{t,R}] \in \mathbb{R}^{256 \times Rd_v}$ initialized at $S_0 = 0$, with slice $S_{t,m} \in \mathbb{R}^{256 \times d_v}$ and output slice $o_{t,m} \in \mathbb{R}^{d_v}$ ($o_t = [o_{t,1}; \dots; o_{t,R}]$) for value head $m \in \{1, \dots, R\}$, SLA decays the state, writes the prediction error, and normalizes each value head:

SLA Token-Level Recurrence ($d_k = 256,\ G = 4,\ b = 64$)

$$S_t = F_t S_{t-1} + \beta_t k_t \big(v_t - (F_t S_{t-1})^\top k_t\big)^\top, \tag{4}$$

$$o_{t,m} = \sigma(z_{t,m}) \odot \operatorname{RMSNorm}\big(S_{t,m}^\top q_t\big). \tag{5}$$

Here $(F_t S_{t-1})^\top k_t$ is the value currently stored for k_t, so Eq. (4) writes only the prediction error, and one rank-one write updates all R value heads.

(ii) Chunk-boundary momentum. Let $\tilde{S}_c$ be the state after unrolling Eq. (4) over the C tokens of chunk c from S_{c-1}, and let $F_c = \prod_{t \in c} F_t$. SLA updates the chunk-boundary state S_c via a velocity buffer $Z_c \in \mathbb{R}^{256 \times R d_v}$ reset to zero at document boundaries:

$$Z_c = \mu Z_{c-1} + (\tilde{S}_c - F_c S_{c-1}), \qquad S_c = F_c S_{c-1} + Z_c, \tag{6}$$

with $\mu = 0.25$ by default, while $\mu = 0$ recovers $S_c = \tilde{S}_c$.

Notation and operator taxonomy. Unless redefined locally, symbols follow Algorithm 1 and Definition 2.1; Table 2 collects them for quick reference (symbols used only inside one section, such as the roofline constants in Section 3.1, are defined where they first appear). Table 3 contrasts the algebraic operators and systolic hardware mappings of GDN [68], KDA [71], GDN-2 [25], **SLA** (`SLA-R2` / `SLA-R4`), and `SPD-SLA-W`, an optional extension of SLA that preconditions chunks with Σ_c to form whitened keys $\tilde{K}_c$ (rows $\tilde{k}_t^\top$) under step-size margin $\varepsilon_\beta > 0$ (Section 4); `SPD-SLA-W` is not used in any core SLA result except the middle bar of Figure 7a and the conditioning analysis of Section 4.

Table 2: Summary of notation used across the paper.

Symbol	Meaning
$T,\ C,\ c$	Sequence length; chunk length ($C = 128$); chunk index $c \in \{1, \dots, T/C\}$.
t, j, s	Token indices inside a chunk, $1 \le j \le t \le C$ (with chunk-entry boundary 0 in $P_{t \leftarrow 0}$).
$D,\ H_k,\ H_v,\ R$	Model width; number of query-key heads; number of value heads; sharing ratio $R = H_v / H_k \in \{2, 4\}$.
$d_k,\ d_v,\ D_{v,\text{pack}},\ D_v$	Key width ($d_k = 256$); per-value-head width; packed value width per key head $D_{v,\text{pack}} = R d_v$; total value width $D_v = H_v d_v$.
$G,\ b,\ g$	Number of decay tiles ($G = 4$), tile width $b = d_k/G = 64$, tile index $g \in \{1, \dots, G\}$.
$q_t, k_t \in \mathbb{R}^{d_k}$, $v_t \in \mathbb{R}^{D_{v,\text{pack}}}$	L_2-normalized query and key of one key head; packed values of its R value heads.
$\beta_t \in (0, s_\beta)$	Write gate $s_\beta \sigma(\tilde{\beta}_t)$ per token and key head, shared across the R packed value heads ($s_\beta = 2$ in SLA; $s_\beta = 1$ in GDN/KDA).
$f_{t,i},\ \gamma_{t,g} \in (0, 1]$	Per-channel decay of channel i (KDA, $G = d_k$); per-tile decay $\gamma_{t,g} = e^{-\operatorname{softplus}(\tilde{f}_{t,g})}$ of tile g (SLA, $f_{t,i} = \gamma_{t,g}$ on tile g).
$F_t,\ P_{t \leftarrow j},\ F_c$	Diagonal forget gate $F_t = \operatorname{diag}(f_t)$; cumulative decay $P_{t \leftarrow j} = \prod_{s=j+1}^{t} F_s$ with $P_{t \leftarrow t} = I_{d_k}$; chunk decay $F_c = P_{C \leftarrow 0}$.
$\Lambda_{t,i}$	Cumulative log-decay of channel i inside a chunk, $\Lambda_{t,i} = \sum_{s=1}^{t} \log f_{s,i}$.
$S_c,\ \tilde{S}_t,\ \Delta S_c$	Chunk-boundary state $S_c \in \mathbb{R}^{d_k \times D_{v,\text{pack}}}$; intra-chunk state after local token t ($\tilde{S}_0 = S_{c-1}$, $\tilde{S}_C$ written $\tilde{S}_c$ in Def. 2.1(ii)); increment $\Delta S_c = \tilde{S}_C - F_c S_{c-1}$.
$Z_c,\ \mu,\ \kappa_c$	Chunk-boundary momentum (velocity) buffer; momentum coefficient ($\mu = 0.25$); reset indicator $\kappa_c \in \{0, 1\}$, equal to 0 whenever any document reset occurs inside chunk c (so Z_{c-1} and cross-document carry are discarded).
$\Psi_g,\ \Gamma_g \in \mathbb{R}^{C \times C}$	Tile-g causal decay masks: $(\Psi_g)_{t,j} = \mathbf{1}[t \ge j] \prod_{s=j+1}^{t} \gamma_{s,g}$ and $(\Gamma_g)_{t,j} = \mathbf{1}[t > j]\, \beta_j (\Psi_g)_{t,j}$.
$M_c,\ L_c,\ A_c \in \mathbb{R}^{C \times C}$	Strictly lower-triangular Gram matrix M_c (Op 3a), lower-triangular query-key score matrix L_c (Op 3b), and WY inverse $A_c = (I_C + M_c)^{-1}$ (Op 4).
$\bar{Q}_c, \bar{K}_c, \hat{K}_c \in \mathbb{R}^{C \times d_k}$	Row-stacked decayed queries $(P_{t \leftarrow 0} q_t)^\top$, decayed keys $(P_{t \leftarrow 0} k_t)^\top$, and carried keys $(P_{C \leftarrow t} k_t)^\top$.
$U_c,\ W_c,\ E_c$	WY factors $U_c = \operatorname{diag}(\beta_c) A_c V_c$, $W_c = \operatorname{diag}(\beta_c) A_c \bar{K}_c$, and chunk writes $E_c = U_c - W_c S_{c-1}$ (Ops 5–6a).
$P,\ p,\ L_{\text{loc}}$	Number of context-parallel shards (devices); shard index; chunks per shard.
$\sigma(\cdot),\ \kappa_2(\cdot)$	Logistic sigmoid; spectral condition number (distinct from the reset indicator κ_c).

Why chunk-level momentum, implementation, and caveats. Unlike token-level momentum [9], applying heavy-ball momentum [44] once per chunk boundary (Op 6b of Algorithm 1) leaves the intra-chunk WY solve unchanged. Specifically, in Op 6b, the net chunk state increment in Eq. (6) is

$$\Delta S_c = \tilde{S}_c - F_c S_{c-1} = \hat{K}_c^\top E_c.$$

Table 3: Side-by-side algebraic and systolic hardware comparison across prior GPU gated-delta recurrences and our SLA family.

Design Axis	**GDN** [68]	**KDA** [71]	**GDN-2** [25]	**SLA** (`SLA-R2` / `SLA-R4`) (Ours)	`SPD-SLA-W` (Section 4)
Key/Value Heads	Unpacked $H_k = H_v$	Unpacked $H_k = H_v$	Unpacked $H_k = H_v$	Packed GVA $H_k = H_v/R$ ($R \in \{2, 4\}$)	Packed GVA $H_k = H_v/R$
State Tile & Dtype	$S^{(h_k)} \in \mathbb{R}^{128\times d_v}$ FP32 (50% pad)	$S^{(h_k)} \in \mathbb{R}^{128\times d_v}$ FP32 (50% pad)	$S^{(h_k)} \in \mathbb{R}^{d_k\times d_v}$ FP32	$S \in \mathbb{R}^{256\times D_{v,\text{pack}}}$ BF16 / Row-FP8 (no padding)	$S \in \mathbb{R}^{256\times D_{v,\text{pack}}}$ BF16 / Row-FP8
Forget Gate F_t	Scalar $\gamma_t I_{d_k}$ ($G = 1$)	Channel $\mathrm{diag}(f_t)$ ($G = d_k$)	Decoupled ch. ($G = d_k$)	Tile-grouped $\bigoplus_{g=1}^{4} \gamma_{t,g} I_{64}$ ($G = 4$)	Tile-grouped $\bigoplus_{g=1}^{4} \gamma_{t,g} I_{64}$ ($G = 4$)
Step Size β_t	$(0, 1)$ per H_v	$(0, 1)$ per H_v	Ch. vector per H_v	$(0, 2)$ scalar per H_k (neg. eig.)	Clamped $(0, (2 - \varepsilon_\beta)/k_t^\top \tilde{k}_t]$
Chunk Gram M_c	1× MXU + scalar mask	16-row VPU pivots ($\lvert\Lambda_{j,i}\rvert > 88.7$)	Decoupled ch. broadcast	$G \times [C, b] \times [b, C]$ pure MXU ($b \geq 64$, 0 ovf.)	Oblique $k_t^\top P_{t\leftarrow j} \tilde{k}_j$ MXU
Chunk Inverse A_c	$H_v\times$ blocked, sequential in-block	$H_v\times$ blocked, sequential in-block ($C \leq 64$ Tokamax)	$H_v\times$ blocked, sequential in-block	$H_k\times$ Exact Inv. (or 1-Step Neumann, $C = 128$)	$H_k\times$ MXU inverse + `Polar-2+` (5 MM)
Momentum / 2nd-Ord.	None ($\mu = 0$, Euclidean)	None ($\mu = 0$, Euclidean)	None ($\mu = 0$, Euclidean)	Chunk velocity $Z_c = \mu\kappa_c Z_{c-1} + \Delta S_c$	Whitened keys $\tilde{K}_c$ (Eq. (14))
Output Norm	Head RMSNorm	Head RMSNorm	Head RMSNorm	Per-value-head $\mathrm{RMSNorm}$	Per-value-head $\mathrm{RMSNorm}$
VMEM / HBM Trips	6 HBM trips (unfused)	6 HBM trips (unfused wrapper)	6 HBM trips (6.8–14.9 GiB)	Fused SLA: 1 kernel in 64 MiB VMEM + custom VJP	Fits 64 MiB VMEM
Context Parallelism	Prefix tree, one $d_k \times d_k$ map per head	Prefix tree, one $d_k \times d_k$ map per head	Prefix tree, one $d_k \times d_k$ map per head	Prefix tree, one $[256, 256]$ map per key head (2× with momentum)	Same as SLA

The velocity buffer $Z_c \in \mathbb{R}^{d_k\times D_{v,\text{pack}}}$ has the same shape as S_c and stays in VMEM inside the fused kernel, adding zero extra HBM state traffic. On packed multi-document streams, intra-chunk Gram and score matrices apply an exact segment mask $\mathbf{1}[\mathrm{seg}_t = \mathrm{seg}_j]$, only the final document segment writes to ΔS_c, and any document reset anywhere inside chunk c zeros both incoming state carry F_c and $\kappa_c = 0$, preventing mid-chunk velocity leakage. Moreover, for any eigenmode with chunk decay $f \in (0, 1)$ and generalized WY survival $j \in [-f, f]$ (including $\beta_t \in (0, 2)$), the lifted 2×2 companion block $\left[\begin{smallmatrix} j & \mu \\ j-f & \mu \end{smallmatrix}\right]$ has characteristic polynomial $p(\lambda) = \lambda^2 - (j + \mu)\lambda + \mu f$ ($j\mu$ cancels in $\det = \mu f$), satisfying Jury's Schur stability conditions $|\det| = \mu f < 1$, $p(1) \geq (1 - f)(1 - \mu) > 0$, and $p(-1) = 1 + j + \mu(1 + f) > 0$ unconditionally for all $\mu \in [0, 1)$. Because Z_c changes only at chunk boundaries, autoregressive decoding reproduces training exactly by updating Z_c once every C tokens. Under context parallelism, momentum couples (S_c, Z_c) into a $2d_k \times 2d_k$ affine map (Eq. (13)); Table 8 therefore reports both the uncoupled ($\mu = 0$, 256^2) and coupled ($\mu = 0.25$, 512^2) tree scans.

Table 4 compares $\mu = 0$ and $\mu = 0.25$ on a 12.5B-total (0.79B-active) MoE trained on 15.1B Dolma 3.5 tokens (7,200 steps, 2.1M tokens/step, 128 TPU v4 chips, seed 42). With 4:1 key sharing ($R = 4$), momentum lowers loss by 0.019–0.023 nats and matches KDA, recovering the capacity lost by sharing each key head across four value heads; with 2:1 sharing ($R = 2$), it changes loss by at most $+0.004$ nats. An input-dependent μ obtained by reusing the β_t projection degrades loss by $+0.026$ nats, so we keep μ fixed. We did not sweep μ beyond $\{0, 0.25\}$. In the TPU v7x ablation ladder (Table 11b), momentum is introduced together with VMEM fusion in `A5`, so its effect is not isolated there. Moreover, the ablation evidence is from a single seed, and momentum has costs: the recurrence now depends on the chunk size C, decoding must track chunk boundaries, and context parallelism must carry a $2\times$ larger transition state (Section 3.3).

2.4 The Three Algorithmic Pillars of SLA

Each pillar targets one of the bottlenecks in Section 2.2 (Figure 1c–e; detailed diagrams in Figure 3).

Table 4: Chunk-boundary momentum ablation: 20-step rolling training cross-entropy at step 7,199 (single pass over Dolma 3.5, 12.5B MoE, 128 TPU v4 chips, seed 42; lower is better). †Trained on an 8,200-step cosine schedule and evaluated at step 7,199.

Configuration	μ	Loss @ 7,199
KDA (per-channel decay, FP32 state)	–	2.2925
SLA, $R = 2$ (two runs)	0	2.2912 / 2.2925
SLA, $R = 2$	0.25	2.2955
SLA, $R = 2$, input-dependent μ	adaptive	2.3216
SLA, $R = 4$ (two runs)	0	2.3103 / 2.3148
SLA, $R = 4^{\dagger}$	0.25	**2.2914**

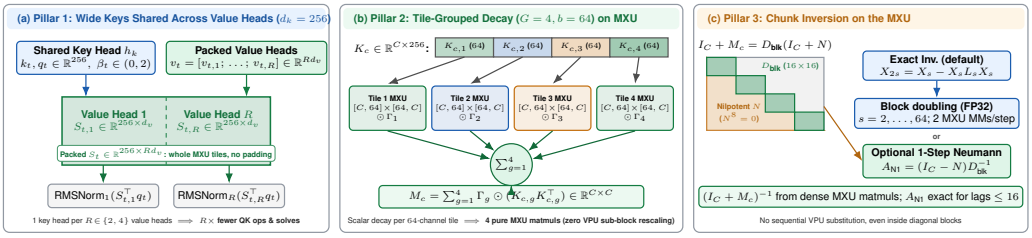

Figure 3: **Detailed architectural diagram of the three mathematical pillars of SLA. (a) Pillar 1:** Widening the key dimension to $d_k = 256$ and sharing each key head h_k and write gate $\beta_t \in (0, 2)$ across $R \in \{2, 4\}$ value heads forms a dense $256 \times D_{v,\text{pack}}$ state that fills whole 256×256 MXU tiles without padding while cutting QK projections, short convolutions, and chunk solves by $R\times$. **(b) Pillar 2:** Partitioning the $d_k = 256$ key channels into $G = 4$ contiguous tiles of $b = 64$ channels decomposes the intra-chunk coupling matrix M_c into four 64-wide MXU matrix multiplications weighted by scalar decay masks Γ_g, eliminating VPU sub-block rescaling. **(c) Pillar 3:** The default `Exact Inv.` computes $(I_C + M_c)^{-1}$ by recursive block doubling on the MXU; the optional 1-step Neumann inverse keeps only the first off-diagonal term of $I_C + M_c = D_{\text{blk}}(I_C + N)$ with 16×16 diagonal blocks, where $N^8 = 0$ at $C = 128$. Both avoid sequential VPU substitution and use dense MXU matrix multiplications (Section 2.4.3).

2.4.1 Pillar 1: Wide Keys with Grouped Value Heads

A key dimension below 256 still occupies a full 256-wide tile on TPU v6e/v7x, so SLA widens keys to $d_k = 256$, which costs no extra cycles in the inter-chunk scan. To keep parameters and non-matmul work in check, SLA adopts linear grouped-value head sharing (GVA) [48, 65, 71] together with gated output normalization [47]: each query-key head and gate β_t is shared across $R = H_v/H_k \in \{2, 4\}$ value heads. Packing these value heads side by side gives a state $S_t \in \mathbb{R}^{d_k \times Rd_v}$ that maps onto whole MXU tiles without padding. Query and key projections, short convolutions, M_c, L_c, and the solve A_c are computed once per key head, cutting their cost by $R\times$. Each value head's output is normalized separately (Eq. (5)). Head sharing itself is prior work; our contribution is choosing d_k and R to match the systolic tile.

Derivation of GVA systolic invariance. Fix total value width $D_v = H_v d_v$ and set $H_k = H_v/R$ with $D_{v,\text{pack}} = Rd_v$. Each key head $h_k \in \{1, \ldots, H_k\}$ maintains state $S^{(h_k)} \in \mathbb{R}^{d_k \times D_{v,\text{pack}}}$, so total state elements are invariant to R:

$$H_k d_k (Rd_v) = H_v d_k d_v.$$

Inter-chunk state scan FLOPs per token are likewise invariant to R:

$$H_k \big(2d_k D_{v,\text{pack}}\ [\text{erase}] + 2d_k D_{v,\text{pack}}\ [\text{write}] + 2d_k D_{v,\text{pack}}\ [\text{readout}]\big) = 6H_v d_k d_v.$$

Because β_t is scalar per key head h_k, QK projections ($2DH_k d_k$ parameters), QK causal `Conv1D` (for kernel width $K_{\text{conv}} = 4$, $2K_{\text{conv}} H_k d_k$ weights), chunk Gram matmuls ($4CH_k d_k$ FLOPs/token), and $[C, C]$ WY/Neumann solves (H_k systems per chunk) scale linearly with $H_k = H_v/R$ and shrink by $R\times$. Reshaping $\tilde{S}_t^\top q_t \in \mathbb{R}^{Rd_v}$ into R slices of width d_v and applying independent RMSNorm to all $H_v = H_k R$ value heads prior to the sigmoid output gate $\sigma(z_t) \in (0, 1)^{Rd_v}$ (slices $\sigma(z_{t,m})$ in

Eq. (5)) retains H_v independent head-scale degrees of freedom while keeping $[256, D_{v,\text{pack}}]$ MXU tiles contiguous. □

2.4.2 Pillar 2: Tile-Grouped Forget Gate and Packed WY Equivalence

Bottleneck 2 arises because the decay cannot be pulled out of key–key inner products when every channel has its own rate. SLA splits each query and key into the same $G = 4$ groups of $b = 64$ channels as the forget gate in Eq. (3), $q_t = [q_{t,1}; \ldots; q_{t,4}]$ and $k_t = [k_{t,1}; \ldots; k_{t,4}]$. Between tokens j and t, group g decays by a single number $\prod_{s=j+1}^{t} \gamma_{s,g} \leq 1$, which moves outside the group's inner product. Stacking group-g queries and keys as $Q_{c,g}, K_{c,g} \in \mathbb{R}^{C\times 64}$ and defining causal masks $(\Psi_g)_{t,j} = \mathbf{1}[t \geq j] \prod_{s=j+1}^{t} \gamma_{s,g}$ and $(\Gamma_g)_{t,j} = \mathbf{1}[t > j]\, \beta_j (\Psi_g)_{t,j}$, both intra-chunk matrices factor across tiles as

$$M_c = \sum_{g=1}^{4} \Gamma_g \odot \big(K_{c,g} K_{c,g}^\top\big), \qquad L_c = \sum_{g=1}^{4} \Psi_g \odot \big(Q_{c,g} K_{c,g}^\top\big),$$

where each term is a $[C, 64] \times [64, C]$ MXU matmul followed by an elementwise scaling by factors no larger than 1, so there is no overflow and no need for VPU sub-blocks. The trade-off relative to GDN's single rate ($G = 1$) is modest: in our sweep (Section 3.5), $G = 4$ lowers in-domain perplexity slightly across all three paired seeds but extrapolates somewhat worse than $G = 1$; architecturally, unpacked GDN ($H_k = 8, G = 1$) provides $H_k G = 8$ decay timescales per layer, whereas sharing key heads by $R = 4$ ($H_k = 2$) with $G = 1$ collapses the layer to only 2 rates; setting $G = 4$ ($b = 64$) restores $H_k G = 8$ (`SLA-R4`) or 16 (`SLA-R2`) independent timescales per layer at zero latency overhead over $G = 1$.

Proof of tile-grouped packed WY equivalence. Consider chunk $c \in \{1, \ldots, T/C\}$ for a single key head $h_k \in \{1, \ldots, H_k\}$ serving $R = H_v/H_k$ packed value heads of total width $D_{v,\text{pack}} = Rd_v$, starting from incoming state $\tilde{S}_0 = S_{c-1} \in \mathbb{R}^{d_k \times D_{v,\text{pack}}}$. Each token $t \in \{1, \ldots, C\}$ has $G = 4$ tile-grouped diagonal decay $F_t = \bigoplus_{g=1}^{4} \gamma_{t,g} I_{64}$ (diagonal entries in $(0, 1]$), query and key $q_t, k_t \in \mathbb{R}^{d_k}$, packed value $v_t \in \mathbb{R}^{D_{v,\text{pack}}}$, and scalar step size $\beta_t \in (0, 2)$ shared across the R packed value heads of key head h_k.

Let $P_{t\leftarrow j} = \prod_{s=j+1}^{t} F_s$ denote the cumulative diagonal transition from step j to step t ($P_{t\leftarrow t} = I_{d_k}$). Unrolling Eq. (4) ($\bar{S}_t = F_t \tilde{S}_{t-1}$, $\hat{v}_t = \bar{S}_t^\top k_t$, $\tilde{S}_t = \bar{S}_t + \beta_t k_t (v_t - \hat{v}_t)^\top$) from $\tilde{S}_0 = S_{c-1}$ gives

$$\tilde{S}_t = P_{t\leftarrow 0} S_{c-1} + \sum_{j=1}^{t} \beta_j \big(P_{t\leftarrow j} k_j\big) u_j^\top, \tag{7}$$

where $u_t = v_t - \hat{v}_t \in \mathbb{R}^{D_{v,\text{pack}}}$ is the post-forget delta residual at local step t. Multiplying $\bar{S}_t = F_t \tilde{S}_{t-1}$ on the left by $k_t^\top$ substitutes $F_t P_{t-1\leftarrow j} = P_{t\leftarrow j}$ into every prior term:

$$u_t^\top = v_t^\top - k_t^\top P_{t\leftarrow 0} S_{c-1} - \sum_{j=1}^{t-1} \beta_j \big(k_t^\top P_{t\leftarrow j} k_j\big) u_j^\top. \tag{8}$$

Whereas per-channel decay ($G = d_k$ in KDA) overflows the naive factorization $e^{\Lambda_{t,i}} \cdot e^{-\Lambda_{j,i}}$ in FP32 once $|\Lambda_{j,i}| > 88.7$ and forces 16-row sub-block VPU pivots, grouping F_t onto $G = 4$ scalar 64-lane tiles gives

$$k_t^\top P_{t\leftarrow j} k_j = \sum_{g=1}^{4} \exp(\Lambda_{t,g} - \Lambda_{j,g}) \langle k_{t,g}, k_{j,g} \rangle, \qquad \Lambda_{t,g} = \sum_{s=1}^{t} \log \gamma_{s,g},$$

where $\Lambda_{t,g}$ is the tile's cumulative log-decay and $q_{t,g}, k_{t,g} \in \mathbb{R}^{64}$ are the tile-g slices of q_t, k_t. Because $\exp(\Lambda_{t,g} - \Lambda_{j,g}) \leq 1$ is formed directly as a scalar mask, this reduces to four direct $[C, 64] \times [64, C]$ MXU matmuls multiplied by $G = 4$ causal scalar decay masks. Because β_j is scalar per key head h_k, the strictly lower-triangular coupling matrix $M_c \in \mathbb{R}^{C\times C}$ defined by

$$(M_c)_{t,j} = \mathbf{1}[t > j]\, \beta_j \big(k_t^\top P_{t\leftarrow j} k_j\big) \tag{9}$$

is shared across all $D_{v,\text{pack}}$ value columns and requires inverting a single $C \times C$ unit lower-triangular system $A_c = (I_C + M_c)^{-1}$ per key head.

Stacking chunk values $V_c \in \mathbb{R}^{C \times D_{v,\text{pack}}}$ (row t equal to $v_t^\top$), decayed keys $\bar{K}_c \in \mathbb{R}^{C \times d_k}$ (row t equal to $(P_{t\leftarrow 0} k_t)^\top$), and residuals $u_{1:C} \in \mathbb{R}^{C \times D_{v,\text{pack}}}$ (row t equal to $u_t^\top$) yields β-scaled WY factors and chunk writes:

$$U_c = \operatorname{diag}(\beta_c) A_c V_c \in \mathbb{R}^{C \times D_{v,\text{pack}}}, \qquad W_c = \operatorname{diag}(\beta_c) A_c \bar{K}_c \in \mathbb{R}^{C \times d_k},$$
$$E_c = \operatorname{diag}(\beta_c)\, u_{1:C} = U_c - W_c S_{c-1}.$$

Letting the carried keys $\hat{K}_c \in \mathbb{R}^{C \times d_k}$ have row j equal to $(P_{C\leftarrow j} k_j)^\top$ and $F_c = P_{C\leftarrow 0}$, the end-of-chunk state and chunk increment satisfy

$$\tilde{S}_C = F_c S_{c-1} + \hat{K}_c^\top (U_c - W_c S_{c-1}), \qquad \Delta S_c = \tilde{S}_C - F_c S_{c-1} = \hat{K}_c^\top E_c \in \mathbb{R}^{d_k \times D_{v,\text{pack}}}.$$

(When optional oblique key whitening $\tilde{k}_j$ is enabled in `SPD-SLA-W`, Section 4, replacing the write key k_j with $\tilde{k}_j$ in M_c and $\hat{K}_c$ yields the identical WY dual equations.) Finally, the chunk-level velocity $Z_c = \mu \kappa_c Z_{c-1} + \Delta S_c$ and $S_c = F_c S_{c-1} + Z_c$ (Eq. (6), with reset indicator $\kappa_c \in \{0, 1\}$ equal to 0 iff chunk c starts a new document) act solely at chunk boundaries, preserving timestep key-value pairing inside every chunk and matching token-by-token decoding that accumulates ΔS_c over C steps before applying Eq. (6). □

Intra-tile conditioning and orthogonal rotation absorption. Because each group's decay is a scalar multiple of the identity, it preserves conditioning within each tile and commutes with rotary rotations inside the group:

Observation 2.2 (Intra-Tile Conditioning and Orthogonal Rotation Absorption)**.** Within each group g, the decay block $\gamma_{t,g} I_{d_k/G}$ has unit condition number $\kappa_2(\gamma_{t,g} I_{d_k/G}) = 1$ (preserving intra-tile cosine similarity) and commutes with any rotation acting within each tile, $\text{Rot} \in \text{SO}(d_k/G)^G$, such as Givens or rotary rotations, leaving the grouped Gram computation of Pillar 2 unchanged.

Proof. For a per-channel diagonal decay $F_t = \operatorname{diag}(f_t)$ with $f_{t,i} \in (0, 1)$ and cumulative log-decay $\Lambda_{t,i} = \sum_{s=1}^{t} \log f_{s,i}$, the chunk decay $F_c = P_{C\leftarrow 0} = \operatorname{diag}\left((\exp(\Lambda_{C,i}))_{i=1}^{d_k}\right)$ has singular values $\exp(\Lambda_{C,i})$ and condition number $\kappa_2(F_c) = \exp(\max_i \Lambda_{C,i} - \min_i \Lambda_{C,i})$; write $\Delta\Lambda_C = \max_i \Lambda_{C,i} - \min_i \Lambda_{C,i}$. In BF16 (7 mantissa bits, machine epsilon $2^{-7} \approx 7.8 \times 10^{-3}$, unit roundoff $2^{-8} \approx 3.9 \times 10^{-3}$, exponent range $\approx [2^{-126}, 2^{127}]$), relative channel scaling across $P_{t\leftarrow 0} P_{j\leftarrow 0}^{-1}$ loses all mantissa significance across fast vs. slow channels once $\Delta\Lambda_C \gtrsim 5.5$ in BF16 ($\kappa_2 \approx 2^8$) or $\Delta\Lambda_C \gtrsim 16$ in FP32 accumulators ($\kappa_2 \approx 8.9 \times 10^6 \approx 2^{23}$), and $\exp(\pm\Lambda_{j,i})$ overflows or underflows FP32 when $|\Lambda_{j,i}| > 88.7$. Under tile-grouped decay $F_t = \bigoplus_{g=1}^{4} \gamma_{t,g} I_{64}$, each 64-channel tile block is a scalar multiple of I_{64} with $\kappa_2(\gamma_{t,g} I_{64}) = 1$, preserving exact intra-tile cosine similarity.

Furthermore, for any 2×2 block-diagonal (Givens/rotary) rotation with per-pair angles $\theta_t = (\theta_{t,1}, \ldots, \theta_{t,d_k/2}) \in \mathbb{R}^{d_k/2}$ indexed by pair $l \in \{1, \ldots, d_k/2\}$,

$$\text{Rot}(\theta_t) = \bigoplus_{l=1}^{d_k/2} \begin{pmatrix} \cos\theta_{t,l} & -\sin\theta_{t,l} \\ \sin\theta_{t,l} & \cos\theta_{t,l} \end{pmatrix} \in \text{SO}(2)^{d_k/2},$$

the cumulative rotation $\prod_{t=1}^{C} \text{Rot}(\theta_t) = \text{Rot}(\Theta_C)$ with $\Theta_t = \sum_{s=1}^{t} \theta_s$ has all singular values equal to 1 ($\kappa_2 = 1$). Since $\text{Rot}(\theta_s)$ commutes with $G = 4$ scalar tile multiples (64 channels = 32 contiguous 2×2 pairs per tile), for tile-grouped F_s we have

$$k_t^\top \Big(\prod_{s=j+1}^{t} F_s \text{Rot}(\theta_s) \Big) k_j = \big(\text{Rot}(\Theta_t)^\top k_t\big)^\top P_{t\leftarrow j} \big(\text{Rot}(\Theta_j)^\top k_j\big).$$

Replacing (q_t, k_t) with the rotated queries and keys $(\text{Rot}(\Theta_t)^\top q_t,\ \text{Rot}(\Theta_t)^\top k_t)$ before the chunk kernel reduces the intra-chunk transition to pure $G = 4$ real tile decay, preserving the unchanged systolic WY kernel up to a final boundary unrotation $\text{Rot}(\Theta_C)$ on outgoing state. □

2.4.3 Pillar 3: Chunk Inversion on the MXU

Existing solvers substitute sequentially inside diagonal blocks on the VPU. SLA instead inverts $I_C + M_c$ by recursive block-doubling down to 2×2 blocks, so the whole inverse uses only matrix products. Because the matrix is lower-triangular, the inverse of two neighboring diagonal blocks can be assembled from the inverse of each block plus one new off-diagonal block, which costs two small matmuls. SLA starts from 2×2 blocks and doubles the block size six times until it covers the whole 128×128 chunk, so every step runs on the MXU (Figure 1e). The result is the exact inverse in FP32, which we call `Exact Inv.`. SLA also offers a cheaper approximation, `1-Step Neumann`, which keeps only the first correction term between blocks; the ablation runs and prefill benchmarks use a matmul-free variant of it. This approximation changes the learned model: in the pre-training ablation, training with it instead of `Exact Inv.` costs 0.22 perplexity, comparing rows `A2` and `A3` in Section 3.5, even though swapping `Exact Inv.` in at test time shifts 64K QA by at most 0.01 points. All multi-scale MoE pre-training and context-parallel runs therefore use `Exact Inv.`, while single-layer and serving benchmarks report both.

Block-doubling and block-Neumann factorization. Let D_{blk} be the block-diagonal part of $I_C + M_c$ (its $C/16$ unit lower-triangular 16×16 diagonal blocks) and let $N = D_{\text{blk}}^{-1}(I_C + M_c - D_{\text{blk}})$ (local to Pillar 3), so that $I_C + M_c = D_{\text{blk}}(I_C + N)$ with N strictly block-lower-triangular and $N^{C/16} = 0$. Algebraically, at $C = 128$, the exact inverse and its 1-step block-Neumann truncation $A_{\text{N1}} \in \mathbb{R}^{C \times C}$ take the form

$$\begin{aligned} A_{\text{N1}} &= (I_C - N)\, D_{\text{blk}}^{-1}, \\ (I_C + M_c)^{-1} &= \sum_{s=0}^{7} (-N)^s D_{\text{blk}}^{-1} = (I_C + N^4)(I_C + N^2)(I_C - N)\, D_{\text{blk}}^{-1}. \end{aligned} \tag{10}$$

Rather than evaluating the expanded polynomial product $(I_C + N^4)(I_C + N^2)(I_C - N)D_{\text{blk}}^{-1}$, which can amplify powers of N in finite precision, our Pallas kernel (`Exact Inv.`) computes $(I_C + M_c)^{-1}$ via the numerically stable **recursive Schur/block-doubling recurrence** in FP32 (`Precision.HIGHEST`):

$$X_{2s} = X_s - X_s L_s X_s \qquad (s = 2, 4, 8, 16, 32, 64),$$

where X_s is the block-diagonal matrix of already-inverted $s \times s$ diagonal blocks of $I_C + M_c$ ($X_{128} = A_c$) and L_s keeps only the strictly lower-left $s \times s$ coupling block inside each $2s \times 2s$ diagonal block (zero elsewhere), so $L_s X_s L_s = 0$ and the update is exact. Concretely, at each step $s \to 2s$, given the inverse of the two $s \times s$ diagonal halves $(X_{\text{top}}, X_{\text{bot}})$ and their strictly lower-left $s \times s$ coupling block L_{off}, the lower-left block of X_{2s} is computed via two MXU matmuls as $-X_{\text{bot}} L_{\text{off}} X_{\text{top}}$.

Variants of `1-Step Neumann` and truncation properties. Two different first-order approximations appear under the name `1-Step Neumann`, and we keep them distinct:

1. the *block* truncation $A_{\text{N1}} = (I_C - N)D_{\text{blk}}^{-1}$ with block width $B_{\text{blk}} = 16$, analyzed below and evaluated in our standalone lag/skew benchmark; and
2. the matmul-free *row-scaled* approximation $A_{\text{N1}}^{\text{row}} \in \mathbb{R}^{C \times C}$, with row-scaling vector $r = (r_1, \ldots, r_C)^\top \in \mathbb{R}^C$, used by the pre-training ablation arms `A3` and `A5` (Table 11b) and the prefill benchmarks:

$$A_{\text{N1}}^{\text{row}} = I_C - \operatorname{diag}(r)^{-1} M_c, \qquad r_i = \max\bigl(1,\, 1.25\, \|(M_c)_{i,:}\|_1\bigr).$$

The exactness property below (zero error for lags $t - j \le 16$) holds for (1) but not for (2). Partition unit lower-triangular $I_C + M_c = D_{\text{blk}} + M_{\text{off}} = D_{\text{blk}}(I_C + N)$, where $D_{\text{blk}} = \operatorname{diag}(D_1, \ldots, D_m)$ consists of $m = C/B_{\text{blk}}$ unit lower-triangular diagonal blocks $D_a \in \mathbb{R}^{B_{\text{blk}} \times B_{\text{blk}}}$ ($a \in \{1, \ldots, m\}$) and M_{off} is zero on and above the block diagonal. Then $N = D_{\text{blk}}^{-1} M_{\text{off}}$ is strictly block-lower-triangular (its (a, b) block satisfies $N_{a,b} = 0$ for $a \le b$, hence $(N^s)_{a,b} = 0$ for all $a - b < s$ and $N^m = 0$). Comparing $(I_C + M_c)^{-1} = \sum_{s=0}^{m-1} (-N)^s D_{\text{blk}}^{-1}$ against $A_{\text{N1}} = (I_C - N)D_{\text{blk}}^{-1} = D_{\text{blk}}^{-1} - D_{\text{blk}}^{-1} M_{\text{off}} D_{\text{blk}}^{-1}$, the residual

$$(I_C + M_c)^{-1} - A_{\text{N1}} = \Bigl(\sum_{s=2}^{m-1} (-1)^s N^s \Bigr) D_{\text{blk}}^{-1}$$

vanishes identically on all block entries (a, b) with $a - b \leq 1$, namely every diagonal block (a, a) and every first sub-diagonal block $(a + 1, a)$. Any two token indices $1 \leq j \leq t \leq C$ with $t - j \leq B_{\text{blk}}$ satisfy $\lceil t/B_{\text{blk}} \rceil - \lceil j/B_{\text{blk}} \rceil \in \{0, 1\}$ and therefore lie in block (a, a) or $(a + 1, a)$, incurring zero truncation error.

In finite-precision BF16/FP32 arithmetic, exact WY training vs. stateful step-by-step decoding achieves measured output KL divergence $< 10^{-6}$ per layer (3.4×10^{-5} across 4-layer BF16 stacks when chunk momentum syncs every C steps), while under `1-Step Neumann` ($B_{\text{blk}} = 16, C = 128$), truncation at lags $t - j > 16$ exhibits relative Frobenius error 1.6×10^{-5} in short-horizon `4L:1A` layers up to 4.9×10^{-2} in long-horizon `15L:1A` layers, yielding single-layer output KL $< 10^{-5}$ against exact recurrent decoding (and zero discrepancy if chunked prefill uses `1-Step Neumann`). Training with the row-scaled approximation (2) nonetheless changes the learned model: in the pre-training ablation (Table 11b, `A2`→`A3`) it raises 11-set perplexity by 0.22. □

2.5 Single-Kernel VMEM Execution and Custom-VJP Backward Pass

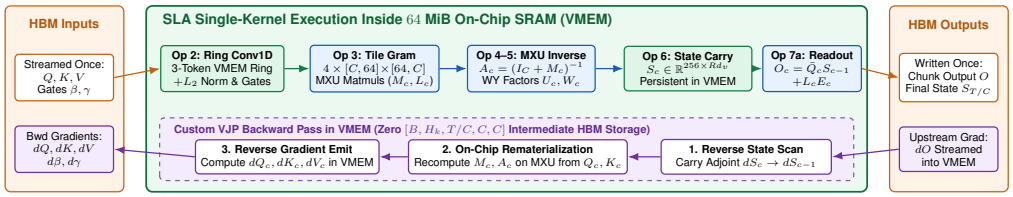

Figure 4: **Single-kernel VMEM execution and custom VJP backward pass in SLA. Top (Forward):** Streamed Q, K, V and gates enter the 64 MiB VMEM scratchpad once; a 3-token ring buffer executes causal `Conv1D` and L_2 normalization (Op 2), followed in-place by 4-tile MXU Gram contractions (Op 3), MXU chunk inversion and WY factors (Ops 4–5), persistent $[256, D_{v,\text{pack}}]$ state updates (Op 6), and chunk readout (Op 7a) before writing O to HBM. **Bottom (Backward):** During the reverse state scan, the custom VJP kernel recomputes intra-chunk matrices M_c, A_c on the MXU directly inside VMEM, eliminating $[B, H_k, T/C, C, C]$ intermediate HBM writes for batch size B.

Algorithm 2 details the custom-VJP backward pass of fused SLA (illustrated in Figure 4, bottom), where for a scalar training loss $\mathcal{L}$ we write $dM_c = \partial\mathcal{L}/\partial M_c$ (and analogously $dQ_c, dK_c, dV_c, \dots$) for the loss gradient (adjoint) with respect to each tensor. Let $\gamma_c \in (0, 1]^{C \times 4}$ stack chunk c's tile decays $\gamma_{t,g}$ (with tile column $\gamma_{c,g} \in (0, 1]^C$), and let $dO_c = \partial\mathcal{L}/\partial O_c \in \mathbb{R}^{C \times Rd_v}$ denote the upstream gradient of chunk c's readout after backpropagating through the per-value-head $\mathrm{RMSNorm}$ and output gate in Eq. (5). For batch size B, a standard autodiff tape stores the intra-chunk matrices $M_c, L_c, A_c \in \mathbb{R}^{C \times C}$ from the forward pass in HBM, consuming $3BH_k(T/C)C^2$ words. Because SLA's Pillar 2 and Pillar 3 evaluate both the $G = 4$ Gram matrices and the chunk inverse $A_c = (I_C + M_c)^{-1}$ as MXU matmuls, our backward kernel instead saves only the chunk-boundary states $S_{c-1} \in \mathbb{R}^{d_k \times Rd_v}$ and recomputes M_c, L_c, A_c, and the unscaled writes $\tilde{E}_c = A_c(V_c - \bar{K}_c S_{c-1}) \in \mathbb{R}^{C \times Rd_v}$ on-the-fly inside VMEM during the reverse scan over $c = T/C, \dots, 1$.

Crucially, whereas unpacked baselines must execute a second sequential VPU triangular solve for the transposed system $(I_C + M_c)^{-\top}$ on every head in the backward pass, SLA already holds the dense inverse A_c in VMEM. Differentiating $E_c = \mathrm{diag}(\beta_c)A_c(V_c - \bar{K}_c S_{c-1}) = \mathrm{diag}(\beta_c)\tilde{E}_c$ through $A_c = (I_C + M_c)^{-1}$ reduces the inverse-solve backpropagation to two MXU matmuls for the scaled adjoint $D_c \in \mathbb{R}^{C \times Rd_v}$ and Gram adjoint $dM_c \in \mathbb{R}^{C \times C}$, shared across all R packed value heads:

$$D_c = A_c^\top \,\mathrm{diag}(\beta_c)\, dE_c \in \mathbb{R}^{C \times Rd_v}, \qquad dM_c = -\,\mathrm{tril}_{-1}\big(D_c \tilde{E}_c^\top\big) \in \mathbb{R}^{C \times C}, \tag{11}$$

where $\mathrm{tril}_{-1}(\cdot)$ and $\mathrm{tril}_0(\cdot)$ mask to strictly $(t > j)$ and weakly $(t \geq j)$ lower-triangular entries, $\mathrm{rowsum}(\cdot)$ and $\mathrm{colsum}(\cdot)$ sum matrix rows and columns, and $p^{\text{in}}_{c,g}, p^{\text{out}}_{c,g} \in (0, 1]^C$ denote the tile-g boundary-decay vectors $(p^{\text{in}}_{c,g})_t = \prod_{s=1}^{t} \gamma_{s,g}$ and $(p^{\text{out}}_{c,g})_t = \prod_{s=t+1}^{C} \gamma_{s,g}$. From D_c, the value and decayed-key gradients follow immediately as $dV_c = D_c$ and $d\bar{K}_c = -D_c S_{c-1}^\top$.

Algorithm 2 Fused SLA Custom-VJP Backward Pass (Per Key Head in VMEM)

Require: Saved boundary states $\{S_{c-1}\}_{c=1}^{T/C}$, chunk inputs $\{Q_c, K_c, V_c, \beta_c, \gamma_c\}$, upstream readout gradients $\{dO_c\}_{c=1}^{T/C}$
Ensure: Input gradients $\{dQ_c, dK_c, dV_c, d\beta_c, d\gamma_c\}_{c=1}^{T/C}$

Initialize reverse carry $dS_{T/C} \leftarrow 0,\ dZ_{T/C} \leftarrow 0 \in \mathbb{R}^{d_k \times Rd_v}$ in VMEM
for $c = T/C$ **down to** 1 **do**
  // 1. On-chip VMEM recomputation (Ops 3–5 on MXU; zero $[C, C]$ HBM reads)
  Recompute $p^{\text{in}}_{c,g}, p^{\text{out}}_{c,g}$, decayed $(\bar{K}_c, \bar{Q}_c, \hat{K}_c)$, $G = 4$ masks (Γ_g, Ψ_g), and M_c, L_c via 64-wide MXU matmuls
  Recompute $A_c \leftarrow (I_C + M_c)^{-1}$ on MXU; form $\tilde{E}_c \leftarrow A_c(V_c - \bar{K}_c S_{c-1})$ and $E_c \leftarrow \mathrm{diag}(\beta_c)\tilde{E}_c$
  // 2. Reverse state and momentum scan adjoints (Ops 6–7a)
  $\widetilde{dZ}_c \leftarrow dZ_c + dS_c; \quad dZ_{c-1} \leftarrow \mu\kappa_c \widetilde{dZ}_c$ ▷ velocity carry ($\kappa_c = 0$ at doc reset)
  $dE_c \leftarrow \hat{K}_c \widetilde{dZ}_c + L_c^\top dO_c \in \mathbb{R}^{C \times Rd_v}; \quad d\hat{K}_c \leftarrow E_c \widetilde{dZ}_c^\top \in \mathbb{R}^{C \times d_k}$ ▷ MXU
  $dL_c \leftarrow \mathrm{tril}_0(dO_c E_c^\top) \in \mathbb{R}^{C \times C}; \quad d\bar{Q}_c \leftarrow dO_c S_{c-1}^\top \in \mathbb{R}^{C \times d_k}$ ▷ MXU
  // 3. Transposed WY inverse and value/key adjoints on MXU (Ops 4–5)
  $D_c \leftarrow A_c^\top \mathrm{diag}(\beta_c)\, dE_c \in \mathbb{R}^{C \times Rd_v}; \quad dV_c \leftarrow D_c; \quad d\bar{K}_c \leftarrow -D_c S_{c-1}^\top$ ▷ MXU
  $dM_c \leftarrow -\mathrm{tril}_{-1}(D_c \tilde{E}_c^\top) \in \mathbb{R}^{C \times C}$ ▷ contracts Rd_v packed value cols on MXU
  $dS_{c-1} \leftarrow F_c^\top dS_c - \bar{K}_c^\top D_c + \bar{Q}_c^\top dO_c \in \mathbb{R}^{d_k \times Rd_v}$ ▷ prev. chunk state carry ($\bar{K}_c^\top D_c = W_c^\top dE_c$)
  // 4. Tile-grouped Q, K and gate adjoints ($G = 4$ tiles of $b = 64$ channels)
  for $g \in \{1, 2, 3, 4\}$ **do**
    $\widetilde{dM}_{c,g} \leftarrow dM_c \odot \Gamma_g; \quad \widetilde{dL}_{c,g} \leftarrow dL_c \odot \Psi_g$
    $dQ_{c,g} \leftarrow \widetilde{dL}_{c,g} K_{c,g} + \mathrm{diag}(p^{\text{in}}_{c,g})(d\bar{Q}_c)_{:,g}$ ▷ 64-wide MXU matmul; $(\cdot)_{:,g}$ is tile-g slice
    $dK_{c,g} \leftarrow (\widetilde{dM}_{c,g} + \widetilde{dM}_{c,g}^\top) K_{c,g} + \widetilde{dL}_{c,g}^\top Q_{c,g} + \mathrm{diag}(p^{\text{in}}_{c,g})(d\bar{K}_c)_{:,g} + \mathrm{diag}(p^{\text{out}}_{c,g})(d\hat{K}_c)_{:,g}$
  end for ▷ concatenate $dQ_c = [dQ_{c,1}, \ldots, dQ_{c,4}],\ dK_c = [dK_{c,1}, \ldots, dK_{c,4}]$
  $d\beta_c \leftarrow \mathrm{rowsum}(dE_c \odot \tilde{E}_c) + \mathrm{colsum}\big(dM_c \odot \sum_{g=1}^{4} \Psi_g \odot (K_{c,g} K_{c,g}^\top)\big)^\top$
  Accumulate columns $d\gamma_{c,g}$ of $d\gamma_c$ via reverse prefix sum over $dM_c, dL_c, d\bar{K}_c, d\bar{Q}_c, d\hat{K}_c$, and $dS_c S_{c-1}^\top$
end for
return $\{dQ_c, dK_c, dV_c, d\beta_c, d\gamma_c\}_{c=1}^{T/C}$

3 Empirical Results

We evaluate SLA on TPU v7x (Ironwood [6, 18]) at four levels: single-layer kernels and multi-generation rooflines (Section 3.1), long-context serving (Section 3.2), context-parallel training and end-to-end 12.5B MoE cube serving (Section 3.3), and multi-scale MoE pre-training throughput and quality (Sections 3.4–3.5). We use two configurations with $d_k = 256$, $G = 4$ decay tiles, chunk size $C = 128$, and $H_v = 8$ value heads: `SLA-R2` shares each query-key head across 2 value heads ($H_q = H_k = 4$, where H_q is the query-head count) and `SLA-R4` across 4 ($H_q = H_k = 2$); layer shapes are written $H_k \times d_k \times Rd_v$. x`L:1A` denotes x linear attention layers per full softmax attention layer, so `0L:1A` is a pure softmax attention Transformer. SLA keeps the baseline's value width, $d_v = 256$ at model width $D = 2048$ and 512 at $D = 1792$, so SLA's BF16 state has twice the entries of FP32 KDA with $d_k = 128$ in the same memory, and the same entries as the 63.1B KDA with $d_k = 256$ in half the memory. `SLA Iso-State` instead halves d_v to match KDA's element count in half the bytes. Part of the measured speedup therefore comes from precision; a BF16-state KDA is only slightly faster than FP32 KDA (Table 6, footnote). All SLA pre-training runs except ablation arms `A3` and `A5` (Table 11b), and `SLA Tree-CP` (Table 8c), use `Exact Inv.`; 32-layer prefill and tree verification (Table 8a–b) use `1-Step Neumann`, and single-layer benchmarks report both (Tables 5–6).

3.1 Kernel Performance and Multi-Generation TPU Rooflines

Fusion and stage breakdown. Because every intra-chunk stage of SLA runs on the MXU, the fused kernel keeps Ops 2–7a in the 64 MiB VMEM scratchpad and recomputes M_c and A_c in the backward pass instead of storing them in HBM (Section 2.5 and Figure 4). Table 5 shows where the time goes. Fusing GPU-shaped KDA gains only $1.03\times$. Reshaping the layer without fusion (unfused `SLA-R4`) gives $1.79\times$, mostly from the Gram/inversion stages and the state scan. Fusion then brings the total to $2.38\times$.

Comparison with open-source TPU baselines. Table 6 compares SLA with MaxText Pallas KDA [1], Tokamax Mosaic KDA [19] (the fastest baseline; it supports only $C \leq 64$), and unpacked scalar GDN. The iso-controls (b) show that the speedup depends on head sharing: without 4:1 sharing,

Table 5: **Seven-operator forward audit (ms) on TPU v7x** ($D = 2048$, $T = 8{,}192$). Unfused columns report isolated per-stage forward medians; fused brackets report single-kernel Ops 2–7a latency.

Stage Latency ($D = 2048, T = 8\text{K}$)	KDA (Unf.)	KDA (Fus.)	SLA-R4 (Unf.)	**SLA-R4 (Fused)**
Op 1: QKV + Gate Proj (MXU)	0.251	0.251	0.168	0.168
Op 2: Causal Conv4 + QK Norm	0.562		0.341	
Op 3–5: Gram + WY Inversion	1.184	Fused	0.418	Fused
Op 6: Chunk State Scan (S_c)	1.248	3.104	0.614	1.112
Op 7a: Intra + State Readout	0.219		0.219	
Op 7b–7c: Out Norm + Out Proj	0.255	0.255	0.245	0.245
Measured Full Layer Fwd (ms)	3.719	3.610	2.005	**1.525**
Analytical Roofline Fwd (ms)	2.890	2.740	1.380	**0.867**
Measured Full Layer F+B (ms)	7.108	6.912	3.982	**2.992**
F+B Speedup vs. KDA (Unf.)	1.00×	1.03×	1.79×	2.38×

Stateless functional harness (`1-Step Neumann` for `SLA-R4`). Fused: one kernel for Ops 2–7a; full NNX module latencies in Table 6c.

Table 6: **Measured TPU v7x (batch size $B = 1$) kernel and full-layer latency (ms)**, 32K temporary HBM (8K is $\sim$0.26×), and F+B speedups normalized to `Open-Source Tokamax Mosaic KDA Base` (1.00×) in each panel across (a) core kernels at $D = 1792$, (b) iso-controlled core kernels at $D = 2048$, and (c1–c2) full linear attention layers (Ops 1–7c, unified harness; Tokamax supports $C \leq 64$ and $d_k \leq 256$); bold marks the `SLA-R4` rows.

Regime	Kernel or Full Layer ($H_k \times d_k \times Rd_v$, Chunk C, State)	8K Fwd	8K Bwd	8K F+B	8K Spd.	32K F+B	32K Spd.	32K Temp HBM
(a) Core Kernel $D = 1792$ (63.1B MoE)	Open-Source MaxText Pallas KDA ($8 \times 256 \times 512, C = 128$, FP32)	3.320	2.000	5.320	0.63×	20.430	0.61×	–
	Open-Source Tokamax Mosaic KDA ($8 \times 256 \times 512, C = 64$, Base)	1.951	1.403	3.354	1.00×	12.449	1.00×	–
	SLA-R2 ($4 \times 256 \times 1024, C = 128$, BF16)	1.960	0.905	2.865	1.17×	10.153	1.23×	–
	SLA-R4 ($2 \times 256 \times 2048, C = 128$, BF16)	**1.371**	**0.671**	**2.042**	1.64×	**6.175**	2.02×	–
(b) Iso-Controls Core Kernel $D = 2048$	Open-Source MaxText Pallas KDA ($8 \times 128 \times 256, C = 128$, FP32)	2.668	1.580	4.248	0.44×	16.455	0.35×	–
	Open-Source Tokamax Mosaic KDA ($8 \times 128 \times 256, C = 64$, Base)	1.072	0.781	1.853	1.00×	5.829	1.00×	–
	SLA Iso-State ($4 \times 256 \times 256, C = 128$, equal state & scan FLOPs)	1.482	0.934	2.416	0.77×	7.565	0.77×	–
	SLA-R2 ($4 \times 256 \times 512, C = 128$, equal $D_v = 2048$; 2× KDA state)	1.571	1.007	2.578	0.72×	8.236	0.71×	–
	SLA Iso-Mixer-Param ($4 \times 256 \times 640, C = 128$, equal mixer params)	1.630	1.056	2.686	0.69×	8.721	0.67×	–
	SLA-R4 ($2 \times 256 \times 1024, C = 128$, equal 1.0 MiB bytes)	**1.051**	**0.652**	**1.703**	1.09×	**5.002**	1.17×	–
(c1) Full Layer $D = 1792$ (Ops 1–7c, Unified)	GDN Scalar Unpacked ($G = 1, 8 \times 256 \times 512, C = 128$, Pallas)	6.305	7.477	13.782	0.96×	48.393	0.88×	9,217 MiB
	Open-Source MaxText Pallas KDA ($8 \times 256 \times 512, C = 128$, FP32)*	6.412	8.043	14.456	0.91×	48.662	0.87×	9,302 MiB
	Open-Source Tokamax Mosaic KDA ($8 \times 256 \times 512, C = 64$, Base)	5.896	7.326	13.222	1.00×	42.378	1.00×	9,483 MiB
	SLA-R2 (Unfused, $G = 1$ Tile, $4 \times 256 \times 1024, C = 128$, BF16)	4.609	6.275	10.884	1.21×	32.459	1.31×	6,922 MiB
	SLA-R2 (Unfused, $G = 4$ Tiles, $4 \times 256 \times 1024, C = 128$, BF16)	4.470	6.088	10.558	1.25×	32.559	1.30×	6,922 MiB
	SLA-R4 (Unfused, $G = 4$ Tiles, $2 \times 256 \times 2048, C = 128$, BF16)	3.605	4.105	7.709	1.72×	26.240	1.62×	5,782 MiB
	SLA-R4 (Exact Inv.) ($2 \times 256 \times 2048, C = 128$, BF16)	**2.882**	**2.686**	**5.568**	2.37×	**20.237**	2.09×	3,738 **MiB**
	SLA-R4 (1-Step Neumann) ($2 \times 256 \times 2048, C = 128$, BF16)	**2.499**	**2.755**	**5.254**	2.52×	**18.786**	2.26×	3,738 **MiB**
(c2) Full Layer $D = 2048$ (Ops 1–7c, Unified)	GDN Scalar Unpacked ($G = 1, 8 \times 128 \times 256, C = 128$, Pallas)	5.006	5.229	10.235	0.60×	31.859	0.76×	4,360 MiB
	Open-Source MaxText Pallas KDA ($8 \times 128 \times 256, C = 128$, FP32)*	4.737	4.794	9.531	0.65×	32.454	0.74×	4,432 MiB
	Open-Source Tokamax Mosaic KDA ($8 \times 128 \times 256, C = 64$, Base)	3.570	2.590	6.160	1.00×	24.060	1.00×	4,727 MiB
	SLA-R2 (Unfused, $G = 1$ Tile, $4 \times 256 \times 512, C = 128$, BF16)	4.285	2.964	7.249	0.85×	25.284	0.95×	4,260 MiB
	SLA-R4 (Unfused, $G = 4$ Tiles, $2 \times 256 \times 1024, C = 128$, BF16)	2.876	2.348	5.224	1.18×	19.680	1.22×	3,472 MiB
	SLA-R2 (1-Step Neumann, $4 \times 256 \times 512, C = 128$, BF16)	2.795	1.970	4.765	1.29×	17.289	1.39×	3,092 MiB
	SLA-R4 (Exact Inv.) ($2 \times 256 \times 1024, C = 128$, BF16)	**2.437**	**2.104**	**4.541**	1.36×	**15.720**	1.53×	2,263 **MiB**
	SLA-R4 (1-Step Neumann) ($2 \times 256 \times 1024, C = 128$, BF16)	**2.119**	**1.867**	**3.985**	1.55×	**14.824**	1.62×	2,263 **MiB**

*Switching MaxText Pallas KDA's state from FP32 to BF16 reduces its latency only slightly, from $14.46/48.66$ ms to $13.91/46.40$ ms ($D = 1792$) and from $9.53/32.45$ ms to $8.80/31.72$ ms ($D = 2048$) at 8K/32K (1.9×–2.6× slower than fused `SLA-R4`). SLA rows change the model (fewer key heads, grouped decay) as well as the kernel; a fused KDA/GDN with the same $d_k = 256$, $R = 4$ sharing was not measured. GDN-2 and Mamba-3 latencies: Section 3.5.

$d_k = 256$ SLA variants run at 0.67×–0.77× of Tokamax, and only `SLA-R4` is faster (1.09×–1.17×). For full layers (c1–c2), fused `SLA-R4` is 1.6×–2.5× faster than Tokamax and needs 52%–61% less temporary HBM; with the exact inverse it is 1.4×–2.4× faster. These comparisons change the model (fewer key heads, BF16 state, grouped decay) as well as the kernel. We did not build a fused kernel for KDA or GDN with the same $d_k = 256$, $R = 4$ head sharing; that control would separate the contribution of the tile-grouped gate from that of head sharing and fusion.

Multi-generation TPU rooflines and FP8 skew. On an identical portable kernel stack, `SLA-R4`'s forward speedup grows with the MXU:VPU ratio from 4.5× on TPU v4 to 6.7× on TPU v7x (Figure 5 and Table 7). Specifically, Table 7 reports measured single-core forward core-kernel latencies ($B = 1, T = 8{,}192, D = 2048$) and out-of-fold leave-one-out (LOO) roofline speedup

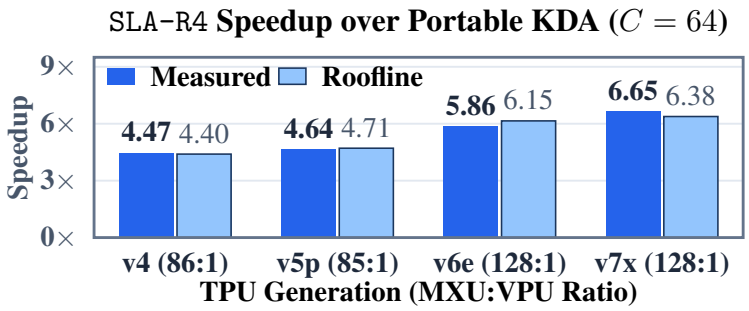

Figure 5: **Forward speedup of SLA-R4 across TPU generations** over a portable FP32 KDA baseline, with leave-one-out roofline predictions (3.1% error). Not comparable to the Tokamax-normalized speedups in Table 6.

predictions across TPU v4, v5p, v6e, and v7x. To keep the software implementation identical across all four hardware generations (since Tokamax Mosaic targets v6e/v7x at $C \leq 64$), Table 7 compares SLA-R4 against the uniform portable JAX/Pallas KDA ($C = 64$, FP32; 6.99 ms on v7x) baseline on every chip, whereas Table 6 compares against optimized open-source MaxText Pallas ($C = 128$, 2.67 ms fwd) and Tokamax Mosaic KDA ($C = 64$, 1.07 ms fwd) on TPU v7x.

On a single TensorCore with systolic tile width $W \in \{128, 256\}$ and single-core MXU and VPU peaks $\mathcal{P}_{\text{MXU}}, \mathcal{P}_{\text{VPU}}$ (per-chip values in Table 1 divided by Cores), let $(\tilde{d}_k, \tilde{D}_{v,\text{pack}})$ denote $(d_k, D_{v,\text{pack}})$ padded to multiples of W, $F_{\text{step}} = 2\max(C, W)\,\tilde{d}_k \tilde{D}_{v,\text{pack}}$ the FLOP count of one per-step scan matmul, $\tau_{\text{pipe}} \approx 1.5$–$2.5\,\mu$s the per-head in-VMEM Pallas scan-pipeline fill/drain overhead (smaller than the top-level XLA lax.scan HBM-barrier cost $\tau_{\text{sync}} \approx 12\,\mu$s per chunk step in Figure 2), $\eta_{\text{sys}} \approx 8.6\%$ the effective single-core scan MXU utilization, γ_{inv} the VPU operation count per entry of the triangular solve (unrelated to the decay rates $\gamma_{t,g}$), and $\Phi_{\text{MXU}}(C) = \mathcal{O}(H_k T C d_k + H_v T C d_v)$ the intra-chunk matmul FLOPs. We model single-core forward latency $T_{\text{fwd}}(C)$ at chunk size C as

$$T_{\text{fwd}}(C) \approx H_k \frac{T}{C}\left(\tau_{\text{pipe}} + \frac{F_{\text{step}}}{\eta_{\text{sys}}\mathcal{P}_{\text{MXU}}}\right) + \frac{\gamma_{\text{inv}} H_k T C}{\mathcal{P}_{\text{VPU}}} + \frac{\Phi_{\text{MXU}}(C)}{\mathcal{P}_{\text{MXU}}}. \tag{12}$$

Across the four hardware generations in Table 7, parameter-free two-cluster comparison directly isolates the 256-tile transition: on 128×128-tile accelerators (0% KDA key-tile padding, $\alpha_{\text{mv}} \approx 85\times$), SLA-R4 already achieves 4.47×–4.64× speedup (TPU v4 and v5p), and stepping to 256 × 256-tile accelerators (50% KDA key-tile padding at $d_k = 128$, $\alpha_{\text{mv}} \approx 128\times$) lifts the speedup by a further 1.3×–1.46× to 5.86×–6.65× (TPU v6e and v7x; reduced-form linearized LOO fit shown for reference in Table 7). Separately, comparing SLA-R4 against the standalone JAX/Pallas $C = 128$ KDA core kernel on the isolated FP8 MXU microbenchmark harness (1.975 ms BF16, 2.168 ms E4M3 FP8, excluding MaxText's segment-reset and L_2-norm wrapper included in the 2.668 ms row of Table 6b) yields a 1.69× forward speedup in native E4M3 FP8 MXU mode (1.280 vs. 2.168 ms at $\alpha_{\text{mv}} = 256\times$), matching BF16's 1.67× (1.181 vs. 1.975 ms) on the same harness.

Table 7: Measured forward core-kernel latency (ms, $B = 1, T = 8{,}192, D = 2048$) vs. portable JAX/Pallas KDA ($C = 64$) and out-of-fold LOO Roofline predictions across four TPU generations.

TPU Generation	Array Tile	MXU:VPU (BF16)	KDA Measured (ms)	SLA-R4 Measured (ms)	Measured Speedup	Out-of-Fold LOO Pred. (Error)
TPU v4	128 × 128	85.9×	14.84	3.32	4.47×	4.40× (−1.6%)
TPU v5p	128 × 128	85.0×	9.19	1.98	4.64×	4.71× (+1.5%)
TPU v6e	256 × 256	127.5×	11.02	1.88	5.86×	6.15× (+5.0%)
TPU v7x	256 × 256	128.2×	6.99	1.05	6.65×	6.38× (−4.1%)

3.2 Long-Context Serving

Prefill latency and prefix caching. In 32-layer hybrid stacks with SwiGLU FFNs [51] (Table 8a), 15L:1A SLA-R4 (one softmax layer per 16) reaches a 131K time-to-first-token (TTFT) of 1.67 s, compared with 2.66 s for the same hybrid layout with KDA and 7.14 s for full grouped-query attention (GQA) [2]. Most of the gain over GQA comes from the hybrid layout itself; the gain attributable to

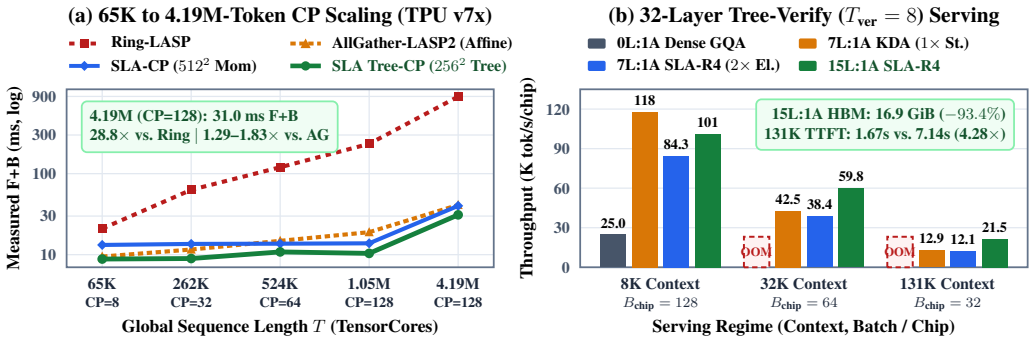

Figure 6: **Context parallelism and serving on TPU v7x. (a)** Forward-backward latency of context-parallel scans across 8–128 TensorCores (Section 3.3). **(b)** Single-chip 32-layer serving throughput and 131K prefill latency (Table 8).

Table 8: **Single-chip 32-layer prefill latency (a)**, 32-layer attention-layer serving throughput with speculative verification of $T_{ver} \in \{8, 16\}$ drafted tokens (b; `OOM` > 190 GiB/chip), and context-parallel F+B scan latency across $P \in \{32, 64, 128\}$ TensorCores (c) on TPU v7x ($D = 2048$).

Regime	Architecture / Method (32L, $D = 2048$)	8K TTFT	32K TTFT	131K TTFT	32K Cache/Lyr
(a) 32L Prefill Full TTFT ($B = 1$)	`0L:1A` Dense GQA ($H_q = 16$, $H_{kv} = 4$ KV heads)	214.7 ms	761.1 ms	7,139 ms (1.00×)	64.0 MiB KV
	`15L:1A` MaxText KDA ($8 \times 128 \times 256$, FP32)	265.7 ms	689.1 ms	2,664 ms (2.68×)	1.0 MiB FP32
	`7L:1A` SLA-R4 ($2 \times 256 \times 1024$, BF16)	204.3 ms	439.8 ms	2,034 ms (3.51×)	1.0 MiB BF16
	`15L:1A` **SLA-R4** ($2 \times 256 \times 1024$, FP8)	203.6 **ms**	416.9 **ms**	1,669 **ms** (4.28×)	0.5 **MiB FP8**
Regime	Attention-Layer Serving (`tok/s/chip`)†	8K, $B=128$	32K, $B=64$	131K, $B=32$	KV+State (131K)
(b) Serving Tree Verify ($T_{ver} \in \{8, 16\}$)	`0L:1A` Dense GQA ($T_{ver} = 8$)	25.0K	OOM	OOM	256.0 GiB
	`7L:1A` MaxText KDA (FP32, $T_{ver} = 8$)	117.7K	42.5K	12.9K	32.88 GiB
	`7L:1A` SLA-R4 (2× El., $T_{ver} = 8$)	84.3K	38.4K	12.1K	32.88 GiB
	`7L:1A` MaxText KDA (FP32, $T_{ver} = 16$)	152.9K	64.3K	OOM	32.88 GiB
	`7L:1A` SLA Iso-State (BF16, $T_{ver} = 16$)	162.1K	68.6K	OOM	32.44 GiB
	`15L:1A` MaxText KDA (FP32, $T_{ver} = 16$)	202.5K	98.0K	34.3K	16.94 GiB
	`15L:1A` SLA Iso-State (BF16, $T_{ver} = 16$)	243.8**K**	119.3**K**	35.2**K**	16.47 **GiB**
	`15L:1A` SLA-R4 FP8 (E4M3, $T_{ver} = 16$)	174.7K	85.6K	33.1K	16.47 **GiB**
Regime	Context-Parallel Scan (F+B ms)	$P=32$ (262K)	$P=64$ (524K)	$P=128$ (1.05M)	$P=128$ (4.19M)
(c) Context Parallelism (32–128 Cores)	Ring-LASP Sequential Ring ($P - 1$ Hops)	63.20	120.12	233.47	894.41
	AllGather-LASP2 Collective Baseline	11.60	14.84	19.00	41.06
	SLA Tree-CP ($\mu = 0.25$, trained, 512^2)	13.59	13.71	13.84	40.08
	SLA Tree-CP ($\mu = 0$, no mom., 256^2)	8.97	10.84	10.40	31.03
	Tree-CP vs. AllGather ($\mu = 0.25$ / $\mu = 0$)	0.85× / 1.29×	1.08× / 1.37×	1.37× / 1.83×	1.02× / 1.32×

†Panel (b) reports attention layers only (2×95 GiB/core HBM cap); with SwiGLU FFNs included at $T_{ver} = 16$, `SLA Iso-State` is 2%–7% faster end-to-end than KDA. Panels (a)–(b) use `1-Step Neumann`; panel (c) uses `Exact Inv.`.

SLA over KDA is 1.6×. A cached SLA state is also much smaller than a KV cache (0.5–1 MiB per layer versus 64 MiB at 32K), which makes multi-turn prefix caching cheap.

Speculative verification. Single-token decoding underuses the MXU, so we verify $T_{ver} \in \{8, 16\}$ drafted tokens per step [35] with the same block inversion (Table 8b). At equal state size (`SLA Iso-State`), SLA is 1.03×–1.22× faster than KDA on the attention layers, and 2%–7% faster once FFNs are included. `SLA-R4`, which holds twice KDA's state elements, is slower than KDA at short context (84.3K vs. 117.7K tok/s at 8K), as expected from its larger state reads. Table 9c reports a 12.5B MoE speculative-decoding setting in which SLA achieves longer greedy acceptance lengths than unclamped KDA on out-of-distribution 32K/64K `HotpotQA` (1.17×–1.33× vs. clamped KDA; 1.06×–1.10× under stochastic sampling, whereas KDA is slightly higher in-domain at 8K).

3.3 Context-Parallel Training and End-to-End 12.5B MoE Cube Serving

Context parallelism with a prefix-tree scan. Head-sharded Ulysses [29] needs at least as many key heads as devices, which SLA's $H_k \in \{2, 4\}$ do not provide, and ring-based LASP [56] passes the state through $P - 1$ sequential hops across P devices. For any gated delta rule with scalar β_t,

a chunk's state update is an affine map $S \mapsto JS + B$ with a $d_k \times d_k$ transition $J \in \mathbb{R}^{d_k \times d_k}$ and offset $B \in \mathbb{R}^{d_k \times D_{v,\text{pack}}}$, so P devices can reduce their local chunks independently and combine them in $\lceil \log_2 P \rceil$ prefix-tree steps [11, 39]. This construction is not specific to SLA, and our KDA baseline uses the same tree. SLA's only advantage here is that one 256×256 transition serves all R value heads of a key head, so fewer transitions are exchanged; chunk momentum ($\mu = 0.25$) doubles the carried state to 512×512. Table 8c reports both settings on 32–128 TensorCores. With momentum, as trained, the tree scan is roughly on par with an AllGather-based LASP-2 baseline [57] ($0.85\times$–$1.37\times$); without momentum it is $1.3\times$–$1.8\times$ faster. Both are much faster than ring-based LASP at large P. In end-to-end 12.5B MoE training at 1M context (Table 9a), `15L:1A SLA-R2` steps 13% faster than KDA using the same tree scan, consistent with its layer-level speedup.

Affine monoid derivation for uncoupled and coupled `SLA Tree-CP`. Shard $p \in \{1, \ldots, P\}$ holds L_{loc} consecutive chunks; quantities of its c-th local chunk ($c \in \{1, \ldots, L_{\text{loc}}\}$) carry the subscript (p, c), with shard input state $S_{p,0}$ and shard output state $S_{p,L_{\text{loc}}}$. Substituting $E_{p,c} = U_{p,c} - W_{p,c} S_{p,c-1}$ (Section 2.4.2) into $\tilde{S}_C = F_{p,c} S_{p,c-1} + \hat{K}_{p,c}^\top E_{p,c}$ gives per-chunk transition $J_{p,c}^{\text{step}} \in \mathbb{R}^{d_k \times d_k}$ and offset $B_{p,c} \in \mathbb{R}^{d_k \times D_{v,\text{pack}}}$ such that, without momentum, each chunk is the affine map

$$\begin{aligned} S_{p,c} &= J_{p,c}^{\text{step}} S_{p,c-1} + B_{p,c}, \\ J_{p,c}^{\text{step}} &= F_{p,c} - \hat{K}_{p,c}^\top W_{p,c}, \qquad B_{p,c} = \hat{K}_{p,c}^\top U_{p,c}; \end{aligned}$$

$J_{p,c}^{\text{step}}$ does not depend on the values, so one transition serves all $D_{v,\text{pack}}$ value columns of a key head (this holds for any gated delta rule, since the transition only involves keys and gates). With $\mu = 0$, the L_{loc} local chunks compose into a single shard map (J_p, B_p) satisfying $S_{p,L_{\text{loc}}} = J_p S_{p,0} + B_p$, where $J_p = \prod_{c=L_{\text{loc}}}^{1} J_{p,c}^{\text{step}}$ (ordered from $c = 1$ on the right to $c = L_{\text{loc}}$ on the left) and B_p is the shard output $S_{p,L_{\text{loc}}}$ obtained from $S_{p,0} = 0$. (With $\mu > 0$, the same construction applies to the lifted maps below.) Writing shard maps as pairs (J_a, B_a) and (J_b, B_b) for shards a and b with composition (apply a, then b)

$$(J_b, B_b) \circ (J_a, B_a) = (J_b J_a,\ J_b B_a + B_b),$$

for any third shard d with map (J_d, B_d), matrix associativity and left-distributivity yield

$$\begin{aligned} \big((J_d, B_d) \circ (J_b, B_b)\big) \circ (J_a, B_a) &= (J_d J_b J_a,\ J_d J_b B_a + J_d B_b + B_d) \\ &= (J_d, B_d) \circ \big((J_b, B_b) \circ (J_a, B_a)\big), \end{aligned}$$

so the P shard maps can be combined by any parallel prefix scan in $\lceil \log_2 P \rceil$ rounds.

When the velocity $Z_c \in \mathbb{R}^{d_k \times D_{v,\text{pack}}}$ is carried across shard boundaries ($\mu > 0$, `SLA Tree-CP+Mom`; shard subscript p dropped, with per-chunk transition J_c^{step} and offset B_c), we have $\Delta S_c = \tilde{S}_C - F_c S_{c-1} = (J_c^{\text{step}} - F_c) S_{c-1} + B_c$. Substituting into Eq. (6) ($Z_c = \mu \kappa_c Z_{c-1} + \Delta S_c$, $S_c = F_c S_{c-1} + Z_c$) defines a lifted $2d_k \times 2d_k$ transition $\tilde{J}_c \in \mathbb{R}^{2d_k \times 2d_k}$ and offset $\tilde{B}_c \in \mathbb{R}^{2d_k \times D_{v,\text{pack}}}$ updating the stacked carry $[S_c; Z_c] \in \mathbb{R}^{2d_k \times D_{v,\text{pack}}}$ via the exact block-affine map

$$\begin{bmatrix} S_c \\ Z_c \end{bmatrix} = \underbrace{\begin{bmatrix} J_c^{\text{step}} & \mu\kappa_c I_{d_k} \\ J_c^{\text{step}} - F_c & \mu\kappa_c I_{d_k} \end{bmatrix}}_{\tilde{J}_c} \begin{bmatrix} S_{c-1} \\ Z_{c-1} \end{bmatrix} + \underbrace{\begin{bmatrix} B_c \\ B_c \end{bmatrix}}_{\tilde{B}_c}. \tag{13}$$

For consecutive shards a and b, the shard-reduced lifted maps $(\tilde{J}_a, \tilde{B}_a)$ and $(\tilde{J}_b, \tilde{B}_b)$ obey the identical associative composition law $(\tilde{J}_b, \tilde{B}_b) \circ (\tilde{J}_a, \tilde{B}_a) = (\tilde{J}_b \tilde{J}_a, \tilde{J}_b \tilde{B}_a + \tilde{B}_b)$, with $[512, 512]$ transitions for $d_k = 256$. □

End-to-end 12.5B sparse MoE 1M-token `SLA Tree-CP`, cube serving, and speculative acceptance. To complement the isolated 128-core `SLA Tree-CP` kernel and single-chip 32-layer dense-FFN serving benchmarks in Table 8, Table 9 benchmarks a full 16-layer 12.5B sparse MoE (0.79B active, $E = 256$ routed experts with $k = 8$ active per token plus one shared expert, $D = 1024$, $D_v = 2048$, vocabulary size $V = 100{,}352$) end-to-end across all 64 chips (128 TensorCores) of a TPU v7x `4x4x4` cube across three workloads:

1. **End-to-End `CP=EP=128` Training Steps up to 1.05M Tokens (Table 9a):** With vocab-chunked cross-entropy across all 24 grid cells, uncoupled `15L:1A SLA Tree-CP R2` ($2\times$-state 4×256^2

Table 9: Measured 64-chip TPU v7x (4x4x4, 128 TensorCores) end-to-end 16-layer 12.5B sparse MoE (0.79B active, 256 experts $k = 8$) (a) fixed-CP=EP=128 F+B latency and tok/s/chip across $T = 128$K–1.05M ($T_{\text{loc}} = T/128$; at $T = 512$K, $T_{\text{loc}} = 4{,}096$ selects a faster XLA EP=128 MoE ragged-all-to-all tile schedule across all linear hybrids relative to $T = 256$K), (b) EP=128 cube verification throughput ($B_{\text{chip}} = 64$, comparing Tree-N1 [= Tree-Neumann-1, $T_{\text{ver}} = 16$] and FP8Fused seq_scan_t16), and (c) real-checkpoint lossless speculative acceptance $\bar{\tau}$ ($K = 16$, $N = 64$, **greedy [stochastic]**).

(a) 12.5**B MoE** CP=EP=128 **F+B**	$T = 128$K (tok/s/ch)	$T = 256$K (tok/s/ch)	$T = 512$K (tok/s/ch)	$T = 1.05$M (tok/s/ch)	1M HBM / Mixer ICI
0L:1A Full-Attn Ring-CP	1,740.4 ms (1,177)	4,797.5 ms (854)	10,712.1 ms (765)	36,725.0 ms (446)	25.8 GiB / 32,512 MiB
7L:1A MaxText KDA Tree ($G = 128$)	1,232.2 ms (1,662)	2,573.3 ms (1,592)	2,704.5 ms (3,029)	7,377.0 ms (2,221)	25.73 GiB / 4,232 MiB
7L:1A **Iso-Param SLA Tree-CP R4**	**1,212.9** ms (**1,689**)	**2,500.8** ms (**1,638**)	**2,542.6** ms (**3,222**)	**7,023.9** ms (**2,333**)	25.92 GiB / 4,736 MiB
15L:1A MaxText KDA Tree ($G = 128$)	1,221.4 ms (1,677)	2,498.2 ms (1,639)	2,249.8 ms (3,641)	5,337.6 ms (3,070)	25.49 GiB / 2,760 MiB
15L:1A Ring-State LASP R2	1,575.5 ms (1,299)	2,744.8 ms (1,492)	2,346.3 ms (3,491)	5,204.6 ms (3,148)	25.46 GiB / 17,392 MiB
15L:1A **SLA Tree-CP R2 NoMom / +Mom**	**1,174**/**1,189** ms	**2,331**/**2,368** ms	**1,965**/**2,003** ms (**5.45**×/ 5.35×)	**4,738**/**4,904** ms (**3,458**/**3,341**)	**25.45** GiB/**2,992** MiB (−**83**%)
(b) 12.5**B MoE** EP=128 **Cube Decode**	8K, $B_{\text{chip}} = 64$ (/cube)	32K, $B_{\text{chip}} = 64$ (/cube)	131K, $B_{\text{chip}} = 64$ (/cube)	8K seq_scan_t16 (/cube)	131K KV+State / Chip
0L:1A Full-Attn GQA (Tree-N1, T_{ver}=16)	15,381 /ch (984.4K)	10,677 /ch (683.4K)	OOM ($B_{\text{chip}} = 8$: 250.7K)	n/a (softmax)	256.0 GiB (OOM)
7L:1A KDA Base FP32 (Tree-N1, T_{ver}=16)	17,212 /ch (1.102M)	16,258 /ch (1.041M)	13,106 /ch (838.8K)	14,180 /ch (907.5K)	32.88 GiB
7L:1A **SLA Iso-State BF16/FP8Fused**	**17,513** /ch (**1.121M**)	**16,544** /ch (**1.059M**)	**13,255** /ch (**848.3K**)	**16,628** /ch (**1,064.2K**)	**32.44**/**32.22** GiB
15L:1A KDA Base FP32 (Tree-N1, T_{ver}=16)	17,513 /ch (1.121M)	17,140 /ch (1.097M)	15,912 /ch (1.018M)	14,366 /ch (919.4K)	16.94 GiB
15L:1A KDA Base FP8 (JAX / FP8Fused)	18,089 /ch (1.158M)	17,792 /ch (1.139M)	16,285 /ch (1.042M)	14,901 /ch (953.6K/1,083K)	16.24 GiB
15L:1A **SLA Iso-State FP8Fused** (T_{ver}=16)	**17,960** /ch (**1.149M**)	**17,510** /ch (**1.121M**)	**16,220** /ch (**1.038M**)	**17,061** /ch (**1,091.9K**, +**18.8**%)	**16.23** GiB (−**93.7**%)
15L:1A **SLA-R2 FP8Fused** (T_{ver}=16)	**17,948** /ch (**1.149M**)	**17,499** /ch (**1.120M**)	**16,200** /ch (**1.037M**)	**16,363** /ch (**1,047.2K**, +**13.9**%)	**16.47** GiB (−**93.6**%)
15L:1A **SPD-SLA-W Polar-2+ FP8** (T_{ver}=16)	**17,994** /ch (**1.152M**)	**17,009** /ch (**1.089M**)	**16,210** /ch (**1.037M**)	12,801 /ch (819.3K JAX)	**16.47** GiB (−**93.6**%)
(c) Lossless Speculative $\bar{\tau}$ ($K = 16$)	Dolma 8K $\bar{\tau}$	Dolma 32K $\bar{\tau}$	SQuAD 32K/64K $\bar{\tau}$	HotpotQA 32K/64K $\bar{\tau}$	Accepted Hotpot@32K / Cube
0.35B KDA → 12.5B 7L KDA (raw)	4.84 [2.53]	3.75 [1.94]	3.72/3.55 [1.49/1.19]	2.63/2.14 [1.36/1.22]	170.7K [88.4K] tok/s
0.35B KDA ($\tau_{\max}$=8K) → 12.5B KDA ($\tau_{\max}$=8K)	**5.30** [**2.49**]	5.14 [3.04]	**4.62**/4.11 [1.65/1.56]	4.42/4.09 [1.72/1.65]	287.6K [111.9K] tok/s
0.35**B SLA** → 12.5**B** 7L **Iso-SLA** (60.4**B**)	4.56 [2.43]	**5.27** [**3.18**]	4.44/**4.20** [**1.69**/**1.71**]	**5.86**/**4.80** [**1.82**/**1.81**] (**1.33**×)	**379.8K** [**118.0K**] (**1.32**×)
0.35**B SLA** → 12.5**B** 7L **Iso-SLA (FP8 PTQ)**	4.67 [2.44]	5.14 [**3.18**]	4.45/4.13 [**1.69**/**1.70**]	**5.78**/**4.66** [**1.82**/**1.82**]	**373.9K** [**117.7K**]
0.35**B SLA** → 63.1**B** 4L **SLA-R4** (120.8**B**)	4.20 [2.46]	2.75 [1.91]	3.66/3.47 [1.71/1.78]	3.17/2.36 [1.32/1.25]	n/a (512c; 10.7× draft span)

Jacobians) completes $T = 1{,}048{,}576$ (per-core local sequence length $T_{\text{loc}} = 8{,}192$) in 4,737.6 ms (3,458 tok/s/chip; 4,904.0 ms and 3,341 tok/s/chip w/ +Mom), running +12.7% faster than our affine-tree port to 1×-state KDA (15L:1A KDA Tree, 8×128^2, 5,337.6 ms, 3,070 tok/s/chip, 2,760 MiB), 7.75× faster than 0L:1A Ring-CP (36,725.0 ms, 446 tok/s/chip), and cutting mixer ICI by −83% vs. 15L:1A Ring-LASP (5,204.6 ms, 2,992 vs. 17,392 MiB/rank). At $T = 524{,}288$ ($T_{\text{loc}} = 4{,}096$/core), 15L:1A SLA Tree-CP R2 reaches 4,170 tok/s/chip (1,965 ms uncoupled, 5.45× vs. 0L:1A Ring-CP 10,712.1 ms; 2,003.1 ms +Mom, 5.35×).

2. **Full 64-Chip EP=128 MoE Cube Serving and FP8Fused Scan (Table 9b):** Across per-chip batch size $B_{\text{chip}} = 64$ (4,096 sequences/cube), our fused Pallas E4M3 kernel (FP8Fused, keeping row-scaled E4M3 DMA, dequantization, 16-token delta scan, and requantization inside VMEM) lifts 16-step sequential scan (seq_scan_t16) at 8K to 1,091.9K tok/s/cube (60.02 ms) on 15L SLA Iso-State (+18.8% vs. 15L KDA FP32 919.4K, +14.5% vs. iso-precision JAX 15L KDA FP8 953.6K, and +14.6% vs. pure-JAX SLA FP8 952.9K) and 1,047.2K–1,052.3K tok/s/cube (+10.3% vs. JAX KDA FP8 and +28.1% vs. pure-JAX SLA FP8 817.6K) on 15L SLA-R2; applying the same FP8Fused recurrent scan kernel to KDA reaches 1,083.0K (within 0.8% of SLA Iso-State, since token-level recurrent decode involves no chunk WY solve). With parallel Tree-Neumann-1 ($T_{\text{ver}} = 16$), raw verification reaches 1.121M–1.158M tok/s/cube at 8K and sustains 1.037M–1.042M tok/s/cube in 16.24–16.47 GiB/chip (−93.6% vs. 0L:1A OOM) at 131K.

3. **Real-Checkpoint Lossless Speculative Acceptance incl. Doubly-Clamped KDA (Table 9c):** Under exact lossless rejection [35] (draft length $K = 16$, $N = 64$ prompts/suite), zero-clamp 0.35B SLA → 12.5B 7L SLA-R4 (iso-param) (60.4B) lifts greedy acceptance $\bar{\tau}$ (mean accepted tokens per draft step) on 32K/64K HotpotQA from 2.63/2.14 (unclamped KDA) to 5.86/4.80 (2.23 × /2.24×; 5.78/4.66 in FP8 PTQ), and still beats *doubly-clamped* KDA (0.35B KDA TrainMaxTau=8K → 12.5B 7L KDA with retention-horizon cap $\tau_{\max} = 8$K, Section 3.5: 4.42/4.09 greedy, 1.72/1.65 stochastic) by 1.33 × /1.17× greedy (1.06 × /1.10× stochastic) on OOD 32K/64K HotpotQA (379.8K vs. 287.6K greedy [1.32×] and 118.0K vs. 111.9K stochastic accepted tok/s/cube at 32K, while KDA edges in-distribution 8K Dolma 4.84–5.30 vs. 4.56).

3.4 MoE Pre-Training Setup, Architecture, and Throughput

All pre-training runs within each cohort share identical optimizer, tokenization, MoE routing, and batch-ordering configurations so that the attention layer design is the sole independent variable:

- **Optimizer and Schedule:** AdamW [38] (moment coefficients $(0.9, 0.95)$, numerical stability constant $\epsilon = 10^{-8}$, weight decay 0.1 with embedding/norm exclusion), global gradient norm clip 1.0, router load-balancing weight 0.01, z-loss 10^{-5}, peak learning rate 3.0×10^{-4} with 5% linear warmup and cosine decay to 3.0×10^{-5} (plus tuned LR $= 4.5 \times 10^{-4}$ with RoPE base frequency $\theta_{\max} = 5 \times 10^5$ [55] for fair `OL:1A` GQA+RoPE baselines across all scales). `+MS` denotes multi-scale depth-staggered decay bias initialization across linear layers.
- **Corpus, FP8 State PTQ, and QA Scoring Protocol:** Real Dolma 3.5 [54, 60, 61] streams (training sequence length $T_{\text{train}} = 8{,}192$) with deterministic `OlmoCoreOrderSampler` batch parity and document-boundary EOS resets (`eos_id = 100257`). In-loop validation tracks geometric-mean PPL across all 11 OLMo v3 validation sets. Post-training evaluation covers: (1) a disjoint 16.78M-token (512×32K) Dolma 3 split (`4K, 8K, 16K, 32K` full and `Tail-4K` PPL; true row-scaled E4M3 FP8 recurrent-state PTQ shifts 8K and 32K NLL by only $+0.00001$ and $+0.00006$ nats with 0.00% NIAH loss across SLA checkpoints, whereas an earlier post-layer activation-quantization probe altered 32K Tail PPL by -0.07 to -0.54 and is excluded from state-PTQ claims), (2) multi-needle Passkey (`MK-NIAH` $4/8$-Needle, $N = 256$–512 evaluation items), (3) 9-task zero-shot `lm-eval` [17] across 11,793 items (1 SE ≈ 0.65 pts), and (4) real-text 4-candidate `RULER` QA [28] on `SQuAD` [49] and `HotpotQA` [69] at 8K, 32K, 64K (25% chance): unless labeled $N = 512$ (1 SE $\approx$ 2.0 pts), tables report raw log-likelihood `cand4_acc` at $N = 128$ (1 SE ≈ 4.0 pts; byte-length-normalized `cand4_acc_norm` agrees within ± 2–5 pts, e.g. $82.0/76.6\%$ vs. $76.6/82.0\%$ on clamped 63.1B KDA).
- **Four Sparse MoE Scales ($10.7\times$ Active Span; 20 tok/active param $\approx 1\times$ Chinchilla [27]):** Across two MXU-aligned backbone width tiers ($D = 1024$ and $D = 1792$), we evaluate (1) 0.35**B active /** 6.5**B total** (decoder depth $L = 10$, $D = 1024$, `4L:1A`, $k = 4$, 15.10B and 30.20B tok), (2) 0.79**B active /** 12.5**B total** ($L = 16$, $D = 1024$, `7L:1A`, `15L:1A`, & `Pure-Linear` ∞:1, $k = 8$, 15.10B and 60.40B tok), (3) 1.80**B active /** 31.5**B total** ($L = 15$, $D = 1792$, `4L:1A`, $k = 8$, 5.03B and 36.20B tok $\approx 1.0\times$ Chinchilla), and (4) 3.74**B active /** 63.1**B total** ($L = 30$, $D = 1792$, `4L:1A` & `15L:1A`, $k = 8$, 5.03B/30.2B RoPE on 64 chips and 120.80B NoPE on 512 chips).

Hybrid `LatentMoE` backbone and multi-scale parameterization. All four pre-training scales are instantiated from a single TPU-aligned hybrid linear-softmax `LatentMoE` decoder family implemented in open-source MaxText [1] (Table 10). Every weight matrix (embeddings, projections, routers, and experts) uses fixed truncated-normal initialization (standard deviation $\sigma_{\text{init}} = 0.02$, clipped at $\pm 3\sigma_{\text{init}}$), and each decoder block applies four $\mathrm{RMSNorm}$ layers ($\epsilon = 10^{-6}$): pre- and post-sublayer norms around both the sequence mixer and the FFN, with post-norms applied to the sublayer output prior to residual addition, followed by a final stack $\mathrm{RMSNorm}$.

Layer 0 uses a dense SwiGLU FFN [51] of intermediate width F_{dense} with no routed experts to prevent early-layer routing collapse. Layers $1, \ldots, L-1$ pair one full-width shared SwiGLU expert of intermediate width F_{sh} with `LatentMoE`: a shared linear down-projection compresses the normalized residual stream from $\mathbb{R}^D$ to an MXU-aligned latent width $\ell < D$, $E = 256$ routed SwiGLU experts (intermediate width $F_{\text{moe}} = 2D$) execute inside $\mathbb{R}^\ell$, and a shared up-projection maps the weighted expert sum from $\mathbb{R}^\ell$ back to $\mathbb{R}^D$. Routing scores are still computed on the uncompressed residual in $\mathbb{R}^D$ in FP32, using top-k softmax weights rescaled to sum to k together with document-pool routing (restricting each document to a pool size in $[k, E]$). Along depth, each $(x{+}1)$-layer cycle places linear delta mixers (4-tap causal depthwise `Conv1D` with SiLU, L_2-normalized queries/keys, low-rank decay and sigmoid output-gate projections of rank r_{proj}, and per-group output $\mathrm{RMSNorm}$ at $\epsilon = 10^{-5}$) on the first x layers and gated Grouped-Query Attention (GQA) [2, 47] ($H_q = 8, H_{kv} = 4$) on every $(x{+}1)$-th layer, with QK-$\mathrm{RMSNorm}$, a $2d_{\text{head}}$-widened query projection (for per-head dimension d_{head}) split into attention queries and an elementwise sigmoid output gate, and either RoPE ($\theta_{\max} = 5 \times 10^5$) [55] or NoPE with Scalable Softmax [40].

Rather than choosing intermediate widths that pad 256×256 systolic tiles, we step across two 256-divisible width tiers ($D = 1024 = 4 \times 256$ for Rungs 1–2; $D = 1792 = 7 \times 256$ for Rungs 3–4, a $1.75\times$ width step) and scale depth L, active expert count k, and shared-expert width F_{sh} within each tier:

1. **Embeddings and untied LM head ($2VD$ linear scaling):** All four scales share a fixed vocabulary $V = 100{,}352$ with untied $V \times D$ input embedding and $D \times V$ output LM-head matrices.

Input token embeddings are multiplied by $\sqrt{D}$ and passed through $\mathrm{RMSNorm}(D)$ before Layer 0 so initial residual-stream RMS remains invariant across D under fixed $\sigma_{\text{init}} = 0.02$ initialization. Because V is fixed, combined embedding and LM-head parameters $2VD$ grow strictly linearly with D (205.5M $\rightarrow$ 359.7M, 1.75×), whereas non-embedding capacity grows with both width and depth; accordingly, $2VD$ drops monotonically from 46.6% of active (2.71% of total) parameters at Rung 1 to 9.6% (0.57%) at Rung 4.

2. **Sublayer widths across tiers (D : 1024 $\rightarrow$ 1792):** Stepping from Tier A to Tier B scales the `LatentMoE` bottleneck 1.5× (ℓ : 512 $\rightarrow$ 768, keeping $\ell/D \in [0.43, 0.50]$ so expert All-to-All dispatch and grouped-matmul K-tiles stay compressed and 256-aligned) and routed expert width 1.75× ($F_{\text{moe}} = 2D$: 2048 $\rightarrow$ 3584). In parallel, global GQA head dimension (d_{head} : 128 $\rightarrow$ 256), linear value width (D_v : 2048 $\rightarrow$ 4096), low-rank gate bottleneck (r_{proj} : 256 $\rightarrow$ 512), and output $\mathrm{RMSNorm}$ groups (8 $\rightarrow$ 16, fixing 256 channels/group) all double (2×). Baseline KDA doubles d_k : 128 $\rightarrow$ 256 ($H_k = 8$), whereas `SLA-R2` ($H_k = 4$) and `SLA-R4` ($H_k = 2$) hold $d_k = 256$ fixed at every scale to preserve unpadded 256 × 256 MXU state updates while doubling packed value width Rd_v.
3. **Within-tier depth, sparsity, and iso-parameter compensation:** Within $D = 1024$ (Rung 1 $\rightarrow$ Rung 2), depth increases from $L = 10$ to 16, active routed experts double (k : 4 $\rightarrow$ 8), and Layer 0 widens (F_{dense} : $4D \rightarrow 8D$). Within $D = 1792$ (Rung 3 $\rightarrow$ Rung 4), depth doubles (L : 15 $\rightarrow$ 30) and shared-expert width doubles (F_{sh} : $D \rightarrow 2D$). Finally, because 4:1 key-head sharing in `SLA-R4` reduces per-layer linear mixer parameters (7.63M $\rightarrow$ 6.31M at $D = 1024$; 27.05M $\rightarrow$ 20.48M at $D = 1792$), `Iso-Param SLA-R4` widens F_{sh} by +256 (1024 $\rightarrow$ 1280 on Rungs 1–2; 1792 $\rightarrow$ 2048 on Rung 3) to match baseline parameter count.

Table 10: Exact model architecture and multi-scale parameterization across our four hybrid `LatentMoE` pre-training rungs ($V = 100{,}352$ untied vocabulary, $E = 256$ routed experts on layers $1, \ldots, L-1$, dense SwiGLU on layer 0).

Architecture / Scaling Parameter	**Rung 1**	**Rung 2**	**Rung 3**	**Rung 4**
Cohort Label (Active / Total)	0.35B/6.5B*	0.79B/12.5B	1.80B/31.5B*	3.74B/63.1B
Exact Total Params (KDA / `SLA-R2`)	7.58B/7.57B	12.49B/12.49B	30.58B/30.52B	63.13B/63.02B
Exact Active: incl. $2VD$ ($1VD$ FLOP)*	0.44B (0.34B)	0.79B (0.69B)	1.91B (1.73B)	3.74B (3.56B)
Decoder Depth L & Hybrid Cycle	10 (4L: 8+2)	16 (7L: 14+2; 15L)	15 (4L: 12+3)	30 (4L: 24+6; 15L)
Model Width D (4 × RMSNorm/block)	1024 (4×256)	1024 (4×256)	1792 (7×256)	1792 (7×256)
Embeddings + Untied LM Head ($2VD$)	205.5M ($\sqrt{D}$+Norm)	205.5M ($\sqrt{D}$+Norm)	359.7M ($\sqrt{D}$+Norm)	359.7M ($\sqrt{D}$+Norm)
$2VD$ Share of Active / Total Params	46.6% / 2.71%	26.0% / 1.65%	18.8% / 1.18%	9.6% / 0.57%
Layer 0 Dense SwiGLU Width F_{dense}	4096 ($4D$)	8192 ($8D$)	14,336 ($8D$)	14,336 ($8D$)
`LatentMoE` Bottleneck ℓ (ℓ/D)	512 ($0.50D$)	512 ($0.50D$)	768 ($0.43D$)	768 ($0.43D$)
Routed Experts (E, Top-k, F_{moe})	256, k=4, 2048	256, k=8, 2048	256, k=8, 3584	256, k=8, 3584
Shared Expert F_{sh} (Base / `Iso-R4`)	1024 / 1280	1024 / 1280	1792 / 2048	3584 ($2D$)
Global Gated GQA (H_q/H_{kv}, d_{head})	8/4, d_{head}=128	8/4, d_{head}=128	8/4, d_{head}=256	8/4, d_{head}=256
Linear D_v, Gate Rank r_{proj}, Norm Grps	2048, r_{proj}=256, 8	2048, r_{proj}=256, 8	4096, r_{proj}=512, 16	4096, r_{proj}=512, 16
Linear Geometry ($H_k \times d_k \times Rd_v$)	KDA 8×128×256; `R2` 4×256×512; `R4` 2×256×1024		KDA 8×256×512; `R2` 4×256×1024; `R4` 2×256×2048	
Mixer Params / Layer (KDA vs. `R4`; GQA)	7.63 vs. 6.31M; 4.20M	7.63 vs. 6.31M; 4.20M	27.1 vs. 20.5M; 14.7M	27.1 vs. 20.5M; 14.7M
Pre-Training Horizons Studied	15.1B, 30.2B	15.1B, 60.4B (+32K)	5.03B, 36.2B	5.03B, 30.2B, 120.8B

*Cohort labels 0.35B/6.5B and 1.80B/31.5B reflect the 8-layer proxy (6.08B total, 0.38B active) and half-63.1B shorthand; extending Rung 1 to $L = 10$ (2 × `4L:1A` cycles) and Rung 3 to $L = 15$ (3 × `4L:1A` cycles, $F_{\text{sh}} = 1792$) ensures exact cycle divisibility and yields the parameter counts above. Active parameters are shown both with $2VD$ (matching the 0.79B and 3.74B convention) and as FLOP-bearing active parameters ($1VD$, excluding input token lookup; pure non-embedding/non-head active parameters are 0.24B, 0.59B, 1.55B, 3.38B).

End-to-end MoE training throughput. We pre-train four token-choice sparse MoE families (6.5B–63.1B total and 0.35B–3.74B active parameters; 256 routed experts, top-8 routing, one shared expert [16, 52]) with AdamW [38] on packed Dolma 3.5 [54, 60, 61] at $T_{\text{train}} = 8{,}192$, using 64–512 TPU v7x chips per run across a 2,048-chip cluster (Table 15). Relative to MaxText Pallas KDA (used because open-source Tokamax Mosaic KDA restricts $C \leq 64$ and omits packed multi-document segment resets; even against an idealized zero-overhead Tokamax $D = 1792$ kernel from Table 6c1, 63.1B `SLA-R4` remains 1.13× faster end-to-end), SLA trains 8%–24% faster across matched scales (1.10×–1.18× at 6.5B/31.5B, 1.16×–1.24× at 12.5B in Table 11a, and 1.11×–1.15× at 63.1B on 512 chips in Table 12). At native $T_{\text{train}} = 32$K with 4.19M tokens/step (`EP=128`), SLA fits in per-core

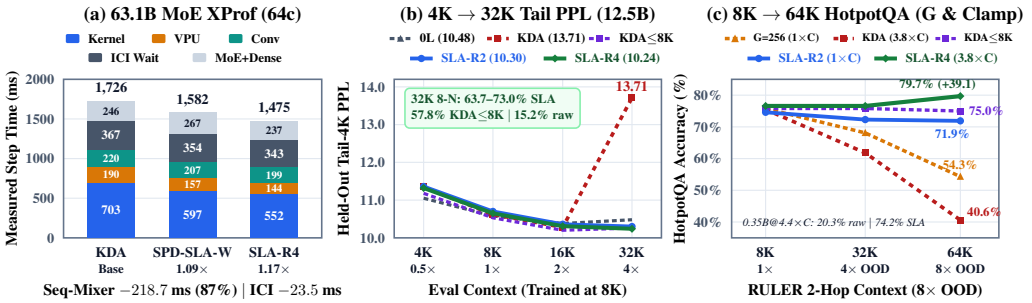

Figure 7: **(a)** Per-step XProf breakdown for 63.1B MoE on 64 TPU v7x chips across `KDA Base`, `SPD-SLA-W (R2)` (`SLA-R2` with optional in-chunk key whitening, Section 4), and `SLA-R4` (Table 13); most of SLA's saving is in linear attention ops. **(b)** Held-out Dolma tail perplexity from 4K to 32K for 12.5B MoEs (Table 11a; seed `s0`). **(c)** Zero-shot 8K → 64K RULER `HotpotQA` for KDA (raw and clamped) and SLA (60.4B-token curves: seed `s0`; 15.1B-token curves: mean of `s0` and `s1`).

Table 11: **Measured** 12.5**B MoE** (0.79**B active,** 16**L,** $D = 1024, D_v = 2048, T_{\text{train}} = 8{,}192$) **pre-training** across (a) 60.40B-token (3.8× Chinchilla [27], 8.39M tok/step) v7x arms and (b) 15.10B-token (1.0× Chinchilla, 2.10M/step) and ∼90-min Iso-Wall ablations ($n \in \{2, 3\}$ random seeds where noted; 4-choice 64K RULER QA, $N = 128$, SE = 4 pts; full controls in Tables 15–16; bold: best within each cohort at equal tokens, with perplexity compared on matched seeds `s0,s1` in (b); GDN-2 uses our unfused TPU port; seeds `s0`/`s1`/`s2` use different data offsets, so perplexities are comparable only on matched seeds).

Cohort	Architecture (7,200 Steps)	Step (ms)	11-Set PPL	8K Tail	32K Tail	32K 8-N	64K SQuAD	64K Hotpot
(a) v7x 60.4B Tok (3.8× Ch., 8.4M/st.)	`0L:1A` Full-Attn GQA+RoPE (`s0`)	2,313.8	16.128	10.57	10.48	12.1%	58.6%	53.1%
	`7L:1A` MaxText KDA Base ($n = 3$)*	2,225 ± 29	15.602 ± 0.582	10.46	12.45 ± 1.11	24.6%	54.9%	50.3%
	+Zero-Retrain $\tau_{\max} = 8$K Clamp (`s0`)	–	16.168	10.53	10.26	57.8%	**78.1%**	75.0%
	`7L:1A` KDA + `TrainMaxTau=8K` ($n = 2$)	2,251 ± 11	15.800 ± 0.332	10.46	10.75 ± 0.21	49.6%	75.8%	75.4%
	`7L:1A` SLA-R2 (BF16, $n = 3$)*	1,919 ± 37	15.580 ± 0.508	10.46	**10.19 ± 0.10**	56.6%	75.5%	**77.9%**
	`7L:1A` SLA-R4 (Iso-Param, $n = 3$)*†	**1,800 ± 39**†	15.664 ± 0.414	10.57	10.27 ± 0.04	56.1%	74.7%	74.5%
	`7L:1A` SLA-R4 (−20% mixer, `s0`)	1,819.2	16.200	10.81	10.41	**73.0%**	72.7%	75.0%
	`15L:1A` MaxText KDA Base (FP32, `s0`)	2,152.0	15.498	10.46	11.27	15.2%	62.5%	60.2%
	`15L:1A` SLA-R2 (Uniform, BF16, `s0`)	1,941.3	**15.436**	**10.43**	**10.19**	41.0%	72.7%	70.3%
(b) 15.1**B** (1.0× Ch., 2.1M/st. & 90m Iso)	`7L:1A` MaxText KDA Base ($n = 3$)§	556.0	17.638 (17.893)	11.57	13.28	18.1%	57.3%	56.0%
	`7L:1A` SLA-R2 (Unf., `Exact Inv.`, $n = 3$)§	497.3	17.533 (17.750)	11.54	**11.23**	52.3%	69.5%	73.7%
	`7L:1A` SLA-R2 (Fus., `Exact Inv.`, $n = 2$)	481.8	17.806	11.70	11.36	58.0%	69.1%	74.6%
	`7L:1A` GDN-2 (Unf., 93m, 15.1B, $n = 2$)	778.1	**17.592**	**11.44**	11.86	19.3%	62.1%	62.1%
	`7L:1A` SLA-R2 (Fus., 90m, 23.1B tok)	488.3	16.660 (+53%)	10.84	10.98	52.3%	65.6%	70.3%
	`7L:1A A1`: Wide-Key R=2 KDA ([4, 256, 512], G=256, n=2)	524.8	17.694	11.61	13.39	23.8%	55.9%	54.3%
	`7L:1A A2`: $+G = 4$ Decay (`Exact Inv.`, $n = 2$)	514.1	17.749	11.65	11.34	46.1%	70.7%	73.8%
	`7L:1A A3`: +1-Step Neumann Inv. ($n = 2$)	497.2	17.970	11.69	11.36	53.5%	71.9%	71.9%
	`7L:1A A5`: SLA-R2 (1-Step, $\mu = 0.25$, `s0`)	**467.1**	18.016	11.87	11.45	**60.9%**	**74.2%**	**83.6%**
	`7L:1A A6`: SLA-R2 + FP8 Out-QAT (`s0`)	520.5	18.006	11.80	11.37	59.8%	**74.2%**	75.8%
	`15L:1A` SLA-R2 (Uniform, BF16, $n = 3$)	486.7	17.609	11.60	11.42	34.8%	69.0%	65.1%

*Paired 3-seed differences vs. KDA: `SLA-R2` −0.02 ± 0.11 PPL, `SLA-R4` +0.06 ± 0.26 PPL (within seed noise). †Deferring the fused `SLA-R4` conv `AllReduce` across gradient accumulation gives 1,812 ms (Table 14). §In parentheses: mean over `s0,s1`, matched to the $n = 2$ rows. On `s0` alone: KDA 18.074, `A1` 17.850, `A2` 17.873, fused `SLA-R2` 17.913, `A3` 18.102, `A5` 18.016, `A6` 18.006.

HBM (≤ 94.74 GiB) while KDA (96.74 GiB) and full attention (95.48 GiB) exceed the 95 GiB cap; in 8K → 32K continued pre-training at 12.5B, where all three fit, SLA is 1.08×–1.15× faster than KDA (Table 15).

Hardware XProf operator attribution on 64-chip TPU v7x. XProf attributes most of SLA's step-time saving to linear attention ops (Figure 7a). Table 13 provides the complete hardware `XProf` HLO category breakdown (`xplane` traces captured over steps 15–17 across all 8 TensorCores of host 0, normalized by 24 core-steps to ms/step) for 63.1B MoE (`FSDP=2, EP=64`) on 64 TPU v7x chips. Mapping directly to Figure 7a (where the top bar groups ICI `async-done` wait with MoE/dense matmuls: 366.6+245.8 = 612.4 ms, 354.1+266.8 = 620.9 ms, and 343.1+237.2 = 580.3 ms), the three direct linear attention op categories save −218.7 ms (87.2%) and reduced mixer memory/sync traffic saves −23.5 ms (9.4%) of exposed ICI `async-done` wait, totaling −242.2 ms (96.6%) of the −250.8 ms step-time reduction between `SLA-R4` and KDA Base (with −8.6 ms [3.4%] in MoE/FFN matmul overlap).

Table 12: **Measured** 512**-chip TPU v7x** 63.1**B MoE** (3.74**B active,** $8 \times 8 \times 8$**,** 16.78**M tok/step) at** 120.80**B tokens** ($\approx 1.6\times$ Chinchilla, NoPE $T_{\text{train}} = 8{,}192$; `SSMax` [40] query scale $\alpha \in \{0.25, 0.5\}$, Section 3.5; 64K RULER QA `cand4_acc` at $N = 512$, $^*N = 128$; full 64-chip XProf and RoPE ablations in Tables 15, 13, and 17).

63.1B MoE (120.8B Tok, 512 Chips, NoPE 8K)	Step (ms)	tok/s/ch	11-Set PPL	32K Tail	32K 8-N	64K `SQuAD`	64K `Hotpot`
`4L:1A` MaxText KDA Base ($8 \times 256 \times 512$, FP32)	2,165.1	15,135	12.596	353.4	15.6%	27.7%	27.0%
+Zero-Retrain $\tau_{\max} = 8$K Clamp (0 retr.)	–	–	12.596	**7.70**	48.4%	76.6%*	82.0%*
$+\tau_{\max} = 8$K Clamp + `SSMax` ($\alpha = 0.25$)	–	–	12.596	8.04	74.2%	**90.6%**	**86.5%**
`4L:1A` SLA-R2 ($4 \times 256 \times 1024$, BF16, 0.5× Bytes)	1,952.9	16,779	12.584	13.78	18.6%	55.7%	56.8%
+Zero-Clamp SLA-R2 + `SSMax` ($\alpha = 0.5$)	–	–	12.584	8.45	**84.4%**	–	–
`4L:1A` SLA-R4 ($2 \times 256 \times 2048$, BF16, 0.5× Bytes)†	**1,878.3**	**17,445**	**12.454**	20.47	18.0%	51.6%	45.1%
+Zero-Clamp SLA-R4 + `SSMax` ($\alpha = 0.5$)	–	–	**12.454**	8.41	**83.2%**	75.6%	79.1%

Single seed (`s0`). Full 64K QA breakdown (`SQuAD`, `HotpotQA`) and 64-chip RoPE controls are in Table 17. †Deferring the fused conv `AllReduce` across gradient accumulation yields 1,843 ms.

Table 14 complements this with a symmetric 4-arm `GA=4` (8.39M tok/step) `XProf` audit on 12.5B MoE (`EP=128`): deferring causal-conv weight `AllReduce` outside the `GA=4` microbatch loop yields 0.0% speedup on MaxText KDA (2,118.9 → 2,119.7 ms) while cutting `SLA-R4` (iso-param) to 1,686.0 ms in 3-step `XProf` (1,661.9 ms unprofiled min; $1{,}812.0 \pm 14.1$ ms over full 7,200-step $n = 2$ training runs due to host checkpoint/validation wrapper hooks) via a 1.96× faster Pallas linear attention kernel (528.4 → 269.1 ms).

Table 13: Hardware `XProf` per-step HLO category breakdown (ms/step) on 64-chip TPU v7x 63.1B MoE (`FSDP=2, EP=64`, 2.10M tok/step).

HLO Category (`xplane` Measured)	`4L:1A` KDA Base	`4L:1A` SPD-SLA-W (R2)	`4L:1A` SLA-R4	Δ (SLA-R4 vs. KDA)
Pallas Linear Attention `custom-call`	703.0 ms	597.2 ms	552.0 ms	−151.0 ms (−21.5%)
VPU Gate / Norm `loop fusion`	190.4 ms	156.8 ms	143.7 ms	−46.7 ms (−24.5%)
Causal Short `convolution fusion`	220.0 ms	207.4 ms	199.0 ms	−21.0 ms (−9.5%)
ICI Collective `async-done` Wait	366.6 ms	354.1 ms	343.1 ms	−23.5 ms (−6.4%)
MoE Expert `ragged-dot` + Dense Matmul	245.8 ms	266.8 ms	237.2 ms	−8.6 ms (−3.5%)
Total Steady-State Step Time	**1,725.8** ms	**1,582.3** ms	**1,475.0** ms	**−250.8** ms (**1.170**×)
Steady-State Throughput (`tok/s/chip`)	18,988	20,710 (1.09×)	**22,216** (**1.17**×)	**+3,228** (**+17.0%**)

At 6.3B tokens (RoPE validation checkpoint on the same 64-chip slice), `4L:1A` KDA Base and `4L:1A SLA-R4` reach 9.90 vs. 10.15 8K Dolma PPL, 10.27 vs. **9.75** 32K Tail PPL, and 14.5% vs. **37.5**% 32K 8-Needle accuracy.

Table 14: Symmetric 4-arm `GA=4` (8.39M tok/step) hardware `XProf` leaf-op breakdown (ms/step) on 64-chip TPU v7x 12.5B MoE (`EP=128`).

12.5B `GA=4` Arm (3 profiled steps)	Step Mean (ms)	`linear_mixer _pallas`	`short_conv _fusion`	Mixer + Conv	`full_attn _splash`	`all_gather _rs`	`moe_ragged _a2a`
`7L:1A` MaxText KDA (default AR)	2,118.9	508.1	91.7	599.8	69.8	191.9	264.7
`7L:1A` MaxText KDA + `defer_ar`	2,119.7 (1.00×)	528.4	91.6	620.0	70.0	191.6	234.3
`7L:1A` **SLA (Fused) R2+Mom +** `defer_ar`	**1,797.4** (1.18×)	**340.9** (−187.5)	**65.4**	**406.3** (−213.7)	69.9	**152.5**	282.5
`7L:1A` **Iso-Param SLA (Fused) R4 +** `defer_ar`	**1,686.0** (1.26×)	**269.1** (1.96×)	**61.7**	**330.8** (−289.2)	69.7	**152.6**	265.3

3.5 Pre-Training Quality and Long-Context Extrapolation

In-domain quality. At 12.5B parameters and 60.4B tokens with three paired seeds (Table 11a), `SLA-R2` and `SLA-R4` are within seed noise of KDA on 11-set Dolma perplexity and on zero-shot `lm-eval` [17] (Table 16); single-seed 63.1B runs are also close (Table 12). The 15.1B-token cohort (Table 11b) shows where this parity comes from. Seeds `s0`, `s1`, and `s2` use different data offsets that shift perplexity by up to 0.9, so we compare only on matched seeds. First, over three seeds, `SLA-R2` (`Exact Inv.`, BF16 state, $\mu = 0.25$) improves on KDA Base on every seed (17.533 vs. 17.638; paired difference -0.10 ± 0.08) and matches the same-kernel $G = 4$ FP32 control without momentum (17.538, Table 18), so BF16 state and momentum cost no measurable quality. On seeds `s0,s1`, widening KDA's keys to $d_k = 256$ with $R = 2$ head sharing (`A1`, $[4, 256, 512], G = 256$: 524.8 ms, 17.694 PPL) isolates head sharing from decay grouping; moving from $G = 256$ to $G = 4$ with MXU block-doubling at identical $R = 2$ sharing cuts step time by a further −27.5 ms (497.3 ms unfused, −15.5 ms further to 481.8 ms fused) while staying within 0.06–0.11 PPL of `A1` and lifting 32K 8-Needle from 23.8% to 52.3%–58.0% and 64K `HotpotQA` from 54.3% to 73.8%–74.6%. The

Table 15: Matched RoPE ($\theta_{\max} = 5 \times 10^5$) 4-rung scaling ladder (0.35B to 63.1B MoE; see Table 11 for 12.5B Rung 2), 0.35B 3-seed & 30.2B (4.4× Chinchilla), 31.5B at 36.2B tok (1.0× Chinchilla), 12.5B 8K → 32K continued pre-training, and 12.5B `15L:1A` dilution control on 64-chip TPU v7x (§under `remat_32k` with 4-step gradient accumulation [`GA=4`], KDA/`0L:1A` hit `NaN` at steps 1–43 and required `noremat` reruns, whereas 2-seed `SLA-R4` (iso-param) completed all 4,300 steps at 3,040 ms with 12.54 → 12.08 Tail PPL, 59.2%/48.2% 32K 4/8-N, and 58.5/62.6% N=512 64K QA, though router-remat spikes raised 8K PPL to 17.99).

MoE Scale (Active / Total, Tokens)	Architecture	Step (ms)	11-Set PPL (or 32K-Ext CE)	8K Tail PPL	32K Tail PPL	32K 4-N / 8-N	8-N Gain vs. KDA
0.35B/6.5B (15.1B tok, 2.2× Chin., `4L:1A` interval 5, `s0`)	`0L:1A` GQA+RoPE (LR $= 4.5 \times 10^{-4}$)	361.6	19.582	12.83	13.18	18.4%/12.9%	0.91×
	`4L:1A` MaxText KDA Base (FP32)	351.8	**18.573**	**12.42**	15.72	32.0%/14.1%	1.00×
	`4L:1A` KDA + TrainMaxTau ($\tau_{\max} = 8$K, FP32)	351.2	19.035 (`s1`; 2.54 CE)	12.43	**12.36**	**56.6%/48.4%**	**3.43×**
	`4L:1A` **SLA-R2** (BF16)	**320.0**	18.773	12.62	**12.30**	**63.3%/55.1%**	**3.91×**
	`4L:1A` **SLA (Fused) R2+Mom** (BF16)	**308.5**	18.890	12.60	**12.28**	**60.9%/51.6%**	**3.66×**
	`4L:1A` **Iso-Param SLA-R4** (BF16)	**302.4**	19.093	12.83	**12.51**	**57.8%/48.8%**	**3.46×**
	`4L:1A` **SPD-SLA-W** (BF16)	344.8	19.104	12.64	**12.32**	**66.4%/56.2%**	**3.99×**
0.35B/6.5B 3-Seed & 30.2B Tok (4.4× Chinchilla)	3-Seed `4L:1A` MaxText KDA (15.1B, $n = 3$)	352.0	18.880 ± 0.267	**12.59**	15.89	32.9%/19.9%	1.00×
	3-**Seed** `4L:1A` **SLA-R2** (15.1**B,** $n = 3$)	**314.8**	19.077 ± 0.265	12.72	**12.77**	**49.9%/37.6%**	**1.89×**
	`0L:1A` GQA+RoPE (30.2B, 4.4× Chin.)	361.5	17.104	11.43	12.02	25.0%/15.2%	0.90×
	`4L:1A` MaxText KDA (30.2B, 4.4× Chin.)	351.5	**16.161**	11.03	620.87	31.6%/16.8%	1.00×
	`4L:1A` **SLA-R2** (30.2**B,** 4.4× **Chin.)**	**315.0**	16.250	**11.02**	**13.53**	35.2%/21.1%	1.26×
	`4L:1A` **Iso-Param SLA-R4** (30.2**B,** 4.4× **Chin.)**	**297.9**	16.326	11.15	**11.09**	**59.4%/45.3%**	**2.70×**
1.80B/31.5B (5.03B & 36.2B tok, `4L:1A`)	`0L:1A` GQA+RoPE (5.03B, LR $= 4.5 \times 10^{-4}$)	867.5	**20.380**	**12.98**	13.73	24.6%/14.8%	1.15×
	`4L:1A` MaxText KDA Base (5.03B, FP32)	807.7	21.228	13.99	47.16	28.9%/12.9%	1.00×
	`4L:1A` **SLA (Fused) R2+Mom** (5.03B, BF16)	**747.2**	21.439	14.19	**13.88**	**48.8%/33.2%**	**2.57×**
	`4L:1A` **SLA (Fused) R4+Mom** (5.03B, BF16)	**707.4**	21.709	14.44	**14.20**	**59.8%/46.1%**	**3.57×**
	`4L:1A` MaxText KDA (36.2B, 4,300 st., `noremat`)§	3,526	15.264	10.57	51.68	24.6%/12.5%	1.00× (39.1/34.4% QA)
	`4L:1A` KDA + TrainMaxTau (36.2B, $\tau_{\max} = 8$K, `noremat`)§	3,307	**15.211**	**10.36**	**16.77**	**27.3%/14.8%**	**1.18× (59.6/53.3%** N=512 QA)
3.74B/63.1B (5.03B tok, RoPE)	`0L:1A` GQA+RoPE (LR $= 4.5 \times 10^{-4}$)	1,749.6	**19.373**	**11.66**	11.94	25.0%/11.7%	0.81×
	`4L:1A` MaxText KDA Base (FP32)	1,746.0	20.076	12.25	18.24	23.1%/14.4%	1.00×
	`4L:1A` GDN-2 + RoPE (FP32, `EP=128`, 14,143 tok/s/ch)	2,316.9	19.731	12.27	14.22	25.8%/13.3%	0.92×
	`4L:1A` **SLA (Fused) R2+Mom** (BF16, 20,386 tok/s/ch)	**1,607.4**	20.201	12.37	**12.47**	**35.2%/22.3%**	**1.55×**
	`4L:1A` **SLA (Fused) R4+Mom** (BF16, 20,847 tok/s/ch)	**1,571.8**	20.423	12.50	**12.80**	**27.0%/13.7%**	0.95×
	`15L:1A` **SLA (Fused) R2+MS** (BF16)	**1,547.2**	20.890	13.17	**13.01**	**59.4%/48.4%**	**3.36×**
	`15L:1A` **SPD-SLA-W+MS** (BF16, 20,473 tok/s/ch)	**1,600.6**	20.506	12.49	**12.42**	**41.4%/27.0%**	**1.88×**
0.79B/12.5B 8K → 32K Ext. (5.0–10.1B at 32K)	`0L:1A` GQA+RoPE (32K-ext, 1.00× speed)	4,697.0	2.2089 CE	11.29	10.95	25.4%/11.7%	0.18×
	`7L:1A` MaxText KDA (32K-ext, 1.53× vs. `0L`)	3,061.4	2.2006 **CE**	11.20	**10.78**	71.5%/63.3%	1.00×
	`7L:1A` **SLA-R2** (32**K-ext,** 1.66× **vs.** `0L`)	**2,828.0**	2.2048 CE	11.26	10.82	**73.8%/64.1%**	**1.01×**
	`7L:1A` **Iso-Param SLA-R4** (32**K-ext,** 1.77×)	**2,661.0**	2.2255 CE	11.46	11.01	60.9%/49.6%	0.78×
	`7L:1A` MaxText KDA (60.4B → 32K, +5.03B)	3,103.6	2.0907 **CE**	**10.38**	**9.98**	55.5%/47.3%	1.00×
	`7L:1A` **Iso-SLA (Fused) R4** (60.4**B** → 32**K,** 1.15×)	**2,688.7**	2.1142 CE	10.60	10.15	**64.8%/57.0%**	**1.21×**
	`15L:1A` MaxText KDA (32K-ext, uniform)	2,827.0	2.2092 CE	**11.10**	10.79	56.2%/40.6%	1.00×
	`15L:1A` SLA-R2+MS / SPD-SLA-W+MS (32K-ext)	**2,608 / 2,715**	2.238/2.241	11.38/11.42	11.17/11.18	23.8%/16.0, 12.9%	0.39×
0.79B/12.5B (15.1B & 60.4B `15L:1A` dilution controls)	`15L:1A` MaxText KDA Base (15.1B, FP32, $n = 3$)	556.4	17.658 ± 0.297	**11.64**	11.57	34.0%/21.4%	1.00×
	`15L:1A` **SLA-R2** (15.1B, uniform, BF16, $n = 3$)	**486.7 ± 11.7**	**17.609 ± 0.246**	11.70	**11.42**	**45.1%/34.8%**	**1.63×**
	`15L:1A` **SPD-SLA-W Polar-2+** (15.1B, uniform, `s0`)	509.3	**17.575**	11.65	**11.38**	**50.0%/39.8%**	**1.86×** (1.29× `s0`)
	`15L:1A` SLA-R2+MS (15.1B, staggered, $n = 3$)	492.2	17.882 ± 0.294	11.93	11.72	23.6%/13.0%	0.61×
	`15L:1A` KDA + `TrainMaxTau=8K` (60.4B, FP32, `s0`)	2,199.0	15.579	10.47	**10.30**	37.5%/26.6%	1.75×
	`15L:1A` **SLA-R2+MS** (60.4B, staggered, `s0`)	**1,934.0**	**15.443**	10.61	**10.38**	27.3%/19.9%	1.31×
	`15L:1A` SPD-SLA-W+MS (60.4B, staggered, `s0`)	2,006.0	15.624	10.67	10.49	23.1%/12.1%	0.79×
0.79B/12.5B (15.1B, additional controls & 128c TPU v4)	`7L:1A` Scalar GDN ($G = 1$, FP32, $n = 2$)	555.9	18.142	11.85	**11.51**	– / **60.8%**	**4.50×**
	`7L:1A` KDA + `TrainMaxTau=8K` (FP32, `s0`)	560.6	17.924	11.66	11.60	58.2%/48.1%	3.56×
	`Pure-Linear` (∞:1) KDA / **SLA-R2+MS** (`s0`)	545.8 / **475.6**	18.096 / **17.642**	12.02 / 11.95	12.05 / **11.80**	n.r. (55.5/59.4% `SQuAD`)	n.r. (56.2/53.9% `Hotpot`)
	TPU v4 (128c) `7L:1A` **SLA-R4** (NoPE, $n = 2$)	**992.0** (1.10×)	−0.017 nats vs. KDA	10.18	**10.05**	– / **66.4%**	**2.42×**

optional 1-step Neumann inverse is the one approximation with a clear cost: `A3` reaches 17.970 (+0.22 over `A2`), and on `s0` the fused Neumann arm `A5` reaches 18.016 versus 17.913 for fused `Exact Inv.`, which is why all flagship pre-training runs use `Exact Inv.`. Second, GDN-2 [25] reaches the best per-token perplexity in this cohort (17.592 on `s0,s1`, versus 17.806 for fused `SLA-R2`). SLA's advantage is throughput: in the same 90-minute budget, fused `SLA-R2` trains on 53% more tokens and reaches 16.660 perplexity, versus 17.459 for GDN-2 on the same seed. Because this comparison uses our unfused TPU port of GDN-2, it reflects kernel maturity as much as architecture; combining GDN-2's decoupled erase and write gates with SLA's TPU mapping is a natural next step.

GDN-2 TPU comparison details. GDN-2 [25] decouples the delta rule's scalar gate into channel-wise erase and write gates and inherits KDA's per-channel decay. On TPU v7x, in our unified single-harness full-layer benchmark ($B = 1$, 8K/32K F+B), our unfused XLA port of per-channel GDN-2 takes 63.86/251.39 ms at $D = 1792$ (15,217 MiB temporary HBM) and 34.15/132.83 ms at $D = 2048$ (Mamba-3 MIMO $R = 4$ isolated kernel for reference: 9.27 ms). End to end, unfused per-channel GDN-2 raises MoE step time by 1.23×–1.59× (1.349 vs. 1.093 s on 128-chip TPU v4; 778.1 ± 3.9 vs. 488.3 ms for fused `SLA-R2` on 64-chip TPU v7x).

On quality, GDN-2 is the strongest arm per token in our 15.1B-token cohort (17.592 ± 0.188 vs. 17.806 for fused `SLA-R2` and 17.893 for KDA on the same seeds `s0,s1`; Table 11b), but extrapolates worse (19.3% vs. 52.3% 32K 8-needle). Under equal pure-step compute on seed `s0` (5,582 s for 7,200 GDN-2 steps vs. 5,371 s for 11,000 fused `SLA-R2` steps), SLA reaches lower perplexity (16.660 vs. 17.459), whereas at equal 15.1B tokens on `s0` GDN-2 leads fused `SLA-R2` by 0.45 (17.459 vs. 17.913). This comparison reflects the lack of a fused Pallas/Mosaic TPU kernel for GDN-2 as much as the architectures; combining GDN-2's decoupled update with SLA's head sharing, tile-grouped gating, and fused chunk inversion is a promising direction for future work.

Length extrapolation. We test zero-shot extrapolation from 8K training to 32K–64K on RULER [28] needle retrieval and 4-choice `SQuAD` [49] and `HotpotQA` [69] (Figure 7b–c). Unclamped KDA degrades beyond the training length, severely at 63.1B, because about 1% of its channels learn horizons far beyond 8K (Table 17). SLA extrapolates without clamping: at 12.5B its 32K tail perplexity stays flat and 64K `HotpotQA` reaches 74%–78%. The advantage is limited, however. Clamping KDA's horizons to 8K, with or without retraining, gives similar results at 12.5B. At 63.1B, clamped KDA with `SSMax` [40] on the softmax layers beats SLA with `SSMax` on 64K QA, although SLA is better on 32K needle retrieval. With 128–512 questions per QA set (standard error 2–4 points), many differences in Tables 11–12 are within noise.

Causal τ_{eff} dispersion and NoPE `SSMax`. Table 17 audits per-channel retention horizons

$$\tau_{\text{eff}}(i) = \mathbb{E}_t\left[-\frac{1}{\ln f_{t,i}}\right]$$

(in tokens, for channel i with learned decay $f_{t,i}$) and Within-Tile Std (the RMS standard deviation of token-mean log-decay $\mathbb{E}_t[\ln f_{t,i}]$ across channels i within each $d_k/4$ tile block, equal to 0 when $G \in \{1, 4\}$) alongside zero-retraining inference interventions across 6.5B, 12.5B, and 63.1B MoEs.

Across all three scales, $G = d_k$ KDA develops 0.8%–1.3% ultra-slow $\tau_{\text{eff}}(i) \approx 131$K channels (within-tile log-decay RMS std = 3.84–4.97) that inject OOD state spikes into hybrid layers: clamping $\tau_{\text{eff}}(i) \le \tau_{\max} = 8{,}192$ at inference with zero retraining cuts 63.1B KDA 32K Tail PPL by 45.9× (353.42 → 7.70), lifts 32K 8-Needle across every scale (14.1% → 69.9% at 6.5B, 15.2% → 57.8% at 12.5B `7L`, 15.2% → 42.2% at 12.5B `15L`, 15.6% → 48.4% at 63.1B), and rescues +14.0 to +56.2 pts on 64K QA, whereas post-hoc $G = 4$ tile-mean projection collapses (57.8–78.1 Tail PPL). Live training-time softplus saturation (`TrainMaxTau=8K`) likewise stabilizes OOD retrieval at both 6.5B (69.5/73.4% QA) and 12.5B (60.4B tok, $n = 2$: 10.75 Tail, 49.6% 32K 8-N, 74.5/72.8% $N = 512$ QA) while retaining $G = 128$'s VPU gate tax and 4× FP32 state.

Conversely, on 63.1B `SLA-R4` (120.8B tok, NoPE), 24-layer state-norm telemetry confirms linear states contract cleanly from 8K to 32K ($\|S_t\|_F = 2.04 \to 1.62$, max row norm 1.59 → 1.15, 0.00% rows near E4M3 448), and rescaling each global softmax query q on the 6 NoPE full-attention layers at evaluation sequence length T via zero-retraining `SSMax` [40],

$$q \leftarrow q\left(1 + \alpha \ln\left(\frac{T}{T_{\text{train}}}\right)\right)$$

Table 16: Zero-shot 9-task `lm-eval` accuracy (11,793 real validation items; `norm9` macro-average using byte-length normalized `acc_norm` on multiple-choice tasks and exact-word `acc` on `WinoGrande`/`LAMBADA`, with zero-shot row-scaled FP8 state `norm9` in parentheses for SLA rows; 1 SE ≈ 0.65 pts) and real-text 4-candidate `RULER` QA accuracy (`SQuAD` 1-hop & `HotpotQA` 2-hop, $N = 128$/length, 25% chance, 1 SE ≈ 4.0 pts) across 8K $\rightarrow$ 32K $\rightarrow$ 64K against `0L:1A` GQA+RoPE and MaxText KDA baselines.

Model, Horizon, & Mixer Architecture	HS	ARC-E/C	PIQA	WG	SciQ	MMLU	BoolQ	LAMB	`lm-eval` (FP8)	`SQuAD` / `HotpotQA` (8K $\rightarrow$ 32K $\rightarrow$ 64K)
6.5B `0L:1A` GQA+RoPE (15.1B, BF16)	34.8	45.8 / 25.8	63.8	50.7	71.7	28.0	62.2	33.4	46.24%	64.8 $\rightarrow$ 54.7 $\rightarrow$ 56.2% / 67.2 $\rightarrow$ 52.3 $\rightarrow$ 50.0%
6.5B `4L:1A` MaxText KDA (15.1B, FP32)	35.4	45.7 / 25.8	64.8	50.5	70.3	27.1	63.8	37.2	46.73%	77.3 $\rightarrow$ 64.1 $\rightarrow$ 52.3% / 78.9 $\rightarrow$ 60.2 $\rightarrow$ 42.2%
6.5B `4L:1A` KDA + TrainMaxTau ($\tau_{\max} = 8$K)	34.9	46.4 / 25.6	65.2	51.1	69.1	27.6	63.6	36.3	46.64%	**75.0 $\rightarrow$ 71.1 $\rightarrow$ 69.5% / 79.7 $\rightarrow$ 74.2 $\rightarrow$ 73.4%**
6.5**B** `4L:1A` **SLA-R2** (15.1**B, BF16/FP8**)	34.6	45.1 / 24.7	65.4	52.2	70.6	28.5	63.8	38.2	**47.01%** (46.92%)	**71.9 $\rightarrow$ 69.5 $\rightarrow$ 71.1% / 71.1 $\rightarrow$ 71.1 $\rightarrow$ 74.2%**
6.5**B** `4L:1A` **SLA (Fused) R2** (15.1**B, BF16/FP8**)	35.8	46.4 / 24.3	64.5	51.6	71.4	28.5	63.8	36.4	46.97% (**47.08%**)	**75.8 $\rightarrow$ 72.7 $\rightarrow$ 68.8% / 72.7 $\rightarrow$ 72.7 $\rightarrow$ 70.3%**
6.5B `0L:1A` GQA+RoPE (30.2B, 4.4$\times$ Chin.)	37.7	47.3 / 26.7	66.3	52.2	73.0	29.6	63.5	35.8	48.01%	72.7 $\rightarrow$ 53.1 $\rightarrow$ 54.7% / 75.8 $\rightarrow$ 55.5 $\rightarrow$ 50.0%
6.5B `4L:1A` MaxText KDA (30.2B, 4.4$\times$ Chin.)	38.8	50.4 / 27.1	66.4	52.2	78.3	30.6	63.5	39.6	**49.66%**	79.7 $\rightarrow$ 53.1 $\rightarrow$ 32.0% / 85.2 $\rightarrow$ 39.8 $\rightarrow$ 20.3%
6.5**B** `4L:1A` **SLA-R2** (30.2**B, BF16/FP8**)	38.4	50.7 / 26.5	67.1	51.7	75.1	28.3	63.9	43.6	49.47% (49.44%)	78.1 $\rightarrow$ 59.4 $\rightarrow$ 56.2% / 79.7 $\rightarrow$ 58.6 $\rightarrow$ 51.6%
6.5**B** `4L:1A` **Iso-SLA-R4** (30.2**B, BF16/FP8**)	37.6	49.4 / 26.1	67.4	49.0	74.8	29.7	63.8	44.4	49.13% (48.97%)	**72.7 $\rightarrow$ 71.9 $\rightarrow$ 70.3% / 82.8 $\rightarrow$ 74.2 $\rightarrow$ 74.2%**
12.5B `0L:1A` GQA+RoPE (15.1B, $n = 3$ mean)	36.1	47.4 / 25.3	65.6	51.5	73.1	29.2	63.3	39.6	47.90 $\pm$ 0.36%	72.9 $\rightarrow$ 54.9 $\rightarrow$ 57.3% / 70.8 $\rightarrow$ 57.0 $\rightarrow$ 51.3%
12.5B `7L:1A` MaxText KDA (15.1B, $n = 2$ mean)	37.2	47.2 / 26.7	66.4	50.4	72.5	29.9	62.6	37.8	47.83 $\pm$ 0.48%	71.9 $\rightarrow$ 62.5 $\rightarrow$ 56.6% / 76.2 $\rightarrow$ 62.9 $\rightarrow$ 57.8%
12.5B `7L:1A` Scalar GDN ($G = 1$, 15.1B, `s1`)	37.0	47.6 / 25.5	65.2	51.5	74.1	28.8	62.8	40.6	48.11%	**74.2 $\rightarrow$ 72.7 $\rightarrow$ 74.2% / 77.3 $\rightarrow$ 76.6 $\rightarrow$ 76.6%**
12.5B `7L:1A` Wide-Key KDA ($G = 256$, $n = 2$ mean)	37.2	47.3 / 26.4	65.8	52.2	73.3	27.8	63.0	41.8	48.30%	74.2 $\rightarrow$ 66.8 $\rightarrow$ 55.9% / 76.2 $\rightarrow$ 68.0 $\rightarrow$ 54.3%
12.5**B** `7L:1A` **A2** ($G = 4$ **Exact Inv., FP32,** `s0`)	38.1	47.9 / 25.7	65.0	51.1	73.8	28.7	59.8	42.3	48.05% (48.04%)	71.9 $\rightarrow$ 65.6 $\rightarrow$ 62.5% / 73.4 $\rightarrow$ 70.3 $\rightarrow$ 68.0%
12.5**B** `7L:1A` **A3** ($G = 4$ `1-Step Neumann`, **FP32,** $n = 2$)	37.0	46.8 / 25.8	65.9	50.0	72.6	28.6	63.6	38.9	47.66% (47.47%)	**76.6 $\rightarrow$ 71.9 $\rightarrow$ 71.9% / 74.6 $\rightarrow$ 72.3 $\rightarrow$ 71.9%**
12.5**B** `7L:1A` **SLA-R2** (15.1**B,** $n = 3$ **mean**)	37.6	47.9 / 26.1	66.2	52.3	72.2	29.3	63.1	42.1	**48.55 $\pm$ 0.48%** (**48.64%**)	**73.7 $\rightarrow$ 68.8 $\rightarrow$ 69.5% / 74.0 $\rightarrow$ 71.9 $\rightarrow$ 73.7%**
12.5**B** `7L:1A` **SLA (Fused) R2** (15.1**B, BF16/FP8,** `s0`)	37.6	47.7 / 26.3	66.3	50.8	71.6	28.0	60.0	45.3	48.18% (48.21%)	**73.4 $\rightarrow$ 67.2 $\rightarrow$ 64.1% / 76.6 $\rightarrow$ 75.8 $\rightarrow$ 75.0%**
12.5**B** `7L:1A` **Iso-SLA-R4** (15.1**B,** $n = 3$ **mean**)	36.3	48.8 / 25.9	66.3	50.6	72.3	28.7	63.0	43.0	48.33 $\pm$ 0.45% (48.39%)	**74.5 $\rightarrow$ 72.7 $\rightarrow$ 73.2% / 74.5 $\rightarrow$ 75.5 $\rightarrow$ 72.9%**
12.5B `15L:1A` MaxText KDA (15.1B, FP32)	36.0	48.6 / 26.6	66.0	50.2	71.9	29.4	64.2	40.7	48.18%	76.6 $\rightarrow$ 68.0 $\rightarrow$ 66.4% / 75.8 $\rightarrow$ 70.3 $\rightarrow$ 68.0%
12.5**B** `15L:1A` **SLA-R2 Uniform** ($n = 3$ **mean**)	37.0	50.1 / 27.0	65.6	50.5	74.5	29.2	63.5	41.6	**48.79 $\pm$ 0.22%** (**48.75%**)	**72.6 $\rightarrow$ 71.9 $\rightarrow$ 69.0% / 72.2 $\rightarrow$ 65.9 $\rightarrow$ 65.1%** (68.2% N=512)
12.5**B** `15L:1A` **SPD-SLA-W Polar-2+ Uniform** (`s0`)	35.9	49.9 / 26.2	65.7	50.2	74.1	29.8	63.8	41.5	48.57% (**48.79%**)	71.9 $\rightarrow$ 71.1 $\rightarrow$ 65.6% / 71.9 $\rightarrow$ 68.8 $\rightarrow$ 65.6% (69.5% N=512)
12.5B `15L:1A` SLA-R2+MS (15.1B, BF16/FP8)	36.6	50.8 / 27.3	65.8	51.1	74.1	27.5	63.5	37.2	48.20% (48.18%)	57.0 $\rightarrow$ 57.0 $\rightarrow$ 57.0% / 67.2 $\rightarrow$ 57.8 $\rightarrow$ 53.9%
12.5B `Pure-Linear` KDA (15.1B, FP32)	37.0	49.1 / 25.5	65.6	51.2	70.1	27.7	63.6	40.1	47.77%	57.8 $\rightarrow$ 57.8 $\rightarrow$ 55.5% / 62.5 $\rightarrow$ 56.2 $\rightarrow$ 56.2%
12.5B `Pure-Linear` SLA-R2+MS (15.1B, BF16/FP8)	36.4	50.3 / 26.7	66.6	49.9	72.6	27.9	63.3	40.2	48.21% (48.30%)	60.2 $\rightarrow$ 57.8 $\rightarrow$ 59.4% / 62.5 $\rightarrow$ 56.2 $\rightarrow$ 53.9%
12.5B `0L:1A` GQA+RoPE (32K-ext, 10.1B)	37.5	49.0 / 24.2	65.8	51.7	73.7	28.9	63.8	40.6	48.35%	75.0 $\rightarrow$ 64.8 $\rightarrow$ 56.2% / 73.4 $\rightarrow$ 65.6 $\rightarrow$ 57.8%
12.5B `7L:1A` MaxText KDA (32K-ext, 10.1B)	37.9	47.0 / 26.6	66.8	51.8	72.0	30.0	59.2	37.2	47.62%	73.4 $\rightarrow$ 73.4 $\rightarrow$ 71.9% / 78.1 $\rightarrow$ 78.1 $\rightarrow$ 72.7%
12.5**B** `7L:1A` **SLA-R2** (32**K-ext, BF16/FP8**)	38.9	48.4 / 26.5	67.1	51.5	71.9	29.4	62.6	42.7	48.78% (48.86%)	74.2 $\rightarrow$ 72.7 $\rightarrow$ 71.9% / 71.9 $\rightarrow$ 77.3 $\rightarrow$ **75.8%**
12.5**B** `7L:1A` **Iso-SLA-R4** (32**K-ext, BF16/FP8**)	37.1	49.8 / 26.0	65.7	50.7	75.5	30.0	61.9	44.7	**49.04%** (**49.18%**)	**78.1 $\rightarrow$ 75.0 $\rightarrow$ 75.8%** / 78.1 $\rightarrow$ 72.7 $\rightarrow$ 74.2%
12.5B `0L:1A` GQA+RoPE (60.4B, BF16)	41.4	49.2 / 27.7	68.3	51.2	76.9	29.6	62.7	43.5	50.06%	70.3 $\rightarrow$ 59.4 $\rightarrow$ 58.6% / 76.6 $\rightarrow$ 58.6 $\rightarrow$ 53.1%
12.5B `7L:1A` MaxText KDA (60.4B, $n = 3$ mean)	41.4	51.4 / 27.4	68.2	51.4	75.7	30.4	62.9	44.6	50.36 $\pm$ 0.53%	79.2 $\rightarrow$ 69.3 $\rightarrow$ 54.9% / 75.0 $\rightarrow$ 65.6 $\rightarrow$ 50.3%
12.5B `7L:1A` KDA + TrainMaxTau (60.4B, $n = 2$)	40.9	51.2 / 27.6	67.8	51.9	75.8	30.2	62.8	44.5	50.30%	**78.1 $\rightarrow$ 75.0 $\rightarrow$ 74.5% / 77.3 $\rightarrow$ 73.8 $\rightarrow$ 72.8%** (N=512)
12.5**B** `7L:1A` **SLA-R2** (60.4**B,** $n = 3$ **mean**)	41.4	52.5 / 27.2	68.6	52.0	76.7	29.9	63.2	46.9	**50.94 $\pm$ 0.47%** (**50.93%**)	**78.1 $\rightarrow$ 75.3 $\rightarrow$ 75.5% / 78.6 $\rightarrow$ 78.7 $\rightarrow$ 77.9%** (77.4% FP8)
12.5**B** `7L:1A` **SLA-R4** (60.4**B, BF16/FP8**)	41.3	49.9 / 25.7	67.6	51.1	76.1	30.3	63.2	47.3	50.28% (50.13%)	**72.7 $\rightarrow$ 75.8 $\rightarrow$ 72.7% / 72.7 $\rightarrow$ 78.9 $\rightarrow$ 75.0%** (76.6% FP8)
12.5**B** `7L:1A` **Iso-SLA-R4** (60.4**B, BF16/FP8**)	41.2	48.9 / 28.1	67.1	47.7	76.7	30.0	62.3	49.3	50.13% (50.31%)	**78.1 $\rightarrow$ 75.0 $\rightarrow$ 76.6% / 78.1 $\rightarrow$ 80.5 $\rightarrow$ 77.3%** (77.3% FP8)
12.5**B** `7L:1A` **Iso-SLA (Fused) R4** (60.4**B,** $n = 3$ **mean**)	41.5	52.6 / 27.8	67.9	51.7	77.5	29.5	63.1	45.8	**50.82 $\pm$ 0.22%** (**50.84%**)	**77.1 $\rightarrow$ 73.2 $\rightarrow$ 74.7% / 75.0 $\rightarrow$ 73.2 $\rightarrow$ 74.5%** (73.2% FP8)
12.5B `15L:1A` MaxText KDA (60.4B, FP32)	41.6	49.5 / 27.2	68.0	51.7	77.2	30.1	62.7	45.6	50.41%	77.3 $\rightarrow$ 69.5 $\rightarrow$ 62.5% / 76.6 $\rightarrow$ 71.1 $\rightarrow$ 60.2%
12.5**B** `15L:1A` **SLA-R2 Uniform** (60.4**B, BF16/FP8**)	41.1	53.3 / 28.4	68.0	52.6	77.6	30.3	63.6	46.3	**51.24%** (**51.08%**)	**78.9 $\rightarrow$ 75.8 $\rightarrow$ 72.7% / 76.6 $\rightarrow$ 72.7 $\rightarrow$ 70.3%** (70.3% FP8)
12.5B `15L:1A` SLA-R2+MS (60.4B, BF16/FP8)	42.8	56.4 / 28.4	68.2	49.8	78.5	28.9	63.8	40.6	50.82% (51.02%)	67.2 $\rightarrow$ 64.1 $\rightarrow$ 64.8% / 71.1 $\rightarrow$ 57.8 $\rightarrow$ 57.0% (57.0% FP8)
12.5B `15L:1A` SPD-SLA-W+MS (60.4B, BF16/FP8)	42.0	53.5 / 28.9	68.5	50.3	77.1	30.1	63.6	40.8	50.54% (50.59%)	60.9 $\rightarrow$ 59.4 $\rightarrow$ 57.8% / 66.4 $\rightarrow$ 58.6 $\rightarrow$ 56.2% (56.2% FP8)
63.1B `0L:1A` GQA+RoPE (30.2B, RoPE)	40.6	50.2 / 27.0	68.4	51.6	74.0	28.9	63.8	42.3	49.64%	71.9 $\rightarrow$ 53.9 $\rightarrow$ 45.3% / 71.9 $\rightarrow$ 56.2 $\rightarrow$ 39.8%
63.1B `4L:1A` GDN-2 (30.2B, RoPE, 2,728 ms)	37.9	50.8 / 26.9	67.7	51.3	72.6	29.7	63.1	40.7	48.96%	69.5 $\rightarrow$ 62.5 $\rightarrow$ 55.5% / 72.7 $\rightarrow$ 66.4 $\rightarrow$ 64.1%
63.1**B** `4L:1A` **SLA (Fused) R4** (30.2**B, RoPE**)	37.7	47.3 / 26.1	67.5	51.7	73.4	29.2	59.6	44.2	48.53% (48.51%)	68.0 $\rightarrow$ 59.4 $\rightarrow$ **59.4%** / 71.1 $\rightarrow$ 59.4 $\rightarrow$ 54.7%
63.1**B** `15L:1A` **SPD-SLA-W+MS** (30.2**B, RoPE**)	38.1	48.4 / 25.2	66.4	50.8	71.5	30.3	63.0	42.0	48.40% (48.65%)	**71.1 $\rightarrow$ 64.1 $\rightarrow$ 57.8% / 72.7 $\rightarrow$ 65.6 $\rightarrow$ 64.8%**
63.1B `4L:1A` MaxText KDA (120.8B, NoPE)	**54.4**	61.5 / **35.4**	**73.2**	**54.6**	86.9	34.1	**64.1**	56.2	**57.83%**	85.9 $\rightarrow$ 39.1 $\rightarrow$ 33.6% / **89.1** $\rightarrow$ 28.9 $\rightarrow$ 27.3%
63.1**B** `4L:1A` **SLA-R2** (120.8**B, NoPE**)	53.4	**62.4** / 34.6	72.2	53.7	**87.5**	**34.4**	**64.1**	55.4	57.52%	**88.3 $\rightarrow$ 60.2 $\rightarrow$ 57.0%** / 84.4 $\rightarrow$ **61.7** $\rightarrow$ **51.6%**
63.1**B** `4L:1A` **SLA-R4** (120.8**B, NoPE**)	53.8	57.1 / 34.6	73.1	51.9	85.5	34.3	63.0	**57.2**	56.72% (56.74%)	85.2 $\rightarrow$ 57.0 $\rightarrow$ 53.9% / 87.5 $\rightarrow$ 51.6 $\rightarrow$ 43.0% (43.8% FP8)

Table 17: Multi-scale causal τ_{eff} dispersion, zero-retraining $\tau_{\max} = 8{,}192$ inference clamping, live `TrainMaxTau=8K` saturation, post-hoc $G = 4$ projection, and NoPE `SSMax` ($\alpha = 0.5$) on 32K Tail PPL, 32K 8-N, and 64K `RULER` QA ($N = 128$ unless noted N=512). The clamp evaluator (`a3_causal_clamp_eval`) scores its own held-out Dolma slice and paired $N = 128$ QA subset, so absolute Tail PPL and QA values differ from the standalone scorers of Tables 11, 12, 15, and 16 (by 1–2 QA items, 0.8–1.6 pts) while within-row intervention deltas remain exact.

Scale, Horizon, & Trained Checkpoint	G	Within-Tile Std	Unclamped 32K Tail / 8-N	Intervened 32K Tail / 8-N (Proj. G=4)	Unclamped 64K `SQuAD / Hotpot`	Intervened 64K QA (Proj. G=4 / Note)
6.5B `4L:1A` MaxText KDA (15.1B)	128	4.41	14.64/14.1%	**11.77**/**69.9**% (70.47)	53.1%/42.2%	**75.0**%/**79.7**% (24.2/39.8%)
6.5B `4L:1A` KDA + TrainMaxTau (15.1B)	128	$\leq \tau_{\max}$	**12.36**/**48.4**%	**12.36**/**48.4**% (co-capped)	**69.5**%/**73.4**%	**69.5**%/**73.4**% (68.8/70.3% N=512)
6.5B `4L:1A` Iso-Param SLA-R4 (15.1B)	4	**0.00**	**11.94**/**48.8**%	**11.94**/**48.8**% (11.94)	**67.2**%/**76.6**%	**67.2**%/**76.6**% (invariant)
12.5B `7L:1A` MaxText KDA (60.4B, `s0`)	128	4.58	12.35/15.2%	**10.26**/**57.8**% (57.85)	47.7%/44.5%	**78.1**%/**75.0**% (28.1/45.3%)
12.5B `7L:1A` KDA + TrainMaxTau (60.4B, $n = 2$)	128	$\leq \tau_{\max}$	**10.75**/**49.6**%	**10.75**/**49.6**% (co-capped)	**74.5**%/**72.8**% (N=512)	**74.5**%/**72.8**% (N=512)
12.5B `15L:1A` MaxText KDA (60.4B)	128	4.97	10.48/15.2%	**9.66**/**42.2**% (63.30)	59.4%/61.7%	**73.4**%/**78.9**% (25.8/34.4%)
12.5**B** `7L:1A` **Iso-SLA (Fused) R4** (60.4**B,** `s0`)	**4**	**0.00**	**9.80**/**63.7**%	**9.80**/**63.7**% (**9.80**)	**73.4**%/**77.3**%	**75.0**%/**77.3**% (+1.6/0.0 pts)
63.1B `4L:1A` MaxText KDA (120.8B, NoPE)	256	3.84	353.42/15.6%	**7.70**/**48.4**% (78.06)	34.4%/25.8%	**76.6**%/**82.0**% (29.7/25.0%)
$+\tau_{\max} = 8\text{K} +$ `SSMax` ($\alpha = 0.25$, N=512)	256	$\leq \tau_{\max}$	380.66/18.8% (`SSMax` only)	**8.04**/**74.2**% ($\tau_{\max}$+`SSMax`)	27.3%/27.0% (`SSMax` only)	**90.6**%/**86.5**% (77.5/76.4% norm)
63.1**B** `4L:1A` **SLA-R2** (120.8**B, NoPE)**	**4**	**0.00**	**13.78**/18.6%	**12.07**/21.9% ($\tau_{\max}$ only)	**57.0**%/**53.1**%	**61.7**%/**60.9**% (55.7/56.8% N=512)
63.1**B** `4L:1A` **SLA-R2 +** `SSMax` (120.8**B**)	**4**	**0.00**	13.78/18.6%	**8.45**/**84.4**% (`SSMax`, 0 clamp)	57.0%/53.1%	– (zero linear clamp)
63.1**B** `4L:1A` **SLA-R4 +** `SSMax` (120.8**B**)	**4**	**0.00**	22.78/18.0%	**8.41**/**83.2**% (`SSMax`, 0 clamp)	51.6%/45.1% (N=512)	**75.6**%/**79.1**% (62.3/66.6% norm)

(with $T_{\text{train}} = 8{,}192$ and SSMax scale $\alpha = 0.5$, unrelated to α_{safe}; zero linear clamp), cuts 32K Tail PPL to 8.45 (`SLA-R2`) and 8.41 (`SLA-R4`) and restores 32K 8-Needle to 84.4% (`SLA-R2`) and 83.2% (`SLA-R4`) with zero linear gate clamping (whereas linear $\tau_{\max} = 8{,}192$ clamping alone without `SSMax` moves 63.1B `SLA-R2` 32K 8-Needle only from 18.6% to 21.9%, and `SSMax` alone on unclamped KDA still diverges at 380.66 Tail PPL and 18.8% 8-Needle, cleanly separating global NoPE softmax dilution from KDA's linear gate drift).

Decay granularity. Sweeping $G \in \{1, 4, 16, 64, 256\}$ with everything else fixed (Table 18), $G = 4$ improves 8K perplexity over $G = 1$ on all three paired seeds, but only by 0.15 on average, while $G = 1$ extrapolates better (59.0% vs. 49.2% 32K 8-needle). Finer groups ($G \geq 16$) extrapolate much worse. A scalar decay per head, as in GDN, is therefore a strong alternative; we choose $G = 4$ for its multiple memory horizons and slightly better in-domain fit, not for long-context robustness.

Specifically, Table 18 dissects the 5-point $G \in \{1, 4, 16, 64, 256\}$ sweep (with decay sub-tile width $b = d_k/G = 256/G$) on the identical $[4, 256, 512]$ `Exact Inv.` FP32 kernel using both matched-seed (`s0,s1`) means and per-seed paired differences against $G = 1$. On matched `s0,s1` seeds, 8K 11-set PPL favors finer G without a monotone trend (17.946, 17.749, 17.873, 17.748, 17.694 for $G = 1, 4, 16, 64, 256$; all $G \geq 4$ lie within 0.18 PPL of each other), and $G = 4$ improves 11-set PPL over $G = 1$ on 3/3 paired seeds (ΔPPL $= -0.34, -0.06, -0.06$; mean -0.151). Conversely, zero-shot OOD retrieval favors larger tile width $b = 256/G$: $G = 1$ ($b = 256$) achieves the highest 32K 8-N (59.0%) and 64K QA (67.8/71.2%, $\tau_{\max} \leq 56$K across seeds), $G = 4$ ($b = 64$) trades a small OOD step (49.2% 8-N; 64.5/68.9% QA, where seed `s2` developed one slow tile with $\tau_{\max} = 124$K, scoring 58.6/64.5% vs. 67.5/71.2% on `s0,s1`) for 4× intra-head timescales, and $G \geq 16$ ($b \leq 16$, $\tau_{\max} = 131$K) degrades sharply to 27.1%/19.0% 8-N and 60.3/61.9% → 53.5/50.8% QA.

Table 18: Unconfounded 5-point $[4, 256, 512]$ decay-granularity sweep on the identical `chunk_kda Exact Inv.` FP32 kernel ($\mu = 0$, 12.5B MoE, `7L:1A`, 15.10B tok), reporting matched `s0,s1` 11-set PPL (valid across all 5 granularities), paired per-seed ΔPPL vs. $G = 1$, 8K → 32K Tail PPL, 32K 4/8-N, `lm-eval norm9`, and $N = 512$ (1 SE = 2.0 pts) 64K `RULER` QA.

$[4, 256, 512]$ `Exact Inv.` FP32 Arm ($\mu = 0$)	G (b)	Within-Tile Std	`s0,s1` Mean PPL (n=3 Mean)*	Paired ΔPPL vs. $G = 1$ (`s0..s2`)	8K → 32K Tail	32K 4-N / 8-N	`norm9`	$N = 512$ 64K `SQuAD / Hotpot`
Full-Tile Scalar Decay ($n = 3$)	**1** (256)	**0.00**	17.946 (17.689)	0.000 (ref)	11.58 → **11.26**	**70.4**%/**59.0**%	48.30%	**67.8** ± **1.6**%/**71.2** ± **1.5**%
Systolic Sub-Tile Default ($n = 3$)	**4** (64)	**0.00**	17.749 (17.538)	**−0.151** (−0.34, −0.06, −0.06)	11.65 → **11.41**	58.6%/**49.2**%	48.10%	**64.5** ± **4.3**%/**68.9** ± **3.3**%*
Intermediate Sub-Tile ($n = 3$)	16 (16)	4.34	17.873 (17.626)	−0.064 (−0.14, −0.00, −0.05)	11.75 → 12.12 ± 0.72	41.4%/27.1%	48.35%	60.3 ± 4.2%/61.9 ± 5.4%
Fine Sub-Tile Decay ($n = 3$)	64 (4)	5.05	17.748 (17.523)	−0.167 (−0.36, −0.04, −0.11)	11.53 → 12.94 ± 0.23	27.7%/19.0%	48.61%	53.5 ± 2.2%/50.8 ± 1.8%
Per-Channel Wide-Key ($n = 2$ PPL, 3 OOD)	256 (1)	4.78	**17.694** (`s0,s1`)	**−0.251** (−0.36, −0.14)	11.53 → 12.86 ± 1.13	41.0%/25.1%	48.54%	56.8%/58.9%†

*Per-seed 11-set PPL (`s0, s1, s2`): $G = 1$ (18.210, 17.681, 17.178), $G = 4$ (17.873, 17.625, 17.117), $G = 16$ (18.067, 17.679, 17.132), $G = 64$ (17.851, 17.644, 17.073), $G = 256$ (17.850, 17.538, `s2` 11-set not run). Seed `s2` used a deterministic `OLMO_DATA_STEP_OFFSET` stream slice (17.07–17.18 PPL across all G). On `s0,s1`, $G = 4$ 64K QA is 67.5 ± 1.8%/71.2 ± 0.1% ($\tau_{\max} \leq 85$K); on `s2`, one $G = 4$ tile hit $\tau_{\max} = 124$K (58.6/64.5%). †For $G = 256$, $N = 128$ 64K QA is 57.3%/59.1%.

FP8 state. SLA states quantize well: zero-retrain row-scaled E4M3 state quantization changes 8K/32K loss by less than 10^{-4} nats (Section 3.4), and FP8 output QAT with `Exact Inv. (A6)` is within 0.09 perplexity of BF16 on `s0` (18.006 vs. 17.913).

4 Conditioning Recurrent State in Low Precision: Oblique Whitening and `Polar-2+`

Diluting softmax attention from `7L:1A` (12.5%) to `15L:1A` (6.25%) cuts 131K KV+state HBM to 16.47 GiB (Table 8) while forcing middle linear layers (`L5, L8, L11`) to expand active retention 10×–35× ($\rho_{\text{eff}} = N_{\text{eff}}/d_k \in [0.101, 0.152]$ vs. ≤ 0.015 at `4L:1A`, where N_{eff} is the number of key-value associations a layer effectively retains, estimated from its decay horizons) under coherent keys ($\kappa_2(I_C + K_c K_c^\top) \in [20.6, 62.9]$; Figure 8c). Whereas GPU second-order memories require 30-step Conjugate Gradient (CG) [64] or diagonal approximations [62], a 256 × 256 MXU runs five [128, 128] matmuls across 64 chunks in 0.62 ms forward (0.90 ms F+B). Here we formalize when key whitening helps in low precision and detail **Polar-2+**, a five-matmul BF16-safe inverse-square-root primitive.

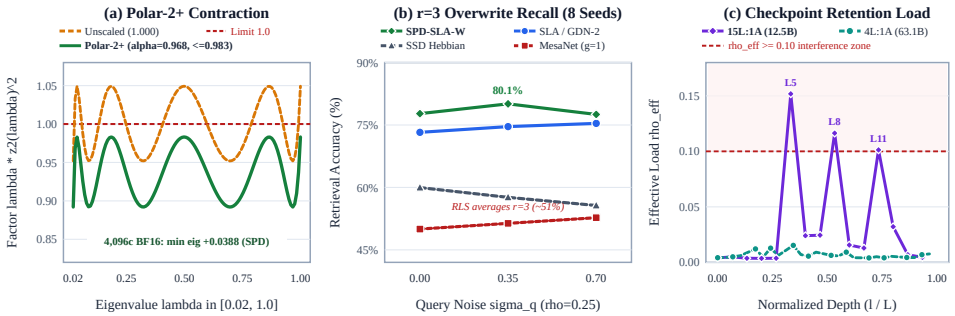

Figure 8: **Polar-2+ minimax contraction, overwrite recall, and checkpoint retention audit. (a)** Scalar inverse-square-root contraction $\lambda z_2(\lambda)^2$ on normalized eigenvalue $\lambda \in [0.02, 1.0]$ (Proposition 4.2): unscaled polynomial factor $\alpha = 1.000$ hits $1.0492 > 1$, whereas `Polar-2+` (safety factor $\alpha_{\text{safe}} = 0.968$) bounds the peak to $0.9831 < 1$, keeping minimum eigenvalue $\lambda_{\min}(\Sigma_c) \geq 0.0388$ of key preconditioner Σ_c across 4,096 BF16 chunks (524K tok). **(b)** 8-seed Hadamard overwrite accuracy (overwrite multiplicity $r = 3$) at memory load $\rho_{\text{eff}} = N/d_k = 0.25$ ($N = 16$ stored keys, $d_k = 64$). **(c)** Trained 63.1B (`4L:1A`) and 12.5B (`15L:1A`) checkpoint load ρ_{eff}.

Proposition 4.1 (Oblique Whitening vs. Euclidean & RLS)**.** *Let $K \in \mathbb{R}^{N \times d_k}$ stack $N \leq d_k$ unit keys $k_1, \ldots, k_N \in \mathbb{R}^{d_k}$ with mutual coherence $\nu_{coh} = \max_{i \neq j} |\langle k_i, k_j \rangle|$ and Gram matrix $\mathcal{G} = KK^\top$ satisfying $\kappa_2(\mathcal{G}) > 1$ (proofs in Appendix A).*

1. [Prior Delta [62, 67]] *For equiangular keys ($|\langle k_i, k_j \rangle| = \nu_{coh}$ for all $i \neq j$), zero-mean values $v_j \in \mathbb{R}^{d_v}$ of expected squared norm $s_v^2 = \mathbb{E}\,\|v_j\|_2^2$, and load $\rho_{eff} = N/d_k$, single-pass Euclidean delta incurs readout error $\Omega(\rho_{eff}(d_k \nu_{coh}^2) s_v^2)$, and diagonal column-energy preconditioning $\text{diag}(p(K))^{-1/2} k_t$ [62] (where $p(K) \in \mathbb{R}^{d_k}$ has m-th entry $p_m(K) = \sum_{i=1}^N K_{i,m}^2$) fails under orthogonal key rotation $K \mapsto KQ$ ($Q \in \text{O}(d_k)$).*

2. [Prior RLS [64]] *MesaNet (with recursive least-squares key covariance $H_t = \gamma H_{t-1} + k_t k_t^\top$ and scalar forget factor $\gamma \in (0, 1]$) averages $r \geq 2$ overwrites $v^{(1)}, \ldots, v^{(r)} \in \mathbb{R}^{d_v}$ of the same key as $\frac{1}{r}\sum_{m=1}^r v^{(m)}$ at $\gamma = 1$, and suppresses stale values at $\gamma < 1$ only by decaying write-once keys stored $t - j$ steps earlier (at step $j < t$) at rate γ^{t-j}.*

3. ***[Ours]** Oblique whitened keys $\tilde{K}_c \in \mathbb{R}^{C \times d_k}$ and causal oblique keys $\tilde{k}_t \in \mathbb{R}^{d_k}$ (Eq. (14)) inside Eq. (4) decorrelate intra-chunk writes while erasing selectively along $\tilde{k}_t$.*

Non-circular normalization and `Polar-2+` formulation. Matrix-sign polynomials (Polar Express [3], CANS [21]) collapse singular values to 1, while coupled `NS-5` [4, 24, 26, 37] uses 10 $[C, C]$ matmuls [Prior Polynomials]. Here $K_c \in \mathbb{R}^{C \times d_k}$ stacks the chunk's keys, $\Sigma_{c-1} \succ 0$ is the $d_k \times d_k$ key preconditioner carried across chunks (Eq. (15); initialized at $\Sigma_0 = I_{d_k}$ and relaxed toward I_{d_k} at rate $\rho_p \in [0, 1]$), $\eta_c \in (0, 1]$ is the whitening step size, and $\varepsilon_\beta > 0$ is the erasure safety margin (Section 2.3). To avoid circular regularization, we form the raw chunk Gram matrix $\mathcal{G}_0 \in \mathbb{R}^{C \times C}$, its

Frobenius normalization $\bar{\mathcal{G}}_0$, spectral upper bound $\hat{u}_0$, and adaptive ridge ε_{reg}:

$$\mathcal{G}_0 = I_C + \eta_c K_c \Sigma_{c-1} K_c^\top, \qquad \bar{\mathcal{G}}_0 = \frac{\mathcal{G}_0}{\|\mathcal{G}_0\|_F}, \qquad \hat{u}_0 = \max\big(\|\mathcal{G}_0\|_F \left\|\bar{\mathcal{G}}_0^2\right\|_F^{1/2}, 1\big) \geq \lambda_{\max}(\mathcal{G}_0),$$

and

$$\varepsilon_{\text{reg}} = \max\left(0.02,\ \frac{0.0205}{0.9795}\hat{u}_0 - 1\right).$$

Setting regularized Gram $\mathcal{G}_{\text{reg}} = \mathcal{G}_0 + \varepsilon_{\text{reg}} I_C$ and scale $u_0 = \hat{u}_0 + \varepsilon_{\text{reg}}$ maps the eigenvalues of $T_0 = \mathcal{G}_{\text{reg}}/u_0$ into $[0.02, 1.0]$ while reusing $\bar{\mathcal{G}}_0^2$ (T_0^2 adds zero matmuls; 5 matmuls total), with step-size clamp

$$\beta_t \leq \frac{2 - \varepsilon_\beta}{\max(k_t^\top \tilde{k}_t,\ \varepsilon_\beta)}.$$

Letting $K_{<t}$ stack the chunk keys strictly preceding token t, $\mathcal{M}_\ell \in \mathbb{R}^{C\times C}$ denote the degree-2 polynomial factors for stages $\ell \in \{0, 1\}$ with $T_1 = \mathcal{M}_0 T_0 \mathcal{M}_0$ and minimax coefficients $(a_0, b_0, e_0) = (4.92864, -11.29513, 7.70099)$ and $(a_1, b_1, e_1) = (2.00410, -1.38700, 0.38894)$, and $\mathcal{W}_{\text{isq}} \in \mathbb{R}^{C\times C}$ denote the safety-scaled ($\alpha_{\text{safe}} = 0.968$) inverse-square-root factor, **Polar-2+** computes:

Low-Precision Conditioning: Oblique Whitening & BF16-Safe `Polar-2+`

$$\tilde{K}_c = \sqrt{\eta_c}\,\mathcal{W}_{\text{isq}} K_c \Sigma_{c-1}, \qquad \tilde{k}_t \propto \big(I_{d_k} - \eta_c \Sigma_{c-1} K_{<t}^\top K_{<t}\big)\Sigma_{c-1} k_t, \tag{14}$$

$$\Sigma_c = (1 - \rho_p)\big(\Sigma_{c-1} - \tilde{K}_c^\top \tilde{K}_c\big) + \rho_p I_{d_k}, \tag{15}$$

$$\mathcal{M}_\ell = a_\ell I_C + b_\ell T_\ell + e_\ell T_\ell^2, \qquad \mathcal{W}_{\text{isq}} = \alpha_{\text{safe}} u_0^{-1/2} \mathcal{M}_1 \mathcal{M}_0. \tag{16}$$

Symmetric $\tilde{K}_c$ runs the chunk-whitening kernels (Table 19), the 4,096-chunk ($\rho_p = 0$) BF16 Woodbury test, and lagged block-32 causal whitening (Appendix C). Since dense $\mathcal{W}_{\text{isq}}$ mixes chunk tokens, packed LM training uses causal Schur-oblique $\tilde{k}_t$ in Eq. (14) ($k_t^\top \tilde{k}_t = 1$) with `Polar-2+` in the chunk-boundary Woodbury tail ($\rho_p = 1$ resets $\Sigma_c = I_{d_k}$ for `SLA Tree-CP`, while unscaled $\alpha = 1.000$ still diverges at step 299 via tail gradients; for $\rho_p = 0.02$ cross-chunk Σ_c, multiplicative Sylvester congruence $\mathcal{B}_c \Sigma_{c-1} \mathcal{B}_c^\top$ with contraction matrix $\mathcal{B}_c \in \mathbb{R}^{d_k \times d_k}$ [Appendix A] replaces additive Eq. (15), which otherwise hits `NaN` at step 282).

Proposition 4.2 (Exact Equioscillation and Safety Margin of Polar-2+)**.** *Writing $m_\ell(x) = a_\ell + b_\ell x + e_\ell x^2$, let*

$$z_2(\lambda) = m_1\big(\lambda\, m_0(\lambda)^2\big)\, m_0(\lambda)$$

be the scalar map that Eq. (16) applies to an eigenvalue λ of T_0 (since $T_1 = \mathcal{M}_0 T_0 \mathcal{M}_0$ has eigenvalue $\lambda\, m_0(\lambda)^2$), and let $\kappa_{post} = \kappa_2(\mathcal{W}_{isq}\mathcal{G}_{reg}\mathcal{W}_{isq})$ be the post-whitening condition number. On $\lambda \in [0.02, 1.0]$, z_2 equioscillates with $\lambda z_2(\lambda)^2 \in [0.9520, 1.0492]$, exceeding 1 for unscaled $\alpha = 1$ and driving recursive Woodbury indefinite. Scaling by $\alpha_{safe} = 0.968$ bounds exact $\lambda(\alpha_{safe} z_2(\lambda))^2 \in [0.8920, 0.9831]$ (1.69% BF16 contractive margin below 1; exact $\kappa_{post} \leq 1.102$, BF16 5-matmul $\kappa_{post} \leq 1.123$).

Table 19: Measured TPU v7x chunk-whitening latency (64 chunks = 8,192 tokens, $C = 128$, $H_k = 4$, $d_k = 256$; MesaNet `CG-10`: 1.74/11.16 ms), exact conditioned Gram κ_{post} (≤ 1.123 in BF16), and 4,096-chunk (524K-token, $\rho_p = 0$) BF16 Woodbury stress-test $\lambda_{\min}$.

Whitening / Second-Order Solver	Ops	Fwd (ms)	F+B (ms)	κ_{post}	$\lambda_{\min}$ (4,096c BF16)
MesaNet Gated `CG-30`	VPU+MXU	3.050	27.640	1.02	– (CG)
Panel-16 Cholesky (VPU)	Tri. Solve	1.412	2.385	1.00	+0.0148 (SPD)
Coupled Newton-Schulz (`NS-5`)	10 MM	0.708	1.024	1.04	+0.0294 (SPD)
Polar-2 Unscaled ($\alpha = 1.000$)	5 MM	0.615	0.895	1.102	−0.0412 (Indef.)
Polar-2+ Safe ($\alpha = 0.968$)	**5** MM	0.617	0.897	1.102	+0.0388 **(SPD)**

When whitening helps vs. unwhitened SLA. `Polar-2+` runs $2.66\times$ faster than Panel-16 Cholesky and $12.4\times$–$30.8\times$ faster than MesaNet `CG-10/30` (0.897 ms F+B; Table 19), keeping 4,096-chunk (524K-token) BF16 Woodbury SPD ($\lambda_{\min} \geq 0.0388$). It trains 2-seed 12.5B `15L:1A` with zero NaNs and matches 10-matmul `NS-5` (`s1`: 17.936 vs. 17.970 PPL), whereas $\alpha = 1.000$ diverges at step 299 and BF16 Cholesky hits NaN at step 54 (Appendix C). Because its language-modeling benefit shrinks to 0.004 nats by 15B tokens and unwhitened $G = 4$ SLA already saturates PPL at 60.4B `15L:1A+MS` (15.443 vs. 15.624 at $+3.7\%$ step time), unwhitened SLA is our default mixer.

`SPD-SLA-W` serves three high-collision regimes:

1. correlated $r \geq 2$ overwrite ($\kappa_2(\mathcal{G}) \approx 48$; Figure 8b, Appendix Table 20), reaching 60.3 dB BF16 SNR and 80.1% noisy $r = 3$ recall vs. 2.67 dB and 50.0% for MesaNet;
2. uniform-decay 12.5B `15L:1A` (15.1B tok, Table 15), edging `15L SLA-R2` on 32K 8-N (39.8% vs. 34.8%) and $N = 512$ 64K `HotpotQA` (69.5% vs. 68.2%); and
3. 63.1B `15L:1A`, reaching 12.42 vs. 13.01 32K Tail PPL at 5.0B tok (Table 15) and 11.82 Tail PPL with $57.8\%/64.8\%$ 64K `SQuAD/HotpotQA` at 30.2B ($+12.5/+25.0$ pts over `0L:1A`; Table 16).

5 Related Work

Decay Granularity in Linear Recurrences. Linear attentions replace the growing softmax KV cache with a fixed-size state updated once per token [12, 33]. Using the decay-group view of Section 2.1, most follow-up models differ in how finely their forget gate is resolved across key channels: RetNet [58], Mamba-2 [14], and GDN [68] use one decay rate per head, whereas GLA [66], RWKV-6/7 [42, 43], KDA [71], and GDN-2 [25] give every key channel its own rate; Mamba [22], HGRN2 [46], xLSTM [7], and Mamba-3 [34] explore further gate and state designs. SLA takes a middle point with four decay rates per 256-channel head. This is coarse enough that the decay factors out of 64-channel MXU matmuls. In our sweep it slightly lowers perplexity relative to one rate per head across all three paired seeds but extrapolates somewhat worse, so the choice between the two is close (Section 3.5).

Delta Rules, Higher-Order Memories, and Momentum. DeltaNet [50, 67] and test-time-trained memories [8, 9, 59] view the state update as one step of online regression, which improves associative recall over decay-only models [5, 30]. Later work adds multiple Householder reflections [53], negative eigenvalues [20], decoupled per-channel erase and write gates [25], second-order preconditioning [62, 64], log-linear state hierarchies [23], and heavy-ball momentum [9, 41, 44]. These works improve what the update can express; SLA instead keeps GDN's single-reflection update and changes its shape for systolic hardware. The two directions are complementary: GDN-2 reaches lower per-token perplexity than SLA in our runs (Section 3.5), making a fused TPU kernel that combines GDN-2's decoupled gates with SLA's systolic mapping a natural next step.

Sharing Key Heads Across Value Heads. GQA [2] shares key and value heads across query heads in softmax attention. For linear attention, Kimi Linear [71], Qwen3-Next [48], and `fla` [65] already share query-key heads across value heads at $d_k = 128$, while Gated Attention [47] establishes sigmoid output gating. SLA uses the same idea at $d_k = 256$ with per-value-head normalization, choosing the sizes so that one key head fills a full MXU tile.

Hardware-Aware Kernels. FlashAttention [13, 15, 70] re-architects softmax attention around the GPU memory hierarchy. For linear recurrences, `fla` [65] provides chunkwise WY kernels [10] for GPU tensor cores, and MaxText and Tokamax [1, 19] port KDA to TPUs. These kernels optimize a fixed model shape; SLA also changes the model so that every intra-chunk stage maps onto the MXU and fits in one fused kernel (distinct from works that abbreviate sparse-linear or softmax-linear hybrids as SLA).

Context Parallelism for Linear Recurrences. Head-sharded Ulysses [29] requires at least as many key heads as devices, ring-based methods [36, 45, 56] pass the state through $\mathcal{O}(P)$ sequential hops, and LASP-2 [57] gathers chunk states with one AllGather. Because every gated delta chunk transition with scalar β_t is affine, chunks can also be combined with a parallel prefix scan [11, 39]; we use this standard construction for both SLA and KDA (Section 3.3).

6 Conclusion

Summary. Recent gated-delta models were shaped around GPUs, where per-channel gate rescaling and sequential in-block triangular solves are cheap relative to matmuls. On systolic accelerators, where matrix units outpace vector lanes by two orders of magnitude, kernel fusion alone cannot hide a vector-bound recurrence. Systolic Linear Attention (SLA) resolves this mismatch by co-designing the recurrence around the systolic array: sharing 256-wide keys across value heads fills full tiles, grouping forget gates into tiles factors decay out of matmuls, and block-doubling keeps chunk inversion and its backward pass on the MXU in one VMEM kernel. On TPU v7x, a fused `SLA-R4` layer runs $1.6\times$–$2.5\times$ faster than Tokamax KDA ($1.4\times$–$2.4\times$ with exact inversion) with 52%–61% less temporary HBM, and across 6.5B–63.1B MoE models on up to 512 chips per run, SLA trains 8%–24% faster than MaxText KDA at comparable perplexity without requiring linear-attention gate clamping at any scale (8K $\rightarrow$ 64K zero-shot on RoPE backbones, and 83.2%–84.4% 32K 8-Needle on 63.1B NoPE with global `SSMax` alone).

Future Work. SLA targets accelerators with wide matrix tiles and a large matrix-to-vector gap; on GPUs, small Tensor Core tiles and fast vector lanes already saturate narrow states. Because our speedups combine tile-grouped decay with head sharing, BF16 state storage, and VMEM fusion while GDN-2 uses an unfused port despite lower per-token perplexity, next steps include building fused TPU kernels for wide-key, shared-head KDA and GDN-2 to isolate each factor and combine GDN-2's decoupled gates with SLA's systolic mapping, closing SLA's ≈ 0.1 perplexity gap to wide-key KDA (0.3 with the 1-step Neumann inverse) on matched seeds, and decoupling chunk-boundary momentum from C. Finally, extending multi-seed runs beyond 8K Dolma, where BF16 MoE routing variance (≈ 0.18 perplexity) masks small gaps, 63.1B runs are single-seed, and 64K QA sets are small, will sharpen long-context comparisons.

References

[1] AI-Hypercomputer Contributors. MaxText: A high performance, highly scalable, open-source LLM library written in pure Python/JAX. https://github.com/AI-Hypercomputer/maxtext, 2026.

[2] J. Ainslie, J. Lee-Thorp, M. de Jong, Y. Zemlyanskiy, F. Lebrón, and S. Sanghai. GQA: Training generalized multi-query transformer models from multi-head checkpoints. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2023. arXiv:2305.13245.

[3] N. Amsel, D. Persson, C. Musco, and R. M. Gower. The Polar Express: Optimal matrix sign methods and their application to the Muon algorithm. *arXiv preprint arXiv:2505.16932*, 2025.

[4] R. Anil, V. Gupta, T. Koren, K. Regan, and Y. Singer. Scalable second order optimization for deep learning. *arXiv preprint arXiv:2002.09018*, 2020.

[5] S. Arora, S. Eyuboglu, A. Timalsina, I. Johnson, M. Poli, J. Zou, A. Rudra, and C. Ré. Zoology: Measuring and improving recall in efficient language models. In *International Conference on Learning Representations (ICLR)*, 2024.

[6] J. Austin, S. Douglas, R. Frostig, A. Levskaya, C. Chen, S. Vikram, F. Lebron, P. Choy, V. Ramasesh, A. Webson, and R. Pope. How to scale your model. https://jax-ml.github.io/scaling-book/, 2025.

[7] M. Beck, K. Pöppel, M. Spanring, A. Auer, O. Prudnikova, M. Kopp, G. Klambauer, J. Brandstetter, and S. Hochreiter. xLSTM: Extended long short-term memory. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

[8] A. Behrouz, Z. Li, P. Kacham, M. Daliri, Y. Deng, P. Zhong, M. Razaviyayn, and V. Mirrokni. Atlas: Learning to optimally memorize the context at test time. *arXiv preprint arXiv:2505.23735*, 2025.

[9] A. Behrouz, P. Zhong, and V. Mirrokni. Titans: Learning to memorize at test time. *arXiv preprint arXiv:2501.00663*, 2025.

[10] C. Bischof and C. Van Loan. The WY representation for products of Householder matrices. *SIAM Journal on Scientific and Statistical Computing*, 8(1):s2–s13, 1987.

[11] G. E. Blelloch. Prefix sums and their applications. Technical Report CMU-CS-90-190, Carnegie Mellon University, 1990.

[12] K. Choromanski, V. Likhosherstov, D. Dohan, X. Song, A. Gane, T. Sarlos, P. Hawkins, J. Davis, A. Mohiuddin, L. Kaiser, D. Belanger, L. Colwell, and A. Weller. Rethinking attention with performers. In *International Conference on Learning Representations (ICLR)*, 2021.

[13] T. Dao. FlashAttention-2: Faster attention with better parallelism and work partitioning. In *International Conference on Learning Representations (ICLR)*, 2024.

[14] T. Dao and A. Gu. Transformers are SSMs: Generalized models and efficient algorithms through structured state space duality. In *International Conference on Machine Learning (ICML)*, 2024.

[15] T. Dao, D. Y. Fu, S. Ermon, A. Rudra, and C. Ré. FlashAttention: Fast and memory-efficient exact attention with IO-awareness. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

[16] W. Fedus, B. Zoph, and N. Shazeer. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *Journal of Machine Learning Research (JMLR)*, 23(120): 1–39, 2022. arXiv:2101.03961.

[17] L. Gao, J. Tow, B. Abbasi, S. Biderman, S. Black, A. DiPofi, C. Foster, L. Golding, J. Hsu, A. L. Noac'h, H. Li, K. McDonell, N. Muennighoff, C. Ociepa, J. Phang, L. Reynolds, H. Schoelkopf, A. Skowron, L. Sutawika, E. Tang, A. Thite, B. Wang, K. Wang, and A. Zou. A framework for few-shot language model evaluation. Zenodo, https://doi.org/10.5281/zenodo.10256836, 2023.

[18] Google Cloud. TPU7x (Ironwood) architecture and performance optimizations. https://docs.cloud.google.com/tpu/docs/ironwood-performance, 2025.

[19] Google DeepMind. Tokamax: Open-source accelerator kernel library for JAX. https://github.com/google-deepmind/tokamax, 2026.

[20] R. Grazzi, J. Siems, J. K. H. Franke, A. Zela, F. Hutter, and M. Pontil. Unlocking state-tracking in linear RNNs through negative eigenvalues. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2411.12537.

[21] E. Grishina, M. Smirnov, and M. Rakhuba. Accelerating Newton-Schulz iteration for orthogonalization via Chebyshev-type polynomials. *arXiv preprint arXiv:2506.10935*, 2025.

[22] A. Gu and T. Dao. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv preprint arXiv:2312.00752*, 2023.

[23] H. Guo, S. Yang, T. Goel, E. P. Xing, T. Dao, and Y. Kim. Log-linear attention. *arXiv preprint arXiv:2506.04761*, 2025.

[24] V. Gupta, T. Koren, and Y. Singer. Shampoo: Preconditioned stochastic tensor optimization. In *International Conference on Machine Learning (ICML)*, 2018.

[25] A. Hatamizadeh, Y. Choi, and J. Kautz. Gated DeltaNet-2: Decoupling erase and write in linear attention. *arXiv preprint arXiv:2605.22791*, 2026.

[26] N. J. Higham. *Functions of Matrices: Theory and Computation*. Society for Industrial and Applied Mathematics (SIAM), 2008.

[27] J. Hoffmann, S. Borgeaud, A. Mensch, E. Buchatskaya, T. Cai, E. Rutherford, D. de Las Casas, L. A. Hendricks, J. Welbl, A. Clark, T. Hennigan, E. Noland, K. Millican, G. van den Driessche, B. Damoc, A. Guy, S. Osindero, K. Simonyan, E. Elsen, J. W. Rae, O. Vinyals, and L. Sifre. Training compute-optimal large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2203.15556.

[28] C.-P. Hsieh, S. Sun, S. Kriman, S. Acharya, D. Rekesh, F. Jia, Y. Zhang, and B. Ginsburg. RULER: What's the real context size of your long-context language models? In *Conference on Language Modeling (COLM)*, 2024.

[29] S. A. Jacobs, M. Tanaka, C. Zhang, M. Zhang, S. L. Song, S. Rajbhandari, and Y. He. DeepSpeed Ulysses: System optimizations for enabling training of extreme long sequence transformer models. *arXiv preprint arXiv:2309.14509*, 2023.

[30] S. Jelassi, D. Brandfonbrener, S. M. Kakade, and E. Malach. Repeat after me: Transformers are better than state space models at copying. In *International Conference on Machine Learning (ICML)*, 2024.

[31] N. P. Jouppi, C. Young, N. Patil, D. Patterson, G. Agrawal, R. Bajwa, S. Bates, S. Bhatia, N. Boden, A. Borchers, R. Boyle, P.-L. Cantin, C. Chao, C. Clark, J. Coriell, M. Daley, M. Dau, J. Dean, B. Gelb, T. V. Ghaemmaghami, R. Gottipati, W. Gulland, R. Hagmann, C. R. Ho, D. Hogberg, J. Hu, R. Hundt, D. Hurt, J. Ibarz, A. Jaffey, A. Jaworski, A. Kaplan, H. Khaitan, A. Koch, N. Kumar, S. Lacy, J. Laudon, J. Law, D. Le, C. Leary, Z. Liu, K. Lucke, A. Lundin, G. MacKean, A. Maggiore, M. Mahony, K. Miller, R. Nagarajan, R. Narayanaswami, R. Ni, K. Nix, T. Norrie, M. Omernick, N. Penukonda, A. Phelps, J. Ross, M. Ross, A. Salek, E. Samadiani, C. Severn, G. Sizikov, M. Snelham, J. Souter, D. Steinberg, A. Swing, M. Tan, G. Thorson, B. Tian, H. Toma, E. Tuttle, V. Vasudevan, R. Walter, W. Wang, E. Wilcox, and D. H. Yoon. In-datacenter performance analysis of a tensor processing unit. In *International Symposium on Computer Architecture (ISCA)*, 2017.

[32] N. P. Jouppi, G. Kurian, S. Li, P. Ma, R. Nagarajan, L. Nai, N. Patil, S. Subramanian, A. Swing, B. Towles, C. Young, X. Zhou, Z. Zhou, and D. Patterson. TPU v4: An optically reconfigurable supercomputer for machine learning with hardware support for embeddings. In *International Symposium on Computer Architecture (ISCA)*, 2023.

[33] A. Katharopoulos, A. Vyas, N. Pappas, and F. Fleuret. Transformers are RNNs: Fast autoregressive transformers with linear attention. In *International Conference on Machine Learning (ICML)*, 2020.

[34] A. Lahoti, K. Y. Li, B. Chen, C. Wang, A. Bick, J. Z. Kolter, T. Dao, and A. Gu. Mamba-3: Improved sequence modeling using state space principles. In *International Conference on Learning Representations (ICLR)*, 2026. arXiv:2603.15569.

[35] Y. Leviathan, M. Kalman, and Y. Matias. Fast inference from transformers via speculative decoding. In *International Conference on Machine Learning (ICML)*, 2023. arXiv:2211.17192.

[36] H. Liu, M. Zaharia, and P. Abbeel. Ring attention with blockwise transformers for near-infinite context. In *International Conference on Learning Representations (ICLR)*, 2024.

[37] J. Liu, J. Su, X. Yao, Z. Jiang, G. Lai, Y. Du, Y. Qin, W. Xu, E. Lu, J. Yan, Y. Chen, H. Zheng, Y. Liu, S. Liu, B. Yin, W. He, H. Zhu, Y. Wang, J. Wang, M. Dong, Z. Zhang, Y. Kang, H. Zhang, X. Xu, Y. Zhang, Y. Wu, X. Zhou, and Z. Yang. Muon is scalable for LLM training. *arXiv preprint arXiv:2502.16982*, 2025.

[38] I. Loshchilov and F. Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations (ICLR)*, 2019. arXiv:1711.05101.

[39] E. Martin and C. Cundy. Parallelizing linear recurrent neural nets over sequence length. In *International Conference on Learning Representations (ICLR)*, 2018.

[40] K. M. Nakanishi. Scalable-softmax is superior for attention. *arXiv preprint arXiv:2501.19399*, 2025.

[41] T. M. Nguyen, R. G. Baraniuk, R. M. Kirby, S. J. Osher, and B. Wang. Momentum transformer: Closing the performance gap between self-attention and its linearization. *arXiv preprint arXiv:2208.00579*, 2022.

[42] B. Peng, D. Goldstein, Q. Anthony, A. Albalak, E. Alcaide, S. Biderman, E. Cheah, X. Du, T. Ferdinan, H. Hou, P. Kazienko, K. K. GV, J. Kocoń, B. Koptyra, S. Krishna, R. McClelland, Jr., J. Lin, N. Muennighoff, F. Obeid, A. Saito, G. Song, H. Tu, C. Wirawan, S. Woźniak, R. Zhang, B. Zhao, Q. Zhao, P. Zhou, J. Zhu, and R.-J. Zhu. Eagle and finch: RWKV with matrix-valued states and dynamic recurrence. In *Conference on Language Modeling (COLM)*, 2024.

[43] B. Peng, R. Zhang, D. Goldstein, E. Alcaide, X. Du, H. Hou, J. Lin, J. Liu, J. Lu, W. Merrill, G. Song, K. Tan, S. Utpala, N. Wilce, J. S. Wind, T. Wu, D. Wuttke, and C. Zhou-Zheng. RWKV-7 "goose" with expressive dynamic state evolution. *arXiv preprint arXiv:2503.14456*, 2025.

[44] B. T. Polyak. Some methods of speeding up the convergence of iteration methods. *USSR Computational Mathematics and Mathematical Physics*, 4(5):1–17, 1964.

[45] Z. Qin, W. Sun, D. Li, X. Shen, W. Sun, and Y. Zhong. Various lengths, constant speed: Efficient language modeling with lightning attention. In *International Conference on Machine Learning (ICML)*, 2024. arXiv:2405.17381.

[46] Z. Qin, S. Yang, W. Sun, X. Shen, D. Li, W. Sun, and Y. Zhong. HGRN2: Gated linear RNNs with state expansion. In *Conference on Language Modeling (COLM)*, 2024.

[47] Z. Qiu, Z. Wang, B. Zheng, Z. Huang, K. Wen, S. Yang, R. Men, L. Yu, F. Huang, S. Huang, D. Liu, J. Zhou, and J. Lin. Gated attention for large language models: Non-linearity, sparsity, and attention-sink-free. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. arXiv:2505.06708.

[48] Qwen Team. Qwen3-Next: Towards ultimate training and inference efficiency. https://qwenlm.github.io/blog/qwen3-next/, 2025.

[49] P. Rajpurkar, J. Zhang, K. Lopyrev, and P. Liang. SQuAD: 100,000+ questions for machine comprehension of text. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2016. arXiv:1606.05250.

[50] I. Schlag, K. Irie, and J. Schmidhuber. Linear transformers are secretly fast weight programmers. In *International Conference on Machine Learning (ICML)*, 2021.

[51] N. Shazeer. GLU variants improve transformer. *arXiv preprint arXiv:2002.05202*, 2020.

[52] N. Shazeer, A. Mirhoseini, K. Maziarz, A. Davis, Q. Le, G. Hinton, and J. Dean. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. In *International Conference on Learning Representations (ICLR)*, 2017. arXiv:1701.03538.

[53] J. Siems, T. Carstensen, A. Zela, F. Hutter, M. Pontil, and R. Grazzi. DeltaProduct: Improving state-tracking in linear RNNs via Householder products. In *International Conference on Learning Representations (ICLR)*, 2026. arXiv:2502.10297.

[54] L. Soldaini, R. Kinney, A. Bhagia, D. Schwenk, D. Atkinson, R. Authur, B. Bogin, K. Chandu, J. Dumas, Y. Elazar, V. Hofmann, A. H. Jha, S. Kumar, L. Lucy, X. Lyu, N. Lambert, I. Magnusson, J. Morrison, N. Muennighoff, A. Naik, C. Nam, M. E. Peters, A. Ravichander, K. Richardson, Z. Shen, E. Strubell, N. Subramani, O. Tafjord, P. Walsh, L. Zettlemoyer, N. A. Smith, H. Hajishirzi, I. Beltagy, D. Groeneveld, J. Dodge, and K. Lo. Dolma: an open corpus of three trillion tokens for language model pretraining research. In *Association for Computational Linguistics (ACL)*, 2024.

[55] J. Su, M. Ahmed, Y. Lu, S. Pan, W. Bo, and Y. Liu. RoFormer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024. arXiv:2104.09864.

[56] W. Sun, Z. Qin, D. Li, X. Shen, Y. Qiao, and Y. Zhong. Linear attention sequence parallelism. *arXiv preprint arXiv:2404.02882*, 2024.

[57] W. Sun, Z. Qin, D. Li, X. Shen, Y. Qiao, and Y. Zhong. LASP-2: Rethinking sequence parallelism for linear attention and its hybrid. *arXiv preprint arXiv:2502.07563*, 2025.

[58] Y. Sun, L. Dong, S. Huang, S. Ma, Y. Xia, J. Xue, J. Wang, and F. Wei. Retentive network: A successor to transformer for large language models. *arXiv preprint arXiv:2307.08621*, 2023.

[59] Y. Sun, X. Li, K. Dalal, J. Xu, A. Vikram, G. Zhang, Y. Dubois, X. Chen, X. Wang, S. Koyejo, T. Hashimoto, and C. Guestrin. Learning to (learn at test time): RNNs with expressive hidden states. In *International Conference on Machine Learning (ICML)*, 2024. arXiv:2407.04620.

[60] Team Olmo. Olmo 3. *arXiv preprint arXiv:2512.13961*, 2025.

[61] Team OLMo, P. Walsh, L. Soldaini, D. Groeneveld, K. Lo, S. Arora, A. Bhagia, Y. Gu, S. Huang, M. Jordan, N. Lambert, D. Schwenk, O. Tafjord, T. Anderson, D. Atkinson, F. Brahman, C. Clark, P. Dasigi, N. Dziri, A. Ettinger, M. Guerquin, D. Heineman, H. Ivison, P. W. Koh, J. Liu, S. Malik, W. Merrill, L. J. V. Miranda, J. Morrison, T. Murray, C. Nam, J. Poznanski, V. Pyatkin, A. Rangapur, M. Schmitz, S. Skjonsberg, D. Wadden, C. Wilhelm, M. Wilson, L. Zettlemoyer, A. Farhadi, N. A. Smith, and H. Hajishirzi. 2 OLMo 2 furious. *arXiv preprint arXiv:2501.00656*, 2025.

[62] N. Tumma, N. Loo, and D. Rus. Preconditioned DeltaNet: Curvature-aware sequence modeling for linear recurrences. *arXiv preprint arXiv:2604.21100*, 2026.

[63] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.

[64] J. von Oswald, N. Scherrer, S. Kobayashi, L. Versari, S. Yang, S. Mittal, M. Schlegel, K. Maile, Y. Schimpf, O. Sieberling, A. Meulemans, R. A. Saurous, G. Lajoie, C. Frenkel, R. Pascanu, B. A. y Arcas, and J. Sacramento. MesaNet: Sequence modeling by locally optimal test-time training. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. arXiv:2506.05233.

[65] S. Yang and Y. Zhang. FLA: A Triton-based library for hardware-efficient implementations of linear attention mechanism. `https://github.com/fla-org/flash-linear-attention`, 2024.

[66] S. Yang, B. Wang, Y. Shen, R. Panda, and Y. Kim. Gated linear attention transformers with hardware-efficient training. In *International Conference on Machine Learning (ICML)*, 2024.

[67] S. Yang, B. Wang, Y. Zhang, Y. Shen, and Y. Kim. Parallelizing linear transformers with the delta rule over sequence length. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

[68] S. Yang, J. Kautz, and A. Hatamizadeh. Gated Delta networks: Improving Mamba2 with Delta rule. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2412.06464.

[69] Z. Yang, P. Qi, S. Zhang, Y. Bengio, W. W. Cohen, R. Salakhutdinov, and C. D. Manning. HotpotQA: A dataset for diverse, explainable multi-hop question answering. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2018. arXiv:1809.09600.

[70] T. Zadouri, M. Hoehnerbach, J. Shah, T. Liu, V. Thakkar, and T. Dao. FlashAttention-4: Algorithm and kernel pipelining co-design for asymmetric hardware scaling. *arXiv preprint arXiv:2603.05451*, 2026.

[71] Y. Zhang, Z. Lin, X. Yao, J. Hu, F. Meng, C. Liu, X. Men, S. Yang, Z. Li, Z. Yang, Y. Du, and Kimi Team. Kimi Linear: An expressive, efficient attention architecture. *arXiv preprint arXiv:2510.26692*, 2025.

A Proofs of Propositions 4.1 and 4.2 (Whitening and Polar-2+)

Part 1 (Euclidean Interference and Diagonal Non-Invariance). Let unit keys $k_1, \dots, k_N \in \mathbb{R}^{d_k}$ satisfy $|\langle k_i, k_j \rangle| = \nu_{\text{coh}}$ for $i \neq j$ (coherence ν_{coh}; e.g., an equiangular tight frame or shared-subspace correlated dictionary) and draw independent zero-mean values $v_j \in \mathbb{R}^{d_v}$ with covariance $\mathbb{E}[v_j v_j^\top] = \frac{s_v^2}{d_v} I_{d_v}$ (for $d_v \times d_v$ identity I_{d_v}), so $s_v^2 = \operatorname{tr}\big(\frac{s_v^2}{d_v} I_{d_v}\big) = \mathbb{E}\,\|v_j\|_2^2$. Under single-pass Euclidean delta updates with $\beta_t = 1$ from $S_0 = 0$, updating $S_t = (I_{d_k} - k_t k_t^\top) S_{t-1} + k_t v_t^\top$ writes later values $j > i$ with projection $\langle k_j, k_i \rangle v_j$ along k_i: for $i = N - 1$,

$$S_N^\top k_{N-1} - v_{N-1} = \langle k_N, k_{N-1} \rangle \big(v_N - S_{N-1}^\top k_N\big),$$

and summing independent cross-key leakage over $\Theta(N)$ interfering keys yields

$$\mathbb{E}\,\big\|S_N^\top k_i - v_i\big\|_2^2 = \Omega(N \nu_{\text{coh}}^2 s_v^2) = \Omega\big(\rho_{\text{eff}} (d_k \nu_{\text{coh}}^2) s_v^2\big)$$

when $N \nu_{\text{coh}}^2 \leq 1/2$ ($\rho_{\text{eff}} = N/d_k$, reducing to $\Omega(\rho_{\text{eff}} s_v^2)$ at isotropic $\nu_{\text{coh}}^2 = 1/d_k$).

For the diagonal column-energy preconditioner $K \operatorname{diag}(p(K))^{-1/2}$ where $p_m(K) = \sum_{i=1}^N K_{i,m}^2$, rotating $K \mapsto KQ$ with $Q = \operatorname{diag}(\xi) Q_H \in \mathrm{O}(d_k)$ for random Rademacher signs $\xi \in \{\pm 1\}^{d_k}$ followed by a normalized Hadamard matrix $Q_H \in \{\pm d_k^{-1/2}\}^{d_k \times d_k}$ equalizes all column energies in expectation,

$$\mathbb{E}_\xi\big[p_m(K \operatorname{diag}(\xi) Q_H)\big] = \frac{1}{d_k} \|K\|_F^2$$

(and by Schur-Horn some rotation $Q \in \mathrm{O}(d_k)$ makes $Q^\top K^\top K Q$ exactly constant-diagonal), reducing $\operatorname{diag}(p(KQ))^{-1/2}$ to a scalar multiple of I_{d_k} that leaves $\kappa_2((KQ)(KQ)^\top) = \kappa_2(KK^\top)$ unchanged; by contrast, $(I_C + \eta_c K_c K_c^\top)^{-1/2} K_c$ transforms equivariantly ($K_c \mapsto K_c Q \implies \tilde{K}_c \mapsto \tilde{K}_c Q$) for all $Q \in \mathrm{O}(d_k)$.

Part 2 (Gated RLS Overwrite-Retention Trade-off). Under gated RLS [64], with key covariance $H_t \in \mathbb{R}^{d_k \times d_k}$ and key-value moment $B_t \in \mathbb{R}^{d_k \times d_v}$ (local to this paragraph), $H_t = \gamma H_{t-1} + k_t k_t^\top$ and $B_t = \gamma B_{t-1} + k_t v_t^\top$ with scalar forget factor $\gamma \in (0, 1]$ and ridge ϵI_{d_k} ($\epsilon > 0$); readout on query $k \in \mathbb{R}^{d_k}$ is $B_t^\top H_t^{-1} k$. If a unit key k is written $r \geq 2$ times, with value $v^{(m)}$ at step t_m ($t_1 < \cdots < t_r = t$) and weights $w_m = \gamma^{t - t_m}$, Sherman-Morrison on $H_t = \epsilon I_{d_k} + (\sum_{m=1}^r w_m) k k^\top$ gives

$$\lim_{\epsilon \to 0^+} B_t^\top H_t^{-1} k = \frac{\sum_{m=1}^r w_m v^{(m)}}{\sum_{m=1}^r w_m}.$$

Letting $\bar{v}_{1:r-1} = \frac{1}{r-1} \sum_{m=1}^{r-1} v^{(m)}$ denote the mean of the first $r - 1$ values, readout at $\gamma = 1$ ($w_m \equiv 1$) equals $\frac{1}{r} \sum_{m=1}^r v^{(m)}$ with error $\frac{r-1}{r} \big\|v^{(r)} - \bar{v}_{1:r-1}\big\|_2$ relative to $v^{(r)}$. Already for $r = 2$ consecutive writes and any target relative tolerance $\varepsilon \in (0, 1/2)$, the error is $\frac{\gamma}{1+\gamma} \big\|v^{(2)} - v^{(1)}\big\|_2$, so making it at most $\varepsilon \big\|v^{(2)} - v^{(1)}\big\|_2$ requires $\gamma \leq \varepsilon/(1 - \varepsilon)$, which simultaneously decays any non-overwritten key stored $t - j$ steps earlier (at step $j < t$) by γ^{t-j} relative to the regularizer ϵI_{d_k}, causing rapid forgetting in finite precision (verified in Table 20). Under oblique whitened delta, intra-chunk Gram whitening orthogonalizes concurrent keys while the oblique erasure $(I_{d_k} - \beta_t \tilde{k}_t k_t^\top)$ subtracts the current state projection $\hat{v}_t = (F_t \tilde{S}_{t-1})^\top k_t$ selectively along $\tilde{k}_t$ without decaying orthogonal associations.

Part 3 (Polar-2+ Non-Circular Normalization, Equioscillation, and BF16 Safety). Let $\Sigma_{c-1} \succ 0$ (initialized from $\Sigma_0 = I_{d_k}$) and let $Y_c = K_c \Sigma_{c-1}^{1/2} \in \mathbb{R}^{C \times d_k}$ have singular values $\sigma_i \geq 0$ (unrelated to the sigmoid σ), so the unregularized Gram $\mathcal{G}_0 = I_C + \eta_c Y_c Y_c^\top$ ($\eta_c \in (0, 1]$) has eigenvalues $\omega_{0,i} = 1 + \eta_c \sigma_i^2$. Computing $\bar{\mathcal{G}}_0 = \mathcal{G}_0 / \|\mathcal{G}_0\|_F$ and

$$\hat{u}_0 = \max\big(\|\mathcal{G}_0\|_F \big\|\bar{\mathcal{G}}_0^2\big\|_F^{1/2}, 1\big) = \max\left(\Big(\sum_j \omega_{0,j}^4\Big)^{1/4}, 1\right) \geq \max_i \omega_{0,i}$$

yields a non-circular upper bound on $\lambda_{\max}(\mathcal{G}_0)$ from one $[C, C]$ matmul $\bar{\mathcal{G}}_0^2$. Setting adaptive ridge $\varepsilon_{\text{reg}} = \max\left(0.02, \frac{0.0205}{0.9795}\hat{u}_0 - 1\right)$, $\mathcal{G}_{\text{reg}} = \mathcal{G}_0 + \varepsilon_{\text{reg}} I_C$ with shifted eigenvalues $\omega_i = \omega_{0,i} + \varepsilon_{\text{reg}}$, and $u_0 = \hat{u}_0 + \varepsilon_{\text{reg}} \geq \max_i \omega_i$ guarantees normalized eigenvalues $\lambda_i = \omega_i / u_0 \in [0.02, 1.0]$ for all $\hat{u}_0 \geq 1$ while expanding

$$T_0^2 = \frac{(\mathcal{G}_0 + \varepsilon_{\text{reg}} I_C)^2}{u_0^2} = \frac{\|\mathcal{G}_0\|_F^2 \bar{\mathcal{G}}_0^2 + 2\varepsilon_{\text{reg}} \mathcal{G}_0 + \varepsilon_{\text{reg}}^2 I_C}{u_0^2}$$

with zero additional matmuls (5 matmuls total across stages $\ell \in \{0, 1\}$, where $\mathcal{M}_\ell = m_\ell(T_\ell)$ evaluates $m_0(\lambda) = a_0 + b_0\lambda + e_0\lambda^2$ and $m_1(x) = a_1 + b_1 x + e_1 x^2$).

Clamping $\beta_t \leq (2 - \varepsilon_\beta)/\max(k_t^\top \tilde{k}_t, \varepsilon_\beta)$ keeps the nontrivial eigenvalue $1 - \beta_t k_t^\top \tilde{k}_t$ of $I_{d_k} - \beta_t \tilde{k}_t k_t^\top$ in $[-1 + \varepsilon_\beta, 1]$ whenever $k_t^\top \tilde{k}_t > 0$ (in packed LM training, intra-chunk causal Schur-oblique projection sets

$$\tilde{k}_t = \frac{(I_{d_k} - \eta_c \Sigma_{c-1} K_{<t}^\top K_{<t}) \Sigma_{c-1} k_t}{k_t^\top (I_{d_k} - \eta_c \Sigma_{c-1} K_{<t}^\top K_{<t}) \Sigma_{c-1} k_t}$$

so $k_t^\top \tilde{k}_t = 1$ depends strictly on rows $< t$, while symmetric `Polar-2+` governs the chunk-boundary Woodbury tail, two-sided chunk-whitening kernels, and lagged block-32 causal whitening). Substituting $\mathcal{W}_{\text{isq}} = \alpha_{\text{safe}} u_0^{-1/2} \mathcal{M}_1 \mathcal{M}_0 = \alpha_{\text{safe}} u_0^{-1/2} z_2(T_0)$ and $\tilde{K}_c = \sqrt{\eta_c} \mathcal{W}_{\text{isq}} K_c \Sigma_{c-1}$ (Eq. (14)) into the unrelaxed Woodbury update $\tilde{\Sigma}_c = \Sigma_{c-1} - \tilde{K}_c^\top \tilde{K}_c$ (Eq. (15) at $\rho_p = 0$) gives $\tilde{\Sigma}_c = \Sigma_{c-1}^{1/2}(I_{d_k} - \eta_c Y_c^\top \mathcal{W}_{\text{isq}}^2 Y_c)\Sigma_{c-1}^{1/2}$, whose i-th singular-direction multiplier is

$$1 - \eta_c \sigma_i^2 u_0^{-1} \alpha_{\text{safe}}^2 z_2(\lambda_i)^2 = 1 - \frac{\eta_c \sigma_i^2}{1 + \varepsilon_{\text{reg}} + \eta_c \sigma_i^2} \lambda_i \left(\alpha_{\text{safe}} z_2(\lambda_i)\right)^2. \tag{17}$$

Direct evaluation of the scalar map z_2 (Proposition 4.2) on $\lambda \in [0.02, 1.0]$ gives minimax equioscillation $\sqrt{\lambda} z_2(\lambda) \in [0.9757, 1.0243]$, hence $\lambda z_2(\lambda)^2 \in [0.9520, 1.0492]$ (and < 0.952 on $(0, 0.02)$ since $z_2(0) = a_1 a_0 \approx 9.877$). Thus unscaled $\alpha = 1$ exceeds 1 by $+4.92\%$ near equioscillation peaks, driving Eq. (17) negative when $\eta_c \sigma_i^2 \gg 1$. Setting $\alpha_{\text{safe}} = 0.968$ ($\alpha_{\text{safe}}^2 = 0.937024$) maps exact $\lambda(\alpha_{\text{safe}} z_2(\lambda))^2$ onto $[0.8920, 0.9831]$ across $[0.02, 1.0]$, giving exact $\kappa_2(\mathcal{W}_{\text{isq}} \mathcal{G}_{\text{reg}} \mathcal{W}_{\text{isq}}) = 0.9831/0.8920 = 1.102$ (≤ 1.123 in 5-matmul BF16 hardware execution). At $\rho_p = 0$, 4,096 correlated BF16 chunks (524,288 tokens) maintain $\lambda_{\min}(\Sigma_c) \geq 0.0388 > 0$; when evolving Σ_c across thousands of packed BF16 LM gradient steps with $\rho_p < 1$, replacing additive subtraction with multiplicative Sylvester congruence $\mathcal{B}_c \Sigma_{c-1} \mathcal{B}_c^\top$ (with contraction matrix $\mathcal{B}_c = I_{d_k} - \frac{\eta_c}{2} \Sigma_{c-1} K_c^\top \mathcal{W}_{\text{isq}}^2 K_c$, or $\mathcal{B}_c = I_{d_k} - \hat{C}_c$ using row-sum–normalized key covariance $\hat{C}_c \propto \eta_c K_c^\top K_c$ in our kernel) enforces PSD preservation unconditionally (Appendix C). □

B Associative Recall, Gated RLS Trade-Off, and Gate Conditioning

Table 20 reports synthetic associative-recall SNR (dB) in both BF16 and FP32 alongside Mod-2 parity accuracy and the chunk forget-gate condition number $\kappa_{\text{forget}} = \kappa_2(F_c)$ ($d_k = 256, C = 128$, 8 seeds, correlated keys with $\kappa_2(\mathcal{G}) \approx 48$). Separating BF16 from FP32 clarifies precision ceilings: exact inverse-square-root whitening reaches > 140 dB ($r = 2, 3$ overwrite) and 51.9 dB (write-once) in FP32, whereas BF16 readout caps SNR near 59–60 dB due to 7-bit mantissa rounding ($20 \log_{10}(2^7\sqrt{12}) \approx 53$–$60$ dB). Notably, unwhitened Euclidean delta rules reach 58.40–60.97 dB on the final $r = 3$ overwrite via recency-weighted last-write erasure, whereas at $r = 2$ KDA's channel decay ($\kappa_{\text{forget}} = 1.1 \times 10^7$) collapses to 0.07 dB (and write-once -0.84 dB) and unwhitened SLA ($\kappa_{\text{forget}} = 1.00$) drops to 24.10 dB vs. 60.33 dB for `SPD-SLA-W`. Sweeping gated MesaNet's forget factor $\gamma \in \{1.00, 0.98, 0.80\}$ directly confirms Proposition 4.1(2): $\gamma = 1.00$ and $\gamma = 0.98$ suffer severe stale-value averaging (1.90 and 4.13 dB at $r = 3$), whereas $\gamma = 0.80$ improves overwrite SNR (29.53 dB) only by degrading BF16 write-once retention from 10.63 dB down to 5.78 dB.

C Negative Results and Reproducibility Details

We record three negative results observed during hardware-algorithm co-design alongside baseline software versions:

Table 20: Measured associative recall SNR (dB, reported as **BF16 [FP32]**) and Mod-2 parity accuracy across linear attention update rules ($d_k = 256, C = 128$, 8 seeds, $\kappa_2(\mathcal{G}) \approx 48$); $\kappa_{\text{forget}} = \kappa_2(F_c)$ is the condition number of the chunk forget gate.

Mechanism / Update Rule	Axis-Aligned (FP32)	Hadamard-Rot. (FP32)	Write-Once BF16 [FP32]	$r = 2$ Overwrite BF16 [FP32]	$r = 3$ Overwrite BF16 [FP32]	Mod-2 Parity (κ_{forget})
Mamba-2 / SSD (Scalar Hebbian)	2.14 dB	2.11 dB	1.85 [2.11]	1.72 [1.89]	1.40 [1.52]	50.2% (1.00)
KDA Real Channel Decay ($\beta \in (0, 1)$)	5.31 dB	5.29 dB	−0.84 [−0.91]	0.07 [0.26]	60.97 [> 140]	51.2% (1.1×10^7)
Diagonal Preconditioned Delta [62]	22.41 dB	6.12 dB	4.92 [6.12]	16.80 [19.05]	18.10 [21.40]	51.4% (1.1×10^7)
MesaNet Unforgetted ($\gamma = 1.00$) [64]	95.08 dB	95.06 dB	**10.63** [22.79]	2.67 [3.23]	1.90 [1.73]	50.8% (1.00)
MesaNet Gated ($\gamma = 0.98$)	91.40 dB	91.38 dB	9.54 [22.56]	7.32 [8.35]	4.13 [4.07]	51.0% (1.00)
MesaNet Aggressively Gated ($\gamma = 0.80$)	74.20 dB	74.18 dB	5.78 [26.09]	47.68 [89.36]	29.53 [30.22]	51.1% (1.00)
SLA ($G = 4$ + SO(2) + $\beta \in (0, 2)$)	7.84 dB	7.82 dB	6.12 [7.82]	24.10 [26.74]	58.40 [> 140]	**82.2%** (**1.00**)
SPD-SLA-W (Polar-2+ Oblique Delta)	**95.05** dB	**95.05** dB	7.01 [**51.93**]	**60.33** [> **140**]	**59.26** [> **140**]	**82.0%** (**1.00**)

1. **Tier-1 VMEM Fusion Without Shape Co-Design** ($1.03\times$)**.** Fusing Ops 2–7a on GPU-shaped $8 \times 128 \times 256$ KDA yields only 3% forward-backward speedup on TPU v7x ($7.108 \rightarrow 6.912$ ms; Table 5) because the sequential in-block substitution of the triangular solve and the 50%-padded FP32 state scans remain VPU- and compute-bound inside VMEM.

2. **Unscaled Minimax Polynomials, BF16 Cholesky, and Additive Woodbury Drift in Live LM Training.** Using exact-arithmetic minimax coefficients ($\alpha = 1.000$) causes $\lambda z_2(\lambda)^2$ to reach $1.0492 > 1$, driving BF16 Woodbury covariance Σ_c indefinite ($\lambda_{\min} = -0.0412$; Table 19) and diverging 12.5B MoE pre-training to 258.5 PPL by step 299, while direct BF16 Cholesky/Woodbury hits `NaN` at step 54. Scaling by $\alpha_{\text{safe}} = 0.968$ bounds the exact peak to 0.9831, preserves $\lambda_{\min} \geq 0.0388$ across 4,096 chunks (524K tokens at $\rho_p = 0$), and trains cleanly for 7,200 steps across 2 seeds at $\rho_p = 1.0$ (0 NaNs; seed-matched `s1`: 17.936 vs. 17.970 PPL for `NS-5`, 2-seed `Polar-2+` mean 17.970 on the `+MS` arms [`s0` 18.004]; the uniform-decay arm reaches 17.575 in Table 15). When evolving Σ_c across thousands of packed BF16 LM gradient steps with $\rho_p = 0.02$, additive subtraction $\Sigma_{c-1} - \tilde{K}_c^\top \tilde{K}_c$ still accumulates BF16 rounding drift (`NaN` at step 282), whereas multiplicative Sylvester congruence $\mathcal{B}_c \Sigma_{c-1} \mathcal{B}_c^\top$ preserves PSD unconditionally (and strictly causal block-32 lagged two-sided `Polar-2+`, `polar2_causal_blk32`, completes all 7,200 steps [15.10B tokens] on 12.5B `15L:1A` with 0 NaNs, compressing $\kappa_2(I_C + M_c)$ from 11.9–79.3 to 1.12–1.40 across all 15 linear layers and reaching 17.684 11-set PPL, 11.49 32K Tail PPL, 49.0% `norm9` [49.1% FP8], and 68.8% 64K `HotpotQA`).

3. **Global First-Order Neumann Inversion at $C \geq 128$.** The unscaled global truncation $(I_C + M_c)^{-1} \approx I_C - M_c$ degrades local delta erasure and incurs $+0.014$ nats validation loss on Dolma 3. The block variant A_{N1} (Section 2.4.3) is exact for all lags $t - j \leq 16$; the row-scaled variant $A_{\text{N1}}^{\text{row}}$ used in `A3`/`A5` is a contractive version of this global truncation, and its cost is the $+0.22$ perplexity `A2`→`A3` step in Table 11b.