arXiv:2611.08363v1 [cs.LG] 5 Nov 2026

STRATAFLUX*: MEMORY-HIERARCHY-NATIVE LOW-BIT LLM INFERENCE

Jake Hyun [1] **Ziyu Zhao** [2] **Jae-sun Seo** [1] **Mohamed S. Abdelfattah** [1]

[1]Cornell University [2]Carnegie Mellon University

ABSTRACT

Quantization reduces the memory and bandwidth demands of large language model (LLM) inference but can degrade accuracy; offloading higher-precision data can instead stall generation. We introduce StrataFlux, a framework that challenges this tradeoff by recovering accuracy lost to low-bit weight and key–value (KV) cache quantization while retaining quantization's memory and speed advantages. StrataFlux co-designs refinement representation, selection, and delivery to make host RAM and solid-state drives useful sources of additional precision. Its *selective precision streaming* reduces transfer volume, overlaps delivery with computation, and reuses fetched precision across decoding steps so generation need not wait for every transfer. We evaluate 2–4-bit weights and 2–3-bit KV-cache on dedicated GPUs and unified-memory platforms. For 2-bit Llama-3.1-8B-Instruct weights, StrataFlux reduces WikiText-2 perplexity from 9.46 to 8.40 while maintaining 95.4% of the quantized base's throughput on RTX 5090. With 2-bit weights and KV, StrataFlux raises Qwen3-8B's AIME 2024/2025 accuracy from 11.0% to 47.1% at unchanged GPU memory use. At 128K context, Llama-3.1-8B-Instruct with 3-bit weights and 2-bit KV achieves comparable perplexity and speed to vLLM's 4-bit AWQ weights and 8-bit KV on RTX 5090, using approximately half the GPU memory. These results show that GPU-resident precision need not set the accuracy ceiling, nor offloading-tier latency dictate decoding speed.

1 INTRODUCTION

LLM inference involves three coupled objectives: accuracy, latency, and memory use. Model weights encode learned parameters, while the key–value (KV) cache retains keys and values from earlier tokens. Their representations determine both the precision available for computation and the volume of data stored and moved during generation.

Quantization and offloading address limited fast-memory capacity differently. Quantization reduces storage and traffic with fewer bits; it can accelerate bandwidth-bound inference but may degrade accuracy (Frantar et al., 2023; Liu et al., 2024). Offloading preserves stored precision by placing data in a larger, slower backing tier, but transfer latency can delay generation (Na et al., 2025). These approaches suggest an apparent choice between accuracy loss and transfer latency.

We argue that this is a false choice imposed by separating representation from placement. Considered alone, quantization discards precision that another tier could retain, while offloading brings in that tier only after choosing the representation. Can we use fewer bits in fast memory without sacrificing accuracy or slowing inference?

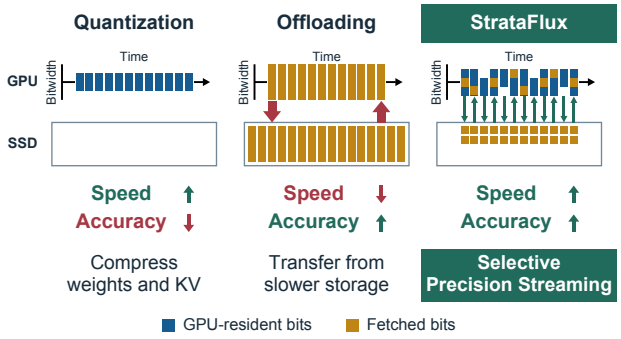

Figure 1. StrataFlux against quantization and offloading. Column height indicates bitwidth; horizontal flow indicates time. Selective precision streaming recovers accuracy while retaining low-bit inference's memory and speed advantages.

We address this question with StrataFlux, which employs *selective precision streaming* to distribute weight and KV precision across tiers (Figure 1). Compact, quantized weights and KV stay in GPU memory as the *resident base*, providing high-bandwidth access. Larger, cheaper tiers—host memory or solid-state drives (SSDs)—retain additional precision, but their lower bandwidth and higher latency call for selective fetching and flexible scheduling.

Recovering accuracy by fetching all additional precision

*StrataFlux is a provisional system name and may change in future revisions.

would erode quantization bandwidth savings. Selectivity is viable because protecting a small subset of weights can improve low-bit accuracy (Lin et al., 2024; Kim et al., 2024), while attention can concentrate on a small, query-dependent subset of tokens (Tang et al., 2024; Ribar et al., 2024). StrataFlux exploits this uneven importance to fetch additional precision selectively for the current computation.

Retaining quantization's speed advantage also requires flexible scheduling: selective transfers can still stall generation when their results are needed immediately. For weights, overlapping a fetch only with the weight-matrix operation it refines limits the time available to hide delivery. StrataFlux widens this window by selecting channels earlier or adding their refinement later, allowing transfers and refinement computation to overlap intervening operations. For both weights and KV, it also reuses fetched precision across decoding steps, avoiding repeated transfers and letting data arriving too late for one step benefit later selections.

StrataFlux makes LLM inference native to the memory hierarchy, recovering accuracy while retaining the memory and speed advantages of quantization. We contribute:

1. **Selective precision streaming.** We augment resident low-bit weights and KV with refinement bits from host RAM or SSD, concentrating transfers on channels and tokens that matter most to each computation.
2. **Asynchronous precision delivery.** By decoupling when precision is fetched from when it is used, StrataFlux hides the latency of fetches without stalling generation.
3. **Hierarchy-native inference.** By co-designing numerical representation with memory placement, StrataFlux removes the choice between quantization and offloading, enabling fast, accurate, and memory-efficient inference on both dedicated GPU and unified memory platforms.

For 2-bit Llama-3.1-8B-Instruct weights, StrataFlux reduces WikiText-2 perplexity from 9.46 to 8.40 while maintaining 95.4% of the quantized base's throughput on RTX 5090. With 2-bit weights and KV, StrataFlux raises Qwen3-8B's AIME 2024/2025 accuracy from 11.0% to 47.1% at unchanged GPU memory use. At 128K context, Llama-3.1-8B-Instruct with 3-bit weights and 2-bit KV achieves comparable perplexity and speed to vLLM's 4-bit AWQ weights and 8-bit KV on RTX 5090, using approximately half the GPU memory.

2 Background and Motivation

2.1 Post-Training Quantization

Interactive on-device LLM inference repeatedly reads model weights and the KV cache to generate new tokens. Their size constrains which models and contexts fit in fast memory, while the traffic per token can make memory bandwidth the bottleneck for generation speed (Pope et al., 2023; Liu et al., 2024). Quantization reduces these demands by representing values with low-bit codes, such as codebook indices (Han et al., 2016) or scaled and offset integers (Jacob et al., 2018). Post-training quantization (PTQ) compresses weights without retraining the model (Frantar et al., 2023), while KV-cache quantization compresses the keys and values generated during inference (Liu et al., 2024).

These savings introduce errors that affect accuracy differently across values. Large-magnitude *outliers* can be difficult to represent alongside typical values (Dettmers et al., 2022), while *saliency* describes how strongly representation errors affect computation or loss. PTQ methods address these differences: SmoothQuant rescales activation and weight channels (Xiao et al., 2023), and AWQ protects salient weight channels through scaling (Lin et al., 2024). SqueezeLLM instead combines sensitivity-weighted codebooks with higher-precision storage for selected sensitive weights and outliers (Kim et al., 2024).

2.2 Dynamic Precision Refinement

These methods account for weight sensitivity when constructing the quantized representation, but that representation remains fixed during decoding. Input activations change with each token, shifting which weight channels contribute most to the projection error. Let W be the original floating-point weights, $\widehat{W}$ their dequantized approximation, and $R = W - \widehat{W}$ the weight residual. For input x, each residual column contributes $x_i R_{:,i}$ to the projection error. Large activation magnitudes amplify these contributions, making $|x_i|$ a proxy for which channels to refine:

$$(W - \widehat{W})x = Rx = \sum_i x_i R_{:,i}. \tag{1}$$

DecDEC (Park et al., 2025) finds that calibration-selected channels often differ from those with the largest activation magnitudes during decoding. At each step, it selects the k largest $|x_i|$, fetches their 4-bit residual columns from host memory, and adds their products with signed x_i to the quantized projection's output. During single-sequence decoding, the memory-bound projection leaves GPU cores available for concurrent refinement without slowing the projection. Fetching awaits the projection's input, and refinement must finish before its output is consumed to avoid stalls, confining overlap to that projection's execution window.

We now turn to the KV cache, where the tokens receiving the most attention can likewise change between decoding steps. Attention compares a *query*, a vector derived from the current token's representation, with cached keys to weigh the cached values. Top-k attention and Quest exploit this

dependence to restrict attention to selected tokens (Gupta et al., 2021; Tang et al., 2024). SpeCache instead retains a full low-bit cache and substitutes selected 16-bit entries fetched from host memory (Jie et al., 2025). It uses a speculative token to prefetch selected entries for the next decoding step while current-step computation continues. However, fetching complete 16-bit replacements overlooks the opportunity to jointly design refinement encoding and attention computation to reduce transfer and execution costs.

2.3 Bandwidth and Latency Constraints

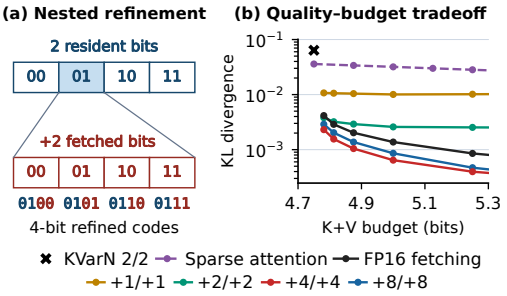

Figure 2. (a) Nested refinement subdivides GPU-resident quantization bins with fetched bits. (b) KL divergence from BF16 Llama-3.1-8B-Instruct on PG-19 at matched K+V budgets, including resident, fetched, and metadata bits per K/V pair. K/V bitwidth labels give resident bits per key/value element, respectively. All schemes use KVarN 2/2 and select tokens by attention weights from resident low-bit keys (Equation 7). Sparse attention uses selected FP16 entries alone; FP16 fetching and nested refinement retain unselected low-bit entries. FP16 fetching replaces selected entries; nested $+b/+b$ adds b bits per key/value element.

Sparse attention reduces KV traffic by omitting unselected tokens from attention. SpeCache preserves their contributions through a low-bit cache and fetches complete 16-bit floating-point (FP16) entries for selected tokens. We instead keep every key and value in low bits on the GPU and, for selected tokens, fetch only the additional bits that refine those resident values, adapting the *nested quantization* principle used for weights in Any-Precision LLM (Park et al., 2024): the fetched bits pick a finer value inside the interval that the resident low-bit value already represents. Figure 2 compares these three representation choices at matched K+V budgets. Nested refinement with 4 or 8 additional bits per selected key/value element yields lower Kullback–Leibler (KL) divergence than both sparse attention and complete FP16 replacement.

Encoding determines how much data each refinement requires; selection determines which errors it addresses. Figure 3 isolates selection by adding selected channels' residual contributions or restoring selected KV entries to FP16. We rank weight channels by input-activation magnitude and KV tokens by attention weights from resident low-bit keys.

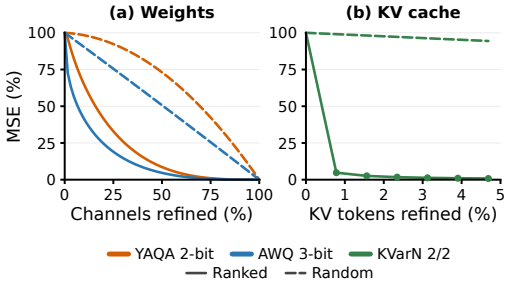

Figure 3. Ranked versus random selection in Llama-3.1-8B-Instruct. (a) With WikiText-2 inputs, weight channels are ranked by input activation magnitude and refined with corresponding columns of R. (b) For PG-19 inputs, KV tokens are ranked by attention weights from resident low-bit keys and restored to FP16. Output mean squared error (MSE) against FP16 is normalized to the error without refinement, then averaged over tokens and Transformer blocks, across projections in (a) and KV heads in (b).

Refining only a small top-ranked fraction removes a disproportionate share of output mean squared error. Activation-magnitude-based selection remains effective even on YAQA, whose rounding compensates for quantization errors across input and output channels (Tseng et al., 2026).

Compact representations and selective refinement reduce traffic, but do not eliminate access latency. DecDEC's overlap window shrinks as lower-bit quantization and higher GPU bandwidth accelerate the underlying projection (Park et al., 2025). On DGX Spark, isolated 2 KB random SSD reads have a median latency of 200 μs (709 μs at the 99th percentile)—far longer than the 21 μs average for Llama-3.1-8B-Instruct's YAQA 2-bit output projection, making long stalls unavoidable with projection-local overlap.

Two prior mechanisms suggest a solution. Each Transformer block's attention and MLP outputs are added to a running activation vector, the residual stream. Deja Vu finds that this vector's direction changes slowly across neighboring OPT blocks and uses earlier activations to predict which attention heads and MLP neurons to compute later (Liu et al., 2023). KTransformers' Expert Deferral adds the outputs of CPU-computed experts in mixture-of-experts (MoE) models to the residual stream one block later, so that their computation overlaps with the GPU's work (Chen et al., 2025).

These mechanisms extend to weight refinement (Section 2.2). Query/key/value and gate/up projections read from the residual stream, allowing earlier activations to predict channels and initiate fetching before those projections execute (*fetch lead*). Refinement is executed once each projection's input becomes available. Attention-output and down projections write to the stream, allowing refinement to be deferred to a later stream update (*fetch lag*). Together, these cases expose overlap opportunities for every atten-

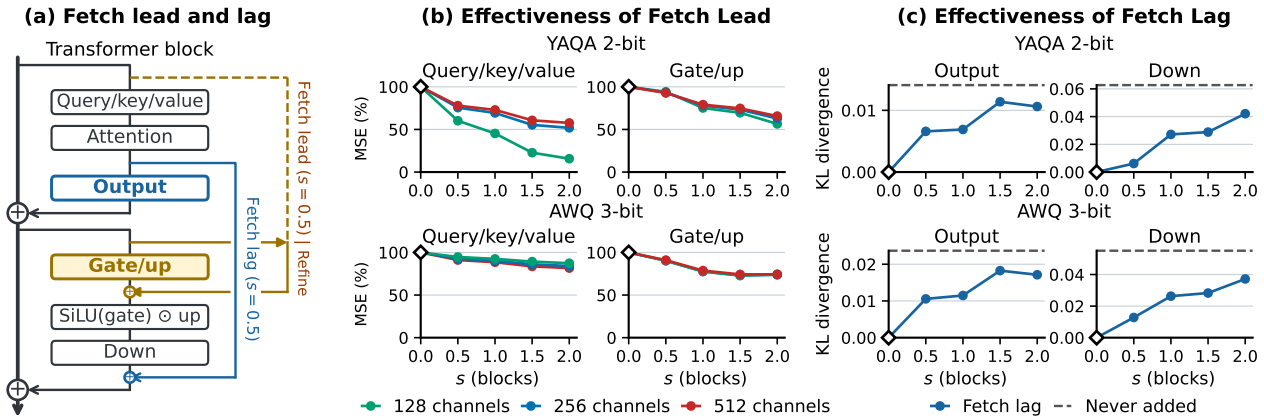

Figure 4. Fetch lead/lag in Llama-3.1-8B-Instruct on WikiText-2. Shifts s are measured in Transformer blocks; $s = 0$ means no shift. (a) Gate/up fetch lead and attention-output fetch lag at $s = 0.5$. The dashed gold path shows early selection; the solid gold input arrow feeds the later refinement step. (b) Projection-output MSE reduction retained by fetch lead, relative to selection at the projection's input (100%), using errors summed over tokens and blocks. (c) Fetch-lag refinement is applied to all input channels to isolate the effect of delayed addition. Only attention-output or down projections are quantized; KL to FP16 is averaged over predicted tokens across all sequences.

tion and MLP projection, extending refinement beyond the execution window of a single projection (Figure 4(a)).

We evaluate the shifts separately: fetch lead tests how well earlier activations predict which channels to refine, while fetch lag tests how delaying refinement addition affects model accuracy. Earlier selection can preserve most of the projection-output mean squared error (MSE) reduction at small shifts (Figure 4(b)); delayed additions continue to reduce KL divergence relative to omitting refinement (Figure 4(c)). The benefit of refinement thus persists beyond the original projection's execution window. In Section 3.3, StrataFlux combines fetch lead and lag across Transformer blocks to overlap refinement transfers and refinement computation with intervening model computation.

An independent opportunity arises from repeated selections across consecutive decoding steps. Figure 5 shows that caching only a fraction of all weight channels or KV tokens can yield very high hit rates. Caching refinement bits for recurring selections avoids repeated transfers. Data arriving too late for one step can also serve subsequent selections. We exploit these two benefits to reduce transfers and waits for weight and KV refinement (Section 3.4).

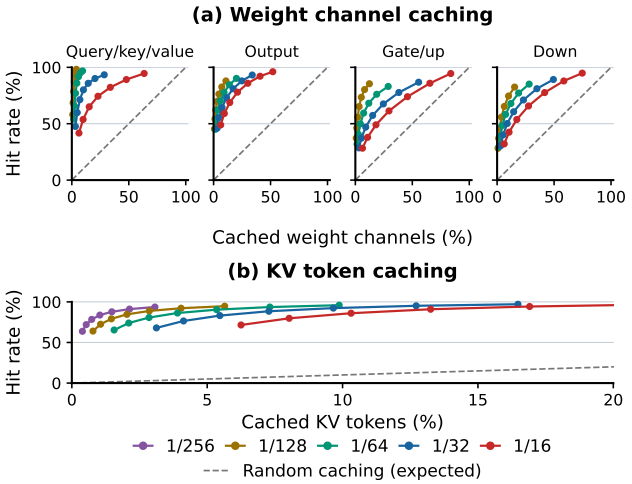

Figure 5. Reuse across decoding steps in Llama-3.1-8B-Instruct on PG-19. Selections rank (a) weight channels by input activation magnitude and (b) KV tokens by attention weight (Equation 7); legend fractions give the share selected per step. Caches contain the union of selections from the preceding 1, 2, 4, 8, 16, 32, or 64 steps. Axes show the cached share of all entries against hit rate, averaged over decoding steps and Transformer blocks. Dashed lines show expected hit rates under equally likely random caching.

3 StrataFlux: Selective Precision Streaming

StrataFlux's selective precision streaming expands the precision available to inference beyond what fits in fast memory or can be fetched within a single projection. Figure 6 outlines its weight and KV computation paths and their shared refinement-cache policy. We tailor refinement delivery to both dedicated-GPU and unified-memory platforms.

3.1 Distributing Precision Across the Memory Hierarchy

On dedicated-GPU platforms, GPU memory holds quantized weights and a compact KV cache covering all tokens, while the GPU fetches refinement bits directly from pinned host RAM. On unified-memory platforms, the CPU and GPU share RAM holding these compact representations, while SSD stores refinement bits that we load into shared RAM for GPU access (Section 3.4). These relatively inex-

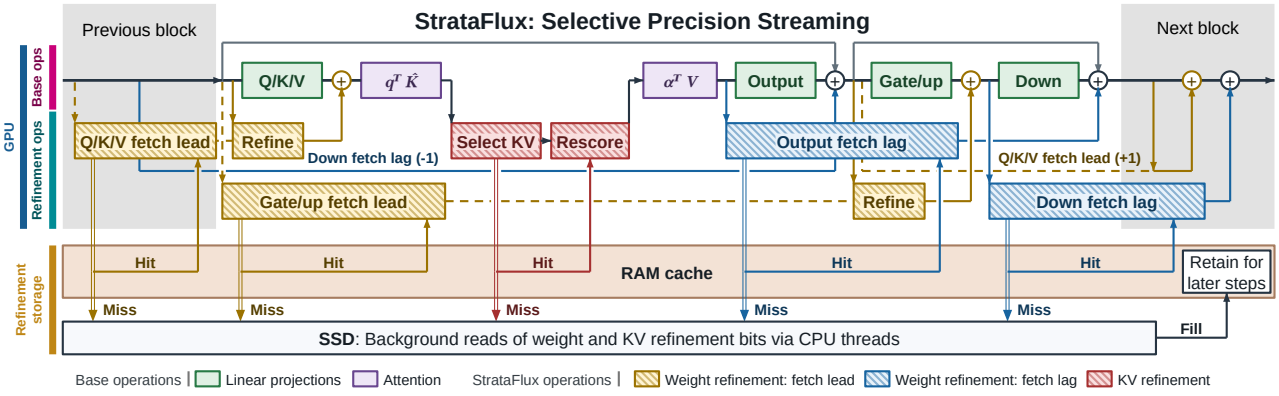

Figure 6. StrataFlux's selective precision streaming in action. Fetch lead/lag overlaps weight refinement with model computation; KV refinement runs within attention. Labeled lines replace refinement boxes for previous (-1) and next ($+1$) Transformer blocks. In this unified-memory configuration, the GPU reads the RAM cache directly; misses retain resident low-bit precision while background SSD reads supply later decoding steps. Dedicated GPUs fetch directly from host RAM, without SSD. Some Transformer operations omitted.

pensive backing tiers let StrataFlux use capacity that would otherwise remain underutilized during LLM inference.

3.2 Designing Precision for Selective Refinement

3.2.1 Weight Refinement

StrataFlux retains the resident weight quantizer's format and projection kernels. We store quantized weight-refinement columns as refinement bits; C is the corresponding matrix of dequantized values. For selected input-channel indices $S(x)$, we add the selected columns' contribution to the base quantized projection:

$$y' = \widehat{W}x + \sum_{i \in S(x)} x_i C_{:,i}. \tag{2}$$

Channel selection. Following DecDEC (Park et al., 2025), we partition each projection's input channels into chunks of 1,024 and select k channels per chunk, targeting the largest $|x_i|$. These indices form $S(x)$; each selected column's refinement bits are fetched contiguously. Separate per-chunk budgets k_{qkv}, k_{o}, k_{gu}, and k_{d} govern four projection sets: query/key/value, attention-output, gate/up, and down, respectively. We encode refinement values as integers with separate scale factors, giving the refinement kernel a common format independent of the resident weight quantizer.

Fitting refinement values. DecDEC quantizes the weight residual R (Section 2.2). StrataFlux uses *refinement fitting* to adjust refinement values so selected columns can compensate for omitted contributions. Let $z_k(x)$ zero all but the k largest-magnitude activations per input chunk. We minimize projection error while penalizing deviations from R with strength λ:

$$C_k^* = \arg\min_C \ \mathrm{E}_x[\|Rx - Cz_k(x)\|_2^2] + \lambda \|C - R\|_F^2. \tag{3}$$

The minimizer of Equation 3 solves Equation 4, where $M_k = \mathrm{E}_x[z_k(x)z_k(x)^\top]$, $G_k = \mathrm{E}_x[xz_k(x)^\top]$, and I is the identity matrix. We estimate the expectations on a small RedPajama calibration set (Weber et al., 2024) and use the mean diagonal of M_k as the default λ:

$$C_k^*(M_k + \lambda I) = RG_k + \lambda R. \tag{4}$$

We quantize the fitted matrix to 4 bits before inference, following DecDEC's 4-bit encoding of weight residuals.

Neuron-wise alternative. For input x, the quantized MLP combines gate and up outputs into neuron activations h, then computes its unrefined output y_{q}:

$$\begin{aligned} g &= \widehat{W}_{\text{gate}}x, & u &= \widehat{W}_{\text{up}}x, \\ h &= \phi(g) \odot u, & y_{\text{q}} &= \widehat{W}_{\text{down}}h. \end{aligned} \tag{5}$$

Here ϕ is the gate's activation function, $\odot$ is elementwise multiplication, and h_j is neuron j's activation after gating. Channel-wise refinement targets large input activations separately: x for gate/up and h for down. Neuron-wise refinement instead selects the k_{n} neurons with largest $|h_j|$ in each chunk of 1,024 and refines their gate/up rows and down columns together. Each selected neuron's *MLP neuron bundle* encodes gate/up refinement rows $c_j^{\text{g}}, c_j^{\text{u}}$, a down refinement column c_j^{d}, and a 4-bit copy d_j of the BF16 down column. For selected neuron indices $N(h)$, we recompute h_j from Equation 5 as h'_j. As in Equation 2, the resulting output refinement is added to y_{q}:

$$\begin{aligned} h'_j &= \phi(g_j + c_j^{\text{g}}x)(u_j + c_j^{\text{u}}x), \\ y' &= y_{\text{q}} + \sum_{j \in N(h)} \left[d_j(h'_j - h_j) + c_j^{\text{d}}h_j \right]. \end{aligned} \tag{6}$$

Neither channel-wise nor neuron-wise MLP refinement consistently offers the better speed–accuracy tradeoff across models, quantizers, platforms, and speed budgets. We therefore retain both options for StrataFlux.

3.2.2 *KV Refinement*

For KV refinement, we tailor nested reconstruction to the resident cache format. Our main KV implementation, StrataFlux-KV, adopts KVarN's (Muller et al., 2026) quantization format for its balance of accuracy and memory footprint. We co-design its low-bit attention kernels and selective refinement. For adaptations of StrataFlux to other KV-cache formats, see Appendix E.

Nested encoding. As illustrated in Figure 2(a), we subdivide each KVarN code's quantization interval into equal sub-intervals and store their indices as refinement bits in the backing tier. The resident code and fetched index identify a finer reconstruction, independently for keys and values.

Token selection. We compute attention weights by applying softmax to scores between the current query and resident low-bit keys. A KV head stores keys and values shared by one or more query heads. Let $\alpha_{h,i}$ denote the resulting attention weight assigned to token i by query head h, and let $H(j)$ be the query heads sharing KV head j. For token i and KV head j, we define the selection score

$$s_j(i) = \max_{h \in H(j)} \alpha_{h,i}. \tag{7}$$

This score captures the strongest attention to the token among query heads sharing its keys and values. For each KV head, we select a configurable number of the highest-scoring tokens for refinement.

Attention refinement. Combining fetched bits with resident codes reconstructs selected tokens' higher-precision keys and values (Figure 6). For one query head, resident scores ℓ_i and values v_i give output o and normalizer $Z = \sum_i e^{\ell_i}$. Refining tokens S gives scores ℓ'_i, values v'_i, and output

$$o' = \frac{Zo + \sum_{i \in S} \left(e^{\ell'_i} v'_i - e^{\ell_i} v_i\right)}{Z + \sum_{i \in S} \left(e^{\ell'_i} - e^{\ell_i}\right)}. \tag{8}$$

In practice, we group consecutive tokens within each KV head and pack each group's nested refinement bits contiguously for larger reads. We score token groups using Equation 7, with each query head's attention weights summed over the group. Because token selection depends on the initial attention scores and the weighted sum depends on refined scores, KV refinement adds a region that cannot overlap the surrounding attention computation. We limit this overhead through contiguous fetches and fusion of value refinement into the weighted sum.

Producing KV refinement bits. During prefill, we compute attention on uncompressed keys and values, then encode resident low-bit codes and offload refinement bits. During decoding, we encode completed blocks of new KV entries in the same format.

3.3 Extending the Weight Delivery Interval

The runtime uses fetch lead/lag to decouple refinement fetching from its application, hiding latency through overlap with base operations (Section 2.3). Appendix G explains the CUDA stream and event scheduling.

Fetch lead. For each query/key/value or gate/up projection, we apply its normalization to a residual-stream activation s Transformer blocks earlier, then select channels and initiate fetching. Once the projection's input is available, the refinement kernel multiplies fetched columns by corresponding activations and adds the summed products to the quantized projection's output (Figure 6).

Fetch lag. We add attention-output and down refinements to the residual stream s Transformer blocks after their original addition point, overlapping their fetching and computation with intervening projections (Figure 6). Neuron-wise MLP refinement likewise supports fetch lag.

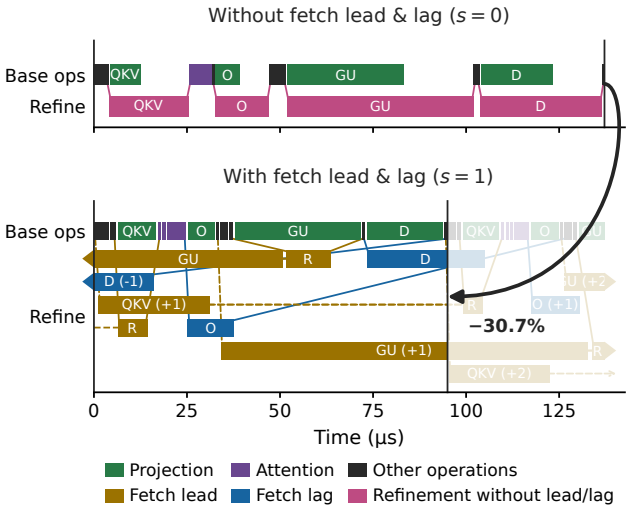

Figure 7. Nsight Systems kernel traces comparing refinement without fetch lead/lag ($s = 0$) and with fetch lead/lag ($s = 1$) on Llama-3.1-8B-Instruct (AWQ 3-bit, RTX 5090, $k = 32$). QKV/GU/O/D denote query/key/value, gate/up, output, and down projections; R denotes refinement after fetch lead. Refinement for the previous and next blocks are marked with an offset (-1 or +1). Dashed lines show fetch-lead's early selection; solid lines show activation reads and writes for refinement.

The Nsight Systems kernel traces of our runtime system in Figure 7 confirm the cross-block overlap illustrated in Figure 6. With the projection-local schedule used by DecDEC ($s = 0$), the main stream waits for refinement after each projection, delaying subsequent operations in the Transformer block's computation graph. Fetch lead/lag at $s = 1$ successfully overlaps refinement with computation across neighboring Transformer blocks, reducing these waits and shortening block execution at matched channel count k.

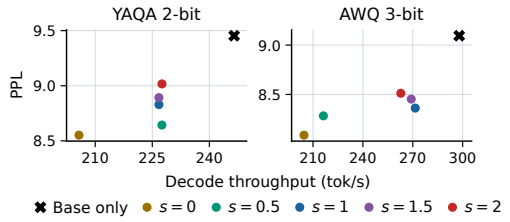

Figure 8. Perplexity–throughput tradeoffs on Llama-3.1-8B-Instruct with $k = 32$ and a common fetch lead/lag distance s across all four projection sets. PPL uses WikiText-2; throughput uses PG-19 on RTX 5090 with refinement bits in host RAM.

However, the shift needed to hide refinement depends on the quantization method. At $k = 32$, Figure 8 shows that slower YAQA 2-bit obtains most of its throughput gain at $s = 0.5$, whereas faster AWQ 3-bit needs $s = 1$ because its shorter projections leave less time to hide refinement. Furthermore, these sweeps apply the same s to all four projection sets within each method, leaving their individual speed–accuracy tradeoffs unexplored. In Section 3.5, we introduce a tuner that directly explores this search space to select s and k per projection set under a configurable speed budget.

3.4 Reusing Precision Across Decoding Steps

For dedicated-GPU platforms, host RAM's lower access latency than SSD allows refinement bits to be fetched directly without an intermediate cache. For unified-memory systems, however, SSD's higher access latency motivates caching fetched refinement bits in GPU-accessible RAM.

Caching fetched precision. Weight and KV selections recur across consecutive decoding steps (Figure 5). To avoid repeated SSD reads and their latency, we retain fetched refinement bits in a configurably sized RAM cache that the GPU reads directly (Figure 6).

Miss policy. Selections can still miss the cache. For weights and KV, the current step can decide to skip missing refinement and use the resident low-bit representation. A CPU loader fetches the bits into the cache for future selections, keeping the transfer useful without making the requesting step wait. Alternatively, the runtime can wait and apply refinement in the current step, trading speed for accuracy.

Cache organization. Weight and KV refinement have separate cache budgets. The weight cache stores refinement bits for all four projection sets and uses least-recently-used replacement. KV token groups share capacity across all KV heads and Transformer blocks.

Populating the KV refinement cache. During prefill, we write KV refinement bits to SSD and cache copies for the most recently encoded token groups in RAM. During decoding, newly produced refinement bits enter the cache directly, serving subsequent selections without SSD readback.

3.5 Tuning Weight Refinement Within a Speed Budget

Search space. The weight tuner selects a configuration subject to a user-specified slowdown limit relative to the quantized base. It chooses channel budgets k, fetch lead/lag distances s for each projection set, and whether the MLP uses channel-wise refinement or neuron-wise refinement with budget k_{n}. Following DecDEC (Park et al., 2025), it also tunes the number of GPU thread blocks (groups of cooperating threads) executing refinement. More thread blocks can accelerate refinement but also slow the concurrent quantized projections by competing for GPU resources.

Candidate selection. For each combination of shift distances and channel-wise or neuron-wise MLP refinement, the tuner searches for the largest common channel budget k that meets the slowdown limit in end-to-end decoding. It selects the retained candidate with the lowest KL divergence from the BF16 reference on held-out text, then increases individual channel budgets or the neuron budget while the measured slowdown remains within the limit. We tune each model, weight quantizer, and platform separately. We then fit refinements at the chosen selection budgets, as detailed in Section 3.2.1.

4 Evaluation

We implement StrataFlux as an end-to-end system in llama.cpp (llama.cpp contributors, n.d.), with CUDA and Metal backends. We evaluate StrataFlux on Llama-3.1-8B-Instruct (Meta, 2024a), Llama-3.3-70B-Instruct (Meta, 2024b), Qwen3-8B, Qwen3-32B, Qwen3-30B-A3B (Qwen Team, 2025), Qwen3.5-9B (Qwen Team, 2026), and DeepSeek-V2-Lite (DeepSeek-AI, 2024). Platforms span dedicated GPUs (RTX 5090, RTX 5070 Ti, and RTX PRO 6000) and unified-memory platforms (DGX Spark and MacBook Air M3). Appendix F reports the Mac evaluation. We measure perplexity and KL divergence from BF16 on 2K-token WikiText-2 windows (Merity et al., 2017) and choose PG-19 (Rae et al., 2020) for long-context perplexity. We measure accuracy on RULER (Hsieh et al., 2024), MATH500 (Lightman et al., 2024), and AIME 2024/2025 (Mathematical Association of America, n.d.); AIME averages eight samples per problem. Timings use llama.cpp except for methods with their own runtimes. Appendix A details protocols and implementations.

In tables and figures, StrataFlux-W and StrataFlux-KV isolate weight and KV refinement, respectively; StrataFlux without a suffix denotes the full system with both enabled.

Unless ablated, DGX Spark uses separate 128 MiB RAM caches for weight and KV refinement. At the small slowdown budgets used in our main results, we find that skipping cache misses gives a better latency–accuracy tradeoff on unified-memory platforms. There, we skip misses and

set $s = 0$: no SSD wait remains for fetch lead/lag to hide. Appendix B evaluates fetch lead/lag with waiting on cache misses at larger budgets.

4.1 Weight Refinement

We compare StrataFlux-W with DecDEC (Park et al., 2025) on YAQA (Tseng et al., 2026) with QTIP's format and kernels (Tseng et al., 2024), LNQ with GuidedQuant (GQ) (Kim et al., 2025), AWQ (Lin et al., 2024), and SqueezeLLM (SqLLM) (Kim et al., 2024) bases.

Table 1. Weight-only refinement: WikiText-2 perplexity and KL divergence with 2–3-bit weights. Platform blocks report StrataFlux-W, alongside DecDEC on RTX 5090, tuned to a 5% slowdown budget. DecDEC does not support SSD fetching on DGX Spark. Bold marks the better score within each RTX 5090 pair.

		Refinement	L3.1-8B-I		Qwen3-8B	
	Weights	method	PPL	KL	PPL	KL
	BF16	–	7.22	0.000	9.73	0.000
Base	YAQA 2-bit	–	9.46	0.257	10.95	0.214
	YAQA 3-bit		7.61	0.047	9.98	0.052
	LNQ+GQ 3-bit		8.52	0.157	10.92	0.184
	AWQ 3-bit		9.10	0.229	11.88	0.284
	SqLLM 3-bit		9.36	0.263	11.71	0.255
DGX Spark	YAQA 2-bit	StrataFlux-W	8.94	0.200	10.77	0.175
	YAQA 3-bit	StrataFlux-W	7.55	0.041	9.89	0.043
	LNQ+GQ 3-bit	StrataFlux-W	8.29	0.135	10.70	0.159
	AWQ 3-bit	StrataFlux-W	8.64	0.179	11.30	0.222
	SqLLM 3-bit	StrataFlux-W	8.56	0.174	11.03	0.201
RTX 5090	YAQA 2-bit	DecDEC	9.03	0.211	10.69	0.184
		StrataFlux-W	**8.40**	**0.139**	**10.51**	**0.132**
	YAQA 3-bit	DecDEC	7.55	0.040	**9.89**	0.044
		StrataFlux-W	**7.49**	**0.033**	**9.89**	**0.033**
	LNQ+GQ 3-bit	DecDEC	8.37	0.142	10.66	0.170
		StrataFlux-W	**8.18**	**0.121**	**10.47**	**0.139**
	AWQ 3-bit	DecDEC	8.80	0.195	11.38	0.231
		StrataFlux-W	**8.38**	**0.145**	**10.71**	**0.167**
	SqLLM 3-bit	DecDEC	9.02	0.225	11.36	0.215
		StrataFlux-W	**8.42**	**0.156**	**10.81**	**0.169**

We use a 5% slowdown budget for main weight comparisons because it is sufficient for large accuracy gains (Figure 9(a)). At matched budgets, StrataFlux-W outperforms DecDEC across quantizers (Table 1). These gains extend to reasoning tasks (Table 2) while retaining quantized speed and low GPU memory use (Figure 9). For Qwen3-8B with YAQA 2-bit weights, AIME accuracy rises from 19.2% to 46.7%. Appendix B gives full 2–4-bit, larger-model, and downstream results; Appendix B.7 lists tuned configurations.

4.2 KV Refinement

KV baselines include KVarN (Muller et al., 2026), KIVI (Liu et al., 2024), KVQuant (Hooper et al., 2024), TurboQuant (Zandieh et al., 2026), and AQUA-KV (Shutova et al., 2025); offloading comparisons include ShadowKV (Sun et al., 2025), ParisKV (Qi et al., 2026), TailorKV (Yao et al., 2025), FreeKV (Liu et al., 2026),

Table 2. Weight-only refinement: Qwen3-8B MATH500/AIME accuracy. Platform columns use StrataFlux-W's respective 5% slowdown configurations. 4-bit results appear in Appendix B.

	Base		StrataFlux-W DGX Spark		StrataFlux-W RTX 5090	
Weights	M500	AIME	M500	AIME	M500	AIME
BF16	97.2	71.0	–	–	–	–
YAQA 2-bit	77.2	19.2	89.4	40.0	92.4	46.7
YAQA 3-bit	95.0	66.7	95.0	66.0	95.2	70.0
LNQ+GQ 3-bit	91.8	39.4	93.0	47.3	93.2	49.2
AWQ 3-bit	91.8	46.5	92.6	51.9	93.8	51.0
SqLLM 3-bit	93.0	47.3	94.2	48.3	93.0	50.8

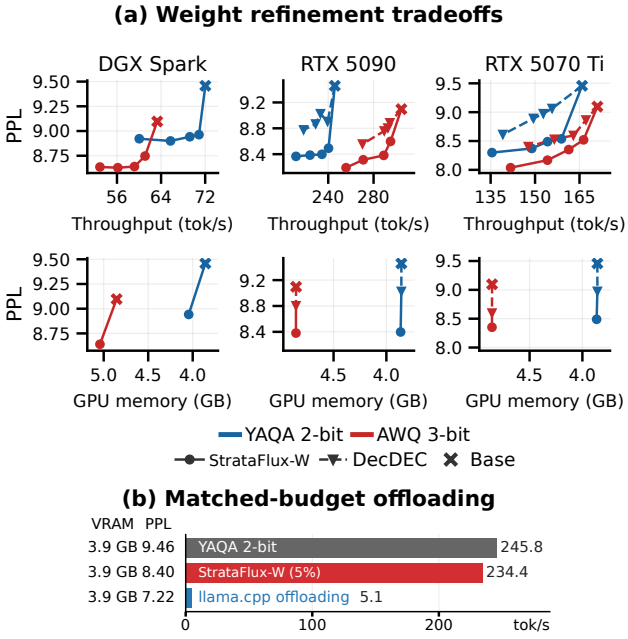

Figure 9. Weight-refinement tradeoffs on Llama-3.1-8B-Instruct on WikiText-2. (a) Throughput plots connect configurations tuned for slowdown budgets of 2.5% / 5% / 10% / 20%; memory plots show the configurations tuned for the 5% budget. DGX Spark memory includes the RAM cache. DecDEC does not support SSD fetching on Spark. (b) StrataFlux-W tuned for 5% slowdown on YAQA 2-bit base, and llama.cpp BF16 partial offloading at a matched 3.9 GB GPU memory budget, RTX 5090.

SpeCache (Jie et al., 2025), and KVSwap (Zhang et al., 2026). StrataFlux is configured to select $1/128$ (0.78%) of token groups for refinement, adding 4 bits per selected key and value element (+4/+4). Groups contain two consecutive tokens on dedicated GPU platforms and four on unified-memory platforms. For other refinement settings, see Appendix C.6; for resident 2/1 KV, see Appendix C.5.

StrataFlux-KV pushes beyond the memory–accuracy and throughput–accuracy frontiers of existing KV quantization and offloading methods, bringing near-BF16 accuracy to 2-bit keys and values without giving up low-bit memory savings or fast decoding (Table 3; Figure 10(a)). Accuracy recovery also extends to alternative attention architectures

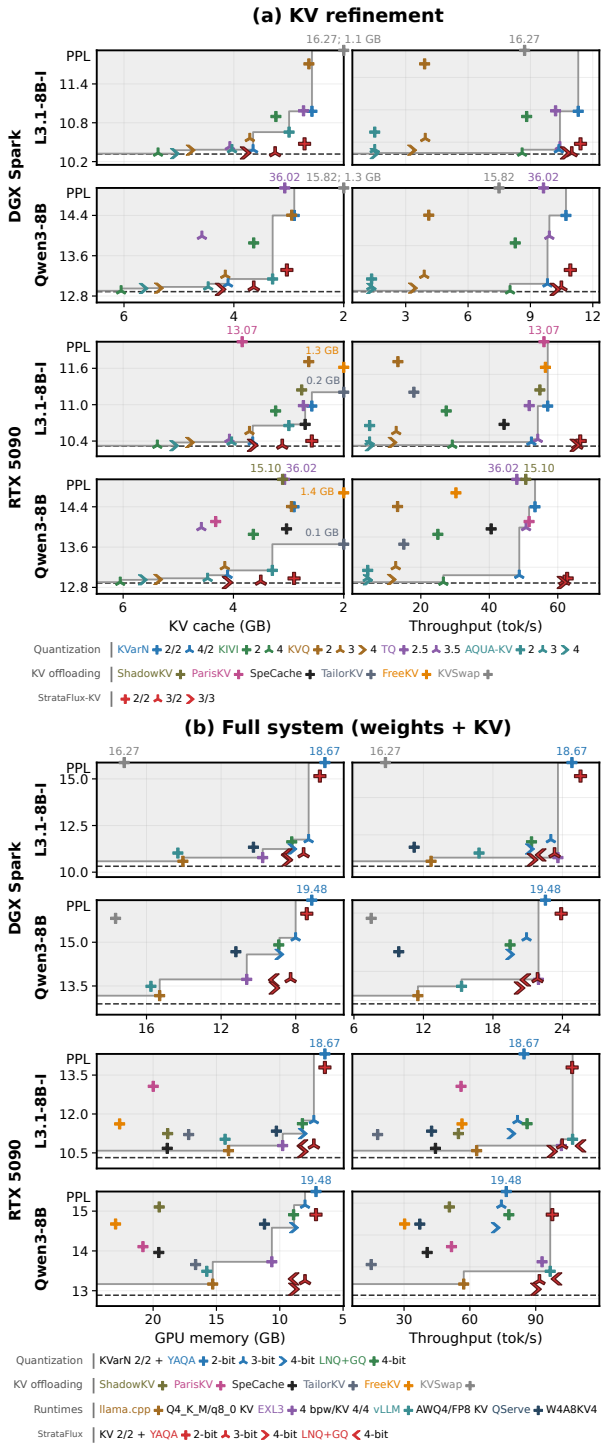

Figure 10. PG-19 perplexity at 128K: (a) KV refinement with BF16 weights; (b) Weight and KV joint refinement. Memory counts KV in (a), weights+KV in (b), including DGX Spark's RAM cache. Out-of-range points are shown beside the corresponding borders with their values. Of the evaluated KV offloading systems, only KVSwap supports DGX Spark. The gray boundary marks the baseline frontier. Dashed lines mark BF16 PPL. W/A/KV: weight/activation/cache bits; bpw: bits/weight; KVQ: KVQuant; TQ: TurboQuant.

Table 3. KV-only refinement: PG-19 perplexity/KL and RULER accuracy at 128K, plus MATH500/AIME accuracy, with BF16 weights. Platform blocks report StrataFlux-KV.

		L3.1-8B-I			Qwen3-8B				
	KV	PPL	KL	R128	PPL	KL	R128	M500	AIME
	BF16	10.32	0.000	77.9	12.89	0.000	74.1	97.2	71.0
Base	2/2	10.98	0.067	68.6	14.40	0.130	66.0	96.2	66.0
	3/2	10.45	0.014	76.2	13.17	0.026	72.3	96.6	69.2
	3/3	10.39	0.008	77.2	13.02	0.016	72.9	95.8	70.6
StrataFlux-KV DGX Spark	2/2	10.47	0.011	75.2	13.32	0.030	71.4	96.0	68.3
	3/2	10.34	0.002	77.6	12.98	0.005	74.0	95.0	70.2
	3/3	10.33	0.001	77.5	12.93	0.004	74.2	96.2	69.4
StrataFlux-KV RTX 5090	2/2	10.40	0.003	78.1	12.98	0.007	74.2	95.6	69.2
	3/2	10.33	0.001	78.3	12.92	0.002	74.7	96.2	70.0
	3/3	10.33	0.001	78.2	12.90	0.001	74.5	97.0	71.5

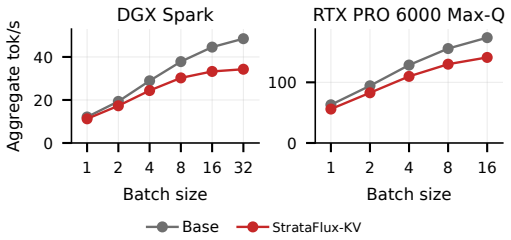

Figure 11. StrataFlux's KV refinement overhead versus batch size on Llama-3.1-8B-Instruct at 128K, using 2/2 KV.

and MoE models (Appendix C, Table A20). On RTX 5090, Llama-3.1-8B-Instruct with 2/2 KV improves from 68.6% to 78.1% RULER accuracy at 128K. For full task and context comparisons, see Appendix C.

At batch sizes 1 and 16, refinement reduces throughput by 11.2% and 18.8% on RTX PRO 6000 Max-Q, versus 7.0% and 25.4% on DGX Spark (Figure 11).

4.3 Full-System Results

We enable both refinement techniques and compare the full StrataFlux system with static weight/KV pairs, vLLM (Kwon et al., 2023), EXL3 (ExLlamaV3 contributors, n.d.), QServe (Lin et al., 2025), llama.cpp, and KV offloading systems. Together, weight and KV refinement let StrataFlux recover long-context and reasoning accuracy with aggressively compressed weights and KV (Table 4), pushing beyond the memory–accuracy and throughput–accuracy frontiers of static quantization pairs, KV offloading systems, and inference runtimes (Figure 10(b)). On RTX 5090, refinement raises Qwen3-8B AIME accuracy from 11.0% to 47.1% with YAQA 2-bit weights and 2/2 KV. Appendix D gives the full configuration comparisons and task results.

4.4 Ablations

Fetch lead/lag and fitting improve accuracy, while the better MLP scheme, channel-wise refinement or neuron bundles,

Table 4. Joint refinement: PG-19 perplexity/KL and RULER accuracy at 128K, plus MATH500/AIME accuracy. Platform blocks report StrataFlux.

			L3.1-8B-I			Qwen3-8B				
	Weights	KV	PPL	KL	R128	PPL	KL	R128	M500	AIME
	BF16		10.32	0.000	77.9	12.89	0.000	74.1	97.2	71.0
Base	YAQA 2-bit	2/2	18.67	0.656	50.4	19.48	0.434	58.8	68.0	11.0
	YAQA 3-bit	2/2	11.75	0.146	67.2	15.15	0.185	63.8	94.6	59.6
	YAQA 4-bit	2/2	11.24	0.090	67.5	14.58	0.143	65.4	95.0	62.5
	LNQ+GQ 4-bit	2/2	11.62	0.129	64.3	14.91	0.167	64.0	95.2	58.3
StrataFlux DGX Spark	YAQA 2-bit	2/2	15.15	0.422	60.5	15.98	0.228	65.2	89.6	40.6
	YAQA 3-bit	2/2	10.99	0.064	73.2	13.75	0.068	68.4	95.2	68.8
	YAQA 4-bit	2/2	10.66	0.026	73.5	13.43	0.040	70.3	96.2	69.0
	LNQ+GQ 4-bit	2/2	10.92	0.055	71.1	13.71	0.062	70.6	96.0	65.6
StrataFlux RTX 5090	YAQA 2-bit	2/2	13.79	0.312	63.8	14.91	0.165	68.3	91.2	47.1
	YAQA 3-bit	2/2	10.80	0.045	77.1	13.25	0.038	71.2	95.4	65.6
	YAQA 4-bit	2/2	10.55	0.015	77.2	13.04	0.015	73.5	96.4	69.2
	LNQ+GQ 4-bit	2/2	10.77	0.038	75.2	13.30	0.035	73.8	94.4	61.5

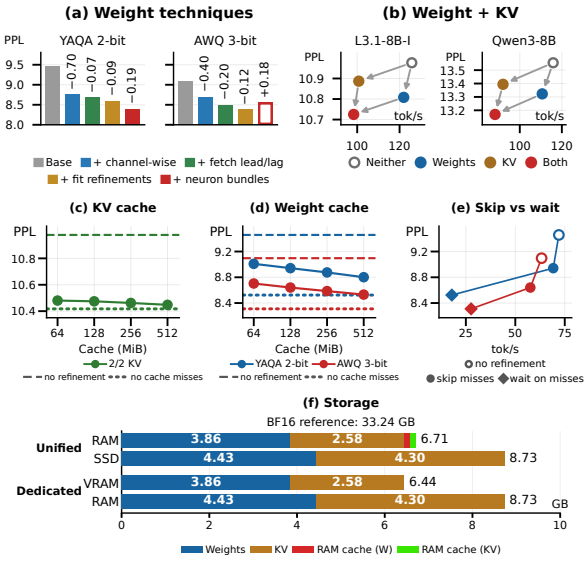

Figure 12. Refinement ablations on Llama-3.1-8B-Instruct, using RTX 5090 in (a,b) and DGX Spark in (c–e). (a) Adding weight-refinement techniques, re-tuned to a 5% slowdown budget. Labels show WikiText-2 PPL changes. The outlined AWQ bar forces neuron bundles. (b) PPL versus throughput for separate and combined weight/KV refinement (also on Qwen3-8B), using PG-19 at 128K, YAQA 3-bit weights, 3/3 KV, and the 5% weight configuration. (c,d) PPL versus RAM-cache capacity with skipped misses: (c) 2/2 KV on PG-19; (d) weight refinement on WikiText-2 at 5% configurations. (e) WikiText-2 PPL versus throughput for skipping and waiting on weight-cache misses with a 128 MiB cache. (f) Storage (GB) by tier: YAQA 2-bit weights and 2/2 KV at 128K.

depends on the quantizer (Figure 12(a)). Weight and KV refinement yield approximately additive accuracy gains and decoding-time overheads, allowing both to run together without negative interference (Figure 12(b)). RAM caches make SSD-backed refinement effective despite skipped misses. Larger caches improve accuracy, while skipping avoids the throughput penalty of waiting (Figure 12(c–e)). Appendix C extends the fetch-budget and cache-capacity sweeps. With 4-bit weight and KV refinement values, backing-tier storage scales roughly with 4-bit weights and KV, plus metadata and values duplicated in neuron bundles (Figure 12(f)). For other KV bases, see Appendix E.

5 Related Work

Heterogeneous-memory inference systems accommodate weights and KV caches beyond GPU memory. FlexGen (Sheng et al., 2023) schedules compressed tensors across GPU, CPU, and disk; InfiniGen (Lee et al., 2024) and FreeKV (Liu et al., 2026) prefetch selected KV entries from host memory. Quest (Tang et al., 2024) and SparQ (Ribar et al., 2024) reduce KV reads through query-dependent selection. SpeCache (Jie et al., 2025) supplements a full low-bit cache with prefetched 16-bit entries. StrataFlux instead co-designs low-bit attention and selectively fetched refinement bits, using caching and skipped misses to recover accuracy without stalling decoding on backing-tier reads.

For weights, DecDEC (Park et al., 2025) selects refinement channels from current activations and overlaps host-memory fetching and refinement with the corresponding quantized projection. Other inference systems overlap work across Transformer blocks: Deja Vu (Liu et al., 2023) uses earlier activations to predict which attention heads and MLP neurons to compute, while KTransformers (Chen et al., 2025) defers expert outputs to overlap CPU and GPU computation. StrataFlux extends weight-refinement overlap beyond DecDEC's projection window through early channel selection and deferred addition, and fits refinement values to compensate for unselected channels.

6 Conclusion

StrataFlux recovers accuracy from low-bit weights and KV caches while retaining quantization's memory and speed advantages. Its selective precision streaming makes host RAM and SSD useful sources of additional precision by coordinating what is fetched, when it arrives, and how it refines computation. Across dedicated GPUs and unified-memory platforms, StrataFlux demonstrates that precision distributed across the memory hierarchy can improve accuracy without letting backing-tier transfers dictate decoding speed.

References

Bai, Y., Lv, X., Zhang, J., Lyu, H., Tang, J., Huang, Z., Du, Z., Liu, X., Zeng, A., Hou, L., Dong, Y., Tang, J., and Li, J. LongBench: A bilingual, multitask benchmark for long context understanding. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 3119–3137. Association

for Computational Linguistics, 2024. doi: 10.18653/v1/2024.acl-long.172. URL https://aclanthology.org/2024.acl-long.172/.

Chen, H., Xie, W., Zhang, B., Tang, J., Wang, J., Dong, J., Chen, S., Yuan, Z., Lin, C., Qiu, C., Zhu, Y., Ou, Q., Liao, J., Chen, X., Ai, Z., Wu, Y., and Zhang, M. KTransformers: Unleashing the full potential of CPU/GPU hybrid inference for MoE models. In *Proceedings of the ACM SIGOPS 31st Symposium on Operating Systems Principles*, SOSP '25. Association for Computing Machinery, 2025. doi: 10.1145/3731569.3764843. URL https://doi.org/10.1145/3731569.3764843.

DeepSeek-AI. DeepSeek-V2-Lite. Model card, 2024. URL https://huggingface.co/deepseek-ai/DeepSeek-V2-Lite. Accessed October 30, 2026.

Dettmers, T., Lewis, M., Belkada, Y., and Zettlemoyer, L. LLM.int8(): 8-bit matrix multiplication for transformers at scale. In *Advances in Neural Information Processing Systems*, volume 35, pp. 30318–30332. Curran Associates, Inc., 2022. doi: 10.52202/068431-2198.

ExLlamaV3 contributors. ExLlamaV3: Quantization and inference library. Software repository, n.d. URL https://github.com/turboderp-org/exllamav3. Accessed October 30, 2026.

Frantar, E., Ashkboos, S., Hoefler, T., and Alistarh, D. OPTQ: Accurate post-training quantization for generative pre-trained transformers. In *The Eleventh International Conference on Learning Representations*, 2023. URL https://openreview.net/forum?id=tcbBPnfwxS.

Gao, L., Tow, J., Abbasi, B., Biderman, S., Black, S., DiPofi, A., Foster, C., Golding, L., Hsu, J., Le Noac'h, A., Li, H., McDonell, K., Muennighoff, N., Ociepa, C., Phang, J., Reynolds, L., Schoelkopf, H., Skowron, A., Sutawika, L., Tang, E., Thite, A., Wang, B., Wang, K., and Zou, A. A framework for few-shot language model evaluation, 12 2023. URL https://zenodo.org/records/10256836.

Gupta, A., Dar, G., Goodman, S., Ciprut, D., and Berant, J. Memory-efficient transformers via top-k attention. In *Proceedings of the Second Workshop on Simple and Efficient Natural Language Processing*, pp. 39–52. Association for Computational Linguistics, 2021. doi: 10.18653/v1/2021.sustainlp-1.5. URL https://aclanthology.org/2021.sustainlp-1.5/.

Han, S., Mao, H., and Dally, W. J. Deep compression: Compressing deep neural networks with pruning, trained quantization and huffman coding. In *The Fourth International Conference on Learning Representations*, 2016. URL https://arxiv.org/abs/1510.00149.

Hooper, C., Kim, S., Mohammadzadeh, H., Mahoney, M. W., Shao, Y. S., Keutzer, K., and Gholami, A. KVQuant: Towards 10 million context length LLM inference with KV cache quantization. In *Advances in Neural Information Processing Systems*, volume 37, pp. 1270–1303. Curran Associates, Inc., 2024. doi: 10.52202/079017-0040. URL https://doi.org/10.52202/079017-0040.

Hsieh, C.-P., Sun, S., Kriman, S., Acharya, S., Rekesh, D., Jia, F., Zhang, Y., and Ginsburg, B. RULER: What's the real context size of your long-context language models? In *Conference on Language Modeling*, 2024. URL https://openreview.net/forum?id=kIoBbc76Sy.

Jacob, B., Kligys, S., Chen, B., Zhu, M., Tang, M., Howard, A., Adam, H., and Kalenichenko, D. Quantization and training of neural networks for efficient integer-arithmetic-only inference. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 2704–2713, 2018. URL https://openaccess.thecvf.com/content_cvpr_2018/html/Jacob_Quantization_and_Training_CVPR_2018_paper.html.

Jie, S., Tang, Y., Han, K., Deng, Z.-H., and Han, J. SpeCache: Speculative key-value caching for efficient generation of LLMs. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 27917–27928. PMLR, 2025. URL https://proceedings.mlr.press/v267/jie25a.html.

Kim, J., El Halabi, M., Park, W., Schaefer, C. J., Lee, D., Park, Y., Lee, J. W., and Song, H. O. GuidedQuant: Large language model quantization via exploiting end loss guidance. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 30011–30037. PMLR, 2025. URL https://proceedings.mlr.press/v267/kim25d.html.

Kim, S., Hooper, C. R. C., Gholami, A., Dong, Z., Li, X., Shen, S., Mahoney, M. W., and Keutzer, K. SqueezeLLM: Dense-and-sparse quantization. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 23901–23923. PMLR, 2024. URL https://proceedings.mlr.press/v235/kim24f.html.

Kwon, W., Li, Z., Zhuang, S., Sheng, Y., Zheng, L., Yu, C. H., Gonzalez, J. E., Zhang, H., and Stoica, I. Efficient memory management for large language model serving with PagedAttention. In *ACM Symposium on Operating Systems Principles*, pp. 611–626,

2023. doi: 10.1145/3600006.3613165. URL https://doi.org/10.1145/3600006.3613165.

Lee, W., Lee, J., Seo, J., and Sim, J. InfiniGen: Efficient generative inference of large language models with dynamic KV cache management. In *18th USENIX Symposium on Operating Systems Design and Implementation*, pp. 155–172. USENIX Association, 2024. URL https://www.usenix.org/conference/osdi24/presentation/lee.

Lightman, H., Kosaraju, V., Burda, Y., Edwards, H., Baker, B., Lee, T., Leike, J., Schulman, J., Sutskever, I., and Cobbe, K. Let's verify step by step. In *International Conference on Learning Representations*, 2024.

Lin, J., Tang, J., Tang, H., Yang, S., Chen, W.-M., Wang, W.-C., Xiao, G., Dang, X., Gan, C., and Han, S. AWQ: Activation-aware weight quantization for on-device LLM compression and acceleration. In *Proceedings of Machine Learning and Systems*, volume 6, pp. 87–100, 2024.

Lin, Y., Tang, H., Yang, S., Zhang, Z., Xiao, G., Gan, C., and Han, S. QServe: W4A8KV4 quantization and system co-design for efficient LLM serving. In *Proceedings of Machine Learning and Systems*, volume 7, 2025. URL https://hanlab.mit.edu/projects/qserve.

Liu, G., Li, C., Ning, Z., Lin, J., Yao, Y., Ke, D., Guo, M., and Zhao, J. FreeKV: Boosting KV cache retrieval for efficient LLM inference. In *International Conference on Learning Representations*, 2026. URL https://openreview.net/pdf?id=wXAn7orB1H.

Liu, Z., Wang, J., Dao, T., Zhou, T., Yuan, B., Song, Z., Shrivastava, A., Zhang, C., Tian, Y., Ré, C., and Chen, B. Deja Vu: Contextual sparsity for efficient LLMs at inference time. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pp. 22137–22176. PMLR, 2023. URL https://proceedings.mlr.press/v202/liu23am.html.

Liu, Z., Yuan, J., Jin, H., Zhong, S., Xu, Z., Braverman, V., Chen, B., and Hu, X. KIVI: A tuning-free asymmetric 2bit quantization for KV cache. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 32332–32344. PMLR, 2024. URL https://proceedings.mlr.press/v235/liu24bz.html.

llama.cpp contributors. llama.cpp: LLM inference in C/C++. Software repository, n.d. URL https://github.com/ggml-org/llama.cpp. Accessed October 30, 2026.

Mathematical Association of America. AIME 2024 and 2025 problems. Competition problems, reproduced with permission by LIVE by Po-Shen Loh, n.d. URL https://live.poshenloh.com/past-contests. AIME I and II, 2024 and 2025. Accessed October 30, 2026.

Merity, S., Xiong, C., Bradbury, J., and Socher, R. Pointer sentinel mixture models. In *International Conference on Learning Representations*, 2017. URL https://openreview.net/forum?id=Byj72udxe.

Meta. Llama-3.1-8B-Instruct. Model card, 2024a. URL https://huggingface.co/meta-llama/Llama-3.1-8B-Instruct. Accessed October 30, 2026.

Meta. Llama-3.3-70B-Instruct. Model card, 2024b. URL https://huggingface.co/meta-llama/Llama-3.3-70B-Instruct. Accessed October 30, 2026.

Muller, L. K., Bich, P., Boretti, C., Chang, H.-M., Zhuang, J., and Cavigelli, L. KVarN: Variance-normalized KV-cache quantization mitigates error accumulation in reasoning tasks, 2026. URL https://arxiv.org/abs/2606.03458.

Na, S., Jeong, G., Ahn, B. H., Jezghani, A., Young, J., Hughes, C. J., Krishna, T., and Kim, H. FlexInfer: Flexible LLM inference with CPU computations. In *Proceedings of Machine Learning and Systems*, volume 7. MLSys, 2025.

Park, Y., Hyun, J., Cho, S., Sim, B., and Lee, J. W. Any-precision LLM: Low-cost deployment of multiple, different-sized LLMs. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 39682–39701. PMLR, 2024. URL https://proceedings.mlr.press/v235/park24e.html.

Park, Y., Hyun, J., Kim, H., and Lee, J. W. DecDEC: A systems approach to advancing low-bit LLM quantization. In *19th USENIX Symposium on Operating Systems Design and Implementation (OSDI 25)*, pp. 803–819. USENIX Association, 2025.

Pope, R., Douglas, S., Chowdhery, A., Devlin, J., Bradbury, J., Levskaya, A., Heek, J., Xiao, K., Agrawal, S., and Dean, J. Efficiently scaling transformer inference. In *Proceedings of Machine Learning and Systems*, volume 5, pp. 606–624, 2023.

Qi, Y., Chen, X., Jiang, H., Wang, Q., Peng, B., and Palpanas, T. ParisKV: Fast and drift-robust KV-cache retrieval for long-context LLMs. In *International Con-*

ference on Machine Learning, volume 306 of *Proceedings of Machine Learning Research*, pp. 100229–100248, 2026. URL https://proceedings.mlr.press/v306/qi26b.html.

Qwen Team. Qwen3 Model Family. Model repository, 2025. URL https://github.com/QwenLM/Qwen3. Accessed October 30, 2026.

Qwen Team. Qwen3.5-9B. Model card, 2026. URL https://huggingface.co/Qwen/Qwen3.5-9B. Accessed October 30, 2026.

Rae, J. W., Potapenko, A., Jayakumar, S. M., Hillier, C., and Lillicrap, T. P. Compressive transformers for long-range sequence modelling. In *International Conference on Learning Representations*, 2020. URL https://openreview.net/forum?id=SylKikSYDH.

Ribar, L., Chelombiev, I., Hudlass-Galley, L., Blake, C., Luschi, C., and Orr, D. SparQ attention: Bandwidth-efficient LLM inference. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 42558–42583. PMLR, 2024. URL https://proceedings.mlr.press/v235/ribar24a.html.

Sheng, Y., Zheng, L., Yuan, B., Li, Z., Ryabinin, M., Chen, B., Liang, P., Ré, C., Stoica, I., and Zhang, C. FlexGen: High-throughput generative inference of large language models with a single GPU. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pp. 31094–31116. PMLR, 2023. URL https://proceedings.mlr.press/v202/sheng23a.html.

Shutova, A., Malinovskii, V., Egiazarian, V., Kuznedelev, D., Mazur, D., Surkov, N., Ermakov, I., and Alistarh, D. Cache me if you must: Adaptive key-value quantization for large language models. In *International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 55451–55473, 2025. URL https://proceedings.mlr.press/v267/shutova25a.html.

Sun, H., Chang, L.-W., Bao, W., Zheng, S., Zheng, N., Liu, X., Dong, H., Chi, Y., and Chen, B. ShadowKV: KV cache in shadows for high-throughput long-context LLM inference. In *International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 57355–57373, 2025. URL https://proceedings.mlr.press/v267/sun25b.html.

Tang, J., Zhao, Y., Zhu, K., Xiao, G., Kasikci, B., and Han, S. QUEST: Query-aware sparsity for efficient long-context LLM inference. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 47901–47911. PMLR, 2024. URL https://proceedings.mlr.press/v235/tang24l.html.

Tseng, A., Sun, Q., Hou, D., and De Sa, C. QTIP: Quantization with trellises and incoherence processing. In *Advances in Neural Information Processing Systems*, volume 37, pp. 59597–59620. Curran Associates, Inc., 2024. doi: 10.52202/079017-1904. URL https://doi.org/10.52202/079017-1904.

Tseng, A., Sun, Z., and De Sa, C. Model-preserving adaptive rounding. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, pp. 122351–122370. PMLR, 2026. URL https://proceedings.mlr.press/v306/tseng26b.html.

Weber, M., Fu, D. Y., Anthony, Q., Oren, Y., Adams, S., Alexandrov, A., Lyu, X., Nguyen, H., Yao, X., Adams, V., Athiwaratkun, B., Chalamala, R., Chen, K., Ryabinin, M., Dao, T., Liang, P., Ré, C., Rish, I., and Zhang, C. RedPajama: an open dataset for training large language models. In *Advances in Neural Information Processing Systems*, volume 37, pp. 116462–116492. Curran Associates, Inc., 2024. doi: 10.52202/079017-3697. URL https://doi.org/10.52202/079017-3697.

Xiao, G., Lin, J., Seznec, M., Wu, H., Demouth, J., and Han, S. SmoothQuant: Accurate and efficient post-training quantization for large language models. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pp. 38087–38099. PMLR, 2023. URL https://proceedings.mlr.press/v202/xiao23c.html.

Yao, D., Shen, B., Lin, Z., Liu, W., Luan, J., Wang, B., and Wang, W. TailorKV: A hybrid framework for long-context inference via tailored KV cache optimization. In *Findings of the Association for Computational Linguistics: ACL 2025*, pp. 20340–20359, 2025. doi: 10.18653/v1/2025.findings-acl.1043. URL https://aclanthology.org/2025.findings-acl.1043/.

Zandieh, A., Daliri, M., Hadian, M., and Mirrokni, V. TurboQuant: Online vector quantization with near-optimal distortion rate. In *International Conference on Learning Representations*, pp. 56418–56439, 2026. URL https://openreview.net/pdf?id=tO3ASKZlok.

Zhang, H., Xia, C., and Wang, Z. KVSwap: Disk-aware KV cache offloading for long-context on-device inference. In *ACM International Conference on Mobile Systems, Applications and Services*, pp. 723–736, 2026. doi: 10.1145/3745756.3809234. URL https://doi.org/10.1145/3745756.3809234.

AI USAGE DISCLOSURE

We used Anthropic's Claude Code and OpenAI's Codex to assist with implementation and debugging, evaluation and visualization scripts, figure and table preparation, and manuscript revision. These tools were used iteratively under author direction. The authors determined the research objectives, directed revisions, and selected the material included in the paper, and remain responsible for the implementation, experimental results, and final manuscript.

APPENDIX CONTENTS

A Evaluation Setup and Implementation Details **16**

B Extended Weight Evaluation **16**

B.1 Accuracy and Downstream Tasks . . . 17
B.2 Long-Context Weight-Only Refinement . . . 17
B.3 Larger Models . . . 18
B.4 Slowdown-Budget Sweeps . . . 18
B.5 YAQA Recovery Fine-Tuning . . . 20
B.6 Refinement and Cache Ablations . . . 20
B.7 Weight Refinement Tuner . . . 21

C Extended KV Evaluation **24**

C.1 Accuracy and Context Length . . . 24
C.2 Alternative Attention Architectures and MoE Models . . . 28
C.3 KV Representation and Selection Controls . . . 29
C.4 Fetch Budgets, RAM Caches, and Miss Policies . . . 29
C.5 Refinement with 2-bit Keys and 1-bit Values . . . 30
C.6 KV Refinement Settings Sweep . . . 31

D Joint Configurations and Task Results **31**

D.1 Configuration Comparisons . . . 31
D.2 Full Task Results . . . 32

E Applicability to Other KV Bases **33**

F Mac Evaluation **34**

G Parallel Weight Refinement with CUDA Streams **34**

A Evaluation Setup and Implementation Details

Perplexity and KL. Weight evaluation uses 2K-token WikiText-2 windows. For long-context evaluation, we sample 17 PG-19 test books with sufficient length and evaluate 16 windows of 128K tokens from each book. We prefill each window except its final 256 tokens, then score those tokens sequentially with teacher forcing. Context-length sweeps shorten the windows to 8K–128K tokens while keeping their endpoints and scored tokens fixed. KL uses the same datasets, windows, and scored tokens as perplexity. We approximate divergence from BF16 by keeping its 1,024 most probable vocabulary tokens separate and grouping all remaining tokens into one category in both distributions.

Evaluation software. Weight evaluation uses the Language Model Evaluation Harness (lm-eval) (Gao et al., 2023). We evaluate LongBench (Bai et al., 2024) using the benchmark's evaluation scripts.

Timing and memory. We implement StrataFlux in llama.cpp, with CUDA 13.4.2 on NVIDIA GPUs and Metal on the Mac. For fair comparisons, we port baselines with author-provided kernels into the same framework. Systems with their own runtimes, including KV offloading systems, vLLM, and ExLlamaV3, are measured in their respective implementations. KV and joint decode timing uses a PG-19 prompt, prefill to $C - 641$ tokens for context length C, 512 untimed steps, and 128 timed steps. For weight-only comparisons on NVIDIA GPUs, we measure decode throughput at empty context for WikiText-2 and at 128K context for PG-19. Memory counts stored representations, metadata, and RAM caches where used; captions specify whether weights, KV, or both are counted. Memory values at 128K context assume 131,072 tokens.

KV refinement settings. Main evaluations use resident key/value bitwidths of 2/2, 3/2, and 3/3, with +4/+4 refinement bits and a 1/128 token-group fetch fraction. Groups contain two tokens on dedicated GPUs and four on unified-memory platforms. Appendix C.6 varies resident and fetched bitwidths, token-group size, and fetch fraction. The first 128 tokens remain in FP16 unless ablated. A 128-token buffer for the current token block also remains in FP16; resident-memory totals include both allocations.

Baseline Kernels Integrated into llama.cpp

YAQA. We integrate QTIP's released decompression and matrix–vector kernels for YAQA's weight format.

SqLLM and LNQ+GQ. SqLLM uses Any-Precision LLM's quantizer and kernels; LNQ+GQ uses GuidedQuant's weights with adapted Any-Precision kernels.

AWQ 4-bit. We integrate the authors' released matrix–vector and matrix–matrix kernels.

KVQuant. We add grouped-query attention support to KVQuant's released quantization and attention kernels.

KIVI. We retain KIVI's released matrix–vector kernels and supply packing and softmax integration in llama.cpp.

KVarN. KVarN baselines use the released Triton kernels.

DecDEC. We port DecDEC's correction kernel into llama.cpp and select its configurations with DecDEC's tuner.

Our Baseline Implementations and Adaptations

AWQ 3-bit. The authors release a 4-bit matrix–vector kernel but no 3-bit kernel, so we implement a group-size-128 3-bit kernel in llama.cpp.

TurboQuant. With no author-released implementation available at integration time, we adapt the community vLLM backend (v0.30.0) to llama.cpp and match the paper's quantization settings.

AQUA-KV. The authors release PyTorch code without CUDA kernels; we adapt it to CUDA, reconstructing cached keys and values during decoding.

SpeCache. With no author-released code available, we implement the paper's low-bit KV cache and selective FP16 fetching in llama.cpp.

Systems Evaluated in Their Own Runtimes

We measure ShadowKV, FreeKV, ParisKV, TailorKV, and KVSwap in their released runtimes. The vLLM, ExLlamaV3, and QServe comparisons also use their respective runtimes; we add Qwen3-8B support to QServe.

StrataFlux Implementation

Weights. We adapt DecDEC's dynamic error-compensation kernel, replacing bucket-based selection with exact radix-sort selection and improving zero-copy request batching at low GPU thread-block counts. We add fetch lead/lag scheduling, RAM caches, and refinement delivery logic, and develop an independent tuner inspired by DecDEC's search procedure.

KV. We adopt KVarN's quantization algorithm and independently develop CUDA and Metal attention kernels with fused refinement operations.

B Extended Weight Evaluation

We extend Sections 4.1 and 4.4 across weight bases, model sizes, slowdown budgets, and cache configurations. We also test weight refinement with YAQA's recovery fine-tuning (Section B.5).

B.1 Accuracy and Downstream Tasks

Tables A1 and A2 extend the 2–4-bit PPL/KL comparisons across platforms; Table A4 adds downstream DecDEC comparisons.

Table A1. Weight-only refinement: WikiText-2 perplexity/KL with DGX Spark configurations at 5% slowdown budgets.

		L3.1-8B-I		Qwen3-8B	
Weights	**Method**	**PPL**	**KL**	**PPL**	**KL**
BF16		7.22	0.000	9.73	0.000
YAQA 2-bit	Base	9.46	0.257	10.95	0.214
	StrataFlux-W	8.94	0.200	10.77	0.175
YAQA 3-bit	Base	7.61	0.047	9.98	0.052
	StrataFlux-W	7.55	0.041	9.89	0.043
YAQA 4-bit	Base	7.33	0.014	9.79	0.016
	StrataFlux-W	7.31	0.012	9.72	0.014
LNQ+GQ 3-bit	Base	8.52	0.157	10.92	0.184
	StrataFlux-W	8.29	0.135	10.70	0.159
LNQ+GQ 4-bit	Base	7.48	0.034	9.95	0.044
	StrataFlux-W	7.43	0.031	9.87	0.038
AWQ 3-bit	Base	9.10	0.229	11.88	0.284
	StrataFlux-W	8.64	0.179	11.30	0.222
AWQ 4-bit	Base	7.59	0.049	10.01	0.059
	StrataFlux-W	7.51	0.039	9.96	0.047
SqLLM 3-bit	Base	9.36	0.263	11.71	0.255
	StrataFlux-W	8.56	0.174	11.03	0.201
SqLLM 4-bit	Base	7.56	0.052	9.98	0.062
	StrataFlux-W	7.48	0.040	9.97	0.048

Table A2. Weight-only refinement: WikiText-2 perplexity/KL with configurations tuned on RTX 5090 at 5% slowdown budgets.

		L3.1-8B-I		Qwen3-8B	
Weights	**Method**	**PPL**	**KL**	**PPL**	**KL**
BF16		7.22	0.000	9.73	0.000
YAQA 2-bit	Base	9.46	0.257	10.95	0.214
	DecDEC	9.03	0.211	10.69	0.184
	StrataFlux-W	8.40	0.139	10.51	0.132
YAQA 3-bit	Base	7.61	0.047	9.98	0.052
	DecDEC	7.55	0.040	9.89	0.044
	StrataFlux-W	7.49	0.033	9.89	0.033
YAQA 4-bit	Base	7.33	0.014	9.79	0.016
	DecDEC	7.30	0.011	9.75	0.013
	StrataFlux-W	7.29	0.010	9.75	0.010
LNQ+GQ 3-bit	Base	8.52	0.157	10.92	0.184
	DecDEC	8.37	0.142	10.66	0.170
	StrataFlux-W	8.18	0.121	10.47	0.139
LNQ+GQ 4-bit	Base	7.48	0.034	9.95	0.044
	DecDEC	7.42	0.031	9.89	0.038
	StrataFlux-W	7.41	0.027	9.88	0.034
AWQ 3-bit	Base	9.10	0.229	11.88	0.284
	DecDEC	8.80	0.195	11.38	0.231
	StrataFlux-W	8.38	0.145	10.71	0.167
AWQ 4-bit	Base	7.59	0.049	10.01	0.059
	DecDEC	7.52	0.042	9.92	0.049
	StrataFlux-W	7.44	0.031	9.88	0.038
SqLLM 3-bit	Base	9.36	0.263	11.71	0.255
	DecDEC	9.02	0.225	11.36	0.215
	StrataFlux-W	8.42	0.156	10.81	0.169
SqLLM 4-bit	Base	7.56	0.052	9.98	0.062
	DecDEC	7.51	0.044	9.86	0.049
	StrataFlux-W	7.44	0.035	9.85	0.038

Table A3. Weight-only refinement on non-Instruct Llama-3.1-8B: WikiText-2 PPL, 2K windows, 5% slowdown budget on RTX 5090.

Weights	**Base**	**DecDEC**	**StrataFlux-W**
BF16	6.24	–	–
LNQ+GQ 3-bit	7.49	7.37	7.21
LNQ+GQ 4-bit	6.53	6.50	6.47
YAQA 2-bit	8.12	7.82	7.38
YAQA 3-bit	6.62	6.56	6.51
YAQA 4-bit	6.35	6.32	6.31

Table A3 reports the non-Instruct Llama comparison.

Table A4. Weight-only refinement on Qwen3-8B: MATH500 and AIME 2024/2025 (avg@8) accuracy at 5% slowdown budgets; StrataFlux-W columns use the configuration tuned on each platform.

	Base		DecDEC		StrataFlux-W DGX Spark		StrataFlux-W RTX 5090	
Weights	**M500**	**AIME**	**M500**	**AIME**	**M500**	**AIME**	**M500**	**AIME**
BF16	97.2	71.0	–	–	–	–	–	–
YAQA 2-bit	77.2	19.2	86.2	30.8	89.4	40.0	92.4	46.7
YAQA 3-bit	95.0	66.7	95.4	66.9	95.0	66.0	95.2	70.0
YAQA 4-bit	95.8	71.5	95.4	67.7	95.6	70.8	97.0	69.8
LNQ+GQ 3-bit	91.8	39.4	93.0	44.4	93.0	47.3	93.2	49.2
LNQ+GQ 4-bit	96.2	64.4	96.4	66.7	95.8	66.0	94.4	67.5
AWQ 3-bit	91.8	46.5	92.8	50.8	92.6	51.9	93.8	51.0
AWQ 4-bit	95.8	64.8	95.2	66.9	96.4	68.5	95.8	65.0
SqLLM 3-bit	93.0	47.3	92.8	49.8	94.2	48.3	93.0	50.8
SqLLM 4-bit	95.2	64.2	96.0	67.1	95.6	66.5	95.2	70.2

B.2 Long-Context Weight-Only Refinement

Table A5 evaluates weight refinement at 128K context with FP16 KV.

Table A5. Weight-only refinement: PG-19 PPL/KL and RTX 5090 tok/s at 128K with FP16 KV, on Llama-3.1-8B-Instruct and Qwen3-8B. StrataFlux-W reuses its empty-context 5% configurations.

		L3.1-8B-I			Qwen3-8B		
Weights	**Method**	**PPL**	**KL**	**tok/s**	**PPL**	**KL**	**tok/s**
BF16		10.32	0.000	47.6	12.89	0.000	OOM
YAQA 2-bit	Base	15.89	0.473	64.2	16.55	0.254	58.2
	StrataFlux-W	13.73	0.311	57.3	14.83	0.162	50.2
YAQA 3-bit	Base	10.87	0.057	62.5	13.38	0.045	57.0
	StrataFlux-W	10.71	0.042	56.0	13.16	0.034	47.2
YAQA 4-bit	Base	10.51	0.016	61.4	13.02	0.014	55.4
	StrataFlux-W	10.47	0.013	60.0	12.95	0.010	48.6
LNQ+GQ 3-bit	Base	13.65	0.303	68.1	18.15	0.328	61.2
	StrataFlux-W	12.96	0.243	67.3	17.21	0.283	60.6
LNQ+GQ 4-bit	Base	10.77	0.043	65.7	13.28	0.037	59.1
	StrataFlux-W	10.69	0.037	65.2	13.22	0.032	58.1

B.3 Larger Models

Table A6 compares weight refinement on Llama-3.3-70B-Instruct and Qwen3-32B.

Table A6. Weight-only refinement: WikiText-2 perplexity and KL divergence on Llama-3.3-70B-Instruct and Qwen3-32B at matched 5% slowdown budgets. Throughput uses empty context: Qwen3-32B on RTX 5090; * Llama-3.3-70B-Instruct on RTX PRO 6000.

		L3.3-70B-I			Qwen3-32B		
Weights	**Method**	**PPL**	**KL**	**tok/s***	**PPL**	**KL**	**tok/s**
BF16	Base	3.86	0.000	OOM	7.61	0.000	OOM
YAQA 2-bit	Base	6.19	0.449	43.5	8.30	0.211	76.2
	DecDEC	5.71	0.373	42.8	8.04	0.164	71.8
	StrataFlux-W	5.52	0.349	41.9	8.04	0.153	72.8
YAQA 3-bit	Base	4.50	0.141	36.9	7.70	0.064	67.7
	DecDEC	4.36	0.111	35.6	7.68	0.049	62.9
	StrataFlux-W	4.33	0.107	35.3	7.69	0.043	62.5
YAQA 4-bit	Base	4.06	0.046	32.7	7.62	0.023	62.9
	DecDEC	4.00	0.035	31.3	7.61	0.016	58.7
	StrataFlux-W	3.99	0.033	31.4	7.62	0.015	59.6
LNQ+GQ 3-bit	Base	5.39	0.320	46.6	8.15	0.173	98.0
	DecDEC	5.15	0.282	43.5	8.10	0.153	92.9
	StrataFlux-W	5.05	0.261	44.9	7.94	0.132	92.6
LNQ+GQ 4-bit	Base	4.26	0.096	36.5	7.72	0.053	78.2
	DecDEC	4.20	0.082	34.0	7.65	0.044	74.4
	StrataFlux-W	4.17	0.078	35.1	7.65	0.039	73.4
AWQ 3-bit	Base	5.75	0.371	35.3	8.56	0.241	83.8
	DecDEC	5.09	0.260	34.0	8.30	0.187	81.2
	StrataFlux-W	5.04	0.244	34.1	8.01	0.146	80.8
AWQ 4-bit	Base	4.40	0.110	37.8	7.84	0.070	77.6
	DecDEC	4.16	0.069	36.0	7.76	0.052	73.8
	StrataFlux-W	4.16	0.066	36.1	7.71	0.043	73.3
SqLLM 3-bit	Base	5.64	0.377	46.3	8.68	0.242	98.1
	DecDEC	5.10	0.272	44.2	8.52	0.211	94.6
	StrataFlux-W	5.07	0.253	44.5	8.09	0.148	94.3
SqLLM 4-bit	Base	4.28	0.105	36.1	7.82	0.071	78.1
	DecDEC	4.17	0.079	34.1	7.71	0.049	74.5
	StrataFlux-W	4.17	0.074	34.9	7.69	0.043	73.8

B.4 Slowdown-Budget Sweeps

Tables A7, A8, and A9 compare weight-refinement accuracy and throughput across slowdown budgets on DGX Spark, RTX 5090, and RTX 5070 Ti.

Table A7. Weight-only refinement on DGX Spark: 2–4-bit bases across slowdown budgets. Columns report WikiText-2 PPL/KL and Llama-3.1-8B-Instruct empty-context tok/s.

DGX Spark

		L3.1-8B-I			Qwen3-8B	
Weights	**Budget / method**	**PPL**	**KL**	**tok/s**	**PPL**	**KL**
BF16	Base	7.22	0.000	15.6	9.73	0.000
YAQA 2-bit	Base	9.46	0.257	72.0	10.95	0.214
	2.5% / StrataFlux-W	8.96	0.202	70.9	10.81	0.178
	5% / StrataFlux-W	8.94	0.200	69.2	10.77	0.175
	10% / StrataFlux-W	8.90	0.194	65.7	10.67	0.174
	20% / StrataFlux-W	8.92	0.198	60.0	10.69	0.174
YAQA 3-bit	Base	7.61	0.047	57.7	9.98	0.052
	2.5% / StrataFlux-W	7.56	0.042	57.4	9.91	0.044
	5% / StrataFlux-W	7.55	0.041	55.8	9.89	0.043
	10% / StrataFlux-W	7.56	0.041	53.1	9.89	0.045
	20% / StrataFlux-W	7.56	0.041	49.0	9.88	0.044

DGX Spark

		L3.1-8B-I			Qwen3-8B	
Weights	**Budget / method**	**PPL**	**KL**	**tok/s**	**PPL**	**KL**
LNQ+GQ 3-bit	Base	8.52	0.157	60.9	10.92	0.184
	2.5% / StrataFlux-W	8.34	0.139	61.4	10.73	0.160
	5% / StrataFlux-W	8.29	0.135	59.8	10.70	0.159
	10% / StrataFlux-W	8.31	0.135	57.1	10.64	0.149
	20% / StrataFlux-W	8.31	0.135	52.8	10.66	0.154
AWQ 3-bit	Base	9.10	0.229	63.3	11.88	0.284
	2.5% / StrataFlux-W	8.75	0.190	61.1	11.33	0.234
	5% / StrataFlux-W	8.64	0.179	59.2	11.30	0.222
	10% / StrataFlux-W	8.63	0.175	56.2	11.16	0.214
	20% / StrataFlux-W	8.64	0.177	52.9	11.12	0.211
SqLLM 3-bit	Base	9.36	0.263	60.8	11.71	0.255
	2.5% / StrataFlux-W	8.72	0.194	60.5	11.19	0.210
	5% / StrataFlux-W	8.56	0.174	59.3	11.03	0.201
	10% / StrataFlux-W	8.51	0.170	57.4	11.14	0.187
	20% / StrataFlux-W	8.51	0.171	53.1	11.11	0.186
YAQA 4-bit	Base	7.33	0.014	48.1	9.79	0.016
	2.5% / StrataFlux-W	7.31	0.012	47.5	9.73	0.014
	5% / StrataFlux-W	7.31	0.012	46.6	9.72	0.014
	10% / StrataFlux-W	7.31	0.012	44.4	9.71	0.014
	20% / StrataFlux-W	7.31	0.012	40.9	9.72	0.014
LNQ+GQ 4-bit	Base	7.48	0.034	49.4	9.95	0.044
	2.5% / StrataFlux-W	7.45	0.031	49.9	9.88	0.038
	5% / StrataFlux-W	7.43	0.031	49.0	9.87	0.038
	10% / StrataFlux-W	7.44	0.031	46.7	9.88	0.039
	20% / StrataFlux-W	7.44	0.031	42.3	9.86	0.038
AWQ 4-bit	Base	7.59	0.049	49.5	10.01	0.059
	2.5% / StrataFlux-W	7.52	0.040	48.8	9.96	0.047
	5% / StrataFlux-W	7.51	0.039	46.1	9.96	0.047
	10% / StrataFlux-W	7.51	0.038	45.4	9.90	0.046
	20% / StrataFlux-W	7.51	0.038	42.4	9.90	0.045
SqLLM 4-bit	Base	7.56	0.052	49.5	9.98	0.062
	2.5% / StrataFlux-W	7.50	0.043	48.1	9.89	0.050
	5% / StrataFlux-W	7.48	0.040	48.6	9.97	0.048
	10% / StrataFlux-W	7.47	0.039	46.6	9.87	0.046
	20% / StrataFlux-W	7.47	0.039	44.0	9.88	0.046

Table A8. Weight-only refinement on RTX 5090: 2–4-bit bases across slowdown budgets. Columns report WikiText-2 PPL/KL and empty-context decode tok/s.

RTX 5090

		L3.1-8B-I			Qwen3-8B		
Weights	**Budget / method**	**PPL**	**KL**	**tok/s**	**PPL**	**KL**	**tok/s**
BF16	Base	7.22	0.000	104.7	9.73	0.000	103.0
YAQA 2-bit	Base	9.46	0.257	245.8	10.95	0.214	232.3
	2.5% / DecDEC	8.90	0.197	239.4	10.57	0.172	223.1
	2.5% / StrataFlux-W	8.49	0.149	240.3	10.66	0.144	227.7
	5% / DecDEC	9.03	0.211	233.2	10.69	0.184	220.8
	5% / StrataFlux-W	8.40	0.139	234.4	10.51	0.132	224.4
	10% / DecDEC	8.87	0.193	228.9	10.53	0.170	217.0
	10% / StrataFlux-W	8.39	0.137	223.9	10.43	0.123	209.6
	20% / DecDEC	8.78	0.183	218.3	10.46	0.163	205.5
	20% / StrataFlux-W	8.37	0.134	211.6	10.37	0.120	192.9
YAQA 3-bit	Base	7.61	0.047	222.5	9.98	0.052	213.1
	2.5% / DecDEC	7.57	0.041	215.1	9.93	0.046	206.1
	2.5% / StrataFlux-W	7.50	0.034	217.5	9.89	0.036	208.3
	5% / DecDEC	7.55	0.040	213.3	9.89	0.044	202.0
	5% / StrataFlux-W	7.49	0.033	212.7	9.89	0.033	202.8
	10% / DecDEC	7.54	0.038	209.3	9.88	0.042	198.8
	10% / StrataFlux-W	7.48	0.032	202.6	9.87	0.032	194.2
	20% / DecDEC	7.52	0.037	197.5	9.87	0.041	191.7
	20% / StrataFlux-W	7.47	0.030	184.8	9.86	0.031	171.5

RTX 5090

		L3.1-8B-I			Qwen3-8B		
Weights	**Budget / method**	**PPL**	**KL**	**tok/s**	**PPL**	**KL**	**tok/s**
	Base	8.52	0.157	314.3	10.92	0.184	290.9
	2.5% / DecDEC	8.40	0.146	312.4	10.78	0.187	290.5
	2.5% / StrataFlux-W	8.20	0.122	306.0	10.51	0.144	284.5
	5% / DecDEC	8.37	0.142	311.5	10.66	0.170	287.1
LNQ+GQ 3-bit	5% / StrataFlux-W	8.18	0.121	299.7	10.47	0.139	281.1
	10% / DecDEC	8.32	0.138	302.8	10.55	0.159	283.0
	10% / StrataFlux-W	8.11	0.112	285.8	10.54	0.130	263.7
	20% / DecDEC	8.39	0.147	287.4	10.85	0.181	265.6
	20% / StrataFlux-W	8.07	0.105	261.6	10.42	0.117	240.4
	Base	9.10	0.229	304.5	11.88	0.284	286.0
	2.5% / DecDEC	8.88	0.204	294.4	11.63	0.265	281.5
	2.5% / StrataFlux-W	8.60	0.170	295.0	11.24	0.208	273.9
	5% / DecDEC	8.80	0.195	292.6	11.38	0.231	276.1
AWQ 3-bit	5% / StrataFlux-W	8.38	0.145	289.0	10.71	0.167	271.5
	10% / DecDEC	8.76	0.189	289.5	11.24	0.220	273.2
	10% / StrataFlux-W	8.31	0.140	270.9	10.49	0.143	259.0
	20% / DecDEC	8.56	0.166	270.2	11.21	0.218	253.8
	20% / StrataFlux-W	8.19	0.120	255.7	10.50	0.144	238.3
	Base	9.36	0.263	315.4	11.71	0.255	292.1
	2.5% / DecDEC	9.10	0.233	312.2	11.42	0.234	291.2
	2.5% / StrataFlux-W	8.47	0.162	307.3	10.85	0.174	267.1
	5% / DecDEC	9.02	0.225	311.3	11.36	0.215	287.3
SqLLM 3-bit	5% / StrataFlux-W	8.42	0.156	299.6	10.81	0.169	280.3
	10% / DecDEC	8.91	0.213	305.2	11.20	0.201	283.8
	10% / StrataFlux-W	8.33	0.143	288.6	10.62	0.145	264.0
	20% / DecDEC	8.86	0.209	289.1	11.41	0.216	266.2
	20% / StrataFlux-W	8.13	0.121	263.6	10.53	0.132	244.9
	Base	7.33	0.014	204.2	9.79	0.016	193.8
	2.5% / DecDEC	7.30	0.012	197.6	9.76	0.015	185.9
	2.5% / StrataFlux-W	7.29	0.010	199.1	9.77	0.011	189.0
	5% / DecDEC	7.30	0.011	195.9	9.75	0.013	183.3
YAQA 4-bit	5% / StrataFlux-W	7.29	0.010	194.3	9.75	0.010	184.1
	10% / DecDEC	7.30	0.011	190.6	9.74	0.013	180.6
	10% / StrataFlux-W	7.29	0.010	185.5	9.75	0.010	176.0
	20% / DecDEC	7.29	0.010	180.5	9.74	0.012	172.5
	20% / StrataFlux-W	7.29	0.009	169.5	9.74	0.010	158.4
	Base	7.48	0.034	264.7	9.95	0.044	248.3
	2.5% / DecDEC	7.43	0.031	263.8	9.89	0.038	247.6
	2.5% / StrataFlux-W	7.42	0.028	258.9	9.88	0.036	242.5
	5% / DecDEC	7.42	0.031	263.1	9.89	0.038	244.9
LNQ+GQ 4-bit	5% / StrataFlux-W	7.41	0.027	254.5	9.88	0.034	238.0
	10% / DecDEC	7.45	0.032	249.9	9.89	0.041	236.7
	10% / StrataFlux-W	7.40	0.027	241.9	9.89	0.033	225.9
	20% / DecDEC	7.45	0.032	238.8	9.96	0.043	224.2
	20% / StrataFlux-W	7.39	0.025	218.9	9.89	0.031	193.4
	Base	7.59	0.049	262.8	10.01	0.059	247.5
	2.5% / DecDEC	7.53	0.043	259.5	9.95	0.055	245.2
	2.5% / StrataFlux-W	7.45	0.033	257.4	9.88	0.040	241.4
	5% / DecDEC	7.52	0.042	257.7	9.92	0.049	242.2
AWQ 4-bit	5% / StrataFlux-W	7.44	0.031	250.4	9.88	0.038	236.9
	10% / DecDEC	7.51	0.041	254.1	9.90	0.047	238.0
	10% / StrataFlux-W	7.43	0.029	237.0	9.81	0.034	223.0
	20% / DecDEC	7.47	0.036	236.2	9.89	0.045	226.1
	20% / StrataFlux-W	7.40	0.028	218.4	9.79	0.031	206.0
	Base	7.56	0.052	265.6	9.98	0.062	249.0
	2.5% / DecDEC	7.51	0.044	262.9	9.87	0.050	246.4
	2.5% / StrataFlux-W	7.46	0.037	257.9	9.89	0.042	239.3
	5% / DecDEC	7.51	0.044	262.2	9.86	0.049	245.0
SqLLM 4-bit	5% / StrataFlux-W	7.44	0.035	254.6	9.85	0.038	235.7
	10% / DecDEC	7.52	0.044	254.6	9.87	0.050	234.8
	10% / StrataFlux-W	7.42	0.032	241.9	9.81	0.038	226.3
	20% / DecDEC	7.52	0.044	238.8	9.88	0.051	222.2
	20% / StrataFlux-W	7.41	0.029	216.4	9.81	0.033	205.8

Table A9. Weight-only refinement on RTX 5070 Ti: 2–4-bit bases across slowdown budgets. Columns report WikiText-2 PPL/KL and empty-context decode tok/s.

RTX 5070 Ti

		L3.1-8B-I			Qwen3-8B		
Weights	**Budget / method**	**PPL**	**KL**	**tok/s**	**PPL**	**KL**	**tok/s**
	Base	9.46	0.257	165.9	10.95	0.214	158.7
	2.5% / DecDEC	9.06	0.215	155.8	10.58	0.175	146.4
	2.5% / StrataFlux-W	8.54	0.153	158.9	10.54	0.136	149.8
	5% / DecDEC	8.97	0.205	152.8	10.56	0.173	144.6
YAQA 2-bit	5% / StrataFlux-W	8.49	0.147	154.2	10.45	0.129	149.4
	10% / DecDEC	8.89	0.196	149.5	10.54	0.171	141.0
	10% / StrataFlux-W	8.37	0.135	148.9	10.36	0.120	143.5
	20% / DecDEC	8.61	0.164	139.1	10.37	0.151	134.1
	20% / StrataFlux-W	8.30	0.127	135.5	10.34	0.115	130.8
	Base	7.61	0.047	149.4	9.98	0.052	145.1
	2.5% / DecDEC	7.57	0.041	141.7	9.90	0.044	136.9
	2.5% / StrataFlux-W	7.51	0.034	144.5	9.90	0.035	138.6
	5% / DecDEC	7.56	0.040	138.4	9.89	0.043	133.8
YAQA 3-bit	5% / StrataFlux-W	7.49	0.032	141.0	9.88	0.034	137.6
	10% / DecDEC	7.51	0.035	133.8	9.88	0.042	129.1
	10% / StrataFlux-W	7.47	0.030	135.5	9.87	0.031	130.0
	20% / DecDEC	7.48	0.032	125.0	9.83	0.037	120.5
	20% / StrataFlux-W	7.45	0.028	124.2	9.85	0.030	120.0
	Base	8.52	0.157	183.1	10.92	0.184	172.5
	2.5% / DecDEC	8.23	0.126	179.4	10.42	0.145	169.4
	2.5% / StrataFlux-W	8.22	0.123	176.3	10.48	0.123	167.0
	5% / DecDEC	8.33	0.141	176.1	10.69	0.167	165.9
LNQ+GQ 3-bit	5% / StrataFlux-W	8.03	0.103	170.7	10.51	0.126	163.3
	10% / DecDEC	8.33	0.141	170.1	10.70	0.171	160.8
	10% / StrataFlux-W	8.06	0.103	165.2	10.42	0.117	157.1
	20% / DecDEC	8.27	0.133	161.0	10.54	0.152	152.4
	20% / StrataFlux-W	8.00	0.098	147.8	10.39	0.109	139.6
	Base	9.10	0.229	171.0	11.88	0.284	163.2
	2.5% / DecDEC	8.86	0.202	167.2	11.53	0.243	156.9
	2.5% / StrataFlux-W	8.52	0.163	166.4	11.14	0.204	156.5
	5% / DecDEC	8.60	0.172	162.9	11.07	0.207	154.5
AWQ 3-bit	5% / StrataFlux-W	8.35	0.142	161.4	10.69	0.153	154.0
	10% / DecDEC	8.53	0.162	156.5	11.04	0.203	148.6
	10% / StrataFlux-W	8.17	0.117	154.2	10.40	0.128	147.4
	20% / DecDEC	8.40	0.146	147.7	10.71	0.170	140.4
	20% / StrataFlux-W	8.04	0.102	141.7	10.33	0.114	134.0
	Base	9.36	0.263	183.1	11.71	0.255	172.6
	2.5% / DecDEC	9.09	0.234	179.2	10.90	0.177	169.4
	2.5% / StrataFlux-W	8.44	0.158	173.0	11.02	0.179	164.1
	5% / DecDEC	8.64	0.183	175.8	11.21	0.192	166.0
SqLLM 3-bit	5% / StrataFlux-W	8.17	0.129	166.7	10.60	0.138	163.6
	10% / DecDEC	8.71	0.191	170.3	11.27	0.199	160.2
	10% / StrataFlux-W	8.08	0.114	166.3	10.50	0.125	153.0
	20% / DecDEC	8.32	0.145	160.9	10.71	0.150	152.5
	20% / StrataFlux-W	7.97	0.102	149.0	10.38	0.114	142.9
	Base	7.33	0.014	132.2	9.79	0.016	128.0
	2.5% / DecDEC	7.30	0.012	126.8	9.75	0.014	120.8
	2.5% / StrataFlux-W	7.28	0.009	127.7	9.76	0.011	123.1
	5% / DecDEC	7.29	0.010	123.5	9.74	0.013	118.1
YAQA 4-bit	5% / StrataFlux-W	7.29	0.009	124.2	9.75	0.010	120.6
	10% / DecDEC	7.28	0.009	118.2	9.73	0.011	113.6
	10% / StrataFlux-W	7.28	0.008	119.2	9.75	0.010	115.8
	20% / DecDEC	7.28	0.009	110.6	9.73	0.011	106.4
	20% / StrataFlux-W	7.28	0.008	109.2	9.75	0.010	106.2

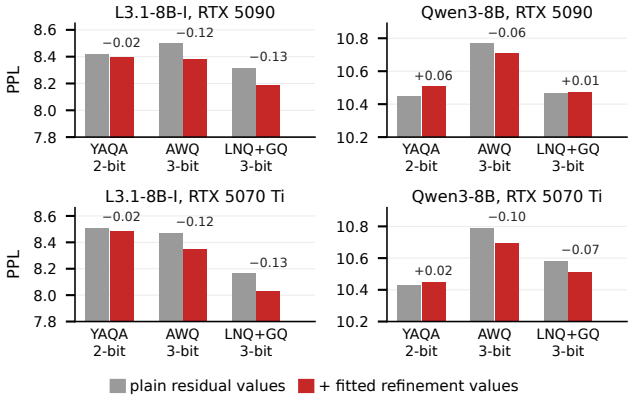

Figure A1. Fitted versus unfitted weight refinement on WikiText-2, at the 5% configurations tuned on RTX 5090 and RTX 5070 Ti. Each pair fixes selection counts, schedule, and MLP scheme: gray bars use plain residuals; red bars use fitted refinement. Labels give the PPL change from fitting.

RTX 5070 Ti

		L3.1-8B-I			Qwen3-8B		
Weights	**Budget / method**	**PPL**	**KL**	**tok/s**	**PPL**	**KL**	**tok/s**
LNQ+GQ 4-bit	Base	7.48	0.034	148.3	9.95	0.044	141.3
	2.5% / DecDEC	7.43	0.030	144.3	9.89	0.041	139.0
	2.5% / StrataFlux-W	7.38	0.024	142.1	9.90	0.032	137.8
	5% / DecDEC	7.42	0.029	142.2	9.85	0.039	136.5
	5% / StrataFlux-W	7.38	0.024	139.0	9.88	0.031	133.8
	10% / DecDEC	7.41	0.028	137.5	9.83	0.037	132.4
	10% / StrataFlux-W	7.38	0.023	133.4	9.88	0.030	129.5
	20% / DecDEC	7.40	0.026	129.0	9.83	0.035	124.5
	20% / StrataFlux-W	7.37	0.022	121.7	9.87	0.029	118.7
AWQ 4-bit	Base	7.59	0.049	152.4	10.01	0.059	144.8
	2.5% / DecDEC	7.47	0.036	148.4	9.88	0.043	140.8
	2.5% / StrataFlux-W	7.41	0.029	147.2	9.86	0.035	139.7
	5% / DecDEC	7.48	0.037	145.7	9.95	0.048	138.6
	5% / StrataFlux-W	7.40	0.026	143.1	9.80	0.032	136.2
	10% / DecDEC	7.44	0.031	140.8	9.83	0.037	133.6
	10% / StrataFlux-W	7.38	0.024	136.5	9.77	0.029	130.6
	20% / DecDEC	7.40	0.026	132.1	9.80	0.034	126.1
	20% / StrataFlux-W	7.36	0.021	126.8	9.84	0.026	119.5
SqLLM 4-bit	Base	7.56	0.052	148.3	9.98	0.062	141.4
	2.5% / DecDEC	7.46	0.037	144.5	9.85	0.048	138.9
	2.5% / StrataFlux-W	7.38	0.028	142.5	9.83	0.036	137.3
	5% / DecDEC	7.44	0.035	142.2	9.79	0.041	136.4
	5% / StrataFlux-W	7.38	0.026	139.4	9.78	0.032	134.3
	10% / DecDEC	7.41	0.031	137.5	9.77	0.038	132.3
	10% / StrataFlux-W	7.37	0.025	134.0	9.79	0.030	129.8
	20% / DecDEC	7.39	0.029	129.0	9.75	0.034	124.4
	20% / StrataFlux-W	7.36	0.023	123.7	9.79	0.027	118.3

B.5 YAQA Recovery Fine-Tuning

StrataFlux-W is compatible with YAQA recovery fine-tuning. We use QTIP's RedPajama fine-tuning recipe, but freeze embeddings, final normalization, and the output head during end-to-end fine-tuning; Qwen3-8B's query/key norms remain frozen during both block-wise and end-to-end fine-tuning. This preserves the BF16 refinement target; normalization changes elsewhere are absorbed by rescaling its projection weights. We find that this preserves the gains of fine-tuning while maintaining compatibility with weight refinement (Table A10).

Table A10. Weight-only refinement after YAQA recovery fine-tuning: WikiText-2 PPL/KL using configurations tuned on RTX 5090. Original and fine-tuned bases are shown for each bitwidth; both DecDEC and StrataFlux-W refine the fine-tuned weights.

			L3.1-8B-I		Qwen3-8B	
Weights	**Method**	**Budget**	**PPL**	**KL**	**PPL**	**KL**
YAQA 2-bit	Original base		9.46	0.257	10.95	0.214
	Fine-tuned base		8.45	0.161	10.24	0.162
	DecDEC	2.5%	8.32	0.142	10.22	0.146
	StrataFlux-W	2.5%	8.24	0.133	10.50	0.137
	DecDEC	5%	8.35	0.149	10.26	0.153
	StrataFlux-W	5%	8.19	0.128	10.40	0.125
	DecDEC	10%	8.32	0.141	10.21	0.146
	StrataFlux-W	10%	8.18	0.126	10.37	0.120
	DecDEC	20%	8.28	0.137	10.14	0.140
	StrataFlux-W	20%	8.22	0.126	10.36	0.117
YAQA 3-bit	Original base		7.61	0.047	9.98	0.052
	Fine-tuned base		7.50	0.040	9.73	0.046
	DecDEC	2.5%	7.48	0.036	9.75	0.043
	StrataFlux-W	2.5%	7.48	0.032	9.78	0.035
	DecDEC	5%	7.48	0.035	9.75	0.041
	StrataFlux-W	5%	7.47	0.032	9.79	0.032
	DecDEC	10%	7.48	0.035	9.74	0.040
	StrataFlux-W	10%	7.47	0.031	9.77	0.032
	DecDEC	20%	7.47	0.033	9.73	0.039
	StrataFlux-W	20%	7.45	0.029	9.76	0.031
YAQA 4-bit	Original base		7.33	0.014	9.79	0.016
	Fine-tuned base		7.28	0.013	9.71	0.015
	DecDEC	2.5%	7.27	0.011	9.70	0.014
	StrataFlux-W	2.5%	7.27	0.010	9.73	0.011
	DecDEC	5%	7.27	0.011	9.69	0.013
	StrataFlux-W	5%	7.28	0.009	9.74	0.010
	DecDEC	10%	7.27	0.011	9.69	0.013
	StrataFlux-W	10%	7.27	0.009	9.74	0.010
	DecDEC	20%	7.26	0.010	9.69	0.012
	StrataFlux-W	20%	7.27	0.009	9.74	0.010

B.6 Refinement and Cache Ablations

Figure A1 compares fitted and unfitted weight refinement at fixed configurations. Figure A2 compares skipping and waiting on weight-cache misses across slowdown budgets. Tables A11 and A12 give the cache measurements and paired technique ablations.

Table A11. Weight-cache capacity and miss policy on DGX Spark: Llama-3.1-8B-Instruct, 5% configuration, YAQA 2-bit and AWQ 3-bit. WikiText-2 PPL/KL and decode tok/s; capacities in MiB. Raw values behind Figure 12(d,e).

		YAQA 2-bit			AWQ 3-bit		
MiB	**Miss policy**	**PPL**	**KL**	**tok/s**	**PPL**	**KL**	**tok/s**
64	skip	9.01	0.207	68.8	8.70	0.186	57.4
128	skip	8.94	0.200	69.2	8.64	0.179	57.6
256	skip	8.88	0.192	69.7	8.59	0.172	57.6
512	skip	8.80	0.183	69.4	8.53	0.166	57.4
128	wait	8.52	0.152	17.9	8.31	0.137	27.7

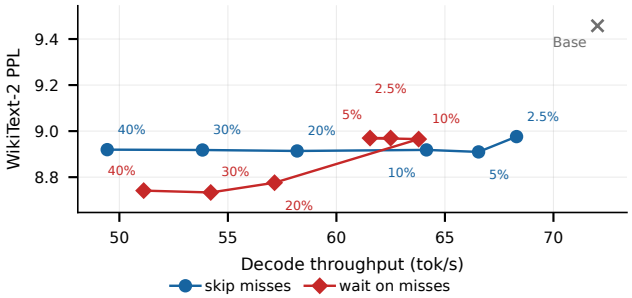

Figure A2. Waiting versus skipping weight-cache misses on Llama-3.1-8B-Instruct with YAQA 2-bit weights, DGX Spark, and a 128 MiB cache. Curves plot WikiText-2 PPL against throughput. Points follow slowdown budgets of 2.5%, 5%, 10%, 20%, 30%, and 40%. Measured wait-policy throughput is non-monotonic across budgets.

Table A12. Paired weight-refinement ablations on Llama-3.1-8B-Instruct, RTX 5090: WikiText-2 PPL/KL and empty-context decode tok/s. Pairs are tuned separately at a 5% slowdown budget, except in the unchanged-settings pair: there, $s = 0$ disables lead/lag in the final configuration while retaining its k, fitted refinements, MLP scheme, and GPU thread-block counts.

		YAQA 2-bit			AWQ 3-bit		
Comparison	**Setting**	**PPL**	**KL**	**tok/s**	**PPL**	**KL**	**tok/s**
Fitting, $s = 0$	Unfitted	8.76	0.181	234.1	8.70	0.182	288.1
	Fitted	8.67	0.170	230.1	8.61	0.171	287.5
Fitting, tuned s	Unfitted	8.68	0.174	234.9	8.50	0.158	290.2
	Fitted	8.59	0.162	234.3	8.38	0.145	289.0
MLP scheme	Channel-wise	8.59	0.162	234.3	8.38	0.145	289.0
	Neuron bundles	8.40	0.139	234.4	8.56	0.165	289.3
Lead/lag, unfitted	$s = 0$	8.76	0.181	234.1	8.70	0.182	288.1
	Tuned s	8.68	0.174	234.9	8.50	0.158	290.2
Lead/lag, full method	$s = 0$	8.57	0.157	232.2	8.57	0.168	286.8
	Tuned s	8.40	0.139	234.4	8.38	0.145	289.0
Lead/lag, unchanged settings	$s = 0$	8.34	0.131	182.2	8.33	0.141	271.1
	Tuned s	8.40	0.139	234.4	8.38	0.145	289.0

B.7 Weight Refinement Tuner

The tuner searches separately for each model, weight quantizer, and platform under a specified slowdown budget relative to the quantized base. It selects channel budgets k, fetch lead/lag distances s, the MLP refinement scheme, and GPU thread-block counts.

Resource allocation. We first profile base projections alongside refinement while varying its thread-block count and selection budget. The resulting timing table identifies launch configurations that balance refinement speed against interference with the base projections.

Budget-constrained search. For each combination of shift distances and channel-wise or neuron-wise MLP refinement, the tuner searches for the largest common channel budget k that satisfies the slowdown limit in end-to-end decoding. For neuron-wise MLP refinement, it sets $k_n = \mathrm{round}(3k/4)$: each increment of k_n fetches $4/3$ as many bytes as each increment of k for channel-wise gate/up and down refinement. It also considers disabling refinement for projection sets when needed to meet the budget. On DGX Spark, we search skip-policy configurations only at $s = 0$: skipping cache misses leaves no SSD wait for fetch lead/lag to hide.

Accuracy selection and budget allocation. The tuner compares feasible candidates by KL divergence from BF16 on held-out text. Starting from the lowest-KL candidate, it increases individual selection budgets while remeasuring end-to-end latency and retaining increases that satisfy the slowdown limit. These final per-projection budgets need not be equal.

Refinement fitting during search. Fitted refinement values depend on the selection budget (Section 3.2.1). We approximate candidate KL during search using refinements pre-fitted at a ladder of power-of-two selection budgets, avoiding a separate fit for each candidate. Once the final per-projection budgets are selected, we fit refinement values at those exact budgets before quantizing them for inference.

Tuned configurations. Table A13 lists tuned configurations across slowdown budgets and miss-policy sweeps. Each row displays one configuration. The Weights column gives the quantizer and bitwidth. Four-entry tuples give the query/key/value, attention-output, gate/up, and down configurations for channel-wise refinement. Three-entry tuples give query/key/value, attention-output, and neuron bundles. This order applies to selection counts (k, or k_n for neuron bundles) and shift distances s. n_{tb} lists the GPU thread-block count shared by channel-wise kernels, followed by the count for the neuron-bundle kernel when used. Counts k are per 1,024 input channels or neurons, and s is in Transformer blocks. SSD-backed configurations use a 128 MiB cache; miss-policy sweeps identify skipping or waiting.

Table A13. Tuned weight-refinement configurations. Model and platform headings identify each configuration block. The main DGX Spark skip-policy search uses only $s = 0$; skip/wait policy sweeps are listed separately.

Llama-3.1-8B-Instruct / DGX Spark				
Weights	**Budget**	$\boldsymbol{k}$	$\boldsymbol{s}$	$\boldsymbol{n_{\mathrm{tb}}}$
AWQ 3	2.5%	3,3,2	0,0,0	7,4
	5%	10,14,8	0,0,0	2,4
	10%	21,21,21,21	0,0,0,0	3
	20%	31,31,23	0,0,0	3,8
AWQ 4	2.5%	5,6,0,5	0,0,0,0	3
	5%	16,16,14,16	0,0,0,0	4
	10%	28,28,0,28	0,0,0,0	3
	20%	34,33,33,33	0,0,0,0	6
LNQ+GQ 3	2.5%	0,4,3	0,0,0	7,4
	5%	11,0,8	0,0,0	6,4
	10%	19,19,14	0,0,0	6,8
	20%	33,32,24	0,0,0	5,16

Llama-3.1-8B-Instruct / DGX Spark

Weights	Budget	k	s	n_{tb}
LNQ+GQ 4	2.5%	8,8,0,8	0,0,0,0	7
	5%	14,0,9	0,0,0	5,4
	10%	23,24,16	0,0,0	7,8
	20%	40,40,29	0,0,0	8,8
SqLLM 3	2.5%	9,8,6	0,0,0	6,4
	5%	10,14,9,9	0,0,0,0	3
	10%	18,18,19,18	0,0,0,0	4
	20%	28,29,28,27	0,0,0,0	3
SqLLM 4	2.5%	0,15,15,15	0,0,0,0	4
	5%	10,9,10,10	0,0,0,0	6
	10%	21,21,21,21	0,0,0,0	5
	20%	33,33,33,33	0,0,0,0	6
YAQA 2	2.5%	8,0,5	0,0,0	1,4
	5%	14,12,8	0,0,0	2,4
	10%	0,26,18	0,0,0	2,8
	20%	32,33,22	0,0,0	3,8
YAQA 3	2.5%	5,5,4	0,0,0	1,8
	5%	16,16,12	0,0,0	2,4
	10%	28,28,21	0,0,0	3,8
	20%	35,35,26	0,0,0	3,4
YAQA 4	2.5%	0,6,4	0,0,0	1,4
	5%	18,18,12	0,0,0	2,4
	10%	0,32,24	0,0,0	2,8
	20%	42,43,30	0,0,0	3,8

Llama-3.1-8B-Instruct / DGX Spark, skip-policy sweep

Weights	Budget	k	s	n_{tb}
YAQA 2	2.5%	6,12,4	0,0,0	3,4
	5%	0,17,12	0,0,0	5,4
	10%	24,23,18	0,0,0	2,8
	20%	0,37,28	0,0,0	4,8
	30%	0,40,30	0,0,0	13,4
	40%	0,46,32	0,0,0	13,4

Llama-3.1-8B-Instruct / DGX Spark, wait-policy sweep

Weights	Budget	k	s	n_{tb}
YAQA 2	2.5%	0,3,1	0,1.5,1.5	2,4
	5%	0,3,1	0,1.5,1.5	4,4
	10%	4,4,3	2.5,2,2	2,4
	20%	10,10,8	1.5,1.5,1.5	3,4
	30%	40,0,9	2,0,1.5	9,4
	40%	16,0,11	1.5,0,1.5	14,4

Llama-3.1-8B-Instruct / RTX 5070 Ti

Weights	Budget	k	s	n_{tb}
AWQ 3	2.5%	3,10,2	0,0.5,0	3,4
	5%	18,0,18,18	1,0,0,0	4
	10%	37,37,37,37	1,0,0,0.5	6
	20%	45,40,39,40	0,0,0,0	7
AWQ 4	2.5%	28,28,34,36	0.5,1,0,0.5	6
	5%	52,53,49,52	1,0.5,0,0.5	24
	10%	45,45,45,45	0,0.5,0,0	24
	20%	53,55,52,52	0,0,0,0	24
LNQ+GQ 3	2.5%	6,8,3	0,1,0.5	4,4
	5%	63,63,33,61	1,0.5,0,0.5	16
	10%	29,28,20	0.5,0,0.5	16,16
	20%	60,61,58,60	1,0.5,0,0.5	16
LNQ+GQ 4	2.5%	45,42,47,45	1,0.5,0,0	16
	5%	48,45,48,49	1,0,0,0	16
	10%	68,59,58,59	1,0,0,0.5	16
	20%	64,64,64,63	1,0,0,0	24

Llama-3.1-8B-Instruct / RTX 5070 Ti

Weights	Budget	k	s	n_{tb}
SqLLM 3	2.5%	10,10,10,10	0.5,0.5,0,0	4
	5%	32,37,31,30	1,0.5,0,0	16
	10%	27,27,26,27	0,0,0,0	16
	20%	43,43,43,43	1,0,0,0	16
SqLLM 4	2.5%	85,46,46,45	1,0.5,0,0	16
	5%	48,48,48,48	1,0,0,0	16
	10%	54,53,53,54	1,0,0,0	16
	20%	59,60,57,58	0,0,0,0	24
YAQA 2	2.5%	15,23,12	1,1,0.5	6,4
	5%	28,28,16	1,1,0.5	6,4
	10%	37,39,25	1,0,0.5	7,16
	20%	97,49,35	0.5,0,0.5	14,32
YAQA 3	2.5%	32,32,31,32	0.5,0.5,0,0	5
	5%	66,50,44,44	0,0.5,0,0.5	7
	10%	62,48,47,48	0,0,0,0	14
	20%	61,61,60,61	0,0,0,0	16
YAQA 4	2.5%	59,60,58,58	1,0,0,0.5	7
	5%	82,60,58,60	0,1,0,0.5	7
	10%	73,77,70,70	1,0,0,0.5	23
	20%	87,87,86,87	0,0,0,0.5	28

Llama-3.1-8B-Instruct / RTX 5090

Weights	Budget	k	s	n_{tb}
AWQ 3	2.5%	8,0,0,8	0.5,0,0,0	8
	5%	17,0,17,13	1,0,0,0	16
	10%	13,13,13,13	0,1,0,0	14
	20%	29,28,22,27	1,0,0,0	14
AWQ 4	2.5%	10,0,10,10	0.5,0,0,0	16
	5%	20,0,18,13	0.5,0,0,0	16
	10%	30,39,26,28	0.5,0.5,0,0.5	24
	20%	53,54,46,46	0.5,1,1,0.5	24
LNQ+GQ 3	2.5%	18,0,13,9	0.5,0,0,0	16
	5%	17,0,13,13	0.5,0,0,0	16
	10%	59,65,5	1,1,0.5	24,4
	20%	29,29,16	1,0,0.5	24,16
LNQ+GQ 4	2.5%	25,0,15,11	0.5,0,0,0	11
	5%	21,0,21,18	0.5,0,0,0	14
	10%	23,25,21,22	1,1,0,0	16
	20%	40,32,31,32	1,0,0,0	16
SqLLM 3	2.5%	17,0,9,9	0.5,0,0,0	16
	5%	17,0,13,13	0.5,0,0,0	16
	10%	10,10,10,10	0.5,0,0,0	16
	20%	26,26,20,22	1,0,0,0	16
SqLLM 4	2.5%	16,0,16,9	1,0,0,0	11
	5%	22,0,21,21	0.5,0,0,0	24
	10%	19,19,19,19	0.5,0,0,0	24
	20%	29,29,28,29	0,0,0,0	38
YAQA 2	2.5%	16,28,12	0,0,0.5	23,8
	5%	39,40,18	0,0,0.5	23,16
	10%	39,39,20	0,0,0.5	19,32
	20%	50,51,24	0.5,0,0.5	38,16
YAQA 3	2.5%	29,28,26,26	1,0,0,0	56
	5%	45,46,20	0,0,0.5	56,16
	10%	51,51,24	0.5,0,0.5	38,32
	20%	53,53,49,49	0,0,0,0.5	38
YAQA 4	2.5%	42,42,34,34	0,0.5,0,0	56
	5%	38,39,38,38	0,0,0,0.5	28
	10%	50,50,50,50	0.5,0.5,0.5,0.5	38
	20%	58,60,57,57	0,0,0,0.5	38

Qwen3-8B / DGX Spark

Weights	Budget	k	s	n_{tb}
AWQ 3	2.5%	0,8,2,2	0,0,0,0	3
	5%	0,16,16,16	0,0,0,0	5
	10%	25,21,20,21	0,0,0,0	3
	20%	28,27,28,27	0,0,0,0	5
AWQ 4	2.5%	10,10,0,10	0,0,0,0	10
	5%	11,11,0,11	0,0,0,0	11
	10%	20,20,20,19	0,0,0,0	11
	20%	37,37,37,37	0,0,0,0	7
LNQ+GQ 3	2.5%	4,0,3	0,0,0	7,4
	5%	13,10,7	0,0,0	3,4
	10%	0,24,18	0,0,0	3,8
	20%	34,34,26	0,0,0	5,8
LNQ+GQ 4	2.5%	0,13,10	0,0,0	5,4
	5%	0,19,12	0,0,0	5,8
	10%	31,31,23	0,0,0	5,8
	20%	44,44,32	0,0,0	7,8
SqLLM 3	2.5%	5,4,4,4	0,0,0,0	5
	5%	8,8,8,8	0,0,0,0	3
	10%	22,23,22,22	0,0,0,0	4
	20%	31,30,30,32	0,0,0,0	9
SqLLM 4	2.5%	0,9,8,9	0,0,0,0	7
	5%	11,12,11,11	0,0,0,0	5
	10%	24,24,24,24	0,0,0,0	6
	20%	38,38,38,38	0,0,0,0	6
YAQA 2	2.5%	5,7,4	0,0,0	2,4
	5%	11,13,8	0,0,0	5,4
	10%	27,29,19	0,0,0	2,8
	20%	32,33,24	0,0,0	3,8
YAQA 3	2.5%	0,9,7	0,0,0	5,4
	5%	0,17,12	0,0,0	6,4
	10%	32,34,22	0,0,0	4,8
	20%	38,40,27	0,0,0	2,8
YAQA 4	2.5%	0,21,15	0,0,0	7,4
	5%	14,14,11	0,0,0	3,4
	10%	30,29,21	0,0,0	3,8
	20%	42,42,32	0,0,0	2,8

Qwen3-8B / RTX 5070 Ti

Weights	Budget	k	s	n_{tb}
AWQ 3	2.5%	0,6,6,6	0,0.5,0,0.5	3
	5%	22,21,14	1,1,0.5	4,4
	10%	44,44,44,44	1,1,0,0.5	6
	20%	54,54,53,55	1,0,0,0.5	6
AWQ 4	2.5%	19,19,19,19	0.5,1,0,0	6
	5%	35,41,35,41	0.5,1,0,0.5	8
	10%	61,61,60,61	1,0.5,0,0.5	24
	20%	63,63,62,62	0,0,1,0	24
LNQ+GQ 3	2.5%	0,16,14	0,1,0.5	5,16
	5%	15,17,7	0.5,1,0	5,32
	10%	32,39,22	1,0.5,0.5	16,16
	20%	44,44,32	0.5,0,0.5	16,32
LNQ+GQ 4	2.5%	27,28,21	0.5,1,0.5	11,16
	5%	46,46,24	0.5,0,0.5	16,16
	10%	44,45,29	0.5,0,0.5	16,16
	20%	54,57,40	0,0,0.5	16,32
SqLLM 3	2.5%	11,11,8	0.5,1,0.5	5,16
	5%	36,40,24,37	0.5,0.5,0,0.5	5
	10%	42,42,42,42	1,0.5,1,0	16
	20%	51,51,50,51	1,0,1,0	16
SqLLM 4	2.5%	38,0,37,37	1,0,0,0	16
	5%	45,42,45,47	1,0.5,0,0	16
	10%	44,45,44,45	0,0,0,0	16
	20%	62,62,61,61	1,0,0,0	16

Qwen3-8B / RTX 5070 Ti

Weights	Budget	k	s	n_{tb}
YAQA 2	2.5%	21,28,14	1,1,0.5	6,4
	5%	27,28,18	1,0,0.5	6,4
	10%	67,74,29	0.5,0.5,0.5	6,32
	20%	52,52,36	0,0,0.5	6,32
YAQA 3	2.5%	26,31,17	0.5,1,0.5	5,4
	5%	32,37,24	0.5,1,0.5	6,16
	10%	72,50,34	1,0,0.5	6,32
	20%	56,57,42	0,0,0.5	16,32
YAQA 4	2.5%	38,39,22	1,0,0.5	7,4
	5%	49,52,35	1,1,0.5	6,32
	10%	57,59,40	0.5,0,0.5	6,32
	20%	67,67,48	0,0,0.5	16,32

Qwen3-8B / RTX 5090

Weights	Budget	k	s	n_{tb}
AWQ 3	2.5%	2,0,2	0,0,0.5	12,32
	5%	17,0,16,10	1,0,0,0	16
	10%	60,60,32,32	1,0.5,1,0.5	12
	20%	48,48,44,48	1,0.5,1,0.5	16
AWQ 4	2.5%	14,0,13,5	0.5,0,0,0	16
	5%	18,0,16,10	0.5,0,0,0	16
	10%	29,29,26,29	0.5,0.5,0,0.5	32
	20%	51,112,50,51	1,0.5,1,0.5	32
LNQ+GQ 3	2.5%	0,16,12,5	0,0.5,0,0	32
	5%	0,11,11,11	0,0.5,0,0	32
	10%	13,13,10	0.5,0.5,0.5	32,16
	20%	27,31,20	0.5,0,0.5	32,16
LNQ+GQ 4	2.5%	16,0,14,14	1,0,1,0	48
	5%	31,0,26,25	1,0,1,0	32
	10%	54,54,9	1,1,0.5	48,4
	20%	31,31,23	0,0,0.5	16,16
SqLLM 3	2.5%	10,0,9,10	1,0,1,0	24
	5%	11,0,11,11	1,0,0,0	32
	10%	25,27,25,26	1,0.5,1,0.5	24
	20%	28,28,27,28	1,0,1,0	32
SqLLM 4	2.5%	16,0,16,16	0.5,0,0,0	48
	5%	34,0,23,21	1,0,0,0	32
	10%	25,25,25,25	1,0.5,0,0.5	48
	20%	37,37,37,37	1,0,1,0	16
YAQA 2	2.5%	11,10,8	0,0,0.5	7,16
	5%	22,20,16	0,0,0.5	7,16
	10%	40,41,22	0,0,0.5	20,32
	20%	47,46,26	0,0,0.5	32,32
YAQA 3	2.5%	17,17,13	0,0,0.5	32,8
	5%	36,36,23	0,0,0.5	24,32
	10%	57,57,25	0.5,0,0.5	32,32
	20%	44,44,33	0.5,0,0.5	32,32
YAQA 4	2.5%	19,18,14	0,0,0.5	5,4
	5%	42,42,26	0,0,0.5	8,16
	10%	44,45,28	0,0,0.5	32,32
	20%	47,48,36	0.5,0,0.5	32,16

Llama-3.3-70B-Instruct / RTX PRO 6000 / 5% slowdown budget

Weights	k	s	n_{tb}
AWQ 3	25,25,25,25	0,0.5,0,0	8
AWQ 4	30,30,31,29	1,0,0,0	40
LNQ+GQ 3	13,13,10	0.5,0.5,0.5	17,32
LNQ+GQ 4	24,24,24,24	0,0.5,0,0	28
SqLLM 3	16,16,15,15	0,0,0,0	16
SqLLM 4	24,24,24,24	1,0,0,0	28
YAQA 2	32,32,24	0.5,0.5,1	38,8
YAQA 3	42,37,33,34	0,0,0,0.5	32
YAQA 4	42,45,38,42	0,0,0,0.5	32

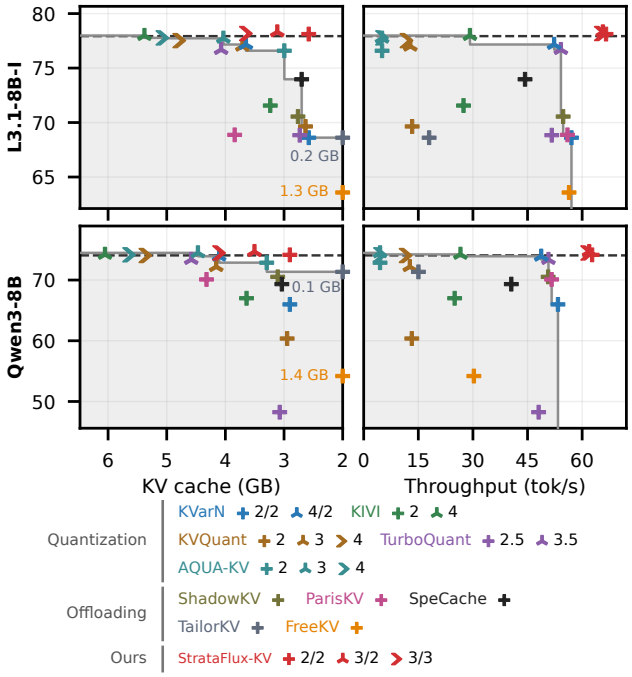

Figure A3. KV-only refinement: RULER accuracy at 128K (13 tasks, 96 samples each) versus GPU-resident KV-cache size and RTX 5090 throughput at the same context. StrataFlux-KV uses its dedicated-GPU configuration. Gray shading extends below the baseline frontiers; dashed lines mark BF16 accuracy. TailorKV's size is labeled at the KV-cache axis's right edge.

Qwen3-32B / RTX 5090 / 5% slowdown budget			
Weights	k	s	n_{tb}
AWQ 3	18,20,18,18	0,0.5,0,0	11
AWQ 4	19,19,19,17	0,0.5,0,0	12
LNQ+GQ 3	16,18,17,18	0,0.5,0,0	14
LNQ+GQ 4	28,26,28,26	1,0,0,0	16
SqLLM 3	16,18,16,16	0,0.5,0,0	14
SqLLM 4	27,27,26,27	1,0,0,0	16
YAQA 2	41,41,15	0.5,0,0.5	50,20
YAQA 3	201,32,19	0.5,0,0.5	50,20
YAQA 4	27,28,20	0.5,0,0.5	25,20

C Extended KV Evaluation

We extend Section 4.2 across KV precisions and context lengths. In configuration labels, K and V give resident bits per key and value element; +4/+4 gives their additional refinement bits. The fetch budget is the fraction of eligible token groups selected. Main host-RAM configurations use two-token groups. Main SSD-backed configurations use four-token groups and skip refinement on RAM-cache misses.

C.1 Accuracy and Context Length

Figure A3 compares KV-cache size and decode throughput against RULER accuracy at 128K.

Table A14. KV-only refinement: GPU-resident KV-cache size and PG-19 PPL/KL at 128K, with RULER accuracy at 8K–128K. *GB counts only KV memory. Baseline sizes are for RTX 5090; StrataFlux-KV headings identify the platform. DGX Spark sizes include its RAM cache. KVQ and TQ denote KVQuant and TurboQuant.

Llama-3.1-8B-Instruct								
Method	**GB***	**PPL**	**KL**	**R8**	**R16**	**R32**	**R64**	**R128**
BF16	17.2	10.32	0.000	95.1	92.9	87.1	85.8	77.9
KVarN 2/2	2.6	10.98	0.067	94.4	90.4	85.5	82.5	68.6
KVQ 2-bit	2.6	11.71	0.125	93.4	88.8	84.8	82.3	69.6
TQ 2.5-bit	2.7	10.98	0.055	93.6	90.8	84.4	82.9	68.8
KIVI-2	3.2	10.90	0.059	93.8	91.7	85.6	83.2	71.6
KVarN 4/2	3.7	10.38	0.008	94.8	92.8	87.1	85.9	77.2
KVQ 3-bit	3.7	10.56	0.024	94.8	92.2	86.9	85.3	77.0
TQ 3.5-bit	4.1	10.43	0.010	94.9	92.6	87.1	85.3	76.7
KVQ 4-bit	4.8	10.38	0.007	95.1	93.0	87.1	85.7	77.5
KIVI-4	5.4	10.33	0.002	95.3	93.1	87.0	85.8	78.0
TailorKV-2	0.2	11.21	0.107	94.1	92.5	88.4	81.7	68.6
FreeKV	1.3	11.62	0.126	93.7	87.8	80.8	75.4	63.6
SpeCache	2.7	10.68	0.035	95.0	92.6	87.1	85.5	74.0
ShadowKV	2.8	11.24	0.086	90.5	88.4	84.4	81.0	70.6
ParisKV	3.8	13.07	0.282	93.1	90.2	85.6	78.5	68.9
StrataFlux-KV / DGX Spark								
2/2	2.7	10.47	0.011	95.2	92.7	87.5	85.1	75.2
3/2	3.2	10.34	0.002	95.0	93.2	87.1	85.8	77.6
3/3	3.8	10.33	0.001	95.2	93.2	87.0	85.8	77.5
StrataFlux-KV / RTX 5090								
2/2	2.6	10.40	0.003	95.1	92.9	87.1	86.0	78.1
3/2	3.1	10.33	0.001	95.1	93.1	87.0	85.8	78.3
3/3	3.7	10.33	0.001	95.0	93.0	86.9	85.8	78.2

Qwen3-8B								
Method	**GB***	**PPL**	**KL**	**R8**	**R16**	**R32**	**R64**	**R128**
BF16	19.3	12.89	0.000	93.8	93.8	92.1	80.6	74.1
KVarN 2/2	2.9	14.40	0.130	93.0	91.9	89.0	72.6	66.0
KVQ 2-bit	2.9	14.40	0.167	91.2	90.2	85.1	67.7	60.4
TQ 2.5-bit	3.1	36.02	1.115	90.2	86.8	81.3	58.8	48.2
KIVI-2	3.6	13.85	0.106	92.9	92.1	88.8	73.0	67.0
KVarN 4/2	4.1	13.04	0.014	93.8	93.6	91.9	79.7	73.9
KVQ 3-bit	4.2	13.21	0.041	93.6	93.1	91.9	79.1	72.2
TQ 3.5-bit	4.6	13.99	0.102	94.0	93.1	92.0	78.6	73.4
KVQ 4-bit	5.4	12.96	0.011	93.8	93.6	91.9	79.9	74.0
KIVI-4	6.1	12.90	0.004	94.0	93.5	92.1	80.8	74.2
TailorKV-2	0.1	13.66	0.098	93.7	93.1	90.5	78.1	71.4
FreeKV	1.4	14.68	0.184	91.1	86.4	82.2	68.7	54.2
SpeCache	3.0	13.96	0.106	93.8	93.6	91.5	77.0	69.3
ShadowKV	3.1	15.10	0.187	89.4	87.3	82.8	74.7	70.5
ParisKV	4.3	14.11	0.154	93.3	91.6	88.5	77.4	70.1
StrataFlux-KV / DGX Spark								
2/2	3.0	13.32	0.030	94.0	93.4	91.6	77.2	71.4
3/2	3.6	12.98	0.005	93.7	93.7	92.0	79.5	74.0
3/3	4.2	12.93	0.004	93.8	93.8	92.2	79.8	74.2
StrataFlux-KV / RTX 5090								
2/2	2.9	12.98	0.007	94.0	93.8	91.8	79.9	74.2
3/2	3.5	12.92	0.002	93.9	93.8	92.0	80.4	74.7
3/3	4.1	12.90	0.001	93.9	93.8	92.1	80.3	74.5

Tables A14 and A15 compare 128K accuracy, memory, and throughput; Tables A16 and A17 report task accuracy. Tables A18 and A19 give context-length sweeps.

Table A15. KV-only refinement: PG-19 PPL/KL and RTX 5090 throughput at 128K. StrataFlux-KV fetches +4/+4 bits from host RAM for 1/128 of two-token groups.

	L3.1-8B-I			Qwen3-8B		
KV configuration	**PPL**	**KL**	**tok/s**	**PPL**	**KL**	**tok/s**
BF16 KV	10.32	0.000	47.6	12.89	0.000	OOM
KVarN 2/2	10.98	0.067	57.1	14.40	0.130	53.4
KIVI-2	10.90	0.059	27.4	13.85	0.106	25.0
KVQuant 2-bit	11.71	0.125	13.3	14.40	0.167	13.2
TurboQuant 2.5-bit	10.98	0.055	51.7	36.02	1.115	48.1
AQUA-KV 2-bit	10.66	0.030	5.1	13.14	0.029	4.5
KVQuant 3-bit	10.56	0.024	12.8	13.21	0.041	12.7
TurboQuant 3.5-bit	10.43	0.010	54.2	13.99	0.102	50.7
AQUA-KV 3-bit	10.39	0.006	5.0	12.98	0.017	4.5
KVarN 4/2	10.38	0.008	52.3	13.04	0.014	48.8
KIVI-4	10.33	0.002	29.2	12.90	0.004	26.6
KVQuant 4-bit	10.38	0.007	11.8	12.96	0.011	11.5
AQUA-KV 4-bit	10.32	0.002	5.0	12.95	0.016	4.5
ShadowKV	11.24	0.086	54.8	15.10	0.187	50.7
ParisKV	13.07	0.282	56.0	14.11	0.154	51.6
TailorKV-2	11.21	0.107	18.0	13.66	0.098	15.0
FreeKV	11.62	0.126	56.4	14.68	0.184	30.3
SpeCache	10.68	0.035	44.3	13.96	0.106	40.5
StrataFlux-KV 2/2	10.40	0.003	66.6	12.98	0.007	62.7
StrataFlux-KV 3/2	10.33	0.001	65.7	12.92	0.002	61.9
StrataFlux-KV 3/3	10.33	0.001	65.1	12.90	0.001	61.4

Table A16. KV-only refinement: quantized KV task accuracy. R4–R128 give RULER accuracy at 4K–128K; LB32 uses LongBench prompts truncated to 32,128 tokens; M500/AIME report MATH500/AIME accuracy. Platform headings distinguish StrataFlux-KV configurations; baseline accuracies are shared across devices.

Llama-3.1-8B-Instruct

Configuration	**R4**	**R8**	**R16**	**R32**	**R64**	**R128**	**LB32**
BF16 KV	95.7	95.1	92.9	87.1	85.8	77.9	49.7
KVarN 2/2	94.4	94.4	90.4	85.5	82.5	68.6	49.3
KIVI-2	94.7	93.8	91.7	85.6	83.2	71.6	49.2
KVQuant 2-bit	94.8	93.4	88.8	84.8	82.3	69.6	48.8
TurboQuant 2.5-bit	94.5	93.6	90.8	84.4	82.9	68.8	49.0
AQUA-KV 2-bit	95.4	94.8	92.8	86.5	85.7	76.6	49.3
KVQuant 3-bit	95.6	94.8	92.2	86.9	85.3	77.0	49.5
TurboQuant 3.5-bit	95.7	94.9	92.6	87.1	85.3	76.7	49.5
AQUA-KV 3-bit	95.6	95.0	92.8	87.0	85.8	77.7	49.6
KVarN 4/2	95.5	94.8	92.8	87.1	85.9	77.2	49.6
KIVI-4	95.6	95.3	93.1	87.0	85.8	78.0	49.6
KVQuant 4-bit	95.7	95.1	93.0	87.1	85.7	77.5	49.6
AQUA-KV 4-bit	95.7	95.1	92.9	87.0	85.8	77.8	49.5
StrataFlux-KV / DGX Spark							
2/2	95.8	95.2	92.7	87.5	85.1	75.2	49.6
3/2	95.7	95.0	93.2	87.1	85.8	77.6	49.5
3/3	95.8	95.2	93.2	87.0	85.8	77.5	49.6
StrataFlux-KV / RTX 5090							
2/2	95.9	95.1	92.9	87.1	86.0	78.1	49.6
3/2	95.7	95.1	93.1	87.0	85.8	78.3	49.5
3/3	95.7	95.0	93.0	86.9	85.8	78.2	49.5

Qwen3-8B

Configuration	**R4**	**R8**	**R16**	**R32**	**R64**	**R128**	**LB32**	**M500**	**AIME**
BF16 KV	95.2	93.8	93.8	92.1	80.6	74.1	49.4	97.2	71.0
KVarN 2/2	94.2	93.0	91.9	89.0	72.6	66.0	48.6	96.2	66.0
KIVI-2	94.0	92.9	92.1	88.8	73.0	67.0	48.8	96.0	64.8
KVQuant 2-bit	93.5	91.2	90.2	85.1	67.7	60.4	48.0	94.4	54.0
TurboQuant 2.5-bit	92.3	90.2	86.8	81.3	58.8	48.2	48.4	79.4	30.0
AQUA-KV 2-bit	94.9	93.7	93.4	91.9	79.8	72.9	49.2	96.0	70.4
KVQuant 3-bit	94.9	93.6	93.1	91.9	79.1	72.2	49.2	96.2	70.8
TurboQuant 3.5-bit	95.1	94.0	93.1	92.0	78.6	73.4	49.4	97.0	69.8
AQUA-KV 3-bit	95.3	94.1	93.6	92.2	80.5	74.5	49.3	97.2	71.9
KVarN 4/2	95.4	93.8	93.6	91.9	79.7	73.9	49.3	96.2	71.0
KIVI-4	95.0	94.0	93.5	92.1	80.8	74.2	49.4	96.0	71.3
KVQuant 4-bit	95.3	93.8	93.6	91.9	79.9	74.0	49.2	96.6	72.1
AQUA-KV 4-bit	95.2	93.9	93.6	92.1	80.7	74.2	49.3	95.6	70.2
StrataFlux-KV / DGX Spark									
2/2	95.4	94.0	93.4	91.6	77.2	71.4	49.0	96.0	68.3
3/2	95.2	93.7	93.7	92.0	79.5	74.0	49.4	95.0	70.2
3/3	95.3	93.8	93.8	92.2	79.8	74.2	49.3	96.2	69.4
StrataFlux-KV / RTX 5090									
2/2	95.5	94.0	93.8	91.8	79.9	74.2	49.3	95.6	69.2
3/2	95.3	93.9	93.8	92.0	80.4	74.7	49.3	96.2	70.0
3/3	95.3	93.9	93.8	92.1	80.3	74.5	49.3	97.0	71.5

Table A17. KV-only refinement: selective KV task accuracy. R8–R128 give RULER accuracy at 8K–128K; LB32 uses LongBench prompts truncated to 32,128 tokens. M500/AIME report Qwen3-8B MATH500/AIME accuracy. ShadowKV and TailorKV do not compress KV produced during decoding and are omitted from the M500/AIME comparison. Platform headings distinguish StrataFlux-KV configurations; baseline accuracies are shared across devices.

Llama-3.1-8B-Instruct

Configuration	**R8**	**R16**	**R32**	**R64**	**R128**	**LB32**
BF16 KV	95.1	92.9	87.1	85.8	77.9	49.7
ShadowKV	90.5	88.4	84.4	81.0	70.6	49.3
ParisKV	93.1	90.2	85.6	78.5	68.9	49.6
TailorKV-2	94.1	92.5	88.4	81.7	68.6	48.8
TailorKV	93.6	91.8	87.3	77.8	65.2	48.6
FreeKV	93.7	87.8	80.8	75.4	63.6	49.6
SpeCache	95.0	92.6	87.1	85.5	74.0	49.5
KVSwap	83.9	78.4	73.3	69.0	49.0	48.5
StrataFlux-KV / DGX Spark						
2/2	95.2	92.7	87.5	85.1	75.2	49.6
3/2	95.0	93.2	87.1	85.8	77.6	49.5
3/3	95.2	93.2	87.0	85.8	77.5	49.6
StrataFlux-KV / RTX 5090						
2/2	95.1	92.9	87.1	86.0	78.1	49.6
3/2	95.1	93.1	87.0	85.8	78.3	49.5
3/3	95.0	93.0	86.9	85.8	78.2	49.5

Qwen3-8B

Configuration	**R8**	**R16**	**R32**	**R64**	**R128**	**LB32**	**M500**	**AIME**
BF16 KV	93.8	93.8	92.1	80.6	74.1	49.4	97.2	71.0
ShadowKV	89.4	87.3	82.8	74.7	70.5	49.3	–	–
ParisKV	93.3	91.6	88.5	77.4	70.1	49.3	95.4	64.2
TailorKV-2	93.7	93.1	90.5	78.1	71.4	49.1	–	–
FreeKV	91.1	86.4	82.2	68.7	54.2	49.1	96.8	63.5
SpeCache	93.8	93.6	91.5	77.0	69.3	49.2	96.6	70.2
KVSwap	88.6	84.0	78.3	68.6	59.6	49.1	94.0	52.7
StrataFlux-KV / DGX Spark								
2/2	94.0	93.4	91.6	77.2	71.4	49.0	96.0	68.3
3/2	93.7	93.7	92.0	79.5	74.0	49.4	95.0	70.2
3/3	93.8	93.8	92.2	79.8	74.2	49.3	96.2	69.4

Qwen3-8B								
Configuration	**R8**	**R16**	**R32**	**R64**	**R128**	**LB32**	**M500**	**AIME**
StrataFlux-KV / RTX 5090								
2/2	94.0	93.8	91.8	79.9	74.2	49.3	95.6	69.2
3/2	93.9	93.8	92.0	80.4	74.7	49.3	96.2	70.0
3/3	93.9	93.8	92.1	80.3	74.5	49.3	97.0	71.5

Table A18. KV-only refinement: PG-19 context sweep on Llama-3.1-8B-Instruct. *GB counts only the resident KV representation and, on DGX Spark, its refinement RAM cache; weights and backing-tier storage are excluded. Throughput is tok/s at the column's context length.

Configuration	Metric	8K	16K	32K	64K	128K
BF16 KV	PPL	10.34	10.29	10.27	10.17	10.32
	KL	0.000	0.000	0.000	0.000	0.000
	GB*	1.07	2.15	4.30	8.59	17.18
	Spark tok/s	14.9	14.1	12.5	9.9	7.3
	5090 tok/s	96.1	90.1	80.2	65.0	47.6
Quantization						
KVarN 2/2	PPL	10.74	10.73	10.74	10.64	10.98
	KL	0.045	0.049	0.051	0.052	0.067
	Spark tok/s	15.3	14.7	13.9	13.2	11.3
	5090 tok/s	92.9	88.0	79.7	72.1	57.1
KIVI-2	PPL	10.82	10.84	10.81	10.69	10.90
	KL	0.055	0.058	0.059	0.059	0.059
	Spark tok/s	14.7	14.0	12.9	11.3	8.8
	5090 tok/s	82.1	72.5	58.9	42.7	27.4
KVQuant 2-bit	PPL	11.15	11.11	11.12	11.04	11.71
	KL	0.084	0.086	0.087	0.089	0.125
	Spark tok/s	12.9	11.6	9.5	7.0	3.9
	5090 tok/s	60.2	48.8	36.2	22.6	13.3
TurboQuant 2.5-bit	PPL	10.77	10.76	10.78	10.67	10.98
	KL	0.040	0.042	0.045	0.045	0.055
	Spark tok/s	15.0	14.5	13.6	12.4	10.2
	5090 tok/s	88.9	85.1	77.2	65.9	51.7
AQUA-KV 2-bit	PPL	10.54	10.51	10.54	10.48	10.66
	KL	0.022	0.021	0.024	0.030	0.030
	Spark tok/s	9.8	6.9	4.5	2.7	1.5
	5090 tok/s	44.4	30.6	17.9	9.1	5.1
KVQuant 3-bit	PPL	10.52	10.47	10.46	10.35	10.56
	KL	0.020	0.021	0.021	0.021	0.024
	Spark tok/s	12.9	11.4	9.1	6.4	3.9
	5090 tok/s	59.5	48.1	35.0	21.9	12.8
TurboQuant 3.5-bit	PPL	10.42	10.37	10.35	10.26	10.43
	KL	0.007	0.007	0.008	0.008	0.010
	Spark tok/s	15.0	14.5	13.7	12.5	10.4
	5090 tok/s	89.1	85.7	78.7	67.9	54.2
AQUA-KV 3-bit	PPL	10.38	10.35	10.34	10.24	10.39
	KL	0.005	0.006	0.006	0.007	0.006
	Spark tok/s	9.7	6.8	4.5	2.6	1.5
	5090 tok/s	44.4	30.5	17.9	9.1	5.0
KVarN 4/2	PPL	10.38	10.34	10.32	10.22	10.38
	KL	0.007	0.007	0.007	0.007	0.008
	Spark tok/s	15.1	14.5	13.5	12.6	10.4
	5090 tok/s	91.8	86.2	76.9	68.3	52.3
KIVI-4	PPL	10.35	10.30	10.28	10.19	10.33
	KL	0.002	0.002	0.002	0.002	0.002
	Spark tok/s	14.7	14.0	12.9	11.1	8.6
	5090 tok/s	83.4	74.4	61.0	44.8	29.2
KVQuant 4-bit	PPL	10.38	10.34	10.32	10.22	10.38
	KL	0.006	0.006	0.006	0.006	0.007
	Spark tok/s	12.5	10.8	8.4	5.8	3.3
	5090 tok/s	57.3	46.0	33.0	20.3	11.8
AQUA-KV 4-bit	PPL	10.35	10.31	10.29	10.19	10.32
	KL	0.002	0.002	0.002	0.002	0.002
	Spark tok/s	9.7	6.8	4.4	2.6	1.5
	5090 tok/s	44.2	30.4	17.8	9.1	5.0
KV offloading / DGX Spark						
KVSwap	PPL	10.87	10.99	11.10	11.09	16.27
	KL	0.056	0.072	0.087	0.097	0.477
	GB*	0.12	0.19	0.32	0.59	1.13
	tok/s	8.7	8.2	9.2	8.8	8.7
KV offloading / RTX 5090						
ShadowKV	PPL	10.69	10.66	10.65	10.60	11.24
	KL	0.042	0.044	0.045	0.048	0.086
	GB*	0.23	0.40	0.74	1.41	2.76
	tok/s	88.3	86.9	84.6	67.9	54.8
ParisKV	PPL	10.59	10.61	10.62	10.60	13.07
	KL	0.045	0.060	0.072	0.085	0.282
	GB*	0.32	0.55	1.02	1.96	3.84
	tok/s	77.6	75.5	72.9	67.2	56.0
TailorKV-2	PPL	10.38	10.36	10.35	10.25	11.21
	KL	0.018	0.027	0.035	0.040	0.107
	GB*	0.03	0.04	0.05	0.09	0.16
	tok/s	36.3	32.3	34.9	23.2	18.0
FreeKV	PPL	10.45	10.50	10.58	10.61	11.62
	KL	0.015	0.026	0.038	0.052	0.126
	GB*	0.33	0.39	0.53	0.79	1.32
	tok/s	58.0	56.3	50.8	49.6	56.4
SpeCache	PPL	10.45	10.42	10.42	10.34	10.68
	KL	0.013	0.016	0.018	0.021	0.035
	GB*	0.18	0.35	0.69	1.36	2.70
	tok/s	78.8	76.4	67.7	58.1	44.3
TailorKV-1	PPL	11.46	11.59	11.58	11.60	13.71
	KL	0.128	0.149	0.157	0.173	0.322
StrataFlux-KV / DGX Spark						
StrataFlux-KV 2/2	PPL	10.42	10.39	10.37	10.27	10.47
	KL	0.007	0.006	0.006	0.006	0.011
	GB*	0.32	0.48	0.80	1.44	2.71
	tok/s	14.9	14.7	14.2	13.4	11.4
StrataFlux-KV 3/2	PPL	10.35	10.31	10.29	10.19	10.34
	KL	0.002	0.002	0.002	0.002	0.002
	GB*	0.36	0.55	0.93	1.71	3.25
	tok/s	14.8	14.6	14.1	13.1	11.0
StrataFlux-KV 3/3	PPL	10.35	10.30	10.28	10.18	10.33
	KL	0.001	0.001	0.001	0.001	0.001
	GB*	0.39	0.61	1.07	1.97	3.79
	tok/s	14.8	14.5	14.0	12.8	10.7
StrataFlux-KV / RTX 5090						
StrataFlux-KV 2/2	PPL	10.41	10.37	10.35	10.25	10.40
	KL	0.004	0.004	0.003	0.002	0.003
	GB*	0.19	0.35	0.67	1.30	2.58
	tok/s	88.0	85.4	81.8	75.7	66.6
StrataFlux-KV 3/2	PPL	10.35	10.30	10.29	10.18	10.33
	KL	0.001	0.001	0.001	0.001	0.001
	GB*	0.22	0.41	0.80	1.57	3.12
	tok/s	87.5	84.9	81.2	75.1	65.7
StrataFlux-KV 3/3	PPL	10.35	10.30	10.28	10.18	10.33
	KL	0.001	0.001	0.001	0.001	0.001
	GB*	0.25	0.48	0.93	1.84	3.65
	tok/s	87.5	84.6	80.8	74.4	65.1

Table A19. KV-only refinement: PG-19 context sweep on Qwen3-8B. *GB counts only the resident KV representation and, on DGX Spark, its refinement RAM cache; weights and backing-tier storage are excluded. Throughput is tok/s at the column's context length.

Configuration	Metric	8K	16K	32K	64K	128K
BF16 KV	PPL	13.23	12.91	12.64	12.99	12.89
	KL	0.000	0.000	0.000	0.000	0.000
	GB*	1.21	2.42	4.83	9.66	19.33
	Spark tok/s	14.7	13.7	12.1	9.4	6.7
	5090 tok/s	93.8	87.5	77.0	61.4	OOM
Quantization						
KVarN 2/2	PPL	13.77	13.47	13.11	14.44	14.40
	KL	0.062	0.065	0.059	0.131	0.130
	Spark tok/s	14.9	14.4	13.6	12.7	10.7
	5090 tok/s	89.9	84.9	76.4	68.5	53.4
KIVI-2	PPL	13.56	13.25	13.01	13.88	13.85
	KL	0.048	0.048	0.051	0.090	0.106
	Spark tok/s	14.4	13.7	12.5	10.7	8.3
	5090 tok/s	78.6	69.1	55.3	39.4	25.0
KVQuant 2-bit	PPL	14.48	14.18	13.97	14.39	14.40
	KL	0.117	0.124	0.129	0.158	0.167
	Spark tok/s	12.8	11.1	8.8	6.3	4.1
	5090 tok/s	57.6	47.3	34.7	21.7	13.2
TurboQuant 2.5-bit	PPL	20.66	21.20	21.82	33.85	36.02
	KL	0.582	0.631	0.685	1.047	1.115
	Spark tok/s	14.7	14.2	13.2	11.9	9.6
	5090 tok/s	86.0	81.9	73.6	62.2	48.1
AQUA-KV 2-bit	PPL	13.40	13.11	12.88	13.21	13.14
	KL	0.012	0.013	0.013	0.028	0.029
	Spark tok/s	9.3	6.5	4.1	2.4	1.4
	5090 tok/s	41.1	27.9	16.2	8.1	4.5
KVQuant 3-bit	PPL	13.47	13.20	12.90	13.29	13.21
	KL	0.026	0.027	0.028	0.039	0.041
	Spark tok/s	12.6	10.8	8.5	6.4	3.9
	5090 tok/s	56.6	45.8	33.5	21.3	12.7
TurboQuant 3.5-bit	PPL	13.80	13.56	13.12	14.32	13.99
	KL	0.057	0.067	0.051	0.118	0.102
	Spark tok/s	14.7	14.2	13.3	12.0	9.9
	5090 tok/s	86.3	82.4	75.4	64.3	50.7
AQUA-KV 3-bit	PPL	13.27	12.96	12.69	13.08	12.98
	KL	0.005	0.006	0.006	0.017	0.017
	Spark tok/s	9.3	6.4	4.1	2.4	1.4
	5090 tok/s	41.0	27.8	16.1	8.1	4.5
KVarN 4/2	PPL	13.32	13.00	12.72	13.15	13.04
	KL	0.008	0.008	0.009	0.014	0.014
	Spark tok/s	14.8	14.2	13.2	12.0	9.8
	5090 tok/s	88.8	82.9	73.3	64.7	48.8
KIVI-4	PPL	13.23	12.92	12.64	13.01	12.90
	KL	0.002	0.002	0.002	0.004	0.004
	Spark tok/s	14.4	13.7	12.5	10.6	8.0
	5090 tok/s	80.1	70.8	57.5	41.4	26.6
KVQuant 4-bit	PPL	13.30	12.97	12.69	13.04	12.96
	KL	0.007	0.007	0.007	0.011	0.011
	Spark tok/s	12.3	10.3	8.2	5.6	3.4
	5090 tok/s	54.5	43.5	31.4	19.3	11.5
AQUA-KV 4-bit	PPL	13.24	12.93	12.65	13.05	12.95
	KL	0.004	0.004	0.004	0.016	0.016
	Spark tok/s	9.3	6.4	4.0	2.4	1.3
	5090 tok/s	41.0	27.8	16.1	8.1	4.5
KV offloading / DGX Spark						
KVSwap	PPL	14.09	14.10	14.21	15.64	15.82
	KL	0.074	0.101	0.133	0.259	0.289
	GB*	0.14	0.21	0.36	0.66	1.27
	tok/s	7.4	7.7	7.7	7.6	7.5
KV offloading / RTX 5090						
ShadowKV	PPL	15.78	14.85	13.92	15.65	15.10
	KL	0.229	0.193	0.145	0.214	0.187
	GB*	0.26	0.45	0.83	1.59	3.11
	tok/s	84.9	83.3	81.1	62.5	50.7
ParisKV	PPL	13.34	13.10	12.87	14.09	14.11
	KL	0.050	0.064	0.070	0.140	0.154
	GB*	0.36	0.62	1.15	2.21	4.32
	tok/s	70.8	71.6	68.9	62.8	51.6
TailorKV-2	PPL	13.09	12.85	12.58	13.62	13.66
	KL	0.017	0.030	0.037	0.082	0.098
	GB*	0.03	0.03	0.04	0.05	0.07
	tok/s	24.5	25.7	21.8	16.7	15.0
FreeKV	PPL	13.45	13.34	13.30	14.37	14.68
	KL	0.021	0.040	0.062	0.146	0.184
	GB*	0.36	0.43	0.58	0.86	1.42
	tok/s	47.3	38.9	35.5	31.5	30.3
SpeCache	PPL	13.43	13.14	12.94	13.74	13.96
	KL	0.026	0.030	0.038	0.080	0.106
	GB*	0.21	0.39	0.77	1.53	3.04
	tok/s	73.4	70.9	63.0	53.1	40.5
StrataFlux-KV / DGX Spark						
StrataFlux-KV 2/2	PPL	13.39	13.09	12.82	13.34	13.32
	KL	0.011	0.010	0.009	0.028	0.030
	GB*	0.35	0.53	0.88	1.60	3.04
	tok/s	14.6	14.4	13.9	13.0	10.9
StrataFlux-KV 3/2	PPL	13.28	12.96	12.69	13.07	12.98
	KL	0.002	0.002	0.002	0.005	0.005
	GB*	0.38	0.60	1.03	1.90	3.64
	tok/s	14.6	14.3	13.7	12.6	10.5
StrataFlux-KV 3/3	PPL	13.25	12.93	12.66	13.03	12.93
	KL	0.001	0.001	0.001	0.003	0.004
	GB*	0.42	0.67	1.18	2.20	4.24
	tok/s	14.5	14.2	13.7	12.4	10.2
StrataFlux-KV / RTX 5090						
StrataFlux-KV 2/2	PPL	13.39	13.07	12.80	13.07	12.98
	KL	0.007	0.005	0.004	0.009	0.007
	GB*	0.21	0.39	0.75	1.47	2.90
	tok/s	84.7	82.2	78.5	72.3	62.7
StrataFlux-KV 3/2	PPL	13.28	12.96	12.68	13.02	12.92
	KL	0.002	0.001	0.001	0.002	0.002
	GB*	0.25	0.47	0.90	1.77	3.50
	tok/s	84.2	81.6	77.9	71.5	61.9
StrataFlux-KV 3/3	PPL	13.25	12.93	12.66	12.99	12.90
	KL	0.001	0.001	0.001	0.001	0.001
	GB*	0.29	0.54	1.05	2.07	4.11
	tok/s	84.1	81.3	77.4	70.9	61.4

C.2 Alternative Attention Architectures and MoE Models

Table A20. KV-only refinement: PG-19 perplexity/KL and RULER accuracy at 32K, plus MATH500 accuracy for Qwen3.5-9B. The comparison covers hybrid linear/full attention (Qwen3.5-9B) and MoE models with standard attention (Qwen3-30B-A3B) or latent attention (DeepSeek-V2-Lite). Refinement applies only to full attention. For DeepSeek, 2/2 and 3/3 indicate 2- and 3-bit latent precision; 3/2 is inapplicable because the shared latent representation uses one bitwidth.

		Qwen3.5-9B				Qwen3-30B-A3B			DeepSeek-V2-Lite		
K/V	**Method**	**PPL**	**KL**	**R32**	**M500**	**PPL**	**KL**	**R32**	**PPL**	**KL**	**R32**
	BF16	9.48	0.000	96.6	85.6	9.29	0.000	94.1	9.12	0.000	67.5
2/2	Base	9.71	0.027	96.3	83.0	10.06	0.086	89.9	9.59	0.138	59.9
	StrataFlux-KV	9.54	0.004	96.5	88.0	9.55	0.021	93.6	9.17	0.017	67.2
3/2	Base	9.62	0.014	96.5	84.6	9.42	0.016	93.9		–	
	StrataFlux-KV	9.51	0.003	96.5	86.6	9.35	0.004	93.9			
3/3	Base	9.50	0.004	96.5	86.6	9.34	0.009	93.9	9.16	0.026	66.5
	StrataFlux-KV	9.48	0.001	96.5	89.6	9.32	0.002	94.0	9.13	0.006	67.2

Table A20 groups the three architectures; Tables A21–A23 give per-model task results.

Table A21. KV-only refinement: accuracy on Qwen3.5-9B. R4–R32 are RULER context lengths in thousands of tokens. M500 is MATH500. KVQuant uses C4 calibration and retains 1% outliers. TurboQuant presets mix bitwidths across channels. StrataFlux-KV uses two-token groups, 1/128 selection, and no cache misses on dedicated GPUs; Base disables refinement.

Method	**Bits**	**R4**	**R8**	**R16**	**R32**	**M500**
BF16 KV	16/16	97.1	97.1	96.6	96.6	85.6
KIVI	2	97.2	97.0	96.7	96.4	81.6
	4	97.1	97.1	96.6	96.5	85.2
KVQuant	2	96.9	97.0	96.6	96.6	94.2
	3	97.1	97.0	96.7	96.7	86.2
	4	97.1	97.1	96.6	96.5	89.0
TurboQuant	2.5	97.1	96.9	96.5	96.4	88.8
	3.5	97.1	97.2	96.6	96.5	86.2
AQUA-KV	2	97.1	97.0	96.9	96.9	87.4
	3	97.1	97.1	96.7	96.7	83.8
	4	97.1	97.1	96.7	96.7	84.6
Base	2/2	97.1	97.0	96.5	96.3	83.0
StrataFlux-KV		97.1	97.2	96.6	96.5	88.0
Base	3/2	97.1	97.1	96.4	96.5	84.6
StrataFlux-KV		97.1	97.2	96.6	96.5	86.6
Base	3/3	97.1	97.2	96.6	96.5	86.6
StrataFlux-KV		97.1	97.1	96.6	96.5	89.6

Table A22. KV-only refinement: accuracy on DeepSeek-V2-Lite. R4–R32 are RULER context lengths in thousands of tokens. KVQuant uses C4 calibration and retains 1% outliers. TurboQuant presets mix bitwidths across channels. StrataFlux-KV uses two-token groups, 1/128 selection, and no cache misses on dedicated GPUs; Base disables refinement. StrataFlux-KV bitwidths give latent precision, rather than separate K/V widths.

Method	**Bits**	**R4**	**R8**	**R16**	**R32**
BF16 KV	16/16	91.1	85.2	77.8	67.5
KIVI	2	86.7	78.6	69.8	58.5
	4	90.4	85.3	77.7	67.3
KVQuant	2	86.4	75.4	65.4	53.2
	3	90.2	83.1	76.2	64.0
	4	90.9	84.4	78.1	66.9
TurboQuant	2.5	84.4	77.3	69.0	52.5
	3.5	90.9	84.3	77.0	66.0
AQUA-KV	2	90.7	82.5	74.8	63.8
	3	91.2	84.8	77.5	65.8
	4	90.9	85.1	78.4	66.7
Base	2	86.8	79.2	70.9	59.9
StrataFlux-KV		90.4	84.5	78.8	67.2
Base	3	90.6	84.1	77.5	66.5
StrataFlux-KV		90.7	85.2	78.2	67.2

Table A23. KV-only refinement: accuracy on Qwen3-30B-A3B. R4–R32 are RULER context lengths in thousands of tokens. KVQuant uses C4 calibration and retains 1% outliers. TurboQuant presets mix bitwidths across channels. StrataFlux-KV uses two-token groups, 1/128 selection, and no cache misses on dedicated GPUs; Base disables refinement.

Method	**Bits**	**R4**	**R8**	**R16**	**R32**
BF16 KV	16/16	95.6	95.8	95.1	94.1
KIVI	2	94.6	95.0	93.1	92.3
	4	95.5	95.9	94.8	94.1
KVQuant	2	93.9	93.9	91.9	90.7
	3	95.3	95.3	94.6	94.2
	4	95.5	95.6	94.8	94.2
TurboQuant	2.5	93.0	91.4	91.1	85.2
	3.5	95.3	95.7	94.7	93.6
AQUA-KV	2	95.6	95.7	94.8	94.1
	3	95.6	95.8	94.7	94.1
	4	95.7	95.8	94.7	94.2
Base	2/2	93.4	93.3	92.0	89.9
StrataFlux-KV		95.6	95.8	95.0	93.6
Base	3/2	95.5	95.6	94.9	93.9
StrataFlux-KV		95.7	95.9	94.8	93.9
Base	3/3	95.6	95.7	95.1	93.9
StrataFlux-KV		95.6	95.8	94.9	94.0

C.3 KV Representation and Selection Controls

Tables A24 and A25 vary the retained FP16 tokens and the attention-score aggregation used for token-group selection.

Table A24. KV representation and selection controls: Llama-3.1-8B-Instruct, PG-19 at 32K and 128K. Four-token groups, 1/128 fetch fraction. We retain the first 128 tokens in FP16 following KVarN (Muller et al., 2026) and select by the maximum score over query heads; controls change one setting.

GPU resident KV	Configuration	32K PPL	KL	128K PPL	KL
2/2	Reference	10.36	0.004	10.42	0.003
	First 128 tokens quantized	10.36	0.004	10.41	0.004
	Mean over query heads	10.36	0.004	10.41	0.003
3/2	Reference	10.29	0.001	10.33	0.001
	First 128 tokens quantized	10.29	0.001	10.33	0.001
	Mean over query heads	10.29	0.001	10.33	0.001
3/3	Reference	10.29	0.001	10.33	0.001
	First 128 tokens quantized	10.29	0.001	10.33	0.001
	Mean over query heads	10.28	0.001	10.33	0.001

Table A25. KV representation and selection controls: Qwen3-8B, PG-19 at 32K and 128K. Four-token groups, 1/128 fetch fraction. Reference shows our configuration, where we retain the first 128 tokens in FP16 following KVarN (Muller et al., 2026) and select token groups using the maximum score over query heads. Each control changes one of these settings.

GPU resident KV	Configuration	32K PPL	KL	128K PPL	KL
2/2	Reference	12.81	0.006	13.01	0.009
	First 128 tokens quantized	12.81	0.006	13.00	0.010
	Mean over query heads	12.80	0.006	13.01	0.009
3/2	Reference	12.69	0.001	12.93	0.002
	First 128 tokens quantized	12.69	0.002	12.93	0.002
	Mean over query heads	12.68	0.001	12.93	0.002
3/3	Reference	12.66	0.001	12.90	0.002
	First 128 tokens quantized	12.66	0.001	12.90	0.002
	Mean over query heads	12.66	0.001	12.90	0.002

C.4 Fetch Budgets, RAM Caches, and Miss Policies

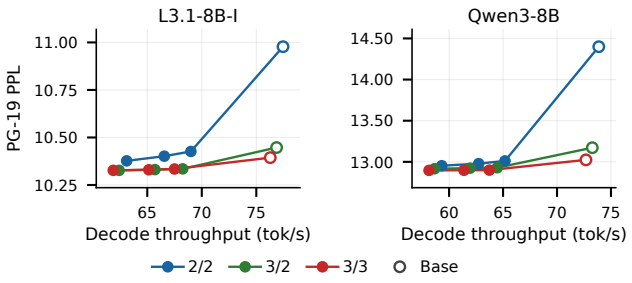

Figure A4. KV fetch-fraction sweep: PG-19 PPL versus RTX 5090 throughput at 128K. Curves identify resident K/V bitwidths; fractions 1/256, 1/128, and 1/64 run from highest to lowest throughput. StrataFlux-KV uses two-token groups and +4/+4 bits. Open markers show the matching bases without refinement.

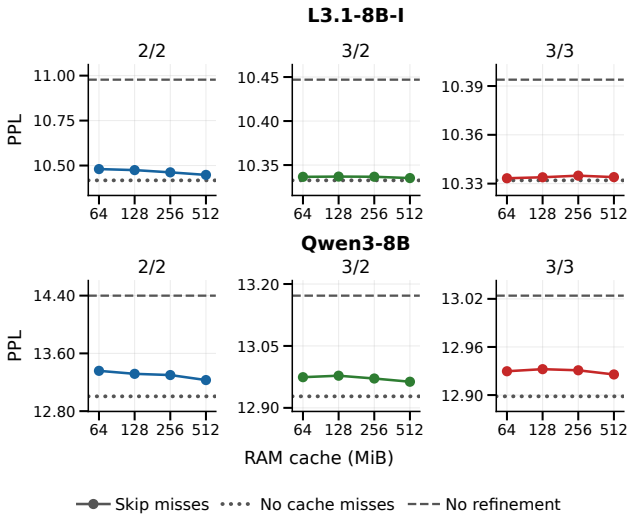

Figure A5. KV RAM-cache capacity by resident bitwidth on DGX Spark: PG-19 PPL at 128K with four-token groups, +4/+4 bits, 1/128 selection, and skipped misses. Dotted lines assume no cache misses; dashed lines show the matching no-refinement bases.

Figure A4 varies the host-RAM fetch fraction on RTX 5090; Figure A5 varies RAM-cache capacity on DGX Spark. Table A31 reports throughput when waiting on cache misses. Table A26 compares SSD-backed configurations; Tables A27–A30 report cache-capacity and no-cache-miss comparisons.

Table A26. KV-only refinement: PG-19 PPL/KL and DGX Spark throughput at 128K. StrataFlux-KV uses +4/+4 bits, 1/128 selection, four-token groups, and an SSD backing tier with a 128 MiB RAM cache and skipped misses.

	L3.1-8B-I			Qwen3-8B		
Configuration	PPL	KL	tok/s	PPL	KL	tok/s
BF16 KV	10.32	0.000	7.3	12.89	0.000	6.7
KVarN 2/2	10.98	0.067	11.3	14.40	0.130	10.7
KIVI-2	10.90	0.059	8.8	13.85	0.106	8.3
KVQuant 2-bit	11.71	0.125	3.9	14.40	0.167	4.1
TurboQuant 2.5-bit	10.98	0.055	10.2	36.02	1.115	9.6
AQUA-KV 2-bit	10.66	0.030	1.5	13.14	0.029	1.4
KVQuant 3-bit	10.56	0.024	3.9	13.21	0.041	3.9
TurboQuant 3.5-bit	10.43	0.010	10.4	13.99	0.102	9.9
AQUA-KV 3-bit	10.39	0.006	1.5	12.98	0.017	1.4
KVarN 4/2	10.38	0.008	10.4	13.04	0.014	9.8
KIVI-4	10.33	0.002	8.6	12.90	0.004	8.0
KVQuant 4-bit	10.38	0.007	3.3	12.96	0.011	3.4
AQUA-KV 4-bit	10.32	0.002	1.5	12.95	0.016	1.3
StrataFlux-KV 2/2	10.47	0.011	11.4	13.32	0.030	10.9
StrataFlux-KV 3/2	10.34	0.002	11.0	12.98	0.005	10.5
StrataFlux-KV 3/3	10.33	0.001	10.7	12.93	0.004	10.2

Table A27. SSD-backed KV refinement on DGX Spark: Llama-3.1-8B-Instruct, PG-19 at 128K. PPL, KL, and decode tok/s by resident bits, fetched group fraction, and RAM-cache capacity. The 1/128, 128 MiB configuration repeats Table A26.

GPU resident	Fetch	64 MiB			128 MiB		
KV	fraction	PPL	KL	tok/s	PPL	KL	tok/s
2/2	1/256	10.50	0.014	11.5	10.50	0.013	11.5
	1/128	10.48	0.013	11.2	10.47	0.011	11.4
	1/64	10.47	0.013	10.6	10.46	0.010	10.8
3/2	1/256	10.34	0.003	11.1	10.34	0.002	11.1
	1/128	10.34	0.002	10.7	10.34	0.002	11.0
	1/64	10.34	0.002	10.3	10.34	0.002	10.4
3/3	1/256	10.33	0.002	10.8	10.34	0.001	10.8
	1/128	10.33	0.001	10.5	10.33	0.001	10.7
	1/64	10.33	0.002	10.0	10.33	0.001	10.2

GPU resident	Fetch	256 MiB			512 MiB		
KV	fraction	PPL	KL	tok/s	PPL	KL	tok/s
2/2	1/256	10.49	0.011	11.5	10.49	0.010	11.5
	1/128	10.46	0.009	11.4	10.45	0.007	11.3
	1/64	10.44	0.008	11.1	10.43	0.006	11.0
3/2	1/256	10.34	0.002	11.1	10.34	0.002	11.1
	1/128	10.34	0.002	10.9	10.34	0.002	10.9
	1/64	10.33	0.002	10.6	10.33	0.001	10.6
3/3	1/256	10.34	0.001	10.8	10.34	0.001	10.7
	1/128	10.33	0.001	10.6	10.33	0.001	10.6
	1/64	10.33	0.001	10.3	10.33	0.001	10.3

Table A28. KV refinement with no cache misses: Llama-3.1-8B-Instruct, PG-19 at 128K. PPL/KL by fetched group fraction, using four-token groups.

	GPU resident KV					
	2/2		3/2		3/3	
Fetch fraction	PPL	KL	PPL	KL	PPL	KL
1/256	10.44	0.005	10.34	0.001	10.34	0.001
1/128	10.42	0.003	10.33	0.001	10.33	0.001
1/64	10.39	0.003	10.33	0.001	10.33	0.001

Table A29. SSD-backed KV refinement on DGX Spark: Qwen3-8B, PG-19 at 128K. PPL, KL, and decode tok/s by resident bits, fetched group fraction, and RAM-cache capacity. The 1/128, 128 MiB configuration repeats Table A26.

GPU resident	Fetch	64 MiB			128 MiB		
KV	fraction	PPL	KL	tok/s	PPL	KL	tok/s
2/2	1/256	13.41	0.038	11.1	13.37	0.033	11.1
	1/128	13.36	0.036	10.6	13.32	0.030	10.9
	1/64	14.28	0.119	10.0	13.28	0.028	10.3
3/2	1/256	12.99	0.007	10.6	12.99	0.006	10.7
	1/128	12.97	0.006	10.1	12.98	0.005	10.5
	1/64	13.14	0.022	9.7	12.97	0.005	9.9
3/3	1/256	12.94	0.005	10.3	12.94	0.004	10.3
	1/128	12.93	0.004	9.8	12.93	0.004	10.2
	1/64	13.00	0.013	9.4	12.93	0.004	9.6

GPU resident	Fetch	256 MiB			512 MiB		
KV	fraction	PPL	KL	tok/s	PPL	KL	tok/s
2/2	1/256	13.34	0.030	11.0	13.28	0.026	11.0
	1/128	13.30	0.026	10.9	13.23	0.021	10.9
	1/64	13.25	0.023	10.5	13.20	0.018	10.4
3/2	1/256	12.98	0.005	10.6	12.97	0.004	10.6
	1/128	12.97	0.005	10.4	12.96	0.004	10.4
	1/64	12.96	0.004	10.1	12.96	0.003	10.0
3/3	1/256	12.94	0.004	10.3	12.93	0.003	10.2
	1/128	12.93	0.003	10.1	12.93	0.003	10.0
	1/64	12.93	0.003	9.8	12.93	0.003	9.7

Table A30. KV refinement with no cache misses: Qwen3-8B, PG-19 at 128K. PPL/KL by fetched group fraction, using four-token groups.

	GPU resident KV					
	2/2		3/2		3/3	
Fetch fraction	PPL	KL	PPL	KL	PPL	KL
1/256	13.05	0.013	12.93	0.002	12.90	0.002
1/128	13.01	0.009	12.93	0.002	12.90	0.002
1/64	12.97	0.006	12.92	0.002	12.90	0.001

Table A31. KV-only refinement: DGX Spark throughput (tok/s) when waiting on cache misses at 128K, with four-token groups. Columns give RAM-cache capacities (MiB). Matching PPL/KL appears in Tables A28 and A30.

GPU resident	Fetch	L3.1-8B-I				Qwen3-8B			
KV	fraction	64	128	256	512	64	128	256	512
2/2	1/256	8.6	8.8	9.1	9.2	7.8	8.5	9.3	8.5
	1/128	6.8	8.1	8.7	9.2	6.5	7.7	8.6	8.4
	1/64	4.0	5.8	7.5	8.3	2.9	5.5	7.1	8.1
3/2	1/256	8.2	8.6	9.0	9.0	7.8	8.4	8.9	8.3
	1/128	7.0	7.9	8.4	8.8	6.5	7.3	8.0	8.2
	1/64	3.9	6.0	7.5	8.0	3.1	5.5	6.9	8.2
3/3	1/256	8.2	8.6	8.9	8.9	7.8	8.2	8.4	8.6
	1/128	6.9	7.8	8.2	8.6	6.2	7.1	7.9	8.5
	1/64	4.0	5.8	7.2	7.9	3.0	5.4	7.0	7.4

C.5 Refinement with 2-bit Keys and 1-bit Values

Table A32 compares Base and StrataFlux-KV with 2-bit keys and 1-bit values.

Table A32. KV refinement with resident 2/1 bits at 128K: PG-19 PPL/KL, RULER accuracy, Qwen3-8B MATH500 accuracy, RTX 5090 decode tok/s at 128K, and memory. *KV memory only. StrataFlux-KV uses +0/+1 bits, two-token groups, and a 1/16 fetch fraction. In our Llama-3.1-8B-Instruct 32K sweep, +0/+1 yields the lowest PPL/KL among tested fetched-bitwidth pairs at matched payload.

	L3.1-8B-I					Qwen3-8B					
Method	PPL	KL	R128	tok/s	GB*	PPL	KL	R128	M500	tok/s	GB*
BF16	10.32	0.000	77.9	47.6	17.2	12.89	0.000	74.1	97.2	OOM	19.3
KVarN 2/1	31.29	1.180	40.7	77.8	2.0	112.6	2.318	18.0	28.6	74.2	2.3
StrataFlux-KV 2/1	11.65	0.135	64.1	63.5	2.0	15.16	0.181	60.4	93.4	59.8	2.3

C.6 KV Refinement Settings Sweep

Figure A6 varies fetched key/value bitwidths and the selected token-group fraction for two- and four-token groups. +4/+4 performs well across evaluated fetch budgets and resident precisions, supporting its use in our main comparisons.

D Joint Configurations and Task Results

These comparisons include the four configurations evaluated in Section 4.3. Table A35 reports all four weight/KV refinement combinations at each resident format; Figure 12 shows the YAQA 3-bit, 3/3 comparison.

D.1 Configuration Comparisons

Tables A33, A34, and A35 compare inference-system configurations, static weight/KV pairs, and the four weight/KV refinement combinations.

Table A33. Joint configurations at 128K: PG-19 PPL/KL, GPU memory (GB), and decode tok/s. *DGX Spark; †RTX 5090. StrataFlux uses RTX 5090's 5% weight configurations. GB counts stored weights and KV. W/A/KV: weight/activation/cache bits; bpw: bits/weight; FP8: 8-bit floating point.

	L3.1-8B-I					Qwen3-8B				
Configuration	**PPL**	**KL**	**GB**	**tok/s***	**tok/s†**	**PPL**	**KL**	**GB**	**tok/s***	**tok/s†**
BF16	10.32	0.000	33.2	7.3	47.6	12.89	0.000	35.7	6.7	OOM
vLLM AWQ4 + FP8 KV	11.03	0.074	14.3	16.8	106.8	13.49	0.078	15.8	15.3	96.6
EXL3 3.0 bpw + 4/4 KV	11.39	0.107	8.9	25.6	99.6	14.00	0.125	9.8	23.6	91.3
EXL3 4.0 bpw + 4/4 KV	10.78	0.051	9.8	23.6	101.7	13.73	0.086	10.6	21.9	92.9
EXL3 4.0 bpw + 2/2 KV	36.26	1.364	7.6	25.7	104.7	46.43	1.394	8.2	24.0	95.6
llama.cpp Q4_K_M + q8_0 KV	10.59	0.027	14.0	12.7	63.1	13.17	0.027	15.3	11.5	57.2
llama.cpp Q4_K_M + q4_0 KV	10.88	0.062	9.8	14.4	64.6	13.94	0.110	10.5	13.0	58.7
llama.cpp Q3_K_M + q8_0 KV	11.13	0.073	13.1	13.0	64.2	13.75	0.092	14.4	11.9	58.1
llama.cpp Q3_K_M + q4_0 KV	11.46	0.109	8.9	15.0	65.9	14.64	0.172	9.6	13.5	59.6
QServe W4A8KV4	11.34	0.105	10.3	11.2	42.6	14.67	0.163	11.2	9.9	37.2

	L3.1-8B-I				Qwen3-8B			
Configuration	**PPL**	**KL**	**GB**	**tok/s†**	**PPL**	**KL**	**GB**	**tok/s†**
BF16	10.32	0.000	33.2	47.6	12.89	0.000	35.7	OOM
StrataFlux / YAQA 2-bit / 2/2	13.79	0.312	6.4	106.6	14.91	0.165	7.1	97.6
StrataFlux / YAQA 2-bit / 3/2	13.76	0.313	7.0	104.2	14.88	0.162	7.7	95.4
StrataFlux / YAQA 2-bit / 3/3	13.74	0.311	7.5	102.1	14.84	0.161	8.3	93.3
StrataFlux / YAQA 3-bit / 2/2	10.80	0.045	7.3	102.0	13.25	0.038	8.0	91.7
StrataFlux / YAQA 3-bit / 3/2	10.74	0.044	7.8	100.3	13.20	0.034	8.6	90.3
StrataFlux / YAQA 3-bit / 3/3	10.72	0.042	8.4	98.3	13.17	0.034	9.2	88.2
StrataFlux / LNQ+GQ 3-bit / 2/2	12.98	0.241	7.3	117.9	17.46	0.292	8.0	105.1
StrataFlux / LNQ+GQ 3-bit / 3/2	12.95	0.242	7.9	115.8	17.29	0.286	8.6	103.3
StrataFlux / LNQ+GQ 3-bit / 3/3	12.95	0.242	8.4	113.1	17.27	0.285	9.2	100.6

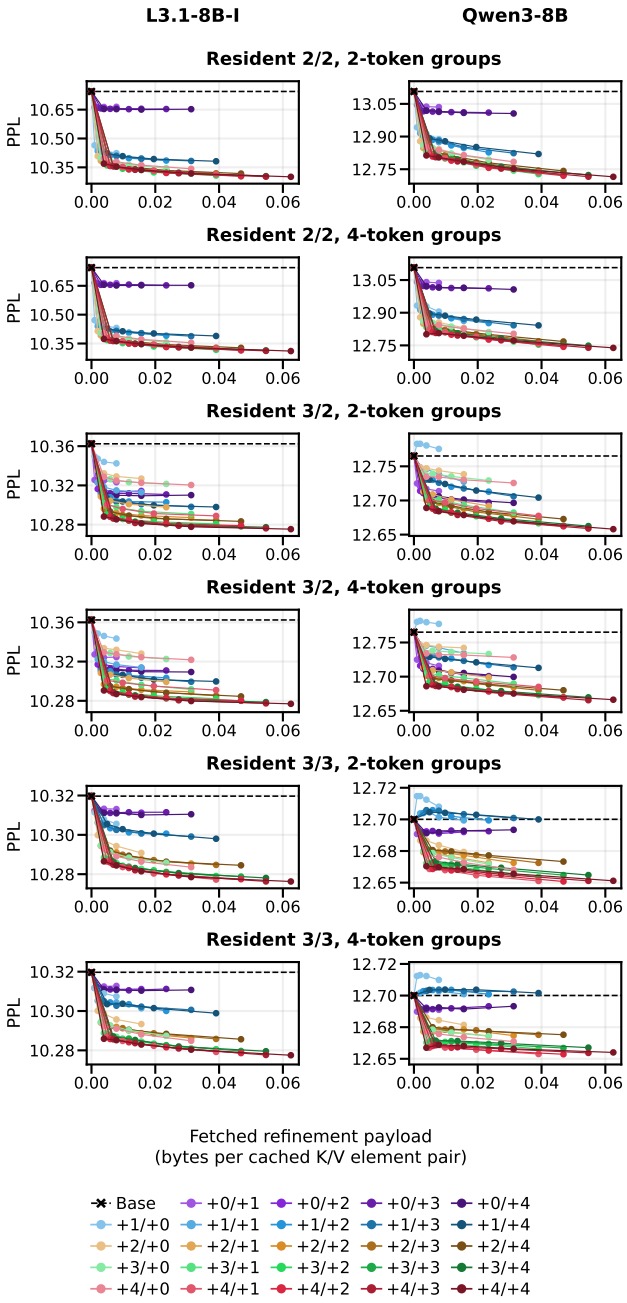

Figure A6. KV refinement payload versus PG-19 PPL at 32K: left, Llama-3.1-8B-Instruct; right, Qwen3-8B. Rows compare resident 2/2, 3/2, and 3/3 with two- and four-token groups. Curves vary fetched key/value bits independently from 0 to 4, excluding 0/0; zero leaves that component unrefined. Color identifies fetched key bits; darker shades indicate more fetched value bits. Each curve shows all registered selected-group fractions from 1/256 to 1/16. Payload bytes per cached K/V element pair equal fetched key plus value bits, multiplied by the selected fraction and divided by eight. Dashed lines mark unrefined bases.

Table A34. Joint refinement: static weight/KV pairs at 128K: PG-19 PPL/KL, GPU memory (GB), and RTX 5090 tokens/s. GB counts stored weights and KV. Neither weight nor KV refinement is enabled.

		L3.1-8B-I				Qwen3-8B			
Weight base	**KV configuration**	**PPL**	**KL**	**GB**	**tok/s**	**PPL**	**KL**	**GB**	**tok/s**
BF16	16/16	10.32	0.000	33.2	47.6	12.89	0.000	35.7	OOM
YAQA w2	KIVI-2	18.55	0.645	7.1	32.9	18.84	0.403	7.9	29.4
	KIVI-4	15.97	0.477	9.2	35.2	16.59	0.257	10.3	31.5
	KVQuant 2bit-1%	19.79	0.708	6.5	13.7	18.92	0.414	7.2	14.2
	KVQuant 3bit-1%	16.64	0.521	7.6	13.2	17.09	0.290	8.4	13.5
	KVQuant 4bit-1%	16.08	0.486	8.6	12.1	16.63	0.262	9.6	12.2
	KVarN 2/2	18.67	0.656	6.4	84.7	19.48	0.434	7.1	76.6
	KVarN 4/2	16.18	0.494	7.5	74.7	16.83	0.271	8.3	67.5
YAQA w3	KIVI-2	11.62	0.130	8.0	32.4	14.60	0.155	8.7	29.1
	KIVI-4	10.90	0.059	10.1	34.6	13.42	0.049	11.2	31.1
	KVQuant 2bit-1%	12.41	0.192	7.4	14.2	14.98	0.196	8.1	14.3
	KVQuant 3bit-1%	11.17	0.084	8.4	13.6	13.68	0.081	9.3	13.6
	KVQuant 4bit-1%	10.95	0.064	9.5	12.5	13.46	0.054	10.5	12.3
	KVarN 2/2	11.75	0.146	7.3	81.7	15.15	0.185	8.0	74.4
	KVarN 4/2	10.98	0.068	8.4	72.4	13.59	0.059	9.2	65.7
YAQA w4	KIVI-2	11.17	0.079	8.8	31.9	14.01	0.118	9.6	28.6
	KIVI-4	10.54	0.018	11.0	34.1	13.04	0.017	12.0	30.6
	KVQuant 2bit-1%	11.88	0.140	8.2	14.2	14.51	0.174	8.9	14.2
	KVQuant 3bit-1%	10.77	0.041	9.3	13.7	13.33	0.051	10.1	13.5
	KVQuant 4bit-1%	10.58	0.023	10.4	12.4	13.08	0.023	11.3	12.2
	KVarN 2/2	11.24	0.090	8.2	79.0	14.58	0.143	8.9	71.8
	KVarN 4/2	10.60	0.025	9.3	70.3	13.18	0.026	10.1	63.7
LNQ+GQ w3	KIVI-2	15.26	0.427	8.0	33.7	20.27	0.452	8.8	29.9
	KIVI-4	13.70	0.306	10.1	36.2	18.20	0.331	11.2	32.3
	KVQuant 2bit-1%	16.52	0.506	7.4	14.4	20.93	0.491	8.1	14.0
	KVQuant 3bit-1%	14.18	0.342	8.5	13.8	18.78	0.366	9.3	13.4
	KVQuant 4bit-1%	13.81	0.315	9.5	12.6	18.34	0.339	10.5	12.2
	KVarN 2/2	15.39	0.442	7.3	90.9	20.76	0.473	8.0	81.6
	KVarN 4/2	13.80	0.317	8.4	79.9	18.38	0.341	9.2	71.6
LNQ+GQ w4	KIVI-2	11.52	0.117	8.9	33.0	14.40	0.143	9.6	29.5
	KIVI-4	10.79	0.046	11.0	35.5	13.30	0.040	12.1	31.7
	KVQuant 2bit-1%	12.34	0.180	8.3	14.5	14.83	0.189	9.0	14.4
	KVQuant 3bit-1%	11.04	0.068	9.3	13.9	13.60	0.073	10.2	13.7
	KVQuant 4bit-1%	10.84	0.050	10.4	12.6	13.35	0.046	11.4	12.4
	KVarN 2/2	11.62	0.129	8.2	86.1	14.91	0.167	8.9	77.8
	KVarN 4/2	10.84	0.053	9.3	76.1	13.46	0.050	10.1	68.7

Table A35. Weight/KV component combinations at 128K: PG-19 PPL/KL, GPU memory (GB), and RTX 5090 tok/s. Each four-row group fixes resident formats and uses the 5% weight configuration. All groups use the same CUDA attention implementation. GB counts weights and KV.

		StrataFlux		L3.1-8B-I				Qwen3-8B			
Weights	**K/V**	**Weight**	**KV**	**PPL**	**KL**	**GB**	**tok/s**	**PPL**	**KL**	**GB**	**tok/s**
YAQA 2-bit	2/2	–	–	18.67	0.656	6.4	137.5	19.48	0.434	7.1	125.7
		✓	–	15.57	0.458		133.5	17.02	0.311		123.3
		–	✓	15.95	0.472		109.2	16.66	0.255		99.2
		✓	✓	13.79	0.312		106.6	14.91	0.165		97.6
YAQA 2-bit	3/2	–	–	16.42	0.511	7.0	134.6	17.06	0.287	7.7	123.0
		✓	–	14.15	0.348		130.0	15.23	0.189		120.0
		–	✓	15.91	0.473		107.2	16.62	0.255		97.4
		✓	✓	13.76	0.313		104.2	14.88	0.162		95.4
YAQA 2-bit	3/3	–	–	16.21	0.496	7.5	132.5	16.85	0.274	8.3	120.8
		✓	–	13.94	0.329		128.1	15.05	0.179		118.3
		–	✓	15.91	0.472		105.2	16.56	0.253		95.4
		✓	✓	13.74	0.311		102.1	14.84	0.161		93.3
YAQA 3-bit	2/2	–	–	11.75	0.146	7.3	130.7	15.15	0.185	8.0	120.2
		✓	–	11.52	0.124		126.6	14.83	0.167		114.6
		–	✓	10.96	0.058		105.2	13.48	0.048		95.8
		✓	✓	10.80	0.045		102.0	13.25	0.038		91.7
YAQA 3-bit	3/2	–	–	11.06	0.076	7.8	128.1	13.71	0.071	8.6	117.8
		✓	–	10.89	0.060		123.5	13.46	0.058		112.2
		–	✓	10.90	0.058		103.0	13.43	0.046		93.6
		✓	✓	10.74	0.044		100.3	13.20	0.034		90.3
YAQA 3-bit	3/3	–	–	10.98	0.067	8.4	125.8	13.56	0.061	9.2	115.9
		✓	–	10.81	0.051		122.0	13.32	0.049		110.7
		–	✓	10.89	0.057		100.8	13.39	0.045		92.0
		✓	✓	10.72	0.042		98.3	13.17	0.034		88.2
YAQA 4-bit	2/2	–	–	11.24	0.090	8.2	123.9	14.58	0.143	8.9	113.8
		✓	–	11.17	0.085		119.8	14.47	0.137		111.1
		–	✓	10.60	0.018		100.3	13.10	0.018		91.2
		✓	✓	10.55	0.015		97.5	13.04	0.015		90.1
YAQA 4-bit	3/2	–	–	10.66	0.031	8.7	121.6	13.31	0.038	9.5	111.4
		✓	–	10.61	0.027		117.1	13.23	0.034		108.3
		–	✓	10.53	0.017		98.3	13.05	0.014		89.6
		✓	✓	10.49	0.014		95.9	12.98	0.011		87.5
YAQA 4-bit	3/3	–	–	10.59	0.025	9.3	119.7	13.16	0.028	10.1	110.0
		✓	–	10.55	0.021		115.7	13.09	0.025		107.0
		–	✓	10.53	0.017		96.6	13.03	0.014		87.9
		✓	✓	10.48	0.013		94.6	12.96	0.011		85.8
LNQ+GQ 3-bit	2/2	–	–	15.39	0.442	7.3	152.2	20.76	0.473	8.0	137.6
		✓	–	14.48	0.372		149.4	19.61	0.423		134.6
		–	✓	13.68	0.301		119.2	18.42	0.336		107.0
		✓	✓	12.98	0.241		117.9	17.46	0.292		105.1
LNQ+GQ 3-bit	3/2	–	–	13.96	0.330	7.9	150.0	18.61	0.354	8.6	135.6
		✓	–	13.23	0.269		147.2	17.63	0.308		132.1
		–	✓	13.64	0.302		117.0	18.23	0.330		104.8
		✓	✓	12.95	0.242		115.8	17.29	0.286		103.3
LNQ+GQ 3-bit	3/3	–	–	13.83	0.318	8.4	147.1	18.41	0.343	9.2	133.1
		✓	–	13.11	0.258		144.4	17.46	0.298		130.1
		–	✓	13.64	0.301		114.1	18.20	0.329		102.6
		✓	✓	12.95	0.242		113.1	17.27	0.285		100.6
LNQ+GQ 4-bit	2/2	–	–	11.62	0.129	8.2	140.2	14.91	0.167	8.9	127.6
		✓	–	11.52	0.120		137.7	14.83	0.161		124.6
		–	✓	10.84	0.044		111.7	13.36	0.040		101.2
		✓	✓	10.77	0.038		110.4	13.30	0.035		99.5
LNQ+GQ 4-bit	3/2	–	–	10.90	0.060	8.8	138.4	13.59	0.062	9.5	125.8
		✓	–	10.82	0.054		135.6	13.53	0.056		122.9
		–	✓	10.78	0.044		110.0	13.31	0.038		99.5
		✓	✓	10.71	0.038		109.2	13.25	0.032		97.7
LNQ+GQ 4-bit	3/3	–	–	10.85	0.053	9.3	136.1	13.44	0.052	10.1	123.7
		✓	–	10.77	0.046		133.6	13.38	0.047		120.8
		–	✓	10.78	0.043		107.7	13.29	0.037		97.4
		✓	✓	10.71	0.037		106.8	13.23	0.032		95.5

D.2 Full Task Results

Tables A36, A37, and A38 report RULER, LongBench, MATH500, and AIME results for the evaluated full-system configurations.

Table A36. Joint refinement at 128K on Llama-3.1-8B-Instruct and Qwen3-8B: RULER/MATH500/AIME accuracy, throughput, and GPU memory. GB includes weights, KV, and DGX Spark's RAM caches. StrataFlux combines the 5% weight configuration with 2/2 KV.

DGX Spark

	L3.1-8B-I			Qwen3-8B				
Configuration	**R128**	**tok/s**	**GB**	**R128**	**M500**	**AIME**	**tok/s**	**GB**
BF16	77.9	7.3	33.2	74.1	97.2	71.0	6.7	35.7
StrataFlux (LNQ+GQ 4-bit)	71.1	22.1	8.5	70.6	96.0	65.6	20.8	9.2
StrataFlux (YAQA 4-bit)	73.5	21.4	8.4	70.3	96.2	69.0	20.3	9.1
StrataFlux (YAQA 3-bit)	73.2	23.3	7.6	68.4	95.2	68.8	21.9	8.3
StrataFlux (YAQA 2-bit)	60.5	25.6	6.7	65.2	89.6	40.6	23.9	7.4

RTX 5090

	L3.1-8B-I			Qwen3-8B				
Configuration	**R128**	**tok/s**	**GB**	**R128**	**M500**	**AIME**	**tok/s**	**GB**
BF16	77.9	47.6	33.2	74.1	97.2	71.0	OOM	35.7
LNQ+GQ 4-bit / 2/2	64.3	86.1	8.2	64.0	95.2	58.3	77.8	8.9
StrataFlux (LNQ+GQ 4-bit)	75.2	110.4	8.2	73.8	94.4	61.5	99.5	8.9
YAQA 4-bit / 2/2	67.5	79.0	8.2	65.4	95.0	62.5	71.8	8.9
StrataFlux (YAQA 4-bit)	77.2	97.5	8.2	73.5	96.4	69.2	90.1	8.9
YAQA 3-bit / 2/2	67.2	81.7	7.3	63.8	94.6	59.6	74.4	8.0
StrataFlux (YAQA 3-bit)	77.1	102.0	7.3	71.2	95.4	65.6	91.7	8.0
YAQA 2-bit / 2/2	50.4	84.7	6.4	58.8	68.0	11.0	76.6	7.1
StrataFlux (YAQA 2-bit)	63.8	106.6	6.4	68.3	91.2	47.1	97.6	7.1
llama.cpp Q4_K_M + q8_0	75.3	63.1	14.0	72.0	95.6	71.9	57.2	15.3
EXL3 4 bpw + 4/4	75.5	101.7	9.8	71.6	94.4	66.9	92.9	10.6
vLLM AWQ4 + FP8 KV	74.7	106.8	14.3	70.0	95.4	68.3	96.6	15.8

Table A37. Joint refinement: DGX Spark task grid. R4–R128 give RULER accuracy at 4K–128K; LB/M500 report LongBench/-MATH500 accuracy. LongBench prompts are truncated to 130,432 tokens here, versus 32,128 in the KV-only comparisons. Throughput uses 128K context. Each pair uses the same weight base and resident 2/2 KV; Base disables both refinements. Configurations match the main full-system table.

DGX Spark / Llama-3.1-8B-Instruct

Weights	**Method**	**R4**	**R8**	**R16**	**R32**	**R64**	**R128**	**LB**	**tok/s**
BF16		95.7	95.1	92.9	87.1	85.8	77.9	49.6	7.3
llama.cpp Q4_K_M + q8_0		95.5	95.0	93.2	87.8	85.3	75.3	49.1	12.7
EXL3 4 bpw + 4/4		95.3	94.7	92.6	88.7	84.9	75.5	48.1	23.6
vLLM AWQ4 + FP8 KV		95.1	94.1	92.1	87.6	83.2	74.7	48.4	16.8
LNQ+GQ 4-bit	Base	93.7	92.2	89.9	84.0	79.8	64.3	47.9	21.3
	StrataFlux	95.4	94.2	92.2	86.1	84.0	71.1	48.3	22.1
YAQA 4-bit	Base	94.2	93.6	91.2	86.0	80.3	67.5	48.7	21.2
	StrataFlux	95.9	94.9	93.1	87.3	84.5	73.5	49.0	21.4
YAQA 3-bit	Base	93.8	92.8	90.4	85.7	80.6	67.2	46.8	23.0
	StrataFlux	95.0	94.7	92.6	86.9	84.0	73.2	47.0	23.3
YAQA 2-bit	Base	91.6	90.5	87.1	80.4	73.0	50.4	46.2	24.8
	StrataFlux	93.9	93.2	90.9	85.5	80.0	60.5	46.5	25.6

DGX Spark / Qwen3-8B

Weights	**Method**	**R4**	**R8**	**R16**	**R32**	**R64**	**R128**	**LB**	**M500**	**AIME**	**tok/s**
BF16		95.2	93.8	93.8	92.1	80.6	74.1	49.4	97.2	71.0	6.7
llama.cpp Q4_K_M + q8_0		95.0	93.9	93.4	91.9	79.5	72.0	49.2	95.6	71.9	11.5
EXL3 4 bpw + 4/4		94.8	94.2	93.7	90.8	78.5	71.6	48.1	94.4	66.9	21.9
vLLM AWQ4 + FP8 KV		94.8	93.9	93.0	91.0	78.8	70.0	48.0	95.4	68.3	15.3
LNQ+GQ 4-bit	Base	94.0	92.9	91.3	87.3	72.4	64.0	48.1	95.2	58.3	19.5
	StrataFlux	95.5	94.2	93.8	90.8	78.8	70.6	48.4	96.0	65.6	20.8
YAQA 4-bit	Base	94.5	93.7	92.1	87.9	71.1	65.4	48.4	95.0	62.5	19.5
	StrataFlux	95.5	94.5	93.2	91.3	76.4	70.3	49.0	96.2	69.0	20.3
YAQA 3-bit	Base	94.4	93.4	92.2	86.6	71.1	63.8	48.1	94.6	59.6	20.9
	StrataFlux	95.3	94.6	93.7	91.4	76.0	68.4	48.9	95.2	68.8	21.9
YAQA 2-bit	Base	92.9	90.9	87.2	77.8	63.8	58.8	43.3	68.0	11.0	22.5
	StrataFlux	94.4	93.0	90.8	83.6	70.2	65.2	44.1	89.6	40.6	23.9

Table A38. Joint refinement: RTX 5090 task grid. R4–R128 give RULER accuracy at 4K–128K; LB/M500 report LongBench/-MATH500 accuracy. LongBench prompts are truncated to 130,432 tokens here, versus 32,128 in the KV-only comparisons. Throughput uses 128K context. Each pair uses the same weight base and resident 2/2 KV; Base disables both refinements. Configurations match the main full-system table.

RTX 5090 / Llama-3.1-8B-Instruct

Weights	**Method**	**R4**	**R8**	**R16**	**R32**	**R64**	**R128**	**LB**	**tok/s**
BF16		95.7	95.1	92.9	87.1	85.8	77.9	49.6	47.6
llama.cpp Q4_K_M + q8_0		95.5	95.0	93.2	87.8	85.3	75.3	49.1	63.1
EXL3 4 bpw + 4/4		95.3	94.7	92.6	88.7	84.9	75.5	48.1	101.7
vLLM AWQ4 + FP8 KV		95.1	94.1	92.1	87.6	83.2	74.7	48.4	106.8
LNQ+GQ 4-bit	Base	93.7	92.2	89.9	84.0	79.8	64.3	47.9	86.1
	StrataFlux	95.4	94.3	91.7	86.2	84.7	75.2	48.4	110.4
YAQA 4-bit	Base	94.2	93.6	91.2	86.0	80.3	67.5	48.7	79.0
	StrataFlux	95.9	94.7	93.5	87.4	85.3	77.2	49.1	97.5
YAQA 3-bit	Base	93.8	92.8	90.4	85.7	80.6	67.2	46.8	81.7
	StrataFlux	94.9	94.7	93.0	86.4	83.8	77.1	47.0	102.0
YAQA 2-bit	Base	91.6	90.5	87.1	80.4	73.0	50.4	46.2	84.7
	StrataFlux	94.1	92.9	91.4	84.2	81.4	63.8	46.5	106.6

RTX 5090 / Qwen3-8B

Weights	**Method**	**R4**	**R8**	**R16**	**R32**	**R64**	**R128**	**LB**	**M500**	**AIME**	**tok/s**
llama.cpp Q4_K_M + q8_0		95.0	93.9	93.4	91.9	79.5	72.0	49.2	95.6	71.9	57.2
EXL3 4 bpw + 4/4		94.8	94.2	93.7	90.8	78.5	71.6	48.1	94.4	66.9	92.9
vLLM AWQ4 + FP8 KV		94.8	93.9	93.0	91.0	78.8	70.0	48.0	95.4	68.3	96.6
LNQ+GQ 4-bit	Base	94.0	92.9	91.3	87.3	72.4	64.0	48.1	95.2	58.3	77.8
	StrataFlux	95.4	94.0	93.8	91.5	81.0	73.8	48.7	94.4	61.5	99.5
YAQA 4-bit	Base	94.5	93.7	92.1	87.9	71.1	65.4	48.4	95.0	62.5	71.8
	StrataFlux	95.5	94.5	93.4	91.1	79.5	73.5	49.0	96.4	69.2	90.1
YAQA 3-bit	Base	94.4	93.4	92.2	86.6	71.1	63.8	48.1	94.6	59.6	74.4
	StrataFlux	95.4	94.4	94.0	91.3	78.9	71.2	48.8	95.4	65.6	91.7
YAQA 2-bit	Base	92.9	90.9	87.2	77.8	63.8	58.8	43.3	68.0	11.0	76.6
	StrataFlux	94.3	92.7	91.2	85.6	72.3	68.3	44.5	91.2	47.1	97.6

E Applicability to Other KV Bases

The main KV implementation adopts KVarN's quantization format. Other bases require a refinement encoding and decoder that combine their resident codes with separately stored refinement bits. The selection of token groups and the full-sequence attention update remain the same, but the refinement levels must match the base's quantization bins—the value ranges mapped to each resident code—and its coordinate transformations and metadata.

Nonuniform quantization bins. Our KVQuant-style adaptation starts from sensitivity-weighted nonuniform codebooks (Hooper et al., 2024) and fits finer levels within each bin using sensitivity-weighted calibration samples. Our TurboQuant adaptation uses its scalar quantizer optimized for mean squared error in the rotated, normalized coordinate system (Zandieh et al., 2026); each bin is refined using levels derived from the base's coordinate distribution restricted to that bin. Both keep the resident code as the identifier of the coarse bin and store the finer-level index separately.

Evaluation scope. Table A39 compares Base and StrataFlux-KV reconstruction accuracy for the KVQuant and TurboQuant adaptations.

Table A39. KV-only refinement: nested refinements for other resident formats, PG-19 at 128K. Base and StrataFlux-KV rows report PPL/KL side by side. Refinement uses +4/+4 bits, two-token groups, and a 1/128 fetch fraction.

		L3.1-8B-I		Qwen3-8B	
Resident base	**Method**	**PPL**	**KL**	**PPL**	**KL**
KVQuant 2-bit + 1% outliers	Base	11.71	0.125	14.40	0.167
	StrataFlux-KV	10.46	0.012	13.07	0.021
KVQuant 3-bit + 1% outliers	Base	10.56	0.024	13.21	0.041
	StrataFlux-KV	10.35	0.003	12.92	0.005
KVQuant 4-bit + 1% outliers	Base	10.38	0.007	12.96	0.011
	StrataFlux-KV	10.33	0.001	12.89	0.002
TurboQuant 2.5-bit	Base	10.98	0.055	36.02	1.115
	StrataFlux-KV	10.47	0.005	14.03	0.091
TurboQuant 3.5-bit	Base	10.43	0.010	13.99	0.102
	StrataFlux-KV	10.34	0.001	12.89	0.004

F Mac Evaluation

Table A40 compares weight-only, KV-only, and joint refinement on a MacBook Air M3 (16 GB, Metal), using Llama-3.1-8B-Instruct with llama.cpp's Q3_K_M weights and KVarN 2/2 KV.

Table A40. MacBook Air M3 weight/KV refinement: WikiText-2 PPL/KL on one sequentially scored 16K-token window. Throughput averages 30 paired 32-token blocks at the listed initial contexts.

	StrataFlux				Throughput (tok/s)	
Configuration	**Weight**	**KV**	**PPL**	**KL**	**Empty context**	**16K context**
Base	–	–	6.71	0.126	14.1	10.6
StrataFlux-W	✓	–	6.66	0.118	15.1	11.3
StrataFlux-KV	–	✓	6.44	0.081	13.8	9.9
StrataFlux	✓	✓	6.38	0.073	14.4	10.5

Weight refinement selects $k = 2$ channels per 1,024 in all four projection sets. Both refinement components use separate 128 MiB SSD-backed RAM caches; KV refinement uses the main unified-memory settings.

G Parallel Weight Refinement with CUDA Streams

Separate CUDA streams do not guarantee parallel execution. If the base projection starts first, its large grid can occupy all streaming multiprocessors (SMs), leaving refinement queued until an SM becomes available. Merely assigning refinement to another stream does not prevent this serialization.

DecDEC's projection-local overlap. DecDEC (Park et al., 2025) keeps refinement on the current stream and places the base projection on another stream behind an event wait. The intent is for faster same-stream scheduling to start refinement before the cross-stream base projection, including during CUDA Graph replay. Refinement's small grid leaves SMs idle, allowing the scheduler to place the projection's thread blocks on those SMs while refinement runs.

The event is recorded before refinement, so the cross-stream dependency does not wait for refinement to finish. A second event joins the projection back into the current stream before the output is consumed; that stream's own ordering also ensures refinement is complete. Both kernels add to the output atomically. The snippets show the stream/event structure captured in the CUDA Graph, with kernel arguments and event/stream allocation omitted.

```
cudaEventRecord(fork, current);
refine<<<small_grid,threads,0,current>>>(...);
cudaStreamWaitEvent(other, fork, 0);
base_gemv(..., other);
cudaEventRecord(done, other);
cudaStreamWaitEvent(current, done, 0);
```

StrataFlux's fork-and-continue scheduling. We continue the graph on another stream while the fetch-lead or fetch-lag kernel remains in flight on the previous execution stream. Each outstanding operation j has its own events and continuation stream. The graph waits for its completion only where its result is needed, allowing projection sets with different shift distances s to overlap independently. For fetch lead, when the projection's input activations become available, we fork again: the refinement branch waits for the earlier fetch and applies refinement using those activations, while the base projection proceeds on another stream. The branches join before the projection output is consumed. Fetch lag instead joins before the deferred residual update is consumed. Contributions use atomic additions to initialized accumulators.

```
cudaEventRecord(fork[j], current);
fetch_or_refine<<<small_grid,threads,0,current>>>(...);
cudaEventRecord(done[j], current);
cudaStreamWaitEvent(next[j], fork[j], 0);
current = next[j];
// Where j's result is consumed:
cudaStreamWaitEvent(current, done[j], 0);
```