XposedMIA: Exploiting Explainable Artificial Intelligence (XAI) for One-Shot Membership Inference Attacks

Omar Hamdache
omar.hamdache@bilkent.edu.tr
Bilkent University
Ankara, Turkey

Serdar Kara
serdar.kara@ug.bilkent.edu.tr
Bilkent University
Ankara, Turkey

Kıvanç Kuzey Dikici
kuzey.dikici@ug.bilkent.edu.tr
Bilkent University
Ankara, Turkey

Abdullah Caglar Oksuz
axo223@case.edu
Case Western Reserve University
Cleveland, Ohio, USA

Anisa Halimi
anisa.halimi@ibm.com
IBM Research
Dublin, Ireland

Erman Ayday
exa208@case.edu
Case Western Reserve University
Cleveland, Ohio, USA

Sinem Sav*
sinem.sav@cs.bilkent.edu.tr
Bilkent University
Ankara, Turkey

arXiv:2611.12067v1 [cs.LG] 9 Nov 2026

Abstract

Machine learning (ML) services increasingly expose explainability outputs alongside model predictions. While such explanations are intended to improve transparency, they may also reveal information about the underlying training data. In this work, we show that the local surrogate geometry employed by Local Interpretable Model-agnostic Explanations (LIME) can be repurposed as a proxy for a sample's distance to the model's decision boundary. We extend this boundary-distance view to tabular counterfactual explanations and, for the first time, to visual counterfactual explanations for image classifiers. We present XposedMIA, a one-shot, explanation-based black-box membership inference attack (MIA) framework that converts the returned explanation into a proxy for the queried sample's distance to the model decision boundary. After a one-time calibration step on a small labeled calibration set of as few as ~ 30 samples, our attack requires only a single query for each target sample. Unlike prior label-only attacks requiring iterative queries or shadow training, our method relies solely on the predicted label and its accompanying explanation. Our evaluation indicates that XposedMIA achieves high effectiveness, with success rates of up to 0.95 using only a single query per target, revealing that releasing explanations introduces an additional channel for membership leakage.

Keywords

membership inference attack, explainable artificial intelligence, black-box attack, decision boundary, machine learning as a service, privacy

1 Introduction

The rapid adoption of machine learning (ML) models for everyday tasks is increasing their use by individuals, businesses, and organizations [4, 28, 48]. Since ML models are generally resource-intensive, require training with large datasets, and are difficult to deploy, machine learning as a service (MLaaS) platforms are becoming more common [45]. These platforms provide models that users can query with their own data and receive the output back [58].

A major problem arises in such MLaaS platforms when providers do not want to disclose the internals of their models or datasets, as these are considered intellectual property [58]. While this protects model internals, the resulting black-box nature introduces risks. Complex models such as neural networks (NNs) offer only the input and output to the user, without revealing what leads to a specific prediction or classification decision [26]. This problem is especially salient in high-stakes domains, including clinical decision support and credit or financial decision-making, where explanations support clinical oversight, accountability, and responsible deployment [6, 9, 27]. Moreover, specific legal regimes require reason-giving or meaningful explanations for certain automated decisions, including adverse credit actions in the US and decisions based on specified high-risk AI systems in the EU [15, 18]. For this reason, some MLaaS platforms offer explanations through explainable artificial intelligence (XAI) tools [1, 2, 14, 51].

XAI enables users to interpret model predictions and gain insight into the factors influencing their decisions. A wide range of explanation techniques has been developed, each adopting a distinct methodological approach and offering various perspectives on how a model arrives at its predictions [29, 36, 46, 50, 56, 57].

While XAI facilitates transparency, the shared information inherently exposes behavioral signals that expand the model's attack surface. ML models are vulnerable to a wide range of security/privacy attacks, including model extraction [5, 11, 41], adversarial attacks [23], model inversion [22, 67], and membership inference attacks (MIAs) [16, 44, 54, 64]. Prior work has demonstrated that the integration of XAI techniques can further increase the vulnerability of ML models to such threats, effectively broadening the attack surface [19, 37, 39, 41, 53, 55, 67]. A successful MIA can lead to serious data leakage, thereby compromising data privacy. For example, if a model is trained to classify a sensitive property such as cancer type, a successful MIA could reveal that a specific individual was included in the training set, exposing sensitive information about

*Corresponding author.

their health. Moreover, since training datasets may be considered intellectual property, a successful MIA on a public dataset could reveal whether the dataset was used to train the model.

In this work, we study whether explanations can be turned into a membership signal in a black-box setting where the adversary receives a predicted label and its accompanying explanation. We present XposedMIA (illustrated in Figure 1), a one-shot explanation-based MIA framework that derives a proxy for the queried sample's distance to the model decision boundary from the returned explanation. This estimated boundary distance is then used to infer the membership status. Prior label-only boundary-distance attacks [13, 31, 33] require a substantial number of queries or shadow model training with large datasets to estimate the position of the decision boundary and do not leverage information from XAI methods. As a result, these approaches incur high query and computational costs, limiting their practicality in real-world black-box settings. In contrast, XposedMIA reduces the attack cost to a single external service query per target sample after a one-time calibration phase. Existing *XAI-based attacks*, on the other hand, typically rely on stronger adversarial assumptions or incur higher attack costs, such as white-box model access, shadow-model training, or repeated queries to the target model [34, 53]. To the best of our knowledge, no prior work has explicitly used LIME's local linear surrogate as a proxy for decision-boundary distance under a black-box setting or investigated visual counterfactual explanations for membership inference on image classifiers.

In addition, XposedMIA is designed to minimize the number of queries, which is crucial in MLaaS settings where application programming interfaces (APIs) often enforce query budgets and detect adaptive query patterns. Existing perturbation-based attacks require multiple queries per sample, increasing their visibility and making them easier to detect. In contrast, XposedMIA uses only a single standard query per target instance after calibration. We thus demonstrate that exposing explanations alongside model predictions introduces a lower-interaction attack surface for MIA, enabling more stealthy attacks under practical API constraints.

XposedMIA consists of two main steps. First, we calibrate the distance threshold τ by issuing queries to the MLaaS platform and collecting predicted labels together with their corresponding explanations. This calibration phase is used only to determine a membership decision threshold; it does not require training shadow models. While the best performance is achieved with approximately 100 calibration queries, the method remains effective with smaller calibration sets as few as 30 samples. Then, the attack requires only a single query per target sample and makes its decision based on the sample's distance to the estimated decision boundary. This enables a one-shot, explanation-guided black-box MIA that leverages explanations without requiring shadow-model training.

We note here that the calibration requirement is substantially weaker than the data requirements of shadow-model-based pipelines. The calibration set is used only to determine a threshold and decision direction, rather than to train surrogate models or approximate the target model. On the other hand, prior attacks require significantly larger auxiliary datasets to train models that mimic the target, sometimes reaching up to 10^6 queries, resulting in much heavier data requirements. We summarize our main contributions below:

- We demonstrate that releasing LIME and counterfactual explanations with model outputs can introduce measurable membership leakage in a label- and explanation-only black-box setting.
- We introduce the first unified **one-shot** explanation-based membership inference framework that leverages explanations as the attack surface: a single query suffices to recover membership information about a target sample.
- We evaluate XposedMIA across four tabular and two image benchmark datasets. Our experimental results, ablation studies, and comparisons with prior work demonstrate its strong effectiveness and efficiency.
- We propose and evaluate, to the best of our knowledge, the first extension of MIA to visual counterfactual explanations on image classifiers using SEDC [59] and DiVE [47].

Scope. In this work, we base our MIA method on LIME [46] as a local surrogate explanation method and extend our method to other XAI methods, including three tabular counterfactual methods, Wachter [60], NICE [8], and Growing Spheres [30], and two visual counterfactual methods, SEDC [59] and DiVE [47]. These methods span local surrogate and counterfactual explanations, providing different forms of information about the model's decision structure. Local surrogate explanations approximate the model's behavior around the queried sample, while counterfactual explanations identify changes to the input that alter the model's prediction. We derive a proxy for the queried sample's distance to the decision boundary from explanation outputs. More details on these XAI methods and the rationale behind targeting them can be found in Section 3.

2 Related Work

We organize prior work on MIAs into five categories: (i) *classical MIAs*, which typically rely on confidence scores, loss values, or internal model signals; (ii) *label-only MIAs*, which operate in black-box settings using only the predicted label; (iii) *low-query label-only MIAs*, which aim to reduce the substantial query cost of earlier label-only attacks; (iv) *feature-attribution-based MIAs*, which exploit information exposed by explanation methods that return feature attributions; and (v) *counterfactual-explanation-based attacks*, which exploit information revealed by counterfactual explanations for membership inference, model extraction, or reconstruction.

Classical Membership Inference Attacks. Classical MIAs are commonly studied in black-box or white-box settings in which the attacker has access to confidence scores or other signals produced by the target model. In black-box classification settings, these scores typically correspond to posterior or confidence vectors over the output classes. Shokri et al. [54] introduced the first MIA against machine-learning classifiers, training shadow models on auxiliary data drawn from the same distribution as the target model's training data. The shadow models are queried with member and non-member samples, and their outputs are subsequently used to train a separate membership classifier. Salem et al. [49] later showed that comparable attacks can be constructed using only a single shadow model.

More recently, Ye et al. [64] formulated membership inference as a hypothesis-testing problem and developed stronger loss-based

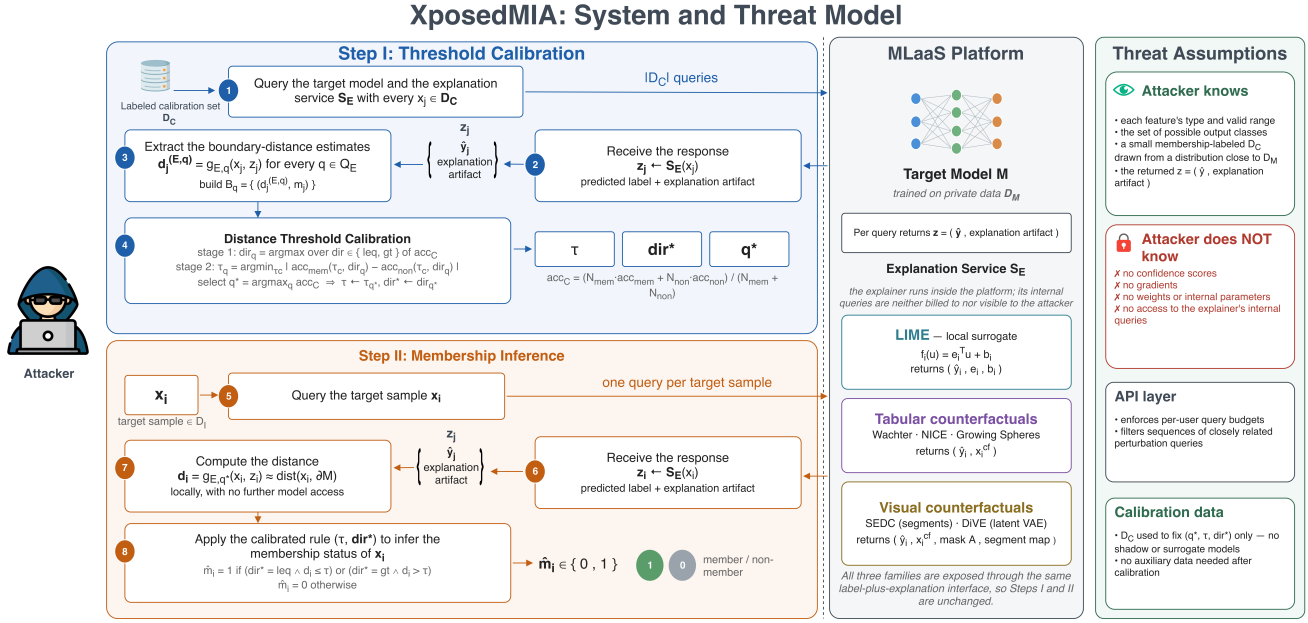

Figure 1: Overview of XposedMIA. The attacker interacts with the MLaaS platform only through the explanation service $\mathcal{S}_\mathcal{E}$, which returns $z = (\hat{y}$, explanation artifact) under Local Interpretable Model-Agnostic Explanations (LIME) [46], tabular counterfactuals (Wachter [60], Nearest Instance Counterfactual Explanations (NICE) [8], and Growing Spheres (GS) [30]), or visual counterfactuals (Search for EviDence Counterfactual (SEDC) [59] and Diverse Valuable Explanations (DiVE) [47]). In Step I, the attacker queries $\mathcal{S}_\mathcal{E}$ with the labeled calibration set $\mathcal{D}_C$, converts each response into a boundary-distance estimate $d_j^{(\mathcal{E},q)} = g_{\mathcal{E},q}(x_j, z_j)$, and fixes (q^*, τ, dir^*) via Algorithm A.1 (Appendix A). In Step II, one query per target sample x_i yields z_i, from which $d_i = g_{\mathcal{E},q^*}(x_i, z_i)$ is computed locally and compared against (τ, dir^*) (Algorithm 1). Only the extractor $g_{\mathcal{E},q}$ is explanation-specific (Sec. 5.2); both steps are otherwise identical across explainers.

attacks using reference models. Carlini et al. [10] further demonstrated that average-case metrics can obscure privacy leakage in operationally relevant regimes and advocated evaluating MIAs at low false-positive rates (FPRs). They introduced the likelihood ratio attack (LiRA), which has since become an important benchmark.

Nasr et al. [40] extended MIA to the white-box setting and showed that access to internal signals, including gradients and intermediate representations, can further amplify membership leakage. Their results demonstrate that stronger access to the target model can reveal membership information beyond that available from prediction outputs alone. These attacks generally rely on confidence or loss information, internal model signals, or auxiliary model training. In contrast, XposedMIA requires none of these forms of access.

Label-only Membership Inference Attacks. Label-only MIAs consider a more restrictive black-box setting in which the attacker observes only the predicted class label. Yeom et al. [65] proposed a simple label-only baseline in which a correctly classified sample is predicted to be a member. This attack, commonly referred to as the *gap attack*, is directly related to the difference between the target model's training and test accuracy and is widely used as a baseline in label-only settings.

Choquette-Choo et al. [13] introduced a stronger label-only MIA based on a sample's distance to the decision boundary. The intuition is that training samples may lie farther from the decision boundary and therefore require larger perturbations to change their predicted label. They estimate this distance using an adversarial sample generation technique called HopSkipJump [12], which searches for a small perturbation that changes the target model's prediction. The magnitude of the resulting perturbation is then used as a proxy for boundary distance and, consequently, as a membership signal.

Li et al. [33] proposed a label-only MIA based on transferability between shadow and target models. They additionally introduced boundary-distance attacks based on HopSkipJump [12] and the query-efficient boundary-based black-box attack [32], removing the dependence on this transferability assumption. More recently, Li et al. [31] proposed an enhanced label-only MIA that learns boundary-distance statistics from shadow models and performs MIA using $\sim 10 - 50$ target-model queries per sample.

A common limitation of these MIAs is therefore their per-sample query cost, since estimating the decision boundary generally requires repeatedly querying the target model. Such repeated probing may be constrained by rate limits or abuse-detection mechanisms in MLaaS environments. XposedMIA instead performs a one-time calibration using a small calibration set and subsequently requires only one explanation request per target sample.

sec:low-query-label-onlyLow-query Label-only MIAs. Peng et al. [44] proposed OSLO, a one-shot label-only MIA based on transferable adversarial examples. OSLO requires only a single hard-label query to the target model for each inference sample; however, this inference relies on an offline stage in which the attacker trains

an ensemble of surrogate models on an auxiliary dataset and uses these models to construct sample-specific adversarial examples. In their CIFAR experiments, OSLO uses 25,000 auxiliary samples to train its surrogate models. Oksuz et al. [42] introduced LoMime, a label-only MIA that reduces the amortized query cost of membership inference through model extraction. LoMime first extracts a surrogate of the target model using hard-label queries and then performs MIA locally on the extracted model, requiring no additional target-model queries for subsequent inference samples. In its strongest configurations, LoMime uses up to 10^6 target-model queries during this one-time extraction stage, making the approach particularly attractive when this cost can be amortized over a large number of membership queries. Overall, OSLO and LoMime reduce the per-sample query cost by shifting substantial work to an initial surrogate-training or model-extraction. Although this amortization is advantageous when attacking very large numbers of samples, the corresponding upfront data, computation, and target-query requirements can be substantial and may be constrained by MLaaS. In contrast, XposedMIA achieves one-shot inference without surrogate-model training, adversarial-example construction, or model extraction: it requires only a small calibration set to fix a scalar threshold and its direction, followed by a single explanation request per target sample. Thus, XposedMIA has a more realistic assumption about the amount of auxiliary data.

Explanation-Based Membership Inference Attacks. A separate line of work considers MLaaS settings in which an explanation is returned alongside the model's prediction and investigates whether the additional information leaks membership.

Shokri et al. [53] studied MIAs in which the adversary obtains both prediction outputs and feature-based explanations, including Integrated Gradients (IG) [57], Layer-wise Relevance Propagation (LRP) [7], LIME [46], and SmoothGrad [56]. They showed that statistics derived from predictions and explanations, including their variance, can reveal membership information. Their attacks use either an optimal threshold, which assumes knowledge unavailable to a practical attacker, or thresholds calibrated using shadow models trained on auxiliary data. Their approach is effective for several backpropagation-based explanation methods, including gradients, IG, and LRP, but does not obtain useful membership signals from perturbation-based methods such as LIME and SmoothGrad.

Liu et al. [34] further showed that feature-attribution explanations can expose robustness differences between members and non-members. Their attack iteratively perturbs the most and least important pixels identified by an attribution map and uses the resulting model behavior as membership features. The attack requires training auxiliary models to approximate the target model's behavior, and additionally relies on posterior outputs. Their approach is primarily evaluated on image classifiers and achieves strong attack performance, but requires substantially richer attacker-side machinery than a direct threshold-based attack.

Few prior works combine explanation access with hard-label prediction access while avoiding confidence scores and model internals. Ma et al. [37] use explanations to guide iterative feature removal until the target model's predicted label changes. The resulting robustness score is then used for membership inference, with shadow models used to calibrate the decision. Although this approach substantially reduces the query cost compared with earlier label-only attacks, obtaining the membership signal still requires repeatedly querying the target model and incurs hundreds of queries per target sample in addition to shadow-model training.

Overall, prior feature-attribution-based MIAs commonly rely on confidence or posterior information, shadow-model training, or repeated target-model queries. XposedMIA instead derives its membership signal directly from the returned explanation and requires only lightweight calibration before one-shot inference.

Counterfactual-Explanation-Based Attacks. Counterfactual explanations expose a different form of information about a model's local decision. Given an input x, a counterfactual explanation produces a modified instance x' whose prediction differs from that of x while seeking to keep the modification small according to a notion of proximity [8, 30, 60]. Consequently, the distance $\|x - x'\|$ provides an explanation-derived measure of the change required to alter the model's decision.

This property has been exploited for model extraction and reconstruction. Aïvodji et al. [5] query a target service for predictions and counterfactual explanations and use the resulting samples to train a surrogate model. Wang et al. [61] improve the efficiency of this strategy by additionally obtaining a counterfactual of each returned counterfactual, reducing the decision-boundary shift. Dissanayake and Dutta [17] formalize this intuition by showing that a sufficient number of optimal counterfactuals can be used to approximate convex decision boundaries. Related reconstruction attacks have also been demonstrated for feature-attribution explanations: Milli et al. [39] reconstruct models from gradient-based explanations, while Oksuz et al. [41] extract models from LIME and SHAP outputs.

Counterfactual explanations have also been shown to support membership inference. Pawelczyk et al. [43] use the recourse cost $c(x, x')$ as a membership signal, and the intuition is that training can move the decision boundary away from member samples, causing members and non-members to exhibit systematically different counterfactual distances. They instantiate this idea through two attacks: a direct counterfactual-distance threshold, which sweeps possible values to select an optimal threshold, and a likelihood-ratio test inspired by Carlini et al. [10] that uses auxiliary data and shadow models to estimate reference distributions. They further observe that the direction of the membership signal can depend on the counterfactual generator, with some methods producing smaller rather than larger distances for members.

Ezzeddine et al. [20] further demonstrated that counterfactual explanations can increase membership leakage. Their approach augments a shadow-model-based MIA with information derived from counterfactual explanations and evaluates the attack on two structured numerical datasets. On the EEG dataset, incorporating counterfactual information increases the reported attack success rate from approximately 58% to 77%. Unlike their learned shadow-model-based attack, XposedMIA performs MIA using a lightweight calibrated distance rule and requires no shadow-model training.

Taken together, prior work establishes that both feature-attribution and counterfactual explanations can expose membership-relevant information. XposedMIA extends this line of work in three directions: it exploits LIME's local surrogate geometry as a membership signal, performs inference without confidence vectors, model internals, or shadow-model training, and, to the best of our knowledge,

provides the first systematic MIA exploiting visual counterfactual explanations for image classifiers.

3 Explainable Artificial Intelligence (XAI)

Explainable Artificial Intelligence (XAI) aims to make the outputs of complex ML models more interpretable [35]. Since modern models, particularly NNs, often provide limited insight into the reasoning behind their predictions, some MLaaS providers expose post-hoc explanation tools alongside model predictions. For instance, Amazon SageMaker Clarify [52] provides SHapley Additive exPlanations (SHAP)-based explanations for model predictions [51]. Post-hoc explanation methods can also be model-agnostic: LIME [46] and model-agnostic variants of SHAP [36], for example, infer feature contributions from model outputs without requiring the user to access the model's internal parameters. In this work, we consider two forms of local explanation: local surrogate explanations and counterfactual explanations. We note that we do not consider SHAP, as it lacks the distance-related features that we employ in our attack.

Local Interpretable Model-Agnostic Explanations (LIME) [46]. LIME explains an individual prediction by learning an interpretable approximation of the model's behavior in the neighborhood of the input being explained. Given an input x and a model $\mathcal{M}$, LIME generates perturbed instances around x and obtains the corresponding predictions from $\mathcal{M}$. These samples are weighted according to their proximity to x and used to fit a local interpretable surrogate model, commonly a sparse linear model. The coefficients of this surrogate represent the local contribution of the input features to the prediction and can be presented as feature-importance scores or visual heat maps. Because the surrogate approximates $\mathcal{M}$ only locally, its decision boundary need not coincide exactly with that of the underlying model, particularly when $\mathcal{M}$ is highly nonlinear in the neighborhood of x. Nevertheless, LIME is post-hoc, model-agnostic, instance- and prediction-specific, making it applicable to a broad range of black-box prediction services.

Counterfactual Explanations. A counterfactual explanation identifies a modified version of an input that changes the model's prediction while remaining close to the original instance according to a chosen notion of proximity [60]. A common formulation considers an input x and a model $\mathcal{M}$ and searches for a counterfactual x_{cf} such that $\mathcal{M}(x_{\text{cf}}) \neq \mathcal{M}(x)$. It does this while minimizing, or otherwise constraining, a distance $d(x, x_{\text{cf}})$. Practical counterfactual generators differ in how this objective is formulated and do not necessarily guarantee a globally minimal counterfactual. Unlike feature-attribution explanations, which characterize the contribution of features to the current prediction, counterfactual explanations identify changes that are sufficient to alter the model's decision [21]. The form of these changes depends on the input domain. For tabular data, counterfactuals typically modify the values of one or more features. For images, visual counterfactual explanations may instead remove or replace image regions, or generate semantically modified images through learned representations [25, 59].

Tabular Counterfactual Explanations. Growing Spheres [30] follows a model-agnostic search procedure that samples progressively expanding regions around the factual instance until it identifies an instance with a different prediction, after which the solution can be further sparsified. Wachter et al. [60], in contrast, formulate counterfactual generation as an optimization problem that balances achieving a desired model output against remaining close to the factual instance. NICE [8] follows an example-based approach, beginning from a nearby instance with a different prediction and iteratively replacing feature values to construct a counterfactual. These methods therefore represent different strategies for obtaining decision-changing instances, including search-based, optimization-based, and example-based approaches.

Search for EviDence Counterfactual (SEDC). SEDC builds on the document-explanation approach introduced by Martens and Provost [38], in which an explanation consists of a small set of present terms whose removal changes the predicted class. Vermeire et al. [59] subsequently adapted this approach to image classification and refer to the resulting method as *Search for EviDence Counterfactual* (SEDC). Image SEDC operates on image segments or superpixels. The image is first partitioned into regions, after which SEDC searches for a set of segments whose removal or replacement changes the predicted class. The method supports different segmentation and replacement strategies; Vermeire et al. evaluate several alternatives and use Gaussian blurring as the segment-removal mechanism in their main experiments. The identified decision-changing regions constitute the evidence counterfactual (EdC), while applying these changes to the input produces the corresponding counterfactual image. Standard SEDC searches for a change to any alternative class, whereas SEDC-T extends the approach by allowing a specific target counterfactual class to be specified. In this work, we consider untargeted SEDC, since any valid change in the predicted class provides the decision-changing counterfactual required by our attack.

Diverse Valuable Explanations (DiVE) [47]. While SEDC constructs visual counterfactuals by modifying discrete image regions, DiVE generates counterfactual images through a learned latent representation. DiVE first maps an input image into a disentangled latent space using a variational autoencoder (VAE) and then searches for perturbations of the latent representation that change the model's prediction while encouraging the resulting image to remain close to the original. The method jointly optimizes multiple latent perturbations and incorporates a diversity objective so that the resulting counterfactuals capture different decision-changing variations. DiVE additionally proposes variants that restrict highly influential latent factors to discourage trivial counterfactuals and expose alternative attributes. Thus, DiVE provides a complementary visual counterfactual mechanism to SEDC: rather than removing predefined image segments, it produces decision-changing images through modifications in a learned generative representation.

4 System and Threat Model

We consider a black-box MLaaS platform akin to commercial platforms such as Amazon Web Services (AWS) SageMaker Clarify, Google Vertex AI, Azure Machine Learning, or IBM Watson OpenScale [1, 2, 14, 52]. A target model $\mathcal{M}$ and an explainability tool $\mathcal{E}$ are provided by the MLaaS for user queries. We denote the training set as D_M, which is used by the model owner to train $\mathcal{M}$. A user can query the service with a sample x_i to obtain the corresponding predicted class $\hat{y}_i$ and its explanation z_i. Users only have access

to the predicted labels and explanations; the corresponding confidence scores and gradient information are not disclosed. We refer to this access model as an explanation-based black-box setting. We consider an API layer that enforces per-user query budgets and filters sequences of closely related perturbation queries.

The attacker has no access to the model's weights or internal parameters, but is assumed to have knowledge of the feature space, including each feature's type, valid range, and the set of possible output classes. The attacker has access to a small membership-labeled dataset $\mathcal{D}_C$ (see Section 5.3), drawn from a distribution close to that of D_M, which is only used for the threshold τ calibration. In particular, $\mathcal{D}_C$ is not used to train shadow or surrogate models. No additional auxiliary data are required after calibration. The attacker could get access to $\mathcal{D}_C$ in many ways. For example if the attacker is participating in training a shared model, they would then have a subset of the training set with high confidence labels. Another way is if a few samples from the training dataset are leaked to the public either through a targeted attack or some internal leak. Given an inference dataset $\mathcal{D}_I$, the attacker's objective is to determine, for each sample $x_i \in D_I$, whether $x_i \in D_M$.

5 XposedMIA's Overview

Our attack, XposedMIA, is a unified framework based on the boundary-distance intuition of [13]. While their method estimates distance to the target model's decision boundary through iterative perturbation queries, XposedMIA uses explanations returned by the MLaaS service (see Section 4) as a proxy for boundary information. It derives a boundary-distance score from each explanation and uses it to infer membership. Thus, XposedMIA is a **one-shot explanation-based black-box MIA framework**: after one-time calibration, each target sample requires only one external service query.

We instantiate the framework with LIME [46] and five counterfactual explainers: the tabular methods of Wachter et al. [60], NICE [8], and Growing Spheres [30]; and the visual methods SEDC [59] and DiVE [47]. The framework is shared across all methods; only the conversion of explanations into boundary-distance approximations differs. Calibration and membership decisions remain unchanged.

The attacker neither trains shadow models nor obtains the target model's parameters, gradients, logits, confidence scores, or internal representations. They observe only the model's output and explanation returned by the MLaaS provider. Requiring one query per target after calibration helps evade automatic detection systems and MLaaS rate limits. Consequently, XposedMIA operates through a label-plus-explanation interface: the prediction remains label-only, while the MLaaS response contains both the label and explanation.

5.1 Attack Intuition

Let $\mathcal{M}$ denote the target classifier and let $\partial\mathcal{M}$ denote its decision boundary. For an explanation method $\mathcal{E}$, the MLaaS explanation service $\mathcal{S}_\mathcal{E}$ returns a response $z_i \leftarrow \mathcal{S}_\mathcal{E}(x_i)$ containing the predicted label $\hat{y}_i$ and the explanation produced for x_i. Depending on $\mathcal{E}$, the explanation may consist of local surrogate parameters, a counterfactual sample, or a counterfactual image and its associated mask. XposedMIA applies an explanation-specific extractor $g_{\mathcal{E},q}$ to convert this response into a scalar score:

$$d_i^{(\mathcal{E},q)} = g_{\mathcal{E},q}(x_i, z_i) \approx \mathrm{dist}\big(x_i, \partial\mathcal{M}\big),$$

where q denotes the distance measure used by the extractor. We refer to $d_i^{(\mathcal{E},q)}$ as a *boundary-distance estimate*. For LIME, it is the geometric distance to a locally fitted surrogate boundary. For counterfactual explanations, it is the amount of change between the factual and counterfactual samples. A valid counterfactual lies on the other side of a decision boundary; therefore, the change required to reach it estimates the distance from the factual sample to that boundary. This is generally a proxy rather than the exact minimum distance because an explanation method is not guaranteed to recover the globally closest boundary point.

More generally, a predicted label alone tells in which region the sample x_i is located, while explanations expose additional geometric information in different forms. Thus XposedMIA does not treat the explanation artifact itself as a membership signal, but instead uses the geometric information exposed and transforms it to a scalar boundary distance that can then be used to distinguish between members and non-members. In this sense, explanations provide a shortcut to information about the proximity of a sample to the decision boundary, which would normally require many queries from a prediction-only attacker. We also note that our attack can therefore be affected by the quality of the generated explanations, as incorrect or inaccurate explanations can cause the attack's performance to degrade.

Prior label-only MIAs commonly assume that members lie farther from the decision boundary than non-members [13]. We do not impose this direction universally. Depending on the target model, dataset, explanation generator, and distance measure, members may instead produce smaller estimates. XposedMIA therefore calibrates both the inequality direction and the corresponding decision threshold using a small set of samples with known membership.

5.2 Estimating the Boundary Distance

The attack framework is initially designed for LIME, but is generalized to work with all explanation methods we cover in this work. The only explanation-specific component is the distance scalar $d_i^{(\mathcal{E},q)}$. Once this is obtained, the attack is explanation-agnostic.

LIME. LIME fits a locally weighted interpretable model around the explained sample to approximate the behavior of the target model [46]. Consider a tabular sample $x_i \in \mathbb{R}^n$ and let $\phi(x_i)$ denote the scaled feature representation used by LIME. The explanation for the predicted class consists of a coefficient vector $e_i \in \mathbb{R}^n$ and an intercept b_i, defining the local surrogate as

$$f_i(u) = e_i^\mathsf{T} u + b_i.$$

We use the hyperplane $f_i(u) = t$ as an operational estimate of the local one-versus-rest decision boundary, with $t = 0.5$. The resulting LIME boundary-distance estimate is

$$d_i^{\mathrm{LIME}} = \frac{\big|f_i\big(\phi(x_i)\big) - t\big|}{\|e_i\|_2} = \frac{\big|e_i^\mathsf{T}\phi(x_i) + b_i - 0.5\big|}{\|e_i\|_2}.$$

Thus, LIME provides an analytical estimate through its local surrogate; XposedMIA does not generate adversarial examples or issue additional attacker-side perturbation queries.

Tabular counterfactual explanations. Given a sample x_i, its counterfactual explanation is a sample x_i^{cf} whose predicted class

differs from the original factual sample as:

$$\mathcal{M}(x_i^{\text{cf}}) = \hat{y_i}, \qquad \mathcal{M}(x_i) = y_i, \qquad \hat{y_i} \neq y_i.$$

Given a valid counterfactual, we estimate the boundary distance as $d_i^{(\mathcal{E},q)} = q(x_i, x_i^{\text{cf}})$ where $\mathcal{E} \in \{\text{Wachter, NICE, GS}\}$. While this measures the distance to the returned counterfactual rather than directly to the decision boundary, the counterfactual is expected to lie close to the boundary, as it crosses the decision boundary by construction under the assumptions of counterfactual XAI methods. It therefore can be considered a proxy for the distance to the decision boundary.

We evaluate three tabular counterfactual generators. Wachter et al. [60] conduct continuous optimization using a target-output loss and an inverse-median-absolute-deviation-weighted L_1 proximity term. It searches over candidate alternative classes and returns the closest counterfactual. NICE [8] starts from the nearest unlike instance and optimizes the explanation for sparsity, using a reference pool to construct the counterfactual. Growing Spheres [30] samples around the factual sample using an increasingly growing sphere. It keeps iterating until it finds an instance that has a different class than the factual sample.

After obtaining the counterfactual, the attacker computes the distance locally without querying the model $\mathcal{M}$. We measure the distance between the factual and counterfactual samples using standard L_p distances and data-aware measures, including Euclidean, Manhattan, Chebyshev, inverse median absolute deviation (MAD) L_1 [60], Mahalanobis, cosine, Canberra, Bray–Curtis, and Gower [24] distances.

SEDC visual counterfactuals. SEDC constructs a counterfactual by identifying a small set of image segments whose removal changes the predicted class [59]. Our implementation reproduces one untargeted counterfactual for every request. We first partition the image using a grid (for the MNIST dataset) or Simple Linear Iterative Clustering (SLIC) [3] (for CIFAR-10 and CIFAR-100). Removing a segment is instantiated by replacing it with a Gaussian-blurred version of itself. It first performs a bounded best-first search over segment subsets and then applies a deterministic reduction step while removing unnecessary segments and maintaining the changed label. We also bound the search by a segment count budget.

The explainer returns the counterfactual image x_i^{cf}, the selected pixel mask, the segment map, and the selected segment identifiers. The attacker does not get access to any of the model's internal parameters that were used during the explanation generation. The boundary distance is then calculated as follows:

$$d_{\text{seg}} = k, \qquad d_{\text{segfrac}} = \frac{k}{K}, \qquad d_{\text{area}} = \frac{|A|}{HW},$$

where k is the number of selected segments, K is the total number of segments, A is the selected pixel mask, and H and W are the image dimensions. Note that d_{segfrac} and d_{area} generally perform better than seg since they are normalized.

DiVE visual counterfactuals. DiVE generates diverse visual counterfactuals by optimizing perturbations in the latent space of a generative model [47]. DiVE is targeted by nature since its counterfactual loss explicitly optimizes the generated image toward a specific alternative class. For a fixed target class y_t, the provider generates a set of candidates satisfying

$$\mathcal{M}(x_{i,j}^{\text{cf}}) = y_t \neq \mathcal{M}(x_i).$$

Among the candidates, the explainer selects one counterfactual sample using DiVE's internal objective that optimizes proximity. This sample is sent to the attacker along with the target class.

Image distance measures. SEDC and DiVE share the same pixel- and structure-based distances. Let $\Delta_i = |x_i - x_i^{\text{cf}}|$ for images represented in $[0, 1]$. We compute

$$d_0 = \frac{1}{HW} \sum_{p=1}^{HW} \mathbf{1}\left[\max_c \Delta_{i,p,c} > \epsilon\right], \tag{1}$$

$$d_1 = \text{mean}(\Delta_i), \tag{2}$$

$$d_2 = \sqrt{\text{mean}(\Delta_i^2)}, \tag{3}$$

$$d_\infty = \max(\Delta_i), \tag{4}$$

$$d_{\text{SSIM}} = 1 - \text{SSIM}(x_i, x_i^{\text{cf}}), \tag{5}$$

where ϵ prevents negligible numerical differences from being counted as changed pixels and structural similarity (SSIM) measures structural similarity [62]. SEDC additionally supports the segment- and mask-based measures defined above. As in the tabular setting, every metric for a sample is computed from the same returned counterfactual.

For DiVE, we additionally evaluate Learned Perceptual Image Patch Similarity (LPIPS) [66], a perceptual distance based on deep feature representations:

$$d_{\text{LPIPS}} = \text{LPIPS}(x_i, x_i^{\text{cf}}).$$

5.3 Distance Threshold Calibration (DTC)

This constitutes **Step I** of XposedMIA as shown in Figure 1 and summarized in Algorithm A.1 in Appendix A. Prior work suggests that the training members are expected to exhibit a larger distance to the decision boundary compared to non-members [13]. We observe, however, that this assumption does not necessarily hold across all datasets. In certain cases, non-members may exhibit a larger distance. This may be caused by differences in the learned boundary geometry, the distribution and difficulty of member and non-member samples, and the approximation objective of the explanation method. To account for this variability, we use **a two-stage calibration procedure** that adapts to the relative distance distributions of members and non-members. We assume that the attacker has access to a small calibration dataset $\mathcal{D}_C = \{(x_j, m_j)\}$, where $m_j \in \{0, 1\}$ is the known membership status of x_j. We refer to Section 6.2 for details on how the size of $\mathcal{D}_C$ affects the results.

For an explanation method $\mathcal{E}$ and a distance measure q, the attacker queries each calibration sample and constructs

$$\mathcal{B}_q = \left\{\left(d_j^{(\mathcal{E},q)}, m_j\right)\right\},$$

using the distance estimates returned by the explainer. Each observed distance is a candidate threshold. For every candidate τ_c and direction $dir \in \{\texttt{leq}, \texttt{gt}\}$, it computes the attack accuracy on the calibration dataset to evaluate the effectiveness of that threshold. The direction `leq` denotes distances less than or equal to the threshold, while `gt` denotes distances greater than the threshold.

The sample-weighted overall accuracy is computed as follows:

$$acc_C = \frac{N_{\text{mem}} acc_{\text{mem}} + N_{\text{non}} acc_{\text{non}}}{N_{\text{mem}} + N_{\text{non}}}.$$

This first stage selects the direction dir_q with the highest calibration accuracy acc_C. The direction determines whether samples closer to the decision boundary should be considered members or non-members. After this direction is determined from the first stage, the second stage is then executed:

$$\tau_q = \arg\min_{\tau_c} \left| acc_{\text{mem}}(\tau_c, dir_q) - acc_{\text{non}}(\tau_c, dir_q) \right|.$$

Thus, the threshold equalizes the accuracy for members and non-members. Since several distance measures are available, DTC is executed on each q, and the measure with the highest calibration accuracy is selected as the best attack measure.

5.4 Membership Inference

This constitutes **Step II** of XposedMIA (as shown in Figure 1). Given an unseen target sample x_i, the attacker sends one request to $\mathcal{S}_\mathcal{E}$ and obtains its predicted label and explanation. The attacker then computes $d_i = g_{\mathcal{E},q^*}(x_i, z_i)$ without further access to the target model. The calibrated decision rule is

$$\widehat{m}_i = \begin{cases} 1, & dir^* = \texttt{leq} \text{ and } d_i \le \tau, \\ 1, & dir^* = \texttt{gt} \text{ and } d_i > \tau, \\ 0, & \text{otherwise}, \end{cases}$$

where $\widehat{m}_i = 1$ denotes a member and $\widehat{m}_i = 0$ denotes a non-member.

Algorithm 1 formalizes our complete end-to-end MIA pipeline. The algorithm is structurally divided into two main phases. The first phase (Lines 1–14) constructs the calibration records $\mathcal{B}_q$, applies DTC independently to each candidate distance measure, and selects the finalized measure q^*, threshold τ, and direction dir^*, as detailed in Section 5.3. The second phase performs the attack on the inference dataset $\mathcal{D}_I$ (Lines 15–29). For each unseen target sample, the pipeline requests an explanation and computes its boundary-distance estimate using $g_{\mathcal{E},q^*}$. It then applies the calibrated decision rule by comparing the distance d with τ according to dir^*. Invalid explanation responses are processed using the method-specific failure rule. Finally, the algorithm returns a mapping $\mathcal{R}$ that associates each target sample with its predicted membership status or, where applicable, an abstention.

Algorithm 1: Unified explanation-distance membership inference pipeline.

Input: Inference dataset $\mathcal{D}_I$; calibration dataset $\mathcal{D}_C$ containing tuples (x, m); explanation service $\mathcal{S}_\mathcal{E}$; candidate distance measures $\mathcal{Q}_\mathcal{E}$
Output: Membership decisions, with abstentions for unsupported image counterfactual responses

1 Initialize $\mathcal{B}_q \leftarrow \emptyset$ for every $q \in \mathcal{Q}_\mathcal{E}$;
2 **foreach** $(x, m) \in \mathcal{D}_C$ **do**
3 $z \leftarrow \mathcal{S}_\mathcal{E}(x)$;
4 **if** *z contains a valid explanation* **then**
5 **foreach** $q \in \mathcal{Q}_\mathcal{E}$ **do**
6 $d \leftarrow g_{\mathcal{E},q}(x, z)$;
7 $\mathcal{B}_q \leftarrow \mathcal{B}_q \cup \{(d, m)\}$;
8 **else**
9 Record an observable explanation failure with membership label m;
10 **foreach** $q \in \mathcal{Q}_\mathcal{E}$ **do**
11 $(\tau_q, dir_q) \leftarrow \text{DTC}(\mathcal{B}_q)$;
12 Select q^* using calibration data only;
13 $\tau \leftarrow \tau_{q^*}$;
14 $dir^* \leftarrow dir_{q^*}$;
15 Initialize $\mathcal{R} \leftarrow \emptyset$;
16 **foreach** $x \in \mathcal{D}_I$ **do**
17 $z \leftarrow \mathcal{S}_\mathcal{E}(x)$;
18 **if** *z contains a valid explanation* **then**
19 $d \leftarrow g_{\mathcal{E},q^*}(x, z)$;
20 $c_1 \leftarrow (dir^* = \texttt{leq}) \wedge (d \le \tau)$;
21 $c_2 \leftarrow (dir^* = \texttt{gt}) \wedge (d > \tau)$;
22 **if** $c_1 \vee c_2$ **then**
23 $\widehat{m} \leftarrow 1$;
24 **else**
25 $\widehat{m} \leftarrow 0$;
26 **else**
27 $\widehat{m} \leftarrow \text{HandleFailure}(\mathcal{E})$;
28 $\mathcal{R} \leftarrow \mathcal{R} \cup \{(x, \widehat{m})\}$;
29 **return** $\mathcal{R}$;

6 Results

In this section, we present our experimental results. Section 6.1 describes the experimental setup. Section 6.2 reports the results using LIME explanations and the corresponding calibration and distance analysis. Section 6.3 presents the counterfactual explanation results across three generators. Section 6.4 reports the extension to visual counterfactual explanations. Finally, Section 6.5 compares XposedMIA with prior work.

6.1 Experimental Setup

In our experiments, we evaluate XposedMIA using both LIME [46] and counterfactual explanations. We use the LIME tabular explainer. We additionally evaluate three counterfactual generators: NICE [8], Wachter [60], Growing Spheres [30]. For a returned counterfactual x', the attack uses its Euclidean displacement $\|x - x'\|$ from the queried sample x as its membership signal. For a counterfactual request, the response contains the predicted label together with either a valid counterfactual x' or an explicit generation failure. Operations required to construct an explanation are performed internally by the MLaaS service and therefore do not increase the attacker's external query count. Consequently, after calibration, XposedMIA requires only one explanation request per target sample, regardless of the explanation mechanism used by the service. More details on the setup of the explainers can be found in Appendix E.1. Details about dataset splits can be found in Table 5 in Appnedix E.2.

Table 1: Statistics of the datasets used in the LIME attack before preprocessing.

Dataset	Dataset Size	# Features
Purchase-100	197,324	600
Texas-100	67,330	6,170
Location	5,010	446
Adult	48,842	14

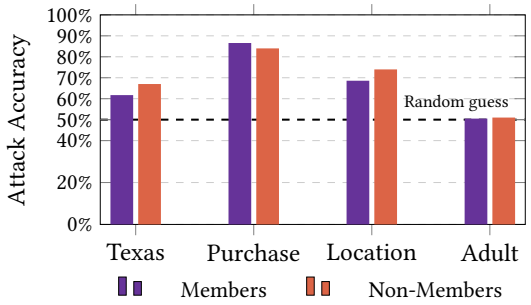

Figure 2: Attack accuracy for members and non-members for Location, Purchase-100, Texas-100, and Adult datasets.

We evaluate XposedMIA on four benchmark datasets introduced by Shokri et al. [54] and widely adopted in subsequent literature: Purchase-100,[1] Location,[1] Texas-100,[1] and Adult. [2] For consistency, we instantiate the target model as a multi-layer perceptron (MLP) and keep the training pipeline fixed except for dataset-dependent choices. Each target MLP uses fully connected layers with 128 neurons and the tanh activation function. We vary the number of hidden layers from 1 to 4 depending on the dataset complexity. We train the models using the Adam optimizer and categorical cross-entropy loss, together with L_2 regularization and learning-rate decay. This setup has been used in Choquette-Choo et al. [13] following the work of Shokri et al. [54]. Table 1 summarizes the dataset statistics.

For the extension of XposedMIA to images in Section 6.4, we use CIFAR-10 [3] and CIFAR-100.[3] For CIFAR-10, the target model is a custom three-block convolutional neural network (CNN) with two 3×3 convolutions per block. For CIFAR-100, we use a RestNet-18 with four residual stages of widths 64, 128, 256, and 512. DiVE uses a BigGAN-inspired β-total correlation variational autoencoder (β-TCVAE) as a generative model for the explainer with a 128-dimensional latent space: width/depth of 4/2 for CIFAR-10 and 6/3 for CIFAR-100.

We evaluate the attack using standard metrics: Area Under the Receiver Operating Characteristic curve (ROC–AUC) and attack success rate (ASR) computed as $ASR = \frac{TP+TN}{TP+TN+FP+FN}$, where TP, FP, TN, FN denote true/false positives and negatives.

6.2 LIME Explanations

We evaluate XposedMIA across multiple datasets and target-model configurations. Figure 2 reports ASRs of 0.85 on Purchase-100, 0.71 on Location, and 0.64 on Texas-100. XposedMIA performs best on Purchase-100 and Location, surpasses the baseline accuracy on Texas-100, and requires only one query after calibration. These results are consistent with the dataset and model characteristics. By contrast, the Adult model leaks less membership information, plausibly because of its smaller feature space and lack of a train-test gap. Adult has only 14 features, compared with 600, 446, and 6,170 for Purchase-100, Location, and Texas-100, respectively. In our experiments, this limited feature space improved generalization and reduced ASR, consistent with prior work [13, 49, 53]. We therefore did not pursue further ablations on Adult. The remainder of this section examines how calibration-set size affects XposedMIA and how sample distances to the LIME-estimated decision boundary are distributed across datasets.

Calibration dataset ($\mathcal{D}_C$) size. To evaluate the optimal calibration-dataset ($\mathcal{D}_C$) size, we conduct ablations on Purchase-100 and Location, where our method performs best. We use a fixed inference dataset ($\mathcal{D}_I$) and vary the size of $\mathcal{D}_C$ to measure its effect on attack performance.

Figures 3a and 3b show the results for Purchase-100 and Location, respectively. In Figure 3a, calibration sets of 20 or fewer samples cause member and non-member ASRs to diverge markedly. With insufficient samples, threshold selection becomes unstable, placing the estimated distance threshold (τ) either too close to zero or too far into the distribution's tail (see Figure 4a).

With only 2 samples in $\mathcal{D}_C$, for example, the algorithm initially sets $\tau = 10.76$. This yields high member ASRs but poor non-member ASRs. The inferred rule assumes that members have smaller distances than non-members, highlighting the need to estimate both the threshold and inequality direction accurately. As $\mathcal{D}_C$ grows, the threshold decreases to $\tau = 10.33$, $\tau = 6.72$, and $\tau = 5.72$ for 10, 20, and 30 samples, respectively. At approximately 30 samples, both the threshold and direction stabilize, producing reliable ASR curves.

Distance to the boundary. We also analyze the distance distributions from data samples to the local surrogate boundary estimated by LIME via XposedMIA. Figures 4 and 5 visualize these distances for the Purchase dataset, illustrating how the distance metric distinguishes members from non-members. Ultimately, ASR is driven by the separation between these distributions: greater separation minimizes classification errors, while substantial overlap makes threshold-based inference unreliable.

Conversely, Figure 4b illustrates a more generalized target model where member and non-member distributions are nearly indistinguishable. Due to this severe overlap, no effective distance threshold can be established to separate the two groups.

Similar patterns across datasets support the same interpretation. Our target model trained on the Location dataset would create two different distance distributions (see Figure 5a). For the Adult dataset, however, the distribution for members and non-members shown in Figure 5b is not practically separable, as the model generalizes well. This is consistent with the lowered ASR observed on Adult both in our experiments and in prior MIA studies [13, 49, 53].

[1] https://github.com/privacytrustlab/datasets
[2] http://archive.ics.uci.edu/ml/datasets/Adult
[3] https://www.cs.toronto.edu/~kriz/cifar.html

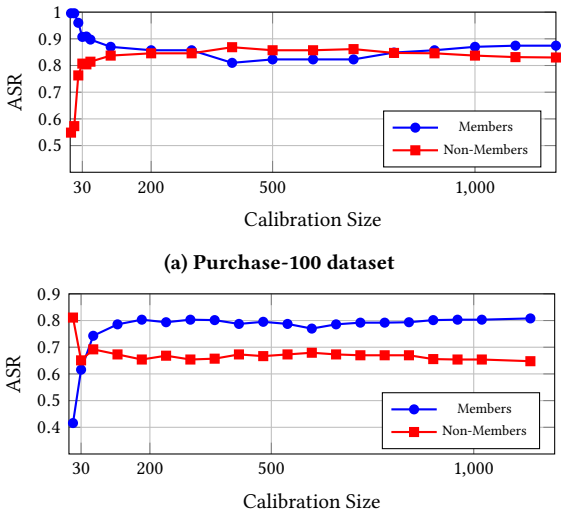

(a) Purchase-100 dataset

(b) Location dataset

Figure 3: Comparison of the attack accuracies on members and non-members for different calibration sizes on the (a) Purchase-100 and (b) Location datasets.

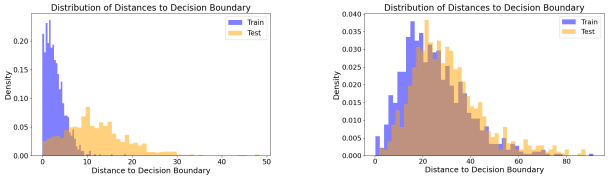

(a) Less generalized model. (b) Well generalized model.

Figure 4: Distance distribution of instances to the LIME-fitted local model based on the model generalization on Purchase-100 dataset.

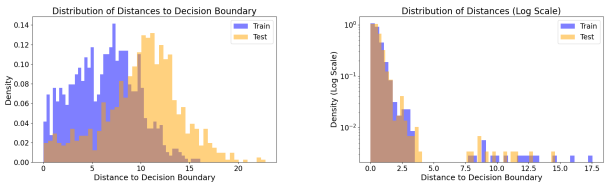

(a) Location dataset (b) Adult dataset

Figure 5: Distance distribution of instances to the LIME-fitted local model on the (a) Location dataset and (b) Adult dataset.

6.3 Counterfactual Explanations

In this section, we evaluate whether the calibrated (τ, dir) decision rule that XposedMIA learns for LIME transfers, without modification, to a structurally different family of explanations. Counterfactual explanations provide a direct model-space displacement: a valid counterfactual x' changes the target model's prediction, allowing us to measure the generator-selected displacement $\|x - x'\|$ without estimating a local surrogate boundary.

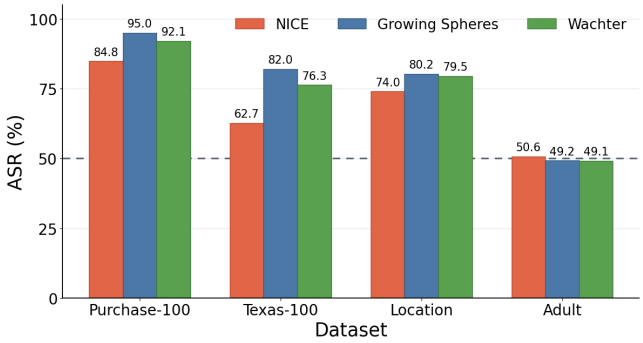

Figure 6: ASR across datasets for XposedMIA using the NICE, Growing Spheres, and Wachter counterfactuals. The dashed line denotes random guessing.

NICE. With NICE explanations, our attack achieves balanced ASRs of 84.8%, 74.0%, and 62.7% on Purchase-100, Location, and Texas-100, respectively, as reported in Figure 6. Figure 13a in Appendix C.1 shows the member and non-member counterfactual-distance distributions on Location, with a calibrated threshold of $\tau = 8.18$. NICE achieves complete counterfactual coverage on all evaluated datasets, so generation failures do not contribute to its membership signal. However, the weaker performance on Texas-100 shows that complete coverage is insufficient when counterfactual distances do not strongly separate members from non-members. Because NICE uses nearest unlike reference instances, this effect may be amplified in the high-dimensional Texas-100 feature space, where reference-set proximity may not closely track distance to the model boundary.

Wachter. With Wachter counterfactuals, our attack achieves balanced ASRs of 92.1%, 79.5%, and 76.3% on Purchase-100, Location, and Texas-100, respectively, as reported in Figure 6. Figure 13b in Appendix C.1 shows the corresponding member and non-member counterfactual-distance distributions on Purchase-100, with $\tau = 1.94$. On Adult, we use a stronger penalty schedule to increase counterfactual coverage, achieving 100.0% member/non-member coverage but only 49.4% ROC-AUC. Thus, reliable counterfactual generation is necessary but insufficient for effective membership separation. This generator-side configuration does not change the attacker's access model: the attacker still receives one counterfactual response and applies the same calibrated (τ, dir) decision rule.

Growing Spheres. The attack on Growing Spheres achieves balanced ASRs of 95.0%, 82.0%, and 80.16% on Purchase-100, Texas-100, and Location, respectively. Figure 6 and 13c, Appendix C.1 illustrate the ASR and member and non-member counterfactual-distance distributions on Texas-100, for which $\tau = 23.65$. For Texas-100, we therefore increase the search radius to recover valid counterfactuals over a broader region of the feature space.

Cross-generator analysis. Table 2 shows that counterfactual coverage alone does not determine attack effectiveness. NICE achieves 100% coverage on every dataset, yet Growing Spheres often yields

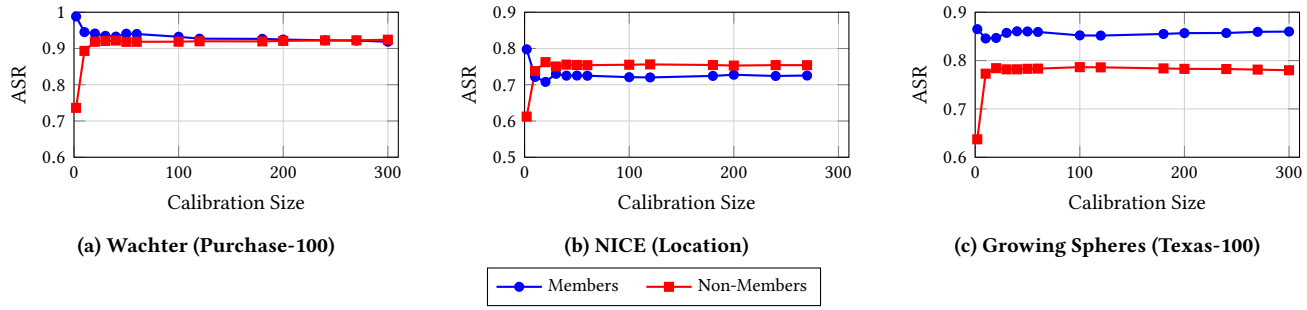

(a) Wachter (Purchase-100) (b) NICE (Location) (c) Growing Spheres (Texas-100)

Figure 7: Calibration-size sensitivity for three counterfactual generators. The panels report member and non-member attack success rates for representative Wachter, NICE, and Growing Spheres runs on Purchase-100, Location, and Texas-100, respectively. Each point is the mean over 30 nested calibration-subset resamples with a fixed inference set.

Table 2: Counterfactual-distance results under the calibrated DTC rule. CF success reports valid label-changing counterfactuals for members (M) and non-members (N); each attack uses one generator request per target.

Dataset	Generator	ASR	ROC–AUC	CF Success	
				M	N
Purchase-100	NICE	84.8%	91.6%	100.0%	100.0%
	Growing Spheres	95.0%	97.6%	88.7%	98.9%
	Wachter	92.1%	96.7%	99.6%	100.0%
Texas-100	NICE	62.7%	67.6%	100.0%	100.0%
	Growing Spheres	82.0%	88.1%	78.4%	95.9%
	Wachter	76.3%	79.8%	85.9%	98.3%
Location	NICE	74.0%	81.2%	100.0%	100.0%
	Growing Spheres	80.2%	87.1%	81.6%	93.2%
	Wachter	79.5%	84.7%	99.7%	99.8%
Adult	NICE	50.6%	52.1%	100.0%	100.0%
	Growing Spheres	49.2%	49.1%	98.9%	99.1%
	Wachter	49.1%	49.4%	100.0%	100.0%

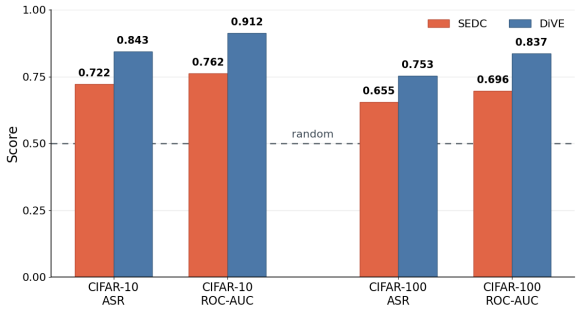

Figure 8: Results using the SEDC and DiVE counterfactual explanations on CIFAR-10 and CIFAR-100. Both ASR and ROC–AUC are reported. The results present substantial membership information leakage above the random baseline.

stronger membership separation despite lower coverage. Thus, MIA effectiveness depends not only on valid counterfactual generation but also on how strongly the induced distance distributions separate members from non-members. Because the generators construct counterfactuals differently, they expose different aspects of the decision geometry: NICE is shaped by its reference data, Wachter by its optimization objective, and Growing Spheres by its radial search.

The Adult results reinforce this point: despite complete coverage in most settings, all three generators remain close to random guessing when member and non-member distances are insufficiently separated. Conversely, in most incomplete-coverage settings, valid counterfactuals are returned less often for members than non-members, consistent with members lying farther from the decision boundary and requiring larger displacements within a finite search budget. Generation failure can therefore provide an additional membership signal alongside the distance of successfully returned counterfactuals.

Calibration-size ablation. Figure 7 repeats the calibration analysis of Section 6.2 for counterfactual distances, resampling only $\mathcal{D}_C$ while keeping the target model, inference records, and generator responses fixed; each point averages 30 nested resamples. As in LIME, very small calibration sets produce unstable estimates of both the threshold τ and the inequality direction, which appears as a pronounced imbalance between member and non-member ASR. For example, with only two calibration records per membership class, Wachter on Purchase-100 achieves mean member and non-member ASRs of 98.8% and 73.6.%, respectively, with $\tau = 1.18 \pm 0.52$. Increasing the calibration size to 10, 20, and 30 records from each membership class yields thresholds of $\tau = 1.72 \pm 0.37$, $\tau = 1.88 \pm 0.18$, and $\tau = 1.90 \pm 0.16$, respectively, while the corresponding member/non-member ASRs become 94.5%/89.3%, 94.1%/91.9%, and 93.5%/92.1.%. At 30 records per membership class, the balanced ASR reaches 92.8%, already close to the full-calibration value of 92.1%, with the threshold approaching its full-calibration value of $\tau = 1.94$. The same stabilization pattern is observed for NICE and Growing Spheres. Across all generators, roughly 20–30 calibration records per class are sufficient for the DTC rule to reach a stable operating point, after which additional data produces only marginal changes. Together with the LIME results in Section 6.2, this suggests that the small calibration requirement is primarily a property of the DTC procedure rather than of the underlying explanation mechanism. We also perform an ablation to study the effect of changing the second stage of DTC with an objective that maximizes balanced

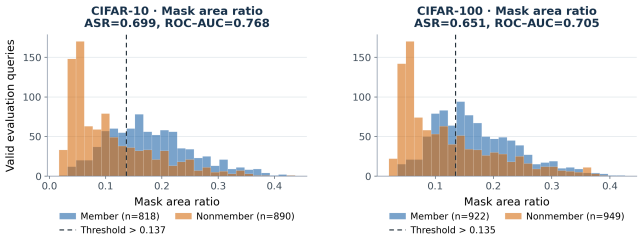

Figure 9: CIFAR-10 and CIFAR-100 mask area ratio distances and the chosen decision threshold for the SEDC attack.

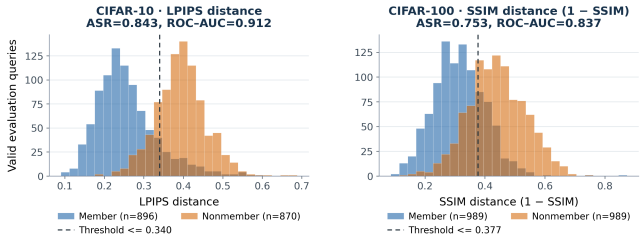

Figure 10: CIFAR-10 Learned Perceptual Image Patch Similarity (LPIPS) and CIFAR-100 SSIM distances and the chosen decision threshold for the DiVE attack.

accuracy instead of minimizing the difference between member and nonmember accuracy. Our results are summarized in Table 3, Appendix C.2.

6.4 Extension to Images

To evaluate whether XposedMIA generalizes to more complex datasets, we extend our experiments to image classifiers. We use CIFAR-10 and CIFAR-100 since they are among the main benchmark datasets for most MIA works [10, 13, 31, 40, 49]. We target two visual counterfactual explanation methods: SEDC [59] and DiVE [47]. We select these two methods to introduce diversity in counterfactual generation and demonstrate the generalizability of our approach. SEDC changes the prediction by removing image segments, while DiVE utilizes a generative model (VAE) and searches in its latent space for counterfactuals. Example SEDC and DiVE counterfactuals can be seen in Appendix B. It is important to note that the attack itself remains unchanged once the distance scores are obtained; we use the same DTC and MIA framework described in Section 5.

Relevance to privacy-sensitive applications. CIFAR provides a controlled benchmark for testing whether visual counterfactuals expose membership-related signals. The same interface can arise in privacy-sensitive image-classification services, such as medical imaging.

SEDC. Figure 8 demonstrates the balanced ASR using SEDC counterfactuals and the ROC–AUC results for the CIFAR-10 and CIFAR-100 datasets. XposedMIA using SEDC reports 0.722 and 0.655 ASR for CIFAR-10 and CIFAR-100, respectively. The reported results show membership information leakage.

Figure 9 visualizes the distance distributions for members and non-members. The DTC-derived threshold from Section 5.3 is marked by a vertical line. The members are skewed to the right of the plot since they have higher distances, while non-members are skewed to the left. The members that are on the left of the line are the false negatives, while the non-members that are on the right of the threshold line are the false positives.

DiVE. DiVE uses a generative model to generate counterfactuals. Since DiVE is targeted by nature, it finds a set of counterfactuals in the latent space of the generative model that would be classified as the target class. We then pick the counterfactual that has the smallest proximity score $s_j = \lambda_{\text{prox}} R_j + \lambda_{\text{latent}} L_j$, where R_j is the sum of the per-channel mean absolute pixel differences between the factual image and candidate j, and $L_j = \|\epsilon_j\|_1$ measures the magnitude of its latent-space perturbation.

Figure 8 demonstrates the results for CIFAR-10 and CIFAR-100. We observe the best results using SSIM distance, with a balanced accuracy of 0.843 and ROC–AUC of 0.912 for CIFAR-10 and a balanced accuracy of 0.753 and ROC–AUC of 0.837 for CIFAR-100.

Figure 10 presents clear separation between members and non-members in terms of their distance to their counterfactuals. Yet, there are some false positives and false negatives, as detailed for SEDC. Since the ASR is higher for DiVE explanations, we observe a cleaner and more separated distribution compared to the SEDC.

6.5 Comparison to Other Works

We compare XposedMIA with seven representative MIA approaches: (i) Gap Attack [65], (ii) Choquette-Choo et al. [13], (iii) Shokri et al. [53], (iv) Ye et al. [64], (v) YOQO [63], (vi) DHAttack [31], and (vii) LoMime [42]. We report the ASR and ROC–AUC across the Purchase-100, Location, and Texas-100 datasets in Table 4, Appendix D. Compared with the Gap Attack [65] and prior explanation-variance attacks [53], XposedMIA achieves stronger results while using only a single query after calibration and requiring no shadow-model training. Unlike shadow-model-based methods, XposedMIA uses its auxiliary data only for threshold calibration rather than for training surrogate models, which substantially reduces the amount of data and computational cost for the attack. Furthermore, compared with perturbation-based boundary-distance attacks such as [13], XposedMIA substantially reduces the attack-time query budget while still achieving strong and comparable attack performance.

Performance under constrained query budgets. Table 4 reports the query budgets associated with prior attacks' reported performance. Choquette-Choo et al. [13] require approximately 2,500 queries per candidate for their boundary-distance attack to match their strongest confidence-based baseline, and approximately 12,500 for their combined attack. With < 100 queries, the improvement over the gap attack is only approximately four percentage points, while below 300 queries the simpler noise-robustness attack performs better because the boundary estimate remains imprecise. LoMime [42] likewise degrades under a reduced one-time model-extraction budget: on Purchase-100, AUC increases from approximately 0.57 at 10^4 queries to 0.87 at 10^6, and on Texas-100 from approximately

0.52 to 0.81 over the same budgets. Thus, the performance of query-intensive attacks can depend substantially on the available budget. This does not apply uniformly to all prior work: DHAttack [31] and YOQO [63] reduce target-model queries but rely on additional reference/shadow models or auxiliary computation.

Future Directions. As future work, we will focus on a more comprehensive set of explanation tools, deeper neural network architectures, and extend our analysis to text data, their corresponding explanations, and LLMs. A more detailed analysis of image XAI leakage across datasets and explainers is another promising direction for future research. Finally we plan on building a unified framework that combines attacks for tabular, image, and text data.

7 Conclusion

In this paper, we presented XposedMIA, a lightweight, **one-shot explanation-based black-box MIA framework**. We instantiate this framework with LIME and extend it to five counterfactual explanation methods spanning tabular and image domains. XposedMIA either exploits the geometry of a local linear surrogate to approximate the decision boundary or uses the displacement between a factual sample and its returned label-changing counterfactual as a proxy for distance to the decision boundary.

Our evaluation across these structurally different explanation mechanisms demonstrates that the boundary-related information exposed through explanations can consistently reveal membership information. Moreover, the effectiveness of XposedMIA across both tabular and image domains shows that this privacy leakage is not limited to a particular explanation method or data modality, making it relevant to many MLaaS platforms that release XAI explanations.

Acknowledgments

This work was supported in part by the National Science Foundation (NSF) under Grants OAC-2112606, 2427505, and 2141622; the National Institutes of Health (NIH) under Grant R01LM014520; and the U.S. Department of Agriculture (USDA) under Grant 2026-67023-46766.

References

[1] 2023. IBM Cloud, Watson Openscale. [Accessed 24-Feb-2024]. https://www.ibm.com/docs/en/cloud-paks/cp-data/4.8.x

[2] 2024. Microsoft Azure, Model interpretability. [Accessed 01-Mar-2024]. https://learn.microsoft.com/en-us/azure/machine-learning/how-to-machine-learning-interpretability

[3] Radhakrishna Achanta, Appu Shaji, Kevin Smith, Aurelien Lucchi, Pascal Fua, and Sabine Süsstrunk. 2012. SLIC Superpixels Compared to State-of-the-Art Superpixel Methods. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 34, 11 (2012), 2274–2282. doi:10.1109/TPAMI.2012.120

[4] Md. Rasel Ahmed, Md Mujahidul Islam, Faysal Ahmed, and Mohammad Anowarul Kabir. 2024. A Systematic Literature Review Of Machine Learning Adoption In Emerging Marketing Applications. *Journal of Machine Learning, Data Engineering and Data Science* 1, 01 (Nov. 2024), 163–180.

[5] Ulrich Aïvodji, Alexandre Bolot, and Sébastien Gambs. 2020. Model extraction from counterfactual explanations. *CoRR* abs/2009.01884 (2020). arXiv:2009.01884

[6] Julia Amann, Alessandro Blasimme, Effy Vayena, Dietmar Frey, Vince I. Madai, and the Precise4Q consortium. 2020. Explainability for artificial intelligence in healthcare: a multidisciplinary perspective. *BMC Medical Informatics and Decision Making* 20, 1 (Nov. 2020), 310. doi:10.1186/s12911-020-01332-6

[7] Sebastian Bach, Alexander Binder, Grégoire Montavon, Frederick Klauschen, Klaus-Robert Müller, and Wojciech Samek. 2015. On Pixel-Wise Explanations for Non-Linear Classifier Decisions by Layer-Wise Relevance Propagation. *PLOS ONE* 10, 7 (07 2015), 1–46. doi:10.1371/journal.pone.0130140

[8] Dieter Brughmans, Pieter Leyman, and David Martens. 2024. NICE: an algorithm for nearest instance counterfactual explanations. *Data Mining and Knowledge Discovery* 38, 5 (Sept. 2024), 2665–2703. doi:10.1007/s10618-023-00930-y

[9] Niklas Bussmann, Paolo Giudici, Dimitri Marinelli, and Jochen Papenbrock. 2021. Explainable Machine Learning in Credit Risk Management. *Computational Economics* 57, 1 (Jan. 2021), 203–216. doi:10.1007/s10614-020-10042-0

[10] Nicholas Carlini, Steve Chien, Milad Nasr, Shuang Song, Andreas Terzis, and Florian Tramèr. 2022. Membership Inference Attacks From First Principles. In *2022 IEEE S&P*. San Francisco, CA, USA, 1897–1914.

[11] Varun Chandrasekaran, Kamalika Chaudhuri, Irene Giacomelli, Somesh Jha, and Songbai Yan. 2020. Exploring Connections Between Active Learning and Model Extraction. In *29th USENIX Security Symposium*. 1309–1326.

[12] Jianbo Chen, Michael I. Jordan, and Martin J. Wainwright. 2020. HopSkipJumpAttack: A Query-Efficient Decision-Based Attack. In *2020 IEEE S&P*. 1277–1294.

[13] Christopher A. Choquette-Choo, Florian Tramer, Nicholas Carlini, and Nicolas Papernot. 2021. Label-Only Membership Inference Attacks. In *Proceedings of the 38th ICML (PMLR, Vol. 139)*. 1964–1974.

[14] Google Cloud. 2021. Vertex Explainable AI. [Accessed 24-Feb-2024]. https://cloud.google.com/vertex-ai/docs/explainable-ai/overview

[15] Consumer Financial Protection Bureau. [n. d.]. 12 CFR 1002.9 – Notifications. https://www.ecfr.gov/current/title-12/part-1002/section-1002.9

[16] Kıvanç Kuzey Dikici, Serdar Kara, Semih Çağlar, Eray Tüzün, and Sinem Sav. 2026. SERSEM: Selective Entropy-Weighted Scoring for Membership Inference in Code Language Models. In *Proceedings of the 34th ACM International Conference on the Foundations of Software Engineering* (Concordia University, Montreal, QC, Canada) *(FSE Companion '26)*. Association for Computing Machinery, New York, NY, USA, 1456–1459. doi:10.1145/3803437.3807734

[17] Pasan Dissanayake and Sanghamitra Dutta. 2024. Model Reconstruction Using Counterfactual Explanations: A Perspective From Polytope Theory. In *NeurIPS 2024*, Amir Globersons, Lester Mackey, Danielle Belgrave, Angela Fan, Ulrich Paquet, Jakub M. Tomczak, and Cheng Zhang (Eds.).

[18] European Parliament and the Council of the European Union. 2024. Regulation (EU) 2024/1689 of the European Parliament and of the Council. Legislative Body: CONSIL, EP. http://data.europa.eu/eli/reg/2024/1689/oj

[19] Fatima Ezzeddine. 2024. Privacy Implications of Explainable AI in Data-Driven Systems. In *Joint Proceedings of the xAI 2024 Late-breaking Work, Demos and Doctoral Consortium co-located with the 2nd World Conference on eXplainable Artificial Intelligence (xAI-2024) (CEUR Workshop Proceedings, Vol. 3793)*. 361–368.

[20] Fatima Ezzeddine, Osama Zammar, Silvia Giordano, and Omran Ayoub. 2026. Explanations Leak: Membership Inference with Differential Privacy and Active Learning Defense. arXiv:2602.03611

[21] Carlos Fernández-Loría, Foster Provost, and Xintian Han. 2021. Explaining Data-Driven Decisions made by AI Systems: The Counterfactual Approach. arXiv:2001.07417

[22] Matt Fredrikson, Somesh Jha, and Thomas Ristenpart. 2015. Model Inversion Attacks that Exploit Confidence Information and Basic Countermeasures. In *Proceedings of the 22nd ACM CCS* (Denver, Colorado, USA) *(CCS '15)*. Association for Computing Machinery, 1322–1333.

[23] Amirata Ghorbani, Abubakar Abid, and James Zou. 2019. Interpretation of neural networks is fragile. In *Proceedings of the AAAI conference on artificial intelligence*, Vol. 33.

[24] J. C. Gower. 1971. A General Coefficient of Similarity and Some of Its Properties. *Biometrics* 27, 4 (1971), 857–871. http://www.jstor.org/stable/2528823

[25] Yash Goyal, Ziyan Wu, Jan Ernst, Dhruv Batra, Devi Parikh, and Stefan Lee. 2019. Counterfactual Visual Explanations. In *Proceedings of the 36th ICML (PMLR, Vol. 97)*, Kamalika Chaudhuri and Ruslan Salakhutdinov (Eds.). 2376–2384.

[26] Riccardo Guidotti, Anna Monreale, Salvatore Ruggieri, Franco Turini, Fosca Giannotti, and Dino Pedreschi. 2018. A Survey of Methods for Explaining Black Box Models. *ACM Comput. Surv.* 51, 5, Article 93 (2018), 42 pages.

[27] Md Imran Hossain, Ghada Zamzmi, Peter R. Mouton, Md Sirajus Salekin, Yu Sun, and Dmitry Goldgof. 2025. Explainable AI for Medical Data: Current Methods, Limitations, and Future Directions. *ACM Comput. Surv.* 57, 6, Article 148 (2025).

[28] M. I. Jordan and T. M. Mitchell. 2015. Machine learning: Trends, perspectives, and prospects. *Science* 349, 6245 (2015), 255–260.

[29] Andrei Kapishnikov, Tolga Bolukbasi, Fernanda Viegas, and Michael Terry. 2019. XRAI: Better Attributions Through Regions. In *2019 IEEE/CVF ICCV*. 4947–4956.

[30] Thibault Laugel, Marie-Jeanne Lesot, Christophe Marsala, Xavier Renard, and Marcin Detyniecki. 2018. Comparison-Based Inverse Classification for Interpretability in Machine Learning. In *IPMU*. Springer International Publishing, Cham, 100–111.

[31] Hao Li, Zheng Li, Siyuan Wu, Yutong Ye, Min Zhang, Dengguo Feng, and Yang Zhang. 2025. Enhanced label-only membership inference attacks with fewer queries. In *34th USENIX Security Symposium* (Seattle, WA, USA). USENIX Association, USA, Article 281, 19 pages.

[32] Huichen Li, Xiaojun Xu, Xiaolu Zhang, Shuang Yang, and Bo Li. 2020. QEBA: Query-Efficient Boundary-Based Blackbox Attack. In *2020 IEEE/CVF CVPR*. 1218–1227. doi:10.1109/CVPR42600.2020.00130

[33] Zheng Li and Yang Zhang. 2021. Membership Leakage in Label-Only Exposures. In *Proceedings of the 2021 ACM CCS* (Virtual Event, Republic of Korea) *(CCS '21)*. Association for Computing Machinery, New York, NY, USA, 880–895.
[34] Han Liu, Yuhao Wu, Zhiyuan Yu, and Ning Zhang. 2024. Please Tell Me More: Privacy Impact of Explainability through the Lens of Membership Inference Attack. In *2024 IEEE S&P*. 4791–4809.
[35] Luca Longo, Mario Brcic, Federico Cabitza, Jaesik Choi, Roberto Confalonieri, et al. 2024. Explainable Artificial Intelligence (XAI) 2.0: A manifesto of open challenges and interdisciplinary research directions. *Information Fusion* 106 (2024). doi:10.1016/j.inffus.2024.102301
[36] Scott M. Lundberg and Su-In Lee. 2017. A unified approach to interpreting model predictions. In *Proceedings of the 31st NeurIPS (NIPS'17)*. Curran Associates Inc., Red Hook, NY, USA, 4768–4777.
[37] Yao Ma, Xurong Zhai, Dan Yu, Yuli Yang, Xingyu Wei, and Yongle Chen. 2024. Label-Only Membership Inference Attack Based on Model Explanation. *Neural Processing Letters* 56, 5 (Sept. 2024), 236. doi:10.1007/s11063-024-11682-1
[38] David Martens and Foster Provost. 2014. Explaining data-driven document classifications. *MIS Q.* 38, 1 (March 2014), 73–100. doi:10.25300/MISQ/2014/38.1.04
[39] Smitha Milli, Ludwig Schmidt, Anca D. Dragan, and Moritz Hardt. 2019. Model Reconstruction from Model Explanations. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* '19)*. 1–9.
[40] Milad Nasr, Reza Shokri, and Amir Houmansadr. 2019. Comprehensive privacy analysis of deep learning: Passive and active white-box inference attacks against centralized and federated learning. In *2019 IEEE S&P*. 739–753.
[41] Abdullah Caglar Oksuz, Anisa Halimi, and Erman Ayday. 2024. AUTOLYCUS: Exploiting Explainable Artificial Intelligence (XAI) for Model Extraction Attacks against Interpretable Models. *PoPETs* 2024, 4 (Oct. 2024), 684–699.
[42] Abdullah Caglar Oksuz, Anisa Halimi, and Erman Ayday. 2026. LoMime: Query-Efficient Membership Inference using Model Extraction in Label-Only Settings. arXiv:2602.18934
[43] Martin Pawelczyk, Himabindu Lakkaraju, and Seth Neel. 2023. On the Privacy Risks of Algorithmic Recourse. In *AISTATS (PMLR, Vol. 206)*, Francisco J. R. Ruiz, Jennifer G. Dy, and Jan-Willem van de Meent (Eds.). PMLR, 9680–9696.
[44] Yuefeng Peng, Jaechul Roh, Subhransu Maji, and Amir Houmansadr. 2024. OSLO: One-Shot Label-Only Membership Inference Attacks. *Advances in Neural Information Processing Systems* 37 (Dec. 2024), 62310–62333. doi:10.52202/079017-1991
[45] Mauro Ribeiro, Katarina Grolinger, and Miriam A.M. Capretz. 2015. MLaaS: Machine Learning as a Service. In *2015 IEEE 14th ICMLA*. 896–902.
[46] Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. 2016. "Why Should I Trust You?": Explaining the Predictions of Any Classifier. In *Proceedings of the 22nd ACM SIGKDD*. San Francisco, CA, USA, 1135–1144.
[47] Pau Rodríguez, Massimo Caccia, Alexandre Lacoste, Lee Zamparo, Issam Laradji, Laurent Charlin, and David Vazquez. 2021. Beyond Trivial Counterfactual Explanations with Diverse Valuable Explanations. In *2021 IEEE/CVF ICCV*. 1036–1045.
[48] Emilia Romeo and Ján Lacko. 2026. Adoption and integration of AI in organizations: a systematic review of challenges and drivers towards future directions of research. *Kybernetes* 55, 3 (04 2026), 1286–1307.
[49] Ahmed Salem, Yang Zhang, Mathias Humbert, Pascal Berrang, Mario Fritz, and Michael Backes. 2019. ML-Leaks: Model and Data Independent Membership Inference Attacks and Defenses on Machine Learning Models. In *NDSS 2019,*. The Internet Society, USA.
[50] Ramprasaath R. Selvaraju, Michael Cogswell, Abhishek Das, Ramakrishna Vedantam, Devi Parikh, and Dhruv Batra. 2017. Grad-CAM: Visual Explanations from Deep Networks via Gradient-Based Localization. In *2017 IEEE ICCV*. 618–626.
[51] Amazon Web Services. 2025. Feature attributions that use Shapley Values - Amazon SageMaker AI. https://docs.aws.amazon.com/sagemaker/latest/dg/clarify-shapley-values.html
[52] Amazon Web Services. 2026. Amazon SageMaker Clarify. [Accessed 24-Feb-2026]. https://aws.amazon.com/sagemaker/clarify/
[53] Reza Shokri, Martin Strobel, and Yair Zick. 2021. On the Privacy Risks of Model Explanations. In *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society (AIES '21)*. Association for Computing Machinery, 231–241.
[54] Reza Shokri, Marco Stronati, Congzheng Song, and Vitaly Shmatikov. 2017. Membership Inference Attacks Against Machine Learning Models. In *2017 IEEE S&P*. USA, 3–18.
[55] Dylan Slack, Sophie Hilgard, Emily Jia, Sameer Singh, and Himabindu Lakkaraju. 2020. Fooling LIME and SHAP: Adversarial Attacks on Post hoc Explanation Methods. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society (AIES '20)*. 180–186.
[56] Daniel Smilkov, Nikhil Thorat, Been Kim, Fernanda B. Viégas, and Martin Wattenberg. 2017. SmoothGrad: removing noise by adding noise. *CoRR* abs/1706.03825 (2017). arXiv:1706.03825 http://arxiv.org/abs/1706.03825
[57] Mukund Sundararajan, Ankur Taly, and Qiqi Yan. 2017. Axiomatic attribution for deep networks. In *International conference on machine learning*. PMLR, 3319–3328.
[58] Florian Tramèr, Fan Zhang, Ari Juels, Michael K. Reiter, and Thomas Ristenpart. 2016. Stealing machine learning models via prediction APIs. In *25th USENIX Security Symposium* (Austin, TX, USA) *(SEC'16)*. USA, 601–618.

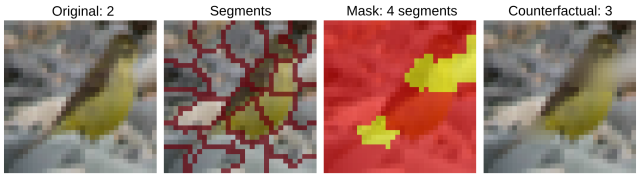

Figure 11: SEDC counterfactual example of a sample from CIFAR-10 whose class changed from the correctly classified class 2 (bird) to class 3 (cat) by applying Gaussian blur to 4 segments.

[59] Tom Vermeire, Dieter Brughmans, Sofie Goethals, Raphael Mazzine Barbossa de Oliveira, and David Martens. 2022. Explainable image classification with evidence counterfactual. *Pattern Analysis and Applications* 25, 2 (May 2022), 315–335. doi:10.1007/s10044-021-01055-y
[60] Sandra Wachter, Brent Mittelstadt, and Chris Russell. 2017. Counterfactual explanations without opening the black box: Automated decisions and the GDPR. *Harv. JL & Tech.* 31 (2017), 841.
[61] Yongjie Wang, Hangwei Qian, and Chunyan Miao. 2022. DualCF: Efficient Model Extraction Attack from Counterfactual Explanations. In *FAccT '22: 2022 ACM Conference on Fairness, Accountability, and Transparency, Seoul, Republic of Korea, June 21 - 24, 2022*. ACM, 1318–1329. doi:10.1145/3531146.3533188
[62] Zhou Wang, A.C. Bovik, H.R. Sheikh, and E.P. Simoncelli. 2004. Image quality assessment: from error visibility to structural similarity. *IEEE Transactions on Image Processing* 13, 4 (2004), 600–612. doi:10.1109/TIP.2003.819861
[63] Yutong Wu, Han Qiu, Shangwei Guo, Jiwei Li, and Tianwei Zhang. 2024. You Only Query Once: An Efficient Label-Only Membership Inference Attack. In *ICLR*, Vol. 2024. 41291–41304.
[64] Jiayuan Ye, Aadyaa Maddi, Sasi Kumar Murakonda, Vincent Bindschaedler, and Reza Shokri. 2022. Enhanced membership inference attacks against machine learning models. In *Proceedings of the 2022 ACM CCS*. 3093–3106.
[65] Samuel Yeom, Irene Giacomelli, Matt Fredrikson, and Somesh Jha. 2018. Privacy Risk in Machine Learning: Analyzing the Connection to Overfitting. In *2018 IEEE 31st CSF*. 268–282.
[66] Richard Zhang, Phillip Isola, Alexei A. Efros, Eli Shechtman, and Oliver Wang. 2018. The Unreasonable Effectiveness of Deep Features as a Perceptual Metric. arXiv:1801.03924
[67] Xuejun Zhao, Wencan Zhang, Xiaokui Xiao, and Brian Lim. 2021. Exploiting Explanations for Model Inversion Attacks. In *IEEE/CVF ICCV*. 682–692.

A Distance Threshold Calibration (DTC)

Algorithm A.1 formalizes the two-stage calibration procedure described in Section 5.3.

B Example Visual Counterfactual Explanations

An example SEDC counterfactual where we used Gaussian blur to remove image segments can be seen in Figure 11. Additionally, a DiVE example is also presented in Figure 12.

C Additional Ablation Studies

C.1 Distance to the Decision Boundary Distribution

Figure 13 presents member and non-member counterfactual-distance distributions, which are detailed in Section 6.3.

C.2 Second Stage DTC Ablation Using a Joint Objective

We provide an ablation study of the second-stage distance-threshold calibration objective in Table 3 which facilitates the discussion in Section 6.3. The objective maximizes balanced accuracy instead

Algorithm A.1: Distance Threshold Calibration (DTC)

Input: Calibration records $\mathcal{B}_q = \{(d_i^{(\mathcal{E},q)}, m_i)\}$ for distance measure q, containing finite boundary-distance estimates $d_i^{(\mathcal{E},q)}$ and known membership labels $m_i \in \{0, 1\}$

Output: Calibrated threshold τ_q and inequality direction dir_q

$\mathcal{T}_q \leftarrow \text{sort}\left(\text{unique}\left(\{d \mid (d, m) \in \mathcal{B}_q\}\right)\right)$;
$N_{\text{mem}} \leftarrow |\{(d, m) \in \mathcal{B}_q \mid m = 1\}|$;
$N_{\text{non}} \leftarrow |\{(d, m) \in \mathcal{B}_q \mid m = 0\}|$;
$best_acc \leftarrow -\infty$;
foreach $\tau_c \in \mathcal{T}_q$ **do**
  foreach $dir \in \{\texttt{leq}, \texttt{gt}\}$ **do**
    Compute $acc_{\text{mem}}(\tau_c, dir)$ and $acc_{\text{non}}(\tau_c, dir)$;
    $acc_C \leftarrow \dfrac{N_{\text{mem}}\, acc_{\text{mem}} + N_{\text{non}}\, acc_{\text{non}}}{N_{\text{mem}} + N_{\text{non}}}$;
    if $acc_C > best_acc$ **then**
      $best_acc \leftarrow acc_C$;
      $dir_q \leftarrow dir$;
$min_diff \leftarrow \infty$;
foreach $\tau_c \in \mathcal{T}_q$ **do**
  Compute $acc_{\text{mem}}(\tau_c, dir_q)$ and $acc_{\text{non}}(\tau_c, dir_q)$;
  $diff \leftarrow |acc_{\text{mem}} - acc_{\text{non}}|$;
  if $diff < min_diff$ **then**
    $min_diff \leftarrow diff$;
    $\tau_q \leftarrow \tau_c$;
return τ_q, dir_q;

of minimizing the difference between member and nonmember accuracy.

D Comparison to Related Work

Table 4 presents a comparison of prior MIA work on Purchase-100, Location, and Texas-100, which facilitates the discussion in Section 6.5.

E Implementation Details

E.1 XAI Setup

For the explanation tools, we use an authors' implementation whenever it is directly applicable to our setting. LIME explanations are generated using the official Python implementation of LIME [46], specifically, we use the LIME Tabular Explainer. Unless otherwise stated, we use LIME's default configuration, including 5,000 perturbed samples and the Euclidean distance metric. For LIME, each response contains the predicted label together with the local surrogate's weights and intercepts.

For NICE we use the official `NICEx` implementation (version 0.2.3). For the remaining methods, a local implementation or adaptation was required to support a common explanation-service interface, identical target-model checkpoints and preprocessing, multiclass

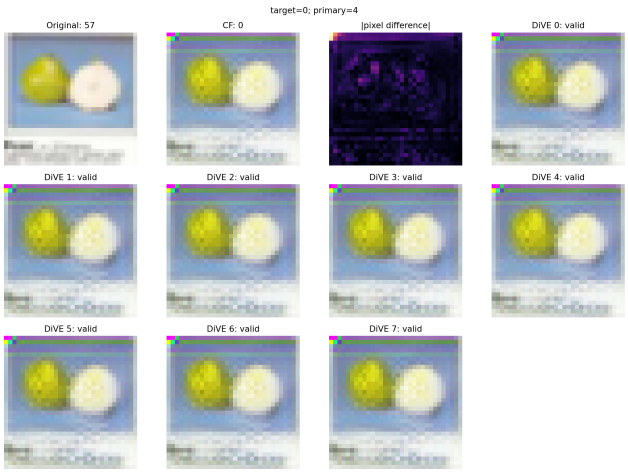

Figure 12: DiVE counterfactual example of a sample from CIFAR-100 whose class changed from the correctly classified class 57 (pear) to class 0 (apple). The primary example was chosen to be the one with index 4. DiVE generates multiple candidates and chooses the best one as previously explained in Section 5.2.

prediction, explicit generation failures, and reproducible control of the generation budget.

Wachter [60] does not provide an official maintained implementation, we therefore implement its original squared target-output loss and inverse-MAD-weighted L_1 proximity objective. We use Adam as the numerical optimizer and extend target selection to multiclass models by optimizing toward the most probable alternative classes. Candidate values are constrained to the feature ranges observed in the provider's training data, and the closest valid candidate is returned.

Although reference code is available for Growing Spheres, we reimplement its two defining stages to integrate it with the same TensorFlow target models, preprocessing pipeline, batched provider-side inference, deterministic random seeding, and explicit failure and budget semantics used by the other explainers. The implementation retains the published algorithmic structure: volume-uniform sampling within expanding Euclidean shells, selection of the nearest differently classified candidate, and subsequent feature sparsification. The search is untargeted and accepts any label change, thereby supporting the multiclass datasets considered in our evaluation. Feasibility clipping is disabled by default so that it does not alter the original unconstrained search.

For the visual methods, the released implementations are specialized to different experimental settings. The released SEDC workflow is centered on ImageNet-scale inputs and MobileNetV2, whereas our evaluation uses MNIST- and CIFAR-scale images and separately trained multiclass target models. We therefore implement the published segment-removal search using grid or SLIC segmentation and Gaussian-blur replacement, which is a replacement operation supported by the authors' implementation. We additionally introduce explicit evaluation, runtime, queue-size, and segment-count limits and perform a final set-minimal reduction; these choices

Table 3: Ablation of the second-stage distance-threshold calibration objective. "Gap" is the original objective that minimizes the member/non-member calibration-accuracy difference; "Joint" maximizes their balanced accuracy. The final column reports the absolute held-out difference between member and non-member accuracies. Directions indicate the membership decision rule.

Method	Dataset (distance)	Threshold: Gap → Joint	ASR_{Gap} (%)	ASR_{Joint} (%)	Δ ASR (pp)	$\lvert Acc_{mem} - Acc_{non} \rvert$
Image attacks						
SEDC	CIFAR-10 (segment count)	$> 2 \rightarrow > 2$	72.17	72.17	0.00	$9.95 \rightarrow 9.95$
SEDC	CIFAR-100 (segment count)	$> 2 \rightarrow > 1$	65.54	68.42	+2.88	$9.70 \rightarrow 52.96$
DiVE	CIFAR-10 (LPIPS)	$\leq 0.33994 \rightarrow \leq 0.30868$	84.27	85.87	+1.61	$6.47 \rightarrow 11.93$
DiVE	CIFAR-100 (SSIM)	$\leq 0.37663 \rightarrow \leq 0.38916$	75.33	75.08	−0.25	$2.02 \rightarrow 11.63$
Tabular attacks						
NICE	Purchase-100	$> 6.1635 \rightarrow > 5.5554$	84.79	87.21	+2.43	$3.00 \rightarrow 13.00$
NICE	Location	$> 8.1828 \rightarrow > 7.9683$	73.97	75.16	+1.19	$2.86 \rightarrow 3.02$
NICE	Texas-100	$> 22.5220 \rightarrow > 15.1152$	62.71	62.29	−0.43	$0.00 \rightarrow 34.57$
Growing Spheres	Purchase-100	$> 6.3567 \rightarrow > 5.3271$	95.00	94.43	−0.57	$2.57 \rightarrow 10.86$
Growing Spheres	Location	$> 7.4628 \rightarrow > 7.6169$	80.16	80.40	+0.24	$9.52 \rightarrow 8.10$
Growing Spheres	Texas-100	$> 23.6475 \rightarrow > 19.3501$	82.00	83.64	+1.64	$8.00 \rightarrow 23.29$
Wachter	Purchase-100	$> 1.9382 \rightarrow > 1.7404$	92.14	94.93	+2.79	$0.57 \rightarrow 8.71$
Wachter	Location	$> 2.6870 \rightarrow > 2.6861$	79.52	79.52	0.00	$3.81 \rightarrow 3.81$
Wachter	Texas-100	$> 2.8330 \rightarrow > 2.0718$	76.29	79.21	+2.93	$6.86 \rightarrow 28.43$

Table 4: Comparison of MIAs on Purchase-100, Location, and Texas-100. ASR and AUC are reported as fractions; "–" denotes an unreported result. Queries count attacker-issued target-model or explanation-service interactions.

Method	Attack signal	Experimental setup	Queries		Purchase		Location		Texas	
			One-time	Per candidate	ASR	AUC	ASR	AUC	ASR	AUC
Gap Attack [65]	Prediction correctness	Evaluated on our target models; no calibration	0	1 label	.692	–	.661	–	.652	–
Choquette-Choo et al. [13]	Robustness to input perturbations	Shadow/source model for threshold selection; hard-label access	0 target	10^4 labels	.874	–	**.892**	–	.803	–
Shokri et al. [53]	Variance of LIME coefficients	Optimal variance threshold; Purchase and Texas only	–	1 expl.[a]	≤ .500[a]	–	–	–	≤ .500[a]	–
Ye et al. [64] (Attack P)	Population-calibrated loss score	1,000 population records per class; score access	10^5 scores[b]	1 score	–	.818	–	–	–	–
YOQO [63]	Shadow-crafted membership query	Online: 16 in/out pairs; offline: 16 out-models	0 target	1 label	.6828/.700[c]	–	.775/.8325[c]	–	–	–
DHAttack [31]	Fixed-direction boundary distance	256 shadow models; 40,000 reference records	$4{\times}10^4$ labels[d]	≤ 30 labels	–	.727	–	–	–	–
LoMime [42]	Boundary distance on an extracted surrogate	Label-only extraction followed by offline MIA	$10^6/10^5/10^6$[e]	0 target	.8495	.8530	.8563	.8582	.7823	.7996
XposedMIA (LIME)	Local boundary distance	Calibrated threshold and direction; no shadows	30 expl.[f]	1 expl.	.850	.914	.710	.780	.640	.688
XposedMIA (NICE)	Counterfactual distance	NICE counterfactual; no shadows	30 expl.[f]	1 expl.	.848	.916	.740	.812	.627	.676
XposedMIA (Growing Spheres)	Counterfactual distance	Growing Spheres counterfactual; no shadows	30 expl.[f]	1 expl.	**.950**	**.976**	.802	**.871**	**.820**	**.881**
XposedMIA (Wachter)	Counterfactual distance	Wachter counterfactual; no shadows	30 expl.[f]	1 expl.	.921	.967	.795	.847	.763	.798

[a] Shokri et al. do not report exact ASRs, calibration costs, or LIME's internal prediction-query count; "1 expl." denotes one released explanation per candidate. [b] Attack P uses 1,000 population records for each of the 100 Purchase classes, giving 10^5 target-score evaluations. [c] Values are YOQO online/offline. [d] DHAttack's $4{\times}10^4$ queries are for Purchase reference-data relabeling; calibrating a fixed binary threshold may require additional queries, whereas the reported Purchase result is AUC. [e] LoMime budgets are one-time extraction queries for Purchase/Location/Texas; membership inference is subsequently offline. [f] XposedMIA uses 30 explanation requests for calibration and one externally issued explanation request per candidate; service-side model evaluations are not counted separately.

make failures observable and prevent unbounded provider-side computation.

DiVE's released implementation and pretrained models target different image domains and binary prediction tasks. We adapt the authors' reference implementation to TensorFlow/NHWC, 28×28 and 32×32 inputs, one-and three-channel images, and multiclass classifiers. We preserve the β-TCVAE architecture and its reconstruction, latent-sparsity, and diversity objectives, replacing only the binary target loss with multiclass target cross-entropy. DiVE generates multiple candidates internally, after which the provider returns the valid candidate having the smallest published proximity objective. These adaptations concern compatibility rather than the membership decision rule.

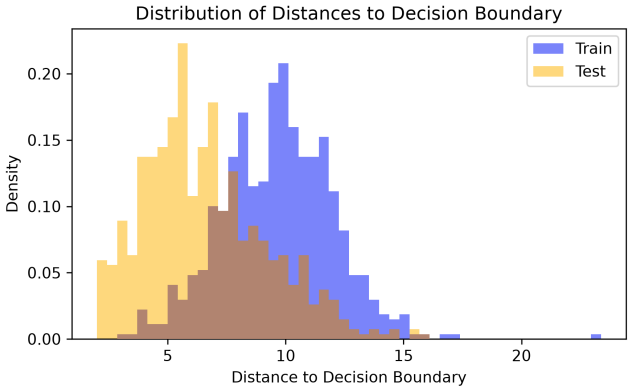

(a) NICE on Location.

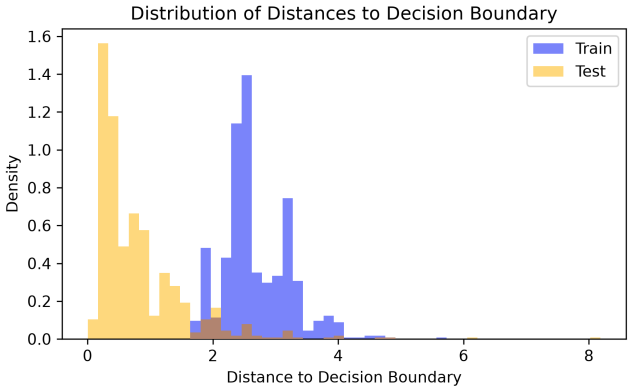

(b) Wachter on Purchase-100.

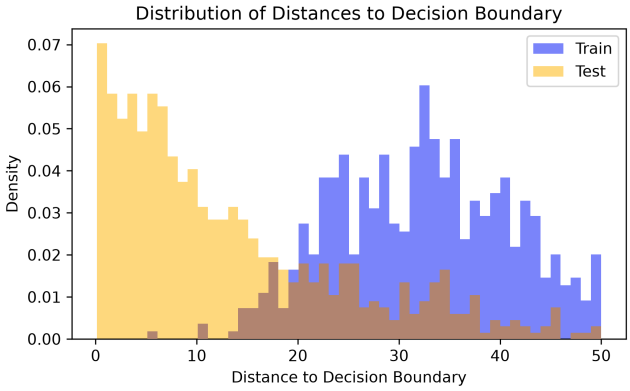

(c) Growing Spheres on Texas-100.

Figure 13: Member and non-member counterfactual-distance distributions for the three counterfactual explanation methods. (a) NICE on Location, (b) Wachter on Purchase-100, and (c) Growing Spheres on Texas-100.

E.2 Dataset Splits Among the Different Attacks

We provide a detailed information on dataset splitting and attack-set construction in Table 5.

Table 5: Dataset splitting and attack-set construction for the selected configurations. T denotes the target-model member-source pool and N the disjoint non-member source pool. "Fit/validation" gives the target model's gradient-training and validation counts. For supervised calibration, c/i denotes the number of calibration and held-out inference examples *per membership group*; thus, each attack uses $2c$ calibration and $2i$ inference queries. For each image dataset, SEDC and DiVE use the identical fixed target model and member-source pool; only their non-member pool sizes and calibration/inference allocations differ.

Attack family	Dataset	Target model and member source	Method	$\lvert N\rvert$	c/i per group	Attack-subset construction
Tabular attacks						
LIME	Purchase-100	$\lvert T\rvert$ = 6,000 4,800 / 1,200 fit/validation	LIME	1,000	300 / 700	Balanced
LIME	Location	$\lvert T\rvert$ = 1,200 960 / 240 fit/validation	LIME	900	270 / 630	Balanced
LIME	Texas-100	$\lvert T\rvert$ = 10,000 8,000 / 2,000 fit/validation	LIME	1,000	300 / 700	Not class-balanced
Counterfactual	Purchase-100	$\lvert T\rvert$ = 10,000 8,000 / 2,000 fit/validation	NICE, Wachter, Growing Spheres	1,000	300 / 700	Not class-balanced
Counterfactual	Location	$\lvert T\rvert$ = 1,200 960 / 240 fit/validation	NICE, Wachter, Growing Spheres	1,200	270 / 630	Balanced
Counterfactual	Texas-100	$\lvert T\rvert$ = 10,000 8,000 / 2,000 fit/validation	NICE, Wachter, Growing Spheres	1,000	300 / 700	Balanced
Counterfactual	Adult	$\lvert T\rvert$ = 10,000 8,000 / 2,000 fit/validation	NICE, Wachter, Growing Spheres	1,000	300 / 700	Balanced
Image attacks: shared target model within each dataset						
Image counterfactual	CIFAR-10	Shared CNN target $\lvert T\rvert$ = 10,000; 9,000 / 1,000 fit/validation	SEDC DiVE	10,000 5,000	500 / 1,000 100 / 1,000	Balanced Balanced
Image counterfactual	CIFAR-100	Shared ResNet-18 target $\lvert T\rvert$ = 50,000; 45,000 / 5,000 fit/validation	SEDC DiVE	10,000 10,000	500 / 1,000 100 / 1,000	Balanced Balanced

Implementation details. All calibration and inference subsets are disjoint. The supervised calibration subsets are used only to select the distance threshold; ASR is computed exclusively on the held-out inference subsets. For tabular target models, Keras reserves 20% of T as a validation split during fitting. For SEDC and DiVE, the 10% target-validation split is excluded from the image-attack member population. DiVE additionally uses a separate explainer validation set: for CIFAR-10, its explainer is trained on 8,100 target-fitting images and validated on 810 target-validation images; for CIFAR-100, its explainer uses 40,500 target-fitting images and 4,500 source-holdout validation images, which are excluded from the DiVE member attack population.